

# GNU Guix Reference Manual

---

Using the GNU Guix Functional Package Manager

The GNU Guix Developers

---

Edition 8b23207  
17 February 2026

Copyright © 2012–2026 Ludovic Courtès  
Copyright © 2013, 2014, 2016, 2024 Andreas Enge  
Copyright © 2013 Nikita Karetnikov  
Copyright © 2014, 2015, 2016 Alex Kost  
Copyright © 2015, 2016 Mathieu Lirzin  
Copyright © 2014 Pierre-Antoine Rault  
Copyright © 2015 Taylan Ulrich Bayırlı/Kammer  
Copyright © 2015, 2016, 2017, 2019, 2020, 2021, 2023, 2025 Leo Famulari  
Copyright © 2015, 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023 Ricardo Wurmus  
Copyright © 2016 Ben Woodcroft  
Copyright © 2016, 2017, 2018, 2021 Chris Marusich  
Copyright © 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2025 Efraim Flashner  
Copyright © 2016 John Darrington  
Copyright © 2016, 2017 Nikita Gillmann  
Copyright © 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2024, 2025 Janneke Nieuwenhuizen  
Copyright © 2016, 2017, 2018, 2019, 2020, 2021 Julien Lepiller  
Copyright © 2016 Alex ter Weele  
Copyright © 2016, 2017, 2018, 2019, 2020, 2021 Christopher Baines  
Copyright © 2017, 2018, 2019 Clément Lassieur  
Copyright © 2017, 2018, 2020, 2021, 2022 Mathieu Othacehe  
Copyright © 2017 Federico Beffa  
Copyright © 2017, 2018, 2024, 2025 Carlo Zancanaro  
Copyright © 2017 Thomas Danckaert  
Copyright © 2017 humanitiesNerd  
Copyright © 2017, 2021 Christine Lemmer-Webber  
Copyright © 2017, 2018, 2019, 2020, 2021, 2022 Marius Bakke  
Copyright © 2017, 2019, 2020, 2022 Hartmut Goebel  
Copyright © 2017, 2019–2025 Maxim Cournoyer  
Copyright © 2017–2022 Tobias Geerinckx-Rice  
Copyright © 2017 George Clemmer  
Copyright © 2017 Andy Wingo  
Copyright © 2017, 2018, 2019, 2020, 2023, 2024, 2025 Arun Isaac  
Copyright © 2017 nee  
Copyright © 2018 Rutger Helling  
Copyright © 2018, 2021, 2023, 2025, 2026 Oleg Pykhalov  
Copyright © 2018 Mike Gerwitz  
Copyright © 2018 Pierre-Antoine Rouby  
Copyright © 2018, 2019 Gábor Boskovits  
Copyright © 2018, 2019, 2020, 2022, 2023, 2024 Florian Pelz  
Copyright © 2018 Laura Lazzati  
Copyright © 2018 Alex Vong

Copyright © 2019 Josh Holland  
Copyright © 2019, 2020 Diego Nicola Barbato  
Copyright © 2019 Ivan Petkov  
Copyright © 2019 Jakob L. Kreuze  
Copyright © 2019 Kyle Andrews  
Copyright © 2019 Alex Griffin  
Copyright © 2019, 2020, 2021, 2022 Guillaume Le Vaillant  
Copyright © 2020 Liliana Marie Prikler  
Copyright © 2019, 2020, 2021, 2022, 2023 Simon Tournier  
Copyright © 2020 Wiktor Żelazny  
Copyright © 2020 Damien Cassou  
Copyright © 2020 Jakub Kądziołka  
Copyright © 2020 Jack Hill  
Copyright © 2020 Naga Malleswari  
Copyright © 2020, 2021 Brice Waegeneire  
Copyright © 2020 R Veera Kumar  
Copyright © 2020, 2021, 2022 Pierre Langlois  
Copyright © 2020 pinoaffe  
Copyright © 2020, 2023 André Batista  
Copyright © 2020, 2021 Alexandru-Sergiu Marton  
Copyright © 2020 raingloom  
Copyright © 2020 Daniel Brooks  
Copyright © 2020 John Soo  
Copyright © 2020 Jonathan Brielmaier  
Copyright © 2020 Edgar Vincent  
Copyright © 2021, 2022 Maxime Devos  
Copyright © 2021 B. Wilson  
Copyright © 2021 Xinglu Chen  
Copyright © 2021 Raghav Gururajan  
Copyright © 2021 Domagoj Stolf  
Copyright © 2021 Hui Lu  
Copyright © 2021 pukkamustard  
Copyright © 2021 Alice Brenon  
Copyright © 2021-2023 Josselin Poiret  
Copyright © 2021, 2023 muradm  
Copyright © 2021, 2022 Andrew Tropin  
Copyright © 2021 Sarah Morgensen  
Copyright © 2022 Remco van 't Veer  
Copyright © 2022 Aleksandr Vityazev  
Copyright © 2022 Philip M<sup>c</sup>Grath  
Copyright © 2022 Karl Hallsby  
Copyright © 2022 Justin Veilleux  
Copyright © 2022 Reily Siegel  
Copyright © 2022 Simon Streit  
Copyright © 2022 (  
Copyright © 2022 John Kehayias  
Copyright © 2022–2023 Bruno Victal

Copyright © 2022 Ivan Vilata-i-Balaguer  
Copyright © 2023-2025 Giacomo Leidi  
Copyright © 2022 Antero Mejr  
Copyright © 2023 Karl Hallsby  
Copyright © 2023 Nathaniel Nicandro  
Copyright © 2023 Tanguy Le Carrou  
Copyright © 2023-2024 Zheng Junjie  
Copyright © 2023 Brian Cully  
Copyright © 2023 Felix Lechner  
Copyright © 2023 Foundation Devices, Inc.  
Copyright © 2023 Thomas Jeong  
Copyright © 2023 Saku Laesvuori  
Copyright © 2023 Graham James Addis  
Copyright © 2023-2025 Tomas Volf  
Copyright © 2024, 2025 Herman Rimm  
Copyright © 2024 Matthew Trzcinski  
Copyright © 2024 Richard Sent  
Copyright © 2024 Dariqq  
Copyright © 2024 Denis 'GNUtoo' Carikli  
Copyright © 2024 Fabio Natali  
Copyright © 2024 Arnaud Daby-Seesaram  
Copyright © 2024 Wilko Meyer  
Copyright © 2024-2025 Nigko Yerden  
Copyright © 2024 Troy Figiel  
Copyright © 2024 Sharlatan Hellseher  
Copyright © 2024 45mg  
Copyright © 2024 Greg Hogan  
Copyright © 2025 Sören Tempel  
Copyright © 2025 Rostislav Svoboda  
Copyright © 2025 Zacchaeus  
Copyright © 2025 Sergio Pastor Pérez  
Copyright © 2024 Evgeny Pisemsky  
Copyright © 2025 jgart  
Copyright © 2025 Artur Wroblewski  
Copyright © 2025 Edouard Klein  
Copyright © 2025 Rodion Goritskov  
Copyright © 2025 dan  
Copyright © 2025 Noé Lopez  
Copyright © 2026 David Elsing

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled “GNU Free Documentation License”.

# Table of Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Managing Software the Guix Way	1
1.2	GNU Distribution	2
<b>2</b>	<b>Installation</b>	<b>4</b>
2.1	Binary Installation	4
2.2	Setting Up the Daemon	5
2.2.1	Build Environment Setup	6
	Daemon Running as Root	6
	Daemon Running Without Privileges	7
	Migrating to the Unprivileged Daemon	7
	The Isolated Build Environment	8
2.2.2	Using the Offload Facility	9
2.2.3	SELinux Support	12
2.2.3.1	Installing the SELinux policy	12
2.2.3.2	Limitations	13
2.2.4	AppArmor Support	13
2.2.4.1	Installing the AppArmor profile	13
2.3	Invoking <code>guix-daemon</code>	14
2.4	Application Setup	18
2.4.1	Locales	18
2.4.2	Name Service Switch	19
2.4.3	X11 Fonts	20
2.4.4	X.509 Certificates	20
2.4.5	Emacs Packages	21
2.5	Upgrading Guix	21
<b>3</b>	<b>System Installation</b>	<b>23</b>
3.1	Limitations	23
3.2	Hardware Considerations	23
3.3	USB Stick and DVD Installation	23
	Copying to a USB Stick	24
	Burning on a DVD	24
	Bootting	24
3.4	Preparing for Installation	25
3.5	Guided Graphical Installation	25
3.6	Manual Installation	27
3.6.1	Keyboard Layout, Networking, and Partitioning	27
3.6.1.1	Keyboard Layout	27
3.6.1.2	Networking	27
3.6.1.3	Disk Partitioning	28
3.6.2	Proceeding with the Installation	30

3.7	After System Installation .....	31
3.8	Installing Guix in a Virtual Machine .....	31
3.9	Building the Installation Image .....	32
3.10	Building the Installation Image for ARM Boards .....	32
<b>4</b>	<b>Getting Started .....</b>	<b>33</b>
<b>5</b>	<b>Package Management .....</b>	<b>36</b>
5.1	Features .....	36
5.2	Invoking <code>guix package</code> .....	37
5.3	Substitutes .....	46
5.3.1	Official Substitute Servers .....	47
5.3.2	Substitute Server Authorization .....	47
5.3.3	Getting Substitutes from Other Servers .....	48
5.3.4	Substitute Authentication .....	49
5.3.5	Proxy Settings .....	50
5.3.6	Substitution Failure .....	50
5.3.7	On Trusting Binaries .....	50
5.4	Packages with Multiple Outputs .....	51
5.5	Invoking <code>guix locate</code> .....	52
5.6	Invoking <code>guix gc</code> .....	53
5.7	Invoking <code>guix pull</code> .....	56
5.8	Invoking <code>guix time-machine</code> .....	60
5.9	Inferiors .....	62
5.10	Invoking <code>guix describe</code> .....	64
5.11	Invoking <code>guix archive</code> .....	65
<b>6</b>	<b>Channels .....</b>	<b>68</b>
6.1	Specifying Additional Channels .....	68
6.2	Using a Custom Guix Channel .....	68
6.3	Replicating Guix .....	69
6.4	Customizing the System-Wide Guix .....	70
6.5	Channel Authentication .....	71
6.6	Channels with Substitutes .....	71
6.7	Creating a Channel .....	72
6.8	Package Modules in a Sub-directory .....	73
6.9	Declaring Channel Dependencies .....	74
6.10	Specifying Channel Authorizations .....	74
6.11	Primary URL .....	76
6.12	Writing Channel News .....	76

<b>7</b>	<b>Development</b>	<b>78</b>
7.1	Invoking <code>guix shell</code>	78
7.2	Invoking <code>guix environment</code>	85
7.3	Invoking <code>guix pack</code>	90
7.4	The GCC toolchain	98
7.5	Invoking <code>guix git authenticate</code>	98
<b>8</b>	<b>Programming Interface</b>	<b>100</b>
8.1	Package Modules	100
8.2	Defining Packages	101
8.2.1	<code>package</code> Reference	104
8.2.2	<code>origin</code> Reference	109
8.3	Defining Package Variants	113
8.4	Writing Manifests	117
8.5	Build Systems	121
8.6	Build Phases	143
8.7	Build Utilities	146
8.7.1	Dealing with Store File Names	147
8.7.2	File Types	147
8.7.3	File Manipulation	147
8.7.4	File Search	149
8.7.5	Program Invocation	150
8.7.6	Build Phases	151
8.7.7	Wrappers	152
8.8	Search Paths	153
8.9	The Store	156
8.10	Derivations	158
8.11	The Store Monad	161
8.12	G-Expressions	165
8.13	Invoking <code>guix repl</code>	175
8.14	Using Guix Interactively	176
<b>9</b>	<b>Utilities</b>	<b>180</b>
9.1	Invoking <code>guix build</code>	180
9.1.1	Common Build Options	180
9.1.2	Package Transformation Options	182
9.1.3	Additional Build Options	189
9.1.4	Debugging Build Failures	193
9.2	Invoking <code>guix edit</code>	195
9.3	Invoking <code>guix download</code>	195
9.4	Invoking <code>guix hash</code>	196
9.5	Invoking <code>guix import</code>	197
9.6	Invoking <code>guix refresh</code>	206
9.7	Invoking <code>guix style</code>	213
9.8	Invoking <code>guix lint</code>	217

9.9	Invoking <code>guix size</code> .....	220
9.10	Invoking <code>guix graph</code> .....	222
9.11	Invoking <code>guix publish</code> .....	227
9.12	Invoking <code>guix challenge</code> .....	231
9.13	Invoking <code>guix copy</code> .....	233
9.14	Invoking <code>guix container</code> .....	234
9.15	Invoking <code>guix weather</code> .....	235
9.16	Invoking <code>guix processes</code> .....	237
<b>10</b>	<b>Foreign Architectures .....</b>	<b>239</b>
10.1	Cross-Compilation .....	239
10.2	Native Builds .....	240
<b>11</b>	<b>System Configuration .....</b>	<b>242</b>
11.1	Getting Started .....	242
11.2	Using the Configuration System .....	244
	Bootloader .....	246
	Globally-Visible Packages .....	246
	System Services .....	247
	Inspecting Services .....	252
	Instantiating the System .....	253
	The Programming Interface .....	253
11.3	<code>operating-system</code> Reference .....	253
11.4	File Systems .....	257
	11.4.1 Btrfs file system .....	261
11.5	Mapped Devices .....	263
11.6	Swap Space .....	266
11.7	User Accounts .....	269
11.8	Keyboard Layout .....	271
11.9	Locales .....	273
	11.9.1 Locale Data Compatibility Considerations .....	274
11.10	Services .....	275
	11.10.1 Base Services .....	276
	11.10.2 Scheduled Job Execution .....	299
	11.10.2.1 Shepherd Timers .....	300
	11.10.2.2 Mcron .....	301
	11.10.3 Log Rotation .....	304
	11.10.4 Networking Setup .....	307
	11.10.5 Networking Services .....	322
	11.10.6 Unattended Upgrades .....	352
	11.10.7 X Window .....	355
	11.10.8 Printing Services .....	365
	11.10.9 Desktop Services .....	377
	11.10.10 Sound Services .....	396
	11.10.11 File Search Services .....	400
	11.10.12 Database Services .....	401



11.10.13	Mail Services	407
11.10.14	Messaging Services	437
11.10.15	Telephony Services	456
11.10.16	File-Sharing Services	463
11.10.17	Monitoring Services	474
11.10.18	Kerberos Services	483
11.10.19	LDAP Services	484
11.10.20	Web Services	494
11.10.21	High Availability Services	517
11.10.22	Certificate Services	518
11.10.23	DNS Services	522
11.10.24	VNC Services	535
11.10.25	VPN Services	538
11.10.26	DAV Services	544
11.10.27	Network File System	545
11.10.28	Samba Services	552
11.10.29	Continuous Integration	554
11.10.30	Power Management Services	560
11.10.31	Audio Services	575
11.10.32	Virtualization Services	581
11.10.33	Version Control Services	608
11.10.34	Game Services	626
11.10.35	PAM Mount Service	628
11.10.36	Guix Services	631
11.10.37	Linux Services	641
11.10.38	Hurd Services	649
11.10.39	Miscellaneous Services	649
11.11	Privileged Programs	678
11.12	X.509 Certificates	679
11.13	Name Service Switch	680
11.14	Initial RAM Disk	682
11.15	Bootloader Configuration	685
11.16	Invoking <code>guix system</code>	692
11.17	Invoking <code>guix deploy</code>	701
11.18	Running Guix in a Virtual Machine	707
11.18.1	Connecting Through SSH	708
11.18.2	Using <code>virt-viewer</code> with Spice	708
11.19	Defining Services	709
11.19.1	Service Composition	709
11.19.2	Service Types and Services	710
11.19.3	Service Reference	712
11.19.4	Shepherd Services	717
11.19.5	Complex Configurations	723

## 12 System Troubleshooting Tips ..... 730

12.1	Chrooting into an existing system	730
------	-----------------------------------	-----

<b>13</b>	<b>Home Configuration</b>	<b>732</b>
13.1	Declaring the Home Environment	732
13.2	Configuring the Shell	734
13.3	Home Services	735
13.3.1	Essential Home Services	735
13.3.2	Shells	740
13.3.3	Scheduled User's Job Execution	745
13.3.4	Power Management Home Services	746
13.3.5	Managing User Daemons	747
13.3.6	Secure Shell	749
13.3.7	GNU Privacy Guard	754
13.3.8	Desktop Home Services	755
13.3.9	Guix Home Services	759
13.3.10	Backup Services	759
13.3.11	Fonts Home Services	761
13.3.12	Sound Home Services	761
13.3.13	Mail Home Services	763
13.3.14	Messaging Home Services	765
13.3.15	Media Home Services	766
13.3.16	mpv Media Player	766
13.3.17	Niri window manager	772
13.3.18	Sway window manager	772
13.3.19	Networking Home Services	779
13.3.20	Secrets Home Services	779
13.3.21	Miscellaneous Home Services	781
13.4	Invoking <code>guix home</code>	784
<b>14</b>	<b>Documentation</b>	<b>790</b>
<b>15</b>	<b>Platforms</b>	<b>791</b>
15.1	<code>platform</code> Reference	791
15.2	Supported Platforms	791
<b>16</b>	<b>Creating System Images</b>	<b>793</b>
16.1	<code>image</code> Reference	793
16.1.1	<code>partition</code> Reference	794
16.2	Instantiate an Image	795
16.3	<code>image-type</code> Reference	798
16.4	Image Modules	800
<b>17</b>	<b>Installing Debugging Files</b>	<b>801</b>
17.1	Separate Debug Info	801
17.2	Rebuilding Debug Info	802

<b>18</b>	<b>Using T<sub>E</sub>X and L<sub>A</sub>T<sub>E</sub>X .....</b>	<b>804</b>
<b>19</b>	<b>Security Updates.....</b>	<b>806</b>
<b>20</b>	<b>Bootstrapping.....</b>	<b>808</b>
20.1	The Full-Source Bootstrap .....	808
20.2	Preparing to Use the Bootstrap Binaries.....	811
	Building the Build Tools.....	812
	Building the Bootstrap Binaries .....	814
	Reducing the Set of Bootstrap Binaries .....	814
<b>21</b>	<b>Porting to a New Platform.....</b>	<b>815</b>
<b>22</b>	<b>Contributing.....</b>	<b>816</b>
22.1	Requirements .....	816
22.2	Building from Git .....	817
22.3	Running the Test Suite .....	819
22.4	Running Guix Before It Is Installed .....	821
22.5	The Perfect Setup.....	822
22.5.1	Viewing Bugs within Emacs .....	823
22.6	Alternative Setups .....	825
22.6.1	Guile Studio .....	825
22.6.2	Vim and NeoVim .....	825
22.7	Source Tree Structure.....	826
22.8	Packaging Guidelines.....	830
22.8.1	Software Freedom .....	831
22.8.2	Package Naming .....	831
22.8.3	Version Numbers.....	832
22.8.4	Synopses and Descriptions.....	833
22.8.5	Snippets versus Phases .....	834
22.8.6	Cyclic Module Dependencies .....	835
22.8.7	Bulk Updates .....	835
22.8.8	Emacs Packages.....	836
22.8.9	Python Modules .....	836
22.8.9.1	Specifying Dependencies.....	837
22.8.10	Perl Modules.....	837
22.8.11	Java Packages.....	838
22.8.12	Rust Crates.....	838
22.8.13	Elm Packages .....	839
22.8.14	Fonts .....	840
22.9	Coding Style.....	841
22.9.1	Programming Paradigm .....	841
22.9.2	Modules.....	841
22.9.3	Data Types and Pattern Matching.....	841

22.9.4	Formatting Code .....	841
22.10	Submitting Patches .....	842
22.10.1	Configuring Git .....	845
22.11	Tracking Bugs and Changes .....	845
22.11.1	The Issue Tracker .....	845
22.11.2	Managing Patches and Branches .....	845
22.11.3	Debbugs User Interfaces .....	847
22.11.3.1	Web interface .....	847
22.11.3.2	Command-Line Interface .....	847
22.11.3.3	Emacs Interface .....	849
22.11.4	Cuirass Build Notifications .....	849
22.12	Teams .....	849
22.13	Making Decisions .....	851
22.14	Commit Access .....	851
22.14.1	Applying for Commit Access .....	851
22.14.2	Commit Policy .....	852
22.14.3	Addressing Issues .....	853
22.14.3.1	Reverting commits .....	853
22.14.4	Commit Revocation .....	854
22.14.5	Helping Out .....	854
22.15	Reviewing the Work of Others .....	854
22.16	Updating the Guix Package .....	856
22.17	Deprecation Policy .....	856
22.18	Writing Documentation .....	859
22.19	Translating Guix .....	859
22.20	Contributing to Guix's Infrastructure .....	864
22.20.1	Coding .....	864
22.20.2	System Administration .....	865
22.20.3	Day-to-Day System Administration .....	865
22.20.4	On-Site Intervention .....	866
22.20.5	Hosting .....	866
22.20.6	Administrative Tasks .....	866
<b>23</b>	<b>Acknowledgments .....</b>	<b>867</b>
 <b>Appendix A GNU Free Documentation</b>		
	<b>License .....</b>	<b>868</b>
 <b>Concept Index .....</b>		
	<b>876</b>	
 <b>Programming Index .....</b>		
	<b>886</b>	

# 1 Introduction

GNU Guix<sup>1</sup> is a package management tool for and distribution of the GNU system. Guix makes it easy for unprivileged users to install, upgrade, or remove software packages, to roll back to a previous package set, to build packages from source, and generally assists with the creation and maintenance of software environments.

You can install GNU Guix on top of an existing GNU/Linux system where it complements the available tools without interference (see Chapter 2 [Installation], page 4), or you can use it as a standalone operating system distribution, *Guix System*<sup>2</sup>. See Section 1.2 [GNU Distribution], page 2.

## 1.1 Managing Software the Guix Way

Guix provides a command-line package management interface (see Chapter 5 [Package Management], page 36), tools to help with software development (see Chapter 7 [Development], page 78), command-line utilities for more advanced usage (see Chapter 9 [Utilities], page 180), as well as Scheme programming interfaces (see Chapter 8 [Programming Interface], page 100). Its *build daemon* is responsible for building packages on behalf of users (see Section 2.2 [Setting Up the Daemon], page 5) and for downloading pre-built binaries from authorized sources (see Section 5.3 [Substitutes], page 46).

Guix includes package definitions for many GNU and non-GNU packages, all of which respect the user’s computing freedom (<https://www.gnu.org/philosophy/free-sw.html>). It is *extensible*: users can write their own package definitions (see Section 8.2 [Defining Packages], page 101) and make them available as independent package modules (see Section 8.1 [Package Modules], page 100). It is also *customizable*: users can *derive* specialized package definitions from existing ones, including from the command line (see Section 9.1.2 [Package Transformation Options], page 182).

Under the hood, Guix implements the *functional package management* discipline pioneered by Nix (see Chapter 23 [Acknowledgments], page 867). In Guix, the package build and installation process is seen as a *function*, in the mathematical sense. That function takes inputs, such as build scripts, a compiler, and libraries, and returns an installed package. As a pure function, its result depends solely on its inputs—for instance, it cannot refer to software or scripts that were not explicitly passed as inputs. A build function always produces the same result when passed a given set of inputs. It cannot alter the environment of the running system in any way; for instance, it cannot create, modify, or delete files outside of its build and installation directories. This is achieved by running build processes in isolated environments (or *containers*), where only their explicit inputs are visible.

The result of package build functions is *cached* in the file system, in a special directory called *the store* (see Section 8.9 [The Store], page 156). Each package is installed in a directory of its own in the store—by default under `/gnu/store`. The directory name contains a hash of all the inputs used to build that package; thus, changing an input yields a different directory name.

<sup>1</sup> “Guix” is pronounced like “geeks”, or “iks” using the international phonetic alphabet (IPA).

<sup>2</sup> We used to refer to Guix System as “Guix System Distribution” or “GuixSD”. We now consider it makes more sense to group everything under the “Guix” banner since, after all, Guix System is readily available through the `guix system` command, even if you’re using a different distro underneath!

This approach is the foundation for the salient features of Guix: support for transactional package upgrade and rollback, per-user installation, and garbage collection of packages (see Section 5.1 [Features], page 36).

## 1.2 GNU Distribution

Guix comes with a distribution of the GNU system consisting entirely of free software<sup>3</sup>. The distribution can be installed on its own (see Chapter 3 [System Installation], page 23), but it is also possible to install Guix as a package manager on top of an installed GNU/Linux system (see Chapter 2 [Installation], page 4). When we need to distinguish between the two, we refer to the standalone distribution as Guix System.

The distribution provides core GNU packages such as GNU libc, GCC, and Binutils, as well as many GNU and non-GNU applications. The complete list of available packages can be browsed on-line (<https://www.gnu.org/software/guix/packages>) or by running `guix package` (see Section 5.2 [Invoking guix package], page 37):

```
guix package --list-available
```

Our goal is to provide a practical 100% free software distribution of Linux-based and other variants of GNU, with a focus on the promotion and tight integration of GNU components, and an emphasis on programs and tools that help users exert that freedom.

Packages are currently available on the following platforms:

### x86\_64-linux

Intel/AMD x86\_64 architecture, Linux-Libre kernel.

### i686-linux

Intel 32-bit architecture (IA32), Linux-Libre kernel.

### armhf-linux

ARMv7-A architecture with hard float, Thumb-2 and NEON, using the EABI hard-float application binary interface (ABI), and Linux-Libre kernel.

### aarch64-linux

little-endian 64-bit ARMv8-A processors, Linux-Libre kernel.

### i586-gnu

GNU/Hurd (<https://hurd.gnu.org>) on the Intel 32-bit architecture (IA32). This configuration is experimental and under development. The easiest way for you to give it a try is by setting up an instance of `hurd-vm-service-type` on your GNU/Linux machine (see [transparent-emulation-qemu], page 588). See Chapter 22 [Contributing], page 816, on how to help!

### x86\_64-gnu

GNU/Hurd (<https://hurd.gnu.org>) on the x86\_64 Intel/AMD 64-bit architecture.

This configuration is even more experimental and under heavy upstream development.

---

<sup>3</sup> The term “free” here refers to the freedom provided to users of that software (<https://www.gnu.org/philosophy/free-sw.html>).

**mips64el-linux (unsupported)**

little-endian 64-bit MIPS processors, specifically the Loongson series, n32 ABI, and Linux-Libre kernel. This configuration is no longer fully supported; in particular, there is no ongoing work to ensure that this architecture still works. Should someone decide they wish to revive this architecture then the code is still available.

**powerpc-linux (unsupported)**

big-endian 32-bit PowerPC processors, specifically the PowerPC G4 with AltiVec support, and Linux-Libre kernel. This configuration is not fully supported and there is no ongoing work to ensure this architecture works.

**powerpc64le-linux**

little-endian 64-bit Power ISA processors, Linux-Libre kernel. This includes POWER9 systems such as the RYF Talos II mainboard (<https://www.fsf.org/news/talos-ii-mainboard-and-talos-ii-lite-mainboard-now-fsf-certified-to-resp>). This platform is available as a "technology preview": although it is supported, substitutes are not yet available from the build farm (see Section 5.3 [Substitutes], page 46), and some packages may fail to build (see Section 22.11 [Tracking Bugs and Changes], page 845). That said, the Guix community is actively working on improving this support, and now is a great time to try it and get involved!

**riscv64-linux**

little-endian 64-bit RISC-V processors, specifically RV64GC, and Linux-Libre kernel. This platform is available as a "technology preview": although it is supported, substitutes are not yet available from the build farm (see Section 5.3 [Substitutes], page 46), and some packages may fail to build (see Section 22.11 [Tracking Bugs and Changes], page 845). That said, the Guix community is actively working on improving this support, and now is a great time to try it and get involved!

With Guix System, you *declare* all aspects of the operating system configuration and Guix takes care of instantiating the configuration in a transactional, reproducible, and stateless fashion (see Chapter 11 [System Configuration], page 242). Guix System uses the Linux-libre kernel, the Shepherd initialization system (see Section “Introduction” in *The GNU Shepherd Manual*), the well-known GNU utilities and tool chain, as well as the graphical environment or system services of your choice.

Guix System is available on all the above platforms except **mips64el-linux**, **powerpc-linux**, **powerpc64le-linux** and **riscv64-linux**.

For information on porting to other architectures or kernels, see Chapter 21 [Porting], page 815.

Building this distribution is a cooperative effort, and you are invited to join! See Chapter 22 [Contributing], page 816, for information about how you can help.

## 2 Installation

You can install the package management tool Guix on top of an existing GNU/Linux or GNU/Hurd system<sup>1</sup>, referred to as a *foreign distro*. If, instead, you want to install the complete, standalone GNU system distribution, *Guix System*, see Chapter 3 [System Installation], page 23. This section is concerned only with the installation of Guix on a foreign distro.

**Important:** This section only applies to systems without Guix. Following it for existing Guix installations will overwrite important system files.

When installed on a foreign distro, GNU Guix complements the available tools without interference. Its data lives exclusively in two directories, usually `/gnu/store` and `/var/guix`; other files on your system, such as `/etc`, are left untouched.

Once installed, Guix can be updated by running `guix pull` (see Section 5.7 [Invoking guix pull], page 56).

### 2.1 Binary Installation

This section describes how to install Guix from a self-contained tarball providing binaries for Guix and for all its dependencies. This is often quicker than installing from source, described later (see Section 22.2 [Building from Git], page 817).

**Important:** This section only applies to systems without Guix. Following it for existing Guix installations will overwrite important system files.

The `guix-install.sh` (<https://guix.gnu.org/guix-install.sh>) shell script automates the binary installation process. Use of `guix-install.sh` requires Bash, GnuPG, GNU tar, wget, and Xz. On systems where systemd is detected, the tool `newgidmap` from the `uidmap` package is also required.

To install Guix on top of your GNU/Linux system, run the following commands as root:

```
cd /tmp
wget https://guix.gnu.org/guix-install.sh
chmod +x guix-install.sh
./guix-install.sh
```

When you're done installing Guix, see Section 2.4 [Application Setup], page 18, for extra configuration you might need, and Chapter 4 [Getting Started], page 33, for your first steps!

The script goes through the following steps:

1. Downloading the binary tarball from '[https://ftpmirror.gnu.org/gnu/guix/guix-binary-8b23207.x86\\_64-linux](https://ftpmirror.gnu.org/gnu/guix/guix-binary-8b23207.x86_64-linux.tar.gz)' (where `x86_64-linux` can be replaced with `i686-linux` for an i686 (32-bits) machine already running the kernel Linux, and so on—see Section 1.2 [GNU Distribution], page 2) and extracting it.
2. Setting up the build daemon (see Section 2.2 [Setting Up the Daemon], page 5).
3. Making the `guix` command available to non-root users.
4. Configuring substitute servers (see Section 5.3.1 [Official Substitute Servers], page 47).

---

<sup>1</sup> Hurd support is currently limited.



**Note:** By default, `guix-install.sh` will configure Guix to download pre-built package binaries, called *substitutes* (see Section 5.3 [Substitutes], page 46), from the project’s build farms. If you choose not to permit this, Guix will build *everything* from source, making each installation and upgrade very expensive. See Section 5.3.7 [On Trusting Binaries], page 50, for a discussion of why you may want to build packages from source.

To use substitutes from `bordeaux.guix.gnu.org`, `ci.guix.gnu.org` or a mirror, you must authorize them. For example,

```
guix archive --authorize < \
  ~root/.config/guix/current/share/guix/bordeaux.guix.gnu.org.pub
guix archive --authorize < \
  ~root/.config/guix/current/share/guix/ci.guix.gnu.org.pub
```

**Note:** The binary installation tarball can be (re)produced and verified simply by running the following command in the Guix source tree:

```
make guix-binary.system.tar.xz
```

... which, in turn, runs:

```
guix pack -s system --localstatedir \
  --profile-name=current-guix guix
```

See Section 7.3 [Invoking `guix pack`], page 90, for more info on this handy tool.

Should you eventually want to uninstall Guix, run the same script with the `--uninstall` flag:

```
./guix-install.sh --uninstall
```

With `--uninstall`, the script irreversibly deletes all the Guix files, configuration, and services.

Some GNU/Linux distributions, such as openSUSE, Alpine and NixOS provide Guix through their own package managers. The version of Guix may be older than 8b23207 but you can update it afterwards by running `guix pull`.

We advise system administrators who install Guix, both from the installation script or *via* the native package manager of their foreign distribution, to also regularly read and follow security notices, as shown by `guix pull`.

## 2.2 Setting Up the Daemon

During installation, the *build daemon* that must be running to use Guix has already been set up and you can run `guix` commands in your terminal program, see Chapter 4 [Getting Started], page 33:

```
guix build hello
```

If this runs through without error, feel free to skip this section. You should continue with the following section, Section 2.4 [Application Setup], page 18.

However, now would be a good time to replace outdated daemon versions, tweak it, perform builds on other machines (see Section 2.2.2 [Daemon Offload Setup], page 9), or start it manually in special environments like “chroots” (see Section 12.1 [Chrooting into an existing system], page 730) or WSL (not needed for WSL images created with Guix, see

Chapter 16 [System Images], page 793). If you want to know more or optimize your system, this section is worth reading.

Operations such as building a package or running the garbage collector are all performed by a specialized process, the build daemon, on behalf of clients. Only the daemon may access the store and its associated database. Thus, any operation that manipulates the store goes through the daemon. For instance, command-line tools such as **guix package** and **guix build** communicate with the daemon (*via* remote procedure calls) to instruct it what to do.

The following sections explain how to prepare the build daemon’s environment. See Section 5.3 [Substitutes], page 46, for how to allow the daemon to download pre-built binaries.

### 2.2.1 Build Environment Setup

In a standard multi-user setup, Guix and its daemon—the **guix-daemon** program—are installed by the system administrator. Unprivileged users may use Guix tools to build packages or otherwise access the store, and the daemon will do it on their behalf, ensuring that the store is kept in a consistent state, and allowing built packages to be shared among users.

There are currently two ways to set up and run the build daemon:

1. running **guix-daemon** as “root”, letting it run build processes as unprivileged users taken from a pool of build users—this is the historical approach;
2. running **guix-daemon** as a separate unprivileged user, relying on Linux’s *unprivileged user namespace* functionality to set up isolated environments—this is the option chosen when installing Guix on a systemd-based distribution with the installation script (see Section 2.1 [Binary Installation], page 4).

The sections below describe each of these two configurations in more detail and summarize the kind of build isolation they provide.

#### Daemon Running as Root

When **guix-daemon** runs as **root**, you may not want package build processes themselves to run as **root** too, for obvious security reasons. To avoid that, a special pool of *build users* should be created for use by build processes started by the daemon. Having several such users allows the daemon to launch distinct build processes under separate UIDs, which guarantees that they do not interfere with each other—an essential feature since builds are regarded as pure functions (see Chapter 1 [Introduction], page 1).

On a GNU/Linux system, a build user pool may be created like this (using Bash syntax and the **shadow** commands):

```
# groupadd --system guixbuild
# for i in $(seq -w 1 10);
do
    useradd -g guixbuild -G guixbuild          \
          -d /var/empty -s $(which nologin)    \
          -c "Guix build user $i" --system    \
          guixbuilder$i;
done
```

The number of build users determines how many build jobs may run in parallel, as specified by the `--max-jobs` option (see Section 2.3 [Invoking guix-daemon], page 14). To use `guix system vm` and related commands, you may need to add the build users to the `kvm` group so they can access `/dev/kvm`, using `-G guixbuild,kvm` instead of `-G guixbuild` (see Section 11.16 [Invoking guix system], page 692).

The `guix-daemon` program may then be run as `root` with the following command<sup>2</sup>:

```
# guix-daemon --build-users-group=guixbuild
```

In this setup, `/gnu/store` is owned by `root`.

## Daemon Running Without Privileges

The second and preferred option is to run `guix-daemon` *as an unprivileged user*. It has the advantage of reducing the harm that can be done should a build process manage to exploit a vulnerability in the daemon. This option requires the use of Linux's unprivileged user namespace mechanism; today it is available and enabled by most GNU/Linux distributions but can still be disabled. The installation script automatically determines whether this option is available on your system (see Section 2.1 [Binary Installation], page 4).

When using this option, you only need to create one user account, and `guix-daemon` will run with the authority of that account:

```
# groupadd --system guix-daemon
# useradd -g guix-daemon -G guix-daemon \
        -d /var/empty -s $(which nologin) \
        -c "Guix daemon privilege separation user" \
        --system guix-daemon
```

**Warning:** If the `nologin` binary is not found by `which nologin`, you may need to specify its path manually. This will usually be `/sbin/nologin` or `/usr/sbin/nologin`.

In this configuration, `/gnu/store` is owned by the `guix-daemon` user.

## Migrating to the Unprivileged Daemon

To switch an existing installation to the unprivileged execution mode, a number of steps must be taken: creating a new dedicated `guix-daemon` user account, changing ownership of the relevant files to `guix-daemon`, and ensuring that the `guix-daemon` program runs as `guix-daemon`.

On Guix System, these steps are carried out automatically when you set the `privileged?` field of the `guix-configuration` record to `#f` and reconfigure (see [guix-configuration-type], page 285).

However, on a foreign distribution, the process is manual. The following paragraphs describe what you need to do.

**Warning:** Follow the instructions below only after making sure you have a recent version of `guix-daemon` with support for unprivileged execution.

---

<sup>2</sup> If your machine uses the `systemd` init system, copying the `prefix/lib/systemd/system/guix-daemon.service` file to `/etc/systemd/system` will ensure that `guix-daemon` is automatically started. Similarly, if your machine uses the `Upstart` init system, copy the `prefix/lib/upstart/system/guix-daemon.conf` file to `/etc/init`.

File ownership can be changed, after stopping the daemon, by running the following commands as root (the `chown` can take a while if there are many files in `/gnu/store`):

```
groupadd --system guix-daemon
useradd -g guix-daemon -G guix-daemon,kvm          \
        -d /var/empty -s $(which nologin)          \
        -c "Guix daemon privilege separation user" \
        --system guix-daemon

# Make the store writable, in case the systemd 'gnu-store.mount'
# unit made it read-only.
mount -o remount,rw /gnu/store

chown -R guix-daemon:guix-daemon \
    /gnu \
    /var/guix/{gc.lock,daemon-socket,db,discover} \
    /var/guix/{gcroots,offload,substitute,temproots} \
    /var/log/guix \
    /etc/guix
```

If your system uses the `systemd` service manager, running the daemon as `guix-daemon` will be a matter of copying the relevant configuration files—make sure to review any changes you might have made in your own `.service` files before overwriting them:

```
cp /var/guix/profiles/per-user/root/current-guix/lib/systemd/system/*.service \
    /etc/systemd/system
systemctl daemon-reload
systemctl start guix-daemon
```

If your system has AppArmor enabled, see Section 2.2.4 [AppArmor Support], page 13.

**Warning:** The commands above assume that `guix pull` was run for the root user. You can check whether this is the case by running this command:

```
grep User=guix-daemon \
    /var/guix/profiles/per-user/root/current-guix/lib/systemd/system/guix-daemon.se
```

If that command does not show the `User=guix-daemon` line, then run `guix pull` as the root user.

## The Isolated Build Environment

In both cases, privileged and unprivileged, the daemon starts build processes without privileges in an *isolated* or *hermetic* build environment—a “chroot”. On GNU/Linux, by default, the build environment contains nothing but:

- a minimal `/dev` directory, created mostly independently from the host `/dev`<sup>3</sup>;
- the `/proc` directory; it only shows the processes of the container since a separate PID name space is used;
- `/etc/passwd` with an entry for the current user and an entry for user `nobody`;
- `/etc/group` with an entry for the user’s group;

<sup>3</sup> “Mostly”, because while the set of files that appear in the chroot’s `/dev` is fixed, most of these files can only be created if the host has them.

- `/etc/hosts` with an entry that maps `localhost` to `127.0.0.1`;
- a writable `/tmp` directory.

The chroot does not contain a `/home` directory, and the `HOME` environment variable is set to the non-existent `/homeless-shelter`. This helps to highlight inappropriate uses of `HOME` in the build scripts of packages.

All this is usually enough to ensure details of the environment do not influence build processes. In some exceptional cases where more control is needed—typically over the date, kernel, or CPU—you can resort to a virtual build machine (see [build-vm], page 591).

You can influence the directory where the daemon stores build trees *via* the `TMPDIR` environment variable. However, the build tree within the chroot is always called `/tmp/guix-build-name.drv-0`, where *name* is the derivation name—e.g., `coreutils-8.24`. This way, the value of `TMPDIR` does not leak inside build environments, which avoids discrepancies in cases where build processes capture the name of their build tree.

The daemon also honors the `http_proxy` and `https_proxy` environment variables for HTTP and HTTPS downloads it performs, be it for fixed-output derivations (see Section 8.10 [Derivations], page 158) or for substitutes (see Section 5.3 [Substitutes], page 46).

### 2.2.2 Using the Offload Facility

When desired, the build daemon can *offload* derivation builds to other machines running Guix, using the `offload build hook`<sup>4</sup>. When that feature is enabled, a list of user-specified build machines is read from `/etc/guix/machines.scm`; every time a build is requested, for instance via `guix build`, the daemon attempts to offload it to one of the machines that satisfy the constraints of the derivation, in particular its system types—e.g., `x86_64-linux`. A single machine can have multiple system types, either because its architecture natively supports it, via emulation (see [transparent-emulation-qemu], page 588), or both. Missing prerequisites for the build are copied over SSH to the target machine, which then proceeds with the build; upon success the output(s) of the build are copied back to the initial machine. The offload facility comes with a basic scheduler that attempts to select the best machine. The best machine is chosen among the available machines based on criteria such as:

1. The availability of a build slot. A build machine can have as many build slots (connections) as the value of the `parallel-builds` field of its `build-machine` object.
2. Its relative speed, as defined via the `speed` field of its `build-machine` object.
3. Its load. The normalized machine load must be lower than a threshold value, configurable via the `overload-threshold` field of its `build-machine` object.
4. Disk space availability. More than 100 MiB must be available.

The `/etc/guix/machines.scm` file typically looks like this:

```
(list (build-machine
      (name "eightysix.example.org")
      (systems (list "x86_64-linux" "i686-linux"))
      (host-key "ssh-ed25519 AAAAC3Nza...")
      (user "bob"))
```

---

<sup>4</sup> This feature is available only when Guile-SSH (<https://github.com/artiom-poptsov/guile-ssh>) is present.

```

(speed 2.))      ;incredibly fast!

(build-machine
  (name "armeight.example.org")
  (systems (list "aarch64-linux"))
  (host-key "ssh-rsa AAAAB3Nza...")
  (user "alice")

  ;; Remember 'guix offload' is spawned by
  ;; 'guix-daemon' as root.
  (private-key "/root/.ssh/identity-for-guix")))

```

In the example above we specify a list of two build machines, one for the `x86_64` and `i686` architectures and one for the `aarch64` architecture.

In fact, this file is—not surprisingly!—a Scheme file that is evaluated when the `offload` hook is started. Its return value must be a list of `build-machine` objects. While this example shows a fixed list of build machines, one could imagine, say, using DNS-SD to return a list of potential build machines discovered in the local network (see Section “Introduction” in *Using Avahi in Guile Scheme Programs*). The `build-machine` data type is detailed below.

#### `build-machine` [Data Type]

This data type represents build machines to which the daemon may offload builds. The important fields are:

<b>name</b>	The host name of the remote machine.
<b>systems</b>	The system types the remote machine supports—e.g., <code>(list "x86_64-linux" "i686-linux")</code> .
<b>user</b>	The user account on the remote machine to use when connecting over SSH. Note that the SSH key pair must <i>not</i> be passphrase-protected, to allow non-interactive logins.
<b>host-key</b>	This must be the machine’s SSH <i>public host key</i> in OpenSSH format. This is used to authenticate the machine when we connect to it. It is a long string that looks like this:

```
ssh-ed25519 AAAAC3NzaC...mde+Uhl hint@example.org
```

If the machine is running the OpenSSH daemon, `sshd`, the host key can be found in a file such as `/etc/ssh/ssh_host_ed25519_key.pub`.

A number of optional fields may be specified:

<b>port</b> (default: 22)	Port number of SSH server on the machine.
<b>private-key</b> (default: <code>~root/.ssh/id_rsa</code> )	The SSH private key file to use when connecting to the machine, in OpenSSH format. This key must not be protected with a passphrase. Note that the default value is the private key <i>of the root account</i> . Make sure it exists if you use the default.

`compression` (default: "zlib@openssh.com,zlib")

`compression-level` (default: 3)

The SSH-level compression methods and compression level requested.

Note that offloading relies on SSH compression to reduce bandwidth usage when transferring files to and from build machines.

`daemon-socket` (default: "/var/guix/daemon-socket/socket")

File name of the Unix-domain socket `guix-daemon` is listening to on that machine.

`overload-threshold` (default: 0.8)

The load threshold above which a potential offload machine is disregarded by the offload scheduler. The value roughly translates to the total processor usage of the build machine, ranging from 0.0 (0%) to 1.0 (100%). It can also be disabled by setting `overload-threshold` to `#f`.

`parallel-builds` (default: 1)

The number of builds that may run in parallel on the machine.

`speed` (default: 1.0)

A “relative speed factor”. The offload scheduler will tend to prefer machines with a higher speed factor.

`features` (default: '()')

A list of strings denoting specific features supported by the machine. An example is "kvm" for machines that have the KVM Linux modules and corresponding hardware support. Derivations can request features by name, and they will be scheduled on matching build machines.

**Note:** On Guix System, instead of managing `/etc/guix/machines.scm` independently, you can choose to specify build machines directly in the `operating-system` declaration, in the `build-machines` field of `guix-configuration`. See [guix-configuration-build-machines], page 288.

The `guix` command must be in the search path on the build machines. You can check whether this is the case by running:

```
ssh build-machine guix repl --version
```

There is one last thing to do once `machines.scm` is in place. As explained above, when offloading, files are transferred back and forth between the machine stores. For this to work, you first need to generate a key pair on each machine to allow the daemon to export signed archives of files from the store (see Section 5.11 [Invoking guix archive], page 65):

```
# guix archive --generate-key
```

**Note:** This key pair is not related to the SSH key pair that was previously mentioned in the description of the `build-machine` data type.

Each build machine must authorize the key of the master machine so that it accepts store items it receives from the master:

```
# guix archive --authorize < master-public-key.txt
```

Likewise, the master machine must authorize the key of each build machine.

All the fuss with keys is here to express pairwise mutual trust relations between the master and the build machines. Concretely, when the master receives files from a build machine (and *vice versa*), its build daemon can make sure they are genuine, have not been tampered with, and that they are signed by an authorized key.

To test whether your setup is operational, run this command on the master node:

```
# guix offload test
```

This will attempt to connect to each of the build machines specified in `/etc/guix/machines.scm`, make sure Guix is available on each machine, attempt to export to the machine and import from it, and report any error in the process.

If you want to test a different machine file, just specify it on the command line:

```
# guix offload test machines-qualif.scm
```

Last, you can test the subset of the machines whose name matches a regular expression like this:

```
# guix offload test machines.scm '\.gnu\.org$'
```

To display the current load of all build hosts, run this command on the main node:

```
# guix offload status
```

## 2.2.3 SELinux Support

Guix includes an SELinux policy file at `etc/guix-daemon.cil` that can be installed on a system where SELinux is enabled, in order to label Guix files and to specify the expected behavior of the daemon. Since Guix System does not provide an SELinux base policy, the daemon policy cannot be used on Guix System.

### 2.2.3.1 Installing the SELinux policy

**Note:** The `guix-install.sh` binary installation script offers to perform the steps below for you (see Section 2.1 [Binary Installation], page 4).

To install the policy run this command as root:

```
semodule -i /var/guix/profiles/per-user/root/current-guix/share/selinux/guix-daemon.cil
```

Then, as root, relabel the file system, possibly after making it writable:

```
mount -o remount,rw /gnu/store
restorecon -R /gnu /var/guix
```

At this point you can start or restart `guix-daemon`; on a distribution that uses `systemd` as its service manager, you can do that with:

```
systemctl restart guix-daemon
```

Once the policy is installed, the file system has been relabeled, and the daemon has been restarted, it should be running in the `guix_daemon_t` context. You can confirm this with the following command:

```
ps -Zax | grep guix-daemon
```

Monitor the SELinux log files as you run a command like `guix build hello` to convince yourself that SELinux permits all necessary operations.



### 2.2.3.2 Limitations

This policy is not perfect. Here is a list of limitations or quirks that should be considered when deploying the provided SELinux policy for the Guix daemon.

1. `guix_daemon_socket_t` isn't actually used. None of the socket operations involve contexts that have anything to do with `guix_daemon_socket_t`. It doesn't hurt to have this unused label, but it would be preferable to define socket rules for only this label.
2. `guix gc` cannot access arbitrary links to profiles. By design, the file label of the destination of a symlink is independent of the file label of the link itself. Although all profiles under `$localstatedir` are labelled, the links to these profiles inherit the label of the directory they are in. For links in the user's home directory this will be `user_home_t`. But for links from the root user's home directory, or `/tmp`, or the HTTP server's working directory, etc, this won't work. `guix gc` would be prevented from reading and following these links.
3. The daemon's feature to listen for TCP connections might no longer work. This might require extra rules, because SELinux treats network sockets differently from files.
4. Currently all files with a name matching the regular expression `/gnu/store/.+-(guix-.|profile)/bin/guix-daemon` are assigned the label `guix_daemon_exec_t`; this means that *any* file with that name in any profile would be permitted to run in the `guix_daemon_t` domain. This is not ideal. An attacker could build a package that provides this executable and convince a user to install and run it, which lifts it into the `guix_daemon_t` domain. At that point SELinux could not prevent it from accessing files that are allowed for processes in that domain.

You will need to relabel the store directory after all upgrades to `guix-daemon`, such as after running `guix pull`. Assuming the store is in `/gnu`, you can do this with `restorecon -vR /gnu`, or by other means provided by your operating system.

We could generate a much more restrictive policy at installation time, so that only the *exact* file name of the currently installed `guix-daemon` executable would be labelled with `guix_daemon_exec_t`, instead of using a broad regular expression. The downside is that root would have to install or upgrade the policy at installation time whenever the Guix package that provides the effectively running `guix-daemon` executable is upgraded.

### 2.2.4 AppArmor Support

Guix includes an AppArmor profile for the build daemon in `etc/apparmor.d/guix-daemon` that can be installed on systems with strict AppArmor policies to allow it to run unprivileged (see Section 2.2.1 [Build Environment Setup], page 6). Indeed, the unprivileged daemon makes use of Linux user namespaces but these are disallowed without an AppArmor policy on some systems like Ubuntu.

To know if this applies to you, check if the `kernel.apparmor_restrict_unprivileged_userns` kernel parameter is enabled.

#### 2.2.4.1 Installing the AppArmor profile

**Note:** The `guix-install.sh` binary installation script offers to perform the steps below for you (see Section 2.1 [Binary Installation], page 4).

Run these commands as root to install the profile:

```
export apparmor_sources=/var/guix/profiles/per-user/root/current-guix/etc/apparmor.d
cp -f -t /etc/apparmor.d/tunables "$apparmor_sources/tunables/guix"
cp -f -t /etc/apparmor.d "$apparmor_sources/guix-daemon"
cp -f -t /etc/apparmor.d "$apparmor_sources/guix"
apparmor_parser -r /etc/apparmor.d/guix-daemon
apparmor_parser -r /etc/apparmor.d/guix
```

After this, the build daemon will be able to function correctly.

## 2.3 Invoking guix-daemon

The `guix-daemon` program implements all the functionality to access the store. This includes launching build processes, running the garbage collector, querying the availability of a build result, etc. It is normally run as `root` like this:

```
# guix-daemon --build-users-group=guixbuild
```

This daemon can also be started following the systemd “socket activation” protocol (see Section “Service De- and Constructors” in *The GNU Shepherd Manual*).

For details on how to set it up, see Section 2.2 [Setting Up the Daemon], page 5.

By default, `guix-daemon` launches build processes under different UIDs, taken from the build group specified with `--build-users-group`. In addition, each build process is run in a chroot environment that only contains the subset of the store that the build process depends on, as specified by its derivation (see Chapter 8 [Programming Interface], page 100), plus a set of specific system directories. By default, the latter contains `/dev` and `/dev/pts`. Furthermore, on GNU/Linux, the build environment is a *container*: in addition to having its own file system tree, it has a separate mount name space, its own PID name space, network name space, etc. This helps achieve reproducible builds (see Section 5.1 [Features], page 36).

When the daemon performs a build on behalf of the user, it creates a build directory under `/tmp` or under the directory specified by its `TMPDIR` environment variable. This directory is shared with the container for the duration of the build, though within the container, the build tree is always called `/tmp/guix-build-name.drv-0`.

The build directory is automatically deleted upon completion, unless the build failed and the client specified `--keep-failed` (see Section 9.1.1 [Common Build Options], page 180).

The daemon listens for connections and spawns one sub-process for each session started by a client (one of the `guix` sub-commands). The `guix processes` command allows you to get an overview of the activity on your system by viewing each of the active sessions and clients. See Section 9.16 [Invoking guix processes], page 237, for more information.

The following command-line options are supported:

**`--build-users-group=group`**

Take users from *group* to run build processes (see Section 2.2 [Setting Up the Daemon], page 5).

**`--no-substitutes`**

Do not use substitutes for build products. That is, always build things locally instead of allowing downloads of pre-built binaries (see Section 5.3 [Substitutes], page 46).

When the daemon runs with `--no-substitutes`, clients can still explicitly enable substitution *via* the `set-build-options` remote procedure call (see Section 8.9 [The Store], page 156).

**--substitute-urls=urls**

Consider *urls* the default whitespace-separated list of substitute source URLs. When this option is omitted, `'https://bordeaux.guix.gnu.org https://ci.guix.gnu.org'` is used.

This means that substitutes may be downloaded from *urls*, as long as they are signed by a trusted signature (see Section 5.3 [Substitutes], page 46).

See Section 5.3.3 [Getting Substitutes from Other Servers], page 48, for more information on how to configure the daemon to get substitutes from other servers.

**--no-offload**

Do not use offload builds to other machines (see Section 2.2.2 [Daemon Offload Setup], page 9). That is, always build things locally instead of offloading builds to remote machines.

**--cache-failures**

Cache build failures. By default, only successful builds are cached.

When this option is used, `guix gc --list-failures` can be used to query the set of store items marked as failed; `guix gc --clear-failures` removes store items from the set of cached failures. See Section 5.6 [Invoking guix gc], page 53.

**--cores=n**

**-c n** Use *n* CPU cores to build each derivation; 0 means as many as available.

The default value is 0, but it may be overridden by clients, such as the `--cores` option of `guix build` (see Section 9.1 [Invoking guix build], page 180).

The effect is to define the `NIX_BUILD_CORES` environment variable in the build process, which can then use it to exploit internal parallelism—for instance, by running `make -j$NIX_BUILD_CORES`.

**--max-jobs=n**

**-M n** Allow at most *n* build jobs in parallel. The default value is 1. Setting it to 0 means that no builds will be performed locally; instead, the daemon will offload builds (see Section 2.2.2 [Daemon Offload Setup], page 9), or simply fail.

**--max-silent-time=seconds**

When the build or substitution process remains silent for more than *seconds*, terminate it and report a build failure.

The default value is 3600 (one hour).

The value specified here can be overridden by clients (see Section 9.1.1 [Common Build Options], page 180).

**--timeout=seconds**

Likewise, when the build or substitution process lasts for more than *seconds*, terminate it and report a build failure.

The default value is 24 hours.

The value specified here can be overridden by clients (see Section 9.1.1 [Common Build Options], page 180).

**--rounds=*N***

Build each derivation *n* times in a row, and raise an error if consecutive build results are not bit-for-bit identical. Note that this setting can be overridden by clients such as **guix build** (see Section 9.1 [Invoking guix build], page 180).

When used in conjunction with **--keep-failed**, the differing output is kept in the store, under **/gnu/store/...-check**. This makes it easy to look for differences between the two results.

**--debug** Produce debugging output.

This is useful to debug daemon start-up issues, but then it may be overridden by clients, for example the **--verbosity** option of **guix build** (see Section 9.1 [Invoking guix build], page 180).

**--chroot-directory=*dir***

Add *dir* to the build chroot.

Doing this may change the result of build processes—for instance if they use optional dependencies found in *dir* when it is available, and not otherwise. For that reason, it is not recommended to do so. Instead, make sure that each derivation declares all the inputs that it needs.

**--disable-chroot**

Disable chroot builds.

**Warning:** Using this option is not recommended since it allows build processes to gain access to undeclared dependencies, to interfere with one another, and more generally to do anything that can be done with the authority of build users or that of the daemon—which includes at least the ability to tamper with any file in the store!

You may find it necessary, though, when support for Linux unprivileged user namespaces is missing (see Section 2.2.1 [Build Environment Setup], page 6). Use at your own risk!

**--isolate-host-loopback**

Prevent fixed-output chroot builds from accessing the host loopback. When this is used, fixed-output builds will have their own loopback interface, and attempts to connect to 127.0.0.1 or ::1 will instead attempt to connect to its local loopback interface.

**--log-compression=*type***

Compress build logs according to *type*, one of **gzip**, **bzip2**, or **none**.

Unless **--lose-logs** is used, all the build logs are kept in the *localstatedir*. To save space, the daemon automatically compresses them with gzip by default.

**--discover[=*yes|no*]**

Whether to discover substitute servers on the local network using mDNS and DNS-SD.

This feature is still experimental. However, here are a few considerations.

1. It might be faster/less expensive than fetching from remote servers;

2. There are no security risks, only genuine substitutes will be used (see Section 5.3.4 [Substitute Authentication], page 49);
3. An attacker advertising `guix publish` on your LAN cannot serve you malicious binaries, but they can learn what software you're installing;
4. Servers may serve substitute over HTTP, unencrypted, so anyone on the LAN can see what software you're installing.

It is also possible to enable or disable substitute server discovery at run-time by running:

```
herd discover guix-daemon on
herd discover guix-daemon off
```

#### `--disable-deduplication`

Disable automatic file “deduplication” in the store.

By default, files added to the store are automatically “deduplicated”: if a newly added file is identical to another one found in the store, the daemon makes the new file a hard link to the other file. This can noticeably reduce disk usage, at the expense of slightly increased input/output load at the end of a build process. This option disables this optimization.

#### `--gc-keep-outputs[=yes|no]`

Tell whether the garbage collector (GC) must keep outputs of live derivations.

When set to `yes`, the GC will keep the outputs of any live derivation available in the store—the `.drv` files. The default is `no`, meaning that derivation outputs are kept only if they are reachable from a GC root. See Section 5.6 [Invoking `guix gc`], page 53, for more on GC roots.

#### `--gc-keep-derivations[=yes|no]`

Tell whether the garbage collector (GC) must keep derivations corresponding to live outputs.

When set to `yes`, as is the case by default, the GC keeps derivations—i.e., `.drv` files—as long as at least one of their outputs is live. This allows users to keep track of the origins of items in their store. Setting it to `no` saves a bit of disk space.

In this way, setting `--gc-keep-derivations` to `yes` causes liveness to flow from outputs to derivations, and setting `--gc-keep-outputs` to `yes` causes liveness to flow from derivations to outputs. When both are set to `yes`, the effect is to keep all the build prerequisites (the sources, compiler, libraries, and other build-time tools) of live objects in the store, regardless of whether these prerequisites are reachable from a GC root. This is convenient for developers since it saves rebuilds or downloads.

#### `--impersonate-linux-2.6`

On Linux-based systems, impersonate Linux 2.6. This means that the kernel's `uname` system call will report 2.6 as the release number.

This might be helpful to build programs that (usually wrongfully) depend on the kernel version number.

**--lose-logs**  
Do not keep build logs. By default they are kept under *localstatedir/guix/log*.

**--system=system**  
Assume *system* as the current system type. By default it is the architecture/kernel pair found at configure time, such as *x86\_64-linux*.

**--listen=endpoint**  
Listen for connections on *endpoint*. *endpoint* is interpreted as the file name of a Unix-domain socket if it starts with / (slash sign). Otherwise, *endpoint* is interpreted as a host name or host name and port to listen to. Here are a few examples:

**--listen=/gnu/var/daemon**  
Listen for connections on the */gnu/var/daemon* Unix-domain socket, creating it if needed.

**--listen=localhost**  
Listen for TCP connections on the network interface corresponding to *localhost*, on port 44146.

**--listen=128.0.0.42:1234**  
Listen for TCP connections on the network interface corresponding to *128.0.0.42*, on port 1234.

This option can be repeated multiple times, in which case **guix-daemon** accepts connections on all the specified endpoints. Users can tell client commands what endpoint to connect to by setting the **GUIX\_DAEMON\_SOCKET** environment variable (see Section 8.9 [The Store], page 156).

**Note:** The daemon protocol is *unauthenticated and unencrypted*. Using **--listen=host** is suitable on local networks, such as clusters, where only trusted nodes may connect to the build daemon. In other cases where remote access to the daemon is needed, we recommend using Unix-domain sockets along with SSH.

When **--listen** is omitted, **guix-daemon** listens for connections on the Unix-domain socket located at *localstatedir/guix/daemon-socket/socket*.

## 2.4 Application Setup

When using Guix on top of GNU/Linux distribution other than Guix System—a so-called *foreign distro*—a few additional steps are needed to get everything in place. Here are some of them.

### 2.4.1 Locales

Packages installed *via* Guix will not use the locale data of the host system. Instead, you must first install one of the locale packages available with Guix and then define the **GUIX\_LOCPATH** environment variable:

```
$ guix install glibc-locales
$ export GUIX_LOCPATH=$HOME/.guix-profile/lib/locale
```

Note that the `glibc-locales` package contains data for all the locales supported by the GNU libc and weighs in at around 930 MiB<sup>5</sup>. If you only need a few locales, you can define your custom locales package via the `make-glibc-utf8-locales` procedure from the `(gnu packages base)` module. The following example defines a package containing the various Canadian UTF-8 locales known to the GNU libc, that weighs around 14 MiB:

```
(use-modules (gnu packages base))

(define my-glibc-locales
  (make-glibc-utf8-locales
    glibc
    #:locales (list "en_CA" "fr_CA" "ik_CA" "iu_CA" "shs_CA")
    #:name "glibc-canadian-utf8-locales"))
```

The `GUIX_LOCPATH` variable plays a role similar to `LOCPATH` (see Section “Locale Names” in *The GNU C Library Reference Manual*). There are two important differences though:

1. `GUIX_LOCPATH` is honored only by the libc in Guix, and not by the libc provided by foreign distros. Thus, using `GUIX_LOCPATH` allows you to make sure the programs of the foreign distro will not end up loading incompatible locale data.
2. libc suffixes each entry of `GUIX_LOCPATH` with `/X.Y`, where `X.Y` is the libc version—e.g., 2.22. This means that, should your Guix profile contain a mixture of programs linked against different libc version, each libc version will only try to load locale data in the right format.

This is important because the locale data format used by different libc versions may be incompatible.

### 2.4.2 Name Service Switch

When using Guix on a foreign distro, we *strongly recommend* that the system run the GNU C library’s *name service cache daemon*, `nsd`, which should be listening on the `/var/run/nsd/socket` socket. Failing to do that, applications installed with Guix may fail to look up host names or user accounts, or may even crash. The next paragraphs explain why.

The GNU C library implements a *name service switch* (NSS), which is an extensible mechanism for “name lookups” in general: host name resolution, user accounts, and more (see Section “Name Service Switch” in *The GNU C Library Reference Manual*).

Being extensible, the NSS supports *plugins*, which provide new name lookup implementations: for example, the `nss-mdns` plugin allow resolution of `.local` host names, the `nis` plugin allows user account lookup using the Network information service (NIS), and so on. These extra “lookup services” are configured system-wide in `/etc/nsswitch.conf`, and all the programs running on the system honor those settings (see Section “NSS Configuration File” in *The GNU C Reference Manual*).

When they perform a name lookup—for instance by calling the `getaddrinfo` function in C—applications first try to connect to the `nsd`; on success, `nsd` performs name lookups on their behalf. If the `nsd` is not running, then they perform the name lookup by themselves,

<sup>5</sup> The size of the `glibc-locales` package is reduced down to about 213 MiB with store deduplication and further down to about 67 MiB when using a `zstd`-compressed Btrfs file system.

by loading the name lookup services into their own address space and running it. These name lookup services—the `libnss_*.so` files—are `dlopen`'d, but they may come from the host system's C library, rather than from the C library the application is linked against (the C library coming from Guix).

And this is where the problem is: if your application is linked against Guix's C library (say, `glibc 2.24`) and tries to load NSS plugins from another C library (say, `libnss_mdns.so` for `glibc 2.22`), it will likely crash or have its name lookups fail unexpectedly.

Running `nsd` on the system, among other advantages, eliminates this binary incompatibility problem because those `libnss_*.so` files are loaded in the `nsd` process, not in applications themselves.

Note that `nsd` is no longer provided on some GNU/Linux distros, such as Arch Linux (as of Dec. 2024). `nsncd` can be used as a drop-in-replacement. See the `nsncd` repository (<https://github.com/twosigma/nsncd>) and this blog post (<https://flokli.de/posts/2022-11-18-nsncd/>) for more information.

### 2.4.3 X11 Fonts

The majority of graphical applications use Fontconfig to locate and load fonts and perform X11-client-side rendering. The `fontconfig` package in Guix looks for fonts in `$HOME/.guix-profile` by default. Thus, to allow graphical applications installed with Guix to display fonts, you have to install fonts with Guix as well. Essential font packages include `font-ghostscript`, `font-dejavu`, and `font-gnu-freefont`.

Once you have installed or removed fonts, or when you notice an application that does not find fonts, you may need to install Fontconfig and to force an update of its font cache by running:

```
guix install fontconfig
fc-cache -rv
```

To display text written in Chinese languages, Japanese, or Korean in graphical applications, consider installing `font-adobe-source-han-sans` or `font-wqy-zenhei`. The former has multiple outputs, one per language family (see Section 5.4 [Packages with Multiple Outputs], page 51). For instance, the following command installs fonts for Chinese languages:

```
guix install font-adobe-source-han-sans:cn
```

Older programs such as `xterm` do not use Fontconfig and instead rely on server-side font rendering. Such programs require to specify a full name of a font using XLFD (X Logical Font Description), like this:

```
-*-dejavu sans-medium-r-normal-***-100-***-***-1
```

To be able to use such full names for the TrueType fonts installed in your Guix profile, you need to extend the font path of the X server:

```
xset +fp $(dirname $(readlink -f ~/.guix-profile/share/fonts/truetype/fonts.dir))■
```

After that, you can run `xlsfonts` (from `xlsfonts` package) to make sure your TrueType fonts are listed there.

### 2.4.4 X.509 Certificates

The `nss-certs` package provides X.509 certificates, which allow programs to authenticate Web servers accessed over HTTPS.



When using Guix on a foreign distro, you can install this package and define the relevant environment variables so that packages know where to look for certificates. See Section 11.12 [X.509 Certificates], page 679, for detailed information.

### 2.4.5 Emacs Packages

When you install Emacs packages with Guix, the Emacs files are placed under the `share/emacs/site-lisp/` directory of the profile in which they are installed. The Emacs libraries are made available to Emacs through the `EMACSLOADPATH` environment variable, which is set when installing Emacs itself.

Additionally, autoload definitions are automatically evaluated at the initialization of Emacs, by the Guix-specific `guix-emacs-autoload-packages` procedure. This procedure can be interactively invoked to have newly installed Emacs packages discovered, without having to restart Emacs. If, for some reason, you want to avoid auto-loading the Emacs packages installed with Guix, you can do so by running Emacs with the `--no-site-file` option (see Section “Init File” in *The GNU Emacs Manual*).

**Note:** Most Emacs variants are now capable of doing native compilation. The approach taken by Guix Emacs however differs greatly from the approach taken upstream.

Upstream Emacs compiles packages just-in-time and typically places shared object files in a special folder within your `user-emacs-directory`. These shared objects within said folder are organized in a flat hierarchy, and their file names contain two hashes to verify the original file name and contents of the source code.

Guix Emacs on the other hand prefers to compile packages ahead-of-time. Shared objects retain much of the original file name and no hashes are added to verify the original file name or the contents of the file. Crucially, this allows Guix Emacs and packages built against it to be grafted (see Chapter 19 [Security Updates], page 806), but at the same time, Guix Emacs lacks the hash-based verification of source code baked into upstream Emacs. As this naming schema is trivial to exploit, we disable just-in-time compilation.

Further note, that `emacs-minimal`—the default Emacs for building packages—has been configured without native compilation. To natively compile your emacs packages ahead of time, use a transformation like `--with-input=emacs-minimal=emacs`.

## 2.5 Upgrading Guix

To upgrade Guix, run:

```
guix pull
```

See Section 5.7 [Invoking `guix pull`], page 56, for more information.

If you installed Guix on a foreign distro with `guix-install.sh`, you can upgrade the build daemon by running:

```
sudo -i guix pull
```

followed by (assuming your distro uses the `systemd` service management tool):

```
systemctl restart guix-daemon.service
```

While this works for `guix-install.sh`, it will most likely not work if you installed Guix through the package manager of another distribution (for instance with `sudo apt install guix`, `pacman -S guix`, or similar commands) because on many distributions `guix-daemon.service` will only run the Guix daemon provided by the distribution.

Waiting for your distribution to provide security fixes sometimes work, but fixing security issues in the Guix package takes times for distributions because Guix doesn't backport the security fixes to the releases used by the distributions. On some distribution versions, the Guix package is never fixed.

However since Guix updates also comes with a new daemon, so you can simply configure your distribution to use the `guix-daemon` provided by Guix instead. The following will assume that your distribution uses `systemd`.

You first need to override `guix-daemon.service` with the following command:

```
sudo install -m 644 \
  /lib/systemd/system/guix-daemon.service \
  /etc/systemd/system/guix-daemon.service
```

You will then need to edit `/etc/systemd/system/guix-daemon.service`. To do that open this file with a text editor that has root permissions:

```
sudo $EDITOR /etc/systemd/system/guix-daemon.service
```

Then try to find a line with `ExecStart=`, it should look a bit like that:

```
...
[Service]
ExecStart=/usr/bin/guix-daemon --build-users-group=_guixbuild
...
```

Then change `/usr/bin/guix-daemon` into `/var/guix/profiles/per-user/root/current-guix/bin/guix-daemon` and save the changes.

Once done you need to tell `systemd` that `guix-daemon.service` was overridden:

```
sudo systemctl daemon-reload
```

Once done you can then update the Guix daemon in the same way than if you installed Guix with `guix-install.sh`:

```
sudo -i guix pull
sudo systemctl restart guix-daemon.service
```

On Guix System, upgrading the daemon is achieved by reconfiguring the system (see Section 11.16 [Invoking guix system], page 692).

To migrate an existing installation to the *unprivileged daemon* where `guix-daemon` does not run as root, see [unprivileged-daemon-migration], page 7.

## 3 System Installation

This section explains how to install Guix System on a machine. Guix, as a package manager, can also be installed on top of a running GNU/Linux system, see Chapter 2 [Installation], page 4.

### 3.1 Limitations

We consider Guix System to be ready for a wide range of “desktop” and server use cases. The reliability guarantees it provides—transactional upgrades and rollbacks, reproducibility—make it a solid foundation.

More and more system services are provided (see Section 11.10 [Services], page 275).

Nevertheless, before you proceed with the installation, be aware that some services you rely on may still be missing from version 8b23207.

More than a disclaimer, this is an invitation to report issues (and success stories!), and to join us in improving it. See Chapter 22 [Contributing], page 816, for more info.

### 3.2 Hardware Considerations

GNU Guix focuses on respecting the user’s computing freedom. It builds around the kernel Linux-libre, which means that only hardware for which free software drivers and firmware exist is supported. Nowadays, a wide range of off-the-shelf hardware is supported on GNU/Linux-libre—from keyboards to graphics cards to scanners and Ethernet controllers. Unfortunately, there are still areas where hardware vendors deny users control over their own computing, and such hardware is not supported on Guix System.

One of the main areas where free drivers or firmware are lacking is WiFi devices. WiFi devices known to work include those using Atheros chips (AR9271 and AR7010), which corresponds to the **ath9k** Linux-libre driver, and those using Broadcom/AirForce chips (BCM43xx with Wireless-Core Revision 5), which corresponds to the **b43-open** Linux-libre driver. Free firmware exists for both and is available out-of-the-box on Guix System, as part of **%base-firmware** (see Section 11.3 [operating-system Reference], page 253).

The installer warns you early on if it detects devices that are known *not* to work due to the lack of free firmware or free drivers.

The Free Software Foundation (<https://www.fsf.org/>) runs *Respects Your Freedom* (<https://www.fsf.org/ryf>) (RYF), a certification program for hardware products that respect your freedom and your privacy and ensure that you have control over your device. We encourage you to check the list of RYF-certified devices.

Another useful resource is the H-Node (<https://www.h-node.org/>) web site. It contains a catalog of hardware devices with information about their support in GNU/Linux.

### 3.3 USB Stick and DVD Installation

An ISO-9660 installation image that can be written to a USB stick or burnt to a DVD can be downloaded from `https://ftpmirror.gnu.org/gnu/guix/guix-system-install-8b23207.x86_64-linux.` where you can replace **x86\_64-linux** with one of:

**x86\_64-linux**

for a GNU/Linux system on Intel/AMD-compatible 64-bit CPUs;

**i686-linux**

for a 32-bit GNU/Linux system on Intel-compatible CPUs.

Make sure to download the associated `.sig` file and to verify the authenticity of the image against it, along these lines:

```
$ wget https://ftpmirror.gnu.org/gnu/guix/guix-system-install-8b23207.x86_64-linux.iso
$ gpg --verify guix-system-install-8b23207.x86_64-linux.iso.sig
```

If that command fails because you do not have the required public key, then run this command to import it:

```
$ wget https://sv.gnu.org/people/viewgpg.php?user_id=15145 \
-q0 - | gpg --import -
```

and rerun the `gpg --verify` command.

Take note that a warning like “This key is not certified with a trusted signature!” is normal.

This image contains the tools necessary for an installation. It is meant to be copied *as is* to a large-enough USB stick or DVD.

## Copying to a USB Stick

Insert a USB stick of 1 GiB or more into your machine, and determine its device name. Assuming that the USB stick is known as `/dev/sdX`, copy the image with:

```
dd if=guix-system-install-8b23207.x86_64-linux.iso of=/dev/sdX status=progress
sync
```

Access to `/dev/sdX` usually requires root privileges.

## Burning on a DVD

Insert a blank DVD into your machine, and determine its device name. Assuming that the DVD drive is known as `/dev/srX`, copy the image with:

```
growisofs -dvd-compat -Z /dev/srX=guix-system-install-8b23207.x86_64-linux.iso
```

Access to `/dev/srX` usually requires root privileges.

## Booting

Once this is done, you should be able to reboot the system and boot from the USB stick or DVD. The latter usually requires you to get in the BIOS or UEFI boot menu, where you can choose to boot from the USB stick. In order to boot from Libreboot, switch to the command mode by pressing the `c` key and type `search_grub usb`.

Sadly, on some machines, the installation medium cannot be properly booted and you only see a black screen after booting even after you waited for ten minutes. This may indicate that your machine cannot run Guix System; perhaps you instead want to install Guix on a foreign distro (see Section 2.1 [Binary Installation], page 4). But don’t give up just yet; a possible workaround is pressing the `e` key in the GRUB boot menu and appending `nomodeset` to the Linux bootline. Sometimes the black screen issue can also be resolved by connecting a different display.

See Section 3.8 [Installing Guix in a VM], page 31, if, instead, you would like to install Guix System in a virtual machine (VM).

### 3.4 Preparing for Installation

Once you have booted, you can use the guided graphical installer, which makes it easy to get started (see Section 3.5 [Guided Graphical Installation], page 25). Alternatively, if you are already familiar with GNU/Linux and if you want more control than what the graphical installer provides, you can choose the “manual” installation process (see Section 3.6 [Manual Installation], page 27).


The graphical installer is available on TTY1. You can obtain root shells on TTYs 3 to 6 by hitting `ctrl-alt-f3`, `ctrl-alt-f4`, etc. TTY2 shows this documentation and you can reach it with `ctrl-alt-f2`. Documentation is browsable using the Info reader commands (see *Stand-alone GNU Info*). The installation system runs the GPM mouse daemon, which allows you to select text with the left mouse button and to paste it with the middle button.

**Note:** Installation requires access to the Internet so that any missing dependencies of your system configuration can be downloaded. See the “Networking” section below.

### 3.5 Guided Graphical Installation

The graphical installer is a text-based user interface. It will guide you, with dialog boxes, through the steps needed to install GNU Guix System.


The first dialog boxes allow you to set up the system as you use it during the installation: you can choose the language, keyboard layout, and set up networking, which will be used during the installation. The image below shows the networking dialog.



Later steps allow you to partition your hard disk, as shown in the image below, to choose whether or not to use encrypted file systems, to enter the host name and root password, and to create an additional account, among other things.



Note that, at any time, the installer allows you to exit the current installation step and resume at a previous step, as show in the image below.



Once you're done, the installer produces an operating system configuration and displays it (see Section 11.2 [Using the Configuration System], page 244). At that point you can hit

“OK” and installation will proceed. On success, you can reboot into the new system and enjoy. See Section 3.7 [After System Installation], page 31, for what’s next!

## 3.6 Manual Installation

This section describes how you would “manually” install GNU Guix System on your machine. This option requires familiarity with GNU/Linux, with the shell, and with common administration tools. If you think this is not for you, consider using the guided graphical installer (see Section 3.5 [Guided Graphical Installation], page 25).

The installation system provides root shells on TTYs 3 to 6; press `ctrl-alt-f3`, `ctrl-alt-f4`, and so on to reach them. It includes many common tools needed to install the system, but is also a full-blown Guix System. This means that you can install additional packages, should you need it, using `guix package` (see Section 5.2 [Invoking guix package], page 37).

### 3.6.1 Keyboard Layout, Networking, and Partitioning

Before you can install the system, you may want to adjust the keyboard layout, set up networking, and partition your target hard disk. This section will guide you through this.

#### 3.6.1.1 Keyboard Layout

The installation image uses the US qwerty keyboard layout. If you want to change it, you can use the `loadkeys` command. For example, the following command selects the Dvorak keyboard layout:

```
loadkeys dvorak
```

See the files under `/run/current-system/profile/share/keymaps` for a list of available keyboard layouts. Run `man loadkeys` for more information.

#### 3.6.1.2 Networking

Run the following command to see what your network interfaces are called:

```
ifconfig -a
```

... or, using the GNU/Linux-specific `ip` command:

```
ip address
```

Wired interfaces have a name starting with ‘e’; for example, the interface corresponding to the first on-board Ethernet controller is called ‘`eno1`’. Wireless interfaces have a name starting with ‘w’, like ‘`w1p2s0`’.

Wired connection

To configure a wired network run the following command, substituting *interface* with the name of the wired interface you want to use.

```
ifconfig interface up
```

... or, using the GNU/Linux-specific `ip` command:

```
ip link set interface up
```

Wireless connection

To configure wireless networking, you can create a configuration file for the `wpa_supplicant` configuration tool (its location is not important) using one of the available text editors such as `nano`:

```
nano wpa_supplicant.conf
```

As an example, the following stanza can go to this file and will work for many wireless networks, provided you give the actual SSID and passphrase for the network you are connecting to:

```
network={
    ssid="my-ssid"
    key_mgmt=WPA-PSK
    psk="the network's secret passphrase"
}
```

Start the wireless service and run it in the background with the following command (substitute *interface* with the name of the network interface you want to use):

```
wpa_supplicant -c wpa_supplicant.conf -i interface -B
```

Run `man wpa_supplicant` for more information.

At this point, you need to acquire an IP address. On a network where IP addresses are automatically assigned *via* DHCP, you can run:

```
dhclient -v interface
```

Try to ping a server to see if networking is up and running:

```
ping -c 3 gnu.org
```

Setting up network access is almost always a requirement because the image does not contain all the software and tools that may be needed.

If you need HTTP and HTTPS access to go through a proxy, run the following command:

```
herd set-http-proxy guix-daemon URL
```

where *URL* is the proxy URL, for example `http://example.org:8118`.

If you want to, you can continue the installation remotely by starting an SSH server:

```
herd start ssh-daemon
```

Make sure to either set a password with `passwd`, or configure OpenSSH public key authentication before logging in.

### 3.6.1.3 Disk Partitioning

Unless this has already been done, the next step is to partition, and then format the target partition(s).

The installation image includes several partitioning tools, including `Parted` (see Section “Overview” in *GNU Parted User Manual*), `fdisk`, and `cfdisk`. Run it and set up your disk with the partition layout you want:

```
cfdisk
```

If your disk uses the GUID Partition Table (GPT) format and you plan to install BIOS-based GRUB (which is the default), make sure a BIOS Boot Partition is available (see Section “BIOS installation” in *GNU GRUB manual*).

If you instead wish to use EFI-based GRUB, a FAT32 *EFI System Partition* (ESP) is required. This partition can be mounted at `/boot/efi` for instance and must have the `esp` flag set. E.g., for `parted`:

```
parted /dev/sda set 1 esp on
```



**Note:** Unsure whether to use EFI- or BIOS-based GRUB? If the directory `/sys/firmware/efi` exists in the installation image, then you should probably perform an EFI installation, using `grub-efi-bootloader`. Otherwise you should use the BIOS-based GRUB, known as `grub-bootloader`. See Section 11.15 [Bootloader Configuration], page 685, for more info on bootloaders.

Once you are done partitioning the target hard disk drive, you have to create a file system on the relevant partition(s)<sup>1</sup>. For the ESP, if you have one and assuming it is `/dev/sda1`, run:

```
mkfs.fat -F32 /dev/sda1
```

For the root file system, `ext4` is the most widely used format. Other file systems, such as `Btrfs`, support compression, which is reported to nicely complement file deduplication that the daemon performs independently of the file system (see Section 2.3 [Invoking `guix-daemon`], page 14).

Preferably, assign file systems a label so that you can easily and reliably refer to them in `file-system` declarations (see Section 11.4 [File Systems], page 257). This is typically done using the `-L` option of `mkfs.ext4` and related commands. So, assuming the target root partition lives at `/dev/sda2`, a file system with the label `my-root` can be created with:

```
mkfs.ext4 -L my-root /dev/sda2
```

If you are instead planning to encrypt the root partition, you can use the `Cryptsetup`/`LUKS` utilities to do that (see `man cryptsetup` for more information).

Assuming you want to store the root partition on `/dev/sda2`, the command sequence to format it as a `LUKS` partition would be along these lines:

```
cryptsetup luksFormat /dev/sda2
cryptsetup open /dev/sda2 my-partition
mkfs.ext4 -L my-root /dev/mapper/my-partition
```

Once that is done, mount the target file system under `/mnt` with a command like (again, assuming `my-root` is the label of the root file system):

```
mount LABEL=my-root /mnt
```

Also mount any other file systems you would like to use on the target system relative to this path. If you have opted for `/boot/efi` as an EFI mount point for example, mount it at `/mnt/boot/efi` now so it is found by `guix system init` afterwards.

Finally, if you plan to use one or more swap partitions (see Section 11.6 [Swap Space], page 266), make sure to initialize them with `mkswap`. Assuming you have one swap partition on `/dev/sda3`, you would run:

```
mkswap /dev/sda3
swapon /dev/sda3
```

Alternatively, you may use a swap file. For example, assuming that in the new system you want to use the file `/swapfile` as a swap file, you would run<sup>2</sup>:

```
# This is 10 GiB of swap space. Adjust "count" to change the size.
```

<sup>1</sup> Currently `Guix System` only supports `ext4`, `btrfs`, `JFS`, `F2FS`, and `XFS` file systems. In particular, code that reads file system UUIDs and labels only works for these file system types.

<sup>2</sup> This example will work for many types of file systems (e.g., `ext4`). However, for copy-on-write file systems (e.g., `btrfs`), the required steps may be different. For details, see the manual pages for `mkswap` and `swapon`.

```
dd if=/dev/zero of=/mnt/swapfile bs=1MiB count=10240
# For security, make the file readable and writable only by root.
chmod 600 /mnt/swapfile
mkswap /mnt/swapfile
swapon /mnt/swapfile
```

Note that if you have encrypted the root partition and created a swap file in its file system as described above, then the encryption also protects the swap file, just like any other file in that file system.

### 3.6.2 Proceeding with the Installation

With the target partitions ready and the target root mounted on `/mnt`, we're ready to go. First, run:

```
herd start cow-store /mnt
```

This makes `/gnu/store` copy-on-write, such that packages added to it during the installation phase are written to the target disk on `/mnt` rather than kept in memory. This is necessary because the first phase of the `guix system init` command (see below) entails downloads or builds to `/gnu/store` which, initially, is an in-memory file system.

Next, you have to edit a file and provide the declaration of the operating system to be installed. To that end, the installation system comes with three text editors. We recommend GNU nano (see *GNU nano Manual*), which supports syntax highlighting and parentheses matching; other editors include `mg` (an Emacs clone), and `nvi` (a clone of the original BSD `vi` editor). We strongly recommend storing that file on the target root file system, say, as `/mnt/etc/config.scm`. Failing to do that, you will have lost your configuration file once you have rebooted into the newly-installed system.

See Section 11.2 [Using the Configuration System], page 244, for an overview of the configuration file. The example configurations discussed in that section are available under `/etc/configuration` in the installation image. Thus, to get started with a system configuration providing a graphical display server (a “desktop” system), you can run something along these lines:

```
# mkdir /mnt/etc
# cp /etc/configuration/desktop.scm /mnt/etc/config.scm
# nano /mnt/etc/config.scm
```

You should pay attention to what your configuration file contains, and in particular:

- Make sure the `bootloader-configuration` form refers to the targets you want to install GRUB on. It should mention `grub-bootloader` if you are installing GRUB in the legacy way, or `grub-efi-bootloader` for newer UEFI systems. For legacy systems, the `targets` field contain the names of the devices, like `(list "/dev/sda")`; for UEFI systems it names the paths to mounted EFI partitions, like `(list "/boot/efi")`; do make sure the paths are currently mounted and a `file-system` entry is specified in your configuration.
- Be sure that your file system labels match the value of their respective `device` fields in your `file-system` configuration, assuming your `file-system` configuration uses the `file-system-label` procedure in its `device` field.
- If there are encrypted or RAID partitions, make sure to add a `mapped-devices` field to describe them (see Section 11.5 [Mapped Devices], page 263).

Once you are done preparing the configuration file, the new system must be initialized (remember that the target root file system is mounted under `/mnt`):

```
guix system init /mnt/etc/config.scm /mnt
```

This copies all the necessary files and installs GRUB on `/dev/sdX`, unless you pass the `--no-bootloader` option. For more information, see Section 11.16 [Invoking guix system], page 692. This command may trigger downloads or builds of missing packages, which can take some time.

Once that command has completed—and hopefully succeeded!—you can run `reboot` and boot into the new system. The `root` password in the new system is initially empty; other users’ passwords need to be initialized by running the `passwd` command as `root`, unless your configuration specifies otherwise (see [user-account-password], page 270). See Section 3.7 [After System Installation], page 31, for what’s next!

### 3.7 After System Installation

Success, you’ve now booted into Guix System! Login to the system using the non-`root` user that you created during installation. You can upgrade the system whenever you want by running:

```
guix pull
sudo guix system reconfigure /etc/config.scm
```

This builds a new system *generation* with the latest packages and services.

Now, see Section 11.1 [Getting Started with the System], page 242, and join us on `#guix` on the Libera.Chat IRC network or on `guix-devel@gnu.org` to share your experience!

### 3.8 Installing Guix in a Virtual Machine

If you’d like to install Guix System in a virtual machine (VM) or on a virtual private server (VPS) rather than on your beloved machine, this section is for you.

To boot a QEMU (<https://qemu.org/>) VM for installing Guix System in a disk image, follow these steps:

1. First, retrieve and decompress the Guix system installation image as described previously (see Section 3.3 [USB Stick and DVD Installation], page 23).
2. Create a disk image that will hold the installed system. To make a qcow2-formatted disk image, use the `qemu-img` command:

```
qemu-img create -f qcow2 guix-system.img 50G
```

The resulting file will be much smaller than 50 GB (typically less than 1 MB), but it will grow as the virtualized storage device is filled up.

3. Boot the USB installation image in a VM:

```
qemu-system-x86_64 -m 1024 -smp 1 -enable-kvm \
  -nic user,model=virtio-net-pci -boot menu=on,order=d \
  -drive file=guix-system.img \
  -drive media=cdrom,readonly=on,file=guix-system-install-8b23207.system.iso
```

`-enable-kvm` is optional, but significantly improves performance, see Section 11.18 [Running Guix in a VM], page 707.

4. You're now root in the VM, proceed with the installation process. See Section 3.4 [Preparing for Installation], page 25, and follow the instructions.

Once installation is complete, you can boot the system that's on your `guix-system.img` image. See Section 11.18 [Running Guix in a VM], page 707, for how to do that.

## 3.9 Building the Installation Image

The installation image described above was built using the `guix system` command, specifically:

```
guix system image -t iso9660 gnu/system/install.scm
```

Have a look at `gnu/system/install.scm` in the source tree, and see also Section 11.16 [Invoking `guix system`], page 692, for more information about the installation image.

## 3.10 Building the Installation Image for ARM Boards

Many ARM boards require a specific variant of the U-Boot (<https://www.denx.de/wiki/U-Boot/>) bootloader.

If you build a disk image and the bootloader is not available otherwise (on another boot drive etc), it's advisable to build an image that includes the bootloader, specifically:

```
guix system image --system=armhf-linux -e '(@ (gnu system install) os-with-u-boot) (@
```

`A20-OLinuxino-Lime2` is the name of the board. If you specify an invalid board, a list of possible boards will be printed.

## 4 Getting Started

Presumably, you’ve reached this section because either you have installed Guix on top of another distribution (see Chapter 2 [Installation], page 4), or you’ve installed the standalone Guix System (see Chapter 3 [System Installation], page 23). It’s time for you to get started using Guix and this section aims to help you do that and give you a feel of what it’s like.

Guix is about installing software, so probably the first thing you’ll want to do is to actually look for software. Let’s say you’re looking for a text editor, you can run:

```
guix search text editor
```

This command shows you a number of matching *packages*, each time showing the package’s name, version, a description, and additional info. Once you’ve found out the one you want to use, let’s say Emacs (ah ha!), you can go ahead and install it (run this command as a regular user, *no need for root privileges!*):

```
guix install emacs
```

You’ve installed your first package, congrats! The package is now visible in your default *profile*, `$HOME/.guix-profile`—a profile is a directory containing installed packages. In the process, you’ve probably noticed that Guix downloaded pre-built binaries; or, if you explicitly chose to *not* use pre-built binaries, then probably Guix is still building software (see Section 5.3 [Substitutes], page 46, for more info).

Unless you’re using Guix System, the `guix install` command must have printed this hint:

```
hint: Consider setting the necessary environment variables by running:
```

```
GUIX_PROFILE="$HOME/.guix-profile"
. "$GUIX_PROFILE/etc/profile"
```

Alternately, see ``guix package --search-paths -p "$HOME/.guix-profile"`.`

Indeed, you must now tell your shell where `emacs` and other programs installed with Guix are to be found. Pasting the two lines above will do just that: it will add `$HOME/.guix-profile/bin`—which is where the installed package is—to the `PATH` environment variable. You can paste these two lines in your shell so they take effect right away, but more importantly you should add them to `~/.bash_profile` (or equivalent file if you do not use Bash) so that environment variables are set next time you spawn a shell. You only need to do this once and other search paths environment variables will be taken care of similarly—e.g., if you eventually install `python` and Python libraries, `GUIX_PYTHONPATH` will be defined.

You can go on installing packages at your will. To list installed packages, run:

```
guix package --list-installed
```

To remove a package, you would unsurprisingly run `guix remove`. A distinguishing feature is the ability to *roll back* any operation you made—installation, removal, upgrade—by simply typing:

```
guix package --roll-back
```

This is because each operation is in fact a *transaction* that creates a new *generation*. These generations and the difference between them can be displayed by running:

```
guix package --list-generations
```

Now you know the basics of package management!

**Going further:** See Chapter 5 [Package Management], page 36, for more about package management. You may like *declarative* package management with `guix package --manifest`, managing separate *profiles* with `--profile`, deleting old generations, collecting garbage, and other nifty features that will come in handy as you become more familiar with Guix. If you are a developer, see Chapter 7 [Development], page 78, for additional tools. And if you're curious, see Section 5.1 [Features], page 36, to peek under the hood.

You can also manage the configuration of your entire *home environment*—your user “dot files”, services, and packages—using Guix Home. See Chapter 13 [Home Configuration], page 732, to learn more about it!

Once you've installed a set of packages, you will want to periodically *upgrade* them to the latest and greatest version. To do that, you will first pull the latest revision of Guix and its package collection:

```
guix pull
```

The end result is a new `guix` command, under `~/.config/guix/current/bin`. Unless you're on Guix System, the first time you run `guix pull`, be sure to follow the hint that the command prints and, similar to what we saw above, paste these two lines in your terminal and `.bash_profile`:

```
GUIX_PROFILE="$HOME/.config/guix/current"
. "$GUIX_PROFILE/etc/profile"
```

You must also instruct your shell to point to this new `guix`:

```
hash guix
```

At this point, you're running a brand new Guix. You can thus go ahead and actually upgrade all the packages you previously installed:

```
guix upgrade
```

As you run this command, you will see that binaries are downloaded (or perhaps some packages are built), and eventually you end up with the upgraded packages. Should one of these upgraded packages not be to your liking, remember you can always roll back!

You can display the exact revision of Guix you're currently using by running:

```
guix describe
```

The information it displays is *all it takes to reproduce the exact same Guix*, be it at a different point in time or on a different machine.

**Going further:** See Section 5.7 [Invoking `guix pull`], page 56, for more information. See Chapter 6 [Channels], page 68, on how to specify additional *channels* to pull packages from, how to replicate Guix, and more. You may also find `time-machine` handy (see Section 5.8 [Invoking `guix time-machine`], page 60).

If you installed Guix System, one of the first things you'll want to do is to upgrade your system. Once you've run `guix pull` to get the latest Guix, you can upgrade the system like this:

```
sudo guix system reconfigure /etc/config.scm
```

Upon completion, the system runs the latest versions of its software packages. Just like for packages, you can always *roll back* to a previous generation *of the whole system*. See

Section 11.1 [Getting Started with the System], page 242, to learn how to manage your system.

Now you know enough to get started!

**Resources:** The rest of this manual provides a reference for all things Guix. Here are some additional resources you may find useful:

- See *The GNU Guix Cookbook*, for a list of “how-to” style of recipes for a variety of applications.
- The GNU Guix Reference Card (<https://guix.gnu.org/guix-refcard.pdf>) lists in two pages most of the commands and options you’ll ever need.
- The web site contains instructional videos (<https://guix.gnu.org/en/videos/>) covering topics such as everyday use of Guix, how to get help, and how to become a contributor.
- See Chapter 14 [Documentation], page 790, to learn how to access documentation on your computer.

We hope you will enjoy Guix as much as the community enjoys building it!

## 5 Package Management

The purpose of GNU Guix is to allow users to easily install, upgrade, and remove software packages, without having to know about their build procedures or dependencies. Guix also goes beyond this obvious set of features.

This chapter describes the main features of Guix, as well as the package management tools it provides. Along with the command-line interface described below (see Section 5.2 [Invoking guix package], page 37), you may also use the Emacs-Guix interface (see *The Emacs-Guix Reference Manual*), after installing `emacs-guix` package (run `M-x guix-help` command to start with it):

```
guix install emacs-guix
```

### 5.1 Features

Here we assume you’ve already made your first steps with Guix (see Chapter 4 [Getting Started], page 33) and would like to get an overview about what’s going on under the hood.

When using Guix, each package ends up in the *package store*, in its own directory—something that resembles `/gnu/store/xxx-package-1.2`, where `xxx` is a base32 string.

Instead of referring to these directories, users have their own *profile*, which points to the packages that they actually want to use. These profiles are stored within each user’s home directory, at `$HOME/.guix-profile`.

For example, `alice` installs GCC 4.7.2. As a result, `/home/alice/.guix-profile/bin/gcc` points to `/gnu/store/...-gcc-4.7.2/bin/gcc`. Now, on the same machine, `bob` had already installed GCC 4.8.0. The profile of `bob` simply continues to point to `/gnu/store/...-gcc-4.8.0/bin/gcc`—i.e., both versions of GCC coexist on the same system without any interference.

The `guix package` command is the central tool to manage packages (see Section 5.2 [Invoking guix package], page 37). It operates on the per-user profiles, and can be used *with normal user privileges*.

The command provides the obvious install, remove, and upgrade operations. Each invocation is actually a *transaction*: either the specified operation succeeds, or nothing happens. Thus, if the `guix package` process is terminated during the transaction, or if a power outage occurs during the transaction, then the user’s profile remains in its previous state, and remains usable.

In addition, any package transaction may be *rolled back*. So, if, for example, an upgrade installs a new version of a package that turns out to have a serious bug, users may roll back to the previous instance of their profile, which was known to work well. Similarly, the global system configuration on Guix is subject to transactional upgrades and roll-back (see Section 11.1 [Getting Started with the System], page 242).

All packages in the package store may be *garbage-collected*. Guix can determine which packages are still referenced by user profiles, and remove those that are provably no longer referenced (see Section 5.6 [Invoking guix gc], page 53). Users may also explicitly remove old generations of their profile so that the packages they refer to can be collected.

Guix takes a *purely functional* approach to package management, as described in the introduction (see Chapter 1 [Introduction], page 1). Each `/gnu/store` package directory



name contains a hash of all the inputs that were used to build that package—compiler, libraries, build scripts, etc. This direct correspondence allows users to make sure a given package installation matches the current state of their distribution. It also helps maximize *build reproducibility*: thanks to the isolated build environments that are used, a given build is likely to yield bit-identical files when performed on different machines (see Section 2.3 [Invoking guix-daemon], page 14).

This foundation allows Guix to support *transparent binary/source deployment*. When a pre-built binary for a `/gnu/store` item is available from an external source—a *substitute*, Guix just downloads it and unpacks it; otherwise, it builds the package from source, locally (see Section 5.3 [Substitutes], page 46). Because build results are usually bit-for-bit reproducible, users do not have to trust servers that provide substitutes: they can force a local build and *challenge* providers (see Section 9.12 [Invoking guix challenge], page 231).

Control over the build environment is a feature that is also useful for developers. The `guix shell` command allows developers of a package to quickly set up the right development environment for their package, without having to manually install the dependencies of the package into their profile (see Section 7.1 [Invoking guix shell], page 78).

All of Guix and its package definitions is version-controlled, and `guix pull` allows you to “travel in time” on the history of Guix itself (see Section 5.7 [Invoking guix pull], page 56). This makes it possible to replicate a Guix instance on a different machine or at a later point in time, which in turn allows you to *replicate complete software environments*, while retaining precise *provenance tracking* of the software.

## 5.2 Invoking guix package

The `guix package` command is the tool that allows users to install, upgrade, and remove packages, as well as rolling back to previous configurations. These operations work on a user *profile*—a directory of installed packages. Each user has a default profile in `$HOME/.guix-profile`. The command operates only on the user’s own profile, and works with normal user privileges (see Section 5.1 [Features], page 36). Its syntax is:

**guix package options**

Primarily, *options* specifies the operations to be performed during the transaction. Upon completion, a new profile is created, but previous *generations* of the profile remain available, should the user want to roll back.

For example, to remove `lua` and install `guile` and `guile-cairo` in a single transaction:

```
guix package -r lua -i guile guile-cairo
```

For your convenience, we also provide the following aliases:

- `guix search` is an alias for `guix package -s`,
- `guix install` is an alias for `guix package -i`,
- `guix remove` is an alias for `guix package -r`,
- `guix upgrade` is an alias for `guix package -u`,
- and `guix show` is an alias for `guix package --show=`.

These aliases are less expressive than `guix package` and provide fewer options, so in some cases you’ll probably want to use `guix package` directly.

`guix package` also supports a *declarative approach* whereby the user specifies the exact set of packages to be available and passes it *via* the `--manifest` option (see [profile-manifest], page 41).

For each user, a symlink to the user's default profile is automatically created in `$HOME/.guix-profile`. This symlink always points to the current generation of the user's default profile. Thus, users can add `$HOME/.guix-profile/bin` to their `PATH` environment variable, and so on. If you are not using Guix System, consider adding the following lines to your `~/.bash_profile` (see Section “Bash Startup Files” in *The GNU Bash Reference Manual*) so that newly-spawned shells get all the right environment variable definitions:

```
GUIX_PROFILE="$HOME/.guix-profile" ; \
source "$GUIX_PROFILE/etc/profile"
```

In a multi-user setup, user profiles are stored in a place registered as a *garbage-collector root*, which `$HOME/.guix-profile` points to (see Section 5.6 [Invoking `guix gc`], page 53). That directory is normally `localstatedir/guix/profiles/per-user/user`, where `localstatedir` is the value passed to `configure` as `--localstatedir`, and `user` is the user name. The `per-user` directory is created when `guix-daemon` is started, and the `user` sub-directory is created by `guix package`.

The *options* can be among the following:

`--install=package ...`

`-i package ...`

Install the specified *packages*.

Each *package* may specify a simple package name, such as `guile`, optionally followed by an at-sign and version number, such as `guile@3.0.7` or simply `guile@3.0`. In the latter case, the newest version prefixed by 3.0 is selected.

If no version number is specified, the newest available version will be selected. In addition, such a *package* specification may contain a colon, followed by the name of one of the outputs of the package, as in `gcc:doc` or `binutils@2.22:lib` (see Section 5.4 [Packages with Multiple Outputs], page 51).

Packages with a corresponding name (and optionally version) are searched for among the GNU distribution modules (see Section 8.1 [Package Modules], page 100).

Alternatively, a *package* can directly specify a store file name such as `/gnu/store/...-guile-3.0.7`, as produced by, e.g., `guix build`.

Sometimes packages have *propagated inputs*: these are dependencies that automatically get installed along with the required package (see [package-propagated-inputs], page 106, for information about propagated inputs in package definitions).

An example is the GNU MPC library: its C header files refer to those of the GNU MPFR library, which in turn refer to those of the GMP library. Thus, when installing MPC, the MPFR and GMP libraries also get installed in the profile; removing MPC also removes MPFR and GMP—unless they had also been explicitly installed by the user.

Besides, packages sometimes rely on the definition of environment variables for their search paths (see explanation of `--search-paths` below). Any missing or possibly incorrect environment variable definitions are reported here.

`--install-from-expression=exp`

`-e exp` Install the package *exp* evaluates to.

*exp* must be a Scheme expression that evaluates to a <package> object. This option is notably useful to disambiguate between same-named variants of a package, with expressions such as `(@ (gnu packages commencement) guile-final)`.

Note that this option installs the first output of the specified package, which may be insufficient when needing a specific output of a multiple-output package.

`--install-from-file=file`

`-f file` Install the package that the code within *file* evaluates to.

As an example, *file* might contain a definition like this (see Section 8.2 [Defining Packages], page 101):

```
(use-modules (guix)
              (guix build-system gnu)
              (guix licenses))

(package
  (name "hello")
  (version "2.10")
  (source (origin
            (method url-fetch)
            (uri (string-append "mirror://gnu/hello/hello-" version
                                ".tar.gz"))
            (sha256
              (base32
                "0ssi1wpaf7plawqqjwigppsg5fyh99vdlb9kz17c9lng89ndq1i"))))
  (build-system gnu-build-system)
  (synopsis "Hello, GNU world: An example GNU package")
  (description "Guess what GNU Hello prints!")
  (home-page "http://www.gnu.org/software/hello/")
  (license gpl3+))
```

Developers may find it useful to include such a `guix.scm` file in the root of their project source tree that can be used to test development snapshots and create reproducible development environments (see Section 7.1 [Invoking guix shell], page 78).

The *file* may also contain a JSON representation of one or more package definitions. Running `guix package -f` on `hello.json` with the following contents would result in installing the package `greeter` after building `myhello`:

```
[
  {
    "name": "myhello",
    "version": "2.10",
    "source": "mirror://gnu/hello/hello-2.10.tar.gz",
    "build-system": "gnu",
    "arguments": {
```

```

    "tests?": false
  },
  "home-page": "https://www.gnu.org/software/hello/",
  "synopsis": "Hello, GNU world: An example GNU package",
  "description": "GNU Hello prints a greeting.",
  "license": "GPL-3.0+",
  "native-inputs": ["gettext"]
},
{
  "name": "greeter",
  "version": "1.0",
  "source": "mirror://gnu/hello/hello-2.10.tar.gz",
  "build-system": "gnu",
  "arguments": {
    "test-target": "foo",
    "parallel-build?": false
  },
  "home-page": "https://example.com/",
  "synopsis": "Greeter using GNU Hello",
  "description": "This is a wrapper around GNU Hello.",
  "license": "GPL-3.0+",
  "inputs": ["myhello", "hello"]
}
]

```

`--remove=package ...`

`-r package ...`

Remove the specified *packages*.

As for `--install`, each *package* may specify a version number and/or output name in addition to the package name. For instance, `'-r glibc:debug'` would remove the debug output of `glibc`.

`--upgrade[=regex ...]`

`-u [regex ...]`

Upgrade all the installed packages. If one or more *regexps* are specified, upgrade only installed packages whose name matches a *regex*. Also see the `--do-not-upgrade` option below.

Note that this upgrades package to the latest version of packages found in the distribution currently installed. To update your distribution, you should regularly run `guix pull` (see Section 5.7 [Invoking `guix pull`], page 56).

When upgrading, package transformations that were originally applied when creating the profile are automatically re-applied (see Section 9.1.2 [Package Transformation Options], page 182). For example, assume you first installed Emacs from the tip of its development branch with:

```
guix install emacs-next --with-branch=emacs-next=master
```

Next time you run `guix upgrade`, Guix will again pull the tip of the Emacs development branch and build `emacs-next` from that checkout.

Note that transformation options such as `--with-branch` and `--with-source` depend on external state; it is up to you to ensure that they work as expected. You can also discard a transformations that apply to a package by running:

```
guix install package
```

`--do-not-upgrade[=regexp ...]`

When used together with the `--upgrade` option, do *not* upgrade any packages whose name matches a *regexp*. For example, to upgrade all packages in the current profile except those containing the substring “emacs”:

```
$ guix package --upgrade . --do-not-upgrade emacs
```

`--manifest=file`

`-m file` Create a new generation of the profile from the manifest object returned by the Scheme code in *file*. This option can be repeated several times, in which case the manifests are concatenated.

This allows you to *declare* the profile’s contents rather than constructing it through a sequence of `--install` and similar commands. The advantage is that *file* can be put under version control, copied to different machines to reproduce the same profile, and so on.

*file* must return a *manifest* object, which is roughly a list of packages:

```
(use-package-modules guile emacs)

(packages->manifest
 (list emacs
       guile-2.0
       ;; Use a specific package output.
       (list guile-2.0 "debug")))
```

See Section 8.4 [Writing Manifests], page 117, for information on how to write a manifest. See [export-manifest], page 45, to learn how to obtain a manifest file from an existing profile.

`--roll-back`

Roll back to the previous *generation* of the profile—i.e., undo the last transaction.

When combined with options such as `--install`, roll back occurs before any other actions.

When rolling back from the first generation that actually contains installed packages, the profile is made to point to the *zeroth generation*, which contains no files apart from its own metadata.

After having rolled back, installing, removing, or upgrading packages overwrites previous future generations. Thus, the history of the generations in a profile is always linear.

`--switch-generation=pattern`

`-S pattern`

Switch to a particular generation defined by *pattern*.

*pattern* may be either a generation number or a number prefixed with “+” or “-”. The latter means: move forward/backward by a specified number of

generations. For example, if you want to return to the latest generation after `--roll-back`, use `--switch-generation=+1`.

The difference between `--roll-back` and `--switch-generation=-1` is that `--switch-generation` will not make a zeroth generation, so if a specified generation does not exist, the current generation will not be changed.

#### `--search-paths[=kind]`

Report environment variable definitions, in Bash syntax, that may be needed in order to use the set of installed packages. These environment variables are used to specify *search paths* for files used by some of the installed packages.

For example, GCC needs the `CPATH` and `LIBRARY_PATH` environment variables to be defined so it can look for headers and libraries in the user's profile (see Section "Environment Variables" in *Using the GNU Compiler Collection (GCC)*). If GCC and, say, the C library are installed in the profile, then `--search-paths` will suggest setting these variables to *profile/include* and *profile/lib*, respectively (see Section 8.8 [Search Paths], page 153, for info on search path specifications associated with packages.)

The typical use case is to define these environment variables in the shell:

```
$ eval $(guix package --search-paths)
```

*kind* may be one of `exact`, `prefix`, or `suffix`, meaning that the returned environment variable definitions will either be exact settings, or prefixes or suffixes of the current value of these variables. When omitted, *kind* defaults to `exact`.

This option can also be used to compute the *combined* search paths of several profiles. Consider this example:

```
$ guix package -p foo -i guile
$ guix package -p bar -i guile-json
$ guix package -p foo -p bar --search-paths
```

The last command above reports about the `GUILE_LOAD_PATH` variable, even though, taken individually, neither `foo` nor `bar` would lead to that recommendation.

#### `--profile=profile`

##### `-p profile`

Use *profile* instead of the user's default profile.

*profile* must be the name of a file that will be created upon completion. Concretely, *profile* will be a mere symbolic link ("symlink") pointing to the actual profile where packages are installed:

```
$ guix install hello -p ~/code/my-profile
...
$ ~/code/my-profile/bin/hello
Hello, world!
```

All it takes to get rid of the profile is to remove this symlink and its siblings that point to specific generations:

```
$ rm ~/code/my-profile ~/code/my-profile-*-link
```

**--list-profiles**

List all the user's profiles:

```
$ guix package --list-profiles
/home/charlie/.guix-profile
/home/charlie/code/my-profile
/home/charlie/code/devel-profile
/home/charlie/tmp/test
```

When running as root, list all the profiles of all the users.

**--allow-collisions**

Allow colliding packages in the new profile. Use at your own risk!

By default, `guix package` reports as an error *collisions* in the profile. Collisions happen when two or more different versions or variants of a given package end up in the profile.

**--bootstrap**

Use the bootstrap Guile to build the profile. This option is only useful to distribution developers.

In addition to these actions, `guix package` supports the following options to query the current state of a profile, or the availability of packages:

**--search=regex**

**-s regex** List the available packages whose name, synopsis, or description matches *regex* (in a case-insensitive fashion), sorted by relevance. Print all the metadata of matching packages in `recutils` format (see *GNU recutils manual*).

This allows specific fields to be extracted using the `recsel` command, for instance:

```
$ guix package -s malloc | recsel -p name,version,relevance
name: jemalloc
version: 4.5.0
relevance: 6

name: glibc
version: 2.25
relevance: 1

name: libgc
version: 7.6.0
relevance: 1
```

Similarly, to show the name of all the packages available under the terms of the GNU LGPL version 3:

```
$ guix package -s "" | recsel -p name -e 'license ~ "LGPL 3"'
name: elfutils

name: gmp
...
```

It is also possible to refine search results using several `-s` flags to `guix package`, or several arguments to `guix search`. For example, the following command returns a list of board games (this time using the `guix search` alias):

```
$ guix search '\<board\>' game | recsel -p name
name: gnubg
...
```

If we were to omit `-s game`, we would also get software packages that deal with printed circuit boards; removing the angle brackets around `board` would further add packages that have to do with keyboards.

And now for a more elaborate example. The following command searches for cryptographic libraries, filters out Haskell, Perl, Python, and Ruby libraries, and prints the name and synopsis of the matching packages:

```
$ guix search crypto library | \
  recsel -e '! (name ~ "^(ghc|perl|python|ruby)")' -p name,synopsis
```

See Section “Selection Expressions” in *GNU recutils manual*, for more information on *selection expressions* for `recsel -e`.

#### `--show=package`

Show details about *package*, taken from the list of available packages, in *recutils* format (see *GNU recutils manual*).

```
$ guix package --show=guile | recsel -p name,version
name: guile
version: 3.0.5

name: guile
version: 3.0.2

name: guile
version: 2.2.7
...
```

You may also specify the full name of a package to only get details about a specific version of it (this time using the `guix show` alias):

```
$ guix show guile@3.0.5 | recsel -p name,version
name: guile
version: 3.0.5
```

#### `--list-installed[=regex]`

##### `-I [regex]`

List the currently installed packages in the specified profile, with the most recently installed packages shown last. When *regex* is specified, list only installed packages whose name matches *regex*.

For each installed package, print the following items, separated by tabs: the package name, its version string, the part of the package that is installed (for instance, `out` for the default output, `include` for its headers, etc.), and the path of this package in the store.



`--list-available[=regex]`

`-A [regex]`

List packages currently available in the distribution for this system (see Section 1.2 [GNU Distribution], page 2). When *regex* is specified, list only available packages whose name matches *regex*.

For each package, print the following items separated by tabs: its name, its version string, the parts of the package (see Section 5.4 [Packages with Multiple Outputs], page 51), and the source location of its definition.

`--list-generations[=pattern]`

`-l [pattern]`

Return a list of generations along with their creation dates; for each generation, show the installed packages, with the most recently installed packages shown last. Note that the zeroth generation is never shown.

For each installed package, print the following items, separated by tabs: the name of a package, its version string, the part of the package that is installed (see Section 5.4 [Packages with Multiple Outputs], page 51), and the location of this package in the store.

When *pattern* is used, the command returns only matching generations. Valid patterns include:

- *Integers and comma-separated integers.* Both patterns denote generation numbers. For instance, `--list-generations=1` returns the first one. And `--list-generations=1,8,2` outputs three generations in the specified order. Neither spaces nor trailing commas are allowed.
- *Ranges.* `--list-generations=2..9` prints the specified generations and everything in between. Note that the start of a range must be smaller than its end.  
It is also possible to omit the endpoint. For example, `--list-generations=2..`, returns all generations starting from the second one.
- *Durations.* You can also get the last *N* days, weeks, or months by passing an integer along with the first letter of the duration. For example, `--list-generations=20d` lists generations that are up to 20 days old.

`--delete-generations[=pattern]`

`-d [pattern]`

When *pattern* is omitted, delete all generations except the current one.

This command accepts the same patterns as `--list-generations`. When *pattern* is specified, delete the matching generations. When *pattern* specifies a duration, generations *older* than the specified duration match. For instance, `--delete-generations=1m` deletes generations that are more than one month old.

If the current generation matches, it is *not* deleted. Also, the zeroth generation is never deleted.

Note that deleting generations prevents rolling back to them. Consequently, this command must be used with care.

**--export-manifest**

Write to standard output a manifest suitable for **--manifest** corresponding to the chosen profile(s).

This option is meant to help you migrate from the “imperative” operating mode—running **guix install**, **guix upgrade**, etc.—to the declarative mode that **--manifest** offers.

Be aware that the resulting manifest *approximates* what your profile actually contains; for instance, depending on how your profile was created, it can refer to packages or package versions that are not exactly what you specified.

Keep in mind that a manifest is purely symbolic: it only contains package names and possibly versions, and their meaning varies over time. If you wish to “pin” channels to the revisions that were used to build the profile(s), see **--export-channels** below.

**--export-channels**

Write to standard output the list of channels used by the chosen profile(s), in a format suitable for **guix pull --channels** or **guix time-machine --channels** (see Chapter 6 [Channels], page 68).

Together with **--export-manifest**, this option provides information allowing you to replicate the current profile (see Section 6.3 [Replicating Guix], page 69).

However, note that the output of this command *approximates* what was actually used to build this profile. In particular, a single profile might have been built from several different revisions of the same channel. In that case, **--export-manifest** chooses the last one and writes the list of other revisions in a comment. If you really need to pick packages from different channel revisions, you can use inferiors in your manifest to do so (see Section 5.9 [Inferiors], page 62).

Together with **--export-manifest**, this is a good starting point if you are willing to migrate from the “imperative” model to the fully declarative model consisting of a manifest file along with a channels file pinning the exact channel revision(s) you want.

Finally, since **guix package** may actually start build processes, it supports all the common build options (see Section 9.1.1 [Common Build Options], page 180). It also supports package transformation options, such as **--with-source**, and preserves them across upgrades (see Section 9.1.2 [Package Transformation Options], page 182).

## 5.3 Substitutes

Guix supports transparent source/binary deployment, which means that it can either build things locally, or download pre-built items from a server, or both. We call these pre-built items *substitutes*—they are substitutes for local build results. In many cases, downloading a substitute is much faster than building things locally.

Substitutes can be anything resulting from a derivation build (see Section 8.10 [Derivations], page 158). Of course, in the common case, they are pre-built package binaries, but source tarballs, for instance, which also result from derivation builds, can be available as substitutes.

### 5.3.1 Official Substitute Servers

`bordeaux.guix.gnu.org` and `ci.guix.gnu.org` are both front-ends to official build farms that build packages from Guix continuously for some architectures, and make them available as substitutes. These are the default source of substitutes; which can be overridden by passing the `--substitute-urls` option either to `guix-daemon` (see `[guix-daemon --substitute-urls]`, page 15) or to client tools such as `guix package` (see `[client --substitute-urls option]`, page 181).

Substitute URLs can be either HTTP or HTTPS. HTTPS is recommended because communications are encrypted; conversely, using HTTP makes all communications visible to an eavesdropper, who could use the information gathered to determine, for instance, whether your system has unpatched security vulnerabilities.

Substitutes from the official build farms are enabled by default when using Guix System (see Section 1.2 [GNU Distribution], page 2). However, they are disabled by default when using Guix on a foreign distribution, unless you have explicitly enabled them via one of the recommended installation steps (see Chapter 2 [Installation], page 4). The following paragraphs describe how to enable or disable substitutes for the official build farm; the same procedure can also be used to enable substitutes for any other substitute server.

### 5.3.2 Substitute Server Authorization

To allow Guix to download substitutes from `bordeaux.guix.gnu.org`, `ci.guix.gnu.org` or a mirror, you must add the relevant public key to the access control list (ACL) of archive imports, using the `guix archive` command (see Section 5.11 [Invoking guix archive], page 65). Doing so implies that you trust the substitute server to not be compromised and to serve genuine substitutes.

**Note:** If you are using Guix System, you can skip this section: Guix System authorizes substitutes from `bordeaux.guix.gnu.org` and `ci.guix.gnu.org` by default.

The public keys for each of the project maintained substitute servers are installed along with Guix, in `prefix/share/guix/`, where *prefix* is the installation prefix of Guix. If you installed Guix from source, make sure you checked the GPG signature of `guix-8b23207.tar.gz`, which contains this public key file. Then, you can run something like this:

```
# guix archive --authorize < prefix/share/guix/bordeaux.guix.gnu.org.pub
# guix archive --authorize < prefix/share/guix/ci.guix.gnu.org.pub
```

Once this is in place, the output of a command like `guix build` should change from something like:

```
$ guix build emacs --dry-run
The following derivations would be built:
  /gnu/store/yr7bnx8xwcayd6j95r2clmmdl1qh688w-emacs-24.3.drv
  /gnu/store/x8qsh1hlhgjx6cwsjyvybnfv2i37z23w-dbus-1.6.4.tar.gz.drv
  /gnu/store/1ixwp12f1950d15h2cj11c73733jay0z-alsa-lib-1.0.27.1.tar.bz2.drv
  /gnu/store/nlma1pw0p603fpfiqy7kn4zm105r5dmw-util-linux-2.21.drv
  ...
```

to something like:

```
$ guix build emacs --dry-run
```

```

112.3 MB would be downloaded:
/gnu/store/pk3n22lbq6ydamyymqkz7i69wiwjiwi-emacs-24.3
/gnu/store/2ygn4ncnhrpr61rssa6z0d9x22si0va3-libjpeg-8d
/gnu/store/71yz6lgx4dazma9dwn2mcjxaah9w77jq-cairo-1.12.16
/gnu/store/7zdhgp0n1518lvfn8mb96sxqfmvqrl7v-libxrender-0.9.7
...

```

The text changed from “The following derivations would be built” to “112.3 MB would be downloaded”. This indicates that substitutes from the configured substitute servers are usable and will be downloaded, when possible, for future builds.

The substitute mechanism can be disabled globally by running `guix-daemon` with `--no-substitutes` (see Section 2.3 [Invoking `guix-daemon`], page 14). It can also be disabled temporarily by passing the `--no-substitutes` option to `guix package`, `guix build`, and other command-line tools.

### 5.3.3 Getting Substitutes from Other Servers

Guix can look up and fetch substitutes from several servers. This is useful when you are using packages from additional channels for which the official server does not have substitutes but another server provides them. Another situation where this is useful is when you would prefer to download from your organization’s substitute server, resorting to the official server only as a fallback or dismissing it altogether.

You can give Guix a list of substitute server URLs and it will check them in the specified order. You also need to explicitly authorize the public keys of substitute servers to instruct Guix to accept the substitutes they sign.

On Guix System, this is achieved by modifying the configuration of the `guix` service. Since the `guix` service is part of the default lists of services, `%base-services` and `%desktop-services`, you can use `modify-services` to change its configuration and add the URLs and substitute keys that you want (see Section 11.19.3 [Service Reference], page 712).

As an example, suppose you want to fetch substitutes from `guix.example.org` and to authorize the signing key of that server, in addition to the default `bordeaux.guix.gnu.org` and `ci.guix.gnu.org`. The resulting operating system configuration will look something like:

```

(operating-system
  ;; ...
  (services
    ;; Assume we're starting from '%desktop-services'. Replace it
    ;; with the list of services you're actually using.
    (modify-services %desktop-services
      (guix-service-type config =>
        (guix-configuration
          (inherit config)
          (substitute-urls
            (append (list "https://guix.example.org")
                     %default-substitute-urls))
          (authorized-keys
            (append (list (local-file "./key.pub"))

```

```
%default-authorized-guix-keys))))))
```

This assumes that the file `key.pub` contains the signing key of `guix.example.org`. With this change in place in your operating system configuration file (say `/etc/config.scm`), you can reconfigure and restart the `guix-daemon` service or reboot so the changes take effect:

```
$ sudo guix system reconfigure /etc/config.scm
$ sudo herd restart guix-daemon
```

If you're running Guix on a "foreign distro", you would instead take the following steps to get substitutes from additional servers:

1. Edit the service configuration file for `guix-daemon`; when using `systemd`, this is normally `/etc/systemd/system/guix-daemon.service`. Add the `--substitute-urls` option on the `guix-daemon` command line and list the URLs of interest (see [daemon-substitute-urls], page 15):

```
... --substitute-urls='https://guix.example.org https://bordeaux.guix.gnu.org
https://ci.guix.gnu.org'
```

2. Restart the daemon. For `systemd`, it goes like this:

```
systemctl daemon-reload
systemctl restart guix-daemon.service
```

3. Authorize the key of the new server (see Section 5.11 [Invoking guix archive], page 65):

```
guix archive --authorize < key.pub
```

Again this assumes `key.pub` contains the public key that `guix.example.org` uses to sign substitutes.

Now you're all set! Substitutes will be preferably taken from `https://guix.example.org`, using `bordeaux.guix.gnu.org` then `ci.guix.gnu.org` as fallback options. Of course you can list as many substitute servers as you like, with the caveat that substitute lookup can be slowed down if too many servers need to be contacted.

**Troubleshooting:** To diagnose problems, you can run `guix weather`. For example, running:

```
guix weather coreutils
```

not only tells you which of the currently-configured servers has substitutes for the `coreutils` package, it also reports whether one of these servers is unauthorized. See Section 9.15 [Invoking guix weather], page 235, for more information.

Note that there are also situations where one may want to add the URL of a substitute server *without* authorizing its key. See Section 5.3.4 [Substitute Authentication], page 49, to understand this fine point.

### 5.3.4 Substitute Authentication

Guix detects and raises an error when attempting to use a substitute that has been tampered with. Likewise, it ignores substitutes that are not signed, or that are not signed by one of the keys listed in the ACL.

There is one exception though: if an unauthorized server provides substitutes that are *bit-for-bit identical* to those provided by an authorized server, then the unauthorized server becomes eligible for downloads. For example, assume we have chosen two substitute servers with this option:

```
--substitute-urls="https://a.example.org https://b.example.org"
```

If the ACL contains only the key for ‘`b.example.org`’, and if ‘`a.example.org`’ happens to serve the *exact same* substitutes, then Guix will download substitutes from ‘`a.example.org`’ because it comes first in the list and can be considered a mirror of ‘`b.example.org`’. In practice, independent build machines usually produce the same binaries, thanks to bit-reproducible builds (see below).

When using HTTPS, the server’s X.509 certificate is *not* validated (in other words, the server is not authenticated), contrary to what HTTPS clients such as Web browsers usually do. This is because Guix authenticates substitute information itself, as explained above, which is what we care about (whereas X.509 certificates are about authenticating bindings between domain names and public keys).

### 5.3.5 Proxy Settings

Substitutes are downloaded over HTTP or HTTPS. The `http_proxy` and `https_proxy` environment variables can be set in the environment of `guix-daemon` and are honored for downloads of substitutes. Note that the value of those environment variables in the environment where `guix build`, `guix package`, and other client commands are run has *absolutely no effect*.

### 5.3.6 Substitution Failure

Even when a substitute for a derivation is available, sometimes the substitution attempt will fail. This can happen for a variety of reasons: the substitute server might be offline, the substitute may recently have been deleted, the connection might have been interrupted, etc.

When substitutes are enabled and a substitute for a derivation is available, but the substitution attempt fails, Guix will attempt to build the derivation locally depending on whether or not `--fallback` was given (see [common build option `--fallback`], page 181). Specifically, if `--fallback` was omitted, then no local build will be performed, and the derivation is considered to have failed. However, if `--fallback` was given, then Guix will attempt to build the derivation locally, and the success or failure of the derivation depends on the success or failure of the local build. Note that when substitutes are disabled or no substitute is available for the derivation in question, a local build will *always* be performed, regardless of whether or not `--fallback` was given.

To get an idea of how many substitutes are available right now, you can try running the `guix weather` command (see Section 9.15 [Invoking `guix weather`], page 235). This command provides statistics on the substitutes provided by a server.

### 5.3.7 On Trusting Binaries

Today, each individual’s control over their own computing is at the mercy of institutions, corporations, and groups with enough power and determination to subvert the computing infrastructure and exploit its weaknesses. While using substitutes can be convenient, we encourage users to also build on their own, or even run their own build farm, such that the project run substitute servers are less of an interesting target. One way to help is by publishing the software you build using `guix publish` so that others have one more choice of server to download substitutes from (see Section 9.11 [Invoking `guix publish`], page 227).

Guix has the foundations to maximize build reproducibility (see Section 5.1 [Features], page 36). In most cases, independent builds of a given package or derivation should yield

bit-identical results. Thus, through a diverse set of independent package builds, we can strengthen the integrity of our systems. The `guix challenge` command aims to help users assess substitute servers, and to assist developers in finding out about non-deterministic package builds (see Section 9.12 [Invoking `guix challenge`], page 231). Similarly, the `--check` option of `guix build` allows users to check whether previously-installed substitutes are genuine by rebuilding them locally (see [build-check], page 192). To force a full rebuild of a package (ignoring security updates via grafts (see Chapter 19 [Security Updates], page 806), if any grafts exist—which is not always the case), use `--check` together with `--no-grafts` (see Section 9.1.1 [Common Build Options], page 180). Because grafts are built as their own derivation, if the package you want to rebuild is subject to being grafted, merely using `--check` will only rebuild the grafting derivation, and not actually recompile the package.

In the future, we want Guix to have support to publish and retrieve binaries to/from other users, in a peer-to-peer fashion. If you would like to discuss this project, join us on `guix-devel@gnu.org`.

## 5.4 Packages with Multiple Outputs

Often, packages defined in Guix have a single *output*—i.e., the source package leads to exactly one directory in the store. When running `guix install glibc`, one installs the default output of the GNU libc package; the default output is called `out`, but its name can be omitted as shown in this command. In this particular case, the default output of `glibc` contains all the C header files, shared libraries, static libraries, Info documentation, and other supporting files.

Sometimes it is more appropriate to separate the various types of files produced from a single source package into separate outputs. For instance, the GLib C library (used by GTK+ and related packages) installs more than 20 MiB of reference documentation as HTML pages. To save space for users who do not need it, the documentation goes to a separate output, called `doc`. To install the main GLib output, which contains everything but the documentation, one would run:

```
guix install glib
```

The command to install its documentation is:

```
guix install glib:doc
```

While the colon syntax works for command-line specification of package outputs, it will not work when using a package *variable* in Scheme code. For example, to add the documentation of `glib` to the globally installed packages of an `operating-system` (see Section 11.3 [operating-system Reference], page 253), a list of two items, the first one being the package *variable* and the second one the name of the output to select (a string), must be used instead:

```
(use-modules (gnu packages glib))
;; glib-with-documentation is the Guile symbol for the glib package
(operating-system
  ...
  (packages
    (append
      (list (list glib-with-documentation "doc"))
      %base-packages)))
```

Some packages install programs with different “dependency footprints”. For instance, the WordNet package installs both command-line tools and graphical user interfaces (GUIs). The former depend solely on the C library, whereas the latter depend on Tcl/Tk and the underlying X libraries. In this case, we leave the command-line tools in the default output, whereas the GUIs are in a separate output. This allows users who do not need the GUIs to save space. The `guix size` command can help find out about such situations (see Section 9.9 [Invoking `guix size`], page 220). `guix graph` can also be helpful (see Section 9.10 [Invoking `guix graph`], page 222).

There are several such multiple-output packages in the GNU distribution. Other conventional output names include `lib` for libraries and possibly header files, `bin` for stand-alone programs, and `debug` for debugging information (see Chapter 17 [Installing Debugging Files], page 801). The outputs of a package are listed in the third column of the output of `guix package --list-available` (see Section 5.2 [Invoking `guix package`], page 37).

## 5.5 Invoking `guix locate`

There’s so much free software out there that sooner or later, you will need to search for packages. The `guix search` command that we’ve seen before (see Section 5.2 [Invoking `guix package`], page 37) lets you search by keywords:

```
guix search video editor
```

Sometimes, you instead want to find which package provides a given file, and this is where `guix locate` comes in. Here is how you can find which package provides the `ls` command:

```
$ guix locate ls
coreutils@9.1      /gnu/store/...-coreutils-9.1/bin/ls
```

Of course the command works for any file, not just commands:

```
$ guix locate unistr.h
icu4c@71.1         /gnu/store/.../include/unicode/unistr.h
libunistring@1.0   /gnu/store/.../include/unistr.h
```

You may also specify *glob patterns* with wildcards. For example, here is how you would search for packages providing `.service` files:

```
$ guix locate -g '*.service'
man-db@2.11.1      .../lib/systemd/system/man-db.service
wpa-supPLICANT@2.10 .../system-services/fi.w1.wpa-supPLICANT1.service
```

The `guix locate` command relies on a database that maps file names to package names. By default, it automatically creates that database if it does not exist yet by traversing packages available *locally*, which can take a few minutes (depending on the size of your store and the speed of your storage device).

**Note:** For now, `guix locate` builds its database based on purely local knowledge—meaning that you will not find packages that never reached your store. Eventually it will support downloading a pre-built database so you can potentially find more packages.

By default, `guix locate` first tries to look for a system-wide database, usually under `/var/cache/guix/locate`; if it does not exist or is too old, it falls back to the per-user database, by default under `~/.cache/guix/locate`. On a multi-user system, administrators



may want to periodically update the system-wide database so that all users can benefit from it, for instance by setting up `package-database-service-type` (see Section 11.10.11 [File Search Services], page 400).

The general syntax is:

```
guix locate [options...] file...
```

... where *file* is the name of a file to search for (specifically, the “base name” of the file: files whose parent directories are called *file* are not matched).

The available options are as follows:

`--glob`

`-g` Interpret *file...* as *glob patterns*—patterns that may include wildcards, such as `*.scm` to denote all files ending in `.scm`.

`--stats` Display database statistics.

`--update`

`-u` Update the file database.

By default, the database is automatically updated when it is too old.

`--clear` Clear the database and re-populate it.

This option lets you start anew, ensuring old data is removed from the database, which also avoids having an endlessly growing database. By default `guix locate` automatically does that periodically, though infrequently.

`--database=file`

Use *file* as the database, creating it if necessary.

By default, `guix locate` picks the database under `~/.cache/guix` or `/var/cache/guix`, whichever is the most recent one.

`--method=method`

`-m method` Use *method* to select the set of packages to index. Possible values are:

**manifests**

This is the default method: it works by traversing profiles on the machine and recording packages it encounters—packages you or other users of the machine installed, directly or indirectly. It is fast but it can miss other packages available in the store but not referred to by any profile.

**store**

This is a slower but more exhaustive method: it checks among all the existing packages those that are available in the store and records them.

## 5.6 Invoking `guix gc`

Packages that are installed, but not used, may be *garbage-collected*. The `guix gc` command allows users to explicitly run the garbage collector to reclaim space from the `/gnu/store` directory. It is the *only* way to remove files from `/gnu/store`—removing files or directories manually may break it beyond repair!

The garbage collector has a set of known *roots*: any file under `/gnu/store` reachable from a root is considered *live* and cannot be deleted; any other file is considered *dead* and may be deleted. The set of garbage collector roots (“GC roots” for short) includes default user profiles; by default, the symlinks under `/var/guix/gcroots` represent these GC roots. New GC roots can be added with `guix build --root`, for example (see Section 9.1 [Invoking guix build], page 180). The `guix gc --list-roots` command lists them.

Prior to running `guix gc --collect-garbage` to make space, it is often useful to remove old generations from user profiles; that way, old package builds referenced by those generations can be reclaimed. This is achieved by running `guix package --delete-generations` (see Section 5.2 [Invoking guix package], page 37).

Our recommendation is to run a garbage collection periodically, or when you are short on disk space. For instance, to guarantee that at least 5 GB are available on your disk, simply run:

```
guix gc -F 5G
```

It is perfectly safe to run as a non-interactive periodic job (see Section 11.10.2 [Scheduled Job Execution], page 299, for how to set up such a job). Running `guix gc` with no arguments will collect as much garbage as it can, but that is often inconvenient: you may find yourself having to rebuild or re-download software that is “dead” from the GC viewpoint but that is necessary to build other pieces of software—e.g., the compiler tool chain.

The `guix gc` command has three modes of operation: it can be used to garbage-collect any dead files (the default), to delete specific files (the `--delete` option), to print garbage-collector information, or for more advanced queries. The garbage collection options are as follows:

`--collect-garbage[=min]`

`-C [min]` Collect garbage—i.e., unreachable `/gnu/store` files and sub-directories. This is the default operation when no option is specified.

When *min* is given, stop once *min* bytes have been collected. *min* may be a number of bytes, or it may include a unit as a suffix, such as MiB for mebibytes and GB for gigabytes (see Section “Block size” in *GNU Coreutils*).

When *min* is omitted, collect all the garbage.

`--free-space=free`

`-F free` Collect garbage until *free* space is available under `/gnu/store`, if possible; *free* denotes storage space, such as 500MiB, as described above.

When *free* or more is already available in `/gnu/store`, do nothing and exit immediately.

`--delete-generations[=duration]`

`-d [duration]`

Before starting the garbage collection process, delete all the generations older than *duration*, for all the user profiles and home environment generations; when run as root, this applies to all the profiles *of all the users*.

For example, this command deletes all the generations of all your profiles that are older than 2 months (except generations that are current), and then proceeds to free space until at least 10 GiB are available:

```
guix gc -d 2m -F 10G
```

**--delete**  
**-D** Attempt to delete all the store files and directories specified as arguments. This fails if some of the files are not in the store, or if they are still live.

**--list-failures**  
 List store items corresponding to cached build failures.  
 This prints nothing unless the daemon was started with **--cache-failures** (see Section 2.3 [Invoking guix-daemon], page 14).

**--list-roots**  
 List the GC roots owned by the user; when run as root, list *all* the GC roots.

**--list-busy**  
 List store items in use by currently running processes. These store items are effectively considered GC roots: they cannot be deleted.

**--clear-failures**  
 Remove the specified store items from the failed-build cache.  
 Again, this option only makes sense when the daemon is started with **--cache-failures**. Otherwise, it does nothing.

**--list-dead**  
 Show the list of dead files and directories still present in the store—i.e., files and directories no longer reachable from any root.

**--list-live**  
 Show the list of live store files and directories.

In addition, the references among existing store files can be queried:

**--references**  
**--referrers**  
 List the references (respectively, the referrers) of store files given as arguments.

**--requisites**  
**-R** List the requisites of the store files passed as arguments. Requisites include the store files themselves, their references, and the references of these, recursively. In other words, the returned list is the *transitive closure* of the store files.  
 See Section 9.9 [Invoking guix size], page 220, for a tool to profile the size of the closure of an element. See Section 9.10 [Invoking guix graph], page 222, for a tool to visualize the graph of references.

**--derivders**  
 Return the derivation(s) leading to the given store items (see Section 8.10 [Derivations], page 158).  
 For example, this command:

```
guix gc --derivders $(guix package -I ^emacs$ | cut -f4)
```

returns the `.drv` file(s) leading to the `emacs` package installed in your profile. Note that there may be zero matching `.drv` files, for instance because these files have been garbage-collected. There can also be more than one matching `.drv` due to fixed-output derivations.

Lastly, the following options allow you to check the integrity of the store and to control disk usage.

#### `--verify[=options]`

Verify the integrity of the store.

By default, make sure that all the store items marked as valid in the database of the daemon actually exist in `/gnu/store`.

When provided, *options* must be a comma-separated list containing one or more of `contents` and `repair`.

When passing `--verify=contents`, the daemon computes the content hash of each store item and compares it against its hash in the database. Hash mismatches are reported as data corruptions. Because it traverses *all the files in the store*, this command can take a long time, especially on systems with a slow disk drive.

Using `--verify=repair` or `--verify=contents,repair` causes the daemon to try to repair corrupt store items by fetching substitutes for them (see Section 5.3 [Substitutes], page 46). Because repairing is not atomic, and thus potentially dangerous, it is available only to the system administrator. A lightweight alternative, when you know exactly which items in the store are corrupt, is `guix build --repair` (see Section 9.1 [Invoking guix build], page 180).

#### `--optimize`

Optimize the store by hard-linking identical files—this is *deduplication*.

The daemon performs deduplication after each successful build or archive import, unless it was started with `--disable-deduplication` (see Section 2.3 [Invoking guix-daemon], page 14). Thus, this option is primarily useful when the daemon was running with `--disable-deduplication`.

#### `--vacuum-database`

Guix uses an sqlite database to keep track of the items in the store (see Section 8.9 [The Store], page 156). Over time it is possible that the database may grow to a large size and become fragmented. As a result, one may wish to clear the freed space and join the partially used pages in the database left behind from removed packages or after running the garbage collector. Running `sudo guix gc --vacuum-database` will lock the database and `VACUUM` the store, defragmenting the database and purging freed pages, unlocking the database when it finishes.

## 5.7 Invoking guix pull

Packages are installed or upgraded to the latest version available in the distribution currently available on your local machine. To update that distribution, along with the Guix tools, you must run `guix pull`: the command downloads the latest Guix source code and package descriptions, and deploys it. Source code is downloaded from a Git (<https://git-scm.com/book/en/>) repository, by default the official GNU Guix repository, though this can be customized. `guix pull` ensures that the code it downloads is *authentic* by verifying that commits are signed by Guix developers.

Specifically, `guix pull` downloads code from the *channels* (see Chapter 6 [Channels], page 68) specified by one of the following, in this order:

1. the `--channels` option;
2. the user's `~/.config/guix/channels.scm` file, unless `-q` is passed;
3. the system-wide `/etc/guix/channels.scm` file, unless `-q` is passed (on Guix System, this file can be declared in the operating system configuration, see [guix-configuration-channels], page 287);
4. the built-in default channels specified in the `%default-channels` variable.

On completion, `guix package` will use packages and package versions from this just-retrieved copy of Guix. Not only that, but all the Guix commands and Scheme modules will also be taken from that latest version. New `guix` sub-commands added by the update also become available.

Any user can update their Guix copy using `guix pull`, and the effect is limited to the user who ran `guix pull`. For instance, when user `root` runs `guix pull`, this has no effect on the version of Guix that user `alice` sees, and vice versa.

The result of running `guix pull` is a *profile* available under `~/.config/guix/current` containing the latest Guix.

The `--list-generations` or `-l` option lists past generations produced by `guix pull`, along with details about their provenance:

```
$ guix pull -l
Generation 1 Jun 10 2018 00:18:18
  guix 65956ad
    repository URL: https://git.guix.gnu.org/guix.git
    branch: origin/master
    commit: 65956ad3526ba09e1f7a40722c96c6ef7c0936fe

Generation 2 Jun 11 2018 11:02:49
  guix e0cc7f6
    repository URL: https://git.guix.gnu.org/guix.git
    branch: origin/master
    commit: e0cc7f669bec22c37481dd03a7941c7d11a64f1d

Generation 3 Jun 13 2018 23:31:07 (current)
  guix 844cc1c
    repository URL: https://git.guix.gnu.org/guix.git
    branch: origin/master
    commit: 844cc1c8f394f03b404c5bb3aee086922373490c
```

See Section 5.10 [Invoking `guix describe`], page 64, for other ways to describe the current status of Guix.

This `~/.config/guix/current` profile works exactly like the profiles created by `guix package` (see Section 5.2 [Invoking `guix package`], page 37). That is, you can list generations, roll back to the previous generation—i.e., the previous Guix—and so on:

```
$ guix pull --roll-back
switched from generation 3 to 2
```

```
$ guix pull --delete-generations=1
deleting /var/guix/profiles/per-user/charlie/current-guix-1-link
```

You can also use `guix package` (see Section 5.2 [Invoking `guix package`], page 37) to manage the profile by naming it explicitly:

```
$ guix package -p ~/.config/guix/current --roll-back
switched from generation 3 to 2
$ guix package -p ~/.config/guix/current --delete-generations=1
deleting /var/guix/profiles/per-user/charlie/current-guix-1-link
```

The `guix pull` command is usually invoked with no arguments, but it supports the following options:

`--url=url`

`--commit=commit`

`--branch=branch`

Download code for the `guix` channel from the specified *url*, at the given *commit* (a valid Git commit ID represented as a hexadecimal string or the name of a tag), or *branch*.

These options are provided for convenience, but you can also specify your configuration in the `~/.config/guix/channels.scm` file or using the `--channels` option (see below).

`--channels=file`

`-C file` Read the list of channels from *file* instead of `~/.config/guix/channels.scm` or `/etc/guix/channels.scm`. *file* must contain Scheme code that evaluates to a list of channel objects. See Chapter 6 [Channels], page 68, for more information.

`--no-channel-files`

`-q` Inhibit loading of the user and system channel files, `~/.config/guix/channels.scm` and `/etc/guix/channels.scm`.

`--news`

`-N` Display news written by channel authors for their users for changes made since the previous generation (see Chapter 6 [Channels], page 68). When `--details` is passed, additionally display new and upgraded packages.

You can view that information for previous generations with `guix pull -l`.

`--list-generations[=pattern]`

`-l [pattern]`

List all the generations of `~/.config/guix/current` or, if *pattern* is provided, the subset of generations that match *pattern*. The syntax of *pattern* is the same as with `guix package --list-generations` (see Section 5.2 [Invoking `guix package`], page 37).

By default, this prints information about the channels used in each revision as well as the corresponding news entries. If you pass `--details`, it will also print the list of packages added and upgraded in each generation compared to the previous one.

`--details`

Instruct `--list-generations` or `--news` to display more information about the differences between subsequent generations—see above.

**--roll-back**

Roll back to the previous *generation* of `~/.config/guix/current`—i.e., undo the last transaction.

**--switch-generation=*pattern*****-S *pattern***

Switch to a particular generation defined by *pattern*.

*pattern* may be either a generation number or a number prefixed with “+” or “-”. The latter means: move forward/backward by a specified number of generations. For example, if you want to return to the latest generation after `--roll-back`, use `--switch-generation=+1`.

**--delete-generations[=*pattern*]****-d [*pattern*]**

When *pattern* is omitted, delete all generations except the current one.

This command accepts the same patterns as `--list-generations`. When *pattern* is specified, delete the matching generations. When *pattern* specifies a duration, generations *older* than the specified duration match. For instance, `--delete-generations=1m` deletes generations that are more than one month old.

If the current generation matches, it is *not* deleted.

Note that deleting generations prevents rolling back to them. Consequently, this command must be used with care.

See Section 5.10 [Invoking `guix describe`], page 64, for a way to display information about the current generation only.

**--profile=*profile*****-p *profile***

Use *profile* instead of `~/.config/guix/current`.

**--dry-run**

**-n** Show which channel commit(s) would be used and what would be built or substituted but do not actually do it.

**--allow-downgrades**

Allow pulling older or unrelated revisions of channels than those currently in use.

By default, `guix pull` protects against so-called “downgrade attacks” whereby the Git repository of a channel would be reset to an earlier or unrelated revision of itself, potentially leading you to install older, known-vulnerable versions of software packages.

**Note:** Make sure you understand its security implications before using `--allow-downgrades`.

**--disable-authentication**

Allow pulling channel code without authenticating it.

By default, `guix pull` authenticates code downloaded from channels by verifying that its commits are signed by authorized developers, and raises an error if this is not the case. This option instructs it to not perform any such verification.

**Note:** Make sure you understand its security implications before using `--disable-authentication`.

**--no-check-certificate**

Do not validate the X.509 certificates of HTTPS servers.

When using this option, you have *absolutely no guarantee* that you are communicating with the authentic server responsible for the given URL. Unless the channel is authenticated, this makes you vulnerable to “man-in-the-middle” attacks.

**--system=system**

**-s system** Attempt to build for *system*—e.g., `i686-linux`—instead of the system type of the build host.

**--bootstrap**

Use the bootstrap Guile to build the latest Guix. This option is only useful to Guix developers.

The *channel* mechanism allows you to instruct `guix pull` which repository and branch to pull from, as well as *additional* repositories containing package modules that should be deployed. See Chapter 6 [Channels], page 68, for more information.

In addition, `guix pull` supports all the common build options (see Section 9.1.1 [Common Build Options], page 180).

## 5.8 Invoking guix time-machine

The `guix time-machine` command provides access to other revisions of Guix, for example to install older versions of packages, or to reproduce a computation in an identical environment. The revision of Guix to be used is defined by a commit or by a channel description file created by `guix describe` (see Section 5.10 [Invoking guix describe], page 64).

Let’s assume that you want to travel to those days of November 2020 when version 1.2.0 of Guix was released and, once you’re there, run the `guile` of that time:

```
guix time-machine --commit=v1.2.0 -- \
  environment -C --ad-hoc guile -- guile
```

The command above fetches Guix 1.2.0 (and possibly other channels specified by your `channels.scm` configuration files—see below) and runs its `guix environment` command to spawn an environment in a container running `guile` (`guix environment` has since been subsumed by `guix shell`; see Section 7.1 [Invoking guix shell], page 78). It’s like driving a DeLorean<sup>1</sup>! The first `guix time-machine` invocation can be expensive: it may have to download or even build a large number of packages; the result is cached though and subsequent commands targeting the same commit are almost instantaneous.

As for `guix pull`, in the absence of any options, `time-machine` fetches the latest commits of the channels specified in `~/.config/guix/channels.scm`, `/etc/guix/channels.scm`, or the default channels; the `-q` option lets you ignore these configuration files. The command:

```
guix time-machine -q -- build hello
```

<sup>1</sup> If you don’t know what a DeLorean is, consider traveling back to the 1980’s. (Back to the Future (1985) (<https://www.imdb.com/title/tt0088763/>))



will thus build the package **hello** as defined in the main branch of Guix, without any additional channel, which is in general a newer revision of Guix than you have installed. Time travel works in both directions!

**Note:** The history of Guix is immutable and **guix time-machine** provides the exact same software as they are in a specific Guix revision. Naturally, no security fixes are provided for old versions of Guix or its channels. A careless use of **guix time-machine** opens the door to security vulnerabilities. See Section 5.7 [Invoking **guix pull**], page 56.

**guix time-machine** raises an error when attempting to travel to commits older than “v0.16.0” (commit ‘4a0b87f0’), dated Dec. 2018. This is one of the oldest commits supporting the channel mechanism that makes “time travel” possible.

**Note:** Although it should technically be possible to travel to such an old commit, the ease to do so will largely depend on the availability of binary substitutes. When traveling to a distant past, some packages may not easily build from source anymore. One such example are old versions of OpenSSL whose tests would fail after a certain date. This particular problem can be worked around by running a *virtual build machine* with its clock set to the right time (see [build-vm], page 591).

The general syntax is:

```
guix time-machine options... -- command arg...
```

where *command* and *arg...* are passed unmodified to the **guix** command of the specified revision. The *options* that define this revision are the same as for **guix pull** (see Section 5.7 [Invoking **guix pull**], page 56):

```
--url=url
```

```
--commit=commit
```

```
--branch=branch
```

Use the **guix** channel from the specified *url*, at the given *commit* (a valid Git commit ID represented as a hexadecimal string or the name of a tag), or *branch*.

```
--channels=file
```

```
-C file
```

Read the list of channels from *file*. *file* must contain Scheme code that evaluates to a list of channel objects. See Chapter 6 [Channels], page 68, for more information.

```
--no-channel-files
```

```
-q
```

Inhibit loading of the user and system channel files, `~/.config/guix/channels.scm` and `/etc/guix/channels.scm`.

Thus, **guix time-machine -q** is equivalent to the following Bash command, using the “process substitution” syntax (see Section “Process Substitution” in *The GNU Bash Reference Manual*):

```
guix time-machine -C <(echo %default-channels) ...
```

Note that **guix time-machine** can trigger builds of channels and their dependencies, and these are controlled by the standard build options (see Section 9.1.1 [Common Build Options], page 180).

If `guix time-machine` is executed without any command, it prints the file name of the profile that would be used to execute the command. This is sometimes useful if you need to get store file name of the profile—e.g., when you want to `guix copy` it.

## 5.9 Inferiors

**Note:** The functionality described here is a “technology preview” as of version 8b23207. As such, the interface is subject to change.

Sometimes you might need to mix packages from the revision of Guix you’re currently running with packages available in a different revision of Guix. Guix *inferiors* allow you to achieve that by composing different Guix revisions in arbitrary ways.

Technically, an “inferior” is essentially a separate Guix process connected to your main Guix process through a REPL (see Section 8.13 [Invoking `guix repl`], page 175). The (`guix inferior`) module allows you to create inferiors and to communicate with them. It also provides a high-level interface to browse and manipulate the packages that an inferior provides—*inferior packages*.

When combined with channels (see Chapter 6 [Channels], page 68), inferiors provide a simple way to interact with a separate revision of Guix. For example, let’s assume you want to install in your profile the current `guile` package, along with the `guile-json` as it existed in an older revision of Guix—perhaps because the newer `guile-json` has an incompatible API and you want to run your code against the old API. To do that, you could write a manifest for use by `guix package --manifest` (see Section 8.4 [Writing Manifests], page 117); in that manifest, you would create an inferior for that old Guix revision you care about, and you would look up the `guile-json` package in the inferior:

```
(use-modules (guix inferior) (guix channels)
             (srfi srfi-1)) ;for 'first'

(define channels
  ;; This is the old revision from which we want to
  ;; extract guile-json.
  (list (channel
        (name 'guix)
        (url "https://git.guix.gnu.org/guix.git")
        (commit
         "65956ad3526ba09e1f7a40722c96c6ef7c0936fe")))))

(define inferior
  ;; An inferior representing the above revision.
  (inferior-for-channels channels))

;; Now create a manifest with the current "guile" package
;; and the old "guile-json" package.
(packages->manifest
 (list (first (lookup-inferior-packages inferior "guile-json"))
       (specification->package "guile")))
```

On its first run, `guix package --manifest` might have to build the channel you specified before it can create the inferior; subsequent runs will be much faster because the Guix revision will be cached.

The `(guix inferior)` module provides the following procedures to open an inferior:

`inferior-for-channels channels` `[#:cache-directory]` `[#:ttl]` [Procedure]

Return an inferior for *channels*, a list of channels. Use the cache at *cache-directory*, where entries can be reclaimed after *ttl* seconds. This procedure opens a new connection to the build daemon.

As a side effect, this procedure may build or substitute binaries for *channels*, which can take time.

`open-inferior directory` `[#:command "bin/guix"]` [Procedure]

Open the inferior Guix in *directory*, running *directory/command repl* or equivalent. Return `#f` if the inferior could not be launched.

The procedures listed below allow you to obtain and manipulate inferior packages.

`inferior-packages inferior` [Procedure]

Return the list of packages known to *inferior*.

`lookup-inferior-packages inferior name` `[version]` [Procedure]

Return the sorted list of inferior packages matching *name* in *inferior*, with highest version numbers first. If *version* is true, return only packages with a version number prefixed by *version*.

`inferior-package? obj` [Procedure]

Return true if *obj* is an inferior package.

`inferior-package-name package` [Procedure]

`inferior-package-version package` [Procedure]

`inferior-package-synopsis package` [Procedure]

`inferior-package-description package` [Procedure]

`inferior-package-home-page package` [Procedure]

`inferior-package-location package` [Procedure]

`inferior-package-inputs package` [Procedure]

`inferior-package-native-inputs package` [Procedure]

`inferior-package-propagated-inputs package` [Procedure]

`inferior-package-transitive-propagated-inputs package` [Procedure]

`inferior-package-native-search-paths package` [Procedure]

`inferior-package-transitive-native-search-paths package` [Procedure]

`inferior-package-search-paths package` [Procedure]

These procedures are the counterpart of package record accessors (see Section 8.2.1 [package Reference], page 104). Most of them work by querying the inferior *package* comes from, so the inferior must still be live when you call these procedures.

Inferior packages can be used transparently like any other package or file-like object in G-expressions (see Section 8.12 [G-Expressions], page 165). They are also transparently handled by the `packages->manifest` procedure, which is commonly used in manifests (see

Section 5.2 [Invoking guix package], page 37). Thus you can insert an inferior package pretty much anywhere you would insert a regular package: in manifests, in the `packages` field of your `operating-system` declaration, and so on.

## 5.10 Invoking guix describe

Often you may want to answer questions like: “Which revision of Guix am I using?” or “Which channels am I using?” This is useful information in many situations: if you want to *replicate* an environment on a different machine or user account, if you want to report a bug or to determine what change in the channels you are using caused it, or if you want to record your system state for reproducibility purposes. The `guix describe` command answers these questions.

When run from a `guix pulled guix`, `guix describe` displays the channel(s) that it was built from, including their repository URL and commit IDs (see Chapter 6 [Channels], page 68):

```
$ guix describe
Generation 10 Sep 03 2018 17:32:44 (current)
guix e0fa68c
  repository URL: https://git.guix.gnu.org/guix.git
  branch: master
  commit: e0fa68c7718fffd33d81af415279d6ddb518f727
```

If you’re familiar with the Git version control system, this is similar in spirit to `git describe`; the output is also similar to that of `guix pull --list-generations`, but limited to the current generation (see Section 5.7 [Invoking guix pull], page 56). Because the Git commit ID shown above unambiguously refers to a snapshot of Guix, this information is all it takes to describe the revision of Guix you’re using, and also to replicate it.

To make it easier to replicate Guix, `guix describe` can also be asked to return a list of channels instead of the human-readable description above:

```
$ guix describe -f channels
(list (channel
      (name 'guix)
      (url "https://git.guix.gnu.org/guix.git")
      (commit
        "e0fa68c7718fffd33d81af415279d6ddb518f727")
      (introduction
        (make-channel-introduction
          "9edb3f66fd807b096b48283debdccddccfea34bad"
          (openpgp-fingerprint
            "BBB0 2DDF 2CEA F6A8 0D1D  E643 A2A0 6DF2 A33A 54FA")))))
```

You can save this to a file and feed it to `guix pull -C` on some other machine or at a later point in time, which will instantiate *this exact Guix revision* (see Section 5.7 [Invoking guix pull], page 56). From there on, since you’re able to deploy the same revision of Guix, you can just as well *replicate a complete software environment*. We humbly think that this is *awesome*, and we hope you’ll like it too!

The details of the options supported by `guix describe` are as follows:

```

--format=format
-f format  Produce output in the specified format, one of:

    human      produce human-readable output;

    channels   produce a list of channel specifications that can be passed to guix
                pull -C or installed as ~/.config/guix/channels.scm (see Sec-
                tion 5.7 [Invoking guix pull], page 56);

    channels-sans-intro
                like channels, but omit the introduction field; use it to produce
                a channel specification suitable for Guix version 1.1.0 or earlier—
                the introduction field has to do with channel authentication (see
                Chapter 6 [Channels], page 68) and is not supported by these older
                versions;

    json       produce a list of channel specifications in JSON format;

    recutils   produce a list of channel specifications in Recutils format.

--list-formats
                Display available formats for --format option.

--profile=profile
-p profile    Display information about profile.

```

## 5.11 Invoking guix archive

The **guix archive** command allows users to *export* files from the store into a single archive, and to later *import* them on a machine that runs Guix. In particular, it allows store files to be transferred from one machine to the store on another machine.

**Note:** If you're looking for a way to produce archives in a format suitable for tools other than Guix, see Section 7.3 [Invoking guix pack], page 90.

To export store files as an archive to standard output, run:

```
guix archive --export options specifications...
```

*specifications* may be either store file names or package specifications, as for **guix package** (see Section 5.2 [Invoking guix package], page 37). For instance, the following command creates an archive containing the **gui** output of the **git** package and the main output of **emacs**:

```
guix archive --export git:gui /gnu/store/...-emacs-24.3 > great.nar
```

If the specified packages are not built yet, **guix archive** automatically builds them. The build process may be controlled with the common build options (see Section 9.1.1 [Common Build Options], page 180).

To transfer the **emacs** package to a machine connected over SSH, one would run:

```
guix archive --export -r emacs | ssh the-machine guix archive --import
```

Similarly, a complete user profile may be transferred from one machine to another like this:

```
guix archive --export -r $(readlink -f ~/.guix-profile) | \
ssh the-machine guix archive --import
```

However, note that, in both examples, all of **emacs** and the profile as well as all of their dependencies are transferred (due to **-r**), regardless of what is already available in the store on the target machine. The **--missing** option can help figure out which items are missing from the target store. The **guix copy** command simplifies and optimizes this whole process, so this is probably what you should use in this case (see Section 9.13 [Invoking **guix copy**], page 233).

Each store item is written in the *normalized archive* or *nar* format (described below), and the output of **guix archive --export** (and input of **guix archive --import**) is a *nar bundle*.

The nar format is comparable in spirit to ‘tar’, but with differences that make it more appropriate for our purposes. First, rather than recording all Unix metadata for each file, the nar format only mentions the file type (regular, directory, or symbolic link); Unix permissions and owner/group are dismissed. Second, the order in which directory entries are stored always follows the order of file names according to the C locale collation order. This makes archive production fully deterministic.

That nar bundle format is essentially the concatenation of zero or more nars along with metadata for each store item it contains: its file name, references, corresponding derivation, and a digital signature.

When exporting, the daemon digitally signs the contents of the archive, and that digital signature is appended. When importing, the daemon verifies the signature and rejects the import in case of an invalid signature or if the signing key is not authorized.

The main options are:

**--export** Export the specified store files or packages (see below). Write the resulting archive to the standard output.

Dependencies are *not* included in the output, unless **--recursive** is passed.

**-r**

**--recursive**

When combined with **--export**, this instructs **guix archive** to include dependencies of the given items in the archive. Thus, the resulting archive is self-contained: it contains the closure of the exported store items.

**--import** Read an archive from the standard input, and import the files listed therein into the store. Abort if the archive has an invalid digital signature, or if it is signed by a public key not among the authorized keys (see **--authorize** below).

**--missing**

Read a list of store file names from the standard input, one per line, and write on the standard output the subset of these files missing from the store.

**--generate-key[=*parameters*]**

Generate a new key pair for the daemon. This is a prerequisite before archives can be exported with **--export**. This operation is usually instantaneous but it can take time if the system’s entropy pool needs to be refilled. On Guix System, **guix-service-type** takes care of generating this key pair the first boot.

The generated key pair is typically stored under **/etc/guix**, in **signing-key.pub** (public key) and **signing-key.sec** (private key, which must be kept

secret). When *parameters* is omitted, an ECDSA key using the Ed25519 curve is generated, or, for Libgcrypt versions before 1.6.0, it is a 4096-bit RSA key. Alternatively, *parameters* can specify **genkey** parameters suitable for Libgcrypt (see Section “General public-key related Functions” in *The Libgcrypt Reference Manual*).

#### **--authorize**

Authorize imports signed by the public key passed on standard input. The public key must be in “s-expression advanced format”—i.e., the same format as the **signing-key.pub** file.

The list of authorized keys is kept in the human-editable file **/etc/guix/acl**. The file contains “advanced-format s-expressions” (<https://people.csail.mit.edu/rivest/Sexp.txt>) and is structured as an access-control list in the Simple Public-Key Infrastructure (SPKI) (<https://theworld.com/~cme/spki.txt>).

#### **--extract=directory**

##### **-x directory**

Read a single-item archive as served by substitute servers (see Section 5.3 [Substitutes], page 46) and extract it to *directory*. This is a low-level operation needed in only very narrow use cases; see below.

For example, the following command extracts the substitute for Emacs served by [bordeaux.guix.gnu.org](https://bordeaux.guix.gnu.org) to **/tmp/emacs**:

```
$ wget -O - \
  https://bordeaux.guix.gnu.org/nar/gzip/...-emacs-24.5 \
  | gunzip | guix archive -x /tmp/emacs
```

Single-item archives are different from multiple-item archives produced by **guix archive --export**; they contain a single store item, and they do *not* embed a signature. Thus this operation does *no* signature verification and its output should be considered unsafe.

The primary purpose of this operation is to facilitate inspection of archive contents coming from possibly untrusted substitute servers (see Section 9.12 [Invoking guix challenge], page 231).

#### **--list**

##### **-t**

Read a single-item archive as served by substitute servers (see Section 5.3 [Substitutes], page 46) and print the list of files it contains, as in this example:

```
$ wget -O - \
  https://bordeaux.guix.gnu.org/nar/lzip/...-emacs-26.3 \
  | lzip -d | guix archive -t
```

## 6 Channels

Guix and its package collection are updated by running `guix pull`. By default `guix pull` downloads and deploys Guix itself from the official GNU Guix repository. This can be customized by providing a file specifying the set of *channels* to pull from (see Section 5.7 [Invoking `guix pull`], page 56). A channel specifies the URL and branch of a Git repository to be deployed, and `guix pull` can be instructed to pull from one or more channels. In other words, channels can be used to *customize* and to *extend* Guix, as we will see below. Guix is able to take into account security concerns and deal with authenticated updates.

### 6.1 Specifying Additional Channels

You can specify *additional channels* to pull from. To use a channel, write `~/.config/guix/channels.scm` to instruct `guix pull` to pull from it *in addition* to the default Guix channel(s):

```
;; Add variant packages to those Guix provides.
(cons (channel
      (name 'variant-packages)
      (url "https://example.org/variant-packages.git"))
      %default-channels)
```

Note that the snippet above is (as always!) Scheme code; we use `cons` to add a channel the list of channels that the variable `%default-channels` is bound to (see Section “Pairs” in *GNU Guile Reference Manual*). With this file in place, `guix pull` builds not only Guix but also the package modules from your own repository. The result in `~/.config/guix/current` is the union of Guix with your own package modules:

```
$ guix describe
Generation 19 Aug 27 2018 16:20:48
guix d894ab8
  repository URL: https://git.guix.gnu.org/guix.git
  branch: master
  commit: d894ab8e9bfabcefa6c49d9ba2e834dd5a73a300
variant-packages dd3df5e
  repository URL: https://example.org/variant-packages.git
  branch: master
  commit: dd3df5e2c8818760a8fc0bd699e55d3b69fef2bb
```

The output of `guix describe` above shows that we’re now running Generation 19 and that it includes both Guix and packages from the `variant-packages` channel (see Section 5.10 [Invoking `guix describe`], page 64).

### 6.2 Using a Custom Guix Channel

The channel called `guix` specifies where Guix itself—its command-line tools as well as its package collection—should be downloaded. For instance, suppose you want to update from another copy of the Guix repository at `example.org`, and specifically the `super-hacks` branch, you can write in `~/.config/guix/channels.scm` this specification:

```
;; Tell 'guix pull' to use another repo.
```



```
(list (channel
      (name 'guix)
      (url "https://example.org/another-guix.git")
      (branch "super-hacks")))
```

From there on, `guix pull` will fetch code from the `super-hacks` branch of the repository at `example.org`. The authentication concern is addressed below (see Section 6.5 [Channel Authentication], page 71).

Note that you can specify a local directory on the `url` field above if the channel that you intend to use resides on a local file system. However, in this case `guix` checks said directory for ownership before any further processing. This means that if the user is not the directory owner, but wants to use it as their default, they will then need to set it as a safe directory in their global git configuration file. Otherwise, `guix` will refuse to even read it. Supposing your system-wide local directory is at `/src/guix.git`, you would then create a git configuration file at `~/.gitconfig` with the following contents:

```
[safe]
    directory = /src/guix.git
```

This also applies to the root user unless when called with `sudo` by the directory owner.

## 6.3 Replicating Guix

The `guix describe` command shows precisely which commits were used to build the instance of Guix we're using (see Section 5.10 [Invoking `guix describe`], page 64). We can replicate this instance on another machine or at a different point in time by providing a channel specification “pinned” to these commits that looks like this:

```
;; Deploy specific commits of my channels of interest.
(list (channel
      (name 'guix)
      (url "https://git.guix.gnu.org/guix.git")
      (commit "6298c3ffd9654d3231a6f25390b056483e8f407c"))
      (channel
      (name 'variant-packages)
      (url "https://example.org/variant-packages.git")
      (commit "dd3df5e2c8818760a8fc0bd699e55d3b69fef2bb")))
```

To obtain this pinned channel specification, the easiest way is to run `guix describe` and to save its output in the `channels` format in a file, like so:

```
guix describe -f channels > channels.scm
```

The resulting `channels.scm` file can be passed to the `-C` option of `guix pull` (see Section 5.7 [Invoking `guix pull`], page 56) or `guix time-machine` (see Section 5.8 [Invoking `guix time-machine`], page 60), as in this example:

```
guix time-machine -C channels.scm -- shell python -- python3
```

Given the `channels.scm` file, the command above will always fetch the *exact same Guix instance*, then use that instance to run the exact same Python (see Section 7.1 [Invoking `guix shell`], page 78). On any machine, at any time, it ends up running the exact same binaries, bit for bit.

Pinned channels address a problem similar to “lock files” as implemented by some deployment tools—they let you pin and reproduce a set of packages. In the case of Guix though, you are effectively pinning the entire package set as defined at the given channel commits; in fact, you are pinning all of Guix, including its core modules and command-line tools. You’re also getting strong guarantees that you are, indeed, obtaining the exact same software.

This gives you super powers, allowing you to track the provenance of binary artifacts with very fine grain, and to reproduce software environments at will—some sort of “meta reproducibility” capabilities, if you will. See Section 5.9 [Inferiors], page 62, for another way to take advantage of these super powers.

## 6.4 Customizing the System-Wide Guix

If you’re running Guix System or building system images with it, maybe you will want to customize the system-wide guix it provides—specifically, `/run/current-system/profile/bin/guix`. For example, you might want to provide additional channels or to pin its revision.

This can be done using the `guix-for-channels` procedure, which returns a package for the given channels, and using it as part of your operating system configuration, as in this example:

```
(use-modules (gnu packages package-management)
             (guix channels))

(define my-channels
  ;; Channels that should be available to
  ;; /run/current-system/profile/bin/guix.
  (append
    (list (channel
           (name 'guix-science)
           (url "https://codeberg.org/guix-science/guix-science")
           (branch "master"))))
    %default-channels))

(operating-system
  ;; ...
  (services
    ;; Change the package used by 'guix-service-type'.
    (modify-services %base-services
      (guix-service-type
        config => (guix-configuration
                    (inherit config)
                    (channels my-channels)
                    (guix (guix-for-channels my-channels)))))))
```

The resulting operating system will have both the `guix` and the `guix-science` channels visible by default. The `channels` field of `guix-configuration` above further ensures that

`/etc/guix/channels.scm`, which is used by `guix pull`, specifies the same set of channels (see [guix-configuration-channels], page 287).

The `(gnu packages package-management)` module exports the `guix-for-channels` procedure, described below.

**guix-for-channels channels** [Procedure]

Return a package corresponding to *channels*.

The result is a “regular” package, which can be used in `guix-configuration` as shown above or in any other place that expects a package.

## 6.5 Channel Authentication

The `guix pull` and `guix time-machine` commands *authenticate* the code retrieved from channels: they make sure each commit that is fetched is signed by an authorized developer. The goal is to protect from unauthorized modifications to the channel that would lead users to run malicious code.

As a user, you must provide a *channel introduction* in your channels file so that Guix knows how to authenticate its first commit. A channel specification, including its introduction, looks something along these lines:

```
(channel
  (name 'some-channel)
  (url "https://example.org/some-channel.git")
  (introduction
    (make-channel-introduction
      "6f0d8cc0d88abb59c324b2990bfee2876016bb86"
      (openpgp-fingerprint
        "CABB A931 COFF EEC6 900D  OCFB 090B 1199 3D9A EBB5")))))
```

The specification above shows the name and URL of the channel. The call to `make-channel-introduction` above specifies that authentication of this channel starts at commit `6f0d8cc...`, which is signed by the OpenPGP key with fingerprint `CABB A931...`

For the main channel, called `guix`, you automatically get that information from your Guix installation. For other channels, include the channel introduction provided by the channel authors in your `channels.scm` file. Make sure you retrieve the channel introduction from a trusted source since that is the root of your trust.

If you’re curious about the authentication mechanics, read on!

## 6.6 Channels with Substitutes

When running `guix pull`, Guix will first compile the definitions of every available package. This is an expensive operation for which substitutes (see Section 5.3 [Substitutes], page 46) may be available. The following snippet in `channels.scm` will ensure that `guix pull` uses the latest commit with available substitutes for the package definitions: this is done by querying the continuous integration server at `https://ci.guix.gnu.org`.

```
(use-modules (guix ci))

(list (channel-with-substitutes-available
```

```
%default-guix-channel
"https://ci.guix.gnu.org"))
```

Note that this does not mean that all the packages that you will install after running `guix pull` will have available substitutes. It only ensures that `guix pull` will not try to compile package definitions. This is particularly useful when using machines with limited resources.

## 6.7 Creating a Channel

Let's say you have a bunch of custom package variants or personal packages that you think would make little sense to contribute to the Guix project, but would like to have these packages transparently available to you at the command line. By creating a *channel*, you can use and publish such a package collection. This involves the following steps:

1. A channel lives in a Git repository so the first step, when creating a channel, is to create its repository:

```
mkdir my-channel
cd my-channel
git init
```

2. The next step is to create files containing package modules (see Section 8.1 [Package Modules], page 100), each of which will contain one or more package definitions (see Section 8.2 [Defining Packages], page 101). A channel can provide things other than packages, such as build systems or services; we're using packages as it's the most common use case.

For example, Alice might want to provide a module called `(alice packages greetings)` that will provide her favorite “hello world” implementations. To do that Alice will create a directory corresponding to that module name.

```
mkdir -p alice/packages
$EDITOR alice/packages/greetings.scm
git add alice/packages/greetings.scm
```

You can name your package modules however you like; the main constraint to keep in mind is to avoid name clashes with other package collections, which is why our hypothetical Alice wisely chose the `(alice packages ...)` name space.

Note that you can also place modules in a sub-directory of the repository; see Section 6.8 [Package Modules in a Sub-directory], page 73, for more info on that.

3. With this first module in place, the next step is to test the packages it provides. This can be done with `guix build`, which needs to be told to look for modules in the Git checkout. For example, assuming `(alice packages greetings)` provides a package called `hi-from-alice`, Alice will run this command from the Git checkout:

```
guix build -L. hi-from-alice
```

... where `-L.` adds the current directory to Guile's load path (see Section “Load Paths” in *GNU Guile Reference Manual*).

4. It might take Alice a few iterations to obtain satisfying package definitions. Eventually Alice will commit this file:

```
git commit
```

As a channel author, consider bundling authentication material with your channel so that users can authenticate it. See Section 6.5 [Channel Authentication], page 71, and Section 6.10 [Specifying Channel Authorizations], page 74, for info on how to do it.

5. To use Alice’s channel, anyone can now add it to their channel file (see Section 6.1 [Specifying Additional Channels], page 68) and run `guix pull` (see Section 5.7 [Invoking `guix pull`], page 56):

```
$EDITOR ~/.config/guix/channels.scm
guix pull
```

Guix will now behave as if the root directory of that channel’s Git repository had been permanently added to the Guile load path. In this example, (`alice packages greetings`) will automatically be found by the `guix` command.

Voilà!

**Warning:** Before you publish your channel, we would like to share a few words of caution:

- Before publishing a channel, please consider contributing your package definitions to Guix proper (see Chapter 22 [Contributing], page 816). Guix as a project is open to free software of all sorts, and packages in Guix proper are readily available to all Guix users and benefit from the project’s quality assurance process.
- Package modules and package definitions are Scheme code that uses various programming interfaces (APIs). Those may occasionally change, in accordance with our deprecation policy (see Section 22.17 [Deprecation Policy], page 856), and it is your responsibility as a channel author to adjust your code accordingly.
- Corollary: if you’re using an external channel and that channel breaks, please *report the issue to the channel authors*, not to the Guix project.

You’ve been warned! Having said this, we believe external channels are a practical way to exert your freedom to augment Guix’ package collection and to share your improvements, which are basic tenets of free software (<https://www.gnu.org/philosophy/free-sw.html>). Please email us at `guix-devel@gnu.org` if you’d like to discuss this.

## 6.8 Package Modules in a Sub-directory

As a channel author, you may want to keep your channel modules in a sub-directory. If your modules are in the sub-directory `guix`, you must add a meta-data file `.guix-channel` that contains:

```
(channel
 (version 0)
 (directory "guix"))
```

The modules must be **underneath** the specified directory, as the `directory` changes Guile’s `load-path`. For example, if `.guix-channel` has `(directory "base")`, then a module defined as `(define-module (gnu packages fun))` must be located at `base/gnu/packages/fun.scm`.

Doing this allows for only parts of a repository to be used as a channel, as Guix expects valid Guile modules when pulling. For instance, `guix deploy` machine configuration files are not valid Guile modules, and treating them as such would make `guix pull` fail.

## 6.9 Declaring Channel Dependencies

Channel authors may decide to augment a package collection provided by other channels. They can declare their channel to be dependent on other channels in a meta-data file `.guix-channel`, which is to be placed in the root of the channel repository.

The meta-data file should contain a simple S-expression like this:

```
(channel
  (version 0)
  (dependencies
    (channel
      (name some-collection)
      (url "https://example.org/first-collection.git")

      ;; The 'introduction' bit below is optional: you would
      ;; provide it for dependencies that can be authenticated.
      (introduction
        (channel-introduction
          (version 0)
          (commit "a8883b58dc82e167c96506cf05095f37c2c2c6cd")
          (signer "CABB A931 COFF EEC6 900D OCFB 090B 1199 3D9A EBB5")))))
    (channel
      (name some-other-collection)
      (url "https://example.org/second-collection.git")
      (branch "testing")))))
```

In the above example this channel is declared to depend on two other channels, which will both be fetched automatically. The modules provided by the channel will be compiled in an environment where the modules of all these declared channels are available.

For the sake of reliability and maintainability, you should avoid dependencies on channels that you don't control, and you should aim to keep the number of dependencies to a minimum.

## 6.10 Specifying Channel Authorizations

As we saw above, Guix ensures the source code it pulls from channels comes from authorized developers. As a channel author, you need to specify the list of authorized developers in the `.guix-authorizations` file in the channel's Git repository. The authentication rule is simple: each commit must be signed by a key listed in the `.guix-authorizations` file of its parent commit(s)<sup>1</sup> The `.guix-authorizations` file looks like this:

```
;; Example '.guix-authorizations' file.
```

---

<sup>1</sup> Git commits form a *directed acyclic graph* (DAG). Each commit can have zero or more parents; “regular” commits have one parent and merge commits have two parent commits. Read *Git for Computer Scientists* (<https://eagain.net/articles/git-for-computer-scientists/>) for a great overview.

```
(authorizations
  (version 0)                                ;current file format version

  (("AD17 A21E F8AE D8F1 CC02 DBD9 F8AE D8F1 765C 61E3"
    (name "alice")))
  ("2A39 3FFF 68F4 EF7A 3D29 12AF 68F4 EF7A 22FB B2D5"
    (name "bob")))
  ("CABB A931 C0FF EEC6 900D 0CFB 090B 1199 3D9A EBB5"
    (name "charlie"))))
```

Each fingerprint is followed by optional key/value pairs, as in the example above. Currently these key/value pairs are ignored.

This authentication rule creates a chicken-and-egg issue: how do we authenticate the first commit? Related to that: how do we deal with channels whose repository history contains unsigned commits and lack `.guix-authorizations`? And how do we fork existing channels?

Channel introductions answer these questions by describing the first commit of a channel that should be authenticated. The first time a channel is fetched with `guix pull` or `guix time-machine`, the command looks up the introductory commit and verifies that it is signed by the specified OpenPGP key. From then on, it authenticates commits according to the rule above. Authentication fails if the target commit is neither a descendant nor an ancestor of the introductory commit.

Additionally, your channel must provide all the OpenPGP keys that were ever mentioned in `.guix-authorizations`, stored as `.key` files, which can be either binary or “ASCII-armored”. By default, those `.key` files are searched for in the branch named `keyring` but you can specify a different branch name in `.guix-channel` like so:

```
(channel
  (version 0)
  (keyring-reference "my-keyring-branch"))
```

To summarize, as the author of a channel, there are three things you have to do to allow users to authenticate your code:

1. Export the OpenPGP keys of past and present committers with `gpg --export` and store them in `.key` files, by default in a branch named `keyring` (we recommend making it an *orphan branch*).
2. Introduce an initial `.guix-authorizations` in the channel’s repository. Do that in a signed commit (see Section 22.14 [Commit Access], page 851, for information on how to sign Git commits).
3. Advertise the channel introduction, for instance on your channel’s web page. The channel introduction, as we saw above, is the commit/key pair—i.e., the commit that introduced `.guix-authorizations`, and the fingerprint of the OpenPGP used to sign it.

Before pushing to your public Git repository, you can run `guix git authenticate` to verify that you did sign all the commits you are about to push with an authorized key:

```
guix git authenticate commit signer
```

where *commit* and *signer* are your channel introduction. See Section 7.5 [Invoking `guix git authenticate`], page 98, for details.

Publishing a signed channel requires discipline: any mistake, such as an unsigned commit or a commit signed by an unauthorized key, will prevent users from pulling from your channel—well, that’s the whole point of authentication! Pay attention to merges in particular: merge commits are considered authentic if and only if they are signed by a key present in the `.guix-authorizations` file of *both* branches.

## 6.11 Primary URL

Channel authors can indicate the primary URL of their channel’s Git repository in the `.guix-channel` file, like so:

```
(channel
  (version 0)
  (url "https://example.org/guix.git"))
```

This allows `guix pull` to determine whether it is pulling code from a mirror of the channel; when that is the case, it warns the user that the mirror might be stale and displays the primary URL. That way, users cannot be tricked into fetching code from a stale mirror that does not receive security updates.

This feature only makes sense for authenticated repositories, such as the official `guix` channel, for which `guix pull` ensures the code it fetches is authentic.

## 6.12 Writing Channel News

Channel authors may occasionally want to communicate to their users information about important changes in the channel. You’d send them all an email, but that’s not convenient.

Instead, channels can provide a *news file*; when the channel users run `guix pull`, that news file is automatically read and `guix pull --news` can display the announcements that correspond to the new commits that have been pulled, if any.

To do that, channel authors must first declare the name of the news file in their `.guix-channel` file:

```
(channel
  (version 0)
  (news-file "etc/news.txt"))
```

The news file itself, `etc/news.txt` in this example, must look something like this:

```
(channel-news
  (version 0)
  (entry (tag "the-bug-fix")
    (title (en "Fixed terrible bug")
      (fr "Oh la la"))
    (body (en "@emph{Good news}! It's fixed!")
      (eo "Certe ĝi pli bone funkcias nun!"))))
  (entry (commit "bdcabe815cd28144a2d2b4bc3c5057b051fa9906")
    (title (en "Added a great package")
      (ca "Què vol dir guix?"))
    (body (en "Don't miss the @code{hello} package!"))))
```



While the news file is using the Scheme syntax, avoid naming it with a `.scm` extension or else it will get picked up when building the channel and yield an error since it is not a valid module. Alternatively, you can move the channel module to a subdirectory and store the news file in another directory.

The file consists of a list of *news entries*. Each entry is associated with a commit or tag: it describes changes made in this commit, possibly in preceding commits as well. Users see entries only the first time they obtain the commit the entry refers to.

The `title` field should be a one-line summary while `body` can be arbitrarily long, and both can contain Texinfo markup (see Section “Overview” in *GNU Texinfo*). Both the title and body are a list of language tag/message tuples, which allows `guix pull` to display news in the language that corresponds to the user’s locale.

If you want to translate news using a gettext-based workflow, you can extract translatable strings with `xgettext` (see Section “xgettext Invocation” in *GNU Gettext Utilities*). For example, assuming you write news entries in English first, the command below creates a PO file containing the strings to translate:

```
xgettext -o news.po -l scheme -ken etc/news.txt
```

To sum up, yes, you could use your channel as a blog. But beware, this is *not quite* what your users might expect.

## 7 Development

If you are a software developer, Guix provides tools that you should find helpful—independently of the language you’re developing in. This is what this chapter is about.

The `guix shell` command provides a convenient way to set up one-off software environments, be it for development purposes or to run a command without installing it in your profile. The `guix pack` command allows you to create *application bundles* that can be easily distributed to users who do not run Guix.

### 7.1 Invoking guix shell

The purpose of `guix shell` is to make it easy to create one-off software environments, without changing one’s profile. It is typically used to create development environments; it is also a convenient way to run applications without “polluting” your profile.

**Note:** The `guix shell` command was recently introduced to supersede `guix environment` (see Section 7.2 [Invoking guix environment], page 85). If you are familiar with `guix environment`, you will notice that it is similar but also—we hope!—more convenient.

The general syntax is:

```
guix shell [options] [package...]
```

Sometimes an interactive shell session is not desired. An arbitrary command may be invoked by placing the `--` token to separate the command from the rest of the arguments.

The following example creates an environment containing Python and NumPy, building or downloading any missing package, and runs the `python3` command in that environment:

```
guix shell python python-numpy -- python3
```

Note that it is necessary to include the main `python` package in this command even if it is already installed into your environment. This is so that the shell environment knows to set `PYTHONPATH` and other related variables. The shell environment cannot check the previously installed environment, because then it would be non-deterministic. This is true for most libraries: their corresponding language package should be included in the shell invocation.

**Note:** `guix shell` can be also be used as a script interpreter, also known as *shebang*. Here is an example self-contained Python script making use of this feature:

```
#!/usr/bin/env -S guix shell python python-numpy -- python3
import numpy
print("This is numpy", numpy.version.version)
```

You may pass any `guix shell` option, but there’s one caveat: the Linux kernel has a limit of 127 bytes on shebang length.

Development environments can be created as in the example below, which spawns an interactive shell containing all the dependencies and environment variables needed to work on Inkscape:

```
guix shell --development inkscape
```

Exiting the shell places the user back in the original environment before `guix shell` was invoked. The next garbage collection (see Section 5.6 [Invoking `guix gc`], page 53) may clean up packages that were installed in the environment and that are no longer used outside of it.

As an added convenience, `guix shell` will try to do what you mean when it is invoked interactively without any other arguments as in:

```
guix shell
```

If it finds a `manifest.scm` in the current working directory or any of its parents, it uses this manifest as though it was given via `--manifest`. Likewise, if it finds a `guix.scm` in the same directories, it uses it to build a development profile as though both `--development` and `--file` were present. In either case, the file will only be loaded if the directory it resides in is listed in `~/.config/guix/shell-authorized-directories`. This provides an easy way to define, share, and enter development environments.

By default, the shell session or command runs in an *augmented* environment, where the new packages are added to search path environment variables such as `PATH`. You can, instead, choose to create an *isolated* environment containing nothing but the packages you asked for. Passing the `--pure` option clears environment variable definitions found in the parent environment<sup>1</sup>; passing `--container` goes one step further by spawning a *container* isolated from the rest of the system:

```
guix shell --container emacs gcc-toolchain
```

The command above spawns an interactive shell in a container where nothing but `emacs`, `gcc-toolchain`, and their dependencies is available. The container lacks network access and shares no files other than the current working directory with the surrounding environment. This is useful to prevent access to system-wide resources such as `/usr/bin` on foreign distros.

This `--container` option can also prove useful if you wish to run a security-sensitive application, such as a web browser, in an isolated environment. For example, the command below launches Ungogled-Chromium in an isolated environment, which:

- shares network access with the host
- inherits host's environment variables `DISPLAY` and `XAUTHORITY`
- has access to host's authentication records from the `XAUTHORITY` file
- has no information about host's current directory

```
guix shell --container --network --no-cwd ungoogled-chromium \
  --preserve='^XAUTHORITY$' --expose="${XAUTHORITY}" \
  --preserve='^DISPLAY$' -- chromium
```

`guix shell` defines the `GUIX_ENVIRONMENT` variable in the shell it spawns; its value is the file name of the profile of this environment. This allows users to, say, define a specific prompt for development environments in their `.bashrc` (see Section “Bash Startup Files” in *The GNU Bash Reference Manual*):

```
if [ -n "$GUIX_ENVIRONMENT" ]
then
  export PS1="\u@\h \w [dev]\$ "
```

<sup>1</sup> Be sure to use the `--check` option the first time you use `guix shell` interactively to make sure the shell does not undo the effect of `--pure`.

```
fi
```

... or to browse the profile:

```
$ ls "$GUIX_ENVIRONMENT/bin"
```

The available options are summarized below.

**--check** Set up the environment and check whether the shell would clobber environment variables. It's a good idea to use this option the first time you run `guix shell` for an interactive session to make sure your setup is correct.

For example, if the shell modifies the `PATH` environment variable, report it since you would get a different environment than what you asked for.

Such problems usually indicate that the shell startup files are unexpectedly modifying those environment variables. For example, if you are using Bash, make sure that environment variables are set or modified in `~/.bash_profile` and *not* in `~/.bashrc`—the former is sourced only by log-in shells. See Section “Bash Startup Files” in *The GNU Bash Reference Manual*, for details on Bash start-up files.

**--development**

**-D** Cause `guix shell` to include in the environment the dependencies of the following package rather than the package itself. This can be combined with other packages. For instance, the command below starts an interactive shell containing the build-time dependencies of GNU Guile, plus Autoconf, Automake, and Libtool:

```
guix shell -D guile autoconf automake libtool
```

**--expression=expr**

**-e expr** Create an environment for the package or list of packages that `expr` evaluates to.

For example, running:

```
guix shell -D -e '(@ (gnu packages maths) petsc-openmpi)'
```

starts a shell with the environment for this specific variant of the PETSc package.

Running:

```
guix shell -e '(@ (gnu) %base-packages)'
```

starts a shell with all the base system packages available.

The above commands only use the default output of the given packages. To select other outputs, two element tuples can be specified:

```
guix shell -e '(list (@ (gnu packages bash) bash) "include")'
```

See [package-development-manifest], page 120, for information on how to write a manifest for the development environment of a package.

**--file=file**

**-f file** Create an environment containing the package or list of packages that the code within `file` evaluates to.

As an example, *file* might contain a definition like this (see Section 8.2 [Defining Packages], page 101):

```
(use-modules (guix)
             (gnu packages gdb)
             (gnu packages autotools)
             (gnu packages texinfo))

;; Augment the package definition of GDB with the build tools
;; needed when developing GDB (and which are not needed when
;; simply installing it.)
(package
  (inherit gdb)
  (native-inputs (modify-inputs (package-native-inputs gdb)
                                (prepend autoconf-2.69 automake texinfo))))
```

With the file above, you can enter a development environment for GDB by running:

```
guix shell -D -f gdb-devel.scm
```

**--manifest=***file*

**-m** *file* Create an environment for the packages contained in the manifest object returned by the Scheme code in *file*. This option can be repeated several times, in which case the manifests are concatenated.

This is similar to the same-named option in `guix package` (see [profile-manifest], page 41) and uses the same manifest files.

See Section 8.4 [Writing Manifests], page 117, for information on how to write a manifest. See **--export-manifest** below on how to obtain a first manifest.

**--export-manifest**

Write to standard output a manifest suitable for **--manifest** corresponding to given command-line options.

This is a way to “convert” command-line arguments into a manifest. For example, imagine you are tired of typing long lines and would like to get a manifest equivalent to this command line:

```
guix shell -D guile git emacs emacs-geiser emacs-geiser-guile
```

Just add **--export-manifest** to the command line above:

```
guix shell --export-manifest \
  -D guile git emacs emacs-geiser emacs-geiser-guile
```

... and you get a manifest along these lines:

```
(concatenate-manifests
  (list (specifications->manifest
        (list "git"
              "emacs"
              "emacs-geiser"
              "emacs-geiser-guile")))
  (package->development-manifest
    (specification->package "guile"))))
```

You can store it into a file, say `manifest.scm`, and from there pass it to `guix shell` or indeed pretty much any `guix` command:

```
guix shell -m manifest.scm
```

Voilà, you’ve converted a long command line into a manifest! That conversion process honors package transformation options (see Section 9.1.2 [Package Transformation Options], page 182) so it should be lossless.

`--profile=profile`

`-p profile`

Create an environment containing the packages installed in *profile*. Use `guix package` (see Section 5.2 [Invoking guix package], page 37) to create and manage profiles.

`--pure`

Unset existing environment variables when building the new environment, except those specified with `--preserve` (see below). This has the effect of creating an environment in which search paths only contain package inputs.

`--preserve=regex`

`-E regex` When used alongside `--pure`, preserve the environment variables matching *regex*—in other words, put them on a “white list” of environment variables that must be preserved. This option can be repeated several times.

```
guix shell --pure --preserve=~SLURM openmpi ... \
-- mpirun ...
```

This example runs `mpirun` in a context where the only environment variables defined are `PATH`, environment variables whose name starts with ‘`SLURM`’, as well as the usual “precious” variables (`HOME`, `USER`, etc.).

`--search-paths`

Display the environment variable definitions that make up the environment.

`--system=system`

`-s system` Attempt to build for *system*—e.g., `i686-linux`.

`--container`

`-C` Run *command* within an isolated container. The current working directory outside the container is mapped inside the container. Additionally, unless overridden with `--user`, a dummy home directory is created that matches the current user’s home directory, and `/etc/passwd` is configured accordingly.

The spawned process runs as the current user outside the container. Inside the container, it has the same UID and GID as the current user, unless `--user` is passed (see below).

`--network`

`-N` For containers, share the network namespace with the host system. Containers created without this flag only have access to the loopback device.

`--link-profile`

`-P` For containers, link the environment profile to `~/.guix-profile` within the container and set `GUIX_ENVIRONMENT` to that. This is equivalent to making `~/.guix-profile` a symlink to the actual profile within the container. Linking

will fail and abort the environment if the directory already exists, which will certainly be the case if `guix shell` was invoked in the user's home directory. Certain packages are configured to look in `~/.guix-profile` for configuration files and data;<sup>2</sup> `--link-profile` allows these programs to behave as expected within the environment.

`--user=user`

`-u user` For containers, use the username *user* in place of the current user. The generated `/etc/passwd` entry within the container will contain the name *user*, the home directory will be `/home/user`, and no user GECOS data will be copied. Furthermore, the UID and GID inside the container are 1000. *user* need not exist on the system.

Additionally, any shared or exposed path (see `--share` and `--expose` respectively) whose target is within the current user's home directory will be remapped relative to `/home/USER`; this includes the automatic mapping of the current working directory.

```
# will expose paths as /home/foo/wd, /home/foo/test, and /home/foo/target
cd $HOME/wd
guix shell --container --user=foo \
  --expose=$HOME/test \
  --expose=/tmp/target=$HOME/target
```

While this will limit the leaking of user identity through home paths and each of the user fields, this is only one useful component of a broader privacy/anonymity solution—not one in and of itself.

`--no-cwd` For containers, the default behavior is to share the current working directory with the isolated container and immediately change to that directory within the container. If this is undesirable, `--no-cwd` will cause the current working directory to *not* be automatically shared and will change to the user's home directory within the container instead. See also `--user`.

`--writable-root`

When using `--container`, this option makes the root file system writable (it is read-only by default).

`--expose=source[=target]`

`--share=source[=target]`

For containers, `--expose` (resp. `--share`) exposes the file system *source* from the host system as the read-only (resp. writable) file system *target* within the container. If *target* is not specified, *source* is used as the target mount point in the container.

The example below spawns a Guile REPL in a container in which the user's home directory is accessible read-only via the `/exchange` directory:

```
guix shell --container --expose=$HOME=/exchange guile -- guile
```

`--symlink=spec`

`-S spec` For containers, create the symbolic links specified by *spec*, as documented in `[pack-symlink-option]`, page 97.

<sup>2</sup> For example, the `fontconfig` package inspects `~/.guix-profile/share/fonts` for additional fonts.

**--emulate-fhs**

**-F** When used with **--container**, emulate a Filesystem Hierarchy Standard (FHS) (<https://refspecs.linuxfoundation.org/fhs.shtml>) configuration within the container, providing **/bin**, **/lib**, and other directories and files specified by the FHS.

As Guix deviates from the FHS specification, this option sets up the container to more closely mimic that of other GNU/Linux distributions. This is useful for reproducing other development environments, testing, and using programs which expect the FHS specification to be followed. With this option, the container will include a version of glibc that will read **/etc/ld.so.cache** within the container for the shared library cache (contrary to glibc in regular Guix usage) and set up the expected FHS directories: **/bin**, **/etc**, **/lib**, and **/usr** from the container's profile.

**--nesting**

**-W** When used with **--container**, provide Guix *inside* the container and arrange so that it can interact with the build daemon that runs outside the container. This is useful if you want, within your isolated container, to create other containers, as in this sample session:

```
$ guix shell -CW coreutils
[env]$ guix shell -C guile -- guile -c '(display "hello!\n")'
hello!
[env]$ exit
```

The session above starts a container with **coreutils** programs available in **PATH**. From there, we spawn **guix shell** to create a *nested* container that provides nothing but Guile.

Another example is evaluating a **guix.scm** file that is untrusted, as shown here:

```
guix shell -CW -- guix build -f guix.scm
```

The **guix build** command as executed above can only access the current directory.

Under the hood, the **-W** option does several things:

- map the daemon's socket (by default **/var/guix/daemon-socket/socket**) inside the container;
- map the whole store (by default **/gnu/store**) inside the container such that store items made available by nested **guix** invocations are visible;
- add the currently-used **guix** command to the profile in the container, such that **guix describe** returns the same state inside and outside the container;
- share the cache (by default **~/.cache/guix**) with the host, to speed up operations such as **guix time-machine** and **guix shell**.

**--rebuild-cache**

In most cases, **guix shell** caches the environment so that subsequent uses are instantaneous. Least-recently used cache entries are periodically removed. The cache is also invalidated, when using **--file** or **--manifest**, anytime the corresponding file is modified.



The `--rebuild-cache` forces the cached environment to be refreshed. This is useful when using `--file` or `--manifest` and the `guix.scm` or `manifest.scm` file has external dependencies, or if its behavior depends, say, on environment variables.

`--root=file`

`-r file` Make *file* a symlink to the profile for this environment, and register it as a garbage collector root.

This is useful if you want to protect your environment from garbage collection, to make it “persistent”.

When this option is omitted, `guix shell` caches profiles so that subsequent uses of the same environment are instantaneous—this is comparable to using `--root` except that `guix shell` takes care of periodically removing the least-recently used garbage collector roots.

In some cases, `guix shell` does not cache profiles—e.g., if transformation options such as `--with-latest` are used. In those cases, the environment is protected from garbage collection only for the duration of the `guix shell` session. This means that next time you recreate the same environment, you could have to rebuild or re-download packages.

See Section 5.6 [Invoking `guix gc`], page 53, for more on GC roots.

`guix shell` also supports all of the common build options that `guix build` supports (see Section 9.1.1 [Common Build Options], page 180) as well as package transformation options (see Section 9.1.2 [Package Transformation Options], page 182).

## 7.2 Invoking `guix environment`

The purpose of `guix environment` is to assist in creating development environments.

**Deprecation warning:** The `guix environment` command is deprecated in favor of `guix shell`, which performs similar functions but is more convenient to use. See Section 7.1 [Invoking `guix shell`], page 78.

Being deprecated, `guix environment` is slated for eventual removal, but the Guix project is committed to keeping it until May 1st, 2023. Please get in touch with us at [guix-devel@gnu.org](mailto:guix-devel@gnu.org) if you would like to discuss it.

The general syntax is:

```
guix environment options package...
```

The following example spawns a new shell set up for the development of GNU Guile:

```
guix environment guile
```

If the needed dependencies are not built yet, `guix environment` automatically builds them. The environment of the new shell is an augmented version of the environment that `guix environment` was run in. It contains the necessary search paths for building the given package added to the existing environment variables. To create a “pure” environment, in which the original environment variables have been unset, use the `--pure` option<sup>3</sup>.

<sup>3</sup> Users sometimes wrongfully augment environment variables such as `PATH` in their `~/.bashrc` file. As a consequence, when `guix environment` launches it, Bash may read `~/.bashrc`, thereby introducing

Exiting from a Guix environment is the same as exiting from the shell, and will place the user back in the old environment before `guix environment` was invoked. The next garbage collection (see Section 5.6 [Invoking guix gc], page 53) will clean up packages that were installed from within the environment and are no longer used outside of it.

`guix environment` defines the `GUIX_ENVIRONMENT` variable in the shell it spawns; its value is the file name of the profile of this environment. This allows users to, say, define a specific prompt for development environments in their `.bashrc` (see Section “Bash Startup Files” in *The GNU Bash Reference Manual*):

```
if [ -n "$GUIX_ENVIRONMENT" ]
then
  export PS1="\u@\h \w [dev]\$ "
fi
```

... or to browse the profile:

```
$ ls "$GUIX_ENVIRONMENT/bin"
```

Additionally, more than one package may be specified, in which case the union of the inputs for the given packages are used. For example, the command below spawns a shell where all of the dependencies of both Guile and Emacs are available:

```
guix environment guile emacs
```

Sometimes an interactive shell session is not desired. An arbitrary command may be invoked by placing the `--` token to separate the command from the rest of the arguments:

```
guix environment guile -- make -j4
```

In other situations, it is more convenient to specify the list of packages needed in the environment. For example, the following command runs `python` from an environment containing Python 3 and NumPy:

```
guix environment --ad-hoc python-numpy python -- python3
```

Furthermore, one might want the dependencies of a package and also some additional packages that are not build-time or runtime dependencies, but are useful when developing nonetheless. Because of this, the `--ad-hoc` flag is positional. Packages appearing before `--ad-hoc` are interpreted as packages whose dependencies will be added to the environment. Packages appearing after are interpreted as packages that will be added to the environment directly. For example, the following command creates a Guix development environment that additionally includes Git and strace:

```
guix environment --pure guix --ad-hoc git strace
```

Sometimes it is desirable to isolate the environment as much as possible, for maximal purity and reproducibility. In particular, when using Guix on a host distro that is not Guix System, it is desirable to prevent access to `/usr/bin` and other system-wide resources from the development environment. For example, the following command spawns a Guile REPL in a “container” where only the store and the current working directory are mounted:

```
guix environment --ad-hoc --container guile -- guile
```

**Note:** The `--container` option requires Linux-libre 3.19 or newer.

---

“impurities” in these environment variables. It is an error to define such environment variables in `.bashrc`; instead, they should be defined in `.bash_profile`, which is sourced only by log-in shells. See Section “Bash Startup Files” in *The GNU Bash Reference Manual*, for details on Bash start-up files.

Another typical use case for containers is to run security-sensitive applications such as a web browser. To run Eolie, we must expose and share some files and directories; we include `nss-certs` and expose `/etc/ssl/certs/` for HTTPS authentication; finally we preserve the `DISPLAY` environment variable since containerized graphical applications won't display without it.

```
guix environment --preserve='^DISPLAY$' --container --network \
  --expose=/etc/machine-id \
  --expose=/etc/ssl/certs/ \
  --share=$HOME/.local/share/eolie/=$HOME/.local/share/eolie/ \
  --ad-hoc eolie nss-certs dbus -- eolie
```

The available options are summarized below.

**--check** Set up the environment and check whether the shell would clobber environment variables. See Section 7.1 [Invoking guix shell], page 78, for more info.

**--root=file**

**-r file** Make *file* a symlink to the profile for this environment, and register it as a garbage collector root.

This is useful if you want to protect your environment from garbage collection, to make it “persistent”.

When this option is omitted, the environment is protected from garbage collection only for the duration of the `guix environment` session. This means that next time you recreate the same environment, you could have to rebuild or re-download packages. See Section 5.6 [Invoking guix gc], page 53, for more on GC roots.

**--expression=expr**

**-e expr** Create an environment for the package or list of packages that *expr* evaluates to.

For example, running:

```
guix environment -e '(@ (gnu packages maths) petsc-openmpi)'
```

starts a shell with the environment for this specific variant of the PETSc package.

Running:

```
guix environment --ad-hoc -e '(@ (gnu) %base-packages)'
```

starts a shell with all the base system packages available.

The above commands only use the default output of the given packages. To select other outputs, two element tuples can be specified:

```
guix environment --ad-hoc -e '(list (@ (gnu packages bash) bash) "include")'
```

**--load=file**

**-l file** Create an environment for the package or list of packages that the code within *file* evaluates to.

As an example, *file* might contain a definition like this (see Section 8.2 [Defining Packages], page 101):

```
(use-modules (guix)
```

```
(gnu packages gdb)
(gnu packages autotools)
(gnu packages texinfo))
```

```
;; Augment the package definition of GDB with the build tools
;; needed when developing GDB (and which are not needed when
;; simply installing it.)
(package
  (inherit gdb)
  (native-inputs (modify-inputs (package-native-inputs gdb)
                                (prepend autoconf-2.69 automake texinfo))))
```

**--manifest=***file*

**-m** *file* Create an environment for the packages contained in the manifest object returned by the Scheme code in *file*. This option can be repeated several times, in which case the manifests are concatenated.

This is similar to the same-named option in `guix package` (see [profile-manifest], page 41) and uses the same manifest files.

See [shell-export-manifest], page 81, for information on how to “convert” command-line options into a manifest.

**--ad-hoc** Include all specified packages in the resulting environment, as if an *ad hoc* package were defined with them as inputs. This option is useful for quickly creating an environment without having to write a package expression to contain the desired inputs.

For instance, the command:

```
guix environment --ad-hoc guile guile-sdl -- guile
```

runs `guile` in an environment where Guile and Guile-SDL are available.

Note that this example implicitly asks for the default output of `guile` and `guile-sdl`, but it is possible to ask for a specific output—e.g., `glib:bin` asks for the `bin` output of `glib` (see Section 5.4 [Packages with Multiple Outputs], page 51).

This option may be composed with the default behavior of `guix environment`. Packages appearing before **--ad-hoc** are interpreted as packages whose dependencies will be added to the environment, the default behavior. Packages appearing after are interpreted as packages that will be added to the environment directly.

**--profile=***profile*

**-p** *profile*

Create an environment containing the packages installed in *profile*. Use `guix package` (see Section 5.2 [Invoking guix package], page 37) to create and manage profiles.

**--pure**

Unset existing environment variables when building the new environment, except those specified with **--preserve** (see below). This has the effect of creating an environment in which search paths only contain package inputs.

**--preserve=regex**

**-E regex** When used alongside **--pure**, preserve the environment variables matching *regex*—in other words, put them on a “white list” of environment variables that must be preserved. This option can be repeated several times.

```
guix environment --pure --preserve=~SLURM --ad-hoc openmpi ... \
-- mpirun ...
```

This example runs **mpirun** in a context where the only environment variables defined are **PATH**, environment variables whose name starts with ‘**SLURM**’, as well as the usual “precious” variables (**HOME**, **USER**, etc.).

**--search-paths**

Display the environment variable definitions that make up the environment.

**--system=system**

**-s system** Attempt to build for *system*—e.g., **i686-linux**.

**--container**

**-C** Run *command* within an isolated container. The current working directory outside the container is mapped inside the container. Additionally, unless overridden with **--user**, a dummy home directory is created that matches the current user’s home directory, and **/etc/passwd** is configured accordingly.

The spawned process runs as the current user outside the container. Inside the container, it has the same UID and GID as the current user, unless **--user** is passed (see below).

**--network**

**-N** For containers, share the network namespace with the host system. Containers created without this flag only have access to the loopback device.

**--link-profile**

**-P** For containers, link the environment profile to **~/.guix-profile** within the container and set **GUIX\_ENVIRONMENT** to that. This is equivalent to making **~/.guix-profile** a symlink to the actual profile within the container. Linking will fail and abort the environment if the directory already exists, which will certainly be the case if **guix environment** was invoked in the user’s home directory.

Certain packages are configured to look in **~/.guix-profile** for configuration files and data;<sup>4</sup> **--link-profile** allows these programs to behave as expected within the environment.

**--user=user**

**-u user** For containers, use the username *user* in place of the current user. The generated **/etc/passwd** entry within the container will contain the name *user*, the home directory will be **/home/user**, and no user GECOS data will be copied. Furthermore, the UID and GID inside the container are 1000. *user* need not exist on the system.

Additionally, any shared or exposed path (see **--share** and **--expose** respectively) whose target is within the current user’s home directory will be remapped

<sup>4</sup> For example, the **fontconfig** package inspects **~/.guix-profile/share/fonts** for additional fonts.

relative to `/home/USER`; this includes the automatic mapping of the current working directory.

```
# will expose paths as /home/foo/wd, /home/foo/test, and /home/foo/target
cd $HOME/wd
guix environment --container --user=foo \
  --expose=$HOME/test \
  --expose=/tmp/target=$HOME/target
```

While this will limit the leaking of user identity through home paths and each of the user fields, this is only one useful component of a broader privacy/anonymity solution—not one in and of itself.

**--no-cwd** For containers, the default behavior is to share the current working directory with the isolated container and immediately change to that directory within the container. If this is undesirable, **--no-cwd** will cause the current working directory to *not* be automatically shared and will change to the user’s home directory within the container instead. See also **--user**.

**--expose=source[=target]**

**--share=source[=target]**

For containers, **--expose** (resp. **--share**) exposes the file system *source* from the host system as the read-only (resp. writable) file system *target* within the container. If *target* is not specified, *source* is used as the target mount point in the container.

The example below spawns a Guile REPL in a container in which the user’s home directory is accessible read-only via the `/exchange` directory:

```
guix environment --container --expose=$HOME=/exchange --ad-hoc guile -- guil
```

**--emulate-fhs**

**-F** For containers, emulate a Filesystem Hierarchy Standard (FHS) configuration within the container, see the official specification (<https://refspecs.linuxfoundation.org/fhs.shtml>). As Guix deviates from the FHS specification, this option sets up the container to more closely mimic that of other GNU/Linux distributions. This is useful for reproducing other development environments, testing, and using programs which expect the FHS specification to be followed. With this option, the container will include a version of `glibc` which will read `/etc/ld.so.cache` within the container for the shared library cache (contrary to `glibc` in regular Guix usage) and set up the expected FHS directories: `/bin`, `/etc`, `/lib`, and `/usr` from the container’s profile.

`guix environment` also supports all of the common build options that `guix build` supports (see Section 9.1.1 [Common Build Options], page 180) as well as package transformation options (see Section 9.1.2 [Package Transformation Options], page 182).

## 7.3 Invoking `guix pack`

Occasionally you want to pass software to people who are not (yet!) lucky enough to be using Guix. You’d tell them to run `guix package -i something`, but that’s not possible in this case. This is where `guix pack` comes in.

**Note:** If you are looking for ways to exchange binaries among machines that already run Guix, see Section 9.13 [Invoking guix copy], page 233, Section 9.11 [Invoking guix publish], page 227, and Section 5.11 [Invoking guix archive], page 65.

The `guix pack` command creates a shrink-wrapped *pack* or *software bundle*: it creates a tarball or some other archive containing the binaries of the software you’re interested in, and all its dependencies. The resulting archive can be used on any machine that does not have Guix, and people can run the exact same binaries as those you have with Guix. The pack itself is created in a bit-reproducible fashion, so anyone can verify that it really contains the build results that you pretend to be shipping.

For example, to create a bundle containing Guile, Emacs, Geiser, and all their dependencies, you can run:

```
$ guix pack guile emacs emacs-geiser
...
/gnu/store/...-pack.tar.gz
```

The result here is a tarball containing a `/gnu/store` directory with all the relevant packages. The resulting tarball contains a *profile* with the three packages of interest; the profile is the same as would be created by `guix package -i`. It is this mechanism that is used to create Guix’s own standalone binary tarball (see Section 2.1 [Binary Installation], page 4).

Users of this pack would have to run `/gnu/store/...-profile/bin/guile` to run Guile, which you may find inconvenient. To work around it, you can create, say, a `/opt/gnu/bin` symlink to the profile:

```
guix pack -S /opt/gnu/bin=bin guile emacs emacs-geiser
```

That way, users can happily type `/opt/gnu/bin/guile` and enjoy.

What if the recipient of your pack does not have root privileges on their machine, and thus cannot unpack it in the root file system? In that case, you will want to use the `--relocatable` option (see below). This option produces *relocatable binaries*, meaning they can be placed anywhere in the file system hierarchy: in the example above, users can unpack your tarball in their home directory and directly run `./opt/gnu/bin/guile`.

Alternatively, you can produce a pack in the Docker image format using the following command:

```
guix pack -f docker -S /bin=bin guile guile-readline
```

The result is a tarball that can be passed to the `docker load` command, followed by `docker run`:

```
docker load < file
docker run -ti guile-guile-readline /bin/guile
```

where *file* is the image returned by `guix pack`, and `guile-guile-readline` is its “image tag”. See the Docker documentation (<https://docs.docker.com/engine/reference/commandline/load/>) for more information.

Yet another option is to produce a SquashFS image with the following command:

```
guix pack -f squashfs bash guile emacs emacs-geiser
```

The result is a SquashFS file system image that can either be mounted or directly be used as a file system container image with the Singularity container execution

environment (<https://www.sylabs.io/docs/>), using commands like `singularity shell` or `singularity exec`.

Another format internally based on SquashFS is AppImage (<https://appimage.org/>). An AppImage file can be created and executed without any special privileges:

```
file=$(guix pack -f appimage --entry-point=bin/guile guile)
$file --help
```

Several command-line options allow you to customize your pack:

`--format=format`

`-f format` Produce a pack in the given *format*.

The available formats are:

**tarball** This is the default format. It produces a tarball containing all the specified binaries and symlinks.

**docker** This produces a tarball that follows the Docker Image Specification (<https://github.com/docker/docker/blob/master/image/spec/v1.2.md>). By default, the “repository name” as it appears in the output of the `docker images` command is computed from package names passed on the command line or in the manifest file. Alternatively, the “repository name” can also be configured via the `--image-tag` option. Refer to `--help-docker-format` for more information on such advanced options.

**squashfs** This produces a SquashFS image containing all the specified binaries and symlinks, as well as empty mount points for virtual file systems like procs.

**Note:** Singularity *requires* you to provide `/bin/sh` in the image. For that reason, `guix pack -f squashfs` always implies `-S /bin=bin`. Thus, your `guix pack` invocation must always start with something like:

```
guix pack -f squashfs bash ...
```

If you forget the `bash` (or similar) package, `singularity run` and `singularity exec` will fail with an unhelpful “no such file or directory” message.

**deb** This produces a Debian archive (a package with the ‘`.deb`’ file extension) containing all the specified binaries and symbolic links, that can be installed on top of any dpkg-based GNU/Linux distribution. Advanced options can be revealed via the `--help-deb-format` option. They allow embedding control files for more fine-grained control, such as activating specific triggers or providing a maintainer configure script to run arbitrary setup code upon installation.

```
guix pack -f deb -C xz -S /usr/bin/hello=bin/hello hello
```

**Note:** Because archives produced with `guix pack` contain a collection of store items and because each `dpkg` package must not have conflicting files, in practice that



means you likely won't be able to install more than one such archive on a given system. You can nonetheless pack as many Guix packages as you want in one such archive.

**Warning:** `dpkg` will assume ownership of any files contained in the pack that it does *not* know about. It is unwise to install Guix-produced `.deb` files on a system where `/gnu/store` is shared by other software, such as a Guix installation or other, non-deb packs.

**rpm** This produces an RPM archive (a package with the `.rpm` file extension) containing all the specified binaries and symbolic links, that can be installed on top of any RPM-based GNU/Linux distribution. The RPM format embeds checksums for every file it contains, which the `rpm` command uses to validate the integrity of the archive.

Advanced RPM-related options are revealed via the `--help-rpm-format` option. These options allow embedding maintainer scripts that can run before or after the installation of the RPM archive, for example.

The RPM format supports relocatable packages via the `--prefix` option of the `rpm` command, which can be handy to install an RPM package to a specific prefix.

```
guix pack -f rpm -R -C xz -S /usr/bin/hello=bin/hello hello
sudo rpm --install --prefix=/opt /gnu/store/...-hello.rpm
```

**Note:** Contrary to Debian packages, conflicting but *identical* files in RPM packages can be installed simultaneously, which means multiple `guix pack`-produced RPM packages can usually be installed side by side without any problem.

**Warning:** `rpm` assumes ownership of any files contained in the pack, which means it will remove `/gnu/store` upon uninstalling a Guix-generated RPM package, unless the RPM package was installed with the `--prefix` option of the `rpm` command. It is unwise to install Guix-produced `.rpm` packages on a system where `/gnu/store` is shared by other software, such as a Guix installation or other, non-rpm packs.

**appimage** This produces an AppImage file (<https://appimage.org/>) with the `.AppImage` extension. AppImage is a SquashFS volume prefixed with a runtime that mounts the SquashFS file system and executes the binary provided with `--entry-point`. This results in a self-contained archive that bundles the software and all its requirements into a single file. When the file is made executable it runs the packaged software.

```
guix pack -f appimage --entry-point=bin/vlc vlc
```

The runtime used by AppImages invokes the `fusermount3` command to mount the image quickly. If that command is unavailable, the AppImage fails to run, but it can still be started by passing the `--appimage-extract-and-run` flag.

**Warning:** When building an AppImage, always *pass* the `--relocatable` option (or `-R`, or `-RR`) to make sure the image can be used on systems where Guix is not installed. A warning is printed when this option is not used.

```
guix pack -f appimage --entry-point=bin/hello --relocatable hello
```

**Note:** The resulting AppImage does not conform to the complete standard as it currently does not contain a `.DirIcon` file. This does not impact functionality of the AppImage itself, but possibly that of software used to manage AppImages.

**Note:** As the generated AppImage packages the complete dependency graph, it will be larger than comparable AppImage files found online, which depend on host system libraries.

#### `--relocatable`

`-R` Produce *relocatable binaries*—i.e., binaries that can be placed anywhere in the file system hierarchy and run from there.

When this option is passed once, the resulting binaries require support for *user namespaces* in the kernel Linux; when passed *twice*<sup>5</sup>, relocatable binaries fall to back to other techniques if user namespaces are unavailable, and essentially work anywhere—see below for the implications.

For example, if you create a pack containing Bash with:

```
guix pack -RR -S /mybin=bin bash
```

... you can copy that pack to a machine that lacks Guix, and from your home directory as a normal user, run:

```
tar xf pack.tar.gz
./mybin/sh
```

In that shell, if you type `ls /gnu/store`, you'll notice that `/gnu/store` shows up and contains all the dependencies of `bash`, even though the machine actually lacks `/gnu/store` altogether! That is probably the simplest way to deploy Guix-built software on a non-Guix machine.

**Note:** By default, relocatable binaries rely on the *user namespace* feature of the kernel Linux, which allows unprivileged users to mount or change root. Old versions of Linux did not support it, and some GNU/Linux distributions turn it off.

To produce relocatable binaries that work even in the absence of user namespaces, pass `--relocatable` or `-R` *twice*. In that case,

<sup>5</sup> Here's a trick to memorize it: `-RR`, which adds PRoot support, can be thought of as the abbreviation of "Really Relocatable". Neat, isn't it?

binaries will try user namespace support and fall back to another *execution engine* if user namespaces are not supported. The following execution engines are supported:

**default** Try user namespaces and fall back to PRoot if user namespaces are not supported (see below).

**performance** Try user namespaces and fall back to Fakechroot if user namespaces are not supported (see below).

**usersns** Run the program through user namespaces and abort if they are not supported.

**proot** Run through PRoot. The PRoot (<https://proot-me.github.io/>) program provides the necessary support for file system virtualization. It achieves that by using the `ptrace` system call on the running program. This approach has the advantage to work without requiring special kernel support, but it incurs run-time overhead every time a system call is made.

**fakechroot** Run through Fakechroot. Fakechroot (<https://github.com/dex4er/fakechroot/>) virtualizes file system accesses by intercepting calls to C library functions such as `open`, `stat`, `exec`, and so on. Unlike PRoot, it incurs very little overhead. However, it does not always work: for example, some file system accesses made from within the C library are not intercepted, and file system accesses made *via* direct syscalls are not intercepted either, leading to erratic behavior.

When running a wrapped program, you can explicitly request one of the execution engines listed above by setting the `GUIX_EXECUTION_ENGINE` environment variable accordingly.

**--entry-point=command**

Use *command* as the *entry point* of the resulting pack, if the pack format supports it—currently `docker`, `appimage`, and `squashfs` (Singularity) support it. *command* must be relative to the profile contained in the pack.

The entry point specifies the command that tools like `docker run` or `singularity run` automatically start by default. For example, you can do:

```
guix pack -f docker --entry-point=bin/guile guile
```

The resulting pack can easily be loaded and `docker run` with no extra arguments will spawn `bin/guile`:

```
docker load -i pack.tar.gz
docker run image-id
```

**--entry-point-argument=command**

**-A command**

Use *command* as an argument to *entry point* of the resulting pack. This option is only valid in conjunction with **--entry-point** and can appear multiple times on the command line.

```
guix pack -f docker --entry-point=bin/guile --entry-point-argument="--help"
```

**--max-layers=n**

Specifies the maximum number of Docker image layers allowed when building an image.

```
guix pack -f docker --max-layers=100 guile
```

This option allows you to limit the number of layers in a Docker image. Docker images are comprised of multiple layers, and each layer adds to the overall size and complexity of the image. By setting a maximum number of layers, you can control the following effects:

- **Disk Usage:** Increasing the number of layers can help optimize the disk space required to store multiple images built with a similar package graph.
- **Pulling:** When transferring images between different nodes or systems, having more layers can reduce the time required to pull the image.

**--expression=expr**

**-e expr** Consider the package *expr* evaluates to.

This has the same purpose as the same-named option in **guix build** (see Section 9.1.3 [Additional Build Options], page 189).

**--file=file**

Build a pack containing the package or other object the code within *file* evaluates to.

This has the same purpose as the same-named option in **guix build** (see Section 9.1.3 [Additional Build Options], page 189), but it has no shorthand, because **-f** already means **--format**.

**--manifest=file**

**-m file** Use the packages contained in the manifest object returned by the Scheme code in *file*. This option can be repeated several times, in which case the manifests are concatenated.

This has a similar purpose as the same-named option in **guix package** (see [profile-manifest], page 41) and uses the same manifest files. It allows you to define a collection of packages once and use it both for creating profiles and for creating archives for use on machines that do not have Guix installed. Note that you can specify *either* a manifest file *or* a list of packages, but not both.

See Section 8.4 [Writing Manifests], page 117, for information on how to write a manifest. See [shell-export-manifest], page 81, for information on how to “convert” command-line options into a manifest.

**--system=system**

**-s system** Attempt to build for *system*—e.g., **i686-linux**—instead of the system type of the build host.

**--target=triplet**  
 Cross-build for *triplet*, which must be a valid GNU triplet, such as "aarch64-linux-gnu" (see Section "Specifying target triplets" in *Autoconf*).

**--compression=tool**  
**-C tool** Compress the resulting tarball using *tool*—one of **gzip**, **zstd**, **bzip2**, **xz**, **lzip**, or **none** for no compression.

**--symlink=spec**  
**-S spec** Add the symlinks specified by *spec* to the pack. This option can appear several times.  
*spec* has the form **source=target**, where *source* is the symlink that will be created and *target* is the symlink target.  
 For instance, **-S /opt/gnu/bin=bin** creates a **/opt/gnu/bin** symlink pointing to the **bin** sub-directory of the profile.

**--save-provenance**  
 Save provenance information for the packages passed on the command line. Provenance information includes the URL and commit of the channels in use (see Chapter 6 [Channels], page 68).  
 Provenance information is saved in the **/gnu/store/...-profile/manifest** file in the pack, along with the usual package metadata—the name and version of each package, their propagated inputs, and so on. It is useful information to the recipient of the pack, who then knows how the pack was (supposedly) obtained.  
 This option is not enabled by default because, like timestamps, provenance information contributes nothing to the build process. In other words, there is an infinity of channel URLs and commit IDs that can lead to the same pack. Recording such “silent” metadata in the output thus potentially breaks the source-to-binary bitwise reproducibility property.

**--root=file**  
**-r file** Make *file* a symlink to the resulting pack, and register it as a garbage collector root.

**--localstatedir**  
**--profile-name=name**  
 Include the “local state directory”, **/var/guix**, in the resulting pack, and notably the **/var/guix/profiles/per-user/root/name** profile—by default *name* is **guix-profile**, which corresponds to **~root/.guix-profile**.  
**/var/guix** contains the store database (see Section 8.9 [The Store], page 156) as well as garbage-collector roots (see Section 5.6 [Invoking guix gc], page 53). Providing it in the pack means that the store is “complete” and manageable by Guix; not providing it in the pack means that the store is “dead”: items cannot be added to it or removed from it after extraction of the pack.  
 One use case for this is the Guix self-contained binary tarball (see Section 2.1 [Binary Installation], page 4).

**--derivation**  
**-d** Print the name of the derivation that builds the pack.

**--bootstrap**

Use the bootstrap binaries to build the pack. This option is only useful to Guix developers.

In addition, **guix pack** supports all the common build options (see Section 9.1.1 [Common Build Options], page 180) and all the package transformation options (see Section 9.1.2 [Package Transformation Options], page 182).

## 7.4 The GCC toolchain

If you need a complete toolchain for compiling and linking C or C++ source code, use the **gcc-toolchain** package. This package provides a complete GCC toolchain for C/C++ development, including GCC itself, the GNU C Library (headers and binaries, plus debugging symbols in the **debug** output), Binutils, and a linker wrapper.

The wrapper's purpose is to inspect the **-L** and **-l** switches passed to the linker, add corresponding **-rpath** arguments, and invoke the actual linker with this new set of arguments. You can instruct the wrapper to refuse to link against libraries not in the store by setting the **GUIX\_LD\_WRAPPER\_ALLOW\_IMPURITIES** environment variable to **no**.

The package **gfortran-toolchain** provides a complete GCC toolchain for Fortran development. For other languages, please use '**guix search gcc toolchain**' (see [Invoking guix package], page 43).

## 7.5 Invoking guix git authenticate

The **guix git authenticate** command authenticates a Git checkout following the same rule as for channels (see [channel-authentication], page 71). That is, starting from a given commit, it ensures that all subsequent commits are signed by an OpenPGP key whose fingerprint appears in the **.guix-authorizations** file of its parent commit(s).

You will find this command useful if you maintain a channel. But in fact, this authentication mechanism is useful in a broader context, so you might want to use it for Git repositories that have nothing to do with Guix.

The general syntax is:

```
guix git authenticate commit signer [options...]
```

By default, this command authenticates the Git checkout in the current directory; it outputs nothing and exits with exit code zero on success and non-zero on failure. *commit* above denotes the first commit where authentication takes place, and *signer* is the OpenPGP fingerprint of public key used to sign *commit*. Together, they form a *channel introduction* (see [channel-authentication], page 71). On your first successful run, the introduction is recorded in the **.git/config** file of your checkout, allowing you to omit them from subsequent invocations:

```
guix git authenticate [options...]
```

Should you have branches that require different introductions, you can specify them directly in **.git/config**. For example, if the branch called **personal-fork** has a different introduction than other branches, you can extend **.git/config** along these lines:

```
[guix "authentication-personal-fork"]
introduction-commit = cabba936fd807b096b48283debdccddccfea3900d
introduction-signer = COFF EECA BBA9 E6A8 0D1D E643 A2A0 6DF2 A33A 54FA
```

```
keyring = keyring
```

The first run also attempts to install pre-push and post-merge hooks, such that `guix git authenticate` is invoked as soon as you run `git push`, `git pull`, and related commands; it does not overwrite preexisting hooks though.

The command-line options described below allow you to fine-tune the process.

```
--repository=directory
```

```
-r directory
```

Open the Git repository in *directory* instead of the current directory.

```
--keyring=reference
```

```
-k reference
```

Load OpenPGP keyring from *reference*, the reference of a branch such as `origin/keyring` or `my-keyring`. The branch must contain OpenPGP public keys in `.key` files, either in binary form or “ASCII-armored”. By default the keyring is loaded from the branch named `keyring`.

```
--end=commit
```

Authenticate revisions up to *commit*.

```
--stats
```

 Display commit signing statistics upon completion.

```
--cache-key=key
```

Previously-authenticated commits are cached in a file under `~/.cache/guix/authentication`. This option forces the cache to be stored in file *key* in that directory.

```
--historical-authorizations=file
```

By default, any commit whose parent commit(s) lack the `.guix-authorizations` file is considered inauthentic. In contrast, this option considers the authorizations in *file* for any commit that lacks `.guix-authorizations`. The format of *file* is the same as that of `.guix-authorizations` (see [channel-authorizations], page 74).

## 8 Programming Interface

GNU Guix provides several Scheme programming interfaces (APIs) to define, build, and query packages. The first interface allows users to write high-level package definitions. These definitions refer to familiar packaging concepts, such as the name and version of a package, its build system, and its dependencies. These definitions can then be turned into concrete build actions.

Build actions are performed by the Guix daemon, on behalf of users. In a standard setup, the daemon has write access to the store—the `/gnu/store` directory—whereas users do not. The recommended setup also has the daemon perform builds in chroots, under specific build users, to minimize interference with the rest of the system.

Lower-level APIs are available to interact with the daemon and the store. To instruct the daemon to perform a build action, users actually provide it with a *derivation*. A derivation is a low-level representation of the build actions to be taken, and the environment in which they should occur—derivations are to package definitions what assembly is to C programs. The term “derivation” comes from the fact that build results *derive* from them.

This chapter describes all these APIs in turn, starting from high-level package definitions. See Section 22.7 [Source Tree Structure], page 826, for a more general overview of the source code.

### 8.1 Package Modules

From a programming viewpoint, the package definitions of the GNU distribution are provided by Guile modules in the `(gnu packages ...)` name space<sup>1</sup> (see Section “Modules” in *GNU Guile Reference Manual*). For instance, the `(gnu packages emacs)` module exports a variable named `emacs`, which is bound to a `<package>` object (see Section 8.2 [Defining Packages], page 101).

The `(gnu packages ...)` module name space is automatically scanned for packages by the command-line tools. For instance, when running `guix install emacs`, all the `(gnu packages ...)` modules are scanned until one that exports a package object whose name is `emacs` is found. This package search facility is implemented in the `(gnu packages)` module.

Users can store package definitions in modules with different names—e.g., `(my-packages emacs)`<sup>2</sup>. There are two ways to make these package definitions visible to the user interfaces:

1. By adding the directory containing your package modules to the search path with the `-L` flag of `guix package` and other commands (see Section 9.1.1 [Common Build Options], page 180), or by setting the `GUIX_PACKAGE_PATH` environment variable described below.
2. By defining a *channel* and configuring `guix pull` so that it pulls from it. A channel is essentially a Git repository containing package modules. See Chapter 6 [Channels], page 68, for more information on how to define and use channels.

<sup>1</sup> Note that packages under the `(gnu packages ...)` module name space are not necessarily “GNU packages”. This module naming scheme follows the usual Guile module naming convention: `gnu` means that these modules are distributed as part of the GNU system, and `packages` identifies modules that define packages.

<sup>2</sup> Note that the file name and module name must match. For instance, the `(my-packages emacs)` module must be stored in a `my-packages/emacs.scm` file relative to the load path specified with `--load-path` or `GUIX_PACKAGE_PATH`. See Section “Modules and the File System” in *GNU Guile Reference Manual*, for details.



`GUIX_PACKAGE_PATH` works similarly to other search path variables:

`GUIX_PACKAGE_PATH` [Environment Variable]

This is a colon-separated list of directories to search for additional package modules. Directories listed in this variable take precedence over the own modules of the distribution.

The distribution is fully *bootstrapped* and *self-contained*: each package is built based solely on other packages in the distribution. The root of this dependency graph is a small set of *bootstrap binaries*, provided by the `(gnu packages bootstrap)` module. For more information on bootstrapping, see Chapter 20 [Bootstrapping], page 808.

## 8.2 Defining Packages

The high-level interface to package definitions is implemented in the `(guix packages)` and `(guix build-system)` modules. As an example, the package definition, or *recipe*, for the GNU Hello package looks like this:

```
(define-module (gnu packages hello)
  #:use-module (guix gexp)
  #:use-module (guix packages)
  #:use-module (guix download)
  #:use-module (guix build-system gnu)
  #:use-module (guix licenses)
  #:use-module (gnu packages gawk))

(define-public hello
  (package
    (name "hello")
    (version "2.10")
    (source (origin
      (method url-fetch)
      (uri (string-append "mirror://gnu/hello/hello-" version
        ".tar.gz"))
      (sha256
        (base32
          "0ssi1wpaf7plawqqjwigppsg5fyh99vdlb9kzl7c9lng89ndqi"))))
    (build-system gnu-build-system)
    (arguments
      (list #:configure-flags #~(list "--enable-silent-rules")))
    (inputs (list gawk))
    (synopsis "Hello, GNU world: An example GNU package")
    (description "Guess what GNU Hello prints!")
    (home-page "https://www.gnu.org/software/hello/")
    (license gpl3+)))
```

Without being a Scheme expert, the reader may have guessed the meaning of the various fields here. This expression binds the variable `hello` to a `<package>` record (see Section “SRFI-9” in *GNU Guile Reference Manual*). This package object can be inspected using

procedures found in the `(guix packages)` module; for instance, `(package-name hello)` returns—surprise!—"hello".

**Quick start:** Impatient? With luck, you may be able to import part or all of the definition of the package you are interested in from another repository, using the `guix import` command (see Section 9.5 [Invoking `guix import`], page 197).

Alternatively, try out the Guix-Packager interface (<https://guix-hpc.gitlabpages.inria.fr/guix-packager/>) in your web browser to produce a preliminary package definition.

In the example above, `hello` is defined in a module of its own, `(gnu packages hello)`. Technically, this is not strictly necessary, but it is convenient to do so: all the packages defined in modules under `(gnu packages ...)` are automatically known to the command-line tools (see Section 8.1 [Package Modules], page 100).

There are a few points worth noting in the above package definition:

- The `source` field of the package is an `<origin>` object (see Section 8.2.2 [origin Reference], page 109, for the complete reference). Here, the `url-fetch` method from `(guix download)` is used, meaning that the source is a file to be downloaded over FTP or HTTP.

The `mirror://gnu` prefix instructs `url-fetch` to use one of the GNU mirrors defined in `(guix download)`.

The `sha256` field specifies the expected SHA256 hash of the file being downloaded. It is mandatory, and allows Guix to check the integrity of the file. The `(base32 ...)` form introduces the base32 representation of the hash. You can obtain this information with `guix download` (see Section 9.3 [Invoking `guix download`], page 195) and `guix hash` (see Section 9.4 [Invoking `guix hash`], page 196).

When needed, the `origin` form can also have a `patches` field listing patches to be applied, and a `snippet` field giving a Scheme expression to modify the source code.

- The `build-system` field specifies the procedure to build the package (see Section 8.5 [Build Systems], page 121). Here, `gnu-build-system` represents the familiar GNU Build System, where packages may be configured, built, and installed with the usual `./configure && make && make check && make install` command sequence.

When you start packaging non-trivial software, you may need tools to manipulate those build phases, manipulate files, and so on. See Section 8.7 [Build Utilities], page 146, for more on this.

- The `arguments` field specifies options for the build system (see Section 8.5 [Build Systems], page 121). Here it is interpreted by `gnu-build-system` as a request run `configure` with the `--enable-silent-rules` flag.

The value of the `arguments` field is a list of arguments passed to the build system and that end up being applied, as with `apply` (see Section “Fly Evaluation” in *GNU Guile Reference Manual*).

The hash-colon (`#:`) sequence defines a Scheme *keyword* (see Section “Keywords” in *GNU Guile Reference Manual*), and `#:configure-flags` is a keyword used to pass a keyword argument to the build system (see Section “Coding With Keywords” in *GNU Guile Reference Manual*).

What about these hash-tilde (`#~`) characters? This is Guix-specific syntax introducing a so-called *G-expression* (see Section 8.12 [G-Expressions], page 165), which denotes code that is *staged* for eventual evaluation: this staged code is only evaluated if and when attempting to build the package, by the build daemon. In this example, the staged code is the `(list ...)` expression, which builds a list of one element, the configure flag.

- The `inputs` field specifies inputs to the build process—i.e., build-time or run-time dependencies of the package. Here, we add an input, a reference to the `gawk` variable; `gawk` is itself bound to a `<package>` object.

Note that GCC, Coreutils, Bash, and other essential tools do not need to be specified as inputs here. Instead, `gnu-build-system` takes care of ensuring that they are present (see Section 8.5 [Build Systems], page 121).

However, any other dependencies need to be specified in the `inputs` field. Any dependency not specified here will simply be unavailable to the build process, possibly leading to a build failure.

See Section 8.2.1 [package Reference], page 104, for a full description of possible fields.

**Going further:** Intimidated by the Scheme language or curious about it? The Cookbook has a short section to get started that recaps some of the things shown above and explains the fundamentals. See Section “A Scheme Crash Course” in *GNU Guix Cookbook*, for more information.

Once a package definition is in place, the package may actually be built using the `guix build` command-line tool (see Section 9.1 [Invoking guix build], page 180), troubleshooting any build failures you encounter (see Section 9.1.4 [Debugging Build Failures], page 193). You can easily jump back to the package definition using the `guix edit` command (see Section 9.2 [Invoking guix edit], page 195). See Section 22.8 [Packaging Guidelines], page 830, for more information on how to test package definitions, and Section 9.8 [Invoking guix lint], page 217, for information on how to check a definition for style conformance. Lastly, see Chapter 6 [Channels], page 68, for information on how to extend the distribution by adding your own package definitions in a “channel”.

Finally, updating the package definition to a new upstream version can be partly automated by the `guix refresh` command (see Section 9.6 [Invoking guix refresh], page 206).

Behind the scenes, a derivation corresponding to the `<package>` object is first computed by the `package-derivation` procedure. That derivation is stored in a `.drv` file under `/gnu/store`. The build actions it prescribes may then be realized by using the `build-derivations` procedure (see Section 8.9 [The Store], page 156).

`package-derivation store package [system]` [Procedure]

Return the `<derivation>` object of *package* for *system* (see Section 8.10 [Derivations], page 158).

*package* must be a valid `<package>` object, and *system* must be a string denoting the target system type—e.g., `"x86_64-linux"` for an x86\_64 Linux-based GNU system. *store* must be a connection to the daemon, which operates on the store (see Section 8.9 [The Store], page 156).

Similarly, it is possible to compute a derivation that cross-builds a package for some other system:

**package-cross-derivation** *store package target* [*system*] [Procedure]

Return the <derivation> object of *package* cross-built from *system* to *target*.

*target* must be a valid GNU triplet denoting the target hardware and operating system, such as "aarch64-linux-gnu" (see Section “Specifying Target Triplets” in *Autoconf*).

Once you have package definitions, you can easily define *variants* of those packages. See Section 8.3 [Defining Package Variants], page 113, for more on that.

## 8.2.1 package Reference

This section summarizes all the options available in **package** declarations (see Section 8.2 [Defining Packages], page 101).

**package** [Data Type]

This is the data type representing a package recipe.

**name** The name of the package, as a string.

**version** The version of the package, as a string. See Section 22.8.3 [Version Numbers], page 832, for guidelines.

**source** An object telling how the source code for the package should be acquired. Most of the time, this is an **origin** object, which denotes a file fetched from the Internet (see Section 8.2.2 [origin Reference], page 109). It can also be any other “file-like” object such as a **local-file**, which denotes a file from the local file system (see Section 8.12 [G-Expressions], page 165).

**build-system**

The build system that should be used to build the package (see Section 8.5 [Build Systems], page 121).

**arguments** (default: '())

The arguments that should be passed to the build system (see Section 8.5 [Build Systems], page 121). This is a list, typically containing sequential keyword-value pairs, as in this example:

```
(package
  (name "example")
  ;; several fields omitted
  (arguments
    (list #:tests? #f ;skip tests
          #:make-flags #~'("VERBOSE=1") ;pass flags to 'make'
          #:configure-flags #~'("--enable-frobbing"))))
```

The exact set of supported keywords depends on the build system (see Section 8.5 [Build Systems], page 121), but you will find that almost all of them honor **#:configure-flags**, **#:make-flags**, **#:tests?**, and **#:phases**. The **#:phases** keyword in particular lets you modify the set of build phases for your package (see Section 8.6 [Build Phases], page 143).

The REPL has dedicated commands to interactively inspect values of some of these arguments, as a convenient debugging aid (see Section 8.14 [Using Guix Interactively], page 176).

**Compatibility Note:** Until version 1.3.0, the `arguments` field would typically use `quote (')` or `quasiquote (`)` and no G-expressions, like so:

```
(package
  ;; several fields omitted
  (arguments ;old-style quoted arguments
    '(:tests? #f
      #:configure-flags '("--enable-frobbing"))))
```

To convert from that style to the one shown above, you can run `guix style -S arguments package` (see Section 9.7 [Invoking guix style], page 213).

`inputs` (default: '())

`native-inputs` (default: '())

`propagated-inputs` (default: '())

These fields list dependencies of the package. Each element of these lists is either a package, origin, or other “file-like object” (see Section 8.12 [G-Expressions], page 165); to specify the output of that file-like object that should be used, pass a two-element list where the second element is the output (see Section 5.4 [Packages with Multiple Outputs], page 51, for more on package outputs). For example, the list below specifies three inputs:

```
(list libffi libunistring
  `(:glib "bin")) ;the "bin" output of GLib
```

In the example above, the “out” output of `libffi` and `libunistring` is used.

**Compatibility Note:** Until version 1.3.0, input lists were a list of tuples, where each tuple has a label for the input (a string) as its first element, a package, origin, or derivation as its second element, and optionally the name of the output thereof that should be used, which defaults to “out”. For example, the list below is equivalent to the one above, but using the *old input style*:

```
;; Old input style (deprecated).
`(("libffi" ,libffi)
  ("libunistring" ,libunistring)
  ("glib:bin" ,glib "bin")) ;the "bin" output of GLib
```

This style is now deprecated; it is still supported but support will be removed in a future version. It should not be used for new package definitions. See Section 9.7 [Invoking guix style], page 213, on how to migrate to the new style.

The distinction between `native-inputs` and `inputs` is necessary when considering cross-compilation. When cross-compiling, dependencies listed in `inputs` are built for the *target* architecture; conversely, dependencies listed in `native-inputs` are built for the architecture of the *build* machine.

**native-inputs** is typically used to list tools needed at build time, but not at run time, such as Autoconf, Automake, pkg-config, Gettext, or Bison. **guix lint** can report likely mistakes in this area (see Section 9.8 [Invoking guix lint], page 217).

Lastly, **propagated-inputs** is similar to **inputs**, but the specified packages will be automatically installed to profiles (see Section 5.1 [Features], page 36) alongside the package they belong to (see [package-cmd-propagated-inputs], page 38, for information on how **guix package** deals with propagated inputs).

For example this is necessary when packaging a C/C++ library that needs headers of another library to compile, or when a pkg-config file refers to another one *via* its **Requires** field.

Another example where **propagated-inputs** is useful is for languages that lack a facility to record the run-time search path akin to the **RUNPATH** of ELF files; this includes Guile, Python, Perl, and more. When packaging libraries written in those languages, ensure they can find library code they depend on at run time by listing run-time dependencies in **propagated-inputs** rather than **inputs**.

**outputs** (default: `'("out")`)

The list of output names of the package. See Section 5.4 [Packages with Multiple Outputs], page 51, for typical uses of additional outputs.

**native-search-paths** (default: `'()`)

**search-paths** (default: `'()`)

A list of **search-path-specification** objects describing search-path environment variables honored by the package. See Section 8.8 [Search Paths], page 153, for more on search path specifications.

As for inputs, the distinction between **native-search-paths** and **search-paths** only matters when cross-compiling. In a cross-compilation context, **native-search-paths** applies exclusively to native inputs whereas **search-paths** applies only to host inputs.

Packages such as cross-compilers care about target inputs—for instance, our (modified) GCC cross-compiler has **CROSS\_C\_INCLUDE\_PATH** in **search-paths**, which allows it to pick `.h` files for the target system and *not* those of native inputs. For the majority of packages though, only **native-search-paths** makes sense.

**replacement** (default: `#f`)

This must be either `#f` or a package object that will be used as a *replacement* for this package. See Chapter 19 [Security Updates], page 806, for details.

**synopsis** A one-line description of the package.

**description**

A more elaborate description of the package, as a string in Texinfo syntax.

**license** The license of the package; a value from (**guix licenses**), or a list of such values.

**home-page**

The URL to the home-page of the package, as a string.

**supported-systems** (default: `%supported-systems`)

The list of systems supported by the package, as strings of the form `architecture-kernel`, for example `"x86_64-linux"`.

**location** (default: source location of the `package` form)

The source location of the package. It is useful to override this when inheriting from another package, in which case this field is not automatically corrected.

**this-package**

[Macro]

When used in the *lexical scope* of a package field definition, this identifier resolves to the package being defined.

The example below shows how to add a package as a native input of itself when cross-compiling:

```
(package
  (name "guile")
  ;; ...

  ;; When cross-compiled, Guile, for example, depends on
  ;; a native version of itself. Add it here.
  (native-inputs (if (%current-target-system)
                     (list this-package)
                     '()))))
```

It is an error to refer to `this-package` outside a package definition.

The following helper procedures are provided to help deal with package inputs.

**lookup-package-input** *package name* [Procedure]

**lookup-package-native-input** *package name* [Procedure]

**lookup-package-propagated-input** *package name* [Procedure]

**lookup-package-direct-input** *package name* [Procedure]

Look up *name* among *package*'s inputs (or native, propagated, or direct inputs). Return it if found, `#f` otherwise.

*name* is the name of a package or the file name of an origin depended on. Here's how you might use it:

```
(use-modules (guix packages) (gnu packages base))

(lookup-package-direct-input coreutils "gmp")
⇒ #<package gmp@6.2.1 ...>
```

In this example we obtain the `gmp` package that is among the direct inputs of `coreutils`.

When looking up an origin, use the name that appears in the origin's `file-name` field or its default file name—e.g., `"foo-1.2.tar.gz"`.

Sometimes you will want to obtain the list of inputs needed to *develop* a package—all the inputs that are visible when the package is compiled. This is what the `package-development-inputs` procedure returns.

`package-development-inputs` *package* [*system*] [#:*target* #*f*] [Procedure]

Return the list of inputs required by *package* for development purposes on *system*. When *target* is true, return the inputs needed to cross-compile *package* from *system* to *target*, where *target* is a triplet such as "aarch64-linux-gnu".

Note that the result includes both explicit inputs and implicit inputs—inputs automatically added by the build system (see Section 8.5 [Build Systems], page 121). Let us take the `hello` package to illustrate that:

```
(use-modules (gnu packages base) (guix packages))

hello
⇒ #<package hello@2.10 gnu/packages/base.scm:79 7f585d4f6790>

(package-direct-inputs hello)
⇒ ()

(package-development-inputs hello)
⇒ (("source" ...) ("tar" #<package tar@1.32 ...>) ...)
```

In this example, `package-direct-inputs` returns the empty list, because `hello` has zero explicit dependencies. Conversely, `package-development-inputs` includes inputs implicitly added by `gnu-build-system` that are required to build `hello`: `tar`, `gzip`, `GCC`, `libc`, `Bash`, and more. To visualize it, `guix graph hello` would show you explicit inputs, whereas `guix graph -t bag hello` would include implicit inputs (see Section 9.10 [Invoking guix graph], page 222).

Because packages are regular Scheme objects that capture a complete dependency graph and associated build procedures, it is often useful to write procedures that take a package and return a modified version thereof according to some parameters. Below are a few examples.

`package-with-c-toolchain` *package* *toolchain* [Procedure]

Return a variant of *package* that uses *toolchain* instead of the default GNU C/C++ toolchain. *toolchain* must be a list of inputs (label/package tuples) providing equivalent functionality, such as the `gcc-toolchain` package.

The example below returns a variant of the `hello` package built with GCC 10.x and the rest of the GNU tool chain (Binutils and the GNU C Library) instead of the default tool chain:

```
(let ((toolchain (specification->package "gcc-toolchain@10")))
  (package-with-c-toolchain hello `(("toolchain" ,toolchain))))
```

The build tool chain is part of the *implicit inputs* of packages—it’s usually not listed as part of the various “inputs” fields and is instead pulled in by the build system. Consequently, this procedure works by changing the build system of *package* so that it pulls in *toolchain* instead of the defaults. See Section 8.5 [Build Systems], page 121, for more on build systems.



### 8.2.2 origin Reference

This section documents *origins*. An **origin** declaration specifies data that must be “produced”—downloaded, usually—and whose content hash is known in advance. Origins are primarily used to represent the source code of packages (see Section 8.2 [Defining Packages], page 101). For that reason, the **origin** form allows you to declare patches to apply to the original source code as well as code snippets to modify it.

**origin** [Data Type]

This is the data type representing a source code origin.

**uri** An object containing the URI of the source. The object type depends on the **method** (see below). For example, when using the *url-fetch* method of (**guix download**), the valid **uri** values are: a URL represented as a string, or a list thereof.

**method** A monadic procedure that handles the given URI. The procedure must accept at least three arguments: the value of the **uri** field and the hash algorithm and hash value specified by the **hash** field. It must return a store item or a derivation in the store monad (see Section 8.11 [The Store Monad], page 161); most methods return a fixed-output derivation (see Section 8.10 [Derivations], page 158).

Commonly used methods include **url-fetch**, which fetches data from a URL, and **git-fetch**, which fetches data from a Git repository (see below).

**sha256** A bytevector containing the SHA-256 hash of the source. This is equivalent to providing a **content-hash** SHA256 object in the **hash** field described below.

**hash** The **content-hash** object of the source—see below for how to use **content-hash**.

You can obtain this information using **guix download** (see Section 9.3 [Invoking guix download], page 195) or **guix hash** (see Section 9.4 [Invoking guix hash], page 196).

**file-name** (default: **#f**)

The file name under which the source code should be saved. When this is **#f**, a sensible default value will be used in most cases. In case the source is fetched from a URL, the file name from the URL will be used. For version control checkouts, it is recommended to provide the file name explicitly because the default is not very descriptive.

**patches** (default: **'()**)

A list of file names, origins, or file-like objects (see Section 8.12 [G-Expressions], page 165) pointing to patches to be applied to the source.

This list of patches must be unconditional. In particular, it cannot depend on the value of **%current-system** or **%current-target-system**.

**snippet** (default: **#f**)

A G-expression (see Section 8.12 [G-Expressions], page 165) or S-expression that will be run in the source directory. This is a convenient way to modify the source, sometimes more convenient than a patch.

**patch-flags** (default: **'("-p1")**)

A list of command-line flags that should be passed to the **patch** command.

**patch-inputs** (default: **#f**)

Input packages or derivations to the patching process. When this is **#f**, the usual set of inputs necessary for patching are provided, such as GNU Patch.

**modules** (default: **'()**)

A list of Guile modules that should be loaded during the patching process and while running the code in the **snippet** field.

**patch-guile** (default: **#f**)

The Guile package that should be used in the patching process. When this is **#f**, a sensible default is used.

**content-hash** *value* [*algorithm*] [Data Type]

Construct a content hash object for the given *algorithm*, and with *value* as its hash value. When *algorithm* is omitted, assume it is **sha256**.

*value* can be a literal string, in which case it is base32-decoded, or it can be a bytevector.

The following forms are all equivalent:

```
(content-hash "05zxkyz9bv3j9h0xyid1rhvh3klhsmrpkf3bcs6frvlgyr2gwilj")■
(content-hash "05zxkyz9bv3j9h0xyid1rhvh3klhsmrpkf3bcs6frvlgyr2gwilj"■
  sha256)
(content-hash (base32
  "05zxkyz9bv3j9h0xyid1rhvh3klhsmrpkf3bcs6frvlgyr2gwilj"))■
(content-hash (base64 "kkb+RPaP7uyMZmu4eXPVkm4BN8yhRd8BTHLslb6f/Rc=")■
  sha256)
```

Technically, **content-hash** is currently implemented as a macro. It performs sanity checks at macro-expansion time, when possible, such as ensuring that *value* has the right size for *algorithm*.

As we have seen above, how exactly the data an origin refers to is retrieved is determined by its **method** field. The (**guix download**) module provides the most common method, **url-fetch**, described below.

**url-fetch** *url* *hash-algo* *hash* [*name*] [#:executable? *#f*] [Procedure]

Return a fixed-output derivation that fetches data from *url* (a string, or a list of strings denoting alternate URLs), which is expected to have hash *hash* of type *hash-algo* (a symbol). By default, the file name is the base name of URL; optionally, *name* can specify a different file name. When *executable?* is true, make the downloaded file executable.

When one of the URL starts with `mirror://`, then its host part is interpreted as the name of a mirror scheme, taken from `%mirror-file`.

Alternatively, when URL starts with `file://`, return the corresponding file name in the store.

Likewise, the `(guix git-download)` module defines the `git-fetch` origin method, which fetches data from a Git version control repository, and the `git-reference` data type to describe the repository and revision to fetch.

**git-fetch** *ref hash-algo hash [name]* [Procedure]

Return a fixed-output derivation that fetches *ref*, a `<git-reference>` object. The output is expected to have recursive hash *hash* of type *hash-algo* (a symbol). Use *name* as the file name, or a generic name if `#f`.

**git-fetch/lfs** *ref hash-algo hash [name]* [Procedure]

This is a variant of the `git-fetch` procedure that supports the Git LFS (Large File Storage) extension. This may be useful to pull some binary test data to run the test suite of a package, for example.

**git-reference** [Data Type]

This data type represents a Git reference for `git-fetch` to retrieve.

**url** The URL of the Git repository to clone.

**commit** This string denotes either the commit to fetch (a hexadecimal string), or the tag to fetch. You can also use a “short” commit ID or a `git describe` style identifier such as `v1.0.1-10-g58d7909c97`.

**recursive?** (default: `#f`)

This Boolean indicates whether to recursively fetch Git sub-modules.

The example below denotes the `v2.10` tag of the GNU Hello repository:

```
(git-reference
 (url "https://git.savannah.gnu.org/git/hello.git")
 (commit "v2.10"))
```

This is equivalent to the reference below, which explicitly names the commit:

```
(git-reference
 (url "https://git.savannah.gnu.org/git/hello.git")
 (commit "dc7dc56a00e48fe6f231a58f6537139fe2908fb9"))
```

For Mercurial repositories, the module `(guix hg-download)` defines the `hg-fetch` origin method and `hg-reference` data type for support of the Mercurial version control system.

**hg-fetch** *ref hash-algo hash [name]* [Procedure]

Return a fixed-output derivation that fetches *ref*, a `<hg-reference>` object. The output is expected to have recursive hash *hash* of type *hash-algo* (a symbol). Use *name* as the file name, or a generic name if `#f`.

**hg-reference** [Data Type]

This data type represents a Mercurial reference for `hg-fetch` to retrieve.

**url** The URL of the Mercurial repository to clone.

**changeset**

This string denotes the changeset to fetch.

For Subversion repositories, the module (`guix svn-download`) defines the **svn-fetch** origin method and **svn-reference** data type for support of the Subversion version control system.

**svn-fetch** *ref hash-algo hash [name]* [Procedure]

Return a fixed-output derivation that fetches *ref*, a **<svn-reference>** object. The output is expected to have recursive hash *hash* of type *hash-algo* (a symbol). Use *name* as the file name, or a generic name if **#f**.

**svn-reference** [Data Type]

This data type represents a Subversion reference for **svn-fetch** to retrieve.

**url** The URL of the Subversion repository to clone.

**revision** This string denotes the revision to fetch specified as a number.

**recursive?** (default: **#f**)

This Boolean indicates whether to recursively fetch Subversion “externals”.

**user-name** (default: **#f**)

The name of an account that has read-access to the repository, if the repository isn’t public.

**password** (default: **#f**)

Password to access the Subversion repository, if required.

For Bazaar repositories, the module (`guix bzz-download`) defines the **bzz-fetch** origin method and **bzz-reference** data type for support of the Bazaar version control system.

**bzz-fetch** *ref hash-algo hash [name]* [Procedure]

Return a fixed-output derivation that fetches *ref*, a **<bzz-reference>** object. The output is expected to have recursive hash *hash* of type *hash-algo* (a symbol). Use *name* as the file name, or a generic name if **#f**.

**bzz-reference** [Data Type]

This data type represents a Bazaar reference for **bzz-fetch** to retrieve.

**url** The URL of the Bazaar repository to clone.

**revision** This string denotes revision to fetch specified as a number.

For CVS repositories, the module (`guix cvs-download`) defines the **cvs-fetch** origin method and **cvs-reference** data type for support of the Concurrent Versions System (CVS).

**cvs-fetch** *ref hash-algo hash [name]* [Procedure]

Return a fixed-output derivation that fetches *ref*, a **<cvs-reference>** object. The output is expected to have recursive hash *hash* of type *hash-algo* (a symbol). Use *name* as the file name, or a generic name if **#f**.

**cvs-reference** [Data Type]

This data type represents a CVS reference for **cvs-fetch** to retrieve.

**root-directory**

The CVS root directory.

**module** Module to fetch.

**revision** Revision to fetch.

The example below denotes a version of **gnu-standards** to fetch:

```
(cvs-reference
  (root-directory ":pserver:anonymous@cvs.savannah.gnu.org:/sources/gnustandards"
  (module "gnustandards")
  (revision "2020-11-25"))
```

## 8.3 Defining Package Variants

One of the nice things with Guix is that, given a package definition, you can easily *derive* variants of that package—for a different upstream version, with different dependencies, different compilation options, and so on. Some of these custom packages can be defined straight from the command line (see Section 9.1.2 [Package Transformation Options], page 182). This section describes how to define package variants in code. This can be useful in “manifests” (see Section 8.4 [Writing Manifests], page 117) and in your own package collection (see Section 6.7 [Creating a Channel], page 72), among others!

As discussed earlier, packages are first-class objects in the Scheme language. The **(guix packages)** module provides the **package** construct to define new package objects (see Section 8.2.1 [package Reference], page 104). The easiest way to define a package variant is using the **inherit** keyword together with **package**. This allows you to inherit from a package definition while overriding the fields you want.

For example, given the **hello** variable, which contains a definition for the current version of GNU Hello, here’s how you would define a variant for version 2.2 (released in 2006, it’s vintage!):

```
(use-modules (gnu packages base)) ;for 'hello'

(define hello-2.2
  (package
    (inherit hello)
    (version "2.2")
    (source (origin
              (method url-fetch)
              (uri (string-append "mirror://gnu/hello/hello-" version
                                   ".tar.gz"))
              (sha256
               (base32
                "0lappv4slgb5spyqbh6yl5r013zv72yqg2pcl30mginf3wdqd8k9"))))))
```

The example above corresponds to what the **--with-version** or **--with-source** package transformations option do. Essentially **hello-2.2** preserves all the fields of **hello**, except

`version` and `source`, which it overrides. Note that the original `hello` variable is still there, in the `(gnu packages base)` module, unchanged. When you define a custom package like this, you are really *adding* a new package definition; the original one remains available.

You can just as well define variants with a different set of dependencies than the original package. For example, the default `gdb` package depends on `guile`, but since that is an optional dependency, you can define a variant that removes that dependency like so:

```
(use-modules (gnu packages gdb)) ;for 'gdb'
```

```
(define gdb-sans-guile
  (package
    (inherit gdb)
    (inputs (modify-inputs (package-inputs gdb)
                          (delete "guile")))))
```

The `modify-inputs` form above removes the `"guile"` package from the `inputs` field of `gdb`. The `modify-inputs` macro is a helper that can prove useful anytime you want to remove, add, or replace package inputs.

**modify-inputs** *inputs clauses* [Macro]

Modify the given package inputs, as returned by `package-inputs` & co., according to the given clauses. Each clause must have one of the following forms:

```
(delete name...)
```

Delete from the inputs packages with the given *names* (strings).

```
(prepend package...)
```

Add *packages* to the front of the input list.

```
(append package...)
```

Add *packages* to the end of the input list.

```
(replace name replacement)
```

Replace the package called *name* with *replacement*.

The example below removes the GMP and ACL inputs of `Coreutils` and adds `libcap` to the front of the input list:

```
(modify-inputs (package-inputs coreutils)
  (delete "gmp" "acl")
  (prepend libcap))
```

The example below replaces the `guile` package from the inputs of `guile-redis` with `guile-2.2`:

```
(modify-inputs (package-inputs guile-redis)
  (replace "guile" guile-2.2))
```

The last type of clause is `append`, to add inputs at the back of the list.

In some cases, you may find it useful to write functions (“procedures”, in Scheme parlance) that return a package based on some parameters. For example, consider the `luasocket` library for the Lua programming language. We want to create `luasocket` packages for major versions of Lua. One way to do that is to define a procedure that takes a Lua package and returns a `luasocket` package that depends on it:

```
(define (make-lua-socket name lua)
```

```
;; Return a luasocket package built with LUA.
(package
  (name name)
  (version "3.0")
  ;; several fields omitted
  (inputs (list lua))
  (synopsis "Socket library for Lua"))

(define-public lua5.1-socket
  (make-lua-socket "lua5.1-socket" lua-5.1))

(define-public lua5.2-socket
  (make-lua-socket "lua5.2-socket" lua-5.2))
```

Here we have defined packages `lua5.1-socket` and `lua5.2-socket` by calling `make-lua-socket` with different arguments. See Section “Procedures” in *GNU Guile Reference Manual*, for more info on procedures. Having top-level public definitions for these two packages means that they can be referred to from the command line (see Section 8.1 [Package Modules], page 100).

These are pretty simple package variants. As a convenience, the `(guix transformations)` module provides a high-level interface that directly maps to the more sophisticated package transformation options (see Section 9.1.2 [Package Transformation Options], page 182):

`options->transformation` *opts* [Procedure]

Return a procedure that, when passed an object to build (package, derivation, etc.), applies the transformations specified by *opts* and returns the resulting objects. *opts* must be a list of symbol/string pairs such as:

```
((with-branch . "guile-gcrypt=master")
 (without-tests . "libgcrypt"))
```

Each symbol names a transformation and the corresponding string is an argument to that transformation.

For instance, a manifest equivalent to this command:

```
guix build guix \
  --with-branch=guile-gcrypt=master \
  --with-debug-info=zlib
```

... would look like this:

```
(use-modules (guix transformations))

(define transform
  ;; The package transformation procedure.
  (options->transformation
    '((with-branch . "guile-gcrypt=master")
      (with-debug-info . "zlib"))))

(packages->manifest
```

```
(list (transform (specification->package "guix"))))
```

The `options->transformation` procedure is convenient, but it's perhaps also not as flexible as you may like. How is it implemented? The astute reader probably noticed that most package transformation options go beyond the superficial changes shown in the first examples of this section: they involve *input rewriting*, whereby the dependency graph of a package is rewritten by replacing specific inputs by others.

Dependency graph rewriting, for the purposes of swapping packages in the graph, is what the `package-input-rewriting` procedure in (`guix packages`) implements.

**package-input-rewriting** *replacements* [*rewrite-name*] [*#:deep?* #f] [*#:recursive?* #f] [Procedure]

Return a procedure that, when passed a package, replaces its direct and indirect dependencies, including implicit inputs when *deep?* is true, according to *replacements*. *replacements* is a list of package pairs; the first element of each pair is the package to replace, and the second one is the replacement.

When *recursive?* is true, apply replacements to the right-hand sides of *replacements* as well, recursively.

Optionally, *rewrite-name* is a one-argument procedure that takes the name of a package and returns its new name after rewrite.

Consider this example:

```
(define libressl-instead-of-openssl
  ;; This is a procedure to replace OPENSSL by LIBRESSL,
  ;; recursively.
  (package-input-rewriting `((,openssl . ,libressl))))

(define git-with-libressl
  (libressl-instead-of-openssl git))
```

Here we first define a rewriting procedure that replaces *openssl* with *libressl*. Then we use it to define a *variant* of the *git* package that uses *libressl* instead of *openssl*. This is exactly what the `--with-input` command-line option does (see Section 9.1.2 [Package Transformation Options], page 182).

The following variant of `package-input-rewriting` can match packages to be replaced by name rather than by identity.

**package-input-rewriting/spec** *replacements* [*#:deep?* #f] [*#:replace-hidden?* #f] [Procedure]

Return a procedure that, given a package, applies the given *replacements* to all the package graph, including implicit inputs unless *deep?* is false.

*replacements* is a list of *spec/procedures* pair; each *spec* is a package specification such as "gcc" or "guile@2", and each procedure takes a matching package and returns a replacement for that package. Matching packages that have the *hidden?* property set are not replaced unless *replace-hidden?* is set to true.

The example above could be rewritten this way:

```
(define libressl-instead-of-openssl
```



```
;; Replace all the packages called "openssl" with LibreSSL.
(package-input-rewriting/spec `(("openssl" . ,(const libressl))))
```

The key difference here is that, this time, packages are matched by spec and not by identity. In other words, any package in the graph that is called `openssl` will be replaced.

A more generic procedure to rewrite a package dependency graph is `package-mapping`: it supports arbitrary changes to nodes in the graph.

`package-mapping` *proc* [*cut?*] [*#:deep?* *#f*] [Procedure]

Return a procedure that, given a package, applies *proc* to all the packages depended on and returns the resulting package. The procedure stops recursion when *cut?* returns true for a given package. When *deep?* is true, *proc* is applied to implicit inputs as well.

**Tips:** Understanding what a variant really looks like can be difficult as one starts combining the tools shown above. There are several ways to inspect a package before attempting to build it that can prove handy:

- You can inspect the package interactively at the REPL, for instance to view its inputs, the code of its build phases, or its configure flags (see Section 8.14 [Using Guix Interactively], page 176).
- When rewriting dependencies, `guix graph` can often help visualize the changes that are made (see Section 9.10 [Invoking guix graph], page 222).

## 8.4 Writing Manifests

`guix` commands let you specify package lists on the command line. This is convenient, but as the command line becomes longer and less trivial, it quickly becomes more convenient to have that package list in what we call a *manifest*. A manifest is some sort of a “bill of materials” that defines a package set. You would typically come up with a code snippet that builds the manifest, store it in a file, say `manifest.scm`, and then pass that file to the `-m` (or `--manifest`) option that many `guix` commands support. For example, here’s what a manifest for a simple package set might look like:

```
;; Manifest for three packages.
(specifications->manifest '("gcc-toolchain" "make" "git"))
```

Once you have that manifest, you can pass it, for example, to `guix package` to install just those three packages to your profile (see [profile-manifest], page 41):

```
guix package -m manifest.scm
```

... or you can pass it to `guix shell` (see [shell-manifest], page 81) to spawn an ephemeral environment:

```
guix shell -m manifest.scm
```

... or you can pass it to `guix pack` in pretty much the same way (see [pack-manifest], page 96). You can store the manifest under version control, share it with others so they can easily get set up, etc.

But how do you write your first manifest? To get started, maybe you’ll want to write a manifest that mirrors what you already have in a profile. Rather than start from a blank page, `guix package` can generate a manifest for you (see [export-manifest], page 45):

```
# Write to 'manifest.scm' a manifest corresponding to the
```

```
# default profile, ~/.guix-profile.
guix package --export-manifest > manifest.scm
```

Or maybe you’ll want to “translate” command-line arguments into a manifest. In that case, `guix shell` can help (see [shell-export-manifest], page 81):

```
# Write a manifest for the packages specified on the command line.
guix shell --export-manifest gcc-toolchain make git > manifest.scm
```

In both cases, the `--export-manifest` option tries hard to generate a faithful manifest; in particular, it takes package transformation options into account (see Section 9.1.2 [Package Transformation Options], page 182).

**Note:** Manifests are *symbolic*: they refer to packages of the channels *currently in use* (see Chapter 6 [Channels], page 68). In the example above, `gcc-toolchain` might refer to version 14 today, but it might refer to version 16 two years from now.

If you want to “pin” your software environment to specific package versions and variants, you need an additional piece of information: the list of channel revisions in use, as returned by `guix describe`. See Section 6.3 [Replicating Guix], page 69, for more information.

Once you’ve obtained your first manifest, perhaps you’ll want to customize it. Since your manifest is code, you now have access to all the Guix programming interfaces!

Let’s assume you want a manifest to deploy a custom variant of GDB, the GNU Debugger, that does not depend on Guile, together with another package. Building on the example seen in the previous section (see Section 8.3 [Defining Package Variants], page 113), you can write a manifest along these lines:

```
(use-modules (guix packages)
              (gnu packages gdb)                ;for 'gdb'
              (gnu packages version-control))    ;for 'git'

;; Define a variant of GDB without a dependency on Guile.
(define gdb-sans-guile
  (package
    (inherit gdb)
    (inputs (modify-inputs (package-inputs gdb)
                           (delete "guile")))))

;; Return a manifest containing that one package plus Git.
(package->manifest (list gdb-sans-guile git))
```

Note that in this example, the manifest directly refers to the `gdb` and `git` variables, which are bound to a `package` object (see Section 8.2.1 [package Reference], page 104), instead of calling `specifications->manifest` to look up packages by name as we did before. The `use-modules` form at the top lets us access the core package interface (see Section 8.2 [Defining Packages], page 101) and the modules that define `gdb` and `git` (see Section 8.1 [Package Modules], page 100). Seamlessly, we’re weaving all this together—the possibilities are endless, unleash your creativity!

The data type for manifests as well as supporting procedures are defined in the (`guix profiles`) module, which is automatically available to code passed to `-m`. The reference follows.

**manifest** [Data Type]

Data type representing a manifest.

It currently has one field:

**entries** This must be a list of `manifest-entry` records—see below.

**manifest-entry** [Data Type]

Data type representing a manifest entry. A manifest entry contains essential meta-data: a name and version string, the object (usually a package) for that entry, the desired output (see Section 5.4 [Packages with Multiple Outputs], page 51), and a number of optional pieces of information detailed below.

Most of the time, you won't build a manifest entry directly; instead, you will pass a package to `package->manifest-entry`, described below. In some unusual cases though, you might want to create manifest entries for things that are *not* packages, as in this example:

```
;; Manually build a single manifest entry for a non-package object.
(let ((hello (program-file "hello" #~(display "Hi!"))))
  (manifest-entry
    (name "foo")
    (version "42")
    (item
      (computed-file "hello-directory"
        #~(let ((bin (string-append #$output "/bin")))
              (mkdir #$output) (mkdir bin)
              (symlink #$hello
                (string-append bin "/hello"))))))))
```

The available fields are the following:

**name**

**version** Name and version string for this entry.

**item** A package or other file-like object (see Section 8.12 [G-Expressions], page 165).

**output** (default: "out")

Output of `item` to use, in case `item` has multiple outputs (see Section 5.4 [Packages with Multiple Outputs], page 51).

**dependencies** (default: '())

List of manifest entries this entry depends on. When building a profile, dependencies are added to the profile.

Typically, the propagated inputs of a package (see Section 8.2.1 [package Reference], page 104) end up having a corresponding manifest entry in among the dependencies of the package's own manifest entry.

**search-paths** (default: '())

The list of search path specifications honored by this entry (see Section 8.8 [Search Paths], page 153).

**properties** (default: '())

List of symbol/value pairs. When building a profile, those properties get serialized.

This can be used to piggyback additional metadata—e.g., the transformations applied to a package (see Section 9.1.2 [Package Transformation Options], page 182).

**parent** (default: (delay #f))

A promise pointing to the “parent” manifest entry.

This is used as a hint to provide context when reporting an error related to a manifest entry coming from a **dependencies** field.

**concatenate-manifests** *lst* [Procedure]

Concatenate the manifests listed in *lst* and return the resulting manifest.

**package->manifest-entry** *package* [*output*] [#:*properties*] [Procedure]

Return a manifest entry for the *output* of package *package*, where *output* defaults to "out", and with the given *properties*. By default *properties* is the empty list or, if one or more package transformations were applied to *package*, it is an association list representing those transformations, suitable as an argument to **options->transformation** (see Section 8.3 [Defining Package Variants], page 113).

The code snippet below builds a manifest with an entry for the default output and the **send-email** output of the *git* package:

```
(use-modules (gnu packages version-control))

(manifest (list (package->manifest-entry git)
                (package->manifest-entry git "send-email")))
```

**packages->manifest** *packages* [Procedure]

Return a list of manifest entries, one for each item listed in *packages*. Elements of *packages* can be either package objects or package/string tuples denoting a specific output of a package.

Using this procedure, the manifest above may be rewritten more concisely:

```
(use-modules (gnu packages version-control))

(packages->manifest (list git `(,git "send-email")))
```

**package->development-manifest** *package* [*system*] [#:*target*] [Procedure]

Return a manifest for the *development inputs* of *package* for *system*, optionally when cross-compiling to *target*. Development inputs include both explicit and implicit inputs of *package*.

Like the **-D** option of **guix shell** (see [shell-development-option], page 80), the resulting manifest describes the environment in which one can develop *package*. For example, suppose you’re willing to set up a development environment for Inkscape,

with the addition of Git for version control; you can describe that “bill of materials” with the following manifest:

```
(use-modules (gnu packages inkscape)           ;for 'inkscape'
              (gnu packages version-control))  ;for 'git'

(concatenate-manifests
 (list (package->development-manifest inkscape)
       (packages->manifest (list git))))
```

In this example, the development manifest that `package->development-manifest` returns includes the compiler (GCC), the many supporting libraries (Boost, GLib, GTK, etc.), and a couple of additional development tools—these are the dependencies `guix show inkscape` lists.

Last, the `(gnu packages)` module provides higher-level facilities to build manifests. In particular, it lets you look up packages by name—see below.

`specifications->manifest specs` [Procedure]

Given *specs*, a list of specifications such as `"emacs@25.2"` or `"guile:debug"`, return a manifest. Specs have the format that command-line tools such as `guix install` and `guix package` understand (see Section 5.2 [Invoking guix package], page 37).

As an example, it lets you rewrite the Git manifest that we saw earlier like this:

```
(specifications->manifest '("git" "git:send-email"))
```

Notice that we do not need to worry about `use-modules`, importing the right set of modules, and referring to the right variables. Instead, we directly refer to packages in the same way as on the command line, which can often be more convenient.

## 8.5 Build Systems

Each package definition specifies a *build system* and arguments for that build system (see Section 8.2 [Defining Packages], page 101). This `build-system` field represents the build procedure of the package, as well as implicit dependencies of that build procedure.

Build systems are `<build-system>` objects. The interface to create and manipulate them is provided by the `(guix build-system)` module, and actual build systems are exported by specific modules.

Under the hood, build systems first compile package objects to *bags*. A *bag* is like a package, but with less ornamentation—in other words, a bag is a lower-level representation of a package, which includes all the inputs of that package, including some that were implicitly added by the build system. This intermediate representation is then compiled to a derivation (see Section 8.10 [Derivations], page 158). The `package-with-c-toolchain` is an example of a way to change the implicit inputs that a package’s build system pulls in (see Section 8.2.1 [package Reference], page 104).

Build systems accept an optional list of *arguments*. In package definitions, these are passed *via* the `arguments` field (see Section 8.2 [Defining Packages], page 101). They are typically keyword arguments (see Section “Optional Arguments” in *GNU Guile Reference Manual*). The value of these arguments is usually evaluated in the *build stratum*—i.e., by a Guile process launched by the daemon (see Section 8.10 [Derivations], page 158).

The main build system is `gnu-build-system`, which implements the standard build procedure for GNU and many other packages. It is provided by the `(guix build-system gnu)` module.

`gnu-build-system` [Variable]

`gnu-build-system` represents the GNU Build System, and variants thereof (see Section “Configuration” in *GNU Coding Standards*).

In a nutshell, packages using it are configured, built, and installed with the usual `./configure && make && make check && make install` command sequence. In practice, a few additional steps are often needed. All these steps are split up in separate *phases*. See Section 8.6 [Build Phases], page 143, for more info on build phases and ways to customize them.

In addition, this build system ensures that the “standard” environment for GNU packages is available. This includes tools such as GCC, libc, Coreutils, Bash, Make, Diffutils, grep, and sed (see the `(guix build-system gnu)` module for a complete list). We call these the *implicit inputs* of a package, because package definitions do not have to mention them.

This build system supports a number of keyword arguments, which can be passed *via* the `arguments` field of a package. Here are some of the main parameters:

`#:phases` This argument specifies build-side code that evaluates to an alist of build phases. See Section 8.6 [Build Phases], page 143, for more information.

`#:configure-flags`  
This is a list of flags (strings) passed to the `configure` script. See Section 8.2 [Defining Packages], page 101, for an example.

`#:make-flags`  
This list of strings contains flags passed as arguments to `make` invocations in the `build`, `check`, and `install` phases.

`#:out-of-source?`  
This Boolean, `#f` by default, indicates whether to run builds in a build directory separate from the source tree.

When it is true, the `configure` phase creates a separate build directory, changes to that directory, and runs the `configure` script from there. This is useful for packages that require it, such as `glibc`.

`#:tests?` This Boolean, `#t` by default, indicates whether the `check` phase should run the package’s test suite.

`#:test-target`  
This string, `"check"` by default, gives the name of the makefile target used by the `check` phase.

`#:parallel-build?`

`#:parallel-tests?`

These Boolean values specify whether to build, respectively run the test suite, in parallel, with the `-j` flag of `make`. When they are true, `make` is passed `-jn`, where `n` is the number specified as the `--cores` option

of `guix-daemon` or that of the `guix` client command (see Section 9.1.1 [Common Build Options], page 180).

**`#:validate-runpath?`**

This Boolean, `#t` by default, determines whether to “validate” the `RUNPATH` of ELF binaries (`.so` shared libraries as well as executables) previously installed by the `install` phase. See [phase-validate-runpath], page 143, for details.

**`#:substitutable?`**

This Boolean, `#t` by default, tells whether the package outputs should be substitutable—i.e., whether users should be able to obtain substitutes for them instead of building locally (see Section 5.3 [Substitutes], page 46).

**`#:allowed-references`**

**`#:disallowed-references`**

When true, these arguments must be a list of dependencies that must not appear among the references of the build results. If, upon build completion, some of these references are retained, the build process fails.

This is useful to ensure that a package does not erroneously keep a reference to some of its build-time inputs, in cases where doing so would, for example, unnecessarily increase its size (see Section 9.9 [Invoking `guix` size], page 220).

Most other build systems support these keyword arguments.

Other `<build-system>` objects are defined to support other conventions and tools used by free software packages. They inherit most of `gnu-build-system`, and differ mainly in the set of inputs implicitly added to the build process, and in the list of phases executed. Some of these build systems are listed below.

**`agda-build-system`**

[Variable]

This variable is exported by `(guix build-system agda)`. It implements a build procedure for Agda libraries.

It adds `agda` to the set of inputs. A different Agda can be specified with the `#:agda` key.

The `#:plan` key is a list of cons cells `(regex . parameters)`, where *regex* is a regexp that should match the `.agda` files to build, and *parameters* is an optional list of parameters that will be passed to `agda` when type-checking it.

When the library uses Haskell to generate a file containing all imports, the convenience `#:gnu-and-haskell?` can be set to `#t` to add `ghc` and the standard inputs of `gnu-build-system` to the input list. You will still need to manually add a phase or tweak the 'build phase, as in the definition of `agda-stdlib`.

**`ant-build-system`**

[Variable]

This variable is exported by `(guix build-system ant)`. It implements the build procedure for Java packages that can be built with Ant build tool (<https://ant.apache.org/>).

It adds both **ant** and the *Java Development Kit* (JDK) as provided by the **icedtea** package to the set of inputs. Different packages can be specified with the **#:ant** and **#:jdk** parameters, respectively.

When the original package does not provide a suitable Ant build file, the parameter **#:jar-name** can be used to generate a minimal Ant build file **build.xml** with tasks to build the specified jar archive. In this case the parameter **#:source-dir** can be used to specify the source sub-directory, defaulting to “src”.

The **#:main-class** parameter can be used with the minimal ant buildfile to specify the main class of the resulting jar. This makes the jar file executable. The **#:test-include** parameter can be used to specify the list of junit tests to run. It defaults to (list “**\*\*/\*Test.java**”). The **#:test-exclude** can be used to disable some tests. It defaults to (list “**\*\*/\*Abstract\*.java**”), because abstract classes cannot be run as tests.

The parameter **#:build-target** can be used to specify the Ant task that should be run during the build phase. By default the “jar” task will be run.

**android-ndk-build-system** [Variable]

This variable is exported by (**guix build-system android-ndk**). It implements a build procedure for Android NDK (native development kit) packages using a Guix-specific build process.

The build system assumes that packages install their public interface (header) files to the subdirectory **include** of the **out** output and their libraries to the subdirectory **lib** the **out** output.

It’s also assumed that the union of all the dependencies of a package has no conflicting files.

For the time being, cross-compilation is not supported - so right now the libraries and header files are assumed to be host tools.

**asdf-build-system/source** [Variable]

**asdf-build-system/sbcl** [Variable]

**asdf-build-system/ecl** [Variable]

**asdf-build-system/clasp** [Variable]

These variables, exported by (**guix build-system asdf**), implement build procedures for Common Lisp packages using “ASDF” (<https://common-lisp.net/project/asdf/>). ASDF is a system definition facility for Common Lisp programs and libraries.

The **asdf-build-system/source** system installs the packages in source form, and can be loaded using any common lisp implementation, via ASDF. The others, such as **asdf-build-system/sbcl**, install binary systems in the format which a particular implementation understands. These build systems can also be used to produce executable programs, or lisp images which contain a set of packages pre-loaded.

The build system uses naming conventions. For binary packages, the package name should be prefixed with the lisp implementation, such as **sbcl-** for **asdf-build-system/sbcl**.



Additionally, the corresponding source package should be labeled using the same convention as Python packages (see Section 22.8.9 [Python Modules], page 836), using the `cl-` prefix.

In order to create executable programs and images, the build-side procedures `build-program` and `build-image` can be used. They should be called in a build phase after the `create-asdf-configuration` phase, so that the system which was just built can be used within the resulting image. `build-program` requires a list of Common Lisp expressions to be passed as the `#:entry-program` argument.

By default, all the `.asd` files present in the sources are read to find system definitions. The `#:asd-files` parameter can be used to specify the list of `.asd` files to read. Furthermore, if the package defines a system for its tests in a separate file, it will be loaded before the tests are run if it is specified by the `#:test-asd-file` parameter. If it is not set, the files `<system>-tests.asd`, `<system>-test.asd`, `tests.asd`, and `test.asd` will be tried if they exist.

If for some reason the package must be named in a different way than the naming conventions suggest, or if several systems must be compiled, the `#:asd-systems` parameter can be used to specify the list of system names.

### `cargo-build-system`

[Variable]

This variable is exported by `(guix build-system cargo)`. It supports builds of packages using Cargo, the build tool of the Rust programming language (<https://www.rust-lang.org>).

**Deprecation warning:** Parameters `#:cargo-inputs` and `#:cargo-development-inputs`, previously used for specifying source inputs, are deprecated. Their support will be removed after Dec. 31, 2026.

It adds `rust` and `rust:cargo` to the set of inputs. A different Rust package can be specified with the `#:rust` parameter.

In its `unpack-rust-crates`, `configure` and `patch-cargo-checksums` phases, this build system will make any source inputs available to Cargo. It will also remove an included `Cargo.lock` file to be recreated by `cargo` during the `build` phase. The `package` phase will run `cargo package` to create a source crate for future use. The `install` phase installs the binaries defined by the crate. Unless `install-source? #f` is defined it will also install a source crate repository of itself and unpacked sources, to ease in future hacking on Rust packages.

This build system supports Cargo workspaces (<https://doc.rust-lang.org/cargo/reference/workspaces.html>). Parameter `#:cargo-package-crates` (default: `'()`) allows specifying names of library crates to package in the `package` phase. Specified crates are packaged from left to right, in case there's dependency among them. For example, specifying `'("pcr2-sys" "pcr2")` will package `"pcr2-sys"` first and then `"pcr2"`. Parameter `#:cargo-install-paths` (default: `'()`) allows specifying paths of binary crates to install in the `install` phase, `'("crates/atuin")`, for example.

See Section 22.8.12 [Rust Crates], page 838, for packaging guidelines.

**chicken-build-system** [Variable]

This variable is exported by `(guix build-system chicken)`. It builds CHICKEN Scheme (<https://call-cc.org/>) modules, also called “eggs” or “extensions”. CHICKEN generates C source code, which then gets compiled by a C compiler, in this case GCC.

This build system adds `chicken` to the package inputs, as well as the packages of `gnu-build-system`.

The build system can’t (yet) deduce the egg’s name automatically, so just like with `go-build-system` and its `#:import-path`, you should define `#:egg-name` in the package’s `arguments` field.

For example, if you are packaging the `srfi-1` egg:

```
(arguments '(:egg-name "srfi-1"))
```

Egg dependencies must be defined in `propagated-inputs`, not `inputs` because CHICKEN doesn’t embed absolute references in compiled eggs. Test dependencies should go to `native-inputs`, as usual.

**copy-build-system** [Variable]

This variable is exported by `(guix build-system copy)`. It supports builds of simple packages that don’t require much compiling, mostly just moving files around.

It adds much of the `gnu-build-system` packages to the set of inputs. Because of this, the `copy-build-system` does not require all the boilerplate code often needed for the `trivial-build-system`.

To further simplify the file installation process, an `#:install-plan` argument is exposed to let the packager specify which files go where. The install plan is a list of `(source target [filters])`. *filters* are optional.

- When *source* matches a file or directory without trailing slash, install it to *target*.
  - If *target* has a trailing slash, install *source* basename beneath *target*.
  - Otherwise install *source* as *target*.
- When *source* is a directory with a trailing slash, or when *filters* are used, the trailing slash of *target* is implied with the same meaning as above.
  - Without *filters*, install the full *source content* to *target*.
  - With *filters* among `#:include`, `#:include-regex`, `#:exclude`, `#:exclude-regex`, only select files are installed depending on the filters. Each filter is specified by a list of strings.
    - With `#:include`, install all the files which the path suffix matches at least one of the elements in the given list.
    - With `#:include-regex`, install all the files which the subpaths match at least one of the regular expressions in the given list.
    - The `#:exclude` and `#:exclude-regex` filters are the complement of their inclusion counterpart. Without `#:include` flags, install all files but those matching the exclusion filters. If both inclusions and exclusions are specified, the exclusions are done on top of the inclusions.
- When a package has multiple outputs, the `#:output` argument can be used to specify which output label the files should be installed to.

In all cases, the paths relative to *source* are preserved within *target*.

Examples:

- (`"foo/bar" "share/my-app/"`): Install `bar` to `share/my-app/bar`.
- (`"foo/bar" "share/my-app/baz"`): Install `bar` to `share/my-app/baz`.
- (`"foo/" "share/my-app"`): Install the content of `foo` inside `share/my-app`, e.g., install `foo/sub/file` to `share/my-app/sub/file`.
- (`"foo/" "share/my-app" #:include ("sub/file")`): Install only `foo/sub/file` to `share/my-app/sub/file`.
- (`"foo/sub" "share/my-app" #:include ("file")`): Install `foo/sub/file` to `share/my-app/file`.
- (`"foo/doc" "share/my-app/doc" #:output "doc"`): Install `"foo/doc"` to `"share/my-app/doc"` within the `"doc"` output.

### `vim-build-system`

[Variable]

This variable is exported by (`guix build-system vim`). It is an extension of the `copy-build-system`, installing Vim and Neovim plugins into locations where these two text editors know to find their plugins, using their packpaths.

Packages which are prefixed with `vim-` will be installed in Vim's packpath, while those prefixed with `neovim-` will be installed in Neovim's packpath. If there is a `doc` directory with the plugin then helptags will be generated automatically.

There are a couple of keywords added with the `vim-build-system`:

- With `plugin-name` it is possible to set the name of the plugin. While by default this is set to the name and version of the package, it is often more helpful to set this to name which the upstream author calls their plugin. This is the name used for `:packadd` from inside Vim.
- With `install-plan` it is possible to augment the built-in install-plan of the `vim-build-system`. This is particularly helpful if you have files which should be installed in other locations. For more information about using the `install-plan`, see the `copy-build-system` (see Section 8.5 [Build Systems], page 121).
- With `#:vim` it is possible to add this package to Vim's packpath, in addition to if it is added automatically because of the `vim-` prefix in the package's name.
- With `#:neovim` it is possible to add this package to Neovim's packpath, in addition to if it is added automatically because of the `neovim-` prefix in the package's name.
- With `#:mode` it is possible to adjust the path which the plugin is installed into. By default the plugin is installed into `start` and other options are available, including `opt`. Adding a plugin into `opt` will mean you will need to run, for example, `:packadd foo` to load the `foo` plugin from inside of Vim.

### `clojure-build-system`

[Variable]

This variable is exported by (`guix build-system clojure`). It implements a simple build procedure for Clojure (<https://clojure.org/>) packages using plain old `compile` in Clojure. Cross-compilation is not supported yet.

It adds `clojure`, `icedtea` and `zip` to the set of inputs. Different packages can be specified with the `#:clojure`, `#:jdk` and `#:zip` parameters, respectively.

A list of source directories, test directories and jar names can be specified with the `#:source-dirs`, `#:test-dirs` and `#:jar-names` parameters, respectively. Compile directory and main class can be specified with the `#:compile-dir` and `#:main-class` parameters, respectively. Other parameters are documented below.

This build system is an extension of `ant-build-system`, but with the following phases changed:

- build** This phase calls `compile` in Clojure to compile source files and runs `jar` to create jars from both source files and compiled files according to the include list and exclude list specified in `#:aot-include` and `#:aot-exclude`, respectively. The exclude list has priority over the include list. These lists consist of symbols representing Clojure libraries or the special keyword `#:all` representing all Clojure libraries found under the source directories. The parameter `#:omit-source?` decides if source should be included into the jars.
- check** This phase runs tests according to the include list and exclude list specified in `#:test-include` and `#:test-exclude`, respectively. Their meanings are analogous to that of `#:aot-include` and `#:aot-exclude`, except that the special keyword `#:all` now stands for all Clojure libraries found under the test directories. The parameter `#:tests?` decides if tests should be run.
- install** This phase installs all jars built previously.

Apart from the above, this build system also contains an additional phase:

- install-doc** This phase installs all top-level files with base name matching `%doc-regex`. A different regex can be specified with the `#:doc-regex` parameter. All files (recursively) inside the documentation directories specified in `#:doc-dirs` are installed as well.

**cmake-build-system** [Variable]

This variable is exported by (`guix build-system cmake`). It implements the build procedure for packages using the CMake build tool (<https://www.cmake.org>).

This build system adds the following keyword parameters to the ones defined by `gnu-build-system`:

- #:cmake** The `cmake` package is added to the set of inputs. Which package is used can be specified with the `#:cmake` parameter.

**#:build-type**

The `#:build-type` parameter specifies in abstract terms the flags passed to the compiler; it defaults to `"RelWithDebInfo"` (short for “release mode with debugging information”), which roughly means that code is compiled with `-O2 -g`, as is the case for Autoconf-based packages by default.

**#:generator**

This parameter specifies the CMake generator (<https://cmake.org/cmake/help/latest/manual/cmake-generators.7.html>) responsible for writing the input files for the native build system.

**#:test-exclude**

Tests matching this regular expression are excluded from testing by ctest (<https://cmake.org/cmake/help/latest/manual/ctest.1.html>).

**#:test-repeat-until-pass?**

Directs ctest (<https://cmake.org/cmake/help/latest/manual/ctest.1.html>) to repeat (<https://cmake.org/cmake/help/latest/manual/ctest.1.html#cmdoption-ctest-repeat>) failed tests up to **#:test-repeat-until-pass-count** times and is enabled by default.

**#:test-repeat-until-pass-count**

When **#:test-repeat-until-pass?** is enabled this parameter sets the maximum number of failures for each test. The default is 5.

**composer-build-system**

[Variable]

This variable is exported by (`guix build-system composer`). It implements the build procedure for packages using Composer (<https://getcomposer.org/>), the PHP package manager.

It automatically adds the `php` package to the set of inputs. Which package is used can be specified with the **#:php** parameter.

The **#:test-target** parameter is used to control which script is run for the tests. By default, the `test` script is run if it exists. If the script does not exist, the build system will run `phpunit` from the source directory, assuming there is a `phpunit.xml` file.

**dune-build-system**

[Variable]

This variable is exported by (`guix build-system dune`). It supports builds of packages using Dune (<https://dune.build/>), a build tool for the OCaml programming language. It is implemented as an extension of the `ocaml-build-system` which is described below. As such, the **#:ocaml** and **#:findlib** parameters can be passed to this build system.

It automatically adds the `dune` package to the set of inputs. Which package is used can be specified with the **#:dune** parameter.

There is no `configure` phase because dune packages typically don't need to be configured. The **#:build-flags** parameter is taken as a list of flags passed to the `dune` command during the build.

The **#:jbuild?** parameter can be passed to use the `jbuild` command instead of the more recent `dune` command while building a package. Its default value is **#f**.

The **#:package** parameter can be passed to specify a package name, which is useful when a package contains multiple packages and you want to build only one of them. This is equivalent to passing the `-p` argument to `dune`.

**elm-build-system** [Variable]

This variable is exported by `(guix build-system elm)`. It implements a build procedure for Elm (<https://elm-lang.org>) packages similar to ‘`elm install`’.

The build system adds an Elm compiler package to the set of inputs. The default compiler package (currently `elm-sans-reactor`) can be overridden using the `#:elm` argument. Additionally, Elm packages needed by the build system itself are added as implicit inputs if they are not already present: to suppress this behavior, use the `#:implicit-elm-package-inputs?` argument, which is primarily useful for bootstrapping.

The `"dependencies"` and `"test-dependencies"` in an Elm package’s `elm.json` file correspond to `propagated-inputs` and `inputs`, respectively.

Elm requires a particular structure for package names: see Section 22.8.13 [Elm Packages], page 839, for more details, including utilities provided by `(guix build-system elm)`.

There are currently a few noteworthy limitations to `elm-build-system`:

- The build system is focused on *packages* in the Elm sense of the word: Elm *projects* which declare `{ "type": "package" }` in their `elm.json` files. Using `elm-build-system` to build Elm *applications* (which declare `{ "type": "application" }`) is possible, but requires ad-hoc modifications to the build phases. For examples, see the definitions of the `elm-todomvc` example application and the `elm` package itself (because the front-end for the ‘`elm reactor`’ command is an Elm application).
- Elm supports multiple versions of a package coexisting simultaneously under `ELM_HOME`, but this does not yet work well with `elm-build-system`. This limitation primarily affects Elm applications, because they specify exact versions for their dependencies, whereas Elm packages specify supported version ranges. As a workaround, the example applications mentioned above use the `patch-application-dependencies` procedure provided by `(guix build-system elm-build-system)` to rewrite their `elm.json` files to refer to the package versions actually present in the build environment. Alternatively, Guix package transformations (see Section 8.3 [Defining Package Variants], page 113) could be used to rewrite an application’s entire dependency graph.
- We are not yet able to run tests for Elm projects because neither `elm-test-rs` (<https://github.com/mpizenberg/elm-test-rs>) nor the Node.js-based `elm-test` (<https://github.com/rtfeldman/node-test-runner>) runner has been packaged for Guix yet.

**go-build-system** [Variable]

This variable is exported by `(guix build-system go)`. It implements a build procedure for Go packages using the standard Go build mechanisms ([https://golang.org/cmd/go/#hdr-Compile\\_packages\\_and\\_dependencies](https://golang.org/cmd/go/#hdr-Compile_packages_and_dependencies)).

The user is expected to provide a value for the key `#:import-path` and, in some cases, `#:unpack-path`. The import path (<https://golang.org/doc/code.html#ImportPaths>) corresponds to the file system path expected by the package’s build scripts and any referring packages, and provides a unique way to refer to a Go package.

It is typically based on a combination of the package source code's remote URI and file system hierarchy structure. In some cases, you will need to unpack the package's source code to a different directory structure than the one indicated by the import path, and `#:unpack-path` should be used in such cases.

Packages that provide Go libraries should install their source code into the built output. The key `#:install-source?`, which defaults to `#t`, controls whether or not the source code is installed. It can be set to `#f` for packages that only provide executable files.

Packages can be cross-built, and if a specific architecture or operating system is desired then the keywords `#:goarch` and `#:goos` can be used to force the package to be built for that architecture and operating system. The combinations known to Go can be found in their documentation (<https://golang.org/doc/install/source#environment>).

The key `#:go` can be used to specify the Go compiler package with which to build the package.

The phase `check` provides a wrapper for `go test` which builds a test binary (or multiple binaries), vets the code and then runs the test binary. Build, test and test binary flags can be provided as `#:test-flags` parameter, default is `'()`. See `go help test` and `go help testflag` for more details.

The key `#:embed-files`, default is `'()`, provides a list of future embedded files or regexps matching files. They will be copied to build directory after `unpack` phase. See <https://pkg.go.dev/embed> for more details.

### `glib-or-gtk-build-system` [Variable]

This variable is exported by `(guix build-system glib-or-gtk)`. It is intended for use with packages making use of GLib or GTK+.

This build system adds the following two phases to the ones defined by `gnu-build-system`:

#### `glib-or-gtk-wrap`

The phase `glib-or-gtk-wrap` ensures that programs in `bin/` are able to find GLib “schemas” and GTK+ modules (<https://developer.gnome.org/gtk3/stable/gtk-running.html>). This is achieved by wrapping the programs in launch scripts that appropriately set the `XDG_DATA_DIRS` and `GTK_PATH` environment variables.

It is possible to exclude specific package outputs from that wrapping process by listing their names in the `#:glib-or-gtk-wrap-excluded-outputs` parameter. This is useful when an output is known not to contain any GLib or GTK+ binaries, and where wrapping would gratuitously add a dependency of that output on GLib and GTK+.

#### `glib-or-gtk-compile-schemas`

The phase `glib-or-gtk-compile-schemas` makes sure that all GSettings schemas (<https://developer.gnome.org/gio/stable/glib-compile-schemas.html>) of GLib are compiled. Compilation is performed by the `glib-compile-schemas` program. It is provided by the package `glib:bin` which is automatically imported by the build

system. The `glib` package providing `glib-compile-schemas` can be specified with the `#:glib` parameter.

Both phases are executed after the `install` phase.

### `guile-build-system` [Variable]

This build system is for Guile packages that consist exclusively of Scheme code and that are so lean that they don't even have a makefile, let alone a `configure` script. It compiles Scheme code using `guild compile` (see Section “Compilation” in *GNU Guile Reference Manual*) and installs the `.scm` and `.go` files in the right place. It also installs documentation.

This build system supports cross-compilation by using the `--target` option of `'guild compile'`.

Packages built with `guile-build-system` must provide a Guile package in their `native-inputs` field.

### `julia-build-system` [Variable]

This variable is exported by `(guix build-system julia)`. It implements the build procedure used by `julia` (<https://julialang.org/>) packages, which essentially is similar to running `'julia -e 'using Pkg; Pkg.add(package)''` in an environment where `JULIA_LOAD_PATH` contains the paths to all Julia package inputs. Tests are run by calling `/test/runtests.jl`.

The Julia package name and uuid is read from the file `Project.toml`. These values can be overridden by passing the argument `#:julia-package-name` (which must be correctly capitalized) or `#:julia-package-uuid`.

Julia packages usually manage their binary dependencies via `JLLWrappers.jl`, a Julia package that creates a module (named after the wrapped library followed by `_jll.jl`).

To add the binary path `_jll.jl` packages, you need to patch the files under `src/wrappers/`, replacing the call to the macro `JLLWrappers.@generate_wrapper_header`, adding as a second argument containing the store path the binary.

As an example, in the `MbedTLS` Julia package, we add a build phase (see Section 8.6 [Build Phases], page 143) to insert the absolute file name of the wrapped `MbedTLS` package:

```
(add-after 'unpack 'override-binary-path
  (lambda* (#:key inputs #:allow-other-keys)
    (for-each (lambda (wrapper)
      (substitute* wrapper
        (("generate_wrapper_header.*")
         (string-append
          "generate_wrapper_header(\"MbedTLS\", \""
          (assoc-ref inputs "mbedtls") "\"")\n"))))
      ;; There's a Julia file for each platform, override them all.
      (find-files "src/wrappers/" "\\.jl$"))))
```

Some older packages that aren't using `Project.toml` yet, will require this file to be created, too. It is internally done if the arguments `#:julia-package-name` and `#:julia-package-uuid` are provided.



**maven-build-system**

[Variable]

This variable is exported by `(guix build-system maven)`. It implements a build procedure for Maven (<https://maven.apache.org>) packages. Maven is a dependency and lifecycle management tool for Java. A user of Maven specifies dependencies and plugins in a `pom.xml` file that Maven reads. When Maven does not have one of the dependencies or plugins in its repository, it will download them and use them to build the package.

The maven build system ensures that maven will not try to download any dependency by running in offline mode. Maven will fail if a dependency is missing. Before running Maven, the `pom.xml` (and subprojects) are modified to specify the version of dependencies and plugins that match the versions available in the guix build environment. Dependencies and plugins must be installed in the fake maven repository at `lib/m2`, and are symlinked into a proper repository before maven is run. Maven is instructed to use that repository for the build and installs built artifacts there. Changed files are copied to the `lib/m2` directory of the package output.

You can specify a `pom.xml` file with the `#:pom-file` argument, or let the build system use the default `pom.xml` file in the sources.

In case you need to specify a dependency's version manually, you can use the `#:local-packages` argument. It takes an association list where the key is the `groupId` of the package and its value is an association list where the key is the `artifactId` of the package and its value is the version you want to override in the `pom.xml`.

Some packages use dependencies or plugins that are not useful at runtime nor at build time in Guix. You can alter the `pom.xml` file to remove them using the `#:exclude` argument. Its value is an association list where the key is the `groupId` of the plugin or dependency you want to remove, and the value is a list of `artifactId` you want to remove.

You can override the default `jdk` and `maven` packages with the corresponding argument, `#:jdk` and `#:maven`.

The `#:maven-plugins` argument is a list of maven plugins used during the build, with the same format as the `inputs` fields of the package declaration. Its default value is `(default-maven-plugins)` which is also exported.

**minetest-mod-build-system**

[Variable]

This variable is exported by `(guix build-system minetest)`. It implements a build procedure for Minetest (<https://www.minetest.net>) mods, which consists of copying Lua code, images and other resources to the location Minetest searches for mods. The build system also minimises PNG images and verifies that Minetest can load the mod without errors.

**minify-build-system**

[Variable]

This variable is exported by `(guix build-system minify)`. It implements a minification procedure for simple JavaScript packages.

It adds `uglify-js` to the set of inputs and uses it to compress all JavaScript files in the `src` directory. A different minifier package can be specified with the `#:uglify-js` parameter, but it is expected that the package writes the minified code to the standard output.

When the input JavaScript files are not all located in the `src` directory, the parameter `#:javascript-files` can be used to specify a list of file names to feed to the minifier.

#### `mix-build-system` [Variable]

This is a variable exported by `guix build-system mix`. It implements the standard build procedure used by Elixir programs with Mix (<https://hexdocs.pm/mix/Mix.html>), which consists of running `mix compile`, thereby creating bytecode.

The API is similar to the *gnu-build-system*, with some specifics for Elixir:

- `#:tests?` is a boolean, `#true` by default, determining whether tests for the current build have to be run or not.
- `#:test-flags` is a list of arguments append to the `mix test` invocation. It defaults to `'()`.
- `#:mix-exs`, changes the full path to the `mix.exs` file. It defaults to `"mix.exs"`. Please refer to upstream (<https://hexdocs.pm/mix/Mix.html#module-environment-variables>) documentation for more details.
- `#:mix-path`, appends extra code paths. It defaults to `#f`. Please refer to upstream (<https://hexdocs.pm/mix/Mix.html#module-environment-variables>) documentation for more details.

#### `mozilla-build-system` [Variable]

This variable is exported by `(guix build-system mozilla)`. It sets the `--target` and `--host` configuration flags to what software developed by Mozilla expects – due to historical reasons, Mozilla software expects `--host` to be the system that is cross-compiled from and `--target` to be the system that is cross-compiled to, contrary to the standard Autotools conventions.

#### `ocaml-build-system` [Variable]

This variable is exported by `(guix build-system ocaml)`. It implements a build procedure for OCaml (<https://ocaml.org>) packages, which consists of choosing the correct set of commands to run for each package. OCaml packages can expect many different commands to be run. This build system will try some of them.

When the package has a `setup.ml` file present at the top-level, it will run `ocaml setup.ml -configure`, `ocaml setup.ml -build` and `ocaml setup.ml -install`. The build system will assume that this file was generated by OASIS (<http://oasis.forge.ocamlcore.org/>) and will take care of setting the prefix and enabling tests if they are not disabled. You can pass configure and build flags with the `#:configure-flags` and `#:build-flags`. The `#:test-flags` key can be passed to change the set of flags used to enable tests. The `#:use-make?` key can be used to bypass this system in the build and install phases.

When the package has a `configure` file, it is assumed that it is a hand-made configure script that requires a different argument format than in the *gnu-build-system*. You can add more flags with the `#:configure-flags` key.

When the package has a `Makefile` file (or `#:use-make?` is `#t`), it will be used and more flags can be passed to the build and install phases with the `#:make-flags` key.

Finally, some packages do not have these files and use a somewhat standard location for its build system. In that case, the build system will run `ocaml pkg/pkg.ml`

or `ocaml pkg/build.ml` and take care of providing the path to the required `findlib` module. Additional flags can be passed via the `#:build-flags` key. Install is taken care of by `opam-installer`. In this case, the `opam` package must be added to the `native-inputs` field of the package definition.

Note that most OCaml packages assume they will be installed in the same directory as OCaml, which is not what we want in `guix`. In particular, they will install `.so` files in their module's directory, which is usually fine because it is in the OCaml compiler directory. In `guix` though, these libraries cannot be found and we use `CAML_LD_LIBRARY_PATH`. This variable points to `lib/ocaml/site-lib/stublibs` and this is where `.so` libraries should be installed.

### `python-build-system` [Variable]

This variable is exported by `(guix build-system python)`. It implements the more or less standard build procedure used by Python packages, which consists in running `python setup.py build` and then `python setup.py install --prefix=/gnu/store/...`

For packages that install stand-alone Python programs under `bin/`, it takes care of wrapping these programs so that their `GUIX_PYTHONPATH` environment variable points to all the Python libraries they depend on.

Which Python package is used to perform the build can be specified with the `#:python` parameter. This is a useful way to force a package to be built for a specific version of the Python interpreter, which might be necessary if the package is only compatible with a single interpreter version.

By default `guix` calls `setup.py` under control of `setuptools`, much like `pip` does. Some packages are not compatible with `setuptools` (and `pip`), thus you can disable this by setting the `#:use-setuptools?` parameter to `#f`.

If a "python" output is available, the package is installed into it instead of the default "out" output. This is useful for packages that include a Python package as only a part of the software, and thus want to combine the phases of `python-build-system` with another build system. Python bindings are a common usecase.

### `pyproject-build-system` [Variable]

This is a variable exported by `guix build-system pyproject`. It is based on `python-build-system`, and adds support for `pyproject.toml` and PEP 517 (<https://peps.python.org/pep-0517/>). It also supports a variety of build backends and test frameworks.

The API is slightly different from `python-build-system`:

- `#:use-setuptools?` and `#:test-target` is removed.
- `#:configure-flags` is changed. Instead of a list this option must be a JSON object, whose interpretation depends on the build backend. For instance the example from PEP 517 (<https://peps.python.org/pep-0517/#config-settings>) should be written as `'(@("CC" "gcc") ("--global-option" ("--some-global-option")) ("--build-option" ("--build-option1" "--build-option2")))`

- `#:backend-path` is added. It defaults to `#false`, but when set to a list it will be appended to Python’s search path and overrides the definition in `pyproject.toml`.
- `#:build-backend` is added. It defaults to `#false` and will try to guess the appropriate backend based on `pyproject.toml`.
- `#:test-backend` is added. It defaults to `#false` and will guess an appropriate test backend based on what is available in package inputs.
- `#:test-flags` is added. The default is `'()`. These flags are passed as arguments to the test command. Note that flags for verbose output is always enabled on supported backends.

It is considered “experimental” in that the implementation details are not set in stone yet, however users are encouraged to try it for new Python projects (even those using `setup.py`). The API is subject to change, but any breaking changes in the Guix channel will be dealt with.

Eventually this build system will be deprecated and merged back into *python-build-system*, probably some time in 2024.

#### `perl-build-system` [Variable]

This variable is exported by `(guix build-system perl)`. It implements the standard build procedure for Perl packages, which either consists in running `perl Build.PL --prefix=/gnu/store/...`, followed by `Build` and `Build install`; or in running `perl Makefile.PL PREFIX=/gnu/store/...`, followed by `make` and `make install`, depending on which of `Build.PL` or `Makefile.PL` is present in the package distribution. Preference is given to the former if both `Build.PL` and `Makefile.PL` exist in the package distribution. This preference can be reversed by specifying `#t` for the `#:make-maker?` parameter.

The initial `perl Makefile.PL` or `perl Build.PL` invocation passes flags specified by the `#:make-maker-flags` or `#:module-build-flags` parameter, respectively.

Which Perl package is used can be specified with `#:perl`.

#### `renpy-build-system` [Variable]

This variable is exported by `(guix build-system renpy)`. It implements the more or less standard build procedure used by Ren’py games, which consists of loading `#:game` once, thereby creating bytecode for it.

It further creates a wrapper script in `bin/` and a desktop entry in `share/applications`, both of which can be used to launch the game.

Which Ren’py package is used can be specified with `#:renpy`. Games can also be installed in outputs other than “out” by using `#:output`.

#### `qt-build-system` [Variable]

This variable is exported by `(guix build-system qt)`. It is intended for use with applications using Qt or KDE.

This build system adds the following two phases to the ones defined by `cmake-build-system`:

**check-setup**

The phase **check-setup** prepares the environment for running the checks as commonly used by Qt test programs. For now this only sets some environment variables: **QT\_QPA\_PLATFORM=offscreen**, **DBUS\_FATAL\_WARNINGS=0** and **CTEST\_OUTPUT\_ON\_FAILURE=1**.

This phase is added before the **check** phase. It's a separate phase to ease adjusting if necessary.

**qt-wrap**

The phase **qt-wrap** searches for Qt5 plugin paths, QML paths and some XDG in the inputs and output. In case some path is found, all programs in the output's **bin/**, **sbin/**, **libexec/** and **lib/libexec/** directories are wrapped in scripts defining the necessary environment variables.

It is possible to exclude specific package outputs from that wrapping process by listing their names in the **#:qt-wrap-excluded-outputs** parameter. This is useful when an output is known not to contain any Qt binaries, and where wrapping would gratuitously add a dependency of that output on Qt, KDE, or such.

This phase is added after the **install** phase.

**r-build-system**

[Variable]

This variable is exported by (**guix build-system r**). It implements the build procedure used by R (<https://r-project.org>) packages, which essentially is little more than running '**R CMD INSTALL --library=/gnu/store/...**' in an environment where **R\_LIBS\_SITE** contains the paths to all R package inputs. Tests are run after installation using the R function **tools::testInstalledPackage**.

**rakudo-build-system**

[Variable]

This variable is exported by (**guix build-system rakudo**). It implements the build procedure used by Rakudo (<https://rakudo.org/>) for Perl6 (<https://perl6.org/>) packages. It installs the package to **/gnu/store/.../NAME-VERSION/share/perl6** and installs the binaries, library files and the resources, as well as wrap the files under the **bin/** directory. Tests can be skipped by passing **#f** to the **tests?** parameter.

Which rakudo package is used can be specified with **rakudo**. Which perl6-tap-harness package used for the tests can be specified with **#:prove6** or removed by passing **#f** to the **with-prove6?** parameter. Which perl6-zef package used for tests and installing can be specified with **#:zef** or removed by passing **#f** to the **with-zef?** parameter.

**rebar-build-system**

[Variable]

This variable is exported by (**guix build-system rebar**). It implements a build procedure around rebar3 (<https://rebar3.org>), a build system for programs written in the Erlang language.

It adds both **rebar3** and the **erlang** to the set of inputs. Different packages can be specified with the **#:rebar** and **#:erlang** parameters, respectively.

This build system is based on **gnu-build-system**, but with the following phases changed:

**unpack**

This phase, after unpacking the source like the **gnu-build-system** does, checks for a file **contents.tar.gz** at the top-level of the source. If this

file exists, it will be unpacked, too. This eases handling of package hosted at <https://hex.pm/>, the Erlang and Elixir package repository.

**bootstrap**  
**configure**

There are no **bootstrap** and **configure** phase because erlang packages typically don't need to be configured.

**build** This phase runs **rebar3 compile** with the flags listed in **#:rebar-flags**.

**check** Unless **#:tests? #f** is passed, this phase runs **rebar3 eunit**, or some other target specified with **#:test-target**, with the flags listed in **#:rebar-flags**,

**install** This installs the files created in the *default* profile, or some other profile specified with **#:install-profile**.

**texlive-build-system**

[Variable]

This variable is exported by (**guix build-system texlive**). It is used to build TeX packages in batch mode with a specified engine. The build system sets the **TEXINPUTS** variable to find all TeX source files in the inputs.

By default it tries to run **luatex** on all **.ins** files, and if it fails to find any, on all **.dtx** files. A different engine and format can be specified with, respectively, the **#:tex-engine** and **#:tex-format** arguments. Different build targets can be specified with the **#:build-targets** argument, which expects a list of file names.

It also generates font metrics (i.e., **.tfm** files) out of Metafont files whenever possible. Likewise, it can also create TeX formats (i.e., **.fmt** files) listed in the **#:create-formats** argument, and generate a symbolic link from **bin/** directory to any script located in **texmf-dist/scripts/**, provided its file name is listed in **#:link-scripts** argument.

The build system adds **texlive-bin** from (**gnu packages tex**) to the native inputs. It can be overridden with the **#:texlive-bin** argument.

The package **texlive-latex-bin**, from the same module, contains most of the tools for building TeX Live packages; for convenience, it is also added by default to the native inputs. However, this can be troublesome when building a dependency of **texlive-latex-bin** itself. In this particular situation, the **#:texlive-latex-bin?** argument should be set to **#f**.

**ruby-build-system**

[Variable]

This variable is exported by (**guix build-system ruby**). It implements the RubyGems build procedure used by Ruby packages, which involves running **gem build** followed by **gem install**.

The **source** field of a package that uses this build system typically references a gem archive, since this is the format that Ruby developers use when releasing their software. The build system unpacks the gem archive, potentially patches the source, runs the test suite, repackages the gem, and installs it. Additionally, directories and tarballs may be referenced to allow building unreleased gems from Git or a traditional source release tarball.

Which Ruby package is used can be specified with the `#:ruby` parameter. A list of additional flags to be passed to the `gem` command can be specified with the `#:gem-flags` parameter.

#### `waf-build-system` [Variable]

This variable is exported by (`guix build-system waf`). It implements a build procedure around the `waf` script. The common phases—`configure`, `build`, and `install`—are implemented by passing their names as arguments to the `waf` script.

The `waf` script is executed by the Python interpreter. Which Python package is used to run the script can be specified with the `#:python` parameter.

#### `zig-build-system` [Variable]

This variable is exported by (`guix build-system zig`). It implements the build procedures for the Zig (<https://ziglang.org/>) build system (`zig build` command).

Selecting this build system adds `zig` to the package inputs, in addition to the packages of `gnu-build-system`.

This build system by default installs package source to output. This behavior can be disabled by setting `#:install-source?` parameter to `#f`.

For packages that don't install anything and don't come with a test suite (likely library packages to be used by other Zig packages), you can set `#:skip-build?` parameter to `#t`, which skips `build` and `check` phases.

The `configure` phase sets up environment for `zig build`. You need to add custom phases after it if you want to invoke `zig`.

The `#:zig-build-flags` parameter is a list of flags that are passed to `zig build` in `build` phase. The `#:zig-test-flags` parameter is a list of flags that are passed to `zig build test` in `check` phase. The default compiler package can be overridden with the `#:zig` parameter.

The optional `#:zig-release-type` parameter declares the type of release. Possible values are: `"safe"`, `"fast"`, or `"small"`. The default value is `#f`, which causes the release flag to be omitted from the `zig` command and results in a `"debug"` build.

#### `scons-build-system` [Variable]

This variable is exported by (`guix build-system scon`s). It implements the build procedure used by the SCons software construction tool. This build system runs `scons` to build the package, `scons test` to run tests, and then `scons install` to install the package.

Additional flags to be passed to `scons` can be specified with the `#:scons-flags` parameter. The default build and install targets can be overridden with `#:build-targets` and `#:install-targets` respectively. The version of Python used to run SCons can be specified by selecting the appropriate SCons package with the `#:scons` parameter.

#### `hare-build-system` [Variable]

This variable is exported by (`guix build-system hare`). It implements the Makefile build process typically used by Hare packages, removing the `configure` phase used by `gnu-build-system` and modifying existing phases to set the Hare architecture tag, module path, and package install destinations appropriately.

This relies on the package’s Makefile supporting, at minimum, the variables `HAREFLAGS` and `THIRDPARTYDIR`. Packages may set `make-flags` to override or supplement these variables.

By default, Hare modules are installed to the `lib` output if it exists, and all other files to the `bin` output. If an output does not exist, the default `out` output is used instead. This behavior may be overridden through the `#:binary-output` and `#:module-output` keyword arguments.

(`guix build-system hare`) also provides utility definitions for creating Hare packages. The variable `%hare-supported-systems` is the list of build targets currently supported by `hare`, for use in `supported-systems` (see Section 8.2.1 [package Reference], page 104). The procedure `target->hare-arch` may be used to get the Hare architecture tag for a given (supported) Guix target or system.

#### `haskell-build-system` [Variable]

This variable is exported by (`guix build-system haskell`). It implements the Cabal build procedure used by Haskell packages, which involves running `runhaskell Setup.hs configure --prefix=/gnu/store/...` and `runhaskell Setup.hs build`. Instead of installing the package by running `runhaskell Setup.hs install`, to avoid trying to register libraries in the read-only compiler store directory, the build system uses `runhaskell Setup.hs copy`, followed by `runhaskell Setup.hs register`. In addition, the build system generates the package documentation by running `runhaskell Setup.hs haddock`, unless `#:haddock? #f` is passed. Optional Haddock parameters can be passed with the help of the `#:haddock-flags` parameter. If the file `Setup.hs` is not found, the build system looks for `Setup.lhs` instead.

Which Haskell compiler is used can be specified with the `#:haskell` parameter which defaults to `ghc`.

#### `dub-build-system` [Variable]

This variable is exported by (`guix build-system dub`). It implements the Dub build procedure used by D packages, which involves running `dub build` and `dub run`. Installation is done by copying the files manually.

Which D compiler is used can be specified with the `#:ldc` parameter which defaults to `ldc`.

#### `emacs-build-system` [Variable]

This variable is exported by (`guix build-system emacs`). It implements an installation procedure similar to the packaging system of Emacs itself (see Section “Packages” in *The GNU Emacs Manual*).

It first creates the `package-autoloads.el` file, then it byte compiles all Emacs Lisp files. Differently from the Emacs packaging system, the Info documentation files are moved to the standard documentation directory and the `dir` file is deleted. The Elisp package files are installed directly under `share/emacs/site-lisp`.

#### `font-build-system` [Variable]

This variable is exported by (`guix build-system font`). It implements an installation procedure for font packages where upstream provides pre-compiled TrueType,



OpenType, etc. font files that merely need to be copied into place. It copies font files to standard locations in the output directory.

**meson-build-system** [Variable]

This variable is exported by (`guix build-system meson`). It implements the build procedure for packages that use Meson (<https://mesonbuild.com>) as their build system.

It adds both Meson and Ninja (<https://ninja-build.org/>) to the set of inputs, and they can be changed with the parameters `#:meson` and `#:ninja` if needed.

This build system is an extension of `gnu-build-system`, but with the following phases changed to some specific for Meson:

**configure**

The phase runs `meson` with the flags specified in `#:configure-flags`. The flag `--buildtype` is always set to `debugoptimized` unless something else is specified in `#:build-type`.

**build**

The phase runs `ninja` to build the package in parallel by default, but this can be changed with `#:parallel-build?`.

**check**

The phase runs `'meson test'` with a base set of options that cannot be overridden. This base set of options can be extended via the `#:test-options` argument, for example to select or skip a specific test suite.

**install**

The phase runs `ninja install` and can not be changed.

Apart from that, the build system also adds the following phases:

**fix-runpath**

This phase ensures that all binaries can find the libraries they need. It searches for required libraries in subdirectories of the package being built, and adds those to `RUNPATH` where needed. It also removes references to libraries left over from the build phase by `meson`, such as test dependencies, that aren't actually required for the program to run.

**glib-or-gtk-wrap**

This phase is the phase provided by `glib-or-gtk-build-system`, and it is not enabled by default. It can be enabled with `#:glib-or-gtk?`.

**glib-or-gtk-compile-schemas**

This phase is the phase provided by `glib-or-gtk-build-system`, and it is not enabled by default. It can be enabled with `#:glib-or-gtk?`.

**linux-module-build-system** [Variable]

`linux-module-build-system` allows building Linux kernel modules.

This build system is an extension of `gnu-build-system`, but with the following phases changed:

**configure**

This phase configures the environment so that the Linux kernel's Makefile can be used to build the external kernel module.

**build** This phase uses the Linux kernel’s Makefile in order to build the external kernel module.

**install** This phase uses the Linux kernel’s Makefile in order to install the external kernel module.

It is possible and useful to specify the Linux kernel to use for building the module (in the **arguments** form of a package using the **linux-module-build-system**, use the key **#:linux** to specify it).

**node-build-system** [Variable]

This variable is exported by (**guix build-system node**). It implements the build procedure used by Node.js (<https://nodejs.org>), which implements an approximation of the **npm install** command, followed by an **npm test** command.

Which Node.js package is used to interpret the **npm** commands can be specified with the **#:node** parameter which defaults to **node**.

**tree-sitter-build-system** [Variable]

This variable is exported by (**guix build-system tree-sitter**). It implements procedures to compile grammars for the Tree-sitter (<https://tree-sitter.github.io/tree-sitter/>) parsing library. It essentially runs **tree-sitter generate** to translate **grammar.js** grammars to JSON and then to C. Which it then compiles to native code.

Tree-sitter packages may support multiple grammars, so this build system supports a **#:grammar-directories** keyword to specify a list of locations where a **grammar.js** file may be found.

Grammars sometimes depend on each other, such as C++ depending on C and TypeScript depending on JavaScript. You may use inputs to declare such dependencies.

Lastly, for packages that do not need anything as sophisticated, a “trivial” build system is provided. It is trivial in the sense that it provides basically no support: it does not pull any implicit inputs, and does not have a notion of build phases.

**trivial-build-system** [Variable]

This variable is exported by (**guix build-system trivial**).

This build system requires a **#:builder** argument. This argument must be a Scheme expression that builds the package output(s)—as with **build-expression->derivation** (see Section 8.10 [Derivations], page 158).

**channel-build-system** [Variable]

This variable is exported by (**guix build-system channel**).

This build system is meant primarily for internal use. A package using this build system must have a channel specification as its **source** field (see Chapter 6 [Channels], page 68); alternatively, its source can be a directory name, in which case an additional **#:commit** argument must be supplied to specify the commit being built (a hexadecimal string).

Optionally, a **#:channels** argument specifying additional channels can be provided. The resulting package is a Guix instance of the given channel(s), similar to how **guix time-machine** would build it.

## 8.6 Build Phases

Almost all package build systems implement a notion *build phases*: a sequence of actions that the build system executes, when you build the package, leading to the installed byproducts in the store. A notable exception is the “bare-bones” **trivial-build-system** (see Section 8.5 [Build Systems], page 121).

As discussed in the previous section, those build systems provide a standard list of phases. For **gnu-build-system**, the main build phases are the following:

### set-paths

Define search path environment variables for all the input packages, including **PATH** (see Section 8.8 [Search Paths], page 153).

### unpack

Unpack the source tarball, and change the current directory to the extracted source tree. If the source is actually a directory, copy it to the build tree, and enter that directory.

### patch-source-shebangs

Patch shebangs encountered in source files so they refer to the right store file names. For instance, this changes `#!/bin/sh` to `#!/gnu/store/...-bash-4.3/bin/sh`.

### configure

Run the **configure** script with a number of default options, such as `--prefix=/gnu/store/...`, as well as the options specified by the `#:configure-flags` argument.

### build

Run **make** with the list of flags specified with `#:make-flags`. If the `#:parallel-build?` argument is true (the default), build with **make -j**.

### check

Run **make check**, or some other target specified with `#:test-target`, unless `#:tests? #f` is passed. If the `#:parallel-tests?` argument is true (the default), run **make check -j**.

### install

Run **make install** with the flags listed in `#:make-flags`.

### patch-shebangs

Patch shebangs on the installed executable files.

### strip

Strip debugging symbols from ELF files (unless `#:strip-binaries?` is false), copying them to the **debug** output when available (see Chapter 17 [Installing Debugging Files], page 801).

### validate-runpath

Validate the **RUNPATH** of ELF binaries, unless `#:validate-runpath?` is false (see Section 8.5 [Build Systems], page 121).

This validation step consists in making sure that all the shared libraries needed by an ELF binary, which are listed as **DT\_NEEDED** entries in its **PT\_DYNAMIC** segment, appear in the **DT\_RUNPATH** entry of that binary. In other words, it ensures that running or using those binaries will not result in a “file not found” error at run time. See Section “Options” in *The GNU Linker*, for more information on **RUNPATH**.

Other build systems have similar phases, with some variations. For example, `cmake-build-system` has same-named phases but its `configure` phases runs `cmake` instead of `./configure`. Others, such as `python-build-system`, have a wholly different list of standard phases. All this code runs on the *build side*: it is evaluated when you actually build the package, in a dedicated build process spawned by the build daemon (see Section 2.3 [Invoking guix-daemon], page 14).

Build phases are represented as association lists or “alists” (see Section “Association Lists” in *GNU Guile Reference Manual*) where each key is a symbol for the name of the phase and the associated value is a procedure that accepts an arbitrary number of arguments. By convention, those procedures receive information about the build in the form of *keyword parameters*, which they can use or ignore.

For example, here is how `(guix build gnu-build-system)` defines `%standard-phases`, the variable holding its alist of build phases<sup>3</sup>:

```
;; The build phases of 'gnu-build-system'.

(define* (unpack #:key source #:allow-other-keys)
  ;; Extract the source tarball.
  (invoke "tar" "xvf" source))

(define* (configure #:key outputs #:allow-other-keys)
  ;; Run the 'configure' script. Install to output "out".
  (let ((out (assoc-ref outputs "out")))
    (invoke "./configure"
             (string-append "--prefix=" out))))

(define* (build #:allow-other-keys)
  ;; Compile.
  (invoke "make"))

(define* (check #:key (test-target "check") (tests? #true)
              #:allow-other-keys)
  ;; Run the test suite.
  (if tests?
      (invoke "make" test-target)
      (display "test suite not run\n")))

(define* (install #:allow-other-keys)
  ;; Install files to the prefix 'configure' specified.
  (invoke "make" "install"))

(define %standard-phases
  ;; The list of standard phases (quite a few are omitted
  ;; for brevity). Each element is a symbol/procedure pair.
  (list (cons 'unpack unpack)
```

---

<sup>3</sup> We present a simplified view of those build phases, but do take a look at `(guix build gnu-build-system)` to see all the details!

```
(cons 'configure configure)
(cons 'build build)
(cons 'check check)
(cons 'install install)))
```

This shows how `%standard-phases` is defined as a list of symbol/procedure pairs (see Section “Pairs” in *GNU Guile Reference Manual*). The first pair associates the `unpack` procedure with the `unpack` symbol—a name; the second pair defines the `configure` phase similarly, and so on. When building a package that uses `gnu-build-system` with its default list of phases, those phases are executed sequentially. You can see the name of each phase started and completed in the build log of packages that you build.

Let’s now look at the procedures themselves. Each one is defined with `define*`: `#:key` lists keyword parameters the procedure accepts, possibly with a default value, and `#:allow-other-keys` specifies that other keyword parameters are ignored (see Section “Optional Arguments” in *GNU Guile Reference Manual*).

The `unpack` procedure honors the `source` parameter, which the build system uses to pass the file name of the source tarball (or version control checkout), and it ignores other parameters. The `configure` phase only cares about the `outputs` parameter, an alist mapping package output names to their store file name (see Section 5.4 [Packages with Multiple Outputs], page 51). It extracts the file name of for `out`, the default output, and passes it to `./configure` as the installation prefix, meaning that `make install` will eventually copy all the files in that directory (see Section “Configuration” in *GNU Coding Standards*). `build` and `install` ignore all their arguments. `check` honors the `test-target` argument, which specifies the name of the Makefile target to run tests; it prints a message and skips tests when `tests?` is false.

The list of phases used for a particular package can be changed with the `#:phases` parameter of the build system. Changing the set of build phases boils down to building a new alist of phases based on the `%standard-phases` alist described above. This can be done with standard alist procedures such as `alist-delete` (see Section “SRFI-1 Association Lists” in *GNU Guile Reference Manual*); however, it is more convenient to do so with `modify-phases` (see Section 8.7 [Build Utilities], page 146).

Here is an example of a package definition that removes the `configure` phase of `%standard-phases` and inserts a new phase before the `build` phase, called `set-prefix-in-makefile`:

```
(define-public example
  (package
    (name "example")
    ;; other fields omitted
    (build-system gnu-build-system)
    (arguments
      (list
        #:phases
        #~(modify-phases %standard-phases
          (delete 'configure)
          (add-before 'build 'set-prefix-in-makefile
            (lambda* (#:key inputs #:allow-other-keys)
```

```
;; Modify the makefile so that its
;; 'PREFIX' variable points to #$output and
;; 'XLLINT' points to the correct path.
(substitute* "Makefile"
  (("PREFIX =.*")
   (string-append "PREFIX = " #$output "\n"))
  (("XLLINT =.*")
   (string-append "XLLINT = "
                   (search-input-file inputs "/bin/xllint")
                   "\n")))))))
```

The new phase that is inserted is written as an anonymous procedure, introduced with `lambda*`; it looks for the `xllint` executable under a `/bin` directory among the package's inputs (see Section 8.2.1 [package Reference], page 104). It also honors the `outputs` parameter we have seen before. See Section 8.7 [Build Utilities], page 146, for more about the helpers used by this phase, and for more examples of `modify-phases`.

**Tip:** You can inspect the code associated with a package's `#:phases` argument interactively, at the REPL (see Section 8.14 [Using Guix Interactively], page 176).

Keep in mind that build phases are code evaluated at the time the package is actually built. This explains why the whole `modify-phases` expression above is quoted (it comes after the `#~` or hash-tilde): it is *staged* for later execution. See Section 8.12 [G-Expressions], page 165, for an explanation of code staging and the *code strata* involved.

## 8.7 Build Utilities

As soon as you start writing non-trivial package definitions (see Section 8.2 [Defining Packages], page 101) or other build actions (see Section 8.12 [G-Expressions], page 165), you will likely start looking for helpers for “shell-like” actions—creating directories, copying and deleting files recursively, manipulating build phases, and so on. The `(guix build utils)` module provides such utility procedures.

Most build systems load `(guix build utils)` (see Section 8.5 [Build Systems], page 121). Thus, when writing custom build phases for your package definitions, you can usually assume those procedures are in scope.

When writing G-expressions, you can import `(guix build utils)` on the “build side” using `with-imported-modules` and then put it in scope with the `use-modules` form (see Section “Using Guile Modules” in *GNU Guile Reference Manual*):

```
(with-imported-modules '((guix build utils)) ;import it
  (computed-file "empty-tree"
    #~(begin
      ;; Put it in scope.
      (use-modules (guix build utils))

      ;; Happily use its 'mkdir-p' procedure.
      (mkdir-p (string-append #$output "/a/b/c")))))
```

The remainder of this section is the reference for most of the utility procedures provided by `(guix build utils)`.

### 8.7.1 Dealing with Store File Names

This section documents procedures that deal with store file names.

**%store-directory** [Procedure]

Return the directory name of the store.

**store-file-name? *file*** [Procedure]

Return true if *file* is in the store.

**strip-store-file-name *file*** [Procedure]

Strip the `/gnu/store` and hash from *file*, a store file name. The result is typically a `"package-version"` string.

**package-name->name+version *name*** [Procedure]

Given *name*, a package name like `"foo-0.9.1b"`, return two values: `"foo"` and `"0.9.1b"`. When the version part is unavailable, *name* and `#f` are returned. The first hyphen followed by a digit is considered to introduce the version part.

### 8.7.2 File Types

The procedures below deal with files and file types.

**directory-exists? *dir*** [Procedure]

Return `#t` if *dir* exists and is a directory.

**executable-file? *file*** [Procedure]

Return `#t` if *file* exists and is executable.

**symbolic-link? *file*** [Procedure]

Return `#t` if *file* is a symbolic link (aka. a “symlink”).

**elf-file? *file*** [Procedure]

**ar-file? *file*** [Procedure]

**gzip-file? *file*** [Procedure]

Return `#t` if *file* is, respectively, an ELF file, an `ar` archive (such as a `.a` static library), or a gzip file.

**reset-gzip-timestamp *file* [`#:keep-mtime?` `#t`]** [Procedure]

If *file* is a gzip file, reset its embedded timestamp (as with `gzip --no-name`) and return true. Otherwise return `#f`. When `keep-mtime?` is true, preserve *file*’s modification time.

### 8.7.3 File Manipulation

The following procedures and macros help create, modify, and delete files. They provide functionality comparable to common shell utilities such as `mkdir -p`, `cp -r`, `rm -r`, and `sed`. They complement Guile’s extensive, but low-level, file system interface (see Section “POSIX” in *GNU Guile Reference Manual*).

**with-directory-excursion *directory* *body* ...** [Macro]

Run *body* with *directory* as the process’s current directory.

Essentially, this macro changes the current directory to *directory* before evaluating *body*, using `chdir` (see Section “Processes” in *GNU Guile Reference Manual*). It changes back to the initial directory when the dynamic extent of *body* is left, be it *via* normal procedure return or *via* a non-local exit such as an exception.

**mkdir-p** *dir* [Procedure]  
Create directory *dir* and all its ancestors.

**install-file** *file directory* [Procedure]  
Create *directory* if it does not exist and copy *file* in there under the same name.

**make-file-writable** *file* [Procedure]  
Make *file* writable for its owner.

**copy-recursively** *source destination* [#:log (*current-output-port*)] [#:follow-symlinks? #f] [#:copy-file *copy-file*] [#:keep-mtime? #f] [#:keep-permissions? #t] [#:select? (*const* #t)] [Procedure]  
Copy *source* directory to *destination*. Follow symlinks if *follow-symlinks?* is true; otherwise, just preserve them. Call *copy-file* to copy regular files. Call *select?*, taking two arguments, *file* and *stat*, for each entry in *source*, where *file* is the entry’s absolute file name and *stat* is the result of `lstat` (or `stat` if *follow-symlinks?* is true); exclude entries for which *select?* does not return true. When *keep-mtime?* is true, keep the modification time of the files in *source* on those of *destination*. When *keep-permissions?* is true, preserve file permissions. Write verbose output to the *log* port.

**delete-file-recursively** *dir* [#:follow-mounts? #f] [Procedure]  
Delete *dir* recursively, like `rm -rf`, without following symlinks. Don’t follow mount points either, unless *follow-mounts?* is true. Report but ignore errors.

**substitute\*** *file* ((*regexp match-var...*) *body...*) ... [Macro]  
Substitute *regexp* in *file* by the string returned by *body*. *body* is evaluated with each *match-var* bound to the corresponding positional regexp sub-expression. For example:

```
(substitute* file
  ("hello")
  "good morning\n")
(("foo([a-z]+)bar(.*)$" all letters end)
 (string-append "baz" letters end)))
```

Here, anytime a line of *file* contains `hello`, it is replaced by `good morning`. Anytime a line of *file* matches the second regexp, *all* is bound to the complete match, *letters* is bound to the first sub-expression, and *end* is bound to the last one.

When one of the *match-var* is `_`, no variable is bound to the corresponding match substring.

Alternatively, *file* may be a list of file names, in which case they are all subject to the substitutions.

Be careful about using `$` to match the end of a line; by itself it won’t match the terminating newline of a line. For example, to match a whole line ending with a backslash, one needs a regexp like `"(.*)\\\\"n$"`.



### 8.7.4 File Search

This section documents procedures to search and filter files.

**file-name-predicate** *regexp* [Procedure]

Return a predicate that returns true when passed a file name whose base name matches *regexp*.

**find-files** *dir* [*pred*] [*#:stat lstat*] [*#:directories? #f*] [*#:fail-on-error? #f*] [Procedure]

Return the lexicographically sorted list of files under *dir* for which *pred* returns true. *pred* is passed two arguments: the absolute file name, and its stat buffer; the default predicate always returns true. *pred* can also be a regular expression, in which case it is equivalent to (**file-name-predicate** *pred*). *stat* is used to obtain file information; using *lstat* means that symlinks are not followed. If *directories?* is true, then directories will also be included. If *fail-on-error?* is true, raise an exception upon error.

Here are a few examples where we assume that the current directory is the root of the Guix source tree:

```
;; List all the regular files in the current directory.
(find-files ".")
⇒ ("./.dir-locals.el" "./.gitignore" ...)

;; List all the .scm files under gnu/services.
(find-files "gnu/services" "\\..scm$")
⇒ ("gnu/services/admin.scm" "gnu/services/audio.scm" ...)

;; List ar files in the current directory.
(find-files "." (lambda (file stat) (ar-file? file)))
⇒ ("./libformat.a" "./libstore.a" ...)
```

**which** *program* [Procedure]

Return the complete file name for *program* as found in \$PATH, or #f if *program* could not be found.

**search-input-file** *inputs name* [Procedure]

**search-input-directory** *inputs name* [Procedure]

Return the complete file name for *name* as found in *inputs*; **search-input-file** searches for a regular file and **search-input-directory** searches for a directory. If *name* could not be found, an exception is raised.

Here, *inputs* must be an association list like **inputs** and **native-inputs** as available to build phases (see Section 8.6 [Build Phases], page 143).

Here is a (simplified) example of how **search-input-file** is used in a build phase of the **wireguard-tools** package:

```
(add-after 'install 'wrap-wg-quick
  (lambda* (#:key inputs outputs #:allow-other-keys)
    (let ((coreutils (string-append (assoc-ref inputs "coreutils"))
```

```

"/bin"))))
(wrap-program (search-input-file outputs "bin/wg-quick")
 #:sh (search-input-file inputs "bin/bash")
 `("PATH" ":" prefix ,(list coreutils))))))

```

### 8.7.5 Program Invocation

You'll find handy procedures to spawn processes in this module, essentially convenient wrappers around Guile's `system*` (see Section “Processes” in *GNU Guile Reference Manual*).

`invoke program args...` [Procedure]

Invoke *program* with the given *args*. Raise an `&invoke-error` exception if the exit code is non-zero; otherwise return `#t`.

The advantage compared to `system*` is that you do not need to check the return value. This reduces boilerplate in shell-script-like snippets for instance in package build phases.

`invoke-error? c` [Procedure]

Return true if *c* is an `&invoke-error` condition.

`invoke-error-program c` [Procedure]

`invoke-error-arguments c` [Procedure]

`invoke-error-exit-status c` [Procedure]

`invoke-error-term-signal c` [Procedure]

`invoke-error-stop-signal c` [Procedure]

Access specific fields of *c*, an `&invoke-error` condition.

`report-invoke-error c [port]` [Procedure]

Report to *port* (by default the current error port) about *c*, an `&invoke-error` condition, in a human-friendly way.

Typical usage would look like this:

```

(use-modules (srfi srfi-34) ;for 'guard'
             (guix build utils))

```

```

(guard (c ((invoke-error? c)
           (report-invoke-error c)))
  (invoke "date" "--imaginary-option"))

```

```

-| command "date" "--imaginary-option" failed with status 1

```

`invoke/quiet program args...` [Procedure]

Invoke *program* with *args* and capture *program*'s standard output and standard error. If *program* succeeds, print nothing and return the unspecified value; otherwise, raise a `&message` error condition that includes the status code and the output of *program*.

Here's an example:

```

(use-modules (srfi srfi-34) ;for 'guard'
             (srfi srfi-35) ;for 'message-condition?'
             (guix build utils))

```

```
(guard (c ((message-condition? c)
          (display (condition-message c))))
  (invoke/quiet "date") ;all is fine
  (invoke/quiet "date" "--imaginary-option"))

+ 'date --imaginary-option' exited with status 1; output follows:

date: unrecognized option '--imaginary-option'
Try 'date --help' for more information.
```

### 8.7.6 Build Phases

The `(guix build utils)` also contains tools to manipulate build phases as used by build systems (see Section 8.5 [Build Systems], page 121). Build phases are represented as association lists or “alists” (see Section “Association Lists” in *GNU Guile Reference Manual*) where each key is a symbol naming the phase and the associated value is a procedure (see Section 8.6 [Build Phases], page 143).

Guile core and the `(srfi srfi-1)` module both provide tools to manipulate alists. The `(guix build utils)` module complements those with tools written with build phases in mind.

**modify-phases** *phases clause...* [Macro]

Modify *phases* sequentially as per each *clause*, which may have one of the following forms:

```
(delete old-phase-name)
(replace old-phase-name new-phase)
(add-before old-phase-name new-phase-name new-phase)
(add-after old-phase-name new-phase-name new-phase)
```

Where every *phase-name* above is an expression evaluating to a symbol, and *new-phase* an expression evaluating to a procedure.

The example below is taken from the definition of the `grep` package. It adds a phase to run after the `install` phase, called `fix-egrep-and-fgrep`. That phase is a procedure (`lambda*` is for anonymous procedures) that takes a `#:outputs` keyword argument and ignores extra keyword arguments (see Section “Optional Arguments” in *GNU Guile Reference Manual*, for more on `lambda*` and optional and keyword arguments.) The phase uses `substitute*` to modify the installed `egrep` and `fgrep` scripts so that they refer to `grep` by its absolute file name:

```
(modify-phases %standard-phases
  (add-after 'install 'fix-egrep-and-fgrep
    ;; Patch 'egrep' and 'fgrep' to execute 'grep' via its
    ;; absolute file name instead of searching for it in $PATH.
    (lambda* (#:key outputs #:allow-other-keys)
      (let* ((out (assoc-ref outputs "out"))
             (bin (string-append out "/bin")))
        (substitute* (list (string-append bin "/egrep"))
```

```

                                (string-append bin "/fgrep"))
    ((("^exec grep")
      (string-append "exec " bin "/grep"))))))))

```

In the example below, phases are modified in two ways: the standard `configure` phase is deleted, presumably because the package does not have a `configure` script or anything similar, and the default `install` phase is replaced by one that manually copies the executable files to be installed:

```

(modify-phases %standard-phases
  (delete 'configure)      ;no 'configure' script
  (replace 'install
    (lambda* (#:key outputs #:allow-other-keys)
      ;; The package's Makefile doesn't provide an "install"
      ;; rule so do it by ourselves.
      (let ((bin (string-append (assoc-ref outputs "out")
                                "/bin")))
        (install-file "footswitch" bin)
        (install-file "scythe" bin))))))

```

### 8.7.7 Wrappers

It is not unusual for a command to require certain environment variables to be set for proper functioning, typically search paths (see Section 8.8 [Search Paths], page 153). Failing to do that, the command might fail to find files or other commands it relies on, or it might pick the “wrong” ones—depending on the environment in which it runs. Examples include:

- a shell script that assumes all the commands it uses are in `PATH`;
- a Guile program that assumes all its modules are in `GUILE_LOAD_PATH` and `GUILE_LOAD_COMPILED_PATH`;
- a Qt application that expects to find certain plugins in `QT_PLUGIN_PATH`.

For a package writer, the goal is to make sure commands always work the same rather than depend on some external settings. One way to achieve that is to *wrap* commands in a thin script that sets those environment variables, thereby ensuring that those run-time dependencies are always found. The wrapper would be used to set `PATH`, `GUILE_LOAD_PATH`, or `QT_PLUGIN_PATH` in the examples above.

To ease that task, the `(guix build utils)` module provides a couple of helpers to wrap commands.

**wrap-program** *program* [*#:sh sh*] [*#:rest variables*] [Procedure]

Make a wrapper for *program*. *variables* should look like this:

```
'(variable delimiter position list-of-directories)
```

where *delimiter* is optional. `:` will be used if *delimiter* is not given.

For example, this call:

```

(wrap-program "foo"
  ("PATH" ":" = ("/gnu/.../bar/bin"))
  ("CERT_PATH" suffix ("/gnu/.../baz/certs"
                       "/qux/certs")))

```

will copy `foo` to `.foo-real` and create the file `foo` with the following contents:

```
#!/location/of/bin/bash
export PATH="/gnu/.../bar/bin"
export CERT_PATH="$CERT_PATH${CERT_PATH:+:}/gnu/.../baz/certs:/qux/certs"
exec -a $0 location/of/.foo-real "$@"
```

If *program* has previously been wrapped by `wrap-program`, the wrapper is extended with definitions for *variables*. If it is not, `sh` will be used as the interpreter.

**wrap-script** *program* [`#:guile guile`] [`#:rest variables`] [Procedure]

Wrap the script *program* such that *variables* are set first. The format of *variables* is the same as in the `wrap-program` procedure. This procedure differs from `wrap-program` in that it does not create a separate shell script. Instead, *program* is modified directly by prepending a Guile script, which is interpreted as a comment in the script's language.

Special encoding comments as supported by Python are recreated on the second line.

Note that this procedure can only be used once per file as Guile scripts are not supported.

## 8.8 Search Paths

Many programs and libraries look for input data in a *search path*, a list of directories: shells like Bash look for executables in the command search path, a C compiler looks for `.h` files in its header search path, the Python interpreter looks for `.py` files in its search path, the spell checker has a search path for dictionaries, and so on.

Search paths can usually be defined or overridden *via* environment variables (see Section “Environment Variables” in *The GNU C Library Reference Manual*). For example, the search paths mentioned above can be changed by defining the `PATH`, `C_INCLUDE_PATH`, `PYTHONPATH` (or `GUIX_PYTHONPATH`), and `DICPATH` environment variables—you know, all these something-PATH variables that you need to get right or things “won’t be found”.

You may have noticed from the command line that Guix “knows” which search path environment variables should be defined, and how. When you install packages in your default profile, the file `~/.guix-profile/etc/profile` is created, which you can “source” from the shell to set those variables. Likewise, if you ask `guix shell` to create an environment containing Python and NumPy, a Python library, and if you pass it the `--search-paths` option, it will tell you about `PATH` and `GUIX_PYTHONPATH` (see Section 7.1 [Invoking guix shell], page 78):

```
$ guix shell python python-numpy --pure --search-paths
export PATH="/gnu/store/...-profile/bin"
export GUIX_PYTHONPATH="/gnu/store/...-profile/lib/python3.9/site-packages"
```

When you omit `--search-paths`, it defines these environment variables right away, such that Python can readily find NumPy:

```
$ guix shell python python-numpy -- python3
Python 3.9.6 (default, Jan 1 1970, 00:00:01)
[GCC 10.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
```

```
>>> import numpy
>>> numpy.version.version
'1.20.3'
```

For this to work, the definition of the `python` package *declares* the search path it cares about and its associated environment variable, `GUIX_PYTHONPATH`. It looks like this:

```
(package
  (name "python")
  (version "3.9.9")
  ;; some fields omitted...
  (native-search-paths
    (list (search-path-specification
            (variable "GUIX_PYTHONPATH")
            (files (list "lib/python/3.9/site-packages"))))))
```

What this `native-search-paths` field says is that, when the `python` package is used, the `GUIX_PYTHONPATH` environment variable must be defined to include all the `lib/python/3.9/site-packages` sub-directories encountered in its environment. (The `native-` bit means that, if we are in a cross-compilation environment, only native inputs may be added to the search path; see Section 8.2.1 [package Reference], page 104.) In the NumPy example above, the profile where `python` appears contains exactly one such sub-directory, and `GUIX_PYTHONPATH` is set to that. When there are several `lib/python/3.9/site-packages`—this is the case in package build environments—they are all added to `GUIX_PYTHONPATH`, separated by colons (:).

**Note:** Notice that `GUIX_PYTHONPATH` is specified as part of the definition of the `python` package, and *not* as part of that of `python-numpy`. This is because this environment variable “belongs” to Python, not NumPy: Python actually reads the value of that variable and honors it.

Corollary: if you create a profile that does not contain `python`, `GUIX_PYTHONPATH` will *not* be defined, even if it contains packages that provide `.py` files:

```
$ guix shell python-numpy --search-paths --pure
export PATH="/gnu/store/...-profile/bin"
```

This makes a lot of sense if we look at this profile in isolation: no software in this profile would read `GUIX_PYTHONPATH`.

Of course, there are many variations on that theme: some packages honor more than one search path, some use separators other than colon, some accumulate several directories in their search path, and so on. A more complex example is the search path of `libxml2`: the value of the `XML_CATALOG_FILES` environment variable is space-separated, it must contain a list of `catalog.xml` files (not directories), which are to be found in `xml` sub-directories—nothing less. The search path specification looks like this:

```
(search-path-specification
  (variable "XML_CATALOG_FILES")
  (separator " ")
  (files '("xml"))
  (file-pattern "^catalog\\.xml$")
  (file-type 'regular))
```

Worry not, search path specifications are usually not this tricky.

The (`guix search-paths`) module defines the data type of search path specifications and a number of helper procedures. Below is the reference of search path specifications.

**search-path-specification** [Data Type]

The data type for search path specifications.

**variable** The name of the environment variable for this search path (a string).

**files** The list of sub-directories (strings) that should be added to the search path.

**separator** (default: `":"`)

The string used to separate search path components.

As a special case, a **separator** value of `#f` specifies a “single-component search path”—in other words, a search path that cannot contain more than one element. This is useful in some cases, such as the `SSL_CERT_DIR` variable (honored by OpenSSL, cURL, and a few other packages) or the `ASPELL_DICT_DIR` variable (honored by the GNU Aspell spell checker), both of which must point to a single directory.

**file-type** (default: `'directory'`)

The type of file being matched—`'directory'` or `'regular'`, though it can be any symbol returned by `stat:type` (see Section “File System” in *GNU Guile Reference Manual*).

In the `XML_CATALOG_FILES` example above, we would match regular files; in the Python example, we would match directories.

**file-pattern** (default: `#f`)

This must be either `#f` or a regular expression specifying files to be matched *within* the sub-directories specified by the `files` field.

Again, the `XML_CATALOG_FILES` example shows a situation where this is needed.

Some search paths are not tied by a single package but to many packages. To reduce duplications, some of them are pre-defined in (`guix search-paths`).

`$SGML_CATALOG_FILES` [Variable]

`$XML_CATALOG_FILES` [Variable]

These two search paths indicate where the TR9401 catalog (<https://www.oasis-open.org/specs/a401.htm>)<sup>4</sup> or XML catalog (<https://www.oasis-open.org/committees/download.php/14809/xml-catalogs.html>) files can be found.

`$SSL_CERT_DIR` [Variable]

`$SSL_CERT_FILE` [Variable]

These two search paths indicate where X.509 certificates can be found (see Section 11.12 [X.509 Certificates], page 679).

<sup>4</sup> Alternatively known as SGML catalog.

These pre-defined search paths can be used as in the following example:

```
(package
  (name "curl")
  ;; some fields omitted ...
  (native-search-paths (list $SSL_CERT_DIR $SSL_CERT_FILE)))
```

How do you turn search path specifications on one hand and a bunch of directories on the other hand in a set of environment variable definitions? That's the job of `evaluate-search-paths`.

**evaluate-search-paths** *search-paths* *directories* [*getenv*] [Procedure]  
 Evaluate *search-paths*, a list of search-path specifications, for *directories*, a list of directory names, and return a list of specification/value pairs. Use *getenv* to determine the current settings and report only settings not already effective.

The (`guix profiles`) provides a higher-level helper procedure, `load-profile`, that sets the environment variables of a profile.

## 8.9 The Store

Conceptually, the *store* is the place where derivations that have been built successfully are stored—by default, `/gnu/store`. Sub-directories in the store are referred to as *store items* or sometimes *store paths*. The store has an associated database that contains information such as the store paths referred to by each store path, and the list of *valid* store items—results of successful builds. This database resides in `localstatedir/guix/db`, where *localstatedir* is the state directory specified *via* `--localstatedir` at configure time, usually `/var`.

The store is *always* accessed by the daemon on behalf of its clients (see Section 2.3 [Invoking `guix-daemon`], page 14). To manipulate the store, clients connect to the daemon over a Unix-domain socket, send requests to it, and read the result—these are remote procedure calls, or RPCs.

**Note:** Users must *never* modify files under `/gnu/store` directly. This would lead to inconsistencies and break the immutability assumptions of Guix's functional model (see Chapter 1 [Introduction], page 1).

See Section 5.6 [Invoking `guix gc`], page 53, for information on how to check the integrity of the store and attempt recovery from accidental modifications.

The (`guix store`) module provides procedures to connect to the daemon, and to perform RPCs. These are described below. By default, `open-connection`, and thus all the `guix` commands, connect to the local daemon or to the URI specified by the `GUIX_DAEMON_SOCKET` environment variable.

**GUIX\_DAEMON\_SOCKET** [Environment Variable]

When set, the value of this variable should be a file name or a URI designating the daemon endpoint. When it is a file name, it denotes a Unix-domain socket to connect to. In addition to file names, the supported URI schemes are:

```
file
unix      These are for Unix-domain sockets.  file:///var/guix/daemon-socket/socket is equivalent to /var/guix/daemon-socket/socket.
```



**guix** These URIs denote connections over TCP/IP, without encryption nor authentication of the remote host. The URI must specify the host name and optionally a port number (by default port 44146 is used):

```
guix://master.guix.example.org:1234
```

This setup is suitable on local networks, such as clusters, where only trusted nodes may connect to the build daemon at `master.guix.example.org`.

The `--listen` option of `guix-daemon` can be used to instruct it to listen for TCP connections (see Section 2.3 [Invoking `guix-daemon`], page 14).

**ssh** These URIs allow you to connect to a remote daemon over SSH. This feature requires Guile-SSH (see Section 22.1 [Requirements], page 816) and a working `guile` binary in `PATH` on the destination machine. It supports public key and GSSAPI authentication. A typical URL might look like this:

```
ssh://charlie@guix.example.org:22
```

As for `guix copy`, the usual OpenSSH client configuration files are honored (see Section 9.13 [Invoking `guix copy`], page 233).

Additional URI schemes may be supported in the future.

**Note:** The ability to connect to remote build daemons is considered experimental as of 8b23207. Please get in touch with us to share any problems or suggestions you may have (see Chapter 22 [Contributing], page 816).

**open-connection** [*uri*] [*#:reserve-space?* *#t*] [Procedure]

Connect to the daemon over the Unix-domain socket at *uri* (a string). When *reserve-space?* is true, instruct it to reserve a little bit of extra space on the file system so that the garbage collector can still operate should the disk become full. Return a server object.

*file* defaults to `%default-socket-path`, which is the normal location given the options that were passed to `configure`.

**close-connection** *server* [Procedure]

Close the connection to *server*.

**current-build-output-port** [Variable]

This variable is bound to a SRFI-39 parameter, which refers to the port where build and error logs sent by the daemon should be written.

Procedures that make RPCs all take a server object as their first argument.

**valid-path?** *server path* [Procedure]

Return `#t` when *path* designates a valid store item and `#f` otherwise (an invalid item may exist on disk but still be invalid, for instance because it is the result of an aborted or failed build).

A `&store-protocol-error` condition is raised if *path* is not prefixed by the store directory (`/gnu/store`).

**add-text-to-store** *server name text* [references] [Procedure]

Add *text* under file *name* in the store, and return its store path. *references* is the list of store paths referred to by the resulting store path.

**build-derivations** *store derivations* [*mode*] [Procedure]

Build *derivations*, a list of <derivation> objects, *.drv* file names, or derivation/output pairs, using the specified *mode*—(build-mode normal) by default.

Note that the (guix monads) module provides a monad as well as monadic versions of the above procedures, with the goal of making it more convenient to work with code that accesses the store (see Section 8.11 [The Store Monad], page 161).

*This section is currently incomplete.*

## 8.10 Derivations

Low-level build actions and the environment in which they are performed are represented by *derivations*. A derivation contains the following pieces of information:

- The outputs of the derivation—derivations produce at least one file or directory in the store, but may produce more.
- The inputs of the derivation—i.e., its build-time dependencies—which may be other derivations or plain files in the store (patches, build scripts, etc.).
- The system type targeted by the derivation—e.g., `x86_64-linux`.
- The file name of a build script in the store, along with the arguments to be passed.
- A list of environment variables to be defined.

Derivations allow clients of the daemon to communicate build actions to the store. They exist in two forms: as an in-memory representation, both on the client- and daemon-side, and as files in the store whose name end in *.drv*—these files are referred to as *derivation paths*. Derivations paths can be passed to the **build-derivations** procedure to perform the build actions they prescribe (see Section 8.9 [The Store], page 156).

Operations such as file downloads and version-control checkouts for which the expected content hash is known in advance are modeled as *fixed-output derivations*. Unlike regular derivations, the outputs of a fixed-output derivation are independent of its inputs—e.g., a source code download produces the same result regardless of the download method and tools being used.

The outputs of derivations—i.e., the build results—have a set of *references*, as reported by the **references** RPC or the **guix gc --references** command (see Section 5.6 [Invoking guix gc], page 53). References are the set of run-time dependencies of the build results. References are a subset of the inputs of the derivation; this subset is automatically computed by the build daemon by scanning all the files in the outputs.

The (guix derivations) module provides a representation of derivations as Scheme objects, along with procedures to create and otherwise manipulate derivations. The lowest-level primitive to create a derivation is the **derivation** procedure:

```

derivation store name builder args [#:outputs '("out")] [#:hash      [Procedure]
      [#f] [#:hash-algo #f] [#:recursive? #f] [#:inputs '()] [#:env-vars '()]
      [#:system (%current-system)] [#:references-graphs #f]
      [#:allowed-references #f] [#:disallowed-references #f] [#:leaked-env-vars
      [#f] [#:local-build? #f] [#:substitutable? #t] [#:properties '()]

```

Build a derivation with the given arguments, and return the resulting `<derivation>` object.

When *hash* and *hash-algo* are given, a *fixed-output derivation* is created—i.e., one whose result is known in advance, such as a file download. If, in addition, *recursive?* is true, then that fixed output may be an executable file or a directory and *hash* must be the hash of an archive containing this output.

When *references-graphs* is true, it must be a list of file name/store path pairs. In that case, the reference graph of each store path is exported in the build environment in the corresponding file, in a simple text format.

When *allowed-references* is true, it must be a list of store items or outputs that the derivation’s output may refer to. Likewise, *disallowed-references*, if true, must be a list of things the outputs may *not* refer to.

When *leaked-env-vars* is true, it must be a list of strings denoting environment variables that are allowed to “leak” from the daemon’s environment to the build environment. This is only applicable to fixed-output derivations—i.e., when *hash* is true. The main use is to allow variables such as `http_proxy` to be passed to derivations that download files.

When *local-build?* is true, declare that the derivation is not a good candidate for offloading and should rather be built locally (see Section 2.2.2 [Daemon Offload Setup], page 9). This is the case for small derivations where the costs of data transfers would outweigh the benefits.

When *substitutable?* is false, declare that substitutes of the derivation’s output should not be used (see Section 5.3 [Substitutes], page 46). This is useful, for instance, when building packages that capture details of the host CPU instruction set.

*properties* must be an association list describing “properties” of the derivation. It is kept as-is, uninterpreted, in the derivation.

Here’s an example with a shell script as its builder, assuming *store* is an open connection to the daemon, and *bash* points to a Bash executable in the store:

```

(use-modules (guix utils)
             (guix store)
             (guix derivations))

(let ((builder ; add the Bash script to the store
      (add-text-to-store store "my-builder.sh"
                          "echo hello world > $out\n" '())))
  (derivation store "foo"
    bash `("-e" ,builder)
    #:inputs `((,bash) (,builder))
    #:env-vars '(("HOME" . "/homeless"))))

```

```
⇒ #<derivation /gnu/store/...-foo.drv => /gnu/store/...-foo>
```

As can be guessed, this primitive is cumbersome to use directly. A better approach is to write build scripts in Scheme, of course! The best course of action for that is to write the build code as a “G-expression”, and to pass it to `gexp->derivation`. For more information, see Section 8.12 [G-Expressions], page 165.

Once upon a time, `gexp->derivation` did not exist and constructing derivations with build code written in Scheme was achieved with `build-expression->derivation`, documented below. This procedure is now deprecated in favor of the much nicer `gexp->derivation`.

```
build-expression->derivation store name exp [#:system [Procedure]
  (%current-system)] [#:inputs '()] [#:outputs '("out")] [#:hash #f]
  [#:hash-algo #f] [#:recursive? #f] [#:env-vars '()] [#:modules '()]
  [#:references-graphs #f] [#:allowed-references #f]
  [#:disallowed-references #f] [#:local-build? #f] [#:substitutable? #t]
  [#:guile-for-build #f]
```

Return a derivation that executes Scheme expression *exp* as a builder for derivation *name*. *inputs* must be a list of (name drv-path sub-drv) tuples; when *sub-drv* is omitted, “out” is assumed. *modules* is a list of names of Guile modules from the current search path to be copied in the store, compiled, and made available in the load path during the execution of *exp*—e.g., ((guix build utils) (guix build gnu-build-system)).

*exp* is evaluated in an environment where `%outputs` is bound to a list of output/path pairs, and where `%build-inputs` is bound to a list of string/output-path pairs made from *inputs*. Optionally, *env-vars* is a list of string pairs specifying the name and value of environment variables visible to the builder. The builder terminates by passing the result of *exp* to `exit`; thus, when *exp* returns `#f`, the build is considered to have failed.

*exp* is built using *guile-for-build* (a derivation). When *guile-for-build* is omitted or is `#f`, the value of the `%guile-for-build` fluid is used instead.

See the `derivation` procedure for the meaning of *references-graphs*, *allowed-references*, *disallowed-references*, *local-build?*, and *substitutable?*.

Here’s an example of a single-output derivation that creates a directory containing one file:

```
(let ((builder '(let ((out (assoc-ref %outputs "out")))
  (mkdir out)      ; create /gnu/store/...-goo
  (call-with-output-file (string-append out "/test")
    (lambda (p)
      (display '(hello guix) p))))))
  (build-expression->derivation store "goo" builder))
```

```
⇒ #<derivation /gnu/store/...-goo.drv => ...>
```

## 8.11 The Store Monad

The procedures that operate on the store described in the previous sections all take an open connection to the build daemon as their first argument. Although the underlying model is functional, they either have side effects or depend on the current state of the store.

The former is inconvenient: the connection to the build daemon has to be carried around in all those functions, making it impossible to compose functions that do not take that parameter with functions that do. The latter can be problematic: since store operations have side effects and/or depend on external state, they have to be properly sequenced.

This is where the (`guix monads`) module comes in. This module provides a framework for working with *monads*, and a particularly useful monad for our uses, the *store monad*. Monads are a construct that allows two things: associating “context” with values (in our case, the context is the store), and building sequences of computations (here computations include accesses to the store). Values in a monad—values that carry this additional context—are called *monadic values*; procedures that return such values are called *monadic procedures*.

Consider this “normal” procedure:

```
(define (sh-symlink store)
  ;; Return a derivation that symlinks the 'bash' executable.
  (let* ((drv (package-derivation store bash))
        (out (derivation->output-path drv))
        (sh (string-append out "/bin/bash")))
    (build-expression->derivation store "sh"
      `(symlink ,sh %output))))
```

Using (`guix monads`) and (`guix gexp`), it may be rewritten as a monadic function:

```
(define (sh-symlink)
  ;; Same, but return a monadic value.
  (mlet %store-monad ((drv (package->derivation bash)))
    (gexp->derivation "sh"
      #~(symlink (string-append #$drv "/bin/bash")
        #$output))))
```

There are several things to note in the second version: the `store` parameter is now implicit and is “threaded” in the calls to the `package->derivation` and `gexp->derivation` monadic procedures, and the monadic value returned by `package->derivation` is *bound* using `mlet` instead of plain `let`.

As it turns out, the call to `package->derivation` can even be omitted since it will take place implicitly, as we will see later (see Section 8.12 [G-Expressions], page 165):

```
(define (sh-symlink)
  (gexp->derivation "sh"
    #~(symlink (string-append #$bash "/bin/bash")
      #$output)))
```

Calling the monadic `sh-symlink` has no effect. As someone once said, “you exit a monad like you exit a building on fire: by running”. So, to exit the monad and get the desired effect, one must use `run-with-store`:

```
(run-with-store (open-connection) (sh-symlink))
```

```
⇒ /gnu/store/...-sh-symlink
```

Note that the `(guix monad-repl)` module extends the Guile REPL with new “commands” to make it easier to deal with monadic procedures: `run-in-store`, and `enter-store-monad` (see Section 8.14 [Using Guix Interactively], page 176). The former is used to “run” a single monadic value through the store:

```
scheme@(guile-user)> ,run-in-store (package->derivation hello)
$1 = #<derivation /gnu/store/...-hello-2.9.drv => ...>
```

The latter enters a recursive REPL, where all the return values are automatically run through the store:

```
scheme@(guile-user)> ,enter-store-monad
store-monad@(guile-user) [1]> (package->derivation hello)
$2 = #<derivation /gnu/store/...-hello-2.9.drv => ...>
store-monad@(guile-user) [1]> (text-file "foo" "Hello!")
$3 = "/gnu/store/...-foo"
store-monad@(guile-user) [1]> ,q
scheme@(guile-user)>
```

Note that non-monadic values cannot be returned in the `store-monad` REPL.

Other meta-commands are available at the REPL, such as `,build` to build a file-like object (see Section 8.14 [Using Guix Interactively], page 176).

The main syntactic forms to deal with monads in general are provided by the `(guix monads)` module and are described below.

**with-monad** *monad body* ... [Macro]

Evaluate any `>>=` or `return` forms in *body* as being in *monad*.

**return** *val* [Macro]

Return a monadic value that encapsulates *val*.

**>>=** *mval mproc* ... [Macro]

Bind monadic value *mval*, passing its “contents” to monadic procedures *mproc*...<sup>5</sup>.

There can be one *mproc* or several of them, as in this example:

```
(run-with-state
  (with-monad %state-monad
    (>>= (return 1)
         (lambda (x) (return (+ 1 x)))
         (lambda (x) (return (* 2 x))))))
'some-state

⇒ 4
⇒ some-state
```

**mlet** *monad ((var mval) ...) body* ... [Macro]

**mlet\*** *monad ((var mval) ...) body* ... [Macro]

Bind the variables *var* to the monadic values *mval* in *body*, which is a sequence of expressions. As with the `bind` operator, this can be thought of as “unpacking” the

---

<sup>5</sup> This operation is commonly referred to as “bind”, but that name denotes an unrelated procedure in Guile. Thus we use this somewhat cryptic symbol inherited from the Haskell language.

raw, non-monadic value “contained” in *mval* and making *var* refer to that raw, non-monadic value within the scope of the *body*. The form *(var -> val)* binds *var* to the “normal” value *val*, as per *let*. The binding operations occur in sequence from left to right. The last expression of *body* must be a monadic expression, and its result will become the result of the *mlet* or *mlet\** when run in the *monad*.

*mlet\** is to *mlet* what *let\** is to *let* (see Section “Local Bindings” in *GNU Guile Reference Manual*).

***mbegin monad mexp ...*** [Macro]

Bind *mexp* and the following monadic expressions in sequence, returning the result of the last expression. Every expression in the sequence must be a monadic expression.

This is akin to *mlet*, except that the return values of the monadic expressions are ignored. In that sense, it is analogous to *begin*, but applied to monadic expressions.

***mwhen condition mexp0 mexp\* ...*** [Macro]

When *condition* is true, evaluate the sequence of monadic expressions *mexp0..mexp\** as in an *mbegin*. When *condition* is false, return *\*unspecified\** in the current monad. Every expression in the sequence must be a monadic expression.

***munless condition mexp0 mexp\* ...*** [Macro]

When *condition* is false, evaluate the sequence of monadic expressions *mexp0..mexp\** as in an *mbegin*. When *condition* is true, return *\*unspecified\** in the current monad. Every expression in the sequence must be a monadic expression.

The *(guix monads)* module provides the *state monad*, which allows an additional value—the state—to be *threaded* through monadic procedure calls.

***%state-monad*** [Variable]

The state monad. Procedures in the state monad can access and change the state that is threaded.

Consider the example below. The *square* procedure returns a value in the state monad. It returns the square of its argument, but also increments the current state value:

```
(define (square x)
  (mlet %state-monad ((count (current-state)))
    (mbegin %state-monad
      (set-current-state (+ 1 count))
      (return (* x x)))))

(run-with-state (sequence %state-monad (map square (iota 3))) 0)
⇒ (0 1 4)
⇒ 3
```

When “run” through *%state-monad*, we obtain that additional state value, which is the number of *square* calls.

***current-state*** [Monadic Procedure]

Return the current state as a monadic value.

**set-current-state** *value* [Monadic Procedure]  
 Set the current state to *value* and return the previous state as a monadic value.

**state-push** *value* [Monadic Procedure]  
 Push *value* to the current state, which is assumed to be a list, and return the previous state as a monadic value.

**state-pop** [Monadic Procedure]  
 Pop a value from the current state and return it as a monadic value. The state is assumed to be a list.

**run-with-state** *mval* [*state*] [Procedure]  
 Run monadic value *mval* starting with *state* as the initial state. Return two values: the resulting value, and the resulting state.

The main interface to the store monad, provided by the (`guix store`) module, is as follows.

**%store-monad** [Variable]  
 The store monad—an alias for **%state-monad**.

Values in the store monad encapsulate accesses to the store. When its effect is needed, a value of the store monad must be “evaluated” by passing it to the **run-with-store** procedure (see below).

**run-with-store** *store mval* [*#:guile-for-build*] [*#:system*] [*#:current-system*] [Procedure]  
 Run *mval*, a monadic value in the store monad, in *store*, an open store connection.

**text-file** *name text* [*references*] [Monadic Procedure]  
 Return as a monadic value the absolute file name in the store of the file containing *text*, a string. *references* is a list of store items that the resulting text file refers to; it defaults to the empty list.

**binary-file** *name data* [*references*] [Monadic Procedure]  
 Return as a monadic value the absolute file name in the store of the file containing *data*, a bytevector. *references* is a list of store items that the resulting binary file refers to; it defaults to the empty list.

**interned-file** *file* [*name*] [*#:recursive?* *#t*] [*#:select?* (*const #t*)] [Monadic Procedure]

Return the name of *file* once interned in the store. Use *name* as its store name, or the basename of *file* if *name* is omitted.

When *recursive?* is true, the contents of *file* are added recursively; if *file* designates a flat file and *recursive?* is true, its contents are added, and its permission bits are kept.

When *recursive?* is true, call (`select? file stat`) for each directory entry, where *file* is the entry’s absolute file name and *stat* is the result of `lstat`; exclude entries for which *select?* does not return true.



The example below adds a file to the store, under two different names:

```
(run-with-store (open-connection)
  (mlet %store-monad ((a (interned-file "README"))
                      (b (interned-file "README" "LEGU-MIN"))))
  (return (list a b))))

⇒ ("/gnu/store/rwm...-README" "/gnu/store/44i...-LEGU-MIN")
```

The `(guix packages)` module exports the following package-related monadic procedures:

`package-file` *package* [*file*] [*#:system*] [Monadic Procedure]  
 (%*current-system*) [*#:target* *#f*] [*#:output* "out"]

Return as a monadic value in the absolute file name of *file* within the *output* directory of *package*. When *file* is omitted, return the name of the *output* directory of *package*. When *target* is true, use it as a cross-compilation target triplet.

Note that this procedure does *not* build *package*. Thus, the result might or might not designate an existing file. We recommend not using this procedure unless you know what you are doing.

`package->derivation` *package* [*system*] [Monadic Procedure]  
`package->cross-derivation` *package* *target* [*system*] [Monadic Procedure]

Monadic version of `package-derivation` and `package-cross-derivation` (see Section 8.2 [Defining Packages], page 101).

## 8.12 G-Expressions

So we have “derivations”, which represent a sequence of build actions to be performed to produce an item in the store (see Section 8.10 [Derivations], page 158). These build actions are performed when asking the daemon to actually build the derivations; they are run by the daemon in a container (see Section 2.3 [Invoking guix-daemon], page 14).

It should come as no surprise that we like to write these build actions in Scheme. When we do that, we end up with two *strata* of Scheme code<sup>6</sup>: the “host code”—code that defines packages, talks to the daemon, etc.—and the “build code”—code that actually performs build actions, such as making directories, invoking `make`, and so on (see Section 8.6 [Build Phases], page 143).

To describe a derivation and its build actions, one typically needs to embed build code inside host code. It boils down to manipulating build code as data, and the homoiconicity of Scheme—code has a direct representation as data—comes in handy for that. But we need more than the normal `quasiquote` mechanism in Scheme to construct build expressions.

The `(guix gexp)` module implements *G-expressions*, a form of S-expressions adapted to build expressions. G-expressions, or *gexps*, consist essentially of three syntactic forms: `gexp`, `ungexp`, and `ungexp-splicing` (or simply: `#~`, `#$`, and `#$@`), which are comparable

<sup>6</sup> The term *stratum* in this context was coined by Manuel Serrano et al. in the context of their work on Hop. Oleg Kiselyov, who has written insightful essays and code on this topic (<http://okmij.org/ftp/meta-programming/#meta-scheme>), refers to this kind of code generation as *staging*.

to `quasiquote`, `unquote`, and `unquote-splicing`, respectively (see Section “Expression Syntax” in *GNU Guile Reference Manual*). However, there are major differences:

- Gexps are meant to be written to a file and run or manipulated by other processes.
- When a high-level object such as a package or derivation is unquoted inside a gexp, the result is as if its output file name had been introduced.
- Gexps carry information about the packages or derivations they refer to, and these dependencies are automatically added as inputs to the build processes that use them.

This mechanism is not limited to package and derivation objects: *compilers* able to “lower” other high-level objects to derivations or files in the store can be defined, such that these objects can also be inserted into gexps. For example, a useful type of high-level objects that can be inserted in a gexp is “file-like objects”, which make it easy to add files to the store and to refer to them in derivations and such (see `local-file` and `plain-file` below).

To illustrate the idea, here is an example of a gexp:

```
(define build-exp
  #~(begin
    (mkdir #$output)
    (chdir #$output)
    (symlink (string-append #$coreutils "/bin/ls")
             "list-files")))
```

This gexp can be passed to `gexp->derivation`; we obtain a derivation that builds a directory containing exactly one symlink to `/gnu/store/...-coreutils-8.22/bin/ls`:

```
(gexp->derivation "the-thing" build-exp)
```

As one would expect, the `/gnu/store/...-coreutils-8.22` string is substituted to the reference to the *coreutils* package in the actual build code, and *coreutils* is automatically made an input to the derivation. Likewise, `#$output` (equivalent to `(ungexp output)`) is replaced by a string containing the directory name of the output of the derivation.

In a cross-compilation context, it is useful to distinguish between references to the *native* build of a package—that can run on the host—versus references to cross builds of a package. To that end, the `#+` plays the same role as `#$`, but is a reference to a native package build:

```
(gexp->derivation "vi"
  #~(begin
    (mkdir #$output)
    (mkdir (string-append #$output "/bin"))
    (system* (string-append #+coreutils "/bin/ln")
             "-s"
             (string-append #$emacs "/bin/emacs")
             (string-append #$output "/bin/vi")))
  #:target "aarch64-linux-gnu")
```

In the example above, the native build of *coreutils* is used, so that `ln` can actually run on the host; but then the cross-compiled build of *emacs* is referenced.

Another gexp feature is *imported modules*: sometimes you want to be able to use certain Guile modules from the “host environment” in the gexp, so those modules should be

imported in the “build environment”. The `with-imported-modules` form allows you to express that:

```
(let ((build (with-imported-modules '((guix build utils))
                                     #~(begin
                                       (use-modules (guix build utils))
                                       (mkdir-p (string-append #output "/bin"))))))
  (gexp->derivation "empty-dir"
    #~(begin
      #build
      (display "success!\n")
      #t)))
```

In this example, the `(guix build utils)` module is automatically pulled into the isolated build environment of our `gexp`, such that `(use-modules (guix build utils))` works as expected.

Usually you want the *closure* of the module to be imported—i.e., the module itself and all the modules it depends on—rather than just the module; failing to do that, attempts to use the module will fail because of missing dependent modules. The `source-module-closure` procedure computes the closure of a module by looking at its source file headers, which comes in handy in this case:

```
(use-modules (guix modules)) ;for 'source-module-closure'

(with-imported-modules (source-module-closure
                        '((guix build utils)
                          (gnu build image)))
  (gexp->derivation "something-with-vms"
    #~(begin
      (use-modules (guix build utils)
                    (gnu build image))
      ...)))
```

In the same vein, sometimes you want to import not just pure-Scheme modules, but also “extensions” such as Guile bindings to C libraries or other “full-blown” packages. Say you need the `guile-json` package available on the build side, here’s how you would do it:

```
(use-modules (gnu packages guile)) ;for 'guile-json'

(with-extensions (list guile-json)
  (gexp->derivation "something-with-json"
    #~(begin
      (use-modules (json))
      ...)))
```

The syntactic form to construct `gexps` is summarized below.

`#~exp` [Macro]

`(gexp exp)` [Macro]

Return a G-expression containing *exp*. *exp* may contain one or more of the following forms:

**#\$obj**

(ungexp *obj*)

Introduce a reference to *obj*. *obj* may have one of the supported types, for example a package or a derivation, in which case the **ungexp** form is replaced by its output file name—e.g., `"/gnu/store/...-coreutils-8.22"`.

If *obj* is a list, it is traversed and references to supported objects are substituted similarly.

If *obj* is another gexp, its contents are inserted and its dependencies are added to those of the containing gexp.

If *obj* is another kind of object, it is inserted as is.

**#\$obj:output**

(ungexp *obj output*)

This is like the form above, but referring explicitly to the *output* of *obj*—this is useful when *obj* produces multiple outputs (see Section 5.4 [Packages with Multiple Outputs], page 51).

Sometimes a gexp unconditionally refers to the "out" output, but the user of that gexp would still like to insert a reference to another output. The **gexp-input** procedure aims to address that. See [gexp-input], page 174.

**#+obj**

**#+obj:output**

(ungexp-native *obj*)

(ungexp-native *obj output*)

Same as **ungexp**, but produces a reference to the *native* build of *obj* when used in a cross compilation context.

**#\$output[:output]**

(ungexp output [*output*])

Insert a reference to derivation output *output*, or to the main output when *output* is omitted.

This only makes sense for gexps passed to **gexp->derivation**.

**#\$@lst**

(ungexp-splicing *lst*)

Like the above, but splices the contents of *lst* inside the containing list.

**#+@lst**

(ungexp-native-splicing *lst*)

Like the above, but refers to native builds of the objects listed in *lst*.

G-expressions created by **gexp** or **#~** are run-time objects of the **gexp?** type (see below).

**with-imported-modules** *modules body...*

[Macro]

Mark the gexps defined in *body...* as requiring *modules* in their execution environment.

Each item in *modules* can be the name of a module, such as `(guix build utils)`, or it can be a module name, followed by an arrow, followed by a file-like object:

```
((guix build utils)
 (guix gcrypt)
 ((guix config) => ,(scheme-file "config.scm"
                                #~(define-module ...))))
```

In the example above, the first two modules are taken from the search path, and the last one is created from the given file-like object.

This form has *lexical* scope: it has an effect on the gexps directly defined in *body...*, but not on those defined, say, in procedures called from *body...*

**with-extensions** *extensions body...* [Macro]

Mark the gexps defined in *body...* as requiring *extensions* in their build and execution environment. *extensions* is typically a list of package objects such as those defined in the `(gnu packages guile)` module.

Concretely, the packages listed in *extensions* are added to the load path while compiling imported modules in *body...*; they are also added to the load path of the gexp returned by *body...*

**gexp?** *obj* [Procedure]

Return `#t` if *obj* is a G-expression.

G-expressions are meant to be written to disk, either as code building some derivation, or as plain files in the store. The monadic procedures below allow you to do that (see Section 8.11 [The Store Monad], page 161, for more information about monads).

**gexp->derivation** *name exp* [#:system (%current-system)] [Monadic Procedure]  
 [#:target #f] [#:graft? #t] [#:hash #f] [#:hash-algo #f] [#:recursive?  
 #f] [#:env-vars '()] [#:modules '()] [#:module-path %load-path]  
 [#:effective-version "2.2"] [#:references-graphs #f] [#:allowed-references  
 #f] [#:disallowed-references #f] [#:leaked-env-vars #f] [#:script-name  
 (string-append *name* "-builder")] [#:deprecation-warnings #f]  
 [#:local-build? #f] [#:substitutable? #t] [#:properties '()]  
 [#:guile-for-build #f]

Return a derivation *name* that runs *exp* (a gexp) with *guile-for-build* (a derivation) on *system*; *exp* is stored in a file called *script-name*. When *target* is true, it is used as the cross-compilation target triplet for packages referred to by *exp*.

*modules* is deprecated in favor of **with-imported-modules**. Its meaning is to make *modules* available in the evaluation context of *exp*; *modules* is a list of names of Guile modules searched in *module-path* to be copied in the store, compiled, and made available in the load path during the execution of *exp*—e.g., `((guix build utils) (guix build gnu-build-system))`.

*effective-version* determines the string to use when adding extensions of *exp* (see **with-extensions**) to the search path—e.g., "2.2".

*graft?* determines whether packages referred to by *exp* should be grafted when applicable.

When *references-graphs* is true, it must be a list of tuples of one of the following forms:

```
(file-name obj)
(file-name obj output)
(file-name gexp-input)
(file-name store-item)
```

The right-hand-side of each element of *references-graphs* is automatically made an input of the build process of *exp*. In the build environment, each *file-name* contains the reference graph of the corresponding item, in a simple text format.

*allowed-references* must be either **#f** or a list of output names and packages. In the latter case, the list denotes store items that the result is allowed to refer to. Any reference to another store item will lead to a build error. Similarly for *disallowed-references*, which can list items that must not be referenced by the outputs.

*deprecation-warnings* determines whether to show deprecation warnings while compiling modules. It can be **#f**, **#t**, or **'detailed**.

The other arguments are as for **derivation** (see Section 8.10 [Derivations], page 158).

The **local-file**, **plain-file**, **computed-file**, **program-file**, and **scheme-file** procedures below return *file-like objects*. That is, when unquoted in a G-expression, these objects lead to a file in the store. Consider this G-expression:

```
#~(system* #$(file-append glibc "/sbin/nscd") "-f"
      #$(local-file "/tmp/my-ns.cd.conf"))
```

The effect here is to “intern” `/tmp/my-ns.cd.conf` by copying it to the store. Once expanded, for instance *via* `gexp->derivation`, the G-expression refers to that copy under `/gnu/store`; thus, modifying or removing the file in `/tmp` does not have any effect on what the G-expression does. **plain-file** can be used similarly; it differs in that the file content is directly passed as a string.

**local-file** *file* [*name*] [*#:recursive?* *#f*] [*#:select?* (*const* *#t*)] [Procedure]

Return an object representing local file *file* to add to the store; this object can be used in a gexp. If *file* is a literal string denoting a relative file name, it is looked up relative to the source file where it appears; if *file* is not a literal string, it is looked up relative to the current working directory at run time. *file* will be added to the store under *name*—by default the base name of *file*.

When *recursive?* is true, the contents of *file* are added recursively; if *file* designates a flat file and *recursive?* is true, its contents are added, and its permission bits are kept.

When *recursive?* is true, call (*select? file stat*) for each directory entry, where *file* is the entry’s absolute file name and *stat* is the result of **lstat**; exclude entries for which *select?* does not return true.

*file* can be wrapped in the **assume-valid-file-name** syntactic keyword. When this is done, there will not be a warning when **local-file** is used with a non-literal path. The path is still looked up relative to the current working directory at run time. Wrapping is done like this:

```
(define alice-key-file-path "alice.pub")
```

```
;; ...
(local-file (assume-valid-file-name alice-key-file-path))
```

*file* can be wrapped in the `assume-source-relative-file-name` syntactic keyword. When this is done, the file name will be looked up relative to the source file where it appears even when it is not a string literal.

This is the declarative counterpart of the `interned-file` monadic procedure (see Section 8.11 [The Store Monad], page 161).

**plain-file** *name content* [Procedure]

Return an object representing a text file called *name* with the given *content* (a string or a bytevector) to be added to the store.

This is the declarative counterpart of `text-file`.

**computed-file** *name gexp* [#:guile] [#:local-build? #t] [#:options] [Procedure]

Return an object representing the store item *name*, a file or directory computed by *gexp*. When *local-build?* is true (the default), the derivation is built locally. *options* is a list of additional arguments to pass to `gexp->derivation`.

This is the declarative counterpart of `gexp->derivation`.

**gexp->script** *name exp* [#:guile (default-guile)] [Monadic Procedure]  
[#:module-path %load-path] [#:system (%current-system)] [#:target  
'current]

Return an executable script *name* that runs *exp* using *guile*, with *exp*'s imported modules in its search path. Look up *exp*'s modules in *module-path*.

The example below builds a script that simply invokes the `ls` command:

```
(use-modules (guix gexp) (gnu packages base))

(gexp->script "list-files"
  #~(execl #$(file-append coreutils "/bin/ls")
    "ls"))
```

When “running” it through the store (see Section 8.11 [The Store Monad], page 161), we obtain a derivation that produces an executable file `/gnu/store/...-list-files` along these lines:

```
#!/gnu/store/...-guile-2.0.11/bin/guile -ds
!#
(execl "/gnu/store/...-coreutils-8.22/bin/ls" "ls")
```

**program-file** *name exp* [#:guile #f] [#:module-path %load-path] [Procedure]

Return an object representing the executable store item *name* that runs *gexp*. *guile* is the Guile package used to execute that script. Imported modules of *gexp* are looked up in *module-path*.

This is the declarative counterpart of `gexp->script`.

**gexp->file** *name exp* [*#:set-load-path? #t*] [Monadic Procedure]  
 [*#:module-path %load-path*] [*#:splice? #f*] [*#:guile (default-guile)*]  
 [*#:system (%current-system)*] [*#:target 'current*]

Return a derivation that builds a file *name* containing *exp*. When *splice?* is true, *exp* is considered to be a list of expressions that will be spliced in the resulting file.

When *set-load-path?* is true, emit code in the resulting file to set *%load-path* and *%load-compiled-path* to honor *exp*'s imported modules. Look up *exp*'s modules in *module-path*.

The resulting file holds references to all the dependencies of *exp* or a subset thereof.

**scheme-file** *name exp* [*#:splice? #f*] [*#:guile #f*] [Procedure]  
 [*#:set-load-path? #t*]

Return an object representing the Scheme file *name* that contains *exp*. *guile* is the Guile package used to produce that file.

This is the declarative counterpart of *gexp->file*.

**text-file\*** *name text* ... [Monadic Procedure]

Return as a monadic value a derivation that builds a text file containing all of *text*. *text* may list, in addition to strings, objects of any type that can be used in a *gexp*: packages, derivations, local file objects, etc. The resulting store file holds references to all these.

This variant should be preferred over *text-file* anytime the file to create will reference items from the store. This is typically the case when building a configuration file that embeds store file names, like this:

```
(define (profile.sh)
  ;; Return the name of a shell script in the store that
  ;; initializes the 'PATH' environment variable.
  (text-file* "profile.sh"
    "export PATH=" coreutils "/bin:"
    grep "/bin:" sed "/bin\n"))
```

In this example, the resulting */gnu/store/...-profile.sh* file will reference *coreutils*, *grep*, and *sed*, thereby preventing them from being garbage-collected during its lifetime.

**mixed-text-file** *name text* ... [Procedure]

Return an object representing store file *name* containing *text*. *text* is a sequence of strings and file-like objects, as in:

```
(mixed-text-file "profile"
  "export PATH=" coreutils "/bin:" grep "/bin")
```

This is the declarative counterpart of *text-file\**.

**file-union** *name files* [*#:dangling-symlinks? #f*] [Procedure]

Return a *<computed-file>* that builds a directory containing all of *files*. Each item in *files* must be a two-element list where the first element is the file name to use in the new directory, and the second element is a *gexp* denoting the target file. *#:dangling-symlinks?* controls if *gexps* must lower to an existing file. Here's an example:

```
(file-union "etc"
```



```

`(("hosts" ,(plain-file "hosts"
                        "127.0.0.1 localhost"))
  ("bashrc" ,(plain-file "bashrc"
                        "alias ls='ls --color=auto'")))))■

```

This yields an `etc` directory containing these two files.

`symlink-to target` [Procedure]

Return a <computed-file> that is a symbolic link to `target`. Note, that the target does not need to exist at the build time. One of uses is to put arbitrary symlinks into user's home:

```

(service home-files-service-type
  `(("Sync" ,(symlink-to "/storage/Sync"))))

```

`directory-union name things` [Procedure]

Return a directory that is the union of *things*, where *things* is a list of file-like objects denoting directories. For example:

```

(directory-union "guile+emacs" (list guile emacs))

```

yields a directory that is the union of the `guile` and `emacs` packages.

`file-append obj suffix ...` [Procedure]

Return a file-like object that expands to the concatenation of *obj* and *suffix*, where *obj* is a lowerable object and each *suffix* is a string.

As an example, consider this gexp:

```

(gexp->script "run-uname"
  #~(system* #$(file-append coreutils
                          "/bin/uname")))

```

The same effect could be achieved with:

```

(gexp->script "run-uname"
  #~(system* (string-append #$(coreutils
                          "/bin/uname")))

```

There is one difference though: in the `file-append` case, the resulting script contains the absolute file name as a string, whereas in the second case, the resulting script contains a `(string-append ...)` expression to construct the file name *at run time*.

`let-system system body...` [Macro]

`let-system (system target) body...` [Macro]

Bind *system* to the currently targeted system—e.g., `"x86_64-linux"`—within *body*.

In the second case, additionally bind *target* to the current cross-compilation target—a GNU triplet such as `"arm-linux-gnueabi"`—or `#f` if we are not cross-compiling.

`let-system` is useful in the occasional case where the object spliced into the gexp depends on the target system, as in this example:

```

#~(system*
  #+(let-system system
    (cond ((string-prefix? "armhf-" system)
      (file-append qemu "/bin/qemu-system-arm"))

```

```

      ((string-prefix? "x86_64-" system)
       (file-append qemu "/bin/qemu-system-x86_64"))
      (else
       (error "dunno!"))))
  "-net" "user" #$image)

```

**with-parameters** ((*parameter value*) ...) *exp* [Macro]

This macro is similar to the **parameterize** form for dynamically-bound *parameters* (see Section “Parameters” in *GNU Guile Reference Manual*). The key difference is that it takes effect when the file-like object returned by *exp* is lowered to a derivation or store item.

A typical use of **with-parameters** is to force the system in effect for a given object:

```

(with-parameters ((%current-system "i686-linux"))
  coreutils)

```

The example above returns an object that corresponds to the i686 build of Coreutils, regardless of the current value of `%current-system`.

**delayed object exp** [Macro]

This macro delays the evaluation of *exp* until the returned object is lowered to a derivation or store item.

Its intended use case is to prevent the use of a non-delayed top-level variable of another module, e.g. when using a **gexp** in the **source** field of a *package*.

**gexp-input obj [output] [#:native? #f]** [Procedure]

Return a *gexp input* record for the given *output* of file-like object *obj*, with `#:native?` determining whether this is a native reference (as with **ungexp-native**) or not.

This procedure is helpful when you want to pass a reference to a specific output of an object to some procedure that may not know about that output. For example, assume you have this procedure, which takes one file-like object:

```

(define (make-symlink target)
  (computed-file "the-symlink"
    #~(symlink #$target #$output)))

```

Here **make-symlink** can only ever refer to the default output of *target*—the “out” output (see Section 5.4 [Packages with Multiple Outputs], page 51). To have it refer to, say, the “lib” output of the **hwloc** package, you can call it like so:

```

(make-symlink (gexp-input hwloc "lib"))

```

You can also compose it like any other file-like object:

```

(make-symlink
  (file-append (gexp-input hwloc "lib") "/lib/libhwloc.so"))

```

Of course, in addition to **gexps** embedded in “host” code, there are also modules containing build tools. To make it clear that they are meant to be used in the build stratum, these modules are kept in the (**guix build** ...) name space.

Internally, high-level objects are *lowered*, using their compiler, to either derivations or store items. For instance, lowering a package yields a derivation, and lowering a **plain-file** yields a store item. This is achieved using the **lower-object** monadic procedure.

`lower-object obj [system] [#:target #f]` [Monadic Procedure]

Return as a value in `%store-monad` the derivation or store item corresponding to *obj* for *system*, cross-compiling for *target* if *target* is true. *obj* must be an object that has an associated gexp compiler, such as a `<package>`.

`gexp->approximate-sexp gexp` [Procedure]

Sometimes, it may be useful to convert a G-exp into a S-exp. For example, some linters (see Section 9.8 [Invoking guix lint], page 217) peek into the build phases of a package to detect potential problems. This conversion can be achieved with this procedure. However, some information can be lost in the process. More specifically, lowerable objects will be silently replaced with some arbitrary object – currently the list `(*approximate*)`, but this may change.

## 8.13 Invoking guix repl

The `guix repl` command makes it easier to program Guix in Guile by launching a Guile *read-eval-print loop* (REPL) for interactive programming (see Section “Using Guile Interactively” in *GNU Guile Reference Manual*), or by running Guile scripts (see Section “Running Guile Scripts” in *GNU Guile Reference Manual*). Compared to just launching the `guile` command, `guix repl` guarantees that all the Guix modules and all its dependencies are available in the search path.

The general syntax is:

```
guix repl options [file args]
```

When a *file* argument is provided, *file* is executed as a Guile script:

```
guix repl my-script.scm
```

To pass arguments to the script, use `--` to prevent them from being interpreted as arguments to `guix repl` itself:

```
guix repl -- my-script.scm --input=foo.txt
```

To make a script executable directly from the shell, using the `guix` executable that is on the user’s search path, add the following two lines at the top of the script:

```
#!/usr/bin/env -S guix repl --
!#
```

To make a script that launches an interactive REPL directly from the shell, use the `--interactive` flag:

```
#!/usr/bin/env -S guix repl --interactive
!#
```

Without a file name argument, a Guile REPL is started, allowing for interactive use (see Section 8.14 [Using Guix Interactively], page 176):

```
$ guix repl
scheme@(guile-user)> ,use (gnu packages base)
scheme@(guile-user)> coreutils
$1 = #<package coreutils@8.29 gnu/packages/base.scm:327 3e28300>
```

In addition, `guix repl` implements a simple machine-readable REPL protocol for use by `(guix inferior)`, a facility to interact with *inferiors*, separate processes running a potentially different revision of Guix.

The available options are as follows:

- `--list-types`  
Display the *TYPE* options for `guix repl --type=TYPE` and exit.
- `--type=type`  
`-t type` Start a REPL of the given *TYPE*, which can be one of the following:
  - `guile` This is default, and it spawns a standard full-featured Guile REPL.
  - `machine` Spawn a REPL that uses the machine-readable protocol. This is the protocol that the (`guix inferior`) module speaks.
- `--listen=endpoint`  
By default, `guix repl` reads from standard input and writes to standard output. When this option is passed, it will instead listen for connections on *endpoint*. Here are examples of valid options:
  - `--listen=tcp:37146`  
Accept connections on localhost on port 37146.
  - `--listen=unix:/tmp/socket`  
Accept connections on the Unix-domain socket `/tmp/socket`.
- `--interactive`  
`-i` Launch the interactive REPL after *file* is executed.
- `--load-path=directory`  
`-L directory`  
Add *directory* to the front of the package module search path (see Section 8.1 [Package Modules], page 100).  
This allows users to define their own packages and make them visible to the script or REPL.
- `-q` Inhibit loading of the `~/.guile` file. By default, that configuration file is loaded when spawning a `guile` REPL.

## 8.14 Using Guix Interactively

The `guix repl` command gives you access to a warm and friendly *read-eval-print loop* (REPL) (see Section 8.13 [Invoking `guix repl`], page 175). If you're getting into Guix programming—defining your own packages, writing manifests, defining services for Guix System or Guix Home, etc.—you will surely find it convenient to toy with ideas at the REPL.

If you use Emacs, the most convenient way to do that is with Geiser (see Section 22.5 [The Perfect Setup], page 822), but you do not have to use Emacs to enjoy the REPL. When using `guix repl` or `guile` in the terminal, we recommend using Readline for completion and Colorized to get colorful output. To do that, you can run:

```
guix install guile guile-readline guile-colorized
```

... and then create a `.guile` file in your home directory containing this:

```
(use-modules (ice-9 readline) (ice-9 colorized))
```

```
(activate-readline)
(activate-colored)
```

The REPL lets you evaluate Scheme code; you type a Scheme expression at the prompt, and the REPL prints what it evaluates to:

```
$ guix repl
scheme@(guix-user)> (+ 2 3)
$1 = 5
scheme@(guix-user)> (string-append "a" "b")
$2 = "ab"
```

It becomes interesting when you start fiddling with Guix at the REPL. The first thing you'll want to do is to “import” the `(guix)` module, which gives access to the main part of the programming interface, and perhaps a bunch of useful Guix modules. You could type `(use-modules (guix))`, which is valid Scheme code to import a module (see Section “Using Guile Modules” in *GNU Guile Reference Manual*), but the REPL provides the `use` command as a shorthand notation (see Section “REPL Commands” in *GNU Guile Reference Manual*):

```
scheme@(guix-user)> ,use (guix)
scheme@(guix-user)> ,use (gnu packages base)
```

Notice that REPL commands are introduced by a leading comma. A REPL command like `use` is not valid Scheme code; it's interpreted specially by the REPL.

Guix extends the Guile REPL with additional commands for convenience. Among those, the `build` command comes in handy: it ensures that the given file-like object is built, building it if needed, and returns its output file name(s). In the example below, we build the `coreutils` and `grep` packages, as well as a “computed file” (see Section 8.12 [G-Expressions], page 165), and we use the `scandir` procedure to list the files in Grep's `/bin` directory:

```
scheme@(guix-user)> ,build coreutils
$1 = "/gnu/store/...-coreutils-8.32-debug"
$2 = "/gnu/store/...-coreutils-8.32"
scheme@(guix-user)> ,build grep
$3 = "/gnu/store/...-grep-3.6"
scheme@(guix-user)> ,build (computed-file "x" #~(mkdir #$output))
building /gnu/store/...-x.drv...
$4 = "/gnu/store/...-x"
scheme@(guix-user)> ,use(ice-9 ftw)
scheme@(guix-user)> (scandir (string-append $3 "/bin"))
$5 = ( "."  ".."  "egrep" "fgrep" "grep")
```

As a packager, you may be willing to inspect the build phases or flags of a given package; this is particularly useful when relying a lot on inheritance to define package variants (see Section 8.3 [Defining Package Variants], page 113) or when package arguments are a result of some computation, both of which can make it harder to foresee what ends up in the package arguments. Additional commands let you inspect those package arguments:

```
scheme@(guix-user)> ,phases grep
$1 = (modify-phases %standard-phases
      (add-after 'install 'fix-egrep-and-fgrep
```

```

(lambda* (#:key outputs #:allow-other-keys)
  (let* ((out (assoc-ref outputs "out"))
        (bin (string-append out "/bin")))
    (substitute* (list (string-append bin "/egrep")
                      (string-append bin "/fgrep"))
      ((("^exec grep")
        (string-append "exec " bin "/grep"))))))
scheme@(guix-user)> ,configure-flags findutils
$2 = (list "--localstatedir=/var")
scheme@(guix-user)> ,make-flags binutils
$3 = '("MAKEINFO=true")

```

At a lower-level, a useful command is `lower`: it takes a file-like object and “lowers” it into a derivation (see Section 8.10 [Derivations], page 158) or a store file:

```

scheme@(guix-user)> ,lower grep
$6 = #<derivation /gnu/store/...-grep-3.6.drv => /gnu/store/...-grep-3.6 7f0e639115f0>
scheme@(guix-user)> ,lower (plain-file "x" "Hello!")
$7 = "/gnu/store/...-x"

```

The full list of REPL commands can be seen by typing `,help guix` and is given below for reference.

**build *object*** [REPL command]

Lower *object* and build it if it’s not already built, returning its output file name(s).

**build-options *options*** [REPL command]

Set build options for rest of the REPL commands to *options*. Must be a list of keywords with values accepted by procedure `(@ (guix store) set-build-options)`.

An example would be:

```
,build-options '(:offload? #f)
```

**graft? *graft?*** [REPL command]

Set whether grafts should be performed as part of the build. Equivalent to `--no-grafts` argument to `guix build`.

**lower *object*** [REPL command]

Lower *object* into a derivation or store file name and return it.

**verbosity *level*** [REPL command]

Change build verbosity to *level*.

This is similar to the `--verbosity` command-line option (see Section 9.1.1 [Common Build Options], page 180): level 0 means total silence, level 1 shows build events only, and higher levels print build logs.

**phases *package*** [REPL command]

**configure-flags *package*** [REPL command]

**make-flags *package*** [REPL command]

These REPL commands return the value of one element of the **arguments** field of *package* (see Section 8.2.1 [package Reference], page 104): the first one show the staged code associated with **#:phases** (see Section 8.6 [Build Phases], page 143), the

second shows the code for `#:configure-flags`, and `,make-flags` returns the code for `#:make-flags`.

**run-in-store *exp*** [REPL command]

Run *exp*, a monadic expression, through the store monad. See Section 8.11 [The Store Monad], page 161, for more information.

**enter-store-monad** [REPL command]

Enter a new REPL to evaluate monadic expressions (see Section 8.11 [The Store Monad], page 161). You can quit this “inner” REPL by typing `,q`.

## 9 Utilities

This section describes Guix command-line utilities. Some of them are primarily targeted at developers and users who write new package definitions, while others are more generally useful. They complement the Scheme programming interface of Guix in a convenient way.

### 9.1 Invoking `guix build`

The `guix build` command builds packages or derivations and their dependencies, and prints the resulting store paths. Note that it does not modify the user's profile—this is the job of the `guix package` command (see Section 5.2 [Invoking `guix package`], page 37). Thus, it is mainly useful for distribution developers.

The general syntax is:

```
guix build options package-or-derivation...
```

As an example, the following command builds the latest versions of Emacs and of Guile, displays their build logs, and finally displays the resulting directories:

```
guix build emacs guile
```

Similarly, the following command builds all the available packages:

```
guix build --quiet --keep-going \  
$(guix package -A | awk '{ print $1 "@" $2 }')
```

*package-or-derivation* may be either the name of a package found in the software distribution such as `coreutils` or `coreutils@8.20`, or a derivation such as `/gnu/store/...-coreutils-8.19.drv`. In the former case, a package with the corresponding name (and optionally version) is searched for among the GNU distribution modules (see Section 8.1 [Package Modules], page 100).

Alternatively, the `--expression` option may be used to specify a Scheme expression that evaluates to a package; this is useful when disambiguating among several same-named packages or package variants is needed.

There may be zero or more *options*. The available options are described in the subsections below.

#### 9.1.1 Common Build Options

A number of options that control the build process are common to `guix build` and other commands that can spawn builds, such as `guix package` or `guix archive`. These are the following:

`--load-path=directory`

`-L directory`

Add *directory* to the front of the package module search path (see Section 8.1 [Package Modules], page 100).

This allows users to define their own packages and make them visible to the command-line tools.

`--keep-failed`

`-K` Keep the build tree of failed builds. Thus, if a build fails, its build tree is kept under `/tmp`, in a directory whose name is shown at the end of the build log.



This is useful when debugging build issues. See Section 9.1.4 [Debugging Build Failures], page 193, for tips and tricks on how to debug build issues.

This option implies `--no-offload`, and it has no effect when connecting to a remote daemon with a `guix://` URI (see Section 8.9 [The Store], page 156).

**`--keep-going`**

**`-k`** Keep going when some of the derivations fail to build; return only once all the builds have either completed or failed.

The default behavior is to stop as soon as one of the specified derivations has failed.

**`--dry-run`**

**`-n`** Do not build the derivations.

**`--fallback`**

When substituting a pre-built binary fails, fall back to building packages locally (see Section 5.3.6 [Substitution Failure], page 50).

**`--substitute-urls=urls`**

Consider *urls* the whitespace-separated list of substitute source URLs, overriding the default list of URLs of `guix-daemon` (see [guix-daemon URLs], page 15).

This means that substitutes may be downloaded from *urls*, provided they are signed by a key authorized by the system administrator (see Section 5.3 [Substitutes], page 46).

When *urls* is the empty string, substitutes are effectively disabled.

**`--no-substitutes`**

Do not use substitutes for build products. That is, always build things locally instead of allowing downloads of pre-built binaries (see Section 5.3 [Substitutes], page 46).

**`--no-grafts`**

Do not “graft” packages. In practice, this means that package updates available as grafts are not applied. See Chapter 19 [Security Updates], page 806, for more information on grafts.

**`--rounds=n`**

Build each derivation *n* times in a row, and raise an error if consecutive build results are not bit-for-bit identical.

This is a useful way to detect non-deterministic builds processes. Non-deterministic build processes are a problem because they make it practically impossible for users to *verify* whether third-party binaries are genuine. See Section 9.12 [Invoking guix challenge], page 231, for more.

When used in conjunction with `--keep-failed`, the differing output is kept in the store, under `/gnu/store/...-check`. This makes it easy to look for differences between the two results.

**`--no-offload`**

Do not use offload builds to other machines (see Section 2.2.2 [Daemon Offload Setup], page 9). That is, always build things locally instead of offloading builds to remote machines.

**--max-silent-time=seconds**

When the build or substitution process remains silent for more than *seconds*, terminate it and report a build failure.

By default, the daemon's setting is honored (see Section 2.3 [Invoking guix-daemon], page 14).

**--timeout=seconds**

Likewise, when the build or substitution process lasts for more than *seconds*, terminate it and report a build failure.

By default, the daemon's setting is honored (see Section 2.3 [Invoking guix-daemon], page 14).

**-v level**

**--verbosity=level**

Use the given verbosity *level*, an integer. Choosing 0 means that no output is produced, 1 is for quiet output; 2 is similar to 1 but it additionally displays download URLs; 3 shows all the build log output on standard error.

**--cores=n**

**-c n** Allow the use of up to *n* CPU cores for the build. The special value 0 means to use as many CPU cores as available.

**--max-jobs=n**

**-M n** Allow at most *n* build jobs in parallel. See Section 2.3 [Invoking guix-daemon], page 14, for details about this option and the equivalent *guix-daemon* option.

**--debug=level**

Produce debugging output coming from the build daemon. *level* must be an integer between 0 and 5; higher means more verbose output. Setting a level of 4 or more may be helpful when debugging setup issues with the build daemon.

Behind the scenes, *guix build* is essentially an interface to the *package-derivation* procedure of the (*guix packages*) module, and to the *build-derivations* procedure of the (*guix derivations*) module.

In addition to options explicitly passed on the command line, *guix build* and other *guix* commands that support building honor the *GUIX\_BUILD\_OPTIONS* environment variable.

**GUIX\_BUILD\_OPTIONS**

[Environment Variable]

Users can define this variable to a list of command line options that will automatically be used by *guix build* and other *guix* commands that can perform builds, as in the example below:

```
$ export GUIX_BUILD_OPTIONS="--no-substitutes -c 2 -L /foo/bar"
```

These options are parsed independently, and the result is appended to the parsed command-line options.

### 9.1.2 Package Transformation Options

Another set of command-line options supported by *guix build* and also *guix package* are *package transformation options*. These are options that make it possible to define *package variants*—for instance, packages built from different source code. This is a convenient way

to create customized packages on the fly without having to type in the definitions of package variants (see Section 8.2 [Defining Packages], page 101).

Package transformation options are preserved across upgrades: `guix upgrade` attempts to apply transformation options initially used when creating the profile to the upgraded packages.

The available options are listed below. Most commands support them and also support a `--help-transform` option that lists all the available options and a synopsis (these options are not shown in the `--help` output for brevity).

`--tune [=cpu]`

Use versions of the packages marked as “tunable” optimized for *cpu*. When *cpu* is `native`, or when it is omitted, tune for the CPU on which the `guix` command is running.

Valid *cpu* names are those recognized by the underlying compiler, by default the GNU Compiler Collection. On x86\_64 processors, this includes CPU names such as `nehalem`, `haswell`, and `skylake` (see Section “x86 Options” in *Using the GNU Compiler Collection (GCC)*).

As new generations of CPUs come out, they augment the standard instruction set architecture (ISA) with additional instructions, in particular instructions for single-instruction/multiple-data (SIMD) parallel processing. For example, while Core2 and Skylake CPUs both implement the x86\_64 ISA, only the latter supports AVX2 SIMD instructions.

The primary gain one can expect from `--tune` is for programs that can make use of those SIMD capabilities *and* that do not already have a mechanism to select the right optimized code at run time. Packages that have the `tunable?` property set are considered *tunable packages* by the `--tune` option; a package definition with the property set looks like this:

```
(package
  (name "hello-simd")
  ;; ...

  ;; This package may benefit from SIMD extensions so
  ;; mark it as "tunable".
  (properties '((tunable? . #t))))
```

Other packages are not considered tunable. This allows Guix to use generic binaries in the cases where tuning for a specific CPU is unlikely to provide any gain.

Tuned packages are built with `-march=CPU`; under the hood, the `-march` option is passed to the actual wrapper by a compiler wrapper. Since the build machine may not be able to run code for the target CPU micro-architecture, the test suite is not run when building a tuned package.

To reduce rebuilds to the minimum, tuned packages are *grafted* onto packages that depend on them (see Chapter 19 [Security Updates], page 806). Thus, using `--no-grafts` cancels the effect of `--tune`.

We call this technique *package multi-versioning*: several variants of tunable packages may be built, one for each CPU variant. It is the coarse-grain coun-

terpart of *function multi-versioning* as implemented by the GNU tool chain (see Section “Function Multiversioning” in *Using the GNU Compiler Collection (GCC)*).

#### `--amd-gpu=targets`

Build relevant packages for the AMD GPUs specified by *targets*. Only packages with an `amd-gpu-targets` property are affected. *targets* must be a comma-separated list of AMD GPU target identifiers (<https://llvm.org/docs/AMDGPUUsage.html#amdgpu-processor-table>). For example, the command below builds the GPU code of `rocm-bandwidth-test` for AMD Instinct MI250 (`gfx90a`) and for AMD Instinct MI300 (`gfx942`):

```
guix build --amd-gpu=gfx90a,gfx942 rocm-bandwidth-test
```

To know the identifier(s) of the GPU(s) available on your machine, run:

```
guix shell rocminfo -- rocm_agent_enumerator
```

Under the hood, these GPU identifiers are passed to Clang’s `--offload-arch` option.

When `--amd-gpu` is omitted, packages are built for a default set of GPUs, that we hope covers common needs. Building for more GPUs increases build time significantly and leads to bigger package binaries.

#### `--with-source=source`

#### `--with-source=package=source`

#### `--with-source=package@version=source`

Use *source* as the source of *package*, and *version* as its version number. *source* must be a file name or a URL, as for `guix download` (see Section 9.3 [Invoking `guix download`], page 195).

When *package* is omitted, it is taken to be the package name specified on the command line that matches the base of *source*—e.g., if *source* is `/src/guile-2.0.10.tar.gz`, the corresponding package is `guile`.

Likewise, when *version* is omitted, the version string is inferred from *source*; in the previous example, it is `2.0.10`.

This option allows users to try out versions of packages other than the one provided by the distribution. The example below downloads `ed-1.7.tar.gz` from a GNU mirror and uses that as the source for the `ed` package:

```
guix build ed --with-source=mirror://gnu/ed/ed-1.4.tar.gz
```

As a developer, `--with-source` makes it easy to test release candidates, and even to test their impact on packages that depend on them:

```
guix build elogind --with-source=.../shepherd-0.9.0rc1.tar.gz
```

... or to build from a checkout in a pristine environment:

```
$ git clone git://git.sv.gnu.org/guix.git
```

```
$ guix build guix --with-source=guix@1.0=./guix
```

#### `--with-input=package=replacement`

Replace dependency on *package* by a dependency on *replacement*. *package* must be a package name, and *replacement* must be a package specification such as `guile` or `guile@1.8`.

For instance, the following command builds Guix, but replaces its dependency on the current stable version of Guile with a dependency on the legacy version of Guile, `guile@2.2`:

```
guix build --with-input=guile=guile@2.2 guix
```

This is a recursive, deep replacement. So in this example, both `guix` and its dependency `guile-json` (which also depends on `guile`) get rebuilt against `guile@2.2`.

This is implemented using the `package-input-rewriting/spec` Scheme procedure (see Section 8.2 [Defining Packages], page 101).

#### `--with-graft=package=replacement`

This is similar to `--with-input` but with an important difference: instead of rebuilding the whole dependency chain, *replacement* is built and then *grafted* onto the binaries that were initially referring to *package*. See Chapter 19 [Security Updates], page 806, for more information on grafts.

For example, the command below grafts version 3.5.4 of GnuTLS onto Wget and all its dependencies, replacing references to the version of GnuTLS they currently refer to:

```
guix build --with-graft=gnutls=gnutls@3.5.4 wget
```

This has the advantage of being much faster than rebuilding everything. But there is a caveat: it works if and only if *package* and *replacement* are strictly compatible—for example, if they provide a library, the application binary interface (ABI) of those libraries must be compatible. If *replacement* is somehow incompatible with *package*, then the resulting package may be unusable. Use with care!

#### `--with-debug-info=package`

Build *package* in a way that preserves its debugging info and graft it onto packages that depend on it. This is useful if *package* does not already provide debugging info as a `debug` output (see Chapter 17 [Installing Debugging Files], page 801).

For example, suppose you’re experiencing a crash in Inkscape and would like to see what’s up in GLib, a library deep down in Inkscape’s dependency graph. GLib lacks a `debug` output, so debugging is tough. Fortunately, you rebuild GLib with debugging info and tack it on Inkscape:

```
guix install inkscape --with-debug-info=glib
```

Only GLib needs to be recompiled so this takes a reasonable amount of time. See Chapter 17 [Installing Debugging Files], page 801, for more info.

**Note:** Under the hood, this option works by passing the `#:strip-binaries? #f` to the build system of the package of interest (see Section 8.5 [Build Systems], page 121). Most build systems support that option but some do not. In that case, an error is raised.

Likewise, if a C/C++ package is built without `-g` (which is rarely the case), debugging info will remain unavailable even when `#:strip-binaries?` is false.

**--with-c-toolchain=package=toolchain**

This option changes the compilation of *package* and everything that depends on it so that they get built with *toolchain* instead of the default GNU tool chain for C/C++.

Consider this example:

```
guix build octave-cli \
  --with-c-toolchain=fftw=gcc-toolchain@10 \
  --with-c-toolchain=fftwf=gcc-toolchain@10
```

The command above builds a variant of the `fftw` and `fftwf` packages using version 10 of `gcc-toolchain` instead of the default tool chain, and then builds a variant of the GNU Octave command-line interface using them. GNU Octave itself is also built with `gcc-toolchain@10`.

This other example builds the Hardware Locality (`hwloc`) library and its dependents up to `intel-mpi-benchmarks` with the Clang C compiler:

```
guix build --with-c-toolchain=hwloc=clang-toolchain \
  intel-mpi-benchmarks
```

**Note:** There can be application binary interface (ABI) incompatibilities among tool chains. This is particularly true of the C++ standard library and run-time support libraries such as that of OpenMP. By rebuilding all dependents with the same tool chain, `--with-c-toolchain` minimizes the risks of incompatibility but cannot entirely eliminate them. Choose *package* wisely.

**--with-git-url=package=url**

Build *package* from the latest commit of the `master` branch of the Git repository at *url*. Git sub-modules of the repository are fetched, recursively, if *package* source is not a Git repository, otherwise it depends on the inherited value of `recursive?`.

For example, the following command builds the NumPy Python library against the latest commit of the master branch of Python itself:

```
guix build python-numpy \
  --with-git-url=python=https://github.com/python/cpython
```

This option can also be combined with `--with-branch` or `--with-commit` (see below).

Obviously, since it uses the latest commit of the given branch, the result of such a command varies over time. Nevertheless it is a convenient way to rebuild entire software stacks against the latest commit of one or more packages. This is particularly useful in the context of continuous integration (CI).

Checkouts are kept in a cache under `~/.cache/guix/checkouts` to speed up consecutive accesses to the same repository. You may want to clean it up once in a while to save disk space.

**--with-branch=package=branch**

Build *package* from the latest commit of *branch*. If the `source` field of *package* is an origin with the `git-fetch` method (see Section 8.2.2 [origin Reference], page 109) or a `git-checkout` object, the repository URL is taken from that

source. Otherwise you have to use `--with-git-url` to specify the URL of the Git repository.

For instance, the following command builds `guile-sqlite3` from the latest commit of its `master` branch, and then builds `guix` (which depends on it) and `cuirass` (which depends on `guix`) against this specific `guile-sqlite3` build:

```
guix build --with-branch=guile-sqlite3=master cuirass
```

`--with-commit=package=commit`

This is similar to `--with-branch`, except that it builds from *commit* rather than the tip of a branch. *commit* must be a valid Git commit SHA1 identifier, a tag, or a `git describe` style identifier such as `1.0-3-gabc123`.

`--with-patch=package=file`

Add *file* to the list of patches applied to *package*, where *package* is a spec such as `python@3.8` or `glibc`. *file* must contain a patch; it is applied with the flags specified in the *origin* of *package* (see Section 8.2.2 [origin Reference], page 109), which by default includes `-p1` (see Section “patch Directories” in *Comparing and Merging Files*).

As an example, the command below rebuilds Coreutils with the GNU C Library (glibc) patched with the given patch:

```
guix build coreutils --with-patch=glibc=./glibc-frob.patch
```

In this example, glibc itself as well as everything that leads to Coreutils in the dependency graph is rebuilt.

`--with-configure-flag=package=flag`

Append *flag* to the configure flags of *package*, where *package* is a spec such as `guile@3.0` or `glibc`. The build system of *package* must support the `#:configure-flags` argument.

For example, the command below builds GNU Hello with the configure flag `--disable-nls`:

```
guix build hello --with-configure-flag=hello=--disable-nls
```

The following command passes an extra flag to `cmake` as it builds `lapack`:

```
guix build lapack \
  --with-configure-flag=lapack=-DBUILD_SHARED_LIBS=OFF
```

**Note:** Under the hood, this option works by passing the `#:configure-flags` argument to the build system of the package of interest (see Section 8.5 [Build Systems], page 121). Most build systems support that option but some do not. In that case, an error is raised.

`--with-latest=package`

`--with-version=package=version`

So you like living on the bleeding edge? The `--with-latest` option is for you! It replaces occurrences of *package* in the dependency graph with its latest upstream version, as reported by `guix refresh` (see Section 9.6 [Invoking guix refresh], page 206).

It does so by determining the latest upstream release of *package* (if possible), downloading it, and authenticating it *if* it comes with an OpenPGP signature. As an example, the command below builds Guix against the latest version of Guile-JSON:

```
guix build guix --with-latest=guile-json
```

The `--with-version` works similarly except that it lets you specify that you want precisely *version*, assuming that version exists upstream. For example, to spawn a development environment with SciPy built against version 1.22.4 of NumPy (skipping its test suite because hey, we're not gonna wait this long), you would run:

```
guix shell python python-scipy --with-version=python-numpy=1.22.4■
```

**Warning:** Because they depend on source code published at a given point in time on upstream servers, deployments made with `--with-latest` and `--with-version` may be non-reproducible: source might disappear or be modified in place on the servers.

To deploy old software versions without compromising on reproducibility, see Section 5.8 [Invoking guix time-machine], page 60.

There are limitations. First, in cases where the tool cannot or does not know how to authenticate source code, you are at risk of running malicious code; a warning is emitted in this case. Second, this option simply changes the source used in the existing package definitions, which is not always sufficient: there might be additional dependencies that need to be added, patches to apply, and more generally the quality assurance work that Guix developers normally do will be missing.

You've been warned! When those limitations are acceptable, it's a snappy way to stay on top. We encourage you to submit patches updating the actual package definitions once you have successfully tested an upgrade with `--with-latest` (see Chapter 22 [Contributing], page 816).

#### `--without-tests=package`

Build *package* without running its tests. This can be useful in situations where you want to skip the lengthy test suite of a intermediate package, or if a package's test suite fails in a non-deterministic fashion. It should be used with care because running the test suite is a good way to ensure a package is working as intended.

Turning off tests leads to a different store item. Consequently, when using this option, anything that depends on *package* must be rebuilt, as in this example:

```
guix install --without-tests=python python-notebook
```

The command above installs `python-notebook` on top of `python` built without running its test suite. To do so, it also rebuilds everything that depends on `python`, including `python-notebook` itself.

Internally, `--without-tests` relies on changing the `#:tests?` option of a package's `check` phase (see Section 8.5 [Build Systems], page 121). Note that some packages use a customized `check` phase that does not respect a `#:tests? #f` setting. Therefore, `--without-tests` has no effect on these packages.



Wondering how to achieve the same effect using Scheme code, for example in your manifest, or how to write your own package transformation? See Section 8.3 [Defining Package Variants], page 113, for an overview of the programming interfaces available.

### 9.1.3 Additional Build Options

The command-line options presented below are specific to `guix build`.

`--quiet`

`-q` Build quietly, without displaying the build log; this is equivalent to `--verbosity=0`. Upon completion, the build log is kept in `/var` (or similar) and can always be retrieved using the `--log-file` option.

`--file=file`

`-f file` Build the package, derivation, or other file-like object that the code within *file* evaluates to (see Section 8.12 [G-Expressions], page 165).

As an example, *file* might contain a package definition like this (see Section 8.2 [Defining Packages], page 101):

```
(use-modules (guix)
              (guix build-system gnu)
              (guix licenses))

(package
  (name "hello")
  (version "2.10")
  (source (origin
            (method url-fetch)
            (uri (string-append "mirror://gnu/hello/hello-" version
                                ".tar.gz"))
            (sha256
              (base32
                "0ssi1wpaf7plaswqqjwigppsg5fyh99vdlb9kz17c9lng89ndq1i"))))
  (build-system gnu-build-system)
  (synopsis "Hello, GNU world: An example GNU package")
  (description "Guess what GNU Hello prints!")
  (home-page "http://www.gnu.org/software/hello/")
  (license gpl3+))
```

The *file* may also contain a JSON representation of one or more package definitions. Running `guix build -f` on `hello.json` with the following contents would result in building the packages `myhello` and `greeter`:

```
[
  {
    "name": "myhello",
    "version": "2.10",
    "source": "mirror://gnu/hello/hello-2.10.tar.gz",
    "build-system": "gnu",
    "arguments": {
      "tests?": false
```

```

    },
    "home-page": "https://www.gnu.org/software/hello/",
    "synopsis": "Hello, GNU world: An example GNU package",
    "description": "GNU Hello prints a greeting.",
    "license": "GPL-3.0+",
    "native-inputs": ["gettext"]
  },
  {
    "name": "greeter",
    "version": "1.0",
    "source": "mirror://gnu/hello/hello-2.10.tar.gz",
    "build-system": "gnu",
    "arguments": {
      "test-target": "foo",
      "parallel-build?": false
    },
    "home-page": "https://example.com/",
    "synopsis": "Greeter using GNU Hello",
    "description": "This is a wrapper around GNU Hello.",
    "license": "GPL-3.0+",
    "inputs": ["myhello", "hello"]
  }
]

```

**--manifest=manifest**

**-m manifest**

Build all packages listed in the given *manifest* (see [profile-manifest], page 41).

**--expression=expr**

**-e expr** Build the package or derivation *expr* evaluates to.

For example, *expr* may be `(@ (gnu packages guile) guile-1.8)`, which unambiguously designates this specific variant of version 1.8 of Guile.

Alternatively, *expr* may be a G-expression, in which case it is used as a build program passed to `gexp->derivation` (see Section 8.12 [G-Expressions], page 165).

Lastly, *expr* may refer to a zero-argument monadic procedure (see Section 8.11 [The Store Monad], page 161). The procedure must return a derivation as a monadic value, which is then passed through **run-with-store**.

**--development**

**-D** Build the “development environment” (build dependencies) of the following package.

For example, the following command builds the inputs of **hello**, but *not* **hello** itself, and also builds **guile**:

```
guix build -D hello guile
```

Notice that **-D** (or **--development**) only applies to the immediately following package on the command line. Under the hood, it uses

`package->development-manifest` (see [package-development-manifest], page 120).

**Note:** The effect of combining `--development` with `--target` (for cross-compilation) may not be what you expect: it will cross-compile all the dependencies of the given package when it is built natively.

`--dependents[=depth]`

`-P depth` Build the dependents of the following package. By default, build all the direct and indirect dependents; when *depth* is provided, limit to dependents at that distance: 1 for direct dependents, 2 for dependents of dependents, and so on.

For example, the command below builds *all* the dependents of `libgit2`:

```
guix build --dependents libgit2
```

To build all the packages that directly depend on NumPy, run:

```
guix build -P1 python-numpy
```

The list of dependents is computed in the same way as with `guix refresh --list-dependent` (see Section 9.6 [Invoking guix refresh], page 206).

`--source`

`-S` Build the source derivations of the packages, rather than the packages themselves.

For instance, `guix build -S gcc` returns something like `/gnu/store/...-gcc-4.7.2.tar.bz2`, which is the GCC source tarball.

The returned source tarball is the result of applying any patches and code snippets specified in the package `origin` (see Section 8.2 [Defining Packages], page 101).

As with other derivations, the result of building a source derivation can be verified using the `--check` option (see [build-check], page 192). This is useful to validate that a (potentially already built or substituted, thus cached) package source matches against its declared hash.

Note that `guix build -S` compiles the sources only of the specified packages. They do not include the sources of statically linked dependencies and by themselves are insufficient for reproducing the packages.

`--sources`

Fetch and return the source of *package-or-derivation* and all their dependencies, recursively. This is a handy way to obtain a local copy of all the source code needed to build *packages*, allowing you to eventually build them even without network access. It is an extension of the `--source` option and can accept one of the following optional argument values:

**package** This value causes the `--sources` option to behave in the same way as the `--source` option.

**all** Build the source derivations of all packages, including any source that might be listed as `inputs`. This is the default value.

```
$ guix build --sources tzdata
```

```
The following derivations will be built:
/gnu/store/...-tzdata2015b.tar.gz.drv
/gnu/store/...-tzcode2015b.tar.gz.drv
```

**transitive**

Build the source derivations of all packages, as well of all transitive inputs to the packages. This can be used e.g. to prefetch package source for later offline building.

```
$ guix build --sources=transitive tzdata
The following derivations will be built:
/gnu/store/...-tzcode2015b.tar.gz.drv
/gnu/store/...-findutils-4.4.2.tar.xz.drv
/gnu/store/...-grep-2.21.tar.xz.drv
/gnu/store/...-coreutils-8.23.tar.xz.drv
/gnu/store/...-make-4.1.tar.xz.drv
/gnu/store/...-bash-4.3.tar.xz.drv
...
```

**--system=system**

**-s system** Attempt to build for *system*—e.g., *i686-linux*—instead of the system type of the build host. The **guix build** command allows you to repeat this option several times, in which case it builds for all the specified systems; other commands ignore extraneous **-s** options.

**Note:** The **--system** flag is for *native* compilation and must not be confused with cross-compilation. See **--target** below for information on cross-compilation.

An example use of this is on Linux-based systems, which can emulate different personalities. For instance, passing **--system=i686-linux** on an *x86\_64-linux* system or **--system=armhf-linux** on an *aarch64-linux* system allows you to build packages in a complete 32-bit environment.

**Note:** Building for an *armhf-linux* system is unconditionally enabled on *aarch64-linux* machines, although certain *aarch64* chipsets do not allow for this functionality, notably the ThunderX.

Similarly, when transparent emulation with QEMU and **binfmt\_misc** is enabled (see Section 11.10.32 [Virtualization Services], page 581), you can build for any system for which a QEMU **binfmt\_misc** handler is installed.

Builds for a system other than that of the machine you are using can also be offloaded to a remote machine of the right architecture. See Section 2.2.2 [Daemon Offload Setup], page 9, for more information on offloading.

**--target=triplet**

Cross-build for *triplet*, which must be a valid GNU triplet, such as "*aarch64-linux-gnu*" (see Section "Specifying Target Triplets" in *Autoconf*).

**--list-systems**

List all the supported systems, that can be passed as an argument to **--system**.

**--list-targets**

List all the supported targets, that can be passed as an argument to **--target**.

**--check** Rebuild *package-or-derivation*, which are already available in the store, and raise an error if the build results are not bit-for-bit identical.

This mechanism allows you to check whether previously installed substitutes are genuine (see Section 5.3 [Substitutes], page 46), or whether the build result of a package is deterministic. See Section 9.12 [Invoking guix challenge], page 231, for more background information and tools.

When used in conjunction with **--keep-failed**, the differing output is kept in the store, under `/gnu/store/...-check`. This makes it easy to look for differences between the two results.

**--repair** Attempt to repair the specified store items, if they are corrupt, by re-downloading or rebuilding them.

This operation is not atomic and thus restricted to **root**.

**--derivations**

**-d** Return the derivation paths, not the output paths, of the given packages.

**--root=file**

**-r file** Make *file* a symlink to the result, and register it as a garbage collector root.

Consequently, the results of this **guix build** invocation are protected from garbage collection until *file* is removed. When that option is omitted, build results are eligible for garbage collection as soon as the build completes. See Section 5.6 [Invoking guix gc], page 53, for more on GC roots.

**--log-file**

Return the build log file names or URLs for the given *package-or-derivation*, or raise an error if build logs are missing.

This works regardless of how packages or derivations are specified. For instance, the following invocations are equivalent:

```
guix build --log-file $(guix build -d guile)
guix build --log-file $(guix build guile)
guix build --log-file guile
guix build --log-file -e '(@ (gnu packages guile) guile-2.0)'
```

If a log is unavailable locally, and unless **--no-substitutes** is passed, the command looks for a corresponding log on one of the substitute servers.

So for instance, imagine you want to see the build log of GDB on **aarch64**, but you are actually on an **x86\_64** machine:

```
$ guix build --log-file gdb -s aarch64-linux
https://bordeaux.guix.gnu.org/log/...-gdb-7.10
```

You can freely access a huge library of build logs!

### 9.1.4 Debugging Build Failures

When defining a new package (see Section 8.2 [Defining Packages], page 101), you will probably find yourself spending some time debugging and tweaking the build until it succeeds. To do that, you need to operate the build commands yourself in an environment as close as possible to the one the build daemon uses.

To that end, the first thing to do is to use the `--keep-failed` or `-K` option of `guix build`, which will keep the failed build tree in `/tmp` or whatever directory you specified as `TMPDIR` (see Section 9.1.1 [Common Build Options], page 180).

From there on, you can `cd` to the failed build tree and source the `environment-variables` file, which contains all the environment variable definitions that were in place when the build failed. So let's say you're debugging a build failure in package `foo`; a typical session would look like this:

```
$ guix build foo -K
... build fails
$ cd /tmp/guix-build-foo.drv-0
$ source ./environment-variables
$ cd foo-1.2
```

Now, you can invoke commands as if you were the daemon (almost) and troubleshoot your build process.

Sometimes it happens that, for example, a package's tests pass when you run them manually but they fail when the daemon runs them. This can happen because the daemon runs builds in containers where, unlike in our environment above, network access is missing, `/bin/sh` does not exist, etc. (see Section 2.2.1 [Build Environment Setup], page 6).

In such cases, you may need to inspect the build process from within a container similar to the one the build daemon creates:

```
$ guix build -K foo
...
$ cd /tmp/guix-build-foo.drv-0
$ guix shell --no-grafts -C -D foo strace gdb
[env]# source ./environment-variables
[env]# cd foo-1.2
```

Here, `guix shell -C` creates a container and spawns a new shell in it (see Section 7.1 [Invoking guix shell], page 78). The `strace gdb` part adds the `strace` and `gdb` commands to the container, which you may find handy while debugging. The `--no-grafts` option makes sure we get the exact same environment, with ungrafted packages (see Chapter 19 [Security Updates], page 806, for more info on grafts).

To get closer to a container like that used by the build daemon, we can remove `/bin/sh` (you'll first need to pass the `--writable-root` option to `guix shell`):

```
[env]# rm /bin/sh
```

(Don't worry, this is harmless: this is all happening in the throw-away container created by `guix shell`.)

The `strace` command is probably not in the search path, but we can run:

```
[env]# $GUIX_ENVIRONMENT/bin/strace -f -o log make check
```

In this way, not only you will have reproduced the environment variables the daemon uses, you will also be running the build process in a container similar to the one the daemon uses.

## 9.2 Invoking guix edit

So many packages, so many source files! The `guix edit` command facilitates the life of users and packagers by pointing their editor at the source file containing the definition of the specified packages. For instance:

```
guix edit gcc@4.9 vim
```

launches the program specified in the `VISUAL` or in the `EDITOR` environment variable to view the recipe of GCC 4.9.3 and that of Vim.

If you are using a Guix Git checkout (see Section 22.2 [Building from Git], page 817), or have created your own packages on `GUIX_PACKAGE_PATH` (see Section 8.1 [Package Modules], page 100), you will be able to edit the package recipes. In other cases, you will be able to examine the read-only recipes for packages currently in the store.

Instead of `GUIX_PACKAGE_PATH`, the command-line option `--load-path=directory` (or in short `-L directory`) allows you to add *directory* to the front of the package module search path and so make your own packages visible.

## 9.3 Invoking guix download

When writing a package definition, developers typically need to download a source tarball, compute its SHA256 hash, and write that hash in the package definition (see Section 8.2 [Defining Packages], page 101). The `guix download` tool helps with this task: it downloads a file from the given URI, adds it to the store, and prints both its file name in the store and its SHA256 hash.

The fact that the downloaded file is added to the store saves bandwidth: when the developer eventually tries to build the newly defined package with `guix build`, the source tarball will not have to be downloaded again because it is already in the store. It is also a convenient way to temporarily stash files, which may be deleted eventually (see Section 5.6 [Invoking guix gc], page 53).

The `guix download` command supports the same URIs as used in package definitions. In particular, it supports `mirror://` URIs. `https` URIs (HTTP over TLS) are supported *provided* the Guile bindings for GnuTLS are available in the user’s environment; when they are not available, an error is raised. See Section “Guile Preparations” in *GnuTLS-Guile*, for more information.

`guix download` verifies HTTPS server certificates by loading the certificates of X.509 authorities from the directory pointed to by the `SSL_CERT_DIR` environment variable (see Section 11.12 [X.509 Certificates], page 679), unless `--no-check-certificate` is used.

Alternatively, `guix download` can also retrieve a Git repository, possibly a specific commit, tag, or branch.

The following options are available:

`--hash=algorithm`

`-H algorithm`

Compute a hash using the specified *algorithm*. See Section 9.4 [Invoking guix hash], page 196, for more information.

`--format=fmt`

`-f fmt` Write the hash in the format specified by *fmt*. For more information on the valid values for *fmt*, see Section 9.4 [Invoking guix hash], page 196.

**--no-check-certificate**

Do not validate the X.509 certificates of HTTPS servers.

When using this option, you have *absolutely no guarantee* that you are communicating with the authentic server responsible for the given URL, which makes you vulnerable to “man-in-the-middle” attacks.

**--output=file**

**-o file** Save the downloaded file to *file* instead of adding it to the store.

**--git**

**-g** Checkout the Git repository at the latest commit on the default branch.

**--commit=commit-or-tag**

Checkout the Git repository at *commit-or-tag*.

*commit-or-tag* can be either a tag or a commit defined in the Git repository.

**--branch=branch**

Checkout the Git repository at *branch*.

The repository will be checked out at the latest commit of *branch*, which must be a valid branch of the Git repository.

**--recursive**

**-r** Recursively clone the Git repository.

## 9.4 Invoking guix hash

The **guix hash** command computes the hash of a file. It is primarily a convenience tool for anyone contributing to the distribution: it computes the cryptographic hash of one or more files, which can be used in the definition of a package (see Section 8.2 [Defining Packages], page 101).

The general syntax is:

```
guix hash option file ...
```

When *file* is - (a hyphen), **guix hash** computes the hash of data read from standard input. **guix hash** has the following options:

**--hash=algorithm****-H algorithm**

Compute a hash using the specified *algorithm*, **sha256** by default.

*algorithm* must be the name of a cryptographic hash algorithm supported by Libgcrypt *via* Guile-Gcrypt—e.g., **sha512** or **sha3-256** (see Section “Hash Functions” in *Guile-Gcrypt Reference Manual*).

**--format=fmt**

**-f fmt** Write the hash in the format specified by *fmt*.

Supported formats: **base64**, **nix-base32**, **base32**, **base16** (**hex** and **hexadecimal** can be used as well).

If the **--format** option is not specified, **guix hash** will output the hash in **nix-base32**. This representation is used in the definitions of packages.



**--recursive**

**-r**        The **--recursive** option is deprecated in favor of **--serializer=nar** (see below); **-r** remains accepted as a convenient shorthand.

**--serializer=type**

**-S type**    Compute the hash on *file* using *type* serialization.  
*type* may be one of the following:

**none**        This is the default: it computes the hash of a file’s contents.

**nar**        Compute the hash of a “normalized archive” (or “nar”) containing *file*, including its children if it is a directory. Some of the metadata of *file* is part of the archive; for instance, when *file* is a regular file, the hash is different depending on whether *file* is executable or not. Metadata such as time stamps have no impact on the hash (see Section 5.11 [Invoking guix archive], page 65, for more info on the nar format).

**git**        Compute the hash of the file or directory as a Git “tree”, following the same method as the Git version control system.

**--exclude-vcs**

**-x**        When combined with **--recursive**, exclude version control system directories (*.bzip*, *.git*, *.hg*, etc.).

As an example, here is how you would compute the hash of a Git checkout, which is useful when using the *git-fetch* method (see Section 8.2.2 [origin Reference], page 109):

```
$ git clone http://example.org/foo.git
$ cd foo
$ guix hash -x --serializer=nar .
```

## 9.5 Invoking guix import

The *guix import* command is useful for people who would like to add a package to the distribution with as little work as possible—a legitimate demand. The command knows of a few repositories from which it can “import” package metadata. The result is a package definition, or a template thereof, in the format we know (see Section 8.2 [Defining Packages], page 101).

The general syntax is:

```
guix import [global-options...] importer package [options...]
```

*importer* specifies the source from which to import package metadata, and *options* specifies a package identifier and other options specific to *importer*. *guix import* itself has the following *global-options*:

**--insert=file**

**-i file**    Insert the package definition(s) that the *importer* generated into the specified *file*, either in alphabetical order among existing package definitions, or at the end of the file otherwise.

Some of the importers rely on the ability to run the `gpgv` command. For these, GnuPG must be installed and in `$PATH`; run `guix install gnupg` if needed.

Currently, the available “importers” are:

**gnu** Import metadata for the given GNU package. This provides a template for the latest version of that GNU package, including the hash of its source tarball, and its canonical synopsis and description.

Additional information such as the package dependencies and its license needs to be figured out manually.

For example, the following command returns a package definition for GNU Hello:

```
guix import gnu hello
```

Specific command-line options are:

```
--key-download=policy
```

As for `guix refresh`, specify the policy to handle missing OpenPGP keys when verifying the package signature. See Section 9.6 [Invoking `guix refresh`], page 206.

**pypi** Import metadata from the Python Package Index (<https://pypi.python.org/>). Information is taken from the JSON-formatted description available at `pypi.python.org` and usually includes all the relevant information, including package dependencies. For maximum efficiency, it is recommended to install the `unzip` utility, so that the importer can unzip Python wheels and gather data from them.

The command below imports metadata for the latest version of the `itsdangerous` Python package:

```
guix import pypi itsdangerous
```

You can also ask for a specific version:

```
guix import pypi itsdangerous@1.1.0
```

```
--recursive
```

```
-r
```

Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**gem** Import metadata from RubyGems (<https://rubygems.org/>). Information is taken from the JSON-formatted description available at `rubygems.org` and includes most relevant information, including runtime dependencies. There are some caveats, however. The metadata doesn’t distinguish between synopses and descriptions, so the same string is used for both fields. Additionally, the details of non-Ruby dependencies required to build native extensions is unavailable and left as an exercise to the packager.

The command below imports metadata for the `rails` Ruby package:

```
guix import gem rails
```

You can also ask for a specific version:

```
guix import gem rails@7.0.4
```

**--recursive**

**-r** Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**nuget** Import metadata from NuGet (<https://www.nuget.org/>), the package manager for .NET. Information is taken from the JSON-formatted metadata provided through NuGet's v3 API at [api.nuget.org](https://api.nuget.org) and includes most relevant information, including package dependencies. There are some caveats, however. The metadata does not always include repository information, in which case the importer attempts to extract it from the symbol package (`.snupkg`) if available. Additionally, dependencies are grouped by target framework in NuGet, but the importer flattens all dependency groups into a single list.

The command below imports metadata for the Avalonia .NET package:

```
guix import nuget Avalonia
```

You can also recursively import all dependencies:

```
guix import nuget -r Avalonia
```

**--archive=repo**

**-a repo** Specify the archive repository. Currently only **nuget** is supported, which uses the official NuGet package repository at [nuget.org](https://www.nuget.org).

**--recursive**

**-r** Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**--license-prefix=prefix**

**-p prefix** Add a custom prefix to license identifiers in the generated package definitions. This can be useful when license identifiers need to be qualified with a module name.

**minetest** Import metadata from ContentDB (<https://content.minetest.net>). Information is taken from the JSON-formatted metadata provided through ContentDB's API (<https://content.minetest.net/help/api/>) and includes most relevant information, including dependencies. There are some caveats, however. The license information is often incomplete. The commit hash is sometimes missing. The descriptions are in the Markdown format, but Guix uses Texinfo instead. Texture packs and subgames are unsupported.

The command below imports metadata for the Mesecons mod by Jeija:

```
guix import minetest Jeija/mesecons
```

The author name can also be left out:

```
guix import minetest mesecons
```

**--recursive**

**-r** Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**cpan** Import metadata from MetaCPAN (<https://www.metacpan.org/>). Information is taken from the JSON-formatted metadata provided through MetaCPAN's API (<https://fastapi.metacpan.org/>) and includes most relevant information, such as module dependencies. License information should be checked closely. If Perl is available in the store, then the **corelist** utility will be used to filter core modules out of the list of dependencies.

The command below imports metadata for the **Acme::Boolean** Perl module:

```
guix import cpan Acme::Boolean
```

Like many other importers, the **cpan** importer supports recursive imports:

```
--recursive
```

```
-r      Traverse the dependency graph of the given upstream package re-
        cursively and generate package expressions for all those packages
        that are not yet in Guix.
```

**cran** Import metadata from CRAN (<https://cran.r-project.org/>), the central repository for the GNU R statistical and graphical environment (<https://r-project.org>).

Information is extracted from the **DESCRIPTION** file of the package.

The command below imports metadata for the **Cairo** R package:

```
guix import cran Cairo
```

You can also ask for a specific version:

```
guix import cran rasterVis@0.50.3
```

When **--recursive** is added, the importer will traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

When **--style=specification** is added, the importer will generate package definitions whose inputs are package specifications instead of references to package variables. This is useful when generated package definitions are to be appended to existing user modules, as the list of used package modules need not be changed. The default is **--style=variable**.

When **--prefix=license:** is added, the importer will prefix license atoms with **license:**, allowing a prefixed import of (**guix licenses**).

When **--archive=bioconductor** is added, metadata is imported from Bioconductor (<https://www.bioconductor.org/>), a repository of R packages for the analysis and comprehension of high-throughput genomic data in bioinformatics.

Information is extracted from the **DESCRIPTION** file contained in the package archive.

The command below imports metadata for the **GenomicRanges** R package:

```
guix import cran --archive=bioconductor GenomicRanges
```

Finally, you can also import R packages that have not yet been published on CRAN or Bioconductor as long as they are in a git repository. Use **--archive=git** followed by the URL of the git repository:

```
guix import cran --archive=git https://github.com/immunogenomics/harmony
```

**texlive** Import TeX package information from the TeX Live package database for TeX packages that are part of the TeX Live distribution (<https://www.tug.org/texlive/>).

Information about the package is obtained from the TeX Live package database, a plain text file that is included in the **texlive-scripts** package. The source code is downloaded from possibly multiple locations in the SVN repository of the TeX Live project. Note that therefore SVN must be installed and in **\$PATH**; run **guix install subversion** if needed.

The command below imports metadata for the **fontspec** TeX package:

```
guix import texlive fontspec
```

Additional options include:

```
--recursive
```

```
-r          Traverse the dependency graph of the given upstream package re-
            cursively and generate package expressions for all those packages
            that are not yet in Guix.
```

**json** Import package metadata from a local JSON file. Consider the following example package definition in JSON format:

```
{
  "name": "hello",
  "version": "2.10",
  "source": "mirror://gnu/hello/hello-2.10.tar.gz",
  "build-system": "gnu",
  "home-page": "https://www.gnu.org/software/hello/",
  "synopsis": "Hello, GNU world: An example GNU package",
  "description": "GNU Hello prints a greeting.",
  "license": "GPL-3.0+",
  "native-inputs": ["gettext"]
}
```

The field names are the same as for the **<package>** record (See Section 8.2 [Defining Packages], page 101). References to other packages are provided as JSON lists of quoted package specification strings such as **guile** or **guile02.0**.

The importer also supports a more explicit source definition using the common fields for **<origin>** records:

```
{
  ...
  "source": {
    "method": "url-fetch",
    "uri": "mirror://gnu/hello/hello-2.10.tar.gz",
    "sha256": {
      "base32": "0ssi1wpaf7plaswqqjwigppsg5fyh99vdlb9kzl7c9lng89ndq1i"
    }
  }
  ...
}
```

The command below reads metadata from the JSON file `hello.json` and outputs a package expression:

```
guix import json hello.json
```

**hackage** Import metadata from the Haskell community’s central package archive Hackage (<https://hackage.haskell.org/>). Information is taken from Cabal files and includes all the relevant information, including package dependencies.

Specific command-line options are:

```
--stdin
-s          Read a Cabal file from standard input.

--no-test-dependencies
-t          Do not include dependencies required only by the test suites.

--cabal-environment=alist
-e alist    alist is a Scheme alist defining the environment in which the Cabal
            conditionals are evaluated. The accepted keys are: os, arch, impl
            and a string representing the name of a flag. The value associated
            with a flag has to be either the symbol true or false. The value
            associated with other keys has to conform to the Cabal file format
            definition. The default value associated with the keys os, arch and
            impl is ‘linux’, ‘x86_64’ and ‘ghc’, respectively.

--recursive
-r          Traverse the dependency graph of the given upstream package re-
            cursively and generate package expressions for all those packages
            that are not yet in Guix.
```

The command below imports metadata for the latest version of the HTTP Haskell package without including test dependencies and specifying the value of the flag ‘`network-uri`’ as `false`:

```
guix import hackage -t -e '"(\\\"network-uri\\\" . false))" HTTP
```

A specific package version may optionally be specified by following the package name by an at-sign and a version number as in the following example:

```
guix import hackage mtl@2.1.3.1
```

**stackage** The `stackage` importer is a wrapper around the `hackage` one. It takes a package name, looks up the package version included in a long-term support (LTS) Stackage (<https://www.stackage.org>) release and uses the `hackage` importer to retrieve its metadata. Note that it is up to you to select an LTS release compatible with the GHC compiler used by Guix.

Specific command-line options are:

```
--no-test-dependencies
-t          Do not include dependencies required only by the test suites.

--lts-version=version
-l version    version is the desired LTS release version. If omitted the latest
            release is used.
```

**--recursive**

**-r** Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

The command below imports metadata for the HTTP Haskell package included in the LTS Stackage release version 7.18:

```
guix import stackage --lts-version=7.18 HTTP
```

**elpa** Import metadata from an Emacs Lisp Package Archive (ELPA) package repository (see Section “Packages” in *The GNU Emacs Manual*).

Specific command-line options are:

**--archive=repo**

**-a repo** *repo* identifies the archive repository from which to retrieve the information. Currently the supported repositories and their identifiers are:

- GNU (<https://elpa.gnu.org/packages>), selected by the **gnu** identifier. This is the default.

Packages from [elpa.gnu.org](https://elpa.gnu.org) are signed with one of the keys contained in the GnuPG keyring at `share/emacs/25.1/etc/package-keyring.gpg` (or similar) in the **emacs** package (see Section “Package Installation” in *The GNU Emacs Manual*).

- GNU-devel (<https://elpa.gnu.org/devel>), selected by the **gnu-devel** identifier. This repository exists for testing development changes before they make it to a stable release.
- NonGNU (<https://elpa.nongnu.org/nongnu/>), selected by the **nongnu** identifier.
- MELPA-Stable (<https://stable.melpa.org/packages>), selected by the **melpa-stable** identifier.
- MELPA (<https://melpa.org/packages>), selected by the **melpa** identifier.

**--recursive**

**-r** Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**crate** Import metadata from the crates.io Rust package repository `crates.io` (<https://crates.io>), as in this example:

```
guix import crate blake2-rfc
```

The crate importer also allows you to specify a version string:

```
guix import crate constant-time-eq@0.1.0
```

Additional options include:

**--allow-yanked**

If no non-yanked version of a crate is available, use the latest yanked version instead instead of aborting.

`--lockfile=file`

`-f file` When `--lockfile` is specified, the importer will ignore other options and won't output package expressions, instead importing source expressions from *file*, a `Cargo.lock` file.

See Section “Packaging Rust Crates” in *GNU Guix Cookbook*, for a packaging workflow utilizing it.

**elm** Import metadata from the Elm package repository `package.elm-lang.org` (<https://package.elm-lang.org>), as in this example:

```
guix import elm elm-explorations/webgl
```

The Elm importer also allows you to specify a version string:

```
guix import elm elm-explorations/webgl@1.1.3
```

Additional options include:

`--recursive`

`-r` Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**npm-binary**

Import metadata from the npm Registry (<https://registry.npmjs.org>), as in this example:

```
guix import npm-binary buffer-crc32
```

The npm-binary importer also allows you to specify a version string:

```
guix import npm-binary buffer-crc32@1.0.0
```

**Note:** Generated package expressions skip the build step of the `node-build-system`. As such, generated package expressions often refer to transpiled or generated files, instead of being built from source.

Additional options include:

`--recursive`

`-r` Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**opam** Import metadata from the OPAM (<https://opam.ocaml.org/>) package repository used by the OCaml community.

Additional options include:

`--recursive`

`-r` Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**composer** Import metadata from the Composer (<https://getcomposer.org/>) package archive used by the PHP community, as in this example:

```
guix import composer phpunit/phpunit
```



Additional options include:

**--recursive**

**-r** Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**--repo** By default, packages are searched in the official OPAM repository. This option, which can be used more than once, lets you add other repositories which will be searched for packages. It accepts as valid arguments:

- the name of a known repository - can be one of `opam`, `coq` (equivalent to `coq-released`), `coq-core-dev`, `coq-extra-dev` or `grew`.
- the URL of a repository as expected by the `opam repository add` command (for instance, the URL equivalent of the above `opam` name would be `https://opam.ocaml.org`).
- the path to a local copy of a repository (a directory containing a `packages/` sub-directory).

Repositories are assumed to be passed to this option by order of preference. The additional repositories will not replace the default `opam` repository, which is always kept as a fallback.

Also, please note that versions are not compared across repositories. The first repository (from left to right) that has at least one version of a given package will prevail over any others, and the version imported will be the latest one found *in this repository only*.

**go** Import metadata for a Go module using proxy.golang.org (`https://proxy.golang.org`).

```
guix import go gopkg.in/yaml.v2
```

It is possible to use a package specification with a `@VERSION` suffix to import a specific version.

Additional options include:

**--recursive**

**-r** Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**--pin-versions**

When using this option, the importer preserves the exact versions of the Go modules dependencies instead of using their latest available versions. This can be useful when attempting to import packages that recursively depend on former versions of themselves to build. When using this mode, the symbol of the package is made by appending the version to its name, so that multiple versions of the same package can coexist.

**egg** Import metadata for CHICKEN eggs (<https://wiki.call-cc.org/eggs>). The information is taken from `PACKAGE.egg` files found in the `eggs-5-all` ([git://code.call-cc.org/eggs-5-all](https://code.call-cc.org/eggs-5-all)) Git repository. However, it does not provide all the information that we need, there is no “description” field, and the licenses used are not always precise (BSD is often used instead of BSD-N).

```
guix import egg sourcehut
```

You can also ask for a specific version:

```
guix import egg arrays@1.0
```

Additional options include:

```
--recursive
```

```
-r
```

Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

**hexpm** Import metadata from the hex.pm Erlang and Elixir package repository `hex.pm` (<https://hex.pm>), as in this example:

```
guix import hexpm stun
```

The importer tries to determine the build system used by the package.

The hexpm importer also allows you to specify a version string:

```
guix import hexpm cf@0.3.0
```

Additional options include:

```
--recursive
```

```
-r
```

Traverse the dependency graph of the given upstream package recursively and generate package expressions for all those packages that are not yet in Guix.

The structure of the `guix import` code is modular. It would be useful to have more importers for other package formats, and your help is welcome here (see Chapter 22 [Contributing], page 816).

## 9.6 Invoking `guix refresh`

The primary audience of the `guix refresh` command is packagers. As a user, you may be interested in the `--with-latest` option, which can bring you package update superpowers built upon `guix refresh` (see Section 9.1.2 [Package Transformation Options], page 182). By default, `guix refresh` reports any packages provided by the distribution that are outdated compared to the latest upstream version, like this:

```
$ guix refresh
gnu/packages/gettext.scm:29:13: gettext would be upgraded from 0.18.1.1 to 0.18.2.1
gnu/packages/glib.scm:77:12: glib would be upgraded from 2.34.3 to 2.37.0
```

Alternatively, one can specify packages to consider, in which case a warning is emitted for packages that lack an updater:

```
$ guix refresh coreutils guile guile-ssh
gnu/packages/ssh.scm:205:2: warning: no updater for guile-ssh
gnu/packages/guile.scm:136:12: guile would be upgraded from 2.0.12 to 2.0.13
```

`guix refresh` browses the upstream repository of each package and determines the highest version number of the releases therein. The command knows how to update specific types of packages: GNU packages, ELPA packages, etc.—see the documentation for `--type` below. There are many packages, though, for which it lacks a method to determine whether a new upstream release is available. However, the mechanism is extensible, so feel free to get in touch with us to add a new method!

If for some reason you don't want to update to the latest version, you can update to a specific version by appending an equal sign and the desired version number to the package specification. Note that not all updaters support this; an error is reported when an updater cannot refresh to the specified version.

```
$ guix refresh -u guile=3.0.4
...
gnu/packages/guile.scm:392:2: guile: updating from version 3.0.3 to version 3.0.4...
...
$ guix refresh -u guile@2.0=2.0.12
...
gnu/packages/guile.scm:147:2: guile: updating from version 2.0.10 to version 2.0.12...
...
```

In some specific cases, you may have many packages specified via a manifest or a module selection which should all be updated together; for these cases, the `--target-version` option can be provided to have them all refreshed to the same version, as shown in the examples below:

```
$ guix refresh qtbase qtdeclarative --target-version=6.5.2
gnu/packages/qt.scm:1248:13: qtdeclarative would be upgraded from 6.3.2 to 6.5.2
gnu/packages/qt.scm:584:2: qtbase would be upgraded from 6.3.2 to 6.5.2
$ guix refresh --manifest=qt5-manifest.scm --target-version=5.15.10
gnu/packages/qt.scm:1173:13: qtxmlpatterns would be upgraded from 5.15.8 to 5.15.10
gnu/packages/qt.scm:1202:13: qtdeclarative would be upgraded from 5.15.8 to 5.15.10
gnu/packages/qt.scm:1762:13: qtserialbus would be upgraded from 5.15.8 to 5.15.10
gnu/packages/qt.scm:2070:13: qtquickcontrols2 would be upgraded from 5.15.8 to 5.15.10
...
$ guix refresh -s module:kde-plasma --target-version=6.3.6
[...] plasma-systemmonitor would be upgraded from 6.3.4 to 6.3.6
[...] plasma-workspace-wallpapers would be upgraded from 6.3.4 to 6.3.6
[...] plasma-browser-integration would be upgraded from 6.3.4 to 6.3.6
...
```

The `--target-version` option accepts partial version prefixes, which can be useful to update to the latest major or major-minor prefixed version:

```
$ guix refresh qtbase@5 qtdeclarative@5 --target-version=5
gnu/packages/qt.scm:1472:13: qtdeclarative would be upgraded from 5.15.8 to 5.15.10
gnu/packages/qt.scm:452:13: qtbase would be upgraded from 5.15.8 to 5.15.10
```

A per-package equivalent to `--target-version` is made available by prefixing the version specification with the `~` (tilde) character. For example:

```
$ guix refresh bash=~5 guile=~3
gnu/packages/guile.scm:354:13: guile would be upgraded from 3.0.9 to 3.0.10
```

```
gnu/packages/bash.scm:150:15: bash would be upgraded from 5.1.16 to 5.2.37■
```

Sometimes the upstream name differs from the package name used in Guix, and `guix refresh` needs a little help. Most updaters honor the `upstream-name` property in package definitions, which can be used to that effect:

```
(define-public network-manager
  (package
    (name "network-manager")
    ;; ...
    (properties '((upstream-name . "NetworkManager")))))
```

When passed `--update`, it modifies distribution source files to update the version numbers and source code hashes of those package definitions, as well as possibly their inputs (see Section 8.2 [Defining Packages], page 101). This is achieved by downloading each package’s latest source tarball and its associated OpenPGP signature, authenticating the downloaded tarball against its signature using `gpgv`, and finally computing its hash—note that GnuPG must be installed and in `$PATH`; run `guix install gnupg` if needed.

You might also want to update definitions that correspond to the packages installed in your profile:

```
$ ./pre-inst-env guix refresh -u \
    $(guix package --list-installed | cut -f1)
```

When considering whether to upgrade a package, it is sometimes convenient to know which packages would be affected by the upgrade and should be checked for compatibility. For this the `--list-dependent` option may be used when passing `guix refresh` one or more package names.

When the public key used to sign the tarball is missing from the user’s keyring, an attempt is made to automatically retrieve it from a public key server; when this is successful, the key is added to the user’s keyring; otherwise, `guix refresh` reports an error.

The following options are supported:

#### `--recursive`

Consider the packages specified, and all the packages upon which they depend.

```
$ guix refresh --recursive coreutils
gnu/packages/acl.scm:40:13: acl would be upgraded from 2.2.53 to 2.3.1■
gnu/packages/m4.scm:30:12: 1.4.18 is already the latest version of m4■
gnu/packages/xml.scm:68:2: warning: no updater for expat
gnu/packages/multiprecision.scm:40:12: 6.1.2 is already the latest version o
...
```

#### `--expression=expr`

`-e expr` Consider the package `expr` evaluates to.

This is useful to precisely refer to a package, as in this example:

```
guix refresh -l -e '(@@ (gnu packages commencement) glibc-final)'■
```

This command lists the dependents of the “final” libc (essentially all the packages).

**--update**

**-u** Update distribution source files (package definitions) in place. This is usually run from a checkout of the Guix source tree (see Section 22.4 [Running Guix Before It Is Installed], page 821):

```
./pre-inst-env guix refresh -s non-core -u
```

See Section 8.2 [Defining Packages], page 101, for more information on package definitions. You can also run it on packages from a third-party channel:

```
guix refresh -L /path/to/channel -u package
```

See Section 6.7 [Creating a Channel], page 72, on how to create a channel.

This command updates the version and source code hash of the package. Depending on the updater being used, it can also update the various ‘inputs’ fields of the package. In some cases, the updater might get inputs wrong—it might not know about an extra input that’s necessary, or it might add an input that should be avoided.

To address that, packagers can add properties stating inputs that should be added to those found by the updater or inputs that should be ignored: the `updater-extra-inputs` and `updater-ignored-inputs` properties pertain to “regular” inputs, and there are equivalent properties for ‘native’ and ‘propagated’ inputs. In the example below, we tell the updater that we need ‘openmpi’ as an additional input:

```
(define-public python-mpi4py
  (package
    (name "python-mpi4py")
    ;; ...
    (inputs (list openmpi))
    (properties
      '((updater-extra-inputs . ("openmpi"))))))
```

That way, `guix refresh -u python-mpi4py` will leave the ‘openmpi’ input, even if it is not among the inputs it would normally add.

**--select=[subset]**

**-s subset** Select all the packages in *subset*, one of `core`, `non-core` or `module:name`.

The `core` subset refers to all the packages at the core of the distribution—i.e., packages that are used to build “everything else”. This includes GCC, libc, Binutils, Bash, etc. Usually, changing one of these packages in the distribution entails a rebuild of all the others. Thus, such updates are an inconvenience to users in terms of build time or bandwidth used to achieve the upgrade.

The `non-core` subset refers to the remaining packages. It is typically useful in cases where an update of the core packages would be inconvenient.

The `module:name` subset refers to all the packages in a specified guile module. The module can be specified as `module:guile` or `module:(gnu packages guile)`, the former is a shorthand for the later.

**--manifest=file**

**-m file** Select all the packages from the manifest in *file*. This is useful to check if any packages of the user manifest can be updated.

`--type=updater`

`-t updater`

Select only packages handled by *updater* (may be a comma-separated list of updaters). Currently, *updater* may be one of:

- `gnu` the updater for GNU packages;
- `savannah` the updater for packages hosted at Savannah (<https://savannah.gnu.org>);
- `sourceforge` the updater for packages hosted at SourceForge (<https://sourceforge.net>);
- `gnome` the updater for GNOME packages;
- `kde` the updater for KDE packages;
- `xorg` the updater for X.org packages;
- `kernel.org` the updater for packages hosted on kernel.org;
- `egg` the updater for Egg (<https://wiki.call-cc.org/eggs/>) packages;
- `elpa` the updater for ELPA (<https://elpa.gnu.org/>) packages;
- `cran` the updater for CRAN (<https://cran.r-project.org/>) packages;
- `bioconductor` the updater for Bioconductor (<https://www.bioconductor.org/>) R packages;
- `cpan` the updater for CPAN (<https://www.cpan.org/>) packages;
- `pypi` the updater for PyPI (<https://pypi.python.org>) packages.
- `gem` the updater for RubyGems (<https://rubygems.org>) packages.
- `github` the updater for GitHub (<https://github.com>) packages.  
The `github` updater uses the GitHub API (<https://developer.github.com/v3/>) to query for new releases. When used repeatedly e.g. when refreshing all packages, GitHub will eventually refuse to answer any further API requests. By default 60 API requests per hour are allowed, and a full refresh on all GitHub packages in Guix requires more than this. Authentication with GitHub through the use of an API token alleviates these limits. To use an API token, set the environment variable `GUIX_GITHUB_TOKEN` to a token procured from <https://github.com/settings/tokens> or otherwise.
- `hackage` the updater for Hackage (<https://hackage.haskell.org>) packages.
- `stackage` the updater for Stackage (<https://www.stackage.org>) packages.

**crate** the updater for Crates (<https://crates.io>) packages.

**launchpad** the updater for Launchpad (<https://launchpad.net>) packages.

**generic-html** a generic updater that crawls, by default, the HTML page where the source tarball of the package is hosted, when applicable. Behavior can be customized with the following package properties:

**release-monitoring-url** an alternate URL to crawl;

**release-file-regexp** an regular expression matching release file names, whose first subexpression must correspond to the version string.

Here is an example package with a custom release monitoring URL and a regexp matching an unconventional release file name (it's unconventional due to the use of upper case letter and the lack of a hyphen before the version string):

```
(package
;; ...
(home-page "http://example.org/software/the-package.html")
(properties
  `((release-monitoring-url . ,home-page)
    (release-file-regexp . "ThePackage([0-9\\.]+)\\.tgz"))))
```

**generic-git** a generic updater for packages hosted on Git repositories. It tries to be smart about parsing Git tag names, but if it is not able to parse the tag name and compare tags correctly, users can define the following properties for a package.

- **release-tag-prefix**: a regular expression for matching a prefix of the tag name.
- **release-tag-suffix**: a regular expression for matching a suffix of the tag name.
- **release-tag-version-delimiter**: a string used as the delimiter in the tag name for separating the numbers of the version.
- **accept-pre-releases**: by default, the updater will ignore pre-releases; to make it also look for pre-releases, set the this property to **#t**.

```
(package
  (name "foo")
  ;; ...
  (properties
    '(release-tag-prefix . "^release0-")
    (release-tag-suffix . "[a-z]?")
    (accept-pre-releases . #t)))
```

```
(release-tag-version-delimiter . ":"))))
```

For instance, the following command only checks for updates of Emacs packages hosted at `elpa.gnu.org` and for updates of CRAN packages:

```
$ guix refresh --type=elpa,cran
gnu/packages/statistics.scm:819:13: r-testthat would be upgraded from 0.10.0
gnu/packages/emacs.scm:856:13: emacs-auctex would be upgraded from 11.88.6 t
```

#### `--list-updaters`

List available updaters and exit (see `--type` above).

For each updater, display the fraction of packages it covers; at the end, display the fraction of packages covered by all these updaters.

#### `--list-dependent`

`-l` List top-level dependent packages that would need to be rebuilt as a result of upgrading one or more packages.

See Section 9.10 [Invoking `guix graph`], page 222, for information on how to visualize the list of dependents of a package.

See [build-dependents], page 191, for a convenient way to build all the dependents of a package.

Be aware that the `--list-dependent` option only *approximates* the rebuilds that would be required as a result of an upgrade. More rebuilds might be required under some circumstances.

```
$ guix refresh --list-dependent flex
Building the following 120 packages would ensure 213 dependent packages are
hop@2.4.0 emacs-geiser@0.13 notmuch@0.18 mu@0.9.9.5 cflow@1.4 idutils@4.6 ..
```

The command above lists a set of packages that could be built to check for compatibility with an upgraded `flex` package.

#### `--list-transitive`

`-T` List all the packages which one or more packages depend upon.

```
$ guix refresh --list-transitive flex
flex@2.6.4 depends on the following 25 packages: perl@5.28.0 help2man@1.47.6
bison@3.0.5 indent@2.2.10 tar@1.30 gzip@1.9 bzip2@1.0.6 xz@5.2.4 file@5.33.1
```

The command above lists a set of packages which, when changed, would cause `flex` to be rebuilt.

#### `--load-path=directory`

`-L directory`

Add *directory* to the front of the package module search path (see Section 8.1 [Package Modules], page 100).

This allows users to define their own packages and make them visible to the command-line tools.

The following options can be used to customize GnuPG operation:

#### `--gpg=command`

Use *command* as the GnuPG 2.x command. *command* is searched for in `$PATH`.



**--keyring=***file*

Use *file* as the keyring for upstream keys. *file* must be in the *keybox* format. Keybox files usually have a name ending in *.kbx* and the GNU Privacy Guard (GPG) can manipulate these files (see Section “*kbxutil*” in *Using the GNU Privacy Guard*, for information on a tool to manipulate keybox files).

When this option is omitted, *guix refresh* uses *~/.config/guix/upstream/trustedkeys.kbx* as the keyring for upstream signing keys. OpenPGP signatures are checked against keys from this keyring; missing keys are downloaded to this keyring as well (see **--key-download** below).

You can fetch keys to a specific keybox file like this:

```
gpg --no-default-keyring --keyring mykeyring.kbx \
  --recv-keys 3CE464558A84FDC69DB40CFB090B11993D9AEBB5
```

See Section “GPG Configuration Options” in *Using the GNU Privacy Guard*, for more information on GPG’s **--keyring** option.

**--key-download=***policy*

Handle missing OpenPGP keys according to *policy*, which may be one of:

**always** Always download missing OpenPGP keys from the key server, and add them to the user’s GnuPG keyring.

**never** Never try to download missing OpenPGP keys. Instead just bail out.

**interactive**

When a package signed with an unknown OpenPGP key is encountered, ask the user whether to download it or not.

**auto** Automatically selects the **interactive** policy when the standard input is connected to a pseudo terminal (TTY), else **always**. This is the default behavior.

**--key-server=***host*

Use *host* as the OpenPGP key server when importing a public key.

## 9.7 Invoking *guix style*

The *guix style* command helps users and packagers alike style their package definitions and configuration files according to the latest fashionable trends. It can either reformat whole files, with the **--whole-file** option, or apply specific *styling rules* to individual package definitions. The command currently provides the following styling rules:

- formatting package definitions according to the project’s conventions (see Section 22.9.4 [Formatting Code], page 841);
- rewriting package inputs to the “new style”, as explained below.

The way package inputs are written is going through a transition (see Section 8.2.1 [package Reference], page 104, for more on package inputs). Until version 1.3.0, package inputs were written using the “old style”, where each input was given an explicit label, most of the time the package name:

```
(package
```

```
;; ...
;; The "old style" (deprecated).
(inputs `(("libunistring" ,libunistring)
          ("libffi" ,libffi)))
```

Today, the old style is deprecated and the preferred style looks like this:

```
(package
  ;; ...
  ;; The "new style".
  (inputs (list libunistring libffi)))
```

Likewise, uses of `alist-delete` and friends to manipulate inputs is now deprecated in favor of `modify-inputs` (see Section 8.3 [Defining Package Variants], page 113, for more info on `modify-inputs`).

In the vast majority of cases, this is a purely mechanical change on the surface syntax that does not even incur a package rebuild. Running `guix style -S inputs` can do that for you, whether you're working on packages in Guix proper or in an external channel.

The general syntax is:

```
guix style [options] package...
```

This causes `guix style` to analyze and rewrite the definition of *package...* or, when *package* is omitted, of *all* the packages. The `--styling` or `-S` option allows you to select the style rule, the default rule being `format`—see below.

To reformat entire source files, the syntax is:

```
guix style --whole-file file...
```

The available options are listed below.

**--dry-run**

**-n** Show source file locations that would be edited but do not modify them.

**--whole-file**

**-f** Reformat the given files in their entirety. In that case, subsequent arguments are interpreted as file names (rather than package names), and the `--styling` option has no effect.

As an example, here is how you might reformat your operating system configuration (you need write permissions for the file):

```
guix style -f /etc/config.scm
```

**--alphabetical-sort**

**-A** Place the top-level package definitions in the given files in alphabetical order. Package definitions with matching names are placed with versions in descending order. This option only has an effect in combination with `--whole-file`.

**--styling=rule**

**-S rule** Apply *rule*, one of the following styling rules:

**format** Format the given package definition(s)—this is the default styling rule. For example, a packager running Guix on a checkout (see Section 22.4 [Running Guix Before It Is Installed], page 821) might want to reformat the definition of the `Coreutils` package like so:

```
./pre-inst-env guix style coreutils
```

**inputs** Rewrite package inputs to the “new style”, as described above. This is how you would rewrite inputs of package `whatnot` in your own channel:

```
guix style -L ~/my/channel -S inputs whatnot
```

Rewriting is done in a conservative way: preserving comments and bailing out if it cannot make sense of the code that appears in an `inputs` field. The `--input-simplification` option described below provides fine-grain control over when inputs should be simplified.

**arguments**

Rewrite package arguments to use G-expressions (see Section 8.12 [G-Expressions], page 165). For example, consider this package definition:

```
(define-public my-package
  (package
    ;; ...
    (arguments      ;old-style quoted arguments
      '(:make-flags '("V=1")
        #:phases (modify-phases %standard-phases
                          (delete 'build)))))
```

Running `guix style -S arguments` on this package would rewrite its `arguments` field like to:

```
(define-public my-package
  (package
    ;; ...
    (arguments
      (list #:make-flags #~'("V=1")
            #:phases #~(modify-phases %standard-phases
                                (delete 'build)))))
```

Note that changes made by the `arguments` rule do not entail a re-build of the affected packages. Furthermore, if a package definition happens to be using G-expressions already, `guix style` leaves it unchanged.

**git-source**

If the `home-page` is a Git repository (as per `git-repository-url?`), and the actual Git repository is tagged with `version` or `(string-append ``v'' version)`, change the package origin to the `git-fetch` method (see Section 8.2.2 [origin Reference], page 109). Consider this example:

```
(define-public guile-json-4
  (package
    (inherit guile-json-3)
    (name "guile-json")
    (version "4.7.3")
    (source (origin
              (method url-fetch)
```

```
(uri (string-append "mirror://savannah/guile-json/guile-json-4.7.3"
                    version ".tar.gz"))
```

```
(sha256
 (base32
  "127k2xc07w1gnyqs40z4865l8p3ra5xgpcn569dz04lksa709"))
```

Running `guix style -S git-source` on this package would rewrite its `source` field like to:

```
(define-public guile-json-4
  (package
    (inherit guile-json-3)
    (name "guile-json")
    (version "4.7.3")
    (source (origin
              (method git-fetch)
              (uri (git-reference (url
                                   "https://github.com/aconchillo/guile-json"
                                   (commit version))))
              (file-name (git-file-name name version))
              (sha256 (base32
                       "0akhm8xjv8fl55fyq0w6c9c6hi5j7mifjx01w07np7")))))
```

`--list-stylings`

`-l` List and describe the available styling rules and exit.

`--load-path=directory`

`-L directory`

Add *directory* to the front of the package module search path (see Section 8.1 [Package Modules], page 100).

`--expression=expr`

`-e expr` Style the package *expr* evaluates to.

For example, running:

```
guix style -e '(@ (gnu packages gcc) gcc-5)'
```

styles the `gcc-5` package definition.

`--input-simplification=policy`

When using the `inputs` styling rule, with `‘-S inputs’`, this option specifies the package input simplification policy for cases where an input label does not match the corresponding package name. *policy* may be one of the following:

**silent** Simplify inputs only when the change is “silent”, meaning that the package does not need to be rebuilt (its derivation is unchanged).

**safe** Simplify inputs only when that is “safe” to do: the package might need to be rebuilt, but the change is known to have no observable effect.

**always** Simplify inputs even when input labels do not match package names, and even if that might have an observable effect.

The default is `silent`, meaning that input simplifications do not trigger any package rebuild.

## 9.8 Invoking `guix lint`

The `guix lint` command is meant to help package developers avoid common errors and use a consistent style. It runs a number of checks on a given set of packages in order to find common mistakes in their definitions. Available *checkers* include (see `--list-checkers` for a complete list):

**synopsis**

**description**

Validate certain typographical and stylistic rules about package descriptions and synopses.

**inputs-should-be-native**

Identify inputs that should most likely be native inputs.

**source**

**home-page**

**mirror-url**

**github-url**

**source-file-name**

Probe **home-page** and **source** URLs and report those that are invalid. Suggest a `mirror://` URL when applicable. If the **source** URL redirects to a GitHub URL, recommend usage of the GitHub URL. Check that the source file name is meaningful, e.g. is not just a version number or “git-checkout”, without a declared **file-name** (see Section 8.2.2 [origin Reference], page 109).

**source-unstable-tarball**

Parse the **source** URL to determine if a tarball from GitHub is autogenerated or if it is a release tarball. Unfortunately GitHub’s autogenerated tarballs are sometimes regenerated.

**misplaced-flags**

Parse the **make-flags** and **configure-flags** to check for any flags which should be set with another flag as part of the build-system. For example,

```
(build-system cmake-build-system)
(arguments
 (list #:configure-flags #~(list "-DCMAKE_BUILD_TYPE=RELEASE"
                                   "-DTESTS=ON")))
```

should instead be written as:

```
(build-system cmake-build-system)
(arguments
 (list #:build-type "RELEASE"
       #:configure-flags #~(list "-DTESTS=ON")))
```

**derivation**

Check that the derivation of the given packages can be successfully computed for all the supported systems (see Section 8.10 [Derivations], page 158).

**profile-collisions**

Check whether installing the given packages in a profile would lead to collisions. Collisions occur when several packages with the same name but a different version or a different store file name are propagated. See Section 8.2.1 [package Reference], page 104, for more information on propagated inputs.

**archival**

Checks whether the package’s source code is archived at Software Heritage (<https://www.softwareheritage.org>).

When the source code that is not archived comes from a version-control system (VCS)—e.g., it’s obtained with `git-fetch`, send Software Heritage a “save” request so that it eventually archives it. This ensures that the source will remain available in the long term, and that Guix can fall back to Software Heritage should the source code disappear from its original host. The status of recent “save” requests can be viewed on-line (<https://archive.softwareheritage.org/save/#requests>).

When source code is a tarball obtained with `url-fetch`, simply print a message when it is not archived. As of this writing, Software Heritage does not allow requests to save arbitrary tarballs; we are working on ways to ensure that non-VCS source code is also archived.

Software Heritage limits the request rate per IP address (<https://archive.softwareheritage.org/api/#rate-limiting>). When the limit is reached, `guix lint` prints a message and the `archival` checker stops doing anything until that limit has been reset.

**cve**

Report known vulnerabilities found in the Common Vulnerabilities and Exposures (CVE) databases of the current and past year published by the US NIST (<https://nvd.nist.gov/vuln/data-feeds>).

To view information about a particular vulnerability, visit pages such as:

- ‘<https://web.nvd.nist.gov/view/vuln/detail?vulnId=CVE-YYYY-ABCD>’
- ‘<https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-YYYY-ABCD>’

where `CVE-YYYY-ABCD` is the CVE identifier—e.g., `CVE-2015-7554`.

Package developers can specify in package recipes the Common Platform Enumeration (CPE) (<https://nvd.nist.gov/products/cpe>) name and version of the package when they differ from the name or version that Guix uses, as in this example:

```
(package
  (name "grub")
  ;; ...
  ;; CPE calls this package "grub2".
  (properties '((cpe-name . "grub2")
                (cpe-version . "2.3"))))
```

A CVE alert can be a false positive when its CPE name matches the one in Guix, while actually referring to a distinct product. These alerts can be addressed by setting the correct CPE vendor, or when no vendors apply, by ignoring alerts from irrelevant vendors, as in these examples:

```
(package
```

```

(name "halibut")
;; ...
(properties '((cpe-vendor . "halibut_project"))))

(package
  (name "cvs")
  ;; ...
  (properties '((lint-hidden-cpe-vendors . ("jenkins"
                                             "vendor2")))))

```

Finally, some entries in the CVE database do not specify which version of a package they apply to, and would thus “stick around” forever. Package developers who found CVE alerts and verified they can be ignored can declare them as in this example:

```

(package
  (name "tllib")
  ;; ...
  ;; These CVEs no longer apply and can be safely ignored.
  (properties `((lint-hidden-cve . ("CVE-2011-0433"
                                     "CVE-2011-1553"
                                     "CVE-2011-1554"
                                     "CVE-2011-5244")))))

```

#### formatting

Warn about obvious source code formatting issues: trailing white space, use of tabulations, etc.

#### input-labels

Report old-style input labels that do not match the name of the corresponding package. This aims to help migrate from the “old input style”. See Section 8.2.1 [package Reference], page 104, for more information on package inputs and input styles. See Section 9.7 [Invoking guix style], page 213, on how to migrate to the new style.

The general syntax is:

```
guix lint options package...
```

If no package is given on the command line, then all packages are checked. The *options* may be zero or more of the following:

#### --list-checkers

-l List and describe all the available checkers that will be run on packages and exit.

#### --checkers

-c Only enable the checkers specified in a comma-separated list using the names returned by --list-checkers.

#### --exclude

-x Only disable the checkers specified in a comma-separated list using the names returned by --list-checkers.

```
--expression=expr
-e expr    Consider the package expr evaluates to.
           This is useful to unambiguously designate packages, as in this example:
               guix lint -c archival -e '(@ (gnu packages guile) guile-3.0)'

--no-network
-n         Only enable the checkers that do not depend on Internet access.

--load-path=directory
-L directory
           Add directory to the front of the package module search path (see Section 8.1
           [Package Modules], page 100).
           This allows users to define their own packages and make them visible to the
           command-line tools.
```

## 9.9 Invoking `guix size`

The `guix size` command helps package developers profile the disk usage of packages. It is easy to overlook the impact of an additional dependency added to a package, or the impact of using a single output for a package that could easily be split (see Section 5.4 [Packages with Multiple Outputs], page 51). Such are the typical issues that `guix size` can highlight.

The command can be passed one or more package specifications such as `gcc@4.8` or `guile:debug`, or a file name in the store. Consider this example:

```
$ guix size coreutils
store item                                total    self
/gnu/store/...-gcc-5.5.0-lib              60.4     30.1  38.1%
/gnu/store/...-glibc-2.27                 30.3     28.8  36.6%
/gnu/store/...-coreutils-8.28             78.9     15.0  19.0%
/gnu/store/...-gmp-6.1.2                  63.1       2.7   3.4%
/gnu/store/...-bash-static-4.4.12          1.5       1.5   1.9%
/gnu/store/...-acl-2.2.52                  61.1       0.4   0.5%
/gnu/store/...-attr-2.4.47                 60.6       0.2   0.3%
/gnu/store/...-libcap-2.25                  60.5       0.2   0.2%
total: 78.9 MiB
```

The store items listed here constitute the *transitive closure* of Coreutils—i.e., Coreutils and all its dependencies, recursively—as would be returned by:

```
$ guix gc -R /gnu/store/...-coreutils-8.23
```

Here the output shows three columns next to store items. The first column, labeled “total”, shows the size in mebibytes (MiB) of the closure of the store item—that is, its own size plus the size of all its dependencies. The next column, labeled “self”, shows the size of the item itself. The last column shows the ratio of the size of the item itself to the space occupied by all the items listed here.

In this example, we see that the closure of Coreutils weighs in at 79 MiB, most of which is taken by libc and GCC’s run-time support libraries. (That libc and GCC’s libraries represent a large fraction of the closure is not a problem *per se* because they are always available on the system anyway.)



Since the command also accepts store file names, assessing the size of a build result is straightforward:

```
guix size $(guix system build config.scm)
```

When the package(s) passed to **guix size** are available in the store<sup>1</sup>, **guix size** queries the daemon to determine its dependencies, and measures its size in the store, similar to **du -ms --apparent-size** (see Section “du invocation” in *GNU Coreutils*).

When the given packages are *not* in the store, **guix size** reports information based on the available substitutes (see Section 5.3 [Substitutes], page 46). This makes it possible to profile the disk usage of store items that are not even on disk, only available remotely.

You can also specify several package names:

```
$ guix size coreutils grep sed bash
store item                                total    self
/gnu/store/...-coreutils-8.24             77.8    13.8  13.4%
/gnu/store/...-grep-2.22                  73.1     0.8   0.8%
/gnu/store/...-bash-4.3.42                72.3     4.7   4.6%
/gnu/store/...-readline-6.3               67.6     1.2   1.2%
...
total: 102.3 MiB
```

In this example we see that the combination of the four packages takes 102.3 MiB in total, which is much less than the sum of each closure since they have a lot of dependencies in common.

When looking at the profile returned by **guix size**, you may find yourself wondering why a given package shows up in the profile at all. To understand it, you can use **guix graph --path -t references** to display the shortest path between the two packages (see Section 9.10 [Invoking guix graph], page 222).

The available options are:

**--substitute-urls=urls**

Use substitute information from *urls*. See [client-substitute-urls], page 181.

**--sort=key**

Sort lines according to *key*, one of the following options:

**self**            the size of each item (the default);

**closure**        the total size of the item’s closure.


**--map-file=file**

Write a graphical map of disk usage in PNG format to *file*.

For the example above, the map looks like this:

---

<sup>1</sup> More precisely, **guix size** looks for the *ungrafted* variant of the given package(s), as returned by **guix build package --no-grafts**. See Chapter 19 [Security Updates], page 806, for information on grafts.



This option requires that Guile-Charting (<https://wingolog.org/software/guile-charting/>) be installed and visible in Guile’s module search path. When that is not the case, **guix size** fails as it tries to load it.

**--system=system**

**-s system** Consider packages for *system*—e.g., `x86_64-linux`.

**--load-path=directory**

**-L directory**

Add *directory* to the front of the package module search path (see Section 8.1 [Package Modules], page 100).

This allows users to define their own packages and make them visible to the command-line tools.

## 9.10 Invoking **guix graph**


Packages and their dependencies form a *graph*, specifically a directed acyclic graph (DAG). It can quickly become difficult to have a mental model of the package DAG, so the **guix graph** command provides a visual representation of the DAG. By default, **guix graph** emits a DAG representation in the input format of Graphviz (<https://www.graphviz.org/>), so its output can be passed directly to the **dot** command of Graphviz. It can also emit an HTML page with embedded JavaScript code to display a “chord diagram” in a Web browser, using the d3.js (<https://d3js.org/>) library, or emit Cypher queries to construct a graph in a graph database supporting the openCypher (<https://www.opencypher.org/>) query language. With **--path**, it simply displays the shortest path between two packages. The general syntax is:

```
guix graph options package...
```

For example, the following command generates a PDF file representing the package DAG for the GNU Core Utilities, showing its build-time dependencies:

```
guix graph coreutils | dot -Tpdf > dag.pdf
```

The output looks like this:



Nice little graph, no?

You may find it more pleasant to navigate the graph interactively with `xdot` (from the `xdot` package):

```
guix graph coreutils | xdot -
```

But there is more than one graph! The one above is concise: it is the graph of package objects, omitting implicit inputs such as GCC, libc, grep, etc. It is often useful to have such a concise graph, but sometimes one may want to see more details. `guix graph` supports several types of graphs, allowing you to choose the level of detail:

**package** This is the default type used in the example above. It shows the DAG of package objects, excluding implicit dependencies. It is concise, but filters out many details.

#### reverse-package

This shows the *reverse* DAG of packages. For example:

```
guix graph --type=reverse-package ocaml
```

... yields the graph of packages that *explicitly* depend on OCaml (if you are also interested in cases where OCaml is an implicit dependency, see **reverse-bag** below).

Note that for core packages this can yield huge graphs. If all you want is to know the number of packages that depend on a given package, use `guix refresh --list-dependent` (see Section 9.6 [Invoking guix refresh], page 206).


#### bag-emerged

This is the package DAG, *including* implicit inputs.

For instance, the following command:

```
guix graph --type=bag-emerged coreutils
```

... yields this bigger graph:



At the bottom of the graph, we see all the implicit inputs of *gnu-build-system* (see Section 8.5 [Build Systems], page 121).

Now, note that the dependencies of these implicit inputs—that is, the *bootstrap dependencies* (see Chapter 20 [Bootstrapping], page 808)—are not shown here, for conciseness.

**bag** Similar to **bag-emerged**, but this time including all the bootstrap dependencies.

**bag-with-origins**

Similar to **bag**, but also showing origins and their dependencies.

**reverse-bag**

This shows the *reverse* DAG of packages. Unlike **reverse-package**, it also takes implicit dependencies into account. For example:

```
guix graph -t reverse-bag dune
```

... yields the graph of all packages that depend on Dune, directly or indirectly. Since Dune is an *implicit* dependency of many packages *via* **dune-build-system**, this shows a large number of packages, whereas **reverse-package** would show very few if any.

**derivation**

This is the most detailed representation: It shows the DAG of derivations (see Section 8.10 [Derivations], page 158) and plain store items. Compared to the above representation, many additional nodes are visible, including build scripts, patches, Guile modules, etc.

For this type of graph, it is also possible to pass a `.drv` file name instead of a package name, as in:

```
guix graph -t derivation $(guix system build -d my-config.scm)
```

**module**

This is the graph of *package modules* (see Section 8.1 [Package Modules], page 100). For example, the following command shows the graph for the package module that defines the `guile` package:

```
guix graph -t module guile | xdot -
```

All the types above correspond to *build-time dependencies*. The following graph type represents the *run-time dependencies*:

**references**

This is the graph of *references* of a package output, as returned by `guix gc --references` (see Section 5.6 [Invoking guix gc], page 53).

If the given package output is not available in the store, `guix graph` attempts to obtain dependency information from substitutes.

Here you can also pass a store file name instead of a package name. For example, the command below produces the reference graph of your profile (which can be big!):

```
guix graph -t references $(readlink -f ~/.guix-profile)
```

**referrers**

This is the graph of the *referrers* of a store item, as returned by `guix gc --referrers` (see Section 5.6 [Invoking guix gc], page 53).

This relies exclusively on local information from your store. For instance, let us suppose that the current Inkscape is available in 10 profiles on your machine; `guix graph -t referrers inkscape` will show a graph rooted at Inkscape and with those 10 profiles linked to it.

It can help determine what is preventing a store item from being garbage collected.

Often, the graph of the package you are interested in does not fit on your screen, and anyway all you want to know is *why* that package actually depends on some seemingly unrelated package. The `--path` option instructs `guix graph` to display the shortest path between two packages (or derivations, or store items, etc.):

```
$ guix graph --path emacs libunistring
emacs@26.3
mailutils@3.9
libunistring@0.9.10
$ guix graph --path -t derivation emacs libunistring
/gnu/store/...-emacs-26.3.drv
/gnu/store/...-mailutils-3.9.drv
```

```

/gnu/store/...-libunistring-0.9.10.drv
$ guix graph --path -t references emacs libunistring
/gnu/store/...-emacs-26.3
/gnu/store/...-libidn2-2.2.0
/gnu/store/...-libunistring-0.9.10

```

Sometimes you still want to visualize the graph but would like to trim it so it can actually be displayed. One way to do it is via the `--max-depth` (or `-M`) option, which lets you specify the maximum depth of the graph. In the example below, we visualize only `libreoffice` and the nodes whose distance to `libreoffice` is at most 2:

```
guix graph -M 2 libreoffice | xdot -f fdp -
```

Mind you, that's still a big ball of spaghetti, but at least `dot` can render it quickly and it can be browsed somewhat.

The available options are the following:

`--type=type`  
`-t type` Produce a graph output of *type*, where *type* must be one of the values listed above.

`--list-types`  
 List the supported graph types.

`--backend=backend`  
`-b backend` Produce a graph using the selected *backend*.

`--list-backends`  
 List the supported graph backends.  
 Currently, the available backends are Graphviz and d3.js.

`--path` Display the shortest path between two nodes of the type specified by `--type`. The example below shows the shortest path between `libreoffice` and `llvm` according to the references of `libreoffice`:

```

$ guix graph --path -t references libreoffice llvm
/gnu/store/...-libreoffice-6.4.2.2
/gnu/store/...-libepoxy-1.5.4
/gnu/store/...-mesa-19.3.4
/gnu/store/...-llvm-9.0.1

```

`--expression=expr`  
`-e expr` Consider the package *expr* evaluates to.  
 This is useful to precisely refer to a package, as in this example:

```
guix graph -e '(@@ (gnu packages commencement) gnu-make-final)'
```

`--system=system`  
`-s system` Display the graph for *system*—e.g., `i686-linux`.

The package dependency graph is largely architecture-independent, but there are some architecture-dependent bits that this option allows you to visualize.

```
--load-path=directory
-L directory
```

Add *directory* to the front of the package module search path (see Section 8.1 [Package Modules], page 100).

This allows users to define their own packages and make them visible to the command-line tools.

On top of that, **guix graph** supports all the usual package transformation options (see Section 9.1.2 [Package Transformation Options], page 182). This makes it easy to view the effect of a graph-rewriting transformation such as **--with-input**. For example, the command below outputs the graph of **git** once **openssl** has been replaced by **libressl** everywhere in the graph:

```
guix graph git --with-input=openssl=libressl
```

So many possibilities, so much fun!

## 9.11 Invoking guix publish

The purpose of **guix publish** is to enable users to easily share their store with others, who can then use it as a substitute server (see Section 5.3 [Substitutes], page 46).

When **guix publish** runs, it spawns an HTTP server which allows anyone with network access to obtain substitutes from it. This means that any machine running Guix can also act as if it were a build farm, since the HTTP interface is compatible with Cuirass, the software behind the **bordeaux.guix.gnu.org** build farm.

For security, each substitute is signed, allowing recipients to check their authenticity and integrity (see Section 5.3 [Substitutes], page 46). Because **guix publish** uses the signing key of the system, which is only readable by the system administrator, it must be started as root; the **--user** option makes it drop root privileges early on.

The signing key pair must be generated before **guix publish** is launched, using **guix archive --generate-key** (see Section 5.11 [Invoking guix archive], page 65).

When the **--advertise** option is passed, the server advertises its availability on the local network using multicast DNS (mDNS) and DNS service discovery (DNS-SD), currently *via* Guile-Avahi (see *Using Avahi in Guile Scheme Programs*).

The general syntax is:

```
guix publish options...
```

Running **guix publish** without any additional arguments will spawn an HTTP server on port 8080:

```
guix publish
```

**guix publish** can also be started following the systemd “socket activation” protocol (see Section “Service De- and Constructors” in *The GNU Shepherd Manual*).

Once a publishing server has been authorized, the daemon may download substitutes from it. See Section 5.3.3 [Getting Substitutes from Other Servers], page 48.

By default, **guix publish** compresses archives on the fly as it serves them. This “on-the-fly” mode is convenient in that it requires no setup and is immediately available. However, when serving lots of clients, we recommend using the **--cache** option, which enables caching of the archives before they are sent to clients—see below for details. The **guix weather**

command provides a handy way to check what a server provides (see Section 9.15 [Invoking guix weather], page 235).

As a bonus, `guix publish` also serves as a content-addressed mirror for source files referenced in `origin` records (see Section 8.2.2 [origin Reference], page 109). For instance, assuming `guix publish` is running on `example.org`, the following URL returns the raw `hello-2.10.tar.gz` file with the given SHA256 hash (represented in `nix-base32` format, see Section 9.4 [Invoking guix hash], page 196):

```
http://example.org/file/hello-2.10.tar.gz/sha256/0ssi1...ndq1i
```

Obviously, these URLs only work for files that are in the store; in other cases, they return 404 (“Not Found”).

Build logs are available from `/log` URLs like:

```
http://example.org/log/gwspk...-guile-2.2.3
```

When `guix-daemon` is configured to save compressed build logs, as is the case by default (see Section 2.3 [Invoking guix-daemon], page 14), `/log` URLs return the compressed log as-is, with an appropriate `Content-Type` and/or `Content-Encoding` header. We recommend running `guix-daemon` with `--log-compression=gzip` since Web browsers can automatically decompress it, which is not the case with Bzip2 compression.

The following options are available:

`--port=port`

`-p port` Listen for HTTP requests on *port*.

`--listen=host`

Listen on the network interface for *host*. The default is to accept connections from any interface.

`--user=user`

`-u user` Change privileges to *user* as soon as possible—i.e., once the server socket is open and the signing key has been read.

`--compression[=method[:level]]`

`-C [method[:level]]`

Compress data using the given *method* and *level*. *method* is one of `lzip`, `zstd`, and `gzip`; when *method* is omitted, `gzip` is used.

When *level* is zero, disable compression. The range 1 to 9 corresponds to different compression levels: 1 is the fastest, and 9 is the best (CPU-intensive). The default is 3.

Usually, `lzip` compresses noticeably better than `gzip` for a small increase in CPU usage; see benchmarks on the `lzip` Web page ([https://nongnu.org/lzip/lzip\\_benchmark.html](https://nongnu.org/lzip/lzip_benchmark.html)). However, `lzip` achieves low decompression throughput (on the order of 50 MiB/s on modern hardware), which can be a bottleneck for someone who downloads over a fast network connection.

The compression ratio of `zstd` is between that of `lzip` and that of `gzip`; its main advantage is a high decompression speed (<https://facebook.github.io/zstd/>).

Unless `--cache` is used, compression occurs on the fly and the compressed streams are not cached. Thus, to reduce load on the machine that runs `guix`



`publish`, it may be a good idea to choose a low compression level, to run `guix publish` behind a caching proxy, or to use `--cache`. Using `--cache` has the advantage that it allows `guix publish` to add `Content-Length` HTTP header to its responses.

This option can be repeated, in which case every substitute gets compressed using all the selected methods, and all of them are advertised. This is useful when users may not support all the compression methods: they can select the one they support.

`--cache=directory`

`-c directory`

Cache archives and meta-data (`.narinfo` URLs) to *directory* and only serve archives that are in cache.

When this option is omitted, archives and meta-data are created on-the-fly. This can reduce the available bandwidth, especially when compression is enabled, since this may become CPU-bound. Another drawback of the default mode is that the length of archives is not known in advance, so `guix publish` does not add a `Content-Length` HTTP header to its responses, which in turn prevents clients from knowing the amount of data being downloaded.

Conversely, when `--cache` is used, the first request for a store item (*via* a `.narinfo` URL) triggers a background process to *bake* the archive—computing its `.narinfo` and compressing the archive, if needed. Once the archive is cached in *directory*, subsequent requests succeed and are served directly from the cache, which guarantees that clients get the best possible bandwidth.

That first `.narinfo` request nonetheless returns 200, provided the requested store item is “small enough”, below the cache bypass threshold—see `--cache-bypass-threshold` below. That way, clients do not have to wait until the archive is baked. For larger store items, the first `.narinfo` request returns 404, meaning that clients have to wait until the archive is baked.

The “baking” process is performed by worker threads. By default, one thread per CPU core is created, but this can be customized. See `--workers` below.

When `--ttl` is used, cached entries are automatically deleted when they have expired.

`--workers=N`

When `--cache` is used, request the allocation of *N* worker threads to “bake” archives.

`--ttl=ttl`

Produce `Cache-Control` HTTP headers that advertise a time-to-live (TTL) of *ttl*. *ttl* must denote a duration: `5d` means 5 days, `1m` means 1 month, and so on.

This allows the user’s Guix to keep substitute information in cache for *ttl*. However, note that `guix publish` does not itself guarantee that the store items it provides will indeed remain available for as long as *ttl*.

Additionally, when `--cache` is used, cached entries that have not been accessed for *ttl* and that no longer have a corresponding item in the store, may be deleted.

**--negative-ttl=ttl**

Similarly produce **Cache-Control** HTTP headers to advertise the time-to-live (TTL) of *negative* lookups—missing store items, for which the HTTP 404 code is returned. By default, no negative TTL is advertised.

This parameter can help adjust server load and substitute latency by instructing cooperating clients to be more or less patient when a store item is missing.

**--cache-bypass-threshold=size**

When used in conjunction with **--cache**, store items smaller than *size* are immediately available, even when they are not yet in cache. *size* is a size in bytes, or it can be suffixed by M for megabytes and so on. The default is 10M.

“Cache bypass” allows you to reduce the publication delay for clients at the expense of possibly additional I/O and CPU use on the server side: depending on the client access patterns, those store items can end up being baked several times until a copy is available in cache.

Increasing the threshold may be useful for sites that have few users, or to guarantee that users get substitutes even for store items that are not popular.

**--nar-path=path**

Use *path* as the prefix for the URLs of “nar” files (see Section 5.11 [Invoking guix archive], page 65).

By default, nars are served at a URL such as `/nar/gzip/...-coreutils-8.25`. This option allows you to change the `/nar` part to *path*.

**--public-key=file****--private-key=file**

Use the specific *files* as the public/private key pair used to sign the store items being published.

The files must correspond to the same key pair (the private key is used for signing and the public key is merely advertised in the signature metadata). They must contain keys in the canonical s-expression format as produced by **guix archive --generate-key** (see Section 5.11 [Invoking guix archive], page 65). By default, `/etc/guix/signing-key.pub` and `/etc/guix/signing-key.sec` are used.

**--repl[=port]**

**-r [port]** Spawn a Guile REPL server (see Section “REPL Servers” in *GNU Guile Reference Manual*) on *port* (37146 by default). This is used primarily for debugging a running **guix publish** server.

Enabling **guix publish** on Guix System is a one-liner: just instantiate a **guix-publish-service-type** service in the **services** field of the **operating-system** declaration (see [guix-publish-service-type], page 292).

If you are instead running Guix on a “foreign distro”, follow these instructions:

- If your host distro uses the systemd init system:

```
# ln -s ~root/.guix-profile/lib/systemd/system/guix-publish.service \
    /etc/systemd/system/
# systemctl start guix-publish && systemctl enable guix-publish
```

- If your host distro uses the Upstart init system:

```
# ln -s ~root/.guix-profile/lib/upstart/system/guix-publish.conf /etc/init/
# start guix-publish
```

- Otherwise, proceed similarly with your distro's init system.

## 9.12 Invoking guix challenge

Do the binaries provided by this server really correspond to the source code it claims to build? Is a package build process deterministic? These are the questions the **guix challenge** command attempts to answer.

The former is obviously an important question: Before using a substitute server (see Section 5.3 [Substitutes], page 46), one had better *verify* that it provides the right binaries, and thus *challenge* it. The latter is what enables the former: If package builds are deterministic, then independent builds of the package should yield the exact same result, bit for bit; if a server provides a binary different from the one obtained locally, it may be either corrupt or malicious.

We know that the hash that shows up in `/gnu/store` file names is the hash of all the inputs of the process that built the file or directory—compilers, libraries, build scripts, etc. (see Chapter 1 [Introduction], page 1). Assuming deterministic build processes, one store file name should map to exactly one build output. **guix challenge** checks whether there is, indeed, a single mapping by comparing the build outputs of several independent builds of any given store item.

The command output looks like this:

```
$ guix challenge \
  --substitute-urls="https://bordeaux.guix.gnu.org https://guix.example.org" \
  openssl git plus coreutils grep
updating substitutes from 'https://bordeaux.guix.gnu.org'... 100.0%
updating substitutes from 'https://guix.example.org'... 100.0%
/gnu/store/...-openssl-1.0.2d contents differ:
  local hash: 0725122r5jnzazaacncwsvp9kgf42266ayyp814v7djsx7nk963q
  https://bordeaux.guix.gnu.org/nar/...-openssl-1.0.2d: 0725122r5jnzazaacncwsvp9kgf42266ayyp814v7djsx7nk9
  https://guix.example.org/nar/...-openssl-1.0.2d: 1zy4fmaaqcjrzazjkd3f5gmjk754b43qkq471lbyak9z0qjyim
  differing files:
    /lib/libcrypto.so.1.1
    /lib/libssl.so.1.1

/gnu/store/...-git-2.5.0 contents differ:
  local hash: 00p3bmryhjxrhpn2gxs2fy0a15lnip05197205pgbk5ra395hyha
  https://bordeaux.guix.gnu.org/nar/...-git-2.5.0: 069nb85bv4d4a6slrwjdy8v1cn4cwsmp3kdbmyb81d6zckj3nq9f
  https://guix.example.org/nar/...-git-2.5.0: 0mdqa9w1p6cmli6976v4wi0sw9r4p5prkj7lzf1877wk11c9c73
  differing file:
    /libexec/git-core/git-fsck

/gnu/store/...-plus-2.1.1 contents differ:
  local hash: 0k4v3m9z1zp8xzzizb7d8kjj72f9172xv078sq4wl73vnq9ig3ax
  https://bordeaux.guix.gnu.org/nar/...-plus-2.1.1: 0k4v3m9z1zp8xzzizb7d8kjj72f9172xv078sq4wl73vnq9ig3ax
  https://guix.example.org/nar/...-plus-2.1.1: 1cy25x1a4fzq5rk0pmvc8xhwyffnqz95h2bpvqsz2mpvlbccy0gs
  differing file:
    /share/man/man1/pius.1.gz

...
```

```

5 store items were analyzed:
- 2 (40.0%) were identical
- 3 (60.0%) differed
- 0 (0.0%) were inconclusive

```

In this example, `guix challenge` queries all the substitute servers for each of the five packages specified on the command line. It then reports those store items for which the servers obtained a result different from the local build (if it exists) and/or different from one another; here, the ‘local hash’ lines indicate that a local build result was available for each of these packages and shows its hash.

As an example, `guix.example.org` always gets a different answer. Conversely, `bordeaux.guix.gnu.org` agrees with local builds, except in the case of Git. This might indicate that the build process of Git is non-deterministic, meaning that its output varies as a function of various things that Guix does not fully control, in spite of building packages in isolated environments (see Section 5.1 [Features], page 36). Most common sources of non-determinism include the addition of timestamps in build results, the inclusion of random numbers, and directory listings sorted by inode number. See <https://reproducible-builds.org/docs/>, for more information.

To find out what is wrong with this Git binary, the easiest approach is to run:

```

guix challenge git \
  --diff=diffoscope \
  --substitute-urls="https://bordeaux.guix.gnu.org https://guix.example.org"

```

This automatically invokes `diffoscope`, which displays detailed information about files that differ.

Alternatively, we can do something along these lines (see Section 5.11 [Invoking guix archive], page 65):

```

$ wget -q -O - https://bordeaux.guix.gnu.org/nar/lzip/...-git-2.5.0 \
  | lzip -d | guix archive -x /tmp/git
$ diff -ur --no-dereference /gnu/store/...-git.2.5.0 /tmp/git

```

This command shows the difference between the files resulting from the local build, and the files resulting from the build on `bordeaux.guix.gnu.org` (see Section “Overview” in *Comparing and Merging Files*). The `diff` command works great for text files. When binary files differ, a better option is `Diffoscope` (<https://diffoscope.org/>), a tool that helps visualize differences for all kinds of files.

Once you have done that work, you can tell whether the differences are due to a non-deterministic build process or to a malicious server. We try hard to remove sources of non-determinism in packages to make it easier to verify substitutes, but of course, this is a process that involves not just Guix, but a large part of the free software community. In the meantime, `guix challenge` is one tool to help address the problem.

If you are writing packages for Guix, you are encouraged to check whether `bordeaux.guix.gnu.org` and other substitute servers obtain the same build result as you did with:

```
guix challenge package
```

The general syntax is:

```
guix challenge options argument...
```

where *argument* is a package specification such as `guile@2.0` or `glibc:debug` or, alternatively, a store file name as returned, for example, by `guix build` or `guix gc --list-live`.

When a difference is found between the hash of a locally-built item and that of a server-provided substitute, or among substitutes provided by different servers, the command displays it as in the example above and its exit code is 2 (other non-zero exit codes denote other kinds of errors).

The one option that matters is:

**--substitute-urls=urls**

Consider *urls* the whitespace-separated list of substitute source URLs to compare to.

**--diff=mode**

Upon mismatches, show differences according to *mode*, one of:

**simple** (the default)

Show the list of files that differ.

**diffoscope**

*command* Invoke Diffoscope (<https://diffoscope.org/>), passing it two directories whose contents do not match.

When *command* is an absolute file name, run *command* instead of Diffoscope.

**none**

Do not show further details about the differences.

Thus, unless **--diff=none** is passed, `guix challenge` downloads the store items from the given substitute servers so that it can compare them.

**--verbose**

**-v** Show details about matches (identical contents) in addition to information about mismatches.

## 9.13 Invoking `guix copy`

The `guix copy` command copies items from the store of one machine to that of another machine over a secure shell (SSH) connection<sup>2</sup>. For example, the following command copies the `coreutils` package, the user's profile, and all their dependencies over to *host*, logged in as *user*:

```
guix copy --to=user@host \
  coreutils $(readlink -f ~/.guix-profile)
```

If some of the items to be copied are already present on *host*, they are not actually sent.

The command below retrieves `libreoffice` and `gimp` from *host*, assuming they are available there:

```
guix copy --from=host libreoffice gimp
```

The SSH connection is established using the Guile-SSH client, which is compatible with OpenSSH: it honors `~/.ssh/known_hosts` and `~/.ssh/config`, and uses the SSH agent for authentication.

<sup>2</sup> This command is available only when Guile-SSH was found. See Section 22.1 [Requirements], page 816, for details.

The key used to sign items that are sent must be accepted by the remote machine. Likewise, the key used by the remote machine to sign items you are retrieving must be in `/etc/guix/acl` so it is accepted by your own daemon. See Section 5.11 [Invoking guix archive], page 65, for more information about store item authentication.

The general syntax is:

```
guix copy [--to=spec|--from=spec] items...
```

You must always specify one of the following options:

`--to=spec`

`--from=spec`

Specify the host to send to or receive from. *spec* must be an SSH spec such as `example.org`, `charlie@example.org`, or `charlie@example.org:2222`.

The *items* can be either package names, such as `gimp`, or store items, such as `/gnu/store/...-idutils-4.6`.

When specifying the name of a package to send, it is first built if needed, unless `--dry-run` was specified. Common build options are supported (see Section 9.1.1 [Common Build Options], page 180).

## 9.14 Invoking guix container

**Note:** As of version 8b23207, this tool is experimental. The interface is subject to radical change in the future.

The purpose of `guix container` is to manipulate processes running within an isolated environment, commonly known as a “container”, typically created by the `guix shell` (see Section 7.1 [Invoking guix shell], page 78) and `guix system container` (see Section 11.16 [Invoking guix system], page 692) commands.

The general syntax is:

```
guix container action options...
```

*action* specifies the operation to perform with a container, and *options* specifies the context-specific arguments for the action.

The following actions are available:

**exec**      Execute a command within the context of a running container.

The syntax is:

```
guix container exec pid program arguments...
```

*pid* specifies the process ID of the running container. *program* specifies an executable file name within the root file system of the container. *arguments* are the additional options that will be passed to *program*.

The following command launches an interactive login shell inside a Guix system container, started by `guix system container`, and whose process ID is 9001:

```
guix container exec 9001 /run/current-system/profile/bin/bash --login
```

Note that the *pid* cannot be the parent process of a container. It must be PID 1 of the container or one of its child processes.

## 9.15 Invoking guix weather

Occasionally you're grumpy because substitutes are lacking and you end up building packages by yourself (see Section 5.3 [Substitutes], page 46). The **guix weather** command reports on substitute availability on the specified servers so you can have an idea of whether you'll be grumpy today. It can sometimes be useful info as a user, but it is primarily useful to people running **guix publish** (see Section 9.11 [Invoking guix publish], page 227). Sometimes substitutes *are* available but they are not authorized on your system; **guix weather** reports it so you can authorize them if you want (see Section 5.3.3 [Getting Substitutes from Other Servers], page 48).

Here's a sample run:

```
$ guix weather --substitute-urls=https://guix.example.org
computing 5,872 package derivations for x86_64-linux...
looking for 6,128 store items on https://guix.example.org..
updating substitutes from 'https://guix.example.org'... 100.0%
https://guix.example.org
  43.4% substitutes available (2,658 out of 6,128)
  7,032.5 MiB of nars (compressed)
 19,824.2 MiB on disk (uncompressed)
   0.030 seconds per request (182.9 seconds in total)
  33.5 requests per second

  9.8% (342 out of 3,470) of the missing items are queued
 867 queued builds
    x86_64-linux: 518 (59.7%)
    i686-linux: 221 (25.5%)
    aarch64-linux: 128 (14.8%)
build rate: 23.41 builds per hour
    x86_64-linux: 11.16 builds per hour
    i686-linux: 6.03 builds per hour
    aarch64-linux: 6.41 builds per hour
```

As you can see, it reports the fraction of all the packages for which substitutes are available on the server—regardless of whether substitutes are enabled, and regardless of whether this server's signing key is authorized. It also reports the size of the compressed archives (“nars”) provided by the server, the size the corresponding store items occupy in the store (assuming deduplication is turned off), and the server's throughput. The second part gives continuous integration (CI) statistics, if the server supports it. In addition, using the **--coverage** option, **guix weather** can list “important” package substitutes missing on the server (see below).

To achieve that, **guix weather** queries over HTTP(S) meta-data (*narinfos*) for all the relevant store items. Like **guix challenge**, it ignores signatures on those substitutes, which is innocuous since the command only gathers statistics and cannot install those substitutes.

The general syntax is:

```
guix weather options... [packages...]
```

When *packages* is omitted, **guix weather** checks the availability of substitutes for *all* the packages, or for those specified with **--manifest**; otherwise it only considers the specified

packages. It is also possible to query specific system types with `--system`. `guix weather` exits with a non-zero code when the fraction of available substitutes is below 100%.

The available options are listed below.

`--substitute-urls=urls`

*urls* is the space-separated list of substitute server URLs to query. When this option is omitted, the URLs specified with the `--substitute-urls` option of `guix-daemon` are used or, as a last resort, the default set of substitute URLs.

`--system=system`

`-s system` Query substitutes for *system*—e.g., `aarch64-linux`. This option can be repeated, in which case `guix weather` will query substitutes for several system types.

`--manifest=file`

Instead of querying substitutes for all the packages, only ask for those specified in *file*. *file* must contain a *manifest*, as with the `-m` option of `guix package` (see Section 5.2 [Invoking guix package], page 37).

This option can be repeated several times, in which case the manifests are concatenated.

`--expression=expr`

`-e expr` Consider the package *expr* evaluates to.

A typical use case for this option is specifying a package that is hidden and thus cannot be referred to in the usual way, as in this example:

```
guix weather -e '(@@ (gnu packages rust) rust-bootstrap)'
```

This option can be repeated.

`--coverage[=count]`

`-c [count]`

Report on substitute coverage for packages: list packages with at least *count* dependents (zero by default) for which substitutes are unavailable. Dependent packages themselves are not listed: if *b* depends on *a* and *a* has no substitutes, only *a* is listed, even though *b* usually lacks substitutes as well. The result looks like this:

```
$ guix weather --substitute-urls=https://bordeaux.guix.gnu.org
https://ci.guix.gnu.org -c 10
computing 8,983 package derivations for x86_64-linux...
looking for 9,343 store items on https://bordeaux.guix.gnu.org
https://ci.guix.gnu.org...
updating substitutes from 'https://bordeaux.guix.gnu.org https://ci.guix.gnu.org
https://bordeaux.guix.gnu.org https://ci.guix.gnu.org
64.7% substitutes available (6,047 out of 9,343)
...
2502 packages are missing from 'https://bordeaux.guix.gnu.org
https://ci.guix.gnu.org' for 'x86_64-linux', among which:
58 kcoreaddons@5.49.0      /gnu/store/...-kcoreaddons-5.49.0
46 qgpgme@1.11.1          /gnu/store/...-qgpgme-1.11.1
```



```
37 perl-http-cookiejar@0.008 /gnu/store/...-perl-http-cookiejar-0.008
...
```

What this example shows is that `kcoreaddons` and presumably the 58 packages that depend on it have no substitutes at `bordeaux.guix.gnu.org`; likewise for `qgpgme` and the 46 packages that depend on it.

If you are a Guix developer, or if you are taking care of this build farm, you'll probably want to have a closer look at these packages: they may simply fail to build.

`--display-missing`

Display the list of store items for which substitutes are missing.

## 9.16 Invoking guix processes

The `guix processes` command can be useful to developers and system administrators, especially on multi-user machines and on build farms: it lists the current sessions (connections to the daemon), as well as information about the processes involved<sup>3</sup>. Here's an example of the information it returns:

```
$ sudo guix processes
SessionPID: 19002
ClientPID: 19090
ClientCommand: guix shell python

SessionPID: 19402
ClientPID: 19367
ClientCommand: guix publish -u guix-publish -p 3000 -C 9 ...

SessionPID: 19444
ClientPID: 19419
ClientCommand: cuirass --cache-directory /var/cache/cuirass ...
LockHeld: /gnu/store/...-perl-ipc-cmd-0.96.lock
LockHeld: /gnu/store/...-python-six-bootstrap-1.11.0.lock
LockHeld: /gnu/store/...-libjpeg-turbo-2.0.0.lock
ChildPID: 20495
ChildCommand: guix offload x86_64-linux 7200 1 28800
ChildPID: 27733
ChildCommand: guix offload x86_64-linux 7200 1 28800
ChildPID: 27793
ChildCommand: guix offload x86_64-linux 7200 1 28800
```

In this example we see that `guix-daemon` has three clients: `guix shell`, `guix publish`, and the Cuirass continuous integration tool; their process identifier (PID) is given by the `ClientPID` field. The `SessionPID` field gives the PID of the `guix-daemon` sub-process of this particular session.

The `LockHeld` fields show which store items are currently locked by this session, which corresponds to store items being built or substituted (the `LockHeld` field is not displayed

<sup>3</sup> Remote sessions, when `guix-daemon` is started with `--listen` specifying a TCP endpoint, are *not* listed.

when `guix processes` is not running as root). Last, by looking at the `ChildPID` and `ChildCommand` fields, we understand that these three builds are being offloaded (see Section 2.2.2 [Daemon Offload Setup], page 9).

The output is in Recutils format so we can use the handy `recsel` command to select sessions of interest (see Section “Selection Expressions” in *GNU recutils manual*). As an example, the command shows the command line and PID of the client that triggered the build of a Perl package:

```
$ sudo guix processes | \
  recsel -p ClientPID,ClientCommand -e 'LockHeld ~ "perl"'
ClientPID: 19419
ClientCommand: cuirass --cache-directory /var/cache/cuirass ...
```

Additional options are listed below.

`--format=format`

`-f format` Produce output in the specified *format*, one of:

**recutils** The default option. It outputs a set of Session recutils records that include each `ChildProcess` as a field.

**normalized**

Normalize the output records into record sets (see Section “Record Sets” in *GNU recutils manual*). Normalizing into record sets allows joins across record types. The example below lists the PID of each `ChildProcess` and the associated PID for `Session` that spawned the `ChildProcess` where the `Session` was started using `guix build`.

```
$ guix processes --format=normalized | \
  recsel \
  -j Session \
  -t ChildProcess \
  -p Session.PID,PID \
  -e 'Session.ClientCommand ~ "guix build"'
PID: 4435
Session_PID: 4278

PID: 4554
Session_PID: 4278

PID: 4646
Session_PID: 4278
```

## 10 Foreign Architectures

You can target computers of different CPU architectures when producing packages (see Section 5.2 [Invoking guix package], page 37), packs (see Section 7.3 [Invoking guix pack], page 90) or full systems (see Section 11.16 [Invoking guix system], page 692).

GNU Guix supports two distinct mechanisms to target foreign architectures:

1. The traditional cross-compilation ([https://en.wikipedia.org/wiki/Cross\\_compiler](https://en.wikipedia.org/wiki/Cross_compiler)) mechanism.
2. The native building mechanism which consists in building using the CPU instruction set of the foreign system you are targeting. It often requires emulation, using the QEMU program for instance.

### 10.1 Cross-Compilation

The commands supporting cross-compilation are proposing the `--list-targets` and `--target` options.

The `--list-targets` option lists all the supported targets that can be passed as an argument to `--target`.

```
$ guix build --list-targets
The available targets are:
```

```
- aarch64-linux-gnu
- arm-linux-gnueabi
- avr
- i586-pc-gnu
- i686-linux-gnu
- i686-w64-mingw32
- loongarch64-linux-gnu
- mips64el-linux-gnu
- or1k-elf
- powerpc-linux-gnu
- powerpc64le-linux-gnu
- riscv64-linux-gnu
- x86_64-linux-gnu
- x86_64-linux-gnux32
- x86_64-w64-mingw32
- xtensa-ath9k-elf
```

Targets are specified as GNU triplets (see Section “Specifying Target Triplets” in *Autoconf*).

Those triplets are passed to GCC and the other underlying compilers possibly involved when building a package, a system image or any other GNU Guix output.

```
$ guix build --target=aarch64-linux-gnu hello
/gnu/store/9926by9qrx91ijkhw9ndgwp4bn24g9h-hello-2.12

$ file /gnu/store/9926by9qrx91ijkhw9ndgwp4bn24g9h-hello-2.12/bin/hello
```

```
/gnu/store/9926by9qrx91ijkh9ndgwp4bn24g9h-hello-2.12/bin/hello: ELF
64-bit LSB executable, ARM aarch64 ...
```

The major benefit of cross-compilation is that there are no performance penalty compared to emulation using QEMU. There are however higher risks that some packages fail to cross-compile because fewer users are using this mechanism extensively.

## 10.2 Native Builds

The commands that support impersonating a specific system have the `--list-systems` and `--system` options.

The `--list-systems` option lists all the supported systems that can be passed as an argument to `--system`.

```
$ guix build --list-systems
The available systems are:
```

```
- x86_64-linux [current]
- aarch64-linux
- armhf-linux
- i586-gnu
- i686-linux
- mips64el-linux
- powerpc-linux
- powerpc64le-linux
- riscv64-linux
```

```
$ guix build --system=i686-linux hello
/gnu/store/cc0km35s8x2z4pmwkrqqjx46i8b1i3gm-hello-2.12
```

```
$ file /gnu/store/cc0km35s8x2z4pmwkrqqjx46i8b1i3gm-hello-2.12/bin/hello
/gnu/store/cc0km35s8x2z4pmwkrqqjx46i8b1i3gm-hello-2.12/bin/hello: ELF
32-bit LSB executable, Intel 80386 ...
```

In the above example, the current system is *x86\_64-linux*. The *hello* package is however built for the *i686-linux* system.

This is possible because the *i686* CPU instruction set is a subset of the *x86\_64*, hence *i686* targeting binaries can be run on *x86\_64*.

Still in the context of the previous example, if picking the *aarch64-linux* system and the `guix build --system=aarch64-linux hello` has to build some derivations, an extra step might be needed.

The *aarch64-linux* targeting binaries cannot directly be run on a *x86\_64-linux* system. An emulation layer is requested. The GNU Guix daemon can take advantage of the Linux kernel `binfmt_misc` ([https://en.wikipedia.org/wiki/Binfmt\\_misc](https://en.wikipedia.org/wiki/Binfmt_misc)) mechanism for that. In short, the Linux kernel can defer the execution of a binary targeting a foreign platform, here *aarch64-linux*, to a userspace program, usually an emulator.

There is a service that registers QEMU as a backend for the `binfmt_misc` mechanism (see Section 11.10.32 [Virtualization Services], page 581). On Debian based foreign distributions, the alternative would be the `qemu-user-static` package.

If the `binfmt_misc` mechanism is not setup correctly, the building will fail this way:

```
$ guix build --system=armhf-linux hello --check
...
unsupported-platform /gnu/store/jjn969pijv7hff62025yxpfmtc8zy0aq0-hello-2.12.drv aarch64-linux
while setting up the build environment: a `aarch64-linux' is required to
build `/gnu/store/jjn969pijv7hff62025yxpfmtc8zy0aq0-hello-2.12.drv', but
I am a `x86_64-linux'...
```

whereas, with the `binfmt_misc` mechanism correctly linked with QEMU, one can expect to see:

```
$ guix build --system=armhf-linux hello --check
/gnu/store/13xz4nghg39wpymivlwghy08yzj97hlj-hello-2.12
```

The main advantage of native building compared to cross-compiling, is that more packages are likely to build correctly. However it comes at a price: compilation backed by QEMU is *way slower* than cross-compilation, because every instruction needs to be emulated.

The availability of substitutes for the architecture targeted by the `--system` option can mitigate this problem. An other way to work around it is to install GNU Guix on a machine whose CPU supports the targeted instruction set, and set it up as an offload machine (see Section 2.2.2 [Daemon Offload Setup], page 9).

## 11 System Configuration

Guix System supports a consistent whole-system configuration mechanism. By that we mean that all aspects of the global system configuration—such as the available system services, timezone and locale settings, user accounts—are declared in a single place. Such a *system configuration* can be *instantiated*—i.e., effected.

One of the advantages of putting all the system configuration under the control of Guix is that it supports transactional system upgrades, and makes it possible to roll back to a previous system instantiation, should something go wrong with the new one (see Section 5.1 [Features], page 36). Another advantage is that it makes it easy to replicate the exact same configuration across different machines, or at different points in time, without having to resort to additional administration tools layered on top of the own tools of the system.

This section describes this mechanism. First we focus on the system administrator’s viewpoint—explaining how the system is configured and instantiated. Then we show how this mechanism can be extended, for instance to support new system services.

### 11.1 Getting Started

You’re reading this section probably because you have just installed Guix System (see Chapter 3 [System Installation], page 23) and would like to know where to go from here. If you’re already familiar with GNU/Linux system administration, the way Guix System is configured is very different from what you’re used to: you won’t install a system service by running `guix install`, you won’t configure services by modifying files under `/etc`, and you won’t create user accounts by invoking `useradd`; instead, all these aspects are spelled out in a *system configuration file*.

The first step with Guix System is thus to write the *system configuration file*; luckily, system installation already generated one for you and stored it under `/etc/config.scm`.

**Note:** You can store your system configuration file anywhere you like—it doesn’t have to be at `/etc/config.scm`. It’s a good idea to keep it under version control, for instance in a Git repository (<https://git-scm.com/book/en/>).

The *entire* configuration of the system—user accounts, system services, timezone, locale settings—is declared in this file, which follows this template:

```
(use-modules (gnu))
(use-package-modules ...)
(use-service-modules ...)

(operating-system
  (host-name ...)
  (timezone ...)
  (locale ...)
  (bootloader ...)
  (file-systems ...)
  (users ...)
  (packages ...)
  (services ...))
```

This configuration file is in fact a Scheme program; the first lines pull in modules providing variables you might need in the rest of the file—e.g., packages, services, etc. The `operating-system` form declares the system configuration as a *record* with a number of *fields*. See Section 11.2 [Using the Configuration System], page 244, to view complete examples and learn what to put in there.

The second step, once you have this configuration file, is to test it. Of course, you can skip this step if you're feeling lucky—you choose! To do that, pass your configuration file to `guix system vm` (no need to be root, you can do that as a regular user):

```
guix system vm /etc/config.scm
```

This command returns the name of a shell script that starts a virtual machine (VM) running the system *as described in the configuration file*:

```
/gnu/store/...-run-vm.sh
```

In this VM, you can log in as `root` with no password. That's a good way to check that your configuration file is correct and that it gives the expected result, without touching your system. See Section 11.16 [Invoking `guix system`], page 692, for more information.

**Note:** When using `guix system vm`, aspects tied to your hardware such as file systems and mapped devices are overridden because they cannot be meaningfully tested in the VM. Other aspects such as static network configuration (see Section 11.10.4 [Networking Setup], page 307) are *not* overridden but they may not work inside the VM.

The third step, once you're happy with your configuration, is to *instantiate* it—make this configuration effective on your system. To do that, run:

```
sudo guix system reconfigure /etc/config.scm
```

This operation is *transactional*: either it succeeds and you end up with an upgraded system, or it fails and nothing has changed. Note that it does *not* restart system services that were already running. Thus, to upgrade those services, you have to reboot or to explicitly restart them; for example, to restart the secure shell (SSH) daemon, you would run:

```
sudo herd restart sshd
```

**Note:** System services are managed by the Shepherd (see Section “Jump Start” in *The GNU Shepherd Manual*). The `herd` command lets you inspect, start, and stop services. To view the status of services, run:

```
sudo herd status
```

To view detailed information about a given service, add its name to the command:

```
sudo herd status sshd
```

See Section 11.10 [Services], page 275, for more information.

The system records its *provenance*—the configuration file and channels that were used to deploy it. You can view it like so:

```
guix system describe
```

Additionally, `guix system reconfigure` preserves previous system generations, which you can list:

```
guix system list-generations
```

Crucially, that means that you can always *roll back* to an earlier generation should something go wrong! When you eventually reboot, you'll notice a sub-menu in the bootloader that reads "Old system generations": it's what allows you to boot *an older generation of your system*, should the latest generation be "broken" or otherwise unsatisfying. You can also "permanently" roll back, like so:

```
sudo guix system roll-back
```

Alternatively, you can use `guix system switch-generation` to switch to a specific generation.

Once in a while, you'll want to delete old generations that you do not need anymore to allow *garbage collection* to free space (see Section 5.6 [Invoking `guix gc`], page 53). For example, to remove generations older than 4 months, run:

```
sudo guix system delete-generations 4m
```

From there on, anytime you want to change something in the system configuration, be it adding a user account or changing parameters of a service, you will first update your configuration file and then run `guix system reconfigure` as shown above. Likewise, to *upgrade* system software, you first fetch an up-to-date Guix and then reconfigure your system with that new Guix:

```
guix pull
sudo guix system reconfigure /etc/config.scm
```

We recommend doing that regularly so that your system includes the latest security updates (see Chapter 19 [Security Updates], page 806).

**Note:** `sudo guix` runs your user's `guix` command and *not* root's, because `sudo` leaves `PATH` unchanged.

The difference matters here, because `guix pull` updates the `guix` command and package definitions only for the user it is run as. This means that if you choose to use `guix system reconfigure` in root's login shell, you'll need to `guix pull` separately.

That's it! If you're getting started with Guix entirely, see Chapter 4 [Getting Started], page 33. The next sections dive in more detail into the crux of the matter: system configuration.

## 11.2 Using the Configuration System

The previous section showed the overall workflow you would follow when administering a Guix System machine (see Section 11.1 [Getting Started with the System], page 242). Let's now see in more detail what goes into the system configuration file.

The operating system is configured by providing an `operating-system` declaration in a file that can then be passed to the `guix system` command (see Section 11.16 [Invoking `guix system`], page 692), as we've seen before. A simple setup, with the default Linux-Libre kernel, initial RAM disk, and a couple of system services added to those provided by default looks like this:

```
;; -*- mode: scheme; -*-
;; This is an operating system configuration template
;; for a "bare bones" setup, with no X11 display server.
```



```

(use-modules (gnu))
(use-service-modules networking ssh)
(use-package-modules screen ssh)

(operating-system
  (host-name "komputilo")
  (timezone "Europe/Berlin")
  (locale "en_US.utf8")

  ;; Boot in "legacy" BIOS mode, assuming /dev/sdX is the
  ;; target hard disk, and "my-root" is the label of the target
  ;; root file system.
  (bootloader (bootloader-configuration
    (bootloader grub-bootloader)
    (targets '("/dev/sdX"))))

  ;; It's fitting to support the equally bare bones '-nographic'
  ;; QEMU option, which also nicely sidesteps forcing QWERTY.
  (kernel-arguments (list "console=ttyS0,115200"))
  (file-systems (cons (file-system
    (device (file-system-label "my-root"))
    (mount-point "/")
    (type "ext4"))
    %base-file-systems))

  ;; This is where user accounts are specified. The "root"
  ;; account is implicit, and is initially created with the
  ;; empty password.
  (users (cons (user-account
    (name "alice")
    (comment "Bob's sister")
    (group "users")

    ;; Adding the account to the "wheel" group
    ;; makes it a sudoer. Adding it to "audio"
    ;; and "video" allows the user to play sound
    ;; and access the webcam.
    (supplementary-groups '("wheel"
      "audio" "video"))

    %base-user-accounts))

  ;; Globally-installed packages.
  (packages (cons screen %base-packages))

  ;; Add services to the baseline: a DHCP client and an SSH
  ;; server. You may wish to add an NTP service here.
  (services (append (list (service dhcpcd-service-type)

```

```
(service openssh-service-type
  (openssh-configuration
    (openssh openssh-sans-x)
    (port-number 2222))))
%base-services)))
```

The configuration is declarative. It is code in the Scheme programming language; the whole `(operating-system ...)` expression produces a *record* with a number of *fields*. Some of the fields defined above, such as `host-name` and `bootloader`, are mandatory. Others, such as `packages` and `services`, can be omitted, in which case they get a default value. See Section 11.3 [operating-system Reference], page 253, for details about all the available fields.

Below we discuss the meaning of some of the most important fields.

**Troubleshooting:** The configuration file is a Scheme program and you might get the syntax or semantics wrong as you get started. Syntactic issues such as misplaced parentheses can often be identified by reformatting your file:

```
guix style -f config.scm
```

The Cookbook has a short section to get started with the Scheme programming language that explains the fundamentals, which you will find helpful when hacking your configuration. See Section “A Scheme Crash Course” in *GNU Guix Cookbook*.

## Bootloader

The `bootloader` field describes the method that will be used to boot your system. Machines based on Intel processors can boot in “legacy” BIOS mode, as in the example above. However, more recent machines rely instead on the *Unified Extensible Firmware Interface* (UEFI) to boot. In that case, the `bootloader` field should contain something along these lines:

```
(bootloader-configuration
  (bootloader grub-efi-bootloader)
  (targets '("/boot/efi")))
```

See Section 11.15 [Bootloader Configuration], page 685, for more information on the available configuration options.

## Globally-Visible Packages

The `packages` field lists packages that will be globally visible on the system, for all user accounts—i.e., in every user’s `PATH` environment variable—in addition to the per-user profiles (see Section 5.2 [Invoking guix package], page 37). The `%base-packages` variable provides all the tools one would expect for basic user and administrator tasks—including the GNU Core Utilities, the GNU Networking Utilities, the `mg` lightweight text editor, `find`, `grep`, etc. The example above adds GNU Screen to those, taken from the `(gnu packages screen)` module (see Section 8.1 [Package Modules], page 100). The `(list package output)` syntax can be used to add a specific output of a package:

```
(use-modules (gnu packages))
(use-modules (gnu packages dns))
```

```
(operating-system
;; ...
(packages (cons (list isc-bind "utils")
                %base-packages)))
```

Referring to packages by variable name, like `isc-bind` above, has the advantage of being unambiguous; it also allows typos and such to be diagnosed right away as “unbound variables”. The downside is that one needs to know which module defines which package, and to augment the `use-package-modules` line accordingly. To avoid that, one can use the `specification->package` procedure of the `(gnu packages)` module, which returns the best package for a given name or name and version:

```
(use-modules (gnu packages))

(operating-system
;; ...
(packages (append (map specification->package
                        '("tcpdump" "htop" "gnupg@2.0"))
                %base-packages)))
```

When a package has more than one output it can be a challenge to refer to a specific output instead of just to the standard out output. For these situations one can use the `specifications->packages` procedure from the `(gnu packages)` module. For example:

```
(use-modules (gnu packages))

(operating-system
;; ...
(packages (append (specifications->packages
                        '("git" "git:send-email"))
                %base-packages)))
```

## System Services

The `services` field lists *system services* to be made available when the system starts (see Section 11.10 [Services], page 275). The `operating-system` declaration above specifies that, in addition to the basic services, we want the OpenSSH secure shell daemon listening on port 2222 (see Section 11.10.5 [Networking Services], page 322). Under the hood, `openssh-service-type` arranges so that `sshd` is started with the right command-line options, possibly with supporting configuration files generated as needed (see Section 11.19 [Defining Services], page 709).

Occasionally, instead of using the base services as is, you will want to customize them. To do this, use `modify-services` (see Section 11.19.3 [Service Reference], page 712) to modify the list.

For example, suppose you want to modify `guix-daemon` and Mingetty (the console login) in the `%base-services` list (see Section 11.10.1 [Base Services], page 276). To do that, you can write the following in your operating system declaration:

```
(define %my-services
;; My very own list of services.
```

```

(modify-services %base-services
  (guix-service-type config =>
    (guix-configuration
      (inherit config)
      ;; Fetch substitutes from example.org.
      (substitute-urls
        (list "https://example.org/guix"
              "https://ci.guix.gnu.org"))))
  (mingetty-service-type config =>
    (mingetty-configuration
      (inherit config)
      ;; Automatically log in as "guest".
      (auto-login "guest")))))

(operating-system
  ;; ...
  (services %my-services))

```

This changes the configuration—i.e., the service parameters—of the `guix-service-type` instance, and that of all the `mingetty-service-type` instances in the `%base-services` list (see Section “Auto-Login to a Specific TTY” in *GNU Guix Cookbook*). Observe how this is accomplished: first, we arrange for the original configuration to be bound to the identifier `config` in the *body*, and then we write the *body* so that it evaluates to the desired configuration. In particular, notice how we use `inherit` to create a new configuration which has the same values as the old configuration, but with a few modifications.

The configuration for a typical “desktop” usage, with an encrypted root partition, a swap file on the root partition, the X11 display server, GNOME and Xfce (users can choose which of these desktop environments to use at the log-in screen by pressing *F1*), network management, power management, and more, would look like this:

```

;; -*- mode: scheme; -*-
;; This is an operating system configuration template
;; for a "desktop" setup with GNOME and Xfce where the
;; root partition is encrypted with LUKS, and a swap file.

(use-modules (gnu) (gnu system nss) (guix utils)
             (guix packages))
(use-service-modules desktop sddm xorg)
(use-package-modules gnome)

(operating-system
  (host-name "antelope")
  (timezone "Europe/Paris")
  (locale "en_US.utf8")

  ;; Choose US English keyboard layout. The "altgr-intl"
  ;; variant provides dead keys for accented characters.
  (keyboard-layout (keyboard-layout "us" "altgr-intl")))

```

```

;; Use the UEFI variant of GRUB with the EFI System
;; Partition mounted on /boot/efi.
(bootloader (bootloader-configuration
  (bootloader grub-efi-bootloader)
  (targets '("/boot/efi"))
  (keyboard-layout keyboard-layout)))

;; Specify a mapped device for the encrypted root partition.
;; The UUID is that returned by 'cryptsetup luksUUID'.
(mapped-devices
  (list (mapped-device
    (source (uuid "12345678-1234-1234-1234-123456789abc"))
    (target "my-root")
    (type luks-device-mapping))))

(file-systems (append
  (list (file-system
    (device (file-system-label "my-root"))
    (mount-point "/")
    (type "ext4")
    (dependencies mapped-devices))
    (file-system
    (device (uuid "1234-ABCD" 'fat))
    (mount-point "/boot/efi")
    (type "vfat"))))
  %base-file-systems))

;; Specify a swap file for the system, which resides on the
;; root file system.
(swap-devices (list (swap-space
  (target "/swapfile"))))

;; Create user `bob' with `alice' as its initial password.
(users (cons (user-account
  (name "bob")
  (comment "Alice's brother")
  (password (crypt "alice" "$6$abc"))
  (group "students")
  (supplementary-groups '("wheel" "netdev"
    "audio" "video")))
  %base-user-accounts))

;; Add the `students' group
(groups (cons* (user-group
  (name "students"))
  %base-groups))

```

```
;; This is where we specify system-wide packages.
(packages (append (list
                    ;; for user mounts
                    gvfs)
                  %base-packages))

;; Add GNOME and Xfce---we can choose at the log-in screen
;; by clicking the gear. Use the "desktop" services, which
;; include the X11 log-in service, networking with
;; NetworkManager, and more.
(services (if (supported-package? gnome
                                (or (and=> (%current-target-system)
                                             gnu-triplet->nix-system)
                                    (%current-system)))
              (append (list (service gnome-desktop-service-type)
                            (service xfce-desktop-service-type)
                            (set-xorg-configuration
                             (xorg-configuration
                              (keyboard-layout keyboard-layout))))
                        %desktop-services)
              ;; FIXME: Since Gnome depends on Rust and Rust is currently
              ;; unavailable on some platforms, we use MATE here instead of
              ;; GNOME.
              (append (list (service mate-desktop-service-type)
                            (service xfce-desktop-service-type)
                            (set-xorg-configuration
                             (xorg-configuration
                              (keyboard-layout keyboard-layout))))
                        %desktop-services)))

;; Allow resolution of '.local' host names with mDNS.
(name-service-switch %mdns-host-lookup-nss))
```

A graphical system with a choice of lightweight window managers instead of full-blown desktop environments would look like this:

```
;; -*- mode: scheme; -*-
;; This is an operating system configuration template
;; for a "desktop" setup without full-blown desktop
;; environments.

(use-modules (gnu) (gnu system nss))
(use-service-modules desktop)
(use-package-modules bootloaders emacs emacs-xyz ratpoison suckless wm
                    xorg)
```

```

(operating-system
  (host-name "antelope")
  (timezone "Europe/Paris")
  (locale "en_US.utf8")

  ;; Use the UEFI variant of GRUB with the EFI System
  ;; Partition mounted on /boot/efi.
  (bootloader (bootloader-configuration
    (bootloader grub-efi-bootloader)
    (targets '("/boot/efi"))))

  ;; Assume the target root file system is labelled "my-root",
  ;; and the EFI System Partition has UUID 1234-ABCD.
  (file-systems (append
    (list (file-system
      (device (file-system-label "my-root"))
      (mount-point "/")
      (type "ext4"))
      (file-system
      (device (uuid "1234-ABCD" 'fat))
      (mount-point "/boot/efi")
      (type "vfat"))))
    %base-file-systems))

  (users (cons (user-account
    (name "alice")
    (comment "Bob's sister")
    (group "users")
    (supplementary-groups '("wheel" "netdev"
      "audio" "video")))
    %base-user-accounts))

  ;; Add a bunch of window managers; we can choose one at
  ;; the log-in screen with F1.
  (packages (append (list
    ;; window managers
    ratpoison i3-wm i3status dmenu
    emacs emacs-exwm emacs-desktop-environment
    ;; terminal emulator
    xterm)
    %base-packages))

  ;; Use the "desktop" services, which include the X11
  ;; log-in service, networking with NetworkManager, and more.
  (services %desktop-services)

  ;; Allow resolution of '.local' host names with mDNS.

```

```
(name-service-switch %mdns-host-lookup-nss))
```

This example refers to the `/boot/efi` file system by its UUID, `1234-ABCD`. Replace this UUID with the right UUID on your system, as returned by the `blkid` command.

See Section 11.10.9 [Desktop Services], page 377, for the exact list of services provided by `%desktop-services`.

Again, `%desktop-services` is just a list of service objects. If you want to remove services from there, you can do so using the procedures for list filtering (see Section “SRFI-1 Filtering and Partitioning” in *GNU Guile Reference Manual*). For instance, the following expression returns a list that contains all the services in `%desktop-services` minus the Avahi service:

```
(remove (lambda (service)
          (eq? (service-kind service) avahi-service-type))
        %desktop-services)
```

Alternatively, the `modify-services` macro can be used:

```
(modify-services %desktop-services
  (delete avahi-service-type))
```

## Inspecting Services

As you work on your system configuration, you might wonder why some system service doesn’t show up or why the system is not as you expected. There are several ways to inspect and troubleshoot problems.

First, you can inspect the dependency graph of Shepherd services like so:

```
guix system shepherd-graph /etc/config.scm | \
guix shell xdot -- xdot -
```

This lets you visualize the Shepherd services as defined in `/etc/config.scm`. Each box is a service as would be shown by `sudo herd status` on the running system, and each arrow denotes a dependency (in the sense that if service *A* depends on *B*, then *B* must be started before *A*).

Not all “services” are Shepherd services though, since Guix System uses a broader definition of the term (see Section 11.10 [Services], page 275). To visualize system services and their relations at a higher level, run:

```
guix system extension-graph /etc/config.scm | \
guix shell xdot -- xdot -
```

This lets you view the *service extension graph*: how services “extend” each other, for instance by contributing to their configuration. See Section 11.19.1 [Service Composition], page 709, to understand the meaning of this graph.

Last, you may also find it useful to inspect your system configuration at the REPL (see Section 8.14 [Using Guix Interactively], page 176). Here is an example session:

```
$ guix repl
scheme@(guix-user)> ,use (gnu)
scheme@(guix-user)> (define os (load "config.scm"))
scheme@(guix-user)> ,pp (map service-kind (operating-system-services os))
$1 = (#<service-type located cabba93>
      ...)
```



See Section 11.19.3 [Service Reference], page 712, to learn about the Scheme interface to manipulate and inspect services.

## Instantiating the System

Assuming the `operating-system` declaration is stored in the `config.scm` file, the `sudo guix system reconfigure config.scm` command instantiates that configuration, and makes it the default boot entry. See Section 11.1 [Getting Started with the System], page 242, for an overview.

The normal way to change the system configuration is by updating this file and re-running `guix system reconfigure`. One should never have to touch files in `/etc` or to run commands that modify the system state such as `useradd` or `grub-install`. In fact, you must avoid that since that would not only void your warranty but also prevent you from rolling back to previous versions of your system, should you ever need to.

## The Programming Interface

At the Scheme level, the bulk of an `operating-system` declaration is instantiated with the following monadic procedure (see Section 8.11 [The Store Monad], page 161):

`operating-system-derivation os` [Monadic Procedure]

Return a derivation that builds `os`, an `operating-system` object (see Section 8.10 [Derivations], page 158).

The output of the derivation is a single directory that refers to all the packages, configuration files, and other supporting files needed to instantiate `os`.

This procedure is provided by the `(gnu system)` module. Along with `(gnu services)` (see Section 11.10 [Services], page 275), this module contains the guts of Guix System. Make sure to visit it!

## 11.3 operating-system Reference

This section summarizes all the options available in `operating-system` declarations (see Section 11.2 [Using the Configuration System], page 244).

`operating-system` [Data Type]

This is the data type representing an operating system configuration. By that, we mean all the global system configuration, not per-user configuration (see Section 11.2 [Using the Configuration System], page 244).

`kernel` (default: `linux-libre`)

The package object of the operating system kernel to use<sup>1</sup>.

`hurd` (default: `#f`)

The package object of the Hurd to be started by the kernel. When this field is set, produce a GNU/Hurd operating system. In that case, `kernel` must also be set to the `gnumach` package—the microkernel the Hurd runs on.

---

<sup>1</sup> Currently only the Linux-libre kernel is fully supported. Using GNU mach with the GNU Hurd is experimental and only available when building a virtual machine disk image.

**Warning:** This feature is experimental and only supported for disk images.

**kernel-loadable-modules** (default: `'()`)

A list of objects (usually packages) to collect loadable kernel modules from—e.g. (list `ddcci-driver-linux`).

**kernel-arguments** (default: `%default-kernel-arguments`)

List of strings or gexps representing additional arguments to pass on the command-line of the kernel—e.g., (`"console=ttyS0"`).

**bootloader**

The system bootloader configuration object. See Section 11.15 [Bootloader Configuration], page 685.

**label**

This is the label (a string) as it appears in the bootloader's menu entry. The default label includes the kernel name and version.

**keyboard-layout** (default: `#f`)

This field specifies the keyboard layout to use in the console. It can be either `#f`, in which case the default keyboard layout is used (usually US English), or a `<keyboard-layout>` record. See Section 11.8 [Keyboard Layout], page 271, for more information.

This keyboard layout is in effect as soon as the kernel has booted. For instance, it is the keyboard layout in effect when you type a passphrase if your root file system is on a `luks-device-mapping` mapped device (see Section 11.5 [Mapped Devices], page 263).

**Note:** This does *not* specify the keyboard layout used by the bootloader, nor that used by the graphical display server. See Section 11.15 [Bootloader Configuration], page 685, for information on how to specify the bootloader's keyboard layout. See Section 11.10.7 [X Window], page 355, for information on how to specify the keyboard layout used by the X Window System.

**initrd-modules** (default: `%base-initrd-modules`)

The list of Linux kernel modules that need to be available in the initial RAM disk. See Section 11.14 [Initial RAM Disk], page 682.

**initrd** (default: `base-initrd`)

A procedure that returns an initial RAM disk for the Linux kernel. This field is provided to support low-level customization and should rarely be needed for casual use. See Section 11.14 [Initial RAM Disk], page 682.

**firmware** (default: `%base-firmware`)

List of firmware packages loadable by the operating system kernel.

The default includes firmware needed for Atheros- and Broadcom-based WiFi devices (Linux-libre modules `ath9k` and `b43-open`, respectively). See Section 3.2 [Hardware Considerations], page 23, for more info on supported hardware.

**host-name**

The host name.

**mapped-devices** (default: '()')

A list of mapped devices. See Section 11.5 [Mapped Devices], page 263.

**file-systems**

A list of file systems. See Section 11.4 [File Systems], page 257.

**swap-devices** (default: '()')

A list of swap spaces. See Section 11.6 [Swap Space], page 266.

**users** (default: %base-user-accounts)

**groups** (default: %base-groups)

List of user accounts and groups. See Section 11.7 [User Accounts], page 269.

If the **users** list lacks a user account with UID 0, a “root” account with UID 0 is automatically added.

**skeletons** (default: (default-skeletons))

A list of target file name/file-like object tuples (see Section 8.12 [G-Expressions], page 165). These are the skeleton files that will be added to the home directory of newly-created user accounts.

For instance, a valid value may look like this:

```
`(("bashrc" ,(plain-file "bashrc" "echo Hello\n"))
  ("guile" ,(plain-file "guile"
                        "(use-modules (ice-9 readline))
                        (activate-readline))))
```

**issue** (default: %default-issue)

A string denoting the contents of the `/etc/issue` file, which is displayed when users log in on a text console.

**packages** (default: %base-packages)

A list of packages to be installed in the global profile, which is accessible at `/run/current-system/profile`. Each element is either a package variable or a package/output tuple. Here’s a simple example of both:

```
(cons* git ; the default "out" output
      (list git "send-email") ; another output of git
      %base-packages) ; the default set
```

The default set includes core utilities and it is good practice to install non-core utilities in user profiles (see Section 5.2 [Invoking guix package], page 37).

**timezone** (default: "Etc/UTC")

A timezone identifying string—e.g., "Europe/Paris".

You can run the `tzselect` command to find out which timezone string corresponds to your region. Choosing an invalid timezone name causes guix system to fail.

- locale** (default: `"en_US.utf8"`)  
 The name of the default locale (see Section “Locale Names” in *The GNU C Library Reference Manual*). See Section 11.9 [Locales], page 273, for more information.
- locale-definitions** (default: `%default-locale-definitions`)  
 The list of locale definitions to be compiled and that may be used at run time. See Section 11.9 [Locales], page 273.
- locale-libcs** (default: `(list glibc)`)  
 The list of GNU libc packages whose locale data and tools are used to build the locale definitions. See Section 11.9 [Locales], page 273, for compatibility considerations that justify this option.
- name-service-switch** (default: `%default-nss`)  
 Configuration of the libc name service switch (NSS)—a `<name-service-switch>` object. See Section 11.13 [Name Service Switch], page 680, for details.
- services** (default: `%base-services`)  
 A list of service objects denoting system services. See Section 11.10 [Services], page 275.
- essential-services** (default: `...`)  
 The list of “essential services”—i.e., things like instances of `system-service-type` (see Section 11.19.3 [Service Reference], page 712) and `host-name-service-type`, which are derived from the operating system definition itself. As a user you should *never* need to touch this field.
- pam-services** (default: `(base-pam-services)`)  
 Linux *pluggable authentication module* (PAM) services.
- privileged-programs** (default: `%default-privileged-programs`)  
 List of `<privileged-program>`. See Section 11.11 [Privileged Programs], page 678, for more information.
- sudoers-file** (default: `%sudoers-specification`)  
 The contents of the `/etc/sudoers` file as a file-like object (see Section 8.12 [G-Expressions], page 165).  
 This file specifies which users can use the `sudo` command, what they are allowed to do, and what privileges they may gain. The default is that only `root` and members of the `wheel` group may use `sudo`.
- this-operating-system** [Macro]  
 When used in the *lexical scope* of an operating system field definition, this identifier resolves to the operating system being defined.  
 The example below shows how to refer to the operating system being defined in the definition of the `label` field:
- ```
(use-modules (gnu) (guix))

(operating-system
```

```
;; ...
(label (package-full-name
        (operating-system-kernel this-operating-system))))
```

It is an error to refer to `this-operating-system` outside an operating system definition.

## 11.4 File Systems

The list of file systems to be mounted is specified in the `file-systems` field of the operating system declaration (see Section 11.2 [Using the Configuration System], page 244). Each file system is declared using the `file-system` form, like this:

```
(file-system
  (mount-point "/home")
  (device "/dev/sda3")
  (type "ext4"))
```

As usual, some of the fields are mandatory—those shown in the example above—while others can be omitted. These are described below.

**file-system** [Data Type]

Objects of this type represent file systems to be mounted. They contain the following members:

**type** This is a string specifying the type of the file system—e.g., `"ext4"`.

**mount-point** This designates the place where the file system is to be mounted.

**device** This names the “source” of the file system. It can be one of three things: a file system label, a file system UUID, or the name of a `/dev` node. Labels and UUIDs offer a way to refer to file systems without having to hard-code their actual device name<sup>2</sup>.

File system labels are created using the `file-system-label` procedure, UUIDs are created using `uuid`, and `/dev` nodes are plain strings. Here’s an example of a file system referred to by its label, as shown by the `e2label` command:

```
(file-system
  (mount-point "/home")
  (type "ext4")
  (device (file-system-label "my-home")))
```

UUIDs are converted from their string representation (as shown by the `tune2fs -l` command) using the `uuid` form<sup>3</sup>, like this:

```
(file-system
```

<sup>2</sup> Note that, while it is tempting to use `/dev/disk/by-uuid` and similar device names to achieve the same result, this is not recommended: These special device nodes are created by the `udev` daemon and may be unavailable at the time the device is mounted.

<sup>3</sup> The `uuid` form expects 16-byte UUIDs as defined in RFC 4122 (<https://tools.ietf.org/html/rfc4122>). This is the form of UUID used by the ext2 family of file systems and others, but it is different from “UUIDs” found in FAT file systems, for instance.

```
(mount-point "/home")
(type "ext4")
(device (uuid "4dab5feb-d176-45de-b287-9b0a6e4c01cb")))
```

When the source of a file system is a mapped device (see Section 11.5 [Mapped Devices], page 263), its `device` field *must* refer to the mapped device name—e.g., `"/dev/mapper/root-partition"`. This is required so that the system knows that mounting the file system depends on having the corresponding device mapping established.

**flags** (default: `'()`)

This is a list of symbols denoting mount flags. Recognized flags include `read-only`, `bind-mount`, `no-dev` (disallow access to special files), `no-suid` (ignore `setuid` and `setgid` bits), `no-atime` (do not update file access times), `no-diratime` (likewise for directories only), `strict-atime` (update file access time), `lazy-time` (only update time on the in-memory version of the file inode), `no-exec` (disallow program execution), `remount` (remount a mounted file system), and `shared` (make the mount shared). See Section “Mount-Unmount-Remount” in *The GNU C Library Reference Manual*, for more information on these flags.

**options** (default: `#f`)

This is either `#f`, or a string denoting mount options passed to the file system driver. See Section “Mount-Unmount-Remount” in *The GNU C Library Reference Manual*, for details.

Run `man 8 mount` for options for various file systems, but beware that what it lists as file-system-independent “mount options” are in fact flags, and belong in the `flags` field described above.

The `file-system-options->alist` and `alist->file-system-options` procedures from `(gnu system file-systems)` can be used to convert file system options given as an association list to the string representation, and vice-versa.

**mount?** (default: `#t`)

This value indicates whether to automatically mount the file system when the system is brought up. When set to `#f`, the file system gets an entry in `/etc/fstab` (read by the `mount` command) but is not automatically mounted.

**needed-for-boot?** (default: `#f`)

This Boolean value indicates whether the file system is needed when booting. If that is true, then the file system is mounted when the initial RAM disk (`initrd`) is loaded. This is always the case, for instance, for the root file system.

**check?** (default: `#t`)

This Boolean indicates whether the file system should be checked for errors before being mounted. How and when this happens can be further adjusted with the following options.

**skip-check-if-clean?** (default: **#t**)

When true, this Boolean indicates that a file system check triggered by **check?** may exit early if the file system is marked as “clean”, meaning that it was previously correctly unmounted and should not contain errors.

Setting this to false will always force a full consistency check when **check?** is true. This may take a very long time and is not recommended on healthy systems—in fact, it may reduce reliability!

Conversely, some primitive file systems like **fat** do not keep track of clean shutdowns and will perform a full scan regardless of the value of this option.

**repair** (default: **'preen'**)

When **check?** finds errors, it can (try to) repair them and continue booting. This option controls when and how to do so.

If false, try not to modify the file system at all. Checking certain file systems like **jfs** may still write to the device to replay the journal. No repairs will be attempted.

If **#t**, try to repair any errors found and assume “yes” to all questions. This will fix the most errors, but may be risky.

If **'preen'**, repair only errors that are safe to fix without human interaction. What that means is left up to the developers of each file system and may be equivalent to “none” or “all”.

**create-mount-point?** (default: **#f**)

When true, the mount point is created if it does not exist yet.

**mount-may-fail?** (default: **#f**)

When true, this indicates that mounting this file system can fail but that should not be considered an error. This is useful in unusual cases; an example of this is **efivarfs**, a file system that can only be mounted on EFI/UEFI systems.

**dependencies** (default: **'()**)

This is a list of **<file-system>** or **<mapped-device>** objects representing file systems that must be mounted or mapped devices that must be opened before (and unmounted or closed after) this one.

As an example, consider a hierarchy of mounts: **/sys/fs/cgroup** is a dependency of **/sys/fs/cgroup/cpu** and **/sys/fs/cgroup/memory**.

Another example is a file system that depends on a mapped device, for example for an encrypted partition (see Section 11.5 [Mapped Devices], page 263).

**shepherd-requirements** (default: **'()**)

This is a list of symbols denoting Shepherd requirements that must be met before mounting the file system.

As an example, an NFS file system would typically have a requirement for **networking**. In some situations, requiring **networking** might not be sufficient as **networking** being marked as started does *not* imply the

network has already been configured and in working order (for example, when using `dhcp-client-service-type`). For such cases, adding a requirement on a custom `network-online` service may be necessary. A sample `network-online` one-shot Shepherd service implementation is provided below:

```
(simple-service 'network-online shepherd-root-service-type
  (list
    (shepherd-service
      (requirement '(networking))
      (provision '(network-online))
      (documentation "Wait for the network to come up.")
      (start #~(lambda _
                  (let* ((cmd
                        "set -eux

c=0
while ! /run/setuid-programs/ping -qc1 -W1 example.org; do
sleep 1
[ \"${c} += 1\")\" -lt 30 ] || exit 1 # Limit the wait time
done
")

(status (system cmd)))
(= 0 (status:exit-val status))))
(one-shot? #t))))
```

Typically, file systems are mounted before most other Shepherd services are started. However, file systems with a non-empty `shepherd-requirements` field are mounted after Shepherd services have begun. Any Shepherd service that depends on a file system with a non-empty `shepherd-requirements` field must depend on it directly and not on the generic symbol `file-systems`.

`file-system-label str`

[Procedure]

This procedure returns an opaque file system label from *str*, a string:

```
(file-system-label "home")
⇒ #<file-system-label "home">
```

File system labels are used to refer to file systems by label rather than by device name. See above for examples.

The `(gnu system file-systems)` exports the following useful variables.

`%base-file-systems`

[Variable]

These are essential file systems that are required on normal systems, such as `%pseudo-terminal-file-system` and `%immutable-store` (see below). Operating system declarations should always contain at least these.

`%pseudo-terminal-file-system`

[Variable]

This is the file system to be mounted as `/dev/pts`. It supports *pseudo-terminals* created *via* `openpty` and similar functions (see Section “Pseudo-Terminals” in *The*



*GNU C Library Reference Manual*). Pseudo-terminals are used by terminal emulators such as `xterm`.

**%shared-memory-file-system** [Variable]

This file system is mounted as `/dev/shm` and is used to support memory sharing across processes (see Section “Memory-mapped I/O” in *The GNU C Library Reference Manual*).

**%immutable-store** [Variable]

This file system performs a read-only “bind mount” of `/gnu/store`, making it read-only for all the users including `root`. This prevents against accidental modification by software running as `root` or by system administrators.

The daemon itself is still able to write to the store: it remounts it read-write in its own “name space.”

**%binary-format-file-system** [Variable]

The `binfmt_misc` file system, which allows handling of arbitrary executable file types to be delegated to user space. This requires the `binfmt.ko` kernel module to be loaded.

**%fuse-control-file-system** [Variable]

The `fusectl` file system, which allows unprivileged users to mount and unmount user-space FUSE file systems. This requires the `fuse.ko` kernel module to be loaded.

The `(gnu system uuid)` module provides tools to deal with file system “unique identifiers” (UUIDs).

**uuid *str* [*type*]** [Procedure]

Return an opaque UUID (unique identifier) object of the given *type* (a symbol) by parsing *str* (a string):

```
(uuid "4dab5feb-d176-45de-b287-9b0a6e4c01cb")
⇒ #<<uuid> type: dce bv: ...>
```

```
(uuid "1234-ABCD" 'fat)
⇒ #<<uuid> type: fat bv: ...>
```

*type* may be one of `dce`, `iso9660`, `fat`, `ntfs`, or one of the commonly found synonyms for these.

UUIDs are another way to unambiguously refer to file systems in operating system configuration. See the examples above.

### 11.4.1 Btrfs file system

The Btrfs has special features, such as subvolumes, that merit being explained in more details. The following section attempts to cover basic as well as complex uses of a Btrfs file system with the Guix System.

In its simplest usage, a Btrfs file system can be described, for example, by:

```
(file-system
  (mount-point "/home"))
```

```
(type "btrfs")
(device (file-system-label "my-home")))
```

The example below is more complex, as it makes use of a Btrfs subvolume, named `rootfs`. The parent Btrfs file system is labeled `my-btrfs-pool`, and is located on an encrypted device (hence the dependency on `mapped-devices`):

```
(file-system
  (device (file-system-label "my-btrfs-pool"))
  (mount-point "/")
  (type "btrfs")
  (options "subvol=rootfs")
  (dependencies mapped-devices))
```

Some bootloaders, for example GRUB, only mount a Btrfs partition at its top level during the early boot, and rely on their configuration to refer to the correct subvolume path within that top level. The bootloaders operating in this way typically produce their configuration on a running system where the Btrfs partitions are already mounted and where the subvolume information is readily available. As an example, `grub-mkconfig`, the configuration generator command shipped with GRUB, reads `/proc/self/mountinfo` to determine the top-level path of a subvolume.

The Guix System produces a bootloader configuration using the operating system configuration as its sole input; it is therefore necessary to extract the subvolume name on which `/gnu/store` lives (if any) from that operating system configuration. To better illustrate, consider a subvolume named `'rootfs'` which contains the root file system data. In such situation, the GRUB bootloader would only see the top level of the root Btrfs partition, e.g.:

```
/                                (top level)
rootfs                          (subvolume directory)
  gnu                            (normal directory)
    store                        (normal directory)
[...]
```

Thus, the subvolume name must be prepended to the `/gnu/store` path of the kernel, `initrd` binaries and any other files referred to in the GRUB configuration that must be found during the early boot.

The next example shows a nested hierarchy of subvolumes and directories:

```
/                                (top level)
rootfs                          (subvolume)
  gnu                            (normal directory)
    store                        (subvolume)
[...]
```

This scenario would work without mounting the `'store'` subvolume. Mounting `'rootfs'` is sufficient, since the subvolume name matches its intended mount point in the file system hierarchy. Alternatively, the `'store'` subvolume could be referred to by setting the `subvol` option to either `/rootfs/gnu/store` or `rootfs/gnu/store`.

Finally, a more contrived example of nested subvolumes:

```
/                                (top level)
```

```

    root-snapshots      (subvolume)
    root-current        (subvolume)
    guix-store          (subvolume)
[...]
```

Here, the 'guix-store' subvolume doesn't match its intended mount point, so it is necessary to mount it. The subvolume must be fully specified, by passing its file name to the `subvol` option. To illustrate, the 'guix-store' subvolume could be mounted on `/gnu/store` by using a file system declaration such as:

```

(file-system
  (device (file-system-label "btrfs-pool-1"))
  (mount-point "/gnu/store")
  (type "btrfs")
  (options "subvol=root-snapshots/root-current/guix-store,\
compress-force=zstd,space_cache=v2"))
```

## 11.5 Mapped Devices

The Linux kernel has a notion of *device mapping*: a block device, such as a hard disk partition, can be *mapped* into another device, usually in `/dev/mapper/`, with additional processing over the data that flows through it<sup>4</sup>. A typical example is encryption device mapping: all writes to the mapped device are encrypted, and all reads are deciphered, transparently. Guix extends this notion by considering any device or set of devices that are *transformed* in some way to create a new device; for instance, RAID devices are obtained by *assembling* several other devices, such as hard disks or partitions, into a new one that behaves as one partition.

Mapped devices are declared using the `mapped-device` form, defined as follows; for examples, see below.

**mapped-device** [Data Type]

Objects of this type represent device mappings that will be made when the system boots up.

|               |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
|---------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>source</b> | This is either a string specifying the name of the block device to be mapped, such as <code>"/dev/sda3"</code> , or a list of such strings when several devices need to be assembled for creating a new one. In case of LVM this is a string specifying name of the volume group to be mapped.                                                                                                                                                                                                                                    |
| <b>target</b> | This string specifies the name of the resulting mapped device. For kernel mappers such as encrypted devices of type <code>luks-device-mapping</code> , specifying <code>"my-partition"</code> leads to the creation of the <code>"/dev/mapper/my-partition"</code> device. For RAID devices of type <code>raid-device-mapping</code> , the full device name such as <code>"/dev/md0"</code> needs to be given. LVM logical volumes of type <code>lvm-device-mapping</code> need to be specified as <code>"VGNAME-LVNAME"</code> . |

<sup>4</sup> Note that the GNU Hurd makes no difference between the concept of a “mapped device” and that of a file system: both boil down to *translating* input/output operations made on a file to operations on its backing store. Thus, the Hurd implements mapped devices, like file systems, using the generic *translator* mechanism (see Section “Translators” in *The GNU Hurd Reference Manual*).

- targets** This list of strings specifies names of the resulting mapped devices in case there are several. The format is identical to *target*.
- type** This must be a `mapped-device-kind` object, which specifies how *source* is mapped to *target*.
- arguments** This is a list of arguments specific to *type* that are passed to its `open`, `close`, and `check` procedures.
- As an example, `luks-device-mapping` supports keyword arguments—see below for details.

### `luks-device-mapping` [Variable]

This defines LUKS block device encryption using the `cryptsetup` command from the package with the same name. It relies on the `dm-crypt` Linux kernel module.

The following options may be passed *via* the `arguments` field of a mapped device:

#### `#:key-file`

If `file` is provided, unlocking is first attempted using that key file. This has an advantage of not requiring a password entry, so it can be used (for example) to unlock RAID arrays automatically on boot. If key file unlock fails, password unlock is attempted as well. Key file is not stored in the store and needs to be available at the given location at the time of the unlock attempt.

```
;; Following definition would be equivalent to running:
;; cryptsetup open --key-file /crypto.key /dev/sdb1 data
(mapped-device
 (source "/dev/sdb1")
 (target "data")
 (type luks-device-mapping)
 (arguments '(:key-file "/crypto.key")))
```

#### `#:allow-discards?`

`allow-discards?` allows the use of discard (TRIM) requests for the underlying device. This is useful for solid state drives. However, this option can have a negative security impact because it can make file system level operations visible on the physical device. For more information, refer to the description of the `--allow-discards` option in the `cryptsetup-open(8)` man page.

#### `#:extra-options`

List of additional command-line options for the `cryptsetup open` command. See the `cryptsetup-open(8)` man page for a list of supported options.

For example, here is how you could specify the `--perf-no_read_workqueue` and `--perf-no_write_workqueue` options, along with `--allow-discards`:

```
(mapped-device
 (source "/dev/sdb1")
```

```
(target "data")
(type luks-device-mapping)
(arguments '(:allow-discards? #t
            #:extra-options
            ("--perf-no_read_workqueue"
             "--perf-no_write_workqueue"))))
```

**raid-device-mapping** [Variable]

This defines a RAID device, which is assembled using the `mdadm` command from the package with the same name. It requires a Linux kernel module for the appropriate RAID level to be loaded, such as `raid456` for RAID-4, RAID-5 or RAID-6, or `raid10` for RAID-10.

**lvm-device-mapping** [Variable]

This defines one or more logical volumes for the Linux Logical Volume Manager (LVM) (<https://www.sourceware.org/lvm2/>). The volume group is activated by the `vgchange` command from the `lvm2` package.

The following example specifies a mapping from `/dev/sda3` to `/dev/mapper/home` using LUKS—the Linux Unified Key Setup (<https://gitlab.com/cryptsetup/cryptsetup>), a standard mechanism for disk encryption. The `/dev/mapper/home` device can then be used as the device of a file-system declaration (see Section 11.4 [File Systems], page 257).

```
(mapped-device
  (source "/dev/sda3")
  (target "home")
  (type luks-device-mapping))
```

Alternatively, to become independent of device numbering, one may obtain the LUKS UUID (*unique identifier*) of the source device by a command like:

```
cryptsetup luksUUID /dev/sda3
```

and use it as follows:

```
(mapped-device
  (source (uuid "cb67fc72-0d54-4c88-9d4b-b225f30b0f44"))
  (target "home")
  (type luks-device-mapping))
```

It is also desirable to encrypt swap space, since swap space may contain sensitive data. One way to accomplish that is to use a swap file in a file system on a device mapped via LUKS encryption. In this way, the swap file is encrypted because the entire device is encrypted. See Section 11.6 [Swap Space], page 266, or See Section 3.4 [Disk Partitioning], page 25, for an example.

A RAID device formed of the partitions `/dev/sda1` and `/dev/sdb1` may be declared as follows:

```
(mapped-device
  (source (list "/dev/sda1" "/dev/sdb1"))
  (target "/dev/md0")
  (type raid-device-mapping))
```

The `/dev/md0` device can then be used as the **device** of a **file-system** declaration (see Section 11.4 [File Systems], page 257). Note that the RAID level need not be given; it is chosen during the initial creation and formatting of the RAID device and is determined automatically later.

LVM logical volumes “alpha” and “beta” from volume group “vg0” can be declared as follows:

```
(mapped-device
 (source "vg0")
 (targets (list "vg0-alpha" "vg0-beta"))
 (type lvm-device-mapping))
```

Devices `/dev/mapper/vg0-alpha` and `/dev/mapper/vg0-beta` can then be used as the **device** of a **file-system** declaration (see Section 11.4 [File Systems], page 257).

## 11.6 Swap Space

Swap space, as it is commonly called, is a disk area specifically designated for paging: the process in charge of memory management (the Linux kernel or Hurd’s default pager) can decide that some memory pages stored in RAM which belong to a running program but are unused should be stored on disk instead. It unloads those from the RAM, freeing up precious fast memory, and writes them to the swap space. If the program tries to access that very page, the memory management process loads it back into memory for the program to use.

A common misconception about swap is that it is only useful when small amounts of RAM are available to the system. However, it should be noted that kernels often use all available RAM for disk access caching to make I/O faster, and thus paging out unused portions of program memory will expand the RAM available for such caching.

For a more detailed description of how memory is managed from the viewpoint of a monolithic kernel, see Section “Memory Concepts” in *The GNU C Library Reference Manual*.

The Linux kernel has support for swap partitions and swap files: the former uses a whole disk partition for paging, whereas the second uses a file on a file system for that (the file system driver needs to support it). On a comparable setup, both have the same performance, so one should consider ease of use when deciding between them. Partitions are “simpler” and do not need file system support, but need to be allocated at disk formatting time (logical volumes notwithstanding), whereas files can be allocated and deallocated at any time.

Swap space is also required to put the system into *hibernation* (also called *suspend to disk*), whereby memory is dumped to swap before shutdown so it can be restored when the machine is eventually restarted. Hibernation uses at most half the size of the RAM in the configured swap space. The Linux kernel needs to know about the swap space to be used to resume from hibernation on boot (*via* a kernel argument). When using a swap file, its offset in the device holding it also needs to be given to the kernel; that value has to be updated if the file is initialized again as swap—e.g., because its size was changed.

Note that swap space is not zeroed on shutdown, so sensitive data (such as passwords) may linger on it if it was paged out. As such, you should consider having your swap reside on an encrypted device (see Section 11.5 [Mapped Devices], page 263).

**swap-space** [Data Type]

Objects of this type represent swap spaces. They contain the following members:

**target** The device or file to use, either a UUID, a **file-system-label** or a string, as in the definition of a **file-system** (see Section 11.4 [File Systems], page 257).

**dependencies** (default: '()')

A list of **file-system** or **mapped-device** objects, upon which the availability of the space depends. Note that just like for **file-system** objects, dependencies which are needed for boot and mounted in early userspace are not managed by the Shepherd, and so automatically filtered out for you.

**priority** (default: #f)

Only supported by the Linux kernel. Either #f to disable swap priority, or an integer between 0 and 32767. The kernel will first use swap spaces of higher priority when paging, and use same priority spaces on a round-robin basis. The kernel will use swap spaces without a set priority after prioritized spaces, and in the order that they appeared in (not round-robin).

**discard?** (default: #f)

Only supported by the Linux kernel. When true, the kernel will notify the disk controller of discarded pages, for example with the TRIM operation on Solid State Drives.

Here are some examples:

```
(swap-space (target (uuid "4dab5feb-d176-45de-b287-9b0a6e4c01cb")))
```

Use the swap partition with the given UUID. You can learn the UUID of a Linux swap partition by running `swaptlabel device`, where *device* is the `/dev` file name of that partition.

```
(swap-space
  (target (file-system-label "swap"))
  (dependencies mapped-devices))
```

Use the partition with label `swap`, which can be found after all the *mapped-devices* mapped devices have been opened. Again, the `swaptlabel` command allows you to view and change the label of a Linux swap partition.

Here's a more involved example with the corresponding **file-systems** part of an **operating-system** declaration.

```
(file-systems
  (list (file-system
        (device (file-system-label "root"))
        (mount-point "/")
        (type "ext4"))
        (file-system
        (device (file-system-label "btrfs"))
        (mount-point "/btrfs"))
```

```

        (type "btrfs")))))

(swap-devices
  (list
    (swap-space
      (target "/btrfs/swapfile")
      (dependencies (filter (file-system-mount-point-predicate "/btrfs")
                            file-systems))))))

```

Use the file `/btrfs/swapfile` as swap space, which depends on the file system mounted at `/btrfs`. Note how we use Guile's filter to select the file system in an elegant fashion!

```

(swap-devices
  (list
    (swap-space
      (target "/dev/mapper/my-swap")
      (dependencies mapped-devices))))

(kernel-arguments
  (cons* "resume=/dev/mapper/my-swap"
        %default-kernel-arguments))

```

The above snippet of an `operating-system` declaration enables the mapped device `/dev/mapper/my-swap` (which may be part of an encrypted device) as swap space, and tells the kernel to use it for hibernation via the `resume` kernel argument (see Section 11.3 [operating-system Reference], page 253, `kernel-arguments`).

```

(swap-devices
  (list
    (swap-space
      (target "/swapfile")
      (dependencies (filter (file-system-mount-point-predicate "/")
                            file-systems))))))

(kernel-arguments
  (cons* "resume=/dev/sda3"           ;device that holds /swapfile
        "resume_offset=92514304"    ;offset of /swapfile on device
        %default-kernel-arguments))

```

This other snippet of `operating-system` enables the swap file `/swapfile` for hibernation by telling the kernel about the partition containing it (`resume` argument) and its offset on that partition (`resume_offset` argument). The latter value can be found in the output of the command `filefrag -e` as the first number right under the `physical_offset` column header (the second command extracts its value directly):

```

$ sudo filefrag -e /swapfile
Filesystem type is: ef53
File size of /swapfile is 2463842304 (601524 blocks of 4096 bytes)
ext:      logical_offset:      physical_offset: length:   expected: flags:
  0:         0..    2047:   92514304..  92516351:   2048:
...
$ sudo filefrag -e /swapfile | grep '^ *0:' | cut -d: -f3 | cut -d. -f1
92514304

```



## 11.7 User Accounts

User accounts and groups are entirely managed through the `operating-system` declaration. They are specified with the `user-account` and `user-group` forms:

```
(user-account
  (name "alice")
  (group "users")
  (supplementary-groups '("wheel"      ;allow use of sudo, etc.
                        "audio"      ;sound card
                        "video"      ;video devices such as webcams
                        "cdrom")) ;the good ol' CD-ROM
  (comment "Bob's sister"))
```

Here's a user account that uses a different shell and a custom home directory (the default would be `"/home/bob"`):

```
(user-account
  (name "bob")
  (group "users")
  (comment "Alice's bro")
  (shell (file-append zsh "/bin/zsh"))
  (home-directory "/home/robert"))
```

When booting or upon completion of `guix system reconfigure`, the system ensures that only the user accounts and groups specified in the `operating-system` declaration exist, and with the specified properties. Thus, account or group creations or modifications made by directly invoking commands such as `useradd` are lost upon reconfiguration or reboot. This ensures that the system remains exactly as declared.

**user-account** [Data Type]

Objects of this type represent user accounts. The following members may be specified:

- name**           The name of the user account.
- group**           This is the name (a string) or identifier (a number) of the user group this account belongs to.
- supplementary-groups** (default: `'()`)  
                     Optionally, this can be defined as a list of group names that this account belongs to.
- uid** (default: `#f`)  
                     This is the user ID for this account (a number), or `#f`. In the latter case, a number is automatically chosen by the system when the account is created.
- comment** (default: `""`)  
                     A comment about the account, such as the account owner's full name.  
                     Note that, for non-system accounts, users are free to change their real name as it appears in `/etc/passwd` using the `chfn` command. When they do, their choice prevails over the system administrator's choice; re-configuring does *not* change their name.

**home-directory**

This is the name of the home directory for the account.

**create-home-directory?** (default: **#t**)

Indicates whether the home directory of this account should be created if it does not exist yet.

**shell** (default: **Bash**)

This is a G-expression denoting the file name of a program to be used as the shell (see Section 8.12 [G-Expressions], page 165). For example, you would refer to the Bash executable like this:

```
(file-append bash "/bin/bash")
```

... and to the Zsh executable like that:

```
(file-append zsh "/bin/zsh")
```

**system?** (default: **#f**)

This Boolean value indicates whether the account is a “system” account. System accounts are sometimes treated specially; for instance, graphical login managers do not list them.

**password** (default: **#f**)

You would normally leave this field to **#f**, initialize user passwords as **root** with the **passwd** command, and then let users change it with **passwd**. Passwords set with **passwd** are of course preserved across reboot and reconfiguration.

If you *do* want to set an initial password for an account, then this field must contain the encrypted password, as a string. You can use the **crypt** procedure for this purpose:

```
(user-account
  (name "charlie")
  (group "users")

  ;; Specify a SHA-512-hashed initial password.
  (password (crypt "InitialPassword!" "$6$abc")))
```

**Note:** The hash of this initial password will be available in a file in **/gnu/store**, readable by all the users, so this method must be used with care.

See Section “Passphrase Storage” in *The GNU C Library Reference Manual*, for more information on password encryption, and Section “Encryption” in *GNU Guile Reference Manual*, for information on Guile’s **crypt** procedure.

User group declarations are even simpler:

```
(user-group (name "students"))
```

**user-group**

[Data Type]

This type is for, well, user groups. There are just a few fields:

**name**        The name of the group.

**id** (default: **#f**)

The group identifier (a number). If **#f**, a new number is automatically allocated when the group is created.

**system?** (default: **#f**)

This Boolean value indicates whether the group is a “system” group. System groups have low numerical IDs.

**password** (default: **#f**)

What, user groups can have a password? Well, apparently yes. Unless **#f**, this field specifies the password of the group.

For convenience, a variable lists all the basic user groups one may expect:

**%base-groups** [Variable]

This is the list of basic user groups that users and/or packages expect to be present on the system. This includes groups such as “root”, “wheel”, and “users”, as well as groups used to control access to specific devices such as “audio”, “disk”, and “cdrom”.

**%base-user-accounts** [Variable]

This is the list of basic system accounts that programs may expect to find on a GNU/Linux system, such as the “nobody” account.

Note that the “root” account is not included here. It is a special-case and is automatically added whether or not it is specified.

The Linux kernel also implements *subordinate user and group IDs*, or “subids”, which are used to map the ID of a user and group to several IDs inside separate name spaces—inside “containers”. See [subordinate-user-group-ids], page 663, for information on how to configure it.

## 11.8 Keyboard Layout

To specify what each key of your keyboard does, you need to tell the operating system what *keyboard layout* you want to use. The default, when nothing is specified, is the US English QWERTY layout for 105-key PC keyboards. However, German speakers will usually prefer the German QWERTZ layout, French speakers will want the AZERTY layout, and so on; hackers might prefer Dvorak or bépo, and they might even want to further customize the effect of some of the keys. This section explains how to get that done.

There are three components that will want to know about your keyboard layout:

- The *bootloader* may want to know what keyboard layout you want to use (see Section 11.15 [Bootloader Configuration], page 685). This is useful if you want, for instance, to make sure that you can type the passphrase of your encrypted root partition using the right layout.
- The *operating system kernel*, Linux, will need that so that the console is properly configured (see Section 11.3 [operating-system Reference], page 253).
- The *graphical display server*, usually Xorg, also has its own idea of the keyboard layout (see Section 11.10.7 [X Window], page 355).

Guix allows you to configure all three separately but, fortunately, it allows you to share the same keyboard layout for all three components.

Keyboard layouts are represented by records created by the `keyboard-layout` procedure of `(gnu system keyboard)`. Following the X Keyboard extension (XKB), each layout has four attributes: a name (often a language code such as “fi” for Finnish or “jp” for Japanese), an optional variant name, an optional keyboard model name, and a possibly empty list of additional options. In most cases the layout name is all you care about.

`keyboard-layout` *name* [*variant*] [*#:model*] [*#:options* '()] [Procedure]

Return a new keyboard layout with the given *name* and *variant*.

*name* must be a string such as “fr”; *variant* must be a string such as “bepo” or “nodeadkeys”. See the `xkeyboard-config` package for valid options.

Here are a few examples:

```
;; The German QWERTZ layout. Here we assume a standard
;; "pc105" keyboard model.
(keyboard-layout "de")

;; The bépo variant of the French layout.
(keyboard-layout "fr" "bepo")

;; The Catalan layout.
(keyboard-layout "es" "cat")

;; Arabic layout with "Alt-Shift" to switch to US layout.
(keyboard-layout "ar,us" #:options '("grp:alt_shift_toggle"))

;; The Latin American Spanish layout. In addition, the
;; "Caps Lock" key is used as an additional "Ctrl" key,
;; and the "Menu" key is used as a "Compose" key to enter
;; accented letters.
(keyboard-layout "latam"
 #:options '("ctrl:nocaps" "compose:menu"))

;; The Russian layout for a ThinkPad keyboard.
(keyboard-layout "ru" #:model "thinkpad")

;; The "US international" layout, which is the US layout plus
;; dead keys to enter accented characters. This is for an
;; Apple MacBook keyboard.
(keyboard-layout "us" "intl" #:model "macbook78")
```

See the `share/X11/xkb` directory of the `xkeyboard-config` package for a complete list of supported layouts, variants, and models.

Let’s say you want your system to use the Turkish keyboard layout throughout your system—bootloader, console, and Xorg. Here’s what your system configuration would look like:

```
;; Using the Turkish layout for the bootloader, the console,
```

```
;; and for Xorg.

(operating-system
  ;; ...
  (keyboard-layout (keyboard-layout "tr")) ;for the console
  (bootloader (bootloader-configuration
    (bootloader grub-efi-bootloader)
    (targets '("/boot/efi"))
    (keyboard-layout keyboard-layout))) ;for GRUB
  (services (cons (set-xorg-configuration
    (xorg-configuration ;for Xorg
      (keyboard-layout keyboard-layout)))
    %desktop-services)))
```

In the example above, for GRUB and for Xorg, we just refer to the `keyboard-layout` field defined above, but we could just as well refer to a different layout. The `set-xorg-configuration` procedure communicates the desired Xorg configuration to the graphical log-in manager, by default GDM.

We’ve discussed how to specify the *default* keyboard layout of your system when it starts, but you can also adjust it at run time:

- If you’re using GNOME, its settings panel has a “Region & Language” entry where you can select one or more keyboard layouts.
- Under Xorg, the `setxkbmap` command (from the same-named package) allows you to change the current layout. For example, this is how you would change the layout to US Dvorak:

```
setxkbmap us dvorak
```

- The `loadkeys` command changes the keyboard layout in effect in the Linux console. However, note that `loadkeys` does *not* use the XKB keyboard layout categorization described above. The command below loads the French bépo layout:

```
loadkeys fr-bepo
```

## 11.9 Locales

A *locale* defines cultural conventions for a particular language and region of the world (see Section “Locales” in *The GNU C Library Reference Manual*). Each locale has a name that typically has the form *language\_territory.codeset*—e.g., `fr_LU.utf8` designates the locale for the French language, with cultural conventions from Luxembourg, and using the UTF-8 encoding.

Usually, you will want to specify the default locale for the machine using the `locale` field of the `operating-system` declaration (see Section 11.3 [operating-system Reference], page 253).

The selected locale is automatically added to the *locale definitions* known to the system if needed, with its codeset inferred from its name—e.g., `bo_CN.utf8` will be assumed to use the UTF-8 codeset. Additional locale definitions can be specified in the `locale-definitions` slot of `operating-system`—this is useful, for instance, if the codeset could not be inferred from the locale name. The default set of locale definitions includes some widely used locales, but not all the available locales, in order to save space.

For instance, to add the North Frisian locale for Germany, the value of that field may be:

```
(cons (locale-definition
      (name "fy_DE.utf8") (source "fy_DE"))
      %default-locale-definitions)
```

Likewise, to save space, one might want `locale-definitions` to list only the locales that are actually used, as in:

```
(list (locale-definition
      (name "ja_JP.eucjp") (source "ja_JP")
      (charset "EUC-JP")))
```

The compiled locale definitions are available at `/run/current-system/locale/X.Y`, where `X.Y` is the libc version, which is the default location where the GNU libc provided by Guix looks for locale data. This can be overridden using the `LOCPATH` environment variable (see [locales-and-locpath], page 18).

The `locale-definition` form is provided by the `(gnu system locale)` module. Details are given below.

**locale-definition** [Data Type]

This is the data type of a locale definition.

**name**        The name of the locale. See Section “Locale Names” in *The GNU C Library Reference Manual*, for more information on locale names.

**source**      The name of the source for that locale. This is typically the *language\_territory* part of the locale name.

**charset** (default: "UTF-8")  
              The “character set” or “code set” for that locale, as defined by IANA (<https://www.iana.org/assignments/character-sets>).

**%default-locale-definitions** [Variable]

A list of commonly used UTF-8 locales, used as the default value of the `locale-definitions` field of `operating-system` declarations.

These locale definitions use the *normalized codeset* for the part that follows the dot in the name (see Section “Using gettextized software” in *The GNU C Library Reference Manual*). So for instance it has `uk_UA.utf8` but *not*, say, `uk_UA.UTF-8`.

### 11.9.1 Locale Data Compatibility Considerations

`operating-system` declarations provide a `locale-libcs` field to specify the GNU libc packages that are used to compile locale declarations (see Section 11.3 [operating-system Reference], page 253). “Why would I care?”, you may ask. Well, it turns out that the binary format of locale data is occasionally incompatible from one libc version to another.

For instance, a program linked against libc version 2.21 is unable to read locale data produced with libc 2.22; worse, that program *aborts* instead of simply ignoring the incompatible locale data<sup>5</sup>. Similarly, a program linked against libc 2.22 can read most, but not

<sup>5</sup> Versions 2.23 and later of GNU libc will simply skip the incompatible locale data, which is already an improvement.

all, of the locale data from libc 2.21 (specifically, `LC_COLLATE` data is incompatible); thus calls to `setlocale` may fail, but programs will not abort.

The “problem” with Guix is that users have a lot of freedom: They can choose whether and when to upgrade software in their profiles, and might be using a libc version different from the one the system administrator used to build the system-wide locale data.

Fortunately, unprivileged users can also install their own locale data and define `GUIX_LOCPATH` accordingly (see [locales-and-locpath], page 18).

Still, it is best if the system-wide locale data at `/run/current-system/locale` is built for all the libc versions actually in use on the system, so that all the programs can access it—this is especially crucial on a multi-user system. To do that, the administrator can specify several libc packages in the `locale-libcs` field of `operating-system`:

```
(use-package-modules base)

(operating-system
 ;; ...
 (locale-libcs (list glibc-2.21 (canonical-package glibc))))
```

This example would lead to a system containing locale definitions for both libc 2.21 and the current version of libc in `/run/current-system/locale`.

## 11.10 Services

An important part of preparing an `operating-system` declaration is listing *system services* and their configuration (see Section 11.2 [Using the Configuration System], page 244). System services are typically daemons launched when the system boots, or other actions needed at that time—e.g., configuring network access.

Guix has a broad definition of “service” (see Section 11.19.1 [Service Composition], page 709), but many services are managed by the GNU Shepherd (see Section 11.19.4 [Shepherd Services], page 717). On a running system, the `herd` command allows you to list the available services, show their status, start and stop them, or do other specific operations (see Section “Jump Start” in *The GNU Shepherd Manual*). For example:

```
# herd status
```

The above command, run as `root`, lists the currently defined services. The `herd doc` command shows a synopsis of the given service and its associated actions:

```
# herd doc nscd
Run libc's name service cache daemon (nscd).

# herd doc nscd action invalidate
invalidate: Invalidate the given cache--e.g., 'hosts' for host name lookups.■
```

The `start`, `stop`, and `restart` sub-commands have the effect you would expect. For instance, the commands below stop the `nscd` service and restart the Xorg display server:

```
# herd stop nscd
Service nscd has been stopped.
# herd restart xorg-server
Service xorg-server has been stopped.
Service xorg-server has been started.
```

For some services, `herd configuration` returns the name of the service’s configuration file, which can be handy to inspect its configuration:

```
# herd configuration sshd
/gnu/store/...-sshd_config
```

The following sections document the available services, starting with the core services, that may be used in an `operating-system` declaration.

### 11.10.1 Base Services

The `(gnu services base)` module provides definitions for the basic services that one expects from the system. The services exported by this module are listed below.

**%base-services** [Variable]

This variable contains a list of basic services (see Section 11.19.2 [Service Types and Services], page 710, for more information on service objects) one would expect from the system: a login service (`mingetty`) on each `tty`, `syslogd`, the `libc` name service cache daemon (`nscd`), the `udev` device manager, and more.

This is the default value of the `services` field of `operating-system` declarations. Usually, when customizing a system, you will want to append services to `%base-services`, like this:

```
(append (list (service avahi-service-type)
              (service openssh-service-type))
        %base-services)
```

**special-files-service-type** [Variable]

This is the service that sets up “special files” such as `/bin/sh`; an instance of it is part of `%base-services`.

The value associated with `special-files-service-type` services must be a list of two-element lists where the first element is the “special file” and the second element is its target. By default it is:

```
`((" /bin/sh" ,(file-append bash "/bin/sh"))
  ("/usr/bin/env" ,(file-append coreutils "/bin/env")))
```

If you want to add, say, `/bin/bash` to your system, you can change it to:

```
`((" /bin/sh" ,(file-append bash "/bin/sh"))
  ("/usr/bin/env" ,(file-append coreutils "/bin/env"))
  ("/bin/bash" ,(file-append bash "/bin/bash")))
```

Since this is part of `%base-services`, you can use `modify-services` to customize the set of special files (see Section 11.19.3 [Service Reference], page 712). But the simple way to add a special file is *via* the `extra-special-file` procedure (see below).

**extra-special-file** *file target* [Procedure]

Use *target* as the “special file” *file*.

For example, adding the following lines to the `services` field of your operating system declaration leads to a `/usr/bin/env` symlink:

```
(extra-special-file "/usr/bin/env"
  (file-append coreutils "/bin/env"))
```



This procedure is meant for `/bin/sh`, `/usr/bin/env` and similar targets. In particular, use for targets under `/etc` might not work as expected if the target is managed by Guix in other ways.

The special files persist on your system even after the `extra-special-file` call has been removed from your system config and the system has been reconfigured. The special files need to be removed manually.

**host-name-service-type** [Variable]

Type of the service that sets the system host name, whose value is a string. This service is included in `operating-system` by default (see [operating-system-essential-services], page 256).

**console-font-service-type** [Variable]

Install the given fonts on the specified ttys (fonts are per virtual console on the kernel Linux). The value of this service is a list of tty/font pairs. The font can be the name of a font provided by the `kbd` package or any valid argument to `setfont`, as in this example:

```
((("tty1" . "LatGrkCyr-8x16")
  ("tty2" . ,(file-append
               font-tamzen
               "/share/kbd/consolefonts/TamzenForPowerline10x20.psf")))
 ("tty3" . ,(file-append
               font-terminus
               "/share/consolefonts/ter-132n"))) ; for HDPI
```

**hosts-service-type** [Variable]

Type of the service that populates the entries for (`/etc/hosts`). This service type can be *extended* by passing it a list of host records.

The example below shows how to add two entries to `/etc/hosts`:

```
(simple-service 'add-extra-hosts
               hosts-service-type
               (list (host "192.0.2.1" "example.com"
                          ("example.net" "example.org"))
                   (host "2001:db8::1" "example.com"
                          ("example.net" "example.org"))))
```

**Note:** By default `/etc/hosts` comes with the following entries:

```
127.0.0.1 localhost host-name
::1      localhost host-name
```

For most setups this is what you want though if you find yourself in the situation where you want to change the default entries, you can do so in `operating-system` via `modify-services` (see Section 11.19.3 [Service Reference], page 712).

The following example shows how to unset *host-name* from being an alias of `localhost`.

```
(operating-system
;; ...
```

```

(essential-services
 (modify-services
  (operating-system-default-essential-services this-operating-system)
  (hosts-service-type config => (list
    (host "127.0.0.1" "localhost")
    (host "::1" "localhost"))))))

```

**host** *address canonical-name* [*aliases*] [Procedure]

Return a new record for the host at *address* with the given *canonical-name* and possibly *aliases*.

*address* must be a string denoting a valid IPv4 or IPv6 address, and *canonical-name* and the strings listed in *aliases* must be valid host names.

**login-service-type** [Variable]

Type of the service that provides a console login service, whose value is a `<login-configuration>` object.

**login-configuration** [Data Type]

Data type representing the configuration of login, which specifies the MOTD (message of the day), among other things.

**motd** A file-like object containing the “message of the day”.

**allow-empty-passwords?** (default: `#t`)

Allow empty passwords by default so that first-time users can log in when the ‘root’ account has just been created.

**mingetty-service-type** [Variable]

Type of the service that runs Mingetty, an implementation of the virtual console log-in. The value for this service is a `<mingetty-configuration>` object.

**mingetty-configuration** [Data Type]

Data type representing the configuration of Mingetty, which specifies the tty to run, among other things.

**tty** The name of the console this Mingetty runs on—e.g., `"tty1"`.

**auto-login** (default: `#f`)

When true, this field must be a string denoting the user name under which the system automatically logs in. When it is `#f`, a user name and password must be entered to log in.

**login-program** (default: `#f`)

This must be either `#f`, in which case the default log-in program is used (`login` from the Shadow tool suite), or a gexp denoting the name of the log-in program.

**login-pause?** (default: `#f`)

When set to `#t` in conjunction with *auto-login*, the user will have to press a key before the log-in shell is launched.

`clear-on-logout?` (default: `#t`)

When set to `#t`, the screen will be cleared before showing the login prompt. The field name is bit unfortunate, since it controls clearing also before the initial login, not just after a logout.

`delay` (default: `#f`)

When set to a number, sleep that many seconds after startup.

`print-issue` (default: `#t`)

When set to `#t`, write out a new line and the content of `/etc/issue`. Value of `'no-nl` can be used to suppress the new line.

`print-hostname` (default: `#t`)

When set to `#t`, print the host name before the login prompt. The host name is printed up to the first dot. Can be set to `'long` to print the full host name.

`nice` (default: `#f`)

When set to a number, change the process priority using `nice`.

`working-directory` (default: `#f`)

When set to a string, change into that directory before calling the login program.

`root-directory` (default: `#f`)

When set to a string, use this directory at the process's root directory.

`shepherd-requirement`

List of shepherd requirements. Unless you know what you are doing, it is recommended to extend the default list instead of overriding it.

As an example, when using auto-login on a system with `elogind`, it is necessary to wait on the `'dbus-system` service:

```
(modify-services %base-services
  (mingetty-service-type config =>
    (mingetty-configuration
      (inherit config)
      ;; Automatically log in as "guest".
      (auto-login "guest")
      (shepherd-requirement
        (cons 'dbus-system
              (mingetty-configuration-shepherd-requirement
                config))))))
```

`mingetty` (default: `mingetty`)

The Mingetty package to use.

`agetty-service-type`

[Variable]

Type of the service that runs `agetty`, which implements virtual and serial console log-in. The value for this service is a `<agetty-configuration>` object.

**agetty-configuration**

[Data Type]

Data type representing the configuration of agetty, which specifies the tty to run, among other things<sup>6</sup>.

**tty** The name of the console this agetty runs on, as a string—e.g., **"ttyS0"**. This argument is optional, it will default to a reasonable default serial port used by the kernel Linux.

For this, if there is a value for an option **agetty.tty** in the kernel command line, agetty will extract the device name of the serial port from it and use that.

If not and if there is a value for an option **console** with a tty in the Linux command line, agetty will extract the device name of the serial port from it and use that.

In both cases, agetty will leave the other serial device settings (baud rate etc.) alone—in the hope that Linux pinned them to the correct values.

**baud-rate** (default: **#f**)

A string containing a comma-separated list of one or more baud rates, in descending order.

**term** (default: **#f**)

A string containing the value used for the **TERM** environment variable.

**eight-bits?** (default: **#f**)

When **#t**, the tty is assumed to be 8-bit clean, and parity detection is disabled.

**auto-login** (default: **#f**)

When passed a login name, as a string, the specified user will be logged in automatically without prompting for their login name or password.

**no-reset?** (default: **#f**)

When **#t**, don't reset terminal cflags (control modes).

**host** (default: **#f**)

This accepts a string containing the "login\_host", which will be written into the **/var/run/utmpx** file.

**remote?** (default: **#f**)

When set to **#t** in conjunction with *host*, this will add an **-r** fakehost option to the command line of the login program specified in *login-program*.

**flow-control?** (default: **#f**)

When set to **#t**, enable hardware (RTS/CTS) flow control.

**no-issue?** (default: **#f**)

When set to **#t**, the contents of the **/etc/issue** file will not be displayed before presenting the login prompt.

**init-string** (default: **#f**)

This accepts a string that will be sent to the tty or modem before sending anything else. It can be used to initialize a modem.

---

<sup>6</sup> See the **agetty(8)** man page for more information.

- no-clear?** (default: **#f**)  
When set to **#t**, *agetty* will not clear the screen before showing the login prompt.
- login-program** (default: (file-append shadow "/bin/login"))  
This must be either a gexp denoting the name of a log-in program, or unset, in which case the default value is the *login* from the Shadow tool suite.
- local-line** (default: **#f**)  
Control the CLOCAL line flag. This accepts one of three symbols as arguments, 'auto', 'always, or 'never. If **#f**, the default value chosen by *agetty* is 'auto.
- extract-baud?** (default: **#f**)  
When set to **#t**, instruct *agetty* to try to extract the baud rate from the status messages produced by certain types of modems.
- skip-login?** (default: **#f**)  
When set to **#t**, do not prompt the user for a login name. This can be used with *login-program* field to use non-standard login systems.
- no-newline?** (default: **#f**)  
When set to **#t**, do not print a newline before printing the */etc/issue* file.
- login-options** (default: **#f**)  
This option accepts a string containing options that are passed to the login program. When used with the *login-program*, be aware that a malicious user could try to enter a login name containing embedded options that could be parsed by the login program.
- login-pause** (default: **#f**)  
When set to **#t**, wait for any key before showing the login prompt. This can be used in conjunction with *auto-login* to save memory by lazily spawning shells.
- chroot** (default: **#f**)  
Change root to the specified directory. This option accepts a directory path as a string.
- hangup?** (default: **#f**)  
Use the Linux system call *vhangup* to do a virtual hangup of the specified terminal.
- keep-baud?** (default: **#f**)  
When set to **#t**, try to keep the existing baud rate. The baud rates from *baud-rate* are used when *agetty* receives a **BREAK** character.
- timeout** (default: **#f**)  
When set to an integer value, terminate if no user name could be read within *timeout* seconds.

**detect-case?** (default: **#f**)

When set to **#t**, turn on support for detecting an uppercase-only terminal. This setting will detect a login name containing only uppercase letters as indicating an uppercase-only terminal and turn on some upper-to-lower case conversions. Note that this will not support Unicode characters.

**wait-cr?** (default: **#f**)

When set to **#t**, wait for the user or modem to send a carriage-return or linefeed character before displaying */etc/issue* or login prompt. This is typically used with the *init-string* option.

**no-hints?** (default: **#f**)

When set to **#t**, do not print hints about Num, Caps, and Scroll locks.

**no-hostname?** (default: **#f**)

By default, the hostname is printed. When this option is set to **#t**, no hostname will be shown at all.

**long-hostname?** (default: **#f**)

By default, the hostname is only printed until the first dot. When this option is set to **#t**, the fully qualified hostname by *gethostname* or *getaddrinfo* is shown.

**erase-characters** (default: **#f**)

This option accepts a string of additional characters that should be interpreted as backspace when the user types their login name.

**kill-characters** (default: **#f**)

This option accepts a string that should be interpreted to mean “ignore all previous characters” (also called a “kill” character) when the user types their login name.

**chdir** (default: **#f**)

This option accepts, as a string, a directory path that will be changed to before login.

**delay** (default: **#f**)

This options accepts, as an integer, the number of seconds to sleep before opening the tty and displaying the login prompt.

**nice** (default: **#f**)

This option accepts, as an integer, the nice value with which to run the login program.

**extra-options** (default: **'()**)

This option provides an “escape hatch” for the user to provide arbitrary command-line arguments to *agetty* as a list of strings.

**shepherd-requirement** (default: **'()**)

The option can be used to provides extra shepherd requirements (for example **'syslogd**) to the respective **'term-\*** shepherd service.

**kmscon-service-type** [Variable]

Type of the service that runs kmscon (<https://www.freedesktop.org/wiki/Software/kmscon>), which implements virtual console log-in. The value for this service is a <kmscon-configuration> object.

**kmscon-configuration** [Data Type]

Data type representing the configuration of Kmscon, which specifies the tty to run, among other things.

**virtual-terminal**

The name of the console this Kmscon runs on—e.g., "tty1".

**login-program** (default: #~(string-append #\$shadow "/bin/login"))

A gexp denoting the name of the log-in program. The default log-in program is `login` from the Shadow tool suite.

**login-arguments** (default: '("-p"))

A list of arguments to pass to `login`.

**auto-login** (default: #f)

When passed a login name, as a string, the specified user will be logged in automatically without prompting for their login name or password.

**hardware-acceleration?** (default: #f)

Whether to use hardware acceleration.

**font-engine** (default: "pango")

Font engine used in Kmscon.

**font-size** (default: 12)

Font size used in Kmscon.

**keyboard-layout** (default: #f)

If this is #f, Kmscon uses the default keyboard layout—usually US English (“qwerty”) for a 105-key PC keyboard.

Otherwise this must be a `keyboard-layout` object specifying the keyboard layout. See Section 11.8 [Keyboard Layout], page 271, for more information on how to specify the keyboard layout.

**kmscon** (default: *kmscon*)

The Kmscon package to use.

**nscd-service-type** [Variable]

Type of the service that runs the libc nscd (name service cache daemon), whose value is an <nscd-configuration> object.

For convenience, the Shepherd service for nscd provides the following actions:

**invalidate**

This invalidate the given cache. For instance, running:

```
herd invalidate nscd hosts
```

invalidates the host name lookup cache of nscd.

**statistics**

Running **herd statistics nsd** displays information about nsd usage and caches.

**nsd-configuration**

[Data Type]

Data type representing the nsd (name service cache daemon) configuration.

**name-services** (default: '() )

List of packages denoting *name services* that must be visible to the nsd—e.g., (list *nss-mdns*).

**glibc** (default: *glibc*)

Package object denoting the GNU C Library providing the **nsd** command.

**log-file** (default: **#f**)

Name of the nsd log file. Debugging output goes to that file when **debug-level** is strictly positive, or to standard error if it is **#f**. Regular messages are written to syslog when **debug-level** is zero, regardless of the value of **log-file**.

**debug-level** (default: 0)

Integer denoting the debugging levels. Higher numbers mean that more debugging output is logged.

**caches** (default: **%nsd-default-caches**)

List of **<nsd-cache>** objects denoting things to be cached; see below.

**nsd-cache**

[Data Type]

Data type representing a cache database of nsd and its parameters.

**database** This is a symbol representing the name of the database to be cached. Valid values are **passwd**, **group**, **hosts**, and **services**, which designate the corresponding NSS database (see Section “NSS Basics” in *The GNU C Library Reference Manual*).

**positive-time-to-live****negative-time-to-live** (default: 20)

A number representing the number of seconds during which a positive or negative lookup result remains in cache.

**check-files?** (default: **#t**)

Whether to check for updates of the files corresponding to *database*.

For instance, when *database* is **hosts**, setting this flag instructs nsd to check for updates in **/etc/hosts** and to take them into account.

**persistent?** (default: **#t**)

Whether the cache should be stored persistently on disk.

**shared?** (default: **#t**)

Whether the cache should be shared among users.

**max-database-size** (default: 32 MiB)

Maximum size in bytes of the database cache.



**%nscd-default-caches** [Variable]

List of `<nscd-cache>` objects used by default by `nscd-configuration` (see above). It enables persistent and aggressive caching of service and host name lookups. The latter provides better host name lookup performance, resilience in the face of unreliable name servers, and also better privacy—often the result of host name lookups is in local cache, so external name servers do not even need to be queried.

**syslog-service-type** [Variable]

Type of the service that runs the syslog daemon (`syslogd`), whose value is a `<syslog-configuration>` object (see below).

**Note:** This service is redundant with and for the most part superseded by `shepherd-system-log-service-type` (see `[shepherd-system-log]`, page 723).

To have a modified `syslog-configuration` come into effect after reconfiguring your system, the ‘`reload`’ action should be preferred to restarting the service, as many services such as the login manager depend on it and would be restarted as well:

```
# herd reload syslog
```

which will cause the running `syslogd` process to reload its configuration.

**syslog-configuration** [Data Type]

Data type representing the configuration of the syslog daemon.

`syslogd` (default: `#~(string-append #$(inetutils) "/libexec/syslogd")`)  
The syslog daemon to use.

`config-file` (default: `%default-syslog.conf`)  
The syslog configuration file to use. See Section “`syslogd` invocation” in *GNU Inetutils*, for more information on the configuration file syntax.

`extra-options` (default: `'()`)  
List of extra command-line options for `syslog`.

**guix-service-type** [Variable]

This is the type of the service that runs the build daemon, `guix-daemon` (see Section 2.3 [Invoking `guix-daemon`], page 14). Its value must be a `guix-configuration` record as described below.

**guix-configuration** [Data Type]

This data type represents the configuration of the Guix build daemon. See Section 2.3 [Invoking `guix-daemon`], page 14, for more information.

`guix` (default: `guix`)  
The Guix package to use. See Section 6.4 [Customizing the System-Wide Guix], page 70, to learn how to provide a package with a pre-configured set of channels.

`privileged?` (default: `#t`)  
Whether to run `guix-daemon` as root.  
When true, `guix-daemon` runs with root privileges and build processes run under unprivileged user accounts as specified by `build-group` and

`build-accounts` (see below); when false, `guix-daemon` run as the `guix-daemon` user, which is unprivileged, and so do build processes. The unprivileged or “rootless” mode can reduce the impact of some classes of vulnerabilities that could affect the daemon.

The default is currently `#t` (`guix-daemon` runs with root privileges) but may eventually be changed to `#f`.

**Warning:** When changing this option, `/gnu/store`, `/var/guix`, and `/etc/guix` have their ownership automatically changed by the `guix-ownership` service to either the `guix-daemon` user or the `root` user (see [unprivileged-daemon-migration], page 7).

This can take a while, especially if `/gnu/store` is big; it cannot be interrupted and `guix-daemon` cannot be used until it has completed.

See Section 2.2.1 [Build Environment Setup], page 6, for more information on the two ways to run `guix-daemon`.

`build-group` (default: "guixbuild")

Name of the group for build user accounts.

`build-accounts` (default: 10)

Number of build user accounts to create.

`chroot?` (default: 'default')

The value should be one of `#t` or `#f`, in which case `chroot` is enabled or disabled, respectively; or it should be 'default', which amounts to `#f` in Docker containers (so that they can be run in non-privileged mode) or `#t` otherwise.

`authorize-key?` (default: `#t`)

Whether to authorize the substitute keys listed in `authorized-keys`—by default that of `bordeaux.guix.gnu.org` and `ci.guix.gnu.org` (see Section 5.3 [Substitutes], page 46).

When `authorize-key?` is true, `/etc/guix/acl` cannot be changed by invoking `guix archive --authorize`. You must instead adjust `guix-configuration` as you wish and reconfigure the system. This ensures that your operating system configuration file is self-contained.

**Note:** When booting or reconfiguring to a system where `authorize-key?` is true, the existing `/etc/guix/acl` file is backed up as `/etc/guix/acl.bak` if it was determined to be a manually modified file. This is to facilitate migration from earlier versions, which allowed for in-place modifications to `/etc/guix/acl`.

`authorized-keys` (default: %default-authorized-guix-keys)

The list of authorized key files for archive imports, as a list of string-valued gexps (see Section 5.11 [Invoking guix archive], page 65). By default, it contains that of `bordeaux.guix.gnu.org` and `ci.guix.gnu.org` (see

Section 5.3 [Substitutes], page 46). See `substitute-urls` below for an example on how to change it.

`use-substitutes?` (default: `#t`)

Whether to use substitutes.

`substitute-urls` (default: `%default-substitute-urls`)

The list of URLs where to look for substitutes by default.

Suppose you would like to fetch substitutes from `guix.example.org` in addition to `bordeaux.guix.gnu.org`. You will need to do two things: (1) add `guix.example.org` to `substitute-urls`, and (2) authorize its signing key, having done appropriate checks (see Section 5.3.2 [Substitute Server Authorization], page 47). The configuration below does exactly that:

```
(guix-configuration
 (substitute-urls
  (append (list "https://guix.example.org")
           %default-substitute-urls))
 (authorized-keys
  (append (list (local-file "./guix.example.org-key.pub"))
           %default-authorized-guix-keys)))
```

This example assumes that the file `./guix.example.org-key.pub` contains the public key that `guix.example.org` uses to sign substitutes.

`generate-substitute-key?` (default: `#t`)

Whether to generate a *substitute key pair* under `/etc/guix/signing-key.pub` and `/etc/guix/signing-key.sec` if there is not already one.

This key pair is used when exporting store items, for instance with `guix publish` (see Section 9.11 [Invoking `guix publish`], page 227) or `guix archive` (see Section 5.11 [Invoking `guix archive`], page 65). Generating a key pair takes a few seconds when enough entropy is available and is only done once; you might want to turn it off for instance in a virtual machine that does not need it and where the extra boot time is a problem.

`channels` (default: `#f`)

List of channels to be specified in `/etc/guix/channels.scm`, which is what `guix pull` uses by default (see Section 5.7 [Invoking `guix pull`], page 56).

**Note:** When reconfiguring a system, the existing `/etc/guix/channels.scm` file is backed up as `/etc/guix/channels.scm.bak` if it was determined to be a manually modified file. This is to facilitate migration from earlier versions, which allowed for in-place modifications to `/etc/guix/channels.scm`.

**max-silent-time** (default: 3600)

**timeout** (default: (\* 3600 24))

The number of seconds of silence and the number of seconds of activity, respectively, after which a build process times out. A value of zero disables the timeout.

**log-compression** (default: 'gzip')

The type of compression used for build logs—one of **gzip**, **bzip2**, or **none**.

**discover?** (default: #f)

Whether to discover substitute servers on the local network using mDNS and DNS-SD.

**build-machines** (default: #f)

This field must be either **#f** or a list of gexps evaluating to a **build-machine** record or to a list of **build-machine** records (see Section 2.2.2 [Daemon Offload Setup], page 9).

When it is **#f**, the `/etc/guix/machines.scm` file is left untouched. Otherwise, the list of gexps is written to `/etc/guix/machines.scm`; if a previously-existing file is found, it is backed up as `/etc/guix/machines.scm.bak`. This allows you to declare build machines for offloading directly in the operating system declaration, like so:

```
(guix-configuration
 (build-machines
  (list #~(build-machine (name "foo.example.org") ...)
        #~(build-machine (name "bar.example.org") ...))))
```

Additional build machines may be added *via* the **guix-extension** mechanism (see below).

**extra-options** (default: '() )

List of extra command-line options for **guix-daemon**.

**log-file** (default: "/var/log/guix-daemon.log")

File where **guix-daemon**'s standard output and standard error are written.

**http-proxy** (default: #f)

The URL of the HTTP and HTTPS proxy used for downloading fixed-output derivations and substitutes.

It is also possible to change the daemon's proxy at run time through the **set-http-proxy** action, which restarts it:

```
herd set-http-proxy guix-daemon http://localhost:8118
```

To clear the proxy settings, run:

```
herd set-http-proxy guix-daemon
```

**tmpdir** (default: #f)

A directory path where the **guix-daemon** will perform builds.

**environment** (default: '()')

Environment variables to be set before starting the daemon, as a list of **key=value** strings.

**socket-directory-permissions** (default: #o755)

Permissions to set for the directory `/var/guix/daemon-socket`. This, together with **socket-directory-group** and **socket-directory-user**, determines who can connect to the build daemon via its Unix socket. TCP socket operation is unaffected by these.

**socket-directory-user** (default: #f)

**socket-directory-group** (default: #f)

User and group owning the `/var/guix/daemon-socket` directory or #f to keep the user or group as root.

**guix-extension**

[Data Type]

This data type represents the parameters of the Guix build daemon that are extendable. This is the type of the object that must be used within a guix service extension. See Section 11.19.1 [Service Composition], page 709, for more information.

**authorized-keys** (default: '()')

A list of file-like objects where each element contains a public key.

**substitute-urls** (default: '()')

A list of strings where each element is a substitute URL.

**build-machines** (default: '()')

A list of gexps that evaluate to **build-machine** records or to a list of **build-machine** records. (see Section 2.2.2 [Daemon Offload Setup], page 9).

Using this field, a service may add new build machines to receive builds offloaded by the daemon. This is useful for a service such as **hurd-vm-service-type**, which can make a GNU/Hurd virtual machine directly usable for offloading (see [hurd-vm], page 593).

**chroot-directories** (default: '()')

A list of file-like objects or strings pointing to additional directories the build daemon can use.

**udev-service-type**

[Variable]

Type of the service that runs udev, a service which populates the `/dev` directory dynamically, whose value is a **<udev-configuration>** object.

Since the file names for udev rules and hardware description files matter, the configuration items for rules and hardware cannot simply be plain file-like objects with the rules content, because the name would be ignored. Instead, they are directory file-like objects that contain optional rules in `lib/udev/rules.d` and optional hardware files in `lib/udev/hwdb.d`. This way, the service can be configured with whole packages from which to take rules and hwdb files.

The **udev-service-type** can be *extended* with file-like directories that respect this hierarchy. For convenience, the **udev-rule** and **file->udev-rule** can be used to

construct udev rules, while `udev-hardware` and `file->udev-hardware` can be used to construct hardware description files.

In an `operating-system` declaration, this service type can be *extended* using procedures `udev-rules-service` and `udev-hardware-service`.

### `udev-configuration`

[Data Type]

Data type representing the configuration of udev.

`udev` (default: `eudev`) (type: `file-like`)

Package object of the udev service. This package is used at run-time, when compiled for the target system. In order to generate the `hwdb.bin` hardware index, it is also used when generating the system definition, compiled for the current system.

`debug?` (default: `#f`) (type: `boolean`)

Whether to enable debug output. The debug output is written to the system log, `/var/log/messages`.

`rules` (default: `'()`) (type: `list-of-file-like`)

List of file-like objects denoting udev rule files under a sub-directory.

`hardware` (default: `'()`) (type: `list-of-file-like`)

List of file-like objects denoting udev hardware description files under a sub-directory.

### `udev-rule file-name contents`

[Procedure]

Return a udev-rule file named *file-name* containing the rules defined by the *contents* literal.

In the following example, a rule for a USB device is defined to be stored in the file `90-usb-thing.rules`. The rule runs a script upon detecting a USB device with a given product identifier.

```
(define %example-udev-rule
  (udev-rule
    "90-usb-thing.rules"
    (string-append "ACTION==" "add", "SUBSYSTEM==" "usb", "
                  "ATTR{product}==" "Example", "
                  "RUN+=" "/path/to/script\"")))

```

### `udev-hardware file-name contents`

[Procedure]

Return a udev hardware description file named *file-name* containing the hardware information *contents*.

### `udev-rules-service name rules` [`#:groups` `'()`]

[Procedure]

Return a service that extends `udev-service-type` with *rules* and `account-service-type` with *groups* as system groups. This works by creating a singleton service type *name-udev-rules*, of which the returned service is an instance.

Here we show how it can be used to extend `udev-service-type` with the previously defined rule `%example-udev-rule`.

```
(operating-system

```

```
;; ...
(services
  (cons (udev-rules-service 'usb-thing %example-udev-rule)
        %desktop-services)))
```

**udev-hardware-service** *name hardware* [Procedure]  
 Return a service that extends `udev-service-type` with *hardware*. The service name is *name-udev-hardware*.

**file->udev-rule** *file-name file* [Procedure]  
 Return a udev-rule file named *file-name* containing the rules defined within *file*, a file-like object.

The following example showcases how we can use an existing rule file.

```
(use-modules (guix download)      ;for url-fetch
             (guix packages)      ;for origin
             ...)

(define %android-udev-rules
  (file->udev-rule
   "51-android-udev.rules"
   (let ((version "20170910"))
     (origin
      (method url-fetch)
      (uri (string-append "https://raw.githubusercontent.com/MORf30/"
                          "android-udev-rules/" version "/51-android.rules"))
      (sha256
       (base32 "0lmmagpyb6xsq6zcr2w1cyx9qmjqmajkvrdbhjx32gqf1d9is003"))))))
```

Since guix package definitions can be included in *rules* in order to use all their rules under the `lib/udev/rules.d` sub-directory, then in lieu of the previous *file->udev-rule* example, we could have used the *android-udev-rules* package which exists in Guix in the `(gnu packages android)` module.

**file->udev-hardware** *file-name file* [Procedure]  
 Return a udev hardware description file named *file-name* containing the rules defined within *file*, a file-like object.

The following example shows how to use the *android-udev-rules* package so that the Android tool `adb` can detect devices without root privileges. It also details how to create the `adbusers` group, which is required for the proper functioning of the rules defined within the *android-udev-rules* package. To create such a group, we must define it both as part of the `supplementary-groups` of our `user-account` declaration, as well as in the *groups* of the `udev-rules-service` procedure.

```
(use-modules (gnu packages android) ;for android-udev-rules
             (gnu system shadow)    ;for user-group
             ...)

(operating-system
```

```
;; ...
(users (cons (user-account
              ;; ...
              (supplementary-groups
                ('("adbusers"    ;for adb
                  "wheel" "netdev" "audio" "video")))))
;; ...
(services
 (cons (udev-rules-service 'android android-udev-rules
                           #:groups '("adbusers"))
       %desktop-services)))
```

**urandom-seed-service-type** [Variable]

Save some entropy in `%random-seed-file` to seed `/dev/urandom` when rebooting. It also tries to seed `/dev/urandom` from `/dev/hwrng` while booting, if `/dev/hwrng` exists and is readable.

**%random-seed-file** [Variable]

This is the name of the file where some random bytes are saved by *urandom-seed-service* to seed `/dev/urandom` when rebooting. It defaults to `/var/lib/random-seed`.

**gpm-service-type** [Variable]

This is the type of the service that runs GPM, the *general-purpose mouse daemon*, which provides mouse support to the Linux console. GPM allows users to use the mouse in the console, notably to select, copy, and paste text.

The value for services of this type must be a **gpm-configuration** (see below). This service is not part of `%base-services`.

**gpm-configuration** [Data Type]

Data type representing the configuration of GPM.

**options** (default: `%default-gpm-options`)

Command-line options passed to **gpm**. The default set of options instruct **gpm** to listen to mouse events on `/dev/input/mice`. See Section “Command Line” in *gpm manual*, for more information.

**gpm** (default: `gpm`)

The GPM package to use.

**guix-publish-service-type** [Variable]

This is the service type for **guix publish** (see Section 9.11 [Invoking **guix publish**], page 227). Its value must be a **guix-publish-configuration** object, as described below.

This assumes that `/etc/guix` already contains a signing key pair as created by **guix archive --generate-key** (see Section 5.11 [Invoking **guix archive**], page 65). If that is not the case, the service will fail to start.

**guix-publish-configuration** [Data Type]

Data type representing the configuration of the **guix publish** service.

**guix** (default: `guix`)

The Guix package to use.



**port** (default: 80)

The TCP port to listen for connections.

**host** (default: "localhost")

The host (and thus, network interface) to listen to. Use "0.0.0.0" to listen on all the network interfaces.

**advertise?** (default: #f)

When true, advertise the service on the local network *via* the DNS-SD protocol, using Avahi.

This allows neighboring Guix devices with discovery on (see **guix-configuration** above) to discover this **guix publish** instance and to automatically download substitutes from it.

**compression** (default: '(("gzip" 3) ("zstd" 3)))

This is a list of compression method/level tuple used when compressing substitutes. For example, to compress all substitutes with *both* lzip at level 7 and gzip at level 9, write:

```
'(("lzip" 7) ("gzip" 9))
```

Level 9 achieves the best compression ratio at the expense of increased CPU usage, whereas level 1 achieves fast compression. See Section 9.11 [Invoking guix publish], page 227, for more information on the available compression methods and the tradeoffs involved.

An empty list disables compression altogether.

**nar-path** (default: "nar")

The URL path at which “nars” can be fetched. See Section 9.11 [Invoking guix publish], page 227, for details.

**cache** (default: #f)

When it is #f, disable caching and instead generate archives on demand. Otherwise, this should be the name of a directory—e.g., `"/var/cache/guix/publish"`—where **guix publish** caches archives and meta-data ready to be sent. See Section 9.11 [Invoking guix publish], page 227, for more information on the tradeoffs involved.

**workers** (default: #f)

When it is an integer, this is the number of worker threads used for caching; when #f, the number of processors is used. See Section 9.11 [Invoking guix publish], page 227, for more information.

**cache-bypass-threshold** (default: 10 MiB)

When **cache** is true, this is the maximum size in bytes of a store item for which **guix publish** may bypass its cache in case of a cache miss. See Section 9.11 [Invoking guix publish], page 227, for more information.

**ttl** (default: #f)

When it is an integer, this denotes the *time-to-live* in seconds of the published archives. See Section 9.11 [Invoking guix publish], page 227, for more information.

**negative-ttl** (default: #f)

When it is an integer, this denotes the *time-to-live* in seconds for the negative lookups. See Section 9.11 [Invoking guix publish], page 227, for more information.

**rngd-service-type** [Variable]

Type of the service that runs `rng-tools rngd`, whose value is an `<rngd-configuration>` object.

**rngd-configuration** [Data Type]

Data type representing the configuration of `rngd`.

**rng-tools** (default: `rng-tools`) (type: file-like)

Package object of the `rng-tools` `rngd`.

**device** (default: `"/dev/hwrng"`) (type: string)

Path of the device to add to the kernel's entropy pool. The service will fail if *device* does not exist.

**pam-limits-service-type** [Variable]

Type of the service that installs a configuration file for the `pam_limits` module ([http://linux-pam.org/Linux-PAM-html/sag-pam\\_limits.html](http://linux-pam.org/Linux-PAM-html/sag-pam_limits.html)). The value for this service type is a list of `pam-limits-entry` values, which can be used to specify `ulimit` limits and `nice` priority limits to user sessions. By default, the value is the empty list.

The following limits definition sets two hard and soft limits for all login sessions of users in the `realtime` group:

```
(service pam-limits-service-type
  (list
    (pam-limits-entry "@realtime" 'both 'rtprio 99)
    (pam-limits-entry "@realtime" 'both 'memlock 'unlimited)))
```

The first entry increases the maximum realtime priority for non-privileged processes; the second entry lifts any restriction of the maximum address space that can be locked in memory. These settings are commonly used for real-time audio systems.

Another useful example is raising the maximum number of open file descriptors that can be used:

```
(service pam-limits-service-type
  (list
    (pam-limits-entry "*" 'both 'nofile 100000)))
```

In the above example, the asterisk means the limit should apply to any user. It is important to ensure the chosen value doesn't exceed the maximum system value visible in the `/proc/sys/fs/file-max` file, else the users would be prevented from login in. For more information about the Pluggable Authentication Module (PAM) limits, refer to the '`pam_limits`' man page from the `linux-pam` package.

**greetd-service-type** [Variable]

`greetd` (<https://git.sr.ht/~kennylevinsen/greetd>) is a minimal and flexible login manager daemon, that makes no assumptions about what you want to launch.

If you can run it from your shell in a TTY, greetd can start it. If it can be taught to speak a simple JSON-based IPC protocol, then it can be a geeter.

`greetd-service-type` provides necessary infrastructure for logging in users, including:

- `greetd` PAM service
- Special variation of `pam-mount` to mount `XDG_RUNTIME_DIR`

Here is an example of switching from `mingetty-service-type` to `greetd-service-type`, and how different terminals could be:

```
(append
(modify-services %base-services
;; greetd-service-type provides "greetd" PAM service
(delete login-service-type)
;; and can be used in place of mingetty-service-type
(delete mingetty-service-type))
(list
(service greetd-service-type
(greetd-configuration
(terminals
(list
;; we can make any terminal active by default
(greetd-terminal-configuration (terminal-vt "1") (terminal-switch
;; we can make environment without XDG_RUNTIME_DIR set
;; even provide our own environment variables
(greetd-terminal-configuration
(terminal-vt "2")
(default-session-command
(greetd-agreety-session
(command
(greetd-user-session
(extra-env '("MY_VAR" . "1")))
(xdg-env? #f))))))
;; we can use different shell instead of default bash
(greetd-terminal-configuration
(terminal-vt "3")
(default-session-command
(greetd-agreety-session
(command
(greetd-user-session
(command (file-append zsh "/bin/zsh"))
(command-args '("-l"))
(extra-env '("MY_VAR" . "1")))
(xdg-env? #f))))))
;; we can use any other executable command as greeter
(greetd-terminal-configuration
(terminal-vt "4")
(default-session-command (program-file "my-noop-greeter" #~(exit
```

```

        (greetd-terminal-configuration (terminal-vt "5"))
        (greetd-terminal-configuration (terminal-vt "6"))))))))■
;; mingetty-service-type can be used in parallel
;; if needed to do so, do not (delete login-service-type)
;; as illustrated above
#| (service mingetty-service-type (mingetty-configuration (tty "tty8"))) |#))

```

**greetd-configuration** [Data Type]

Configuration record for the **greetd-service-type**.

**motd** A file-like object containing the “message of the day”.

**allow-empty-passwords?** (default: **#t**)  
 Allow empty passwords by default so that first-time users can log in when the ‘root’ account has just been created.

**terminals** (default: '())  
 List of **greetd-terminal-configuration** per terminal for which **greetd** should be started.

**greeter-supplementary-groups** (default: '())  
 List of groups which should be added to **greeter** user. For instance:  
 (**greeter-supplementary-groups** '("seat" "video"))  
 Note that this example will fail if **seat** group does not exist.

**greetd-terminal-configuration** [Data Type]

Configuration record for per terminal **greetd** daemon service.

**greetd** (default: **greetd**)  
 The **greetd** package to use.

**extra-shepherd-requirement** (default: '())  
 This option can be used to provide a list of *additional* symbols naming Shepherd services that this service will depend on, such as '**seatd** or '**elogind** when the terminal session is a graphical greeter.

**config-file-name**  
 Configuration file name to use for **greetd** daemon. Generally, autogenerated derivation based on **terminal-vt** value.

**log-file-name**  
 Log file name to use for **greetd** daemon. Generally, autogenerated name based on **terminal-vt** value.

**terminal-vt** (default: “7”)  
 The VT to run on. Use of a specific VT with appropriate conflict avoidance is recommended.

**terminal-switch** (default: **#f**)  
 Make this terminal active on start of **greetd**.

**source-profile?** (default: **#t**)  
 Whether to source **/etc/profile** and **~/.profile**, when they exist.

**initial-session-user** (default: `#f`)

The user to use for running the initial session, if any.

**initial-session-command** (default: `#f`)

The command to run the first time the greeter starts up, to auto-login. For that to be useful, you should also set **initial-session-user**. Then, **greetd** will automatically log in with that user and execute that command. When you eventually log out again, the regular greeter will appear.

**default-session-user** (default: `"greeter"`)

The user to use for running the greeter.

**default-session-command** (default: `(greetd-agreety-session)`)

Can be either **greetd-agreety-session**, **greetd-wlgreet-sway-session**, **greetd-gtkgreet-sway-session** or a file-like object to use as greeter.

## **greetd-user-session**

[Data Type]

Configuration record for the user session command. Greeters require the user command to be specified in some or another way. **greetd-user-session** provides a common command for that. Users should prefer POSIX shell commands like **bash**, which can start an actual user terminal shell, window manager or desktop environment with their own mechanism, for example via `~/.bashrc` in the case of Bash.

**command** (default: `(file-append bash "/bin/bash")`)

Command to be started by **agreety** on successful login.

**command-args** (default: `'("-l")`)

Command arguments to pass to command.

**extra-env** (default: `'()`)

Extra environment variables to set on login.

**xdg-session-type** (default: `"tty"`)

Specify the value of `XDG_SESSION_TYPE`. The user environment may adapt depending on its value (normally by using `.bashrc` or similar).

**xdg-env?** (default: `#t`)

If true `XDG_RUNTIME_DIR` and `XDG_SESSION_TYPE` will be set before starting command. One should note that, **extra-env** variables are set right after mentioned variables, so that they can be overridden.

## **greetd-agreety-session**

[Data Type]

Configuration record for the **agreety** **greetd** greeter.

**agreety** (default: `greetd`)

The package providing the **agreety** command.

**command** (default: `(greetd-user-session)`)

Command to be started by **agreety** on successful login, an instance of **greetd-user-session**.

**greetd-wlgreet-configuration** [Data Type]

Generic configuration record for the wlgreet greetd greeter.

**output-mode** (default: "all")

Option to use for **outputMode** in the TOML configuration file.

**scale** (default: 1)

Option to use for **scale** in the TOML configuration file.

**background** (default: '(0 0 0 0.9))

RGBA list to use as the background colour of the login prompt.

**headline** (default: '(1 1 1 1))

RGBA list to use as the headline colour of the UI popup.

**prompt** (default: '(1 1 1 1))

RGBA list to use as the prompt colour of the UI popup.

**prompt-error** (default: '(1 1 1 1))

RGBA list to use as the error colour of the UI popup.

**border** (default: '(1 1 1 1))

RGBA list to use as the border colour of the UI popup.

**greetd-wlgreet-sway-session** [Data Type]

Sway-specific configuration record for the wlgreet greetd greeter.

**wlgreet-session** (default: (greetd-wlgreet-session))

A **greetd-wlgreet-session** record for generic wlgreet configuration, on top of the Sway-specific **greetd-wlgreet-sway-session**.

**sway** (default: sway)

The package providing the **sway** command.

**sway-configuration** (default: #f)

File-like object providing an additional Sway configuration file to be prepended to the mandatory part of the configuration.

**wlgreet** (default: wlgreet)

The package providing the **wlgreet** command.

**wlgreet-configuration** (default: (greetd-wlgreet-configuration))

Configuration of **wlgreet** represented by **greetd-wlgreet-configuration**.

**command** (default: (greetd-user-session))

Command to be started by **wlgreet** on successful login, an instance of **greetd-user-session**.

Here is an example of a greetd configuration that uses wlgreet and Sway:

```
(greetd-configuration
;; The graphical greeter requires additional group membership.
(greeter-supplementary-groups (list "video" "input" "seat"))
(terminals
(list (greetd-terminal-configuration
```

```
;; Sway requires seatd service.
(extra-shepherd-requirement '(seatd))
(terminal-vt "1")
(terminal-switch #t)
(default-session-command
 (greetd-wlgreet-sway-session
  (sway-configuration
   ;; Optional extra sway configuration.
   (local-file "sway-greetd.conf"))))))))
```

**greetd-gtkgreet-sway-session** [Data Type]

Configuration record for the gtkgreet greetd greeter. It can be used as follows:

```
(greetd-configuration
 ;; The graphical greeter requires additional groups membership.
 (greeter-supplementary-groups (list "video" "input" "seat"))
 (terminals
  (list (greetd-terminal-configuration
        ;; Sway requires the seatd service.
        (extra-shepherd-requirement '(seatd))
        (terminal-vt "1")
        (terminal-switch #t)
        (default-session-command
         (greetd-gtkgreet-sway-session
          (command
           (greetd-user-session
            ;; Optionally signal to .bashrc that we want the Wayland
            ;; compositor.
            (xdg-session-type "wayland"))))))))
```

**sway** (default: sway)

The package providing the sway and swaymsg commands.

**sway-configuration** (default: #f)

Extra file-like configuration for sway to be included before executing the greeter.

**gtkgreet** (default: gtkgreet)

The package providing the gtkgreet command.

**gtkgreet-style** (default: #f)

Extra file-like CSS stylesheet to customize the GTK look.

**command** (default: (greetd-user-session))

The command to be started by gtkgreet on successful login, an instance of greetd-user-session.

### 11.10.2 Scheduled Job Execution

It is often useful to have system tasks run periodically. This can be achieved by defining Shepherd *timers* or by using the historical and somewhat less flexible mcron service.

### 11.10.2.1 Shepherd Timers

The Shepherd supports running jobs periodically by defining *timers*, a special kind of service. A Shepherd timer can be defined like another Shepherd service, but with specific **start** and **stop** methods (see Section 11.19.4 [Shepherd Services], page 717):

```
;; Defining a timer, verbosely.
(shepherd-service
 (provision '(updatedb))
 (requirement '(user-processes))
 (modules '((shepherd service timer)))
 (start #~(make-timer-constructor
           ;; Everyday at 3AM (using Vixie cron syntax).
           (cron-string->calendar-event "0 3 * * *")
           (command '($$(file-append findutils "/bin/updatedb")))
           #:wait-for-termination? #t))
 (stop #~(make-timer-destructor))
 (documentation "Periodically run 'updatedb'.")
 (actions (list shepherd-trigger-action)))
```

This is quite some boilerplate so you can instead use this equivalent shorthand notation:

```
;; Equivalent timer definition, but shorter.
(shepherd-timer '(updatedb) "0 3 * * *"
                #~($$(file-append findutils "/bin/updatedb"))
                #:requirement '(user-processes))
```

This procedure is defined as follows.

```
shepherd-timer provision schedule command [#:requirement '()] [Procedure]
               [#:documentation %default-timer-documentation]
```

Return a Shepherd service with the given *provision* periodically running *command*, a list-valued gexp, according to *schedule*, a string in Vixie cron syntax or a gexp providing a Shepherd calendar event. *documentation* is the string that appears when running `herd doc service`.

The example below shows how to define a timer and to add it to your operating system configuration.

```
(use-modules (gnu))
(use-service-modules shepherd)
(use-package-modules base)

(define updatedb-timer
  ;; Run 'updatedb' at 3AM everyday.
  (shepherd-timer '(updatedb)
                  "0 3 * * *" ;Vixie cron syntax
                  #~($$(file-append findutils "/bin/updatedb")
                    "--prunepaths=/tmp /var/tmp /gnu/store")
                  #:requirement '(user-processes)))

(define garbage-collection-timer
```



```

;; Run 'guix gc' everyday at 5AM.
(shepherd-timer '(garbage-collection)
  #~(calendar-event #:hours '(5) #:minutes '(0))
  #~("/run/current-system/profile/bin/guix"
    "gc" "-F" "10G")
  #:requirement '(guix-daemon)))

(operating-system
  ;; ...

  ;; Extend the Shepherd service with additional timers
  ;; using 'simple-service'.
  (services (cons (simple-service 'my-timers
                                shepherd-root-service-type
                                (list garbage-collection-timer
                                      updatedb-timer))
                  %base-services)))

```

The resulting system contains these two services, which can be inspected with `herd status`. They can also be triggered manually:

```
herd trigger garbage-collection
```

See Section “Timers” in *The GNU Shepherd Manual* for more information on Shepherd timers.

### 11.10.2.2 Mcron

Alternatively, the `(gnu services mcron)` module provides an interface to GNU mcron, a daemon to run jobs at scheduled times (see *GNU mcron*). GNU mcron is similar to the traditional Unix `cron` daemon; the main difference is that it is implemented in Guile Scheme. It is overall less convenient than Shepherd timers: individual jobs cannot depend on specific Shepherd services, logging is coarse-grain (one file for all the jobs), jobs may not be inspected, updated, or triggered manually.

**Note:** For the reasons given above, we recommend using Shepherd timers rather than mcron for your periodic tasks.

The example below defines an operating system that runs the `updatedb` (see Section “Invoking `updatedb`” in *Finding Files*) and the `guix gc` commands (see Section 5.6 [Invoking `guix gc`], page 53) daily, as well as the `mkid` command on behalf of an unprivileged user (see Section “`mkid` invocation” in *ID Database Utilities*). It uses `gexps` to introduce job definitions that are passed to mcron (see Section 8.12 [G-Expressions], page 165).

```

(use-modules (guix) (gnu) (gnu services mcron))
(use-package-modules base idutils)

(define updatedb-job
  ;; Run 'updatedb' at 3AM every day. Here we write the
  ;; job's action as a Scheme procedure.
  #~(job '(next-hour '(3))
    (lambda ()

```

```

        (system* (string-append #$findutils "/bin/updatedb")
                  "--prunepaths=/tmp /var/tmp /gnu/store"))
    "updatedb"))

(define garbage-collector-job
  ;; Collect garbage 5 minutes after midnight every day.
  ;; The job's action is a shell command.
  #~(job "5 0 * * *" ;Vixie cron syntax
        "guix gc -F 1G"))

(define idutils-job
  ;; Update the index database as user "charlie" at 12:15PM
  ;; and 19:15PM. This runs from the user's home directory.
  #~(job '(next-minute-from (next-hour '(12 19)) '(15))
        (string-append #$idutils "/bin/mkid src")
        #:user "charlie"))

(operating-system
  ;; ...

  ;; %BASE-SERVICES already includes an instance of
  ;; 'mcron-service-type', which we extend with additional
  ;; jobs using 'simple-service'.
  (services (cons (simple-service 'my-cron-jobs
                                mcron-service-type
                                (list garbage-collector-job
                                      updatedb-job
                                      idutils-job))
                  %base-services)))

```

**Tip:** When providing the action of a job specification as a procedure, you should provide an explicit name for the job via the optional 3rd argument as done in the `updatedb-job` example above. Otherwise, the job would appear as “Lambda function” in the output of `herd schedule mcron`, which is not nearly descriptive enough!

**Tip:** Avoid calling the Guile procedures `execl`, `execle` or `execlp` inside a job specification, else `mcron` won’t be able to output the completion status of the job.

For more complex jobs defined in Scheme where you need control over the top level, for instance to introduce a `use-modules` form, you can move your code to a separate program using the `program-file` procedure of the `(guix gexp)` module (see Section 8.12 [G-Expressions], page 165). The example below illustrates that.

```

(define %battery-alert-job
  ;; Beep when the battery percentage falls below %MIN-LEVEL.
  #~(job
    '(next-minute (range 0 60 1))
    #$(program-file

```

```

"battery-alert.scm"
(with-imported-modules (source-module-closure
                        '((guix build utils)))

  #~(begin
    (use-modules (guix build utils)
                 (ice-9 popen)
                 (ice-9 regex)
                 (ice-9 textual-ports)
                 (srfi srfi-2))

    (define %min-level 20)

    (setenv "LC_ALL" "C") ;ensure English output
    (and-let* ((input-pipe (open-pipe*
                              OPEN_READ
                              #$(file-append acpi "/bin/acpi"))))
      (output (get-string-all input-pipe))
      (m (string-match "Discharging, ([0-9]+)%" output))
      (level (string->number (match:substring m 1)))
      ((< level %min-level)))
      (format #t "warning: Battery level is low (~a%)~%" level)
      (invoke #$(file-append beep "/bin/beep") "-r5")))))

```

See Section “Guile Syntax” in *GNU mcron*, for more information on mcron job specifications. Below is the reference of the mcron service.

On a running system, you can use the `schedule` action of the service to visualize the mcron jobs that will be executed next:

```
# herd schedule mcron
```

The example above lists the next five tasks that will be executed, but you can also specify the number of tasks to display:

```
# herd schedule mcron 10
```

**mcron-service-type** [Variable]

This is the type of the mcron service, whose value is an `mcron-configuration` object.

This service type can be the target of a service extension that provides additional job specifications (see Section 11.19.1 [Service Composition], page 709). In other words, it is possible to define services that provide additional mcron jobs to run.

**mcron-configuration** [Data Type]

Available `mcron-configuration` fields are:

**mcron** (default: `mcron`) (type: file-like)

The mcron package to use.

**jobs** (default: `'()`) (type: list-of-gexps)

This is a list of gexps (see Section 8.12 [G-Expressions], page 165), where each gexp corresponds to an mcron job specification (see Section “Syntax” in *GNU mcron*).

**log?** (default: `#t`) (type: boolean)  
 Log messages to standard output.

**log-file** (default: `"/var/log/mcron.log"`) (type: string)  
 Log file location.

**log-format** (default: `"~1@*~a ~a: ~a~%"`) (type: string)  
 (ice-9 format) format string for log messages. The default value produces messages like `'pid name: message'` (see Section “Invoking mcron” in *GNU mcron*). Each message is also prefixed by a timestamp by GNU Shepherd.

**date-format** (type: maybe-string)  
 (srfi srfi-19) format string for date.

### 11.10.3 Log Rotation

Log files such as those found in `/var/log` tend to grow endlessly, so it’s a good idea to *rotate* them once in a while—i.e., archive their contents in separate files, possibly compressed. The (`gnu services admin`) module provides an interface to the log rotation service provided by the Shepherd (see Section “Log Rotation Service” in *The GNU Shepherd Manual*).

This log rotation service is made available through `log-rotation-service-type`, which takes a `log-rotation-configuration` record as its value. By default, this provides `log-rotation`, a Shepherd “timed service” that runs periodically—once a week by default. It automatically knows about the log files produced by Shepherd services and can be taught about external log files. You can inspect the service and see when it’s going to run the usual way:

```
$ sudo herd status log-rotation
Status of log-rotation:
  It is running since Mon 09 Dec 2024 03:27:47 PM CET (2 days ago).
  ...

Upcoming timer alarms:
  Sun 15 Dec 2024 10:00:00 PM CET (in 4 days)
  Sun 22 Dec 2024 10:00:00 PM CET (in 11 days)
  Sun 29 Dec 2024 10:00:00 PM CET (in 18 days)
```

You can also list files subject to rotation with `herd files log-rotation` and trigger rotation manually with `herd trigger log-rotation`.

This service is part of `%base-services`, and thus enabled by default, with the default settings.

**log-rotation-service-type** [Variable]  
 This is the type of the log rotation service. Its associated value must be a `log-rotation-configuration` record, as discussed below.

**log-rotation-configuration** [Data Type]  
 Available `log-rotation-configuration` fields are:

**provision** (default: `(log-rotation)`) (type: list-of-symbols)  
 The name(s) of the log rotation Shepherd service.

- requirement** (default: `(user-processes)`) (type: list-of-symbols)  
Dependencies of the log rotation Shepherd service.
- calendar-event** (type: gexp)  
Gexp containing the *calendar event* when log rotation occurs. See Section “Timers” in *The GNU Shepherd Manual*, for more information on calendar events.
- external-log-files** (default: `()`) (type: list-of-strings)  
List of file names, external log files that should also be rotated.
- compression** (default: `zstd`) (type: gexp-or-symbol)  
The compression method used for rotated log files, one of `'none`, `'gzip`, and `'zstd`. Alternatively, it can be a gexp that evaluates to a procedure; that procedure gets called with the file to be rotated.
- expiry** (type: gexp-or-integer)  
Age in seconds after which a log file is deleted.
- size-threshold** (type: gexp-or-integer)  
Size in bytes below which a log file is *not* rotated.

Some log files just need to be deleted periodically once they are old, without any other criterion and without any archival step. This is the case of build logs stored by `guix-daemon` under `/var/log/guix/drvs` (see Section 2.3 [Invoking guix-daemon], page 14). The `log-cleanup` service addresses this use case. For example, `%base-services` (see Section 11.10.1 [Base Services], page 276) includes the following:

```
;; Periodically delete old build logs.
(service log-cleanup-service-type
  (log-cleanup-configuration
    (directory "/var/log/guix/drvs")))
```

That ensures build logs do not accumulate endlessly.

**log-cleanup-service-type** [Variable]  
This is the type of the service to delete old logs. Its value must be a `log-cleanup-configuration` record as described below.

**log-cleanup-configuration** [Data Type]  
Data type representing the log cleanup configuration

- directory**  
Name of the directory containing log files.
- expiry** (default: `(* 6 30 24 3600)`)  
Age in seconds after which a file is subject to deletion (six months by default).
- schedule** (default: `"30 12 01,08,15,22 * *"`)  
Schedule of the log cleanup job written either as a string in traditional cron syntax or as a gexp representing a Shepherd calendar event (see Section “Timers” in *The GNU Shepherd Manual*).

## Anonip Service

Anonip is a privacy filter that removes IP address from web server logs. This service creates a FIFO and filters any written lines with anonip before writing the filtered log to a target file.

The following example sets up the FIFO `/var/run/anonip/https.access.log` and writes the filtered log file `/var/log/anonip/https.access.log`.

```
(service anonip-service-type
  (anonip-configuration
    (input  "/var/run/anonip/https.access.log")
    (output "/var/log/anonip/https.access.log")))
```

Configure your web server to write its logs to the FIFO at `/var/run/anonip/https.access.log` and collect the anonymized log file at `/var/web-logs/https.access.log`.

**anonip-configuration** [Data Type]

This data type represents the configuration of anonip. It has the following parameters:

**anonip** (default: `anonip`)  
The anonip package to use.

**input**        The file name of the input log file to process. The service creates a FIFO of this name. The web server should write its logs to this FIFO.

**output**       The file name of the processed log file. This file is subject to log rotation *via* the Shepherd's log rotation service (see Section 11.10.3 [Log Rotation], page 304).

The following optional settings may be provided:

**debug?**        Print debug messages when `#true`.

**skip-private?**  
When `#true` do not mask addresses in private ranges.

**column**        A 1-based indexed column number. Assume IP address is in the specified column (default is 1).

**replacement**  
Replacement string in case address parsing fails, e.g. `"0.0.0.0"`.

**ipv4mask**      Number of bits to mask in IPv4 addresses.

**ipv6mask**      Number of bits to mask in IPv6 addresses.

**increment**  
Increment the IP address by the given number. By default this is zero.

**delimiter**  
Log delimiter string.

**regex**         Regular expression for detecting IP addresses. Use this instead of `column`.

### 11.10.4 Networking Setup

The (`gnu services networking`) module provides services to configure network interfaces and set up networking on your machine. Those services provide different ways for you to set up your machine: by declaring a static network configuration, by running a Dynamic Host Configuration Protocol (DHCP) client, or by running daemons such as NetworkManager and Connman that automate the whole process, automatically adapt to connectivity changes, and provide a high-level user interface.

On a laptop, NetworkManager and Connman are by far the most convenient options, which is why the default desktop services include NetworkManager (see Section 11.10.9 [Desktop Services], page 377). For a server, or for a virtual machine or a container, static network configuration or a simple DHCP client are often more appropriate.

This section describes the various network setup services available, starting with static network configuration.

**static-networking-service-type** [Variable]

This is the type for statically-configured network interfaces. Its value must be a list of `static-networking` records. Each of them declares a set of *addresses*, *routes*, and *links*, as shown below.

Here is the simplest configuration, with only one network interface controller (NIC) and only IPv4 connectivity:

```
;; Static networking for one NIC, IPv4-only.
(service static-networking-service-type
  (list (static-networking
        (addresses
         (list (network-address
                (device "eno1")
                (value "10.0.2.15/24")))))
        (routes
         (list (network-route
                (destination "default")
                (gateway "10.0.2.2"))))
        (name-servers '("10.0.2.3")))))
```

The snippet above can be added to the `services` field of your operating system configuration (see Section 11.2 [Using the Configuration System], page 244). It will configure your machine to have 10.0.2.15 as its IP address, with a 24-bit netmask for the local network—meaning that any 10.0.2.x address is on the local area network (LAN). Traffic to addresses outside the local network is routed *via* 10.0.2.2. Host names are resolved by sending domain name system (DNS) queries to 10.0.2.3.

**static-networking** [Data Type]

This is the data type representing a static network configuration.

As an example, here is how you would declare the configuration of a machine with a single network interface controller (NIC) available as `eno1`, and with one IPv4 and one IPv6 address:

```
;; Network configuration for one NIC, IPv4 + IPv6.
(static-networking
```

```
(addresses (list (network-address
                  (device "eno1")
                  (value "10.0.2.15/24"))
                (network-address
                  (device "eno1")
                  (value "2001:123:4567:101::1/64"))))
(routes (list (network-route
               (destination "default")
               (gateway "10.0.2.2"))
             (network-route
               (destination "default")
               (gateway "2020:321:4567:42::1"))))
(name-servers '("10.0.2.3"))
```

If you are familiar with the `ip` command of the `iproute2` package (<https://wiki.linuxfoundation.org/networking/iproute2>) found on Linux-based systems, the declaration above is equivalent to typing:

```
ip address add 10.0.2.15/24 dev eno1
ip address add 2001:123:4567:101::1/64 dev eno1
ip route add default via inet 10.0.2.2
ip route add default via inet6 2020:321:4567:42::1
```

Run `man 8 ip` for more info. Venerable GNU/Linux users will certainly know how to do it with `ifconfig` and `route`, but we'll spare you that.

The available fields of this data type are as follows:

**addresses**

**links** (default: '()')

**routes** (default: '()')

The list of `network-address`, `network-link`, and `network-route` records for this network (see below).

**name-servers** (default: '()')

The list of IP addresses (strings) of domain name servers. These IP addresses go to `/etc/resolv.conf`.

**provision** (default: '(networking)')

If true, this should be a list of symbols for the Shepherd service corresponding to this network configuration.

**requirement** (default '()')

The list of Shepherd services depended on.

**network-address**

[Data Type]

This is the data type representing the IP address of a network interface.

**device** The name of the network interface for this address—e.g., `"eno1"`.

**value** The actual IP address and network mask, in CIDR (Classless Inter-Domain Routing) notation ([https://en.wikipedia.org/wiki/CIDR#CIDR\\_notation](https://en.wikipedia.org/wiki/CIDR#CIDR_notation)), as a string.



For example, "10.0.2.15/24" denotes IPv4 address 10.0.2.15 on a 24-bit sub-network—all 10.0.2.x addresses are on the same local network.

**ipv6?** Whether **value** denotes an IPv6 address. By default this is automatically determined.

**network-route** [Data Type]

This is the data type representing a network route.

**destination**

The route destination (a string), either an IP address and network mask or "default" to denote the default route.

**source** (default: #f)

The route source.

**device** (default: #f)

The device used for this route—e.g., "eno2".

**ipv6?** (default: auto)

Whether this is an IPv6 route. By default this is automatically determined based on **destination** or **gateway**.

**gateway** (default: #f)

IP address (a string) through which traffic is routed.

**scope** (default: 'universe')

Scope of the route, as a symbol:

**universe** global route;

**site** route within the local autonomous system;

**link** route on this link;

**host** route on the local host;

**nowhere** means the destination does not exist.

**network-link** [Data Type]

Data type for a network link (see Section “Link” in *Guile-Netlink Manual*). During startup, network links are employed to construct or modify existing or virtual ethernet links. These ethernet links can be identified by their *name* or *mac-address*. If there is a need to create virtual interface, *name* and *type* fields are required.

**name** The name of the link—e.g., "v0p0" (default: #f).

**type** A symbol denoting the type of the link—e.g., 'veth' (default: #f).

**mac-address**

The mac-address of the link—e.g., "98:11:22:33:44:55" (default: #f).

**arguments**

List of arguments for this type of link.

Consider a scenario where a server equipped with a network interface which has multiple ports. These ports are connected to a switch, which supports link aggregation ([https://en.wikipedia.org/wiki/Link\\_aggregation](https://en.wikipedia.org/wiki/Link_aggregation)) (also known as bonding or NIC teaming). The switch uses port channels to consolidate multiple physical interfaces into one logical interface to provide higher bandwidth, load balancing, and link redundancy. When a port is added to a LAG (or link aggregation group), it inherits the properties of the port-channel. Some of these properties are VLAN membership, trunk status, and so on.

VLAN ([https://en.wikipedia.org/wiki/Virtual\\_LAN](https://en.wikipedia.org/wiki/Virtual_LAN)) (or virtual local area network) is a logical network that is isolated from other VLANs on the same physical network. This can be used to segregate traffic, improve security, and simplify network management.

With all that in mind let's configure our static network for the server. We will bond two existing interfaces together using 802.3ad schema and on top of it, build a VLAN interface with id 1055. We assign a static ip to our new VLAN interface.

```
(static-networking
  (links (list (network-link
    (name "bond0")
    (type 'bond)
    (arguments '((mode . "802.3ad")
      (miimon . 100)
      (lacp-active . "on")
      (lacp-rate . "fast")))))

    (network-link
      (mac-address "98:11:22:33:44:55")
      (arguments '((master . "bond0"))))

    (network-link
      (mac-address "98:11:22:33:44:56")
      (arguments '((master . "bond0"))))

    (network-link
      (name "bond0.1055")
      (type 'vlan)
      (arguments '((id . 1055)
        (link . "bond0")))))
  (addresses (list (network-address
    (value "192.168.1.4/24")
    (device "bond0.1055")))))
```

`%loopback-static-networking` [Variable]

This is the `static-networking` record representing the “loopback device”, `lo`, for IP addresses `127.0.0.1` and `::1`, and providing the `loopback` Shepherd service.

**%qemu-static-networking** [Variable]

This is the **static-networking** record representing network setup when using QEMU's user-mode network stack on **eth0** (see Section "Using the user mode network stack" in *QEMU Documentation*).

**dhcp-client-service-type** [Variable]

This is the type of services that run *dhclient*, the ISC Dynamic Host Configuration Protocol (DHCP) client.

**dhcp-client-configuration** [Data Type]

Data type representing the configuration of the ISC DHCP client service.

**package** (default: **isc-dhcp**)

The ISC DHCP client package to use.

**interfaces** (default: **'all'**)

Either **'all'** or the list of interface names that the ISC DHCP client should listen on—e.g., **'("eno1")'**.

When set to **'all'**, the ISC DHCP client listens on all the available non-loopback interfaces that can be activated. Otherwise the ISC DHCP client listens only on the specified interfaces.

**config-file** (default: **#f**)

The configuration file for the ISC DHCP client.

**version** (default: **"4"**)

The DHCP protocol version to use, as a string. Accepted values are **"4"** or **"6"** for DHCPv4 or DHCPv6, respectively, as well as **"4o6"**, for DHCPv4 over DHCPv6 (as specified by RFC 7341).

**shepherd-requirement** (default: **'()**)

**shepherd-provision** (default: **'(networking)'**)

This option can be used to provide a list of symbols naming Shepherd services that this service will depend on, such as **'wpa-supPLICANT'** or **'iwd'** if you require authenticated access for encrypted WiFi or Ethernet networks.

Likewise, **shepherd-provision** is a list of Shepherd service names (symbols) provided by this service. You might want to change the default value if you intend to run several ISC DHCP clients, only one of which provides the **networking** Shepherd service.

**dhcpcd-service-type** [Variable]

This the type for a service running *dhcpcd*, a DHCP (Dynamic Host Configuration Protocol) client that can be used as a replacement for the historical ISC client supported by **dhcp-client-service-type**.

Its value must be a **dhcpcd-configuration** record, as described below. As an example, consider the following setup which runs *dhcpcd* with a local DNS (Domain Name System) resolver:

```
(service dhcpcd-service-type
  (dhcpcd-configuration
```

```
(option '("rapid_commit" "interface_mtu"))
(no-option '("nd_rdnss"
             "dhcp6_name_servers"
             "domain_name_servers"
             "domain_name"
             "domain_search"))
(static '("domain_name_servers=127.0.0.1"))
(no-hook '("hostname")))
```

**dhcpcd-configuration**

[Data Type]

Available **dhcpcd-configuration** fields are:

**interfaces** (default: ()) (type: list)

List of networking interfaces—e.g., **eth0**—to start a DHCP client for. If no interface is specified (i.e., the list is empty) then **dhcpcd** discovers available Ethernet interfaces, that can be configured, automatically.

**command-arguments** (default: ("-q" "-q")) (type: list)

List of additional command-line options.

**host-name** (default: "") (type: maybe-string)

Host name to send via DHCP, defaults to the current system host name.

**duid** (default: "") (type: maybe-string)

DHCPv4 clients require a unique client identifier, this option uses the DHCPv6 Unique Identifier as a DHCPv4 client identifier as well. For more information, refer to RFC 4361 (<https://www.rfc-editor.org/rfc/rfc4361>) and **dhcpcd.conf**(5).

**persistent?** (default: **#t**) (type: boolean)

When true, automatically de-configure the interface when **dhcpcd** exits.

**option** (default: ("rapid\_commit" "domain\_name\_servers" "domain\_name" "domain\_search" "host\_name" "classless\_static\_routes" "interface\_mtu")) (type: list-of-strings)

List of options to request from the server.

**require** (default: ("dhcp\_server\_identifier")) (type: list-of-strings)

List of options to require in responses.

**slaac** (default: "private") (type: maybe-string)

Interface identifier used for SLAAC generated IPv6 addresses.

**no-option** (default: ()) (type: list-of-strings)

List of options to remove from the message before it's processed.

**no-hook** (default: ()) (type: list-of-strings)

List of hook script which should not be invoked.

**static** (default: ()) (type: list-of-strings)

DHCP client can request different options from a DHCP server, through **static** it is possible to configure static values for selected options. For example, **"domain\_name\_servers=127.0.0.1"**.

**vendor-class-id** (type: maybe-string)  
Set the DHCP Vendor Class (e.g., MSFT). For more information, refer to RFC 2132 (<https://www.rfc-editor.org/rfc/rfc2132#section-9.13>).

**client-id** (type: maybe-string)  
Use the interface hardware address or the given string as a client identifier, this is mutually exclusive with the **duid** option.

**extra-content** (type: maybe-string)  
Extra content to append to the configuration as-is.

**shepherd-provision** (default: (**networking**)) (type: list-of-symbols)  
This is a list of symbols naming Shepherd services provided by this service.

**shepherd-requirement** (default: ()) (type: list-of-symbols)  
This is a list of symbols naming Shepherd services that this service will depend on.

**network-manager-service-type** [Variable]  
This is the service type for the NetworkManager (<https://wiki.gnome.org/Projects/NetworkManager>) service. The value for this service type is a **network-manager-configuration** record.  
This service is part of **%desktop-services** (see Section 11.10.9 [Desktop Services], page 377).

**network-manager-configuration** [Data Type]  
Data type representing the configuration of NetworkManager.

**network-manager** (default: **network-manager**)  
The NetworkManager package to use.

**shepherd-requirement** (default: '(wireless-daemon))  
This option can be used to provide a list of symbols naming Shepherd services that this service will depend on. The default **'wireless-daemon** is provided by **'wpa-supPLICANT-service-type** and **'iwd-service-type**. Make sure you have one of these configured. Note, that **%desktop-services** already include **wpa-supPLICANT-service-type**.

**dns** (default: "default")  
Processing mode for DNS, which affects how NetworkManager uses the **resolv.conf** configuration file.

**'default'** NetworkManager will update **resolv.conf** to reflect the nameservers provided by currently active connections.

**'dnsmasq'** NetworkManager will run **dnsmasq** as a local caching name-server, using a *conditional forwarding* configuration if you are connected to a VPN, and then update **resolv.conf** to point to the local nameserver.

With this setting, you can share your network connection. For example when you want to share your network connection

to another laptop *via* an Ethernet cable, you can open `nm-connection-editor` and configure the Wired connection's method for IPv4 and IPv6 to be "Shared to other computers" and reestablish the connection (or reboot).

You can also set up a *host-to-guest connection* to QEMU VMs (see Section 3.8 [Installing Guix in a VM], page 31). With a host-to-guest connection, you can e.g. access a Web server running on the VM (see Section 11.10.20 [Web Services], page 494) from a Web browser on your host system, or connect to the VM *via* SSH (see Section 11.10.5 [Networking Services], page 322). To set up a host-to-guest connection, run this command once:

```
nmcli connection add type tun \
  connection.interface-name tap0 \
  tun.mode tap tun.owner $(id -u) \
  ipv4.method shared \
  ipv4.addresses 172.28.112.1/24
```

Then each time you launch your QEMU VM (see Section 11.18 [Running Guix in a VM], page 707), pass `-nic tap,ifname=tap0,script=no,downscript=no` to `qemu-system-....`

`'none'` NetworkManager will not modify `resolv.conf`.

`vpn-plugins` (default: `'()'`)

This is the list of available plugins for virtual private networks (VPNs). An example of this is the `network-manager-openvpn` package, which allows NetworkManager to manage VPNs *via* OpenVPN.

`extra-configuration-files` (default: `'()'`)

A list of two-element lists; the first element of each list is a file name (as a string), and the second is a file-like object. Used to specify configuration files which will be added to the `/etc/NetworkManager/conf.d`. NetworkManager will read additional configuration from this directory. For details on configuration file precedence and the configuration file format, see `'man 5 NetworkManager.conf'`.

For example, to add two files named `001-basic.conf` and `002-unmanaged.conf`:

```
(service network-manager-service-type
  (network-manager-configuration
    (extra-configuration-files
      `(("001-basic.conf" ,(local-file "basic.conf"))
        ("002-unmanaged.conf" ,(plain-file "constructed-unmanaged.
          "\
[keyfile]
unmanaged-devices=interface-name:wlo1_ap\n"))))))
```

**connman-service-type** [Variable]

This is the service type to run Connman (<https://01.org/connman>), a network connection manager.

Its value must be a **connman-configuration** record as in this example:

```
(service connman-service-type
  (connman-configuration
    (disable-vpn? #t)))
```

See below for details about **connman-configuration**.

**connman-configuration** [Data Type]

Data Type representing the configuration of connman.

**connman** (default: *connman*)

The connman package to use.

**shepherd-requirement** (default: '())

This option can be used to provide a list of symbols naming Shepherd services that this service will depend on, such as **'wpa-supPLICANT** or **'iwd** if you require authenticated access for encrypted WiFi or Ethernet networks.

**disable-vpn?** (default: **#f**)

When true, disable connman's vpn plugin.

**general-configuration** (default: (**connman-general-configuration**))

Configuration serialized to **main.conf** and passed as **--config** to **connmand**.

**connman-general-configuration** [Data Type]

Available **connman-general-configuration** fields are:

**input-request-timeout** (type: maybe-number)

Set input request timeout. Default is 120 seconds. The request for inputs like passphrase will timeout after certain amount of time. Use this setting to increase the value in case of different user interface designs.

**browser-launch-timeout** (type: maybe-number)

Set browser launch timeout. Default is 300 seconds. The request for launching a browser for portal pages will timeout after certain amount of time. Use this setting to increase the value in case of different user interface designs.

**background-scanning?** (type: maybe-boolean)

Enable background scanning. Default is true. If wifi is disconnected, the background scanning will follow a simple back off mechanism from 3s up to 5 minutes. Then, it will stay in 5 minutes unless user specifically asks for scanning through a D-Bus call. If so, the mechanism will start again from 3s. This feature activates also the background scanning while being connected, which is required for roaming on wifi. When **background-scanning?** is false, ConnMan will not perform any scan regardless of wifi is connected or not, unless it is requested by the user through a D-Bus call.

- use-gateways-as-timeservers?** (type: maybe-boolean)  
Assume that service gateways also function as timeservers. Default is false.
- fallback-timeservers** (type: maybe-list)  
List of Fallback timeservers. These timeservers are used for NTP sync when there are no timeservers set by the user or by the service, and when **use-gateways-as-timeservers?** is **#f**. These can contain a mixed combination of fully qualified domain names, IPv4 and IPv6 addresses.
- fallback-nameservers** (type: maybe-list)  
List of fallback nameservers appended to the list of nameservers given by the service. The nameserver entries must be in numeric format, host names are ignored.
- default-auto-connect-technologies** (type: maybe-list)  
List of technologies that are marked autoconnectable by default. The default value for this entry when empty is "ethernet", "wifi", "cellular". Services that are automatically connected must have been set up and saved to storage beforehand.
- default-favourite-technologies** (type: maybe-list)  
List of technologies that are marked favorite by default. The default value for this entry when empty is "ethernet". Connects to services from this technology even if not setup and saved to storage.
- always-connected-technologies** (type: maybe-list)  
List of technologies which are always connected regardless of preferred-technologies setting (**auto-connect? #t**). The default value is empty and this feature is disabled unless explicitly enabled.
- preferred-technologies** (type: maybe-list)  
List of preferred technologies from the most preferred one to the least preferred one. Services of the listed technology type will be tried one by one in the order given, until one of them gets connected or they are all tried. A service of a preferred technology type in state 'ready' will get the default route when compared to another preferred type further down the list with state 'ready' or with a non-preferred type; a service of a preferred technology type in state 'online' will get the default route when compared to either a non-preferred type or a preferred type further down in the list.
- network-interface-blacklist** (type: maybe-list)  
List of blacklisted network interfaces. Found interfaces will be compared to the list and will not be handled by ConnMan, if their first characters match any of the list entries. Default value is "vmnet", "vboxnet", "virbr", "ifb".
- allow-hostname-updates?** (type: maybe-boolean)  
Allow ConnMan to change the system hostname. This can happen for example if we receive DHCP hostname option. Default value is **#t**.



**allow-domainname-updates?** (type: maybe-boolean)

Allow connman to change the system domainname. This can happen for example if we receive DHCP domainname option. Default value is **#t**.

**single-connected-technology?** (type: maybe-boolean)

Keep only a single connected technology at any time. When a new service is connected by the user or a better one is found according to preferred-technologies, the new service is kept connected and all the other previously connected services are disconnected. With this setting it does not matter whether the previously connected services are in 'online' or 'ready' states, the newly connected service is the only one that will be kept connected. A service connected by the user will be used until going out of network coverage. With this setting enabled applications will notice more network breaks than normal. Note this options can't be used with VPNs. Default value is **#f**.

**tethering-technologies** (type: maybe-list)

List of technologies that are allowed to enable tethering. The default value is "wifi", "bluetooth", "gadget". Only those technologies listed here are used for tethering. If one wants to tether ethernet, then add "ethernet" in the list. Note that if ethernet tethering is enabled, then a DHCP server is started on all ethernet interfaces. Tethered ethernet should never be connected to corporate or home network as it will disrupt normal operation of these networks. Due to this ethernet is not tethered by default. Do not activate ethernet tethering unless you really know what you are doing.

**persistent-tethering-mode?** (type: maybe-boolean)

Restore earlier tethering status when returning from offline mode, re-enabling a technology, and after restarts and reboots. Default value is **#f**.

**enable-6to4?** (type: maybe-boolean)

Automatically enable anycast 6to4 if possible. This is not recommended, as the use of 6to4 will generally lead to a severe degradation of connection quality. See RFC6343. Default value is **#f** (as recommended by RFC6343 section 4.1).

**vendor-class-id** (type: maybe-string)

Set DHCP option 60 (Vendor Class ID) to the given string. This option can be used by DHCP servers to identify specific clients without having to rely on MAC address ranges, etc.

**enable-online-check?** (type: maybe-boolean)

Enable or disable use of HTTP GET as an online status check. When a service is in a READY state, and is selected as default, ConnMan will issue an HTTP GET request to verify that end-to-end connectivity is successful. Only then the service will be transitioned to ONLINE state. If this setting is false, the default service will remain in READY state. Default value is **#t**.

- online-check-ipv4-url** (type: maybe-string)  
IPv4 URL used during the online status check. Please refer to the README for more detailed information. Default value is `http://ipv4.connman.net/online/status.html`.
- online-check-ipv6-url** (type: maybe-string)  
IPv6 URL used during the online status check. Please refer to the README for more detailed information. Default value is `http://ipv6.connman.net/online/status.html`.
- online-check-initial-interval** (type: maybe-number)  
Range of intervals between two online check requests. Please refer to the README for more detailed information. Default value is '1'.
- online-check-max-interval** (type: maybe-number)  
Range of intervals between two online check requests. Please refer to the README for more detailed information. Default value is '1'.
- enable-online-to-ready-transition?** (type: maybe-boolean)  
WARNING: This is an experimental feature. In addition to **enable-online-check** setting, enable or disable use of HTTP GET to detect the loss of end-to-end connectivity. If this setting is **#f**, when the default service transitions to ONLINE state, the HTTP GET request is no more called until next cycle, initiated by a transition of the default service to DISCONNECT state. If this setting is **#t**, the HTTP GET request keeps being called to guarantee that end-to-end connectivity is still successful. If not, the default service will transition to READY state, enabling another service to become the default one, in replacement. Default value is **#f**.
- auto-connect-roaming-services?** (type: maybe-boolean)  
Automatically connect roaming services. This is not recommended unless you know you won't have any billing problem. Default value is **#f**.
- address-conflict-detection?** (type: maybe-boolean)  
Enable or disable the implementation of IPv4 address conflict detection according to RFC5227. ConnMan will send probe ARP packets to see if an IPv4 address is already in use before assigning the address to an interface. If an address conflict occurs for a statically configured address, an IPv4LL address will be chosen instead (according to RFC3927). If an address conflict occurs for an address offered via DHCP, ConnMan sends a DHCP DECLINE once and for the second conflict resorts to finding an IPv4LL address. Default value is **#f**.
- localtime** (type: maybe-string)  
Path to localtime file. Defaults to `/etc/localtime`.
- regulatory-domain-follows-timezone?** (type: maybe-boolean)  
Enable regulatory domain to be changed along timezone changes. With this option set to true each time the timezone changes the first present ISO3166 country code is read from `/usr/share/zoneinfo/zone1970.tab` and set as regulatory domain value. Default value is **#f**.

`resolve-conf` (type: maybe-string)

Path to `resolve.conf` file. If the file does not exist, but intermediate directories exist, it will be created. If this option is not set, it tries to write into `/var/run/connman/resolve.conf` if it fails (`/var/run/connman` does not exist or is not writeable). If you do not want to update `resolve.conf`, you can set `/dev/null`.

`wpa-supPLICANT-service-type` [Variable]

This is the service type to run WPA supplicant ([https://w1.fi/wpa\\_supplicant/](https://w1.fi/wpa_supplicant/)), an authentication daemon required to authenticate against encrypted WiFi or ethernet networks.

`wpa-supPLICANT-configuration` [Data Type]

Data type representing the configuration of WPA Supplicant.

It takes the following parameters:

`wpa-supPLICANT` (default: `wpa-supPLICANT`)

The WPA Supplicant package to use.

`shepherd-requirement` (default: `(user-processes loopback syslogd)`)

List of services that should be started before WPA Supplicant starts.

`shepherd-provision` (default: `(wpa-supPLICANT wireless-daemon)`) (type: list-of-symbols)

The name(s) of the service.

`dbus?` (default: `#t`)

Whether to listen for requests on D-Bus.

`pid-file` (default: `"/var/run/wpa_supplicant.pid"`)

Where to store the PID file.

`interface` (default: `#f`)

If this is set, it must specify the name of a network interface that WPA supplicant will control.

`config-file` (default: `#f`)

Optional configuration file to use.

`extra-options` (default: `'()`)

List of additional command-line arguments to pass to the daemon.

`iwd-service-type` [Variable]

This is the service type to run `iwd`, the iNet Wireless Daemon (<https://iwd.wiki.kernel.org/>), a daemon required to authenticate against encrypted WiFi networks.

**Warning:** By default `iwd` removes and re-creates interfaces it manages. It doesn't play nicely with `dhcpcd-client-service-type` that enumerates wireless interfaces before starting the Shepherd service. Use either `dhcpcd-service-type` or the `iwd`'s built-in DHCP client (see `enable-network-configuration` option below).

**Warning:** `iwd-service-type` conflicts with `wpa-supPLICANT-service-type`. Make sure only one is configured.

**iwd-configuration** [Data Type]

Available **iwd-configuration** fields are:

**iwd** (default: **iwd**) (type: file-like)

The IWD package to use.

**interfaces** (default: ()) (type: list-of-strings)

If this is set, it must specify *glob patterns* matching network interfaces that IWD will control.

**ignored-interfaces** (default: ()) (type: list-of-strings)

If this is set, it must specify *glob patterns* matching network interfaces that IWD will not manage.

**phys** (default: ()) (type: list-of-strings)

If this is set, it must specify *glob patterns* matching network PHYs names that IWD will control.

**ignored-phys** (default: ()) (type: list-of-strings)

If this is set, it must specify *glob patterns* matching network PHYs names that IWD will not manage.

**shepherd-requirement** (default: ()) (type: list-of-symbols)

Shepherd requirements the service should depend on.

**shepherd-provision** (default: (**iwd wireless-daemon**)) (type: list-of-symbols)

The name(s) of the service.

**config** (type: **iwd-settings**)

Configuration settings.

**iwd-settings** [Data Type]

Available **iwd-settings** fields are:

**general** (type: **iwd-general-settings**)

General settings.

**network** (default: (**iwd-network-settings**)) (type: maybe-iwd-network-settings)

Network settings.

**scan** (type: maybe-iwd-scan-settings)

Scan settings.

**extra-config** (default: ()) (type: list-of-strings)

Extra configuration values to append to the IWD configuration file.

**iwd-general-settings** [Data Type]

Available **iwd-general-settings** fields are:

**enable-network-configuration?** (default: **#f**) (type: boolean)

Setting this option to true enables **iwd** to configure the network interfaces with the IP addresses.

**extra-options** (default: ()) (type: alist)

An association list of option symbols/strings to string values to be appended to the General settings group.

**iwid-network-settings** [Data Type]

Available **iwid-network-settings** fields are:

**enable-ipv6?** (default: **#t**) (type: boolean)

Sets the global default that tells **iwid** whether it should configure IPv6 addresses and routes

**name-resolving-service** (default: **none**) (type: resolving-service)

Configures a DNS resolution method used by the system.

**extra-options** (default: **()**) (type: alist)

An association list of option symbols/strings to string values to be appended to the Network settings group.

**iwid-scan-settings** [Data Type]

Available **iwid-scan-settings** fields are:

**disable-periodic-scan?** (default: **#f**) (type: boolean)

Setting this option to **#t** will prevent **iwid** from issuing the periodic scans for the available networks while disconnected.

**initial-periodic-scan-interval** (default: 10) (type: number)

The initial periodic scan interval upon disconnect (in seconds).

**maximum-periodic-scan-interval** (default: 300) (type: number)

The maximum periodic scan interval (in seconds).

**disable-roaming-scan?** (default: **#f**) (type: boolean)

Setting this option to **#t** will prevent **iwid** from trying to scan when roaming decisions are activated.

**extra-options** (default: **()**) (type: alist)

An association list of option symbols/strings to string values to be appended to the Scan settings group.

Some networking devices such as modems require special care, and this is what the services below focus on.

**modem-manager-service-type** [Variable]

This is the service type for the ModemManager (<https://wiki.gnome.org/Projects/ModemManager>) service. The value for this service type is a **modem-manager-configuration** record.

This service is part of **%desktop-services** (see Section 11.10.9 [Desktop Services], page 377).

**modem-manager-configuration** [Data Type]

Data type representing the configuration of ModemManager.

**modem-manager** (default: **modem-manager**)

The ModemManager package to use.

**usb-modeswitch-service-type** [Variable]

This is the service type for the USB\_ModeSwitch ([https://www.draisberghof.de/usb\\_modeswitch/](https://www.draisberghof.de/usb_modeswitch/)) service. The value for this service type is a `usb-modeswitch-configuration` record.

When plugged in, some USB modems (and other USB devices) initially present themselves as a read-only storage medium and not as a modem. They need to be *modemswitched* before they are usable. The USB\_ModeSwitch service type installs udev rules to automatically modeswitch these devices when they are plugged in.

This service is part of `%desktop-services` (see Section 11.10.9 [Desktop Services], page 377).

**usb-modeswitch-configuration** [Data Type]

Data type representing the configuration of USB\_ModeSwitch.

`usb-modeswitch` (default: `usb-modeswitch`)

The USB\_ModeSwitch package providing the binaries for modeswitching.

`usb-modeswitch-data` (default: `usb-modeswitch-data`)

The package providing the device data and udev rules file used by USB\_ModeSwitch.

`config-file` (default: `#~(string-append #$usb-modeswitch:dispatcher "/etc/usb_modeswitch.conf")`)

Which config file to use for the USB\_ModeSwitch dispatcher. By default the config file shipped with USB\_ModeSwitch is used which disables logging to `/var/log` among other default settings. If set to `#f`, no config file is used.

### 11.10.5 Networking Services

The (`gnu services networking`) module discussed in the previous section provides services for more advanced setups: providing a DHCP service for others to use, filtering packets with iptables or nftables, running a WiFi access point with `hostapd`, running the `inetd` “superdaemon”, and more. This section describes those.

**dhcpd-service-type** [Variable]

This type defines a service that runs a DHCP daemon. To create a service of this type, you must supply a `<dhcpd-configuration>`. For example:

```
(service dhcpd-service-type
  (dhcpd-configuration
    (config-file (local-file "my-dhcpd.conf"))
    (interfaces '("enp0s25"))))
```

**dhcpd-configuration** [Data Type]

`package` (default: `isc-dhcp`)

The package that provides the DHCP daemon. This package is expected to provide the daemon at `sbin/dhcpd` relative to its output directory. The default package is the ISC’s DHCP server (<https://www.isc.org/dhcp/>).

**config-file** (default: `#f`)

The configuration file to use. This is required. It will be passed to `dhcpcd` via its `-cf` option. This may be any “file-like” object (see Section 8.12 [G-Expressions], page 165). See `man dhcpcd.conf` for details on the configuration file syntax.

**version** (default: `"4"`)

The DHCP version to use. The ISC DHCP server supports the values “4”, “6”, and “4o6”. These correspond to the `dhcpcd` program options `-4`, `-6`, and `-4o6`. See `man dhcpcd` for details.

**run-directory** (default: `"/run/dhcpcd"`)

The run directory to use. At service activation time, this directory will be created if it does not exist.

**pid-file** (default: `"/run/dhcpcd/dhcpcd.pid"`)

The PID file to use. This corresponds to the `-pf` option of `dhcpcd`. See `man dhcpcd` for details.

**interfaces** (default: `'()`)

The names of the network interfaces on which `dhcpcd` should listen for broadcasts. If this list is not empty, then its elements (which must be strings) will be appended to the `dhcpcd` invocation when starting the daemon. It may not be necessary to explicitly specify any interfaces here; see `man dhcpcd` for details.

**hostapd-service-type**

[Variable]

This is the service type to run the `hostapd` (<https://w1.fi/hostapd/>) daemon to set up WiFi (IEEE 802.11) access points and authentication servers. Its associated value must be a `hostapd-configuration` as shown below:

```
;; Use wlan1 to run the access point for "My Network".
(service hostapd-service-type
  (hostapd-configuration
    (interface "wlan1")
    (ssid "My Network")
    (channel 12)))
```

**hostapd-configuration**

[Data Type]

This data type represents the configuration of the `hostapd` service, with the following fields:

**package** (default: `hostapd`)

The `hostapd` package to use.

**interface** (default: `"wlan0"`)

The network interface to run the WiFi access point.

**ssid** The SSID (*service set identifier*), a string that identifies this network.

**broadcast-ssid?** (default: `#t`)

Whether to broadcast this SSID.

**channel** (default: 1)

The WiFi channel to use.

**driver** (default: "nl80211")

The driver interface type. "nl80211" is used with all Linux mac80211 drivers. Use "none" if building hostapd as a standalone RADIUS server that does not control any wireless/wired driver.

**extra-settings** (default: "")

Extra settings to append as-is to the hostapd configuration file. See <https://w1.fi/cgit/hostap/plain/hostapd/hostapd.conf> for the configuration file reference.

**simulated-wifi-service-type** [Variable]

This is the type of a service to simulate WiFi networking, which can be useful in virtual machines for testing purposes. The service loads the Linux kernel `mac80211_hwsim` module ([https://www.kernel.org/doc/html/latest/networking/mac80211\\_hwsim/mac80211\\_hwsim.html](https://www.kernel.org/doc/html/latest/networking/mac80211_hwsim/mac80211_hwsim.html)) and starts hostapd to create a pseudo WiFi network that can be seen on `wlan0`, by default.

The service's value is a `hostapd-configuration` record.

**iptables-service-type** [Variable]

This is the service type to set up an iptables configuration. iptables is a packet filtering framework supported by the Linux kernel. This service supports configuring iptables for both IPv4 and IPv6. A simple example configuration rejecting all incoming connections except those to the ssh port 22 is shown below.

```
(service iptables-service-type
  (iptables-configuration
    (ipv4-rules (plain-file "iptables.rules" "*filter
:INPUT ACCEPT
:FORWARD ACCEPT
:OUTPUT ACCEPT
-A INPUT -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
-A INPUT -p tcp --dport 22 -j ACCEPT
-A INPUT -j REJECT --reject-with icmp-port-unreachable
COMMIT
"))
    (ipv6-rules (plain-file "ip6tables.rules" "*filter
:INPUT ACCEPT
:FORWARD ACCEPT
:OUTPUT ACCEPT
-A INPUT -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
-A INPUT -p tcp --dport 22 -j ACCEPT
-A INPUT -j REJECT --reject-with icmp6-port-unreachable
COMMIT
")))))
```

**iptables-configuration** [Data Type]

The data type representing the configuration of iptables.



**iptables** (default: **iptables**)  
 The iptables package that provides **iptables-restore** and **ip6tables-restore**.

**ipv4-rules** (default: **%iptables-accept-all-rules**)  
 The iptables rules to use. It will be passed to **iptables-restore**. This may be any “file-like” object (see Section 8.12 [G-Expressions], page 165).

**ipv6-rules** (default: **%iptables-accept-all-rules**)  
 The ip6tables rules to use. It will be passed to **ip6tables-restore**. This may be any “file-like” object (see Section 8.12 [G-Expressions], page 165).

**nftables-service-type** [Variable]

This is the service type to set up a nftables configuration. nftables is a netfilter project that aims to replace the iptables, ip6tables, arptables and ebtables framework. It provides a new packet filtering framework, a new user-space utility **nft**, and a compatibility layer for iptables. This service comes with a default ruleset, **%default-nftables-ruleset**, that rejects all incoming connections except those to the SSH port 22 (TCP). To use it, simply write:

```
(service nftables-service-type)
```

**nftables-configuration** [Data Type]

Available **nftables-configuration** fields are:

**package** (default: **nftables**) (type: file-like)  
 The nftables package to use.

**debug-levels** (type: maybe-list-of-debug-levels)  
 A list of debug levels, for enabling debugging output. Valid debug level values are the ‘**scanner**’, ‘**parser**’, ‘**eval**’, ‘**netlink**’, ‘**mnl**’, ‘**proto-ctx**’, ‘**segtree**’ or ‘**all**’ symbols.

**ruleset** (type: file-like)  
 A file-like object containing the complete nftables ruleset. The default ruleset rejects all incoming connections except those to TCP port 22, with connections from the loopback interface are allowed.

**ntp-service-type** [Variable]

This is the type of the service running the Network Time Protocol (NTP) (<https://www.ntp.org>) daemon, **ntpd**. The daemon will keep the system clock synchronized with that of the specified NTP servers.

The value of this service is an **ntpd-configuration** object, as described below.

**ntp-configuration** [Data Type]

This is the data type for the NTP service configuration.

**servers** (default: **%ntp-servers**)  
 This is the list of servers (<**ntp-server**> records) with which **ntpd** will be synchronized. See the **ntp-server** data type definition below.

**allow-large-adjustment?** (default: **#t**)  
 This determines whether **ntpd** is allowed to make an initial adjustment of more than 1,000 seconds.

**ntp** (default: **ntp**)  
The NTP package to use.

**%ntp-servers** [Variable]  
List of host names used as the default NTP servers. These are servers of the NTP Pool Project (<https://www.ntppool.org/en/>).

**ntp-server** [Data Type]  
The data type representing the configuration of a NTP server.

**type** (default: **'server'**)  
The type of the NTP server, given as a symbol. One of **'pool'**, **'server'**, **'peer'**, **'broadcast'** or **'manycastclient'**.

**address** The address of the server, as a string.

**options** NTPD options to use with that specific server, given as a list of option names and/or of option names and values tuples. The following example define a server to use with the options **iburst** and **prefer**, as well as **version 3** and a **maxpoll** time of 16 seconds.

```
(ntp-server
  (type 'server)
  (address "some.ntp.server.org")
  (options `(iburst (version 3) (maxpoll 16) prefer))))
```

**openntpd-service-type** [Variable]  
Run the **ntpd**, the Network Time Protocol (NTP) daemon, as implemented by OpenNTPD (<http://www.openntpd.org>). The daemon will keep the system clock synchronized with that of the given servers.

```
(service
  openntpd-service-type
  (openntpd-configuration
    (listen-on '("127.0.0.1" ":::1"))
    (sensor '("udcf0 correction 70000"))
    (constraint-from '("www.gnu.org"))
    (constraints-from '("https://www.google.com/"))))
```

**%openntpd-servers** [Variable]  
This variable is a list of the server addresses defined in **%ntp-servers**.

**openntpd-configuration** [Data Type]

**openntpd** (default: **openntpd**)  
The openntpd package to use.

**listen-on** (default: **'("127.0.0.1" ":::1")'**)  
A list of local IP addresses or hostnames the ntpd daemon should listen on.

**query-from** (default: **'()**)  
A list of local IP address the ntpd daemon should use for outgoing queries.

- sensor** (default: '()')  
Specify a list of timedelta sensor devices `ntpd` should use. `ntpd` will listen to each sensor that actually exists and ignore non-existent ones. See upstream documentation (<https://man.openbsd.org/ntpd.conf>) for more information.
- server** (default: '()')  
Specify a list of IP addresses or hostnames of NTP servers to synchronize to.
- servers** (default: `%openntp-servers`)  
Specify a list of IP addresses or hostnames of NTP pools to synchronize to.
- constraint-from** (default: '()')  
`ntpd` can be configured to query the 'Date' from trusted HTTPS servers via TLS. This time information is not used for precision but acts as an authenticated constraint, thereby reducing the impact of unauthenticated NTP man-in-the-middle attacks. Specify a list of URLs, IP addresses or hostnames of HTTPS servers to provide a constraint.
- constraints-from** (default: '()')  
As with `constraint from`, specify a list of URLs, IP addresses or hostnames of HTTPS servers to provide a constraint. Should the hostname resolve to multiple IP addresses, `ntpd` will calculate a median constraint from all of them.

### **inetd-service-type** [Variable]

This service runs the `inetd` (see Section “`inetd` invocation” in *GNU Inetutils*) daemon. `inetd` listens for connections on internet sockets, and lazily starts the specified server program when a connection is made on one of these sockets.

The value of this service is an `inetd-configuration` object. The following example configures the `inetd` daemon to provide the built-in `echo` service, as well as an `smtp` service which forwards `smtp` traffic over `ssh` to a server `smtp-server` behind a gateway `hostname`:

```
(service
  inetd-service-type
  (inetd-configuration
    (entries (list
      (inetd-entry
        (name "echo")
        (socket-type 'stream)
        (protocol "tcp")
        (wait? #f)
        (user "root")))
      (inetd-entry
        (node "127.0.0.1")
        (name "smtp")
        (socket-type 'stream))
```

```
(protocol "tcp")
(wait? #f)
(user "root")
(program (file-append openssh "/bin/ssh"))
(arguments
  '("ssh" "-qT" "-i" "/path/to/ssh_key"
    "-W" "smtp-server:25" "user@hostname")))))))
```

See below for more details about `inetd`-configuration.

**inetd-configuration** [Data Type]

Data type representing the configuration of `inetd`.

**program** (default: `(file-append inetutils "/libexec/inetd")`)

The `inetd` executable to use.

**entries** (default: `'()`)

A list of `inetd` service entries. Each entry should be created by the `inetd-entry` constructor.

**inetd-entry** [Data Type]

Data type representing an entry in the `inetd` configuration. Each entry corresponds to a socket where `inetd` will listen for requests.

**node** (default: `#f`)

Optional string, a comma-separated list of local addresses `inetd` should use when listening for this service. See Section “Configuration file” in *GNU Inetutils* for a complete description of all options.

**name** A string, the name must correspond to an entry in `/etc/services`.

**socket-type**

One of `'stream`, `'dgram`, `'raw`, `'rdm` or `'seqpacket`.

**protocol** A string, must correspond to an entry in `/etc/protocols`.

**wait?** (default: `#t`)

Whether `inetd` should wait for the server to exit before listening to new service requests.

**user** A string containing the user (and, optionally, group) name of the user as whom the server should run. The group name can be specified in a suffix, separated by a colon or period, i.e. `"user"`, `"user:group"` or `"user.group"`.

**program** (default: `"internal"`)

The server program which will serve the requests, or `"internal"` if `inetd` should use a built-in service.

**arguments** (default: `'()`)

A list strings or file-like objects, which are the server program’s arguments, starting with the zeroth argument, i.e. the name of the program itself. For `inetd`’s internal services, this entry must be `'()` or `'("internal")`.

See Section “Configuration file” in *GNU Inetutils* for a more detailed discussion of each configuration field.

**opendht-service-type** [Variable]

This is the type of the service running a OpenDHT (<https://opendht.net>) node, **dhtnode**. The daemon can be used to host your own proxy service to the distributed hash table (DHT), for example to connect to with Jami, among other applications.

**Important:** When using the OpenDHT proxy server, the IP addresses it “sees” from the clients should be addresses reachable from other peers. In practice this means that a publicly reachable address is best suited for a proxy server, outside of your private network. For example, hosting the proxy server on a IPv4 private local network and exposing it via port forwarding could work for external peers, but peers local to the proxy would have their private addresses shared with the external peers, leading to connectivity problems.

The value of this service is a **opendht-configuration** object, as described below.

**opendht-configuration** [Data Type]

Available **opendht-configuration** fields are:

**opendht** (default: **opendht**) (type: file-like)

The **opendht** package to use.

**peer-discovery?** (default: **#f**) (type: boolean)

Whether to enable the multicast local peer discovery mechanism.

**enable-logging?** (default: **#f**) (type: boolean)

Whether to enable logging messages to syslog. It is disabled by default as it is rather verbose.

**debug?** (default: **#f**) (type: boolean)

Whether to enable debug-level logging messages. This has no effect if logging is disabled.

**bootstrap-host** (default: **"bootstrap.jami.net:4222"**) (type: maybe-string)

The node host name that is used to make the first connection to the network. A specific port value can be provided by appending the **:PORT** suffix. By default, it uses the Jami bootstrap nodes, but any host can be specified here. It’s also possible to disable bootstrapping by explicitly setting this field to the **%unset-value** value.

**port** (default: **4222**) (type: maybe-number)

The UDP port to bind to. When left unspecified, an available port is automatically selected.

**proxy-server-port** (type: maybe-number)

Spawn a proxy server listening on the specified port.

**proxy-server-port-tls** (type: maybe-number)

Spawn a proxy server listening to TLS connections on the specified port.

**tor-service-type** [Variable]

Type for a service that runs the Tor (<https://torproject.org>) anonymous networking daemon. The service is configured using a `<tor-configuration>` record. By default, the Tor daemon runs as the `tor` unprivileged user, which is a member of the `tor` group.

Services of this type can be extended by other services to specify *onion services* (in addition to those already specified in `tor-configuration`) as in this example:

```
(simple-service 'my-extra-onion-service tor-service-type
  (list (tor-onion-service-configuration
        (name "extra-onion-service")
        (mapping '((80 "127.0.0.1:8080"))))))
```

**tor-configuration** [Data Type]

**tor** (default: `tor`)

The package that provides the Tor daemon. This package is expected to provide the daemon at `bin/tor` relative to its output directory. The default package is the Tor Project's (<https://www.torproject.org>) implementation.

**config-file** (default: `(plain-file "empty" "")`)

The configuration file to use. It will be appended to a default configuration file, and the final configuration file will be passed to `tor` via its `-f` option. This may be any “file-like” object (see Section 8.12 [G-Expressions], page 165). See `man tor` for details on the configuration file syntax.

**hidden-services** (default: `'()`)

The list of `<tor-onion-service-configuration>` records to use. For any onion service you include in this list, appropriate configuration to enable the onion service will be automatically added to the default configuration file.

**socks-socket-type** (default: `'tcp`)

The default socket type that Tor should use for its SOCKS socket. This must be either `'tcp` or `'unix`. If it is `'tcp`, then by default Tor will listen on TCP port 9050 on the loopback interface (i.e., localhost). If it is `'unix`, then Tor will listen on the UNIX domain socket `/var/run/tor/socks-sock`, which will be made writable by members of the `tor` group.

If you want to customize the SOCKS socket in more detail, leave `socks-socket-type` at its default value of `'tcp` and use `config-file` to override the default by providing your own `SocksPort` option.

**control-socket?** (default: `#f`)

Whether or not to provide a “control socket” by which Tor can be controlled to, for instance, dynamically instantiate tor onion services. If `#t`, Tor will listen for control commands on the UNIX domain socket `/var/run/tor/control-sock`, which will be made writable by members of the `tor` group.

**transport-plugins** (default: '()')

The list of `<tor-transport-plugin>` records to use. For any transport plugin you include in this list, appropriate configuration line to enable transport plugin will be automatically added to the default configuration file.

**auto-start?** (default: `#t`)

Whether or not to start the Tor daemon on boot.

**tor-onion-service-configuration** [Data Type]

Data Type representing a Tor *Onion Service* configuration. See the Tor project's documentation (<https://community.torproject.org/onion-services/>) for more information. Available `tor-onion-service-configuration` fields are:

**name** (type: string)

Name for this Onion Service. This creates a `/var/lib/tor/hidden-services/name` directory, where the `hostname` file contains the `'onion'` host name for this Onion Service.

**mapping** (type: alist)

Association list of port to address mappings. The following example:

```
'((22 "127.0.0.1:22")
  (80 "127.0.0.1:8080"))
```

maps ports 22 and 80 of the Onion Service to the local ports 22 and 8080.

**tor-transport-plugin** [Data Type]

Data type representing a Tor pluggable transport plugin in `tor-configuration`. Pluggable transports are programs that disguise Tor traffic, which can be useful in case Tor is censored. See the the Tor project's documentation (<https://tb-manual.torproject.org/circumvention/>) and specification (<https://spec.torproject.org/pt-spec/index.html>) for more information.

Each transport plugin corresponds either to `ClientTransportPlugin ...` or to `ServerTransportPlugin ...` line in the default configuration file, see `man tor`. Available `tor-transport-plugin` fields are:

**role** (default: `'client`)

This must be either `'client` or `'server`. Otherwise, an error is raised. Set the `'server` value if you want to run a bridge to help censored users connect to the Tor network, see the Tor project's bridge guide (<https://community.torproject.org/relay/setup/bridge/>). Set the `'client` value if you want to connect to somebody else's bridge, see the Tor project's "Get Bridges" page (<https://bridges.torproject.org/>). In both cases the required additional configuration should be provided via `#:config-file` option of `tor-configuration`.

**protocol** (default: `"obfs4"`)

A string that specifies a pluggable transport protocol.

**program** This must be a "file-like" object or a string pointing to the pluggable transport plugin executable. This option allows the Tor daemon run

inside the container to access the executable and all the references (e.g. package dependencies) attached to it.

Suppose you would like Tor daemon to use obfs4 type obfuscation and to connect to Tor network via obfs4 bridge (a nonpublic Tor relay with support for obfs4 type obfuscation). Then you may go to <https://bridges.torproject.org/> (<https://bridges.torproject.org/>) and get there a couple of bridge lines (each starts with obfs4 ...) and use these lines in tor-service-type configuration as follows:

```
(service tor-service-type
  (tor-configuration
    (config-file (plain-file "torrc"
      "\
UseBridges 1
Bridge obfs4 ...
Bridge obfs4 ..."))
    (transport-plugins
      (list (tor-transport-plugin
        (program (file-append lyrebird "/bin/lyrebird"))))))))
```

The (gnu services rsync) module provides the following services:

You might want an rsync daemon if you have files that you want available so anyone (or just yourself) can download existing files or upload new files.

**rsync-service-type** [Variable]

This is the service type for the rsync (<https://rsync.samba.org>) daemon. The value for this service type is a **rsync-configuration** record as in this example:

```
;; Export two directories over rsync. By default rsync listens on
;; all the network interfaces.
(service rsync-service-type
  (rsync-configuration
    (modules (list (rsync-module
      (name "music")
      (file-name "/srv/zik")
      (read-only? #f))
      (rsync-module
        (name "movies")
        (file-name "/home/charlie/movies"))))))))
```

See below for details about **rsync-configuration**.

**rsync-configuration** [Data Type]

Data type representing the configuration for **rsync-service**.

**package** (default: *rsync*)  
rsync package to use.

**address** (default: **#f**)  
IP address on which **rsync** listens for incoming connections. If unspecified, it defaults to listening on all available addresses.



**port-number** (default: 873)  
 TCP port on which **rsync** listens for incoming connections. If port is less than 1024 **rsync** needs to be started as the **root** user and group.

**pid-file** (default: `"/var/run/rsyncd/rsyncd.pid"`)  
 Name of the file where **rsync** writes its PID.

**lock-file** (default: `"/var/run/rsyncd/rsyncd.lock"`)  
 Name of the file where **rsync** writes its lock file.

**log-file** (default: `"/var/log/rsyncd.log"`)  
 Name of the file where **rsync** writes its log file.

**user** (default: `"root"`)  
 Owner of the **rsync** process.

**group** (default: `"root"`)  
 Group of the **rsync** process.

**uid** (default: `"rsyncd"`)  
 User name or user ID that file transfers to and from that module should take place as when the daemon was run as **root**.

**gid** (default: `"rsyncd"`)  
 Group name or group ID that will be used when accessing the module.

**modules** (default: `%default-modules`)  
 List of “modules”—i.e., directories exported over **rsync**. Each element must be a **rsync-module** record, as described below.

#### **rsync-module** [Data Type]

This is the data type for **rsync** “modules”. A module is a directory exported over the **rsync** protocol. The available fields are as follows:

**name**           The module name. This is the name that shows up in URLs. For example, if the module is called **music**, the corresponding URL will be `rsync://host.example.org/music`.

**file-name**  
 Name of the directory being exported.

**comment** (default: `""`)  
 Comment associated with the module. Client user interfaces may display it when they obtain the list of available modules.

**read-only?** (default: `#t`)  
 Whether or not client will be able to upload files. If this is false, the uploads will be authorized if permissions on the daemon side permit it.

**chroot?** (default: `#t`)  
 When this is true, the **rsync** daemon changes root to the module’s directory before starting file transfers with the client. This improves security, but requires **rsync** to run as **root**.

**timeout** (default: 300)  
 Idle time in seconds after which the daemon closes a connection with the client.

The (`gnu services syncthing`) module provides the following services:

You might want a Syncthing daemon if you have files between two or more computers and want to sync them in real time, safely protected from prying eyes.

**syncthing-service-type** [Variable]

This is the service type for the syncthing (<https://syncthing.net/>) daemon, The value for this service type is a `syncthing-configuration` record as in this example:

```
(service syncthing-service-type
  (syncthing-configuration (user "alice")))
```

**Note:** This service is also available for Guix Home, where it runs directly with your user privileges (see Section 13.3.19 [Networking Home Services], page 779).

See below for details about `syncthing-configuration`.

**syncthing-configuration** [Data Type]

Data type representing the configuration for `syncthing-service-type`.

**syncthing** (default: `syncthing`)  
`syncthing` package to use.

**arguments** (default: `'()`)  
 List of command-line arguments passing to `syncthing` binary.

**logflags** (default: `0`)  
 Sum of logging flags, see Syncthing documentation logflags (<https://docs.syncthing.net/users/syncthing.html#cmdoption-logflags>).

**user** (default: `#f`)  
 The user as which the Syncthing service is to be run. This assumes that the specified user exists.

**group** (default: `"users"`)  
 The group as which the Syncthing service is to be run. This assumes that the specified group exists.

**home** (default: `#f`)  
 Sets the `HOME` variable for the Syncthing daemon. The default is `$HOME` of the specified Syncthing user.

**config-file** (default: `#f`)  
 Either a file-like object that resolves to a Syncthing configuration XML file, or a `syncthing-config-file` record (see below). If set to `#f`, Guix will not try to generate a config file, and Syncthing will generate a default configuration which will not be touched on reconfigure. Specifying this in a system service moves Syncthing's common configuration and data directory (`--home` in <https://docs.syncthing.net/users/syncthing.html>) to `/var/lib/syncthing-<user>`.

When used as a home service, syncthing configuration will be populated in `~/.local/state/syncthing/`. Old Syncthing installs may have files in `$XDG_CONFIG_HOME/syncthing/` (usually `~/.config/syncthing/`),

which will override configuration in `~/.local/state/syncthing` managed by this service. In such cases, you should move the contents of `$XDG_CONFIG_HOME/syncthing/` into `~/.local/state/syncthing/`.

This section documents a subset of the Syncthing configuration options—specifically those related to Guix or those affecting how your computer will connect to other computers over the network (such as Syncthing relays or discovery servers). The configuration is fully documented in the upstream Syncthing config documentation (<https://docs.syncthing.net/users/config.html>); camelCase there is converted to kebab-case here. If you are migrating from a Syncthing-managed configuration to one managed by Guix, you can check what changes were introduced by diffing the respective `config.xml` files. Note that you will need to add whitespace with 4-space indentation to the file generated by Guix, using the `xmllint` program from the `libxml2` package like so:

```
XMLLINT_INDENT="    " xmllint --format /path/to/new/config.xml | diff /path/to/old/con
```

When generating a configuration file through Guix, you can still temporarily modify Syncthing from the GUI or through `introducer` and `autoAcceptFolders` mechanisms, but such changes will be reset on reconfigure.

### `syncthing-config-file`

[Data Type]

Data type representing the configuration file read by the Syncthing daemon.

```
folders (default: (list (syncthing-folder (id "default") (label "Default
Folder") (path "~/Sync"))))
```

The default here is the same as Syncthing’s default. The value should be a list of `syncthing-folders`.

```
devices (default: '())
```

This should be a list of `syncthing-devices`. Guix will automatically add any devices specified in any ‘folders’ to this list. There are instances when you want to connect to a device despite not (initially) sharing any folders (such as a device with `autoAcceptFolders`). In such instances, you should specify those devices here. If multiple versions of the same device (as determined by comparing device ID) are discovered, the one in this list is prioritized. Otherwise, the first instance in the first folder is used.

```
gui-enabled? (default: #t)
```

By default, any user on the computer can access the GUI and make changes to Syncthing. If you leave this enabled, you should probably set `gui-user` and `gui-password` (see below).

```
gui-tls? (default: #f)
```

```
gui-debugging? (default: #f)
```

```
gui-send-basic-authorization-prompt? (default: #f)
```

```
gui-address (default: "127.0.0.1:8384")
```

```
gui-user (default: #f)
```

```
gui-password (default: #f)
```

A bcrypt hash of the GUI password. Remember that this will be globally exposed in `/gnu/store`.

`gui-apikey` (default: `#f`)

You must specify this to use the Syncthing REST interface. This key is kept in `/gnu/store` and is accessible to all users of the system.

`gui-theme` (default: `"default"`)

`ldap-enabled?` (default: `#f`)

`ldap-address` (default: `""`)

`ldap-bind-dn` (default: `""`)

`ldap-transport` (default: `""`)

`ldap-insecure-skip-verify` (default: `""`)

`ldap-search-base-dn` (default: `""`)

`ldap-search-filter` (default: `""`)

`listen-address` (default: `"default"`)

`global-announce-server` (default: `"default"`)

`global-announce-enabled?` (default: `#t`)

Global discovery servers can be used to help connect devices at unknown IP addresses by storing the last known IP address.

`local-announce-enabled?` (default: `#t`)

This makes devices find each other very easily on the same LAN. Often, this will allow you to just plug an Ethernet between two devices, or connect one device to the other's hotspot and start syncing.

`local-announce-port` (default: 21027)

`local-announce-mac-address` (default: `"[ff12::8384]:21027"`)

`max-send-kbps` (default: 0)

`max-receive-kbps` (default: 0)

`reconnection-interval-seconds` (default: 60)

`relays-enabled?` (default: `#t`)

This option allows your Syncthing instance to use a global network of relays (<https://docs.syncthing.net/users/relaying.html>) to enable syncing between devices when all other methods fail. As always, Syncthing traffic is encrypted in transport and the relays are unable to decrypt it.

`relay-reconnect-interval-minutes` (default: 10)

`start-browser?` (default: `#t`)

`nat-enabled?` (default: `#t`)

`nat-lease-minutes` (default: 60)

`nat-renewal-minutes` (default: 30)

`nat-timeout-seconds` (default: 10)

`usage-reporting-accepted` (default: 0)

Options whose names begin with 'usage-reporting-' control usage reporting. Set to -1 to disable, or to a positive value to enable. The default (0) disables reporting, but causes a usage reporting consent prompt to be displayed in the Syncthing GUI.

```

usage-reporting-seen (default: 0)
usage-reporting-unique-id (default: "")
usage-reporting-url (default: "https://data.syncthing.net/newdata")
usage-reporting-post-insecurely? (default: #f)
usage-reporting-initial-delay-seconds (default: 1800)
auto-upgrade-interval-hours (default: 12)
upgrade-to-pre-releases? (default: #f)
keep-temporaries-hours (default: 24)
cache-ignored-files? (default: #f)
progress-update-interval-seconds (default: 5)
limit-bandwidth-in-lan? (default: #f)
min-home-disk-free-unit (default: "%")
min-home-disk-free (default: 1)
releases-url (default: "https://upgrades.syncthing.net/meta.json")
overwrite-remote-device-names-on-connect? (default: #f)
temp-index-min-blocks (default: 10)
unacked-notification-id (default: "authenticationUserAndPassword")
traffic-class (default: 0)
set-low-priority? (default: #t)
max-folder-concurrency (default: 0)
crash-reporting-url (default: "https://crash.syncthing.net/newcrash")
crash-reporting-enabled? (default: #t)
stun-keepalive-start-seconds (default: 180)
stun-keepalive-min-seconds (default: 20)
stun-server (default: "default")
database-tuning (default: "auto")
max-concurrent-incoming-request-kib (default: 0)
announce-lan-addresses? (default: #t)
send-full-index-on-upgrade? (default: #f)
connection-limit-enough (default: 0)
connection-limit-max (default: 0)
insecure-allow-old-tls-versions? (default: #f)
connection-priority-tcp-lan (default: 10)
connection-priority-quic-lan (default: 20)
connection-priority-tcp-wan (default: 30)
connection-priority-quic-wan (default: 40)
connection-priority-relay (default: 50)
connection-priority-upgrade-threshold (default: 0)
default-folder (default: (syncthing-folder (label "")))
default-device (default: (syncthing-device (id "")))
default-ignores (default: "")

```

Options whose names begin with ‘default-’ above do not affect folders and devices added through the Guix configuration interface. They will, however, affect folders and devices that are added through the Syncthing GUI, by an **introducer**, or a device with **auto-accept-folders**.

**syncthing-folder**

[Data Type]

Data type representing a folder to be synchronized.

**id** (default: **#f**)

This ID cannot match the ID of any other folder on this device. If left unspecified, it will default to the label (see below).

**label**      A human readable label for the folder.

**path**      The path at which to store this folder.

**type** (default: "sendreceive")

**rescan-interval-seconds** (default: 3600)

**file-system-watcher-enabled?** (default: **#t**)

**file-system-watcher-delay-seconds** (default: 10)

**file-system-watcher-timeout-seconds** (default: 0)

**ignore-permissions?** (default: **#f**)

**auto-normalize?** (default: **#t**)

**devices** (default: '() )

This should be a list of other Syncthing devices. You do not need to specify the current device. Each device can be listed as a **syncthing-device** record or a **syncthing-folder-device** record if you want files to be encrypted on disk. See below.

```

filesystem-type (default: "basic")
min-disk-free-unit (default: "%")
min-disk-free (default: 1)
versioning-type (default: #f)
versioning-file-system-path (default: "")
versioning-file-system-type (default: "basic")
versioning-cleanup-interval-seconds (default: 3600)
versioning-cleanout-days (default: #f)
versioning-keep (default: #f)
versioning-max-age (default: #f)
versioning-command (default: #f)
copiers (default: 0)
puller-max-pending-kib (default: 0)
hashers (default: 0)
order (default: "random")
ignore-delete? (default: #f)
scan-progress-interval-seconds (default: 0)
puller-pause-seconds (default: 0)
max-conflicts (default: 10)
disable-sparse-files? (default: #f)
disable-temp-indexes? (default: #f)
paused? (default: #f)
weak-hash-threshold-percentage (default: 25)
marker-name (default: ".stfolder")
copy-ownership-from-parent? (default: #f)
mod-time-window-seconds (default: 0)
max-concurrent-writes (default: 2)
disable-fsync? (default: #f)
block-pull-order (default: "standard")
copy-range-method (default: "standard")
case-sensitive-file-system? (default: #f)
junctions-as-dirs? (default: #f)
sync-ownership? (default: #f)
send-ownership? (default: #f)
sync-xattrs? (default: #f)
send-xattrs? (default: #f)
xattr-filter-max-single-entry-size (default: 1024)
xattr-filter-max-total-size (default: 4096)

```

**syncthing-device**

[Data Type]

Data type representing a device to synchronize folders with.

**id** A long hash representing the keys generated by Syncthing on the first launch. You can obtain this from the Syncthing GUI or by inspecting an existing Syncthing configuration file.

**name** (default: "")

A human readable device name for viewing in the GUI or in Scheme.

```

compression (default: "metadata")
introducer? (default: #f)
skip-introduction-removals? (default: #f)
introduced-by (default: "")
addresses (default: '("dynamic"))
    List of addresses at which to search for this device. When the special
    value "dynamic" is included, Syncthing will search for the device locally
    as well as via the Syncthing project's global discovery (https://docs.
    syncthing.net/users/security.html#global-discovery) servers.

paused? (default: #f)
auto-accept-folders? (default: #f)
max-send-kbps (default: 0)
max-receive-kbps (default: 0)
max-request-kib (default: 0)
untrusted? (default: #f)
remote-gui-port (default: 0)
num-connections (default: 0)

```

#### syncthing-folder-device [Data Type]

This data type offers two folder-specific device options: `introduced-by` and `encryption-password`.

`syncthing-folder-device` corresponds to the 'device' (<https://docs.syncthing.net/users/config.html#config-option-folder.device>) option in the upstream 'folder' element.

If you don't need to use these options, then you can just use `syncthing-device` instead of `syncthing-folder-device` in the `devices` field of a `syncthing-folder` instance.

**device** The `syncthing-device` for which this configuration applies.

**introduced-by** (default: "")  
 The name of the device that "introduced" our device to the device sharing this folder. This is only used when "introduced" devices are removed by the introducer. See Syncthing introductions (<https://docs.syncthing.net/users/introducer.html>).

**encryption-password** (default: "")  
 The password used to encrypt data that is synchronized to untrusted devices.

Beware: specifying this field will include this password as plain text (not encrypted) and globally visible in `/gnu/store/`. If the `encryption-password` is non-empty, then it will be used as a password to encrypt file chunks as they are synchronized to untrusted devices. For more information on syncing to devices you don't totally trust, see Syncthing's documentation on Untrusted (Encrypted) Devices (<https://docs.syncthing.net/users/untrusted.html>). Note that data transfer is always encrypted while in transport ("end-to-end encryption"), regardless of this setting.



Here is a more complex example configuration for illustrative purposes:

```
(service syncthing-service-type
  (let ((laptop (syncthing-device (id "VHOD2D6-...-7XRMDEN")))
        (desktop (syncthing-device (id "64SAZ37-...-FZJ5GUA")
                                     (addresses '("tcp://example.com"))))
        (bob-desktop (syncthing-device (id "KYIMEGO-...-FT77EAO"))))
    (syncthing-configuration
     (user "alice")
     (config-file
      (syncthing-config-file
       (folders (list (syncthing-folder
                        (label "some-files")
                        (path "~/data")
                        (devices (list desktop laptop)))
                       (syncthing-folder
                        (label "critical-files")
                        (path "~/secrets")
                        (devices
                         (list desktop
                              laptop
                              (syncthing-folder-device
                               (device bob-desktop)
                               (encryption-password "mypassword"))))))))))))
```

Furthermore, (gnu services ssh) provides the following services.

**openssh-service-type** [Variable]

This is the type for the OpenSSH (<http://www.openssh.org>) secure shell daemon, `sshd`. Its value must be an `openssh-configuration` record as in this example:

```
(service openssh-service-type
  (openssh-configuration
   (x11-forwarding? #t)
   (permit-root-login 'prohibit-password)
   (authorized-keys
    `(("alice" ,(local-file "alice.pub"))
      ("bob" ,(local-file "bob.pub")))))
```

See below for details about `openssh-configuration`.

This service can be extended with extra authorized keys, as in this example:

```
(service-extension openssh-service-type
  (const `(("charlie"
            ,(local-file "charlie.pub")))))
```

**openssh-configuration** [Data Type]

This is the configuration record for OpenSSH's `sshd`.

**openssh** (default *openssh*)  
The OpenSSH package to use.

- pid-file** (default: `"/var/run/sshd.pid"`)  
Name of the file where `sshd` writes its PID.
- port-number** (default: 22)  
TCP port on which `sshd` listens for incoming connections.
- max-connections** (default: 200)  
Hard limit on the maximum number of simultaneous client connections, enforced by the inetd-style Shepherd service (see Section “Service De- and Constructors” in *The GNU Shepherd Manual*).
- permit-root-login** (default: `#f`)  
This field determines whether and when to allow logins as root. If `#f`, root logins are disallowed; if `#t`, they are allowed. If it’s the symbol `'prohibit-password`, then root logins are permitted but not with password-based authentication.
- allow-empty-passwords?** (default: `#f`)  
When true, users with empty passwords may log in. When false, they may not.
- password-authentication?** (default: `#t`)  
When true, users may log in with their password. When false, they have other authentication methods.
- public-key-authentication?** (default: `#t`)  
When true, users may log in using public key authentication. When false, users have to use other authentication method.  
Authorized public keys are stored in `~/.ssh/authorized_keys`. This is used only by protocol version 2.
- x11-forwarding?** (default: `#f`)  
When true, forwarding of X11 graphical client connections is enabled—in other words, `ssh` options `-X` and `-Y` will work.
- allow-agent-forwarding?** (default: `#t`)  
Whether to allow agent forwarding.
- allow-tcp-forwarding?** (default: `#t`)  
Whether to allow TCP forwarding.
- gateway-ports?** (default: `#f`)  
Whether to allow gateway ports.
- challenge-response-authentication?** (default: `#f`)  
Specifies whether challenge response authentication is allowed (e.g. via PAM).
- use-pam?** (default: `#t`)  
Enables the Pluggable Authentication Module interface. If set to `#t`, this will enable PAM authentication using `challenge-response-authentication?` and `password-authentication?`, in addition to PAM account and session module processing for all authentication types.

Because PAM challenge response authentication usually serves an equivalent role to password authentication, you should disable either `challenge-response-authentication?` or `password-authentication?`.

`print-last-log?` (default: `#t`)

Specifies whether `sshd` should print the date and time of the last user login when a user logs in interactively.

`subsystems` (default: `('("sftp" "internal-sftp"))`)

Configures external subsystems (e.g. file transfer daemon).

This is a list of two-element lists, each of which containing the subsystem name and a command (with optional arguments) to execute upon subsystem request.

The command `internal-sftp` implements an in-process SFTP server. Alternatively, one can specify the `sftp-server` command:

```
(service openssh-service-type
  (openssh-configuration
    (subsystems
      `(("sftp" ,(file-append openssh "/libexec/sftp-server")))))
```

`accepted-environment` (default: `()`)

List of strings describing which environment variables may be exported.

Each string gets on its own line. See the `AcceptEnv` option in `man sshd_config`.

This example allows ssh-clients to export the `COLORTERM` variable. It is set by terminal emulators, which support colors. You can use it in your shell's resource file to enable colors for the prompt and commands if this variable is set.

```
(service openssh-service-type
  (openssh-configuration
    (accepted-environment '("COLORTERM")))
```

`authorized-keys` (default: `()`)

This is the list of authorized keys. Each element of the list is a user name followed by one or more file-like objects that represent SSH public keys. For example:

```
(openssh-configuration
  (authorized-keys
    `(("rekado" ,(local-file "rekado.pub"))
      ("chris" ,(local-file "chris.pub"))
      ("root" ,(local-file "rekado.pub") ,(local-file "chris.pub")))))
```

registers the specified public keys for user accounts `rekado`, `chris`, and `root`.

Additional authorized keys can be specified *via* `service-extension`.

Note that this does *not* interfere with the use of `~/.ssh/authorized_keys`.

**generate-host-keys?** (default: **#t**)

Whether to generate host key pairs with **ssh-keygen -A** under **/etc/ssh** if there are none.

Generating key pairs takes a few seconds when enough entropy is available and is only done once. You might want to turn it off for instance in a virtual machine that does not need it because host keys are provided in some other way, and where the extra boot time is a problem.

**log-level** (default: **'info'**)

This is a symbol specifying the logging level: **quiet**, **fatal**, **error**, **info**, **verbose**, **debug**, etc. See the man page for **sshd\_config** for the full list of level names.

**extra-content** (default: **""**)

This field can be used to append arbitrary text to the configuration file. It is especially useful for elaborate configurations that cannot be expressed otherwise. This configuration, for example, would generally disable root logins, but permit them from one specific IP address:

```
(openssh-configuration
 (extra-content "\
Match Address 192.168.0.1
  PermitRootLogin yes"))
```

**dropbear-service-type** [Variable]

Type of the service that runs the Dropbear SSH daemon (<https://matt.ucc.asn.au/dropbear/dropbear.html>), whose value is a **<dropbear-configuration>** object.

For example, to specify a Dropbear service listening on port 1234:

```
(service dropbear-service-type (dropbear-configuration
                                (port-number 1234)))
```

**dropbear-configuration** [Data Type]

This data type represents the configuration of a Dropbear SSH daemon.

**dropbear** (default: *dropbear*)

The Dropbear package to use.

**port-number** (default: 22)

The TCP port where the daemon waits for incoming connections.

**syslog-output?** (default: **#t**)

Whether to enable syslog output.

**pid-file** (default: **"/var/run/dropbear.pid"**)

File name of the daemon's PID file.

**root-login?** (default: **#f**)

Whether to allow root logins.

**allow-empty-passwords?** (default: **#f**)

Whether to allow empty passwords.

**password-authentication?** (default: **#t**)

Whether to enable password-based authentication.

**endlessh-service-type** [Variable]

This is the type for the EndleSSH (<https://github.com/skeeto/endlessh>) program that runs an SSH tar pit. By very slowly sending an SSH banner, this program keeps unwanted visitors locked away from the actual SSH daemon.

For example, to specify a service running Endlessh on port 2222, add this call to the operating system's `services` field:

```
(service endlessh-service-type
  (endlessh-configuration
    (port-number 2222)))
```

**endlessh-configuration** [Data Type]

This data type represents the configuration of an EndleSSH service.

**endlessh** (default: *endlessh*)

The EndleSSH package to use.

**port-number** (default: 22)

The TCP port where the daemon waits for incoming connections.

**log-level** (default: 1)

The log level. 0 is quiet, 2 is very noisy.

**syslog-output?** (default: *#t*)

Whether to enable syslog output.

**pid-file** (default: *"/var/run/endlessh.pid"*)

File name of the daemon's PID file.

**message-delay** (default: 10000)

The endless banner is sent one line at a time. This is the delay in milliseconds between individual lines.

**max-banner-length** (default: 32)

The length of each line is randomized. This controls the maximum length of each line. Shorter lines may keep clients on for longer if they give up after a certain number of bytes.

**max-clients** (default: 4096)

Maximum number of connections to accept at a time. Connections beyond this are not immediately rejected, but will wait in the queue.

**autossh-service-type** [Variable]

This is the type for the AutoSSH (<https://www.harding.motd.ca/autossh>) program that runs a copy of `ssh` and monitors it, restarting it as necessary should it die or stop passing traffic. AutoSSH can be run manually from the command-line by passing arguments to the binary `autossh` from the package `autossh`, but it can also be run as a Guix service. This latter use case is documented here.

AutoSSH can be used to forward local traffic to a remote machine using an SSH tunnel, and it respects the `~/.ssh/config` of the user it is run as.

For example, to specify a service running `autossh` as the user `pino` and forwarding all local connections to port 8081 to `remote:8081` using an SSH tunnel, add this call to the operating system's `services` field:

```
(service autossh-service-type
  (autossh-configuration
    (user "pino")
    (ssh-options (list "-T" "-N" "-L" "8081:localhost:8081" "remote.net"))))
```

**autossh-configuration** [Data Type]

This data type represents the configuration of an AutoSSH service.

**user** (default "autossh")

The user as which the AutoSSH service is to be run. This assumes that the specified user exists.

**poll** (default 600)

Specifies the connection poll time in seconds.

**first-poll** (default #f)

Specifies how many seconds AutoSSH waits before the first connection test. After this first test, polling is resumed at the pace defined in `poll`. When set to #f, the first poll is not treated specially and will also use the connection poll specified in `poll`.

**gate-time** (default 30)

Specifies how many seconds an SSH connection must be active before it is considered successful.

**log-level** (default 1)

The log level, corresponding to the levels used by syslog—so 0 is the most silent while 7 is the chattiest.

**max-start** (default #f)

The maximum number of times SSH may be (re)started before AutoSSH exits. When set to #f, no maximum is configured and AutoSSH may restart indefinitely.

**message** (default "")

The message to append to the echo message sent when testing connections.

**port** (default "0")

The ports used for monitoring the connection. When set to "0", monitoring is disabled. When set to "`n`" where `n` is a positive integer, ports `n` and `n+1` are used for monitoring the connection, such that port `n` is the base monitoring port and `n+1` is the echo port. When set to "`n:m`" where `n` and `m` are positive integers, the ports `n` and `m` are used for monitoring the connection, such that port `n` is the base monitoring port and `m` is the echo port.

**ssh-options** (default '()')

The list of command-line arguments to pass to `ssh` when it is run. Options `-f` and `-M` are reserved for AutoSSH and may cause undefined behaviour.

**webssh-service-type** [Variable]

This is the type for the WebSSH (<https://webssh.huashengdun.org/>) program that runs a web SSH client. WebSSH can be run manually from the command-line by passing arguments to the binary `wssh` from the package `webssh`, but it can also be run as a Guix service. This latter use case is documented here.

For example, to specify a service running WebSSH on loopback interface on port 8888 with reject policy with a list of allowed to connection hosts, and NGINX as a reverse-proxy to this service listening for HTTPS connection, add this call to the operating system's `services` field:

```
(service webssh-service-type
  (webssh-configuration (address "127.0.0.1")
                        (port 8888)
                        (policy 'reject)
                        (known-hosts '("localhost ecdsa-sha2-nistp256 AAAA..."
                                       "127.0.0.1 ecdsa-sha2-nistp256 AAAA..."))))

(service nginx-service-type
  (nginx-configuration
    (server-blocks
      (list
        (nginx-server-configuration
          (inherit %webssh-configuration-nginx)
          (server-name '("webssh.example.com"))
          (listen '("443 ssl"))
          (ssl-certificate (letsencrypt-certificate "webssh.example.com"))
          (ssl-certificate-key (letsencrypt-key "webssh.example.com"))
          (locations
            (cons (nginx-location-configuration
                  (uri "/.well-known")
                  (body '("root /var/www;"))
                  (nginx-server-configuration-locations %webssh-configuration-n
```

**webssh-configuration** [Data Type]

Data type representing the configuration for `webssh-service`.

**package** (default: `webssh`)  
`webssh` package to use.

**user-name** (default: `"webssh"`)  
 User name or user ID that file transfers to and from that module should take place.

**group-name** (default: `"webssh"`)  
 Group name or group ID that will be used when accessing the module.

**address** (default: `#f`)  
 IP address on which `webssh` listens for incoming connections.

**port** (default: `8888`)  
 TCP port on which `webssh` listens for incoming connections.

**policy** (default: `#f`)

Connection policy. *reject* policy requires to specify *known-hosts*.

**known-hosts** (default: `'()`)

List of hosts which allowed for SSH connection from `webssh`.

**log-file** (default: `"/var/log/webssh.log"`)

Name of the file where `webssh` writes its log file.

**log-level** (default: `#f`)

Logging level.

**block-facebook-hosts-service-type** [Variable]

This service type adds a list of known Facebook hosts to the `/etc/hosts` file. (see Section “Host Names” in *The GNU C Library Reference Manual*) Each line contains an entry that maps a known server name of the Facebook on-line service—e.g., `www.facebook.com`—to unroutable IPv4 and IPv6 addresses.

This mechanism can prevent programs running locally, such as Web browsers, from accessing Facebook.

The `(gnu services avahi)` provides the following definition.

**avahi-service-type** [Variable]

This is the service that runs `avahi-daemon`, a system-wide mDNS/DNS-SD responder that allows for service discovery and “zero-configuration” host name lookups (see <https://avahi.org/>). Its value must be an `avahi-configuration` record—see below.

This service extends the name service cache daemon (`nscd`) so that it can resolve `.local` host names using `nss-mdns` (<https://0pointer.de/lennart/projects/nss-mdns/>). See Section 11.13 [Name Service Switch], page 680, for information on host name resolution.

Additionally, add the `avahi` package to the system profile so that commands such as `avahi-browse` are directly usable.

**avahi-configuration** [Data Type]

Data type representation the configuration for Avahi.

**host-name** (default: `#f`)

If different from `#f`, use that as the host name to publish for this machine; otherwise, use the machine’s actual host name.

**publish?** (default: `#t`)

When true, allow host names and services to be published (broadcast) over the network.

**publish-workstation?** (default: `#t`)

When true, `avahi-daemon` publishes the machine’s host name and IP address via mDNS on the local network. To view the host names published on your local network, you can run:

```
avahi-browse _workstation._tcp
```



**wide-area?** (default: **#f**)

When true, DNS-SD over unicast DNS is enabled.

**ipv4?** (default: **#t**)

**ipv6?** (default: **#t**)

These fields determine whether to use IPv4/IPv6 sockets.

**domains-to-browse** (default: **'()**)

This is a list of domains to browse.

**openvswitch-service-type**

[Variable]

This is the type of the Open vSwitch (<https://www.openvswitch.org>) service, whose value should be an **openvswitch-configuration** object.

**openvswitch-configuration**

[Data Type]

Data type representing the configuration of Open vSwitch, a multilayer virtual switch which is designed to enable massive network automation through programmatic extension.

**package** (default: *openvswitch*)

Package object of the Open vSwitch.

**pagekite-service-type**

[Variable]

This is the service type for the PageKite (<https://pagekite.net>) service, a tunneling solution for making localhost servers publicly visible, even from behind restrictive firewalls or NAT without forwarded ports. The value for this service type is a **pagekite-configuration** record.

Here's an example exposing the local HTTP and SSH daemons:

```
(service pagekite-service-type
  (pagekite-configuration
    (kites '("http:@kitename:localhost:80:@kitesecret"
            "raw/22:@kitename:localhost:22:@kitesecret"))
    (extra-file "/etc/pagekite.rc")))
```

**pagekite-configuration**

[Data Type]

Data type representing the configuration of PageKite.

**package** (default: *pagekite*)

Package object of PageKite.

**kitename** (default: **#f**)

PageKite name for authenticating to the frontend server.

**kitesecret** (default: **#f**)

Shared secret for authenticating to the frontend server. You should probably put this inside **extra-file** instead.

**frontend** (default: **#f**)

Connect to the named PageKite frontend server instead of the [pagekite.net](https://pagekite.net) service.

**kites** (default: `'("http:@kitename:localhost:80:@kitesecret")'`)  
 List of service kites to use. Exposes HTTP on port 80 by default. The format is `proto:kitename:host:port:secret`.

**extra-file** (default: `#f`)  
 Extra configuration file to read, which you are expected to create manually. Use this to add additional options and manage shared secrets out-of-band.

**yggdrasil-service-type** [Variable]

The service type for connecting to the Yggdrasil network (<https://yggdrasil-network.github.io/>), an early-stage implementation of a fully end-to-end encrypted IPv6 network.

Yggdrasil provides name-independent routing with cryptographically generated addresses. Static addressing means you can keep the same address as long as you want, even if you move to a new location, or generate a new address (by generating new keys) whenever you want. <https://yggdrasil-network.github.io/2018/07/28/addressing.html>

Pass it a value of `yggdrasil-configuration` to connect it to public peers and/or local peers.

Here is an example using public peers and a static address. The static signing and encryption keys are defined in `/etc/yggdrasil-private.conf` (the default value for `config-file`).

```
;; part of the operating-system declaration
(service yggdrasil-service-type
  (yggdrasil-configuration
    (autoconf? #f) ;; use only the public peers
    (json-config
      ;; choose one from
      ;; https://github.com/yggdrasil-network/public-peers
      '((peers . #("tcp://1.2.3.4:1337"))))
    ;; /etc/yggdrasil-private.conf is the default value for config-file
  ))

# sample content for /etc/yggdrasil-private.conf
{
  # Your private key. DO NOT share this with anyone!
  PrivateKey: 5c750...
}
```

**yggdrasil-configuration** [Data Type]

Data type representing the configuration of Yggdrasil.

**package** (default: `yggdrasil`)  
 Package object of Yggdrasil.

**json-config** (default: `'()`)  
 Contents of `/etc/yggdrasil.conf`. Will be merged with `/etc/yggdrasil-private.conf`. Note that these settings are stored in

the Guix store, which is readable to all users. **Do not store your private keys in it.** See the output of `yggdrasil -genconf` for a quick overview of valid keys and their default values.

`autoconf?` (default: `#f`)

Whether to use automatic mode. Enabling it makes Yggdrasil use a dynamic IP and peer with IPv6 neighbors.

`log-level` (default: `'info'`)

How much detail to include in logs. Use `'debug'` for more detail.

`log-to` (default: `'stdout'`)

Where to send logs. By default, the service logs standard output to `/var/log/yggdrasil.log`. The alternative is `'syslog'`, which sends output to the running syslog service.

`config-file` (default: `"/etc/yggdrasil-private.conf"`)

What HJSON file to load sensitive data from. This is where private keys should be stored, which are necessary to specify if you don't want a randomized address after each restart. Use `#f` to disable. Options defined in this file take precedence over `json-config`. Use the output of `yggdrasil -genconf` as a starting point. To configure a static address, delete everything except PrivateKey option.

`ipfs-service-type` [Variable]

The service type for connecting to the IPFS network (<https://ipfs.io>), a global, versioned, peer-to-peer file system. Pass it a `ipfs-configuration` to change the ports used for the gateway and API.

Here's an example configuration, using some non-standard ports:

```
(service ipfs-service-type
  (ipfs-configuration
    (gateway "/ip4/127.0.0.1/tcp/8880")
    (api "/ip4/127.0.0.1/tcp/8881"))))
```

`ipfs-configuration` [Data Type]

Data type representing the configuration of IPFS.

`package` (default: `go-ipfs`)

Package object of IPFS.

`gateway` (default: `"/ip4/127.0.0.1/tcp/8082"`)

Address of the gateway, in 'multiaddress' format.

`api` (default: `"/ip4/127.0.0.1/tcp/5001"`)

Address of the API endpoint, in 'multiaddress' format.

`keepalived-service-type` [Variable]

This is the type for the Keepalived (<https://www.keepalived.org/>) routing software, `keepalived`. Its value must be an `keepalived-configuration` record as in this example for master machine:

```
(service keepalived-service-type
```

```

        (keepalived-configuration
         (config-file (local-file "keepalived-master.conf"))))
where keepalived-master.conf:
    vrrp_instance my-group {
        state MASTER
        interface enp9s0
        virtual_router_id 100
        priority 100
        unicast_peer { 10.0.0.2 }
        virtual_ipaddress {
            10.0.0.4/24
        }
    }
and for backup machine:
    (service keepalived-service-type
      (keepalived-configuration
       (config-file (local-file "keepalived-backup.conf"))))
where keepalived-backup.conf:
    vrrp_instance my-group {
        state BACKUP
        interface enp9s0
        virtual_router_id 100
        priority 99
        unicast_peer { 10.0.0.3 }
        virtual_ipaddress {
            10.0.0.4/24
        }
    }

```

### 11.10.6 Unattended Upgrades

The (`gnu services admin`) module provides a service to perform *unattended upgrades*: periodically, the system automatically reconfigures itself from the latest Guix. Guix System has several properties that make unattended upgrades safe:

- upgrades are transactional (either the upgrade succeeds or it fails, but you cannot end up with an “in-between” system state);
- the upgrade log is kept—you can view it with `guix system list-generations`—and you can roll back to any previous generation, should the upgraded system fail to behave as intended;
- channel code is authenticated so you know you can only run genuine code (see Chapter 6 [Channels], page 68);
- `guix system reconfigure` prevents downgrades, which makes it immune to *downgrade attacks*.

To set up unattended upgrades, add an instance of `unattended-upgrade-service-type` like the one below to the list of your operating system services:

```
(service unattended-upgrade-service-type)
```

The defaults above set up weekly upgrades: every Sunday at midnight. You do not need to provide the operating system configuration file: it uses `/run/current-system/configuration.scm`, which ensures it always uses your latest configuration—see [provenance-service-type], page 716, for more information about this file.

There are several things that can be configured, in particular the periodicity and services (daemons) to be restarted upon completion. When the upgrade is successful, the service takes care of deleting system generations older than some threshold, as per `guix system delete-generations`. See the reference below for details.

To ensure that upgrades are actually happening, you can run `guix system describe`. To investigate upgrade failures, visit the unattended upgrade log file (see below).

**unattended-upgrade-service-type** [Variable]

This is the service type for unattended upgrades. It sets up a Shepherd Timer (see Section 11.10.2 [Scheduled Job Execution], page 299) that runs `guix system reconfigure` from the latest version of the specified channels.

Its value must be a `unattended-upgrade-configuration` record (see below).

**unattended-upgrade-configuration** [Data Type]

This data type represents the configuration of the unattended upgrade service. The following fields are available:

**schedule** (default: `"30 01 * * 0"`)

This is the schedule of upgrades, expressed as a string in traditional cron syntax or as a gexp evaluating to a Shepherd calendar event (see Section “Timers” in *The GNU Shepherd Manual*).

**channels** (default: `#%default-channels`)

This gexp specifies the channels to use for the upgrade (see Chapter 6 [Channels], page 68). By default, the tip of the official `guix` channel is used.

**operating-system-file** (default: `"/run/current-system/configuration.scm"`)

This field specifies the operating system configuration file to use. The default is to reuse the config file of the current configuration.

There are cases, though, where referring to `/run/current-system/configuration.scm` is not enough, for instance because that file refers to extra files (SSH public keys, extra configuration files, etc.) *via* `local-file` and similar constructs. For those cases, we recommend something along these lines:

```
(unattended-upgrade-configuration
  (operating-system-file
    (file-append (local-file "." "config-dir" #:recursive? #t)
                  "/config.scm")))
```

The effect here is to import all of the current directory into the store, and to refer to `config.scm` within that directory. Therefore, uses of `local-file` within `config.scm` will work as expected. See Section 8.12 [G-Expressions], page 165, for information about `local-file` and `file-append`.

**operating-system-expression** (default: `#f`)

This field specifies an expression that evaluates to the operating system to use for the upgrade. If no value is provided the **operating-system-file** field value is used.

```
(unattended-upgrade-configuration
 (operating-system-expression
  #~(@ (guix system install) installation-os)))
```

**reboot?** (default: `#f`)

This field specifies whether the system should reboot after completing an unattended upgrade.

When **reboot?** is `#t`, services are not restarted before rebooting. This means that the value for **services-to-restart** is ignored. The updated services will be started after the system reboots.

**services-to-restart** (default: `'(unattended-upgrade)`)

This field specifies the Shepherd services to restart when the upgrade completes.

Those services are restarted right away upon completion, as with **herd restart**, which ensures that the latest version is running—remember that by default **guix system reconfigure** only restarts services that are not currently running, which is conservative: it minimizes disruption but leaves outdated services running.

Use **herd status** to find out candidates for restarting. See Section 11.10 [Services], page 275, for general information about services. Common services to restart would include **ntpd** and **ssh-daemon**.

By default, the **unattended-upgrade** service is restarted. This ensures that the latest version of the unattended upgrade job will be used next time.

**system-expiration** (default: `(* 3 30 24 3600)`)

This is the expiration time in seconds for system generations. System generations older than this amount of time are deleted with **guix system delete-generations** when an upgrade completes.

**Note:** The unattended upgrade service does not run the garbage collector. You will probably want to set up your own **mcron** job to run **guix gc** periodically.

**maximum-duration** (default: `3600`)

Maximum duration in seconds for the upgrade; past that time, the upgrade aborts.

This is primarily useful to ensure the upgrade does not end up rebuilding or re-downloading “the world”.

**log-file** (default: `"/var/log/unattended-upgrade.log"`)

File where unattended upgrades are logged.

### 11.10.7 X Window

Support for the X Window graphical display system—specifically Xorg—is provided by the (`gnu services xorg`) module. Note that there is no `xorg-service` procedure. Instead, the X server is started by the *login manager*, by default the GNOME Display Manager (GDM).

GDM of course allows users to log in into window managers and desktop environments other than GNOME; for those using GNOME, GDM is required for features such as automatic screen locking.

To use X11, you must install at least one *window manager*—for example the `windowmaker` or `openbox` packages—preferably by adding it to the `packages` field of your operating system definition (see Section 11.3 [operating-system Reference], page 253).

GDM also supports Wayland: it can itself use Wayland instead of X11 for its user interface, and it can also start Wayland sessions. Wayland support is enabled by default. To disable it, set `wayland?` to `#f` in `gdm-configuration`.

`gdm-service-type` [Variable]

This is the type for the GNOME Desktop Manager (<https://wiki.gnome.org/Projects/GDM/>) (GDM), a program that manages graphical display servers and handles graphical user logins. Its value must be a `gdm-configuration` (see below).

GDM looks for *session types* described by the `.desktop` files in `/run/current-system/profile/share/xsessions` (for X11 sessions) and `/run/current-system/profile/share/wayland-sessions` (for Wayland sessions) and allows users to choose a session from the log-in screen. Packages such as `gnome`, `xfce`, `i3` and `sway` provide `.desktop` files; adding them to the system-wide set of packages automatically makes them available at the log-in screen.

In addition, `~/.xsession` files are honored. When available, `~/.xsession` must be an executable that starts a window manager and/or other X clients.

`gdm-configuration` [Data Type]

`auto-login?` (default: `#f`)

`default-user` (default: `#f`)

When `auto-login?` is false, GDM presents a log-in screen.

When `auto-login?` is true, GDM logs in directly as `default-user`.

`auto-suspend?` (default `#t`)

When true, GDM will automatically suspend to RAM when nobody is physically connected. When a machine is used via remote desktop or SSH, this should be set to false to avoid GDM interrupting remote sessions or rendering the machine unavailable.

`debug?` (default: `#f`)

When true, GDM writes debug messages to its log.

`gnome-shell-assets` (default: ...)

List of GNOME Shell assets needed by GDM: icon theme, fonts, etc.

`xorg-configuration` (default: `(xorg-configuration)`)

Configuration of the Xorg graphical server.

**x-session** (default: `(xinitrc)`)  
 Script to run before starting a X session.

**xmcp?** (default: `#f`)  
 When true, enable the X Display Manager Control Protocol (XDMCP). This should only be enabled in trusted environments, as the protocol is not secure. When enabled, GDM listens for XDMCP queries on the UDP port 177.

**dbus-daemon** (default: `dbus-daemon-wrapper`)  
 File name of the `dbus-daemon` executable.

**gdm** (default: `gdm`)  
 The GDM package to use.

**wayland?** (default: `#t`)  
 When true, enables Wayland in GDM, necessary to use Wayland sessions.

**wayland-session** (default: `gdm-wayland-session-wrapper`)  
 The Wayland session wrapper to use, needed to setup the environment.

**slim-service-type** [Variable]

This is the type for the SLiM graphical login manager for X11.

Like GDM, SLiM looks for session types described by `.desktop` files and allows users to choose a session from the log-in screen using **F1**. It also honors `~/.xsession` files.

Unlike GDM, SLiM does not spawn the user session on a different VT after logging in, which means that you can only start one graphical session. If you want to be able to run multiple graphical sessions at the same time you have to add multiple SLiM services to your system services. The following example shows how to replace the default GDM service with two SLiM services on `tty7` and `tty8`.

```
(use-modules (gnu services)
             (gnu services desktop)
             (gnu services xorg))

(operating-system
 ;; ...
 (services (cons* (service slim-service-type (slim-configuration
  (display ":0")
  (vt "vt7")))
                  (service slim-service-type (slim-configuration
  (display ":1")
  (vt "vt8"))))
            (modify-services %desktop-services
                              (delete gdm-service-type))))))
```

**slim-configuration** [Data Type]

Data type representing the configuration of `slim-service-type`.

**allow-empty-passwords?** (default: `#t`)  
 Whether to allow logins with empty passwords.



`gnupg?` (default: `#f`)

If enabled, `pam-gnupg` will attempt to automatically unlock the user's GPG keys with the login password via `gpg-agent`. The keygrips of all keys to be unlocked should be written to `~/.pam-gnupg`, and can be queried with `gpg -K --with-keygrip`. Presetting passphrases must be enabled by adding `allow-preset-passphrase` in `~/.gnupg/gpg-agent.conf`.

`auto-login?` (default: `#f`)

`default-user` (default: `"`)

When `auto-login?` is false, SLiM presents a log-in screen.

When `auto-login?` is true, SLiM logs in directly as `default-user`.

`theme` (default: `%default-slim-theme`)

`theme-name` (default: `%default-slim-theme-name`)

The graphical theme to use and its name.

`auto-login-session` (default: `#f`)

If true, this must be the name of the executable to start as the default session—e.g., `(file-append windowmaker "/bin/windowmaker")`.

If false, a session described by one of the available `.desktop` files in `/run/current-system/profile` and `~/.guix-profile` will be used.

**Note:** You must install at least one window manager in the system profile or in your user profile. Failing to do that, if `auto-login-session` is false, you will be unable to log in.

`xorg-configuration` (default `(xorg-configuration)`)

Configuration of the Xorg graphical server.

`display` (default `":0"`)

The display on which to start the Xorg graphical server.

`vt` (default `"vt7"`)

The VT on which to start the Xorg graphical server.

`xauth` (default: `xauth`)

The XAuth package to use.

`shepherd` (default: `shepherd`)

The Shepherd package used when invoking `halt` and `reboot`.

`sessreg` (default: `sessreg`)

The `sessreg` package used in order to register the session.

`slim` (default: `slim`)

The SLiM package to use.

`%default-theme`

[Variable]

`%default-theme-name`

[Variable]

The default SLiM theme and its name.

**sddm-service-type** [Variable]

This is the type of the service to run the SDDM display manager (<https://github.com/sddm/sddm>). Its value must be a **sddm-configuration** record (see below).

Here's an example use:

```
(service sddm-service-type
  (sddm-configuration
    (auto-login-user "alice")
    (auto-login-session "xfce.desktop")))
```

**sddm-configuration** [Data Type]

This data type represents the configuration of the SDDM login manager. The available fields are:

**sddm** (default: **sddm**)

The SDDM package to use.

**Note:** **sddm** has Qt6 enabled by default. If you want to still use a Qt5 theme, you need to set it to **sddm-qt5**.

**display-server** (default: "x11")

Select display server to use for the greeter. Valid values are "x11" or "wayland".

**numlock** (default: "on")

Valid values are "on", "off" or "none".

**halt-command** (default #~(string-append #\$/shepherd "/sbin/halt"))

Command to run when halting.

**reboot-command** (default #~(string-append #\$/shepherd "/sbin/reboot"))

Command to run when rebooting.

**theme** (default "maldives")

Theme to use. Default themes provided by SDDM are "elarun", "maldives" or "maya".

**themes-directory** (default "/run/current-system/profile/share/sddm/themes")

Directory to look for themes.

**faces-directory** (default "/run/current-system/profile/share/sddm/faces")

Directory to look for faces.

**default-path** (default "/run/current-system/profile/bin")

Default PATH to use.

**minimum-uid** (default: 1000)

Minimum UID displayed in SDDM and allowed for log-in.

**maximum-uid** (default: 2000)

Maximum UID to display in SDDM.

**remember-last-user?** (default #t)

Remember last user.

`remember-last-session?` (default `#t`)  
Remember last session.

`hide-users` (default `""`)  
Usernames to hide from SDDM greeter.

`hide-shells` (default `#~(string-append #$shadow "/sbin/nologin")`)  
Users with shells listed will be hidden from the SDDM greeter.

`session-command` (default `#~(string-append #$sddm  
"/share/sddm/scripts/wayland-session")`)  
Script to run before starting a wayland session.

`sessions-directory` (default  
`"/run/current-system/profile/share/wayland-sessions"`)  
Directory to look for desktop files starting wayland sessions.

`xorg-configuration` (default `(xorg-configuration)`)  
Configuration of the Xorg graphical server.

`xauth-path` (default `#~(string-append #$xauth "/bin/xauth")`)  
Path to xauth.

`xephyr-path` (default `#~(string-append #$xorg-server "/bin/Xephyr")`)  
Path to Xephyr.

`xdisplay-start` (default `#~(string-append #$sddm  
"/share/sddm/scripts/Xsetup")`)  
Script to run after starting xorg-server.

`xdisplay-stop` (default `#~(string-append #$sddm  
"/share/sddm/scripts/Xstop")`)  
Script to run before stopping xorg-server.

`xsession-command` (default: `xinitrc`)  
Script to run before starting a X session.

`xsessions-directory` (default: `"/run/current-system/profile/share/xsessions"`)  
Directory to look for desktop files starting X sessions.

`minimum-vt` (default: 7)  
Minimum VT to use.

`auto-login-user` (default `""`)  
User account that will be automatically logged in. Setting this to the empty string disables auto-login.

`auto-login-session` (default `""`)  
The `.desktop` file name to use as the auto-login session, or the empty string.

`relogin?` (default `#f`)  
Relogin after logout.

**lightdm-service-type** [Variable]

This is the type of the service to run the LightDM display manager (<https://github.com/canonical/lightdm>). Its value must be a **lightdm-configuration** record, which is documented below. Among its distinguishing features are TigerVNC integration for easily remotng your desktop as well as support for the XDMCP protocol, which can be used by remote clients to start a session from the login manager.

In its most basic form, it can be used simply as:

```
(service lightdm-service-type)
```

A more elaborate example making use of the VNC capabilities and enabling more features and verbose logs could look like:

```
(service lightdm-service-type
  (lightdm-configuration
    (allow-empty-passwords? #t)
    (xdmcp? #t)
    (vnc-server? #t)
    (vnc-server-command
      (file-append tigervnc-server "/bin/Xvnc"
        " -SecurityTypes None"))
    (seats
      (list (lightdm-seat-configuration
        (name "*")
        (user-session "ratpoison"))))))
```

**lightdm-configuration** [Data Type]

Available **lightdm-configuration** fields are:

**lightdm** (default: **lightdm**) (type: file-like)

The lightdm package to use.

**allow-empty-passwords?** (default: **#f**) (type: boolean)

Whether users not having a password set can login.

**debug?** (default: **#f**) (type: boolean)

Enable verbose output.

**xorg-configuration** (type: xorg-configuration)

The default Xorg server configuration to use to generate the Xorg server start script. It can be refined per seat via the **xserver-command** of the **<lightdm-seat-configuration>** record, if desired.

**greeters** (type: list-of-greeter-configurations)

The LightDM greeter configurations specifying the greeters to use.

**seats** (type: list-of-seat-configurations)

The seat configurations to use. A LightDM seat is akin to a user.

**xdmcp?** (default: **#f**) (type: boolean)

Whether a XDMCP server should listen on port UDP 177.

**xdmcp-listen-address** (type: maybe-string)

The host or IP address the XDMCP server listens for incoming connections. When unspecified, listen on for any hosts/IP addresses.

**vnc-server?** (default: `#f`) (type: boolean)  
Whether a VNC server is started.

**vnc-server-command** (type: file-like)  
The Xvnc command to use for the VNC server, it's possible to provide extra options not otherwise exposed along the command, for example to disable security:

```
(vnc-server-command (file-append tigervnc-server "/bin/Xvnc"
  " -SecurityTypes None" ))
```

Or to set a PasswordFile for the classic (unsecure) VncAuth mechanism:

```
(vnc-server-command (file-append tigervnc-server "/bin/Xvnc"
  " -PasswordFile /var/lib/lightdm/.vnc/
```

The password file should be manually created using the `vncpasswd` command. Note that LightDM will create new sessions for VNC users, which means they need to authenticate in the same way as local users would.

**vnc-server-listen-address** (type: maybe-string)  
The host or IP address the VNC server listens for incoming connections. When unspecified, listen for any hosts/IP addresses.

**vnc-server-port** (default: 5900) (type: number)  
The TCP port the VNC server should listen to.

**extra-config** (default: '()') (type: list-of-strings)  
Extra configuration values to append to the LightDM configuration file.

**lightdm-gtk-greeter-configuration** [Data Type]

Available `lightdm-gtk-greeter-configuration` fields are:

**lightdm-gtk-greeter** (default: `lightdm-gtk-greeter`) (type: file-like)  
The `lightdm-gtk-greeter` package to use.

**assets** (default: (`adwaita-icon-theme` `gnome-themes-extra` `hicolor-icon-theme`)) (type: list-of-file-likes)  
The list of packages complementing the greeter, such as package providing icon themes.

**theme-name** (default: `"Adwaita"`) (type: string)  
The name of the theme to use.

**icon-theme-name** (default: `"Adwaita"`) (type: string)  
The name of the icon theme to use.

**cursor-theme-name** (default: `"Adwaita"`) (type: string)  
The name of the cursor theme to use.

**cursor-theme-size** (default: 16) (type: number)  
The size to use for the cursor theme.

**allow-debugging?** (type: maybe-boolean)  
Set to `#t` to enable debug log level.

**background** (type: file-like)  
The background image to use.

**at-spi-enabled?** (default: `#f`) (type: boolean)  
 Enable accessibility support through the Assistive Technology Service Provider Interface (AT-SPI).

**a11y-states** (default: `(contrast font keyboard reader)`) (type: list-of-a11y-states)  
 The accessibility features to enable, given as list of symbols.

**reader** (type: maybe-file-like)  
 The command to use to launch a screen reader.

**extra-config** (default: `'()`) (type: list-of-strings)  
 Extra configuration values to append to the LightDM GTK Greeter configuration file.

### **lightdm-seat-configuration** [Data Type]

Available **lightdm-seat-configuration** fields are:

**name** (type: seat-name)  
 The name of the seat. An asterisk (\*) can be used in the name to apply the seat configuration to all the seat names it matches.

**user-session** (type: maybe-string)  
 The session to use by default. The session name must be provided as a lowercase string, such as "gnome", "ratpoison", etc.

**type** (default: `local`) (type: seat-type)  
 The type of the seat, either the `local` or `xremote` symbol.

**autologin-user** (type: maybe-string)  
 The username to automatically log in with by default.

**greeter-session** (default: `lightdm-gtk-greeter`) (type: greeter-session)  
 The greeter session to use, specified as a symbol. Currently, only `lightdm-gtk-greeter` is supported.

**xserver-command** (type: maybe-file-like)  
 The Xorg server command to run.

**session-wrapper** (type: file-like)  
 The xinitrc session wrapper to use.

**extra-config** (default: `'()`) (type: list-of-strings)  
 Extra configuration values to append to the seat configuration section.

### **xorg-configuration** [Data Type]

This data type represents the configuration of the Xorg graphical display server. Note that there is no Xorg service; instead, the X server is started by a “display manager” such as GDM, SDDM, LightDM or SLiM. Thus, the configuration of these display managers aggregates an **xorg-configuration** record.

**modules** (default: `%default-xorg-modules`)  
 This is a list of *module packages* loaded by the Xorg server—e.g., `xf86-video-vesa`, `xf86-input-libinput`, and so on.

**fonts** (default: `%default-xorg-fonts`)

This is a list of font directories to add to the server's *font path*.

**drivers** (default: `'()`)

This must be either the empty list, in which case Xorg chooses a graphics driver automatically, or a list of driver names that will be tried in this order—e.g., `'("modesetting" "vesa")`.

**resolutions** (default: `'()`)

When **resolutions** is the empty list, Xorg chooses an appropriate screen resolution. Otherwise, it must be a list of resolutions—e.g., `'((1024 768) (640 480))`.

**keyboard-layout** (default: `#f`)

If this is `#f`, Xorg uses the default keyboard layout—usually US English (“qwerty”) for a 105-key PC keyboard.

Otherwise this must be a **keyboard-layout** object specifying the keyboard layout in use when Xorg is running. See Section 11.8 [Keyboard Layout], page 271, for more information on how to specify the keyboard layout.

**extra-config** (default: `'()`)

This is a list of strings or objects appended to the configuration file. It is used to pass extra text to be added verbatim to the configuration file.

**server** (default: `xorg-server`)

This is the package providing the Xorg server.

**server-arguments** (default: `%default-xorg-server-arguments`)

This is the list of command-line arguments to pass to the X server. The default is `-nolisten tcp`.

**set-xorg-configuration** *config* [*login-manager-service-type*] [Procedure]

Tell the log-in manager (of type *login-manager-service-type*) to use *config*, an `<xorg-configuration>` record.

Since the Xorg configuration is embedded in the log-in manager's configuration—e.g., `gdm-configuration`—this procedure provides a shorthand to set the Xorg configuration.

**xorg-start-command** [*config*] [Procedure]

Return a `startx` script in which the modules, fonts, etc. specified in *config*, are available. The result should be used in place of `startx`.

Usually the X server is started by a login manager.

**xorg-start-command-xinit** [*config*] [Procedure]

Return a `startx` script in which the modules, fonts, etc. specified in *config* are available. The result should be used in place of `startx` and should be invoked by the user from a `tty` after login. Unlike **xorg-start-command**, this script calls `xinit`. Therefore it works well when executed from a `tty`. This script can be set up as `startx` using `startx-command-service-type` or `home-startx-command-service-type`. If you are using a desktop environment, you are unlikely to need this procedure.

**screen-locker-service-type** [Variable]

Type for a service that adds a package for a screen locker or screen saver to the set of privileged programs and/or add a PAM entry for it. The value for this service is a `<screen-locker-configuration>` object.

While the default behavior is to setup both a privileged program and PAM entry, these two methods are redundant. Screen locker programs may not execute when PAM is configured and `setuid` is set on their executable. In this case, `using-setuid?` can be set to `#f`.

For example, to make XlockMore usable:

```
(service screen-locker-service-type
  (screen-locker-configuration
    (name "xlock")
    (program (file-append xlockmore "/bin/xlock"))))
```

makes the good ol' XlockMore usable.

For example, swaylock fails to execute when compiled with PAM support and `setuid` enabled. One can thus disable `setuid`:

```
(service screen-locker-service-type
  (screen-locker-configuration
    (name "swaylock")
    (program (file-append swaylock "/bin/swaylock"))
    (using-pam? #t)
    (using-setuid? #f)))
```

**screen-locker-configuration** [Data Type]

Available `screen-locker-configuration` fields are:

`name` (type: string)

Name of the screen locker.

`program` (type: file-like)

Path to the executable for the screen locker as a G-Expression.

`allow-empty-password?` (default: `#f`) (type: boolean)

Whether to allow empty passwords.

`using-pam?` (default: `#t`) (type: boolean)

Whether to setup PAM entry.

`using-setuid?` (default: `#t`) (type: boolean)

Whether to setup program as `setuid` binary.

**startx-command-service-type** [Variable]

Add `startx` to the system profile putting it onto `PATH`.

The value for this service is a `<xorg-configuration>` object which is passed to the `xorg-start-command-xinit` procedure producing the `startx` used. Default value is `(xorg-configuration)`.



### 11.10.8 Printing Services

The (gnu services cups) module provides a Guix service definition for the CUPS printing service. To add printer support to a Guix system, add a cups-service to the operating system definition:

**cups-service-type** [Variable]

The service type for the CUPS print server. Its value should be a valid CUPS configuration (see below). To use the default settings, simply write:

```
(service cups-service-type)
```

The CUPS configuration controls the basic things about your CUPS installation: what interfaces it listens on, what to do if a print job fails, how much logging to do, and so on. To actually add a printer, you have to visit the <http://localhost:631> URL, or use a tool such as GNOME's printer configuration services. By default, configuring a CUPS service will generate a self-signed certificate if needed, for secure connections to the print server.

Suppose you want to enable the Web interface of CUPS and also add support for Epson printers *via* the `epson-inkjet-printer-escpr` package and for HP printers *via* the `hplip-minimal` package. You can do that directly, like this (you need to use the (gnu packages cups) module):

```
(service cups-service-type
  (cups-configuration
    (web-interface? #t)
    (extensions
      (list cups-filters epson-inkjet-printer-escpr hplip-minimal))))
```

**Note:** If you wish to use the Qt5 based GUI which comes with the `hplip` package then it is suggested that you install the `hplip` package, either in your OS configuration file or as your user.

The available configuration parameters follow. Each parameter definition is preceded by its type; for example, 'string-list foo' indicates that the `foo` parameter should be specified as a list of strings. There is also a way to specify the configuration as a string, if you have an old `cupsd.conf` file that you want to port over from some other system; see the end for more details.

Available cups-configuration fields are:

**package cups** [cups-configuration parameter]

The CUPS package.

**package-list extensions** (default: (list [cups-configuration parameter]

brlaser cups-filters epson-inkjet-printer-escpr  
foomatic-filters hplip-minimal splix))

Drivers and other extensions to the CUPS package.

**files-configuration files-configuration** [cups-configuration parameter]

Configuration of where to write logs, what directories to use for print spools, and related privileged configuration parameters.

Available files-configuration fields are:

- log-location access-log** [files-configuration parameter]  
 Defines the access log filename. Specifying a blank filename disables access log generation. The value **stderr** causes log entries to be sent to the standard error file when the scheduler is running in the foreground, or to the system log daemon when run in the background. The value **syslog** causes log entries to be sent to the system log daemon. The server name may be included in filenames using the string **%s**, as in **/var/log/cups/%s-access\_log**.  
 Defaults to **"/var/log/cups/access\_log"**.
- file-name cache-dir** [files-configuration parameter]  
 Where CUPS should cache data.  
 Defaults to **"/var/cache/cups"**.
- string config-file-perm** [files-configuration parameter]  
 Specifies the permissions for all configuration files that the scheduler writes.  
 Note that the permissions for the **printers.conf** file are currently masked to only allow access from the scheduler user (typically root). This is done because printer device URIs sometimes contain sensitive authentication information that should not be generally known on the system. There is no way to disable this security feature.  
 Defaults to **"0640"**.
- log-location error-log** [files-configuration parameter]  
 Defines the error log filename. Specifying a blank filename disables error log generation. The value **stderr** causes log entries to be sent to the standard error file when the scheduler is running in the foreground, or to the system log daemon when run in the background. The value **syslog** causes log entries to be sent to the system log daemon. The server name may be included in filenames using the string **%s**, as in **/var/log/cups/%s-error\_log**.  
 Defaults to **"/var/log/cups/error\_log"**.
- string fatal-errors** [files-configuration parameter]  
 Specifies which errors are fatal, causing the scheduler to exit. The kind strings are:
- none** No errors are fatal.
  - all** All of the errors below are fatal.
  - browse** Browsing initialization errors are fatal, for example failed connections to the DNS-SD daemon.
  - config** Configuration file syntax errors are fatal.
  - listen** Listen or Port errors are fatal, except for IPv6 failures on the loop-back or **any** addresses.
  - log** Log file creation or write errors are fatal.
  - permissions** Bad startup file permissions are fatal, for example shared TLS certificate and key files with world-read permissions.

Defaults to `"all -browse"`.

**boolean file-device?** [files-configuration parameter]

Specifies whether the file pseudo-device can be used for new printer queues.

The URI `file:///dev/null` is always allowed.

Defaults to `#f`.

**string group** [files-configuration parameter]

Specifies the group name or ID that will be used when executing external programs.

Defaults to `"lp"`.

**string log-file-group** [files-configuration parameter]

Specifies the group name or ID that will be used for log files.

Defaults to `"lpadmin"`.

**string log-file-perm** [files-configuration parameter]

Specifies the permissions for all log files that the scheduler writes.

Defaults to `"0644"`.

**log-location page-log** [files-configuration parameter]

Defines the page log filename. Specifying a blank filename disables page log generation. The value `stderr` causes log entries to be sent to the standard error file when the scheduler is running in the foreground, or to the system log daemon when run in the background. The value `syslog` causes log entries to be sent to the system log daemon. The server name may be included in filenames using the string `%s`, as in `/var/log/cups/%s-page_log`.

Defaults to `"/var/log/cups/page_log"`.

**string remote-root** [files-configuration parameter]

Specifies the username that is associated with unauthenticated accesses by clients claiming to be the root user. The default is `remroot`.

Defaults to `"remroot"`.

**file-name request-root** [files-configuration parameter]

Specifies the directory that contains print jobs and other HTTP request data.

Defaults to `"/var/spool/cups"`.

**sandboxing sandboxing** [files-configuration parameter]

Specifies the level of security sandboxing that is applied to print filters, backends, and other child processes of the scheduler; either `relaxed` or `strict`. This directive is currently only used/supported on macOS.

Defaults to `'strict'`.

**file-name server-keychain** [files-configuration parameter]

Specifies the location of TLS certificates and private keys. CUPS will look for public and private keys in this directory: `.crt` files for PEM-encoded certificates and corresponding `.key` files for PEM-encoded private keys.

Defaults to `"/etc/cups/ssl"`.

- file-name server-root** [files-configuration parameter]  
 Specifies the directory containing the server configuration files.  
 Defaults to `"/etc/cups"`.
- boolean sync-on-close?** [files-configuration parameter]  
 Specifies whether the scheduler calls `fsync(2)` after writing configuration or state files.  
 Defaults to `#f`.
- space-separated-string-list system-group** [files-configuration parameter]  
 Specifies the group(s) to use for `@SYSTEM` group authentication.
- file-name temp-dir** [files-configuration parameter]  
 Specifies the directory where temporary files are stored.  
 Defaults to `"/var/spool/cups/tmp"`.
- string user** [files-configuration parameter]  
 Specifies the user name or ID that is used when running external programs.  
 Defaults to `"lp"`.
- string set-env** [files-configuration parameter]  
 Set the specified environment variable to be passed to child processes.  
 Defaults to `"variable value"`.
- access-log-level access-log-level** [cups-configuration parameter]  
 Specifies the logging level for the AccessLog file. The `config` level logs when printers and classes are added, deleted, or modified and when configuration files are accessed or updated. The `actions` level logs when print jobs are submitted, held, released, modified, or canceled, and any of the conditions for `config`. The `all` level logs all requests.  
 Defaults to `'actions'`.
- boolean auto-purge-jobs?** [cups-configuration parameter]  
 Specifies whether to purge job history data automatically when it is no longer required for quotas.  
 Defaults to `#f`.
- comma-separated-string-list browse-dns-sd-sub-types** [cups-configuration parameter]  
 Specifies a list of DNS-SD sub-types to advertise for each shared printer.  
 The default `'(list "_cups" "_print" "_universal")'` tells clients that CUPS sharing, IPP Everywhere, AirPrint, and Mopria are supported.
- browse-local-protocols browse-local-protocols** [cups-configuration parameter]  
 Specifies which protocols to use for local printer sharing.  
 Defaults to `'dnssd'`.

- boolean browse-web-if?** [cups-configuration parameter]  
Specifies whether the CUPS web interface is advertised.  
Defaults to `'#f'`.
- boolean browsing?** [cups-configuration parameter]  
Specifies whether shared printers are advertised.  
Defaults to `'#f'`.
- default-auth-type default-auth-type** [cups-configuration parameter]  
Specifies the default type of authentication to use.  
Defaults to `'Basic'`.
- default-encryption default-encryption** [cups-configuration parameter]  
Specifies whether encryption will be used for authenticated requests.  
Defaults to `'Required'`.
- string default-language** [cups-configuration parameter]  
Specifies the default language to use for text and web content.  
Defaults to `"en"`.
- string default-paper-size** [cups-configuration parameter]  
Specifies the default paper size for new print queues. `"Auto"` uses a locale-specific default, while `"None"` specifies there is no default paper size. Specific size names are typically `"Letter"` or `"A4"`.  
Defaults to `"Auto"`.
- string default-policy** [cups-configuration parameter]  
Specifies the default access policy to use.  
Defaults to `"default"`.
- boolean default-shared?** [cups-configuration parameter]  
Specifies whether local printers are shared by default.  
Defaults to `'#t'`.
- non-negative-integer dirty-clean-interval** [cups-configuration parameter]  
Specifies the delay for updating of configuration and state files, in seconds. A value of 0 causes the update to happen as soon as possible, typically within a few milliseconds.  
Defaults to `'30'`.
- error-policy error-policy** [cups-configuration parameter]  
Specifies what to do when an error occurs. Possible values are `abort-job`, which will discard the failed print job; `retry-job`, which will retry the job at a later time; `retry-current-job`, which retries the failed job immediately; and `stop-printer`, which stops the printer.  
Defaults to `'stop-printer'`.

**non-negative-integer filter-limit** [cups-configuration parameter]

Specifies the maximum cost of filters that are run concurrently, which can be used to minimize disk, memory, and CPU resource problems. A limit of 0 disables filter limiting. An average print to a non-PostScript printer needs a filter limit of about 200. A PostScript printer needs about half that (100). Setting the limit below these thresholds will effectively limit the scheduler to printing a single job at any time.

Defaults to '0'.

**non-negative-integer filter-nice** [cups-configuration parameter]

Specifies the scheduling priority of filters that are run to print a job. The nice value ranges from 0, the highest priority, to 19, the lowest priority.

Defaults to '0'.

**host-name-lookups host-name-lookups** [cups-configuration parameter]

Specifies whether to do reverse lookups on connecting clients. The **double** setting causes **cupsd** to verify that the hostname resolved from the address matches one of the addresses returned for that hostname. Double lookups also prevent clients with unregistered addresses from connecting to your server. Only set this option to **#t** or **double** if absolutely required.

Defaults to '#f'.

**non-negative-integer job-kill-delay** [cups-configuration parameter]

Specifies the number of seconds to wait before killing the filters and backend associated with a canceled or held job.

Defaults to '30'.

**non-negative-integer job-retry-interval** [cups-configuration parameter]

Specifies the interval between retries of jobs in seconds. This is typically used for fax queues but can also be used with normal print queues whose error policy is **retry-job** or **retry-current-job**.

Defaults to '30'.

**non-negative-integer job-retry-limit** [cups-configuration parameter]

Specifies the number of retries that are done for jobs. This is typically used for fax queues but can also be used with normal print queues whose error policy is **retry-job** or **retry-current-job**.

Defaults to '5'.

**boolean keep-alive?** [cups-configuration parameter]

Specifies whether to support HTTP keep-alive connections.

Defaults to '#t'.

**non-negative-integer limit-request-body** [cups-configuration parameter]

Specifies the maximum size of print files, IPP requests, and HTML form data. A limit of 0 disables the limit check.

Defaults to '0'.

**multiline-string-list listen** [cups-configuration parameter]

Listens on the specified interfaces for connections. Valid values are of the form *address:port*, where *address* is either an IPv6 address enclosed in brackets, an IPv4 address, or *\** to indicate all addresses. Values can also be file names of local UNIX domain sockets. The Listen directive is similar to the Port directive but allows you to restrict access to specific interfaces or networks.

**location-access-control-list** [cups-configuration parameter]

**location-access-controls**

Specifies a set of additional access controls.

Available **location-access-controls** fields are:

**file-name path** [location-access-controls parameter]

Specifies the URI path to which the access control applies.

**access-control-list** [location-access-controls parameter]

**access-controls**

Access controls for all access to this path, in the same format as the **access-controls** of **operation-access-control**.

Defaults to **'()'**.

**method-access-control-list** [location-access-controls parameter]

**method-access-controls**

Access controls for method-specific access to this path.

Defaults to **'()'**.

Available **method-access-controls** fields are:

**boolean reverse?** [method-access-controls parameter]

If **#t**, apply access controls to all methods except the listed methods. Otherwise apply to only the listed methods.

Defaults to **#f**.

**method-list methods** [method-access-controls parameter]

Methods to which this access control applies.

Defaults to **'()'**.

**access-control-list** [method-access-controls parameter]

**access-controls**

Access control directives, as a list of strings. Each string should be one directive, such as **"Order allow,deny"**.

Defaults to **'()'**.

**non-negative-integer log-debug-history** [cups-configuration parameter]

Specifies the number of debugging messages that are retained for logging if an error occurs in a print job. Debug messages are logged regardless of the LogLevel setting.

Defaults to **'100'**.

- log-level log-level** [cups-configuration parameter]  
Specifies the level of logging for the ErrorLog file. The value **none** stops all logging while **debug2** logs everything.  
Defaults to 'info'.
- log-time-format log-time-format** [cups-configuration parameter]  
Specifies the format of the date and time in the log files. The value **standard** logs whole seconds while **usecs** logs microseconds.  
Defaults to 'standard'.
- non-negative-integer max-clients** [cups-configuration parameter]  
Specifies the maximum number of simultaneous clients that are allowed by the scheduler.  
Defaults to '100'.
- non-negative-integer max-clients-per-host** [cups-configuration parameter]  
Specifies the maximum number of simultaneous clients that are allowed from a single address.  
Defaults to '100'.
- non-negative-integer max-copies** [cups-configuration parameter]  
Specifies the maximum number of copies that a user can print of each job.  
Defaults to '9999'.
- non-negative-integer max-hold-time** [cups-configuration parameter]  
Specifies the maximum time a job may remain in the **indefinite** hold state before it is canceled. A value of 0 disables cancellation of held jobs.  
Defaults to '0'.
- non-negative-integer max-jobs** [cups-configuration parameter]  
Specifies the maximum number of simultaneous jobs that are allowed. Set to 0 to allow an unlimited number of jobs.  
Defaults to '500'.
- non-negative-integer max-jobs-per-printer** [cups-configuration parameter]  
Specifies the maximum number of simultaneous jobs that are allowed per printer. A value of 0 allows up to **max-jobs** per printer.  
Defaults to '0'.
- non-negative-integer max-jobs-per-user** [cups-configuration parameter]  
Specifies the maximum number of simultaneous jobs that are allowed per user. A value of 0 allows up to **max-jobs** per user.  
Defaults to '0'.



- non-negative-integer max-job-time** [cups-configuration parameter]  
Specifies the maximum time a job may take to print before it is canceled, in seconds.  
Set to 0 to disable cancellation of “stuck” jobs.  
Defaults to ‘10800’.
- non-negative-integer max-log-size** [cups-configuration parameter]  
Specifies the maximum size of the log files before they are rotated, in bytes. The value 0 disables log rotation.  
Defaults to ‘1048576’.
- non-negative-integer max-subscriptions** [cups-configuration parameter]  
Specifies the maximum number of simultaneous event subscriptions that are allowed.  
Set to ‘0’ to allow an unlimited number of subscriptions.  
Defaults to ‘0’.
- non-negative-integer max-subscriptions-per-job** [cups-configuration parameter]  
Specifies the maximum number of simultaneous event subscriptions that are allowed per job. A value of ‘0’ allows up to **max-subscriptions** per job.  
Defaults to ‘0’.
- non-negative-integer max-subscriptions-per-printer** [cups-configuration parameter]  
Specifies the maximum number of simultaneous event subscriptions that are allowed per printer. A value of ‘0’ allows up to **max-subscriptions** per printer.  
Defaults to ‘0’.
- non-negative-integer max-subscriptions-per-user** [cups-configuration parameter]  
Specifies the maximum number of simultaneous event subscriptions that are allowed per user. A value of ‘0’ allows up to **max-subscriptions** per user.  
Defaults to ‘0’.
- non-negative-integer multiple-operation-timeout** [cups-configuration parameter]  
Specifies the maximum amount of time to allow between files in a multiple file print job, in seconds.  
Defaults to ‘900’.
- environment-variables environment-variables** [cups-configuration parameter]  
Passes the specified environment variable(s) to child processes; a list of strings.  
Defaults to ‘()’.
- policy-configuration-list policies** [cups-configuration parameter]  
Specifies named access control policies.  
Available **policy-configuration** fields are:

- string name** [policy-configuration parameter]  
Name of the policy.
- string job-private-access** [policy-configuration parameter]  
Specifies an access list for a job's private values. @ACL maps to the printer's requesting-user-name-allowed or requesting-user-name-denied values. @OWNER maps to the job's owner. @SYSTEM maps to the groups listed for the system-group field of the files-configuration, which is reified into the cups-files.conf(5) file. Other possible elements of the access list include specific user names, and @group to indicate members of a specific group. The access list may also be simply all or default.  
Defaults to "@OWNER @SYSTEM".
- string job-private-values** [policy-configuration parameter]  
Specifies the list of job values to make private, or all, default, or none.  
Defaults to "job-name job-originating-host-name job-originating-user-name phone".
- string subscription-private-access** [policy-configuration parameter]  
Specifies an access list for a subscription's private values. @ACL maps to the printer's requesting-user-name-allowed or requesting-user-name-denied values. @OWNER maps to the job's owner. @SYSTEM maps to the groups listed for the system-group field of the files-configuration, which is reified into the cups-files.conf(5) file. Other possible elements of the access list include specific user names, and @group to indicate members of a specific group. The access list may also be simply all or default.  
Defaults to "@OWNER @SYSTEM".
- string subscription-private-values** [policy-configuration parameter]  
Specifies the list of job values to make private, or all, default, or none.  
Defaults to "notify-events notify-pull-method notify-recipient-uri notify-subscriber-user-name notify-user-data".
- operation-access-control-list access-controls** [policy-configuration parameter]  
Access control by IPP operation.  
Defaults to '()'.
- boolean-or-non-negative-integer preserve-job-files** [cups-configuration parameter]  
Specifies whether job files (documents) are preserved after a job is printed. If a numeric value is specified, job files are preserved for the indicated number of seconds after printing. Otherwise a boolean value applies indefinitely.  
Defaults to '86400'.

**boolean-or-non-negative-integer** [cups-configuration parameter]  
**preserve-job-history**

Specifies whether the job history is preserved after a job is printed. If a numeric value is specified, the job history is preserved for the indicated number of seconds after printing. If **#t**, the job history is preserved until the MaxJobs limit is reached.

Defaults to '**#t**'.

**comma-separated-string-list-or-#f** [cups-configuration parameter]  
**ready-paper-sizes**

Specifies a list of potential paper sizes that are reported as ready, that is: loaded. The actual list will contain only the sizes that each printer supports.

The default value of **#f** is a special case: CUPS will use '**(list "Letter" "Legal" "Tabloid" "4x6" "Env10")**' if the default paper size is '**Letter**', and '**(list "A3" "A4" "A5" "A6" "EnvDL")**' otherwise.

**non-negative-integer reload-timeout** [cups-configuration parameter]

Specifies the amount of time to wait for job completion before restarting the scheduler.

Defaults to '**30**'.

**string server-admin** [cups-configuration parameter]

Specifies the email address of the server administrator.

Defaults to '**root@localhost.localdomain**'.

**host-name-list-or-*server-alias*** [cups-configuration parameter]

The **ServerAlias** directive is used for HTTP Host header validation when clients connect to the scheduler from external interfaces. Using the special name **\*** can expose your system to known browser-based DNS rebinding attacks, even when accessing sites through a firewall. If the auto-discovery of alternate names does not work, we recommend listing each alternate name with a **ServerAlias** directive instead of using **\***.

Defaults to '**\***'.

**string server-name** [cups-configuration parameter]

Specifies the fully-qualified host name of the server.

Defaults to '**localhost**'.

**server-tokens server-tokens** [cups-configuration parameter]

Specifies what information is included in the Server header of HTTP responses. **None** disables the Server header. **ProductOnly** reports CUPS. **Major** reports CUPS 2. **Minor** reports CUPS 2.0. **Minimal** reports CUPS 2.0.0. **OS** reports CUPS 2.0.0 (**uname**) where **uname** is the output of the **uname** command. **Full** reports CUPS 2.0.0 (**uname**) IPP/2.0.

Defaults to '**Minimal**'.

**multiline-string-list ssl-listen** [cups-configuration parameter]

Listens on the specified interfaces for encrypted connections. Valid values are of the form **address:port**, where **address** is either an IPv6 address enclosed in brackets, an IPv4 address, or **\*** to indicate all addresses.

Defaults to '**()**'.

**ssl-options ssl-options** [cups-configuration parameter]  
 Sets encryption options. By default, CUPS only supports encryption using TLS v1.0 or higher using known secure cipher suites. Security is reduced when **Allow** options are used, and enhanced when **Deny** options are used. The **AllowRC4** option enables the 128-bit RC4 cipher suites, which are required for some older clients. The **AllowSSL3** option enables SSL v3.0, which is required for some older clients that do not support TLS v1.0. The **DenyCBC** option disables all CBC cipher suites. The **DenyTLS1.0** option disables TLS v1.0 support - this sets the minimum protocol version to TLS v1.1.

Defaults to ‘‘()’.

**boolean strict-conformance?** [cups-configuration parameter]  
 Specifies whether the scheduler requires clients to strictly adhere to the IPP specifications.

Defaults to ‘#f’.

**non-negative-integer timeout** [cups-configuration parameter]  
 Specifies the HTTP request timeout, in seconds.

Defaults to ‘900’.

**boolean web-interface?** [cups-configuration parameter]  
 Specifies whether the web interface is enabled.

Defaults to ‘#f’.

At this point you’re probably thinking “oh dear, Guix manual, I like you but you can stop already with the configuration options”. Indeed. However, one more point: it could be that you have an existing `cupsd.conf` that you want to use. In that case, you can pass an `opaque-cups-configuration` as the configuration of a `cups-service-type`.

Available `opaque-cups-configuration` fields are:

**package cups** [opaque-cups-configuration parameter]  
 The CUPS package.

**string cupsd.conf** [opaque-cups-configuration parameter]  
 The contents of the `cupsd.conf`, as a string.

**string cups-files.conf** [opaque-cups-configuration parameter]  
 The contents of the `cups-files.conf` file, as a string.

For example, if your `cupsd.conf` and `cups-files.conf` are in strings of the same name, you could instantiate a CUPS service like this:

```
(service cups-service-type
  (opaque-cups-configuration
    (cupsd.conf cupsd.conf)
    (cups-files.conf cups-files.conf)))
```

### 11.10.9 Desktop Services

The (`gnu services desktop`) module provides services that are usually useful in the context of a “desktop” setup—that is, on a machine running a graphical display server, possibly with graphical user interfaces, etc. It also defines services that provide specific desktop environments like GNOME, Xfce or MATE.

To simplify things, the module defines a variable containing the set of services that users typically expect on a machine with a graphical environment and networking:

**%desktop-services** [Variable]

This is a list of services that builds upon **%base-services** and adds or adjusts services for a typical “desktop” setup.

In particular, it adds a graphical login manager (see Section 11.10.7 [X Window], page 355), screen lockers, a network management tool (see Section 11.10.5 [Networking Services], page 322) with modem support (see Section 11.10.5 [Networking Services], page 322), energy and color management services, the `eelogind` login and seat manager, the Polkit privilege service, the GeoClue location service, the AccountsService daemon that allows authorized users change system passwords, a NTP client (see Section 11.10.5 [Networking Services], page 322) and the Avahi daemon.

The **%desktop-services** variable can be used as the **services** field of an **operating-system** declaration (see Section 11.3 [operating-system Reference], page 253).

Additionally, the following procedures add one (or more!) desktop environments to a system.

- `gnome-desktop-service-type` adds GNOME,
- `plasma-desktop-service-type` adds KDE Plasma,
- `enlightenment-desktop-service-type` adds Enlightenment,
- `lxqt-desktop-service-type` adds LXQt,
- `mate-desktop-service-type` adds MATE, and
- `xfce-desktop-service-type` adds Xfce.

These service types add “metapackages” such as `gnome` or `plasma` to the system profile, but most of them also set up other useful services that mere packages can’t do.

For example, they may elevate privileges on a limited number of special-purpose system interfaces and programs. This allows backlight adjustment helpers, power management utilities, screen lockers, and other integrated functionality to work as expected.

**gnome-desktop-service-type** [Variable]

This is the type of the service that adds the GNOME (<https://www.gnome.org>) desktop environment. Its value is a `gnome-desktop-configuration` object (see below).

This service adds the `gnome` package to the system profile, and extends polkit with the actions from `gnome-settings-daemon`.

**gnome-desktop-configuration** [Data Type]

Configuration record for the GNOME desktop environment. Available `gnome-desktop-configuration` fields are:

**core-services** (type: list-of-packages)

A list of packages that the GNOME Shell and applications may rely on.

**shell** (type: list-of-packages)

A list of packages that constitute the GNOME Shell, without applications.

**utilities** (type: list-of-packages)

A list of packages that serve as applications to use on top of the GNOME Shell.

**gnome** (type: maybe-package)

This field used to be the only configuration point and specified a GNOME meta-package to install system-wide. Since the meta-package itself provides neither sources nor the actual packages and is only used to propagate them, this field is deprecated.

**extra-packages** (type: list-of-packages)

A list of GNOME-adjacent packages to also include. This field is intended for users to add their own packages to their GNOME experience. Note, that it already includes some packages that are considered essential by some (most?) GNOME users.

**udev-ignorelist** (default: ()) (type: list-of-strings)

A list of regular expressions denoting udev rules or hardware file names provided by any package that should not be installed. By default, every udev rule and hardware file specified by any package referenced in the other fields are installed.

**polkit-ignorelist** (default: ()) (type: list-of-strings)

A list of regular expressions denoting polkit rules provided by any package that should not be installed. By default, every polkit rule added by any package referenced in the other fields are installed.

**keyring** (type: gnome-keyring-configuration-or-#f)

A `<gnome-keyring-configuration>` record used to better integrate the GNOME keyring with the system. Refer to the documentation of the `gnome-keyring-service-type` for more information. If you'd rather avoid integrating the GNOME keyring, you can set this to `#f`.

**plasma-desktop-service-type** [Variable]

This is the type of the service that adds the Plasma (<https://kde.org/plasma-desktop/>) desktop environment. Its value is a `plasma-desktop-configuration` object (see below).

This service adds the `plasma` package to the system profile.

**plasma-desktop-configuration** [Data Type]

Configuration record for the Plasma desktop environment.

**plasma** (default: `plasma`)

The Plasma package to use.

**xfce-desktop-service-type** [Variable]

This is the type of a service to run the Xfce (<https://xfce.org/>) desktop environment. Its value is an **xfce-desktop-configuration** object (see below).

This service adds the **xfce** package to the system profile, and extends polkit with the ability for **thunar** to manipulate the file system as root from within a user session, after the user has authenticated with the administrator's password.

Note that **xfce4-panel** and its plugin packages should be installed in the same profile to ensure compatibility. When using this service, you should add extra plugins (**xfce4-whiskermenu-plugin**, **xfce4-weather-plugin**, etc.) to the **packages** field of your **operating-system**.

**xfce-desktop-configuration** [Data Type]

Configuration record for the Xfce desktop environment.

**xfce** (default: **xfce**)

The Xfce package to use.

**mate-desktop-service-type** [Variable]

This is the type of the service that runs the MATE desktop environment (<https://mate-desktop.org/>). Its value is a **mate-desktop-configuration** object (see below).

This service adds the **mate** package to the system profile, and extends polkit with the actions from **mate-settings-daemon**.

**mate-desktop-configuration** [Data Type]

Configuration record for the MATE desktop environment.

**mate** (default: **mate**)

The MATE package to use.

**lxqt-desktop-service-type** [Variable]

This is the type of the service that runs the LXQt desktop environment ([https://lxqt-project.org](https://lxqt-project.org/)). Its value is a **lxqt-desktop-configuration** object (see below).

This service adds the **lxqt** package to the system profile.

**lxqt-desktop-configuration** [Data Type]

Configuration record for the LXQt desktop environment.

**lxqt** (default: **lxqt**)

The LXQT package to use.

**sugar-desktop-service-type** [Variable]

This is the type of the service that runs the Sugar desktop environment (<https://www.sugarlabs.org>). Its value is a **sugar-desktop-configuration** object (see below).

This service adds the **sugar** package to the system profile, as well as any selected Sugar activities. By default it only includes a minimal set of activities.

**sugar-desktop-configuration** [Data Type]

Configuration record for the Sugar desktop environment.

**sugar** (default: **sugar**)

The Sugar package to use.

**gobject-introspection** (default: **gobject-introspection**)

The **gobject-introspection** package to use. This package is used to access libraries installed as dependencies of Sugar activities.

**activities** (default: (list **sugar-help-activity**))

A list of Sugar activities to install.

The following example configures the Sugar desktop environment with a number of useful activities:

```
(use-modules (gnu))
(use-package-modules sugar)
(use-service-modules desktop)
(operating-system
  ...
  (services (cons* (service sugar-desktop-service-type
                           (sugar-desktop-configuration
                            (activities (list sugar-browse-activity
  sugar-help-activity
  sugar-jukebox-activity
  sugar-typing-turtle-activity))))
                %desktop-services))
  ...)
```

**enlightenment-desktop-service-type** [Variable]

Return a service that adds the **enlightenment** package to the system profile, and extends dbus with actions from **efl**.

**enlightenment-desktop-service-configuration** [Data Type]

**enlightenment** (default: **enlightenment**)

The enlightenment package to use.

Because the GNOME, Xfce and MATE desktop services pull in so many packages, the default **%desktop-services** variable doesn't include any of them by default. To add GNOME, Xfce or MATE, just cons them onto **%desktop-services** in the **services** field of your **operating-system**:

```
(use-modules (gnu))
(use-service-modules desktop)
(operating-system
  ...
  ;; cons* adds items to the list given as its last argument.
  (services (cons* (service gnome-desktop-service-type)
                   (service xfce-desktop-service-type)
                   %desktop-services))
  ...)
```



...)

These desktop environments will then be available as options in the graphical login window.

The actual service definitions included in `%desktop-services` and provided by `(gnu services dbus)` and `(gnu services desktop)` are described below.

**dbus-root-service-type** [Variable]

Type for a service that runs the D-Bus “system bus”.<sup>7</sup>

The value for this service type is a `<dbus-configuration>` record.

**dbus-configuration** [Data Type]

Data type representing the configuration for `dbus-root-service-type`.

**dbus** (default: `dbus`) (type: file-like)

Package object for dbus.

**services** (default: `'()`) (type: list)

List of packages that provide an `etc/dbus-1/system.d` directory containing additional D-Bus configuration and policy files. For example, to allow `avahi-daemon` to use the system bus, `services` must be equal to `(list avahi)`.

**verbose?** (default: `#f`) (type: boolean)

When `#t`, D-Bus is launched with environment variable `‘DBUS_VERBOSE’` set to `‘1’`. A verbose-enabled D-Bus package such as `dbus-verbose` should be provided to `dbus` in this scenario. The verbose output is logged to `/var/log/dbus-daemon.log`.

## Elogind

Elogind (<https://github.com/elogind/elogind>) is a login and seat management daemon that also handles most system-level power events for a computer, for example suspending the system when a lid is closed, or shutting it down when the power button is pressed.

It also provides a D-Bus interface that can be used to know which users are logged in, know what kind of sessions they have open, suspend the system, inhibit system suspend, reboot the system, and other tasks.

**elogind-service-type** [Variable]

Type of the service that runs `elogind`, a login and seat management daemon. The value for this service is a `<elogind-configuration>` object.

**elogind-configuration** [Data Type]

Available `elogind-configuration` fields are:

**elogind** (default: `elogind`) (type: file-like)

The `elogind` package to use.

<sup>7</sup> D-Bus (<https://dbus.freedesktop.org/>) is an inter-process communication facility. Its system bus is used to allow system services to communicate and to be notified of system-wide events.

- system-sleep-hook-files** (default: ()) (type: list-of-file-likes)  
A list of executables (file-like objects) that will be installed into the `/etc/elogind/system-sleep` hook directory. See ‘Hook directories’ in the ‘`loginctl(1)`’ man page for more information.
- system-shutdown-hook-files** (default: ()) (type: list-of-file-likes)  
A list of executables (file-like objects) that will be installed into the `/etc/elogind/system-shutdown/` hook directory.
- allow-power-off-interrupts?** (default: **#f**) (type: maybe-boolean)  
Whether the executables in `elogind`’s hook directories (see above) can cause a power-off action to be cancelled (interrupted) by printing an appropriate error message to stdout.
- allow-suspend-interrupts?** (default: **#f**) (type: maybe-boolean)  
Likewise as the `allow-power-off-interrupts?` option, but for the suspend action.
- broadcast-power-off-interrupts?** (default: **#f**) (type: maybe-boolean)  
Whether an interrupt of a power-off action is broadcasted.
- broadcast-suspend-interrupts?** (default: **#f**) (type: maybe-boolean)  
Whether an interrupt of a suspend action is broadcasted.
- kill-user-processes?** (default: **#f**) (type: maybe-boolean)  
Whether the processes of a user should be killed when the user logs out.
- kill-only-users** (type: maybe-list-of-user-names)  
Usernames whose processes should be killed, regardless the value of `kill-user-processes?`.
- kill-exclude-users** (default: ("root")) (type: maybe-list-of-user-names)  
Usernames whose processes should *not* be killed, regardless the value of `kill-user-processes?`.
- inhibit-delay-max-seconds** (default: 5) (type: maybe-non-negative-integer)  
The maximum time a system shutdown or sleep request is delayed due to an inhibitor lock of type `delay` being active before the inhibitor is ignored and the operation executes anyway.
- handle-power-key** (default: **poweroff**) (type: maybe-action)  
The action done when the power key is pressed. The compiled default is `'poweroff'`.
- handle-suspend-key** (default: **suspend**) (type: maybe-action)  
The action done when the suspend key is pressed. The
- handle-hibernate-key** (default: **hibernate**) (type: maybe-action)  
The action done when the hibernate key is pressed.
- handle-lid-switch** (default: **suspend**) (type: maybe-action)  
The action done when the lid is closed.
- handle-lid-switch-docked** (default: **ignore**) (type: maybe-action)  
The action done when the lid is closed and the device docked.

**handle-lid-switch-external-power** (default: **suspend**) (type: maybe-action)  
 The action done when the lid is closed and the device is externally powered.

**power-key-ignore-inhibited?** (default: **#f**) (type: maybe-boolean)  
 Whether to ignore high-level inhibitor locks (shutdown, reboot, sleep or idle) when the power key is pressed.

**suspend-key-ignore-inhibited?** (default: **#f**) (type: maybe-boolean)  
 Whether to ignore high-level inhibitor locks (shutdown, reboot, sleep or idle) when the suspend key is pressed.

**hibernate-key-ignore-inhibited?** (default: **#f**) (type: maybe-boolean)  
 Whether to ignore high-level inhibitor locks (shutdown, reboot, sleep or idle) when the hibernate key is pressed.

**lid-switch-ignore-inhibited?** (default: **#f**) (type: maybe-boolean)  
 Whether to ignore high-level inhibitor locks (shutdown, reboot, sleep or idle) when the lid is closed.

**holdoff-timeout-seconds** (default: 30) (type: maybe-non-negative-integer)  
 Specifies the number of seconds after system startup or system resume during which elogind will hold off on reacting to lid events.

**idle-action** (default: **ignore**) (type: maybe-action)  
 Action to take when the system is idle.

**idle-action-seconds** (type: maybe-non-negative-integer)  
 The delay after which the action configured in **idle-action** is taken after the system is idle.

**runtime-directory-size-percent** (type: maybe-percent)  
 Set the size limit, in percent, on the `XDG_RUNTIME_DIR` runtime directory for each user who logs in. This specifies the per-user size limit relative to the amount of physical RAM (read access memory). This value takes precedence over that specified via **runtime-directory-size**.

**runtime-directory-size** (type: maybe-non-negative-integer)  
 Set the size limit, in bytes, on the `XDG_RUNTIME_DIR` runtime directory for each user who logs in.

**remove-ipc?** (default: **#t**) (type: maybe-boolean)  
 Whether IPC (inter-process communication) objects belonging to the user shall be removed when the user fully logs out.

**suspend-state** (default: (**mem standby freeze**)) (type: maybe-list-of-suspend-states)  
 The suspend state values to be write to `/sys/power/state` by elogind when suspending the system. They will be tried in turn, until one is written without error.

**suspend-mode** (type: maybe-list-of-suspend-modes)  
 The suspend mode values to write to `/sys/power/mem_sleep` by elogind when suspending the system.

**suspend-estimation-seconds** (default: 3600) (type: maybe-non-negative-integer)  
 Cause the RTC alarm to wake the system after the specified time span to measure the system battery capacity level and estimate the battery discharging rate, which is used for estimating the time span until the system battery charge level goes down to 5%. This option is only used by elogind when using the '**suspend-then-hibernate**' action.

**hibernate-mode** (default: (platform shutdown)) (type: maybe-list-of-hibernation-modes)  
 The hibernation mode values to write to `/sys/power/disk` by elogind when hibernating the system.

**hibernate-delay-seconds** (type: maybe-non-negative-integer)  
 The amount of time the system spends in suspend mode before the system is automatically put into hibernate mode.

**hibernate-state** (type: maybe-list-of-strings)  
 Deprecated option.

**hybrid-sleep-state** (type: maybe-list-of-strings)  
 Deprecated option.

**hybrid-sleep-mode** (type: maybe-list-of-strings)  
 Deprecated option.

**accountsservice-service-type** [Variable]  
 Type for the service that runs AccountsService, a system service that can list available accounts, change their passwords, and so on. AccountsService integrates with PolicyKit to enable unprivileged users to acquire the capability to modify their system configuration. See AccountsService (<https://www.freedesktop.org/wiki/Software/AccountsService/>) for more information.  
 The value for this service is a file-like object, by default it is set to `accountsservice` (the package object for AccountsService).

**polkit-service-type** [Variable]  
 Type for the service that runs the Polkit privilege management service (<https://www.freedesktop.org/wiki/Software/polkit/>), which allows system administrators to grant access to privileged operations in a structured way. By querying the Polkit service, a privileged system component can know when it should grant additional capabilities to ordinary users. For example, an ordinary user can be granted the capability to suspend the system if the user is logged in locally.  
 The value for this service is a `<polkit-configuration>` object.

**polkit-wheel-service** [Variable]  
 Service that adds the `wheel` group as admins to the Polkit service. This makes it so that users in the `wheel` group are queried for their own passwords when performing administrative actions instead of `root`'s, similar to the behaviour used by `sudo`.

**upower-service-type** [Variable]  
 Service that runs `upowerd` (<https://upower.freedesktop.org/>), a system-wide monitor for power consumption and battery levels, with the given configuration settings.

It implements the `org.freedesktop.UPower` D-Bus interface, and is notably used by GNOME.

**upower-configuration** [Data Type]

Data type representation the configuration for UPower.

**upower** (default: *upower*)

Package to use for **upower**.

**watts-up-pro?** (default: *#f*)

Enable the Watts Up Pro device.

**poll-batteries?** (default: *#t*)

Enable polling the kernel for battery level changes.

**ignore-lid?** (default: *#f*)

Ignore the lid state, this can be useful if it's incorrect on a device.

**use-percentage-for-policy?** (default: *#t*)

Whether to use a policy based on battery percentage rather than on estimated time left. A policy based on battery percentage is usually more reliable.

**percentage-low** (default: 20)

When **use-percentage-for-policy?** is *#t*, this sets the percentage at which the battery is considered low.

**percentage-critical** (default: 5)

When **use-percentage-for-policy?** is *#t*, this sets the percentage at which the battery is considered critical.

**percentage-action** (default: 2)

When **use-percentage-for-policy?** is *#t*, this sets the percentage at which action will be taken.

**time-low** (default: 1200)

When **use-time-for-policy?** is *#f*, this sets the time remaining in seconds at which the battery is considered low.

**time-critical** (default: 300)

When **use-time-for-policy?** is *#f*, this sets the time remaining in seconds at which the battery is considered critical.

**time-action** (default: 120)

When **use-time-for-policy?** is *#f*, this sets the time remaining in seconds at which action will be taken.

**critical-power-action** (default: *'hybrid-sleep'*)

The action taken when **percentage-action** or **time-action** is reached (depending on the configuration of **use-percentage-for-policy?**).

Possible values are:

- *'power-off'*
- *'hibernate'*
- *'hybrid-sleep'*.

**udisks-service-type** [Variable]

Type for the service that runs UDisks (<https://udisks.freedesktop.org/docs/latest/>), a *disk management* daemon that provides user interfaces with notifications and ways to mount/unmount disks. Programs that talk to UDisks include the `udisksctl` command, part of UDisks, and GNOME Disks. Note that Udisks relies on the `mount` command, so it will only be able to use the file-system utilities installed in the system profile. For example if you want to be able to mount NTFS file-systems in read and write fashion, you'll need to have `ntfs-3g` installed system-wide.

The value for this service is a `<udisks-configuration>` object.

**udisks-configuration** [Data Type]

Data type representing the configuration for `udisks-service-type`.

`udisks` (default: `udisks`) (type: file-like)

Package object for UDisks.

**gvfs-service-type** [Variable]

Type for the service that provides virtual file systems for GIO applications, which enables support for `trash://`, `ftp://`, `sftp://` and many other location schemas in file managers like Nautilus (GNOME Files) and Thunar.

The value for this service is a `<gvfs-configuration>` object.

**gvfs-configuration** [Data Type]

Data type representing the configuration for `gvfs-service-type`.

`gvfs` (default: `gvfs`) (type: file-like)

Package object for GVfs.

**colord-service-type** [Variable]

This is the type of the service that runs `colord`, a system service with a D-Bus interface to manage the color profiles of input and output devices such as screens and scanners. It is notably used by the GNOME Color Manager graphical tool. See the `colord` web site (<https://www.freedesktop.org/software/colord/>) for more information.

**sane-service-type** [Variable]

This service provides access to scanners *via* SANE (<http://www.sane-project.org>) by installing the necessary udev rules and pluggable backends. It is included in `%desktop-services` (see Section 11.10.9 [Desktop Services], page 377) and relies by default on `sane-backends` package (see below) for hardware support.

**sane-configuration** [Data Type]

Data type representing the configuration for SANE.

`sane` (default: `sane`)

Package containing SANE library.

`backends` (default: `(sane-backends)`)

List of packages with pluggable SANE backends:

- `sane-backends`: The default backend collection which supports many recent scanners,

- **sane-airscan**: A backend that enables network scanners supporting eSCL (Apple) or WSD, (Microsoft) protocols
- **hplip**: A backend containing drivers for older Hewlett-Packard scanners,
- **utsushi**: A backend containing drivers for older Epson devices.

In order to use additional backends on a system which relies on `%desktop-services`, you may use `modify-services` (see Section 11.19.3 [Service Reference], page 712) as illustrated below:

```
(use-modules (gnu))
(use-service-modules
  ...
  desktop)
(use-package-modules
  ...
  scanner)

(define %my-desktop-services
  ;; List of desktop services that supports a broader range of scanners.
  (modify-services %desktop-services
    (sane-service-type _ =>
      (sane-configuration
        (backends (list sane-backends sane-airscan))))))

(operating-system
  ...
  (services %my-desktop-services))
```

**geoclue-application** *name* [*#:allowed?* *#t*] [*#:system?* *#f*] [*#:users* '()] [Procedure]

Return a configuration allowing an application to access GeoClue location data. *name* is the Desktop ID of the application, without the `.desktop` part. If *allowed?* is true, the application will have access to location information by default. The boolean *system?* value indicates whether an application is a system component or not. Finally *users* is a list of UIDs of all users for which this application is allowed location info access. An empty *users* list means that all users are allowed.

**%standard-geoclue-applications** [Variable]

The standard list of well-known GeoClue application configurations, granting authority to the GNOME date-and-time utility to ask for the current location in order to set the time zone, and allowing the IceCat and Epiphany web browsers to request location information. IceCat and Epiphany both query the user before allowing a web page to know the user's location.

**geoclue-service-type** [Variable]

Type for the service that runs the GeoClue (<https://wiki.freedesktop.org/www/Software/GeoClue/>) location service. This service provides a D-Bus interface to al-

low applications to request access to a user's physical location, and optionally to add information to online location databases.

The value for this service is a `<geoclue-configuration>` object.

#### `iio-sensor-proxy-service-type` [Variable]

Type for the service that runs IIO sensor proxy (<https://gitlab.freedesktop.org/hadess/iio-sensor-proxy>), a daemon that proxies low-level sensor data from Industrial I/O (IIO) devices into a D-Bus interface. It enables user-space applications, such as desktop environments like GNOME and KDE, and services like GeoClue, to access data from various sensors commonly found in devices like ultrabooks and 2-in-1s. This allows for:

- Automatic screen rotation based on accelerometer data.
- Ambient brightness adjustment using light sensor data.
- Compass/direction information for location services (via GeoClue).

The value for this service is a `<iio-sensor-proxy-configuration>` object.

#### `iio-sensor-proxy-configuration` [Data Type]

Data type representing the configuration for `iio-sensor-proxy-service-type`.

`iio-sensor-proxy` (default: `iio-sensor-proxy`) (type: file-like)  
Package object of `iio-sensor-proxy`.

#### `bluetooth-service-type` [Variable]

This is the type for the Linux Bluetooth Protocol Stack (<https://bluez.org/>) (BlueZ) system, which generates the `/etc/bluetooth/main.conf` configuration file. The value for this type is a `bluetooth-configuration` record as in this example:

```
(service bluetooth-service-type)
```

See below for details about `bluetooth-configuration`.

#### `bluetooth-configuration` [Data Type]

Data type representing the configuration for `bluetooth-service`.

`bluez` (default: `bluez`)  
`bluez` package to use.

`name` (default: "BlueZ")  
Default adapter name.

`class` (default: `#x000000`)  
Default device class. Only the major and minor device class bits are considered.

`discoverable-timeout` (default: 180)  
How long to stay in discoverable mode before going back to non-discoverable. The value is in seconds.

`always-pairable?` (default: `#f`)  
Always allow pairing even if there are no agents registered.



`pairable-timeout` (default: 0)

How long to stay in pairable mode before going back to non-discoverable. The value is in seconds.

`device-id` (default: `#f`)

Use vendor id source (assigner), vendor, product and version information for DID profile support. The values are separated by ":" and *assigner*, *VID*, *PID* and *version*.

Possible values are:

- `#f` to disable it,
- `"assigner:1234:5678:abcd"`, where *assigner* is either `usb` (default) or `bluetooth`.

`reverse-service-discovery?` (default: `#t`)

Do reverse service discovery for previously unknown devices that connect to us. For BR/EDR this option is really only needed for qualification since the BITE tester doesn't like us doing reverse SDP for some test cases, for LE this disables the GATT client functionally so it can be used in system which can only operate as peripheral.

`name-resolving?` (default: `#t`)

Enable name resolving after inquiry. Set it to `#f` if you don't need remote devices name and want shorter discovery cycle.

`debug-keys?` (default: `#f`)

Enable runtime persistency of debug link keys. Default is false which makes debug link keys valid only for the duration of the connection that they were created for.

`controller-mode` (default: `'dual`)

Restricts all controllers to the specified transport. `'dual` means both BR/EDR and LE are enabled (if supported by the hardware).

Possible values are:

- `'dual`
- `'bredr`
- `'le`

`multi-profile` (default: `'off`)

Enables Multi Profile Specification support. This allows to specify if system supports only Multiple Profiles Single Device (MPSD) configuration or both Multiple Profiles Single Device (MPSD) and Multiple Profiles Multiple Devices (MPMD) configurations.

Possible values are:

- `'off`
- `'single`
- `'multiple`

**fast-connectable?** (default: #f)

Permanently enables the Fast Connectable setting for adapters that support it. When enabled other devices can connect faster to us, however the tradeoff is increased power consumptions. This feature will fully work only on kernel version 4.1 and newer.

**privacy** (default: 'off')

Default privacy settings.

- 'off: Disable local privacy
- 'network/on: A device will only accept advertising packets from peer devices that contain private addresses. It may not be compatible with some legacy devices since it requires the use of RPA(s) all the time
- 'device: A device in device privacy mode is only concerned about the privacy of the device and will accept advertising packets from peer devices that contain their Identity Address as well as ones that contain a private address, even if the peer device has distributed its IRK in the past

and additionally, if *controller-mode* is set to 'dual:

- 'limited-network: Apply Limited Discoverable Mode to advertising, which follows the same policy as to BR/EDR that publishes the identity address when discoverable, and Network Privacy Mode for scanning
- 'limited-device: Apply Limited Discoverable Mode to advertising, which follows the same policy as to BR/EDR that publishes the identity address when discoverable, and Device Privacy Mode for scanning.

**just-works-repairing** (default: 'never')

Specify the policy to the JUST-WORKS repairing initiated by peer.

Possible values:

- 'never
- 'confirm
- 'always

**temporary-timeout** (default: 30)

How long to keep temporary devices around. The value is in seconds. 0 disables the timer completely.

**refresh-discovery?** (default: #t)

Enables the device to issue an SDP request to update known services when profile is connected.

**experimental** (default: #f)

Enables experimental features and interfaces, alternatively a list of UUIDs can be given.

Possible values:

- #t

- `#f`
- `(list (uuid <uuid-1>) (uuid <uuid-2>) ...)`.

List of possible UUIDs:

- `d4992530-b9ec-469f-ab01-6c481c47da1c`: BlueZ Experimental Debug,
- `671b10b5-42c0-4696-9227-eb28d1b049d6`: BlueZ Experimental Simultaneous Central and Peripheral,
- `15c0a148-c273-11ea-b3de-0242ac130004`: BlueZ Experimental LL privacy,
- `330859bc-7506-492d-9370-9a6f0614037f`: BlueZ Experimental Bluetooth Quality Report,
- `a6695ace-ee7f-4fb9-881a-5fac66c629af`: BlueZ Experimental Offload Codecs.

`remote-name-request-retry-delay` (default: 300)

The duration to avoid retrying to resolve a peer's name, if the previous try failed.

`page-scan-type` (default: `#f`)

BR/EDR Page scan activity type.

`page-scan-interval` (default: `#f`)

BR/EDR Page scan activity interval.

`page-scan-window` (default: `#f`)

BR/EDR Page scan activity window.

`inquiry-scan-type` (default: `#f`)

BR/EDR Inquiry scan activity type.

`inquiry-scan-interval` (default: `#f`)

BR/EDR Inquiry scan activity interval.

`inquiry-scan-window` (default: `#f`)

BR/EDR Inquiry scan activity window.

`link-supervision-timeout` (default: `#f`)

BR/EDR Link supervision timeout.

`page-timeout` (default: `#f`)

BR/EDR Page timeout.

`min-sniff-interval` (default: `#f`)

BR/EDR minimum sniff interval.

`max-sniff-interval` (default: `#f`)

BR/EDR maximum sniff interval.

`min-advertisement-interval` (default: `#f`)

LE minimum advertisement interval (used for legacy advertisement only).

`max-advertisement-interval` (default: `#f`)

LE maximum advertisement interval (used for legacy advertisement only).

**multi-advertisement-rotation-interval** (default: #f)  
LE multiple advertisement rotation interval.

**scan-interval-auto-connect** (default: #f)  
LE scanning interval used for passive scanning supporting auto connect.

**scan-window-auto-connect** (default: #f)  
LE scanning window used for passive scanning supporting auto connect.

**scan-interval-suspend** (default: #f)  
LE scanning interval used for active scanning supporting wake from suspend.

**scan-window-suspend** (default: #f)  
LE scanning window used for active scanning supporting wake from suspend.

**scan-interval-discovery** (default: #f)  
LE scanning interval used for active scanning supporting discovery.

**scan-window-discovery** (default: #f)  
LE scanning window used for active scanning supporting discovery.

**scan-interval-adv-monitor** (default: #f)  
LE scanning interval used for passive scanning supporting the advertisement monitor APIs.

**scan-window-adv-monitor** (default: #f)  
LE scanning window used for passive scanning supporting the advertisement monitor APIs.

**scan-interval-connect** (default: #f)  
LE scanning interval used for connection establishment.

**scan-window-connect** (default: #f)  
LE scanning window used for connection establishment.

**min-connection-interval** (default: #f)  
LE default minimum connection interval. This value is superseded by any specific value provided via the Load Connection Parameters interface.

**max-connection-interval** (default: #f)  
LE default maximum connection interval. This value is superseded by any specific value provided via the Load Connection Parameters interface.

**connection-latency** (default: #f)  
LE default connection latency. This value is superseded by any specific value provided via the Load Connection Parameters interface.

**connection-supervision-timeout** (default: #f)  
LE default connection supervision timeout. This value is superseded by any specific value provided via the Load Connection Parameters interface.

**autoconnect-timeout** (default: #f)  
LE default autoconnect timeout. This value is superseded by any specific value provided via the Load Connection Parameters interface.

- adv-mon-allowlist-scan-duration** (default: 300)  
Allowlist scan duration during interleaving scan. Only used when scanning for ADV monitors. The units are msec.
- adv-mon-no-filter-scan-duration** (default: 500)  
No filter scan duration during interleaving scan. Only used when scanning for ADV monitors. The units are msec.
- enable-adv-mon-interleave-scan?** (default: #t)  
Enable/Disable Advertisement Monitor interleave scan for power saving.
- cache** (default: 'always)  
GATT attribute cache.  
Possible values are:
- 'always: Always cache attributes even for devices not paired, this is recommended as it is best for interoperability, with more consistent reconnection times and enables proper tracking of notifications for all devices
  - 'yes: Only cache attributes of paired devices
  - 'no: Never cache attributes.
- key-size** (default: 0)  
Minimum required Encryption Key Size for accessing secured characteristics.  
Possible values are:
- 0: Don't care
  - $7 \leq N \leq 16$
- exchange-mtu** (default: 517)  
Exchange MTU size. Possible values are:
- $23 \leq N \leq 517$
- att-channels** (default: 3)  
Number of ATT channels. Possible values are:
- 1: Disables EATT
  - $2 \leq N \leq 5$
- session-mode** (default: 'basic)  
AVDTP L2CAP signalling channel mode.  
Possible values are:
- 'basic: Use L2CAP basic mode
  - 'ertm: Use L2CAP enhanced retransmission mode.
- stream-mode** (default: 'basic)  
AVDTP L2CAP transport channel mode.  
Possible values are:
- 'basic: Use L2CAP basic mode
  - 'streaming: Use L2CAP streaming mode.

**reconnect-uuids** (default: '()')

The ReconnectUUIDs defines the set of remote services that should try to be reconnected to in case of a link loss (link supervision timeout). The policy plugin should contain a sane set of values by default, but this list can be overridden here. By setting the list to empty the reconnection feature gets disabled.

Possible values:

- '()
- (list (uuid <uuid-1>) (uuid <uuid-2>) ...).

**reconnect-attempts** (default: 7)

Defines the number of attempts to reconnect after a link lost. Setting the value to 0 disables reconnecting feature.

**reconnect-intervals** (default: '(1 2 4 8 16 32 64)')

Defines a list of intervals in seconds to use in between attempts. If the number of attempts defined in *reconnect-attempts* is bigger than the list of intervals the last interval is repeated until the last attempt.

**auto-enable?** (default: #f)

Defines option to enable all controllers when they are found. This includes adapters present on start as well as adapters that are plugged in later on.

**resume-delay** (default: 2)

Audio devices that were disconnected due to suspend will be reconnected on resume. *resume-delay* determines the delay between when the controller resumes from suspend and a connection attempt is made. A longer delay is better for better co-existence with Wi-Fi. The value is in seconds.

**rss-sampling-period** (default: #xFF)

Default RSSI Sampling Period. This is used when a client registers an advertisement monitor and leaves the RSSISamplingPeriod unset.

Possible values are:

- #x0: Report all advertisements
- N = #xXX: Report advertisements every N x 100 msec (range: #x01 to #xFE)
- #xFF: Report only one advertisement per device during monitoring period.

**gnome-keyring-service-type**

[Variable]

This is the type of the service that adds the GNOME Keyring (<https://wiki.gnome.org/Projects/GnomeKeyring>). Its value is a **gnome-keyring-configuration** object (see below). Note that there is no need to use this service when using **gnome-desktop-service-type**, which includes it.

This service adds the **gnome-keyring** package to the system profile and extends PAM with entries using **pam\_gnome\_keyring.so**, unlocking a user's login keyring when they log in or setting its password with **passwd**.

**gnome-keyring-configuration** [Data Type]

Configuration record for the GNOME Keyring service.

**keyring** (default: **gnome-keyring**)

The GNOME keyring package to use.

**pam-services**

A list of (**service** . **kind**) pairs denoting PAM services to extend, where **service** is the name of an existing service to extend and **kind** is one of **login** or **passwd**.

If **login** is given, it adds an optional **pam\_gnome\_keyring.so** to the auth block without arguments and to the session block with **auto\_start**. If **passwd** is given, it adds an optional **pam\_gnome\_keyring.so** to the password block without arguments.

By default, this field contains “gdm-password” with the value **login** and “passwd” is with the value **passwd**.

**kwallet-service-type** [Variable]

This is the type of the service that adds the KWallet keyring (<https://invent.kde.org/plasma/kwallet-pam>). Its value is a **kwallet-configuration** object (see below). Note that, contrary to **gnome-desktop-service-type**, which includes the respective keyring for that service, **gnome-keyring-service-type**, **plasma-desktop-service-type** does not include **kwallet-service-type**.

This service adds the **kwallet-pam** package to the system profile and extends PAM with entries using **pam\_kwallet5.so**. It can unlock the user’s login keyring or set their password with **passwd** at the time they log in.

**kwallet-configuration** [Data Type]

Configuration record for the KWallet Keyring service.

**keyring** (default: **kwallet-pam**)

The KWallet keyring package to use.

**pam-services**

A list of (**service** . **kind**) pairs denoting PAM services to extend, where **service** is the name of an existing service to extend and **kind** is one of **login** or **passwd** symbols.

If **login** is given, it adds an optional **pam\_kwallet5.so** to the auth block without arguments and to the session block with **auto\_start**. If **passwd** is given, it adds an optional **pam\_kwallet5.so** to the password block without arguments.

By default, it takes for value an alist associating the “sddm” key with the ‘login value, and the “passwd” key with the ‘passwd value.

**seatd-service-type** [Variable]

**seatd** (<https://sr.ht/~kennylevinsen/seatd/>) is a minimal seat management daemon.

Seat management takes care of mediating access to shared devices (graphics, input), without requiring the applications needing access to be root.

(append

```
(list
;; make sure seatd is running
(service seatd-service-type))

;; normally one would want %base-services
%base-services)
```

**seatd** operates over a UNIX domain socket, with **libseat** providing the client side of the protocol. Applications that acquire access to the shared resources via **seatd** (e.g. **sway**) need to be able to talk to this socket. This can be achieved by adding the user they run under to the group owning **seatd**'s socket (usually "seat"), like so:

```
(user-account
  (name "alice")
  (group "users")
  (supplementary-groups '("wheel"      ; allow use of sudo, etc.
                          "seat"      ; seat management
                          "audio"     ; sound card
                          "video"     ; video devices such as webcams
                          "cdrom")) ; the good ol' CD-ROM
  (comment "Bob's sister"))
```

Depending on your setup, you will have to not only add regular users, but also system users to this group. For instance, some greetd greeters require graphics and therefore also need to negotiate with **seatd**.

### **seatd-configuration**

[Data Type]

Configuration record for the **seatd** daemon service.

```
seatd (default: seatd)
    The seatd package to use.

group (default: "seat")
    Group to own the seatd socket.

socket (default: "/run/seatd.sock")
    Where to create the seatd socket.

logfile (default: "/var/log/seatd.log")
    Log file to write to.

loglevel (default: "error")
    Log level to output logs. Possible values: "silent", "error", "info"
    and "debug".
```

#### **11.10.10 Sound Services**

The (**gnu services sound**) module provides a service to configure the Advanced Linux Sound Architecture (ALSA) system, which makes PulseAudio the preferred ALSA output driver.



**alsa-service-type** [Variable]

This is the type for the Advanced Linux Sound Architecture (<https://alsa-project.org/>) (ALSA) system, which generates the `/etc/asound.conf` configuration file. The value for this type is a `alsa-configuration` record as in this example:

```
(service alsa-service-type)
```

See below for details about `alsa-configuration`.

**alsa-configuration** [Data Type]

Data type representing the configuration for `alsa-service`.

**alsa-plugins** (default: *alsa-plugins*)  
 alsa-plugins package to use.

**pulseaudio?** (default: *#t*)  
 Whether ALSA applications should transparently be made to use the PulseAudio (<https://www.pulseaudio.org/>) sound server.  
 Using PulseAudio allows you to run several sound-producing applications at the same time and to individual control them *via* `pavucontrol`, among other things.

**extra-options** (default: *""*)  
 String to append to the `/etc/asound.conf` file.

Individual users who want to override the system configuration of ALSA can do it with the `~/.asoundrc` file:

```
# In guix, we have to specify the absolute path for plugins.
pcm_type.jack {
  lib "/home/alice/.guix-profile/lib/alsa-lib/libasound_module_pcm_jack.so"
}

# Routing ALSA to jack:
# <http://jackaudio.org/faq/routing_alsa.html>.
pcm.rawjack {
  type jack
  playback_ports {
    0 system:playback_1
    1 system:playback_2
  }

  capture_ports {
    0 system:capture_1
    1 system:capture_2
  }
}

pcm.!default {
  type plug
```

```

    slave {
        pcm "rawjack"
    }
}

```

See <https://www.alsa-project.org/main/index.php/Asoundrc> for the details.

**pulseaudio-service-type** [Variable]

This is the type for the PulseAudio (<https://www.pulseaudio.org/>) sound server. It exists to allow system overrides of the default settings via **pulseaudio-configuration**, see below.

**Warning:** This service overrides per-user configuration files. If you want PulseAudio to honor configuration files in `~/.config/pulse`, you have to unset the environment variables `PULSE_CONFIG` and `PULSE_CLIENTCONFIG` in your `~/.bash_profile`.

**Warning:** This service on its own does not ensure, that the **pulseaudio** package exists on your machine. It merely adds configuration files for it, as detailed below. In the (admittedly unlikely) case, that you find yourself without a **pulseaudio** package, consider enabling it through the **alsa-service-type** above.

**pulseaudio-configuration** [Data Type]

Data type representing the configuration for **pulseaudio-service**.

**client-conf** (default: '()')

List of settings to set in `client.conf`. Accepts a list of strings or symbol-value pairs. A string will be inserted as-is with a newline added. A pair will be formatted as “key = value”, again with a newline added.

**daemon-conf** (default: '((flat-volumes . no)))

List of settings to set in `daemon.conf`, formatted just like *client-conf*.

**script-file** (default: (file-append pulseaudio "/etc/pulse/default.pa"))

Script file to use as `default.pa`. In case the **extra-script-files** field below is used, an `.include` directive pointing to `/etc/pulse/default.pa.d` is appended to the provided script.

**extra-script-files** (default: '()')

A list of file-like objects defining extra PulseAudio scripts to run at the initialization of the **pulseaudio** daemon, after the main **script-file**. The scripts are deployed to the `/etc/pulse/default.pa.d` directory; they should have the `.pa` file name extension. For a reference of the available commands, refer to `man pulse-cli-syntax`.

**system-script-file** (default: (file-append pulseaudio  
"/etc/pulse/system.pa"))

Script file to use as `system.pa`.

The example below sets the default PulseAudio card profile, the default sink and the default source to use for a old SoundBlaster Audigy sound card:

```
(pulseaudio-configuration
```

```
(extra-script-files
  (list (plain-file "audigy.pa"
    (string-append "\
set-card-profile alsa_card.pci-0000_01_01.0 \
  output:analog-surround-40+input:analog-mono
set-default-source alsa_input.pci-0000_01_01.0.analog-mono
set-default-sink alsa_output.pci-0000_01_01.0.analog-surround-40\n")))))■
```

Note that `pulseaudio-service-type` is part of `%desktop-services`; if your operating system declaration was derived from one of the desktop templates, you'll want to adjust the above example to modify the existing `pulseaudio-service-type` via `modify-services` (see Section 11.19.3 [Service Reference], page 712), instead of defining a new one.

**ladspa-service-type** [Variable]

This service sets the `LADSPA_PATH` variable, so that programs, which respect it, e.g. PulseAudio, can load LADSPA plugins.

The following example will setup the service to enable modules from the `swh-plugins` package:

```
(service ladspa-service-type
  (ladspa-configuration (plugins (list swh-plugins))))
```

See <http://plugin.org.uk/ladspa-swh/docs/ladspa-swh.html> for the details.

## Speaker Safety Daemon System Service

Speaker Safety Daemon (<https://github.com/AsahiLinux/speakersafetyd>) is a user-space daemon that implements an analogue of the Texas Instruments Smart Amp speaker protection model. It can be used to protect the speakers on Apple Silicon devices.

**speakersafetyd-service-type** [Variable]

This is the type for the `speakersafetyd` system service, whose value is a `speakersafetyd-configuration` record.

```
(service speakersafetyd-service-type)
```

See below for details about `speakersafetyd-configuration`.

**speakersafetyd-configuration** [Data Type]

Available `speakersafetyd-configuration` fields are:

**blackbox-directory** (default: `"/var/lib/speakersafetyd/blackbox"`) (type: string)

The directory to which blackbox files are written when the speakers are getting too hot. The blackbox files contain audio and debug information which the developers of `speakersafetyd` might ask for when reporting bugs.

**configuration-directory** (type: file-like)

The base directory as a G-expression (see Section 8.12 [G-Expressions], page 165) that contains the configuration files of the speaker models.

**group** (default: `"speakersafetyd"`) (type: string)

The group to run the Speaker Safety Daemon as.

**log-file** (default: `"/var/log/speakersafetyd.log"`) (type: string)  
 The file name to the Speaker Safety Daemon log file.

**maximum-gain-reduction** (default: 7) (type: integer)  
 Maximum gain reduction before panicking, useful for debugging.

**speakersafetyd** (default: `speakersafetyd`) (type: file-like)  
 The Speaker Safety Daemon package to use.

**user** (default: `"speakersafetyd"`) (type: string)  
 The user to run the Speaker Safety Daemon as.

### 11.10.11 File Search Services

The services in this section populate *file databases* that let you search for files on your machine. These services are provided by the (`gnu services admin`) module.

The first one, `file-database-service-type`, periodically runs the venerable `updatedb` command (see Section “Invoking `updatedb`” in *GNU Findutils*). That command populates a database of file names that you can then search with the `locate` command (see Section “Invoking `locate`” in *GNU Findutils*), as in this example:

```
locate important-notes.txt
```

You can enable this service with its default settings by adding this snippet to your operating system services:

```
(service file-database-service-type)
```

This updates the database once a week, excluding files from `/gnu/store`—these are more usefully handled by `guix locate` (see Section 5.5 [Invoking `guix locate`], page 52). You can of course provide a custom configuration, as described below.

**file-database-service-type** [Variable]

This is the type of the file database service, which runs `updatedb` periodically. Its associated value must be a `file-database-configuration` record, as described below.

**file-database-configuration** [Data Type]

Record type for the `file-database-service-type` configuration, with the following fields:

**package** (default: `findutils`)  
 The package from which the `updatedb` command is taken. Examples of such packages are GNU Findutils and Plocate.

**schedule** (default: `%default-file-database-update-schedule`)  
 This is the schedule of database updates, expressed as a string in traditional cron syntax or as a gexp evaluating to a Shepherd calendar event (see Section “Timers” in *The GNU Shepherd Manual*).

**excluded-directories** (default `%default-file-database-excluded-directories`)  
 List of regular expressions of directories to ignore when building the file database. By default, this includes `/tmp` and `/gnu/store`; the latter should instead be indexed by `guix locate` (see Section 5.5 [Invoking `guix`]).

locate], page 52). This list is passed to the `--prunepaths` option of `updatedb` (see Section “Invoking updatedb” in *GNU Findutils*).

The second service, `package-database-service-type`, builds the database used by `guix locate`, which lets you search for packages that contain a given file (see Section 5.5 [Invoking `guix locate`], page 52). The service periodically updates a system-wide database, which will be readily available to anyone running `guix locate` on the system. To use this service with its default settings, add this snippet to your service list:

```
(service package-database-service-type
```

This will run `guix locate --update` once a week.

`package-database-service-type` [Variable]

This is the service type for periodic `guix locate` updates (see Section 5.5 [Invoking `guix locate`], page 52). Its value must be a `package-database-configuration` record, as shown below.

`package-database-configuration` [Data Type]

Data type to configure periodic package database updates. It has the following fields:

`package` (default: `guix`)

The Guix package to use.

`schedule` (default: `%default-package-database-update-schedule`)

This is the schedule of the periodic `guix locate --update` job, expressed as a string in traditional cron syntax or as a gexp evaluating to a Shepherd calendar event (see Section “Timers” in *The GNU Shepherd Manual*).

`method` (default: `'store`)

Indexing method for `guix locate`. The default value, `'store`, yields a more complete database but is relatively expensive in terms of CPU and input/output.

`channels` (default: `#~%default-channels`)

G-exp denoting the channels to use when updating the database (see Chapter 6 [Channels], page 68).

### 11.10.12 Database Services

The `(gnu services databases)` module provides the following services.

#### PostgreSQL

`postgresql-service-type` [Variable]

The service type for the PostgreSQL database server. Its value should be a valid `postgresql-configuration` object, documented below. The following example describes a PostgreSQL service with the default configuration.

```
(service postgresql-service-type
  (postgresql-configuration
    (postgresql postgresql)))
```

If the services fails to start, it may be due to an incompatible cluster already present in *data-directory*. Adjust it (or, if you don't need the cluster anymore, delete *data-directory*), then restart the service.

Peer authentication is used by default and the `postgres` user account has no shell (unless `allow-login?` is `#t`), which prevents the direct execution of `psql` commands as this user. To use `psql`, you can temporarily log in as `postgres` using a shell, create a PostgreSQL superuser with the same name as one of the system users and then create the associated database.

```
sudo -u postgres -s /bin/sh
createuser --interactive
createdb $MY_USER_LOGIN      # Replace appropriately.
```

`postgresql-configuration` [Data Type]

Data type representing the configuration for the `postgresql-service-type`.

`postgresql`

PostgreSQL package to use for the service.

`port` (default: 5432)

Port on which PostgreSQL should listen.

`locale` (default: "en\_US.utf8")

Locale to use as the default when creating the database cluster.

`config-file` (default: (`postgresql-config-file`))

The configuration file to use when running PostgreSQL. The default behaviour uses the `postgresql-config-file` record with the default values for the fields.

`log-directory` (default: "/var/log/postgresql")

The directory where `pg_ctl` output will be written in a file named "`pg_ctl.log`". This file can be useful to debug PostgreSQL configuration errors for instance.

`data-directory` (default: "/var/lib/postgresql/data")

Directory in which to store the data.

`extension-packages` (default: '()')

Additional extensions are loaded from packages listed in *extension-packages*. Extensions are available at runtime. For instance, to create a geographic database using the `postgis` extension, a user can configure the `postgresql-service` as in this example:

```
(use-package-modules databases geo)
```

```
(operating-system
```

```
...
```

```
;; postgresql is required to run `psql' but postgis is not required f
```

```
;; proper operation.
```

```
(packages (cons* postgresql %base-packages))
```

```
(services
```

```
(cons*
  (service postgresql-service-type
    (postgresql-configuration
      (postgresql postgresql)
      (extension-packages (list postgis))))
  %base-services)))
```

Then the extension becomes visible and you can initialise an empty geographic database in this way:

```
psql -U postgres
> create database postgistest;
> \connect postgistest;
> create extension postgis;
> create extension postgis_topology;
```

There is no need to add this field for contrib extensions such as hstore or dblink as they are already loadable by postgresql. This field is only required to add extensions provided by other packages.

**create-account?** (default: **#t**)

Whether or not the **postgres** user and group should be created.

**allow-login?** (default: **#f**)

Whether or not to allow login into the created account.

**home-directory** (default: **"/var/empty"**)

The home directory of the user. It is strongly advised to change this if you set **allow-login?** to **#t**.

**uid** (default: **#f**)

Explicitly specify the UID of the **postgres** daemon account. You normally do not need to specify this, in which case a free UID will be automatically assigned.

One situation where this option might be useful is if the *data-directory* is located on a mounted network share.

**gid** (default: **#f**)

Explicitly specify the GID of the **postgres** group.

**postgresql-config-file**

[Data Type]

Data type representing the PostgreSQL configuration file. As shown in the following example, this can be used to customize the configuration of PostgreSQL. Note that you can use any G-expression or filename in place of this record, if you already have a configuration file you'd like to use for example.

```
(service postgresql-service-type
  (postgresql-configuration
    (config-file
      (postgresql-config-file
        (log-destination "stderr")
        (hba-file
          (plain-file "pg_hba.conf"))
```

```

"
local all all trust
host all all 127.0.0.1/32 md5
host all all ::1/128 md5"))
    (extra-config
      '(("session_preload_libraries"      "auto_explain")
        ("random_page_cost"              2)
        ("auto_explain.log_min_duration" "100 ms")
        ("work_mem"                      "500 MB")
        ("logging_collector"              #t)
        ("log_directory"                  "/var/log/postgresql"))))))

```

**log-destination** (default: "syslog")

The logging method to use for PostgreSQL. Multiple values are accepted, separated by commas.

**hba-file** (default: %default-postgres-hba)

Filename or G-expression for the host-based authentication configuration.

**ident-file** (default: %default-postgres-ident)

Filename or G-expression for the user name mapping configuration.

**socket-directory** (default: "/var/run/postgresql")

Specifies the directory of the Unix-domain socket(s) on which PostgreSQL is to listen for connections from client applications. If set to "" PostgreSQL does not listen on any Unix-domain sockets, in which case only TCP/IP sockets can be used to connect to the server.

By default, the **#false** value means the PostgreSQL default value will be used, which is currently **‘/tmp’**.

**extra-config** (default: '()')

List of additional keys and values to include in the PostgreSQL config file. Each entry in the list should be a list where the first element is the key, and the remaining elements are the values.

The values can be numbers, booleans or strings and will be mapped to PostgreSQL parameters types **Boolean**, **String**, **Numeric**, **Numeric with Unit** and **Enumerated** described here (<https://www.postgresql.org/docs/current/config-setting.html>).

**postgresql-role-service-type** [Variable]

This service allows to create PostgreSQL roles and databases after PostgreSQL service start. Here is an example of its use.

```

(service postgresql-role-service-type
  (postgresql-role-configuration
    (roles
      (list (postgresql-role
              (name "test")
              (create-database? #t)))))

```

This service can be extended with extra roles, as in this example:

```

(service-extension postgresql-role-service-type

```



```
(const (list
  (postgresql-role
    (name "alice")
    (create-database? #t))))
```

**postgresql-role** [Data Type]

PostgreSQL manages database access permissions using the concept of roles. A role can be thought of as either a database user, or a group of database users, depending on how the role is set up. Roles can own database objects (for example, tables) and can assign privileges on those objects to other roles to control who has access to which objects.

**name** The role name.

**permissions** (default: '(createdb login))  
The role permissions list. Supported permissions are `bypassrls`, `createdb`, `createrole`, `login`, `replication` and `superuser`.

**create-database?** (default: #f)  
whether to create a database with the same name as the role.

**password-file** (default: #f)  
A string representing the path of a file that contains the password to be set for the role.

**encoding** (default: "UTF8")  
The character set to use for storing text in the database.

**collation** (default: "en\_US.utf8")  
The string sort order locale setting.

**ctype** (default: "en\_US.utf8")  
The character classification locale setting.

**template** (default: "template1")  
The default template to copy the new database from when creating it. Use "template0" for a pristine database with no system-local modifications.

**postgresql-role-configuration** [Data Type]

Data type representing the configuration of *postgresql-role-service-type*.

**host** (default: "/var/run/postgresql")  
The PostgreSQL host to connect to.

**log** (default: "/var/log/postgresql\_roles.log")  
File name of the log file.

**shepherd-requirement** (default: '(user-processes postgres))  
The Shepherd services dependencies to use. Add extra dependencies to `%default-postgresql-role-shepherd-requirement` to extend its value.

**roles** (default: '())  
The initial PostgreSQL roles to create.

## MariaDB/MySQL

**mysql-service-type** [Variable]

This is the service type for a MySQL or MariaDB database server. Its value is a **mysql-configuration** object that specifies which package to use, as well as various settings for the **mysqld** daemon.

**mysql-configuration** [Data Type]

Data type representing the configuration of *mysql-service-type*.

**mysql** (default: *mariadb*)

Package object of the MySQL database server, can be either *mariadb* or *mysql*.

For MySQL, a temporary root password will be displayed at activation time. For MariaDB, the root password is empty.

**bind-address** (default: "127.0.0.1")

The IP on which to listen for network connections. Use "0.0.0.0" to bind to all available network interfaces.

**port** (default: 3306)

TCP port on which the database server listens for incoming connections.

**socket** (default: "/run/mysqld/mysqld.sock")

Socket file to use for local (non-network) connections.

**extra-content** (default: "")

Additional settings for the *my.cnf* configuration file.

**extra-environment** (default: #~'())

List of environment variables passed to the **mysqld** process.

**auto-upgrade?** (default: #t)

Whether to automatically run **mysql\_upgrade** after starting the service. This is necessary to upgrade the *system schema* after “major” updates (such as switching from MariaDB 10.4 to 10.5), but can be disabled if you would rather do that manually.

## Memcached

**memcached-service-type** [Variable]

This is the service type for the Memcached (<https://memcached.org/>) service, which provides a distributed in memory cache. The value for the service type is a **memcached-configuration** object.

(service memcached-service-type)

**memcached-configuration** [Data Type]

Data type representing the configuration of memcached.

**memcached** (default: **memcached**)

The Memcached package to use.

**interfaces** (default: `'("0.0.0.0")`)  
 Network interfaces on which to listen.

**tcp-port** (default: `11211`)  
 Port on which to accept connections.

**udp-port** (default: `11211`)  
 Port on which to accept UDP connections on, a value of 0 will disable listening on a UDP socket.

**additional-options** (default: `'()`)  
 Additional command line options to pass to `memcached`.

## Redis

Redis (<https://redis.io/>) is an in-memory data store used by millions of developers as a cache, vector database, document database, streaming engine, and message broker. Redis has built-in replication and different levels of on-disk persistence. It supports complex data types (for example, strings, hashes, lists, sets, sorted sets, and JSON), with atomic operations defined on those data types.

**redis-service-type** [Variable]  
 Type of the service that runs `redis`, an in-memory data store. The value for this service is a `<redis-configuration>` object.

**redis-configuration** [Data Type]  
 Available `redis-configuration` fields are:

**redis** (default: `redis`) (type: package)  
 The Redis package to use.

**bind** (default: `"127.0.0.1"`) (type: string)  
 Network interface on which to listen.

**port** (default: `6379`) (type: number)  
 Port on which to accept connections on, a value of 0 will disable listening on a TCP socket.

**working-directory** (default: `"/var/lib/redis"`) (type: string)  
 Directory in which to store the database and related files.

**config-file** (type: maybe-string)  
 Default location for config file.

### 11.10.13 Mail Services

The `(gnu services mail)` module provides Guix service definitions for email services: IMAP, POP3, and LMTP servers, as well as mail transport agents (MTAs). Lots of acronyms! These services are detailed in the subsections below.

## Dovecot Service

**dovecot-service-type** [Variable]  
 Type for the service that runs the Dovecot IMAP/POP3/LMTP mail server, whose value is a `<dovecot-configuration>` object.

By default, Dovecot does not need much configuration; the default configuration object created by `(dovecot-configuration)` will suffice if your mail is delivered to `~/Maildir`. A self-signed certificate will be generated for TLS-protected connections, though Dovecot will also listen on cleartext ports by default. There are a number of options, though, which mail administrators might need to change, and as is the case with other services, Guix allows the system administrator to specify these parameters via a uniform Scheme interface.

For example, to specify that mail is located at `maildir~/.mail`, one would instantiate the Dovecot service like this:

```
(service dovecot-service-type
  (dovecot-configuration
    (mail-location "maildir:~/.mail"))))
```

The available configuration parameters follow. Each parameter definition is preceded by its type; for example, `'string-list foo'` indicates that the `foo` parameter should be specified as a list of strings. There is also a way to specify the configuration as a string, if you have an old `dovecot.conf` file that you want to port over from some other system; see the end for more details.

Available `dovecot-configuration` fields are:

**package dovecot** [dovecot-configuration parameter]  
The dovecot package.

**package-list extensions** [dovecot-configuration parameter]  
A list of additional Dovecot plugin packages to make available at runtime. During service activation, the `lib/dovecot` directory from each specified package is combined with Dovecot's core modules into a unified module directory.

For example, to enable Sieve filtering:

```
(extensions (list dovecot-pigeonhole))
```

Each package in the list must provide its modules at `lib/dovecot`, as this is where `make-dovecot-moduledir` expects to find its extensions. The service combines these directories to create a unified module structure.

The default value is an empty list, providing only core Dovecot functionality.

**comma-separated-string-list listen** [dovecot-configuration parameter]  
A list of IPs or hosts where to listen for connections. `'*'` listens on all IPv4 interfaces, `'::'` listens on all IPv6 interfaces. If you want to specify non-default ports or anything more complex, customize the address and port fields of the `'inet-listener'` of the specific services you are interested in.

**protocol-configuration-list protocols** [dovecot-configuration parameter]  
List of protocols we want to serve. Available protocols include `'imap'`, `'pop3'`, and `'lmtp'`.

Available `protocol-configuration` fields are:

**string name** [protocol-configuration parameter]  
The name of the protocol.

**string auth-socket-path** [protocol-configuration parameter]  
 UNIX socket path to the master authentication server to find users.  
 This is used by `imap` (for shared users) and `lda`. It defaults to  
`"/var/run/dovecot/auth-userdb"`.

**boolean imap-metadata?** [protocol-configuration parameter]  
 Whether to enable the IMAP METADATA extension as defined in RFC 5464  
 (<https://tools.ietf.org/html/rfc5464>), which provides a means for  
 clients to set and retrieve per-mailbox, per-user metadata and annotations  
 over IMAP.  
 If this is `#t`, you must also specify a dictionary *via* the `mail-attribute-dict`  
 setting.  
 Defaults to `#f`.

**space-separated-string-list** [protocol-configuration parameter]  
**managesieve-notify-capabilities**  
 Which NOTIFY capabilities to report to clients that first connect to the Man-  
 ageSieve service, before authentication. These may differ from the capabilities  
 offered to authenticated users. If this field is left empty, report what the Sieve  
 interpreter supports by default.  
 Defaults to `'()'`.

**space-separated-string-list** [protocol-configuration parameter]  
**managesieve-sieve-capability**  
 Which SIEVE capabilities to report to clients that first connect to the Mana-  
 geSieve service, before authentication. These may differ from the capabilities  
 offered to authenticated users. If this field is left empty, report what the Sieve  
 interpreter supports by default.  
 Defaults to `'()'`.

**space-separated-string-list** [protocol-configuration parameter]  
**mail-plugins**  
 Space separated list of plugins to load.

**non-negative-integer** [protocol-configuration parameter]  
**mail-max-userip-connections**  
 Maximum number of IMAP connections allowed for a user from each IP address.  
 NOTE: The username is compared case-sensitively. Defaults to `'10'`.

**service-configuration-list services** [dovecot-configuration parameter]  
 List of services to enable. Available services include `'imap'`, `'imap-login'`, `'pop3'`,  
`'pop3-login'`, `'auth'`, and `'lmtp'`.  
 Available `service-configuration` fields are:

**string kind** [service-configuration parameter]  
 The service kind. Valid values include `director`, `imap-login`, `pop3-login`,  
`lmtp`, `imap`, `pop3`, `auth`, `auth-worker`, `dict`, `tcpwrap`, `quota-warning`, or any-  
 thing else.

**listener-configuration-list** [service-configuration parameter]  
**listeners**

Listeners for the service. A listener is either a **unix-listener-configuration**, a **fifo-listener-configuration**, or an **inet-listener-configuration**. Defaults to **'()'**.

Available **unix-listener-configuration** fields are:

**string path** [unix-listener-configuration parameter]  
 Path to the file, relative to **base-dir** field. This is also used as the section name.

**string mode** [unix-listener-configuration parameter]  
 The access mode for the socket. Defaults to **"0600"**.

**string user** [unix-listener-configuration parameter]  
 The user to own the socket. Defaults to **""**.

**string group** [unix-listener-configuration parameter]  
 The group to own the socket. Defaults to **""**.

Available **fifo-listener-configuration** fields are:

**string path** [fifo-listener-configuration parameter]  
 Path to the file, relative to **base-dir** field. This is also used as the section name.

**string mode** [fifo-listener-configuration parameter]  
 The access mode for the socket. Defaults to **"0600"**.

**string user** [fifo-listener-configuration parameter]  
 The user to own the socket. Defaults to **""**.

**string group** [fifo-listener-configuration parameter]  
 The group to own the socket. Defaults to **""**.

Available **inet-listener-configuration** fields are:

**string protocol** [inet-listener-configuration parameter]  
 The protocol to listen for.

**string address** [inet-listener-configuration parameter]  
 The address on which to listen, or empty for all addresses. Defaults to **""**.

**non-negative-integer port** [inet-listener-configuration parameter]  
 The port on which to listen.

**boolean ssl?** [inet-listener-configuration parameter]  
 Whether to use SSL for this service; **'yes'**, **'no'**, or **'required'**. Defaults to **'#t'**.

**non-negative-integer** [service-configuration parameter]  
**client-limit**

Maximum number of simultaneous client connections per process. Once this number of connections is received, the next incoming connection will prompt Dovecot to spawn another process. If set to 0, `default-client-limit` is used instead.

Defaults to '0'.

**non-negative-integer** [service-configuration parameter]  
**service-count**

Number of connections to handle before starting a new process. Typically the only useful values are 0 (unlimited) or 1. 1 is more secure, but 0 is faster. <doc/wiki/LoginProcess.txt>. Defaults to '1'.

**non-negative-integer** [service-configuration parameter]  
**process-limit**

Maximum number of processes that can exist for this service. If set to 0, `default-process-limit` is used instead.

Defaults to '0'.

**non-negative-integer** [service-configuration parameter]  
**process-min-avail**

Number of processes to always keep waiting for more connections. Defaults to '0'.

**non-negative-integer** **vsz-limit** [service-configuration parameter]  
 If you set 'service-count 0', you probably need to grow this. Defaults to '256000000'.

**dict-configuration** **dict** [dovecot-configuration parameter]  
 Dict configuration, as created by the `dict-configuration` constructor.

Available `dict-configuration` fields are:

**free-form-fields** **entries** [dict-configuration parameter]  
 A list of key-value pairs that this dict should hold. Defaults to '()'.

**passdb-configuration-list** **passdbs** [dovecot-configuration parameter]  
 A list of passdb configurations, each one created by the `passdb-configuration` constructor.

Available `passdb-configuration` fields are:

**string** **driver** [passdb-configuration parameter]  
 The driver that the passdb should use. Valid values include 'pam', 'passwd', 'shadow', 'bsdauth', and 'static'. Defaults to "pam".

**space-separated-string-list** **args** [passdb-configuration parameter]  
 Space separated list of arguments to the passdb driver. Defaults to "".

**userdb-configuration-list userdbs** [dovecot-configuration parameter]  
 List of userdb configurations, each one created by the **userdb-configuration** constructor.

Available **userdb-configuration** fields are:

**string driver** [userdb-configuration parameter]  
 The driver that the userdb should use. Valid values include ‘passwd’ and ‘static’. Defaults to “passwd”.

**space-separated-string-list args** [userdb-configuration parameter]  
 Space separated list of arguments to the userdb driver. Defaults to “”.

**free-form-args override-fields** [userdb-configuration parameter]  
 Override fields from passwd. Defaults to ‘()’.

**plugin-configuration** [dovecot-configuration parameter]  
**plugin-configuration**  
 Plug-in configuration, created by the **plugin-configuration** constructor.

**list-of-namespace-configuration namespaces** [dovecot-configuration parameter]  
**namespaces**  
 List of namespaces. Each item in the list is created by the **namespace-configuration** constructor.

Available **namespace-configuration** fields are:

**string name** [namespace-configuration parameter]  
 Name for this namespace.

**string type** [namespace-configuration parameter]  
 Namespace type: ‘private’, ‘shared’ or ‘public’. Defaults to “private”.

**string separator** [namespace-configuration parameter]  
 Hierarchy separator to use. You should use the same separator for all namespaces or some clients get confused. ‘/’ is usually a good one. The default however depends on the underlying mail storage format. Defaults to “”.

**string prefix** [namespace-configuration parameter]  
 Prefix required to access this namespace. This needs to be different for all namespaces. For example ‘Public/’. Defaults to “”.

**string location** [namespace-configuration parameter]  
 Physical location of the mailbox. This is in the same format as **mail\_location**, which is also the default for it. Defaults to “”.

**boolean inbox?** [namespace-configuration parameter]  
 There can be only one INBOX, and this setting defines which namespace has it. Defaults to ‘#f’.



**boolean hidden?** [namespace-configuration parameter]

If namespace is hidden, it's not advertised to clients via NAMESPACE extension. You'll most likely also want to set 'list? #f'. This is mostly useful when converting from another server with different namespaces which you want to deprecate but still keep working. For example you can create hidden namespaces with prefixes '~/'mail/', '~%u/mail/' and 'mail/'. Defaults to '#f'.

**boolean list?** [namespace-configuration parameter]

Show the mailboxes under this namespace with the LIST command. This makes the namespace visible for clients that do not support the NAMESPACE extension. The special **children** value lists child mailboxes, but hides the namespace prefix. Defaults to '#t'.

**boolean subscriptions?** [namespace-configuration parameter]

Namespace handles its own subscriptions. If set to #f, the parent namespace handles them. The empty prefix should always have this as #t). Defaults to '#t'.

**mailbox-configuration-list** [namespace-configuration parameter]  
**mailboxes**

List of predefined mailboxes in this namespace. Defaults to '()'.

Available mailbox-configuration fields are:

**string name** [mailbox-configuration parameter]  
Name for this mailbox.

**string auto** [mailbox-configuration parameter]  
'create' will automatically create this mailbox. 'subscribe' will both create and subscribe to the mailbox. Defaults to "no".

**space-separated-string-list** [mailbox-configuration parameter]  
**special-use**

List of IMAP SPECIAL-USE attributes as specified by RFC 6154. Valid values are \All, \Archive, \Drafts, \Flagged, \Junk, \Sent, and \Trash. Defaults to '()'.

**file-name base-dir** [dovecot-configuration parameter]  
Base directory where to store runtime data. Defaults to "/var/run/dovecot/".

**string login-greeting** [dovecot-configuration parameter]  
Greeting message for clients. Defaults to "Dovecot ready.".

**space-separated-string-list** [dovecot-configuration parameter]  
**login-trusted-networks**

List of trusted network ranges. Connections from these IPs are allowed to override their IP addresses and ports (for logging and for authentication checks). 'disable-plaintext-auth' is also ignored for these networks. Typically you would specify your IMAP proxy servers here. Defaults to '()'.

- space-separated-string-list** [dovecot-configuration parameter]  
**login-access-sockets**  
List of login access check sockets (e.g. tcpwrap). Defaults to ‘‘()’.
- boolean verbose-proctitle?** [dovecot-configuration parameter]  
Show more verbose process titles (in ps). Currently shows user name and IP address. Useful for seeing who is actually using the IMAP processes (e.g. shared mailboxes or if the same uid is used for multiple accounts). Defaults to ‘#f’.
- boolean shutdown-clients?** [dovecot-configuration parameter]  
Should all processes be killed when Dovecot master process shuts down. Setting this to #f means that Dovecot can be upgraded without forcing existing client connections to close (although that could also be a problem if the upgrade is e.g. due to a security fix). Defaults to ‘#t’.
- non-negative-integer** [dovecot-configuration parameter]  
**doveadm-worker-count**  
If non-zero, run mail commands via this many connections to doveadm server, instead of running them directly in the same process. Defaults to ‘0’.
- string doveadm-socket-path** [dovecot-configuration parameter]  
UNIX socket or host:port used for connecting to doveadm server. Defaults to “doveadm-server”.
- space-separated-string-list** [dovecot-configuration parameter]  
**import-environment**  
List of environment variables that are preserved on Dovecot startup and passed down to all of its child processes. You can also give key=value pairs to always set specific settings.
- boolean disable-plaintext-auth?** [dovecot-configuration parameter]  
Disable LOGIN command and all other plaintext authentications unless SSL/TLS is used (LOGINDISABLED capability). Note that if the remote IP matches the local IP (i.e. you’re connecting from the same computer), the connection is considered secure and plaintext authentication is allowed. See also the ‘ssl=required’ setting. Defaults to ‘#t’.
- non-negative-integer auth-cache-size** [dovecot-configuration parameter]  
Authentication cache size (e.g. ‘#e10e6’). 0 means it’s disabled. Note that bsdauth, PAM and vpopmail require ‘cache-key’ to be set for caching to be used. Defaults to ‘0’.
- string auth-cache-ttl** [dovecot-configuration parameter]  
Time to live for cached data. After TTL expires the cached record is no longer used, \*except\* if the main database lookup returns internal failure. We also try to handle password changes automatically: If user’s previous authentication was successful, but this one wasn’t, the cache isn’t used. For now this works only with plaintext authentication. Defaults to “1 hour”.

**string auth-cache-negative-ttl** [dovecot-configuration parameter]  
 TTL for negative hits (user not found, password mismatch). 0 disables caching them completely. Defaults to `"1 hour"`.

**space-separated-string-list auth-realms** [dovecot-configuration parameter]  
 List of realms for SASL authentication mechanisms that need them. You can leave it empty if you don't want to support multiple realms. Many clients simply use the first one listed here, so keep the default realm first. Defaults to `'()'`.

**string auth-default-realm** [dovecot-configuration parameter]  
 Default realm/domain to use if none was specified. This is used for both SASL realms and appending `@domain` to username in plaintext logins. Defaults to `""`.

**string auth-username-chars** [dovecot-configuration parameter]  
 List of allowed characters in username. If the user-given username contains a character not listed in here, the login automatically fails. This is just an extra check to make sure user can't exploit any potential quote escaping vulnerabilities with SQL/LDAP databases. If you want to allow all characters, set this value to empty. Defaults to `"abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0123456789.-_@"`. ■

**string auth-username-translation** [dovecot-configuration parameter]  
 Username character translations before it's looked up from databases. The value contains series of from -> to characters. For example `'#@/'` means that `#` and `/` characters are translated to `@`. Defaults to `""`.

**string auth-username-format** [dovecot-configuration parameter]  
 Username formatting before it's looked up from databases. You can use the standard variables here, e.g. `%Lu` would lowercase the username, `%n` would drop away the domain if it was given, or `'%n-AT-%d'` would change the `@` into `-AT-`. This translation is done after `'auth-username-translation'` changes. Defaults to `"%Lu"`.

**string auth-master-user-separator** [dovecot-configuration parameter]  
 If you want to allow master users to log in by specifying the master username within the normal username string (i.e. not using SASL mechanism's support for it), you can specify the separator character here. The format is then `<username><separator><master username>`. UW-IMAP uses `*` as the separator, so that could be a good choice. Defaults to `""`.

**string auth-anonymous-username** [dovecot-configuration parameter]  
 Username to use for users logging in with ANONYMOUS SASL mechanism. Defaults to `"anonymous"`.

**non-negative-integer auth-worker-max-count** [dovecot-configuration parameter]  
 Maximum number of dovecot-auth worker processes. They're used to execute blocking passdb and userdb queries (e.g. MySQL and PAM). They're automatically created and destroyed as needed. Defaults to `'30'`.

- string auth-gssapi-hostname** [dovecot-configuration parameter]  
Host name to use in GSSAPI principal names. The default is to use the name returned by `gethostname()`. Use `'$ALL'` (with quotes) to allow all keytab entries. Defaults to `''`.
- string auth-krb5-keytab** [dovecot-configuration parameter]  
Kerberos keytab to use for the GSSAPI mechanism. Will use the system default (usually `/etc/krb5.keytab`) if not specified. You may need to change the auth service to run as root to be able to read this file. Defaults to `''`.
- boolean auth-use-winbind?** [dovecot-configuration parameter]  
Do NTLM and GSS-SPNEGO authentication using Samba's winbind daemon and `'ntlm-auth'` helper. <doc/wiki/Authentication/Mechanisms/Winbind.txt>. Defaults to `#f`.
- file-name auth-winbind-helper-path** [dovecot-configuration parameter]  
Path for Samba's `'ntlm-auth'` helper binary. Defaults to `"/usr/bin/ntlm_auth"`.
- string auth-failure-delay** [dovecot-configuration parameter]  
Time to delay before replying to failed authentications. Defaults to `"2 secs"`.
- boolean auth-ssl-require-client-cert?** [dovecot-configuration parameter]  
Require a valid SSL client certificate or the authentication fails. Defaults to `#f`.
- boolean auth-ssl-username-from-cert?** [dovecot-configuration parameter]  
Take the username from client's SSL certificate, using `X509_NAME_get_text_by_NID()` which returns the subject's DN's CommonName. Defaults to `#f`.
- space-separated-string-list** [dovecot-configuration parameter]  
**auth-mechanisms**  
List of wanted authentication mechanisms. Supported mechanisms are: `'plain'`, `'login'`, `'digest-md5'`, `'cram-md5'`, `'ntlm'`, `'rpa'`, `'apop'`, `'anonymous'`, `'gssapi'`, `'otp'`, `'skey'`, and `'gss-spnego'`. See also the `'disable-plaintext-auth'` setting.
- space-separated-string-list** [dovecot-configuration parameter]  
**director-servers**  
List of IPs or hostnames to all director servers, including ourself. Ports can be specified as `ip:port`. The default port is the same as what director service's `'inet-listener'` is using. Defaults to `'()'`.
- space-separated-string-list** [dovecot-configuration parameter]  
**director-mail-servers**  
List of IPs or hostnames to all backend mail servers. Ranges are allowed too, like `10.0.0.10-10.0.0.30`. Defaults to `'()'`.
- string director-user-expire** [dovecot-configuration parameter]  
How long to redirect users to a specific server after it no longer has any connections. Defaults to `"15 min"`.

- string director-username-hash** [dovecot-configuration parameter]  
How the username is translated before being hashed. Useful values include %Ln if user can log in with or without @domain, %Ld if mailboxes are shared within domain. Defaults to "%Lu".
- string log-path** [dovecot-configuration parameter]  
Log file to use for error messages. 'syslog' logs to syslog, '/dev/stderr' logs to stderr. Defaults to "syslog".
- string info-log-path** [dovecot-configuration parameter]  
Log file to use for informational messages. Defaults to 'log-path'. Defaults to "".
- string debug-log-path** [dovecot-configuration parameter]  
Log file to use for debug messages. Defaults to 'info-log-path'. Defaults to "".
- string syslog-facility** [dovecot-configuration parameter]  
Syslog facility to use if you're logging to syslog. Usually if you don't want to use 'mail', you'll use local0..local7. Also other standard facilities are supported. Defaults to "mail".
- boolean auth-verbose?** [dovecot-configuration parameter]  
Log unsuccessful authentication attempts and the reasons why they failed. Defaults to '#f'.
- string auth-verbose-passwords** [dovecot-configuration parameter]  
In case of password mismatches, log the attempted password. Valid values are no, plain and sha1. sha1 can be useful for detecting brute force password attempts vs. user simply trying the same password over and over again. You can also truncate the value to n chars by appending ":n" (e.g. sha1:6). Defaults to "no".
- boolean auth-debug?** [dovecot-configuration parameter]  
Even more verbose logging for debugging purposes. Shows for example SQL queries. Defaults to '#f'.
- boolean auth-debug-passwords?** [dovecot-configuration parameter]  
In case of password mismatches, log the passwords and used scheme so the problem can be debugged. Enabling this also enables 'auth-debug'. Defaults to '#f'.
- boolean mail-debug?** [dovecot-configuration parameter]  
Enable mail process debugging. This can help you figure out why Dovecot isn't finding your mails. Defaults to '#f'.
- boolean verbose-ssl?** [dovecot-configuration parameter]  
Show protocol level SSL errors. Defaults to '#f'.
- string log-timestamp** [dovecot-configuration parameter]  
Prefix for each line written to log file. % codes are in strftime(3) format. Defaults to "%b %d %H:%M:%S \".

**space-separated-string-list** [dovecot-configuration parameter]  
**login-log-format-elements**

List of elements we want to log. The elements which have a non-empty variable value are joined together to form a comma-separated string.

**string login-log-format** [dovecot-configuration parameter]  
 Login log format. %s contains 'login-log-format-elements' string, %\$ contains the data we want to log. Defaults to "%\$: %s".

**string mail-log-prefix** [dovecot-configuration parameter]  
 Log prefix for mail processes. See doc/wiki/Variables.txt for list of possible variables you can use. Defaults to "\"%s(%u)<{%pid}><{%session}>: \"".

**string deliver-log-format** [dovecot-configuration parameter]  
 Format to use for logging mail deliveries. You can use variables:

%\$          Delivery status message (e.g. 'saved to INBOX')

%m          Message-ID

%s          Subject

%f          From address

%p          Physical size

%w          Virtual size.

Defaults to "msgid=%m: %\$".

**string mail-location** [dovecot-configuration parameter]

Location for users' mailboxes. The default is empty, which means that Dovecot tries to find the mailboxes automatically. This won't work if the user doesn't yet have any mail, so you should explicitly tell Dovecot the full location.

If you're using mbox, giving a path to the INBOX file (e.g. /var/mail/%u) isn't enough. You'll also need to tell Dovecot where the other mailboxes are kept. This is called the *root mail directory*, and it must be the first path given in the 'mail-location' setting.

There are a few special variables you can use, e.g.:

'%u'          username

'%n'          user part in user@domain, same as %u if there's no domain

'%d'          domain part in user@domain, empty if there's no domain

'%h'          home directory

See doc/wiki/Variables.txt for full list. Some examples:

'maildir:~/Maildir'

'mbox:~/mail:INBOX=/var/mail/%u'

'mbox:/var/mail/%d/%1n/%n:INDEX=/var/indexes/%d/%1n/%'

Defaults to "".

- string mail-uid** [dovecot-configuration parameter]  
 System user and group used to access mails. If you use multiple, userdb can override these by returning uid or gid fields. You can use either numbers or names. <doc/wiki/UserIds.txt>. Defaults to `""`.
- string mail-gid** [dovecot-configuration parameter]  
 Defaults to `""`.
- string mail-privileged-group** [dovecot-configuration parameter]  
 Group to enable temporarily for privileged operations. Currently this is used only with INBOX when either its initial creation or dotlocking fails. Typically this is set to `"mail"` to give access to `/var/mail`. Defaults to `""`.
- string mail-access-groups** [dovecot-configuration parameter]  
 Grant access to these supplementary groups for mail processes. Typically these are used to set up access to shared mailboxes. Note that it may be dangerous to set these if users can create symlinks (e.g. if `'mail'` group is set here, `ln -s /var/mail ~/mail/var` could allow a user to delete others' mailboxes, or `ln -s /secret/shared/box ~/mail/mybox` would allow reading it). Defaults to `""`.
- string mail-attribute-dict** [dovecot-configuration parameter]  
 The location of a dictionary used to store IMAP METADATA as defined by RFC 5464 (<https://tools.ietf.org/html/rfc5464>).  
 The IMAP METADATA commands are available only if the "imap" protocol configuration's `imap-metadata?` field is `#t`.  
 Defaults to `""`.
- boolean mail-full-filesystem-access?** [dovecot-configuration parameter]  
 Allow full file system access to clients. There's no access checks other than what the operating system does for the active UID/GID. It works with both maildir and mboxes, allowing you to prefix mailboxes names with e.g. `/path/` or `~user/`. Defaults to `#f`.
- boolean mmap-disable?** [dovecot-configuration parameter]  
 Don't use `mmap()` at all. This is required if you store indexes to shared file systems (NFS or clustered file system). Defaults to `#f`.
- boolean dotlock-use-excl?** [dovecot-configuration parameter]  
 Rely on `'O_EXCL'` to work when creating dotlock files. NFS supports `'O_EXCL'` since version 3, so this should be safe to use nowadays by default. Defaults to `#t`.
- string mail-fsync** [dovecot-configuration parameter]  
 When to use `fsync()` or `fdatasync()` calls:
- optimized** Whenever necessary to avoid losing important data
  - always** Useful with e.g. NFS when `write()`s are delayed
  - never** Never use it (best performance, but crashes can lose data).
- Defaults to `"optimized"`.

**boolean mail-nfs-storage?** [dovecot-configuration parameter]  
Mail storage exists in NFS. Set this to yes to make Dovecot flush NFS caches whenever needed. If you're using only a single mail server this isn't needed. Defaults to '#f'.

**boolean mail-nfs-index?** [dovecot-configuration parameter]  
Mail index files also exist in NFS. Setting this to yes requires 'mmap-disable? #t' and 'fsync-disable? #f'. Defaults to '#f'.

**string lock-method** [dovecot-configuration parameter]  
Locking method for index files. Alternatives are fcntl, flock and dotlock. Dotlocking uses some tricks which may create more disk I/O than other locking methods. NFS users: flock doesn't work, remember to change 'mmap-disable'. Defaults to "fcntl".

**file-name mail-temp-dir** [dovecot-configuration parameter]  
Directory in which LDA/LMTP temporarily stores incoming mails >128 kB. Defaults to "/tmp".

**non-negative-integer first-valid-uid** [dovecot-configuration parameter]  
Valid UID range for users. This is mostly to make sure that users can't log in as daemons or other system users. Note that denying root logins is hardcoded to dovecot binary and can't be done even if 'first-valid-uid' is set to 0. Defaults to '500'.

**non-negative-integer last-valid-uid** [dovecot-configuration parameter]  
Defaults to '0'.

**non-negative-integer first-valid-gid** [dovecot-configuration parameter]  
Valid GID range for users. Users having non-valid GID as primary group ID aren't allowed to log in. If user belongs to supplementary groups with non-valid GIDs, those groups are not set. Defaults to '1'.

**non-negative-integer last-valid-gid** [dovecot-configuration parameter]  
Defaults to '0'.

**non-negative-integer mail-max-keyword-length** [dovecot-configuration parameter]  
Maximum allowed length for mail keyword name. It's only forced when trying to create new keywords. Defaults to '50'.

**colon-separated-file-name-list valid-chroot-dirs** [dovecot-configuration parameter]  
List of directories under which chrooting is allowed for mail processes (i.e. /var/mail will allow chrooting to /var/mail/foo/bar too). This setting doesn't affect 'login-chroot' 'mail-chroot' or auth chroot settings. If this setting is empty, './.' in home dirs are ignored. WARNING: Never add directories here which local users can modify, that may lead to root exploit. Usually this should be done only if you don't allow shell access for users. <doc/wiki/Chrooting.txt>. Defaults to '()'.  
</div>



- string mail-chroot** [dovecot-configuration parameter]  
 Default chroot directory for mail processes. This can be overridden for specific users in user database by giving ‘./.’ in user’s home directory (e.g. ‘/home/./user’ chroots into /home). Note that usually there is no real need to do chrooting, Dovecot doesn’t allow users to access files outside their mail directory anyway. If your home directories are prefixed with the chroot directory, append ‘./.’ to ‘mail-chroot’. <doc/wiki/Chrooting.txt>. Defaults to ‘’.
- file-name auth-socket-path** [dovecot-configuration parameter]  
 UNIX socket path to master authentication server to find users. This is used by imap (for shared users) and lda. Defaults to ‘/var/run/dovecot/auth-userdb’.
- file-name mail-plugin-dir** [dovecot-configuration parameter]  
 Directory where to look up mail plugins. Defaults to ‘/usr/lib/dovecot’.
- space-separated-string-list mail-plugins** [dovecot-configuration parameter]  
 List of plugins to load for all services. Plugins specific to IMAP, LDA, etc. are added to this list in their own .conf files. Defaults to ‘()’.
- non-negative-integer mail-cache-min-mail-count** [dovecot-configuration parameter]  
 The minimum number of mails in a mailbox before updates are done to cache file. This allows optimizing Dovecot’s behavior to do less disk writes at the cost of more disk reads. Defaults to ‘0’.
- string mailbox-idle-check-interval** [dovecot-configuration parameter]  
 When IDLE command is running, mailbox is checked once in a while to see if there are any new mails or other changes. This setting defines the minimum time to wait between those checks. Dovecot can also use dnotify, inotify and kqueue to find out immediately when changes occur. Defaults to ‘30 secs’.
- boolean mail-save-crlf?** [dovecot-configuration parameter]  
 Save mails with CR+LF instead of plain LF. This makes sending those mails take less CPU, especially with sendfile() syscall with Linux and FreeBSD. But it also creates a bit more disk I/O which may just make it slower. Also note that if other software reads the mboxs/maildirs, they may handle the extra CRs wrong and cause problems. Defaults to ‘#f’.
- boolean maildir-stat-dirs?** [dovecot-configuration parameter]  
 By default LIST command returns all entries in maildir beginning with a dot. Enabling this option makes Dovecot return only entries which are directories. This is done by stat()ing each entry, so it causes more disk I/O. (For systems setting struct ‘dirent->d\_type’ this check is free and it’s done always regardless of this setting). Defaults to ‘#f’.
- boolean maildir-copy-with-hardlinks?** [dovecot-configuration parameter]  
 When copying a message, do it with hard links whenever possible. This makes the performance much better, and it’s unlikely to have any side effects. Defaults to ‘#t’.

**boolean maildir-very-dirty-syncs?** [dovecot-configuration parameter]  
 Assume Dovecot is the only MUA accessing Maildir: Scan cur/ directory only when its mtime changes unexpectedly or when we can't find the mail otherwise. Defaults to '#f'.

**space-separated-string-list** [dovecot-configuration parameter]  
**mbox-read-locks**

Which locking methods to use for locking mbox. There are four available:

**dotlock** Create <mailbox>.lock file. This is the oldest and most NFS-safe solution. If you want to use /var/mail/ like directory, the users will need write access to that directory.

**dotlock-try** Same as dotlock, but if it fails because of permissions or because there isn't enough disk space, just skip it.

**fcntl** Use this if possible. Works with NFS too if lockd is used.

**flock** May not exist in all systems. Doesn't work with NFS.

**lockf** May not exist in all systems. Doesn't work with NFS.

You can use multiple locking methods; if you do the order they're declared in is important to avoid deadlocks if other MTAs/MUAs are using multiple locking methods as well. Some operating systems don't allow using some of them simultaneously.

**space-separated-string-list** [dovecot-configuration parameter]  
**mbox-write-locks**

**string mbox-lock-timeout** [dovecot-configuration parameter]  
 Maximum time to wait for lock (all of them) before aborting. Defaults to "5 mins".

**string mbox-dotlock-change-timeout** [dovecot-configuration parameter]  
 If dotlock exists but the mailbox isn't modified in any way, override the lock file after this much time. Defaults to "2 mins".

**boolean mbox-dirty-syncs?** [dovecot-configuration parameter]  
 When mbox changes unexpectedly we have to fully read it to find out what changed. If the mbox is large this can take a long time. Since the change is usually just a newly appended mail, it'd be faster to simply read the new mails. If this setting is enabled, Dovecot does this but still safely fallbacks to re-reading the whole mbox file whenever something in mbox isn't how it's expected to be. The only real downside to this setting is that if some other MUA changes message flags, Dovecot doesn't notice it immediately. Note that a full sync is done with SELECT, EXAMINE, EXPUNGE and CHECK commands. Defaults to '#t'.

**boolean mbox-very-dirty-syncs?** [dovecot-configuration parameter]  
 Like 'mbox-dirty-syncs', but don't do full syncs even with SELECT, EXAMINE, EXPUNGE or CHECK commands. If this is set, 'mbox-dirty-syncs' is ignored. Defaults to '#f'.

- boolean mbox-lazy-writes?** [dovecot-configuration parameter]  
 Delay writing mbox headers until doing a full write sync (EXPUNGE and CHECK commands and when closing the mailbox). This is especially useful for POP3 where clients often delete all mails. The downside is that our changes aren't immediately visible to other MUAs. Defaults to '#t'.
- non-negative-integer** [dovecot-configuration parameter]  
**mbox-min-index-size**  
 If mbox size is smaller than this (e.g. 100k), don't write index files. If an index file already exists it's still read, just not updated. Defaults to '0'.
- non-negative-integer** [dovecot-configuration parameter]  
**mdbox-rotate-size**  
 Maximum mbox file size until it's rotated. Defaults to '10000000'.
- string mdbox-rotate-interval** [dovecot-configuration parameter]  
 Maximum mbox file age until it's rotated. Typically in days. Day begins from midnight, so 1d = today, 2d = yesterday, etc. 0 = check disabled. Defaults to '"1d"'.
- boolean mdbox-preallocate-space?** [dovecot-configuration parameter]  
 When creating new mdbox files, immediately preallocate their size to 'mdbox-rotate-size'. This setting currently works only in Linux with some file systems (ext4, xfs). Defaults to '#f'.
- string mail-attachment-dir** [dovecot-configuration parameter]  
 sdbox and mdbox support saving mail attachments to external files, which also allows single instance storage for them. Other backends don't support this for now.  
 WARNING: This feature hasn't been tested much yet. Use at your own risk.  
 Directory root where to store mail attachments. Disabled, if empty. Defaults to '""'.
- non-negative-integer** [dovecot-configuration parameter]  
**mail-attachment-min-size**  
 Attachments smaller than this aren't saved externally. It's also possible to write a plugin to disable saving specific attachments externally. Defaults to '128000'.
- string mail-attachment-fs** [dovecot-configuration parameter]  
 File system backend to use for saving attachments:
- posix** No SiS done by Dovecot (but this might help FS's own deduplication)
  - sis posix** SiS with immediate byte-by-byte comparison during saving
  - sis-queue posix** SiS with delayed comparison and deduplication.
- Defaults to '"sis posix"'.
- string mail-attachment-hash** [dovecot-configuration parameter]  
 Hash format to use in attachment filenames. You can add any text and variables: %{md4}, %{md5}, %{sha1}, %{sha256}, %{sha512}, %{size}. Variables can be truncated, e.g. %{sha256:80} returns only first 80 bits. Defaults to '"%{sha1}"'.

|                                                                                                                                                                                                                                                                                 |                                   |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------|
| <b>non-negative-integer</b><br><b>default-process-limit</b><br>Defaults to '100'.                                                                                                                                                                                               | [dovecot-configuration parameter] |
| <b>non-negative-integer</b><br><b>default-client-limit</b><br>Defaults to '1000'.                                                                                                                                                                                               | [dovecot-configuration parameter] |
| <b>non-negative-integer</b><br><b>default-vsz-limit</b><br>Default VSZ (virtual memory size) limit for service processes. This is mainly intended to catch and kill processes that leak memory before they eat up everything. Defaults to '256000000'.                          | [dovecot-configuration parameter] |
| <b>string default-login-user</b><br>Login user is internally used by login processes. This is the most untrusted user in Dovecot system. It shouldn't have access to anything at all. Defaults to "dovenull".                                                                   | [dovecot-configuration parameter] |
| <b>string default-internal-user</b><br>Internal user is used by unprivileged processes. It should be separate from login user, so that login processes can't disturb other processes. Defaults to "dovecot".                                                                    | [dovecot-configuration parameter] |
| <b>string ssl?</b><br>SSL/TLS support: yes, no, required. <doc/wiki/SSL.txt>. Defaults to "required".                                                                                                                                                                           | [dovecot-configuration parameter] |
| <b>string ssl-cert</b><br>PEM encoded X.509 SSL/TLS certificate (public key). Defaults to "</etc/dovecot/default.pem".                                                                                                                                                          | [dovecot-configuration parameter] |
| <b>string ssl-key</b><br>PEM encoded SSL/TLS private key. The key is opened before dropping root privileges, so keep the key file unreadable by anyone but root. Defaults to "</etc/dovecot/private/default.pem".                                                               | [dovecot-configuration parameter] |
| <b>string ssl-key-password</b><br>If key file is password protected, give the password here. Alternatively give it when starting dovecot with -p parameter. Since this file is often world-readable, you may want to place this setting instead to a different. Defaults to "". | [dovecot-configuration parameter] |
| <b>string ssl-ca</b><br>PEM encoded trusted certificate authority. Set this only if you intend to use 'ssl-verify-client-cert? #t'. The file should contain the CA certificate(s) followed by the matching CRL(s). (e.g. 'ssl-ca </etc/ssl/certs/ca.pem'). Defaults to "".      | [dovecot-configuration parameter] |
| <b>boolean ssl-require-crl?</b><br>Require that CRL check succeeds for client certificates. Defaults to '#t'.                                                                                                                                                                   | [dovecot-configuration parameter] |
| <b>boolean ssl-verify-client-cert?</b><br>Request client to send a certificate. If you also want to require it, set 'auth-ssl-require-client-cert? #t' in auth section. Defaults to '#f'.                                                                                       | [dovecot-configuration parameter] |

**string ssl-cert-username-field** [dovecot-configuration parameter]  
Which field from certificate to use for username. `commonName` and `x500UniqueIdentifier` are the usual choices. You'll also need to set `'auth-ssl-username-from-cert? #t'`. Defaults to `"commonName"`.

**string ssl-min-protocol** [dovecot-configuration parameter]  
Minimum SSL protocol version to accept. Defaults to `"TLSv1"`.

**string ssl-cipher-list** [dovecot-configuration parameter]  
SSL ciphers to use. Defaults to `"ALL:!kRSA:!SRP:!kDHd:!DSS:!aNULL:!eNULL:!EXPORT:!DES:!3DES:"`.

**string ssl-crypto-device** [dovecot-configuration parameter]  
SSL crypto device to use, for valid values run `"openssl engine"`. Defaults to `""`.

**string postmaster-address** [dovecot-configuration parameter]  
Address to use when sending rejection mails. `%d` expands to recipient domain. Defaults to `"postmaster%d"`.

**string hostname** [dovecot-configuration parameter]  
Hostname to use in various parts of sent mails (e.g. in Message-Id) and in LMTP replies. Default is the system's real hostname@domain. Defaults to `""`.

**boolean quota-full-tempfail?** [dovecot-configuration parameter]  
If user is over quota, return with temporary failure instead of bouncing the mail. Defaults to `#f`.

**file-name sendmail-path** [dovecot-configuration parameter]  
Binary to use for sending mails. Defaults to `"/usr/sbin/sendmail"`.

**string submission-host** [dovecot-configuration parameter]  
If non-empty, send mails via this SMTP host[:port] instead of sendmail. Defaults to `""`.

**string rejection-subject** [dovecot-configuration parameter]  
Subject: header to use for rejection mails. You can use the same variables as for `'rejection-reason'` below. Defaults to `"Rejected: %s"`.

**string rejection-reason** [dovecot-configuration parameter]  
Human readable error message for rejection mails. You can use variables:

|                 |                  |
|-----------------|------------------|
| <code>%n</code> | CRLF             |
| <code>%r</code> | reason           |
| <code>%s</code> | original subject |
| <code>%t</code> | recipient        |

Defaults to `"Your message to <%t> was automatically rejected:%n%r"`.

**string recipient-delimiter** [dovecot-configuration parameter]  
Delimiter character between local-part and detail in email address. Defaults to `"+"`.

**string lda-original-recipient-header** [dovecot-configuration parameter]  
Header where the original recipient address (SMTP's RCPT TO: address) is taken from if not available elsewhere. With `dovecot-lda -a` parameter overrides this. A commonly used header for this is X-Original-To. Defaults to `""`.

**boolean lda-mailbox-autocreate?** [dovecot-configuration parameter]  
Should saving a mail to a nonexistent mailbox automatically create it?. Defaults to `#f`.

**boolean lda-mailbox-autosubscribe?** [dovecot-configuration parameter]  
Should automatically created mailboxes be also automatically subscribed?. Defaults to `#f`.

**non-negative-integer** [dovecot-configuration parameter]  
**imap-max-line-length**  
Maximum IMAP command line length. Some clients generate very long command lines with huge mailboxes, so you may need to raise this if you get "Too long argument" or "IMAP command line too large" errors often. Defaults to `64000`.

**string imap-logout-format** [dovecot-configuration parameter]  
IMAP logout format string:  
**%i** total number of bytes read from client  
**%o** total number of bytes sent to client.  
See `doc/wiki/Variables.txt` for a list of all the variables you can use. Defaults to `"in=%i out=%o deleted=%{deleted} expunged=%{expunged} trashed=%{trashed} hdr_count=%{fetch_hdr_count} hdr_bytes=%{fetch_hdr_bytes} body_count=%{fetch_body_count} body_bytes=%{fetch_body_bytes}"`.

**string imap-capability** [dovecot-configuration parameter]  
Override the IMAP CAPABILITY response. If the value begins with `'+'`, add the given capabilities on top of the defaults (e.g. `+XFOO XBAR`). Defaults to `""`.

**string imap-idle-notify-interval** [dovecot-configuration parameter]  
How long to wait between "OK Still here" notifications when client is IDLEing. Defaults to `"2 mins"`.

**string imap-id-send** [dovecot-configuration parameter]  
ID field names and values to send to clients. Using `*` as the value makes Dovecot use the default value. The following fields have default values currently: name, version, os, os-version, support-url, support-email. Defaults to `""`.

**string imap-id-log** [dovecot-configuration parameter]  
ID fields sent by client to log. `*` means everything. Defaults to `""`.

**space-separated-string-list** [dovecot-configuration parameter]  
**imap-client-workarounds**  
Workarounds for various client bugs:  
**delay-newmail**  
Send EXISTS/RECENT new mail notifications only when replying to NOOP and CHECK commands. Some clients ignore them otherwise, for

example OSX Mail (<v2.1). Outlook Express breaks more badly though, without this it may show user "Message no longer in server" errors. Note that OE6 still breaks even with this workaround if synchronization is set to "Headers Only".

#### **tb-extra-mailbox-sep**

Thunderbird gets somehow confused with LAYOUT=fs (mbox and mbox) and adds extra '/' suffixes to mailbox names. This option causes Dovecot to ignore the extra '/' instead of treating it as invalid mailbox name.

#### **tb-lsub-flags**

Show \Noselect flags for LSUB replies with LAYOUT=fs (e.g. mbox). This makes Thunderbird realize they aren't selectable and show them greyed out, instead of only later giving "not selectable" popup error.

Defaults to '()'.

**string imap-urlauth-host** [dovecot-configuration parameter]  
Host allowed in URLAUTH URLs sent by client. "\*" allows all. Defaults to '""'.

Whew! Lots of configuration options. The nice thing about it though is that Guix has a complete interface to Dovecot's configuration language. This allows not only a nice way to declare configurations, but also offers reflective capabilities as well: users can write code to inspect and transform configurations from within Scheme.

However, it could be that you just want to get a `dovecot.conf` up and running. In that case, you can pass an `opaque-dovecot-configuration` as the `#:config` parameter to `dovecot-service`. As its name indicates, an opaque configuration does not have easy reflective capabilities.

Available `opaque-dovecot-configuration` fields are:

**package dovecot** [opaque-dovecot-configuration parameter]  
The dovecot package.

**string string** [opaque-dovecot-configuration parameter]  
The contents of the `dovecot.conf`, as a string.

For example, if your `dovecot.conf` is just the empty string, you could instantiate a dovecot service like this:

```
(dovecot-service #:config
  (opaque-dovecot-configuration
    (string "")))
```

## OpenSMTPD Service

**opensmtpd-service-type** [Variable]  
This is the type of the OpenSMTPD (<https://www.opensmtpd.org>) service, whose value should be an `opensmtpd-configuration` object as in this example:

```
(service opensmtpd-service-type
  (opensmtpd-configuration
    (config-file (local-file "./my-smtpd.conf"))))
```

**opensmtpd-configuration** [Data Type]

Data type representing the configuration of opensmtpd.

**package** (default: *opensmtpd*)

Package object of the OpenSMTPD SMTP server.

**shepherd-requirement** (default: '()')

This option can be used to provide a list of symbols naming Shepherd services that this service will depend on, such as '**networking**' if you want to configure OpenSMTPD to listen on non-loopback interfaces.

**config-file** (default: %default-opensmtpd-config-file)

File-like object of the OpenSMTPD configuration file to use. By default it listens on the loopback network interface, and allows for mail from users and daemons on the local machine, as well as permitting email to remote servers. Run **man smtpd.conf** for more information.

**log-file** (default: "/var/log/mail.log")

The file location for the OpenSMTPD log file. Logging occurs in the foreground via the Shepherd, i.e. OpenSMTPD does not detach. When set to **#f** OpenSMTPD will instead detach and log to the syslog.

**setgid-commands?** (default: **#t**)

Make the following commands setgid to **smtpq** so they can be executed: **smtpctl**, **sendmail**, **send-mail**, **makemap**, **mailq**, and **newaliases**. See Section 11.11 [Privileged Programs], page 678, for more information on setgid programs.

## Exim Service

**exim-service-type** [Variable]

This is the type of the Exim (<https://exim.org>) mail transfer agent (MTA), whose value should be an **exim-configuration** object as in this example:

```
(service exim-service-type
  (exim-configuration
    (config-file (local-file "./my-exim.conf"))))
```

In order to use an **exim-service-type** service you must also have a **mail-aliases-service-type** service present in your **operating-system** (even if it has no aliases).

**exim-configuration** [Data Type]

Data type representing the configuration of exim.

**package** (default: *exim*)

Package object of the Exim server.

**config-file** (default: **#f**)

File-like object of the Exim configuration file to use. If its value is **#f** then use the default configuration file from the package provided in **package**. The resulting configuration file is loaded after setting the **exim\_user** and **exim\_group** configuration variables.



## Getmail service

**getmail-service-type** [Variable]

This is the type of the Getmail (<http://pyropus.ca/software/getmail/>) mail retriever, whose value should be a **getmail-configuration**.

Available **getmail-configuration** fields are:

**symbol name** [getmail-configuration parameter]

A symbol to identify the getmail service.

Defaults to `"unset"`.

**package package** [getmail-configuration parameter]

The getmail package to use.

**string user** [getmail-configuration parameter]

The user to run getmail as.

Defaults to `"getmail"`.

**string group** [getmail-configuration parameter]

The group to run getmail as.

Defaults to `"getmail"`.

**string directory** [getmail-configuration parameter]

The getmail directory to use.

Defaults to `"/var/lib/getmail/default"`.

**getmail-configuration-file rcfile** [getmail-configuration parameter]

The getmail configuration file to use.

Available **getmail-configuration-file** fields are:

**getmail-retriever-configuration** [getmail-configuration-file parameter]  
**retriever**

What mail account to retrieve mail from, and how to access that account.

Available **getmail-retriever-configuration** fields are:

**string type** [getmail-retriever-configuration parameter]

The type of mail retriever to use. Valid values include `'passwd'` and `'static'`.

Defaults to `"SimpleIMAPSSLRetriever"`.

**string server** [getmail-retriever-configuration parameter]

Username to login to the mail server with.

Defaults to `'unset'`.

**string username** [getmail-retriever-configuration parameter]

Username to login to the mail server with.

Defaults to `'unset'`.

```

non-negative-integer [getmail-retriever-configuration parameter]
    port
    Port number to connect to.
    Defaults to '#f'.

string password [getmail-retriever-configuration parameter]
    Override fields from passwd.
    Defaults to ''.

list [getmail-retriever-configuration parameter]
    password-command
    Override fields from passwd.
    Defaults to '()'.

string keyfile [getmail-retriever-configuration parameter]
    PEM-formatted key file to use for the TLS negotiation.
    Defaults to ''.

string certfile [getmail-retriever-configuration parameter]
    PEM-formatted certificate file to use for the TLS negotiation.
    Defaults to ''.

string ca-certs [getmail-retriever-configuration parameter]
    CA certificates to use.
    Defaults to ''.

parameter-alist [getmail-retriever-configuration parameter]
    extra-parameters
    Extra retriever parameters.
    Defaults to '()'.

getmail-destination-configuration [getmail-configuration-file parameter]
    destination
    What to do with retrieved messages.
    Available getmail-destination-configuration fields are:

string type [getmail-destination-configuration parameter]
    The type of mail destination. Valid values include 'Maildir', 'Mboxrd'
    and 'MDA_external'.
    Defaults to 'unset'.

string-or-filelike [getmail-destination-configuration parameter]
    path
    The path option for the mail destination. The behaviour depends on the
    chosen type.
    Defaults to ''.

```

**parameter-alist** [getmail-destination-configuration parameter]  
**extra-parameters**  
 Extra destination parameters  
 Defaults to ‘()’.

**getmail-options-configuration** [getmail-configuration-file parameter]  
**options**  
 Configure getmail.

Available getmail-options-configuration fields are:

**non-negative-integer** [getmail-options-configuration parameter]  
**verbose**  
 If set to ‘0’, getmail will only print warnings and errors. A value of ‘1’ means that messages will be printed about retrieving and deleting messages. If set to ‘2’, getmail will print messages about each of its actions.  
 Defaults to ‘1’.

**boolean read-all** [getmail-options-configuration parameter]  
 If true, getmail will retrieve all available messages. Otherwise it will only retrieve messages it hasn’t seen previously.  
 Defaults to ‘#t’.

**boolean delete** [getmail-options-configuration parameter]  
 If set to true, messages will be deleted from the server after retrieving and successfully delivering them. Otherwise, messages will be left on the server.  
 Defaults to ‘#f’.

**non-negative-integer** [getmail-options-configuration parameter]  
**delete-after**  
 Getmail will delete messages this number of days after seeing them, if they have been delivered. This means messages will be left on the server this number of days after delivering them. A value of ‘0’ disabled this feature.  
 Defaults to ‘0’.

**non-negative-integer** [getmail-options-configuration parameter]  
**delete-bigger-than**  
 Delete messages larger than this of bytes after retrieving them, even if the delete and delete-after options are disabled. A value of ‘0’ disables this feature.  
 Defaults to ‘0’.

**non-negative-integer** [getmail-options-configuration parameter]  
**max-bytes-per-session**  
 Retrieve messages totalling up to this number of bytes before closing the session with the server. A value of ‘0’ disables this feature.  
 Defaults to ‘0’.

**non-negative-integer** [getmail-options-configuration parameter]  
**max-message-size**

Don't retrieve messages larger than this number of bytes. A value of '0' disables this feature.

Defaults to '0'.

**boolean delivered-to** [getmail-options-configuration parameter]

If true, getmail will add a Delivered-To header to messages.

Defaults to '#t'.

**boolean received** [getmail-options-configuration parameter]

If set, getmail adds a Received header to the messages.

Defaults to '#t'.

**string message-log** [getmail-options-configuration parameter]

Getmail will record a log of its actions to the named file. A value of '' disables this feature.

Defaults to ''.

**boolean** [getmail-options-configuration parameter]

**message-log-syslog**

If true, getmail will record a log of its actions using the system logger.

Defaults to '#f'.

**boolean** [getmail-options-configuration parameter]

**message-log-verbose**

If true, getmail will log information about messages not retrieved and the reason for not retrieving them, as well as starting and ending information lines.

Defaults to '#f'.

**parameter-alist** [getmail-options-configuration parameter]

**extra-parameters**

Extra options to include.

Defaults to '()'.

**list idle** [getmail-configuration parameter]

A list of mailboxes that getmail should wait on the server for new mail notifications. This depends on the server supporting the IDLE extension.

Defaults to '()'.

**list environment-variables** [getmail-configuration parameter]

Environment variables to set for getmail.

Defaults to '()'.

## Mail Aliases Service

**mail-aliases-service-type** [Variable]

This is the type of the service which provides `/etc/aliases`, specifying how to deliver mail to users on this system.

```
(service mail-aliases-service-type
  ' (("postmaster" "bob")
    ("bob" "bob@example.com" "bob@example2.com")))
```

The configuration for a `mail-aliases-service-type` service is an association list denoting how to deliver mail that comes to this system. Each entry is of the form `(alias addresses ...)`, with `alias` specifying the local alias and `addresses` specifying where to deliver this user's mail.

The aliases aren't required to exist as users on the local system. In the above example, there doesn't need to be a `postmaster` entry in the operating-system's `user-accounts` in order to deliver the `postmaster` mail to bob (which subsequently would deliver mail to `bob@example.com` and `bob@example2.com`).

## GNU Mailutils IMAP4 Daemon

**imap4d-service-type** [Variable]

This is the type of the GNU Mailutils IMAP4 Daemon (see Section “imap4d” in *GNU Mailutils Manual*), whose value should be an `imap4d-configuration` object as in this example:

```
(service imap4d-service-type
  (imap4d-configuration
    (config-file (local-file "imap4d.conf"))))
```

**imap4d-configuration** [Data Type]

Data type representing the configuration of `imap4d`.

**package** (default: `mailutils`)

The package that provides `imap4d`.

**config-file** (default: `%default-imap4d-config-file`)

File-like object of the configuration file to use, by default it will listen on TCP port 143 of `localhost`. See Section “Conf-`imap4d`” in *GNU Mailutils Manual*, for details.

## Radicale Service

**radicale-service-type** [Variable]

This is the type of the Radicale (<https://radicale.org>) CalDAV/CardDAV server whose value should be a `radicale-configuration`. The default configuration matches the upstream documentation (<https://radicale.org/v3.html#configuration>).

**radicale-configuration** [Data Type]

Data type representing the configuration of `radicale`. Available `radicale-configuration` fields are:

**package** (default: **radicale**) (type: package)  
 Package that provides **radicale**.

**auth** (default: '()') (type: radicale-auth-configuration)  
 Configuration for auth-related variables.

**radicale-auth-configuration** [Data Type]  
 Data type representing the **auth** section of a **radicale** configuration file. Available **radicale-auth-configuration** fields are:

**type** (default: 'none') (type: symbol)  
 The method to verify usernames and passwords. Options are **none**, **htpasswd**, **remote-user**, and **http-x-remote-user**. This value is tied to **htpasswd-filename** and **htpasswd-encryption**.

**htpasswd-filename** (default: "/etc/radicale/users") (type: file-name)  
 Path to the htpasswd file. Use htpasswd or similar to generate this file.

**htpasswd-encryption** (default: 'md5') (type: symbol)  
 Encryption method used in the htpasswd file. Options are **plain**, **bcrypt**, and **md5**.

**delay** (default: 1) (type: non-negative-integer)  
 Average delay after failed login attempts in seconds.

**realm** (default: "Radicale - Password Required") (type: string)  
 Message displayed in the client when a password is needed.

**encoding** (default: '()') (type: radicale-encoding-configuration)  
 Configuration for encoding-related variables.

**radicale-encoding-configuration** [Data Type]  
 Data type representing the **encoding** section of a **radicale** configuration file. Available **radicale-encoding-configuration** fields are:

**request** (default: 'utf-8') (type: symbol)  
 Encoding for responding requests.

**stock** (default: 'utf-8') (type: symbol)  
 Encoding for storing local collections.

**headers-file** (default: none) (type: file-like)  
 Custom HTTP headers.

**logging** (default: '()') (type: radicale-logging-configuration)  
 Configuration for logging-related variables.

**radicale-logging-configuration** [Data Type]

Data type representing the **logging** section of a **radicale** configuration file. Available **radicale-logging-configuration** fields are:

**level** (default: 'warning') (type: symbol)

Set the logging level. One of **debug**, **info**, **warning**, **error**, or **critical**.

**mask-passwords?** (default: #t) (type: boolean)

Whether to include passwords in logs.

**rights** (default: '()') (type: radicale-rights-configuration)

Configuration for rights-related variables. This should be a **radicale-rights-configuration**.

**radicale-rights-configuration** [Data Type]

Data type representing the **rights** section of a **radicale** configuration file. Available **radicale-rights-configuration** fields are:

**type** (default: 'owner-only') (type: symbol)

Backend used to check collection access rights. The recommended backend is **owner-only**. If access to calendars and address books outside the home directory of users is granted, clients won't detect these collections and will not show them to the user. Choosing any other method is only useful if you access calendars and address books directly via URL. Options are **authenticate**, **owner-only**, **owner-write**, and **from-file**.

**file** (default: "") (type: file-name)

File for the rights backend **from-file**.

**server** (default: '()') (type: radicale-server-configuration)

Configuration for server-related variables. Ignored if WSGI is used.

**radicale-server-configuration** [Data Type]

Data type representing the **server** section of a **radicale** configuration file. Available **radicale-server-configuration** fields are:

**hosts** (default: (list "localhost:5232")) (type: list-of-ip-addresses)

List of IP addresses that the server will bind to.

**max-connections** (default: 8) (type: non-negative-integer)

Maximum number of parallel connections. Set to 0 to disable the limit.

**max-content-length** (default: 100000000) (type: non-negative-integer)

Maximum size of the request body in bytes.

**timeout** (default: 30) (type: non-negative-integer)  
 Socket timeout in seconds.

**ssl?** (default: #f) (type: boolean)  
 Whether to enable transport layer encryption.

**certificate** (default: "/etc/ssl/radicale.cert.pem") (type: file-name)  
 Path of the SSL certificate.

**key** (default: "/etc/ssl/radicale.key.pem") (type: file-name)  
 Path to the private key for SSL. Only effective if **ssl?** is #t.

**certificate-authority** (default: "") (type: file-name)  
 Path to CA certificate for validating client certificates. This can be used to secure TCP traffic between Radicale and a reverse proxy. If you want to authenticate users with client-side certificates, you also have to write an authentication plugin that extracts the username from the certificate.

**storage** (default: '()') (type: radicale-storage-configuration)  
 Configuration for storage-related variables.

**radicale-storage-configuration** [Data Type]  
 Data type representing the **storage** section of a **radicale** configuration file. Available **radicale-storage-configuration** fields are:

**type** (default: 'multifilesystem') (type: symbol)  
 Backend used to store data. Options are **multifilesystem** and **multifilesystem-nolock**.

**filesystem-folder** (default: "/var/lib/radicale/collections") (type: file-name)  
 Folder for storing local collections. Created if not present.

**max-sync-token-age** (default: 2592000) (type: non-negative-integer)  
 Delete sync-tokens that are older than the specified time in seconds.

**hook** (default: "") (type: string)  
 Command run after changes to storage.

**web-interface?** (default: #t) (type: boolean)  
 Whether to use Radicale's built-in web interface.



## Rspamd Service

**rspamd-service-type** [Variable]

This is the type of the Rspamd (<https://rspamd.com/>) filtering system whose value should be a **rspamd-configuration**.

**rspamd-configuration** [Data Type]

Available **rspamd-configuration** fields are:

**package** (default: **rspamd**) (type: file-like)

The package that provides rspamd.

**config-file** (default: **%default-rspamd-config-file**) (type: file-like)

File-like object of the configuration file to use. By default all workers are enabled except fuzzy and they are bound to their usual ports, e.g localhost:11334, localhost:11333 and so on

**local.d-files** (default: **()**) (type: directory-tree)

Configuration files in local.d, provided as a list of two element lists where the first element is the filename and the second one is a file-like object. Settings in these files will be merged with the defaults.

**override.d-files** (default: **()**) (type: directory-tree)

Configuration files in override.d, provided as a list of two element lists where the first element is the filename and the second one is a file-like object. Settings in these files will override the defaults.

**user** (default: **%default-rspamd-account**) (type: user-account)

The user to run rspamd as.

**group** (default: **%default-rspamd-group**) (type: user-group)

The group to run rspamd as.

**debug?** (default: **#f**) (type: boolean)

Force debug output.

**insecure?** (default: **#f**) (type: boolean)

Ignore running workers as privileged users.

**skip-template?** (default: **#f**) (type: boolean)

Do not apply Jinja templates.

**shepherd-requirements** (default: **(user-processes loopback)**) (type: list-of-symbols)

This is a list of symbols naming Shepherd services that this service will depend on.

### 11.10.14 Messaging Services

The (**gnu services messaging**) module provides Guix service definitions for messaging services. Currently it provides the following services:

## Prosody Service

**prosody-service-type** [Variable]

This is the type for the Prosody XMPP communication server (<https://prosody.im>). Its value must be a **prosody-configuration** record as in this example:

```
(service prosody-service-type
  (prosody-configuration
    (modules-enabled (cons* "groups" "mam" %default-modules-enabled)))
    (int-components
      (list
        (int-component-configuration
          (hostname "conference.example.net")
          (plugin "muc")
          (mod-muc (mod-muc-configuration))))))
    (virtualhosts
      (list
        (virtualhost-configuration
          (domain "example.net"))))))
```

See below for details about **prosody-configuration**.

By default, Prosody does not need much configuration. Only one **virtualhosts** field is needed: it specifies the domain you wish Prosody to serve.

You can perform various sanity checks on the generated configuration with the **prosodyctl check** command.

Prosodyctl will also help you to import certificates from the **letsencrypt** directory so that the **prosody** user can access them. See <https://prosody.im/doc/letsencrypt>.

```
prosodyctl --root cert import /etc/certs
```

The available configuration parameters follow. Each parameter definition is preceded by its type; for example, ‘**string-list foo**’ indicates that the **foo** parameter should be specified as a list of strings. Types starting with **maybe-** denote parameters that won’t show up in **prosody.cfg.lua** when their value is left unspecified.

There is also a way to specify the configuration as a string, if you have an old **prosody.cfg.lua** file that you want to port over from some other system; see the end for more details.

The **file-object** type designates either a file-like object (see Section 8.12 [G-Expressions], page 165) or a file name.

Available **prosody-configuration** fields are:

**package prosody** [prosody-configuration parameter]  
The Prosody package.

**file-name data-path** [prosody-configuration parameter]  
Location of the Prosody data storage directory. See <https://prosody.im/doc/configure>. Defaults to ‘**"/var/lib/prosody"**’.

**file-object-list plugin-paths** [prosody-configuration parameter]  
 Additional plugin directories. They are searched in all the specified paths in order.  
 See [https://prosody.im/doc/plugins\\_directory](https://prosody.im/doc/plugins_directory). Defaults to `'()'`.

**file-name certificates** [prosody-configuration parameter]  
 Every virtual host and component needs a certificate so that clients and servers can securely verify its identity. Prosody will automatically load certificates/keys from the directory specified here. Defaults to `"/etc/prosody/certs"`.

**string-list admins** [prosody-configuration parameter]  
 This is a list of accounts that are admins for the server. Note that you must create the accounts separately. See <https://prosody.im/doc/admins> and [https://prosody.im/doc/creating\\_accounts](https://prosody.im/doc/creating_accounts). Example: `(admins '("user1@example.com" "user2@example.net"))` Defaults to `'()'`.

**boolean use-libevent?** [prosody-configuration parameter]  
 Enable use of libevent for better performance under high load. See <https://prosody.im/doc/libevent>. Defaults to `#f`.

**module-list modules-enabled** [prosody-configuration parameter]  
 This is the list of modules Prosody will load on startup. It looks for `mod_modulename.lua` in the plugins folder, so make sure that exists too. Documentation on modules can be found at: <https://prosody.im/doc/modules>. Defaults to `'("roster" "saslauth" "tls" "dialback" "disco" "carbons" "private" "blocklist" "vcard" "version" "uptime" "time" "ping" "pep" "register" "admin_adhoc")'`.

**string-list modules-disabled** [prosody-configuration parameter]  
`"offline"`, `"c2s"` and `"s2s"` are auto-loaded, but should you want to disable them then add them to this list. Defaults to `'()'`.

**file-object groups-file** [prosody-configuration parameter]  
 Path to a text file where the shared groups are defined. If this path is empty then `'mod_groups'` does nothing. See [https://prosody.im/doc/modules/mod\\_groups](https://prosody.im/doc/modules/mod_groups). Defaults to `"/var/lib/prosody/sharedgroups.txt"`.

**boolean allow-registration?** [prosody-configuration parameter]  
 Disable account creation by default, for security. See [https://prosody.im/doc/creating\\_accounts](https://prosody.im/doc/creating_accounts). Defaults to `#f`.

**maybe-ssl-configuration ssl** [prosody-configuration parameter]  
 These are the SSL/TLS-related settings. Most of them are disabled so to use Prosody's defaults. If you do not completely understand these options, do not add them to your config, it is easy to lower the security of your server using them. See [https://prosody.im/doc/advanced\\_ssl\\_config](https://prosody.im/doc/advanced_ssl_config).

Available `ssl-configuration` fields are:

**maybe-string protocol** [ssl-configuration parameter]  
 This determines what handshake to use.

|                                                                                                                                                                                                             |                                   |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------|
| <b>maybe-file-name key</b>                                                                                                                                                                                  | [ssl-configuration parameter]     |
| Path to your private key file.                                                                                                                                                                              |                                   |
| <b>maybe-file-name certificate</b>                                                                                                                                                                          | [ssl-configuration parameter]     |
| Path to your certificate file.                                                                                                                                                                              |                                   |
| <b>file-object capath</b>                                                                                                                                                                                   | [ssl-configuration parameter]     |
| Path to directory containing root certificates that you wish Prosody to trust when verifying the certificates of remote servers. Defaults to <code>"/etc/ssl/certs"</code> .                                |                                   |
| <b>maybe-file-object cafile</b>                                                                                                                                                                             | [ssl-configuration parameter]     |
| Path to a file containing root certificates that you wish Prosody to trust. Similar to <code>capath</code> but with all certificates concatenated together.                                                 |                                   |
| <b>maybe-string-list verify</b>                                                                                                                                                                             | [ssl-configuration parameter]     |
| A list of verification options (these mostly map to OpenSSL's <code>set_verify()</code> flags).                                                                                                             |                                   |
| <b>maybe-string-list options</b>                                                                                                                                                                            | [ssl-configuration parameter]     |
| A list of general options relating to SSL/TLS. These map to OpenSSL's <code>set_options()</code> . For a full list of options available in LuaSec, see the LuaSec source.                                   |                                   |
| <b>maybe-non-negative-integer depth</b>                                                                                                                                                                     | [ssl-configuration parameter]     |
| How long a chain of certificate authorities to check when looking for a trusted root certificate.                                                                                                           |                                   |
| <b>maybe-string ciphers</b>                                                                                                                                                                                 | [ssl-configuration parameter]     |
| An OpenSSL cipher string. This selects what ciphers Prosody will offer to clients, and in what order.                                                                                                       |                                   |
| <b>maybe-file-name dhparam</b>                                                                                                                                                                              | [ssl-configuration parameter]     |
| A path to a file containing parameters for Diffie-Hellman key exchange. You can create such a file with: <code>openssl dhparam -out /etc/prosody/certs/dh-2048.pem 2048</code>                              |                                   |
| <b>maybe-string curve</b>                                                                                                                                                                                   | [ssl-configuration parameter]     |
| Curve for Elliptic curve Diffie-Hellman. Prosody's default is <code>"secp384r1"</code> .                                                                                                                    |                                   |
| <b>maybe-string-list verifyext</b>                                                                                                                                                                          | [ssl-configuration parameter]     |
| A list of "extra" verification options.                                                                                                                                                                     |                                   |
| <b>maybe-string password</b>                                                                                                                                                                                | [ssl-configuration parameter]     |
| Password for encrypted private keys.                                                                                                                                                                        |                                   |
| <b>boolean c2s-require-encryption?</b>                                                                                                                                                                      | [prosody-configuration parameter] |
| Whether to force all client-to-server connections to be encrypted or not. See <a href="https://prosody.im/doc/modules/mod_tls">https://prosody.im/doc/modules/mod_tls</a> . Defaults to <code>'#f'</code> . |                                   |
| <b>string-list disable-sasl-mechanisms</b>                                                                                                                                                                  | [prosody-configuration parameter] |
| Set of mechanisms that will never be offered. See <a href="https://prosody.im/doc/modules/mod_saslauth">https://prosody.im/doc/modules/mod_saslauth</a> . Defaults to <code>'('("DIGEST-MD5")'</code> .     |                                   |

- string-list insecure-sasl-mechanisms** [prosody-configuration parameter]  
Set of mechanisms that will not be offered on unencrypted connections. See [https://prosody.im/doc/modules/mod\\_saslauth](https://prosody.im/doc/modules/mod_saslauth). Defaults to `'("PLAIN" "LOGIN")'`.
- boolean s2s-require-encryption?** [prosody-configuration parameter]  
Whether to force all server-to-server connections to be encrypted or not. See [https://prosody.im/doc/modules/mod\\_tls](https://prosody.im/doc/modules/mod_tls). Defaults to `#f`.
- boolean s2s-secure-auth?** [prosody-configuration parameter]  
Whether to require encryption and certificate authentication. This provides ideal security, but requires servers you communicate with to support encryption AND present valid, trusted certificates. See <https://prosody.im/doc/s2s#security>. Defaults to `#f`.
- string-list s2s-insecure-domains** [prosody-configuration parameter]  
Many servers don't support encryption or have invalid or self-signed certificates. You can list domains here that will not be required to authenticate using certificates. They will be authenticated using DNS. See <https://prosody.im/doc/s2s#security>. Defaults to `'()'`.
- string-list s2s-secure-domains** [prosody-configuration parameter]  
Even if you leave `s2s-secure-auth?` disabled, you can still require valid certificates for some domains by specifying a list here. See <https://prosody.im/doc/s2s#security>. Defaults to `'()'`.
- string authentication** [prosody-configuration parameter]  
Select the authentication backend to use. The default provider stores passwords in plaintext and uses Prosody's configured data storage to store the authentication data. If you do not trust your server please see [https://prosody.im/doc/modules/mod\\_auth\\_internal\\_hashed](https://prosody.im/doc/modules/mod_auth_internal_hashed) for information about using the hashed backend. See also <https://prosody.im/doc/authentication>. Defaults to `"internal_plain"`.
- maybe-string log** [prosody-configuration parameter]  
Set logging options. Advanced logging configuration is not yet supported by the Prosody service. See <https://prosody.im/doc/logging>. Defaults to `"*syslog"`.
- file-name pidfile** [prosody-configuration parameter]  
File to write pid in. See [https://prosody.im/doc/modules/mod\\_posix](https://prosody.im/doc/modules/mod_posix). Defaults to `"/var/run/prosody/prosody.pid"`.
- maybe-non-negative-integer http-max-content-size** [prosody-configuration parameter]  
Maximum allowed size of the HTTP body (in bytes).
- maybe-string http-external-url** [prosody-configuration parameter]  
Some modules expose their own URL in various ways. This URL is built from the protocol, host and port used. If Prosody sits behind a proxy, the public URL will be `http-external-url` instead. See [https://prosody.im/doc/http#external\\_url](https://prosody.im/doc/http#external_url).

**virtualhost-configuration-list** [prosody-configuration parameter]  
**virtualhosts**

A host in Prosody is a domain on which user accounts can be created. For example if you want your users to have addresses like `"john.smith@example.com"` then you need to add a host `"example.com"`. All options in this list will apply only to this host.

**Note:** The name *virtual* host is used in configuration to avoid confusion with the actual physical host that Prosody is installed on. A single Prosody instance can serve many domains, each one defined as a `VirtualHost` entry in Prosody's configuration. Conversely a server that hosts a single domain would have just one `VirtualHost` entry.

See [https://prosody.im/doc/configure#virtual\\_host\\_settings](https://prosody.im/doc/configure#virtual_host_settings).

Available **virtualhost-configuration** fields are:

all these **prosody-configuration** fields: `admins`, `use-libevent?`, `modules-enabled`, `modules-disabled`, `groups-file`, `allow-registration?`, `ssl`, `c2s-require-encryption?`, `disable-sasl-mechanisms`, `insecure-sasl-mechanisms`, `s2s-require-encryption?`, `s2s-secure-auth?`, `s2s-insecure-domains`, `s2s-secure-domains`, `authentication`, `log`, `http-max-content-size`, `http-external-url`, `raw-content`, plus:

**string domain** [virtualhost-configuration parameter]  
 Domain you wish Prosody to serve.

**int-component-configuration-list** [prosody-configuration parameter]  
**int-components**

Components are extra services on a server which are available to clients, usually on a subdomain of the main server (such as `"mycomponent.example.com"`). Example components might be chatroom servers, user directories, or gateways to other protocols.

Internal components are implemented with Prosody-specific plugins. To add an internal component, you simply fill the `hostname` field, and the plugin you wish to use for the component.

See <https://prosody.im/doc/components>. Defaults to `'()'`.

Available **int-component-configuration** fields are:

all these **prosody-configuration** fields: `admins`, `use-libevent?`, `modules-enabled`, `modules-disabled`, `groups-file`, `allow-registration?`, `ssl`, `c2s-require-encryption?`, `disable-sasl-mechanisms`, `insecure-sasl-mechanisms`, `s2s-require-encryption?`, `s2s-secure-auth?`, `s2s-insecure-domains`, `s2s-secure-domains`, `authentication`, `log`, `http-max-content-size`, `http-external-url`, `raw-content`, plus:

**string hostname** [int-component-configuration parameter]  
 Hostname of the component.

**string plugin** [int-component-configuration parameter]  
 Plugin you wish to use for the component.

**maybe-mod-muc-configuration** [int-component-configuration parameter]  
**mod-muc**

Multi-user chat (MUC) is Prosody's module for allowing you to create hosted chatrooms/conferences for XMPP users.

General information on setting up and using multi-user chatrooms can be found in the "Chatrooms" documentation (<https://prosody.im/doc/chatrooms>), which you should read if you are new to XMPP chatrooms.

See also [https://prosody.im/doc/modules/mod\\_muc](https://prosody.im/doc/modules/mod_muc).

Available **mod-muc-configuration** fields are:

**string name** [mod-muc-configuration parameter]  
 The name to return in service discovery responses. Defaults to "Prosody Chatrooms".

**string-or-boolean restrict-room-creation** [mod-muc-configuration parameter]

If '#t', this will only allow admins to create new chatrooms. Otherwise anyone can create a room. The value "local" restricts room creation to users on the service's parent domain. E.g. 'user@example.com' can create rooms on 'rooms.example.com'. The value "admin" restricts to service administrators only. Defaults to '#f'.

**non-negative-integer max-history-messages** [mod-muc-configuration parameter]

Maximum number of history messages that will be sent to the member that has just joined the room. Defaults to '20'.

**ext-component-configuration-list ext-components** [prosody-configuration parameter]

External components use XEP-0114, which most standalone components support. To add an external component, you simply fill the hostname field. See <https://prosody.im/doc/components>. Defaults to '()'.

Available **ext-component-configuration** fields are:

all these **prosody-configuration** fields: **admins**, **use-libevent?**, **modules-enabled**, **modules-disabled**, **groups-file**, **allow-registration?**, **ssl**, **c2s-require-encryption?**, **disable-sasl-mechanisms**, **insecure-sasl-mechanisms**, **s2s-require-encryption?**, **s2s-secure-auth?**, **s2s-insecure-domains**, **s2s-secure-domains**, **authentication**, **log**, **http-max-content-size**, **http-external-url**, **raw-content**, plus:

**string component-secret** [ext-component-configuration parameter]  
 Password which the component will use to log in.

**string hostname** [ext-component-configuration parameter]  
 Hostname of the component.

**non-negative-integer-list component-ports** [prosody-configuration parameter]

Port(s) Prosody listens on for component connections. Defaults to '(5347)'.

**string component-interface** [prosody-configuration parameter]  
Interface Prosody listens on for component connections. Defaults to "127.0.0.1".

**maybe-raw-content raw-content** [prosody-configuration parameter]  
Raw content that will be added to the configuration file.

It could be that you just want to get a `prosody.cfg.lua` up and running. In that case, you can pass an `opaque-prosody-configuration` record as the value of `prosody-service-type`. As its name indicates, an opaque configuration does not have easy reflective capabilities. Available `opaque-prosody-configuration` fields are:

**package prosody** [opaque-prosody-configuration parameter]  
The prosody package.

**string prosody.cfg.lua** [opaque-prosody-configuration parameter]  
The contents of the `prosody.cfg.lua` to use.

For example, if your `prosody.cfg.lua` is just the empty string, you could instantiate a prosody service like this:

```
(service prosody-service-type
  (opaque-prosody-configuration
    (prosody.cfg.lua "")))
```

## BitlBee Service

BitlBee (<https://bitlbee.org>) is a gateway that provides an IRC interface to a variety of messaging protocols such as XMPP.

**bitlbee-service-type** [Variable]  
This is the service type for the BitlBee (<https://bitlbee.org>) IRC gateway daemon. Its value is a `bitlbee-configuration` (see below).

To have BitlBee listen on port 6667 on localhost, add this line to your services:

```
(service bitlbee-service-type)
```

**bitlbee-configuration** [Data Type]  
This is the configuration for BitlBee, with the following fields:

**interface** (default: "127.0.0.1")

**port** (default: 6667)

Listen on the network interface corresponding to the IP address specified in *interface*, on *port*.

When *interface* is 127.0.0.1, only local clients can connect; when it is 0.0.0.0, connections can come from any networking interface.

**bitlbee** (default: bitlbee)

The BitlBee package to use.

**plugins** (default: '()')

List of plugin packages to use—e.g., `bitlbee-discord`.

**extra-settings** (default: "")

Configuration snippet added as-is to the BitlBee configuration file.



## ngIRCd service

ngIRCd (<https://ngircd.barton.de/>), is a lightweight IRCd (Internet Relay Chat daemon), which can be used to host your own IRC server.

**ngircd-service-type** [Variable]

The service type for ngIRCd. Its value is a **ngircd-configuration** object, documented below.

In its most simple form, the service can be used with its default configuration, like this:

```
(service ngircd-service-type)
```

A slightly more complex example configuration could look like:

```
(service ngircd-service-type
  (ngircd-configuration
    (channels
      (list (ngircd-channel
              (name "#fruits")
              (topic "All things fruits -- veggies are off-topic"))))
    (operators
      (list (ngircd-operator
              (name "mikan")
              (password "tomatoes-are-fruits/carrots-are-not")))))
```

**ngircd-configuration** [Data Type]

Available **ngircd-configuration** fields are:

**ngircd** (default: **ngircd**) (type: file-like)

The **ngircd** package to use.

**debug?** (default: **#f**) (type: boolean)

Turn on debugging messages.

**shepherd-requirement** (default: (**user-processes networking**)) (type: list-of-symbols)

Shepherd requirements the service should depend on.

**global** (type: **ngircd-global**)

A **ngircd-global** record object used to specify global options.

**limits** (type: **maybe-ngircd-limits**)

The **ngircd-limits** record object used to specify limits options.

**options** (type: **ngircd-options**)

The **ngircd-options** record object used to specify optional features and configuration options.

**ssl** (type: **maybe-ngircd-ssl**)

The **ngircd-ssl** record object used to specify the SSL-related options.

**operators** (type: **maybe-list-of-ngircd-operators**)

A list of **ngircd-operator** record objects used to specify the operators.

**servers** (type: maybe-list-of-ngircd-servers)  
A list of ngircd-server record objects used to specify other remote servers to connect to.

**channels** (type: maybe-list-of-ngircd-channels)  
A list of ngircd-channels record objects specifying pre-defined channels to be created by the server when starting up.

**ngircd-global** [Data Type]

Available **ngircd-global** fields are:

**name** (type: maybe-string)  
Server name in the IRC network. This is an individual name of the IRC server, it is not related to the DNS host name. It must be unique in the IRC network and must contain at least one dot (‘.’) character. When not set, ngIRCd tries to deduce a valid IRC server name from the local host name.

**admin-info-1** (type: maybe-string)  
First administrator information.

**admin-info-2** (type: maybe-string)  
Second administrator information.

**admin-email** (type: maybe-string)  
Email to reach administrators.

**help-file** (type: maybe-file-like)  
File-like containing the ngIRCd help text.

**info** (type: maybe-string)  
Info text of the server. This will be shown by WHOIS and LINKS requests for example.

**listen** (default: ("::" "0.0.0.0")) (type: list-of-strings)  
A list of IP address on which the server should listen. By default it listens on all configured IP addresses and interfaces.

**motd-file** (type: file-like)  
Text file with the *message of the day* (MOTD). This message will be shown to all users connecting to the server.

**motd-phrase** (type: maybe-string)  
A simple phrase (<127 chars) to use if you don’t want to use a MOTD file.

**network** (type: maybe-string)  
The name of the IRC network to which this server belongs. This name is optional, should only contain ASCII characters, and can’t contain spaces. It is only used to inform clients.

**password** (type: maybe-string)  
Global password or all users needed to connect to the server. By default, no password is required. PAM must be disabled for this option to have an effect.

**pid-file** (default: `"/run/ngircd/ngircd.pid"`) (type: string)  
 The file name where the PID of ngIRCd should be written after it starts.

**ports** (default: `(6667)`) (type: list-of-ports)  
 Port number(s) on which the server should listen for *unencrypted* connections.

**server-uid** (default: `"ngircd"`) (type: string-or-number)  
 The user that the `ngircd` command should run as.

**server-gid** (default: `"ngircd"`) (type: string-or-number)  
 The group that the `ngircd` command should run as.

### **ngircd-limits** [Data Type]

Available `ngircd-limits` fields are:

**connect-retry** (default: `60`) (type: maybe-number)  
 The number of seconds the server should wait before re-attempting to establish a link to not yet (or no longer) connected servers.

**max-connections** (default: `0`) (type: maybe-number)  
 Maximum number of simultaneous in- and outbound connections the server is allowed to accept. There is no limit by default.

**max-connections-ip** (default: `5`) (type: maybe-number)  
 Maximum number of simultaneous connections from a single IP address that the server will accept. This configuration options lowers the risk of denial of service attacks (DoS). Set to `0` to remove the limit.

**max-joins** (default: `10`) (type: maybe-number)  
 Maximum number of channels a user can be member of. Set to `0` to remove the limit.

**max-list-size** (default: `100`) (type: maybe-number)  
 Maximum number of channels returned in response to a `LIST` command.

**ping-timeout** (default: `120`) (type: maybe-number)  
 Number of seconds of inactivity after which the server will send a `PING` to the peer to test whether it is alive or not.

**pong-timeout** (default: `20`) (type: maybe-number)  
 If a client fails to answer a `PING` with a `PONG` within this amount of seconds, it will be disconnected by the server.

### **ngircd-options** [Data Type]

Available `ngircd-options` fields are:

**allowed-channel-types** (default: `"#&+"`) (type: maybe-string)  
 List of allowed channel types (channel prefixes) for newly created channels on the local server. By default, all supported channel types are allowed.

**allow-remote-oper?** (default: `#f`) (type: maybe-boolean)  
 If this option is active, IRC operators connected to remote servers are allowed to control this local server using administrative commands, for example like `CONNECT`, `DIE`, `SQUIT`, etc.

- connect-ipv4?** (default: **#t**) (type: maybe-boolean)  
Set to **#f** to prevent ngIRCd from connecting to other IRC servers using the IPv4 protocol, allowed by default.
- connect-ipv6?** (default: **#t**) (type: maybe-boolean)  
Set to **#f** to prevent ngIRCd from connecting to other IRC servers using the IPv6 protocol, allowed by default.
- dns?** (default: **#t**) (type: maybe-boolean)  
Set to **#f** to disable DNS lookups when clients connect. If you configure the daemon to connect to other servers, ngIRCd may still perform a DNS lookup if required.
- more-privacy?** (default: **#f**) (type: maybe-boolean)  
Set this to **#t** to have ngIRCd censor user idle time, logon time as well as the PART/QUIT messages (that sometimes used to inform everyone about which client software is being used). WHOWAS requests are also silently ignored, and NAMES output doesn't list any clients for non-members. This option is most useful when ngIRCd is being used together with anonymizing software such as TOR or I2P and one does not wish to make it too easy to collect statistics on the users.
- notice-before-registration?** (default: **#f**) (type: maybe-boolean)  
Normally ngIRCd doesn't send any messages to a client until it is registered. Enable this option to let the daemon send 'NOTICE \*' messages to clients while connecting.
- oper-can-use-mode?** (default: **#f**) (type: maybe-boolean)  
Should IRC Operators be allowed to use the MODE command even if they are not(!) channel-operators?
- oper-chan-p-auto-op?** (default: **#t**) (type: maybe-boolean)  
Should IRC Operators get AutoOp (+o) in persistent (+P) channels?
- oper-server-mode?** (default: **#f**) (type: maybe-boolean)  
If **open-can-use-mode?** is **#t**, this may lead the compatibility problems with servers that run the ircd-irc2 software. This option masks mode requests by non-chanops as if they were coming from the server. Only enable this if you have ircd-irc2 servers in your IRC network.
- pam?** (default: **#f**) (type: boolean)  
Set to **#t** to enable calls to the PAM library at runtime; all users connecting without password are allowed to connect, all passwords given will fail. Users identified without PAM are registered with a tilde ('~') prepended to their user name. This defaults to **#f** in Guix because the service runs as an unprivileged user and thus cannot authenticate other users via the **pam\_unix** PAM module.
- pam-is-optional?** (default: **#f**) (type: maybe-boolean)  
Set to **#t** to make PAM authentication optional, causing clients not sending a password to still be able to connect, but won't become identified and keep the tilde ('~') character prepended to their supplied user name.

**require-auth-ping?** (default: **#f**) (type: maybe-boolean)  
 Set to **#t** to have ngIRCd send an authentication PING when a new client connects, and register this client only after receiving the corresponding PONG reply.

#### **ngircd-ssl** [Data Type]

Available **ngircd-ssl** fields are:

**cert-file** (type: maybe-string)  
 SSL certificate file of the private server key.

**key-file** (type: maybe-string)  
 File name of the SSL Server Key to be used for SSL connections, which is required for SSL/TLS support.

**ca-file** (default: **"/etc/ssl/certs/ca-certificates.crt"**) (type: maybe-string)  
 A file listing all the certificates of the trusted Certificate Authorities.

**ports** (type: maybe-list-of-ports)  
 Like the global configuration's **port** option, except that ngIRCd will expect incoming connections to be SSL/TLS encrypted. Common port numbers for SSL-encrypted IRC are 6669 and 6697.

**cipher-list** (type: maybe-string)  
 The GnuTLS cipher suites allowed for SSL/TLS connections, a value such as **"SECURE128:-VERS-SSL3.0"**. Refer to **'man 3 gnutls\_priority\_init'** for details.

**dh-file** (type: maybe-file-like)  
 A file-like containing the Diffie-Hellman parameters, which can be created with GnuTLS via **'certtool --generate-dh-params'**. If this file is not present, the Diffie-Hellman parameters will be computed on startup, which may take some time.

#### **ngircd-operator** [Data Type]

Available **ngircd-operator** fields are:

**name** (type: string)  
 ID of the operator (may be different of the nickname).

**password** (type: string)  
 Password of the IRC operator.

**mask** (type: maybe-string)  
 Mask that is to be checked before an **/OPER** for this account is accepted, for example: **"nick!ident@\*.example.com"**.

#### **ngircd-server** [Data Type]

Available **ngircd-server** fields are:

**name** (type: string)  
 IRC name of the remote server.

**host** (type: string)  
 Internet host name (or IP address) of the peer.

- my-password** (type: string)  
Own password for this connection. This password has to be configured as **peer-password** on the other server and must not have ‘:’ as first character.
- peer-password** (type: string)  
Foreign password for this connection. This password has to be configured as **my-password** on the other server.
- bind** (type: maybe-string)  
IP address to use as source IP for the outgoing connection. The default is to let the operating system decide.
- port** (type: maybe-port)  
Port of the remote server to which ngIRCd should connect (active). If no port is assigned to a configured server, the daemon only waits for incoming connections (passive, which is the default).
- group** (type: maybe-number)  
Group of this server.
- passive?** (default: #f) (type: maybe-boolean)  
Set to #t to disable automatic connection even if the port value is specified.
- ssl-connect?** (default: #f) (type: maybe-boolean)  
Set to #t to enforce the use of TLS to connect to the remote server.
- ssl-verify?** (default: #t) (type: maybe-boolean)  
Set to #f to disable the verification of the TLS certificate presented by the remote server. This can be useful if the remote server uses a self-signed certificate. Beware that this degrades security by enabling man-in-the-middle attacks, as the *remote site is not verified at all*.

**ngircd-channel** [Data Type]

Available **ngircd-channel** fields are:

- name** (type: string)  
Name of the channel, including channel prefix (“#” or “&”).
- topic** (type: maybe-string)  
Topic for this channel.
- modes** (type: maybe-list-of-strings)  
Initial channel modes, as used in MODE commands. Modifying lists (ban list, invite list, exception list) is supported. If multiple MODE strings are specified, they are evaluated in the order listed (left to right).
- key-file** (type: maybe-file-like)  
Path and file name of a ngIRCd key file containing individual channel keys for different users. Refer to ‘**man 5 ngircd.conf**’ for more details.

## Pounce Service

**pounce** (<https://git.causal.agency/pounce/about/>) is a multi-client, TLS-only IRC bouncer. It maintains a persistent connection to an IRC server, acting as a proxy and buffer for a number of clients.

**pounce-service-type** [Variable]

This is the service type for the pounce IRC bouncer. Its value is a **pounce-configuration** configuration instance, which is documented below.

The following example configures pounce to act as an IRC bouncer for the Libera.Chat (<https://libera.chat>) server, using CertFP (client certificate fingerprint) authentication to avoid leaking a sensitive password to the publicly readable store. The equally sensitive TLS certificate file should be created in-place or transferred using a secure means such as SSH, prior to deploying the service. The service activation will ensure the ownership and permissions of the certificate/key files are set correctly. In the below example, it is placed at `/etc/pounce/libera.pem` on the target machine. Pounce itself can be used to generate a TLS certificate, using the `'pounce -g libera.pem'` command, which concatenates both the private key and the public certificate in the specified file name. For more information regarding CertFP authentication, refer to `'man pounce'` or the Libera.Chat guide at <https://libera.chat/guides/certfp>.

```
(service pounce-service-type
  (pounce-configuration
    (host "irc.libera.chat")
    (client-cert "/etc/pounce/libera.pem")
    (nick "hannah")
    (join (list "#gnu" "#guix" "#guile" "#hurd"))))
```

Once deployed on the target machine, pounce will act as an IRC server listening for TLS connections on the 6697 TCP port of the `'localhost'` address of that machine. By default, a self-signed certificate for pounce is created at `/var/lib/pounce/.config/pounce/localhost.cert`. If you plan to expose the bouncer to the public Internet, it is advisable to use a CA (Certificate Authority)-signed certificate, as can be obtained using a certificate service (see Section 11.10.22 [Certificate Services], page 518), so that IRC clients can verify the certificate out of the box. If you instead plan to connect to the bouncer strictly via a secure connection, for example using a VPN (Virtual Private Network) or SSH (Secure Shell), then it is acceptable to simply let your IRC client trust the auto-generated, self-signed pounce certificate or even disable TLS certificate verification in your client.

To connect to a second server, a second pounce instance is needed, taking care to specify the **provision** field of its **pounce-configuration** to avoid a name clash with the previous service, along with a distinct **local-port** and **log-file**. The following example shows how to configure another bouncer, this time for the OFTC (<https://www.oftc.net>) IRC server. Like in the previous example, CertFP authentication is used, which can be configured similarly. For more details about using CertFP with the OFTC IRC server, refer to <https://www.oftc.net/NickServ/CertFP/>.

```
(service pounce-service-type
```

```
(pounce-configuration
  (shepherd-provision '(pounce-oftc))
  (local-port 6698)
  (log-file "/var/log/pounce-oftc.log")
  (host "irc.oftc.net")
  (client-cert "/etc/pounce/oftc.pem")
  (nick "sena")
  (join (list "#gcc" "#glibc"))))
```

**pounce-configuration** [Data Type]

Available pounce-configuration fields are:

**pounce** (default: **pounce**) (type: file-like)

The pounce package to use.

**shepherd-provision** (default: (**pounce**)) (type: list-of-symbols)

The name(s) of the service.

**shepherd-requirement** (default: (**user-processes networking**)) (type: list-of-symbols)

Shepherd requirements the service should depend on.

**log-file** (default: **"/var/log/pounce.log"**) (type: string)

The log file name to use.

**verbose?** (type: maybe-boolean)

When true, log IRC messages to standard output.

**local-host** (default: **"localhost"**) (type: maybe-string)

The host to bind to.

**local-port** (default: **6697**) (type: maybe-port)

The port to bind to.

**local-ca** (type: maybe-string)

Require clients to authenticate using a TLS client certificate either contained in or signed by a certificate in the file loaded from **local-ca**, a file name. The file is reloaded when the SIGUSR1 signal is received.

**local-cert** (type: maybe-string)

File name of the TLS certificate to load. The file is reloaded when the SIGUSR1 signal is received. Unless specified, a self-signed certificate is generated at **/var/lib/pounce/.config/pounce/host.pem**, where *host* corresponds to the value of the **local-host** field.

**local-priv** (type: maybe-string)

File name of the private TLS key to load. Unless specified, a key is generated at **/var/lib/pounce/.config/pounce/host.key**, where *host* corresponds to the value of the **local-host** field.

**local-pass** (type: maybe-string)

Require the server password pass for clients to connect. The pass string must be hashed using **'pounce -x'**.



- size** (default: 4096) (type: maybe-power-of-two)  
Set the number of messages contained in the buffer to *size*. This sets the maximum number of recent messages which can be relayed to a reconnecting client. The size must be a power of two.
- bind** (type: maybe-string)  
Host to bind the *source* address to when connecting to the server. To connect from any address over IPv4 only, use '0.0.0.0'. To connect from any address over IPv6 only, use '::'.
- host** (type: string)  
The host name to connect to.
- port** (type: maybe-port)  
The port number to connect to.
- pass** (type: maybe-string)  
Password to use to log in with the server. The password must have been hashed via the '**pounce -x**' command.
- join** (type: maybe-list-of-strings)  
The list of channels to join.
- mode** (type: maybe-string)  
The user mode.
- user** (type: maybe-string)  
To set the username. The default username is the same as the nickname.
- nick** (default: "pounce") (type: maybe-string)  
Set nickname to *nick*.
- real** (type: maybe-string)  
Set the real name. The default is **nick**.
- away** (type: maybe-string)  
The away status to use when no clients are connected and no other away status has been set.
- quit** (type: maybe-string)  
The message to use when quitting.
- no-names?** (type: maybe-boolean)  
Do not request 'NAMES' for each channel when a client connects. This avoids already connected clients receiving unsolicited responses but prevents new clients from populating user lists.
- queue-interval** (default: 200) (type: maybe-number)  
Set the server send queue interval in milliseconds. The queue is used to send automated messages from pounce to the server. Messages from clients are sent to the server directly.
- trust** (type: maybe-string)  
File name of a certificate to trust. When used, server name verification is disabled.

**client-cert** (type: maybe-string)

The file name of the TLS client. If the private key is in a separate file, it is loaded with **client-priv**. With **sasl-external?**, authenticate using SASL EXTERNAL. Certificates can be generated with ‘**pounce -g**’. For more details, refer to “Generating Client Certificates” in ‘**man 1 pounce**’.

**client-priv** (type: maybe-string)

The file name of the TLS client private key.

**sasl-plain** (type: maybe-pair)

A pair of the username and password in plain text to authenticate using SASL PLAIN. Since this method requires the account password in plain text, it is recommended to use CertFP instead with **sasl-external**.

**sasl-external?** (type: maybe-boolean)

Authenticate using SASL EXTERNAL, also known as CertFP. The TLS client certificate is loaded from **client-cert**.

## Quassel Service

Quassel (<https://quassel-irc.org/>) is a distributed IRC client, meaning that one or more clients can attach to and detach from the central core.

**quassel-service-type** [Variable]

This is the service type for the Quassel (<https://quassel-irc.org/>) IRC backend daemon. Its value is a **quassel-configuration** (see below).

**quassel-configuration** [Data Type]

This is the configuration for Quassel, with the following fields:

**quassel** (default: **quassel**)

The Quassel package to use.

**interface** (default: **"::,0.0.0.0"**)

**port** (default: **4242**)

Listen on the network interface(s) corresponding to the IPv4 or IPv6 interfaces specified in the comma delimited *interface*, on *port*.

**loglevel** (default: **"Info"**)

The level of logging desired. Accepted values are Debug, Info, Warning and Error.

## Snuik Service

The Snuik IRC bot (<https://gitlab.com/janneke/snuik>) can be run as a daemon to aid talking to users that are currently offline. With the (**gnu services messaging**) service, you can configure it by adding a **snuik-service-type** service to the **services** field of your operating system declaration.

**snuik-service-type** [Variable]

This is the type of the Snuik service, whose value is a **snuik-configuration** object.

Here is an example of a service and its configuration:

```
(service snuik-service-type
  (snuik-configuration
    (password-file "/var/run/snuik/.password.snuik")
    (channels '("#bootstrappable"
                "#dezyne"
                "#guix-risc-v"
                "#lilypond"))))
```

### snuik-configuration

[Data Type]

Available **snuik-configuration** fields are:

**snuik** (default: **snuik**) (type: package)

The Snuik package to use.

**server** (default: **"irc.libera.chat"**)

The IRC server to connect to.

**port** (default: **6665**) (type: integer)

Port number used by the IRC server.

**nick** (default: **"snuik"**) (type: string)

The nickname for snuik to use.

**password** (default: **#f**) (type: string)

The password to use when logging in.

**password-file** (default: **".password.<nick>"**)

The file to read the password from to use when logging in.

**channels** (default: **'("#botchat")'**)

The channels for snuik to join, a list of strings.

**extra-options** (default: **'()'**)

Extra options will be passed to **snuik**, please run **snuik --help** for more information.

## Mosquitto Service

Mosquitto (<https://mosquitto.org/>) is a lightweight message broker that implements the MQTT protocol versions 5.0, 3.1.1 and 3.1. It is suitable for use on all devices from low power single board computers to full servers.

### mosquitto-service-type

[Variable]

This is the service type for the Mosquitto (<https://mosquitto.org/>) MQTT broker. Its value is a **mosquitto-configuration** (see below).

### mosquitto-configuration

[Data Type]

This is the configuration for Mosquitto, with the following fields:

**package** (default: **mosquitto**)

The Mosquitto package to use.

**config-file** (default: `#f`)  
 The Mosquitto configuration file as a file-like object or the value `#f` to use the default configuration.

**user** (default: `"mosquitto"`)  
 Owner of the broker process.

**group** (default: `"mosquitto"`)  
 Owner's group of the broker process.

## Mollysocket Service

Mollysocket (<https://github.com/mollyim/mollysocket>) is a UnifiedPush provider for the messaging app Molly (<https://molly.im>), a fork of Signal.

**mollysocket-service-type** [Variable]  
 The Mollysocket service. Takes **mollysocket-configuration** as its value. By default, this service listens on port 8020 for any Molly client.

**mollysocket-configuration** [Data Type]  
 The configuration for Mollysocket, with the following fields:

**package** (default: `mollysocket`)  
 The Mollysocket package to use.

**host** (default: `"0.0.0.0"`)  
 Address for the webserver to listen on. It is recommended that this be set to localhost, and instead configure a TLS-supporting reverse proxy in front of it.

**port** (default: `8020`)  
 Port for the webserver to listen on.

**webserver** (default: `#t`)  
 Whether the webserver should be enabled. If disabled, the server is considered to be in airgapped mode, and clients must be configured likewise.

**allowed-endpoints** (default: `'("*")'`)  
 Allowlist of UnifiedPush distributor endpoints, represented as IP addresses or domain names. Addresses on the local network must be specified explicitly.

**allowed-uuids** (default: `'("*")'`)  
 Allowlist of Signal account UUIDs.

**db** (default: `"/var/lib/mollysocket/db.sqlite"`)  
 Path to the Mollysocket database. Will be created if it doesn't exist.

**vapid-key-file** (default: `"/var/lib/mollysocket/vapid.key"`)  
 Path to the VAPID key used to authorize requests to the UnifiedPush distributor. Will be created if it doesn't exist.

### 11.10.15 Telephony Services

The (`gnu services telephony`) module contains Guix service definitions for telephony services. Currently it provides the following services:

## Jami

`jami-service-type` [Variable]

The service type for running Jami (<https://jami.net/>) as a service. It takes a `jami-configuration` object as a value, documented below. This section describes how to configure a Jami server that can be used to host video (or audio) conferences, among other uses. The following example demonstrates how to specify Jami account archives (backups) to be provisioned automatically:

```
(service jami-service-type
  (jami-configuration
    (accounts
      (list (jami-account
              (archive "/etc/jami/unencrypted-account-1.gz"))
            (jami-account
              (archive "/etc/jami/unencrypted-account-2.gz"))))))■
```

When the `accounts` field is specified, the Jami account files of the service found under `/var/lib/jami` are recreated every time the service starts.

Jami accounts and their corresponding backup archives can be generated using the `jami` or `jami-gnome` Jami clients. The accounts should not be password-protected, but it is wise to ensure their files are only readable by `'root'`.

The next example shows how to declare that only some contacts should be allowed to communicate with a given account:

```
(service jami-service-type
  (jami-configuration
    (accounts
      (list (jami-account
              (archive "/etc/jami/unencrypted-account-1.gz")
              (peer-discovery? #t)
              (rendezvous-point? #t)
              (allowed-contacts
                '("1dbcb0f5f37324228235564b79f2b9737e9a008f"
                  "2dbcb0f5f37324228235564b79f2b9737e9a008f"))))))■
```

In this mode, only the declared `allowed-contacts` can initiate communication with the Jami account. This can be used, for example, with rendezvous point accounts to create a private video conferencing space.

To put the system administrator in full control of the conferences hosted on their system, the Jami service supports the following actions:

```
# herd doc jami list-actions
(list-accounts
 list-account-details
 list-banned-contacts
 list-contacts
 list-moderators
 add-moderator
 ban-contact
```

```
enable-account
disable-account)
```

The above actions aim to provide the most valuable actions for moderation purposes, not to cover the whole Jami API. Users wanting to interact with the Jami daemon from Guile may be interested in experimenting with the `(gnu build jami-service)` module, which powers the above Shepherd actions.

The `add-moderator` and `ban-contact` actions accept a contact *fingerprint* (40 characters long hash) as first argument and an account fingerprint or username as second argument:

```
# herd add-moderator jami 1dbcb0f5f37324228235564b79f2b9737e9a008f \
f3345f2775ddfe07a4b0d95daea111d15fbc1199

# herd list-moderators jami
Moderators for account f3345f2775ddfe07a4b0d95daea111d15fbc1199:
- 1dbcb0f5f37324228235564b79f2b9737e9a008f
```

In the case of `ban-contact`, the second username argument is optional; when omitted, the account is banned from all Jami accounts:

```
# herd ban-contact jami 1dbcb0f5f37324228235564b79f2b9737e9a008f

# herd list-banned-contacts jami
Banned contacts for account f3345f2775ddfe07a4b0d95daea111d15fbc1199:
- 1dbcb0f5f37324228235564b79f2b9737e9a008f
```

Banned contacts are also stripped from their moderation privileges.

The `disable-account` action allows to completely disconnect an account from the network, making it unreachable, while `enable-account` does the inverse. They accept a single account username or fingerprint as first argument:

```
# herd disable-account jami f3345f2775ddfe07a4b0d95daea111d15fbc1199

# herd list-accounts jami
The following Jami accounts are available:
- f3345f2775ddfe07a4b0d95daea111d15fbc1199 (dummy) [disabled]
```

The `list-account-details` action prints the detailed parameters of each accounts in the Recutils format, which means the `recsel` command can be used to select accounts of interest (see Section “Selection Expressions” in *GNU recutils manual*). Note that period characters (‘.’) found in the account parameter keys are mapped to underscores (‘\_’) in the output, to meet the requirements of the Recutils format. The following example shows how to print the account fingerprints for all accounts operating in the rendezvous point mode:

```
# herd list-account-details jami | \
recsel -p Account.username -e 'Account.rendezVous ~ "true"'
Account_username: f3345f2775ddfe07a4b0d95daea111d15fbc1199
```

The remaining actions should be self-explanatory.

The complete set of available configuration options is detailed below.

## **jami-configuration** [Data Type]

Available **jami-configuration** fields are:

**libjami** (default: **libjami**) (type: package)

The Jami daemon package to use.

**dbus** (default: **dbus-for-jami**) (type: package)

The D-Bus package to use to start the required D-Bus session.

**nss-certs** (default: **nss-certs**) (type: package)

The nss-certs package to use to provide TLS certificates.

**enable-logging?** (default: **#t**) (type: boolean)

Whether to enable logging to `/var/log/jami`.

**debug?** (default: **#f**) (type: boolean)

Whether to enable debug level messages.

**auto-answer?** (default: **#f**) (type: boolean)

Whether to force automatic answer to incoming calls.

**accounts** (type: maybe-jami-account-list)

A list of Jami accounts to be (re-)provisioned every time the Jami daemon service starts. When providing this field, the account directories under `/var/lib/jami/` are recreated every time the service starts, ensuring a consistent state.

## **jami-account** [Data Type]

Available **jami-account** fields are:

**archive** (type: string-or-computed-file)

The account archive (backup) file name of the account. This is used to provision the account when the service starts. The account archive should *not* be encrypted. It is highly recommended to make it readable only to the ‘root’ user (i.e., not in the store), to guard against leaking the secret key material of the Jami account it contains.

**allowed-contacts** (type: maybe-account-fingerprint-list)

The list of allowed contacts for the account, entered as their 40 characters long fingerprint. Messages or calls from accounts not in that list will be rejected. When left specified, the configuration of the account archive is used as-is with respect to contacts and public inbound calls/messaging allowance, which typically defaults to allow any contact to communicate with the account.

**moderators** (type: maybe-account-fingerprint-list)

The list of contacts that should have moderation privileges (to ban, mute, etc. other users) in rendezvous conferences, entered as their 40 characters long fingerprint. When left unspecified, the configuration of the account archive is used as-is with respect to moderation, which typically defaults to allow anyone to moderate.

**rendezvous-point?** (type: maybe-boolean)

Whether the account should operate in the rendezvous mode. In this mode, all the incoming audio/video calls are mixed into a conference. When left unspecified, the value from the account archive prevails.

**peer-discovery?** (type: maybe-boolean)

Whether peer discovery should be enabled. Peer discovery is used to discover other OpenDHT nodes on the local network, which can be useful to maintain communication between devices on such network even when the connection to the Internet has been lost. When left unspecified, the value from the account archive prevails.

**bootstrap-hostnames** (type: maybe-list-of-strings)

A list of hostnames or IPs pointing to OpenDHT nodes, that should be used to initially join the OpenDHT network. When left unspecified, the value from the account archive prevails.

**name-server-uri** (type: maybe-string)

The URI of the name server to use, that can be used to retrieve the account fingerprint for a registered username.

## Mumble server

This section describes how to set up and run a Mumble (<https://mumble.info>) server (formerly known as Murmur).

**mumble-server-service-type** [Variable]

This is the service to run a Mumble server. It takes a **mumble-server-configuration** object as its value, defined below.

**mumble-server-configuration** [Data Type]

The service type for the Mumble server. An example configuration can look like this:

```
(service mumble-server-service-type
  (mumble-server-configuration
    (welcome-text
      "Welcome to this Mumble server running on Guix!")
    (cert-required? #t) ;disallow text password logins
    (ssl-cert "/etc/certs/mumble.example.com/fullchain.pem")
    (ssl-key "/etc/certs/mumble.example.com/privkey.pem")))
```

After reconfiguring your system, you can manually set the mumble-server **SuperUser** password with the command that is printed during the activation phase.

It is recommended to register a normal Mumble user account and grant it admin or moderator rights. You can use the **mumble** client to login as new normal user, register yourself, and log out. For the next step login with the name **SuperUser** use the **SuperUser** password that you set previously, and grant your newly registered mumble user administrator or moderator rights and create some channels.

Available **mumble-server-configuration** fields are:

**package** (default: **mumble**)

Package that contains **bin/mumble-server**.



**user** (default: "mumble-server")  
User who will run the Mumble-Server server.

**group** (default: "mumble-server")  
Group of the user who will run the mumble-server server.

**port** (default: 64738)  
Port on which the server will listen.

**welcome-text** (default: "")  
Welcome text sent to clients when they connect.

**server-password** (default: "")  
Password the clients have to enter in order to connect.

**max-users** (default: 100)  
Maximum of users that can be connected to the server at once.

**max-user-bandwidth** (default: #f)  
Maximum voice traffic a user can send per second.

**database-file** (default: "/var/lib/mumble-server/db.sqlite")  
File name of the sqlite database. The service's user will become the owner of the directory.

**log-file** (default: "/var/log/mumble-server/mumble-server.log")  
File name of the log file. The service's user will become the owner of the directory.

**autoban-attempts** (default: 10)  
Maximum number of logins a user can make in **autoban-timeframe** without getting auto banned for **autoban-time**.

**autoban-timeframe** (default: 120)  
Timeframe for autoban in seconds.

**autoban-time** (default: 300)  
Amount of time in seconds for which a client gets banned when violating the autoban limits.

**opus-threshold** (default: 100)  
Percentage of clients that need to support opus before switching over to opus audio codec.

**channel-nesting-limit** (default: 10)  
How deep channels can be nested at maximum.

**channelname-regex** (default: #f)  
A string in form of a Qt regular expression that channel names must conform to.

**username-regex** (default: #f)  
A string in form of a Qt regular expression that user names must conform to.

**text-message-length** (default: 5000)  
Maximum size in bytes that a user can send in one text chat message.

**image-message-length** (default: (\* 128 1024))  
Maximum size in bytes that a user can send in one image message.

**cert-required?** (default: #f)  
If it is set to #t clients that use weak password authentication will not be accepted. Users must have completed the certificate wizard to join.

**remember-channel?** (default: #f)  
Should mumble-server remember the last channel each user was in when they disconnected and put them into the remembered channel when they rejoin.

**allow-html?** (default: #f)  
Should html be allowed in text messages, user comments, and channel descriptions.

**allow-ping?** (default: #f)  
Setting to true exposes the current user count, the maximum user count, and the server's maximum bandwidth per client to unauthenticated users. In the Mumble client, this information is shown in the Connect dialog. Disabling this setting will prevent public listing of the server.

**bonjour?** (default: #f)  
Should the server advertise itself in the local network through the bonjour protocol.

**send-version?** (default: #f)  
Should the mumble-server server version be exposed in ping requests.

**log-days** (default: 31)  
Mumble also stores logs in the database, which are accessible via RPC. The default is 31 days of months, but you can set this setting to 0 to keep logs forever, or -1 to disable logging to the database.

**obfuscate-ips?** (default: #t)  
Should logged ips be obfuscated to protect the privacy of users.

**ssl-cert** (default: #f)  
File name of the SSL/TLS certificate used for encrypted connections.  
(ssl-cert "/etc/certs/example.com/fullchain.pem")

**ssl-key** (default: #f)  
Filepath to the ssl private key used for encrypted connections.  
(ssl-key "/etc/certs/example.com/privkey.pem")

**ssl-dh-params** (default: #f)  
File name of a PEM-encoded file with Diffie-Hellman parameters for the SSL/TLS encryption. Alternatively you set it to "@ffdhe2048", "@ffdhe3072", "@ffdhe4096", "@ffdhe6144" or "@ffdhe8192" to use bundled parameters from RFC 7919.

**ssl-ciphers** (default: #f)  
The **ssl-ciphers** option chooses the cipher suites to make available for use in SSL/TLS.

This option is specified using OpenSSL cipher list notation (<https://www.openssl.org/docs/apps/ciphers.html#CIPHER-LIST-FORMAT>).

It is recommended that you try your cipher string using 'openssl ciphers <string>' before setting it here, to get a feel for which cipher suites you will get. After setting this option, it is recommended that you inspect your Mumble server log to ensure that Mumble is using the cipher suites that you expected it to.

**Note:** Changing this option may impact the backwards compatibility of your Mumble-Server server, and can remove the ability for older Mumble clients to be able to connect to it.

**public-registration** (default: #f)

Must be a <mumble-server-public-registration-configuration> record or #f.

You can optionally register your server in the public server list that the mumble client shows on startup. You cannot register your server if you have set a **server-password**, or set **allow-ping** to #f.

It might take a few hours until it shows up in the public list.

**file** (default: #f)

Optional alternative override for this configuration.

**mumble-server-public-registration-configuration** [Data Type]

Configuration for public registration of a mumble-server service.

**name** This is a display name for your server. Not to be confused with the hostname.

**password** A password to identify your registration. Subsequent updates will need the same password. Don't lose your password.

**url** This should be a <http://> or <https://> link to your web site.

**hostname** (default: #f)

By default your server will be listed by its IP address. If it is set your server will be linked by this host name instead.

**Deprecation notice:** Due to historical reasons, all of the above **mumble-server-**procedures are also exported with the **murmur-** prefix. It is recommended that you switch to using **mumble-server-** going forward.

### 11.10.16 File-Sharing Services

The (gnu services file-sharing) module provides services that assist with transferring files over peer-to-peer file-sharing networks.

#### Transmission Daemon Service

Transmission (<https://transmissionbt.com/>) is a flexible BitTorrent client that offers a variety of graphical and command-line interfaces. A **transmission-daemon-service-type** service provides Transmission's headless variant, **transmission-daemon**, as a system service, allowing users to share files via BitTorrent even when they are not logged in.

**transmission-daemon-service-type** [Variable]

The service type for the Transmission Daemon BitTorrent client. Its value must be a **transmission-daemon-configuration** object as in this example:

```
(service transmission-daemon-service-type
  (transmission-daemon-configuration
    ;; Restrict access to the RPC ("control") interface
    (rpc-authentication-required? #t)
    (rpc-username "transmission")
    (rpc-password
      (transmission-password-hash
        "transmission" ; desired password
        "uKd1uMs9")) ; arbitrary salt value

    ;; Accept requests from this and other hosts on the
    ;; local network
    (rpc-whitelist-enabled? #t)
    (rpc-whitelist '("::1" "127.0.0.1" "192.168.0.*"))

    ;; Limit bandwidth use during work hours
    (alt-speed-down (* 1024 2)) ; 2 MB/s
    (alt-speed-up 512) ; 512 kB/s

    (alt-speed-time-enabled? #t)
    (alt-speed-time-day 'weekdays)
    (alt-speed-time-begin
      (+ (* 60 8) 30)) ; 8:30 am
    (alt-speed-time-end
      (+ (* 60 (+ 12 5)) 30))) ; 5:30 pm
```

Once the service is started, users can interact with the daemon through its Web interface (at <http://localhost:9091/>) or by using the **transmission-remote** command-line tool, available in the **transmission** package. (Emacs users may want to also consider the **emacs-transmission** package.) Both communicate with the daemon through its remote procedure call (RPC) interface, which by default is available to all users on the system; you may wish to change this by assigning values to the **rpc-authentication-required?**, **rpc-username** and **rpc-password** settings, as shown in the example above and documented further below.

The value for **rpc-password** must be a password hash of the type generated and used by Transmission clients. This can be copied verbatim from an existing **settings.json** file, if another Transmission client is already being used. Otherwise, the **transmission-password-hash** and **transmission-random-salt** procedures provided by this module can be used to obtain a suitable hash value.

**transmission-password-hash** *password salt* [Procedure]

Returns a string containing the result of hashing *password* together with *salt*, in the format recognized by Transmission clients for their **rpc-password** configuration setting.

*salt* must be an eight-character string. The `transmission-random-salt` procedure can be used to generate a suitable salt value at random.

**transmission-random-salt** [Procedure]

Returns a string containing a random, eight-character salt value of the type generated and used by Transmission clients, suitable for passing to the `transmission-password-hash` procedure.

These procedures are accessible from within a Guile REPL started with the `guix repl` command (see Section 8.13 [Invoking `guix repl`], page 175). This is useful for obtaining a random salt value to provide as the second parameter to ‘`transmission-password-hash`’, as in this example session:

```
$ guix repl
scheme@(guix-user)> ,use (gnu services file-sharing)
scheme@(guix-user)> (transmission-random-salt)
$1 = "uKd1uMs9"
```

Alternatively, a complete password hash can be generated in a single step:

```
scheme@(guix-user)> (transmission-password-hash "transmission"
(transmission-random-salt))
$2 = "{c8bbc6d1740cd8dc819a6e25563b67812c1c19c9VtFPfdsX"
```

The resulting string can be used as-is for the value of `rpc-password`, allowing the password to be kept hidden even in the operating-system configuration.

Torrent files downloaded by the daemon are directly accessible only to users in the “transmission” user group, who receive read-only access to the directory specified by the `download-dir` configuration setting (and also the directory specified by `incomplete-dir`, if `incomplete-dir-enabled?` is `#t`). Downloaded files can be moved to another directory or deleted altogether using `transmission-remote` with its `--move` and `--remove-and-delete` options.

If the `watch-dir-enabled?` setting is set to `#t`, users in the “transmission” group are able also to place `.torrent` files in the directory specified by `watch-dir` to have the corresponding torrents added by the daemon. (The `trash-original-torrent-files?` setting controls whether the daemon deletes these files after processing them.)

Some of the daemon’s configuration settings can be changed temporarily by `transmission-remote` and similar tools. To undo these changes, use the service’s `reload` action to have the daemon reload its settings from disk:

```
# herd reload transmission-daemon
```

The full set of available configuration settings is defined by the `transmission-daemon-configuration` data type.

**transmission-daemon-configuration** [Data Type]

The data type representing configuration settings for Transmission Daemon. These correspond directly to the settings recognized by Transmission clients in their `settings.json` file.

Available `transmission-daemon-configuration` fields are:

**package transmission** [transmission-daemon-configuration parameter]  
The Transmission package to use.

**non-negative-integer stop-wait-period** [transmission-daemon-configuration parameter]  
The period, in seconds, to wait when stopping the service for **transmission-daemon** to exit before killing its process. This allows the daemon time to complete its house-keeping and send a final update to trackers as it shuts down. On slow hosts, or hosts with a slow network connection, this value may need to be increased.  
Defaults to '10'.

**string download-dir** [transmission-daemon-configuration parameter]  
The directory to which torrent files are downloaded.  
Defaults to `"/var/lib/transmission-daemon/downloads"`.

**boolean incomplete-dir-enabled?** [transmission-daemon-configuration parameter]  
If **#t**, files will be held in **incomplete-dir** while their torrent is being downloaded, then moved to **download-dir** once the torrent is complete. Otherwise, files for all torrents (including those still being downloaded) will be placed in **download-dir**.  
Defaults to **#f**.

**maybe-string incomplete-dir** [transmission-daemon-configuration parameter]  
The directory in which files from incompletely downloaded torrents will be held when **incomplete-dir-enabled?** is **#t**.  
Defaults to **'disabled'**.

**umask umask** [transmission-daemon-configuration parameter]  
The file mode creation mask used for downloaded files. (See the **umask** man page for more information.)  
Defaults to **'18'**.

**boolean rename-partial-files?** [transmission-daemon-configuration parameter]  
When **#t**, ".part" is appended to the name of partially downloaded files.  
Defaults to **#t**.

**preallocation-mode preallocation** [transmission-daemon-configuration parameter]  
The mode by which space should be preallocated for downloaded files, one of **none**, **fast** (or **sparse**) and **full**. Specifying **full** will minimize disk fragmentation at a cost to file-creation speed.  
Defaults to **'fast'**.

**boolean watch-dir-enabled?** [transmission-daemon-configuration parameter]  
If **#t**, the directory specified by **watch-dir** will be watched for new **.torrent** files and the torrents they describe added automatically (and the original files removed, if **trash-original-torrent-files?** is **#t**).  
Defaults to **#f**.

**maybe-string watch-dir** [transmission-daemon-configuration parameter]  
The directory to be watched for `.torrent` files indicating new torrents to be added, when `watch-dir-enabled` is `#t`.  
Defaults to `'disabled'`.

**boolean trash-original-torrent-files?** [transmission-daemon-configuration parameter]  
When `#t`, `.torrent` files will be deleted from the watch directory once their torrent has been added (see `watch-directory-enabled?`).  
Defaults to `'#f'`.

**boolean speed-limit-down-enabled?** [transmission-daemon-configuration parameter]  
When `#t`, the daemon's download speed will be limited to the rate specified by `speed-limit-down`.  
Defaults to `'#f'`.

**non-negative-integer speed-limit-down** [transmission-daemon-configuration parameter]  
The default global-maximum download speed, in kilobytes per second.  
Defaults to `'100'`.

**boolean speed-limit-up-enabled?** [transmission-daemon-configuration parameter]  
When `#t`, the daemon's upload speed will be limited to the rate specified by `speed-limit-up`.  
Defaults to `'#f'`.

**non-negative-integer speed-limit-up** [transmission-daemon-configuration parameter]  
The default global-maximum upload speed, in kilobytes per second.  
Defaults to `'100'`.

**boolean alt-speed-enabled?** [transmission-daemon-configuration parameter]  
When `#t`, the alternate speed limits `alt-speed-down` and `alt-speed-up` are used (in place of `speed-limit-down` and `speed-limit-up`, if they are enabled) to constrain the daemon's bandwidth usage. This can be scheduled to occur automatically at certain times during the week; see `alt-speed-time-enabled?`.  
Defaults to `'#f'`.

**non-negative-integer alt-speed-down** [transmission-daemon-configuration parameter]  
The alternate global-maximum download speed, in kilobytes per second.  
Defaults to `'50'`.

**non-negative-integer alt-speed-up** [transmission-daemon-configuration parameter]  
The alternate global-maximum upload speed, in kilobytes per second.  
Defaults to `'50'`.

**boolean** [transmission-daemon-configuration parameter]

**alt-speed-time-enabled?**

When **#t**, the alternate speed limits **alt-speed-down** and **alt-speed-up** will be enabled automatically during the periods specified by **alt-speed-time-day**, **alt-speed-time-begin** and **alt-time-speed-end**.

Defaults to **'#f'**.

**day-list** [transmission-daemon-configuration parameter]

**alt-speed-time-day**

The days of the week on which the alternate-speed schedule should be used, specified either as a list of days (**sunday**, **monday**, and so on) or using one of the symbols **weekdays**, **weekends** or **all**.

Defaults to **'all'**.

**non-negative-integer** [transmission-daemon-configuration parameter]

**alt-speed-time-begin**

The time of day at which to enable the alternate speed limits, expressed as a number of minutes since midnight.

Defaults to **'540'**.

**non-negative-integer** [transmission-daemon-configuration parameter]

**alt-speed-time-end**

The time of day at which to disable the alternate speed limits, expressed as a number of minutes since midnight.

Defaults to **'1020'**.

**string bind-address-ipv4** [transmission-daemon-configuration parameter]

The IP address at which to listen for peer connections, or **"0.0.0.0"** to listen at all available IP addresses.

Defaults to **"0.0.0.0"**.

**string bind-address-ipv6** [transmission-daemon-configuration parameter]

The IPv6 address at which to listen for peer connections, or **"::"** to listen at all available IPv6 addresses.

Defaults to **"::"**.

**boolean** [transmission-daemon-configuration parameter]

**peer-port-random-on-start?**

If **#t**, when the daemon starts it will select a port at random on which to listen for peer connections, from the range specified (inclusively) by **peer-port-random-low** and **peer-port-random-high**. Otherwise, it listens on the port specified by **peer-port**.

Defaults to **'#f'**.

**port-number** [transmission-daemon-configuration parameter]

**peer-port-random-low**

The lowest selectable port number when **peer-port-random-on-start?** is **#t**.

Defaults to **'49152'**.



**port-number** [transmission-daemon-configuration parameter]

**peer-port-random-high**

The highest selectable port number when **peer-port-random-on-start** is **#t**.

Defaults to '65535'.

**port-number peer-port** [transmission-daemon-configuration parameter]

The port on which to listen for peer connections when **peer-port-random-on-start?** is **#f**.

Defaults to '51413'.

**boolean** [transmission-daemon-configuration parameter]

**port-forwarding-enabled?**

If **#t**, the daemon will attempt to configure port-forwarding on an upstream gateway automatically using UPnP and NAT-PMP.

Defaults to '**#t**'.

**encryption-mode encryption** [transmission-daemon-configuration parameter]

The encryption preference for peer connections, one of **prefer-unencrypted-connections**, **prefer-encrypted-connections** or **require-encrypted-connections**.

Defaults to '**prefer-encrypted-connections**'.

**maybe-string** [transmission-daemon-configuration parameter]

**peer-congestion-algorithm**

The TCP congestion-control algorithm to use for peer connections, specified using a string recognized by the operating system in calls to **setsockopt**. When left unspecified, the operating-system default is used.

Note that on GNU/Linux systems, the kernel must be configured to allow processes to use a congestion-control algorithm not in the default set; otherwise, it will deny these requests with "Operation not permitted". To see which algorithms are available on your system and which are currently permitted for use, look at the contents of the files **tcp\_available\_congestion\_control** and **tcp\_allowed\_congestion\_control** in the **/proc/sys/net/ipv4** directory.

As an example, to have Transmission Daemon use the TCP Low Priority congestion-control algorithm (<http://www-ece.rice.edu/networks/TCP-LP/>), you'll need to modify your kernel configuration to build in support for the algorithm, then update your operating-system configuration to allow its use by adding a **sysctl-service-type** service (or updating the existing one's configuration) with lines like the following:

```
(service sysctl-service-type
  (sysctl-configuration
    (settings
      ("net.ipv4.tcp_allowed_congestion_control" .
        "reno cubic lp"))))
```

The Transmission Daemon configuration can then be updated with

```
(peer-congestion-algorithm "lp")
```

and the system reconfigured to have the changes take effect.

Defaults to '**disabled**'.

- tcp-type-of-service** [transmission-daemon-configuration parameter]  
**peer-socket-tos**  
The type of service to request in outgoing TCP packets, one of **default**, **low-cost**, **throughput**, **low-delay** and **reliability**.  
Defaults to **'default'**.
- non-negative-integer** [transmission-daemon-configuration parameter]  
**peer-limit-global**  
The global limit on the number of connected peers.  
Defaults to **'200'**.
- non-negative-integer** [transmission-daemon-configuration parameter]  
**peer-limit-per-torrent**  
The per-torrent limit on the number of connected peers.  
Defaults to **'50'**.
- non-negative-integer** [transmission-daemon-configuration parameter]  
**upload-slots-per-torrent**  
The maximum number of peers to which the daemon will upload data simultaneously for each torrent.  
Defaults to **'14'**.
- non-negative-integer** [transmission-daemon-configuration parameter]  
**peer-id-ttl-hours**  
The maximum lifespan, in hours, of the peer ID associated with each public torrent before it is regenerated.  
Defaults to **'6'**.
- boolean** **blocklist-enabled?** [transmission-daemon-configuration parameter]  
When **#t**, the daemon will ignore peers mentioned in the blocklist it has most recently downloaded from **blocklist-url**.  
Defaults to **'#f'**.
- maybe-string** **blocklist-url** [transmission-daemon-configuration parameter]  
The URL of a peer blocklist (in P2P-plaintext or eMule **.dat** format) to be periodically downloaded and applied when **blocklist-enabled?** is **#t**.  
Defaults to **'disabled'**.
- boolean** [transmission-daemon-configuration parameter]  
**download-queue-enabled?**  
If **#t**, the daemon will be limited to downloading at most **download-queue-size** non-stalled torrents simultaneously.  
Defaults to **'#t'**.
- non-negative-integer** [transmission-daemon-configuration parameter]  
**download-queue-size**  
The size of the daemon's download queue, which limits the number of non-stalled torrents it will download at any one time when **download-queue-enabled?** is **#t**.  
Defaults to **'5'**.

**boolean** [transmission-daemon-configuration parameter]  
**seed-queue-enabled?**

If **#t**, the daemon will be limited to seeding at most **seed-queue-size** non-stalled torrents simultaneously.

Defaults to '**#f**'.

**non-negative-integer** [transmission-daemon-configuration parameter]  
**seed-queue-size**

The size of the daemon's seed queue, which limits the number of non-stalled torrents it will seed at any one time when **seed-queue-enabled?** is **#t**.

Defaults to '**10**'.

**boolean** [transmission-daemon-configuration parameter]  
**queue-stalled-enabled?**

When **#t**, the daemon will consider torrents for which it has not shared data in the past **queue-stalled-minutes** minutes to be stalled and not count them against its **download-queue-size** and **seed-queue-size** limits.

Defaults to '**#t**'.

**non-negative-integer** [transmission-daemon-configuration parameter]  
**queue-stalled-minutes**

The maximum period, in minutes, a torrent may be idle before it is considered to be stalled, when **queue-stalled-enabled?** is **#t**.

Defaults to '**30**'.

**boolean** [transmission-daemon-configuration parameter]  
**ratio-limit-enabled?**

When **#t**, a torrent being seeded will automatically be paused once it reaches the ratio specified by **ratio-limit**.

Defaults to '**#f**'.

**non-negative-rational** [transmission-daemon-configuration parameter]  
**ratio-limit**

The ratio at which a torrent being seeded will be paused, when **ratio-limit-enabled?** is **#t**.

Defaults to '**2.0**'.

**boolean** [transmission-daemon-configuration parameter]  
**idle-seeding-limit-enabled?**

When **#t**, a torrent being seeded will automatically be paused once it has been idle for **idle-seeding-limit** minutes.

Defaults to '**#f**'.

**non-negative-integer** [transmission-daemon-configuration parameter]  
**idle-seeding-limit**

The maximum period, in minutes, a torrent being seeded may be idle before it is paused, when **idle-seeding-limit-enabled?** is **#t**.

Defaults to '**30**'.

**boolean dht-enabled?** [transmission-daemon-configuration parameter]  
Enable the distributed hash table (DHT) protocol ([http://bittorrent.org/beps/bep\\_0005.html](http://bittorrent.org/beps/bep_0005.html)), which supports the use of trackerless torrents.  
Defaults to '#t'.

**boolean lpd-enabled?** [transmission-daemon-configuration parameter]  
Enable local peer discovery ([https://en.wikipedia.org/wiki/Local\\_Peer\\_Discovery](https://en.wikipedia.org/wiki/Local_Peer_Discovery)) (LPD), which allows the discovery of peers on the local network and may reduce the amount of data sent over the public Internet.  
Defaults to '#f'.

**boolean pex-enabled?** [transmission-daemon-configuration parameter]  
Enable peer exchange ([https://en.wikipedia.org/wiki/Peer\\_exchange](https://en.wikipedia.org/wiki/Peer_exchange)) (PEX), which reduces the daemon's reliance on external trackers and may improve its performance.  
Defaults to '#t'.

**boolean utp-enabled?** [transmission-daemon-configuration parameter]  
Enable the micro transport protocol ([http://bittorrent.org/beps/bep\\_0029.html](http://bittorrent.org/beps/bep_0029.html)) (uTP), which aims to reduce the impact of BitTorrent traffic on other users of the local network while maintaining full utilization of the available bandwidth.  
Defaults to '#t'.

**boolean rpc-enabled?** [transmission-daemon-configuration parameter]  
If #t, enable the remote procedure call (RPC) interface, which allows remote control of the daemon via its Web interface, the `transmission-remote` command-line client, and similar tools.  
Defaults to '#t'.

**string rpc-bind-address** [transmission-daemon-configuration parameter]  
The IP address at which to listen for RPC connections, or "0.0.0.0" to listen at all available IP addresses.  
Defaults to "0.0.0.0".

**port-number rpc-port** [transmission-daemon-configuration parameter]  
The port on which to listen for RPC connections.  
Defaults to '9091'.

**string rpc-url** [transmission-daemon-configuration parameter]  
The path prefix to use in the RPC-endpoint URL.  
Defaults to "/transmission/".

**boolean rpc-authentication-required?** [transmission-daemon-configuration parameter]  
When #t, clients must authenticate (see `rpc-username` and `rpc-password`) when using the RPC interface. Note this has the side effect of disabling host-name whitelisting (see `rpc-host-whitelist-enabled?`).  
Defaults to '#f'.

**maybe-string rpc-username** [transmission-daemon-configuration parameter]  
 The username required by clients to access the RPC interface when `rpc-authentication-required?` is `#t`.

Defaults to 'disabled'.

**maybe-transmission-password-hash** [transmission-daemon-configuration parameter]  
**rpc-password**

The password required by clients to access the RPC interface when `rpc-authentication-required?` is `#t`. This must be specified using a password hash in the format recognized by Transmission clients, either copied from an existing `settings.json` file or generated using the `transmission-password-hash` procedure.

Defaults to 'disabled'.

**boolean** [transmission-daemon-configuration parameter]  
**rpc-whitelist-enabled?**

When `#t`, RPC requests will be accepted only when they originate from an address specified in `rpc-whitelist`.

Defaults to '#t'.

**string-list rpc-whitelist** [transmission-daemon-configuration parameter]  
 The list of IP and IPv6 addresses from which RPC requests will be accepted when `rpc-whitelist-enabled?` is `#t`. Wildcards may be specified using '\*'.

Defaults to '("127.0.0.1" ":::1)".

**boolean** [transmission-daemon-configuration parameter]  
**rpc-host-whitelist-enabled?**

When `#t`, RPC requests will be accepted only when they are addressed to a host named in `rpc-host-whitelist`. Note that requests to "localhost" or "localhost.", or to a numeric address, are always accepted regardless of these settings.

Note also this functionality is disabled when `rpc-authentication-required?` is `#t`.

Defaults to '#t'.

**string-list** [transmission-daemon-configuration parameter]  
**rpc-host-whitelist**

The list of host names recognized by the RPC server when `rpc-host-whitelist-enabled?` is `#t`.

Defaults to '()'.

**message-level** [transmission-daemon-configuration parameter]  
**message-level**

The minimum severity level of messages to be logged (to `/var/log/transmission.log`) by the daemon, one of `none` (no logging), `error`, `info` and `debug`.

Defaults to 'info'.

**boolean** `start-added-torrents?` [transmission-daemon-configuration parameter]

When `#t`, torrents are started as soon as they are added; otherwise, they are added in “paused” state.

Defaults to ‘`#t`’.

**boolean** `script-torrent-done-enabled?` [transmission-daemon-configuration parameter]

When `#t`, the script specified by `script-torrent-done-filename` will be invoked each time a torrent completes.

Defaults to ‘`#f`’.

**maybe-file-object** `script-torrent-done-filename` [transmission-daemon-configuration parameter]

A file name or file-like object specifying a script to run each time a torrent completes, when `script-torrent-done-enabled?` is `#t`.

Defaults to ‘`disabled`’.

**boolean** `scrape-paused-torrents-enabled?` [transmission-daemon-configuration parameter]

When `#t`, the daemon will scrape trackers for a torrent even when the torrent is paused.

Defaults to ‘`#t`’.

**non-negative-integer** `cache-size-mb` [transmission-daemon-configuration parameter]

The amount of memory, in megabytes, to allocate for the daemon’s in-memory cache. A larger value may increase performance by reducing the frequency of disk I/O.

Defaults to ‘`4`’.

**boolean** `prefetch-enabled?` [transmission-daemon-configuration parameter]

When `#t`, the daemon will try to improve I/O performance by hinting to the operating system which data is likely to be read next from disk to satisfy requests from peers.

Defaults to ‘`#t`’.

### 11.10.17 Monitoring Services

#### Darkstat Service

Darkstat is a packet sniffer that captures network traffic, calculates statistics about usage, and serves reports over HTTP.

**darkstat-service-type** [Variable]

This is the service type for the darkstat (<https://unix4lyfe.org/darkstat/>) service, its value must be a `darkstat-configuration` record as in this example:

```
(service darkstat-service-type
  (darkstat-configuration
    (interface "eno1")))
```

**darkstat-configuration** [Data Type]

Data type representing the configuration of **darkstat**.

**package** (default: **darkstat**)

The darkstat package to use.

**interface**

Capture traffic on the specified network interface.

**port** (default: "667")

Bind the web interface to the specified port.

**bind-address** (default: "127.0.0.1")

Bind the web interface to the specified address.

**base** (default: "/")

Specify the path of the base URL. This can be useful if **darkstat** is accessed via a reverse proxy.

## Prometheus Node Exporter Service

The Prometheus “node exporter” makes hardware and operating system statistics provided by the Linux kernel available for the Prometheus monitoring system. This service should be deployed on all physical nodes and virtual machines, where monitoring these statistics is desirable.

**prometheus-node-exporter-service-type** [Variable]

This is the service type for the prometheus-node-exporter ([https://github.com/prometheus/node\\_exporter/](https://github.com/prometheus/node_exporter/)) service, its value must be a **prometheus-node-exporter-configuration**.

(service prometheus-node-exporter-service-type)

**prometheus-node-exporter-configuration** [Data Type]

Data type representing the configuration of **node\_exporter**.

**package** (default: **go-github-com-prometheus-node-exporter**)

The prometheus-node-exporter package to use.

**web-listen-address** (default: ":9100")

Bind the web interface to the specified address.

**textfile-directory** (default: **"/var/lib/prometheus/node-exporter"**)

This directory can be used to export metrics specific to this machine. Files containing metrics in the text format, with the filename ending in **.prom** should be placed in this directory.

**extra-options** (default: '()')

Extra options to pass to the Prometheus node exporter.

## vnStat Network Traffic Monitor

vnStat is a network traffic monitor that uses interface statistics provided by the kernel rather than traffic sniffing. This makes it a light resource monitor, regardless of network traffic rate.

**vnstat-service-type** [Variable]

This is the service type for the vnStat (<https://humdi.net/vnstat/>) daemon and accepts a `vnstat-configuration` value.

The following example will configure the service with default values:

```
(service vnstat-service-type)
```

**vnstat-configuration** [Data Type]

Available `vnstat-configuration` fields are:

**package** (default: `vnstat`) (type: file-like)

The vnstat package.

**database-directory** (default: `"/var/lib/vnstat"`) (type: string)

Specifies the directory where the database is to be stored. A full path must be given and a leading `'/'` isn't required.

**5-minute-hours** (default: `48`) (type: maybe-integer)

Data retention duration for the 5 minute resolution entries. The configuration defines for how many past hours entries will be stored. Set to `-1` for unlimited entries or to `0` to disable the data collection of this resolution.

**64bit-interface-counters** (default: `-2`) (type: maybe-integer)

Select interface counter handling. Set to `1` for defining that all interfaces use 64-bit counters on the kernel side and `0` for defining 32-bit counter. Set to `-1` for using the old style logic used in earlier versions where counter values within 32-bits are assumed to be 32-bit and anything larger is assumed to be a 64-bit counter. This may produce false results if a 64-bit counter is reset within the 32-bits. Set to `-2` for using automatic detection based on available kernel datastructures.

**always-add-new-interfaces?** (default: `#t`) (type: maybe-boolean)

Enable or disable automatic creation of new database entries for interfaces not currently in the database even if the database file already exists when the daemon is started. New database entries will also get created for new interfaces seen while the daemon is running. Pseudo interfaces `'lo'`, `'lo0'` and `'sit0'` are always excluded from getting added.

**bandwidth-detection?** (default: `#t`) (type: maybe-boolean)

Try to automatically detect *max-bandwidth* value for each monitored interface. Mostly only ethernet interfaces support this feature. *max-bandwidth* will be used as fallback value if detection fails. Any interface specific *max-BW* configuration will disable the detection for the specified interface. In Linux, the detection is disabled for tun interfaces due to the Linux kernel always reporting 10 Mbit regardless of the used real interface.

**bandwidth-detection-interval** (default: `5`) (type: maybe-integer)

How often in minutes interface specific detection of *max-bandwidth* is done for detecting possible changes when *bandwidth-detection* is enabled. Can be disabled by setting to `0`. Value range: `'0'..'30'`



- boot-variation** (default: 15) (type: maybe-integer)  
Time in seconds how much the boot time reported by system kernel can vary between updates. Value range: '0'..'300'
- check-disk-space?** (default: #t) (type: maybe-boolean)  
Enable or disable the availability check of at least some free disk space before a database write.
- create-directories?** (default: #t) (type: maybe-boolean)  
Enable or disable the creation of directories when a configured path doesn't exist. This includes *database-directory*.
- daemon-group** (type: maybe-user-group)  
Specify the group to which the daemon process should switch during startup. Set to %unset-value to disable group switching.
- daemon-user** (type: maybe-user-account)  
Specify the user to which the daemon process should switch during startup. Set to %unset-value to disable user switching.
- daily-days** (default: 62) (type: maybe-integer)  
Data retention duration for the one day resolution entries. The configuration defines for how many past days entries will be stored. Set to -1 for unlimited entries or to 0 to disable the data collection of this resolution.
- database-synchronous** (default: -1) (type: maybe-integer)  
Change the setting of the SQLite "synchronous" flag which controls how much care is taken to ensure disk writes have fully completed when writing data to the database before continuing other actions. Higher values take extra steps to ensure data safety at the cost of slower performance. A value of 0 will result in all handling being left to the filesystem itself. Set to -1 to select the default value according to database mode controlled by *database-write-ahead-logging* setting. See SQLite documentation for more details regarding values from 1 to 3. Value range: '-1'..'3'
- database-write-ahead-logging?** (default: #f) (type: maybe-boolean)  
Enable or disable SQLite Write-Ahead Logging mode for the database. See SQLite documentation for more details and note that support for read-only operations isn't available in older SQLite versions.
- hourly-days** (default: 4) (type: maybe-integer)  
Data retention duration for the one hour resolution entries. The configuration defines for how many past days entries will be stored. Set to -1 for unlimited entries or to 0 to disable the data collection of this resolution.
- log-file** (type: maybe-string)  
Specify log file path and name to be used if *use-logging* is set to 1.
- max-bandwidth** (type: maybe-integer)  
Maximum bandwidth for all interfaces. If the interface specific traffic exceeds the given value then the data is assumed to be invalid and rejected. Set to 0 in order to disable the feature. Value range: '0'..'50000'

**max-bw** (type: maybe-alist)

Same as *max-bandwidth* but can be used for setting individual limits for selected interfaces. This is an association list of interfaces as strings to integer values. For example,

```
(max-bw `(("eth0" . 15000)
          ("ppp0" . 10000)))
```

*bandwidth-detection* is disabled on an interface specific level for each *max-bw* configuration. Value range: '0'..'50000'

**monthly-months** (default: 25) (type: maybe-integer)

Data retention duration for the one month resolution entries. The configuration defines for how many past months entries will be stored. Set to -1 for unlimited entries or to 0 to disable the data collection of this resolution.

**month-rotate** (default: 1) (type: maybe-integer)

Day of month that months are expected to change. Usually set to 1 but can be set to alternative values for example for tracking monthly billed traffic where the billing period doesn't start on the first day. For example, if set to 7, days of February up to and including the 6th will count for January. Changing this option will not cause existing data to be recalculated. Value range: '1'..'28'

**month-rotate-affects-years?** (default: #f) (type: maybe-boolean)

Enable or disable *month-rotate* also affecting yearly data. Applicable only when *month-rotate* has a value greater than one.

**offline-save-interval** (default: 30) (type: maybe-integer)

How often in minutes cached interface data is saved to file when all monitored interfaces are offline. Value range: *save-interval*..'60'

**pid-file** (default: "/var/run/vnstatd.pid") (type: maybe-string)

Specify pid file path and name to be used.

**poll-interval** (default: 5) (type: maybe-integer)

How often in seconds interfaces are checked for status changes. Value range: '2'..'60'

**rescan-database-on-save?** (type: maybe-boolean)

Automatically discover added interfaces from the database and start monitoring. The rescan is done every *save-interval* or *offline-save-interval* minutes depending on the current activity state.

**save-interval** (default: 5) (type: maybe-integer)

How often in minutes cached interface data is saved to file. Value range: ( *update-interval* / 60 )..'60'

**save-on-status-change?** (default: #t) (type: maybe-boolean)

Enable or disable the additional saving to file of cached interface data when the availability of an interface changes, i.e., when an interface goes offline or comes online.

- time-sync-wait** (default: 5) (type: maybe-integer)  
How many minutes to wait during daemon startup for system clock to sync if most recent database update appears to be in the future. This may be needed in systems without a real-time clock (RTC) which require some time after boot to query and set the correct time. 0 = wait disabled. Value range: '0'..'60'
- top-day-entries** (default: 20) (type: maybe-integer)  
Data retention duration for the top day entries. The configuration defines how many of the past top day entries will be stored. Set to -1 for unlimited entries or to 0 to disable the data collection of this resolution.
- trafficless-entries?** (default: #t) (type: maybe-boolean)  
Create database entries even when there is no traffic during the entry's time period.
- update-file-owner?** (default: #t) (type: maybe-boolean)  
Enable or disable the update of file ownership during daemon process startup. During daemon startup, only database, log and pid files will be modified if the user or group change feature ( *daemon-user* or *daemon-group* ) is enabled and the files don't match the requested user or group. During manual database creation, this option will cause file ownership to be inherited from the database directory if the directory already exists. This option only has effect when the process is started as root or via sudo.
- update-interval** (default: 20) (type: maybe-integer)  
How often in seconds the interface data is updated. Value range: *poll-interval*..'300'
- use-logging** (default: 2) (type: maybe-integer)  
Enable or disable logging. Accepted values are: 0 = disabled, 1 = logfile and 2 = syslog.
- use-utc?** (type: maybe-boolean)  
Enable or disable using UTC as timezone in the database for all entries. When enabled, all entries added to the database will use UTC regardless of the configured system timezone. When disabled, the configured system timezone will be used. Changing this setting will not result in already existing data to be modified.
- yearly-years** (default: -1) (type: maybe-integer)  
Data retention duration for the one year resolution entries. The configuration defines for how many past years entries will be stored. Set to -1 for unlimited entries or to 0 to disable the data collection of this resolution.

## Zabbix server

Zabbix is a high performance monitoring system that can collect data from a variety of sources and provide the results in a web-based interface. Alerting and reporting is built-in, as well as *templates* for common operating system metrics such as network utilization, CPU load, and disk space consumption.

This service provides the central Zabbix monitoring service; you also need [zabbix-front-end], page 482, to configure Zabbix and display results, and optionally [zabbix-agent], page 481, on machines that should be monitored (other data sources are supported, such as [prometheus-node-exporter], page 475).

**zabbix-server-service-type** [Variable]

This is the service type for the Zabbix server service. Its value must be a **zabbix-server-configuration** record, shown below.

**zabbix-server-configuration** [Data Type]

Available **zabbix-server-configuration** fields are:

**zabbix-server** (default: **zabbix-server**) (type: file-like)

The zabbix-server package.

**user** (default: **"zabbix"**) (type: string)

User who will run the Zabbix server.

**group** (default: **"zabbix"**) (type: string)

Group who will run the Zabbix server.

**db-host** (default: **"127.0.0.1"**) (type: string)

Database host name.

**db-name** (default: **"zabbix"**) (type: string)

Database name.

**db-user** (default: **"zabbix"**) (type: string)

Database user.

**db-password** (default: **""**) (type: string)

Database password. Please, use **include-files** with **DBPassword=SECRET** inside a specified file instead.

**db-port** (default: **5432**) (type: number)

Database port.

**log-type** (default: **""**) (type: string)

Specifies where log messages are written to:

- **system** - syslog.
- **file** - file specified with **log-file** parameter.
- **console** - standard output.

**log-file** (default: **"/var/log/zabbix/server.log"**) (type: string)

Log file name for **log-type** file parameter.

**pid-file** (default: **"/var/run/zabbix/zabbix\_server.pid"**) (type: string)

Name of PID file.

**ssl-ca-location** (default: **"/etc/ssl/certs/ca-certificates.crt"**) (type: string)

The location of certificate authority (CA) files for SSL server certificate verification.

**ssl-cert-location** (default: `"/etc/ssl/certs"`) (type: string)  
 Location of SSL client certificates.

**extra-options** (default: `""`) (type: extra-options)  
 Extra options will be appended to Zabbix server configuration file.

**include-files** (default: `'()`) (type: include-files)  
 You may include individual files or all files in a directory in the configuration file.

## Zabbix agent

The Zabbix agent gathers information about the running system for the Zabbix monitoring server. It has a variety of built-in checks, and can be extended with custom *user parameters* (<https://www.zabbix.com/documentation/current/en/manual/config/items/userparameters>).

**zabbix-agent-service-type** [Variable]  
 This is the service type for the Zabbix agent service. Its value must be a **zabbix-agent-configuration** record, shown below.

**zabbix-agent-configuration** [Data Type]  
 Available **zabbix-agent-configuration** fields are:

**zabbix-agent** (default: `zabbix-agentd`) (type: file-like)  
 The zabbix-agent package.

**user** (default: `"zabbix"`) (type: string)  
 User who will run the Zabbix agent.

**group** (default: `"zabbix"`) (type: string)  
 Group who will run the Zabbix agent.

**hostname** (default: `""`) (type: string)  
 Unique, case sensitive hostname which is required for active checks and must match hostname as configured on the server.

**log-type** (default: `""`) (type: string)  
 Specifies where log messages are written to:

- **system** - syslog.
- **file** - file specified with **log-file** parameter.
- **console** - standard output.

**log-file** (default: `"/var/log/zabbix/agent.log"`) (type: string)  
 Log file name for **log-type** file parameter.

**pid-file** (default: `"/var/run/zabbix/zabbix_agent.pid"`) (type: string)  
 Name of PID file.

**server** (default: `'("127.0.0.1")`) (type: list)  
 List of IP addresses, optionally in CIDR notation, or hostnames of Zabbix servers and Zabbix proxies. Incoming connections will be accepted only from the hosts listed here.

**server-active** (default: '("127.0.0.1")') (type: list)  
 List of IP:port (or hostname:port) pairs of Zabbix servers and Zabbix proxies for active checks. If port is not specified, default port is used. If this parameter is not specified, active checks are disabled.

**extra-options** (default: "") (type: extra-options)  
 Extra options will be appended to Zabbix server configuration file.

**include-files** (default: '()') (type: include-files)  
 You may include individual files or all files in a directory in the configuration file.

## Zabbix front-end

The Zabbix front-end provides a web interface to Zabbix. It does not need to run on the same machine as the Zabbix server. This service works by extending the [PHP-FPM], page 509, and [NGINX], page 497, services with the configuration necessary for loading the Zabbix user interface.

**zabbix-front-end-service-type** [Variable]  
 This is the service type for the Zabbix web frontend. Its value must be a **zabbix-front-end-configuration** record, shown below.

**zabbix-front-end-configuration** [Data Type]

Available **zabbix-front-end-configuration** fields are:

**zabbix-server** (default: **zabbix-server**) (type: file-like)  
 The Zabbix server package to use.

**nginx** (default: '()') (type: list)  
 List of [nginx-server-configuration], page 500, blocks for the Zabbix front-end. When empty, a default that listens on port 80 is used.

**db-host** (default: "localhost") (type: string)  
 Database host name.

**db-port** (default: 5432) (type: number)  
 Database port.

**db-name** (default: "zabbix") (type: string)  
 Database name.

**db-user** (default: "zabbix") (type: string)  
 Database user.

**db-password** (default: "") (type: string)  
 Database password. Please, use **db-secret-file** instead.

**db-secret-file** (default: "") (type: string)  
 Secret file which will be appended to **zabbix.conf.php** file. This file contains credentials for use by Zabbix front-end. You are expected to create it manually.

**zabbix-host** (default: "localhost") (type: string)  
 Zabbix server hostname.

**zabbix-port** (default: 10051) (type: number)  
Zabbix server port.

### 11.10.18 Kerberos Services

The (`gnu services kerberos`) module provides services relating to the authentication protocol *Kerberos*.

#### Krb5 Service

Programs using a Kerberos client library normally expect a configuration file in `/etc/krb5.conf`. This service generates such a file from a definition provided in the operating system declaration. It does not cause any daemon to be started.

No “keytab” files are provided by this service—you must explicitly create them. This service is known to work with the MIT client library, `mit-krb5`. Other implementations have not been tested.

**krb5-service-type** [Variable]  
A service type for Kerberos 5 clients.

Here is an example of its use:

```
(service krb5-service-type
  (krb5-configuration
    (default-realm "EXAMPLE.COM")
    (allow-weak-crypto? #t)
    (realms (list
      (krb5-realm
        (name "EXAMPLE.COM")
        (admin-server "groucho.example.com")
        (kdc "karl.example.com")))
      (krb5-realm
        (name "ARGRX.EDU")
        (admin-server "kerb-admin.argrx.edu")
        (kdc "keys.argrx.edu"))))))
```

This example provides a Kerberos 5 client configuration which:

- Recognizes two realms, *viz*: “EXAMPLE.COM” and “ARGRX.EDU”, both of which have distinct administration servers and key distribution centers;
- Will default to the realm “EXAMPLE.COM” if the realm is not explicitly specified by clients;
- Accepts services which only support encryption types known to be weak.

The `krb5-realm` and `krb5-configuration` types have many fields. Only the most commonly used ones are described here. For a full list, and more detailed explanation of each, see the MIT `krb5.conf` documentation.

**krb5-realm** [Data Type]

**name** This field is a string identifying the name of the realm. A common convention is to use the fully qualified DNS name of your organization, converted to upper case.

**admin-server**

This field is a string identifying the host where the administration server is running.

**kdc**

This field is a string identifying the key distribution center for the realm.

**krb5-configuration**

[Data Type]

**allow-weak-crypto?** (default: **#f**)

If this flag is **#t** then services which only offer encryption algorithms known to be weak will be accepted.

**default-realm** (default: **#f**)

This field should be a string identifying the default Kerberos realm for the client. You should set this field to the name of your Kerberos realm. If this value is **#f** then a realm must be specified with every Kerberos principal when invoking programs such as **kinit**.

**realms**

This should be a non-empty list of **krb5-realm** objects, which clients may access. Normally, one of them will have a **name** field matching the **default-realm** field.

## PAM krb5 Service

The **pam-krb5** service allows for login authentication and password management via Kerberos. You will need this service if you want PAM enabled applications to authenticate users using Kerberos.

**pam-krb5-service-type**

[Variable]

A service type for the Kerberos 5 PAM module.

**pam-krb5-configuration**

[Data Type]

Data type representing the configuration of the Kerberos 5 PAM module. This type has the following parameters:

**pam-krb5** (default: **pam-krb5**)

The **pam-krb5** package to use.

**minimum-uid** (default: 1000)

The smallest user ID for which Kerberos authentications should be attempted. Local accounts with lower values will silently fail to authenticate.

### 11.10.19 LDAP Services

#### Authentication against LDAP with **nsld**

The **(gnu services authentication)** module provides the **nsld-service-type**, which can be used to authenticate against an LDAP server. In addition to configuring the service itself, you may want to add **ldap** as a name service to the Name Service Switch. See Section 11.13 [Name Service Switch], page 680, for detailed information.



Here is a simple operating system declaration with a default configuration of the `nsld-service-type` and a Name Service Switch configuration that consults the `ldap` name service last:

```
(use-service-modules authentication)
(use-modules (gnu system nss))
...
(operating-system
  ...
  (services
    (cons*
      (service nsld-service-type)
      (service dhcp-client-service-type)
      %base-services))
  (name-service-switch
    (let ((services (list (name-service (name "db"))
                          (name-service (name "files"))
                          (name-service (name "ldap")))))
      (name-service-switch
        (inherit %mdns-host-lookup-nss)
        (password services)
        (shadow services)
        (group services)
        (netgroup services)
        (gshadow services))))))
```

Available `nsld`-configuration fields are:

- |                                                                                                                                                                                                                                       |                                               |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------|
| <b>package</b> <code>nss-pam-ldapd</code>                                                                                                                                                                                             | [ <code>nsld</code> -configuration parameter] |
| The <code>nss-pam-ldapd</code> package to use.                                                                                                                                                                                        |                                               |
| <b>maybe-number</b> <code>threads</code>                                                                                                                                                                                              | [ <code>nsld</code> -configuration parameter] |
| The number of threads to start that can handle requests and perform LDAP queries. Each thread opens a separate connection to the LDAP server. The default is to start 5 threads.                                                      |                                               |
| Defaults to ‘disabled’.                                                                                                                                                                                                               |                                               |
| <b>string</b> <code>uid</code>                                                                                                                                                                                                        | [ <code>nsld</code> -configuration parameter] |
| This specifies the user id with which the daemon should be run.                                                                                                                                                                       |                                               |
| Defaults to “ <code>nsld</code> ”.                                                                                                                                                                                                    |                                               |
| <b>string</b> <code>gid</code>                                                                                                                                                                                                        | [ <code>nsld</code> -configuration parameter] |
| This specifies the group id with which the daemon should be run.                                                                                                                                                                      |                                               |
| Defaults to “ <code>nsld</code> ”.                                                                                                                                                                                                    |                                               |
| <b>log-option</b> <code>log</code>                                                                                                                                                                                                    | [ <code>nsld</code> -configuration parameter] |
| This option controls the way logging is done via a list containing SCHEME and LEVEL. The SCHEME argument may either be the symbols ‘none’ or ‘syslog’, or an absolute file name. The LEVEL argument is optional and specifies the log |                                               |

level. The log level may be one of the following symbols: ‘crit’, ‘error’, ‘warning’, ‘notice’, ‘info’ or ‘debug’. All messages with the specified log level or higher are logged.

Defaults to ‘‘("/var/log/nslcd" info)’’.

**list uri** [nslcd-configuration parameter]

The list of LDAP server URIs. Normally, only the first server will be used with the following servers as fall-back.

Defaults to ‘‘("ldap://localhost:389/")’’.

**maybe-string ldap-version** [nslcd-configuration parameter]

The version of the LDAP protocol to use. The default is to use the maximum version supported by the LDAP library.

Defaults to ‘disabled’.

**maybe-string binddn** [nslcd-configuration parameter]

Specifies the distinguished name with which to bind to the directory server for lookups. The default is to bind anonymously.

Defaults to ‘disabled’.

**maybe-string bindpw** [nslcd-configuration parameter]

Specifies the credentials with which to bind. This option is only applicable when used with binddn.

Defaults to ‘disabled’.

**maybe-string rootpwmoddn** [nslcd-configuration parameter]

Specifies the distinguished name to use when the root user tries to modify a user’s password using the PAM module.

Defaults to ‘disabled’.

**maybe-string rootpwmodpw** [nslcd-configuration parameter]

Specifies the credentials with which to bind if the root user tries to change a user’s password. This option is only applicable when used with rootpwmoddn.

Defaults to ‘disabled’.

**maybe-string sasl-mech** [nslcd-configuration parameter]

Specifies the SASL mechanism to be used when performing SASL authentication.

Defaults to ‘disabled’.

**maybe-string sasl-realm** [nslcd-configuration parameter]

Specifies the SASL realm to be used when performing SASL authentication.

Defaults to ‘disabled’.

**maybe-string sasl-authcid** [nslcd-configuration parameter]

Specifies the authentication identity to be used when performing SASL authentication.

Defaults to ‘disabled’.

- maybe-string sasl-authzid** [nslcd-configuration parameter]  
Specifies the authorization identity to be used when performing SASL authentication.  
Defaults to 'disabled'.
- maybe-boolean sasl-canonicalize?** [nslcd-configuration parameter]  
Determines whether the LDAP server host name should be canonicalised. If this is enabled the LDAP library will do a reverse host name lookup. By default, it is left up to the LDAP library whether this check is performed or not.  
Defaults to 'disabled'.
- maybe-string krb5-ccname** [nslcd-configuration parameter]  
Set the name for the GSS-API Kerberos credentials cache.  
Defaults to 'disabled'.
- string base** [nslcd-configuration parameter]  
The directory search base.  
Defaults to "dc=example,dc=com".
- scope-option scope** [nslcd-configuration parameter]  
Specifies the search scope (subtree, onelevel, base or children). The default scope is subtree; base scope is almost never useful for name service lookups; children scope is not supported on all servers.  
Defaults to '(subtree)'.
- maybe-deref-option deref** [nslcd-configuration parameter]  
Specifies the policy for dereferencing aliases. The default policy is to never dereference aliases.  
Defaults to 'disabled'.
- maybe-boolean referrals** [nslcd-configuration parameter]  
Specifies whether automatic referral chasing should be enabled. The default behaviour is to chase referrals.  
Defaults to 'disabled'.
- list-of-map-entries maps** [nslcd-configuration parameter]  
This option allows for custom attributes to be looked up instead of the default RFC 2307 attributes. It is a list of maps, each consisting of the name of a map, the RFC 2307 attribute to match and the query expression for the attribute as it is available in the directory.  
Defaults to '()'.
- list-of-filter-entries filters** [nslcd-configuration parameter]  
A list of filters consisting of the name of a map to which the filter applies and an LDAP search filter expression.  
Defaults to '()'.

- maybe-number bind-timelimit** [nslcd-configuration parameter]  
Specifies the time limit in seconds to use when connecting to the directory server.  
The default value is 10 seconds.  
Defaults to 'disabled'.
- maybe-number timelimit** [nslcd-configuration parameter]  
Specifies the time limit (in seconds) to wait for a response from the LDAP server. A value of zero, which is the default, is to wait indefinitely for searches to be completed.  
Defaults to 'disabled'.
- maybe-number idle-timelimit** [nslcd-configuration parameter]  
Specifies the period of inactivity (in seconds) after which the connection to the LDAP server will be closed. The default is not to time out connections.  
Defaults to 'disabled'.
- maybe-number reconnect-sleeptime** [nslcd-configuration parameter]  
Specifies the number of seconds to sleep when connecting to all LDAP servers fails. By default one second is waited between the first failure and the first retry.  
Defaults to 'disabled'.
- maybe-number reconnect-retrytime** [nslcd-configuration parameter]  
Specifies the time after which the LDAP server is considered to be permanently unavailable. Once this time is reached retries will be done only once per this time period. The default value is 10 seconds.  
Defaults to 'disabled'.
- maybe-ssl-option ssl** [nslcd-configuration parameter]  
Specifies whether to use SSL/TLS or not (the default is not to). If 'start-tls' is specified then StartTLS is used rather than raw LDAP over SSL.  
Defaults to 'disabled'.
- maybe-tls-reqcert-option tls-reqcert** [nslcd-configuration parameter]  
Specifies what checks to perform on a server-supplied certificate. The meaning of the values is described in the `ldap.conf(5)` manual page.  
Defaults to 'disabled'.
- maybe-string tls-cacertdir** [nslcd-configuration parameter]  
Specifies the directory containing X.509 certificates for peer authentication. This parameter is ignored when using GnuTLS.  
Defaults to 'disabled'.
- maybe-string tls-cacertfile** [nslcd-configuration parameter]  
Specifies the path to the X.509 certificate for peer authentication.  
Defaults to 'disabled'.
- maybe-string tls-randfile** [nslcd-configuration parameter]  
Specifies the path to an entropy source. This parameter is ignored when using GnuTLS.  
Defaults to 'disabled'.

- maybe-string tls-ciphers** [nslcd-configuration parameter]  
Specifies the ciphers to use for TLS as a string.  
Defaults to 'disabled'.
- maybe-string tls-cert** [nslcd-configuration parameter]  
Specifies the path to the file containing the local certificate for client TLS authentication.  
Defaults to 'disabled'.
- maybe-string tls-key** [nslcd-configuration parameter]  
Specifies the path to the file containing the private key for client TLS authentication.  
Defaults to 'disabled'.
- maybe-number pagesize** [nslcd-configuration parameter]  
Set this to a number greater than 0 to request paged results from the LDAP server in accordance with RFC2696. The default (0) is to not request paged results.  
Defaults to 'disabled'.
- maybe-ignore-users-option** [nslcd-configuration parameter]  
**nss-initgroups-ignoreusers**  
This option prevents group membership lookups through LDAP for the specified users. Alternatively, the value 'all-local' may be used. With that value nslcd builds a full list of non-LDAP users on startup.  
Defaults to 'disabled'.
- maybe-number nss-min-uid** [nslcd-configuration parameter]  
This option ensures that LDAP users with a numeric user id lower than the specified value are ignored.  
Defaults to 'disabled'.
- maybe-number nss-uid-offset** [nslcd-configuration parameter]  
This option specifies an offset that is added to all LDAP numeric user ids. This can be used to avoid user id collisions with local users.  
Defaults to 'disabled'.
- maybe-number nss-gid-offset** [nslcd-configuration parameter]  
This option specifies an offset that is added to all LDAP numeric group ids. This can be used to avoid user id collisions with local groups.  
Defaults to 'disabled'.
- maybe-boolean nss-nested-groups** [nslcd-configuration parameter]  
If this option is set, the member attribute of a group may point to another group. Members of nested groups are also returned in the higher level group and parent groups are returned when finding groups for a specific user. The default is not to perform extra searches for nested groups.  
Defaults to 'disabled'.

**maybe-boolean nss-getgrent-skipmembers** [nslcd-configuration parameter]

If this option is set, the group member list is not retrieved when looking up groups. Lookups for finding which groups a user belongs to will remain functional so the user will likely still get the correct groups assigned on login.

Defaults to 'disabled'.

**maybe-boolean nss-disable-enumeration** [nslcd-configuration parameter]

If this option is set, functions which cause all user/group entries to be loaded from the directory will not succeed in doing so. This can dramatically reduce LDAP server load in situations where there are a great number of users and/or groups. This option is not recommended for most configurations.

Defaults to 'disabled'.

**maybe-string validnames** [nslcd-configuration parameter]

This option can be used to specify how user and group names are verified within the system. This pattern is used to check all user and group names that are requested and returned from LDAP.

Defaults to 'disabled'.

**maybe-boolean ignorecase** [nslcd-configuration parameter]

This specifies whether or not to perform searches using case-insensitive matching. Enabling this could open up the system to authorization bypass vulnerabilities and introduce nscd cache poisoning vulnerabilities which allow denial of service.

Defaults to 'disabled'.

**maybe-boolean pam-authc-ppolicy** [nslcd-configuration parameter]

This option specifies whether password policy controls are requested and handled from the LDAP server when performing user authentication.

Defaults to 'disabled'.

**maybe-string pam-authc-search** [nslcd-configuration parameter]

By default nslcd performs an LDAP search with the user's credentials after BIND (authentication) to ensure that the BIND operation was successful. The default search is a simple check to see if the user's DN exists. A search filter can be specified that will be used instead. It should return at least one entry.

Defaults to 'disabled'.

**maybe-string pam-authz-search** [nslcd-configuration parameter]

This option allows flexible fine tuning of the authorisation check that should be performed. The search filter specified is executed and if any entries match, access is granted, otherwise access is denied.

Defaults to 'disabled'.

**maybe-string pam-password-prohibit-message** [nslcd-configuration parameter]

If this option is set password modification using pam\_ldap will be denied and the specified message will be presented to the user instead. The message can be used to direct the user to an alternative means of changing their password.

Defaults to 'disabled'.

**list pam-services** [nslcd-configuration parameter]  
 List of pam service names for which LDAP authentication should suffice.  
 Defaults to ‘‘()’.

## LDAP Directory Server

The (gnu services ldap) module provides the `directory-server-service-type`, which can be used to create and launch an LDAP server instance.

Here is an example configuration of the `directory-server-service-type`:

```
(use-service-modules ldap)

...
(operating-system
  ...
  (services
    (cons
      (service directory-server-service-type
        (directory-server-instance-configuration
          (slapd
            (slapd-configuration
              (root-password "{PBKDF2_SHA256}AAAgAG...ABSOLUTELYSECRET")))))
      %base-services)))
```

The root password should be generated with the `pwdhash` utility that is provided by the `389-ds-base` package.

Note that changes to the directory server configuration will not be applied to existing instances. You will need to back up and restore server data manually. Only new directory server instances will be created upon system reconfiguration.

**directory-server-instance-configuration** [Data Type]

Available `directory-server-instance-configuration` fields are:

**package** (default: `389-ds-base`) (type: file-like)

The `389-ds-base` package.

**config-version** (default: 2) (type: number)

Sets the format version of the configuration file. To use the INF file with `dscreate`, this parameter must be 2.

**full-machine-name** (default: `"localhost"`) (type: string)

Sets the fully qualified hostname (FQDN) of this system.

**selinux** (default: `#false`) (type: boolean)

Enables SELinux detection and integration during the installation of this instance. If set to `#true`, `dscreate` auto-detects whether SELinux is enabled.

**strict-host-checking** (default: `#true`) (type: boolean)

Sets whether the server verifies the forward and reverse record set in the `full-machine-name` parameter. When installing this instance with

GSSAPI authentication behind a load balancer, set this parameter to **#false**.

**systemd** (default: **#false**) (type: boolean)

Enables systemd platform features. If set to **#true**, **dscreate** auto-detects whether systemd is installed.

**slapd** (type: slapd-configuration)

Configuration of slapd.

**slapd-configuration** [Data Type]

Available **slapd-configuration** fields are:

**instance-name** (default: "localhost") (type: string)

Sets the name of the instance. You can refer to this value in other parameters of this INF file using the **{instance\_name}** variable. Note that this name cannot be changed after the installation!

**user** (default: "dirsrv") (type: string)

Sets the user name the ns-slapd process will use after the service started.

**group** (default: "dirsrv") (type: string)

Sets the group name the ns-slapd process will use after the service started.

**port** (default: 389) (type: number)

Sets the TCP port the instance uses for LDAP connections.

**secure-port** (default: 636) (type: number)

Sets the TCP port the instance uses for TLS-secured LDAP connections (LDAPS).

**root-dn** (default: "cn=Directory Manager") (type: string)

Sets the *Distinguished Name* (DN) of the administrator account for this instance.

**root-password** (default: "{invalid}YOU-SHOULD-CHANGE-THIS") (type: string)

Sets the password of the account specified in the **root-dn** parameter. You can either set this parameter to a plain text password **dscreate** hashes during the installation or to a "{algorithm}hash" string generated by the **pwdhash** utility. Note that setting a plain text password can be a security risk if unprivileged users can read this INF file!

**self-sign-cert** (default: **#true**) (type: boolean)

Sets whether the setup creates a self-signed certificate and enables TLS encryption during the installation.



This is not suitable for production, but it enables administrators to use TLS right after the installation. You can replace the self-signed certificate with a certificate issued by a certificate authority.

**self-sign-cert-valid-months** (default: 24) (type: number)  
Set the number of months the issued self-signed certificate will be valid.

**backup-dir** (default: `"/var/lib/dirsrv/slapd-{instance_name}/bak"`) (type: string)  
Set the backup directory of the instance.

**cert-dir** (default: `"/etc/dirsrv/slapd-{instance_name}"`) (type: string)  
Sets the directory of the instance's Network Security Services (NSS) database.

**config-dir** (default: `"/etc/dirsrv/slapd-{instance_name}"`) (type: string)  
Sets the configuration directory of the instance.

**db-dir** (default: `"/var/lib/dirsrv/slapd-{instance_name}/db"`) (type: string)  
Sets the database directory of the instance.

**initconfig-dir** (default: `"/etc/dirsrv/registry"`) (type: string)  
Sets the directory of the operating system's rc configuration directory.

**ldif-dir** (default: `"/var/lib/dirsrv/slapd-{instance_name}/ldif"`) (type: string)  
Sets the LDIF export and import directory of the instance.

**lock-dir** (default: `"/var/lock/dirsrv/slapd-{instance_name}"`) (type: string)  
Sets the lock directory of the instance.

**log-dir** (default: `"/var/log/dirsrv/slapd-{instance_name}"`) (type: string)  
Sets the log directory of the instance.

**run-dir** (default: `"/run/dirsrv"`) (type: string)  
Sets PID directory of the instance.

**schema-dir** (default: `"/etc/dirsrv/slapd-{instance_name}/schema"`) (type: string)  
Sets schema directory of the instance.

**tmp-dir** (default: `"/tmp"`) (type: string)  
Sets the temporary directory of the instance.

**backend-userroot** (type: backend-userroot-configuration)

Configuration of the userroot backend.

**backend-userroot-configuration** [Data Type]

Available **backend-userroot-configuration** fields are:

**create-suffix-entry?** (default: **#false**) (type: boolean)

Set this parameter to **#true** to create a generic root node entry for the suffix in the database.

**require-index?** (default: **#false**) (type: boolean)

Set this parameter to **#true** to refuse unindexed searches in this database.

**sample-entries** (default: **"no"**) (type: string)

Set this parameter to **"yes"** to add latest version of sample entries to this database. Or, use **"001003006"** to use the 1.3.6 version sample entries. Use this option, for example, to create a database for testing purposes.

**suffix** (type: maybe-string)

Sets the root suffix stored in this database. If you do not set the suffix attribute the install process will not create the backend/suffix. You can also create multiple backends/suffixes by duplicating this section.

### 11.10.20 Web Services

The (**gnu services web**) module provides the Apache HTTP Server, the nginx web server, and also a fastcgi wrapper daemon.

#### Apache HTTP Server

**httpd-service-type** [Variable]

Service type for the Apache HTTP (<https://httpd.apache.org/>) server (*httpd*). The value for this service type is a **httpd-configuration** record.

A simple example configuration is given below.

```
(service httpd-service-type
  (httpd-configuration
    (config
      (httpd-config-file
        (server-name "www.example.com")
        (document-root "/srv/http/www.example.com")))))
```

Other services can also extend the **httpd-service-type** to add to the configuration.

```
(simple-service 'www.example.com-server httpd-service-type
  (list
    (httpd-virtualhost
      "*:80"
      (list (string-join '("ServerName www.example.com"
        "DocumentRoot /srv/http/www.example.com"
        "\n")))))
```

The details for the `httpd-configuration`, `httpd-module`, `httpd-config-file` and `httpd-virtualhost` record types are given below.

**httpd-configuration** [Data Type]

This data type represents the configuration for the httpd service.

**package** (default: `httpd`)

The httpd package to use.

**pid-file** (default: `"/var/run/httpd"`)

The pid file used by the shepherd-service.

**config** (default: `(httpd-config-file)`)

The configuration file to use with the httpd service. The default value is a `httpd-config-file` record, but this can also be a different G-expression that generates a file, for example a `plain-file`. A file outside of the store can also be specified through a string.

**httpd-module** [Data Type]

This data type represents a module for the httpd service.

**name** The name of the module.

**file** The file for the module. This can be relative to the httpd package being used, the absolute location of a file, or a G-expression for a file within the store, for example `(file-append mod-wsgi "/modules/mod_wsgi.so")`.

**%default-httpd-modules** [Variable]

A default list of `httpd-module` objects.

**httpd-config-file** [Data Type]

This data type represents a configuration file for the httpd service.

**modules** (default: `%default-httpd-modules`)

The modules to load. Additional modules can be added here, or loaded by additional configuration.

For example, in order to handle requests for PHP files, you can use Apache's `mod_proxy_fcgi` module along with `php-fpm-service-type`:

```
(service httpd-service-type
  (httpd-configuration
    (config
      (httpd-config-file
        (modules (cons*
                  (httpd-module
                    (name "proxy_module")
                    (file "modules/mod_proxy.so")))
                  (httpd-module
                    (name "proxy_fcgi_module")
                    (file "modules/mod_proxy_fcgi.so")))
                  %default-httpd-modules))
        (extra-config (list "\
```

```

<FilesMatch \\.php$>
    SetHandler \"proxy:unix:/var/run/php-fpm.sock|fcgi://localhost/\"
</FilesMatch>\"))))))
(service php-fpm-service-type
  (php-fpm-configuration
    (socket \"/var/run/php-fpm.sock\")
    (socket-group \"httpd\")))

```

**server-root** (default: `httpd`)

The **ServerRoot** in the configuration file, defaults to the `httpd` package. Directives including **Include** and **LoadModule** are taken as relative to the server root.

**server-name** (default: `#f`)

The **ServerName** in the configuration file, used to specify the request scheme, hostname and port that the server uses to identify itself.

This doesn't need to be set in the server config, and can be specified in virtual hosts. The default is `#f` to not specify a **ServerName**.

**document-root** (default: `"/srv/http"`)

The **DocumentRoot** from which files will be served.

**listen** (default: `'(\"80\")'`)

The list of values for the **Listen** directives in the config file. The value should be a list of strings, when each string can specify the port number to listen on, and optionally the IP address and protocol to use.

**pid-file** (default: `"/var/run/httpd"`)

The **PidFile** to use. This should match the **pid-file** set in the `httpd-configuration` so that the Shepherd service is configured correctly.

**error-log** (default: `"/var/log/httpd/error_log"`)

The **ErrorLog** to which the server will log errors.

**user** (default: `"httpd"`)

The **User** which the server will answer requests as.

**group** (default: `"httpd"`)

The **Group** which the server will answer requests as.

**extra-config** (default: `(list \"TypesConfig etc/httpd/mime.types\")`)

A flat list of strings and G-expressions which will be added to the end of the configuration file.

Any values which the service is extended with will be appended to this list.

**httpd-virtualhost**

[Data Type]

This data type represents a virtualhost configuration block for the `httpd` service.

These should be added to the extra-config for the `httpd-service`.

```

(simple-service 'www.example.com-server httpd-service-type
  (list

```

```
(httpd-virtualhost
  " *:80"
  (list (string-join '("ServerName www.example.com"
    "DocumentRoot /srv/http/www.example.com"
    "\n")))))
```

**addresses-and-ports**

The addresses and ports for the `VirtualHost` directive.

**contents** The contents of the `VirtualHost` directive, this should be a list of strings and G-expressions.

## NGINX

**nginx-service-type** [Variable]

Service type for the NGinx (<https://nginx.org/>) web server. The value for this service type is a `<nginx-configuration>` record.

A simple example configuration is given below.

```
(service nginx-service-type
  (nginx-configuration
    (server-blocks
      (list (nginx-server-configuration
        (server-name '("www.example.com"))
        (root "/srv/http/www.example.com")))))
```

In addition to adding server blocks to the service configuration directly, this service can be extended by other services to add server blocks, as in this example:

```
(simple-service 'my-extra-server nginx-service-type
  (list (nginx-server-configuration
    (root "/srv/http/extra-website")
    (try-files (list "$uri" "$uri/index.html")))))
```

At startup, `nginx` has not yet read its configuration file, so it uses a default file to log error messages. If it fails to load its configuration file, that is where error messages are logged. After the configuration file is loaded, the default error log file changes as per configuration. In our case, startup error messages can be found in `/var/run/nginx/logs/error.log`, and after configuration in `/var/log/nginx/error.log`. The second location can be changed with the `log-directory` configuration option.

**nginx-configuration** [Data Type]

This data type represents the configuration for NGinx. Some configuration can be done through this and the other provided record types, or alternatively, a config file can be provided.

**nginx** (default: `nginx`)

The `nginx` package to use.

**shepherd-requirement** (default: `'()`)

This is a list of symbols naming Shepherd services the `nginx` service will depend on.

This is useful if you would like `nginx` to be started after a back-end web server or a logging service such as Anonip has been started.

`log-directory` (default: `"/var/log/nginx"`)

The directory to which NGinx will write log files.

`log-format` (default: `'combined'`) (type: symbol)

Logging format for the access log. Nginx defaults to the httpd-like format named `'combined'`. Other formats can be defined using `log-formats`.

`log-formats` (default: `'()'`)

A list of additional *log formats* to define, the elements should be of type `<nginx-log-format-configuration>`.

`log-level` (default: `'error'`) (type: symbol)

Logging level for the error log, which can be any of the following values: `'debug'`, `'info'`, `'notice'`, `'warn'`, `'error'`, `'crit'`, `'alert'`, or `'emerg'`.

`run-directory` (default: `"/var/run/nginx"`)

The directory in which NGinx will create a pid file, and write temporary files.

`server-blocks` (default: `'()'`)

A list of *server blocks* to create in the generated configuration file, the elements should be of type `<nginx-server-configuration>`.

The following example would setup NGinx to serve `www.example.com` from the `/srv/http/www.example.com` directory, without using HTTPS.

```
(service nginx-service-type
  (nginx-configuration
    (server-blocks
      (list (nginx-server-configuration
              (server-name '("www.example.com"))
              (root "/srv/http/www.example.com"))))))■
```

`upstream-blocks` (default: `'()'`)

A list of *upstream blocks* to create in the generated configuration file, the elements should be of type `<nginx-upstream-configuration>`.

Configuring upstreams through the `upstream-blocks` can be useful when combined with `locations` in the `<nginx-server-configuration>` records. The following example creates a server configuration with one location configuration, that will proxy requests to a upstream configuration, which will handle requests with two servers.

```
(service
  nginx-service-type
  (nginx-configuration
    (server-blocks
      (list (nginx-server-configuration
              (server-name '("www.example.com"))
              (root "/srv/http/www.example.com")
              (locations
```

```

        (list
          (nginx-location-configuration
            (uri "/path1")
            (body '("proxy_pass http://server-proxy;"))))))))
      (upstream-blocks
        (list (nginx-upstream-configuration
          (name "server-proxy")
          (servers (list "server1.example.com"
            "server2.example.com"))))))))

```

**file** (default: #f)

If a configuration *file* is provided, this will be used, rather than generating a configuration file from the provided *log-directory*, *run-directory*, *server-blocks* and *upstream-blocks*. For proper operation, these arguments should match what is in *file* to ensure that the directories are created when the service is activated.

This can be useful if you have an existing configuration file, or it's not possible to do what is required through the other parts of the nginx-configuration record.

**server-names-hash-bucket-size** (default: #f)

Bucket size for the server names hash tables, defaults to #f to use the size of the processors cache line.

**server-names-hash-bucket-max-size** (default: #f)

Maximum bucket size for the server names hash tables.

**stream** (default: #f)

<nginx-stream-configuration> object describing stream server directives.

**modules** (default: '())

List of nginx dynamic modules to load. This should be a list of file names of loadable modules, as in this example:

```

(modules
  (list
    (file-append nginx-accept-language-module "\
/etc/nginx/modules/nginx_http_accept_language_module.so")
    (file-append nginx-lua-module "\
/etc/nginx/modules/nginx_http_lua_module.so")))

```

**lua-package-path** (default: '())

List of nginx lua packages to load. This should be a list of package names of loadable lua modules, as in this example:

```

(lua-package-path (list lua-resty-core
  lua-resty-lrucache
  lua-resty-signal
  lua-tablepool
  lua-resty-shell))

```

**lua-package-cpath** (default: '()')

List of nginx lua C packages to load. This should be a list of package names of loadable lua C modules, as in this example:

```
(lua-package-cpath (list lua-resty-signal))
```

**global-directives** (default: '((events . ())))

Association list of global directives for the top level of the nginx configuration. Values may themselves be association lists.

```
(global-directives
  `((worker_processes . 16)
    (pcre_jit . on)
    (events . ((worker_connections . 1024)))))
```

**extra-content** (default: "")

Additional content to be appended to the `http` block. Can either be a value that can be lowered into a string or a list of such values. In the former case, it is inserted directly. In the latter, it is prefixed with indentation and suffixed with a newline. Nested lists are flattened into one line.

```
(extra-content "include /etc/nginx/custom-config.conf;")
(extra-content `("include /etc/nginx/custom-config.conf;"
  ("include " ,%custom-config.conf ";")))
```

**nginx-log-format-configuration** [Data Type]

Data type representing the configuration of a custom Nginx log format. This type has the following parameters:

**name** (type: symbol)

The name of the log format as a symbol.

**escape** (default: 'default') (type: symbol)

The mode used to escape values of variables in the format. The supported values are:

- 'default, to escape all characters outside of the ASCII printable range.
- 'json, to escape all characters invalid in JSON strings.
- 'none, to disable escaping.

**format** (type: string)

The format as accepted by the `log_format` directive.

**nginx-server-configuration** [Data Type]

Data type representing the configuration of an nginx server block. This type has the following parameters:

**listen** (default: '("80" "443 ssl"))

Each `listen` directive sets the address and port for IP, or the path for a UNIX-domain socket on which the server will accept requests. Both address and port, or only address or only port can be specified. An address may also be a hostname, for example:

```
'("127.0.0.1:8000" "127.0.0.1" "8000" " *:8000" "localhost:8000")'■
```



**server-name** (default: (list 'default))  
 A list of server names this server represents. **'default'** represents the default server for connections matching no other server.

**root** (default: "/srv/http")  
 Root of the website nginx will serve.

**locations** (default: '()')  
 A list of *nginx-location-configuration* or *nginx-named-location-configuration* records to use within this server block.

**index** (default: (list "index.html"))  
 Index files to look for when clients ask for a directory. If it cannot be found, Nginx will send the list of files in the directory.

**try-files** (default: '()')  
 A list of files whose existence is checked in the specified order. **nginx** will use the first file it finds to process the request.

**ssl-certificate** (default: #f)  
 Where to find the certificate for secure connections. Set it to **#f** if you don't have a certificate or you don't want to use HTTPS.

**ssl-certificate-key** (default: #f)  
 Where to find the private key for secure connections. Set it to **#f** if you don't have a key or you don't want to use HTTPS.

**server-tokens?** (default: #f)  
 Whether the server should add its configuration to response.

**raw-content** (default: '()')  
 A list of strings or file-like objects to be appended to the server block. Each item is prefixed with indentation and suffixed with a new line. Nested lists are flattened.

**nginx-upstream-configuration** [Data Type]  
 Data type representing the configuration of an nginx **upstream** block. This type has the following parameters:

**name**        Name for this group of servers.

**servers**     Specify the addresses of the servers in the group. The address can be specified as a IP address (e.g. '127.0.0.1'), domain name (e.g. 'backend1.example.com') or a path to a UNIX socket using the prefix 'unix:'. For addresses using an IP address or domain name, the default port is 80, and a different port can be specified explicitly.

**extra-content**  
 Additional content to be appended to the upstream block. Can be a string or file-like object or list of thereof. In case of list, each item is prefixed with indentation and suffixed with a new line. Nested lists are flattened.

```
(extra-content "include /etc/nginx/custom-config.conf;")
```

```
(extra-content `("include /etc/nginx/custom-config.conf;"
                ("include " ,%custom-config.conf " ;")))
```

### **nginx-location-configuration** [Data Type]

Data type representing the configuration of an nginx location block. This type has the following parameters:

- uri**           URI which this location block matches.
- body**           Body of the location block, specified as a list of strings or file-like objects. Each item is prefixed with indentation and suffixed with a new line. Nested lists are flattened.  
  
For example, to pass requests to a upstream server group defined using an **nginx-upstream-configuration** block, the following directive would be specified in the body `(list "proxy_pass http://upstream-name;")`.

### **nginx-named-location-configuration** [Data Type]

Data type representing the configuration of an nginx named location block. Named location blocks are used for request redirection, and not used for regular request processing. This type has the following parameters:

- name**           Name to identify this location block.
- body**           See [nginx-location-configuration body], page 502, as the body for named location blocks can be used in a similar way to the **nginx-location-configuration** body. One restriction is that the body of a named location block cannot contain location blocks.

### **nginx-stream-configuration** [Data Type]

Data type representing the configuration file context in which stream directives are specified. This type has the following parameters:

- upstream-blocks** (default: '()')  
A list of upstream blocks to create in the generated configuration file, the elements should be of type **<nginx-upstream-configuration>**.
- server-blocks** (default: '()')  
A list of server blocks to create in the generated configuration file, the elements should be of type **<nginx-server-configuration>**.
- extra-content** (default: '()')  
Additional content to be appended to the **stream** block. Can either be a value that can be lowered into a string or a list of such values. In the former case, it is inserted directly. In the latter, it is prefixed with indentation and suffixed with a newline. Nested lists are flattened into one line.

## **Varnish Cache**

Varnish is a fast cache server that sits in between web applications and end users. It proxies requests from clients and caches the accessed URLs such that multiple requests for the same resource only creates one request to the back-end.

**varnish-service-type** [Variable]

Service type for the Varnish daemon.

**varnish-configuration** [Data Type]

Data type representing the **varnish** service configuration. This type has the following parameters:

**package** (default: **varnish**)

The Varnish package to use.

**name** (default: **"default"**)

A name for this Varnish instance. Varnish will create a directory in **/var/varnish/** with this name and keep temporary files there. If the name starts with a forward slash, it is interpreted as an absolute directory name.

Pass the **-n** argument to other Varnish programs to connect to the named instance, e.g. **varnishncsa -n default**.

**backend** (default: **"localhost:8080"**)

The backend to use. This option has no effect if **vcl** is set.

**vcl** (default: **#f**)

The **VCL** (Varnish Configuration Language) program to run. If this is **#f**, Varnish will proxy **backend** using the default configuration. Otherwise this must be a file-like object with valid VCL syntax.

For example, to mirror [www.gnu.org](https://www.gnu.org) (<https://www.gnu.org>) with VCL you can do something along these lines:

```
(define %gnu-mirror
  (plain-file "gnu.vcl"
    "vcl 4.1;
    backend gnu { .host = \"www.gnu.org\"; }"))
```

```
(operating-system
```

```
;; ...
```

```
(services (cons (service varnish-service-type
                          (varnish-configuration
                           (listen '(":80"))
                           (vcl %gnu-mirror)))
                %base-services)))
```

The configuration of an already running Varnish instance can be inspected and changed using the **varnishadm** program.

Consult the Varnish User Guide (<https://varnish-cache.org/docs/>) and Varnish Book (<https://book.varnish-software.com/4.0/>) for comprehensive documentation on Varnish and its configuration language.

**listen** (default: **'("localhost:80")'**)

List of addresses Varnish will listen on.

**storage** (default: **'("malloc,128m")'**)

List of storage backends that will be available in VCL.

**parameters** (default: '()')  
List of run-time parameters in the form '("parameter" . "value").

**extra-options** (default: '()')  
Additional arguments to pass to the **varnishd** process.

## Whoogle Search

Whoogle Search (<https://github.com/benbusby/whoogle-search>) is a self-hosted, ad-free, privacy-respecting meta search engine that collects and displays Google search results. By default, you can configure it by adding this line to the **services** field of your operating system declaration:

```
(service whoogle-service-type)
```

As a result, Whoogle Search runs as local Web server, which you can access by opening 'http://localhost:5000' in your browser. The configuration reference is given below.

**whoogle-service-type** [Variable]  
Service type for Whoogle Search. Its value must be a **whoogle-configuration** record—see below.

**whoogle-configuration** [Data Type]  
Data type representing Whoogle Search service configuration.

**package** (default: **whoogle-search**)  
The Whoogle Search package to use.

**host** (default: "127.0.0.1")  
The host address to run Whoogle on.

**port** (default: 5000)  
The port where Whoogle will be exposed.

**environment-variables** (default: '()')  
A list of strings with the environment variables to configure Whoogle. You can consult its environment variables template (<https://github.com/benbusby/whoogle-search/blob/main/whoogle.template.env>) for the list of available options.

## go-wbedav

go-webdav (<https://github.com/emersion/go-webdav>) is a server for the WebDAV (<https://www.rfc-editor.org/rfc/rfc4918>) protocol.

**go-webdav-service-type** [Variable]  
This is the service type for go-webdav. Since go-webdav does not have a configuration file, its value must be a list of command-line arguments that should be passed to the **webdav-server** command. For example, the following configuration will serve all files in `/srv/http/` via port 8080 on 127.0.0.1:

```
(service go-webdav-service-type
  '("-addr" "127.0.0.1:8080" "/srv/http/"))
```

## Patchwork

Patchwork is a patch tracking system. It can collect patches sent to a mailing list, and display them in a web interface.

**patchwork-service-type** [Variable]  
Service type for Patchwork.

The following example is an example of a minimal service for Patchwork, for the `patchwork.example.com` domain.

```
(service patchwork-service-type
  (patchwork-configuration
    (domain "patchwork.example.com")
    (settings-module
      (patchwork-settings-module
        (allowed-hosts (list domain))
        (default-from-email "patchwork@patchwork.example.com"))))
  (getmail-retriever-config
    (getmail-retriever-configuration
      (type "SimpleIMAPSSLRetriever")
      (server "imap.example.com")
      (port 993)
      (username "patchwork")
      (password-command
        (list (file-append coreutils "/bin/cat")
              "/etc/getmail-patchwork-imap-password"))
      (extra-parameters
        '((mailboxes . ("Patches"))))))))
```

There are three records for configuring the Patchwork service. The `<patchwork-configuration>` relates to the configuration for Patchwork within the HTTPD service.

The `settings-module` field within the `<patchwork-configuration>` record can be populated with the `<patchwork-settings-module>` record, which describes a settings module that is generated within the Guix store.

For the `database-configuration` field within the `<patchwork-settings-module>`, the `<patchwork-database-configuration>` must be used.

**patchwork-configuration** [Data Type]  
Data type representing the Patchwork service configuration. This type has the following parameters:

**patchwork** (default: `patchwork`)  
The Patchwork package to use.

**domain** The domain to use for Patchwork, this is used in the HTTPD service virtual host.

**settings-module**

The settings module to use for Patchwork. As a Django application, Patchwork is configured with a Python module containing the settings. This can either be an instance of the `<patchwork-settings-module>` record, any other record that represents the settings in the store, or a directory outside of the store.

**static-path** (default: `"/static/"`)

The path under which the HTTPD service should serve the static files.

**getmail-retriever-config**

The getmail-retriever-configuration record value to use with Patchwork. Getmail will be configured with this value, the messages will be delivered to Patchwork.

**patchwork-settings-module**

[Data Type]

Data type representing a settings module for Patchwork. Some of these settings relate directly to Patchwork, but others relate to Django, the web framework used by Patchwork, or the Django Rest Framework library. This type has the following parameters:

**database-configuration** (default: `(patchwork-database-configuration)`)

The database connection settings used for Patchwork. See the `<patchwork-database-configuration>` record type for more information.

**secret-key-file** (default: `"/etc/patchwork/django-secret-key"`)

Patchwork, as a Django web application uses a secret key for cryptographically signing values. This file should contain a unique unpredictable value.

If this file does not exist, it will be created and populated with a random value by the patchwork-setup shepherd service.

This setting relates to Django.

**allowed-hosts**

A list of valid hosts for this Patchwork service. This should at least include the domain specified in the `<patchwork-configuration>` record.

This is a Django setting.

**default-from-email**

The email address from which Patchwork should send email by default.

This is a Patchwork setting.

**static-url** (default: `#f`)

The URL to use when serving static assets. It can be part of a URL, or a full URL, but must end in a `/`.

If the default value is used, the `static-path` value from the `<patchwork-configuration>` record will be used.

This is a Django setting.

**admins** (default: '()')  
 Email addresses to send the details of errors that occur. Each value should be a list containing two elements, the name and then the email address.  
 This is a Django setting.

**debug?** (default: #f)  
 Whether to run Patchwork in debug mode. If set to #t, detailed error messages will be shown.  
 This is a Django setting.

**enable-rest-api?** (default: #t)  
 Whether to enable the Patchwork REST API.  
 This is a Patchwork setting.

**enable-xmlrpc?** (default: #t)  
 Whether to enable the XML RPC API.  
 This is a Patchwork setting.

**force-https-links?** (default: #t)  
 Whether to use HTTPS links on Patchwork pages.  
 This is a Patchwork setting.

**extra-settings** (default: "")  
 Extra code to place at the end of the Patchwork settings module.

**patchwork-database-configuration** [Data Type]

Data type representing the database configuration for Patchwork.

**engine** (default: "django.db.backends.postgresql\_psycopg2")  
 The database engine to use.

**name** (default: "patchwork")  
 The name of the database to use.

**user** (default: "httpd")  
 The user to connect to the database as.

**password** (default: "")  
 The password to use when connecting to the database.

**host** (default: "")  
 The host to make the database connection to.

**port** (default: "")  
 The port on which to connect to the database.

## Mumi

Mumi (<https://codeberg.org/guix/mumi>) is a Web interface to the Debbugs bug tracker, by default for the GNU instance (<https://bugs.gnu.org>). Mumi is a Web server, but it also fetches and indexes mail retrieved from Debbugs.

**mumi-service-type** [Variable]

This is the service type for Mumi.

**mumi-configuration** [Data Type]

Data type representing the Mumi service configuration. This type has the following fields:

**mumi** (default: `mumi`)

The Mumi package to use.

**data-directory** (default: `"/var/mumi/data"`)

Data directory to store Debbugs data from upstream Debbugs instance.

**rsync-remote**

Remote rsync path from which to download Debbugs data. The remote is polled every 5 minutes, unless a previous instance of rsync is already running.

**rsync-flags** (default: `'()`)

Additional flags to pass to `rsync`.

**mailer?** (default: `#true`)

Whether to enable or disable the mailer component.

**sender** The email address used as the sender for comments.

**smtp** A URI to configure the SMTP settings for Mailutils. This could be something like `sendmail:///path/to/bin/msmtp` or any other URI supported by Mailutils. See Section “SMTP Mailboxes” in *GNU Mailutils*.

**packages** List of `<mumi-package-configuration>` objects representing packages configured on Mumi.

**mumi-package-configuration** [Data Type]

Data type representing the configuration of a package on Mumi. This type has the following fields:

**name** Name of the package.

**display-name**

Display name used for this package in the web interface.

**debbugs-packages**

List of Debbugs packages this package refers to.

**bug-email-address**

Email address to send bug reports to.

**patch-email-address**

Email address to send patches to, if that is different from `bug-email-address`.

**file-tags**

Association list mapping tags to filename regular expressions. Keys are tag names, and values are lists of regular expressions. For each entry of the association list, the tag is applied if a filename referred to in a patch matches any of the regular expressions.



## FastCGI

FastCGI is an interface between the front-end and the back-end of a web service. It is a somewhat legacy facility; new web services should generally just talk HTTP between the front-end and the back-end. However there are a number of back-end services such as PHP or the optimized HTTP Git repository access that use FastCGI, so we have support for it in Guix.

To use FastCGI, you configure the front-end web server (e.g., nginx) to dispatch some subset of its requests to the `fastcgi` backend, which listens on a local TCP or UNIX socket. There is an intermediary `fcgiwrap` program that sits between the actual backend process and the web server. The front-end indicates which backend program to run, passing that information to the `fcgiwrap` process.

**fcgiwrap-service-type** [Variable]

A service type for the `fcgiwrap` FastCGI proxy.

**fcgiwrap-configuration** [Data Type]

Data type representing the configuration of the `fcgiwrap` service. This type has the following parameters:

**package** (default: `fcgiwrap`)

The `fcgiwrap` package to use.

**socket** (default: `tcp:127.0.0.1:9000`)

The socket on which the `fcgiwrap` process should listen, as a string. Valid `socket` values include `unix:/path/to/unix/socket`, `tcp:dot.ted.qu.ad:port` and `tcp6:[ipv6_addr]:port`.

**user** (default: `fcgiwrap`)

**group** (default: `fcgiwrap`)

The user and group names, as strings, under which to run the `fcgiwrap` process. The `fastcgi` service will ensure that if the user asks for the specific user or group names `fcgiwrap` that the corresponding user and/or group is present on the system.

It is possible to configure a FastCGI-backed web service to pass HTTP authentication information from the front-end to the back-end, and to allow `fcgiwrap` to run the back-end process as a corresponding local user. To enable this capability on the back-end, run `fcgiwrap` as the `root` user and group. Note that this capability also has to be configured on the front-end as well.

## PHP-FPM

PHP-FPM (FastCGI Process Manager) is an alternative PHP FastCGI implementation with some additional features useful for sites of any size.

These features include:

- Adaptive process spawning
- Basic statistics (similar to Apache's `mod_status`)
- Advanced process management with graceful stop/start

- Ability to start workers with different uid/gid/chroot/environment and different php.ini (replaces safe\_mode)
- Stdout & stderr logging
- Emergency restart in case of accidental opcode cache destruction
- Accelerated upload support
- Support for a "slowlog"
- Enhancements to FastCGI, such as `fastcgi_finish_request()` - a special function to finish request & flush all data while continuing to do something time-consuming (video converting, stats processing, etc.)

... and much more.

**php-fpm-service-type** [Variable]  
A Service type for php-fpm.

**php-fpm-configuration** [Data Type]  
Data Type for php-fpm service configuration.

**php** (default: `php`)

The php package to use.

**socket** (default: `(string-append "/var/run/php" (version-major (package-version php)) "-fpm.sock")`)

The address on which to accept FastCGI requests. Valid syntaxes are:

`"ip.add.re.ss:port"`

Listen on a TCP socket to a specific address on a specific port.

`"port"` Listen on a TCP socket to all addresses on a specific port.

`"/path/to/unix/socket"`

Listen on a unix socket.

**user** (default: `php-fpm`)

User who will own the php worker processes.

**group** (default: `php-fpm`)

Group of the worker processes.

**socket-user** (default: `php-fpm`)

User who can speak to the php-fpm socket.

**socket-group** (default: `nginx`)

Group that can speak to the php-fpm socket.

**pid-file** (default: `(string-append "/var/run/php" (version-major (package-version php)) "-fpm.pid")`)

The process id of the php-fpm process is written to this file once the service has started.

**log-file** (default: `(string-append "/var/log/php" (version-major (package-version php)) "-fpm.log")`)

Log for the php-fpm master process.

**process-manager** (default:  
 (php-fpm-dynamic-process-manager-configuration))  
 Detailed settings for the php-fpm process manager. Must be one of:

```
<php-fpm-dynamic-process-manager-configuration>
<php-fpm-static-process-manager-configuration>
<php-fpm-on-demand-process-manager-configuration>
```

**display-errors** (default #f)  
 Determines whether php errors and warning should be sent to clients and displayed in their browsers. This is useful for local php development, but a security risk for public sites, as error messages can reveal passwords and personal data.

**timezone** (default #f)  
 Specifies `php_admin_value[date.timezone]` parameter.

**workers-logfile** (default (string-append "/var/log/php" (version-major (package-version php)) "-fpm.www.log"))  
 This file will log the `stderr` outputs of php worker processes. Can be set to #f to disable logging.

**file** (default #f)  
 An optional override of the whole configuration. You can use the `mixed-text-file` function or an absolute filepath for it.

**php-ini-file** (default #f)  
 An optional override of the default php settings. It may be any “file-like” object (see Section 8.12 [G-Expressions], page 165). You can use the `mixed-text-file` function or an absolute filepath for it.

For local development it is useful to set a higher timeout and memory limit for spawned php processes. This be accomplished with the following operating system configuration snippet:

```
(define %local-php-ini
  (plain-file "php.ini"
    "memory_limit = 2G
    max_execution_time = 1800"))

(operating-system
  ;; ...
  (services (cons (service php-fpm-service-type
    (php-fpm-configuration
      (php-ini-file %local-php-ini))))
    %base-services)))
```

Consult the core php.ini directives (<https://www.php.net/manual/en/ini.core.php>) for comprehensive documentation on the acceptable php.ini directives.

**php-fpm-dynamic-process-manager-configuration** [Data type]  
 Data Type for the dynamic php-fpm process manager. With the dynamic process manager, spare worker processes are kept around based on its configured limits.

**max-children** (default: 5)  
Maximum of worker processes.

**start-servers** (default: 2)  
How many worker processes should be started on start-up.

**min-spare-servers** (default: 1)  
How many spare worker processes should be kept around at minimum.

**max-spare-servers** (default: 3)  
How many spare worker processes should be kept around at maximum.

**php-fpm-static-process-manager-configuration** [Data type]  
Data Type for the **static** php-fpm process manager. With the **static** process manager, an unchanging number of worker processes are created.

**max-children** (default: 5)  
Maximum of worker processes.

**php-fpm-on-demand-process-manager-configuration** [Data type]  
Data Type for the **on-demand** php-fpm process manager. With the **on-demand** process manager, worker processes are only created as requests arrive.

**max-children** (default: 5)  
Maximum of worker processes.

**process-idle-timeout** (default: 10)  
The time in seconds after which a process with no requests is killed.

**nginx-php-location** [*#:nginx-package nginx*] [*socket*] [Procedure]  
(*string-append* *"/var/run/php"* (*version-major* (*package-version php*))  
*"-fpm.sock"*)]  
A helper function to quickly add php to an **nginx-server-configuration**.

A simple services setup for nginx with php can look like this:

```
(services (cons* (service dhcp-client-service-type)
                 (service php-fpm-service-type)
                 (service nginx-service-type
                           (nginx-server-configuration
                            (server-name '("example.com"))
                            (root "/srv/http/")
                            (locations
                             (list (nginx-php-location)))
                             (listen '("80"))
                             (ssl-certificate #f)
                             (ssl-certificate-key #f)))
                 %base-services))
```

The cat avatar generator is a simple service to demonstrate the use of php-fpm in **Nginx**. It is used to generate cat avatar from a seed, for instance the hash of a user's email address.

**cat-avatar-generator-service** [*#:cache-dir*] [*#:package cat-avatar-generator*] [*#:configuration (nginx-server-configuration)*] [Procedure]

Returns an `nginx-server-configuration` that inherits `configuration`. It extends the `nginx` configuration to add a server block that serves `package`, a version of `cat-avatar-generator`. During execution, `cat-avatar-generator` will be able to use `cache-dir` as its cache directory.

A simple setup for `cat-avatar-generator` can look like this:

```
(services (cons* (cat-avatar-generator-service
                  #:configuration
                  (nginx-server-configuration
                   (server-name '("example.com"))))
               ...
               %base-services))
```

## Hpcguix-web

The `hpcguix-web` (<https://github.com/UMCUGenetics/hpcguix-web/>) program is a customizable web interface to browse Guix packages, initially designed for users of high-performance computing (HPC) clusters.

**hpcguix-web-service-type** [Variable]  
The service type for `hpcguix-web`.

**hpcguix-web-configuration** [Data Type]  
Data type for the `hpcguix-web` service configuration.

**specs** (default: `#f`)  
Either `#f` or a `gexp` (see Section 8.12 [G-Expressions], page 165) specifying the `hpcguix-web` service configuration as an `hpcguix-web-configuration` record. The main fields of that record type are:

**title-prefix** (default: `"hpcguix | "`)  
The page title prefix.

**guix-command** (default: `"guix"`)  
The `guix` command to use in examples that appear on HTML pages.

**package-filter-proc** (default: `(const #t)`)  
A procedure specifying how to filter packages that are displayed.

**package-page-extension-proc** (default: `(const '())`)  
Extension package for `hpcguix-web`.

**menu** (default: `'()`)  
Additional entry in page menu.

**channels** (default: `%default-channels`)  
List of channels from which the package list is built (see Chapter 6 [Channels], page 68).

```
package-list-expiration (default: (* 12 3600))
```

The expiration time, in seconds, after which the package list is rebuilt from the latest instances of the given channels.

See the `hpcguix-web` repository for a complete example (<https://github.com/UMCUGenetics/hpcguix-web/blob/master/hpcweb-configuration.scm>).

```
package (default: hpcguix-web)
```

The `hpcguix-web` package to use.

```
address (default: "127.0.0.1")
```

The IP address to listen to.

```
port (default: 5000)
```

The port number to listen to.

A typical `hpcguix-web` service declaration looks like this:

```
(service hpcguix-web-service-type
  (hpcguix-web-configuration
    (specs
      #~(hpcweb-configuration
        (title-prefix "Guix-HPC - ")
        (menu '("/about" "ABOUT"))))))
```

**Note:** The `hpcguix-web` service periodically updates the package list it publishes by pulling channels from Git. To that end, it needs to access X.509 certificates so that it can authenticate Git servers when communicating over HTTPS, and it assumes that `/etc/ssl/certs` contains those certificates.

A certificate package, `nss-certs`, is provided by default as part of `%base-packages`. Section 11.12 [X.509 Certificates], page 679, for more information on X.509 certificates.

## gmnisrv

The `gmnisrv` (<https://git.sr.ht/~sircmpwn/gmnisrv>) program is a simple Gemini (<https://gemini.circumlunar.space/>) protocol server.

**gmnisrv-service-type** [Variable]

This is the type of the `gmnisrv` service, whose value should be a `gmnisrv-configuration` object, as in this example:

```
(service gmnisrv-service-type
  (gmnisrv-configuration
    (config-file (local-file "./my-gmnisrv.ini"))))
```

**gmnisrv-configuration** [Data Type]

Data type representing the configuration of `gmnisrv`.

```
package (default: gmnisrv)
```

Package object of the `gmnisrv` server.

`config-file` (default: `%default-gmnisrv-config-file`)

File-like object of the gmnisrv configuration file to use. The default configuration listens on port 1965 and serves files from `/srv/gemini`. Certificates are stored in `/var/lib/gemini/certs`. For more information, run `man gmnisrv` and `man gmnisrv.ini`.

## Agate

The Agate ([gemini://qwertqwefsdays.eu/agate.gmi](https://github.com/mbrubeck/agate)) (GitHub page over HTTPS (<https://github.com/mbrubeck/agate>)) program is a simple Gemini (<https://gemini.circumlunar.space/>) protocol server written in Rust.

`agate-service-type` [Variable]

This is the type of the agate service, whose value should be an `agate-service-type` object, as in this example:

```
(service agate-service-type
  (agate-configuration
    (content "/srv/gemini")
    (certificates "/srv/gemini-certs")))
```

The example above represents the minimal tweaking necessary to get Agate up and running. Specifying the path to the certificate and key directory is always necessary, as the Gemini protocol requires TLS by default.

If the specified `certificates` path is writable by Agate, and contains no valid pre-generated key and certificate, Agate will try to generate them on the first start. In this case you should pass at least one hostname using the `hostnames` option. If the specified directory is read-only, key and certificate should be pre-generated by the user.

To obtain a certificate and a key in DER format, you could, for example, use OpenSSL, running commands similar to the following example:

```
openssl genpkey -out key.der -outform DER -algorithm RSA \
  -pkeyopt rsa_keygen_bits:4096
openssl req -x509 -key key.der -outform DER -days 3650 -out cert.der \
  -subj "/CN=example.com"
```

Of course, you'll have to replace *example.com* with your own domain name, and then point the Agate configuration towards the path of the directory with the generated key and certificate using the `certificates` option.

`agate-configuration` [Data Type]

Data type representing the configuration of Agate.

`package` (default: `agate`)

The package object of the Agate server.

`content` (default: `"/srv/gemini"`)

The directory from which Agate will serve files.

`certificates` (default: `"/srv/gemini-certs"`)

Root of the certificate directory. Must be filled in with a value from the user.

**addresses** (default: '("[:]:1965" "0.0.0.0:1965"))  
A list of the addresses to listen on.

**hostnames** (default: '()')  
Virtual hosts for the Gemini server. If multiple values are specified, corresponding directory names should be present in the **content** directory. Optional.

**languages** (default: **#f**)  
RFC 4646 language code(s) for text/gemini documents. Optional.

**only-tls13?** (default: **#f**)  
Set to **#t** to disable support for TLSv1.2.

**serve-secret?** (default: **#f**)  
Set to **#t** to serve secret files (files/directories starting with a dot).

**central-configuration?** (default: **#f**)  
Set to **#t** to look for the **.meta** configuration file in the **content** root directory and will ignore **.meta** files in other directories

**ed25519?** (default: **#f**)  
Set to **#t** to generate keys using the Ed25519 signature algorithm instead of the default ECDSA.

**skip-port-check?** (default: **#f**)  
Set to **#t** to skip URL port check even when a **hostname** is specified.

**log-ip?** (default: **#t**)  
Whether or not to output IP addresses when logging.

**user** (default: "agate")  
Owner of the **agate** process.

**group** (default: "agate")  
Owner's group of the **agate** process.

**log-file** (default: "/var/log/agate.log")  
The file which should store the logging output of Agate.

## Miniflux

The Miniflux (<https://miniflux.app/>) is a minimalist RSS feed reader with a web interface.

Depending on the configuration, an initial administrator user can be pre-created on startup. To enable this, **create-admin?** should be set to **#t**, and both **admin-username-file** and **admin-password-file** should point to files containing the username and password, respectively. However, it is recommended to manually change the password to a secure one via the web UI after the initial service startup.

**miniflux-service-type** [Variable]

This is the type of the Miniflux service. Its value must be a **miniflux-configuration** record as in this example:

```
(service miniflux-service-type
```



```
(miniflux-configuration
 (listen-address "0.0.0.0:8080")
 (base-url "http://my-news-source.test")
 (create-administrator-account? #t)
 (administrator-account-name "/var/miniflux/initial-admin-username")
 (administrator-account-password "/var/miniflux/initial-admin-password"))
```

The details of the miniflux-configuration record type are given below.

**miniflux-configuration** [Data Type]

Available miniflux-configuration fields are:

**listen-address** (default: "127.0.0.1:8080") (type: string)  
Address to listen on. Use absolute path like `"/var/run/miniflux/miniflux.sock"` for a Unix socket.

**base-url** (default: "http://127.0.0.1/") (type: string)  
Base URL to generate HTML links and base path for cookies.

**create-administrator-account?** (default: #f) (type: boolean)  
Create an initial administrator account.

**administrator-account-name** (type: maybe-string-or-file-path)  
Initial administrator account name as a string or an absolute path to a file with an account name inside.

**administrator-account-password** (type: maybe-string-or-file-path)  
Initial administrator account password as a string or an absolute path to a file with a password inside.

**run-migrations?** (default: #t) (type: boolean)  
Run database migrations during application startup.

**database-url** (default: "host=/var/run/postgresql") (type: string)  
PostgreSQL connection string.

**user** (default: "miniflux") (type: string)  
User name for Postgresql and system account.

**group** (default: "miniflux") (type: string)  
Group for the system account.

**log-file** (default: "/var/log/miniflux.log") (type: string)  
Path to the log file.

**extra-settings** (type: maybe-list)  
Extra configuration parameters as a list of strings.

### 11.10.21 High Availability Services

The (gnu services high-availability) module provides RabbitMQ message broker service.

## RabbitMQ

**rabbitmq-service-type** [Variable]

Service type for the RabbitMQ (<https://www.rabbitmq.com/>) web server. The value for this service type is a <rabbitmq-configuration> record.

A simple example configuration is given below.

```
(service rabbitmq-service-type
  (rabbitmq-configuration
    (rabbitmq-configuration
      (plugins '("rabbitmq_stream"
                 "rabbitmq_management"
                 "rabbitmq_prometheus")))))
```

At startup, RabbitMQ broker logs its initialization messages into `/var/log/messages` file. Once running, the logging messages can be found in a logging file of a RabbitMQ node in `/var/log/rabbitmq` directory.

**Note:** The default configuration of the RabbitMQ service enables the RabbitMQ broker to accept connections only on loopback interfaces. However, enabling certain plugins may open ports on all network interfaces. See also RabbitMQ Port Access (<https://www.rabbitmq.com/docs/networking#ports>).

**rabbitmq-configuration** [Data Type]

This data type represents the configuration for RabbitMQ.

**rabbitmq** (default: `rabbitmq`)

The RabbitMQ package to use.

**config-file** (default: `%default-rabbitmq-config-file`)

Configuration file for the RabbitMQ broker. See also RabbitMQ Configuration (<https://www.rabbitmq.com/docs/configure>).

**data-directory** (default: `"/var/lib/rabbitmq/data"`)

Directory with RabbitMQ broker data - a schema database, message stores, cluster member information, and other persistent node state.

**plugins** (default: `'()`)

A list of RabbitMQ plugins to enable. View all available plugins using the `rabbitmq-plugins list` command.

### 11.10.22 Certificate Services

The `(gnu services certbot)` module provides a service to automatically obtain a valid TLS certificate from the Let's Encrypt certificate authority. These certificates can then be used to serve content securely over HTTPS or other TLS-based protocols, with the knowledge that the client will be able to verify the server's authenticity.

Let's Encrypt (<https://letsencrypt.org/>) provides the `certbot` tool to automate the certification process. This tool first securely generates a key on the server. It then makes a request to the Let's Encrypt certificate authority (CA) to sign the key. The CA checks

that the request originates from the host in question by using a challenge-response protocol, requiring the server to provide its response over HTTP. If that protocol completes successfully, the CA signs the key, resulting in a certificate. That certificate is valid for a limited period of time, and therefore to continue to provide TLS services, the server needs to periodically ask the CA to renew its signature.

The certbot service automates this process: the initial key generation, the initial certification request to the Let's Encrypt service, the web server challenge/response integration, writing the certificate to disk, the automated periodic renewals, and the deployment tasks associated with the renewal (e.g. reloading services, copying keys with different permissions).

Certbot is run twice a day, at a random minute within the hour. It won't do anything until your certificates are due for renewal or revoked, but running it regularly would give your service a chance of staying online in case a Let's Encrypt-initiated revocation happened for some reason.

By using this service, you agree to the ACME Subscriber Agreement, which can be found there: <https://acme-v02.api.letsencrypt.org/directory>.

**certbot-service-type** [Variable]

A service type for the `certbot` Let's Encrypt client. Its value must be a `certbot-configuration` record as in this example:

```
(service certbot-service-type
  (certbot-configuration
    (email "foo@example.net")
    (certificates
      (list
        (certificate-configuration
          (domains '("example.net" "www.example.net")))
        (certificate-configuration
          (domains '("bar.example.net"))))))))
```

See below for details about `certbot-configuration`.

**certbot-configuration** [Data Type]

Data type representing the configuration of the `certbot` service. This type has the following parameters:

**package** (default: `certbot`)

The `certbot` package to use.

**webroot** (default: `/var/www`)

The directory from which to serve the Let's Encrypt challenge/response files.

**certificates** (default: `'()`)

A list of `certificates-configurations` for which to generate certificates and request signatures. Each certificate has a **name** and several **domains**.

**email** (default: `#f`)

Optional email address used for registration and recovery contact. Setting this is encouraged as it allows you to receive important notifications about the account and issued certificates.

**server** (default: **#f**)

Optional URL of ACME server. Setting this overrides certbot's default, which is the Let's Encrypt server.

**rsa-key-size** (default: 2048)

Size of the RSA key.

**default-location** (default: *see below*)

The default **nginx-location-configuration**. Because **certbot** needs to be able to serve challenges and responses, it needs to be able to run a web server. It does so by extending the **nginx** web service with an **nginx-server-configuration** listening on the *domains* on port 80, and which has a **nginx-location-configuration** for the */.well-known/* URI path subspace used by Let's Encrypt. See Section 11.10.20 [Web Services], page 494, for more on these **nginx** configuration data types.

Requests to other URL paths will be matched by the **default-location**, which if present is added to all **nginx-server-configurations**.

By default, the **default-location** will issue a redirect from **http://domain/...** to **https://domain/...**, leaving you to define what to serve on your site via **https**.

Pass **#f** to not issue a default location.

**certificate-configuration**

[Data Type]

Data type representing the configuration of a certificate. This type has the following parameters:

**name** (default: *see below*)

This name is used by Certbot for housekeeping and in file paths; it doesn't affect the content of the certificate itself. To see certificate names, run **certbot certificates**.

Its default is the first provided domain.

**domains** (default: '()')

The first domain provided will be the subject CN of the certificate, and all domains will be Subject Alternative Names on the certificate.

**challenge** (default: **#f**)

The challenge type that has to be run by certbot. If **#f** is specified, default to the HTTP challenge. If a value is specified, defaults to the manual plugin (see **authentication-hook**, **cleanup-hook** and the documentation at <https://certbot.eff.org/docs/using.html#hooks>), and gives Let's Encrypt permission to log the public IP address of the requesting machine.

**csr** (default: **#f**)

File name of Certificate Signing Request (CSR) in DER or PEM format. If **#f** is specified, this argument will not be passed to certbot. If a value is specified, certbot will use it to obtain a certificate, instead of using a self-generated CSR. The domain-name(s) mentioned in **domains**, must be consistent with the domain-name(s) mentioned in CSR file.

**authentication-hook** (default: #f)

Command to be run in a shell once for each certificate challenge to be answered. For this command, the shell variable `$CERTBOT_DOMAIN` will contain the domain being authenticated, `$CERTBOT_VALIDATION` contains the validation string and `$CERTBOT_TOKEN` contains the file name of the resource requested when performing an HTTP-01 challenge.

**cleanup-hook** (default: #f)

Command to be run in a shell once for each certificate challenge that have been answered by the **authentication-hook**. For this command, the shell variables available in the **authentication-hook** script are still available, and additionally `$CERTBOT_AUTH_OUTPUT` will contain the standard output of the **authentication-hook** script.

**deploy-hook** (default: #f)

Command to be run in a shell once for each successfully issued certificate. For this command, the shell variable `$RENEWED_LINEAGE` will point to the config live subdirectory (for example, `"/etc/letsencrypt/live/example.com"`) containing the new certificates and keys; the shell variable `$RENEWED_DOMAINS` will contain a space-delimited list of renewed certificate domains (for example, `"example.com www.example.com"`).

**dry-run?** (default: #f)

Communicate with the ACME server but do not update certificates nor trigger **deploy-hook**. This is useful as a temporary setting to test the challenge procedure, especially the **authentication-hook** and **cleanup-hook** while working on them. It's also a good idea to use the Let's Encrypt staging server at <https://acme-staging-v02.api.letsencrypt.org/directory> while testing, which allows for higher rate limits, but with which **certbot** will helpfully refuse to update certificates and recommend the **dry-run?** option. For example:

```
(define %authentication-hook
  (program-file "authentication-hook"
    #~(let ((domain (getenv "CERTBOT_DOMAIN"))
            (token (getenv "CERTBOT_TOKEN")))
        (format #t "Hey, can you authenticate ~a with ~a for me?"
                 domain token))))

(define %cleanup-hook
  (program-file "authentication-hook"
    #~(display "Bye")))

(service certbot-service-type
  (certbot-configuration
    (server "https://acme-staging-v02.api.letsencrypt.org/directo
    (certificates
    (list
```

```

(certificate-configuration
  (dry-run? #t)
  (authentication-hook %authentication-hook)
  (cleanup-hook %cleanup-hook)
  (domains '("example.net" "www.example.net")))))))■
start-self-signed? (default: #t)
  Whether to generate an initial self-signed certificate during system activation. This option is particularly useful to allow nginx to start before certbot has run, because certbot relies on nginx running to perform HTTP challenges.

```

For each `certificate-configuration`, the certificate is saved to `/etc/certs/name/fullchain.pem` and the key is saved to `/etc/certs/name/privkey.pem`.

### 11.10.23 DNS Services

The `(gnu services dns)` module provides services related to the *domain name system* (DNS). It provides a server service for hosting an *authoritative* DNS server for multiple zones, slave or master. This service uses Knot DNS (<https://www.knot-dns.cz/>). And also a caching and forwarding DNS server for the LAN, which uses `dnsmasq` (<http://www.thekeleys.org.uk/dnsmasq/doc.html>).

#### Knot Service

An example configuration of an authoritative server for two zones, one master and one slave, is:

```

(define-zone-entries example.org.zone
;; Name TTL Class Type Data
  ("@" "" "IN" "A" "127.0.0.1")
  ("@" "" "IN" "NS" "ns")
  ("ns" "" "IN" "A" "127.0.0.1"))

(define master-zone
  (knot-zone-configuration
    (domain "example.org")
    (zone (zone-file
            (origin "example.org")
            (entries example.org.zone)))))

(define slave-zone
  (knot-zone-configuration
    (domain "plop.org")
    (dnssec-policy "default")
    (master (list "plop-master"))))

(define plop-master
  (knot-remote-configuration
    (id "plop-master")

```

```

        (address (list "208.76.58.171"))))

(operating-system
;; ...
  (services (cons* (service knot-service-type
                          (knot-configuration
                           (remotes (list plop-master))
                           (zones (list master-zone slave-zone))))
;; ...
    %base-services)))

```

**knot-service-type** [Variable]

This is the type for the Knot DNS server.

Knot DNS is an authoritative DNS server, meaning that it can serve multiple zones, that is to say domain names you would buy from a registrar. This server is not a resolver, meaning that it can only resolve names for which it is authoritative. This server can be configured to serve zones as a master server or a slave server as a per-zone basis. Slave zones will get their data from masters, and will serve it as an authoritative server. From the point of view of a resolver, there is no difference between master and slave.

The following data types are used to configure the Knot DNS server:

**knot-key-configuration** [Data Type]

Data type representing a key. This type has the following parameters:

```

id (default: "")
    An identifier for other configuration fields to refer to this key. IDs must
    be unique and must not be empty.

algorithm (default: #f)
    The algorithm to use. Choose between #f, 'hmac-md5, 'hmac-sha1,
    'hmac-sha224, 'hmac-sha256, 'hmac-sha384 and 'hmac-sha512.

secret (default: "")
    The secret key itself.

```

**knot-acl-configuration** [Data Type]

Data type representing an Access Control List (ACL) configuration. This type has the following parameters:

```

id (default: "")
    An identifier for other configuration fields to refer to this key. IDs must
    be unique and must not be empty.

address (default: '())
    An ordered list of IP addresses, network subnets, or network ranges rep-
    resented with strings. The query must match one of them. Empty value
    means that address match is not required.

```

**key** (default: '()')

An ordered list of references to keys represented with strings. The string must match a key ID defined in a **knot-key-configuration**. No key means that a key is not required to match that ACL.

**action** (default: '()')

An ordered list of actions that are permitted or forbidden by this ACL. Possible values are lists of zero or more elements from **'transfer'**, **'notify'** and **'update'**.

**deny?** (default: **#f**)

When true, the ACL defines restrictions. Listed actions are forbidden. When false, listed actions are allowed.

**zone-entry** [Data Type]

Data type representing a record entry in a zone file. This type has the following parameters:

**name** (default: "@")

The name of the record. "@" refers to the origin of the zone. Names are relative to the origin of the zone. For example, in the **example.org** zone, **"ns.example.org"** actually refers to **ns.example.org.example.org**. Names ending with a dot are absolute, which means that **"ns.example.org."** refers to **ns.example.org**.

**ttl** (default: "")

The Time-To-Live (TTL) of this record. If not set, the default TTL is used.

**class** (default: "IN")

The class of the record. Knot currently supports only "IN" and partially "CH".

**type** (default: "A")

The type of the record. Common types include A (IPv4 address), AAAA (IPv6 address), NS (Name Server) and MX (Mail eXchange). Many other types are defined.

**data** (default: "")

The data contained in the record. For instance an IP address associated with an A record, or a domain name associated with an NS record. Remember that domain names are relative to the origin unless they end with a dot.

**zone-file** [Data Type]

Data type representing the content of a zone file. This type has the following parameters:

**entries** (default: '()')

The list of entries. The SOA record is taken care of, so you don't need to put it in the list of entries. This list should probably contain an entry for your primary authoritative DNS server. Other than using a list of entries



directly, you can use `define-zone-entries` to define a object containing the list of entries more easily, that you can later pass to the `entries` field of the `zone-file`.

`origin` (default: "")

The name of your zone. This parameter cannot be empty.

`ns` (default: "ns")

The domain of your primary authoritative DNS server. The name is relative to the origin, unless it ends with a dot. It is mandatory that this primary DNS server corresponds to an NS record in the zone and that it is associated to an IP address in the list of entries.

`mail` (default: "hostmaster")

An email address people can contact you at, as the owner of the zone. This is translated as `<mail>@<origin>`.

`serial` (default: 1)

The serial number of the zone. As this is used to keep track of changes by both slaves and resolvers, it is mandatory that it *never* decreases. Always increment it when you make a change in your zone.

`refresh` (default: (\* 2 24 3600))

The frequency at which slaves will do a zone transfer. This value is a number of seconds. It can be computed by multiplications or with `(string->duration)`.

`retry` (default: (\* 15 60))

The period after which a slave will retry to contact its master when it fails to do so a first time.

`expiry` (default: (\* 14 24 3600))

Default TTL of records. Existing records are considered correct for at most this amount of time. After this period, resolvers will invalidate their cache and check again that it still exists.

`nx` (default: 3600)

Default TTL of inexistent records. This delay is usually short because you want your new domains to reach everyone quickly.

## `knot-remote-configuration`

[Data Type]

Data type representing a remote configuration. This type has the following parameters:

`id` (default: "")

An identifier for other configuration fields to refer to this remote. IDs must be unique and must not be empty.

`address` (default: '()')

An ordered list of destination IP addresses. Addresses are tried in sequence. An optional port can be given with the @ separator. For instance: `(list "1.2.3.4" "2.3.4.5@53")`. Default port is 53.

**via** (default: '()')

An ordered list of source IP addresses. An empty list will have Knot choose an appropriate source IP. An optional port can be given with the @ separator. The default is to choose at random.

**key** (default: #f)

A reference to a key, that is a string containing the identifier of a key defined in a `knot-key-configuration` field.

### **knot-keystore-configuration**

[Data Type]

Data type representing a keystore to hold dnssec keys. This type has the following parameters:

**id** (default: "")

The id of the keystore. It must not be empty.

**backend** (default: 'pem')

The backend to store the keys in. Can be 'pem' or 'pkcs11'.

**config** (default: "/var/lib/knot/keys/keys")

The configuration string of the backend. An example for the PKCS#11 is: "pkcs11:token=knot;pin-value=1234/gnu/store/.../lib/pkcs11/libsoftsm2.so". For the pem backend, the string represents a path in the file system.

### **knot-policy-configuration**

[Data Type]

Data type representing a dnssec policy. Knot DNS is able to automatically sign your zones. It can either generate and manage your keys automatically or use keys that you generate.

Dnssec is usually implemented using two keys: a Key Signing Key (KSK) that is used to sign the second, and a Zone Signing Key (ZSK) that is used to sign the zone. In order to be trusted, the KSK needs to be present in the parent zone (usually a top-level domain). If your registrar supports dnssec, you will have to send them your KSK's hash so they can add a DS record in their zone. This is not automated and need to be done each time you change your KSK.

The policy also defines the lifetime of keys. Usually, ZSK can be changed easily and use weaker cryptographic functions (they use lower parameters) in order to sign records quickly, so they are changed often. The KSK however requires manual interaction with the registrar, so they are changed less often and use stronger parameters because they sign only one record.

This type has the following parameters:

**id** (default: "")

The id of the policy. It must not be empty.

**keystore** (default: "default")

A reference to a keystore, that is a string containing the identifier of a keystore defined in a `knot-keystore-configuration` field. The "default" identifier means the default keystore (a kasp database that was setup by this service).

**manual?** (default: **#f**)  
Whether the key management is manual or automatic.

**single-type-signing?** (default: **#f**)  
When **#t**, use the Single-Type Signing Scheme.

**algorithm** (default: "ecdsap256sha256")  
An algorithm of signing keys and issued signatures.

**ksk-size** (default: 256)  
The length of the KSK. Note that this value is correct for the default algorithm, but would be unsecure for other algorithms.

**zsk-size** (default: 256)  
The length of the ZSK. Note that this value is correct for the default algorithm, but would be unsecure for other algorithms.

**dnskey-ttl** (default: 'default')  
The TTL value for DNSKEY records added into zone apex. The special 'default' value means same as the zone SOA TTL.

**zsk-lifetime** (default: (\* 30 24 3600))  
The period between ZSK publication and the next rollover initiation.

**propagation-delay** (default: (\* 24 3600))  
An extra delay added for each key rollover step. This value should be high enough to cover propagation of data from the master server to all slaves.

**rrsig-lifetime** (default: (\* 14 24 3600))  
A validity period of newly issued signatures.

**rrsig-refresh** (default: (\* 7 24 3600))  
A period how long before a signature expiration the signature will be refreshed.

**nsec3?** (default: **#f**)  
When **#t**, NSEC3 will be used instead of NSEC.

**nsec3-iterations** (default: 5)  
The number of additional times the hashing is performed.

**nsec3-salt-length** (default: 8)  
The length of a salt field in octets, which is appended to the original owner name before hashing.

**nsec3-salt-lifetime** (default: (\* 30 24 3600))  
The validity period of newly issued salt field.

**knot-zone-configuration** [Data Type]

Data type representing a zone served by Knot. This type has the following parameters:

**domain** (default: "")  
The domain served by this configuration. It must not be empty.

- file** (default: "")  
The file where this zone is saved. This parameter is ignored by master zones. Empty means default location that depends on the domain name.
- zone** (default: (zone-file))  
The content of the zone file. This parameter is ignored by slave zones. It must contain a zone-file record.
- master** (default: '()')  
A list of master remotes. When empty, this zone is a master. When set, this zone is a slave. This is a list of remotes identifiers.
- ddns-master** (default: #f)  
The main master. When empty, it defaults to the first master in the list of masters.
- notify** (default: '()')  
A list of slave remote identifiers.
- acl** (default: '()')  
A list of acl identifiers.
- semantic-checks?** (default: #f)  
When set, this adds more semantic checks to the zone.
- zonefile-sync** (default: 0)  
The delay between a modification in memory and on disk. 0 means immediate synchronization.
- zonefile-load** (default: #f)  
The way the zone file contents are applied during zone load. Possible values are:
- #f for using the default value from Knot,
  - 'none for not using the zone file at all,
  - 'difference for computing the difference between already available contents and zone contents and applying it to the current zone contents,
  - 'difference-no-serial for the same as 'difference, but ignoring the SOA serial in the zone file, while the server takes care of it automatically.
  - 'whole for loading zone contents from the zone file.
- journal-content** (default: #f)  
The way the journal is used to store zone and its changes. Possible values are 'none to not use it at all, 'changes to store changes and 'all to store contents. #f does not set this option, so the default value from Knot is used.
- max-journal-usage** (default: #f)  
The maximum size for the journal on disk. #f does not set this option, so the default value from Knot is used.

**max-journal-depth** (default: **#f**)

The maximum size of the history. **#f** does not set this option, so the default value from Knot is used.

**max-zone-size** (default: **#f**)

The maximum size of the zone file. This limit is enforced for incoming transfer and updates. **#f** does not set this option, so the default value from Knot is used.

**dnssec-policy** (default: **#f**)

A reference to a **knot-policy-configuration** record, or the special name "default". If the value is **#f**, there is no dnssec signing on this zone.

**serial-policy** (default: 'increment')

A policy between 'increment and 'unixtime.

**knot-configuration**

[Data Type]

Data type representing the Knot configuration. This type has the following parameters:

**knot** (default: **knot**)

The Knot package.

**run-directory** (default: "/var/run/knot")

The run directory. This directory will be used for pid file and sockets.

**includes** (default: '()')

A list of strings or file-like objects denoting other files that must be included at the top of the configuration file.

This can be used to manage secrets out-of-band. For example, secret keys may be stored in an out-of-band file not managed by Guix, and thus not visible in **/gnu/store**—e.g., you could store secret key configuration in **/etc/knot/secrets.conf** and add this file to the **includes** list.

One can generate a secret tsig key (for nsupdate and zone transfers with the **keymgr** command from the knot package. Note that the package is not automatically installed by the service. The following example shows how to generate a new tsig key:

```
keymgr -t mysecret > /etc/knot/secrets.conf
chmod 600 /etc/knot/secrets.conf
```

Also note that the generated key will be named *mysecret*, so it is the name that needs to be used in the **key** field of the **knot-acl-configuration** record and in other places that need to refer to that key.

It can also be used to add configuration not supported by this interface.

**listen-v4** (default: "0.0.0.0")

An ip address on which to listen.

**listen-v6** (default: ":::")

An ip address on which to listen.

**listen-port** (default: 53)  
A port on which to listen.

**keys** (default: '()')  
The list of knot-key-configuration used by this configuration.

**acls** (default: '()')  
The list of knot-acl-configuration used by this configuration.

**remotes** (default: '()')  
The list of knot-remote-configuration used by this configuration.

**zones** (default: '()')  
The list of knot-zone-configuration used by this configuration.

## Knot Resolver Service

**knot-resolver-service-type** [Variable]

This is the type of the knot resolver service, whose value should be a **knot-resolver-configuration** object as in this example:

```
(service knot-resolver-service-type
  (knot-resolver-configuration
    (kresd-config-file (plain-file "kresd.conf" "
net.listen('192.168.0.1', 5353)
user('knot-resolver', 'knot-resolver')
modules = { 'hints > iterate', 'stats', 'predict' }
cache.size = 100 * MB
""))))
```

For more information, refer its manual (<https://knot-resolver.readthedocs.io/en/stable/config-overview.html>).

**knot-resolver-configuration** [Data Type]

Data type representing the configuration of knot-resolver.

**package** (default: *knot-resolver*)  
Package object of the knot DNS resolver.

**kresd-config-file** (default: %kresd.conf)  
File-like object of the kresd configuration file to use, by default it will listen on 127.0.0.1 and ::1.

**garbage-collection-interval** (default: 1000)  
Number of milliseconds for **kres-cache-gc** to periodically trim the cache.

## Dnsmasq Service

**dnsmasq-service-type** [Variable]

This is the type of the dnsmasq service, whose value should be a **dnsmasq-configuration** object as in this example:

```
(service dnsmasq-service-type
  (dnsmasq-configuration
```

```
(no-resolv? #t)
(servers '("192.168.1.1"))))
```

`dnsmasq-service-type` provides two actions:

**reload** Send a `SIGHUP` signal to `dnsmasq` to clear cache and reload hosts files which triggers the following actions:

- Clear the DNS cache
- Reload `/etc/hosts` and `/etc/ethers`
- Reload files specified by `--dhcp-hostsfile`, `--dhcp-hostsdir`, `--dhcp-optsfile`, `--dhcp-optsdir`, `--addn-hosts`, or `--hostsdir`
- Execute DHCP lease change scripts for all existing leases
- Reread `/etc/resolv.conf` (only if `--no-poll` is configured).

**Note:** The `SIGHUP` signal does not trigger a reload of the main configuration file.

Example usage:

```
herd reload dnsmasq
```

**stats** When `dnsmasq` receives a `SIGUSR1` signal, it writes statistics to the system log including:

- Cache size and number of evicted entries
- Cache hits, misses, and authoritative queries answered
- For each upstream server: queries sent and errors received
- TCP connection fork statistics.

Example usage:

```
herd stats dnsmasq
```

## `dnsmasq-configuration`

[Data Type]

Data type representing the configuration of `dnsmasq`.

**package** (default: `dnsmasq`)

Package object of the `dnsmasq` server.

**shepherd-provision** (default: `'(dnsmasq)'`)

**shepherd-requirement** (default: `'(user-processes networking)'`)

This option can be used to provide a list of Shepherd service names (symbols) provided by this service. You might want to change the default value if you intend to run several `dnsmasq` instances.

Likewise, **shepherd-requirement** is a list of Shepherd service names (symbols) that this service will depend on.

**pid-file** (default: `"/run/dnsmasq.pid"`)

Specify an alternate path for `dnsmasq` to record its process-id in.

**no-hosts?** (default: `#f`)

When true, don't read the hostnames in `/etc/hosts`.

**port** (default: 53)  
 The port to listen on. Setting this to zero completely disables DNS responses, leaving only DHCP and/or TFTP functions.

**local-service?** (default: #t)  
 Accept DNS queries only from hosts whose address is on a local subnet, ie a subnet for which an interface exists on the server.

**listen-addresses** (default: '())  
 Listen on the given IP addresses.

**resolve-file** (default: "/etc/resolv.conf")  
 The file to read the IP address of the upstream nameservers from.

**no-resolve?** (default: #f)  
 When true, don't read *resolve-file*.

**forward-private-reverse-lookup?** (default: #t)  
 When false, all reverse lookups for private IP ranges are answered with "no such domain" rather than being forwarded upstream.

**query-servers-in-order?** (default: #f)  
 When true, dnsmasq queries the servers in the same order as they appear in *servers*.

**servers** (default: '())  
 Specify IP address of upstream servers directly.

**servers-file** (default: #f)  
 Specify file containing upstream servers. This file is re-read when dnsmasq receives SIGHUP. Could be either a string or a file-like object.

**addresses** (default: '())  
 For each entry, specify an IP address to return for any host in the given domains. Queries in the domains are never forwarded and always replied to with the specified IP address.  
 This is useful for redirecting hosts locally, for example:

```
(service dnsmasq-service-type
  (dnsmasq-configuration
    (addresses
      '(; Redirect to a local web-server.
        "/example.org/127.0.0.1"
        ; Redirect subdomain to a specific IP.
        "/subdomain.example.org/192.168.1.42")))))
```

Note that rules in */etc/hosts* take precedence over this.

**cache-size** (default: 150)  
 Set the size of dnsmasq's cache. Setting the cache size to zero disables caching.

**negative-cache?** (default: #t)  
 When false, disable negative caching.



- cpe-id** (default: **#f**)  
If set, add a CPE (Customer-Premises Equipment) identifier to DNS queries which are forwarded upstream.
- tftp-enable?** (default: **#f**)  
Whether to enable the built-in TFTP server.
- tftp-no-fail?** (default: **#f**)  
If true, does not fail dnsmasq if the TFTP server could not start up.
- tftp-single-port?** (default: **#f**)  
Whether to use only one single port for TFTP.
- tftp-secure?** (default: **#f**)  
If true, only files owned by the user running the dnsmasq process are accessible.  
If dnsmasq is being run as root, different rules apply: **tftp-secure?** has no effect, but only files which have the world-readable bit set are accessible.
- tftp-max** (default: **#f**)  
If set, sets the maximal number of concurrent connections allowed.
- tftp-mtu** (default: **#f**)  
If set, sets the MTU for TFTP packets to that value.
- tftp-no-blocksize?** (default: **#f**)  
If true, stops the TFTP server from negotiating the blocksize with a client.
- tftp-lowercase?** (default: **#f**)  
Whether to convert all filenames in TFTP requests to lowercase.
- tftp-port-range** (default: **#f**)  
If set, fixes the dynamical ports (one per client) to the given range ("**<start>**,**<end>**").
- tftp-root** (default: **/var/empty,lo**)  
Look for files to transfer using TFTP relative to the given directory. When this is set, TFTP paths which include **'..'** are rejected, to stop clients getting outside the specified root. Absolute paths (starting with **'/'**) are allowed, but they must be within the TFTP-root. If the optional interface argument is given, the directory is only used for TFTP requests via that interface.
- tftp-unique-root** (default: **#f**)  
If set, add the IP or hardware address of the TFTP client as a path component on the end of the TFTP-root. Only valid if a TFTP root is set and the directory exists. Defaults to adding IP address (in standard dotted-quad format).  
For instance, if **--tftp-root** is **'/tftp'** and client **'1.2.3.4'** requests file **myfile** then the effective path will be **/tftp/1.2.3.4/myfile** if

`/tftp/1.2.3.4` exists or `/tftp/myfile` otherwise. When `'=mac'` is specified it will append the MAC address instead, using lowercase zero padded digits separated by dashes, e.g.: `'01-02-03-04-aa-bb'`. Note that resolving MAC addresses is only possible if the client is in the local network or obtained a DHCP lease from `dnsmasq`.

`conf-file` (default: `'()`)

File names of configuration files to use. Strings and file-like objects are accepted.

`conf-dir` (default: `#f`)

Read all the files in the given directory as configuration files. Support for `dnsmasq` field extensions is not implemented; users can use `computed-file` to build a directory of files instead. The configuration files are loaded in alphabetical order of file name.

`extra-options` (default: `'()`)

This option provides an “escape hatch” for the user to provide arbitrary command-line arguments to `dnsmasq` as a list of strings.

## Unbound Service

`unbound-service-type`

[Variable]

This is the type of the service to run Unbound (<https://www.unbound.net>), a validating, recursive, and caching DNS resolver. Its value must be a `unbound-configuration` object as in this example:

```
(service unbound-service-type
  (unbound-configuration
    (forward-zone
      (list
        (unbound-zone
          (name ".")
          (forward-addr '("149.112.112.112#dns.quad9.net"
                        "2620:fe::9#dns.quad9.net"))
          (forward-tls-upstream #t))))))
```

`unbound-configuration`

[Data Type]

Available `unbound-configuration` fields are:

`server` (type: `unbound-server`)

General options for the Unbound server.

`remote-control` (type: `unbound-remote`)

Remote control options for the daemon.

`forward-zone` (default: `()`) (type: `list-of-unbound-zone`)

A zone for which queries should be forwarded to another resolver.

`extra-content` (type: `maybe-string`)

Raw content to add to the configuration file.

**unbound-server** [Data Type]

Available **unbound-server** fields are:

- interface** (type: maybe-list-of-strings)  
Interfaces listened on for queries from clients.
- hide-version** (type: maybe-boolean)  
Refuse the version.server and version.bind queries.
- hide-identity** (type: maybe-boolean)  
Refuse the id.server and hostname.bind queries.
- tls-cert-bundle** (type: maybe-string)  
Certificate bundle file, used for DNS over TLS.
- extra-options** (default: ()) (type: alist)  
An association list of options to append.

**unbound-remote** [Data Type]

Available **unbound-remote** fields are:

- control-enable** (type: maybe-boolean)  
Enable remote control.
- control-interface** (type: maybe-string)  
IP address or local socket path to listen on for remote control.
- extra-options** (default: ()) (type: alist)  
An association list of options to append.

**unbound-zone** [Data Type]

Available **unbound-zone** fields are:

- name** (type: string)  
Zone name.
- forward-addr** (type: maybe-list-of-strings)  
IP address of server to forward to.
- forward-tls-upstream** (type: maybe-boolean)  
Whether the queries to this forwarder use TLS for transport.
- extra-options** (default: ()) (type: alist)  
An association list of options to append.

**11.10.24 VNC Services**

The (`gnu services vnc`) module provides services related to *Virtual Network Computing* (VNC), which makes it possible to locally use graphical Xorg applications running on a remote machine. Combined with a graphical manager that supports the *X Display Manager Control Protocol*, such as GDM (see [gdm], page 355) or LightDM (see [lightdm], page 360), it is possible to remote an entire desktop for a multi-user environment.

## Xvnc

Xvnc is a VNC server that spawns its own X window server; which means it can run on headless servers. The Xvnc implementations provided by the `tigervnc-server` and `turbovnc` aim to be fast and efficient.

`xvnc-service-type` [Variable]

The `xvnc-service-type` service can be configured via the `xvnc-configuration` record, documented below. A second virtual display could be made available on a remote machine via the following configuration:

```
(service xvnc-service-type
  (xvnc-configuration (display-number 10)))
```

As a demonstration, the `xclock` command could then be started on the remote machine on display number 10, and it could be displayed locally via the `vncviewer` command:

```
# Start xclock on the remote machine.
ssh -L5910:localhost:5910 your-host -- guix shell xclock \
  -- env DISPLAY=:10 xclock
# Access it via VNC.
guix shell tigervnc-client -- vncviewer localhost:5910
```

The following configuration combines XDMCP and Inetd to allow multiple users to concurrently use the remote system and login graphically via the GDM display manager:

```
(operating-system
  [...]
  (services (cons*
    [...]
    (service xvnc-service-type (xvnc-configuration
      (display-number 5)
      (localhost? #f)
      (xdmcp? #t)
      (inetd? #t)))
    (modify-services %desktop-services
      (gdm-service-type config => (gdm-configuration
        (inherit config)
        (auto-suspend? #f)
        (xdmcp? #t)))))))
```

A remote user could then connect to it by using the `vncviewer` command or a compatible VNC client and start a desktop session of their choosing:

```
vncviewer remote-host:5905
```

**Warning:** Unless your machine is in a controlled environment, for security reasons, the `localhost?` configuration of the `xvnc-configuration` record should be left to its default `#t` value and exposed via a secure means such as an SSH port forward. The XDMCP port, UDP 177 should also be blocked from the outside by a firewall, as it is not a secure protocol and can expose login credentials in clear.

**xvnc-configuration** [Data Type]

Available **xvnc-configuration** fields are:

**xvnc** (default: **tigervnc-server**) (type: file-like)

The package that provides the Xvnc binary.

**display-number** (default: 0) (type: number)

The display number used by Xvnc. You should set this to a number not already used a Xorg server.

**geometry** (default: "1024x768") (type: string)

The size of the desktop to be created.

**depth** (default: 24) (type: color-depth)

The pixel depth in bits of the desktop to be created. Accepted values are 16, 24 or 32.

**port** (type: maybe-port)

The port on which to listen for connections from viewers. When left unspecified, it defaults to 5900 plus the display number.

**ipv4?** (default: **#t**) (type: boolean)

Use IPv4 for incoming and outgoing connections.

**ipv6?** (default: **#t**) (type: boolean)

Use IPv6 for incoming and outgoing connections.

**password-file** (type: maybe-string)

The password file to use, if any. Refer to `vnpasswd(1)` to learn how to generate such a file.

**xdmcp?** (default: **#f**) (type: boolean)

Query the XDMCP server for a session. This enables users to log in a desktop session from the login manager screen. For a multiple users scenario, you'll want to enable the **inetd?** option as well, so that each connection to the VNC server is handled separately rather than shared.

**inetd?** (default: **#f**) (type: boolean)

Use an Inetd-style service, which runs the Xvnc server on demand.

**frame-rate** (default: 60) (type: number)

The maximum number of updates per second sent to each client.

**security-types** (default: **'("None")**) (type: security-types)

The allowed security schemes to use for incoming connections. The default is "None", which is safe given that Xvnc is configured to authenticate the user via the display manager, and only for local connections. Accepted values are any of the following: ("None" "VncAuth" "Plain" "TLSNone" "TLSVnc" "TLSPlain" "X509None" "X509Vnc")

**localhost?** (default: **#t**) (type: boolean)

Only allow connections from the same machine. It is set to **#true** by default for security, which means SSH or another secure means should be used to expose the remote port.

- log-level** (default: 30) (type: log-level)  
 The log level, a number between 0 and 100, 100 meaning most verbose output. The log messages are output to syslog.
- extra-options** (default: '()') (type: strings)  
 This can be used to provide extra Xvnc options not exposed via this <xvnc-configuration> record.

### 11.10.25 VPN Services

The (`gnu services vpn`) module provides services related to *virtual private networks* (VPNs).

#### Bitmask

- bitmask-service-type** [Variable]  
 A service type for the Bitmask (<https://bitmask.net>) VPN client. It makes the client available in the system and loads its polkit policy. Please note that the client expects an active polkit-agent, which is either run by your desktop-environment or should be run manually.

#### OpenVPN

It provides a *client* service for your machine to connect to a VPN, and a *server* service for your machine to host a VPN. Both `openvpn-client-service-type` and `openvpn-server-service-type` can be run simultaneously.

- openvpn-client-service-type** [Variable]  
 Type of the service that runs `openvpn`, a VPN daemon, as a client.  
 The value for this service is a <openvpn-client-configuration> object.

- openvpn-server-service-type** [Variable]  
 Type of the service that runs `openvpn`, a VPN daemon, as a server.  
 The value for this service is a <openvpn-server-configuration> object.

- openvpn-client-configuration** [Data Type]  
 Available `openvpn-client-configuration` fields are:

- openvpn** (default: `openvpn`) (type: file-like)  
 The OpenVPN package.
- pid-file** (default: `"/var/run/openvpn/openvpn.pid"`) (type: string)  
 The OpenVPN pid file.
- proto** (default: `udp`) (type: proto)  
 The protocol (UDP or TCP) used to open a channel between clients and servers.
- dev** (default: `tun`) (type: dev)  
 The device type used to represent the VPN connection.
- ca** (default: `"/etc/openvpn/ca.crt"`) (type: maybe-string)  
 The certificate authority to check connections against.

**cert** (default: `"/etc/openvpn/client.crt"`) (type: maybe-string)  
 The certificate of the machine the daemon is running on. It should be signed by the authority given in **ca**.

**key** (default: `"/etc/openvpn/client.key"`) (type: maybe-string)  
 The key of the machine the daemon is running on. It must be the key whose certificate is **cert**.

**comp-lzo?** (default: `#t`) (type: boolean)  
 Whether to use the lzo compression algorithm.

**persist-key?** (default: `#t`) (type: boolean)  
 Don't re-read key files across SIGUSR1 or `-ping-restart`.

**persist-tun?** (default: `#t`) (type: boolean)  
 Don't close and reopen TUN/TAP device or run up/down scripts across SIGUSR1 or `-ping-restart` restarts.

**fast-io?** (default: `#f`) (type: boolean)  
 (Experimental) Optimize TUN/TAP/UDP I/O writes by avoiding a call to poll/epoll/select prior to the write operation.

**verbosity** (default: 3) (type: number)  
 Verbosity level.

**tls-auth** (default: `#f`) (type: tls-auth-client)  
 Add an additional layer of HMAC authentication on top of the TLS control channel to protect against DoS attacks.

**auth-user-pass** (type: maybe-string)  
 Authenticate with server using username/password. The option is a file containing username/password on 2 lines. Do not use a file-like object as it would be added to the store and readable by any user.

**verify-key-usage?** (default: `#t`) (type: key-usage)  
 Whether to check the server certificate has server usage extension.

**bind?** (default: `#f`) (type: bind)  
 Bind to a specific local port number.

**resolv-retry?** (default: `#t`) (type: resolv-retry)  
 Retry resolving server address.

**remote** (default: `'()`) (type: openvpn-remote-list)  
 A list of remote servers to connect to.

**openvpn-remote-configuration** [Data Type]  
 Available **openvpn-remote-configuration** fields are:

**name** (default: `"my-server"`) (type: string)  
 Server name.

**port** (default: 1194) (type: number)  
 Port number the server listens to.

**openvpn-server-configuration** [Data Type]

Available **openvpn-server-configuration** fields are:

**openvpn** (default: **openvpn**) (type: file-like)

The OpenVPN package.

**pid-file** (default: **"/var/run/openvpn/openvpn.pid"**) (type: string)

The OpenVPN pid file.

**proto** (default: **udp**) (type: proto)

The protocol (UDP or TCP) used to open a channel between clients and servers.

**dev** (default: **tun**) (type: dev)

The device type used to represent the VPN connection.

**ca** (default: **"/etc/openvpn/ca.crt"**) (type: maybe-string)

The certificate authority to check connections against.

**cert** (default: **"/etc/openvpn/client.crt"**) (type: maybe-string)

The certificate of the machine the daemon is running on. It should be signed by the authority given in **ca**.

**key** (default: **"/etc/openvpn/client.key"**) (type: maybe-string)

The key of the machine the daemon is running on. It must be the key whose certificate is **cert**.

**comp-lzo?** (default: **#t**) (type: boolean)

Whether to use the lzo compression algorithm.

**persist-key?** (default: **#t**) (type: boolean)

Don't re-read key files across SIGUSR1 or **-ping-restart**.

**persist-tun?** (default: **#t**) (type: boolean)

Don't close and reopen TUN/TAP device or run up/down scripts across SIGUSR1 or **-ping-restart** restarts.

**fast-io?** (default: **#f**) (type: boolean)

(Experimental) Optimize TUN/TAP/UDP I/O writes by avoiding a call to poll/epoll/select prior to the write operation.

**verbosity** (default: **3**) (type: number)

Verbosity level.

**tls-auth** (default: **#f**) (type: tls-auth-server)

Add an additional layer of HMAC authentication on top of the TLS control channel to protect against DoS attacks.

**port** (default: **1194**) (type: number)

Specifies the port number on which the server listens.

**server** (default: **"10.8.0.0 255.255.255.0"**) (type: ip-mask)

An ip and mask specifying the subnet inside the virtual network.

**server-ipv6** (default: **#f**) (type: cidr6)

A CIDR notation specifying the IPv6 subnet inside the virtual network.



**dh** (default: `"/etc/openvpn/dh2048.pem"`) (type: string)  
 The Diffie-Hellman parameters file.

**ifconfig-pool-persist** (default: `"/etc/openvpn/ipp.txt"`) (type: string)  
 The file that records client IPs.

**redirect-gateway?** (default: `#f`) (type: gateway)  
 When true, the server will act as a gateway for its clients.

**client-to-client?** (default: `#f`) (type: boolean)  
 When true, clients are allowed to talk to each other inside the VPN.

**keepalive** (default: `(10 120)`) (type: keepalive)  
 Causes ping-like messages to be sent back and forth over the link so that each side knows when the other side has gone down. **keepalive** requires a pair. The first element is the period of the ping sending, and the second element is the timeout before considering the other side down.

**max-clients** (default: `100`) (type: number)  
 The maximum number of clients.

**status** (default: `"/var/run/openvpn/status"`) (type: string)  
 The status file. This file shows a small report on current connection. It is truncated and rewritten every minute.

**client-config-dir** (default: `'()`) (type: openvpn-ccd-list)  
 The list of configuration for some clients.

## strongSwan

Currently, the strongSwan service only provides legacy-style configuration with `ipsec.conf` and `ipsec.secrets` files.

**strongswan-service-type** [Variable]  
 A service type for configuring strongSwan for IPsec VPN (Virtual Private Networking). Its value must be a **strongswan-configuration** record as in this example:

```
(service strongswan-service-type
  (strongswan-configuration
    (ipsec-conf "/etc/ipsec.conf")
    (ipsec-secrets "/etc/ipsec.secrets"))))
```

**strongswan-configuration** [Data Type]  
 Data type representing the configuration of the StrongSwan service.

**strongswan**  
 The strongSwan package to use for this service.

**ipsec-conf** (default: `#f`)  
 The file name of your `ipsec.conf`. If not `#f`, then this and **ipsec-secrets** must both be strings.

**ipsec-secrets** (default `#f`)  
 The file name of your `ipsec.secrets`. If not `#f`, then this and **ipsec-conf** must both be strings.

## Wireguard

**wireguard-service-type** [Variable]

A service type for a Wireguard tunnel interface. Its value must be a **wireguard-configuration** record as in this example:

```
(service wireguard-service-type
  (wireguard-configuration
    (peers
      (list
        (wireguard-peer
          (name "my-peer")
          (endpoint "my.wireguard.com:51820")
          (public-key "hzipKg9X1yqu1axN6iJp0mWf6BZGo8m1wteKwtTmDGF4=")
          (allowed-ips '("10.0.0.2/32"))))))))
```

**wireguard-configuration** [Data Type]

Data type representing the configuration of the Wireguard service.

**wireguard**

The wireguard package to use for this service.

**interface** (default: "wg0")

The interface name for the VPN.

**addresses** (default: '("10.0.0.1/32"))

List of strings or G-expressions which represent the IP addresses to be assigned to the above interface.

**port** (default: 51820)

The port on which to listen for incoming connections.

**dns** (default: '())

List of strings or G-expressions which represent the DNS server(s) to announce to VPN clients via DHCP.

**monitor-ips?** (default: #f)

Whether to monitor the resolved Internet addresses (IPs) of the endpoints of the configured peers, resetting the peer endpoints using an IP address that no longer correspond to their freshly resolved host name. Set this to **#t** if one or more endpoints use host names provided by a dynamic DNS service to keep the sessions alive.

**monitor-ips-interval** (default: "\*/5 \* \* \* \*")

This is the monitoring schedule, expressed as a string in traditional cron syntax or as a gexp evaluating to a Shepherd calendar event (see Section “Timers” in *The GNU Shepherd Manual*).

**private-key** (default: "/etc/wireguard/private.key")

The private key file for the interface. It is automatically generated if the file does not exist. If this field is **#f**, a private key is not automatically created and the path is not serialized to the configuration file.

**bootstrap-private-key?** (default: #t)

Whether or not the private key should be generated automatically if it does not exist.

Setting this to #f allows one to set the private key using command substitution. One example shown in the `wg-quick(8)` manual is retrieving a private key using `password-store`. This can be achieved with the following code:

```
(wireguard-configuration
  (private-key
    #~(string-append "<("
                      #$(file-append password-store "/bin/pass")
                      ;; Wireguard replaces %i with the interface name.
                      " WireGuard/private-keys/%i"))))
```

**peers** (default: '())

The authorized peers on this interface. This is a list of *wireguard-peer* records.

**pre-up** (default: '())

List of strings or G-expressions. These are script snippets which will be executed before setting up the interface.

**post-up** (default: '())

List of strings or G-expressions. These are script snippets which will be executed after setting up the interface.

**pre-down** (default: '())

List of strings or G-expressions. These are script snippets which will be executed before tearing down the interface.

**post-down** (default: '())

List of strings or G-expressions. These are script snippets which will be executed after tearing down the interface.

**shepherd-requirement** (default: '()) (type: list-of-symbols)

A list of symbols naming Shepherd services that this service will depend on.

**table** (default: "auto")

The routing table to which routes are added, as a string. There are two special values: "off" that disables the creation of routes altogether, and "auto" (the default) that adds routes to the default table and enables special handling of default routes.

**auto-start?** (default: #t)

Whether the Wireguard network should be started automatically by the Shepherd. If it is #f the service has to be started manually with `herd start wireguard-$interface` (for example: `herd start wireguard-wg0`).

**wireguard-peer** [Data Type]  
 Data type representing a Wireguard peer attached to a given interface.

**name** The peer name.

**endpoint** (default: #f)  
 The optional endpoint for the peer, such as "demo.wireguard.com:51820".

**public-key**  
 The peer public-key represented as a base64 string.

**preshared-key** (default: #f)  
 An optional pre-shared key file for this peer that can be either a string or a G-expression. The given file will not be autogenerated.

**allowed-ips**  
 A list of IP addresses from which incoming traffic for this peer is allowed and to which incoming traffic for this peer is directed.

**keep-alive** (default: #f)  
 An optional time interval in seconds. A packet will be sent to the server endpoint once per time interval. This helps receiving incoming connections from this peer when you are behind a NAT or a firewall.

### 11.10.26 DAV Services

**xandikos-service-type** [Variable]  
 This service starts **xandikos**, a lightweight CardDAV/CalDAV server backed by a Git repository.  
 The service's value is a **xandikos-configuration** record.

**xandikos-configuration** [Data Type]  
 This is the data type representing the configuration for the **xandikos-shepherd-service**.  
 It has the following parameters:

**package** (default: **xandikos**)  
 The **xandikos** package to use.

**directory** (default: **"/var/xandikos/dav"**)  
 The directory to serve from.

**listen-address** (default: **127.0.0.1**)  
 The address **xandikos** listens on.

**port** (default: **8080**)  
 The port to run **xandikos** on.

**current-user-principal** (default: **"/user/"**)  
 Path to current user principal.

**route-prefix** (default: **"/"**)  
 Path to **xandikos** (useful when Xandikos is behind a reverse proxy).

**defaults?** (default: **#t**)  
 Create initial calendar and address book. Implies **--autocreate**.

**dump?** (default: **#f**)  
 Print DAV XML request/responses.

**avahi?** (default: **#f**)  
 Announce services with avahi.

**autocreate?** (default: **#f**)  
 Automatically create necessary directories.

**no-strict?** (default: **#f**)  
 Enable workarounds for buggy CalDAV/CardDAV client implementations.

### 11.10.27 Network File System

The (**gnu services nfs**) module provides the following services, which are most commonly used in relation to mounting or exporting directory trees as *network file systems* (NFS).

While it is possible to use the individual components that together make up a Network File System service, we recommended to configure an NFS server with the **nfs-service-type**.

#### NFS Service

The NFS service takes care of setting up all NFS component services, kernel configuration file systems, and installs configuration files in the locations that NFS expects.

**nfs-service-type** [Variable]  
 A service type for a complete NFS server.

**nfs-configuration** [Data Type]  
 This data type represents the configuration of the NFS service and all of its subsystems.

It has the following parameters:

**nfs-utils** (default: **nfs-utils**)  
 The nfs-utils package to use.

**nfs-versions** (default: **'("4.2" "4.1" "4.0")**)  
 If a list of string values is provided, the **rpc.nfsd** daemon will be limited to supporting the given versions of the NFS protocol.

**exports** (default: **'()**)  
 This is a list of directories the NFS server should export. Each entry is a list consisting of two elements: a directory name and a string containing all options. This is an example in which the directory **/export** is served to all NFS clients as a read-only share:

```
(nfs-configuration
 (exports
  '("/export"
    "(ro,insecure,no_subtree_check,crossmnt,fsid=0)"))))■
```

**rpcmountd-port** (default: **#f**)  
 The network port that the **rpc.mountd** daemon should use.

**rpcstatd-port** (default: **#f**)  
 The network port that the **rpc.statd** daemon should use.

**rpcbind** (default: **rpcbind**)  
 The rpcbind package to use.

**idmap-domain** (default: "localdomain")  
 The local NFSv4 domain name.

**nfsd-port** (default: 2049)  
 The network port that the **nfsd** daemon should use.

**nfsd-threads** (default: 8)  
 The number of threads used by the **nfsd** daemon.

**nfsd-tcp?** (default: **#t**)  
 Whether the **nfsd** daemon should listen on a TCP socket.

**nfsd-udp?** (default: **#f**)  
 Whether the **nfsd** daemon should listen on a UDP socket.

**pipefs-directory** (default: "/var/lib/nfs/rpc\_pipefs")  
 The directory where the pipefs file system is mounted.

**debug** (default: "()")  
 A list of subsystems for which debugging output should be enabled. This is a list of symbols. Any of these symbols are valid: **nfsd**, **nfs**, **rpc**, **idmap**, **statd**, or **mountd**.

If you don't need a complete NFS service or prefer to build it yourself you can use the individual component services that are documented below.

## RPC Bind Service

The RPC Bind service provides a facility to map program numbers into universal addresses. Many NFS related services use this facility. Hence it is automatically started when a dependent service starts.

**rpcbind-service-type** [Variable]  
 A service type for the RPC portmapper daemon.

**rpcbind-configuration** [Data Type]  
 Data type representing the configuration of the RPC Bind Service. This type has the following parameters:

**rpcbind** (default: **rpcbind**)  
 The rpcbind package to use.

**warm-start?** (default: **#t**)  
 If this parameter is **#t**, then the daemon will read a state file on startup thus reloading state information saved by a previous instance.

## Pipefs Pseudo File System

The pipefs file system is used to transfer NFS related data between the kernel and user space programs.

**pipefs-service-type** [Variable]

A service type for the pipefs pseudo file system.

**pipefs-configuration** [Data Type]

Data type representing the configuration of the pipefs pseudo file system service. This type has the following parameters:

**mount-point** (default: `"/var/lib/nfs/rpc_pipefs"`)

The directory to which the file system is to be attached.

## GSS Daemon Service

The *global security system* (GSS) daemon provides strong security for RPC based protocols. Before exchanging RPC requests an RPC client must establish a security context. Typically this is done using the Kerberos command `kinit` or automatically at login time using PAM services (see Section 11.10.18 [Kerberos Services], page 483).

**gss-service-type** [Variable]

A service type for the Global Security System (GSS) daemon.

**gss-configuration** [Data Type]

Data type representing the configuration of the GSS daemon service. This type has the following parameters:

**nfs-utils** (default: `nfs-utils`)

The package in which the `rpc.gssd` command is to be found.

**pipefs-directory** (default: `"/var/lib/nfs/rpc_pipefs"`)

The directory where the pipefs file system is mounted.

## IDMAP Daemon Service

The idmap daemon service provides mapping between user IDs and user names. Typically it is required in order to access file systems mounted via NFSv4.

**idmap-service-type** [Variable]

A service type for the Identity Mapper (IDMAP) daemon.

**idmap-configuration** [Data Type]

Data type representing the configuration of the IDMAP daemon service. This type has the following parameters:

**nfs-utils** (default: `nfs-utils`)

The package in which the `rpc.idmapd` command is to be found.

**pipefs-directory** (default: `"/var/lib/nfs/rpc_pipefs"`)

The directory where the pipefs file system is mounted.

**domain** (default: `#f`)

The local NFSv4 domain name. This must be a string or `#f`. If it is `#f` then the daemon will use the host's fully qualified domain name.

**verbosity** (default: 0)  
The verbosity level of the daemon.

## Autofs Service

The autofs service automatically mounts and unmounts filesystems. It can mount both network and local storage, but is most often used to manage NFS mounts.

A typical configuration might look like:

```
(operating-system
  (services
    (append
      (list
        (service nfs-service-type (nfs-configuration))
        (service autofs-service-type
          (autofs-configuration
            (mounts
              (list
                (autofs-indirect-map
                  (mount-point "/net/home")
                  (entries
                    (list
                      ;; Accessing /net/home/foo will mount
                      ;; big.network.server:/home/foo.
                      (autofs-map-entry
                        (type "nfs")
                        (device "big.network.server:/home/&")
                        (mount-point "*"))))))))))))
    %base-services)))
```

**autofs-indirect-map** [Data Type]

Available **autofs-indirect-map** fields are:

**mount-point** (type: string)  
Where to mount the indirect map.

**entries** (default: ()) (type: autofs-map-entries)  
Entries in this map.

**autofs-map-entry** [Data Type]

Available **autofs-map-entry** fields are:

**type** (default: "auto") (type: string)  
The type of the filesystem.

**device** (type: string)  
Device or remote host to mount. May contain special character `&`, which can be referenced in the *mount-point* field.

**mount-point** (type: string)  
Directory to mount this device on. Map entries come in two flavors: direct and indirect. Direct entries map a single device to a single mountpoint,



while indirect entries can map multiple devices to multiple mountpoints. A direct entry has a *mount-point* beginning with /, representing the absolute path of the directory to mount the device on. For example:

```
(autofs-map-entry
 (type "ext4")
 (device "/dev/sdb1")
 (mount-point "/mnt/external-disk"))
```

An indirect entry has a *mount-point* not beginning with /, representing the subdirectory within the parent indirect map for this entry. Indirect maps may also use the special character \*, which will be replaced with the value of special character & in the *device* field of this entry. For example:

```
(autofs-indirect-map
 (mount-point "/devices")
 (entries
 (list
 ;; Automount any block device r/o by ID.
 (autofs-map-entry
 (type "auto")
 (mount-point "ro/uuid/*")
 (device "/dev/disk/by-id/&")
 (options '(ro)))
 ;; Automount any block device by UUID.
 (autofs-map-entry
 (type "auto")
 (mount-point "rw/uuid/*")
 (device "/dev/disk/by-uuid/&")))))
```

**options** (default: ()) (type: file-system-options)

List of mount options. Some options are simple flags, such as `ro`, `noexec`, `nosuid`, etc. These may be expressed as strings or symbols. Other options also accept a value. These are expressed as pairs of (`option . value`). `option` may be a string or symbol, as with flags. `value` may be a string, symbol, or number. Example: (`ro (errors . remount-ro) noexec`)

## autofs-configuration

[Data Type]

Available **autofs-configuration** fields are:

**autofs** (default: **autofs**) (type: package)

The autofs package to use.

**mounts** (default: ()) (type: autofs-mount-maps)

Mount maps to manage. This is a list of either direct map entries or indirect mount maps.

**timeout** (default: 300) (type: conf-integer)

Sets the default mount timeout in seconds.

**master-wait** (default: 10) (type: conf-integer)

Sets the default maximum number of retries (actual iterations is half this, each is delayed by 2 seconds before retrying) waiting for the master map

to become available if it cannot be read at program start. This can be longer if the map source itself waits for availability (such as sss).

**negative-timeout** (default: 60) (type: conf-integer)

Set the default timeout for caching failed key lookups.

**mount-verbose?** (default: #f) (type: conf-boolean)

Use the verbose flag when spawning mount(8), and log some process info about the requestor and its parent.

**mount-wait** (default: -1) (type: conf-integer)

Set the default time to wait for a response from a spawned mount(8) before sending it a SIGTERM. Note that we still need to wait for the RPC layer to timeout before the sub-process exits so this isn't ideal but it is the best we can do. The default is to wait until mount(8) returns without intervention.

**umount-wait** (default: 12) (type: conf-integer)

Set the default time to wait for a response from a spawned umount(8) before sending it a SIGTERM. Note that we still need to wait for the RPC layer to timeout before the sub-process exits so this isn't ideal but it is the best we can do.

**browse-mode?** (default: #t) (type: conf-boolean)

Should maps be browsable by default?

**mount-nfs-default-protocol** (default: 3) (type: conf-integer)

Set the default protocol that mount.nfs(8) uses when performing a mount. Autofs needs to know the default NFS protocol that mount.nfs(8) uses so it can do special case handling for its availability probe for different NFS protocols. Since we can't identify this default automatically we need to set it in the autofs configuration.

**append-options?** (default: #t) (type: conf-boolean)

When #t, global options are appended to map entry options. When #f, map entry options replace the global options.

**logging** (type: conf-log-level)

Default log level. May be #:none, #:verbose, or #:debug.

**force-standard-program-map-env?** (default: #f) (type: conf-boolean)

Override the use of a prefix with standard environment variables when a program map is executed. Since program maps are run as the privileged user setting these standard environment variables opens automount(8) to potential user privilege escalation when the program map is written in a language that can load components from, for example, a user home directory.

**map-hash-table-size** (default: 1024) (type: conf-integer)

This configuration option may be used to change the number of hash table slots. This configuration option affects the overhead of searching the map entry cache for map entries when there are a large number of entries. It

affects the number of entries that must be looked at to locate a map entry in the map entry cache. For example, the default of 1024 and a direct map with 8000 entries would result in each slot containing an average of 8 entries, which should be acceptable. However, if excessive CPU usage is observed during automount lookups increasing this option can reduce the CPU overhead considerably because it reduces the length of the search chains. Note that the number of entries in a map doesn't necessarily relate to the number of entries used in the map entry cache. There are two distinct cases where the map hash table can make a significant difference, direct maps and indirect maps that use the "browse" option. For indirect maps that do not use the "browse" option entries are added to the map entry cache at lookup so the number of active cache entries, in this case, is usually much less than the number of entries in the map. In this last case it would be unusual for the map entry cache to grow large enough to warrant increasing the default before an event that cleans stale entries, a map re-read for example.

**use-hostname-for-mounts?** (default: #f) (type: conf-boolean)

NFS mounts where the host name resolves to more than one IP address are probed for availability and to establish the order in which mounts to them should be tried. To ensure that mount attempts are made only to hosts that are responding and are tried in the order of hosts with the quickest response the IP address of the host needs to be used for the mount. If it is necessary to use the hostname given in the map entry for the mount regardless, then set this option to #t. Be aware that if this is done there is no defense against the host name resolving to one that isn't responding and while the number of attempts at a successful mount will correspond to the number of addresses the host name resolves to the order will also not correspond to fastest responding hosts.

**disable-not-found-message?** (default: #f) (type: conf-boolean)

The original request to add this log message needed it to be unconditional. That produces, IMHO, unnecessary noise in the log so a configuration option has been added to provide the ability to turn it off.

**use-ignore-mount-option?** (default: #f) (type: conf-boolean)

An option to enable the use of autofs pseudo option "disable". This option is used as a hint to user space that the mount entry should be omitted from mount table listings. The default is #f to avoid unexpected changes in behaviour and so is an opt-in setting.

**sss-master-map-wait** (default: 0) (type: conf-integer)

Set the time to wait and retry if sssd is unable to read the master map at program start. Default is 0 (don't wait) or 10 if sss supports returning EHSTDOWN when the provider isn't available. If the sss library supports returning EHOSTDOWN when the provider is down then this value is how long to wait between retries reading the master map. When reading dependent maps or looking up a map key this value is multiplied by the number of retries that would be used when reading the master map.

**use-mount-request-log-id?** (default: **#f**) (type: **conf-boolean**)  
 Set whether to use a mount request log id so that log entries for specific mount requests can be easily identified in logs that have multiple concurrent requests.

### 11.10.28 Samba Services

The (**gnu services samba**) module provides service definitions for Samba as well as additional helper services. Currently it provides the following services.

#### Samba

Samba (<https://www.samba.org>) provides network shares for folders and printers using the SMB/CIFS protocol commonly used on Windows. It can also act as an Active Directory Domain Controller (AD DC) for other hosts in an heterogeneous network with different types of Computer systems.

**samba-service-type** [Variable]

The service type to enable the samba services **samba**, **nmbd**, **smbd** and **winbindd**. By default this service type does not run any of the Samba daemons; they must be enabled individually.

Below is a basic example that configures a simple, anonymous (unauthenticated) Samba file share exposing the **/public** directory.

**Tip:** The **/public** directory and its contents must be world readable/writable, so you'll want to run **'chmod -R 777 /public'** on it.

**Caution:** Such a Samba configuration should only be used in controlled environments, and you should not share any private files using it, as anyone connecting to your network would be able to access them.

```
(service samba-service-type (samba-configuration
                             (enable-smbd? #t)
                             (config-file (plain-file "smb.conf" "\
[global]
map to guest = Bad User
logging = syslog@1

[public]
browsable = yes
path = /public
read only = no
guest ok = yes
guest only = yes\n")))))
```

**samba-service-configuration** [Data Type]

Configuration record for the Samba suite.

**package** (default: **samba**)  
 The samba package to use.

**config-file** (default: **#f**)  
 The config file to use. To learn about its syntax, run **'man smb.conf'**.

**enable-samba?** (default: **#f**)  
 Enable the **samba** daemon.

**enable-smbd?** (default: **#f**)  
 Enable the **smbd** daemon.

**enable-nmbd?** (default: **#f**)  
 Enable the **nmbd** daemon.

**enable-winbindd?** (default: **#f**)  
 Enable the **winbindd** daemon.

## Web Service Discovery Daemon

The WSD (Web Service Discovery daemon) implements the Web Services Dynamic Discovery (<http://docs.oasis-open.org/ws-dd/discovery/1.1/os/wsdd-discovery-1.1-spec-os.html>) protocol that enables host discovery over Multicast DNS, similar to what Avahi does. It is a drop-in replacement for SMB hosts that have had SMBv1 disabled for security reasons.

**wsdd-service-type** [Variable]  
 Service type for the WSD host daemon. The value for this service type is a **wsdd-configuration** record. The details for the **wsdd-configuration** record type are given below.

**wsdd-configuration** [Data Type]  
 This data type represents the configuration for the wsdd service.

**package** (default: **wsdd**)  
 The wsdd package to use.

**ipv4only?** (default: **#f**)  
 Only listen to IPv4 addresses.

**ipv6only** (default: **#f**)  
 Only listen to IPv6 addresses. Please note: Activating both options is not possible, since there would be no IP versions to listen to.

**chroot** (default: **#f**)  
 Chroot into a separate directory to prevent access to other directories. This is to increase security in case there is a vulnerability in **wsdd**.

**hop-limit** (default: **1**)  
 Limit to the level of hops for multicast packets. The default is *1* which should prevent packets from leaving the local network.

**interface** (default: **'()**)  
 Limit to the given list of interfaces to listen to. By default wsdd will listen to all interfaces. Except the loopback interface is never used.

**uuid-device** (default: **#f**)  
 The WSD protocol requires a device to have a UUID. Set this to manually assign the service a UUID.

- domain** (default: **#f**)  
 Notify this host is a member of an Active Directory.
- host-name** (default: **#f**)  
 Manually set the hostname rather than letting **wsdd** inherit this host's hostname. Only the host name part of a possible FQDN will be used in the default case.
- preserve-case?** (default: **#f**)  
 By default **wsdd** will convert the hostname in workgroup to all uppercase. The opposite is true for hostnames in domains. Setting this parameter will preserve case.
- workgroup** (default: **"WORKGROUP"**)  
 Change the name of the workgroup. By default **wsdd** reports this host being member of a workgroup.

### 11.10.29 Continuous Integration

Cuirass (<https://guix.gnu.org/cuirass/>) is a continuous integration tool for Guix. It can be used both for development and for providing substitutes to others (see Section 5.3 [Substitutes], page 46).

The (**gnu services cuirass**) module provides the following service.

**cuirass-service-type** [Procedure]  
 The type of the Cuirass service. Its value must be a **cuirass-configuration** object, as described below.

To add build jobs, you have to set the **specifications** field of the configuration. For instance, the following example will build all the packages provided by the **my-channel** channel.

```
(define %cuirass-specs
  #~(list (specification
    (name 'my-channel)
    (build '(channels my-channel))
    (channels
      (cons (channel
        (name 'my-channel)
        (url "https://my-channel.git"))
        %default-channels))))))

(service cuirass-service-type
  (cuirass-configuration
    (specifications %cuirass-specs)))
```

To build the **linux-libre** package defined by the default Guix channel, one can use the following configuration.

```
(define %cuirass-specs
  #~(list (specification
    (name 'my-linux)
```

```
(build '(packages "linux-libre")))))

(service cuirass-service-type
  (cuirass-configuration
    (specifications %cuirass-specs)))
```

The other configuration possibilities, as well as the specification record itself are described in the Cuirass manual (see Section “Specifications” in *Cuirass*).

While information related to build jobs is located directly in the specifications, global settings for the `cuirass` process are accessible in other `cuirass-configuration` fields.

**cuirass-configuration** [Data Type]

Data type representing the configuration of Cuirass.

**cuirass** (default: `cuirass`)

The Cuirass package to use.

**log-file** (default: `"/var/log/cuirass.log"`)

Location of the log file.

**web-log-file** (default: `"/var/log/cuirass-web.log"`)

Location of the log file used by the web interface.

**cache-directory** (default: `"/var/cache/cuirass"`)

Location of the repository cache.

**user** (default: `"cuirass"`)

Owner of the `cuirass` process.

**group** (default: `"cuirass"`)

Owner's group of the `cuirass` process.

**interval** (default: 300)

Number of seconds between the poll of the repositories followed by the Cuirass jobs, by default.

Note that individual jobset specifications may override this in their `period` field; the value specified here is only used for specifications that use the default `period` value, which is zero.

**parallel-evaluations** (default: `#f`)

Maximum number of jobset evaluations that may run concurrently. If `#f`, that number is determined based on a number of CPU cores, with an upper bound of 8.

**ttl** (default: 2592000)

Duration to keep build results' GC roots alive, in seconds.

**evaluation-ttl** (default: 24 months)

Number of seconds after which an evaluation may be deleted from the database.

**build-expiry** (default: 4 months)

Duration in seconds after which pending builds are canceled. This helps ensure that the backlog does not grow indefinitely.

**threads** (default: `#f`)

**web-threads** (default: `#f`)  
 Number of kernel threads to use for Cuirass (main process) and for the web interface (the `cuirass web` program). The default value should be appropriate for most cases.

**parameters** (default: `#f`)  
 Read parameters from the given *parameters* file. The supported parameters are described here (see Section “Parameters” in *Cuirass*).

**remote-server** (default: `#f`)  
 A `cuirass-remote-server-configuration` record to use the build remote mechanism or `#f` to use the default build mechanism.

**database** (default: `"dbname=cuirass host=/var/run/postgresql"`)  
 Use *database* as the database containing the jobs and the past build results. Since Cuirass uses PostgreSQL as a database engine, *database* must be a string such as `"dbname=cuirass host=localhost"`.

**port** (default: `8081`)  
 Port number used by the HTTP server.

**host** (default: `"localhost"`)  
 Listen on the network interface for *host*. The default is to accept connections from localhost.

**specifications** (default: `#~'()`)  
 A gexp (see Section 8.12 [G-Expressions], page 165) that evaluates to a list of specifications records. The specification record is described in the Cuirass manual (see Section “Specifications” in *Cuirass*).

**one-shot?** (default: `#f`)  
 Only evaluate specifications and build derivations once.

**fallback?** (default: `#f`)  
 When substituting a pre-built binary fails, fall back to building packages locally.

**extra-options** (default: `'()`)  
 Extra options to pass when running the `cuirass register` process.

**web-extra-options** (default: `'()`)  
 Extra options to pass when running the `cuirass web` process.

## Cuirass remote building

Cuirass supports two mechanisms to build derivations.

- Using the local Guix daemon. This is the default build mechanism. Once the build jobs are evaluated, they are sent to the local Guix daemon. Cuirass then listens to the Guix daemon output to detect the various build events.
- Using the remote build mechanism. The build jobs are not submitted to the local Guix daemon. Instead, one “remote server” dispatches build requests to the connected “remote workers”, according to the build priorities.



To enable this build mode a `cuirass-remote-server-configuration` record must be passed as `remote-server` argument of the `cuirass-configuration` record. The `cuirass-remote-server-configuration` record is described below.

This build mode scales way better than the default build mode. This is the build mode that is used on the GNU Guix build farm at <https://ci.guix.gnu.org>. It should be preferred when using Cuirass to build large amount of packages.

**cuirass-remote-server-configuration** [Data Type]

Data type representing the configuration of the Cuirass remote-server.

**backend-port** (default: 5555)

The TCP port for communicating with `remote-worker` processes using ZMQ. It defaults to 5555.

**log-port** (default: 5556)

The TCP port of the log server. It defaults to 5556.

**publish-port** (default: 5557)

The TCP port of the publish server. It defaults to 5557.

**log-file** (default: `"/var/log/cuirass-remote-server.log"`)

Location of the log file.

**cache** (default: `"/var/cache/cuirass/remote"`)

Use *cache* directory to cache build log files.

**log-expiry** (default: 6 months)

The duration in seconds after which build logs collected by `cuirass remote-worker` may be deleted.

**trigger-url** (default: `#f`)

Once a substitute is successfully fetched, trigger substitute baking at *trigger-url*.

**publish?** (default: `#t`)

If set to false, do not start a publish server and ignore the `publish-port` argument. This can be useful if there is already a standalone publish server standing next to the remote server.

**public-key**

**private-key**

Use the specific *files* as the public/private key pair used to sign the store items being published.

At least one remote worker must also be started on any machine of the local network to actually perform the builds and report their status.

**cuirass-remote-worker-configuration** [Data Type]

Data type representing the configuration of the Cuirass remote-worker.

**cuirass** (default: `cuirass`)

The Cuirass package to use.

**workers** (default: 1)  
Start *workers* parallel workers.

**server** (default: #f)  
Do not use Avahi discovery and connect to the given **server** IP address instead.

**systems** (default: (list (%current-system)))  
Only request builds for the given *systems*.

**log-file** (default: "/var/log/cuirass-remote-worker.log")  
Location of the log file.

**publish-port** (default: 5558)  
The TCP port of the publish server. It defaults to 5558.

**substitute-urls** (default: %default-substitute-urls)  
The list of URLs where to look for substitutes by default.

**public-key**  
**private-key**  
Use the specific *files* as the public/private key pair used to sign the store items being published.

## Laminar

Laminar (<https://laminar.ohwg.net/>) is a lightweight and modular Continuous Integration service. It doesn't have a configuration web UI instead uses version-controllable configuration files and scripts.

Laminar encourages the use of existing tools such as bash and cron instead of reinventing them.

**laminar-service-type** [Variable]  
The type of the Laminar service. Its value must be a **laminar-configuration** object, as described below.

All configuration values have defaults, a minimal configuration to get Laminar running is shown below. By default, the web interface is available on port 8080.

```
(service laminar-service-type)
```

**laminar-configuration** [Data Type]  
Data type representing the configuration of Laminar.

**laminar** (default: laminar)  
The Laminar package to use.

**home-directory** (default: "/var/lib/laminar")  
The directory for job configurations and run directories.

**supplementary-groups** (default: ())  
Supplementary groups for the Laminar user account.

**bind-http** (default: "\*:8080")  
The interface/port or unix socket on which laminard should listen for incoming connections to the web frontend.

**bind-rpc** (default: "unix-abstract:laminar")  
 The interface/port or unix socket on which laminard should listen for incoming commands such as build triggers.

**title** (default: "Laminar")  
 The page title to show in the web frontend.

**keep-rundirs** (default: 0)  
 Set to an integer defining how many rundirs to keep per job. The lowest-numbered ones will be deleted. The default is 0, meaning all run dirs will be immediately deleted.

**archive-url** (default: #f)  
 The web frontend served by laminard will use this URL to form links to artefacts archived jobs.

**base-url** (default: #f)  
 Base URL to use for links to laminar itself.

## Forgejo Runner

The (gnu services ci) also provides a service for Forgejo Runner (<https://code.forgejo.org/forgejo/runner>), a daemon that connects to an instance of the Forgejo code collaboration tool (<https://forgejo.org>) and runs jobs for continuous integration.

A minimal configuration mostly with default values that can be added to the **services** field of your operating system looks like this:

```
(service forgejo-runner-service-type
  (forgejo-runner-configuration
    (name "my-runner")
    (labels '("guix" "linux"))))
```

This provides a **forgejo-runner** Shepherd service. That service will initially fail to start; you will have to manually *register* the runner against the Forgejo server by running a command like:

```
herd register forgejo-runner url token
```

... where the arguments are as follows:

**url** the URL of the Forgejo server—e.g., 'https://codeberg.org';

**token** an access token provided by the Forgejo server (<https://forgejo.org/docs/latest/admin/runner-installation/#standard-registration>).

Once registration has succeeded, you can start the runner:

```
herd enable forgejo-runner
herd start forgejo-runner
```

The runner then receives orders from the Forgejo server to execute *actions*. Actions are commands and workflows specified by YAML files in the **.forgejo/workflows** directory of source code repositories—see the Forgejo Action documentation (<https://forgejo.org/docs/v7.0/user/actions/>) for more info.

Note that at the moment **forgejo-runner-service-type** lets you run only one runner. Details about the configuration of this service follow.

**forgejo-runner-service-type** [Variable]

This is the service type for Forgejo Runner. Its value must be a **forgejo-runner-configuration** record, documented below.

**forgejo-runner-configuration** [Data Type]

This data type represents the configuration of an instance of **forgejo-runner-service-type**. It contains the following fields:

**package** (default: **forgejo-runner**)

The Forgejo Runner package to use.

**name** (default: **#~(gethostname)**)

Name of the runner as will be shown in the runner management interface of Forgejo.

**labels** (default: **'("guix")'**)

List of labels (<https://forgejo.org/docs/latest/admin/actions/choosing-labels>) representing the type of environment the runner provides and that actions may refer to.

**capacity** (default: **1**)

Number of tasks to be executed concurrently.

**timeout** (default: **(\* 3 3600)**)

Maximum duration of a job, in seconds.

**fetch-timeout** (default: **5**)

Maximum duration for fetching the job from the Forgejo server, in seconds.

**fetch-interval** (default: **2**)

Interval (in seconds) for fetching the job from the Forgejo server.

**report-interval** (default: **1**)

Interval (in seconds) for reporting the job status and log to the Forgejo server.

**data-directory** (default: **"/var/lib/forgejo-runner"**)

Directory where **forgejo-runner** will store persistent data such as its configuration and access token.

**run-directory** (default: **"/var/run/forgejo-runner"**)

Directory where **forgejo-runner** stores cached data.

**log-file** (default: **"/var/log/forgejo-runner.log"**)

Name of the **forgejo-runner** log file.

### 11.10.30 Power Management Services

#### Power Profiles Daemon

The (**gnu services pm**) module provides a Guix service definition for the Linux Power Profiles Daemon, which makes power profiles handling available over D-Bus.

The available profiles consist of the default ‘**balanced**’ mode, a ‘**power-saver**’ mode and on supported systems a ‘**performance**’ mode.

**Important:** The `power-profiles-daemon` conflicts with other power management tools like `tlp`. Using both together is not recommended.

`power-profiles-daemon-service-type` [Variable]

This is the service type for the Power Profiles Daemon (<https://gitlab.freedesktop.org/upower/power-profiles-daemon/>). The value for this service is a `power-profiles-daemon-configuration`.

To enable the Power Profiles Daemon with default configuration add this line to your services:

```
(service power-profiles-daemon-service-type)
```

`power-profiles-daemon-configuration` [Data Type]

Data type representing the configuration of `power-profiles-daemon-service-type`.

`power-profiles-daemon` (default: `power-profiles-daemon`) (type: file-like)

Package object of `power-profiles-daemon`.

## TLP daemon

The `(gnu services pm)` module provides a Guix service definition for the Linux power management tool TLP.

TLP enables various powersaving modes in userspace and kernel. Contrary to `upower-service`, it is not a passive, monitoring tool, as it will apply custom settings each time a new power source is detected. More information can be found at TLP home page (<https://linrunner.de/en/tlp/tlp.html>).

`tlp-service-type` [Variable]

The service type for the TLP tool. The default settings are optimised for battery life on most systems, but you can tweak them to your heart's content by adding a valid `tlp-configuration`:

```
(service tlp-service-type
  (tlp-configuration
    (cpu-scaling-governor-on-ac (list "performance"))
    (sched-powersave-on-bat? #t)))
```

Each parameter definition is preceded by its type; for example, `'boolean foo'` indicates that the `foo` parameter should be specified as a boolean. Types starting with `maybe-` denote parameters that won't show up in TLP config file when their value is left unset, or is explicitly set to the `%unset-value` value.

Available `tlp-configuration` fields are:

`package tlp` [tlp-configuration parameter]

The TLP package.

`boolean tlp-enable?` [tlp-configuration parameter]

Set to true if you wish to enable TLP.

Defaults to `'#t'`.

- string tlp-default-mode** [tlp-configuration parameter]  
Default mode when no power supply can be detected. Alternatives are AC and BAT.  
Defaults to 'AC'.
- non-negative-integer disk-idle-secs-on-ac** [tlp-configuration parameter]  
Number of seconds Linux kernel has to wait after the disk goes idle, before syncing on AC.  
Defaults to '0'.
- non-negative-integer disk-idle-secs-on-bat** [tlp-configuration parameter]  
Same as disk-idle-ac but on BAT mode.  
Defaults to '2'.
- non-negative-integer max-lost-work-secs-on-ac** [tlp-configuration parameter]  
Dirty pages flushing periodicity, expressed in seconds.  
Defaults to '15'.
- non-negative-integer max-lost-work-secs-on-bat** [tlp-configuration parameter]  
Same as max-lost-work-secs-on-ac but on BAT mode.  
Defaults to '60'.
- maybe-space-separated-string-list cpu-scaling-governor-on-ac** [tlp-configuration parameter]  
CPU frequency scaling governor on AC mode. With intel\_pstate driver, alternatives are powersave and performance. With acpi-cpufreq driver, alternatives are ondemand, powersave, performance and conservative.  
Defaults to 'disabled'.
- maybe-space-separated-string-list cpu-scaling-governor-on-bat** [tlp-configuration parameter]  
Same as cpu-scaling-governor-on-ac but on BAT mode.  
Defaults to 'disabled'.
- maybe-non-negative-integer cpu-scaling-min-freq-on-ac** [tlp-configuration parameter]  
Set the min available frequency for the scaling governor on AC.  
Defaults to 'disabled'.
- maybe-non-negative-integer cpu-scaling-max-freq-on-ac** [tlp-configuration parameter]  
Set the max available frequency for the scaling governor on AC.  
Defaults to 'disabled'.

- maybe-non-negative-integer** [tlp-configuration parameter]  
**cpu-scaling-min-freq-on-bat**  
Set the min available frequency for the scaling governor on BAT.  
Defaults to 'disabled'.
- maybe-non-negative-integer** [tlp-configuration parameter]  
**cpu-scaling-max-freq-on-bat**  
Set the max available frequency for the scaling governor on BAT.  
Defaults to 'disabled'.
- maybe-non-negative-integer** [tlp-configuration parameter]  
**cpu-min-perf-on-ac**  
Limit the min P-state to control the power dissipation of the CPU, in AC mode.  
Values are stated as a percentage of the available performance.  
Defaults to 'disabled'.
- maybe-non-negative-integer** [tlp-configuration parameter]  
**cpu-max-perf-on-ac**  
Limit the max P-state to control the power dissipation of the CPU, in AC mode.  
Values are stated as a percentage of the available performance.  
Defaults to 'disabled'.
- maybe-non-negative-integer** [tlp-configuration parameter]  
**cpu-min-perf-on-bat**  
Same as **cpu-min-perf-on-ac** on BAT mode.  
Defaults to 'disabled'.
- maybe-non-negative-integer** [tlp-configuration parameter]  
**cpu-max-perf-on-bat**  
Same as **cpu-max-perf-on-ac** on BAT mode.  
Defaults to 'disabled'.
- maybe-boolean** **cpu-boost-on-ac?** [tlp-configuration parameter]  
Enable CPU turbo boost feature on AC mode.  
Defaults to 'disabled'.
- maybe-boolean** **cpu-boost-on-bat?** [tlp-configuration parameter]  
Same as **cpu-boost-on-ac?** on BAT mode.  
Defaults to 'disabled'.
- boolean** **sched-powersave-on-ac?** [tlp-configuration parameter]  
Allow Linux kernel to minimize the number of CPU cores/hyper-threads used under light load conditions.  
Defaults to '#f'.
- boolean** **sched-powersave-on-bat?** [tlp-configuration parameter]  
Same as **sched-powersave-on-ac?** but on BAT mode.  
Defaults to '#t'.

<b>boolean</b> <code>nmi-watchdog?</code>	[tlp-configuration parameter]
Enable Linux kernel NMI watchdog. Defaults to ‘#f’.	
<b>maybe-string</b> <code>phc-controls</code>	[tlp-configuration parameter]
For Linux kernels with PHC patch applied, change CPU voltages. An example value would be “F:V F:V F:V F:V”. Defaults to ‘disabled’.	
<b>string</b> <code>energy-perf-policy-on-ac</code>	[tlp-configuration parameter]
Set CPU performance versus energy saving policy on AC. Alternatives are performance, normal, powersave. Defaults to “performance”.	
<b>string</b> <code>energy-perf-policy-on-bat</code>	[tlp-configuration parameter]
Same as <code>energy-perf-policy-ac</code> but on BAT mode. Defaults to “powersave”.	
<b>space-separated-string-list</b> <code>disks-devices</code>	[tlp-configuration parameter]
Hard disk devices.	
<b>space-separated-string-list</b> <code>disk-apm-level-on-ac</code>	[tlp-configuration parameter]
Hard disk advanced power management level.	
<b>space-separated-string-list</b> <code>disk-apm-level-on-bat</code>	[tlp-configuration parameter]
Same as <code>disk-apm-bat</code> but on BAT mode.	
<b>maybe-space-separated-string-list</b> <code>disk-spindown-timeout-on-ac</code>	[tlp-configuration parameter]
Hard disk spin down timeout. One value has to be specified for each declared hard disk. Defaults to ‘disabled’.	
<b>maybe-space-separated-string-list</b> <code>disk-spindown-timeout-on-bat</code>	[tlp-configuration parameter]
Same as <code>disk-spindown-timeout-on-ac</code> but on BAT mode. Defaults to ‘disabled’.	
<b>maybe-space-separated-string-list</b> <code>disk-iosched</code>	[tlp-configuration parameter]
Select IO scheduler for disk devices. One value has to be specified for each declared hard disk. Example alternatives are cfq, deadline and noop. Defaults to ‘disabled’.	
<b>string</b> <code>sata-linkpwr-on-ac</code>	[tlp-configuration parameter]
SATA aggressive link power management (ALPM) level. Alternatives are min_power, medium_power, max_performance. Defaults to “max_performance”.	



<b>string</b> <code>sata-linkpwr-on-bat</code>	[tlp-configuration parameter]
Same as <code>sata-linkpwr-ac</code> but on BAT mode.	
Defaults to <code>"min_power"</code> .	
<b>maybe-string</b> <code>sata-linkpwr-blacklist</code>	[tlp-configuration parameter]
Exclude specified SATA host devices for link power management.	
Defaults to <code>'disabled'</code> .	
<b>maybe-on-off-boolean</b> <code>ahci-runtime-pm-on-ac?</code>	[tlp-configuration parameter]
Enable Runtime Power Management for AHCI controller and disks on AC mode.	
Defaults to <code>'disabled'</code> .	
<b>maybe-on-off-boolean</b> <code>ahci-runtime-pm-on-bat?</code>	[tlp-configuration parameter]
Same as <code>ahci-runtime-pm-on-ac</code> on BAT mode.	
Defaults to <code>'disabled'</code> .	
<b>non-negative-integer</b> <code>ahci-runtime-pm-timeout</code>	[tlp-configuration parameter]
Seconds of inactivity before disk is suspended.	
Defaults to <code>'15'</code> .	
<b>string</b> <code>pcie-aspm-on-ac</code>	[tlp-configuration parameter]
PCI Express Active State Power Management level. Alternatives are <code>default</code> , <code>performance</code> , <code>powersave</code> .	
Defaults to <code>"performance"</code> .	
<b>string</b> <code>pcie-aspm-on-bat</code>	[tlp-configuration parameter]
Same as <code>pcie-aspm-ac</code> but on BAT mode.	
Defaults to <code>"powersave"</code> .	
<b>maybe-non-negative-integer</b> <code>start-charge-thresh-bat0</code>	[tlp-configuration parameter]
Percentage when battery 0 should begin charging. Only supported on some laptops.	
Defaults to <code>'disabled'</code> .	
<b>maybe-non-negative-integer</b> <code>stop-charge-thresh-bat0</code>	[tlp-configuration parameter]
Percentage when battery 0 should stop charging. Only supported on some laptops.	
Defaults to <code>'disabled'</code> .	
<b>maybe-non-negative-integer</b> <code>start-charge-thresh-bat1</code>	[tlp-configuration parameter]
Percentage when battery 1 should begin charging. Only supported on some laptops.	
Defaults to <code>'disabled'</code> .	

**maybe-non-negative-integer** [tlp-configuration parameter]  
    **stop-charge-thresh-bat1**  
    Percentage when battery 1 should stop charging. Only supported on some laptops.  
    Defaults to 'disabled'.

**string radeon-power-profile-on-ac** [tlp-configuration parameter]  
    Radeon graphics clock speed level. Alternatives are low, mid, high, auto, default.  
    Defaults to "high".

**string radeon-power-profile-on-bat** [tlp-configuration parameter]  
    Same as **radeon-power-ac** but on BAT mode.  
    Defaults to "low".

**string radeon-dpm-state-on-ac** [tlp-configuration parameter]  
    Radeon dynamic power management method (DPM). Alternatives are battery, performance.  
    Defaults to "performance".

**string radeon-dpm-state-on-bat** [tlp-configuration parameter]  
    Same as **radeon-dpm-state-ac** but on BAT mode.  
    Defaults to "battery".

**string radeon-dpm-perf-level-on-ac** [tlp-configuration parameter]  
    Radeon DPM performance level. Alternatives are auto, low, high.  
    Defaults to "auto".

**string radeon-dpm-perf-level-on-bat** [tlp-configuration parameter]  
    Same as **radeon-dpm-perf-ac** but on BAT mode.  
    Defaults to "auto".

**on-off-boolean wifi-pwr-on-ac?** [tlp-configuration parameter]  
    Wifi power saving mode.  
    Defaults to '#f'.

**on-off-boolean wifi-pwr-on-bat?** [tlp-configuration parameter]  
    Same as **wifi-power-ac?** but on BAT mode.  
    Defaults to '#t'.

**y-n-boolean wol-disable?** [tlp-configuration parameter]  
    Disable wake on LAN.  
    Defaults to '#t'.

**non-negative-integer** [tlp-configuration parameter]  
    **sound-power-save-on-ac**  
    Timeout duration in seconds before activating audio power saving on Intel HDA and AC97 devices. A value of 0 disables power saving.  
    Defaults to '0'.

- non-negative-integer** [tlp-configuration parameter]  
    **sound-power-save-on-bat**  
    Same as **sound-powersave-ac** but on BAT mode.  
    Defaults to '1'.
- y-n-boolean sound-power-save-controller?** [tlp-configuration parameter]  
    Disable controller in powersaving mode on Intel HDA devices.  
    Defaults to '#t'.
- boolean bay-poweroff-on-bat?** [tlp-configuration parameter]  
    Enable optical drive in UltraBay/MediaBay on BAT mode. Drive can be powered on again by releasing (and reinserting) the eject lever or by pressing the disc eject button on newer models.  
    Defaults to '#f'.
- string bay-device** [tlp-configuration parameter]  
    Name of the optical drive device to power off.  
    Defaults to "sr0".
- string runtime-pm-on-ac** [tlp-configuration parameter]  
    Runtime Power Management for PCI(e) bus devices. Alternatives are on and auto.  
    Defaults to "on".
- string runtime-pm-on-bat** [tlp-configuration parameter]  
    Same as **runtime-pm-ac** but on BAT mode.  
    Defaults to "auto".
- boolean runtime-pm-all?** [tlp-configuration parameter]  
    Runtime Power Management for all PCI(e) bus devices, except blacklisted ones.  
    Defaults to '#t'.
- maybe-space-separated-string-list** [tlp-configuration parameter]  
    **runtime-pm-blacklist**  
    Exclude specified PCI(e) device addresses from Runtime Power Management.  
    Defaults to 'disabled'.
- space-separated-string-list** [tlp-configuration parameter]  
    **runtime-pm-driver-blacklist**  
    Exclude PCI(e) devices assigned to the specified drivers from Runtime Power Management.
- boolean usb-autosuspend?** [tlp-configuration parameter]  
    Enable USB autosuspend feature.  
    Defaults to '#t'.
- maybe-string usb-blacklist** [tlp-configuration parameter]  
    Exclude specified devices from USB autosuspend.  
    Defaults to 'disabled'.

**boolean** `usb-blacklist-wwan?` [tlp-configuration parameter]  
 Exclude WWAN devices from USB autosuspend.  
 Defaults to `'#t'`.

**maybe-string** `usb-whitelist` [tlp-configuration parameter]  
 Include specified devices into USB autosuspend, even if they are already excluded by the driver or via `usb-blacklist-wwan?`.  
 Defaults to `'disabled'`.

**maybe-boolean** `usb-autosuspend-disable-on-shutdown?` [tlp-configuration parameter]  
 Enable USB autosuspend before shutdown.  
 Defaults to `'disabled'`.

**boolean** `restore-device-state-on-startup?` [tlp-configuration parameter]  
 Restore radio device state (bluetooth, wifi, wwan) from previous shutdown on system startup.  
 Defaults to `'#f'`.

**string** `cpu-energy-perf-policy-on-ac` [tlp-configuration parameter]  
 Set CPU energy/performance policy when on AC mode. Possible values, in order of increasing power saving, are performance, balance\_performance, default, balance\_power and power.

**string** `cpu-energy-perf-policy-on-bat` [tlp-configuration parameter]  
 Set CPU energy/performance policy when on BAT mode. Possible values, in order of increasing power saving, are performance, balance\_performance, default, balance\_power and power.

## Thermal daemon

The (`gnu services pm`) module provides an interface to `thermald`, a CPU frequency scaling service which helps prevent overheating.

**thermald-service-type** [Variable]  
 This is the service type for `thermald` (<https://01.org/linux-thermal-daemon/>), the Linux Thermal Daemon, which is responsible for controlling the thermal state of processors and preventing overheating.

**thermald-configuration** [Data Type]  
 Data type representing the configuration of `thermald-service-type`.

**adaptive?** (default: `#f`)  
 Use DPTF (Dynamic Power and Thermal Framework) adaptive tables when present.

**ignore-cpuid-check?** (default: `#f`)  
 Ignore cpuid check for supported CPU models.

**thermald** (default: `thermald`)  
 Package object of `thermald`.

## PowerTOP

The (gnu services pm) module also provides a service definition for PowerTOP (<https://01.org/powertop/>), a power consumption analysis and tuning tool. When started, it tunes Linux kernel settings to reduce power consumption.

**powertop-service-type** [Variable]

The service type for PowerTOP. No configuration is necessary.

When the service starts, it executes **powertop --auto-tune**.

(service powertop-service-type)

Available **powertop-configuration** fields are:

**package powertop** [powertop-configuration parameter]

The PowerTOP package. Defaults to **powertop** in the (gnu packages linux) module

The (gnu services power) module provides a service definition for **apcupsd** (<http://www.apcupsd.org/>), a utility to interact with APC (APC by Schneider Electric or formerly American Power Conversion Corporation) UPS (Uninterruptible Power Supply) devices. **Apcupsd** also works with some OEM (Original Equipment Manufacturer)-branded products manufactured by APC.

**apcupsd-service-type** [Variable]

The service type for **apcupsd**. For USB UPSes no configuration is necessary, however tweaking some fields to better suit your needs might be desirable. The defaults are taken from the upstream configuration and they are not very conservative (for example, the default **battery-level** of 5% may be considered too low by some).

The default event handlers do send emails, read more in [apcupsd-event-handlers], page 573.

(service apcupsd-service-type)

**apcupsd-configuration** [Data Type]

Available **apcupsd-configuration** fields are:

**apcupsd** (default: **apcupsd**) (type: package)

The **apcupsd** package to use.

**shepherd-service-name** (default: **apcupsd**) (type: symbol)

The name of the shepherd service. You can add the service multiple times with different names to manage multiple UPSes.

**auto-start?** (default: **#t**) (type: boolean)

Should the shepherd service auto-start?

**pid-file** (default: **"/var/run/apcupsd.pid"**) (type: string)

The file name of the pid file.

**debug-level** (default: **0**) (type: integer)

The logging verbosity. Bigger number means more logs. The source code uses up to 300 as debug level value, so a value of 999 seems reasonable to enable all the logs.

**run-dir** (default: `"/var/run/apcupsd"`) (type: string)

The directory containing runtime information. You need to change this if you desire to run multiple instances of the daemon.

**name** (type: maybe-string)

Use this to give your UPS a name in log files and such. This is particularly useful if you have multiple UPSes. This does not set the EEPROM. It should be 8 characters or less.

**cable** (default: `usb`) (type: enum-cable)

The type of a cable connecting the UPS to your computer. Possible generic choices are `'simple`, `'smart`, `'ether` and `'usb`.

Alternatively, a specific cable model number may be used: `'940-0119A`, `'940-0127A`, `'940-0128A`, `'940-0020B`, `'940-0020C`, `'940-0023A`, `'940-0024B`, `'940-0024C`, `'940-1524C`, `'940-0024G`, `'940-0095A`, `'940-0095B`, `'940-0095C`, `'940-0625A`, `'M-04-02-2000`.

**type** (default: `usb`) (type: enum-type)

The type of the UPS you have.

**apcsmart** Newer serial character device, appropriate for SmartUPS models using a serial cable (not an USB).

**usb** Most new UPSes are an USB.

**net** Network link to a master apcupsd through apcupsd's Network Information Server. This is used if the UPS powering your computer is connected to a different computer for monitoring.

**snmp** SNMP network link to an SNMP-enabled UPS device.

**netsnmp** Same as the SNMP above but requires use of the net-snmp library. Unless you have a specific need for this old driver, you should use the `'snmp` instead.

**dumb** An old serial character device for use with simple-signaling UPSes.

**pcnet** A PowerChute Network Shutdown protocol which can be used as an alternative to an SNMP with the AP9617 family of smart slot cards.

**modbus** A serial device for use with newest SmartUPS models supporting the MODBUS protocol.

**device** (default: `""`) (type: string)

For USB UPSes, usually you want to set this to an empty string (the default). For other UPS types, you must specify an appropriate port or address.

**apcsmart** Set to the appropriate `/dev/tty**` device.

**usb** A null string setting enables auto-detection, which is the best choice for most installations.

<b>net</b>	Set to <i>hostname:port</i> .
<b>snmp</b>	Set to <i>hostname:port:vendor:community</i> . The <i>hostname</i> is the ip address or hostname of the UPS on the network. The <i>vendor</i> can be "APC" or "APC_NOTRAP". "APC_NOTRAP" will disable SNMP trap catching; you usually want "APC". The <i>port</i> is usually 161. The <i>community</i> is usually "private".
<b>netsnmp</b>	Same as the 'snmp'.
<b>dumb</b>	Set to the appropriate <i>/dev/tty**</i> device.
<b>pcnet</b>	Set to <i>ipaddr:username:passphrase:port</i> . The <i>ipaddr</i> is the IP address of the UPS management card. The <i>username</i> and the <i>passphrase</i> are the credentials for which the card has been configured. The <i>port</i> is the port number on which to listen for messages from the UPS, normally 3052. If this parameter is empty or missing, the default of 3052 will be used.
<b>modbus</b>	Set to the appropriate <i>/dev/tty**</i> device. You can also leave it empty for MODBUS over USB or set to the serial number of the UPS.

**poll-time** (default: 60) (type: integer)

The interval (in seconds) at which apcupsd polls the UPS for status. This setting applies both to directly-attached UPSes (apcsmart, usb, dumb) and networked UPSes (net, snmp). Lowering this setting will improve the apcupsd's responsiveness to certain events at the cost of higher CPU utilization.

**on-battery-delay** (default: 6) (type: integer)

The time in seconds from when a power failure is detected until we react to it with an onbattery event. The 'powerout event will be triggered immediately when a power failure is detected. However, the 'onbattery event will be trigger only after this delay.

**battery-level** (default: 5) (type: integer)

If during a power failure, the remaining battery percentage (as reported by the UPS) is below or equal to this value, the apcupsd will initiate a system shutdown.

**Note:** **battery-level**, **remaining-minutes**, and **timeout** work in a conjunction, so the first that occurs will cause the initiation of a shutdown.

**remaining-minutes** (default: 3) (type: integer)

If during a power failure, the remaining runtime in minutes (as calculated internally by the UPS) is below or equal to this value, apcupsd will initiate a system shutdown.

**Note:** `battery-level`, `remaining-minutes`, and `timeout` work in a conjunction, so the first that occurs will cause the initiation of a shutdown.

`timeout` (default: 0) (type: integer)

If during a power failure, the UPS has run on batteries for this many seconds or longer, `apcupsd` will initiate a system shutdown. The value of 0 disables this timer.

**Note:** `battery-level`, `remaining-minutes`, and `timeout` work in a conjunction, so the first that occurs will cause the initiation of a shutdown.

`annoy-interval` (default: 300) (type: integer)

The time in seconds between annoying users (via the '`annoyme`' event) to sign off prior to system shutdown. 0 disables.

`annoy-delay` (default: 60) (type: integer)

The initial delay in seconds after a power failure before warning users to get off the system.

`no-logon` (default: `disable`) (type: enum-no-logon)

The condition which determines when users are prevented from logging in during a power failure.

`kill-delay` (default: 0) (type: integer)

If this is non-zero, the `apcupsd` will continue running after a shutdown has been requested, and after the specified time in seconds attempt to kill the power. This is for use on systems where `apcupsd` cannot regain control after a shutdown.

`net-server` (default: `#f`) (type: boolean)

If enabled, a network information server process will be started.

`net-server-ip` (default: "127.0.0.1") (type: string)

An IP address on which the NIS server will listen for incoming connections.

`net-server-port` (default: 3551) (type: integer)

An IP port on which the NIS server will listen for incoming connections.

`net-server-events-file` (type: maybe-string)

If you want the last few EVENTS to be available over the network by the network information server, you must set this to a file name.

`net-server-events-file-max-size` (default: 10) (type: integer)

The maximum size of the events file in kilobytes.

`class` (default: `standalone`) (type: enum-class)

Normally `standalone` unless you share an UPS using an APC ShareUPS card.

`mode` (default: `disable`) (type: enum-mode)

Normally `disable` unless you share an UPS using an APC ShareUPS card.



**stat-time** (default: 0) (type: integer)  
 The time interval in seconds between writing the status file, 0 disables.

**log-stats** (default: #f) (type: boolean)  
 Also write the stats as a logs. This generates a lot of output.

**data-time** (default: 0) (type: integer)  
 The time interval in seconds between writing the data records to the log file, 0 disables.

**facility** (type: maybe-string)  
 The logging facility for the syslog.

**event-handlers** (type: apcupsd-event-handlers)  
 Handlers for events produced by apcupsd.

### **apcupsd-event-handlers** [Data Type]

For a description of the events please refer to ‘**man 8 apccontrol**’.

Each handler shall be a gexp. It is spliced into the control script for the daemon. In addition to the standard Guile programming environment, the following procedures and variables are also available:

**conf**            Variable containing the file name of the configuration file.

**powerfail-file**  
 Variable containing the file name of the powerfail file.

**cmd**            The event currently being handled.

**name**           The name of the UPS as specified in the configuration file.

**connected?**  
 Is #t if **apcupsd** is connected to the UPS via a serial port (or a USB port). In most configurations, this will be the case. In the case of a Slave machine where **apcupsd** is not directly connected to the UPS, this value will be #f.

**powered?**    Is #t if the computer on which **apcupsd** is running is powered by the UPS and #f if not. At the moment, this value is unimplemented and always #f.

**(err fmt args...)**  
 Wrapper around **format** outputting to (**current-error-port**).

**(wall fmt args...)**  
 Wrapper around **format** outputting via **wall**.

**(apcupsd args...)**  
 Call **apcupsd** while passing the correct configuration file and all the arguments.

**(mail-to-root subject body)**  
 Send an email to the local administrator. This procedure assumes the **sendmail** is located at **/run/privileged/bin/sendmail** (as would be the case with **opensmtpd-service-type**).

Available `apcupsd-event-handlers` fields are:

`modules` (type: gexp)  
Additional modules to import into the generated handler script.

`killpower` (type: gexp)  
The handler for the killpower event.

`commfailure` (type: gexp)  
The handler for the commfailure event.

`commok` (type: gexp)  
The handler for the commfailure event.

`powerout` (type: gexp)  
The handler for the powerout event.

`onbattery` (type: gexp)  
The handler for the onbattery event.

`offbattery` (type: gexp)  
The handler for the offbattery event.

`mainsback` (type: gexp)  
The handler for the mainsback event.

`failing` (type: gexp)  
The handler for the failing event.

`timeout` (type: gexp)  
The handler for the timeout event.

`loadlimit` (type: gexp)  
The handler for the loadlimit event.

`runlimit` (type: gexp)  
The handler for the runlimit event.

`doreboot` (type: gexp)  
The handler for the doreboot event.

`doshutdown` (type: gexp)  
The handler for the doshutdown event.

`annoyme` (type: gexp)  
The handler for the annoyme event.

`emergency` (type: gexp)  
The handler for the emergency event.

`changeme` (type: gexp)  
The handler for the changeme event.

`remotedown` (type: gexp)  
The handler for the remotedown event.

`startselftest` (type: gexp)  
The handler for the startselftest event.

```
endselftest (type: gexp)
    The handler for the endselftest event.

battdetach (type: gexp)
    The handler for the battdetach event.

battattach (type: gexp)
    The handler for the battattach event.
```

### 11.10.31 Audio Services

The (`gnu services audio`) module provides a service to start MPD (the Music Player Daemon).

#### Music Player Daemon

The Music Player Daemon (MPD) is a service that can play music while being controlled from the local machine or over the network by a variety of clients.

The following example shows the simplest configuration to locally expose, via PulseAudio, a music collection kept at `/srv/music`, with `mpd` running as the default ‘`mpd`’ user. This user will spawn its own PulseAudio daemon, which may compete for the sound card access with that of your own user. In this configuration, you may have to stop the playback of your user audio applications to hear MPD’s output and vice-versa.

```
(service mpd-service-type
  (mpd-configuration
    (music-directory "/srv/music"))))
```

**Important:** The music directory must be readable to the MPD user, by default, ‘`mpd`’. Permission problems will be reported via ‘`Permission denied`’ errors in the MPD logs, which appear in `/var/log/messages` by default.

Most MPD clients will trigger a database update upon connecting, but you can also use the `update` action do to so:

```
herd update mpd
```

All the MPD configuration fields are documented below, and a more complex example follows.

**mpd-service-type** [Variable]  
The service type for `mpd`

**mpd-configuration** [Data Type]  
Available `mpd-configuration` fields are:

**package** (default: `mpd`) (type: file-like)  
The MPD package.

**user** (type: user-account)  
The user to run `mpd` as.

**group** (type: user-group)  
The group to run `mpd` as.  
The default `%mpd-group` is a system group with name “`mpd`”.

**shepherd-requirement** (default: '()') (type: list-of-symbols)  
A list of symbols naming Shepherd services that this service will depend on.

**environment-variables** (default: '("PULSE\_CLIENTCONFIG=/etc/pulse/client.conf" "PULSE\_CONFIG=/etc/pulse/daemon.conf")') (type: list-of-strings)  
A list of strings specifying environment variables.

**log-file** (type: maybe-string)  
The location of the log file. Unless specified, logs are sent to the local syslog daemon. Alternatively, a log file name can be specified, for example `/var/log/mpd.log`.

**log-level** (type: maybe-string)  
Suppress any messages below this threshold. The available values, in decreasing order of verbosity, are: **verbose**, **info**, **notice**, **warning** and **error**.

**music-directory** (type: maybe-string)  
The directory to scan for music files.

**music-dir** (type: maybe-string)  
The directory to scan for music files.

**playlist-directory** (type: maybe-string)  
The directory to store playlists.

**playlist-dir** (type: maybe-string)  
The directory to store playlists.

**db-file** (type: maybe-string)  
The location of the music database. When left unspecified, `~/.cache/db` is used.

**state-file** (type: maybe-string)  
The location of the file that stores current MPD's state.

**sticker-file** (type: maybe-string)  
The location of the sticker database.

**default-port** (default: 6600) (type: maybe-port)  
The default port to run mpd on.

**endpoints** (type: maybe-list-of-strings)  
The addresses that mpd will bind to. A port different from *default-port* may be specified, e.g. `localhost:6602` and IPv6 addresses must be enclosed in square brackets when a different port is used. To use a Unix domain socket, an absolute path or a path starting with `~` can be specified here.

**address** (type: maybe-string)  
The address that mpd will bind to. To use a Unix domain socket, an absolute path can be specified here.

**database** (type: maybe-mpd-plugin)  
MPD database plugin configuration.

**partitions** (default: '()') (type: list-of-mpd-partition)  
List of MPD "partitions".

**neighbors** (default: '()') (type: list-of-mpd-plugin)  
List of MPD neighbor plugin configurations.

**inputs** (default: '()') (type: list-of-mpd-plugin)  
List of MPD input plugin configurations.

**archive-plugins** (default: '()') (type: list-of-mpd-plugin)  
List of MPD archive plugin configurations.

**auto-update?** (type: maybe-boolean)  
Whether to automatically update the music database when files are changed in the *music-directory*.

**input-cache-size** (type: maybe-string)  
MPD input cache size.

**decoders** (default: '()') (type: list-of-mpd-plugin)  
List of MPD decoder plugin configurations.

**resampler** (type: maybe-mpd-plugin)  
MPD resampler plugin configuration.

**filters** (default: '()') (type: list-of-mpd-plugin)  
List of MPD filter plugin configurations.

**outputs** (type: list-of-mpd-plugin-or-output)  
The audio outputs that MPD can use. By default this is a single output using pulseaudio.

**playlist-plugins** (default: '()') (type: list-of-mpd-plugin)  
List of MPD playlist plugin configurations.

**extra-options** (default: '()') (type: alist)  
An association list of option symbols/strings to string values to be appended to the configuration.

**mpd-plugin**

[Data Type]

Data type representing a mpd plugin.

**plugin** (type: maybe-string)  
Plugin name.

**name** (type: maybe-string)  
Name.

**enabled?** (type: maybe-boolean)  
Whether the plugin is enabled/disabled.

**extra-options** (default: '()') (type: alist)  
An association list of option symbols/strings to string values to be appended to the plugin configuration. See MPD plugin reference (<https://mpd.readthedocs.io/en/latest/plugins.html>) for available options.

**mpd-partition** [Data Type]

Data type representing a **mpd** partition.

**name** (type: string)

Partition name.

**extra-options** (default: '()') (type: alist)

An association list of option symbols/strings to string values to be appended to the partition configuration. See [Configuring Partitions \(https://mpd.readthedocs.io/en/latest/user.html#configuring-partitions\)](https://mpd.readthedocs.io/en/latest/user.html#configuring-partitions) for available options.

**mpd-output** [Data Type]

Available **mpd-output** fields are:

**name** (default: "MPD") (type: string)

The name of the audio output.

**type** (default: "pulse") (type: string)

The type of audio output.

**enabled?** (default: #t) (type: boolean)

Specifies whether this audio output is enabled when MPD is started. By default, all audio outputs are enabled. This is just the default setting when there is no state file; with a state file, the previous state is restored.

**format** (type: maybe-string)

Force a specific audio format on output. See [Global Audio Format \(https://mpd.readthedocs.io/en/latest/user.html#audio-output-format\)](https://mpd.readthedocs.io/en/latest/user.html#audio-output-format) for a more detailed description.

**tags?** (default: #t) (type: boolean)

If set to #f, then MPD will not send tags to this output. This is only useful for output plugins that can receive tags, for example the **httpd** output plugin.

**always-on?** (default: #f) (type: boolean)

If set to #t, then MPD attempts to keep this audio output always open. This may be useful for streaming servers, when you don't want to disconnect all listeners even when playback is accidentally stopped.

**mixer-type** (type: maybe-string)

This field accepts a string that specifies which mixer should be used for this audio output: the **hardware** mixer, the **software** mixer, the **null** mixer (allows setting the volume, but with no effect; this can be used as a trick to implement an external mixer External Mixer) or no mixer (**none**). When left unspecified, a **hardware** mixer is used for devices that support it.

**replay-gain-handler** (type: maybe-string)

This field accepts a string that specifies how Replay Gain (<https://mpd.readthedocs.io/en/latest/user.html#replay-gain>) is to be applied.

**software** uses an internal software volume control, **mixer** uses the configured (hardware) mixer control and **none** disables replay gain on this audio output.

**extra-options** (default: '()') (type: alist)

An association list of option symbols/strings to string values to be appended to the audio output configuration.

The following example shows a configuration of **mpd** that configures some of its plugins and provides a HTTP audio streaming output.

```
(service mpd-service-type
  (mpd-configuration
    (outputs
      (list (mpd-output
              (name "streaming")
              (type "httpd")
              (mixer-type 'null)
              (extra-options
                `((encoder . "vorbis")
                  (port . "8080"))))))
      (decoders
        (list (mpd-plugin
                (plugin "mikmod")
                (enabled? #f))
              (mpd-plugin
                (plugin "openmpt")
                (enabled? #t)
                (extra-options `((repeat-count . -1)
                                (interpolation-filter . 1)))))
              (resampler (mpd-plugin
                          (plugin "libsamplerate")
                          (extra-options `((type . 0)))))
```

## myMPD

myMPD (<https://jcorporation.github.io/myMPD/>) is a web server frontend for MPD that provides a mobile friendly web client for MPD.

The following example shows a myMPD instance listening on port 80, with album cover caching disabled.

```
(service mympd-service-type
  (mympd-configuration
    (port 80)
    (covercache-ttl 0)))
```

**mympd-service-type**

[Variable]

The service type for mympd.

**mympd-configuration**

[Data Type]

Available mympd-configuration fields are:

**package** (default: `mympd`) (type: file-like)  
The package object of the myMPD server.

**shepherd-requirement** (default: `'()`) (type: list-of-symbols)  
This is a list of symbols naming Shepherd services that this service will depend on.

**user** (default: `%mympd-user`) (type: user-account)  
Owner of the mympd process.  
The default `%mympd-user` is a system user with the name “mympd”, who is a part of the group *group* (see below).

**group** (default: `%mympd-group`) (type: user-group)  
Owner group of the mympd process.  
The default `%mympd-group` is a system group with name “mympd”.

**work-directory** (default: `"/var/lib/mympd"`) (type: string)  
Where myMPD will store its data.

**cache-directory** (default: `"/var/cache/mympd"`) (type: string)  
Where myMPD will store its cache.

**acl** (type: maybe-mympd-ip-acl)  
ACL to access the myMPD webserver.

**covercache-ttl** (default: `31`) (type: maybe-integer)  
How long to keep cached covers, 0 disables cover caching.

**http?** (default: `#t`) (type: boolean)  
HTTP support.

**host** (default: `"[::]"`) (type: string)  
Host name to listen on.

**port** (default: `80`) (type: maybe-port)  
HTTP port to listen on.

**log-level** (default: `5`) (type: integer)  
How much detail to include in logs, possible values: 0 to 7.

**log-to** (type: maybe-string)  
Where to send logs. Unless specified, the service logs to the local syslog service under the ‘daemon’ facility. Alternatively, a log file name can be specified, for example `/var/log/mympd.log`.

**lualibs** (default: `"all"`) (type: maybe-string)  
See <https://jcorporation.github.io/myMPD/scripting/#lua-standard-libraries>.

**uri** (type: maybe-string)  
Override URI to myMPD. See <https://github.com/jcorporation/myMPD/issues/950>.

**script-acl** (default: `(mympd-ip-acl (allow '("127.0.0.1")))`) (type: maybe-mympd-ip-acl)  
ACL to access the myMPD script backend.



**ssl?** (default: **#f**) (type: boolean)  
SSL/TLS support.

**ssl-port** (default: **443**) (type: maybe-port)  
Port to listen for HTTPS.

**ssl-cert** (type: maybe-string)  
Path to PEM encoded X.509 SSL/TLS certificate (public key).

**ssl-key** (type: maybe-string)  
Path to PEM encoded SSL/TLS private key.

**pin-hash** (type: maybe-string)  
SHA-256 hashed pin used by myMPD to control settings access by prompting a pin from the user.

**save-caches?** (type: maybe-boolean)  
Whether to preserve caches between service restarts.

**mympd-ip-acl** [Data Type]

Available **mympd-ip-acl** fields are:

**allow** (default: **'()**) (type: list-of-strings)  
Allowed IP addresses.

**deny** (default: **'()**) (type: list-of-strings)  
Disallowed IP addresses.

### 11.10.32 Virtualization Services

The (**gnu services virtualization**) module provides services for the libvirt and virtlog daemons, as well as other virtualization-related services.

#### Libvirt daemon

**libvirtd** is the server side daemon component of the libvirt virtualization management system. This daemon runs on host servers and performs required management tasks for virtualized guests. To connect to the libvirt daemon as an unprivileged user, it must be added to the 'libvirt' group, as shown in the example below.

**libvirt-service-type** [Variable]

This is the type of the libvirt daemon (<https://libvirt.org>). Its value must be a **libvirt-configuration**.

```
(users (cons (user-account
              (name "user")
              (group "users")
              (supplementary-groups '("libvirt"
                                     "audio" "video" "wheel"))))
        %base-user-accounts))
(service libvirt-service-type
  (libvirt-configuration
    (tls-port "16555")))
```

**libvirt-configuration** [Data Type]

Available **libvirt-configuration** fields are:

**libvirt** (default: **libvirt**) (type: file-like)

Libvirt package.

**qemu** (default: **qemu**) (type: file-like)

The QEMU package to use.

**dmidecode** (default: **dmidecode**) (type: file-like)

The Dmidecode package to use.

**requirement** (default: '()') (type: list-of-symbols)

Shepherd services dependencies to the provisioned Shepherd service.

**firmwares** (default: (**ovmf-x86-64**)) (type: list-of-file-likes)

List of UEFI/BIOS firmware packages to make available. Each firmware package should contain a **share/qemu/firmware/NAME.json** QEMU firmware metadata file.

**listen-tls?** (default: **#t**) (type: boolean)

Flag listening for secure TLS connections on the public TCP/IP port. must set **listen** for this to have any effect. It is necessary to setup a CA and issue server certificates before using this capability.

**listen-tcp?** (default: **#f**) (type: boolean)

Listen for unencrypted TCP connections on the public TCP/IP port. must set **listen** for this to have any effect. Using the TCP socket requires SASL authentication by default. Only SASL mechanisms which support data encryption are allowed. This is **DIGEST\_MD5** and **GSSAPI** (Kerberos5).

**tls-port** (default: "16514") (type: string)

Port for accepting secure TLS connections This can be a port number, or service name.

**tcp-port** (default: "16509") (type: string)

Port for accepting insecure TCP connections This can be a port number, or service name.

**listen-addr** (default: "0.0.0.0") (type: string)

IP address or hostname used for client connections.

**mdns-adv?** (default: **#f**) (type: boolean)

Flag toggling mDNS advertisement of the libvirt service.

**mdns-name** (default: "Virtualization Host terra") (type: string)

Default mDNS advertisement name. This must be unique on the immediate broadcast network.

**unix-sock-group** (default: "libvirt") (type: string)

UNIX domain socket group ownership. This can be used to allow a trusted set of users access to management capabilities without becoming root.

- unix-sock-ro-perms** (default: "0777") (type: string)  
UNIX socket permissions for the R/O socket. This is used for monitoring VM status only.
- unix-sock-rw-perms** (default: "0770") (type: string)  
UNIX socket permissions for the R/W socket. Default allows only root. If PolicyKit is enabled on the socket, the default will change to allow everyone (e.g., "0777").
- unix-sock-admin-perms** (default: "0777") (type: string)  
UNIX socket permissions for the admin socket. Default allows only owner (root), do not change it unless you are sure to whom you are exposing the access to.
- unix-sock-dir** (default: "/var/run/libvirt") (type: string)  
The directory in which sockets will be found/created.
- auth-unix-ro** (default: "polkit") (type: string)  
Authentication scheme for UNIX read-only sockets. By default socket permissions allow anyone to connect.
- auth-unix-rw** (default: "polkit") (type: string)  
Authentication scheme for UNIX read-write sockets. By default socket permissions only allow root. If PolicyKit support was compiled into libvirt, the default will be to use 'polkit' auth.
- auth-tcp** (default: "sasl") (type: string)  
Authentication scheme for TCP sockets. If you don't enable SASL, then all TCP traffic is cleartext. Don't do this outside of a development/test scenario.
- auth-tls** (default: "none") (type: string)  
Authentication scheme for TLS sockets. TLS sockets already have encryption provided by the TLS layer, and limited authentication is done by certificates. It is possible to make use of any SASL authentication mechanism as well, by using "sasl" for this option
- access-drivers** (default: ()) (type: optional-list)  
API access control scheme. By default an authenticated user is allowed access to all APIs. Access drivers can place restrictions on this.
- key-file** (default: "") (type: string)  
Server key file path. If set to an empty string, then no private key is loaded.
- cert-file** (default: "") (type: string)  
Server key file path. If set to an empty string, then no certificate is loaded.
- ca-file** (default: "") (type: string)  
Server key file path. If set to an empty string, then no CA certificate is loaded.

- crl-file** (default: "") (type: string)  
Certificate revocation list path. If set to an empty string, then no CRL is loaded.
- tls-no-sanity-cert** (default: #f) (type: boolean)  
Disable verification of our own server certificates. When `libvirtd` starts it performs some sanity checks against its own certificates.
- tls-no-verify-cert** (default: #f) (type: boolean)  
Disable verification of client certificates. Client certificate verification is the primary authentication mechanism. Any client which does not present a certificate signed by the CA will be rejected.
- tls-allowed-dn-list** (default: ()) (type: optional-list)  
Whitelist of allowed x509 Distinguished Name.
- sasl-allowed-usernames** (default: ()) (type: optional-list)  
Whitelist of allowed SASL usernames. The format for username depends on the SASL authentication mechanism.
- tls-priority** (default: "NORMAL") (type: string)  
Override the compile time default TLS priority string. The default is usually "NORMAL" unless overridden at build time. Only set this if it is desired for libvirt to deviate from the global default settings.
- max-clients** (default: 5000) (type: integer)  
Maximum number of concurrent client connections to allow over all sockets combined.
- max-queued-clients** (default: 1000) (type: integer)  
Maximum length of queue of connections waiting to be accepted by the daemon. Note, that some protocols supporting re-transmission may obey this so that a later reattempt at connection succeeds.
- max-anonymous-clients** (default: 20) (type: integer)  
Maximum length of queue of accepted but not yet authenticated clients. Set this to zero to turn this feature off.
- min-workers** (default: 5) (type: integer)  
Number of workers to start up initially.
- max-workers** (default: 20) (type: integer)  
Maximum number of worker threads. If the number of active clients exceeds **min-workers**, then more threads are spawned, up to **max\_workers** limit. Typically you'd want **max\_workers** to equal maximum number of clients allowed.
- prio-workers** (default: 5) (type: integer)  
Number of priority workers. If all workers from above pool are stuck, some calls marked as high priority (notably `domainDestroy`) can be executed in this pool.
- max-requests** (default: 20) (type: integer)  
Total global limit on concurrent RPC calls.

- max-client-requests** (default: 5) (type: integer)  
Limit on concurrent requests from a single client connection. To avoid one client monopolizing the server this should be a small fraction of the global `max-requests` and `max-workers` parameter.
- admin-min-workers** (default: 1) (type: integer)  
Same as `min-workers` but for the admin interface.
- admin-max-workers** (default: 5) (type: integer)  
Same as `max-workers` but for the admin interface.
- admin-max-clients** (default: 5) (type: integer)  
Same as `max-clients` but for the admin interface.
- admin-max-queued-clients** (default: 5) (type: integer)  
Same as `max-queued-clients` but for the admin interface.
- admin-max-client-requests** (default: 5) (type: integer)  
Same as `max-client-requests` but for the admin interface.
- log-level** (default: 3) (type: integer)  
Semi-deprecated option for the logging level: using the `log-filters` option instead is recommend, as it provides finer control. The log level can be set to 4 for errors, 3 for warnings, 2 for information or 1 for debug. Note that since `log-filters` and `log-outputs` take precedence over this option, you will need to also adjust their logging levels to avoid filtering out messages.
- log-filters** (default: "3:remote 4:event") (type: string)  
Logging filters. A filter allows selecting a different logging level for a given category of logs. The format for a filter is either '`x:name`' or '`x:+name`', where `name` is a string which is matched against the category given in the `VIR_LOG_INIT()` at the top of each libvirt source file, e.g., "`remote`", "`qemu`", or "`util.json`". `name` can be a substring of the full category name, in order to match multiple similar categories. The optional '+' prefix tells libvirt to log stack traces for each message matching `name`. `x` is the log level value used to filter the associated message category. `x` can be set to '4' for errors, '3' for warnings, '2' for information, or '1' for debug. Multiple filters can be defined in a single filters statement, as space-separated values. Since `log-outputs` also include a level filter, you may need to also adjust its value to see all the filtered messages.
- log-outputs** (default: "3:syslog:libvirtd") (type: string)  
Logging outputs. An output is one of the places to save logging information. The format for an output has the form "`x:var{output}`", where `output` can be '`stderr`', '`syslog:name`', where `name` is the syslog `ident` value to use, or '`file:file_name`', where `file_name` is the file name of the file to output to. `x` is the minimal level, which acts as a filter. `x` can be set to `x` can be set to '4' for errors, '3' for warnings, '2' for information, or '1' for debug. Multiple filters can be defined in a single filters statement, as space-separated values.

- audit-level** (default: 1) (type: integer)  
 Modify the behavior of the auditing subsystem. ‘0’ disables all auditing, `samp{1}` enables auditing only if enabled on the host and ‘2’ enables auditing but exits if it is disabled on the host.
- audit-logging** (default: #f) (type: boolean)  
 Send audit messages via libvirt logging infrastructure.
- host-uuid** (default: "") (type: optional-string)  
 Host UUID. UUID must not have all digits be the same.
- host-uuid-source** (default: "smbios") (type: string)  
 Source to read host UUID. Use "smbios" to fetch the UUID via `dmidecode -s system-uuid`, or "machine-id" to fetch the UUID from `/etc/machine-id`. If `dmidecode` does not provide a valid UUID a temporary UUID will be generated.
- keepalive-interval** (default: 5) (type: integer)  
 A keepalive message is sent to a client after `keepalive_interval` seconds of inactivity to check if the client is still responding. If set to -1, `libvirtd` won’t send keepalive requests; however clients can still send them and the daemon will send responses.
- keepalive-count** (default: 5) (type: integer)  
 Maximum number of keepalive messages that are allowed to be sent to the client without getting any response before the connection is considered broken. In other words, the connection is automatically closed approximately after ‘`keepalive_interval * (keepalive_count + 1)`’ seconds since the last message received from the client. When `keepalive-count` is set to 0, connections will be automatically closed after `keepalive-interval` seconds of inactivity without sending any keepalive messages.
- admin-keepalive-interval** (default: 5) (type: integer)  
 Same as above but for admin interface.
- admin-keepalive-count** (default: 5) (type: integer)  
 Same as above but for admin interface.
- ovs-timeout** (default: 5) (type: integer)  
 Timeout for Open vSwitch calls. The `ovs-vsctl` utility is used for the configuration and its timeout option is set by default to 5 seconds to avoid potential infinite waits blocking libvirt.

## Virtlog daemon

The `virtlogd` service is a server side daemon component of libvirt that is used to manage logs from virtual machine consoles.

This daemon is not used directly by libvirt client applications, rather it is called on their behalf by `libvirtd`. By maintaining the logs in a standalone daemon, the main `libvirtd` daemon can be restarted without risk of losing logs. The `virtlogd` daemon has the ability to `re-exec()` itself upon receiving `SIGUSR1`, to allow live upgrades without downtime.

**virtlog-service-type** [Variable]

This is the type of the virtlog daemon. Its value must be a **virtlog-configuration**.

```
(service virtlog-service-type
  (virtlog-configuration
    (max-clients 1000)))
```

**package libvirt** [libvirt parameter]

Libvirt package.

**integer log-level** [virtlog-configuration parameter]

Logging level. 4 errors, 3 warnings, 2 information, 1 debug.

Defaults to '3'.

**string log-filters** [virtlog-configuration parameter]

Logging filters.

A filter allows to select a different logging level for a given category of logs The format for a filter is one of:

- x:name
- x:+name

where **name** is a string which is matched against the category given in the **VIR\_LOG\_INIT()** at the top of each libvirt source file, e.g., "remote", "qemu", or "util.json" (the name in the filter can be a substring of the full category name, in order to match multiple similar categories), the optional "+" prefix tells libvirt to log stack trace for each message matching name, and **x** is the minimal level where matching messages should be logged:

- 1: DEBUG
- 2: INFO
- 3: WARNING
- 4: ERROR

Multiple filters can be defined in a single filters statement, they just need to be separated by spaces.

Defaults to "3:remote 4:event".

**string log-outputs** [virtlog-configuration parameter]

Logging outputs.

An output is one of the places to save logging information The format for an output can be:

**x:stderr** output goes to stderr

**x:syslog:name**

use syslog for the output and use the given name as the ident

**x:file:file\_path**

output to a file, with the given filepath

**x:journald**

output to journald logging system

In all case the x prefix is the minimal level, acting as a filter

- 1: DEBUG
- 2: INFO
- 3: WARNING
- 4: ERROR

Multiple outputs can be defined, they just need to be separated by spaces.

Defaults to `"3:stderr"`.

**integer max-clients** [virtlog-configuration parameter]  
Maximum number of concurrent client connections to allow over all sockets combined.  
Defaults to `'1024'`.

**integer max-size** [virtlog-configuration parameter]  
Maximum file size before rolling over.  
Defaults to `'2MB'`

**integer max-backups** [virtlog-configuration parameter]  
Maximum number of backup files to keep.  
Defaults to `'3'`

## Transparent Emulation with QEMU

**qemu-binfmt-service-type** provides support for transparent emulation of program binaries built for different architectures—e.g., it allows you to transparently execute an ARMv7 program on an x86\_64 machine. It achieves this by combining the QEMU (<https://www.qemu.org>) emulator and the `binfmt_misc` feature of the kernel Linux. This feature only allows you to emulate GNU/Linux on a different architecture, but see below for GNU/Hurd support.

**qemu-binfmt-service-type** [Variable]  
This is the type of the QEMU/binfmt service for transparent emulation. Its value must be a **qemu-binfmt-configuration** object, which specifies the QEMU package to use as well as the architecture we want to emulate:

```
(service qemu-binfmt-service-type
  (qemu-binfmt-configuration
    (platforms (lookup-qemu-platforms "arm" "aarch64"))))
```

In this example, we enable transparent emulation for the ARM and aarch64 platforms. Running `herd stop qemu-binfmt` turns it off, and running `herd start qemu-binfmt` turns it back on (see Section “Invoking herd” in *The GNU Shepherd Manual*).

**qemu-binfmt-configuration** [Data Type]  
This is the configuration for the **qemu-binfmt** service.

**platforms** (default: `'()`)  
The list of emulated QEMU platforms. Each item must be a *platform object* as returned by `lookup-qemu-platforms` (see below).



For example, let's suppose you're on an x86\_64 machine and you have this service:

```
(service qemu-binfmt-service-type
  (qemu-binfmt-configuration
    (platforms (lookup-qemu-platforms "arm"))))
```

You can run:

```
guix build -s armhf-linux inkscape
```

and it will build Inkscape for ARMv7 *as if it were a native build*, transparently using QEMU to emulate the ARMv7 CPU. Pretty handy if you'd like to test a package build for an architecture you don't have access to!

`qemu` (default: `qemu`)

The QEMU package to use.

`lookup-qemu-platforms platforms...` [Procedure]

Return the list of QEMU platform objects corresponding to *platforms...* *platforms* must be a list of strings corresponding to platform names, such as "arm", "sparc", "mips64el", and so on.

`qemu-platform? obj` [Procedure]

Return true if *obj* is a platform object.

`qemu-platform-name platform` [Procedure]

Return the name of *platform*—a string such as "arm".

## QEMU Guest Agent

The QEMU guest agent provides control over the emulated system to the host. The `qemu-guest-agent` service runs the agent on Guix guests. To control the agent from the host, open a socket by invoking QEMU with the following arguments:

```
qemu-system-x86_64 \
-chardev socket,path=/tmp/qga.sock,server=on,wait=off,id=qga0 \
-device virtio-serial \
-device virtserialport,chardev=qga0,name=org.qemu.guest_agent.0 \
...
```

This creates a socket at `/tmp/qga.sock` on the host. Once the guest agent is running, you can issue commands with `socat`:

```
$ guix shell socat -- socat unix-connect:/tmp/qga.sock stdio
{"execute": "guest-get-host-name"}
{"return": {"host-name": "guix"}}
```

See QEMU guest agent documentation (<https://wiki.qemu.org/Features/GuestAgent>) for more options and commands.

`qemu-guest-agent-service-type` [Variable]

Service type for the QEMU guest agent service.

**qemu-guest-agent-configuration** [Data Type]

Configuration for the **qemu-guest-agent** service.

**qemu** (default: **qemu-minimal**)

The QEMU package to use.

**device** (default: **"**)

File name of the device or socket the agent uses to communicate with the host. If empty, QEMU uses a default file name.

## Xen Guest Agent

The Xen guest agent allows a Xen host to control the emulated system. The **xe-guest-utilities** service runs the agent on Guix guests.

**xe-guest-utilities** (<https://github.com/xenserver/xe-guest-utilities>) collects information about the running virtualized guest. This includes:

- static information
  - The operating system running
  - The Linux kernel version
- dynamic information
  - Network interfaces (devices) being added/removed
  - Network connections being unplugged/plugged-in
  - CPUs being added or removed
  - The guest migrating, being paused/resumed, etc.
- ephemeral information
  - The amount of memory currently in-use and free
  - The amount of disk-space used

**Note:** The default Linux-libre kernel that Guix ships already enables the necessary paravirtualization features for a guest. There is nothing you need to do for the guest to support Xen's paravirtualization features.

The guest utilities are used to have the guest report information *back* to the virtualizing host and support tasks that require cooperation between host and guest, like CPU hotplugging.

**xe-guest-utilities-service-type** [Variable]

Service type for the Xen guest utilities service.

**xe-guest-utilities-configuration** [Data Type]

Configuration for the **xe-guest-utilities** service.

**package** (default: **xe-guest-utilities**)

The Xen guest utilities package to use.

**pid-file** (default: **`/var/run/xe-daemon.pid`**)

Name of the file holding the PID of **xe-daemon**.

**log-file** (default: **`/var/log/xe-guest-utilities.log`**)

Name of the **xe-guest-utilities** log file.

`xe-guest-utilities` is the standard guest utilities used for Xen guests. More recently the Xen Project has been working to develop `xen-guest-agent`, a modern drop-in replacement for `xe-guest-utilities`. While they nearly have feature-parity, `xen-guest-agent` currently lacks some of the features of its predecessor, namely disk metrics and “PV drivers version”.

`xen-guest-agent-service-type` [Variable]  
Service type for the Xen guest agent service.

`xen-guest-agent-configuration` [Data Type]  
Configuration for the `xen-guest-agent` service.  
`package` (default: `xen-guest-agent`)  
The Xen guest agent package to use.

**Warning:** `xe-guest-utilities` and `xen-guest-agent` are mutually exclusive.

## Virtual Build Machines

*Virtual build machines* or “build VMs” let you offload builds to a fully controlled environment. “How can it be more controlled than regular builds? And why would it be useful?”, you ask. Good questions.

Builds spawned by `guix-daemon` indeed run in a controlled environment; specifically the daemon spawns build processes in separate namespaces and in a chroot, such as that build processes only see their declared dependencies and a well-defined subset of the file system tree (see Section 2.2.1 [Build Environment Setup], page 6, for details). A few aspects of the environments are not controlled though: the operating system kernel, the CPU model, and the date. Most of the time, these aspects have no impact on the build process: the level of isolation `guix-daemon` provides is “good enough”.

However, there are occasionally cases where those aspects *do* influence the build process. A typical example is *time traps*: build processes that stop working after a certain date<sup>8</sup>. Another one is software that optimizes for the CPU microarchitecture it is built on or, worse, bugs that manifest only on specific CPUs.

To address that, `virtual-build-machine-service-type` lets you add a virtual build machine on your system, as in this example:

```
(use-modules (gnu services virtualization))

(operating-system
  ;; ...
  (services (append (list (service virtual-build-machine-service-type))
                    %base-services)))
```

By default, you have to explicitly start the build machine when you need it, at which point builds may be offloaded to it (see Section 2.2.2 [Daemon Offload Setup], page 9):

```
herd start build-vm
```

---

<sup>8</sup> The most widespread example of time traps is test suites that involve checking the expiration date of a certificate. Such tests exist in TLS implementations such as OpenSSL and GnuTLS, but also in high-level software such as Python.

With the default setting shown above, the build VM runs with its clock set to a date several years in the past, and on a CPU model that corresponds to that date—a model possibly older than that of your machine. This lets you rebuild today software from the past that would otherwise fail to build due to a time trap or other issues in its build process. You can view the VM’s config like this:

```
herd configuration build-vm
```

You can configure the build VM, as in this example:

```
(service virtual-build-machine-service-type
  (virtual-build-machine
    (cpu "Westmere")
    (cpu-count 8)
    (memory-size (* 1 1024))
    (auto-start? #t)))
```

The available options are shown below.

**virtual-build-machine-service-type** [Variable]

This is the service type to run *virtual build machines*. Virtual build machines are configured so that builds are offloaded to them when they are running.

**virtual-build-machine** [Data Type]

This is the data type specifying the configuration of a build machine. It contains the fields below:

- name** (default: 'build-vm')  
The name of this build VM. It is used to construct the name of its Shepherd service.
- image**  
The image of the virtual machine (see Chapter 16 [System Images], page 793). This notably specifies the virtual disk size and the operating system running into it (see Section 11.3 [operating-system Reference], page 253). The default value is a minimal operating system image.
- qemu** (default: `qemu-minimal`)  
The QEMU package to run the image.
- cpu**  
The CPU model being emulated as a string denoting a model known to QEMU.  
The default value is a model that matches **date** (see below). To see what CPU models are available, run, for example:  
`qemu-system-x86_64 -cpu help`
- cpu-count** (default: 4)  
The number of CPUs emulated by the virtual machine.
- memory-size** (default: 2048)  
Size in mebibytes (MiB) of the virtual machine’s main memory (RAM).
- date** (default: a few years ago)  
Date inside the virtual machine when it starts; this must be a SRFI-19 date object (see Section “SRFI-19 Date” in *GNU Guile Reference Manual*).

```
port-forwardings (default: 11022 and 11004)
    TCP ports of the virtual machine forwarded to the host. By default, the
    SSH and secrets ports are forwarded into the host.

systems (default: (list (%current-system)))
    List of system types supported by the build VM—e.g., "x86_64-linux".

auto-start? (default: #f)
    Whether to start the virtual machine when the system boots.
```

In the next section, you'll find a variant on this theme: GNU/Hurd virtual machines!

## The Hurd in a Virtual Machine

Service `hurd-vm` provides support for running GNU/Hurd in a virtual machine (VM), a so-called *childhurd*. This service is meant to be used on GNU/Linux and the given GNU/Hurd operating system configuration is cross-compiled. The virtual machine is a Shepherd service that can be referred to by the names `hurd-vm` and `childhurd` and be controlled with commands such as:

```
herd start hurd-vm
herd stop childhurd
```

When the service is running, you can view its console by connecting to it with a VNC client, for example with:

```
guix shell tigervnc-client -- vncviewer localhost:5900
```

The default configuration (see `hurd-vm-configuration` below) spawns a secure shell (SSH) server in your GNU/Hurd system, which QEMU (the virtual machine emulator) redirects to port 10022 on the host. By default, the service enables *offloading* such that the host `guix-daemon` automatically offloads GNU/Hurd builds to the `childhurd` (see Section 2.2.2 [Daemon Offload Setup], page 9). This is what happens when running a command like the following one, where `i586-gnu` is the system type of 32-bit GNU/Hurd:

```
guix build emacs-minimal -s i586-gnu
```

The `childhurd` is volatile and stateless: it starts with a fresh root file system every time you restart it. By default though, all the files under `/etc/childhurd` on the host are copied as is to the root file system of the `childhurd` when it boots. This allows you to initialize “secrets” inside the VM: SSH host keys, authorized substitute keys, and so on—see the explanation of `secret-root` below.

You will probably find it useful to create an account for you in the GNU/Hurd virtual machine and to authorize logins with your SSH key. To do that, you can define the GNU/Hurd system in the usual way (see Section 11.2 [Using the Configuration System], page 244), and then pass that operating system as the `os` field of `hurd-vm-configuration`, as in this example:

```
(define childhurd-os
  ;; Definition of my GNU/Hurd system, derived from the default one.
  (operating-system
    (inherit %hurd-vm-operating-system)

    ;; Add a user account.
```

```

(users (cons (user-account
              (name "charlie")
              (comment "This is me!")
              (group "users")
              (supplementary-groups '("wheel"))) ;for 'sudo'
            %base-user-accounts))

(services
;; Modify the SSH configuration to allow login as "root"
;; and as "charlie" using public key authentication.
(modify-services (operating-system-user-services
                  %hurd-vm-operating-system)
  (openssh-service-type
   config => (openssh-configuration
              (inherit config)
              (authorized-keys
               `(("root"
                 ,(local-file
                   "/home/charlie/.ssh/id_rsa.pub"))
                ("charlie"
                 ,(local-file
                   "/home/charlie/.ssh/id_rsa.pub"))))))))

(operating-system
;; ...
(services
;; Add the 'hurd-vm' service, configured to use the
;; operating system configuration above.
(append (list (service hurd-vm-service-type
                      (hurd-vm-configuration
                       (os %childhurd-os))))
        %base-services)))

```

That's it! The remainder of this section provides the reference of the service configuration.

#### **hurd-vm-service-type** [Variable]

This is the type of the Hurd in a Virtual Machine service. Its value must be a **hurd-vm-configuration** object, which specifies the operating system (see Section 11.3 [operating-system Reference], page 253) and the disk size for the Hurd Virtual Machine, the QEMU package to use as well as the options for running it.

For example:

```

(service hurd-vm-service-type
  (hurd-vm-configuration
   (disk-size (* 5000 (expt 2 20))) ;5G
   (memory-size 1024))) ;1024MiB

```

would create a disk image big enough to build GNU Hello, with some extra memory.

**hurd-vm-configuration** [Data Type]

The data type representing the configuration for **hurd-vm-service-type**.

**os** (default: *%hurd-vm-operating-system*)

The operating system to instantiate. This default is bare-bones with a permissive OpenSSH secure shell daemon listening on port 2222 (see Section 11.10.5 [Networking Services], page 322).

**qemu** (default: *qemu-minimal*)

The QEMU package to use.

**type** (default: *'hurd-qcow2'*)

The image type name. Use *'hurd-qcow2'* for a 32-bit image or *'hurd64-qcow2'* for a 64-bit image.

**image** (default: *hurd-vm-disk-image*)

The image object representing the disk image of this virtual machine (see Chapter 16 [System Images], page 793).

Alternatively, this can be the file name (a string) of an image, which is useful if you want the image to persist across service restarts—see example below.

**disk-size** (default: *'guess'*)

The size of the disk image.

**memory-size** (default: *512*)

The memory size of the Virtual Machine in mebibytes.

**options** (default: *'( "--snapshot")'*)

The extra options for running QEMU.

**id** (default: *#f*)

If set, a non-zero positive integer used to parameterize Childhurd instances. It is appended to the service's name, e.g. **childhurd1**.

**net-options** (default: *hurd-vm-net-options*)

The procedure used to produce the list of QEMU networking options.

By default, it produces

```
'( "--device" "rtl8139,netdev=net0"
  "--netdev" (string-append
    "user,id=net0,"
    "hostfwd=tcp:127.0.0.1:secrets-port-:1004,"
    "hostfwd=tcp:127.0.0.1:ssh-port-:2222,"
    "hostfwd=tcp:127.0.0.1:vnc-port-:5900"))
```

with forwarded ports:

```
secrets-port: (+ 11004 (* 1000 ID))
ssh-port: (+ 10022 (* 1000 ID))
vnc-port: (+ 15900 (* 1000 ID))
```

**offloading?** (default: *#t*)

Whether to automatically set up offloading of builds to the childhurd.

When enabled, this lets you run GNU/Hurd builds on the host and have them transparently offloaded to the VM, for instance when running a command like this:

```
guix build coreutils -s i586-gnu
```

This option automatically sets up offloading like so:

1. Authorizing the childhurd's key on the host so that the host accepts build results coming from the childhurd, which can be done like so (see Section 5.11 [Invoking guix archive], page 65, for more on that).
2. Creating a user account called **offloading** dedicated to offloading in the childhurd.
3. Creating an SSH key pair on the host and making it an authorized key of the **offloading** account in the childhurd.
4. Adding the childhurd to `/etc/guix/machines.scm` (see Section 2.2.2 [Daemon Offload Setup], page 9).

**secret-root** (default: `/etc/childhurd`)

The root directory with out-of-band secrets to be installed into the childhurd once it runs. Childhurds are volatile which means that on every startup, secrets such as the SSH host keys and Guix signing key are recreated.

If the `/etc/childhurd` directory does not exist, the **secret-service** running in the Childhurd will be sent an empty list of secrets.

By default, the service automatically populates `/etc/childhurd` with the following non-volatile secrets, unless they already exist:

```
/etc/childhurd/etc/guix/acl
/etc/childhurd/etc/guix/signing-key.pub
/etc/childhurd/etc/guix/signing-key.sec
/etc/childhurd/etc/ssh/authorized_keys.d/offloading
/etc/childhurd/etc/ssh/ssh_host_ed25519_key
/etc/childhurd/etc/ssh/ssh_host_ecdsa_key
/etc/childhurd/etc/ssh/ssh_host_ed25519_key.pub
/etc/childhurd/etc/ssh/ssh_host_ecdsa_key.pub
```

Note that by default the VM image is volatile—i.e., once the service is stopped, the contents are lost. If you want a persistent image instead, override the configuration's **image** and **options** without the `--snapshot` flag using something along these lines:

```
(service hurd-vm-service-type
  (hurd-vm-configuration
    ;; Specify an out-of-store, persistent image.
    (image "/out/of/store/writable/hurd.img")
    (options '())))
```

## Ganeti

**Note:** This service is considered experimental. Configuration options may be changed in a backwards-incompatible manner, and not all features have been



thoroughly tested. Users of this service are encouraged to share their experience at [guix-devel@gnu.org](mailto:guix-devel@gnu.org).

Ganeti is a virtual machine management system. It is designed to keep virtual machines running on a cluster of servers even in the event of hardware failures, and to make maintenance and recovery tasks easy. It consists of multiple services which are described later in this section. In addition to the Ganeti service, you will need the OpenSSH service (see Section 11.10.5 [Networking Services], page 322), and update the `/etc/hosts` file (see Section 11.19.3 [Service Reference], page 712) with the cluster name and address (or use a DNS server).

All nodes participating in a Ganeti cluster should have the same Ganeti and `/etc/hosts` configuration. Here is an example configuration for a Ganeti cluster node that supports multiple storage backends, and installs the `debootstrap` and `guix OS providers`:

```
(use-package-modules virtualization)
(use-service-modules base ganeti networking ssh)
(operating-system
  ;; ...
  (host-name "node1")

  ;; Install QEMU so we can use KVM-based instances, and LVM, DRBD and Ceph
  ;; in order to use the "plain", "drbd" and "rbd" storage backends.
  (packages (append (map specification->package
    '("qemu" "lvm2" "drbd-utils" "ceph"
      ;; Add the debootstrap and guix OS providers.
      "ganeti-instance-guix" "ganeti-instance-debootstrap"))
    %base-packages))

  (services
    (append (list (service static-networking-service-type
      (list (static-networking
        (addresses
          (list (network-address
            (device "eth0")
            (value "192.168.1.201/24")))))
        (routes
          (list (network-route
            (destination "default")
            (gateway "192.168.1.254"))))
        (name-servers '("192.168.1.252"
          "192.168.1.253")))))
      (service openssh-service-type
        (openssh-configuration
          (permit-root-login 'prohibit-password)))

      (simple-service 'ganeti-hosts-entries hosts-service-type
        (list
```

```

        (host "192.168.1.200" "ganeti.example.com")
        (host "192.168.1.201" "node1.example.com"
          ("node1"))
        (host "192.168.1.202" "node2.example.com"
          ("node2"))))

(service ganeti-service-type
  (ganeti-configuration
    ;; This list specifies allowed file system paths
    ;; for storing virtual machine images.
    (file-storage-paths '("/srv/ganeti/file-storage"))
    ;; This variable configures a single "variant" for
    ;; both Debootstrap and Guix that works with KVM.
    (os %default-ganeti-os)))
%base-services)))

```

Users are advised to read the Ganeti administrators guide (<https://docs.ganeti.org/docs/ganeti/3.0/html/admin.html>) to learn about the various cluster options and day-to-day operations. There is also a blog post (<https://guix.gnu.org/blog/2020/running-a-ganeti-cluster-on-guix/>) describing how to configure and initialize a small cluster.

**ganeti-service-type** [Variable]

This is a service type that includes all the various services that Ganeti nodes should run.

Its value is a **ganeti-configuration** object that defines the package to use for CLI operations, as well as configuration for the various daemons. Allowed file storage paths and available guest operating systems are also configured through this data type.

**ganeti-configuration** [Data Type]

The **ganeti** service takes the following configuration options:

**ganeti** (default: **ganeti**)

The **ganeti** package to use. It will be installed to the system profile and make **gnt-cluster**, **gnt-instance**, etc available. Note that the value specified here does not affect the other services as each refer to a specific **ganeti** package (see below).

```

noded-configuration (default: (ganeti-noded-configuration))
confd-configuration (default: (ganeti-confd-configuration))
wconfd-configuration (default: (ganeti-wconfd-configuration))
luxid-configuration (default: (ganeti-luxid-configuration))
rapi-configuration (default: (ganeti-rapi-configuration))
kvmd-configuration (default: (ganeti-kvmd-configuration))
mond-configuration (default: (ganeti-mond-configuration))
metad-configuration (default: (ganeti-metad-configuration))
watcher-configuration (default: (ganeti-watcher-configuration))
cleaner-configuration (default: (ganeti-cleaner-configuration))

```

These options control the various daemons and cron jobs that are distributed with Ganeti. The possible values for these are described in detail below. To override a setting, you must use the configuration type for that service:

```

(service ganeti-service-type
  (ganeti-configuration
    (rapi-configuration
      (ganeti-rapi-configuration
        (interface "eth1"))))
    (watcher-configuration
      (ganeti-watcher-configuration
        (rapi-ip "10.0.0.1"))))

```

```
file-storage-paths (default: '())
```

List of allowed directories for file storage backend.

```
hooks (default: #f)
```

When set, this should be a file-like object containing a directory with cluster execution hooks (<https://docs.ganeti.org/docs/ganeti/3.0/html/hooks.html>).

```
os (default: %default-ganeti-os)
```

List of <ganeti-os> records.

In essence `ganeti-service-type` is shorthand for declaring each service individually:

```

(service ganeti-noded-service-type)
(service ganeti-confd-service-type)
(service ganeti-wconfd-service-type)
(service ganeti-luxid-service-type)
(service ganeti-kvmd-service-type)
(service ganeti-mond-service-type)
(service ganeti-metad-service-type)
(service ganeti-watcher-service-type)
(service ganeti-cleaner-service-type)

```

Plus a service extension for `etc-service-type` that configures the file storage backend and OS variants.

**ganeti-os** [Data Type]

This data type is suitable for passing to the `os` parameter of `ganeti-configuration`. It takes the following parameters:

**name** The name for this OS provider. It is only used to specify where the configuration ends up. Setting it to “debootstrap” will create `/etc/ganeti/instance-debootstrap`.

**extension** (default: `#f`)  
The file extension for variants of this OS type. For example `.conf` or `.scm`. It will be appended to the variant file name if set.

**variants** (default: `'()`)  
This must be either a list of `ganeti-os-variant` objects for this OS, or a “file-like” object (see Section 8.12 [G-Expressions], page 165) representing the variants directory.

To use the Guix OS provider with variant definitions residing in a local directory instead of declaring individual variants (see `guix-variants` below), you can do:

```
(ganeti-os
 (name "guix")
 (variants (local-file "ganeti-guix-variants"
                      #:recursive? #true)))
```

Note that you will need to maintain the `variants.list` file (see `ganeti-os-interface(7)` (<https://docs.ganeti.org/docs/ganeti/3.0/man/ganeti-os-interface.html>)) manually in this case.

**ganeti-os-variant** [Data Type]

This is the data type for a Ganeti OS variant. It takes the following parameters:

**name** The name of this variant.

**configuration**  
A configuration file for this variant.

**%default-debootstrap-hooks** [Variable]

This variable contains hooks to configure networking and the GRUB bootloader.

**%default-debootstrap-extra-pkgs** [Variable]

This variable contains a list of packages suitable for a fully-virtualized guest.

**debootstrap-configuration** [Data Type]

This data type creates configuration files suitable for the debootstrap OS provider.

**hooks** (default: `%default-debootstrap-hooks`)

When not `#f`, this must be a G-expression that specifies a directory with scripts that will run when the OS is installed. It can also be a list of `(name . file-like)` pairs. For example:

```
`((99-hello-world . ,(plain-file "#!/bin/sh\nnecho Hello, World")))
```

That will create a directory with one executable named `99-hello-world` and run it every time this variant is installed. If set to `#f`, hooks in `/etc/ganeti/instance-debootstrap/hooks` will be used, if any.

**proxy** (default: **#f**)  
Optional HTTP proxy to use.

**mirror** (default: **#f**)  
The Debian mirror. Typically something like `http://ftp.no.debian.org/debian.`  
The default varies depending on the distribution.

**arch** (default: **#f**)  
The dpkg architecture. Set to **armhf** to debootstrap an ARMv7 instance on an AArch64 host. Default is to use the current system architecture.

**suite** (default: **"stable"**)  
When set, this must be a Debian distribution “suite” such as **buster** or **focal**. If set to **#f**, the default for the OS provider is used.

**extra-pkgs** (default: **%default-debootstrap-extra-pkgs**)  
List of extra packages that will get installed by dpkg in addition to the minimal system.

**components** (default: **#f**)  
When set, must be a list of Debian repository “components”. For example **'("main" "contrib")**.

**generate-cache?** (default: **#t**)  
Whether to automatically cache the generated debootstrap archive.

**clean-cache** (default: **14**)  
Discard the cache after this amount of days. Use **#f** to never clear the cache.

**partition-style** (default: **'msdos**)  
The type of partition to create. When set, it must be one of **'msdos**, **'none** or a string.

**partition-alignment** (default: **2048**)  
Alignment of the partition in sectors.

**debootstrap-variant** *name configuration* [Procedure]  
This is a helper procedure that creates a **ganeti-os-variant** record. It takes two parameters: a name and a **debootstrap-configuration** object.

**debootstrap-os** *variants...* [Procedure]  
This is a helper procedure that creates a **ganeti-os** record. It takes a list of variants created with **debootstrap-variant**.

**guix-variant** *name configuration* [Procedure]  
This is a helper procedure that creates a **ganeti-os-variant** record for use with the Guix OS provider. It takes a name and a G-expression that returns a “file-like” (see Section 8.12 [G-Expressions], page 165) object containing a Guix System configuration.

**guix-os** *variants...* [Procedure]  
This is a helper procedure that creates a **ganeti-os** record. It takes a list of variants produced by **guix-variant**.

**%default-debootstrap-variants** [Variable]

This is a convenience variable to make the debootstrap provider work “out of the box” without users having to declare variants manually. It contains a single debootstrap variant with the default configuration:

```
(list (debootstrap-variant
      "default"
      (debootstrap-configuration)))
```

**%default-guix-variants** [Variable]

This is a convenience variable to make the Guix OS provider work without additional configuration. It creates a virtual machine that has an SSH server, a serial console, and authorizes the Ganeti hosts SSH keys.

```
(list (guix-variant
      "default"
      (file-append ganeti-instance-guix
                    "/share/doc/ganeti-instance-guix/examples/dynamic.scm")))
```

Users can implement support for OS providers unbeknownst to Guix by extending the `ganeti-os` and `ganeti-os-variant` records appropriately. For example:

```
(ganeti-os
 (name "custom")
 (extension ".conf")
 (variants
  (list (ganeti-os-variant
        (name "foo")
        (configuration (plain-file "bar" "this is fine"))))))
```

That creates `/etc/ganeti/instance-custom/variants/foo.conf` which points to a file in the store with contents `this is fine`. It also creates `/etc/ganeti/instance-custom/variants/variants.list` with contents `foo`.

Obviously this may not work for all OS providers out there. If you find the interface limiting, please reach out to [guix-devel@gnu.org](mailto:guix-devel@gnu.org).

The rest of this section documents the various services that are included by `ganeti-service-type`.

**ganeti-noded-service-type** [Variable]

`ganeti-noded` is the daemon responsible for node-specific functions within the Ganeti system. The value of this service must be a `ganeti-noded-configuration` object.

**ganeti-noded-configuration** [Data Type]

This is the configuration for the `ganeti-noded` service.

**ganeti** (default: `ganeti`)

The `ganeti` package to use for this service.

**port** (default: 1811)

The TCP port on which the node daemon listens for network requests.

**address** (default: "0.0.0.0")

The network address that the daemon will bind to. The default address means bind to all available addresses.

**interface** (default: **#f**)

When this is set, it must be a specific network interface (e.g. **eth0**) that the daemon will bind to.

**max-clients** (default: 20)

This sets a limit on the maximum number of simultaneous client connections that the daemon will handle. Connections above this count are accepted, but no responses will be sent until enough connections have closed.

**ssl?** (default: **#t**)

Whether to use SSL/TLS to encrypt network communications. The certificate is automatically provisioned by the cluster and can be rotated with **gnt-cluster renew-crypto**.

**ssl-key** (default: **"/var/lib/ganeti/server.pem"**)

This can be used to provide a specific encryption key for TLS communications.

**ssl-cert** (default: **"/var/lib/ganeti/server.pem"**)

This can be used to provide a specific certificate for TLS communications.

**debug?** (default: **#f**)

When true, the daemon performs additional logging for debugging purposes. Note that this will leak encryption details to the log files, use with caution.

**ganeti-confd-service-type** [Variable]

**ganeti-confd** answers queries related to the configuration of a Ganeti cluster. The purpose of this daemon is to have a highly available and fast way to query cluster configuration values. It is automatically active on all *master candidates*. The value of this service must be a **ganeti-confd-configuration** object.

**ganeti-confd-configuration** [Data Type]

This is the configuration for the **ganeti-confd** service.

**ganeti** (default: **ganeti**)

The **ganeti** package to use for this service.

**port** (default: 1814)

The UDP port on which to listen for network requests.

**address** (default: **"0.0.0.0"**)

Network address that the daemon will bind to.

**debug?** (default: **#f**)

When true, the daemon performs additional logging for debugging purposes.

**ganeti-wconfd-service-type** [Variable]

**ganeti-wconfd** is the daemon that has authoritative knowledge about the cluster configuration and is the only entity that can accept changes to it. All jobs that need

to modify the configuration will do so by sending appropriate requests to this daemon. It only runs on the *master node* and will automatically disable itself on other nodes. The value of this service must be a `ganeti-wconfd-configuration` object.

**ganeti-wconfd-configuration** [Data Type]

This is the configuration for the `ganeti-wconfd` service.

`ganeti` (default: `ganeti`)

The `ganeti` package to use for this service.

`no-voting?` (default: `#f`)

The daemon will refuse to start if the majority of cluster nodes does not agree that it is running on the master node. Set to `#t` to start even if a quorum can not be reached (dangerous, use with caution).

`debug?` (default: `#f`)

When true, the daemon performs additional logging for debugging purposes.

**ganeti-luxid-service-type** [Variable]

`ganeti-luxid` is a daemon used to answer queries related to the configuration and the current live state of a Ganeti cluster. Additionally, it is the authoritative daemon for the Ganeti job queue. Jobs can be submitted via this daemon and it schedules and starts them.

It takes a `ganeti-luxid-configuration` object.

**ganeti-luxid-configuration** [Data Type]

This is the configuration for the `ganeti-luxid` service.

`ganeti` (default: `ganeti`)

The `ganeti` package to use for this service.

`no-voting?` (default: `#f`)

The daemon will refuse to start if it cannot verify that the majority of cluster nodes believes that it is running on the master node. Set to `#t` to ignore such checks and start anyway (this can be dangerous).

`debug?` (default: `#f`)

When true, the daemon performs additional logging for debugging purposes.

**ganeti-rapi-service-type** [Variable]

`ganeti-rapi` provides a remote API for Ganeti clusters. It runs on the master node and can be used to perform cluster actions programmatically via a JSON-based RPC protocol.

Most query operations are allowed without authentication (unless `require-authentication?` is set), whereas write operations require explicit authorization via the `/var/lib/ganeti/rapi/users` file. See the Ganeti Remote API documentation (<https://docs.ganeti.org/docs/ganeti/3.0/html/rapi.html>) for more information.

The value of this service must be a `ganeti-rapi-configuration` object.



**ganeti-rapi-configuration** [Data Type]

This is the configuration for the **ganeti-rapi** service.

**ganeti** (default: **ganeti**)

The **ganeti** package to use for this service.

**require-authentication?** (default: **#f**)

Whether to require authentication even for read-only operations.

**port** (default: 5080)

The TCP port on which to listen to API requests.

**address** (default: "0.0.0.0")

The network address that the service will bind to. By default it listens on all configured addresses.

**interface** (default: **#f**)

When set, it must specify a specific network interface such as **eth0** that the daemon will bind to.

**max-clients** (default: 20)

The maximum number of simultaneous client requests to handle. Further connections are allowed, but no responses are sent until enough connections have closed.

**ssl?** (default: **#t**)

Whether to use SSL/TLS encryption on the RAPI port.

**ssl-key** (default: "/var/lib/ganeti/server.pem")

This can be used to provide a specific encryption key for TLS communications.

**ssl-cert** (default: "/var/lib/ganeti/server.pem")

This can be used to provide a specific certificate for TLS communications.

**debug?** (default: **#f**)

When true, the daemon performs additional logging for debugging purposes. Note that this will leak encryption details to the log files, use with caution.

**ganeti-kvmd-service-type** [Variable]

**ganeti-kvmd** is responsible for determining whether a given KVM instance was shut down by an administrator or a user. Normally Ganeti will restart an instance that was not stopped through Ganeti itself. If the cluster option **user\_shutdown** is true, this daemon monitors the QMP socket provided by QEMU and listens for shutdown events, and marks the instance as *USER\_down* instead of *ERROR\_down* when it shuts down gracefully by itself.

It takes a **ganeti-kvmd-configuration** object.

**ganeti-kvmd-configuration** [Data Type]

**ganeti** (default: **ganeti**)

The **ganeti** package to use for this service.

**debug?** (default: **#f**)

When true, the daemon performs additional logging for debugging purposes.

**ganeti-mond-service-type**

[Variable]

**ganeti-mond** is an optional daemon that provides Ganeti monitoring functionality. It is responsible for running data collectors and publish the collected information through a HTTP interface.

It takes a **ganeti-mond-configuration** object.

**ganeti-mond-configuration**

[Data Type]

**ganeti** (default: **ganeti**)

The **ganeti** package to use for this service.

**port** (default: 1815)

The port on which the daemon will listen.

**address** (default: "0.0.0.0")

The network address that the daemon will bind to. By default it binds to all available interfaces.

**debug?** (default: **#f**)

When true, the daemon performs additional logging for debugging purposes.

**ganeti-metad-service-type**

[Variable]

**ganeti-metad** is an optional daemon that can be used to provide information about the cluster to instances or OS install scripts.

It takes a **ganeti-metad-configuration** object.

**ganeti-metad-configuration**

[Data Type]

**ganeti** (default: **ganeti**)

The **ganeti** package to use for this service.

**port** (default: 80)

The port on which the daemon will listen.

**address** (default: **#f**)

If set, the daemon will bind to this address only. If left unset, the behavior depends on the cluster configuration.

**debug?** (default: **#f**)

When true, the daemon performs additional logging for debugging purposes.

**ganeti-watcher-service-type**

[Variable]

**ganeti-watcher** is a script designed to run periodically and ensure the health of a cluster. It will automatically restart instances that have stopped without Ganeti's consent, and repairs DRBD links in case a node has rebooted. It also archives old cluster jobs and restarts Ganeti daemons that are not running. If the cluster parameter **ensure\_node\_health** is set, the watcher will also shutdown instances and DRBD devices if the node it is running on is declared offline by known master candidates.

It can be paused on all nodes with `gnt-cluster watcher pause`.

The service takes a `ganeti-watcher-configuration` object.

**ganeti-watcher-configuration** [Data Type]

`ganeti` (default: `ganeti`)

The `ganeti` package to use for this service.

`schedule` (default: `"*/5 * * * *"`)

When to run the script, expressed either as a string in traditional cron syntax or as a gexp representing a Shepherd calendar event (see Section “Timers” in *The GNU Shepherd Manual*). The default is every five minutes.

`rapi-ip` (default: `#f`)

This option needs to be specified only if the RAPI daemon is configured to use a particular interface or address. By default the cluster address is used.

`job-age` (default: `(* 6 3600)`)

Archive cluster jobs older than this age, specified in seconds. The default is 6 hours. This keeps `gnt-job list` manageable.

`verify-disks?` (default: `#t`)

If this is `#f`, the watcher will not try to repair broken DRBD links automatically. Administrators will need to use `gnt-cluster verify-disks` manually instead.

`debug?` (default: `#f`)

When `#t`, the script performs additional logging for debugging purposes.

**ganeti-cleaner-service-type** [Variable]

`ganeti-cleaner` is a script designed to run periodically and remove old files from the cluster. This service type controls two *cron jobs*: one intended for the master node that permanently purges old cluster jobs, and one intended for every node that removes expired X509 certificates, keys, and outdated `ganeti-watcher` information. Like all Ganeti services, it is safe to include even on non-master nodes as it will disable itself as necessary.

It takes a `ganeti-cleaner-configuration` object.

**ganeti-cleaner-configuration** [Data Type]

`ganeti` (default: `ganeti`)

The `ganeti` package to use for the `gnt-cleaner` command.

`master-schedule` (default: `"45 1 * * *"`)

How often to run the master cleaning job. The default is once per day, at 01:45:00.

`node-schedule` (default: `"45 2 * * *"`)

How often to run the node cleaning job. The default is once per day, at 02:45:00.

### 11.10.33 Version Control Services

The (`gnu services version-control`) module provides a service to allow remote access to local Git repositories. There are three options: the `git-daemon-service-type`, which provides access to repositories via the `git://` unsecured TCP-based protocol, extending the `nginx` web server to proxy some requests to `git-http-backend`, or providing a web interface with `cgit-service-type`.

**git-daemon-service-type** [Variable]

Type for a service that runs `git daemon`, a simple TCP server to expose repositories over the Git protocol for anonymous access.

The value for this service type is a `<git-daemon-configuration>` record, by default it allows read-only access to exported<sup>9</sup> repositories under `/srv/git`.

**git-daemon-configuration** [Data Type]

Data type representing the configuration for `git-daemon-service-type`.

**package** (default: `git`)

Package object of the Git distributed version control system.

**export-all?** (default: `#f`)

Whether to allow access for all Git repositories, even if they do not have the `git-daemon-export-ok` file.

**base-path** (default: `/srv/git`)

Whether to remap all the path requests as relative to the given path. If you run `git daemon` with (`base-path "/srv/git"`) on `'example.com'`, then if you later try to pull `'git://example.com/hello.git'`, `git daemon` will interpret the path as `/srv/git/hello.git`.

**user-path** (default: `#f`)

Whether to allow `~user` notation to be used in requests. When specified with empty string, requests to `'git://host/~alice/foo'` is taken as a request to access `foo` repository in the home directory of user `alice`. If (`user-path "path"`) is specified, the same request is taken as a request to access `path/foo` repository in the home directory of user `alice`.

**listen** (default: `'()`)

Whether to listen on specific IP addresses or hostnames, defaults to all.

**port** (default: `#f`)

Whether to listen on an alternative port, which defaults to 9418.

**whitelist** (default: `'()`)

If not empty, only allow access to this list of directories.

**extra-options** (default: `'()`)

Extra options that will be passed to `git daemon`.<sup>10</sup>

<sup>9</sup> By creating the magic file `git-daemon-export-ok` in the repository directory.

<sup>10</sup> Run `man git-daemon` for more information.

The `git://` protocol lacks authentication. When you pull from a repository fetched via `git://`, you don't know whether the data you receive was modified or is even coming from the specified host, and your connection is subject to eavesdropping. It's better to use an authenticated and encrypted transport, such as `https`. Although Git allows you to serve repositories using unsophisticated file-based web servers, there is a faster protocol implemented by the `git-http-backend` program. This program is the back-end of a proper Git web service. It is designed to sit behind a FastCGI proxy. See Section 11.10.20 [Web Services], page 494, for more on running the necessary `fcgiwrap` daemon.

Guix has a separate configuration data type for serving Git repositories over HTTP.

**git-http-configuration** [Data Type]

Data type representing the configuration for a future `git-http-service-type`; can currently be used to configure Nginx through `git-http-nginx-location-configuration`.

**package** (default: `git`)

Package object of the Git distributed version control system.

**git-root** (default: `/srv/git`)

Directory containing the Git repositories to expose to the world.

**export-all?** (default: `#f`)

Whether to expose access for all Git repositories in `git-root`, even if they do not have the `git-daemon-export-ok` file.

**uri-path** (default: `"/git/"`)

Path prefix for Git access. With the default `"/git/"` prefix, this will map `'http://server/git/repo.git'` to `/srv/git/repo.git`. Requests whose URI paths do not begin with this prefix are not passed on to this Git instance.

**fcgiwrap-socket** (default: `127.0.0.1:9000`)

The socket on which the `fcgiwrap` daemon is listening. See Section 11.10.20 [Web Services], page 494.

There is no `git-http-service-type`, currently; instead you can create an `nginx-location-configuration` from a `git-http-configuration` and then add that location to a web server.

**git-http-nginx-location-configuration** [Procedure]

`[config=(git-http-configuration)]`

Compute an `nginx-location-configuration` that corresponds to the given Git http configuration. An example nginx service definition to serve the default `/srv/git` over HTTPS might be:

```
(service nginx-service-type
  (nginx-configuration
    (server-blocks
      (list
        (nginx-server-configuration
          (listen '("443 ssl"))
```

```
(server-name "git.my-host.org")
(ssl-certificate
  "/etc/certs/git.my-host.org/fullchain.pem")
(ssl-certificate-key
  "/etc/certs/git.my-host.org/privkey.pem")
(locations
  (list
    (git-http-nginx-location-configuration
      (git-http-configuration (uri-path "/"))))))))
```

This example assumes that you are using Let's Encrypt to get your TLS certificate. See Section 11.10.22 [Certificate Services], page 518. The default `certbot` service will redirect all HTTP traffic on `git.my-host.org` to HTTPS. You will also need to add an `fcgiwrap` proxy to your system services. See Section 11.10.20 [Web Services], page 494.

## Cgit Service

Cgit (<https://git.zx2c4.com/cgit/>) is a web frontend for Git repositories written in C.

The following example will configure the service with default values. By default, Cgit can be accessed on port 80 (<http://localhost:80>).

```
(service cgit-service-type)
```

The `file-object` type designates either a file-like object (see Section 8.12 [G-Expressions], page 165) or a string.

Available `cgit-configuration` fields are:

<b>package package</b>	[ <code>cgit-configuration</code> parameter]
The CGIT package.	

<b>nginx-server-configuration-list nginx</b>	[ <code>cgit-configuration</code> parameter]
NGINX configuration.	

<b>file-object about-filter</b>	[ <code>cgit-configuration</code> parameter]
Specifies a command which will be invoked to format the content of about pages (both top-level and for each repository).	
Defaults to <code>""</code> .	

<b>string agefile</b>	[ <code>cgit-configuration</code> parameter]
Specifies a path, relative to each repository path, which can be used to specify the date and time of the youngest commit in the repository.	
Defaults to <code>""</code> .	

<b>file-object auth-filter</b>	[ <code>cgit-configuration</code> parameter]
Specifies a command that will be invoked for authenticating repository access.	
Defaults to <code>""</code> .	

<b>string branch-sort</b>	[ <code>cgit-configuration</code> parameter]
Flag which, when set to <code>'age'</code> , enables date ordering in the branch ref list, and when set <code>'name'</code> enables ordering by branch name.	
Defaults to <code>"name"</code> .	

- string cache-root** [cgit-configuration parameter]  
Path used to store the cgit cache entries.  
Defaults to `"/var/cache/cgit"`.
- integer cache-static-ttl** [cgit-configuration parameter]  
Number which specifies the time-to-live, in minutes, for the cached version of repository pages accessed with a fixed SHA1.  
Defaults to `'-1'`.
- integer cache-dynamic-ttl** [cgit-configuration parameter]  
Number which specifies the time-to-live, in minutes, for the cached version of repository pages accessed without a fixed SHA1.  
Defaults to `'5'`.
- integer cache-repo-ttl** [cgit-configuration parameter]  
Number which specifies the time-to-live, in minutes, for the cached version of the repository summary page.  
Defaults to `'5'`.
- integer cache-root-ttl** [cgit-configuration parameter]  
Number which specifies the time-to-live, in minutes, for the cached version of the repository index page.  
Defaults to `'5'`.
- integer cache-scanrc-ttl** [cgit-configuration parameter]  
Number which specifies the time-to-live, in minutes, for the result of scanning a path for Git repositories.  
Defaults to `'15'`.
- integer cache-about-ttl** [cgit-configuration parameter]  
Number which specifies the time-to-live, in minutes, for the cached version of the repository about page.  
Defaults to `'15'`.
- integer cache-snapshot-ttl** [cgit-configuration parameter]  
Number which specifies the time-to-live, in minutes, for the cached version of snapshots.  
Defaults to `'5'`.
- integer cache-size** [cgit-configuration parameter]  
The maximum number of entries in the cgit cache. When set to `'0'`, caching is disabled.  
Defaults to `'0'`.
- boolean case-sensitive-sort?** [cgit-configuration parameter]  
Sort items in the repo list case sensitively.  
Defaults to `'#t'`.

- list clone-prefix** [cgit-configuration parameter]  
List of common prefixes which, when combined with a repository URL, generates valid clone URLs for the repository.  
Defaults to ‘‘()’.
- list clone-url** [cgit-configuration parameter]  
List of clone-url templates.  
Defaults to ‘‘()’.
- file-object commit-filter** [cgit-configuration parameter]  
Command which will be invoked to format commit messages.  
Defaults to ‘‘’’.
- string commit-sort** [cgit-configuration parameter]  
Flag which, when set to ‘date’, enables strict date ordering in the commit log, and when set to ‘topo’ enables strict topological ordering.  
Defaults to ‘‘git log’’.
- file-object css** [cgit-configuration parameter]  
URL which specifies the css document to include in all cgit pages.  
Defaults to ‘‘/share/cgit/cgit.css’’.
- file-object email-filter** [cgit-configuration parameter]  
Specifies a command which will be invoked to format names and email address of committers, authors, and taggers, as represented in various places throughout the cgit interface.  
Defaults to ‘‘’’.
- boolean embedded?** [cgit-configuration parameter]  
Flag which, when set to ‘#t’, will make cgit generate a HTML fragment suitable for embedding in other HTML pages.  
Defaults to ‘#f’.
- boolean enable-commit-graph?** [cgit-configuration parameter]  
Flag which, when set to ‘#t’, will make cgit print an ASCII-art commit history graph to the left of the commit messages in the repository log page.  
Defaults to ‘#f’.
- boolean enable-filter-overrides?** [cgit-configuration parameter]  
Flag which, when set to ‘#t’, allows all filter settings to be overridden in repository-specific cgitrc files.  
Defaults to ‘#f’.
- boolean enable-follow-links?** [cgit-configuration parameter]  
Flag which, when set to ‘#t’, allows users to follow a file in the log view.  
Defaults to ‘#f’.



**boolean enable-http-clone?** [cgit-configuration parameter]  
If set to '#t', cgit will act as an dumb HTTP endpoint for Git clones.  
Defaults to '#t'.

**boolean enable-index-links?** [cgit-configuration parameter]  
Flag which, when set to '#t', will make cgit generate extra links "summary", "commit", "tree" for each repo in the repository index.  
Defaults to '#f'.

**boolean enable-index-owner?** [cgit-configuration parameter]  
Flag which, when set to '#t', will make cgit display the owner of each repo in the repository index.  
Defaults to '#t'.

**boolean enable-log-filecount?** [cgit-configuration parameter]  
Flag which, when set to '#t', will make cgit print the number of modified files for each commit on the repository log page.  
Defaults to '#f'.

**boolean enable-log-linecount?** [cgit-configuration parameter]  
Flag which, when set to '#t', will make cgit print the number of added and removed lines for each commit on the repository log page.  
Defaults to '#f'.

**boolean enable-remote-branches?** [cgit-configuration parameter]  
Flag which, when set to #t, will make cgit display remote branches in the summary and refs views.  
Defaults to '#f'.

**boolean enable-subject-links?** [cgit-configuration parameter]  
Flag which, when set to 1, will make cgit use the subject of the parent commit as link text when generating links to parent commits in commit view.  
Defaults to '#f'.

**boolean enable-html-serving?** [cgit-configuration parameter]  
Flag which, when set to '#t', will make cgit use the subject of the parent commit as link text when generating links to parent commits in commit view.  
Defaults to '#f'.

**boolean enable-tree-linenumbers?** [cgit-configuration parameter]  
Flag which, when set to '#t', will make cgit generate linenummer links for plaintext blobs printed in the tree view.  
Defaults to '#t'.

**boolean enable-git-config?** [cgit-configuration parameter]  
Flag which, when set to '#f', will allow cgit to use Git config to set any repo specific settings.  
Defaults to '#f'.

- file-object favicon** [cgit-configuration parameter]  
URL used as link to a shortcut icon for cgit.  
Defaults to `"/favicon.ico"`.
- string footer** [cgit-configuration parameter]  
The content of the file specified with this option will be included verbatim at the bottom of all pages (i.e. it replaces the standard "generated by..." message).  
Defaults to `""`.
- string head-include** [cgit-configuration parameter]  
The content of the file specified with this option will be included verbatim in the HTML HEAD section on all pages.  
Defaults to `""`.
- string header** [cgit-configuration parameter]  
The content of the file specified with this option will be included verbatim at the top of all pages.  
Defaults to `""`.
- file-object include** [cgit-configuration parameter]  
Name of a configfile to include before the rest of the current config- file is parsed.  
Defaults to `""`.
- string index-header** [cgit-configuration parameter]  
The content of the file specified with this option will be included verbatim above the repository index.  
Defaults to `""`.
- string index-info** [cgit-configuration parameter]  
The content of the file specified with this option will be included verbatim below the heading on the repository index page.  
Defaults to `""`.
- boolean local-time?** [cgit-configuration parameter]  
Flag which, if set to `#t`, makes cgit print commit and tag times in the servers timezone.  
Defaults to `#f`.
- file-object logo** [cgit-configuration parameter]  
URL which specifies the source of an image which will be used as a logo on all cgit pages.  
Defaults to `"/share/cgit/cgit.png"`.
- string logo-link** [cgit-configuration parameter]  
URL loaded when clicking on the cgit logo image.  
Defaults to `""`.

- file-object owner-filter** [cgit-configuration parameter]  
Command which will be invoked to format the Owner column of the main page.  
Defaults to `''`.
- integer max-atom-items** [cgit-configuration parameter]  
Number of items to display in atom feeds view.  
Defaults to `'10'`.
- integer max-commit-count** [cgit-configuration parameter]  
Number of entries to list per page in "log" view.  
Defaults to `'50'`.
- integer max-message-length** [cgit-configuration parameter]  
Number of commit message characters to display in "log" view.  
Defaults to `'80'`.
- integer max-repo-count** [cgit-configuration parameter]  
Specifies the number of entries to list per page on the repository index page.  
Defaults to `'50'`.
- integer max-repodesc-length** [cgit-configuration parameter]  
Specifies the maximum number of repo description characters to display on the repository index page.  
Defaults to `'80'`.
- integer max-blob-size** [cgit-configuration parameter]  
Specifies the maximum size of a blob to display HTML for in KBytes.  
Defaults to `'0'`.
- string max-stats** [cgit-configuration parameter]  
Maximum statistics period. Valid values are `'week'`, `'month'`, `'quarter'` and `'year'`.  
Defaults to `''`.
- mimetype-alist mimetype** [cgit-configuration parameter]  
Mimetype for the specified filename extension.  
Defaults to `'((gif "image/gif") (html "text/html") (jpg "image/jpeg") (jpeg "image/jpeg") (pdf "application/pdf") (png "image/png") (svg "image/svg+xml"))'`.
- file-object mimetype-file** [cgit-configuration parameter]  
Specifies the file to use for automatic mimetype lookup.  
Defaults to `''`.
- string module-link** [cgit-configuration parameter]  
Text which will be used as the formatstring for a hyperlink when a submodule is printed in a directory listing.  
Defaults to `''`.

- boolean nocache?** [cgit-configuration parameter]  
If set to the value `#t` caching will be disabled.  
Defaults to `#f`.
- boolean noplainemail?** [cgit-configuration parameter]  
If set to `#t` showing full author email addresses will be disabled.  
Defaults to `#f`.
- boolean noheader?** [cgit-configuration parameter]  
Flag which, when set to `#t`, will make cgit omit the standard header on all pages.  
Defaults to `#f`.
- project-list project-list** [cgit-configuration parameter]  
A list of subdirectories inside of `repository-directory`, relative to it, that should loaded as Git repositories. An empty list means that all subdirectories will be loaded.  
Defaults to `'()'`.
- file-object readme** [cgit-configuration parameter]  
Text which will be used as default `repository-cgit-configuration` readme.  
Defaults to `""`.
- boolean remove-suffix?** [cgit-configuration parameter]  
If set to `#t` and `repository-directory` is enabled, if any repositories are found with a suffix of `.git`, this suffix will be removed for the URL and name.  
Defaults to `#f`.
- integer renamelimit** [cgit-configuration parameter]  
Maximum number of files to consider when detecting renames.  
Defaults to `-1`.
- string repository-sort** [cgit-configuration parameter]  
The way in which repositories in each section are sorted.  
Defaults to `""`.
- robots-list robots** [cgit-configuration parameter]  
Text used as content for the `robots` meta-tag.  
Defaults to `'("noindex" "nofollow")'`.
- string root-desc** [cgit-configuration parameter]  
Text printed below the heading on the repository index page.  
Defaults to `"a fast webinterface for the git dscm"`.
- string root-readme** [cgit-configuration parameter]  
The content of the file specified with this option will be included verbatim below the "about" link on the repository index page.  
Defaults to `""`.

- string root-title** [cgit-configuration parameter]  
Text printed as heading on the repository index page.  
Defaults to `""`.
- boolean scan-hidden-path** [cgit-configuration parameter]  
If set to `#t` and repository-directory is enabled, repository-directory will recurse into directories whose name starts with a period. Otherwise, repository-directory will stay away from such directories, considered as “hidden”. Note that this does not apply to the `.git` directory in non-bare repos.  
Defaults to `#f`.
- list snapshots** [cgit-configuration parameter]  
Text which specifies the default set of snapshot formats that cgit generates links for.  
Defaults to `'()'`.
- repository-directory** [cgit-configuration parameter]  
**repository-directory**  
Name of the directory to scan for repositories (represents `scan-path`).  
Defaults to `"/srv/git"`.
- string section** [cgit-configuration parameter]  
The name of the current repository section - all repositories defined after this option will inherit the current section name.  
Defaults to `""`.
- string section-sort** [cgit-configuration parameter]  
Flag which, when set to `'1'`, will sort the sections on the repository listing by name.  
Defaults to `""`.
- integer section-from-path** [cgit-configuration parameter]  
A number which, if defined prior to repository-directory, specifies how many path elements from each repo path to use as a default section name.  
Defaults to `'0'`.
- boolean side-by-side-diffs?** [cgit-configuration parameter]  
If set to `#t` shows side-by-side diffs instead of unidiffs per default.  
Defaults to `#f`.
- file-object source-filter** [cgit-configuration parameter]  
Specifies a command which will be invoked to format plaintext blobs in the tree view.  
Defaults to `""`.
- integer summary-branches** [cgit-configuration parameter]  
Specifies the number of branches to display in the repository “summary” view.  
Defaults to `'10'`.
- integer summary-log** [cgit-configuration parameter]  
Specifies the number of log entries to display in the repository “summary” view.  
Defaults to `'10'`.

**integer summary-tags** [cgit-configuration parameter]  
Specifies the number of tags to display in the repository “summary” view.  
Defaults to ‘10’.

**string strict-export** [cgit-configuration parameter]  
Filename which, if specified, needs to be present within the repository for cgit to allow access to that repository.  
Defaults to ‘’.

**string virtual-root** [cgit-configuration parameter]  
URL which, if specified, will be used as root for all cgit links.  
Defaults to ‘/’.

**repository-cgit-configuration-list** [cgit-configuration parameter]  
**repositories**  
A list of repository-cgit-configuration records.  
Defaults to ‘()’.  
Available repository-cgit-configuration fields are:

**repo-list snapshots** [repository-cgit-configuration parameter]  
A mask of snapshot formats for this repo that cgit generates links for, restricted by the global **snapshots** setting.  
Defaults to ‘()’.

**repo-file-object** [repository-cgit-configuration parameter]  
**source-filter**  
Override the default **source-filter**.  
Defaults to ‘’.

**repo-string url** [repository-cgit-configuration parameter]  
The relative URL used to access the repository.  
Defaults to ‘’.

**repo-file-object** [repository-cgit-configuration parameter]  
**about-filter**  
Override the default **about-filter**.  
Defaults to ‘’.

**repo-string branch-sort** [repository-cgit-configuration parameter]  
Flag which, when set to ‘age’, enables date ordering in the branch ref list, and when set to ‘name’ enables ordering by branch name.  
Defaults to ‘’.

**repo-list clone-url** [repository-cgit-configuration parameter]  
A list of URLs which can be used to clone repo.  
Defaults to ‘()’.

`repo-file-object` [repository-cgit-configuration parameter]  
    `commit-filter`  
    Override the default `commit-filter`.  
    Defaults to `""`.

`repo-string commit-sort` [repository-cgit-configuration parameter]  
    Flag which, when set to `'date'`, enables strict date ordering in the commit log,  
    and when set to `'topo'` enables strict topological ordering.  
    Defaults to `""`.

`repo-string defbranch` [repository-cgit-configuration parameter]  
    The name of the default branch for this repository. If no such branch exists in  
    the repository, the first branch name (when sorted) is used as default instead.  
    By default branch pointed to by HEAD, or "master" if there is no suitable  
    HEAD.  
    Defaults to `""`.

`repo-string desc` [repository-cgit-configuration parameter]  
    The value to show as repository description.  
    Defaults to `""`.

`repo-string homepage` [repository-cgit-configuration parameter]  
    The value to show as repository homepage.  
    Defaults to `""`.

`repo-file-object` [repository-cgit-configuration parameter]  
    `email-filter`  
    Override the default `email-filter`.  
    Defaults to `""`.

`maybe-repo-boolean` [repository-cgit-configuration parameter]  
    `enable-commit-graph?`  
    A flag which can be used to disable the global setting `enable-commit-graph?`.  
    Defaults to `'disabled'`.

`maybe-repo-boolean` [repository-cgit-configuration parameter]  
    `enable-log-filecount?`  
    A flag which can be used to disable the global setting `enable-log-filecount?`.  
    Defaults to `'disabled'`.

`maybe-repo-boolean` [repository-cgit-configuration parameter]  
    `enable-log-linecount?`  
    A flag which can be used to disable the global setting `enable-log-linecount?`.  
    Defaults to `'disabled'`.

`maybe-repo-boolean` [repository-cgit-configuration parameter]  
    `enable-remote-branches?`  
    Flag which, when set to `#t`, will make cgit display remote branches in the  
    summary and refs views.  
    Defaults to `'disabled'`.

**maybe-repo-boolean enable-subject-links?** [repository-cgit-configuration parameter]  
A flag which can be used to override the global setting `enable-subject-links?`.  
Defaults to `'disabled'`.

**maybe-repo-boolean enable-html-serving?** [repository-cgit-configuration parameter]  
A flag which can be used to override the global setting `enable-html-serving?`.  
Defaults to `'disabled'`.

**repo-boolean hide?** [repository-cgit-configuration parameter]  
Flag which, when set to `#t`, hides the repository from the repository index.  
Defaults to `#f`.

**repo-boolean ignore?** [repository-cgit-configuration parameter]  
Flag which, when set to `#t`, ignores the repository.  
Defaults to `#f`.

**repo-file-object logo** [repository-cgit-configuration parameter]  
URL which specifies the source of an image which will be used as a logo on this repo's pages.  
Defaults to `""`.

**repo-string logo-link** [repository-cgit-configuration parameter]  
URL loaded when clicking on the cgit logo image.  
Defaults to `""`.

**repo-file-object owner-filter** [repository-cgit-configuration parameter]  
Override the default `owner-filter`.  
Defaults to `""`.

**repo-string module-link** [repository-cgit-configuration parameter]  
Text which will be used as the formatstring for a hyperlink when a submodule is printed in a directory listing. The arguments for the formatstring are the path and SHA1 of the submodule commit.  
Defaults to `""`.

**module-link-path module-link-path** [repository-cgit-configuration parameter]  
Text which will be used as the formatstring for a hyperlink when a submodule with the specified subdirectory path is printed in a directory listing.  
Defaults to `'()'`.

**repo-string max-stats** [repository-cgit-configuration parameter]  
Override the default maximum statistics period.  
Defaults to `""`.



**repo-string name** [repository-cgit-configuration parameter]  
 The value to show as repository name.  
 Defaults to `""`.

**repo-string owner** [repository-cgit-configuration parameter]  
 A value used to identify the owner of the repository.  
 Defaults to `""`.

**repo-string path** [repository-cgit-configuration parameter]  
 An absolute path to the repository directory.  
 Defaults to `""`.

**repo-string readme** [repository-cgit-configuration parameter]  
 A path (relative to repo) which specifies a file to include verbatim as the  
 “About” page for this repo.  
 Defaults to `""`.

**repo-string section** [repository-cgit-configuration parameter]  
 The name of the current repository section - all repositories defined after this  
 option will inherit the current section name.  
 Defaults to `""`.

**repo-list extra-options** [repository-cgit-configuration parameter]  
 Extra options will be appended to cgitrc file.  
 Defaults to `'()'`.

**list extra-options** [cgit-configuration parameter]  
 Extra options will be appended to cgitrc file.  
 Defaults to `'()'`.

However, it could be that you just want to get a `cgitrc` up and running. In that case, you can pass an `opaque-cgit-configuration` as a record to `cgit-service-type`. As its name indicates, an opaque configuration does not have easy reflective capabilities.

Available `opaque-cgit-configuration` fields are:

**package cgit** [opaque-cgit-configuration parameter]  
 The cgit package.

**string string** [opaque-cgit-configuration parameter]  
 The contents of the `cgitrc`, as a string.

For example, if your `cgitrc` is just the empty string, you could instantiate a cgit service like this:

```
(service cgit-service-type
  (opaque-cgit-configuration
    (cgitrc "")))
```

## Gitolite Service

Gitolite (<https://gitolite.com/gitolite/>) is a tool for hosting Git repositories on a central server.

Gitolite can handle multiple repositories and users, and supports flexible configuration of the permissions for the users on the repositories.

The following example will configure Gitolite using the default `git` user, and the provided SSH public key.

```
(service gitolite-service-type
  (gitolite-configuration
    (admin-pubkey (plain-file
                  "yourname.pub"
                  "ssh-rsa AAAA... guix@example.com")))))
```

Gitolite is configured through a special admin repository which you can clone, for example, if you setup Gitolite on `example.com`, you would run the following command to clone the admin repository.

```
git clone git@example.com:gitolite-admin
```

When the Gitolite service is activated, the provided `admin-pubkey` will be inserted in to the `keydir` directory in the `gitolite-admin` repository. If this results in a change in the repository, it will be committed using the message “gitolite setup by GNU Guix”.

**gitolite-configuration** [Data Type]

Data type representing the configuration for `gitolite-service-type`.

**package** (default: `gitolite`)

Gitolite package to use. There are optional Gitolite dependencies that are not included in the default package, such as Redis and `git-annex`. These features can be made available by using the `make-gitolite` procedure in the `(gnu packages version-control)` module to produce a variant of Gitolite with the desired additional dependencies.

The following code returns a package in which the Redis and `git-annex` programs can be invoked by Gitolite’s scripts:

```
(use-modules (gnu packages databases)
             (gnu packages haskell-apps)
             (gnu packages version-control))
(make-gitolite (list redis git-annex))
```

**user** (default: `"git"`)

User to use for Gitolite. This will be user that you use when accessing Gitolite over SSH.

**group** (default: `"git"`)

Group to use for Gitolite.

**home-directory** (default: `"/var/lib/gitolite"`)

Directory in which to store the Gitolite configuration and repositories.

**rc-file** (default: `(gitolite-rc-file)`)

A “file-like” object (see Section 8.12 [G-Expressions], page 165), representing the configuration for Gitolite.

**git-config** (default: (gitolite-git-configuration))

A “file-like” object (see Section 8.12 [G-Expressions], page 165), representing the git configuration for Gitolite.

**admin-pubkey**

A “file-like” object (see Section 8.12 [G-Expressions], page 165) used to setup Gitolite. This will be inserted in to the **keydir** directory within the gitolite-admin repository.

To specify the SSH key as a string, use the **plain-file** function.

```
(plain-file "yourname.pub" "ssh-rsa AAAA... guix@example.com")■
```

The file has to have **.pub** extension.

**admin-name** (default: #f)

Usually the key name is derived by taking a base name of the **admin-pubkey**. The resulting file name is the name of the gitolite user. You can use this field to override the automatic name.

**gitolite-rc-file**

[Data Type]

Data type representing the Gitolite RC file.

**umask** (default: #o0077)

This controls the permissions Gitolite sets on the repositories and their contents.

A value like **#o0027** will give read access to the group used by Gitolite (by default: **git**). This is necessary when using Gitolite with software like **cg**it or **gitweb**.

**local-code** (default: "\$rc{GL\_ADMIN\_BASE}/local")

Allows you to add your own non-core programs, or even override the shipped ones with your own.

Please supply the FULL path to this variable. By default, directory called "local" in your gitolite clone is used, providing the benefits of versioning them as well as making changes to them without having to log on to the server.

**unsafe-pattern** (default: #f)

An optional Perl regular expression for catching unsafe configurations in the configuration file. See Gitolite’s documentation ([https://gitolite.com/gitolite/git-config.html#compensating-for-unsafe\\_patt](https://gitolite.com/gitolite/git-config.html#compensating-for-unsafe_patt)) for more information.

When the value is not **#f**, it should be a string containing a Perl regular expression, such as "[`~#\\$&()|;<>]", which is the default value used by gitolite. It rejects any special character in configuration that might be interpreted by a shell, which is useful when sharing the administration burden with other people that do not otherwise have shell access on the server.

**git-config-keys** (default: "")

Gitolite allows you to set git config values using the ‘**config**’ keyword. This setting allows control over the config keys to accept.

**log-extra** (default: **#f**)  
 Whether gitolite should log extra details.  
 For historical reasons, this field defaults to **#f**. The default value from gitolite however is 1 (written as **#t** in this configuration).

**host-name** (default: **#f**)  
 The host name of this gitolite instance. It is used by the mirroring feature.

**roles** (default: `'(("READERS" . 1) ("WRITERS" . ))`)  
 Set the role names allowed to be used by users running the perms command.

**enable** (default: `'("help" "desc" "info" "perms" "writable" "ssh-authkeys" "git-config" "daemon" "gitweb")`)  
 This setting controls the commands and features to enable within Gitolite.

**extra-content** (default: `"`)  
 Extra content to add verbatim into the `%RC` hash.

**gitolite-git-configuration** [Data Type]  
 Data type representing the git configuration file for gitolite.

**name** (default: `"GNU GNU"`)  
 User name used for commits (e.g. during setting up the admin repository).

**email** (default: `"guix@localhost"`)  
 Email used for commits (e.g. during setting up the admin repository).

**default-branch** (default: **#f**)  
 If true, set the `init.defaultBranch` option to that value. Common values are `"master"` and `"main"`.

**receive-fsck-objects** (default: **#f**)  
 If it is set to true, git-receive-pack will check all received objects.

**extra-content** (default: `"`)  
 Extra content to add verbatim into the git configuration file.

## Gitile Service

Gitile (<https://git.lepiller.eu/gitile>) is a Git forge for viewing public git repository contents from a web browser.

Gitile works best in collaboration with Gitolite, and will serve the public repositories from Gitolite by default. The service should listen only on a local port, and a webserver should be configured to serve static resources. The gitile service provides an easy way to extend the Nginx service for that purpose (see [NGINX], page 497).

The following example will configure Gitile to serve repositories from a custom location, with some default messages for the home page and the footers.

```
(service gitile-service-type
  (gitile-configuration
    (repositories "/srv/git")
```

```
(base-git-url "https://myweb.site/git")
(index-title "My git repositories")
(intro '((p "This is all my public work!"))))
(footer '((p "This is the end"))))
(nginx
  (nginx-server-configuration
    (ssl-certificate
      "/etc/certs/myweb.site/fullchain.pem")
    (ssl-certificate-key
      "/etc/certs/myweb.site/privkey.pem")
    (listen '("443 ssl http2" "[:]:443 ssl http2"))
    (locations
      (list
        ;; Allow for https anonymous fetch on /git/ urls.
        (git-http-nginx-location-configuration
          (git-http-configuration
            (uri-path "/git/")
            (git-root "/var/lib/gitolite/repositories"))))))))■
```

In addition to the configuration record, you should configure your git repositories to contain some optional information. First, your public repositories need to contain the `git-daemon-export-ok` magic file that allows Git to export the repository. Gitile uses the presence of this file to detect public repositories it should make accessible. To do so with Gitolite for instance, modify your `conf/gitolite.conf` to include this in the repositories you want to make public:

```
repo foo
  R = daemon
```

In addition, Gitile can read the repository configuration to display more information on the repository. Gitile uses the `gitweb` namespace for its configuration. As an example, you can use the following in your `conf/gitolite.conf`:

```
repo foo
  R = daemon
  desc = A long description, optionally with <i>HTML</i>, shown on the index page■
  config gitweb.name = The Foo Project
  config gitweb.synopsis = A short description, shown on the main page of the projec
```

Do not forget to commit and push these changes once you are satisfied. You may need to change your gitolite configuration to allow the previous configuration options to be set. One way to do that is to add the following service definition:

```
(service gitolite-service-type
  (gitolite-configuration
    (admin-pubkey (local-file "key.pub"))
    (rc-file
      (gitolite-rc-file
        (umask #o0027)
        ;; Allow to set any configuration key
        (git-config-keys ".*"))
```

```
;; Allow any text as a valid configuration value
(unsafe-patt "~$")))))
```

**gitile-configuration** [Data Type]

Data type representing the configuration for **gitile-service-type**.

**package** (default: *gitile*)

Gitile package to use.

**host** (default: "localhost")

The host on which gitile is listening.

**port** (default: 8080)

The port on which gitile is listening.

**database** (default: "/var/lib/gitile/gitile-db.sql")

The location of the database.

**repositories** (default: "/var/lib/gitolite/repositories")

The location of the repositories. Note that only public repositories will be shown by Gitile. To make a repository public, add an empty **git-daemon-export-ok** file at the root of that repository.

**git-owner-validation?** (default: #t)

Whether to check for the service user's ownership of repositories. By default, Git will refuse to process a repository owned by another user.

**base-git-url**

The base git url that will be used to show clone commands.

**index-title** (default: "Index")

The page title for the index page that lists all the available repositories.

**intro** (default: '()')

The intro content, as a list of sxml expressions. This is shown above the list of repositories, on the index page.

**footer** (default: '()')

The footer content, as a list of sxml expressions. This is shown on every page served by Gitile.

**nginx**

An nginx server block that will be extended and used as a reverse proxy by Gitile to serve its pages, and as a normal web server to serve its assets. You can use this block to add more custom URLs to your domain, such as a **/git/** URL for anonymous clones, or serving any other files you would like to serve.

### 11.10.34 Game Services

#### Joycond service

The joycond service allows the pairing of Nintendo joycon game controllers over Bluetooth. (see Section 11.10.9 [Desktop Services], page 377, for setting up Bluetooth.)

**joycond-configuration** [Data Type]

Data type representing the configuration of `joycond`.

**package** (default: `joycond`)  
The `joycond` package to use.

**joycond-service-type** [Variable]

Service type for the `joycond` service. It also extends the `udev-service-type` with the `joycond` package (provided via the `joycond-configuration` configuration), so that `joycond` controllers can be detected and used by an unprivileged user.

## Luanti service

Luanti (<https://www.luanti.org/en/>) is a voxel game engine that powers many games. This service is for hosting a Luanti server. The various options can be configured via the `luanti-configuration` record, documented below:

**luanti-configuration** [Data Type]

Available `luanti-configuration` fields are:

**luanti** (default: `luanti-server`) (type: file-like)  
The Luanti package to use.

**game** (default: `luanti-mineclonia`) (type: file-like)  
The Luanti game package to serve.

**game-configuration** (type: maybe-file-like)  
A configuration file to use for the selected Luanti game, which corresponds to the `minetest.conf` file.

**mods** (type: maybe-list-of-file-likes)  
A list of Luanti mod packages to use. Note that using mods is complicated by the requirements of Luanti to 1) manually enable the mod and any of its dependent mods in the `world.rt` file of the world used and 2) to register the mod names and those of its dependents via a `'secure.trusted_mods'` `game-configuration` directive. Consult the example below for more precise directions.

**log-file** (default: `"/var/log/luanti.log"`) (type: maybe-string)  
The log file to log to. To disable logging, set this to `%unset-value`.

**verbose?** (default: `#f`) (type: boolean)  
Print more detailed information.

**port** (default: `30000`) (type: port)  
The UDP port the server should listen to.

**world** (type: maybe-string)  
An existing Luanti world directory to serve. If omitted, a new world is created under the `/var/lib/luanti/.minetest/worlds/world` directory. If an absolute file name is provided, it is used directly. Otherwise, it is expected to be a directory under `/var/lib/luanti/.minetest/worlds/`.

Here's the simplest example of a Luanti server, which in its default configuration serves the `luanti-mineclonia` game.

```
(service luanti-service-type)
```

Here's a slightly more elaborate one, which adds the `luanti-whitelist` mod. Embedded are comments explaining extra needed steps when using mods. Failing to do these steps will cause the service to fail to start.

```
(service luanti-service-type
  (luanti-configuration
    (game luanti-mineclonia)
    (game-configuration
      (plain-file
        "minetest.conf"
        ;; lib_chatcmdbuilder is a dependency of the whitelist mod
        "secure.trusted_mods = whitelist,lib_chatcmdbuilder\n"))
    ;; The
    ;; '/var/lib/luanti/.minetest/worlds/world/world.mt'
    ;; file needs to be hand-edited to add:
    ;; load_mod_whitelist = true
    ;; load_mod_lib_chatcmdbuilder = true
    (mods (list luanti-whitelist))))
```

## The Battle for Wesnoth Service

The Battle for Wesnoth (<https://wesnoth.org>) is a fantasy, turn based tactical strategy game, with several single player campaigns, and multiplayer games (both networked and local).

**wesnothd-service-type** [Variable]  
 Service type for the wesnothd service. Its value must be a `wesnothd-configuration` object. To run wesnothd in the default configuration, instantiate it as:

```
(service wesnothd-service-type)
```

**wesnothd-configuration** [Data Type]  
 Data type representing the configuration of wesnothd.

**package** (default: `wesnoth-server`)  
 The wesnoth server package to use.

**port** (default: 15000)  
 The port to bind the server to.

### 11.10.35 PAM Mount Service

The (`gnu services pam-mount`) module provides a service allowing users to mount volumes when they log in. It should be able to mount any volume format supported by the system.

**pam-mount-service-type** [Variable]  
 Service type for PAM Mount support.



**pam-mount-configuration**

[Data Type]

Data type representing the configuration of PAM Mount.

It takes the following parameters:

**rules** The configuration rules that will be used to generate `/etc/security/pam_mount.conf.xml`.

The configuration rules are SXML elements (see Section “SXML” in *GNU Guile Reference Manual*), and the default ones don’t mount anything for anyone at login:

```

`((debug (@ (enable "0"))))
  (mntoptions (@ (allow ,(string-join
                        '("nosuid" "nodev" "loop"
                          "encryption" "fsck" "nonempty"
                          "allow_root" "allow_other")
                        ","))))
  (mntoptions (@ (require "nosuid,nodev"))))
  (logout (@ (wait "0")
             (hup "0")
             (term "no")
             (kill "no"))))
  (mkmountpoint (@ (enable "1")
                  (remove "true"))))

```

Some `volume` elements must be added to automatically mount volumes at login. Here’s an example allowing the user `alice` to mount her encrypted `HOME` directory and allowing the user `bob` to mount the partition where he stores his data:

```

(define pam-mount-rules
`((debug (@ (enable "0"))))
  (volume (@ (user "alice")
             (fstype "crypt")
             (path "/dev/sda2")
             (mountpoint "/home/alice"))))
  (volume (@ (user "bob")
             (fstype "auto")
             (path "/dev/sdb3")
             (mountpoint "/home/bob/data")
             (options "defaults,autodefrag,compress"))))
  (mntoptions (@ (allow ,(string-join
                        '("nosuid" "nodev" "loop"
                          "encryption" "fsck" "nonempty"
                          "allow_root" "allow_other")
                        ","))))
  (mntoptions (@ (require "nosuid,nodev"))))
  (logout (@ (wait "0")
             (hup "0")
             (term "no"))))

```

```

                                (kill "no")))
    (mkmountpoint (@ (enable "1")
                     (remove "true")))))

(service pam-mount-service-type
  (pam-mount-configuration
    (rules pam-mount-rules)))

```

The complete list of possible options can be found in the man page for `pam_mount.conf` ([http://pam-mount.sourceforge.net/pam\\_mount.conf.5.html](http://pam-mount.sourceforge.net/pam_mount.conf.5.html)).

## PAM Mount Volume Service

PAM mount volumes are automatically mounted at login by the PAM login service according to a set of per-volume rules. Because they are mounted by PAM the password entered during login may be used directly to mount authenticated volumes, such as `cifs`, using the same credentials.

These volumes will be added in addition to any volumes directly specified in `pam-mount-rules`.

Here is an example of a rule which will mount a remote CIFS share from `//remote-server/share` into a sub-directory of `/shares` named after the user logging in:

```

(simple-service 'pam-mount-remote-share pam-mount-volume-service-type
  (list (pam-mount-volume
    (secondary-group "users")
    (file-system-type "cifs")
    (server "remote-server")
    (file-name "share")
    (mount-point "/shares/%(USER)")
    (options "nosuid,nodev,seal,cifsacl")))))

```

**pam-mount-volume-service-type** [Data Type]

Configuration for a single volume to be mounted. Any fields not specified will be omitted from the run-time PAM configuration. See the man page ([http://pam-mount.sourceforge.net/pam\\_mount.conf.5.html](http://pam-mount.sourceforge.net/pam_mount.conf.5.html)) for the default values when unspecified.

**user-name** (type: maybe-string)

Mount the volume for the given user.

**user-id** (type: maybe-integer-or-range)

Mount the volume for the user with this ID. This field may also be specified as a pair of (`start` . `end`) indicating a range of user IDs for whom to mount the volume.

**primary-group** (type: maybe-string)

Mount the volume for users with this primary group name.

**group-id** (type: maybe-integer-or-range)

Mount the volume for the users with this primary group ID. This field may also be specified as a cons cell of (`start` . `end`) indicating a range of group ids for whom to mount the volume.

- secondary-group** (type: maybe-string)  
Mount the volume for users who are members of this group as either a primary or secondary group.
- file-system-type** (type: maybe-string)  
The file system type for the volume being mounted (e.g., `cifs`)
- no-mount-as-root?** (type: maybe-boolean)  
Whether or not to mount the volume with root privileges. This is normally disabled, but may be enabled for mounts of type `fuse`, or other user-level mounts.
- server** (type: maybe-string)  
The name of the remote server to mount the volume from, when necessary.
- file-name** (type: maybe-string)  
The location of the volume, either local or remote, depending on the **file-system-type**.
- mount-point** (type: maybe-string)  
Where to mount the volume in the local file-system. This may be set to `~` to indicate the home directory of the user logging in. If this field is omitted then `/etc/fstab` is consulted for the mount destination.
- options** (type: maybe-string)  
The options to be passed as-is to the underlying mount program.
- ssh?** (type: maybe-boolean)  
Enable this option to pass the login password to SSH for use with mounts involving SSH (e.g., `sshfs`).
- cipher** (type: maybe-string)  
Cryptsetup cipher name for the volume. To be used with the `crypt` **file-system-type**.
- file-system-key-cipher** (type: maybe-string)  
Cipher name used by the target volume.
- file-system-key-hash** (type: maybe-string)  
SSL hash name used by the target volume.
- file-system-key-file-name** (type: maybe-string)  
File name of the file system key for the target volume.

### 11.10.36 Guix Services

#### Shared Cache Service

Running `guix pull`, `guix time-machine`, `guix system reconfigure`, and `guix home reconfigure` pulls the complete Git version control history of your channels in `$HOME/.cache/guix/checkouts`. If you have multiple user accounts (typically `root` and a main user account for your personal computer), or if multiple users are using the system (typically on a shared server), then this can quickly waste a lot of space.

The `shared-cache` service lets you prevent this waste by either sharing or exposing a specific user's cache to the other users.

In the context of this service, *sharing* and *exposing* have the same meaning as with `guix shell`: making the shared cache directory read-only if you use mode `'expose` (the default), and writable if you use mode `'share`, (see Section 7.1 [Invoking guix shell], page 78).

Whatever mode you choose, from the point of view of the user, the cache appears writable. With `'expose`, any write will be stored in a user-owned overlay directory next to mountpoint, whereas with `'share`, writes will be propagated back to the sharing user's cache.

**Warning:** Use the `'share` option *if and only if* users on the machine trust each other. This is a hard requirement since, with this option, any user could “poison” the shared cache to indirectly gain access to other user's data and privileges.

**Warning:** Using `'share` is meant for a single-seat computer where the single user has unprivileged user accounts—e.g. one for work and one for gaming. Even then, channels from one will leak into the other's cache. Use with caution.

**Warning:** `'share` will not work for multiple concurrent invocation of `guix`, e.g. if two user accounts run `guix time machine` at the same time, the results will be fun but not useful.

Here is an example configuration that exposes `root`'s `guix` checkout cache to `alice` and `bob`:

```
(operating-system
  ;; ...
  (services
    (cons*
      (service shared-cache-service-type
        (shared-cache-configuration
          (users (list (user-cache (user "alice"))
                        (user-cache (user "bob"))))))
      %desktop-services)))
```

The service can also be extended by providing a list of `user-cache`, allowing for easier splitting of configuration.

```
(operating-system
  ;; ...
  (services
    (cons*
      (simple-service 'bob-shared-guix-cache shared-cache-service-type
        (list (user-cache (user "bob"))))
      %desktop-services)))
```

`shared-cache-service-type` [Variable]

Service type for sharing a subdirectory of `~/.cache/` between multiple users.

This service allows different system users to share the contents of the user cache, avoiding duplication and saving disk space. It can be configured in either “share” mode,

where all users write to the same cache, or “expose” mode, where users get a read-only view with a writable overlay. Its value must be a **shared-cache-configuration** object.

**shared-cache-configuration** [Data Type]

Data type representing the configuration of the shared cache.

**mode** (default: **'expose'**)

Either **'share'** (read-write, dangerous) or **'expose'** (read-only with per-user overlay, the default). In share mode, all writes go to the shared directory. In expose mode, only reads use the shared directory; writes go to a user-owned overlay.

**shared-directory** (default: **"/root/.cache"**)

The source cache directory that is shared between users. Each user will see some or all of its subdirectories mounted into their own **~/.cache**.

**users** (default: **'()'**)

The list of users that will benefit from the shared directory. This must be a list of **user-info** records.

**user-info** [Data Type]

Data type representing the configuration for a single user benefitting from a shared cache.

**user** The username of the user.

**location** (default: **"/home/user/.cache"**)

The cache directory of the user.

**group** (default: **"users"**)

The primary group of the user.

**directories** (default: **'("guix/checkouts")'**)

List of subdirectories of the **shared-cache-configuration**'s **shared-directory** (whose default is **/root/.cache**) to mount under **location** (whose default is **~/.cache/**) for the user.

## Build Farm Front-End (BFFE)

The Build Farm Front-End (<https://git.cbaines.net/guix/bffe/>) assists with building Guix packages in bulk. It's responsible for submitting builds and displaying the status of the build farm.

**bffe-service-type** [Variable]

Service type for the Build Farm Front-End. Its value must be a **bffe-configuration** object.

**bffe-configuration** [Data Type]

Data type representing the configuration of the Build Farm Front-End.

**package** (default: **bffe**)

The Build Farm Front-End package to use.

```

user (default: "bffe")
    The system user to run the service as.

group (default: "bffe")
    The system group to run the service as.

arguments
    A list of arguments to the Build Farm Front-End. These are passed to
    the run-bffe-service procedure when starting the service.
    For example, the following value directs the Build Farm Front-End to
    submit builds for derivations available from data.guix.gnu.org to the
    Build Coordinator instance assumed to be running on the same machine.
    (list
      #:build
      (list
        (build-from-guix-data-service
          (data-service-url "https://data.guix.gnu.org")
          (build-coordinator-url "http://127.0.0.1:8746")
          (branches '("master"))
          (systems '("x86_64-linux" "i686-linux"))
          (systems-and-targets
            (map (lambda (target)
                  (cons "x86_64-linux" target))
                  '("aarch64-linux-gnu"
                    "i586-pc-gnu")))))
        #:web-server-args
        '(:event-source "https://example.com"
          #:controller-args
          (:title "example.com build farm")))
      extra-environment-variables (default: '())
    Extra environment variables to set via the shepherd service.

```

## Guix Build Coordinator

The Guix Build Coordinator (<https://git.cbaines.net/guix/build-coordinator/>) aids in distributing derivation builds among machines running an *agent*. The build daemon is still used to build the derivations, but the Guix Build Coordinator manages allocating builds and working with the results.

The Guix Build Coordinator consists of one *coordinator*, and one or more connected *agent* processes. The coordinator process handles clients submitting builds, and allocating builds to agents. The agent processes talk to a build daemon to actually perform the builds, then send the results back to the coordinator.

There is a script to run the coordinator component of the Guix Build Coordinator, but the Guix service uses a custom Guile script instead, to provide better integration with G-expressions used in the configuration.

```

guix-build-coordinator-service-type [Variable]
    Service type for the Guix Build Coordinator. Its value must be a guix-build-coordinator-configuration object.

```

**guix-build-coordinator-configuration** [Data Type]

Data type representing the configuration of the Guix Build Coordinator.

**package** (default: `guix-build-coordinator`)

The Guix Build Coordinator package to use.

**user** (default: `"guix-build-coordinator"`)

The system user to run the service as.

**group** (default: `"guix-build-coordinator"`)

The system group to run the service as.

**database-uri-string** (default:

`"sqlite:///var/lib/guix-build-coordinator/guix_build_coordinator.db"`)

The URI to use for the database.

**agent-communication-uri** (default: `"http://0.0.0.0:8745"`)

The URI describing how to listen to requests from agent processes.

**client-communication-uri** (default: `"http://127.0.0.1:8746"`)

The URI describing how to listen to requests from clients. The client API allows submitting builds and currently isn't authenticated, so take care when configuring this value.

**allocation-strategy** (default: `#~basic-build-allocation-strategy`)

A G-expression for the allocation strategy to be used. This is a procedure that takes the datastore as an argument and populates the allocation plan in the database.

**hooks** (default: `'()`)

An association list of hooks. These provide a way to execute arbitrary code upon certain events, like a build result being processed.

**parallel-hooks** (default: `'()`)

Hooks can be configured to run in parallel. This parameter is an association list of hooks to do in parallel, where the key is the symbol for the hook and the value is the number of threads to run.

**listen-repl** (default: `#f`)

Port or filename to spawn a REPL listening on, pass `#t` to listen on the default port or `#f` to disable.

**guile** (default: `guile-3.0-latest`)

The Guile package with which to run the Guix Build Coordinator.

**extra-environment-variables** (default: `'()`)

Extra environment variables to set via the shepherd service.

**guix-build-coordinator-agent-service-type** [Variable]

Service type for a Guix Build Coordinator agent. Its value must be a `guix-build-coordinator-agent-configuration` object.

**guix-build-coordinator-agent-configuration** [Data Type]

Data type representing the configuration a Guix Build Coordinator agent.

**package** (default: `guix-build-coordinator/agent-only`)

The Guix Build Coordinator package to use.

**user** (default: "guix-build-coordinator-agent")  
 The system user to run the service as.

**coordinator** (default: "http://localhost:8745")  
 The URI to use when connecting to the coordinator.

**authentication**  
 Record describing how this agent should authenticate with the coordinator. Possible record types are described below.

**systems** (default: #f)  
 The systems for which this agent should fetch builds. The agent process will use the current system it's running on as the default.

**max-parallel-builds** (default: #f)  
 The number of builds to perform in parallel.

**max-parallel-uploads** (default: #f)  
 The number of uploads to perform in parallel.

**max-allocated-builds** (default: #f)  
 The maximum number of builds this agent can be allocated.

**max-1min-load-average** (default: #f)  
 Load average value to look at when considering starting new builds, if the 1 minute load average exceeds this value, the agent will wait before starting new builds.  
 This will be unspecified if the value is #f, and the agent will use the number of cores reported by the system as the max 1 minute load average.

**derivation-substitute-urls** (default: #f)  
 URLs from which to attempt to fetch substitutes for derivations, if the derivations aren't already available.

**non-derivation-substitute-urls** (default: #f)  
 URLs from which to attempt to fetch substitutes for build inputs, if the input store items aren't already available.

**extra-options** (default: '())  
 Extra command line options for **guix-build-coordinator-agent**.

**guix-build-coordinator-agent-password-auth** [Data Type]  
 Data type representing an agent authenticating with a coordinator via a UUID and password.

**uuid** The UUID of the agent. This should be generated by the coordinator process, stored in the coordinator database, and used by the intended agent.

**password** The password to use when connecting to the coordinator.

**guix-build-coordinator-agent-password-file-auth** [Data Type]  
 Data type representing an agent authenticating with a coordinator via a UUID and password read from a file.



**uuid**            The UUID of the agent. This should be generated by the coordinator process, stored in the coordinator database, and used by the intended agent.

**password-file**  
A file containing the password to use when connecting to the coordinator.

**guix-build-coordinator-agent-dynamic-auth** [Data Type]  
Data type representing an agent authenticating with a coordinator via a dynamic auth token and agent name.

**agent-name**  
Name of an agent, this is used to match up to an existing entry in the database if there is one. When no existing entry is found, a new entry is automatically added.

**token**            Dynamic auth token, this is created and stored in the coordinator database, and is used by the agent to authenticate.

**guix-build-coordinator-agent-dynamic-auth-with-file** [Data Type]  
Data type representing an agent authenticating with a coordinator via a dynamic auth token read from a file and agent name.

**agent-name**  
Name of an agent, this is used to match up to an existing entry in the database if there is one. When no existing entry is found, a new entry is automatically added.

**token-file**  
File containing the dynamic auth token, this is created and stored in the coordinator database, and is used by the agent to authenticate.

## Guix Data Service

The Guix Data Service (<http://data.guix.gnu.org>) processes, stores and provides data about GNU Guix. This includes information about packages, derivations and lint warnings.

The data is stored in a PostgreSQL database, and available through a web interface.

**guix-data-service-type** [Variable]  
Service type for the Guix Data Service. Its value must be a **guix-data-service-configuration** object. The service optionally extends the getmail service, as the guix-commits mailing list is used to find out about changes in the Guix git repository.

**guix-data-service-configuration** [Data Type]  
Data type representing the configuration of the Guix Data Service.

**package** (default: **guix-data-service**)  
The Guix Data Service package to use.

**user** (default: **"guix-data-service"**)  
The system user to run the service as.

**group** (default: **"guix-data-service"**)  
The system group to run the service as.

**port** (default: 8765)  
The port to bind the web service to.

**host** (default: "127.0.0.1")  
The host to bind the web service to.

**getmail-idle-mailboxes** (default: #f)  
If set, this is the list of mailboxes that the getmail service will be configured to listen to.

**commits-getmail-retriever-configuration** (default: #f)  
If set, this is the `getmail-retriever-configuration` object with which to configure getmail to fetch mail from the guix-commits mailing list.

**extra-options** (default: '())  
Extra command line options for `guix-data-service`.

**extra-process-jobs-options** (default: '())  
Extra command line options for `guix-data-service-process-jobs`.

**git-repositories** (default: #f)  
List of git-repository information to insert into the database.

**build-servers** (default: #f)  
List of build-server information to insert into the database.

## Guix Home Service

The Guix Home service is a way to let Guix System deploy the home environment of one or more users (see Chapter 13 [Home Configuration], page 732, for more on Guix Home). That way, the system configuration embeds declarations of the home environment of those users and can be used to deploy everything consistently at once, saving users the need to run `guix home reconfigure` independently.

`guix home` can operate on the home environment configured for current user via this service, but its effect won't last beyond the next system activation. This can be used for quick testing and may avoid the need of keeping a separate home configuration.

**guix-home-service-type** [Variable]

Service type for the Guix Home service. Its value must be a list of lists containing user and home environment pairs. The key of each pair is a string representing the user to deploy the configuration under and the value is a home-environment configuration.

```
(use-modules (gnu services guix))
```

```
(define my-home
  (home-environment
    ...))
```

```
(operating-system
  (services (append (list (service guix-home-service-type
                                   `(("alice" ,my-home))))
                    %base-services)))
```

This service can be extended by other services to add additional home environments, as in this example:

```
(simple-service 'my-extra-home guix-home-service-type
  `(("bob" ,my-extra-home)))
```

## Nar Herder

The Nar Herder (<https://git.cbaines.net/guix/nar-herder/about/>) is a utility for managing a collection of nars.

**nar-herder-type** [Variable]

Service type for the Guix Data Service. Its value must be a **nar-herder-configuration** object. The service optionally extends the getmail service, as the guix-commits mailing list is used to find out about changes in the Guix git repository.

**nar-herder-configuration** [Data Type]

Data type representing the configuration of the Guix Data Service.

**package** (default: **nar-herder**)

The Nar Herder package to use.

**user** (default: **"nar-herder"**)

The system user to run the service as.

**group** (default: **"nar-herder"**)

The system group to run the service as.

**port** (default: **8734**)

The port to bind the server to.

**host** (default: **"127.0.0.1"**)

The host to bind the server to.

**control-port** (default: **8734**)

The port to bind the control server to.

**control-host** (default: **"127.0.0.1"**)

The host to bind the control server to.

**mirror** (default: **#f**)

Optional URL of the other Nar Herder instance which should be mirrored. This means that this Nar Herder instance will download it's database, and keep it up to date.

**database** (default: **"/var/lib/nar-herder/nar\_herder.db"**)

Location for the database. If this Nar Herder instance is mirroring another, the database will be downloaded if it doesn't exist. If this Nar Herder instance isn't mirroring another, an empty database will be created.

**database-dump** (default: **"/var/lib/nar-herder/nar\_herder\_dump.db"**)

Location of the database dump. This is created and regularly updated by taking a copy of the database. This is the version of the database that is available to download.

- storage** (default: `#f`)  
Optional location in which to store nars.
- storage-limit** (default: `"none"`)  
Limit in bytes for the nars stored in the storage location. This can also be set to “none” so that there is no limit.  
When the storage location exceeds this size, nars are removed according to the nar removal criteria.
- storage-nar-removal-criteria** (default: `'()`)  
Criteria used to remove nars from the storage location. These are used in conjunction with the storage limit.  
When the storage location exceeds the storage limit size, nars will be checked against the nar removal criteria and if any of the criteria match, they will be removed. This will continue until the storage location is below the storage limit size.  
Each criteria is specified by a string, then an equals sign, then another string. Currently, only one criteria is supported, checking if a nar is stored on another Nar Herder instance.
- ttl** (default: `#f`)  
Produce **Cache-Control** HTTP headers that advertise a time-to-live (TTL) of *ttl*. *ttl* must denote a duration: `5d` means 5 days, `1m` means 1 month, and so on.  
This allows the user’s Guix to keep substitute information in cache for *ttl*.
- new-ttl** (default: `#f`)  
If specified, this will override the **ttl** setting when used for the **Cache-Control** headers, but this value will be used when scheduling the removal of nars.  
Use this setting when the TTL is being reduced to avoid removing nars while clients still have cached narinfos.
- negative-ttl** (default: `#f`)  
Similarly produce **Cache-Control** HTTP headers to advertise the time-to-live (TTL) of *negative* lookups—missing store items, for which the HTTP 404 code is returned. By default, no negative TTL is advertised.
- log-level** (default: `'DEBUG`)  
Log level to use, specify a log level like `'INFO` to stop logging individual requests.
- cached-compressions** (default: `'()`)  
Activate generating cached nars with different compression details from the stored nars. This is a list of nar-herder-cached-compression-configuration records.
- min-uses** (default: `3`)  
When cached-compressions are enabled, generate cached nars when at least this number of requests are made for a nar.

**workers** (default: 2)  
 Number of cached nars to generate at a time.

**nar-source** (default: #f)  
 Location to fetch nars from when computing cached compressions. By default, the storage location will be used.

**extra-environment-variables** (default: '() )  
 Extra environment variables to set via the shepherd service.

**nar-herder-cached-compression-configuration** [Data Type]

Data type representing the cached compression configuration.

**type** Type of compression to use, e.g. 'zstd.

**workers** (default: #f)  
 Level of the compression to use.

**directory** (default: #f)  
 Location to store the cached nars. If unspecified, they will be stored in /var/cache/nar-herder/nar/TYPE.

**directory-max-size** (default: #f)  
 Maximum size in bytes of the directory.

**unused-removal-duration** (default: #f)  
 If a cached nar isn't used for **unused-removal-duration**, it will be scheduled for removal.  
*unused-removal-duration* must denote a duration: 5d means 5 days, 1m means 1 month, and so on.

**ttl** (default: #f)  
 If specified this overrides the **ttl** used for narinfos when this cached compression is available.

**new-ttl** (default: #f)  
 As with the **new-ttl** option for **nar-herder-configuration**, this value will override the **ttl** when used for narinfo requests.

### 11.10.37 Linux Services

#### VFS Mapping Service

The VFS (Virtual File System) Mapping service allows you to bind a file name under a different name in the global namespace. For example, this is used by the Shared Cache Service (see Section 11.10.36 [Guix Services], page 631) to bind the 'root' user's cache over another user's cache. Another use case may be to expose game data from the store over your home directory, and apply modifications there.

The service has three *policies* to choose from with respect to access rights on the bound over path:

**'bind** Use a bind mount; the ownership and file permissions of the destination are the same as the source.

- 'translate' Use `bindfs` to make the destination appear as owned by the target user and group. Writes will be propagated back to the source as if made by the original owner.
- 'overlay' Use an overlay to make the destination appear as owned by the target user and group. Writes will not appear in the source, but will be stored in a target-user-owned `‘-overlay’` suffixed directory near the destination.

Here is an example configuration that exposes the default data directory of `tuxpaint` under the home directory of `alice`. Any modification made by Alice to the brushes or stamps will not propagate back to the store (which is read-only anyway).

```
(operating-system
;; ...
(services
 (cons*
  (service vfs-mapping-service-type
    (vfs-mapping-configuration
      (bindings (list (vfs-mapping
                       (source      #~(string-append #$tuxpaint
   "/share/tuxpaint"))
                       (destination "/home/alice/.tuxpaint")
                       (user        "alice")
                       (policy      'overlay)
                       (name        "alice-tuxpaint-overlay"))))))
    %desktop-services)))
```

The service can also be extended by providing a list of `vfs-mapping`, allowing for easier splitting of configuration.

```
(operating-system
;; ...
(services
 (cons*
  (simple-service 'bob-too-wants-to-hack-on-tuxpaint
    vfs-mapping-service-type
    (list (vfs-mapping
          (source      #~(string-append #$tuxpaint "/share/tuxpaint")
          (destination "/home/bob/.tuxpaint")
          (user        "bob")
          (policy      'overlay)
          (name        "bob-tuxpaint-overlay")))))
    %desktop-services)))
```

The `source` file name must exist for the service to start. The `destination` directory will be created if it does not exist.

**vfs-mapping-service-type** [Variable]

Service type for binding a directory in multiple places on the file system.

The access rights are either the same in source and destination (`'bind`), or writes are translated back to the source as if made by the source's owner (`'translate`), or kept

in an overlay directory near the destination (`'overlay'`). The service's value must be a `vfs-mapping-configuration` object.

**vfs-mapping-configuration** [Data Type]

Data type representing the configuration of the `vfs-mapping` service.

**bindfs** (default: `#~(string-append #$bindfs "/bin/bindfs")`)  
The `bindfs` command to use for mounting.

**fusermount** (default: `#~(string-append #$fuse-2 "/bin/fusermount")`)  
The `fusermount` command to use.

**umount** (default: `#~(string-append #$util-linux+udev "/bin/umount")`)  
The `umount` command to use.

**bindings** (default: `'()`)  
A list of `vfs-mapping` records.

**vfs-mapping** [Data Type]

Data type representing the configuration for a single shared directory.

**source** The source file name to be shared.

**destination**  
The destination at which the contents of `source` will be exposed.

**policy** (default: `'translate'`)  
Either `'bind'` (same ownership and access rights for `source` and `destination`), `'translate'` (read-write, with writes translated as if made by `source`'s owner), or `'overlay'` (read-only with `user`-owned writable overlay).

**user** (default: `#f`)  
The user that will own the `destination` directory, and appear to own its content. It must be kept at its default `#f` value when using the `'bind'` policy.

**group** (default: `"users"`)  
The group that will own the `destination` directory, and appear to own its content.

**name** (default: `"<src>-[<policy>]-><dst>"`)  
The name of the shepherd service to mount and unmount the binding.

**requirement** (default: `'(file-systems user-homes)'`)  
The list of services that Shepherd ought to provision before trying to mount.

## Early OOM Service

Early OOM (<https://github.com/rfjakob/earlyoom>), also known as Earlyoom, is a minimalist out of memory (OOM) daemon that runs in user space and provides a more responsive and configurable alternative to the in-kernel OOM killer. It is useful to prevent the system from becoming unresponsive when it runs out of memory.

**earlyoom-service-type** [Variable]

The service type for running **earlyoom**, the Early OOM daemon. Its value must be a **earlyoom-configuration** object, described below. The service can be instantiated in its default configuration with:

```
(service earlyoom-service-type)
```

**earlyoom-configuration** [Data Type]

This is the configuration record for the **earlyoom-service-type**.

**earlyoom** (default: *earlyoom*)

The Earlyoom package to use.

**minimum-available-memory** (default: 10)

The threshold for the minimum *available* memory, in percentages.

**minimum-free-swap** (default: 10)

The threshold for the minimum free swap memory, in percentages.

**prefer-regex** (default: #f)

A regular expression (as a string) to match the names of the processes that should be preferably killed.

**avoid-regex** (default: #f)

A regular expression (as a string) to match the names of the processes that should *not* be killed.

**memory-report-interval** (default: 0)

The interval in seconds at which a memory report is printed. It is disabled by default.

**ignore-positive-oom-score-adj?** (default: #f)

A boolean indicating whether the positive adjustments set in `/proc/*/oom_score_adj` should be ignored.

**show-debug-messages?** (default: #f)

A boolean indicating whether debug messages should be printed. The logs are saved at `/var/log/earlyoom.log`.

**send-notification-command** (default: #f)

This can be used to provide a custom command used for sending notifications.

## fstrim Service

The command **fstrim** can be used to discard (or *trim*) unused blocks on a mounted file system.

**Warning:** Running **fstrim** frequently, or even using `mount -o discard`, might negatively affect the lifetime of poor-quality SSD devices. For most desktop and server systems a sufficient trimming frequency is once a week. Note that not all devices support a queued trim, so each trim command incurs a performance penalty on whatever else might be trying to use the disk at the time.



**fstrim-service-type** [Variable]

Type for a service that periodically runs **fstrim**, whose value must be an **<fstrim-configuration>** object. The service can be instantiated in its default configuration with:

```
(service fstrim-service-type)
```

**fstrim-configuration** [Data Type]

Available **fstrim-configuration** fields are:

**package** (default: **util-linux**) (type: file-like)

The package providing the **fstrim** command.

**schedule** (default: **"0 0 \* \* 0"**) (type: Shepherd calendar event)

Schedule for launching **fstrim**, expressed as a string in traditional cron syntax or as a gexp evaluating to a Shepherd calendar event (see Section “Timers” in *The GNU Shepherd Manual*). By default this is set to run weekly on Sunday at 00:00.

**listed-in** (default: **'("/etc/fstab" "/proc/self/mountinfo")'**) (type: maybe-list-of-strings)

List of files in **fstab** or kernel **mountinfo** format. All missing or empty files are silently ignored. The evaluation of the list *stops* after the first non-empty file. File systems with **X-fstrim.trim** mount option in **fstab** are skipped.

**verbose?** (default: **#t**) (type: boolean)

Verbose execution.

**quiet-unsupported?** (default: **#t**) (type: boolean)

Suppress error messages if trim operation (**ioctl**) is unsupported.

**extra-arguments** (type: maybe-list-of-strings)

Extra options to append to **fstrim** (run **‘man fstrim’** for more information).

## Kernel Module Loader Service

The kernel module loader service allows one to load loadable kernel modules at boot. This is especially useful for modules that don’t autoload and need to be manually loaded, as is the case with **ddcci**.

**kernel-module-loader-service-type** [Variable]

The service type for loading loadable kernel modules at boot with **modprobe**. Its value must be a list of strings representing module names. For example loading the drivers provided by **ddcci-driver-linux**, in debugging mode by passing some module parameters, can be done as follow:

```
(use-modules (gnu) (gnu services))
(use-package-modules linux)
(use-service-modules linux)
```

```
(define ddcci-config
```

```

(plain-file "ddcci.conf"
  "options ddcci dyndbg delay=120"))

(operating-system
  ...
  (services (cons* (service kernel-module-loader-service-type
    '("ddcci" "ddcci_backlight"))
    (simple-service 'ddcci-config etc-service-type
      (list `("modprobe.d/ddcci.conf"
        ,ddcci-config)))
    %base-services))
  (kernel-loadable-modules (list ddcci-driver-linux)))

```

## Cachefilesd Service

The Cachefilesd service starts a daemon that caches network file system data locally. It is especially useful for NFS and AFS shares, where it reduces latencies for repeated access when reading files.

The daemon can be configured as follows:

```

(service cachefilesd-service-type
  (cachefilesd-configuration
    (cache-directory "/var/cache/fscache")))

```

**cachefilesd-service-type** [Variable]

The service type for starting `cachefilesd`. The value for this service type is a `cachefilesd-configuration`, whose only required field is *cache-directory*.

**cachefilesd-configuration** [Data Type]

Available `cachefilesd-configuration` fields are:

**cachefilesd** (default: `cachefilesd`) (type: file-like)  
The `cachefilesd` package to use.

**debug-output?** (default: `#f`) (type: boolean)  
Print debugging output to stderr.

**use-syslog?** (default: `#t`) (type: boolean)  
Log to syslog facility instead of stdout.

**scan?** (default: `#t`) (type: boolean)  
Scan for cacheable objects.

**cache-directory** (type: maybe-string)  
Location of the cache directory.

**cache-name** (default: `"CacheFiles"`) (type: maybe-string)  
Name of cache (keep unique).

**security-context** (type: maybe-string)  
SELinux security context.

**pause-culling-for-block-percentage** (default: 7) (type: maybe-non-negative-integer)  
Pause culling when available blocks exceed this percentage.

**pause-culling-for-file-percentage** (default: 7) (type: maybe-non-negative-integer)  
 Pause culling when available files exceed this percentage.

**resume-culling-for-block-percentage** (default: 5) (type: maybe-non-negative-integer)  
 Start culling when available blocks drop below this percentage.

**resume-culling-for-file-percentage** (default: 5) (type: maybe-non-negative-integer)  
 Start culling when available files drop below this percentage.

**pause-caching-for-block-percentage** (default: 1) (type: maybe-non-negative-integer)  
 Pause further allocations when available blocks drop below this percentage.

**pause-caching-for-file-percentage** (default: 1) (type: maybe-non-negative-integer)  
 Pause further allocations when available files drop below this percentage.

**log2-table-size** (default: 12) (type: maybe-non-negative-integer)  
 Size of tables holding cullable objects in logarithm of base 2.

**cull?** (default: #t) (type: boolean)  
 Create free space by culling (consumes system load).

**trace-function-entry-in-kernel-module?** (default: #f) (type: boolean)  
 Trace function entry in the kernel module (for debugging).

**trace-function-exit-in-kernel-module?** (default: #f) (type: boolean)  
 Trace function exit in the kernel module (for debugging).

**trace-internal-checkpoints-in-kernel-module?** (default: #f) (type: boolean)  
 Trace internal checkpoints in the kernel module (for debugging).

## Rasdaemon Service

The Rasdaemon service provides a daemon which monitors platform RAS (Reliability, Availability, and Serviceability) reports from Linux kernel trace events, logging them to syslogd.

Reliability, Availability and Serviceability is a concept used on servers meant to measure their robustness.

**Reliability** is the probability that a system will produce correct outputs:

- Generally measured as Mean Time Between Failures (MTBF), and
- Enhanced by features that help to avoid, detect and repair hardware faults

**Availability** is the probability that a system is operational at a given time:

- Generally measured as a percentage of downtime per a period of time, and
- Often uses mechanisms to detect and correct hardware faults in runtime.

**Serviceability** is the simplicity and speed with which a system can be repaired or maintained:

- Generally measured on Mean Time Between Repair (MTBR).

Among the monitoring measures, the most usual ones include:

- CPU – detect errors at instruction execution and at L1/L2/L3 caches;
- Memory – add error correction logic (ECC) to detect and correct errors;
- I/O – add CRC checksums for transferred data;
- Storage – RAID, journal file systems, checksums, Self-Monitoring, Analysis and Reporting Technology (SMART).

By monitoring the number of occurrences of error detections, it is possible to identify if the probability of hardware errors is increasing, and, on such case, do a preventive maintenance to replace a degraded component while those errors are correctable.

For detailed information about the types of error events gathered and how to make sense of them, see the kernel administrator's guide at <https://www.kernel.org/doc/html/latest/admin-guide/ras.html>.

**rasdaemon-service-type** [Variable]

Service type for the **rasdaemon** service. It accepts a **rasdaemon-configuration** object. Instantiating like

```
(service rasdaemon-service-type)
```

will load with a default configuration, which monitors all events and logs to syslogd.

**rasdaemon-configuration** [Data Type]

The data type representing the configuration of **rasdaemon**.

**record?** (default: **#f**)

A boolean indicating whether to record the events in an SQLite database. This provides a more structured access to the information contained in the log file. The database location is hard-coded to `/var/lib/rasdaemon/ras-mc_event.db`.

## Zram Device Service

The Zram device service provides a compressed swap device in system memory. The Linux Kernel documentation has more information about zram (<https://www.kernel.org/doc/html/latest/admin-guide/blockdev/zram.html>) devices.

**zram-device-service-type** [Variable]

This service creates the zram block device, formats it as swap and enables it as a swap device. The service's value is a **zram-device-configuration** record.

**zram-device-configuration** [Data Type]

This is the data type representing the configuration for the zram-device service.

**size** (default "1G")

This is the amount of space you wish to provide for the zram device. It accepts a string and can be a number of bytes or use a suffix, eg.: "512M" or 1024000.

**compression-algorithm** (default 'lzo')

This is the compression algorithm you wish to use. It is difficult to list all the possible compression options, but common ones supported by Guix's Linux Libre Kernel include 'lzo', 'lz4' and 'zstd'.

**memory-limit** (default 0)

This is the maximum amount of memory which the zram device can use. Setting it to '0' disables the limit. While it is generally expected that compression will be 2:1, it is possible that uncompressable data can be written to swap and this is a method to limit how much memory can be used. It accepts a string and can be a number of bytes or use a suffix, eg.: "2G".

**priority** (default #f)

This is the priority of the swap device created from the zram device. See Section 11.6 [Swap Space], page 266, for a description of swap priorities. You might want to set a specific priority for the zram device, otherwise it could end up not being used much for the reasons described there.

### 11.10.38 Hurd Services

**hurd-console-service-type** [Variable]

This service starts the fancy VGA console client on the Hurd.

The service's value is a **hurd-console-configuration** record.

**hurd-console-configuration** [Data Type]

This is the data type representing the configuration for the **hurd-console-service**.

**hurd** (default: *hurd*)

The Hurd package to use.

**hurd-getty-service-type** [Variable]

This service starts a tty using the Hurd **getty** program.

The service's value is a **hurd-getty-configuration** record.

**hurd-getty-configuration** [Data Type]

This is the data type representing the configuration for the **hurd-getty-service**.

**hurd** (default: *hurd*)

The Hurd package to use.

**tty** The name of the console this Getty runs on—e.g., "**tty1**".

**baud-rate** (default: 38400)

An integer specifying the baud rate of the tty.

### 11.10.39 Miscellaneous Services

#### Fingerprint Service

The (**gnu services authentication**) module provides a DBus service to read and identify fingerprints via a fingerprint sensor.

**fprind-service-type** [Variable]

The service type for **fprind**, which provides the fingerprint reading capability.

(service **fprind-service-type**)

## System Control Service

The (gnu services `sysctl`) provides a service to configure kernel parameters at boot.

**sysctl-service-type** [Variable]

The service type for `sysctl`, which modifies kernel parameters under `/proc/sys/`. To enable IPv4 forwarding, it can be instantiated as:

```
(service sysctl-service-type
  (sysctl-configuration
    (settings '("net.ipv4.ip_forward" . "1"))))
```

Since `sysctl-service-type` is used in the default lists of services, `%base-services` and `%desktop-services`, you can use `modify-services` to change its configuration and add the kernel parameters that you want (see Section 11.19.3 [Service Reference], page 712).

```
(modify-services %base-services
  (sysctl-service-type config =>
    (sysctl-configuration
      (settings (append '("net.ipv4.ip_forward" . "1"))
        %default-sysctl-settings))))
```

**sysctl-configuration** [Data Type]

The data type representing the configuration of `sysctl`.

**sysctl** (default: (file-append `procs` `"/sbin/sysctl"`))

The `sysctl` executable to use.

**settings** (default: `%default-sysctl-settings`)

An association list specifies kernel parameters and their values.

**%default-sysctl-settings** [Variable]

An association list specifying the default `sysctl` parameters on Guix System.

## PC/SC Smart Card Daemon Service

The (gnu services `security-token`) module provides the following service to run `pcscd`, the PC/SC Smart Card Daemon. `pcscd` is the daemon program for `pcsc-lite` and the `MuscleCard` framework. It is a resource manager that coordinates communications with smart card readers, smart cards and cryptographic tokens that are connected to the system.

**pcscd-service-type** [Variable]

Service type for the `pcscd` service. Its value must be a `pcscd-configuration` object.

To run `pcscd` in the default configuration, instantiate it as:

```
(service pcscd-service-type)
```

**pcscd-configuration** [Data Type]

The data type representing the configuration of `pcscd`.

**pcsc-lite** (default: `pcsc-lite`)

The `pcsc-lite` package that provides `pcscd`.

**usb-drivers** (default: (list `ccid`))

List of packages that provide USB drivers to `pcscd`. Drivers are expected to be under `pcsc/drivers` in the store directory of the package.

## LIRC Service

The (gnu services lirc) module provides the following service.

**lirc-service-type** [Variable]

Type for a service that runs LIRC (<http://www.lirc.org>), a daemon that decodes infrared signals from remote controls.

The value for this service is a <lirc-configuration> object.

**lirc-configuration** [Data Type]

Data type representing the configuration of `lircd`.

`lirc` (default: `lirc`) (type: file-like)

Package object for `lirc`.

`device` (default: `#f`) (type: string)

`driver` (default: `#f`) (type: string)

`config-file` (default: `#f`) (type: string-or-file-like)

TODO. See `lircd` manual for details.

`extra-options` (default: `'()`) (type: list-of-string)

Additional command-line options to pass to `lircd`.

## SPICE Service

The (gnu services spice) module provides the following service.

**spice-vdagent-service-type** [Variable]

Type of the service that runs VDAGENT (<https://www.spice-space.org>), a daemon that enables sharing the clipboard with a vm and setting the guest display resolution when the graphical console window resizes.

**spice-vdagent-configuration** [Data Type]

Data type representing the configuration of `spice-vdagent-service-type`.

`spice-vdagent` (default: `spice-vdagent`) (type: file-like)

Package object for VDAGENT.

## inputattach Service

The `inputattach` (<https://linuxwacom.github.io/>) service allows you to use input devices such as Wacom tablets, touchscreens, or joysticks with the Xorg display server.

**inputattach-service-type** [Variable]

Type of a service that runs `inputattach` on a device and dispatches events from it.

**inputattach-configuration** [Data Type]

`device-type` (default: `"wacom"`)

The type of device to connect to. Run `inputattach --help`, from the `inputattach` package, to see the list of supported device types.

`device` (default: `"/dev/ttyS0"`)

The device file to connect to the device.

**baud-rate** (default: **#f**)  
 Baud rate to use for the serial connection. Should be a number or **#f**.

**log-file** (default: **#f**)  
 If true, this must be the name of a file to log messages to.

## Dictionary Service

The (`gnu services dict`) module provides the following service:

**dicod-service-type** [Variable]

This is the type of the service that runs the **dicod** daemon, an implementation of DICT server (see Section “Dicod” in *GNU Dico Manual*).

You can add **open localhost** to your `~/.dico` file to make **localhost** the default server for **dico** client (see Section “Initialization File” in *GNU Dico Manual*).

**Note:** This service is also available for Guix Home, where it runs directly with your user privileges (see Section 13.3.21 [Miscellaneous Home Services], page 781).

**dicod-configuration** [Data Type]

Data type representing the configuration of **dicod**.

**dico** (default: *dico*)  
 Package object of the GNU Dico dictionary server.

**interfaces** (default: `'("localhost")`)  
 This is the list of IP addresses and ports and possibly socket file names to listen to (see Section “Server Settings” in *GNU Dico Manual*).

**handlers** (default: `'()`)  
 List of **<dicod-handler>** objects denoting handlers (module instances).

**databases** (default: `(list %dicod-database:gcode)`)  
 List of **<dicod-database>** objects denoting dictionaries to be served.

**dicod-handler** [Data Type]

Data type representing a dictionary handler (module instance).

**name** Name of the handler (module instance).

**module** (default: **#f**)  
 Name of the **dicod** module of the handler (instance). If it is **#f**, the module has the same name as the handler. (see Section “Modules” in *GNU Dico Manual*).

**options** List of strings or gexps representing the arguments for the module handler

**dicod-database** [Data Type]

Data type representing a dictionary database.

**name** Name of the database, will be used in DICT commands.

**handler** Name of the **dicod** handler (module instance) used by this database (see Section “Handlers” in *GNU Dico Manual*).



**complex?** (default: *#f*)  
 Whether the database configuration complex. The complex configuration will need a corresponding `<dicod-handler>` object, otherwise not.

**options** List of strings or gexps representing the arguments for the database (see Section “Databases” in *GNU Dico Manual*).

**%dicod-database:gcide** [Variable]  
 A `<dicod-database>` object serving the GNU Collaborative International Dictionary of English using the `gcide` package.

**dicod-freedict-database** *dictionary-name* [Procedure]  
 Returns a record of type `dicod-database` for purpose of configuring a database for the FreeDict multilingual dictionary named by *dictionary-name*, a string such as “kur-eng” for the Kurdish-to-English dictionary, into a service of type `dicod-service-type`.

**%dicod-databases:freedict** [Variable]  
 A relatively large list of records of type `dicod-database`, made available for users who wish to configure all the FreeDict multilingual dictionaries provided by the `freedict-dictionaries` package into a service of type `dicod-service-type`.

The following is an example `dicod-service-type` configuration, extending a `dicod` daemon serving databases for the wordnet and `gcide` dictionaries, as well as a selection of FreeDict multilingual dictionaries.

```
(service dicod-service-type
  (dicod-configuration
    (handlers (list
      (dicod-handler
        (name "wordnet")
        (module "wordnet")
        (options
          (list #~(string-append "wnhome=" #$wordnet))))))
    (databases (cons*
      (dicod-database
        (name "wordnet")
        (complex? #t)
        (handler "wordnet"))
      %dicod-database:gcide
      (map
        dicod-freedict-database
        '("fra-eng" "eng-fra" "eng-spa" "spa-eng"))))))
```

## Docker Service

The `(gnu services docker)` module provides the following services.

**containerd-service-type** [Variable]  
 This service type operates containerd containerd (<https://containerd.io>), a daemon responsible for overseeing the entire container lifecycle on its host system. This

includes image handling, storage management, container execution, supervision, low-level storage operations, network connections, and more.

**containerd-configuration** [Data Type]

This is the data type representing the configuration of containerd.

**containerd** (default: **containerd**)

The containerd daemon package to use.

**debug?** (default **#f**)

Enable or disable debug output.

**environment-variables** (default: '()')

List of environment variables to set for **containerd**.

This must be a list of strings where each string has the form ‘**key=value**’ as in this example:

```
(list "HTTP_PROXY=socks5://127.0.0.1:9150"
      "HTTPS_PROXY=socks5://127.0.0.1:9150")
```

**docker-service-type** [Variable]

This is the type of the service that runs Docker (<https://www.docker.com>), a daemon that can execute application bundles (sometimes referred to as “containers”) in isolated environments.

The **containerd-service-type** service need to be added to a system configuration, otherwise a message about not any service provides **containerd** will be displayed during **guix system reconfigure**.

**docker-configuration** [Data Type]

This is the data type representing the configuration of Docker and Containerd.

**docker** (default: **docker**)

The Docker daemon package to use.

**docker-cli** (default: **docker-cli**)

The Docker client package to use.

**containerd** (default: **containerd**)

This field is deprecated in favor of **containerd-service-type** service.

**proxy** (default **docker-libnetwork-cmd-proxy**)

The Docker user-land networking proxy package to use.

**enable-proxy?** (default **#t**)

Enable or disable the use of the Docker user-land networking proxy.

**debug?** (default **#f**)

Enable or disable debug output.

**enable-iptables?** (default **#t**)

Enable or disable the addition of iptables rules.

**environment-variables** (default: '()')

List of environment variables to set for **dockerd**.

This must be a list of strings where each string has the form ‘**key=value**’ as in this example:

```
(list "LANGUAGE=eo:ca:eu"
      "TMPDIR=/tmp/dockerd")
```

**config-file** (type: maybe-file-like)

JSON configuration file pass to **dockerd**.

**singularity-service-type** [Variable]

This is the type of the service that allows you to run Singularity (<https://www.sylabs.io/singularity/>), a Docker-style tool to create and run application bundles (aka. “containers”). The value for this service is the Singularity package to use.

The service does not install a daemon; instead, it installs helper programs as **setuid-root** (see Section 11.11 [Privileged Programs], page 678) such that unprivileged users can invoke **singularity run** and similar commands.

## Rootless Podman Service

The (**gnu services containers**) module provides the following service.

**rootless-podman-service-type** [Variable]

This is the service type for Podman (<https://podman.io>) is a container management tool.

In addition to providing a drop-in replacement for Docker, Podman offers the ability to run containers in “root-less” mode, meaning that regular users can deploy containers without elevated privileges. It does so mainly by leveraging two Linux kernel features: unprivileged user namespaces, and subordinate user and group IDs (see [subordinate-user-group-ids], page 663).

The **rootless-podman-service-type** sets up the system to allow unprivileged users to run **podman** commands:

```
(use-service-modules containers dbus desktop networking ...)
(use-modules (gnu system accounts)) ;for 'subid-range'

(operating-system
  ;; ...
  (users (cons (user-account
                (name "alice")
                (comment "Bob's sister")
                (group "users"))

              ;; Adding the account to the "cgroup" group
              ;; makes it possible to run podman commands.
              (supplementary-groups ("cgroup" "wheel"
                                     "audio" "video"))))

  %base-user-accounts))

(services
  (cons* (service dbus-root-service-type)
```

```
(service dhcpd-service-type)
(service elogind-service-type)
(service iptables-service-type)
(service rootless-podman-service-type
  (rootless-podman-configuration
    (subgids
      (list (subid-range (name "alice")))))
    (subuids
      (list (subid-range (name "alice"))))))
%base-services)))
```

The `iptables-service-type` is required for Podman to be able to setup its own networks. Due to the change in user groups and file systems it is recommended to reboot (or at least logout), before trying to run Podman commands.

To test your installation you can run:

```
$ podman run -it --rm docker.io/alpine cat /etc/*release*
NAME="Alpine Linux"
ID=alpine
VERSION_ID=3.20.2
PRETTY_NAME="Alpine Linux v3.20"
HOME_URL="https://alpinelinux.org/"
BUG_REPORT_URL="https://gitlab.alpinelinux.org/alpine/aports/-/issues"■
```

`rootless-podman-configuration` [Data Type]

Available `rootless-podman-configuration` fields are:

`podman` (default: `podman`) (type: `package-or-#f`)

The Podman package that will be installed in the system profile. Pass `#f` to not install Podman.

`group-name` (default: `"cgroup"`) (type: `string`)

The name of the group that will own `/sys/fs/cgroup` resources. Users that want to use rootless Podman have to be in this group.

`containers-registries` (type: `lowerable`)

A string or a gexp evaluating to the path of Podman's `containers/registries.conf` configuration file.

`containers-storage` (type: `lowerable`)

A string or a gexp evaluating to the path of Podman's `containers/storage.conf` configuration file.

`containers-policy` (type: `lowerable`)

A string or a gexp evaluating to the path of Podman's `containers/policy.json` configuration file.

`pam-limits` (type: `list-of-pam-limits-entries`)

The PAM limits to be set for rootless Podman.

`subgids` (default: `()`) (type: `list-of-subid-ranges`)

A list of subid ranges representing the subgids that will be available for each configured user.

`subuids` (default: ()) (type: list-of-subid-ranges)

A list of subid ranges representing the subuids that will be available for each configured user.

## OCI backed services

Should you wish to manage your Open Container Initiative (OCI) containers with the same consistent interface you use for your other Shepherd services, *oci-service-type* is the tool to use: given an OCI container image, it will run it in a Shepherd service. One example where this is useful: it lets you run services that are available as OCI images but not yet packaged for Guix.

**oci-service-type** [Variable]

This is a thin wrapper around Docker's or Podman's CLI that executes OCI images backed processes as Shepherd Services.

```
(simple-service 'oci-provisioning
  oci-service-type
  (oci-extension
    (networks
      (list
        (oci-network-configuration (name "monitoring"))))
    (containers
      (list
        (oci-container-configuration
          (network "monitoring")
          (image
            (oci-image
              (repository "guile")
              (tag "3")
              (value (specifications->manifest '("guile"))
                (pack-options '(:symlinks (("/bin/guile" -> "bin/guile")
                  #:max-layers 2))))
            (entrypoint "/bin/guile")
            (command
              '("-c" "(display \"hello!\n\")"))
            (oci-container-configuration
              (image "prom/prometheus")
              (network "host")
              (ports
                '(("9000" . "9000")
                  ("9090" . "9090"))))
            (oci-container-configuration
              (image "grafana/grafana:10.0.1")
              (network "host")
              (volumes
                '("/var/lib/grafana:/var/lib/grafana"))))))))
```

In this example three different Shepherd services are going to be added to the system. Each *oci-container-configuration* record translates to a ‘docker run’ or

'podman run' invocation and its fields directly map to options. You can refer to the Docker (<https://docs.docker.com/engine/reference/commandline/run>) or Podman (<https://docs.podman.io/en/stable/markdown/podman-run.1.html>) upstream documentation for semantics of each value. If the images are not found, they will be pulled. You can refer to the Docker (<https://docs.docker.com/engine/reference/commandline/pull/>) or Podman (<https://docs.podman.io/en/stable/markdown/podman-pull.1.html>) upstream documentation for semantics. The services with (`network "host"`) are going to be attached to the host network and are supposed to behave like native processes with regard to networking.

**oci-configuration** [Data Type]

Available **oci-configuration** fields are:

**runtime** (default: 'docker') (type: symbol)

The OCI runtime to use to run commands. It can be either 'docker' or 'podman'.

**runtime-cli** (type: maybe-package-or-string)

The OCI runtime command line to be installed in the system profile and used to provision OCI resources, it can be either a package or a string representing an absolute file name to the runtime binary entrypoint. When unset it will default to **docker-cli** package for the 'docker' runtime or to **podman** package for the 'podman' runtime.

**runtime-extra-arguments** (default: '()') (type: list)

A list of strings, gexps or file-like objects that will be placed after each **docker** or **podman** invocation.

**user** (type: maybe-string)

The user name under whose authority OCI commands will be run. This field will override the **user** field of **oci-configuration**.

**group** (type: maybe-string)

The group name under whose authority OCI commands will be run. When using the 'podman' OCI runtime, this field will be ignored and the default group of the user configured in the **user** field will be used. This field will override the **group** field of **oci-configuration**.

**subuids-range** (type: maybe-subid-range)

An optional **subid-range** record allocating subuids for the user from the **user** field. When unset, with the rootless Podman OCI runtime, it defaults to (**subid-range** (name "oci-container")).

**subgids-range** (type: maybe-subid-range)

An optional **subid-range** record allocating subgids for the user from the **user** field. When unset, with the rootless Podman OCI runtime, it defaults to (**subid-range** (name "oci-container")).

**containers** (default: '()') (type: list-of-oci-containers)

The list of **oci-container-configuration** records representing the containers to provision. The use of the **oci-extension** record should be preferred for most cases.

**networks** (default: '()') (type: list-of-oci-networks)  
 The list of **oci-network-configuration** records representing the containers to provision. The use of the **oci-extension** record should be preferred for most cases.

**volumes** (default: '()') (type: list-of-oci-volumes)  
 The list of **oci-volumes-configuration** records representing the containers to provision. The use of the **oci-extension** record should be preferred for most cases.

**verbose?** (default: **#f**) (type: boolean)  
 When true, additional output will be printed, allowing to better follow the flow of execution.

**oci-extension** [Data Type]

Available **oci-extension** fields are:

**containers** (default: '()') (type: list-of-oci-containers)  
 The list of **oci-container-configuration** records representing the containers to provision.

**networks** (default: '()') (type: list-of-oci-networks)  
 The list of **oci-network-configuration** records representing the containers to provision.

**volumes** (default: '()') (type: list-of-oci-volumes)  
 The list of **oci-volumes-configuration** records representing the containers to provision.

**oci-container-configuration** [Data Type]

Available **oci-container-configuration** fields are:

**user** (default: "oci-container") (type: string)  
 The user under whose authority docker commands will be run.

**group** (default: "docker") (type: string)  
 The group under whose authority docker commands will be run.

**command** (default: '()') (type: list-of-strings)  
 Overwrite the default command (CMD) of the image.

**entrypoint** (default: "") (type: string)  
 Overwrite the default entrypoint (ENTRYPOINT) of the image.

**host-environment** (default: '()') (type: list)  
 Set environment variables in the host environment where 'docker run' or 'podman run' are invoked. This is especially useful to pass secrets from the host to the container without having them on the OCI runtime command line, for example: by setting the **MYSQL\_PASSWORD** on the host and by passing **--env MYSQL\_PASSWORD** through the **extra-arguments** field, it is possible to securely set values in the container environment. This field's value can be a list of pairs or strings, even mixed:

```
(list '("LANGUAGE" . "eo:ca:eu")
```

```
"JAVA_HOME=/opt/java")
```

Pair members can be strings, gexps or file-like objects. Strings are passed directly to `make-forkexec-creator`.

**environment** (default: '()') (type: list)

Set environment variables inside the container. This can be a list of pairs or strings, even mixed:

```
(list '("LANGUAGE" . "eo:ca:eu")
      "JAVA_HOME=/opt/java")
```

Pair members can be strings, gexps or file-like objects. Strings are passed directly to the OCI runtime CLI. You can refer to the Docker (<https://docs.docker.com/engine/reference/commandline/run/#env>) or Podman (<https://docs.podman.io/en/stable/markdown/podman-run.1.html#env-e-env>) upstream documentation for semantics.

**image** (type: string-or-oci-image)

The image used to build the container. It can be a string or an `oci-image` record. Strings are resolved by the OCI runtime, and follow the usual format `myregistry.local:5000/testing/test-image:tag`.

**provision** (default: "") (type: string)

Set the name of the provisioned Shepherd service.

**requirement** (default: '()') (type: list-of-symbols)

Set additional Shepherd services dependencies to the provisioned Shepherd service.

**log-file** (type: maybe-string)

When `log-file` is set, it names the file to which the service's standard output and standard error are redirected. `log-file` is created if it does not exist, otherwise it is appended to.

**auto-start?** (default: `#t`) (type: boolean)

Whether this service should be started automatically by the Shepherd. If it is `#f`, the service has to be started manually with `herd start`.

**respawn?** (default: `#f`) (type: boolean)

Whether to have Shepherd restart the service when it stops, for instance when the underlying process dies.

**shepherd-actions** (default: '()') (type: list-of-shepherd-actions)

This is a list of `shepherd-action` records defining actions supported by the service.

**network** (default: "") (type: string)

Set an OCI network for the spawned container.

**ports** (default: '()') (type: list)

Set the port or port ranges to expose from the spawned container. This can be a list of pairs or strings, even mixed:

```
(list '("8080" . "80")
```



```
"10443:443")
```

Pair members can be strings, gexps or file-like objects. Strings are passed directly to the OCI runtime CLI. You can refer to the Docker (<https://docs.docker.com/engine/reference/commandline/run/#publish>) or Podman (<https://docs.podman.io/en/stable/markdown/podman-run.1.html#publish-p-ip-hostport-containerport-protocol>) upstream documentation for semantics.

**volumes** (default: '()') (type: list)

Set volume mappings for the spawned container. This can be a list of pairs or strings, even mixed:

```
(list '("/root/data/grafana" . "/var/lib/grafana")
      "/gnu/store:/gnu/store")
```

Pair members can be strings, gexps or file-like objects. Strings are passed directly to the OCI runtime CLI. You can refer to the Docker (<https://docs.docker.com/engine/reference/commandline/run/#volume>) or Podman (<https://docs.podman.io/en/stable/markdown/podman-run.1.html#volume-v-source-volume-host-dir-container-dir-options>) upstream documentation for semantics.

**container-user** (default: "") (type: string)

Set the current user inside the spawned container. You can refer to the Docker (<https://docs.docker.com/engine/reference/run/#user>) or Podman (<https://docs.podman.io/en/stable/markdown/podman-run.1.html#user-u-user-group>) upstream documentation for semantics.

**workdir** (default: "") (type: string)

Set the current working directory for the spawned Shepherd service. You can refer to the Docker (<https://docs.docker.com/engine/reference/run/#workdir>) or Podman (<https://docs.podman.io/en/stable/markdown/podman-run.1.html#workdir-w-dir>) upstream documentation for semantics.

**extra-arguments** (default: '()') (type: list)

A list of strings, gexps or file-like objects that will be directly passed to the ‘docker run’ or ‘podman run’ invocation.

**oci-network-configuration**

[Data Type]

Available **oci-network-configuration** fields are:

**name** (type: string)

The name of the OCI network to provision.

**driver** (type: maybe-string)

The driver to manage the network.

**gateway** (type: maybe-string)

IPv4 or IPv6 gateway for the subnet.

**internal?** (default: **#f**) (type: boolean)  
 Restrict external access to the network

**ip-range** (type: maybe-string)  
 Allocate container ip from a sub-range in CIDR format.

**ipam-driver** (type: maybe-string)  
 IP Address Management Driver.

**ipv6?** (default: **#f**) (type: boolean)  
 Enable IPv6 networking.

**subnet** (type: maybe-string)  
 Subnet in CIDR format that represents a network segment.

**labels** (default: '()') (type: list)  
 The list of labels that will be used to tag the current volume.

**extra-arguments** (default: '()') (type: list)  
 A list of strings, gexps or file-like objects that will be directly passed to the ‘**docker network create**’ or ‘**podman network create**’ invocation.

**oci-volume-configuration** [Data Type]

Available **oci-volume-configuration** fields are:

**name** (type: string)  
 The name of the OCI volume to provision.

**labels** (default: '()') (type: list)  
 The list of labels that will be used to tag the current volume.

**extra-arguments** (default: '()') (type: list)  
 A list of strings, gexps or file-like objects that will be directly passed to the ‘**docker volume create**’ or ‘**podman volume create**’ invocation.

**oci-image** [Data Type]

Available **oci-image** fields are:

**repository** (type: string)  
 A string like **myregistry.local:5000/testing/test-image** that names the OCI image.

**tag** (default: **"latest"**) (type: string)  
 A string representing the OCI image tag. Defaults to **latest**.

**value** (type: oci-lowerable-image)  
 A **manifest** or **operating-system** record that will be lowered into an OCI compatible tarball. Otherwise this field’s value can be a file-like object that evaluates to an OCI compatible tarball.

**pack-options** (default: '()') (type: list)  
 An optional set of keyword arguments that will be passed to the **docker-image** procedure from **guix scripts pack**. They can be used to replicate **guix pack** behavior:

(oci-image

```
(repository "guile")
(tag "3")
(value
  (specifications->manifest '("guile")))
(pack-options '(:symlinks ("/bin/guile" -> "bin/guile"))
  #:max-layers 2)))
```

If the `value` field is an `operating-system` record, this field's value will be ignored.

```
system (default: "") (type: string)
  Attempt to build for a given system, e.g. "i686-linux"

target (default: "") (type: string)
  Attempt to cross-build for a given triple, e.g. "aarch64-linux-gnu"

grafts? (default: #f) (type: boolean)
  Whether to allow grafting or not in the pack build.
```

## Subordinate User and Group ID Service

Among the virtualization facilities implemented by the Linux kernel is the concept of *subordinate IDs*. Subordinate IDs allow for mapping user and group IDs inside process namespaces to user and group IDs of the host system. Subordinate user ID ranges (subuids) allow users to map virtual user IDs inside containers to the user ID of an unprivileged user of the host system. Subordinate group ID ranges (subgids), instead map virtual group IDs to the group ID of an unprivileged user on the host system. You can access `subuid(5)` and `subgid(5)` Linux man pages for more details.

The (`gnu system shadow`) module exposes the `subids-service-type`, its configuration record `subids-configuration` and its extension record `subids-extension`.

With `subids-service-type`, subuids and subgids ranges can be reserved for users that desire so:

```
(use-modules (gnu system shadow)      ;for 'subids-service-type'
             (gnu system accounts)    ;for 'subid-range'
             ...)

(operating-system
 ;; ...
 (services
  (list
   (simple-service 'alice-bob-subids
                  subids-service-type
                  (subids-extension
                   (subgids
                    (list
                     (subid-range (name "alice")))))
                  (subuids
                   (list
                    (subid-range (name "alice"))
                    (subid-range (name "bob"))
```

```
(start 100700))))))))))
```

Users (definitely other services), usually, are supposed to extend the service instead of adding subids directly to `subids-configuration`, unless the want to change the default behavior for root. With default settings the `subids-service-type` adds, if it's not already there, a configuration for the root account to both `/etc/subuid` and `/etc/subgid`, possibly starting at the minimum possible subid. Otherwise the root subuids and subgids ranges are fitted wherever possible.

The above configuration will yield the following:

```
# cat /etc/subgid
root:100000:65536
alice:165536:65536
# cat /etc/subuid
root:100000:700
bob:100700:65536
alice:166236:65536
```

### `subids-configuration`

[Data Type]

With default settings the `subids-service-type` adds, if it's not already there, a configuration for the root account to both `/etc/subuid` and `/etc/subgid`, possibly starting at the minimum possible subid. To disable the default behavior and provide your own definition for the root subid ranges you can set to `#f` the `add-root?` field:

```
(use-modules (gnu system shadow)      ;for 'subids-service-type'
              (gnu system accounts)    ;for 'subid-range'
              ...)

```

```
(operating-system
;; ...
(services
 (list
  (service subids-service-type
            (subids-configuration
              (add-root? #f)
              (subgids
                (subid-range (name "root")
                             (start 120000)
                             (count 100)))
              (subuids
                (subid-range (name "root")
                             (start 120000)
                             (count 100))))))
  (simple-service 'alice-bob-subids
                  subids-service-type
                  (subids-extension
                    (subgids
                     (list
                      (subid-range (name "alice")))))
```

```
(subuids
  (list
    (subid-range (name "alice"))
    (subid-range (name "bob")
      (start 100700))))))
```

Available **subids-configuration** fields are:

**add-root?** (default: **#t**) (type: boolean)

Whether to automatically configure subuids and subgids for root.

**subgids** (default: '()) (type: list-of-subid-ranges)

The list of **subid-ranges** that will be serialized to **/etc/subgid**. If a range doesn't specify a start it will be fitted based on its number of requested subids. If a range doesn't specify a count the default size of 65536 will be assumed.

**subuids** (default: '()) (type: list-of-subid-ranges)

The list of **subid-ranges** that will be serialized to **/etc/subuid**. If a range doesn't specify a start it will be fitted based on its number of requested subids. If a range doesn't specify a count the default size of 65536 will be assumed.

## **subids-extension**

[Data Type]

Available **subids-extension** fields are:

**subgids** (default: '()) (type: list-of-subid-ranges)

The list of **subid-ranges** that will be appended to **subids-configuration-subgids**. Entries with the same name are deduplicated upon merging.

**subuids** (default: '()) (type: list-of-subid-ranges)

The list of **subid-ranges** that will be appended to **subids-configuration-subuids**. Entries with the same name are deduplicated upon merging.

## **subid-range**

[Data Type]

The **subid-range** record is defined at (**gnu system accounts**). Available fields are:

**name** (type: string)

The name of the user or group that will own this range.

**start** (default: **#f**) (type: integer)

The first requested subid. When false the first available subid with enough contiguous subids will be assigned.

**count** (default: **#f**) (type: integer)

The number of total allocated subids. When **#f** the default of 65536 will be assumed .

## **Auditd Service**

The (**gnu services auditd**) module provides the following service.

**auditd-service-type** [Variable]

This is the type of the service that runs `auditd` (<https://people.redhat.com/sgrubb/audit/>), a daemon that tracks security-relevant information on your system.

Examples of things that can be tracked:

1. File accesses
2. System calls
3. Invoked commands
4. Failed login attempts
5. Firewall filtering
6. Network access

`auditctl` from the `audit` package can be used in order to add or remove events to be tracked (until the next reboot). In order to permanently track events, put the command line arguments of `auditctl` into a file called `audit.rules` in the configuration directory (see below). `aureport` from the `audit` package can be used in order to view a report of all recorded events. The audit daemon by default logs into the file `/var/log/audit.log`.

**auditd-configuration** [Data Type]

This is the data type representing the configuration of `auditd`.

`audit` (default: `audit`)

The audit package to use.

`configuration-directory` (default:

`%default-auditd-configuration-directory`)

The directory containing the configuration file for the audit package, which must be named `auditd.conf`, and optionally some audit rules to instantiate on startup.

**R-Shiny service**

The (`gnu services science`) module provides the following service.

**rshiny-service-type** [Variable]

This is a type of service which is used to run a webapp created with `r-shiny`. This service sets the `R_LIBS_USER` environment variable and runs the provided script to call `runApp`.

**rshiny-configuration** [Data Type]

This is the data type representing the configuration of `rshiny`.

`package` (default: `r-shiny`)

The package to use.

`binary` (default `"rshiny"`)

The name of the binary or shell script located at `package/bin/` to run when the service is run.

The common way to create this file is as follows:

...

```

        (let* ((out      (assoc-ref %outputs "out"))
               (targetdir (string-append out "/share/" ,name)))
          (app      (string-append out "/bin/" ,name))
          (Rbin      (search-input-file %build-inputs "/bin/Rscript"))

          ;; ...
          (mkdir-p (string-append out "/bin"))
          (call-with-output-file app
            (lambda (port)
              (format port
                "#!~a\n"
                (library(shiny)
                  setwd("~/~a")
                  runApp(launch.browser=0, port=4202)~%\n"
                  Rbin targetdir))))))

```

## Nix service

The (gnu services nix) module provides the following service.

### nix-service-type

[Variable]

This is the type of the service that runs build daemon of the Nix (<https://nixos.org/nix/>) package manager. Here is an example showing how to use it:

```

(use-modules (gnu))
(use-service-modules nix)
(use-package-modules package-management)

(operating-system
  ;; ...
  (packages (append (list nix)
                    %base-packages))

  (services (append (list (service nix-service-type))
                    %base-services)))

```

After `guix system reconfigure` configure Nix for your user:

- Add a Nix channel and update it. See Nix channels ([https://wiki.nixos.org/wiki/Nix\\_channels](https://wiki.nixos.org/wiki/Nix_channels)) for more information about the available channels. If you would like to use the unstable Nix channel you can do this by running:

```

$ nix-channel --add https://nixos.org/channels/nixpkgs-unstable
$ nix-channel --update

```

- Create your Nix profile directory:

```

$ sudo mkdir -p /nix/var/nix/profiles/per-user/$USER
$ sudo chown $USER:root /nix/var/nix/profiles/per-user/$USER

```

- Create a symlink to your profile and activate Nix profile:

```

$ ln -s "/nix/var/nix/profiles/per-user/$USER/profile" ~/.nix-profile
$ source /run/current-system/profile/etc/profile.d/nix.sh

```

**nix-configuration** [Data Type]

This data type represents the configuration of the Nix daemon.

**nix** (default: **nix**)

The Nix package to use.

**sandbox** (default: **#t**)

Specifies whether builds are sandboxed by default.

**build-directory** (default: **"/tmp"**)

The directory where build directory are stored during builds. This is useful to change if, for example, the default location does not have enough space to hold build trees for big packages.

This is similar to setting the **TMPDIR** environment variable for **guix-daemon**. Section 2.2.1 [Build Environment Setup], page 6, for more info.

**build-sandbox-items** (default: **'()**)

This is a list of strings or objects appended to the **build-sandbox-items** field of the configuration file.

**extra-config** (default: **'()**)

This is a list of strings or objects appended to the configuration file. It is used to pass extra text to be added verbatim to the configuration file.

**extra-options** (default: **'()**)

Extra command line options for **nix-service-type**.

## Fail2Ban service

**fail2ban** (<http://www.fail2ban.org/>) scans log files (e.g. **/var/log/apache/error\_log**) and bans IP addresses that show malicious signs – repeated password failures, attempts to make use of exploits, etc.

**fail2ban-service-type** service type is provided by the **(gnu services security)** module.

This service type runs the **fail2ban** daemon. It can be configured in various ways, which are:

### Basic configuration

The basic parameters of the Fail2Ban service can be configured via its **fail2ban** configuration, which is documented below.

### User-specified jail extensions

The **fail2ban-jail-service** function can be used to add new Fail2Ban jails.

### Shepherd extension mechanism

Service developers can extend the **fail2ban-service-type** service type itself via the usual service extension mechanism.

**fail2ban-service-type** [Variable]

This is the type of the service that runs **fail2ban** daemon. Below is an example of a basic, explicit configuration:

```
(append
```



```
(list
  (service fail2ban-service-type
    (fail2ban-configuration
      (extra-jails
        (list
          (fail2ban-jail-configuration
            (name "sshd")
            (enabled? #t))))))
  ;; There is no implicit dependency on an actual SSH
  ;; service, so you need to provide one.
  (service openssh-service-type))
%base-services)
```

**fail2ban-jail-service** *svc-type jail* [Procedure]

Extend *svc-type*, a `<service-type>` object with *jail*, a `fail2ban-jail-configuration` object.

For example:

```
(append
  (list
    (service
      ;; The 'fail2ban-jail-service' procedure can extend any service type
      ;; with a fail2ban jail. This removes the requirement to explicitly
      ;; extend services with fail2ban-service-type.
      (fail2ban-jail-service
        openssh-service-type
        (fail2ban-jail-configuration
          (name "sshd")
          (enabled? #t)))
        (openssh-configuration ...))))
```

Below is the reference for the different `jail-service-type` configuration records.

**fail2ban-configuration** [Data Type]

Available `fail2ban-configuration` fields are:

**fail2ban** (default: `fail2ban`) (type: `package`)

The `fail2ban` package to use. It is used for both binaries and as base default configuration that is to be extended with `<fail2ban-jail-configuration>` objects.

**run-directory** (default: `"/var/run/fail2ban"`) (type: `string`)

The state directory for the `fail2ban` daemon.

**jails** (default: `'()`) (type: `list-of-fail2ban-jail-configurations`)

Instances of `<fail2ban-jail-configuration>` collected from extensions.

**extra-jails** (default: `'()`) (type: `list-of-fail2ban-jail-configurations`)

Instances of `<fail2ban-jail-configuration>` explicitly provided.

**extra-content** (default: '()') (type: text-config)

Extra raw content to add to the end of the `jail.local` file, provided as a list of file-like objects.

**fail2ban-ignore-cache-configuration** [Data Type]

Available `fail2ban-ignore-cache-configuration` fields are:

**key** (type: string)

Cache key.

**max-count** (type: integer)

Cache size.

**max-time** (type: integer)

Cache time.

**fail2ban-jail-action-configuration** [Data Type]

Available `fail2ban-jail-action-configuration` fields are:

**name** (type: string)

Action name.

**arguments** (default: '()') (type: list-of-arguments)

Action arguments.

**fail2ban-jail-configuration** [Data Type]

Available `fail2ban-jail-configuration` fields are:

**name** (type: string)

Required name of this jail configuration.

**enabled?** (default: `#t`) (type: boolean)

Whether this jail is enabled.

**backend** (type: maybe-symbol)

Backend to use to detect changes in the `log-path`. The default is 'auto'. To consult the defaults of the jail configuration, refer to the `/etc/fail2ban/jail.conf` file of the `fail2ban` package.

**max-retry** (type: maybe-integer)

The number of failures before a host gets banned (e.g. `(max-retry 5)`).

**max-matches** (type: maybe-integer)

The number of matches stored in ticket (resolvable via tag `<matches>`) in action.

**find-time** (type: maybe-string)

The time window during which the maximum retry count must be reached for an IP address to be banned. A host is banned if it has generated `max-retry` during the last `find-time` seconds (e.g. `(find-time "10m")`). It can be provided in seconds or using Fail2Ban's "time abbreviation format", as described in `man 5 jail.conf`.

**ban-time** (type: maybe-string)

The duration, in seconds or time abbreviated format, that a ban should last. (e.g. `(ban-time "10m")`).

**ban-time-increment?** (type: maybe-boolean)  
Whether to consider past bans to compute increases to the default ban time of a specific IP address.

**ban-time-factor** (type: maybe-string)  
The coefficient to use to compute an exponentially growing ban time.

**ban-time-formula** (type: maybe-string)  
This is the formula used to calculate the next value of a ban time.

**ban-time-multipliers** (type: maybe-string)  
Used to calculate next value of ban time instead of formula.

**ban-time-max-time** (type: maybe-string)  
The maximum number of seconds a ban should last.

**ban-time-rnd-time** (type: maybe-string)  
The maximum number of seconds a randomized ban time should last. This can be useful to stop “clever” botnets calculating the exact time an IP address can be unbanned again.

**ban-time-overall-jails?** (type: maybe-boolean)  
When true, it specifies the search of an IP address in the database should be made across all jails. Otherwise, only the current jail of the ban IP address is considered.

**ignore-self?** (type: maybe-boolean)  
Never ban the local machine’s own IP address.

**ignore-ip** (default: '()') (type: list-of-strings)  
A list of IP addresses, CIDR masks or DNS hosts to ignore. **fail2ban** will not ban a host which matches an address in this list.

**ignore-cache** (type: maybe-fail2ban-ignore-cache-configuration)  
Provide cache parameters for the ignore failure check.

**filter** (type: maybe-fail2ban-jail-filter-configuration)  
The filter to use by the jail, specified via a **<fail2ban-jail-filter-configuration>** object. By default, jails have names matching their filter name.

**log-time-zone** (type: maybe-string)  
The default time zone for log lines that do not have one.

**log-encoding** (type: maybe-symbol)  
The encoding of the log files handled by the jail. Possible values are: 'ascii', 'utf-8' and 'auto'.

**log-path** (default: '()') (type: list-of-strings)  
The file names of the log files to be monitored.

**action** (default: '()') (type: list-of-fail2ban-jail-actions)  
A list of **<fail2ban-jail-action-configuration>**.

**extra-content** (default: '()') (type: text-config)  
Extra content for the jail configuration, provided as a list of file-like objects.

**fail2ban-jail-filter-configuration** [Data Type]

Available **fail2ban-jail-filter-configuration** fields are:

**name** (type: string)  
Filter to use.

**mode** (type: maybe-string)  
Mode for filter.

## OpenSnitch Service

OpenSnitch (<https://github.com/evilsocket/opensnitch>) is an application-level firewall that monitors outbound connections and prompts users to allow or deny them on a per-application basis.

**opensnitch-service-type** is provided by the (**gnu services opensnitch**) module.

**opensnitch-service-type** [Variable]

This is the service type for the OpenSnitch application firewall daemon. Its value must be an **opensnitch-configuration** record.

Below is an example configuration:

```
(service opensnitch-service-type)
```

This service depends on the **networking** service.

**opensnitch-configuration** [Data Type]

Available **opensnitch-configuration** fields are:

**opensnitch** (default: **opensnitch-daemon**) (type: package)  
The **opensnitch-daemon** package to use.

**server-address** (default: **"unix:///tmp/osui.sock"**) (type: string)  
Address for the UI to connect to the daemon.

**server-log-file** (default: **"/var/log/opensnitchd.log"**) (type: string)  
Path to the daemon log file.

**authentication-type** (default: **"simple"**) (type: string)  
Authentication type for UI-daemon communication.

**tls-ca-cert** (default: **""**) (type: string)  
Path to TLS CA certificate.

**tls-server-cert** (default: **""**) (type: string)  
Path to TLS server certificate.

**tls-client-cert** (default: **""**) (type: string)  
Path to TLS client certificate.

**tls-client-key** (default: **""**) (type: string)  
Path to TLS client key.

**tls-skip-verify?** (default: **#f**) (type: boolean)  
Whether to skip TLS verification.

**tls-client-auth-type** (default: **"no-client-cert"**) (type: string)  
TLS client authentication type.

**default-action** (default: "allow") (type: string)  
Default action for connections: "allow" or "deny".

**default-duration** (default: "once") (type: string)  
Default duration for rules: "once", "until-restart", "always", etc.

**intercept-unknown?** (default: #f) (type: boolean)  
Whether to intercept connections from unknown processes.

**proc-monitor-method** (default: "ebpf") (type: string)  
Method for monitoring processes: "ebpf", "proc", or "audit".

**log-level** (default: 2) (type: integer)  
Log level: 0=silent, 1=error, 2=warning, 3=important, 4=debug.

**log-utc?** (default: #t) (type: boolean)  
Whether to log timestamps in UTC.

**log-micro?** (default: #f) (type: boolean)  
Whether to include microseconds in log timestamps.

**firewall** (default: "nftables") (type: string)  
Firewall backend: "nftables" or "iptables".

**fw-config-path** (default: "/etc/opensnitchd/system-fw.json") (type: string)  
Path to the system firewall configuration file.

**fw-monitor-interval** (default: "15s") (type: string)  
Interval for monitoring firewall rules.

**fw-queue-bypass?** (default: #t) (type: boolean)  
Whether to bypass the queue when the daemon is not running.

**rules-path** (default: "/etc/opensnitchd/rules/") (type: string)  
Directory where firewall rules are stored.

**rules-enable-checksums?** (default: #f) (type: boolean)  
Whether to enable checksums for rules.

**ebpf-events-workers** (default: 8) (type: integer)  
Number of eBPF event worker threads.

**ebpf-queue-events-size** (default: 0) (type: integer)  
Size of the eBPF events queue (0 = default).

**stats-max-events** (default: 250) (type: integer)  
Maximum number of events to keep in statistics.

**stats-max-stats** (default: 25) (type: integer)  
Maximum number of statistics entries.

**stats-workers** (default: 6) (type: integer)  
Number of statistics worker threads.

**internal-gc-percent** (default: 100) (type: integer)  
Go garbage collector percentage.

**internal-flush-conns-on-start?** (default: #t) (type: boolean)  
Whether to flush existing connections on daemon start.

## Resize File System Service

This service type lets you resize a live file system during boot, which can be convenient if a Guix image is flashed on an SD Card (e.g. for an embedded device) or uploaded to a VPS. In both cases the medium the image will reside upon may be larger than the image you want to produce.

For an embedded device booting from an SD card you may use something like:

```
(service resize-file-system-service-type
  (resize-file-system-configuration
    (file-system
      (file-system (device (file-system-label "root"))
        (type "ext4")))))
```

**Warning:** Be extra cautious to use the correct device and type. The service has little error handling of its own and relies on the underlying tools. Wrong use could end in loss of data or the corruption of the operating system.

Partitions and file systems are grown to the maximum size available. File systems can only grow when they are on the last partition on a device and have empty space available.

This service supports the ext2, ext3, ext4, btrfs, and bcachefs file systems.

**file-system** (type: file-system)

The file-system object to resize (see Section 11.4 [File Systems], page 257). This object must have the **device** and **type** fields set. Other fields are ignored.

**cloud-utils** (default: cloud-utils) (type: file-like)

The cloud-utils package to use. This package is used for the **growpart** command.

**e2fsprogs** (default: e2fsprogs) (type: file-like)

The e2fsprogs package to use, used for resizing ext2, ext3, and ext4 file systems.

**btrfs-progs** (default: btrfs-progs) (type: file-like)

The btrfs-progs package to use, used for resizing the btrfs file system.

**bcachefs-tools** (default: bcachefs-tools) (type: file-like)

The bcachefs-tools package to use, used for resizing the bcachefs file system.

## Backup Services

The (gnu services backup) module offers services for backing up file system trees. For now, it provides the **restic-backup-service-type**.

With **restic-backup-service-type**, you can periodically back up directories and files with Restic (<https://restic.net/>), which supports end-to-end encryption and deduplication. Consider the following configuration:

```
(use-service-modules backup ...) ;for 'restic-backup-service-type'
(use-package-modules sync ...)   ;for 'rclone'

(operating-system
  ;; ...
  (packages (append (list rclone) ;for use by restic
    %base-packages))
```

```
(services
  (list
    (service restic-backup-service-type
      (restic-backup-configuration
        (jobs
          (list (restic-backup-job
            (name "remote-ftp")
            (repository "rclone:remote-ftp:backup/restic")
            (password-file "/root/.restic")
            ;; Every day at 23.
            (schedule "0 23 * * *")
            (files '("/root/.restic"
              "/root/.config/rclone"
              "/etc/ssh/ssh_host_rsa_key"
              "/etc/ssh/ssh_host_rsa_key.pub"
              "/etc/guix/signing-key.pub"
              "/etc/guix/signing-key.sec")))))))))))
```

Each `restic-backup-job` translates to a Shepherd timer which sets the `RESTIC_PASSWORD` environment variable by reading the first line of `password-file` and runs `restic backup`, creating backups using `rclone` of all the files listed in the `files` field.

The `restic-backup-service-type` provides the ability to instantaneously trigger a backup with the `trigger` Shepherd action:

```
sudo herd trigger remote-ftp
```

**restic-backup-configuration** [Data Type]

Available `restic-backup-configuration` fields are:

**jobs** (default: '()) (type: list-of-restic-backup-jobs)

The list of backup jobs for the current system.

**restic-backup-job** [Data Type]

Available `restic-backup-job` fields are:

**restic** (default: `restic`) (type: package)

The `restic` package to be used for the current job.

**user** (default: `"root"`) (type: string)

The user used for running the current job.

**group** (default: `"root"`) (type: string)

The group used for running the current job.

**log-file** (type: maybe-string)

The file system path to the log file for this job. By default the file will have be `/var/log/restic-backup/job-name.log`, where *job-name* is the name defined in the `name` field. For Guix Home services it defaults to `$XDG_STATE_HOME/shepherd/restic-backup/job-name.log`.

**max-duration** (type: maybe-number)

The maximum duration in seconds that a job may last. Past `max-duration` seconds, the job is forcefully terminated.

**wait-for-termination?** (default: **#f**) (type: boolean)  
 Wait until the job has finished before considering executing it again; otherwise, perform it strictly on every occurrence of event, at the risk of having multiple instances running concurrently.

**repository** (type: string)  
 The restic repository target of this job.

**name** (type: string)  
 A string denoting a name for this job.

**password-file** (type: string)  
 Name of the password file, readable by the configured **user**, that will be used to set the **RESTIC\_PASSWORD** environment variable for the current job.

**schedule** (type: gexp-or-string)  
 A string or a gexp representing the frequency of the backup. Gexp must evaluate to **calendar-event** records or to strings. Strings must contain Vixie cron date lines.

**requirement** (type: maybe-list-of-symbols)  
 The list of Shepherd services that this backup job depends upon. When unset it defaults to '()', for Guix Home. Otherwise to '(**user-processes file-systems**)'.

**files** (default: '()') (type: list-of-lowerables)  
 The list of files or directories to be backed up. It must be a list of values that can be lowered to strings.

**verbose?** (default: **#f**) (type: boolean)  
 Whether to enable verbose output for the current backup job.

**extra-flags** (default: '()') (type: list-of-lowerables)  
 A list of values that are lowered to strings. These will be passed as command-line arguments to the current job **restic backup** invocation.

## DLNA/UPnP Services

The (**gnu services upnp**) module offers services related to UPnP (Universal Plug and Play) and DLNA (Digital Living Network Alliance), networking protocols that can be used for media streaming and device interoperability within a local network. For now, this module provides the **readymedia-service-type**.

ReadyMedia (<https://minidlna.sourceforge.net/>) (formerly known as MiniDLNA) is a DLNA/UPnP-AV media server. The project's daemon, **minidlnad**, can serve media files (audio, pictures, and video) to DLNA/UPnP-AV clients available on the network. **readymedia-service-type** is a Guix service that wraps around ReadyMedia's **minidlnad**.

Consider the following configuration:

```
(use-service-modules upnp ...)
```

```
(operating-system
```



```

...
(services
  (list (service readymedia-service-type
    (readymedia-configuration
      (media-directories
        (list (readymedia-media-directory
          (path "/media/audio")
          (types '(A)))
          (readymedia-media-directory
            (path "/media/video")
            (types '(V)))
          (readymedia-media-directory
            (path "/media/misc")))))
        (extra-config '(("notify_interval" . "60")))))
    ...)))

```

This sets up the ReadyMedia daemon to serve files from the media folders specified in `media-directories`. The `media-directories` field is mandatory. All other fields (such as network ports and the server name) come with a predefined default and can be omitted.

**Note:** This service is also available for Guix Home, where it runs directly with your user privileges (see Section 13.3.21 [Miscellaneous Home Services], page 781).

`readymedia-configuration` [Data Type]

Available `readymedia-configuration` fields are:

`readymedia` (default: `readymedia`) (type: package)

The ReadyMedia package to be used for the service.

`friendly-name` (default: `#f`) (type: maybe-string)

A custom name that will be displayed on connected clients.

`media-directories` (type: list)

The list of media folders to serve content from. Each item is a `readymedia-media-directory`.

`cache-directory` (default: `"/var/cache/readymedia"`) (type: string)

A folder for ReadyMedia's cache files. If not existing already, the folder will be created as part of the service activation and the ReadyMedia user will be assigned ownership.

`log-directory` (default: `"/var/log/readymedia"`) (type: string)

A folder for ReadyMedia's log files. If not existing already, the folder will be created as part of the service activation and the ReadyMedia user will be assigned ownership.

`port` (default: `#f`) (type: maybe-integer)

A custom port that the service will be listening on.

`extra-config` (default: `'()`) (type: alist)

An association list of further options, as accepted by ReadyMedia.

**readymedia-media-directory** [Data Type]

A **media-directories** entry includes a folder **path** and, optionally, the **types** of media files included within the folder.

**path** (type: string)

The media folder location.

**types** (default: '()') (type: list)

A list indicating the types of file included in the media folder. Valid values are combinations of individual media types, i.e. symbol **A** for audio, **P** for pictures, **V** for video. An empty list means that no type is specified.

## 11.11 Privileged Programs

Some programs need to run with elevated privileges, even when they are launched by unprivileged users. A notorious example is the **passwd** program, which users can run to change their password, and which needs to access the **/etc/passwd** and **/etc/shadow** files—something normally restricted to root, for obvious security reasons. To address that, **passwd** should be *setuid-root*, meaning that it always runs with root privileges (see Section “How Change Persona” in *The GNU C Library Reference Manual*, for more info about the *setuid* mechanism).

The store itself *cannot* contain privileged programs: that would be a security issue since any user on the system can write derivations that populate the store (see Section 8.9 [The Store], page 156). Thus, a different mechanism is used: instead of directly granting permissions to files that are in the store, we let the system administrator *declare* which programs should be entrusted with these additional privileges.

The **privileged-programs** field of an **operating-system** declaration contains a list of **<privileged-program>** denoting the names of programs to have a *setuid* or *setgid* bit set (see Section 11.2 [Using the Configuration System], page 244). For instance, the **mount.nfs** program, which is part of the **nfs-utils** package, with a *setuid* root can be designated like this:

```
(privileged-program
  (program (file-append nfs-utils "/sbin/mount.nfs"))
  (setuid? #t))
```

And then, to make **mount.nfs** *setuid* on your system, add the previous example to your operating system declaration by appending it to **%default-privileged-programs** like this:

```
(operating-system
  ;; Some fields omitted...
  (privileged-programs
    (append (list (privileged-program
                  (program (file-append nfs-utils "/sbin/mount.nfs"))
                  (setuid? #t)))
            %default-privileged-programs)))
```

**privileged-program** [Data Type]

This data type represents a program with special privileges, such as *setuid*

**program** A file-like object to which all given privileges should apply.

**setuid?** (default: **#f**)  
Whether to set user setuid bit.

**setgid?** (default: **#f**)  
Whether to set group setgid bit.

**user** (default: 0)  
UID (integer) or user name (string) for the user owner of the program, defaults to root.

**group** (default: 0)  
GID (integer) group name (string) for the group owner of the program, defaults to root.

**capabilities** (default: **#f**)  
A string representing the program's POSIX capabilities, as described by the `cap_to_text(3)` man page from the libcap package, or **#f** to make no changes.

A default set of privileged programs is defined by the `%default-privileged-programs` variable of the `(gnu system)` module.

**Scheme Variable** `%default-privileged-programs` [Variable]  
A list of `<privileged-program>` denoting common programs with elevated privileges.  
The list includes commands such as `passwd`, `ping`, `su`, and `sudo`.

Under the hood, the actual privileged programs are created in the `/run/privileged/bin` directory at system activation time. The files in this directory refer to the “real” binaries, which are in the store.

## 11.12 X.509 Certificates

Web servers available over HTTPS (that is, HTTP over the transport-layer security mechanism, TLS) send client programs an *X.509 certificate* that the client can then use to *authenticate* the server. To do that, clients verify that the server's certificate is signed by a so-called *certificate authority* (CA). But to verify the CA's signature, clients must have first acquired the CA's certificate.

Web browsers such as GNU IceCat include their own set of CA certificates, such that they are able to verify CA signatures out-of-the-box.

However, most other programs that can talk HTTPS—`wget`, `git`, `w3m`, etc.—need to be told where CA certificates can be found.

For users of Guix System, this is done by adding a package that provides certificates to the `packages` field of the `operating-system` declaration (see Section 11.3 [operating-system Reference], page 253). Guix includes one such package, `nss-certs`, which is a set of CA certificates provided as part of Mozilla's Network Security Services.

This package is part of `%base-packages`, so there is no need to explicitly add it. The `/etc/ssl/certs` directory, which is where most applications and libraries look for certificates by default, points to the certificates installed globally.

Unprivileged users, including users of Guix on a foreign distro, can also install their own certificate package in their profile. A number of environment variables need to be defined

so that applications and libraries know where to find them. Namely, the OpenSSL library honors the `SSL_CERT_DIR` and `SSL_CERT_FILE` variables. Some applications add their own environment variables; for instance, the Git version control system honors the certificate bundle pointed to by the `GIT_SSL_CAINFO` environment variable. Thus, you would typically run something like:

```
guix install nss-certs
export SSL_CERT_DIR="$HOME/.guix-profile/etc/ssl/certs"
export SSL_CERT_FILE="$HOME/.guix-profile/etc/ssl/certs/ca-certificates.crt"
export GIT_SSL_CAINFO="$SSL_CERT_FILE"
```

As another example, R requires the `CURL_CA_BUNDLE` environment variable to point to a certificate bundle, so you would have to run something like this:

```
guix install nss-certs
export CURL_CA_BUNDLE="$HOME/.guix-profile/etc/ssl/certs/ca-certificates.crt"
```

For other applications you may want to look up the required environment variable in the relevant documentation.

### 11.13 Name Service Switch

The (`gnu system nss`) module provides bindings to the configuration file of the libc *name service switch* or *NSS* (see Section “NSS Configuration File” in *The GNU C Library Reference Manual*). In a nutshell, the NSS is a mechanism that allows libc to be extended with new “name” lookup methods for system databases, which includes host names, service names, user accounts, and more (see Section “Name Service Switch” in *The GNU C Library Reference Manual*).

The NSS configuration specifies, for each system database, which lookup method is to be used, and how the various methods are chained together—for instance, under which circumstances NSS should try the next method in the list. The NSS configuration is given in the `name-service-switch` field of `operating-system` declarations (see Section 11.3 [operating-system Reference], page 253).

As an example, the declaration below configures the NSS to use the `nss-mdns` back-end (<https://0pointer.de/lennart/projects/nss-mdns/>), which supports host name lookups over multicast DNS (mDNS) for host names ending in `.local`:

```
(name-service-switch
 (hosts (list %files      ;first, check /etc/hosts

;; If the above did not succeed, try
;; with 'mdns_minimal'.
(name-service
 (name "mdns_minimal")

;; 'mdns_minimal' is authoritative for
;; '.local'. When it returns "not found",
;; no need to try the next methods.
(reaction (lookup-specification
            (not-found => return))))
```

```
;; Then fall back to DNS.
(name-service
  (name "dns"))

;; Finally, try with the "full" 'mdns'.
(name-service
  (name "mdns"))))
```

Do not worry: the `%mdns-host-lookup-nss` variable (see below) contains this configuration, so you will not have to type it if all you want is to have `.local` host lookup working.

Note that, in this case, in addition to setting the `name-service-switch` of the `operating-system` declaration, you also need to use `avahi-service-type` (see Section 11.10.5 [Networking Services], page 322), or `%desktop-services`, which includes it (see Section 11.10.9 [Desktop Services], page 377). Doing this makes `nss-mdns` accessible to the name service cache daemon (see Section 11.10.1 [Base Services], page 276).

For convenience, the following variables provide typical NSS configurations.

**%default-nss** [Variable]  
This is the default name service switch configuration, a `name-service-switch` object.

**%mdns-host-lookup-nss** [Variable]  
This is the name service switch configuration with support for host name lookup over multicast DNS (mDNS) for host names ending in `.local`.

Furthermore, mDNS lookups could be restricted to IPv4 or IPv6 protocols using the procedure `mdns-host-lookup-nss`.

**mdns-host-lookup-nss** `[#:ipv6? #t] [#:ipv4? #t]` [Procedure]  
Return a `name-service-switch` object with mDNS lookup restricted to specified protocols.

The reference for name service switch configuration is given below. It is a direct mapping of the configuration file format of the C library, so please refer to the C library manual for more information (see Section “NSS Configuration File” in *The GNU C Library Reference Manual*). Compared to the configuration file format of libc NSS, it has the advantage not only of adding this warm parenthetic feel that we like, but also static checks: you will know about syntax errors and typos as soon as you run `guix system`.

**name-service-switch** [Data Type]  
This is the data type representation the configuration of libc’s name service switch (NSS). Each field below represents one of the supported system databases.

```

aliases
ethers
group
gshadow
hosts
initgroups
netgroup
networks
password
public-key
rpc
services
shadow    The system databases handled by the NSS. Each of these fields must be
           a list of <name-service> objects (see below).

```

**name-service** [Data Type]

This is the data type representing an actual name service and the associated lookup action.

**name** A string denoting the name service (see Section “Services in the NSS configuration” in *The GNU C Library Reference Manual*).

Note that name services listed here must be visible to `nscd`. This is achieved by passing the `#:name-services` argument to `nscd-service` the list of packages providing the needed name services (see Section 11.10.1 [Base Services], page 276).

**reaction** An action specified using the `lookup-specification` macro (see Section “Actions in the NSS configuration” in *The GNU C Library Reference Manual*). For example:

```

(lookup-specification (unavailable => continue)
 (success => return))

```

## 11.14 Initial RAM Disk

For bootstrapping purposes, the Linux-Libre kernel is passed an *initial RAM disk*, or *initrd*. An *initrd* contains a temporary root file system as well as an initialization script. The latter is responsible for mounting the real root file system, and for loading any kernel modules that may be needed to achieve that.

The `initrd-modules` field of an `operating-system` declaration allows you to specify Linux-libre kernel modules that must be available in the *initrd*. In particular, this is where you would list modules needed to actually drive the hard disk where your root partition is—although the default value of `initrd-modules` should cover most use cases. For example, assuming you need the `megaraid_sas` module in addition to the default modules to be able to access your root file system, you would write:

```

(operating-system
 ;; ...
 (initrd-modules (cons "megaraid_sas" %base-initrd-modules)))

```

**%base-initrd-modules**

[Variable]

This is the list of kernel modules included in the `initrd` by default.

Furthermore, if you need lower-level customization, the `initrd` field of an `operating-system` declaration allows you to specify which `initrd` you would like to use. The (`gnu system linux-initrd`) module provides three ways to build an `initrd`: the high-level `base-initrd` procedure and the low-level `raw-initrd` and `expression->initrd` procedures.

The `base-initrd` procedure is intended to cover most common uses. For example, if you want to add a bunch of kernel modules to be loaded at boot time, you can define the `initrd` field of the operating system declaration like this:

```
(initrd (lambda (file-systems . rest)
  ;; Create a standard initrd but set up networking
  ;; with the parameters QEMU expects by default.
  (apply base-initrd file-systems
    #:qemu-networking? #t
    rest)))
```

The `base-initrd` procedure also handles common use cases that involves using the system as a QEMU guest, or as a “live” system with volatile root file system.

The `base-initrd` procedure is built from `raw-initrd` procedure. Unlike `base-initrd`, `raw-initrd` doesn’t do anything high-level, such as trying to guess which kernel modules and packages should be included to the `initrd`. An example use of `raw-initrd` is when a user has a custom Linux kernel configuration and default kernel modules included by `base-initrd` are not available.

The initial RAM disk produced by `base-initrd` or `raw-initrd` honors several options passed on the Linux kernel command line (that is, arguments passed *via* the `linux` command of GRUB, or the `-append` option of QEMU), notably:

**gnu.load=boot**

Tell the initial RAM disk to load *boot*, a file containing a Scheme program, once it has mounted the root file system.

Guix uses this option to yield control to a boot program that runs the service activation programs and then spawns the GNU Shepherd, the initialization system.

**root=root**

Mount *root* as the root file system. *root* can be a device name like `/dev/sda1`, a file system label, or a file system UUID. When unspecified, the device name from the root file system of the operating system declaration is used.

**rootfstype=type**

Set the type of the root file system. It overrides the `type` field of the root file system specified via the `operating-system` declaration, if any.

**rootflags=options**

Set the mount *options* of the root file system. It overrides the `options` field of the root file system specified via the `operating-system` declaration, if any.

**fsck.mode=mode**

Whether to check the *root* file system for errors before mounting it. *mode* is one of **skip** (never check), **force** (always check), or **auto** to respect the root <file-system> object's **check?** setting (see Section 11.4 [File Systems], page 257) and run a full scan only if the file system was not cleanly shut down.

**auto** is the default if this option is not present or if *mode* is not one of the above.

**fsck.repair=level**

The level of repairs to perform automatically if errors are found in the *root* file system. *level* is one of **no** (do not write to *root* at all if possible), **yes** (repair as much as possible), or **preen** to repair problems considered safe to repair automatically.

**preen** is the default if this option is not present or if *level* is not one of the above.

**gnu.system=system**

Have */run/booted-system* and */run/current-system* point to *system*.

**modprobe.blacklist=modules...**

Instruct the initial RAM disk as well as the **modprobe** command (from the **kmod** package) to refuse to load *modules*. *modules* must be a comma-separated list of module names—e.g., **usbkbd,9pnet**.

**gnu.repl** Start a read-eval-print loop (REPL) from the initial RAM disk before it tries to load kernel modules and to mount the root file system. Our marketing team calls it *boot-to-Guile*. The Schemer in you will love it. See Section “Using Guile Interactively” in *GNU Guile Reference Manual*, for more information on Guile's REPL.

Now that you know all the features that initial RAM disks produced by **base-initrd** and **raw-initrd** provide, here is how to use it and customize it further.

```
raw-initrd file-systems [#:linux-modules '()'] [#:pre-mount #t] [Procedure]
            [#:mapped-devices '()'] [#:keyboard-layout #f] [#:helper-packages '()']
            [#:qemu-networking? #f] [#:volatile-root? #f]
```

Return a derivation that builds a raw **initrd**. *file-systems* is a list of file systems to be mounted by the **initrd**, possibly in addition to the root file system specified on the kernel command line via **root**. *linux-modules* is a list of kernel modules to be loaded at boot time. *mapped-devices* is a list of device mappings to realize before *file-systems* are mounted (see Section 11.5 [Mapped Devices], page 263). *pre-mount* is a G-expression to evaluate before realizing *mapped-devices*. *helper-packages* is a list of packages to be copied in the **initrd**. It may include **e2fsck/static** or other packages needed by the **initrd** to check the root file system.

When true, *keyboard-layout* is a <keyboard-layout> record denoting the desired console keyboard layout. This is done before *mapped-devices* are set up and before *file-systems* are mounted such that, should the user need to enter a passphrase or use the REPL, this happens using the intended keyboard layout.

When *qemu-networking?* is true, set up networking with the standard QEMU parameters.



When *volatile-root?* is true, the root file system is writable but any changes to it are lost.

```
base-initrd file-systems [#:mapped-devices '()] [Procedure]
      [#:keyboard-layout #f] [#:qemu-networking? #f] [#:volatile-root? #f]
      [#:linux-modules '()]
```

Return as a file-like object a generic initrd, with kernel modules taken from *linux*. *file-systems* is a list of file-systems to be mounted by the initrd, possibly in addition to the root file system specified on the kernel command line via **root**. *mapped-devices* is a list of device mappings to realize before *file-systems* are mounted.

When true, *keyboard-layout* is a <**keyboard-layout**> record denoting the desired console keyboard layout. This is done before *mapped-devices* are set up and before *file-systems* are mounted such that, should the user need to enter a passphrase or use the REPL, this happens using the intended keyboard layout.

*qemu-networking?* and *volatile-root?* behaves as in **raw-initrd**.

The initrd is automatically populated with all the kernel modules necessary for *file-systems* and for the given options. Additional kernel modules can be listed in *linux-modules*. They will be added to the initrd, and loaded at boot time in the order in which they appear.

Needless to say, the initrds we produce and use embed a statically-linked Guile, and the initialization program is a Guile program. That gives a lot of flexibility. The **expression->initrd** procedure builds such an initrd, given the program to run in that initrd.

```
expression->initrd exp [#:guile %guile-static-initrd] [#:name [Procedure]
      "guile-initrd"]
```

Return as a file-like object a Linux initrd (a gzipped cpio archive) containing *guile* and that evaluates *exp*, a G-expression, upon booting. All the derivations referenced by *exp* are automatically copied to the initrd.

## 11.15 Bootloader Configuration

The operating system supports multiple bootloaders. The bootloader is configured using **bootloader-configuration** declaration. All the fields of this structure are bootloader agnostic except for one field, **bootloader** that indicates the bootloader to be configured and installed.

Some of the bootloaders do not honor every field of **bootloader-configuration**. For instance, the extlinux bootloader does not support themes and thus ignores the **theme** field.

```
bootloader-configuration [Data Type]
```

The type of a bootloader configuration declaration.

```
bootloader
```

The bootloader to use, as a **bootloader** object. For now **grub-bootloader**, **grub-efi-bootloader**, **grub-efi-removable-bootloader**, **grub-efi-netboot-bootloader**, **grub-efi-netboot-removable-bootloader**, **extlinux-bootloader** and **u-boot-bootloader** are supported.

Available bootloaders are described in (`gnu bootloader ...`) modules. In particular, (`gnu bootloader u-boot`) contains definitions of bootloaders for a wide range of ARM and AArch64 systems, using the U-Boot bootloader (<https://www.denx.de/wiki/U-Boot/>).

`grub-bootloader` allows you to boot in particular Intel-based machines in “legacy” BIOS mode.

`grub-efi-bootloader` allows to boot on modern systems using the *Unified Extensible Firmware Interface* (UEFI). This is what you should use if the installation image contains a `/sys/firmware/efi` directory when you boot it on your system.

`grub-efi-removable-bootloader` allows you to boot your system from removable media by writing the GRUB file to the UEFI-specification location of `/EFI/BOOT/BOOTX64.efi` of the boot directory, usually `/boot/efi`. This is also useful for some UEFI firmwares that “forget” their configuration from their non-volatile storage. Like `grub-efi-bootloader`, this can only be used if the `/sys/firmware/efi` directory is available.

**Note:** This *will* overwrite the GRUB file from any other operating systems that also place their GRUB file in the UEFI-specification location; making them unbootable.

`grub-efi-netboot-bootloader` allows you to boot your system over network through TFTP. In combination with an NFS root file system this allows you to build a diskless Guix system.

The installation of the `grub-efi-netboot-bootloader` generates the content of the TFTP root directory at `targets` (see Section 11.15 [Bootloader Configuration], page 685) below the sub-directory `efi/Guix`, to be served by a TFTP server. You may want to mount your TFTP server directories onto the `targets` to move the required files to the TFTP server automatically during installation.

If you plan to use an NFS root file system as well (actually if you mount the store from an NFS share), then the TFTP server needs to serve the file `/boot/grub/grub.cfg` and other files from the store (like GRUBs background image, the kernel (see Section 11.3 [operating-system Reference], page 253) and the `initrd` (see Section 11.3 [operating-system Reference], page 253)), too. All these files from the store will be accessed by GRUB through TFTP with their normal store path, for example as `tftp://tftp-server/gnu/store/...-initrd/initrd.cpio.gz`.

Two symlinks are created to make this possible. For each target in the `targets` field, the first symlink is `‘target’/efi/Guix/boot/grub/grub.cfg` pointing to `../.../boot/grub/grub.cfg`, where `‘target’` may be `/boot`. In this case the link is not leaving the served TFTP root directory, but otherwise it does. The second link is `‘target’/gnu/store` and points to `../gnu/store`. This link is leaving the served TFTP root directory.

The assumption behind all this is that you have an NFS server exporting the root file system for your Guix system, and additionally a TFTP server exporting your **targets** directories—usually a single **/boot**—from that same root file system for your Guix system. In this constellation the symlinks will work.

For other constellations you will have to program your own bootloader installer, which then takes care to make necessary files from the store accessible through TFTP, for example by copying them into the TFTP root directory for your **targets**.

It is important to note that symlinks pointing outside the TFTP root directory may need to be allowed in the configuration of your TFTP server. Further the store link exposes the whole store through TFTP. Both points need to be considered carefully for security aspects. It is advised to disable any TFTP write access!

Please note, that this bootloader will not modify the ‘UEFI Boot Manager’ of the system.

Beside the **grub-efi-netboot-bootloader**, the already mentioned TFTP and NFS servers, you also need a properly configured DHCP server to make the booting over netboot possible. For all this we can currently only recommend you to look for instructions about PXE (Preboot eXecution Environment).

If a local EFI System Partition (ESP) or a similar partition with a FAT file system is mounted in **targets**, then symlinks cannot be created. In this case everything will be prepared for booting from local storage, matching the behavior of **grub-efi-bootloader**, with the difference that all GRUB binaries are copied to **targets**, necessary for booting over the network.

**grub-efi-netboot-removable-bootloader** is identical to **grub-efi-netboot-bootloader** with the exception that the sub-directory **efi/boot** will be used instead of **efi/Guix** to comply with the UEFI specification for removable media.

**Note:** This *will* overwrite the GRUB file from any other operating systems that also place their GRUB file in the UEFI-specification location; making them unbootable.

**targets** This is a list of strings denoting the targets onto which to install the bootloader.

The interpretation of targets depends on the bootloader in question. For **grub-bootloader**, for example, they should be device names understood by the bootloader **installer** command, such as **/dev/sda** or **(hd0)** (see Section “Invoking grub-install” in *GNU GRUB Manual*). For **grub-efi-bootloader** and **grub-efi-removable-bootloader** they should be mount points of the EFI file system, usually **/boot/efi**. For **grub-efi-netboot-bootloader**, **targets** should be the mount points corresponding to TFTP root directories served by your TFTP server.

**menu-entries** (default: '()')

A possibly empty list of **menu-entry** objects (see below), denoting entries to appear in the bootloader menu, in addition to the current system entry and the entry pointing to previous system generations.

**default-entry** (default: 0)

The index of the default boot menu entry. Index 0 is for the entry of the current system.

**timeout** (default: 5)

The number of seconds to wait for keyboard input before booting. Set to 0 to boot immediately, and to -1 to wait indefinitely.

**keyboard-layout** (default: #f)

If this is #f, the bootloader's menu (if any) uses the default keyboard layout, usually US English ("qwerty").

Otherwise, this must be a **keyboard-layout** object (see Section 11.8 [Keyboard Layout], page 271).

**Note:** This option is currently ignored by bootloaders other than **grub** and **grub-efi**.

**theme** (default: #f)

The bootloader theme object describing the theme to use. If no theme is provided, some bootloaders might use a default theme, that's true for GRUB.

**terminal-outputs** (default: '(gfxterm)')

The output terminals used for the bootloader boot menu, as a list of symbols. GRUB accepts the values: **console**, **serial**, **serial\_{0-3}**, **gfxterm**, **vga\_text**, **mda\_text**, **morse**, and **pkmodem**. This field corresponds to the GRUB variable **GRUB\_TERMINAL\_OUTPUT** (see Section "Simple configuration" in *GNU GRUB manual*).

**terminal-inputs** (default: '()')

The input terminals used for the bootloader boot menu, as a list of symbols. For GRUB, the default is the native platform terminal as determined at run-time. GRUB accepts the values: **console**, **serial**, **serial\_{0-3}**, **at\_keyboard**, and **usb\_keyboard**. This field corresponds to the GRUB variable **GRUB\_TERMINAL\_INPUT** (see Section "Simple configuration" in *GNU GRUB manual*).

**serial-unit** (default: #f)

The serial unit used by the bootloader, as an integer from 0 to 3. For GRUB, it is chosen at run-time; currently GRUB chooses 0, which corresponds to COM1 (see Section "Serial terminal" in *GNU GRUB manual*).

**serial-speed** (default: #f)

The speed of the serial interface, as an integer. For GRUB, the default value is chosen at run-time; currently GRUB chooses 9600 bps (see Section "Serial terminal" in *GNU GRUB manual*).

**device-tree-support?** (default: **#t**)

Whether to support Linux device tree (<https://en.wikipedia.org/wiki/Devicetree>) files loading.

This option is enabled by default. In some cases involving the **u-boot** bootloader, where the device tree has already been loaded in RAM, it can be handy to disable the option by setting it to **#f**.

**extra-initrd** (default: **#f**)

File name of an additional **initrd** to load during the boot. It may or may not point to a file in the store, but the main use case is for out-of-store files containing secrets.

In order to be able to provide decryption keys for LUKS devices, they need to be available in the initial RAM disk. However they cannot be stored inside the usual **initrd**, since it is stored in the store and being world-readable (as files in the store are) is not a desired property for an **initrd** containing decryption keys. You can therefore use this field to instruct GRUB to also load a manually created **initrd** not stored in the store.

For any use case not involving secrets, you should use regular **initrd** (see Section 11.3 [operating-system Reference], page 253) instead.

Suitable image can be created for example like this:

```
echo /key-file.bin | cpio -oH newc >/key-file.cpio
chmod 0000 /key-file.cpio
```

After it is created, you can use it in this manner:

```
;; Operating system with encrypted boot partition
(operating-system
  ...
  (bootloader (bootloader-configuration
    (bootloader grub-efi-bootloader)
    (targets '("/boot/efi"))
    ;; Load the initrd with a key file...
    (extra-initrd "/key-file.cpio")))
  (mapped-devices
    (list (mapped-device
      (source (uuid "12345678-1234-1234-1234-123456789abc"))
      (target "my-root")
      (type luks-device-mapping)
      ;; ... and use it to unlock the root device.
      (arguments '(:key-file "/key-file.bin"))))))
```

Be careful when using this option, since pointing to a file that is not readable by the grub while booting will cause the boot to fail and require a manual edit of the **initrd** line in the grub menu.

Currently only supported by GRUB.

Should you want to list additional boot menu entries *via* the **menu-entries** field above, you will need to create them with the **menu-entry** form. For example, imagine you want to

be able to boot another distro (hard to imagine!), you can define a menu entry along these lines:

```
(menu-entry
  (label "The Other Distro")
  (linux "/boot/old/vmlinuz-2.6.32")
  (linux-arguments '("root=/dev/sda2"))
  (initrd "/boot/old/initrd"))
```

Details below.

**menu-entry** [Data Type]

The type of an entry in the bootloader menu.

**label**        The label to show in the menu—e.g., "GNU".

**linux** (default: **#f**)

The Linux kernel image to boot, for example:

```
(file-append linux-libre "/bzImage")
```

For GRUB, it is also possible to specify a device explicitly in the file path using GRUB's device naming convention (see Section “Naming convention” in *GNU GRUB manual*), for example:

```
"(hd0,msdos1)/boot/vmlinuz"
```

If the device is specified explicitly as above, then the **device** field is ignored entirely.

**linux-arguments** (default: '() )

The list of extra Linux kernel command-line arguments—e.g., '("console=ttyS0")'.

**initrd** (default: **#f**)

A G-Expression or string denoting the file name of the initial RAM disk to use (see Section 8.12 [G-Expressions], page 165).

**device** (default: **#f**)

The device where the kernel and initrd are to be found—i.e., for GRUB, *root* for this menu entry (see Section “root” in *GNU GRUB manual*).

This may be a file system label (a string), a file system UUID (a bytevector, see Section 11.4 [File Systems], page 257), or **#f**, in which case the bootloader will search the device containing the file specified by the **linux** field (see Section “search” in *GNU GRUB manual*). It must *not* be an OS device name such as */dev/sda1*.

**multiboot-kernel** (default: **#f**)

The kernel to boot in Multiboot-mode (see Section “multiboot” in *GNU GRUB manual*). When this field is set, a Multiboot menu-entry is generated. For example:

```
(file-append mach "/boot/gnumach")
```

**multiboot-arguments** (default: '() )

The list of extra command-line arguments for the multiboot-kernel.

For example, when running in QEMU it can be useful to use a text-based console (use options `--nographic --serial mon:stdio`):

```
'("console=com0")
```

To use the new and still experimental rumpdisk user-level disk driver ([https://darnassus.sceen.net/~hurd-web/rump\\_kernel/](https://darnassus.sceen.net/~hurd-web/rump_kernel/)) instead of GNU Mach's in-kernel IDE driver, set `kernel-arguments` to:

```
'("noide")
```

Of course, these options can be combined:

```
'("console=com0" "noide")
```

`multiboot-modules` (default: `'()`)

The list of commands for loading Multiboot modules. For example:

```
(list (list (file-append hurd "/hurd/ext2fs.static") "ext2fs"
  ...)
  (list (file-append libc "/lib/ld.so.1") "exec"
  ...))
```

`chain-loader` (default: `#f`)

A string that can be accepted by grub's `chainloader` directive. This has no effect if either `linux` or `multiboot-kernel` fields are specified. The following is an example of chainloading a different GNU/Linux system.

```
(bootloader
  (bootloader-configuration
    ;; ...
    (menu-entries
      (list
        (menu-entry
          (label "GNU/Linux")
          (device (uuid "1C31-A17C" 'fat))
          (chain-loader "/EFI/GNULinux/grubx64.efi"))))))
```

For now only GRUB has theme support. GRUB themes are created using the `grub-theme` form, which is not fully documented yet.

**grub-theme** [Data Type]

Data type representing the configuration of the GRUB theme.

`gfxmode` (default: `'("auto")`)

The GRUB `gfxmode` to set (a list of screen resolution strings, see Section “`gfxmode`” in *GNU GRUB manual*).

**grub-theme** [Procedure]

Return the default GRUB theme used by the operating system if no `theme` field is specified in `bootloader-configuration` record.

It comes with a fancy background image displaying the GNU and Guix logos.

For example, to override the default resolution, you may use something like

```
(bootloader
```

```
(bootloader-configuration
;; ...
(theme (grub-theme
        (inherit (grub-theme))
        (gfxmode '("1024x786x32" "auto"))))))
```

## 11.16 Invoking guix system

Once you have written an operating system declaration as seen in the previous section, it can be *instantiated* using the `guix system` command. The synopsis is:

```
guix system options... action file
```

*file* must be the name of a file containing an `operating-system` declaration. *action* specifies how the operating system is instantiated. Currently the following values are supported:

**search**     Display available service type definitions that match the given regular expressions, sorted by relevance:

```
$ guix system search console
name: console-fonts
location: gnu/services/base.scm:806:2
extends: shepherd-root
description: Install the given fonts on the specified ttys (fonts are per
+ virtual console on GNU/Linux). The value of this service is a list of
+ tty/font pairs. The font can be the name of a font provided by the `kbd'
+ package or any valid argument to `setfont', as in this example:
+
+      '(("tty1" . "LatGrkCyr-8x16")
+        ("tty2" . (file-append
+                   font-tamzen
+                   "/share/kbd/consolefonts/TamzenForPowerline10x20.psf"
+                   ("tty3" . (file-append
+                              font-terminus
+                              "/share/consolefonts/ter-132n")))) ; for HDPI
relevance: 9

name: mingetty
location: gnu/services/base.scm:1190:2
extends: shepherd-root
description: Provide console login using the `mingetty' program.
relevance: 2

name: login
location: gnu/services/base.scm:860:2
extends: pam
description: Provide a console log-in service as specified by its
+ configuration value, a `login-configuration' object.
relevance: 2
```



...

As for `guix package --search`, the result is written in `recutils` format, which makes it easy to filter the output (see *GNU recutils manual*).

**edit** Edit or view the definition of the given service types.

For example, the command below opens your editor, as specified by the `EDITOR` environment variable, on the definition of the `openssh` service type:

```
guix system edit openssh
```

**reconfigure**

Build the operating system described in *file*, activate it, and switch to it<sup>11</sup>.

**Note:** It is highly recommended to run `guix pull` once before you run `guix system reconfigure` for the first time (see Section 5.7 [Invoking `guix pull`], page 56). Failing to do that you would see an older version of Guix once `reconfigure` has completed.

This effects all the configuration specified in *file*: user accounts, system services, global package list, privileged programs, etc. The command starts system services specified in *file* that are not currently running; if a service is currently running this command will arrange for it to be upgraded the next time it is stopped (e.g. by `herd stop X` or `herd restart X`).

This command creates a new generation whose number is one greater than the current generation (as reported by `guix system list-generations`). If that generation already exists, it will be overwritten. This behavior mirrors that of `guix package` (see Section 5.2 [Invoking `guix package`], page 37).

It adds a bootloader menu entry for the new OS configuration, —unless `--no-bootloader` is passed. For GRUB, it moves entries for older configurations to a submenu, allowing you to choose an older system generation at boot time should you need it.

On Linux, `guix system reconfigure` also loads the new system for fast reboot *via* `kexec`: running `reboot --kexec` will boot the new system by directly executing its kernel, thus bypassing the BIOS initialization phase and bootloader (see Section “Invoking reboot” in *The GNU Shepherd Manual*). You can avoid this behavior by passing the `--no-kexec` option.

Upon completion, the new system is deployed under `/run/current-system`. This directory contains *provenance meta-data*: the list of channels in use (see Chapter 6 [Channels], page 68) and *file* itself, when available. You can view it by running:

```
guix system describe
```

This information is useful should you later want to inspect how this particular generation was built. In fact, assuming *file* is self-contained, you can later rebuild generation *n* of your operating system with:

```
guix time-machine \
```

<sup>11</sup> This action (and the related actions `switch-generation` and `roll-back`) are usable only on systems already running Guix System.

```
-C /var/guix/profiles/system-n-link/channels.scm -- \
system reconfigure \
/var/guix/profiles/system-n-link/configuration.scm
```

You can think of it as some sort of built-in version control! Your system is not just a binary artifact: *it carries its own source*. See Section 11.19.3 [Service Reference], page 712, for more information on provenance tracking.

By default, **reconfigure** *prevents you from downgrading your system*, which could (re)introduce security vulnerabilities and also cause problems with “stateful” services such as database management systems. You can override that behavior by passing **--allow-downgrades**.

### switch-generation

Switch to an existing system generation. This action atomically switches the system profile to the specified system generation. It also rearranges the system’s existing bootloader menu entries. It makes the menu entry for the specified system generation the default, and it moves the entries for the other generations to a submenu, if supported by the bootloader being used. The next time the system boots, it will use the specified system generation.

The bootloader itself is not being reinstalled when using this command. Thus, the installed bootloader is used with an updated configuration file.

The target generation can be specified explicitly by its generation number. For example, the following invocation would switch to system generation 7:

```
guix system switch-generation 7
```

The target generation can also be specified relative to the current generation with the form **+N** or **-N**, where **+3** means “3 generations ahead of the current generation,” and **-1** means “1 generation prior to the current generation.” When specifying a negative value such as **-1**, you must precede it with **--** to prevent it from being parsed as an option. For example:

```
guix system switch-generation -- -1
```

Currently, the effect of invoking this action is *only* to switch the system profile to an existing generation and rearrange the bootloader menu entries. To actually start using the target system generation, you must reboot after running this action. In the future, it will be updated to do the same things as **reconfigure**, like activating and deactivating services.

This action will fail if the specified generation does not exist.

### roll-back

Switch to the preceding system generation. The next time the system boots, it will use the preceding system generation. This is the inverse of **reconfigure**, and it is exactly the same as invoking **switch-generation** with an argument of **-1**.

Currently, as with **switch-generation**, you must reboot after running this action to actually start using the preceding system generation.

**delete-generations**

Delete system generations, making them candidates for garbage collection (see Section 5.6 [Invoking `guix gc`], page 53, for information on how to run the “garbage collector”).

This works in the same way as ‘`guix package --delete-generations`’ (see Section 5.2 [Invoking `guix package`], page 37). With no arguments, all system generations but the current one are deleted:

```
guix system delete-generations
```

You can also select the generations you want to delete. The example below deletes all the system generations that are more than two months old:

```
guix system delete-generations 2m
```

Running this command automatically reinstalls the bootloader with an updated list of menu entries—e.g., the “old generations” sub-menu in GRUB no longer lists the generations that have been deleted.

**build**

Build the derivation of the operating system, which includes all the configuration files and programs needed to boot and run the system. This action does not actually install anything.

**init**

Populate the given directory with all the files necessary to run the operating system specified in *file*. This is useful for first-time installations of Guix System. For instance:

```
guix system init my-os-config.scm /mnt
```

copies to `/mnt` all the store items required by the configuration specified in `my-os-config.scm`. This includes configuration files, packages, and so on. It also creates other essential files needed for the system to operate correctly—e.g., the `/etc`, `/var`, and `/run` directories, and the `/bin/sh` file.

This command also installs bootloader on the targets specified in `my-os-config`, unless the `--no-bootloader` option was passed.

**installer**

Run the installer. Usually the installer is built as an `iso` image, copied to a USB Stick or DVD, and booted from (Section 3.3 [USB Stick and DVD Installation], page 23). If your machine already runs Guix and you still want to run the installer, e.g., for testing purposes, you can skip the step of creating an `iso` and run for instance:

```
guix system installer --dry-run
```

**Note:** If you do not use `--dry-run` then you need to run as root. Be *very careful* when running the installer as root, it can cause data loss or render your system unbootable!

**vm**

Build a virtual machine (VM) that contains the operating system declared in *file*, and return a script to run that VM.

**Note:** The `vm` action and others below can use KVM support in the Linux-libre kernel. Specifically, if the machine has hardware virtualization support, the corresponding KVM kernel module should be loaded, and the `/dev/kvm` device node must exist and be readable

and writable by the user and by the build users of the daemon (see Section 2.2.1 [Build Environment Setup], page 6).

Arguments given to the script are passed to QEMU as in the example below, which enables networking and requests 1 GiB of RAM for the emulated machine:

```
$ /gnu/store/...-run-vm.sh -m 1024 -smp 2 -nic user,model=virtio-net-pci
```

It's possible to combine the two steps into one:

```
$ $(guix system vm my-config.scm) -m 1024 -smp 2 -nic user,model=virtio-net-
```

The VM shares its store with the host system.

By default, the root file system of the VM is mounted volatile; the `--persistent` option can be provided to make it persistent instead. In that case, the VM disk-image file will be copied from the store to the `TMPDIR` directory to make it writable.

Additional file systems can be shared between the host and the VM using the `--share` and `--expose` command-line options: the former specifies a directory to be shared with write access, while the latter provides read-only access to the shared directory.

The example below creates a VM in which the user's home directory is accessible read-only, and where the `/exchange` directory is a read-write mapping of `$HOME/tmp` on the host:

```
guix system vm my-config.scm \
  --expose=$HOME --share=$HOME/tmp=/exchange
```

On GNU/Linux, the default is to boot directly to the kernel; this has the advantage of requiring only a very tiny root disk image since the store of the host can then be mounted.

The `--full-boot` option forces a complete boot sequence, starting with the bootloader. This requires more disk space since a root image containing at least the kernel, `initrd`, and bootloader data files must be created.

The `--image-size` option can be used to specify the size of the image.

The `--no-graphic` option will instruct `guix system` to spawn a headless VM that will use the invoking `tty` for IO. Among other things, this enables copy-pasting, and scrollback. Use the `ctrl-a` prefix to issue QEMU commands; e.g. `ctrl-a h` prints a help, `ctrl-a x` quits the VM, and `ctrl-a c` switches between the QEMU monitor and the VM.

## image

The `image` command can produce various image types. The image type can be selected using the `--image-type` option. It defaults to `mbr-hybrid-raw`. When its value is `iso9660`, the `--label` option can be used to specify a volume ID with `image`. By default, the root file system of a disk image is mounted non-volatile; the `--volatile` option can be provided to make it volatile instead. When using `image`, the bootloader installed on the generated image is taken from the provided `operating-system` definition. The following example demonstrates how to generate an image that uses the `grub-efi-bootloader` bootloader and boot it with QEMU:

```
image=$(guix system image --image-type=qcow2 \
```

```

        gnu/system/examples/lightweight-desktop.tpl)
cp $image /tmp/my-image.qcow2
chmod +w /tmp/my-image.qcow2
qemu-system-x86_64 -enable-kvm -hda /tmp/my-image.qcow2 -m 1000 \
    -bios $(guix build ovmf-x86-64)/share/firmware/ovmf_x64.b

```

When using the `mbr-hybrid-raw` image type, a raw disk image is produced; it can be copied as is to a USB stick, for instance. Assuming `/dev/sdc` is the device corresponding to a USB stick, one can copy the image to it using the following command:

```
# dd if=$(guix system image my-os.scm) of=/dev/sdc status=progress
```

The `--list-image-types` command lists all the available image types.

When using the `qcow2` image type, the returned image is in `qcow2` format, which the QEMU emulator can efficiently use. See Section 11.18 [Running Guix in a VM], page 707, for more information on how to run the image in a virtual machine. The `grub-bootloader` bootloader is always used independently of what is declared in the `operating-system` file passed as argument. This is to make it easier to work with QEMU, which uses the SeaBIOS BIOS by default, expecting a bootloader to be installed in the Master Boot Record (MBR). In case the virtual machine is going to be AArch64, you might want to take a look at `qcow2-gpt` image type that installs bootloader only in EFI.

When using the `docker` image type, a Docker image is produced. Guix builds the image from scratch, not from a pre-existing Docker base image. As a result, it contains *exactly* what you define in the operating system configuration file. You can then load the image and launch a Docker container using commands like the following:

```

image_id="$(docker load < guix-system-docker-image.tar.gz)"
container_id="$(docker create $image_id)"
docker start $container_id

```

This command starts a new Docker container from the specified image. It will boot the Guix system in the usual manner, which means it will start any services you have defined in the operating system configuration. You can get an interactive shell running in the container using `docker exec`:

```
docker exec -ti $container_id /run/current-system/profile/bin/bash --login
```

Depending on what you run in the Docker container, it may be necessary to give the container additional permissions. For example, if you intend to build software using Guix inside of the Docker container, you may need to pass the `--privileged` option to `docker create`.

Last, the `--network` option applies to `guix system docker-image`: it produces an image where network is supposedly shared with the host, and thus without services like `nsd` or `NetworkManager`.

## container

Return a script to run the operating system declared in *file* within a container. Containers are a set of lightweight isolation mechanisms provided by the kernel Linux-libre. Containers are substantially less resource-demanding than full

virtual machines since the kernel, shared objects, and other resources can be shared with the host system; this also means they provide thinner isolation.

Currently, the script must be run as root in order to support more than a single user and group. The container shares its store with the host system.

As with the `vm` action (see [guix system vm], page 695), additional file systems to be shared between the host and container can be specified using the `--share` and `--expose` options:

```
guix system container my-config.scm \
  --expose=$HOME --share=$HOME/tmp=/exchange
```

The `--share` and `--expose` options can also be passed to the generated script to bind-mount additional directories into the container.

The generated script also supports the `--pid-file` option to write the container process ID to a file. This can be used, along with `guix container exec` to execute scripts within the container (see Section 9.14 [Invoking guix container], page 234).

**Note:** The `container` action requires Linux-libre 3.19 or newer.

*options* can contain any of the common build options (see Section 9.1.1 [Common Build Options], page 180). In addition, *options* can contain one of the following:

`--expression=expr`

`-e expr` Consider the operating-system *expr* evaluates to. This is an alternative to specifying a file which evaluates to an operating system. This is used to generate the Guix system installer see Section 3.9 [Building the Installation Image], page 32).

`--system=system`

`-s system` Attempt to build for *system* instead of the host system type. This works as per `guix build` (see Section 9.1 [Invoking guix build], page 180).

`--target=triplet`

Cross-build for *triplet*, which must be a valid GNU triplet, such as "aarch64-linux-gnu" (see Section "Specifying target triplets" in *Autoconf*).

`--derivation`

`-d` Return the derivation file name of the given operating system without building anything.

`--save-provenance`

As discussed above, `guix system init` and `guix system reconfigure` always save provenance information *via* a dedicated service (see Section 11.19.3 [Service Reference], page 712). However, other commands don't do that by default. If you wish to, say, create a virtual machine image that contains provenance information, you can run:

```
guix system image -t qcow2 --save-provenance config.scm
```

That way, the resulting image will effectively "embed its own source" in the form of meta-data in `/run/current-system`. With that information, one can rebuild the image to make sure it really contains what it pretends to contain; or they could use that to derive a variant of the image.

**--image-type=type**

**-t type** For the **image** action, create an image with given *type*.  
 When this option is omitted, **guix system** uses the **mbr-hybrid-raw** image type.  
**--image-type=iso9660** produces an ISO-9660 image, suitable for burning on CDs and DVDs.

**--image-size=size**  
 For the **image** action, create an image of the given *size*. *size* may be a number of bytes, or it may include a unit as a suffix (see Section “Block size” in *GNU Coreutils*).  
 When this option is omitted, **guix system** computes an estimate of the image size as a function of the size of the system declared in *file*.

**--network**

**-N** For the **container** action, allow containers to access the host network, that is, do not create a network namespace.

**--root=file**

**-r file** Make *file* a symlink to the result, and register it as a garbage collector root.

**--skip-checks**  
 Skip pre-installation safety checks.  
 By default, **guix system init** and **guix system reconfigure** perform safety checks: they make sure the file systems that appear in the **operating-system** declaration actually exist (see Section 11.4 [File Systems], page 257), and that any Linux kernel modules that may be needed at boot time are listed in **initrd-modules** (see Section 11.14 [Initial RAM Disk], page 682). Passing this option skips these tests altogether.

**--allow-downgrades**  
 Instruct **guix system reconfigure** to allow system downgrades.  
 By default, **reconfigure** prevents you from downgrading your system. It achieves that by comparing the provenance info of your system (shown by **guix system describe**) with that of your **guix** command (shown by **guix describe**). If the commits for **guix** are not descendants of those used for your system, **guix system reconfigure** errors out. Passing **--allow-downgrades** allows you to bypass these checks.  
**Note:** Make sure you understand its security implications before using **--allow-downgrades**.

**--on-error=strategy**  
 Apply *strategy* when an error occurs when reading *file*. *strategy* may be one of the following:

**nothing-special**  
 Report the error concisely and exit. This is the default strategy.

**backtrace**  
 Likewise, but also display a backtrace.

**debug** Report the error and enter Guile’s debugger. From there, you can run commands such as `,bt` to get a backtrace, `,locals` to display local variable values, and more generally inspect the state of the program. See Section “Debug Commands” in *GNU Guile Reference Manual*, for a list of available debugging commands.

Once you have built, configured, re-configured, and re-re-configured your Guix installation, you may find it useful to list the operating system generations available on disk—and that you can choose from the bootloader boot menu:

**describe** Describe the running system generation: its file name, the kernel and bootloader used, etc., as well as provenance information when available.

The `--list-installed` flag is available, with the same syntax that is used in `guix package --list-installed` (see Section 5.2 [Invoking `guix package`], page 37). When the flag is used, the description will include a list of packages that are currently installed in the system profile, with optional filtering based on a regular expression.

**Note:** The *running* system generation—referred to by `/run/current-system`—is not necessarily the *current* system generation—referred to by `/var/guix/profiles/system`: it differs when, for instance, you chose from the bootloader menu to boot an older generation.

It can also differ from the *booted* system generation—referred to by `/run/booted-system`—for instance because you reconfigured the system in the meantime.

#### **list-generations**

List a summary of each generation of the operating system available on disk, in a human-readable way. This is similar to the `--list-generations` option of `guix package` (see Section 5.2 [Invoking `guix package`], page 37).

Optionally, one can specify a pattern, with the same syntax that is used in `guix package --list-generations`, to restrict the list of generations displayed. For instance, the following command displays generations that are up to 10 days old:

```
$ guix system list-generations 10d
```

The `--list-installed` flag may also be specified, with the same syntax that is used in `guix package --list-installed`. This may be helpful if trying to determine when a package was added to the system.

The `guix system` command has even more to offer! The following sub-commands allow you to visualize how your system services relate to each other:

#### **extension-graph**

Emit to standard output the *service extension graph* of the operating system defined in *file* (see Section 11.19.1 [Service Composition], page 709, for more information on service extensions). By default the output is in Dot/Graphviz format, but you can choose a different format with `--graph-backend`, as with `guix graph` (see Section 9.10 [Invoking `guix graph`], page 222):



The command:

```
$ guix system extension-graph file | xdot -
```

shows the extension relations among services.

**Note:** The `dot` program is provided by the `graphviz` package.

#### shepherd-graph

Emit to standard output the *dependency graph* of shepherd services of the operating system defined in *file*. See Section 11.19.4 [Shepherd Services], page 717, for more information and for an example graph.

Again, the default output format is Dot/Graphviz, but you can pass `--graph-backend` to select a different one.

## 11.17 Invoking guix deploy

We’ve already seen `operating-system` declarations used to manage a machine’s configuration locally. Suppose you need to configure multiple machines, though—perhaps you’re managing a service on the web that’s comprised of several servers. `guix deploy` enables you to use those same `operating-system` declarations to manage multiple remote hosts at once as a logical “deployment”.

**Note:** The functionality described in this section is still under development and is subject to change. Get in touch with us on [guix-devel@gnu.org](mailto:guix-devel@gnu.org)!

`guix deploy file`

Such an invocation will deploy the machines that the code within *file* evaluates to. As an example, *file* might contain a definition like this:

```
;; This is a Guix deployment of a "bare bones" setup, with
;; no X11 display server, to a machine with an SSH daemon
;; listening on localhost:2222. A configuration such as this
;; may be appropriate for virtual machine with ports
;; forwarded to the host's loopback interface.
```

```
(use-service-modules networking ssh)
```

```
(use-package-modules bootloaders)
```

```
(define %system
  (operating-system
    (host-name "gnu-deployed")
    (timezone "Etc/UTC")
    (bootloader (bootloader-configuration
                  (bootloader grub-bootloader)
                  (targets '("/dev/vda"))
                  (terminal-outputs '(console))))
    (file-systems (cons (file-system
                          (mount-point "/")
                          (device "/dev/vda1")
                          (type "ext4"))
                        %base-file-systems)))
```

```

(services
  (append (list (service dhcp-client-service-type)
                (service openssh-service-type
                  (openssh-configuration
                    (permit-root-login #t)
                    (allow-empty-passwords? #t))))
    %base-services))))

(list (machine
      (operating-system %system)
      (environment managed-host-environment-type)
      (configuration (machine-ssh-configuration
                      (host-name "localhost")
                      (system "x86_64-linux")
                      (user "alice")
                      (identity "./id_rsa")
                      (port 2222)))))

```

The file should evaluate to a list of *machine* objects. This example, upon being deployed, will create a new generation on the remote system realizing the **operating-system** declaration `%system`. **environment** and **configuration** specify how the machine should be provisioned—that is, how the computing resources should be created and managed. The above example does not create any resources, as a **'managed-host** is a machine that is already running the Guix system and available over the network. This is a particularly simple case; a more complex deployment may involve, for example, starting virtual machines through a Virtual Private Server (VPS) provider. In such a case, a different *environment* type would be used.

Do note that you first need to generate a key pair on the coordinator machine to allow the daemon to export signed archives of files from the store (see Section 5.11 [Invoking guix archive], page 65), though this step is automatic on Guix System:

```
# guix archive --generate-key
```

Each target machine must authorize the key of the master machine so that it accepts store items it receives from the coordinator:

```
# guix archive --authorize < coordinator-public-key.txt
```

**user**, in this example, specifies the name of the user account to log in as to perform the deployment. Its default value is **root**, but root login over SSH may be forbidden in some cases. To work around this, **guix deploy** can log in as an unprivileged user and employ **sudo** to escalate privileges. This will only work if **sudo** is currently installed on the remote and can be invoked non-interactively as **user**. That is, the line in **sudoers** granting **user** the ability to use **sudo** must contain the **NOPASSWD** tag. This can be accomplished with the following operating system configuration snippet:

```

(use-modules ...
  (gnu system))                                ;for %sudoers-specification

(define %user "username")

```

```
(operating-system
  ...
  (sudoers-file
    (plain-file "sudoers"
      (string-append (plain-file-content %sudoers-specification)
        (format #f "~a ALL = NOPASSWD: ALL~%"
          %user))))))
```

For more information regarding the format of the `sudoers` file, consult `man sudoers`.

Once you've deployed a system on a set of machines, you may find it useful to run a command on all of them. The `--execute` or `-x` option lets you do that; the example below runs `uname -a` on all the machines listed in the deployment file:

```
guix deploy file -x -- uname -a
```

One thing you may often need to do after deployment is restart specific services on all the machines, which you can do like so:

```
guix deploy file -x -- herd restart service
```

The `guix deploy -x` command returns zero if and only if the command succeeded on all the machines.

You may also wish to roll back configurations on machines to a previous generation. You can do that using the `--roll-back` or `-r` option like so:

```
guix deploy --roll-back file
```

Below are the data types you need to know about when writing a deployment file.

**machine** [Data Type]

This is the data type representing a single machine in a heterogeneous Guix deployment.

**operating-system**

The object of the operating system configuration to deploy.

**environment**

An **environment-type** describing how the machine should be provisioned.

**configuration** (default: `#f`)

An object describing the configuration for the machine's **environment**. If the **environment** has a default configuration, `#f` may be used. If `#f` is used for an environment with no default configuration, however, an error will be thrown.

**machine-ssh-configuration** [Data Type]

This is the data type representing the SSH client parameters for a machine with an **environment** of `managed-host-environment-type`.

**host-name**

**build-locally?** (default: `#t`)

If false, system derivations will be built on the machine being deployed to.

**system** The system type describing the architecture of the machine being deployed to—e.g., "x86\_64-linux".

**authorize?** (default: #t)

If true, the coordinator's signing key will be added to the remote's ACL keyring.

**port** (default: 22)

**user** (default: "root")

**identity** (default: #f)

If specified, the path to the SSH private key to use to authenticate with the remote host.

**host-key** (default: #f)

This should be the SSH host key of the machine, which looks like this:

```
ssh-ed25519 AAAAC3Nz... root@example.org
```

When **host-key** is #f, the server is authenticated against the `~/.ssh/known_hosts` file, just like the OpenSSH `ssh` client does.

**allow-downgrades?** (default: #f)

Whether to allow potential downgrades.

Like `guix system reconfigure`, `guix deploy` compares the channel commits currently deployed on the remote host (as returned by `guix system describe`) to those currently in use (as returned by `guix describe`) to determine whether commits currently in use are descendants of those deployed. When this is not the case and **allow-downgrades?** is false, it raises an error. This ensures you do not accidentally downgrade remote machines.

**safety-checks?** (default: #t)

Whether to perform “safety checks” before deployment. This includes verifying that devices and file systems referred to in the operating system configuration actually exist on the target machine, and making sure that Linux modules required to access storage devices at boot time are listed in the `initrd-modules` field of the operating system.

These safety checks ensure that you do not inadvertently deploy a system that would fail to boot. Be careful before turning them off!

**digital-ocean-configuration** [Data Type]

This is the data type describing the Droplet that should be created for a machine with an `environment` of `digital-ocean-environment-type`.

**ssh-key** The path to the SSH private key to use to authenticate with the remote host. In the future, this field may not exist.

**tags** A list of string “tags” that uniquely identify the machine. Must be given such that no two machines in the deployment have the same set of tags.

**region** A Digital Ocean region slug, such as "nyc3".

**size** A Digital Ocean size slug, such as "s-1vcpu-1gb"

`enable-ipv6?`

Whether or not the droplet should be created with IPv6 networking.

`hetzner-configuration`

[Data Type]

This is the data type describing the server that should be created for a machine with an environment of `hetzner-environment-type`. It allows you to configure deployment to a VPS (virtual private server) hosted by Hetzner (<https://www.hetzner.com>).

`allow-downgrades?` (default: `#f`)

Whether to allow potential downgrades.

`authorize?` (default: `#t`)

If true, the public signing key `"/etc/guix/signing-key.pub"` of the machine that invokes `guix deploy` will be added to the operating system ACL keyring of the target machine.

`build-locally?` (default: `#t`)

If true, system derivations will be built on the machine that invokes `guix deploy`, otherwise derivations are build on the target machine. Set this to `#f` if the machine you are deploying from has a different architecture than the target machine and you can't build derivations for the target architecture by other means, like offloading (see Section 2.2.2 [Daemon Offload Setup], page 9) or emulation (see [transparent-emulation-qemu], page 588).

`delete?` (default: `#t`)

If true, the server will be deleted when an error happens in the provisioning phase. If false, the server will be kept in order to debug any issues.

`ipv4` (default: `#t`)

`ipv6` (default: `#t`)

When false, no public IP address is attached. Specify the name of an existing primary ip to attach it to the machine. Other values would create a new address automatically.

`labels` (default: `'()`)

A user defined alist of key/value pairs attached to the SSH key and the server on the Hetzner API. Keys and values must be strings, e.g. `'(("environment" . "development"))`. For more information, see Labels (<https://docs.hetzner.cloud/#labels>).

`location` (default: `"fsn1"`)

The name of a location (<https://docs.hetzner.com/cloud/general/locations>) to create the server in. For example, `"fsn1"` corresponds to the Hetzner site in Falkenstein, Germany, while `"sin"` corresponds to its site in Singapore.

`server-type` (default: `"cx42"`)

The name of the server type (<https://docs.hetzner.com/cloud/servers/overview#server-types>) this virtual server should be created

with. For example, "cx42" corresponds to a x86\_64 server that has 8 VCPUs, 16 GB of memory and 160 GB of storage, while "cax31" to the AArch64 equivalent. Other server types and their current prices can be found here (<https://www.hetzner.com/cloud/#pricing>). The "cpx11" server type is currently not supported, since its rescue system is too small to bootstrap a Guix system from.

**ssh-key** (default: #f)

If specified, the file name of the SSH private key to use to authenticate with the remote host.

**ssh-public-key** (default: extracted from **ssh-key**)

If specified, either a public key as returned by **string->public-key** or the path to the SSH public key to use to authenticate with the remote host.

When deploying a machine for the first time, the following steps are taken to provision a server for the machine on the Hetzner Cloud (<https://www.hetzner.com/cloud>) service:

- Create the SSH key of the machine on the Hetzner API.
- Create a server for the machine on the Hetzner API.
- Format the root partition of the disk using the file system of the machine's operating system. Supported file systems are btrfs and ext4.
- Install a minimal Guix operating system on the server using the rescue mode (<https://docs.hetzner.com/cloud/servers/getting-started/rescue-system>). This minimal system is used to install the machine's operating system, after rebooting.
- Reboot the server and apply the machine's operating system on the server.

Once the server has been provisioned and SSH is available, deployment continues by delegating it to the **managed-host-environment-type**.

Servers on the Hetzner Cloud service can be provisioned on the AArch64 architecture using UEFI boot mode, or on the x86\_64 architecture using BIOS boot mode. The (gnu machine hetzner) module exports the **%hetzner-os-arm** and **%hetzner-os-x86** operating systems that are compatible with those two architectures, and can be used as a base for defining your custom operating system.

The following example shows the definition of two machines that are deployed on the Hetzner Cloud service. The first one uses the **%hetzner-os-arm** operating system to run a server with 16 shared vCPUs and 32 GB of RAM on the **aarch64** architecture, the second one uses the **%hetzner-os-x86** operating system on a server with 16 shared vCPUs and 32 GB of RAM on the **x86\_64** architecture.

```
(use-modules (gnu machine)
             (gnu machine hetzner))

(list (machine
      (operating-system %hetzner-os-arm)
      (environment hetzner-environment-type))
```

```
(configuration (hetzner-configuration
  (server-type "cax41")
  (ssh-key "/home/charlie/.ssh/id_rsa"))))
(machine
  (operating-system %hetzner-os-x86)
  (environment hetzner-environment-type)
  (configuration (hetzner-configuration
    (server-type "cpx51")
    (ssh-public-key "/home/charlie/.ssh/id_rsa.pub")))))■
```

Passing this file to `guix deploy` with the environment variable `GUIX_HETZNER_API_TOKEN` set to a valid Hetzner API key (<https://docs.hetzner.com/cloud/api/getting-started/generating-api-token>) should provision two machines for you.

## 11.18 Running Guix in a Virtual Machine

To run Guix in a virtual machine (VM), one can use the pre-built Guix VM image distributed at [https://ftpmirror.gnu.org/gnu/guix/guix-system-vm-image-8b23207.x86\\_64-linux.qcow2](https://ftpmirror.gnu.org/gnu/guix/guix-system-vm-image-8b23207.x86_64-linux.qcow2). This image is a compressed image in QCOW format. You can pass it to an emulator such as QEMU (<https://qemu.org/>) (see below for details).

This image boots the Xfce graphical environment and it contains some commonly used tools. You can install more software in the image by running `guix package` in a terminal (see Section 5.2 [Invoking `guix package`], page 37). You can also reconfigure the system based on its initial configuration file available as `/run/current-system/configuration.scm` (see Section 11.2 [Using the Configuration System], page 244).

Instead of using this pre-built image, one can also build their own image using `guix system image` (see Section 11.16 [Invoking `guix system`], page 692).

If you built your own image, you must copy it out of the store (see Section 8.9 [The Store], page 156) and give yourself permission to write to the copy before you can use it. When invoking QEMU, you must choose a system emulator that is suitable for your hardware platform. Here is a minimal QEMU invocation that will boot the result of `guix system image -t qcow2` on x86\_64 hardware:

```
$ qemu-system-x86_64 \
  -nic user,model=virtio-net-pci \
  -enable-kvm -m 2048 \
  -device virtio-blk,drive=myhd \
  -drive if=none,file=guix-system-vm-image-8b23207.x86_64-linux.qcow2,id=myhd■
```

Here is what each of these options means:

`qemu-system-x86_64`

This specifies the hardware platform to emulate. This should match the host.

`-nic user,model=virtio-net-pci`

Enable the unprivileged user-mode network stack. The guest OS can access the host but not vice versa. This is the simplest way to get the guest OS online. `model` specifies which network device to emulate: `virtio-net-pci` is a special device made for virtualized operating systems and recommended for most uses.

Assuming your hardware platform is x86\_64, you can get a list of available NIC models by running `qemu-system-x86_64 -nic model=help`.

**-enable-kvm**

If your system has hardware virtualization extensions, enabling the virtual machine support (KVM) of the Linux kernel will make things run faster.

**-m 2048** RAM available to the guest OS, in mebibytes. Defaults to 128 MiB, which may be insufficient for some operations.

**-device virtio-blk,drive=myhd**

Create a `virtio-blk` drive called “myhd”. `virtio-blk` is a “paravirtualization” mechanism for block devices that allows QEMU to achieve better performance than if it were emulating a complete disk drive. See the QEMU and KVM documentation for more info.

**-drive if=none,file=/tmp/qemu-image,id=myhd**

Use our QCOW image, the `guix-system-vm-image-8b23207.x86_64-linux.qcow2` file, as the backing store of the “myhd” drive.

The default `run-vm.sh` script that is returned by an invocation of `guix system vm` does not add a `-nic user` flag by default. To get network access from within the vm add the `(dhcp-client-service)` to your system definition and start the VM using `$(guix system vm config.scm) -nic user`. An important caveat of using `-nic user` for networking is that `ping` will not work, because it uses the ICMP protocol. You’ll have to use a different command to check for network connectivity, for example `guix download`.

### 11.18.1 Connecting Through SSH

To enable SSH inside a VM you need to add an SSH server like `openssh-service-type` to your VM (see Section 11.10.5 [Networking Services], page 322). In addition you need to forward the SSH port, 22 by default, to the host. You can do this with

```
$(guix system vm config.scm) -nic user,model=virtio-net-pci,hostfwd=tcp::10022-:22
```

To connect to the VM you can run

```
ssh -o UserKnownHostsFile=/dev/null -o StrictHostKeyChecking=no -p 10022 localhost
```

The `-p` tells `ssh` the port you want to connect to. `-o UserKnownHostsFile=/dev/null` prevents `ssh` from complaining every time you modify your `config.scm` file and the `-o StrictHostKeyChecking=no` prevents you from having to allow a connection to an unknown host every time you connect.

**Note:** If you find the above ‘`hostfwd`’ example not to be working (e.g., your SSH client hangs attempting to connect to the mapped port of your VM), make sure that your Guix System VM has networking support, such as by using the `dhcp-client-service-type` service type.

### 11.18.2 Using virt-viewer with Spice

As an alternative to the default `qemu` graphical client you can use the `remote-viewer` from the `virt-viewer` package. To connect pass the `-spice port=5930,disable-ticketing` flag to `qemu`. See previous section for further information on how to do this.



Spice also allows you to do some nice stuff like share your clipboard with your VM. To enable that you'll also have to pass the following flags to `qemu`:

```
-device virtio-serial-pci,id=virtio-serial0,max_ports=16,bus=pci.0,addr=0x5
-chardev spicevmc,name=vdagent,id=vdagent
-device virtserialport,nr=1,bus=virtio-serial0.0,chardev=vdagent,\
name=com.redhat.spice.0
```

You'll also need to add the (`spice-vdagent-service`) to your system definition (see Section 11.10.39 [Miscellaneous Services], page 649).

## 11.19 Defining Services

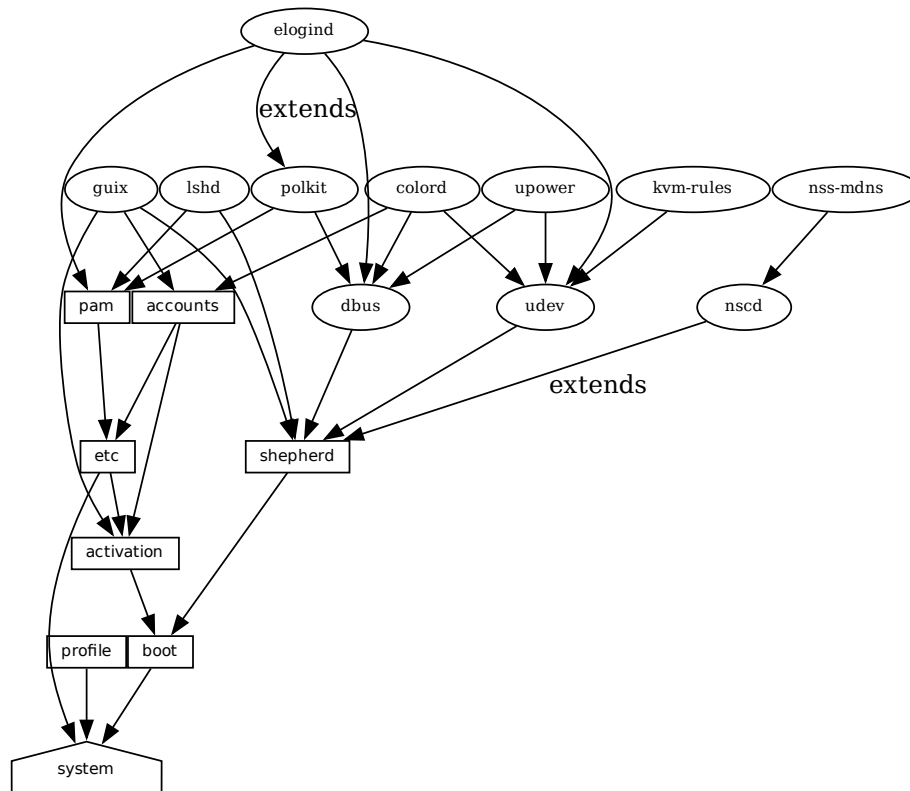
The previous sections show the available services and how one can combine them in an `operating-system` declaration. But how do we define them in the first place? And what is a service anyway?

### 11.19.1 Service Composition

Here we define a *service* as, broadly, something that extends the functionality of the operating system. Often a service is a process—a *daemon*—started when the system boots: a secure shell server, a Web server, the Guix build daemon, etc. Sometimes a service is a daemon whose execution can be triggered by another daemon—e.g., an FTP server started by `inetd` or a D-Bus service activated by `dbus-daemon`. Occasionally, a service does not map to a daemon. For instance, the “account” service collects user accounts and makes sure they exist when the system runs; the “udev” service collects device management rules and makes them available to the `eudev` daemon; the `/etc` service populates the `/etc` directory of the system.

Guix system services are connected by *extensions*. For instance, the secure shell service *extends* the Shepherd—the initialization system, running as PID 1—by giving it the command lines to start and stop the secure shell daemon (see Section 11.10.5 [Networking Services], page 322); the `UPower` service extends the D-Bus service by passing it its `.service` specification, and extends the `udev` service by passing it device management rules (see Section 11.10.9 [Desktop Services], page 377); the Guix daemon service extends the Shepherd by passing it the command lines to start and stop the daemon, and extends the account service by passing it a list of required build user accounts (see Section 11.10.1 [Base Services], page 276).

All in all, services and their “extends” relations form a directed acyclic graph (DAG). If we represent services as boxes and extensions as arrows, a typical system might provide something like this:



At the bottom, we see the *system service*, which produces the directory containing everything to run and boot the system, as returned by the `guix system build` command. See Section 11.19.3 [Service Reference], page 712, to learn about the other service types shown here. See [system-extension-graph], page 700, for information on how to generate this representation for a particular operating system definition.

Technically, developers can define *service types* to express these relations. There can be any number of services of a given type on the system, for instance when running two instances with different parameters.

The following section describes the programming interface for service types and services.

### 11.19.2 Service Types and Services

A *service type* is a node in the DAG described above. Let us start with a simple example, the service type for the Guix build daemon (see Section 2.3 [Invoking guix-daemon], page 14):

```

(define guix-service-type
  (service-type
    (name 'guix)
    (extensions
      (list (service-extension shepherd-root-service-type guix-shepherd-service))
```



```
(compose concatenate)          ;concatenate the list of rules
(extend (lambda (config rules)
  (udev-configuration
    (inherit config)
    (rules (append (udev-configuration-rules config)
                  rules))))))
```

This is the service type for the eudev device management daemon (<https://github.com/eudev-project/eudev>). Compared to the previous example, in addition to an extension of `shepherd-root-service-type`, we see two new fields:

**compose** This is the procedure to *compose* the list of extensions to services of this type. Services can extend the udev service by passing it lists of rules; we compose those extensions simply by concatenating them.

**extend** This procedure defines how the value of the service is *extended* with the composition of the extensions.

Udev extensions are composed into a list of rules, but the udev service value is itself a `<udev-configuration>` record. So here, we extend that record by appending the list of rules it contains to the list of contributed rules.

**description**

This is a string giving an overview of the service type. The string can contain Texinfo markup (see Section “Overview” in *GNU Texinfo*). The `guix system search` command searches these strings and displays them (see Section 11.16 [Invoking guix system], page 692).

There can be only one instance of an extensible service type such as `udev-service-type`. If there were more, the `service-extension` specifications would be ambiguous.

Still here? The next section provides a reference of the programming interface for services.

### 11.19.3 Service Reference

We have seen an overview of service types (see Section 11.19.2 [Service Types and Services], page 710). This section provides a reference on how to manipulate services and service types. This interface is provided by the `(gnu services)` module.

**service type** [*value*] [Procedure]

Return a new service of *type*, a `<service-type>` object (see below). *value* can be any object; it represents the parameters of this particular service instance.

When *value* is omitted, the default value specified by *type* is used; if *type* does not specify a default value, an error is raised.

For instance, this:

```
(service openssh-service-type)
```

is equivalent to this:

```
(service openssh-service-type
  (openssh-configuration))
```

In both cases the result is an instance of `openssh-service-type` with the default configuration.

`service? obj` [Procedure]

Return true if *obj* is a service.

`service-kind service` [Procedure]

Return the type of *service*—i.e., a `<service-type>` object.

`service-value service` [Procedure]

Return the value associated with *service*. It represents its parameters.

Here is an example of how a service is created and manipulated:

```
(define s
  (service nginx-service-type
    (nginx-configuration
      (nginx nginx)
      (log-directory log-directory)
      (run-directory run-directory)
      (file config-file))))
```

```
(service? s)
```

```
⇒ #t
```

```
(eq? (service-kind s) nginx-service-type)
```

```
⇒ #t
```

The `modify-services` form provides a handy way to change the parameters of some of the services of a list such as `%base-services` (see Section 11.10.1 [Base Services], page 276). It evaluates to a list of services. Of course, you could always use standard list combinators such as `map` and `fold` to do that (see Section “SRFI-1” in *GNU Guile Reference Manual*); `modify-services` simply provides a more concise form for this common pattern.

`modify-services services (type variable => body) ...` [Special Form]

Modify the services listed in *services* according to the given clauses. Each clause has the form:

```
(type variable => body)
```

where *type* is a service type—e.g., `guix-service-type`—and *variable* is an identifier that is bound within the *body* to the service parameters—e.g., a `guix-configuration` instance—of the original service of that *type*.

The *body* should evaluate to the new service parameters, which will be used to configure the new service. This new service will replace the original in the resulting list. Because a service’s service parameters are created using `define-record-type*`, you can write a succinct *body* that evaluates to the new service parameters by using the `inherit` feature that `define-record-type*` provides.

Clauses can also have the following form:

```
(delete type)
```

Such a clause removes all services of the given *type* from *services*.

See Section 11.2 [Using the Configuration System], page 244, for example usage.

Next comes the programming interface for service types. This is something you want to know when writing new service definitions, but not necessarily when simply looking for ways to customize your `operating-system` declaration.

**service-type** [Data Type]

This is the representation of a *service type* (see Section 11.19.2 [Service Types and Services], page 710).

**name** This is a symbol, used only to simplify inspection and debugging.

**extensions**

A non-empty list of `<service-extension>` objects (see below).

**compose** (default: `#f`)

If this is `#f`, then the service type denotes services that cannot be extended—i.e., services that do not receive “values” from other services.

Otherwise, it must be a one-argument procedure. The procedure is called by `fold-services` and is passed a list of values collected from extensions. It may return any single value.

**extend** (default: `#f`)

If this is `#f`, services of this type cannot be extended.

Otherwise, it must be a two-argument procedure: `fold-services` calls it, passing it the initial value of the service as the first argument and the result of applying `compose` to the extension values as the second argument. It must return a value that is a valid parameter value for the service instance.

**description**

This is a string, possibly using Texinfo markup, describing in a couple of sentences what the service is about. This string allows users to find about the service through `guix system search` (see Section 11.16 [Invoking guix system], page 692).

**default-value** (default: `&no-default-value`)

The default value associated for instances of this service type. This allows users to use the `service` form without its second argument:

`(service type)`

The returned service in this case has the default value specified by *type*.

See Section 11.19.2 [Service Types and Services], page 710, for examples.

**service-extension** *target-type compute* [Procedure]

Return a new extension for services of type *target-type*. *compute* must be a one-argument procedure: `fold-services` calls it, passing it the value associated with the service that provides the extension; it must return a valid value for the target service.

**service-extension?** *obj* [Procedure]

Return true if *obj* is a service extension.

Occasionally, you might want to simply extend an existing service. This involves creating a new service type and specifying the extension of interest, which can be verbose; the `simple-service` procedure provides a shorthand for this.

**simple-service** *name target value* [Procedure]

Return a service that extends *target* with *value*. This works by creating a singleton service type *name*, of which the returned service is an instance.

For example, this extends `mcron` (see Section 11.10.2 [Scheduled Job Execution], page 299) with an additional job:

```
(simple-service 'my-mcron-job mcron-service-type
  #~(job '(next-hour (3)) "guix gc -F 2G"))
```

At the core of the service abstraction lies the `fold-services` procedure, which is responsible for “compiling” a list of services down to a single directory that contains everything needed to boot and run the system—the directory shown by the `guix system build` command (see Section 11.16 [Invoking `guix system`], page 692). In essence, it propagates service extensions down the service graph, updating each node parameters on the way, until it reaches the root node.

**fold-services** *services* [*#:target-type system-service-type*] [Procedure]

Fold *services* by propagating their extensions down to the root of type *target-type*; return the root service adjusted accordingly.

Lastly, the `(gnu services)` module also defines several essential service types, some of which are listed below.

**system-service-type** [Variable]

This is the root of the service graph. It produces the system directory as returned by the `guix system build` command.

**boot-service-type** [Variable]

The type of the “boot service”, which produces the *boot script*. The boot script is what the initial RAM disk runs when booting.

**etc-service-type** [Variable]

The type of the `/etc` service. This service is used to create files under `/etc` and can be extended by passing it name/file tuples such as:

```
(list `("issue" ,(plain-file "issue" "Welcome!\n")))
```

In this example, the effect would be to add an `/etc/issue` file pointing to the given file.

**etc-profile-d-service-type** [Variable]

The type of the `/etc/profile.d` service. This service is used to create files under `/etc/profile.d`. It takes as value a list of file-like objects, as can be produced with `local-file`, `plain-file`, etc. Note that provided files whose file names do not end with `.sh` are *not* added to `/etc/profile.d/` and are silently dropped. Package objects can also be provided directly to have their `etc/profile.d/*.sh` prefixed files added. A simple usage may look like:

```
(simple-service 'my-important-profile-file etc-profile-d-service-type
  (list (plain-file "youppi.sh" "HOW_IMPORTANT=very")))
```

**etc-bashrc-d-service-type** [Variable]

The type of the `/etc/bashrc.d` service. This service is used to create files under `/etc/bashrc.d`. It takes as value a list of file-like objects, as can be produced with `local-file`, `plain-file`, etc. Note that provided files whose file names do not end with `.sh` are *not* added to `/etc/profile.d/` and are silently dropped. The default value is made available via the `%default-etc-bashrc-d-files` variable for users to extend. Package objects can also be provided directly to have their `etc/bashrc.d/*.sh` prefixed files added. Since the service is part `%base-services`, it can be extended via `simple-service` like so:

```
(use-package-modules gnome)      ;for the `vte' package

(simple-service 'vte-integration etc-bashrc-d-service-type
  (list (file-append vte "/etc/profile.d/vte.sh")))
```

**vte-integration-service-type** [Variable]

This service adds the `/etc/bashrc.d/vte.sh` to your system, which improves the Bash and Zsh experience when using VTE-powered terminal emulators. This causes for example the current directory to be displayed in the terminal emulator's tab title, and the current directory to be preserved when creating a new tab, among other features. The value of the service is the `vte` package to use.

```
(use-package-modules gnome)      ;for the `vte' package

(service vte-integration-service-type vte)
```

**privileged-program-service-type** [Variable]

Type for the “privileged-program service”. This service collects lists of executable file names, passed as `gexps`, and adds them to the set of privileged programs on the system (see Section 11.11 [Privileged Programs], page 678).

**profile-service-type** [Variable]

Type of the service that populates the *system profile*—i.e., the programs under `/run/current-system/profile`. Other services can extend it by passing it lists of packages to add to the system profile.

**provenance-service-type** [Variable]

This is the type of the service that records *provenance meta-data* in the system itself. It creates several files under `/run/current-system`:

**channels.scm**

This is a “channel file” that can be passed to `guix pull -C` or `guix time-machine -C`, and which describes the channels used to build the system, if that information was available (see Chapter 6 [Channels], page 68).

**configuration.scm**

This is the file that was passed as the value for this `provenance-service-type` service. By default, `guix system reconfigure` automatically passes the OS configuration file it received on the command line.



**provenance**

This contains the same information as the two other files but in a format that is more readily processable.

In general, these two pieces of information (channels and configuration file) are enough to reproduce the operating system “from source”.

**Caveats:** This information is necessary to rebuild your operating system, but it is not always sufficient. In particular, `configuration.scm` itself is insufficient if it is not self-contained—if it refers to external Guile modules or to extra files. If you want `configuration.scm` to be self-contained, we recommend that modules or files it refers to be part of a channel.

Besides, provenance meta-data is “silent” in the sense that it does not change the bits contained in your system, *except for the meta-data bits themselves*. Two different OS configurations or sets of channels can lead to the same system, bit-for-bit; when `provenance-service-type` is used, these two systems will have different meta-data and thus different store file names, which makes comparison less trivial.

This service is automatically added to your operating system configuration when you use `guix system reconfigure`, `guix system init`, or `guix deploy`.

**linux-loadable-module-service-type** [Variable]

Type of the service that collects lists of packages containing kernel-loadable modules, and adds them to the set of kernel-loadable modules.

This service type is intended to be extended by other service types, such as below:

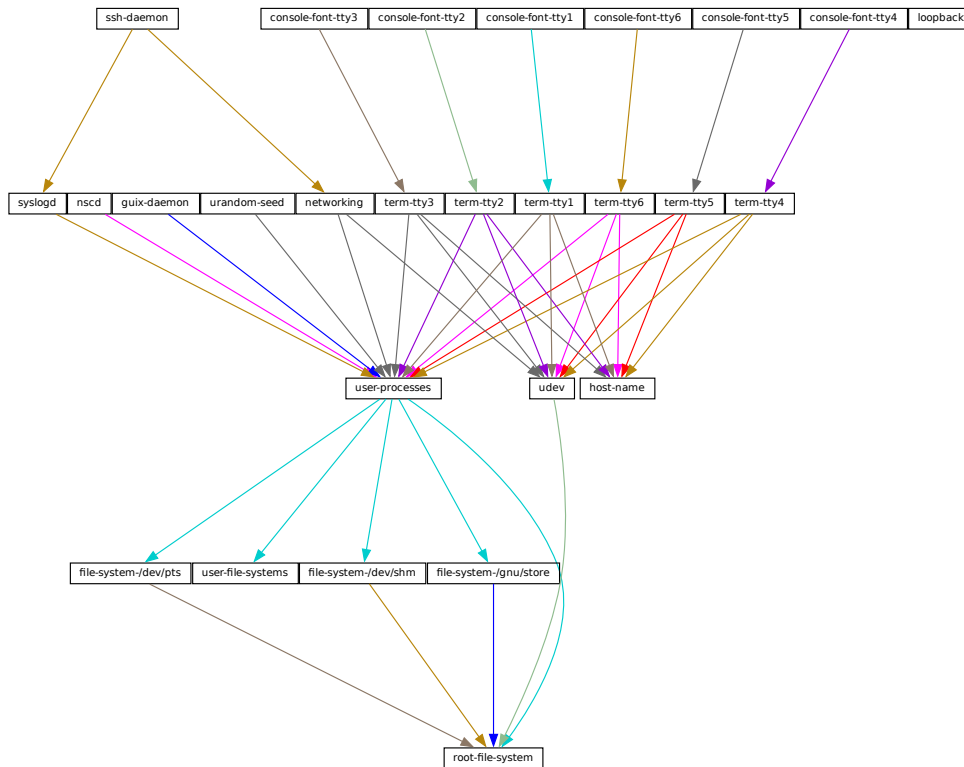
```
(simple-service 'installing-module
                linux-loadable-module-service-type
                (list module-to-install-1
                      module-to-install-2))
```

This does not actually load modules at bootup, only adds it to the kernel profile so that it *can* be loaded by other means.

**11.19.4 Shepherd Services**

The `(gnu services shepherd)` module provides a way to define services managed by the GNU Shepherd, which is the initialization system—the first process that is started when the system boots, also known as PID 1 (see Section “Introduction” in *The GNU Shepherd Manual*).

Services in the Shepherd can depend on each other. For instance, the SSH daemon may need to be started after the syslog daemon has been started, which in turn can only happen once all the file systems have been mounted. The simple operating system defined earlier (see Section 11.2 [Using the Configuration System], page 244) results in a service graph like this:



You can actually generate such a graph for any operating system definition using the `guix system shepherd-graph` command (see [system-shepherd-graph], page 701).

The `%shepherd-root-service` is a service object representing PID 1, of type `shepherd-root-service-type`; it can be extended by passing it lists of `<shepherd-service>` objects.

### shepherd-service

[Data Type]

The data type representing a service managed by the Shepherd.

#### provision

This is a list of symbols denoting what the service provides.

These are the names that may be passed to `herd start`, `herd status`, and similar commands (see Section “Invoking herd” in *The GNU Shepherd Manual*). See Section “Defining Services” in *The GNU Shepherd Manual*, for details.

#### requirement (default: '())

List of symbols denoting the Shepherd services this one depends on.

With very few exceptions, services that start a daemon must depend at least on the `user-processes` service; this ensures that these daemons

are started after all the file systems have been mounted and that, during shutdown, they are stopped before attempting to terminate remaining user processes.

Similarly, services that do make mounts during their startup or operation, should have `user-file-systems` as a requirement. `user-processes` does depend on `user-file-systems` as well. So direct requirement only on `user-processes` is sufficient.

`one-shot?` (default: `#f`)

Whether this service is *one-shot*. One-shot services stop immediately after their `start` action has completed. See Section “Slots of services” in *The GNU Shepherd Manual*, for more info.

`respawn?` (default: `#t`)

Whether to restart the service when it stops, for instance when the underlying process dies.

`respawn-limit` (default: `#f`)

Set a limit on how many times and how frequently a service may be restarted by Shepherd before it is disabled. See Section “Defining Services” in *The GNU Shepherd Manual*, for details.

`respawn-delay` (default: `#f`)

When true, this is the delay in seconds before restarting a failed service.

`start` (default: `#~(const #t)`)

`stop` (default: `#~(const #f)`)

The `start` and `stop` fields refer to the Shepherd’s facilities to start and stop processes (see Section “Service De- and Constructors” in *The GNU Shepherd Manual*). They are given as G-expressions that get expanded in the Shepherd configuration file (see Section 8.12 [G-Expressions], page 165).

`actions` (default: `'()`)

This is a list of `shepherd-action` objects (see below) defining *actions* supported by the service, in addition to the standard `start` and `stop` actions. Actions listed here become available as `herd` sub-commands:

```
herd action service [arguments...]
```

`free-form` (default: `#f`)

When set, this field replaces the `start`, `stop`, and `actions` fields. It is meant to be used when the service definition comes from some other source, typically the service collection provided by the Shepherd proper (see Section “Service Collection” in *The GNU Shepherd Manual*).

For example, the snippet below defines a service for the Shepherd’s built-in REPL (read-eval-print loop) service (see Section “REPL Service” in *The GNU Shepherd Manual*):

```
(shepherd-service
  (provision '(repl))
  (modules '((shepherd service repl)))
```

```
(free-form #~(repl-service)))
```

In this case, the service object is returned by the `repl-service` procedure of the Shepherd, so all the `free-form` G-expression does is call that procedure. Note that the `provision` field must be consistent with the actual service provision.

`auto-start?` (default: `#t`)

Whether this service should be started automatically by the Shepherd. If it is `#f` the service has to be started manually with `herd start`.

`documentation`

A documentation string, as shown when running:

```
herd doc service-name
```

where `service-name` is one of the symbols in `provision` (see Section “Invoking herd” in *The GNU Shepherd Manual*).

`modules` (default: `%default-modules`)

This is the list of modules that must be in scope when `start` and `stop` are evaluated.

The example below defines a Shepherd service that spawns `syslogd`, the system logger from the GNU Networking Utilities (see Section “`syslogd` invocation” in *GNU Inetutils*):

```
(let ((config (plain-file "syslogd.conf" "...")))
  (shepherd-service
    (documentation "Run the syslog daemon (syslogd).")
    (provision '(syslogd))
    (requirement '(user-processes))
    (start #~(make-forkexec-creator
      (list #$(file-append inetutils "/libexec/syslogd")
        "--rcfile" #$(config)
        #:pid-file "/var/run/syslog.pid")))
    (stop #~(make-kill-destructor))))
```

Key elements in this example are the `start` and `stop` fields: they are *staged* code snippets that use the `make-forkexec-creator` procedure provided by the Shepherd and its dual, `make-kill-destructor` (see Section “Service De- and Constructors” in *The GNU Shepherd Manual*). The `start` field will have `shepherd` spawn `syslogd` with the given option; note that we pass `config` after `--rcfile`, which is a configuration file declared above (contents of this file are omitted). Likewise, the `stop` field tells how this service is to be stopped; in this case, it is stopped by making the `kill` system call on its PID. Code staging is achieved using G-expressions: `#~` stages code, while `#$` “escapes” back to host code (see Section 8.12 [G-Expressions], page 165).

`shepherd-action`

[Data Type]

This is the data type that defines additional actions implemented by a Shepherd service (see above).

`name`            Symbol naming the action.

**documentation**

This is a documentation string for the action. It can be viewed by running:

```
herd doc service action action
```

**procedure**

This should be a gexp that evaluates to a procedure of at least one argument, which is the “running value” of the service (see Section “Slots of services” in *The GNU Shepherd Manual*).

The following example defines an action called **say-hello** that kindly greets the user:

```
(shepherd-action
  (name 'say-hello)
  (documentation "Say hi!")
  (procedure #~(lambda (running . args)
    (format #t "Hello, friend! arguments: ~s\n"
      args)
    #t))))
```

Assuming this action is added to the **example** service, then you can do:

```
# herd say-hello example
Hello, friend! arguments: ()
# herd say-hello example a b c
Hello, friend! arguments: ("a" "b" "c")
```

This, as you can see, is a fairly sophisticated way to say hello. See Section “Defining Services” in *The GNU Shepherd Manual*, for more info on actions.

**shepherd-configuration-action**

[Procedure]

Return a **configuration** action to display *file*, which should be the name of the service’s configuration file.

It can be useful to equip services with that action. For example, the service for the Tor anonymous router (see Section 11.10.5 [Networking Services], page 322) is defined roughly like this:

```
(let ((torrc (plain-file "torrc" ...)))
  (shepherd-service
    (provision '(tor))
    (requirement '(user-processes loopback syslogd))

    (start #~(make-forkexec-constructor
      (list #$(file-append tor "/bin/tor") "-f" #$torrc)
      #:user "tor" #:group "tor"))
    (stop #~(make-kill-destructor))
    (actions (list (shepherd-configuration-action torrc)))
    (documentation "Run the Tor anonymous network overlay.")))
```

Thanks to this action, administrators can inspect the configuration file passed to **tor** with this shell command:

```
cat $(herd configuration tor)
```

This can come in as a handy debugging tool!

**shepherd-root-service-type** [Variable]

The service type for the Shepherd “root service”—i.e., PID 1.

This is the service type that extensions target when they want to create shepherd services (see Section 11.19.2 [Service Types and Services], page 710, for an example). Each extension must pass a list of `<shepherd-service>`. Its value must be a `shepherd-configuration`, as described below.

**shepherd-configuration** [Data Type]

This data type represents the Shepherd’s configuration.

**shepherd** (default: `shepherd`)

The Shepherd package to use.

**services** (default: `'()`)

A list of `<shepherd-service>` to start. You should probably use the service extension mechanism instead (see Section 11.19.4 [Shepherd Services], page 717).

The following example specifies the Shepherd package for the operating system:

```
(operating-system
;; ...
(services (append (list openssh-service-type)
;; ...
%desktop-services)
;; ...
;; Use own Shepherd package.
(essential-services
(modify-services (operating-system-default-essential-services
this-operating-system)
(shepherd-root-service-type config => (shepherd-configuration
(inherit config)
(shepherd my-shepherd))))))
```

**shepherd-transient-service-type** [Variable]

This service type represents the Shepherd’s **transient** service, which lets you spawn commands in the background and interact with them as regular Shepherd services; it is similar to `systemd-run`.

For example, the command below spawns `rsync` in the background, in an environment where the `SSH_AUTH_SOCK` environment variable has the given value:

```
herd spawn transient -E SSH_AUTH_SOCK=$SSH_AUTH_SOCK -- \
rsync -e ssh -vur . backup.example.org:
```

See Section “Transient Service Maker” in *The GNU Shepherd Manual*, for more info on the **transient** service.

**shepherd-timer-service-type** [Variable]

This is the service type representing the Shepherd’s **timer** service, which lets you schedule the execution of commands, similar to the venerable `at` command. Here is an example:

```
herd schedule timer at 07:00 -- mpg123 Music/alarm.mp3
```

See Section “Timer Service” in *The GNU Shepherd Manual*, for more info on the `timer` service.

**shepherd-system-log-service-type** [Variable]

This service implements a *system log*, reading messages applications send to `/dev/log` and writing them to files or terminals according to user-defined rules. It provides functionality traditionally implemented by `syslogd` programs.

The value of services of this type must be a `system-log-configuration` record, as described below.

This service is part of `%base-services` (see Section 11.10.1 [Base Services], page 276).

**system-log-configuration** [Data Type]

Available `system-log-configuration` fields are:

**provision** (default: `(system-log syslogd)`) (type: list-of-symbols)

The name(s) of the system log service.

**requirement** (default: `(root-file-system)`) (type: list-of-symbols)

Dependencies of the system log service.

**kernel-log-file** (type: gexp-or-string-or-false)

File from which kernel messages are read, `/dev/kmsg` by default.

**message-destination** (default: `#f`) (type: gexp-or-false)

This gexp must evaluate to a procedure that, when passed a log message, returns the list of files to write it to; `#f` is equivalent to using `(default-message-destination-procedure)`. See Section “System Log Service” in *The GNU Shepherd Manual*, for information on how to write that procedure.

**date-format** (type: gexp-or-string)

String or string-valued gexp specifying how to format timestamps in log files. It must be a valid string for `strftime` (see Section “Time” in *GNU Guile Reference Manual*), including delimiting space—e.g., `"%c "` for a format identical to that of traditional `syslogd` implementations.

**history-size** (type: gexp-or-integer)

Number of logging messages kept in memory for the purposes of making them available to `herd status system-log`.

**max-silent-time** (type: gexp-or-integer-or-false)

Time after which a mark is written to log files if nothing was logged during that time frame. When set to `#f`, this feature is disabled.

**%shepherd-root-service** [Variable]

This service represents PID 1.

### 11.19.5 Complex Configurations

Some programs might have rather complex configuration files or formats, and to make it easier to create Scheme bindings for these configuration files, you can use the utilities defined in the `(gnu services configuration)` module.

The main utility is the `define-configuration` macro, a helper used to define a Scheme record type (see Section “Record Overview” in *GNU Guile Reference Manual*). The fields from this Scheme record can be serialized using *serializers*, which are procedures that take some kind of Scheme value and translates them into another Scheme value or Section 8.12 [G-Expressions], page 165.

**define-configuration** *name clause1 clause2 ...* [Macro]

Create a record type named *name* that contains the fields found in the clauses.

A clause has the following form:

```
(field-name
 type-decl
 documentation
 option*
 ...)
```

*field-name* is an identifier that denotes the name of the field in the generated record.

*type-decl* is either *type* for fields that require a value to be set or (*type default-value*) otherwise.

*type* is the type of the value corresponding to *field-name*; since Guile is untyped, a predicate procedure—*type?*—will be called on the value corresponding to the field to ensure that the value is of the correct type. This means that if say, *type* is *package*, then a procedure named *package?* will be applied on the value to make sure that it is indeed a *<package>* object.

*default-value* is the default value corresponding to the field; if none is specified, the user is forced to provide a value when creating an object of the record type.

*documentation* is a string formatted with Texinfo syntax which should provide a description of what setting this field does.

*option\** is one of the following subclauses:

**empty-serializer**

Exclude this field from serialization.

(**serializer** *serializer*)

*serializer* is the name of a procedure which takes two arguments, the first is the name of the field, and the second is the value corresponding to the field. The procedure should return a string or Section 8.12 [G-Expressions], page 165, that represents the content that will be serialized to the configuration file. If none is specified, a procedure of the name *serialize-type* will be used.

An example of a simple serializer procedure:

```
(define (serialize-boolean field-name value)
  (let ((value (if value "true" "false")))
    #~(string-append '#$field-name " = " #$value)))
```

(**sanitizer** *sanitizer*)

*sanitizer* is a procedure which takes one argument, a user-supplied value, and returns a “sanitized” value for the field. If no sanitizer is specified, a



default sanitizer is used, which raises an error if the value is not of type *type*.

An example of a sanitizer for a field that accepts both strings and symbols looks like this:

```
(define (sanitize-foo value)
  (cond ((string? value) value)
        ((symbol? value) (symbol->string value))
        (else (error "bad value"))))
```

In some cases multiple different configuration records might be defined in the same file, but their serializers for the same type might have to be different, because they have different configuration formats. For example, the `serialize-boolean` procedure for the Getmail service would have to be different from the one for the Transmission service. To make it easier to deal with this situation, one can specify a serializer prefix by using the `prefix` literal in the `define-configuration` form. This means that one doesn't have to manually specify a custom *serializer* for every field.

```
(define (foo-serialize-string field-name value)
  ...)

(define (bar-serialize-string field-name value)
  ...)

(define-configuration foo-configuration
  (label
   string
   "The name of label.")
  (prefix foo-))

(define-configuration bar-configuration
  (ip-address
   string
   "The IPv4 address for this device.")
  (prefix bar-))
```

However, in some cases you might not want to serialize any of the values of the record, to do this, you can use the `no-serialization` literal. There is also the `define-configuration/no-serialization` macro which is a shorthand of this.

```
;; Nothing will be serialized to disk.
(define-configuration foo-configuration
  (field
   (string "test")
   "Some documentation.")
  (no-serialization))

;; The same thing as above.
(define-configuration/no-serialization bar-configuration
  (field
```

```
(string "test")
"Some documentation."))
```

**define-maybe *type*** [Macro]

Sometimes a field should not be serialized if the user doesn't specify a value. To achieve this, you can use the **define-maybe** macro to define a “maybe type”; if the value of a maybe type is left unset, or is set to the **%unset-value** value, then it will not be serialized.

When defining a “maybe type”, the corresponding serializer for the regular type will be used by default. For example, a field of type **maybe-string** will be serialized using the **serialize-string** procedure by default, you can of course change this by specifying a custom serializer procedure. Likewise, the type of the value would have to be a string, or left unspecified.

```
(define-maybe string)

(define (serialize-string field-name value)
  ...)

(define-configuration baz-configuration
  (name
   ;; If set to a string, the `serialize-string' procedure will be used
   ;; to serialize the string. Otherwise this field is not serialized.
   maybe-string
   "The name of this module."))
```

Like with **define-configuration**, one can set a prefix for the serializer name by using the **prefix** literal.

```
(define-maybe integer
  (prefix baz-))

(define (baz-serialize-integer field-name value)
  ...)
```

There is also the **no-serialization** literal, which when set means that no serializer will be defined for the “maybe type”, regardless of whether its value is set or not. **define-maybe/no-serialization** is a shorthand for specifying the **no-serialization** literal.

```
(define-maybe/no-serialization symbol)

(define-configuration/no-serialization test-configuration
  (mode
   maybe-symbol
   "Docstring."))
```

**maybe-value-set? *value*** [Procedure]

Predicate to check whether a user explicitly specified the value of a maybe field.

**serialize-configuration** *configuration fields* [Procedure]

Return a G-expression that contains the values corresponding to the *fields* of *configuration*, a record that has been generated by **define-configuration**. The G-expression can then be serialized to disk by using something like **mixed-text-file**.

Once you have defined a configuration record, you will most likely also want to document it so that other people know to use it. To help with that, there are two procedures, both of which are documented below.

**generate-documentation** *documentation documentation-name* [Procedure]

Generate a Texinfo fragment from the docstrings in *documentation*, a list of (*label fields sub-documentation ...*). *label* should be a symbol and should be the name of the configuration record. *fields* should be a list of all the fields available for the configuration record.

*sub-documentation* is a (*field-name configuration-name*) tuple. *field-name* is the name of the field which takes another configuration record as its value, and *configuration-name* is the name of that configuration record. The same value may be used for multiple *field-names*, in case a field accepts different types of configurations.

*sub-documentation* is only needed if there are nested configuration records. For example, the **getmail-configuration** record (see Section 11.10.13 [Mail Services], page 407) accepts a **getmail-configuration-file** record in one of its **rcfile** field, therefore documentation for **getmail-configuration-file** is nested in **getmail-configuration**.

```
(generate-documentation
  `((getmail-configuration ,getmail-configuration-fields
    (rcfile getmail-configuration-file))
    ...)
  'getmail-configuration)
```

*documentation-name* should be a symbol and should be the name of the configuration record.

**configuration->documentation** *configuration-symbol* [Procedure]

Take *configuration-symbol*, the symbol corresponding to the name used when defining a configuration record with **define-configuration**, and print the Texinfo documentation of its fields. This is useful if there aren't any nested configuration records since it only prints the documentation for the top-level fields.

As of right now, there is no automated way to generate documentation for configuration records and put them in the manual. Instead, every time you make a change to the docstrings of a configuration record, you have to manually call **generate-documentation** or **configuration->documentation**, and paste the output into the **doc/guix.texi** file.

Below is an example of a record type created using **define-configuration** and friends.

```
(use-modules (gnu services)
  (guix gexp)
  (gnu services configuration)
  (srfi srfi-26)
  (srfi srfi-1))
```

```

;; Turn field names, which are Scheme symbols into strings
(define (uglify-field-name field-name)
  (let ((str (symbol->string field-name)))
    ;; field? -> is-field
    (if (string-suffix? "?" str)
        (string-append "is-" (string-drop-right str 1))
        str)))

(define (serialize-string field-name value)
  #~(string-append #$(uglify-field-name field-name) " = " #$(value) "\n"))

(define (serialize-integer field-name value)
  (serialize-string field-name (number->string value)))

(define (serialize-boolean field-name value)
  (serialize-string field-name (if value "true" "false")))

(define (serialize-contact-name field-name value)
  #~(string-append "\n[" #$(value) "]\n"))

(define (list-of-contact-configurations? lst)
  (every contact-configuration? lst))

(define (serialize-list-of-contact-configurations field-name value)
  #~(string-append #$(map (cut serialize-configuration <>
                           contact-configuration-fields)
                          value)))

(define (serialize-contacts-list-configuration configuration)
  (mixed-text-file
   "contactrc"
   #~(string-append "[Owner]\n"
                     #$(serialize-configuration
                        configuration contacts-list-configuration-fields))))■

(define-maybe integer)
(define-maybe string)

(define-configuration contact-configuration
  (name
   string
   "The name of the contact."
   serialize-contact-name)
  (phone-number
   maybe-integer
   "The person's phone number.")

```

```

(email
  maybe-string
  "The person's email address.")
(married?
  boolean
  "Whether the person is married.))

(define-configuration contacts-list-configuration
  (name
    string
    "The name of the owner of this contact list.")
  (email
    string
    "The owner's email address.")
  (contacts
    (list-of-contact-configurations '())
    "A list of @code{contact-configuration} records which contain
    information about all your contacts.))

```

A contacts list configuration could then be created like this:

```

(define my-contacts
  (contacts-list-configuration
    (name "Alice")
    (email "alice@example.org")
    (contacts
      (list (contact-configuration
        (name "Bob")
        (phone-number 1234)
        (email "bob@gnu.org")
        (married? #f))
        (contact-configuration
        (name "Charlie")
        (phone-number 0000)
        (married? #t))))))

```

After serializing the configuration to disk, the resulting file would look like this:

```

[owner]
name = Alice
email = alice@example.org

[Bob]
phone-number = 1234
email = bob@gnu.org
is-married = false

[Charlie]
phone-number = 0
is-married = true

```

## 12 System Troubleshooting Tips

Guix System allows rebooting into a previous generation should the last one be malfunctioning, which makes it quite robust against being broken irreversibly. This feature depends on GRUB being correctly functioning though, which means that if for whatever reasons your GRUB installation becomes corrupted during a system reconfiguration, you may not be able to easily boot into a previous generation. A technique that can be used in this case is to *chroot* into your broken system and reconfigure it from there. Such technique is explained below.

### 12.1 Chrooting into an existing system

This section details how to *chroot* to an already installed Guix System with the aim of reconfiguring it, for example to fix a broken GRUB installation. The process is similar to how it would be done on other GNU/Linux systems, but there are some Guix System particularities such as the daemon and profiles that make it worthy of explaining here.

1. Obtain a bootable image of Guix System. It is recommended the latest development snapshot so the kernel and the tools used are at least as new as those of your installed system; it can be retrieved from the <https://ci.guix.gnu.org> ([https://ci.guix.gnu.org/search/latest/ISO-9660?query=spec:images+status:success+system:x86\\_64-linux+image.iso](https://ci.guix.gnu.org/search/latest/ISO-9660?query=spec:images+status:success+system:x86_64-linux+image.iso)) URL. Follow the see Section 3.3 [USB Stick and DVD Installation], page 23, section for copying it to a bootable media.
2. Boot the image, and proceed with the graphical text-based installer until your network is configured. Alternatively, you could configure the network manually by following the [manual-installation-networking], page 27, section. If you get the error ‘RTNETLINK answers: Operation not possible due to RF-kill’, try ‘rfkill list’ followed by ‘rfkill unblock 0’, where ‘0’ is your device identifier (ID).
3. Switch to a virtual console (tty) if you haven’t already by pressing simultaneously the *Control + Alt + F4* keys. Mount your file system at */mnt*. Assuming your root partition is */dev/sda2*, you would do:

```
mount /dev/sda2 /mnt
```

4. Mount special block devices and Linux-specific directories:

```
mount --rbind /proc /mnt/proc
mount --rbind /sys /mnt/sys
mount --rbind /dev /mnt/dev
```

If your system is EFI-based, you must also mount the ESP partition. Assuming it is */dev/sda1*, you can do so with:

```
mount /dev/sda1 /mnt/boot/efi
```

5. Enter your system via *chroot*:

```
chroot /mnt /bin/sh
```

6. Source the system profile as well as your *user* profile to setup the environment, where *user* is the user name used for the Guix System you are attempting to repair:

```
source /var/guix/profiles/system/profile/etc/profile
```

```
source /home/user/.guix-profile/etc/profile
```

To ensure you are working with the Guix revision you normally would as your normal user, also source your current Guix profile:

```
source /home/user/.config/guix/current/etc/profile
```

7. Start a minimal `guix-daemon` in the background:

```
guix-daemon --build-users-group=guixbuild --disable-chroot &
```

8. Edit your Guix System configuration if needed, then reconfigure with:

```
guix system reconfigure your-config.scm
```

9. Finally, you should be good to reboot the system to test your fix.

## 13 Home Configuration

Guix supports declarative configuration of *home environments* by utilizing the configuration mechanism described in the previous chapter (see Section 11.19 [Defining Services], page 709), but for user’s dotfiles and packages. It works both on Guix System and foreign distros and allows users to declare all the packages and services that should be installed and configured for the user. Once a user has written a file containing a **home-environment** record, such a configuration can be *instantiated* by an unprivileged user with the **guix home** command (see Section 13.4 [Invoking guix home], page 784).

The user’s home environment usually consists of three basic parts: software, configuration, and state. Software in mainstream distros are usually installed system-wide, but with GNU Guix most software packages can be installed on a per-user basis without needing root privileges, and are thus considered part of the user’s *home environment*. Packages on their own are not very useful in many cases, because often they require some additional configuration, usually config files that reside in `XD_CONFIG_HOME` (`~/.config` by default) or other directories. Everything else can be considered state, like media files, application databases, and logs.

Using Guix for managing home environments provides a number of advantages:

- All software can be configured in one language (Guile Scheme), this gives users the ability to share values between configurations of different programs.
- A well-defined home environment is self-contained and can be created in a declarative and reproducible way—there is no need to grab external binaries or manually edit some configuration file.
- After every **guix home reconfigure** invocation, a new home environment generation will be created. This means that users can rollback to a previous home environment generation so they don’t have to worry about breaking their configuration.
- It is possible to manage stateful data with Guix Home, this includes the ability to automatically clone Git repositories on the initial setup of the machine, and periodically running commands like **rsync** to sync data with another host. This functionality is still in an experimental stage, though.

### 13.1 Declaring the Home Environment

The home environment is configured by providing a **home-environment** declaration in a file that can be passed to the **guix home** command (see Section 13.4 [Invoking guix home], page 784). The easiest way to get started is by generating an initial configuration with **guix home import**:

```
guix home import ~/src/guix-config
```

The **guix home import** command reads some of the “dot files” such as `~/.bashrc` found in your home directory and copies them to the given directory, `~/src/guix-config` in this case; it also reads the contents of your profile, `~/.guix-profile`, and, based on that, it populates `~/src/guix-config/home-configuration.scm` with a Home configuration that resembles your current configuration.

A simple setup can include Bash and a custom text configuration, like in the example below. Don’t be afraid to declare home environment parts, which overlaps with your current



dot files: before installing any configuration files, Guix Home will back up existing config files to a separate place in the home directory.

**Note:** It is highly recommended that you manage your shell or shells with Guix Home, because it will make sure that all the necessary scripts are sourced by the shell configuration file. Otherwise you will need to do it manually. (see Section 13.2 [Configuring the Shell], page 734).

```
(use-modules (gnu home)
             (gnu home services)
             (gnu home services shells)
             (gnu services)
             (gnu packages admin)
             (guix gexp))

(home-environment
 (packages (list htop))
 (services
  (append (list
            (service home-bash-service-type
                     (home-bash-configuration
                      (guix-defaults? #t)
                      (variables
                       `(("HISTFILE" . "$XDG_CACHE_HOME/.bash_history")
                        ("HISTSIZE" . "50000")))))
            (simple-service 'test-config
                           home-xdg-configuration-files-service-type
                           (list `("test.conf"
                                   ,(plain-file "tmp-file.txt"
  "the content of
~/.config/test.conf")))))
          %base-home-services)))
```

The `packages` field should be self-explanatory, it will install the list of packages into the user's profile. The most important field is `services`, it contains a list of *home services*, which are the basic building blocks of a home environment.

There is no daemon (at least not necessarily) related to a home service, a home service is just an element that is used to declare part of home environment and extend other parts of it. The extension mechanism discussed in the previous chapter (see Section 11.19 [Defining Services], page 709) should not be confused with Shepherd services (see Section 11.19.4 [Shepherd Services], page 717). Using this extension mechanism and some Scheme code that glues things together gives the user the freedom to declare their own, very custom, home environments.

Once the configuration looks good, you can first test it in a throw-away “container”:

```
guix home container config.scm
```

The command above spawns a shell where your home environment is running. The shell runs in a container, meaning it's isolated from the rest of the system, so it's a good way to try out your configuration—you can see if configuration bits are missing or misbehaving, if daemons get started, and so on. Once you exit that shell, you're back to the prompt of your original shell “in the real world”.

Once you have a configuration file that suits your needs, you can reconfigure your home by running:

```
guix home reconfigure config.scm
```

This “builds” your home environment and creates `~/.guix-home` pointing to it. Voilà!

**Note:** Make sure the operating system has `elogind`, `systemd`, or a similar mechanism to create the XDG run-time directory and has the `XDGRUNTIMEDIR` variable set. Failing that, the `on-first-login` script will not execute anything, and processes like user Shepherd and its descendants will not start.

If you're using Guix System, you can embed your home configuration in your system configuration such that `guix system reconfigure` will deploy both the system *and* your home at once! See [guix-home-service-type], page 638, for how to do that.

## 13.2 Configuring the Shell

This section is safe to skip if your shell or shells are managed by Guix Home. Otherwise, read it carefully.

There are a few scripts that must be evaluated by a login shell to activate the home environment. The shell startup files only read by login shells often have `profile` suffix. For more information about login shells see Section “Invoking Bash” in *The GNU Bash Reference Manual* and see Section “Bash Startup Files” in *The GNU Bash Reference Manual*.

The first script that needs to be sourced is `setup-environment`, which sets all the necessary environment variables (including variables declared by the user) and the second one is `on-first-login`, which starts Shepherd for the current user and performs actions declared by other home services that extends `home-run-on-first-login-service-type`.

Guix Home will always create `~/.profile`, which contains the following lines:

```
HOME_ENVIRONMENT=$HOME/.guix-home
. $HOME_ENVIRONMENT/setup-environment
$HOME_ENVIRONMENT/on-first-login
```

This makes POSIX compliant login shells activate the home environment. However, in most cases this file won't be read by most modern shells, because they are run in non POSIX mode by default and have their own `*profile` startup files. For example Bash will prefer `~/.bash_profile` in case it exists and only if it doesn't will it fallback to `~/.profile`. Zsh (if no additional options are specified) will ignore `~/.profile`, even if `~/.zprofile` doesn't exist.

To make your shell respect `~/.profile`, add `. ~/.profile` or `source ~/.profile` to the startup file for the login shell. In case of Bash, it is `~/.bash_profile`, and in case of Zsh, it is `~/.zprofile`.

**Note:** This step is only required if your shell is *not* managed by Guix Home. Otherwise, everything will be done automatically.

## 13.3 Home Services

A *home service* is not necessarily something that has a daemon and is managed by Shepherd (see Section “Jump Start” in *The GNU Shepherd Manual*), in most cases it doesn’t. It’s a simple building block of the home environment, often declaring a set of packages to be installed in the home environment profile, a set of config files to be symlinked into `XDG_CONFIG_HOME` (`~/.config` by default), and environment variables to be set by a login shell.

There is a service extension mechanism (see Section 11.19.1 [Service Composition], page 709) which allows home services to extend other home services and utilize capabilities they provide; for example: declare `mcron` jobs (see *GNU Mcron*) by extending Section 13.3.3 [Mcron Home Service], page 745; declare daemons by extending Section 13.3.5 [Shepherd Home Service], page 747; add commands, which will be invoked on by the Bash by extending Section 13.3.2 [Shells Home Services], page 740.

A good way to discover available home services is using the `guix home search` command (see Section 13.4 [Invoking guix home], page 784). After the required home services are found, include its module with the `use-modules` form (see Section “Using Guile Modules” in *The GNU Guile Reference Manual*), or the `#:use-modules` directive (see Section “Creating Guile Modules” in *The GNU Guile Reference Manual*) and declare a home service using the `service` function, or extend a service type by declaring a new service with the `simple-service` procedure from (`gnu services`).

### 13.3.1 Essential Home Services

There are a few essential home services defined in (`gnu home services`), they are mostly for internal use and are required to build a home environment, but some of them will be useful for the end user.

`home-environment-variables-service-type` [Variable]

The service of this type will be instantiated by every home environment automatically by default, there is no need to define it, but someone may want to extend it with a list of pairs to set some environment variables.

```
(list ("ENV_VAR1" . "value1")
      ("ENV_VAR2" . "value2"))
```

The easiest way to extend a service type, without defining a new service type is to use the `simple-service` helper from (`gnu services`).

```
(simple-service 'some-useful-env-vars-service
home-environment-variables-service-type
`(("LESSHISTFILE" . "$XDG_CACHE_HOME/.lesshist")
  ("SHELL" . ,(file-append zsh "/bin/zsh"))
  ("USELESS_VAR" . #f)
  ("_JAVA_AWT_WM_NONREParenting" . #t)
  ("LITERAL_VALUE" . ,(literal-string "${abc}"))))
```

If you include such a service in you home environment definition, it will add the following content to the `setup-environment` script (which is expected to be sourced by the login shell):

```
export LESSHISTFILE="$XDG_CACHE_HOME/.lesshist"
export SHELL="/gnu/store/2hsg15n644f0glrcbkb1kqknmmqdar03-zsh-5.8/bin/zsh"■
```

```
export _JAVA_AWT_WM_NONREParenting
export LITERAL_VALUE='${abc}'
```

Notice that `literal-string` above lets us declare that a value is to be interpreted as a *literal string*, meaning that “special characters” such as the dollar sign will not be interpreted by the shell.

**Note:** Make sure that module (`gnu packages shells`) is imported with `use-modules` or any other way, this namespace contains the definition of the `zsh` package, which is used in the example above.

The association list (see Section “Association Lists” in *The GNU Guile Reference manual*) is a data structure containing key-value pairs, for `home-environment-variables-service-type` the key is always a string, the value can be a string, string-valued gexp (see Section 8.12 [G-Expressions], page 165), file-like object (see Section 8.12 [G-Expressions], page 165) or boolean. For gexps, the variable will be set to the value of the gexp; for file-like objects, it will be set to the path of the file in the store (see Section 8.9 [The Store], page 156); for `#t`, it will export the variable without any value; and for `#f`, it will omit variable.

**home-profile-service-type** [Variable]

The service of this type will be instantiated by every home environment automatically, there is no need to define it, but you may want to extend it with a list of packages if you want to install additional packages into your profile. Other services, which need to make some programs available to the user will also extend this service type.

The extension value is just a list of packages:

```
(list http vim emacs)
```

The same approach as `simple-service` (see Section 11.19.3 [Service Reference], page 712) for `home-environment-variables-service-type` can be used here, too. Make sure that modules containing the specified packages are imported with `use-modules`. To find a package or information about its module use `guix search` (see Section 5.2 [Invoking guix package], page 37). Alternatively, `specification->package` can be used to get the package record from a string without importing its related module.

There are few more essential services, but users are not expected to extend them.

**home-service-type** [Variable]

The root of home services DAG, it generates a folder, which later will be symlinked to `~/.guix-home`, it contains configurations, profile with binaries and libraries, and some necessary scripts to glue things together.

**home-run-on-first-login-service-type** [Variable]

The service of this type generates a Guile script, which is expected to be executed by the login shell. It is only executed if the special flag file inside `XDGRUNTIME_DIR` hasn’t been created, this prevents redundant executions of the script if multiple login shells are spawned.

It can be extended with a gexp. However, to autostart an application, users *should not* use this service, in most cases it’s better to extend `home-shepherd-service-type` with a Shepherd service (see Section 11.19.4 [Shepherd Services], page 717), or extend the shell’s startup file with the required command using the appropriate service type.

**home-files-service-type** [Variable]

The service of this type allows to specify a list of files, which will go to `~/.guix-home/files`, usually this directory contains configuration files (to be more precise it contains symlinks to files in `/gnu/store`), which should be placed in `$XDG_CONFIG_HOME` or in rare cases in `$HOME`. It accepts extension values in the following format:

```
((".sway/config" ,sway-file-like-object)
  (".tmux.conf" ,(local-file "./tmux.conf")))
```

Each nested list contains two values: a subdirectory and file-like object. After building a home environment `~/.guix-home/files` will be populated with appropriate content and all nested directories will be created accordingly, however, those files won't go any further until some other service will do it. By default a **home-symlink-manager-service-type**, which creates necessary symlinks in home folder to files from `~/.guix-home/files` and backs up already existing, but clashing configs and other things, is a part of essential home services (enabled by default), but it's possible to use alternative services to implement more advanced use cases like read-only home. Feel free to experiment and share your results.

It is often the case that Guix Home users already have a setup for versioning their user configuration files (also known as *dot files*) in a single directory, and some way of automatically deploy changes to their user home.

The **home-dotfiles-service-type** from (`gnu home services dotfiles`) is designed to ease the way into using Guix Home for this kind of users, allowing them to point the service to their dotfiles directory without migrating them to Guix native configurations.

Please keep in mind that it is advisable to keep your dotfiles directories under version control, for example in the same repository where you'd track your Guix Home configuration.

There are two supported dotfiles directory layouts, for now. The **'plain** layout, which is structured as follows:

```
~$ tree -a ./dotfiles/
dotfiles/
.gitconfig
.gnupg
  gpg-agent.conf
  gpg.conf
.guile
.config
  guix
    channels.scm
  nixpkgs
    config.nix
.nix-channels
.tmux.conf
.vimrc
```

This tree structure is installed as is to the home directory upon **guix home reconfigure**.

The **'stow** layout, which must follow the layout suggested by GNU Stow (<https://www.gnu.org/software/stow/>) presents an additional application specific directory layer, just like:

```

~$ tree -a ./dotfiles/
dotfiles/
  git
    .gitconfig
  gpg
    .gnupg
      gpg-agent.conf
      gpg.conf
  guile
    .guile
  guix
    .config
      guix
        channels.scm
  nix
    .config
      nixpkgs
        config.nix
    .nix-channels
  tmux
    .tmux.conf
  vim
    .vimrc

```

13 directories, 10 files

For an informal specification please refer to the Stow manual (see Section “Introduction” in *The GNU Stow Manual*). This tree structure is installed following GNU Stow’s logic to the home directory upon **guix home reconfigure**.

A suitable configuration with a 'plain layout could be:

```

(home-environment
  ;; ...
  (services
    (service home-dotfiles-service-type
      (home-dotfiles-configuration
        (directories '("./dotfiles"))))))

```

The expected home directory state would then be:

```

.
  .config
    guix
      channels.scm
    nixpkgs
      config.nix
  .gitconfig
  .gnupg
    gpg-agent.conf
    gpg.conf

```

```
.guile
.nix-channels
.tmux.conf
.vimrc
```

**home-dotfiles-service-type** [Variable]

Return a service which is very similar to **home-files-service-type** (and actually extends it), but designed to ease the way into using Guix Home for users that already track their dotfiles under some kind of version control. This service allows users to point Guix Home to their dotfiles directory and have their files automatically provisioned to their home directory, without migrating all of their dotfiles to Guix native configurations.

**home-dotfiles-configuration** [Data Type]

Available **home-dotfiles-configuration** fields are:

**source-directory** (default: (**current-source-directory**)) (type: string)

The path where dotfile directories are resolved. By default dotfile directories are resolved relative the source location where **home-dotfiles-configuration** appears.

**layout** (default: 'plain) (type: symbol)

The intended layout of the specified directory. It can be either 'stow or 'plain.

**directories** (default: '()) (type: list-of-strings)

The list of dotfiles directories where **home-dotfiles-service-type** will look for application dotfiles.

**packages** (type: maybe-list-of-strings)

The names of a subset of the GNU Stow package layer directories. When provided the **home-dotfiles-service-type** will only provision dotfiles from this subset of applications. This field will be ignored if **layout** is set to 'plain.

**excluded** (default: '(".".\*~" ".\*\.\.swp" "\.\.git" "\.\.gitignore")) (type: list-of-strings)

The list of file patterns **home-dotfiles-service-type** will exclude while visiting each one of the **directories**.

**home-xdg-configuration-files-service-type** [Variable]

The service is very similar to **home-files-service-type** (and actually extends it), but used for defining files, which will go to ~/.guix-home/files/.config, which will be symlinked to \$XDG\_CONFIG\_HOME by **home-symlink-manager-service-type** (for example) during activation. It accepts extension values in the following format:

```
`(("sway/config" ,sway-file-like-object)
;; -> ~/.guix-home/files/.config/sway/config
;; -> $XDG_CONFIG_HOME/sway/config (by symlink-manager)
("tmux/tmux.conf" ,(local-file "/tmux.conf")))
```

**home-activation-service-type** [Variable]

The service of this type generates a guile script, which runs on every **guix home reconfigure** invocation or any other action, which leads to the activation of the home environment.

**home-symlink-manager-service-type** [Variable]

The service of this type generates a guile script, which will be executed during activation of home environment, and do a few following steps:

1. Reads the content of **files/** directory of current and pending home environments.
2. Cleans up all symlinks created by **symlink-manager** on previous activation. Also, sub-directories, which become empty also will be cleaned up.
3. Creates new symlinks the following way: It looks **files/** directory (usually defined with **home-files-service-type**, **home-xdg-configuration-files-service-type** and maybe some others), takes the files from **files/.config/** subdirectory and put respective links in **XDG\_CONFIG\_HOME**. For example symlink for **files/.config/sway/config** will end up in **\$XDG\_CONFIG\_HOME/sway/config**. The rest files in **files/** outside of **files/.config/** subdirectory will be treated slightly different: symlink will just go to **\$HOME**. **files/.some-program/config** will end up in **\$HOME/.some-program/config**.
4. If some sub-directories are missing, they will be created.
5. If there is a clashing files on the way, they will be backed up.

**symlink-manager** is a part of essential home services and is enabled and used by default.

### 13.3.2 Shells

Shells play a quite important role in the environment initialization process, you can configure them manually as described in section Section 13.2 [Configuring the Shell], page 734, but the recommended way is to use home services listed below. It's both easier and more reliable.

Each home environment instantiates **home-shell-profile-service-type**, which creates a **~/.profile** startup file for all POSIX-compatible shells. This file contains all the necessary steps to properly initialize the environment, but many modern shells like Bash or Zsh prefer their own startup files, that's why the respective home services (**home-bash-service-type** and **home-zsh-service-type**) ensure that **~/.profile** is sourced by **~/.bash\_profile** and **~/.zprofile**, respectively.

### Shell Profile Service

**home-shell-profile-configuration** [Data Type]

Available **home-shell-profile-configuration** fields are:

**profile** (default: '()') (type: text-config)

**home-shell-profile** is instantiated automatically by **home-environment**, DO NOT create this service manually, it can only be extended. **profile** is a list of file-like objects, which will go to **~/.profile**. By default **~/.profile** contains the initialization code



which must be evaluated by the login shell to make home-environment's profile available to the user, but other commands can be added to the file if it is really necessary. In most cases shell's configuration files are preferred places for user's customizations. Extend home-shell-profile service only if you really know what you do.

## Bash Home Service

**home-bash-configuration** [Data Type]

Available **home-bash-configuration** fields are:

**package** (default: **bash**) (type: package)

The Bash package to use.

**guix-defaults?** (default: **#t**) (type: boolean)

Add sane defaults like setting **PS1**, **SHELL**, and ensuring **/etc/profile** is sourced for non-interactive SSH shells. If you use Guix System, is it safe to set this to **#f**, as in this case this is already taken care of by the globally installed **/etc/bashrc**.

**environment-variables** (default: **()**) (type: alist)

Association list of environment variables to set for the Bash session. The rules for the **home-environment-variables-service-type** apply here (see Section 13.3.1 [Essential Home Services], page 735). The contents of this field will be added after the contents of the **bash-profile** field.

**variables** (default: **()**) (type: alist)

Bash variables (not to be confused with environment variables) to be set in **.bashrc** for use by interactive shells. A typical example of such variables is **HISTSIZE** and related history variables (see Section “Bash Variables” in *Bash Reference Manual*).

**aliases** (type: alist)

Association list of aliases to set for the Bash session. The aliases will be defined after the contents of the **bashrc** field has been put in the **.bashrc** file. The alias will automatically be quoted, so something like this:

```
'(("ls" . "ls -a1F"))
```

turns into

```
alias ls="ls -a1F"
```

**bash-profile** (default: **()**) (type: text-config)

List of file-like objects, which will be added to **.bash\_profile**. Used for executing user's commands at start of login shell (In most cases the shell started on **tty** just after login). **.bash\_login** won't be ever read, because **.bash\_profile** always present.

**bashrc** (default: **()**) (type: text-config)

List of file-like objects, which will be added to **.bashrc**. Used for executing user's commands at start of interactive shell (The shell for interactive usage started by typing **bash** or by terminal app or any other program).

**bash-logout** (default: ()) (type: text-config)

List of file-like objects, which will be added to `.bash_logout`. Used for executing user's commands at the exit of login shell. It won't be read in some cases (if the shell terminates by exec'ing another process for example).

You can extend the Bash service by using the `home-bash-extension` configuration record, whose fields must mirror that of `home-bash-configuration` (see [home-bash-configuration], page 741). The contents of the extensions will be added to the end of the corresponding Bash configuration files (see Section “Bash Startup Files” in *The GNU Bash Reference Manual*).

For example, here is how you would define a service that extends the Bash service such that `~/.bash_profile` defines an additional shell variable, `PS1`:

```
(define bash-fancy-prompt-service
  (simple-service 'bash-fancy-prompt
    home-bash-service-type
    (home-bash-extension
      (variables
        '(("PS1" . "\\u \\wλ ") )))))
```

You would then add `bash-fancy-prompt-service` to the list in the `services` field of your `home-environment`. The reference of `home-bash-extension` follows.

**home-bash-extension**

[Data Type]

Available `home-bash-extension` fields are:

**environment-variables** (default: ()) (type: alist)

Additional environment variables to set. These will be combined with the environment variables from other extensions and the base service to form one coherent block of environment variables.

**variables** (default: ()) (type: alist)

Extra Bash variables (not to be confused with environment variables) to be set in `.bashrc` for use by interactive shells.

**aliases** (default: ()) (type: alist)

Additional aliases to set. These will be combined with the aliases from other extensions and the base service.

**bash-profile** (default: ()) (type: text-config)

Additional text blocks to add to `.bash_profile`, which will be combined with text blocks from other extensions and the base service.

**bashrc** (default: ()) (type: text-config)

Additional text blocks to add to `.bashrc`, which will be combined with text blocks from other extensions and the base service.

**bash-logout** (default: ()) (type: text-config)

Additional text blocks to add to `.bash_logout`, which will be combined with text blocks from other extensions and the base service.

## Zsh Home Service

**home-zsh-configuration** [Data Type]

Available **home-zsh-configuration** fields are:

**package** (default: **zsh**) (type: package)

The Zsh package to use.

**xdg-flavor?** (default: **#t**) (type: boolean)

Place all the configs to **\$XDG\_CONFIG\_HOME/zsh**. Makes **~/ .zshenv** to set **ZDOTDIR** to **\$XDG\_CONFIG\_HOME/zsh**. Shell startup process will continue with **\$XDG\_CONFIG\_HOME/zsh/.zshenv**.

**environment-variables** (default: '()') (type: alist)

Association list of environment variables to set for the Zsh session.

**zshenv** (default: '()') (type: text-config)

List of file-like objects, which will be added to **.zshenv**. Used for setting user's shell environment variables. Must not contain commands assuming the presence of **tty** or producing output. Will be read always. Will be read before any other file in **ZDOTDIR**.

**zprofile** (default: '()') (type: text-config)

List of file-like objects, which will be added to **.zprofile**. Used for executing user's commands at start of login shell (In most cases the shell started on **tty** just after login). Will be read before **.zlogin**.

**zshrc** (default: '()') (type: text-config)

List of file-like objects, which will be added to **.zshrc**. Used for executing user's commands at start of interactive shell (The shell for interactive usage started by typing **zsh** or by terminal app or any other program).

**zlogin** (default: '()') (type: text-config)

List of file-like objects, which will be added to **.zlogin**. Used for executing user's commands at the end of starting process of login shell.

**zlogout** (default: '()') (type: text-config)

List of file-like objects, which will be added to **.zlogout**. Used for executing user's commands at the exit of login shell. It won't be read in some cases (if the shell terminates by exec'ing another process for example).

## Inputrc Profile Service

The GNU Readline package (<https://tiswww.cwru.edu/php/chet/readline/rltop.html>) includes Emacs and vi editing modes, with the ability to customize the configuration with settings in the **~/ .inputrc** file. With the **gnu home services shells** module, you can setup your readline configuration in a predictable manner, as shown below. For more information about configuring an **~/ .inputrc** file, see Section "Readline Init File" in *GNU Readline*.

**home-inputrc-service-type** [Variable]

This is the service to setup various **.inputrc** configurations. The settings in **.inputrc** are read by all programs which are linked with GNU Readline.

Here is an example of a service and its configuration that you could add to the `services` field of your `home-environment`:

```
(service home-inputrc-service-type
  (home-inputrc-configuration
    (key-bindings
      `(("Control-l" . "clear-screen")))
    (variables
      `(("bell-style" . "visible")
        ("colored-completion-prefix" . #t)
        ("editing-mode" . "vi")
        ("show-mode-in-prompt" . #t)))
    (conditional-constructs
      `(("$if mode=vi" .
        , (home-inputrc-configuration
          (variables
            `(("colored-stats" . #t)
              ("enable-bracketed-paste" . #t))))
        ("$else" .
          , (home-inputrc-configuration
            (variables
              `(("show-all-if-ambiguous" . #t))))
        ("endif" . #t)
        ("$include" . "/etc/inputrc")
        ("$include" . ,(file-append
          (specification->package "readline")
          "/etc/inputrc"))))))))
```

The example above starts with a combination of `key-bindings` and `variables`. The `conditional-constructs` show how it is possible to add conditionals and includes. In the example above `colored-stats` is only enabled if the editing mode is `vi` style, and it also reads any additional configuration located in `/etc/inputrc` or in `/gnu/store/...-readline/etc/inputrc`.

The value associated with a `home-inputrc-service-type` instance must be a `home-inputrc-configuration` record, as described below.

### `home-inputrc-configuration`

[Data Type]

Available `home-inputrc-configuration` fields are:

**key-bindings** (default: '()) (type: alist)

Association list of readline key bindings to be added to the `~/.inputrc` file.

```
'(("Control-l" . "clear-screen"))
```

turns into

```
Control-l: clear-screen
```

**variables** (default: '()) (type: alist)

Association list of readline variables to set.

```
'(("bell-style" . "visible"))
```

```

        (\"colored-completion-prefix\" . #t))
turns into
        set bell-style visible
        set colored-completion-prefix on
conditional-constructs (default: '()) (type: alist)
    Association list of conditionals to add to the initialization file. This in-
    cludes $if, else, endif and include and they receive a value of another
    home-inputrc-configuration.
        (conditional-constructs
          `((\"$if mode=vi\" .
            , (home-inputrc-configuration
              (variables
                `((\"show-mode-in-prompt\" . #t))))))
          (\"$else\" .
            , (home-inputrc-configuration
              (key-bindings
                `((\"Control-l\" . \"clear-screen\"))))))
          (\"$endif\" . #t)))
turns into
        $if mode=vi
        set show-mode-in-prompt on
        $else
        Control-l: clear-screen
        $endif
extra-content (default: "") (type: text-config)
    Extra content appended as-is to the configuration file. Run man readline
    for more information about all the configuration options.

```

### 13.3.3 Scheduled User's Job Execution

The (`gnu home services mcron`) module provides an interface to GNU `mcron`, a daemon to run jobs at scheduled times (see *GNU mcron*). The information about system's `mcron` is applicable here (see [mcron-service], page 301), the only difference for home services is that they have to be declared in a `home-environment` record instead of an `operating-system` record.

**Note:** We recommend defining periodic tasks as Shepherd timers, which provide more flexibility than `mcron`. See Section 13.3.5 [Shepherd Home Service], page 747, for more info.

**home-mcron-service-type** [Variable]

This is the type of the `mcron` home service, whose value is a `home-mcron-configuration` object. It allows to manage scheduled tasks.

This service type can be the target of a service extension that provides additional job specifications (see Section 11.19.1 [Service Composition], page 709). In other words, it is possible to define services that provide additional `mcron` jobs to run.

**home-mcron-configuration** [Data Type]

Available **home-mcron-configuration** fields are:

- mcron** (default: **mcron**) (type: file-like)  
The mcron package to use.
- jobs** (default: '()') (type: list-of-gexps)  
This is a list of gexps (see Section 8.12 [G-Expressions], page 165), where each gexp corresponds to an mcron job specification (see Section “Syntax” in *GNU mcron*).
- log?** (default: **#t**) (type: boolean)  
Log messages to standard output.
- log-format** (default: "**~1@\*~a ~a: ~a~%**") (type: string)  
(**ice-9 format**) format string for log messages. The default value produces messages like "**pid name: message**" (see Section “Invoking mcron” in *GNU mcron*). Each message is also prefixed by a timestamp by GNU Shepherd.

### 13.3.4 Power Management Home Services

The (**gnu home services pm**) module provides home services pertaining to battery power.

**home-batsignal-service-type** [Variable]

Service for **batsignal**, a program that monitors battery levels and warns the user through desktop notifications when their battery is getting low. You can also configure a command to be run when the battery level passes a point deemed “dangerous”. This service is configured with the **home-batsignal-configuration** record.

**home-batsignal-configuration** [Data Type]

Data type representing the configuration for **batsignal**.

- warning-level** (default: 15)  
The battery level to send a warning message at.
- warning-message** (default: **#f**)  
The message to send as a notification when the battery level reaches the **warning-level**. Setting to **#f** uses the default message.
- critical-level** (default: 5)  
The battery level to send a critical message at.
- critical-message** (default: **#f**)  
The message to send as a notification when the battery level reaches the **critical-level**. Setting to **#f** uses the default message.
- danger-level** (default: 2)  
The battery level to run the **danger-command** at.
- danger-command** (default: **#f**)  
The command to run when the battery level reaches the **danger-level**. Setting to **#f** disables running the command entirely.

**full-level** (default: **#f**)  
 The battery level to send a full message at. Setting to **#f** disables sending the full message entirely.

**full-message** (default: **#f**)  
 The message to send as a notification when the battery level reaches the **full-level**. Setting to **#f** uses the default message.

**batteries** (default: **'()**)  
 The batteries to monitor. Setting to **'()** tries to find batteries automatically.

**poll-delay** (default: **60**)  
 The time in seconds to wait before checking the batteries again.

**icon** (default: **#f**)  
 A file-like object to use as the icon for battery notifications. Setting to **#f** disables notification icons entirely.

**notifications?** (default: **#t**)  
 Whether to send any notifications.

**notifications-expire?** (default: **#f**)  
 Whether notifications sent expire after a time.

**notification-command** (default: **#f**)  
 Command to use to send messages. Setting to **#f** sends a notification through **libnotify**.

**ignore-missing?** (default: **#f**)  
 Whether to ignore missing battery errors.

### 13.3.5 Managing User Daemons

The (**gnu home services shepherd**) module supports the definitions of per-user Shepherd services (see Section “Introduction” in *The GNU Shepherd Manual*). You extend **home-shepherd-service-type** with new services; Guix Home then takes care of starting the **shepherd** daemon for you when you log in, which in turns starts the services you asked for.

**home-shepherd-service-type** [Variable]

The service type for the userland Shepherd, which allows one to manage long-running processes or one-shot tasks. User’s Shepherd is not an init process (PID 1), but almost all other information described in (see Section 11.19.4 [Shepherd Services], page 717) is applicable here too.

This is the service type that extensions target when they want to create shepherd services (see Section 11.19.2 [Service Types and Services], page 710, for an example). Each extension must pass a list of **<shepherd-service>**. Its value must be a **home-shepherd-configuration**, as described below.

**home-shepherd-configuration** [Data Type]

This data type represents the Shepherd’s configuration.

**shepherd** (default: **shepherd**)  
 The Shepherd package to use.

```

auto-start? (default: #t)
    Whether or not to start Shepherd on first login.

daemonize? (default: #t)
    Whether or not to run Shepherd in the background.

silent? (default: #t)
    When true, the shepherd process does not write anything to standard
    output when started automatically.

services (default: '())
    A list of <shepherd-service> to start. You should probably use the
    service extension mechanism instead (see Section 11.19.4 [Shepherd Ser-
    vices], page 717).

```

The Shepherd allows you to define *timers*, a special type of service that performs a given task periodically. Just like you can define timers at the system level (see Section 11.10.2 [Scheduled Job Execution], page 299), you can do so in your home environment.

The example below defines a home environment where a Shepherd timer runs the **mkid** command twice per day (see Section “mkid invocation” in *ID Database Utilities*). It does so by extending **home-shepherd-service-type** with **simple-service**; the Shepherd timer itself is produced by the **shepherd-timer** procedure (see [shepherd-timer-procedure], page 300), which is given the service name, a gexp specifying its schedule, and a gexp specifying the command to run.

```

(use-modules (gnu) (guix)
             (gnu home services shepherd)
             (gnu packages idutils))

(define idutils-service
  ;; Index my 'doc' directory everyday at 12:15PM and 7:15PM.
  (simple-service
   'update-idutils-database home-shepherd-service-type
   (list (shepherd-timer '(update-idutils-database)
                         #~(calendar-event #:hours '(12 19)
   #:minutes '(15))
                         #~(file-append idutils "/bin/mkid")
                         "doc")))))

(home-environment
 ;; ...
 (services (append (list idutils-service)
                   %base-home-services)))

```

See Section “Timers” in *The GNU Shepherd Manual* for more information on Shepherd timers.

The Shepherd also comes with a *log rotation service*, which compresses and then deletes old log files produced by services and daemons that it runs. This service is made available through **home-log-rotation-service-type** as described below.



**home-log-rotation-service-type** [Variable]

This is the service type for the user Shepherd log rotation service (see Section “Log Rotation Service” in *The GNU Shepherd Manual*). Its value must be a **log-rotation-configuration** record, exactly as for its system-wide counterpart. See [log-rotation-configuration], page 304, for its reference.

This service is part of **%base-home-services**.

**home-shepherd-transient-service-type** [Variable]

**home-shepherd-timer-service-type** [Variable]

These are the **timer** and **transient** Shepherd services. The former lets you schedule command execution for later, while the latter can run commands in the background as a regular service.

See [shepherd-transient-timer-services], page 722, which are their Guix System counterparts, for more info.

### 13.3.6 Secure Shell

The OpenSSH package (<https://www.openssh.com>) includes a client, the **ssh** command, that allows you to connect to remote machines using the SSH (secure shell) protocol. With the (**gnu home services ssh**) module, you can set up OpenSSH so that it works in a predictable fashion, almost independently of state on the local machine. To do that, you instantiate **home-openssh-service-type** in your Home configuration, as explained below.

**home-openssh-service-type** [Variable]

This is the type of the service to set up the OpenSSH client. It takes care of several things:

- providing a `~/.ssh/config` file based on your configuration so that **ssh** knows about hosts you regularly connect to and their associated parameters;
- providing a `~/.ssh/authorized_keys`, which lists public keys that the local SSH server, **sshd**, may accept to connect to this user account;
- optionally providing a `~/.ssh/known_hosts` file so that **ssh** can authenticate hosts you connect to.

Here is an example of a service and its configuration that you could add to the **services** field of your **home-environment**:

```
(service home-openssh-service-type
  (home-openssh-configuration
    (hosts
      (list (openssh-host (name "ci.guix.gnu.org")
                          (user "charlie"))
            (openssh-host (name "chbouib")
                          (host-name "chbouib.example.org")
                          (user "supercharlie")
                          (port 10022))))
    (authorized-keys (list (local-file "alice.pub")))))
```

The example above lists two hosts and their parameters. For instance, running **ssh chbouib** will automatically connect to **chbouib.example.org** on port 10022, log-

ging in as user ‘supercharlie’. Further, it marks the public key in `alice.pub` as authorized for incoming connections.

The value associated with a `home-openssh-service-type` instance must be a `home-openssh-configuration` record, as describe below.

**home-openssh-configuration** [Data Type]

This is the datatype representing the OpenSSH client and server configuration in one’s home environment. It contains the following fields:

**hosts** (default: ‘()’)

A list of `openssh-host` records specifying host names and associated connection parameters (see below). This host list goes into `~/.ssh/config`, which `ssh` reads at startup.

**known-hosts** (default: `*unspecified*`)

This must be either:

- `*unspecified*`, in which case `home-openssh-service-type` leaves it up to `ssh` and to the user to maintain the list of known hosts at `~/.ssh/known_hosts`, or
- a list of file-like objects, in which case those are concatenated and emitted as `~/.ssh/known_hosts`.

The `~/.ssh/known_hosts` contains a list of host name/host key pairs that allow `ssh` to authenticate hosts you connect to and to detect possible impersonation attacks. By default, `ssh` updates it in a *TOFU*, *trust-on-first-use* fashion, meaning that it records the host’s key in that file the first time you connect to it. This behavior is preserved when `known-hosts` is set to `*unspecified*`.

If you instead provide a list of host keys upfront in the `known-hosts` field, your configuration becomes self-contained and stateless: it can be replicated elsewhere or at another point in time. Preparing this list can be relatively tedious though, which is why `*unspecified*` is kept as a default.

**known-hosts2** (default: `*unspecified*`)

Either `*unspecified*` or a list of file-like objects containing known host keys. When specified, these files will be concatenated and installed as `~/.ssh/known_hosts2`. This allows for declaratively manage known hosts (stored in `known_hosts2`), while leaving the primary `known_hosts` file unmanaged by Guix.

The SSH client by default reads both `known-hosts` and `known-hosts2` for host key verification. If only `known-hosts2` is specified in Guix Home, `known_hosts` remains as a regular file that SSH can modify, while `known_hosts2` contains the declaratively managed host keys.

**authorized-keys** (default: `#false`)

The default `#false` value means: Leave any `~/.ssh/authorized_keys` file alone. Otherwise, this must be a list of file-like objects, each of which

containing an SSH public key that should be authorized to connect to this machine.

Concretely, these files are concatenated and made available as `~/.ssh/authorized_keys`. If an OpenSSH server, `sshd`, is running on this machine, then it *may* take this file into account: this is what `sshd` does by default, but be aware that it can also be configured to ignore it.

**add-keys-to-agent** (default: no)

This string specifies whether keys should be automatically added to a running `ssh-agent`. If this option is set to **yes** and a key is loaded from a file, the key and its passphrase are added to the agent with the default lifetime, as if by `ssh-add`. If this option is set to **ask**, `ssh` will require confirmation. If this option is set to **confirm**, each use of the key must be confirmed. If this option is set to **no**, no keys are added to the agent. Alternately, this option may be specified as a time interval to specify the key's lifetime in `ssh-agent`, after which it will automatically be removed. The argument must be **no**, **yes**, **confirm** (optionally followed by a time interval), **ask** or a time interval.

**openssh-host**

[Data Type]

Available **openssh-host** fields are:

**name** (type: string)

Name of this host declaration. A **openssh-host** must define only **name** or **match-criteria**. Use host-name `\"*\"` for top-level options.

**host-name** (type: maybe-string)

Host name—e.g., `"foo.example.org"` or `"192.168.1.2"`.

**match-criteria** (type: maybe-match-criteria)

When specified, this string denotes the set of hosts to which the entry applies, superseding the **host-name** field. Its first element must be all or one of **ssh-match-keywords**. The rest of the elements are arguments for the keyword, or other criteria. A **openssh-host** must define only **name** or **match-criteria**. Other host configuration options will apply to all hosts matching **match-criteria**.

**address-family** (type: maybe-address-family)

Address family to use when connecting to this host: one of **AF\_INET** (for IPv4 only), **AF\_INET6** (for IPv6 only). Additionally, the field can be left unset to allow any address family.

**identity-file** (type: maybe-string)

The identity file to use—e.g., `"/home/charlie/.ssh/id_ed25519"`.

**port** (type: maybe-natural-number)

TCP port number to connect to.

**user** (type: maybe-string)

User name on the remote host.

**forward-x11?** (type: maybe-boolean)  
Whether to forward remote client connections to the local X11 graphical display.

**forward-x11-trusted?** (type: maybe-boolean)  
Whether remote X11 clients have full access to the original X11 graphical display.

**forward-agent?** (type: maybe-boolean)  
Whether the authentication agent (if any) is forwarded to the remote machine.

**compression?** (type: maybe-boolean)  
Whether to compress data in transit.

**proxy** (type: maybe-proxy-command-or-jump-list)  
The command to use to connect to the server or a list of SSH hosts to jump through before connecting to the server. The field may be set to either a **proxy-command** or a list of **proxy-jump** records.  
As an example, a **proxy-command** to connect via an HTTP proxy at 192.0.2.0 would be constructed with: (**proxy-command** "nc -X connect -x 192.0.2.0:8080 %h %p").

**proxy-jump** [Data Type]  
Available **proxy-jump** fields are:

**user** (type: maybe-string)  
User name on the remote host.

**host-name** (type: string)  
Host name—e.g., `foo.example.org` or `192.168.1.2`.

**port** (type: maybe-natural-number)  
TCP port number to connect to.

**host-key-algorithms** (type: maybe-string-list)  
The list of accepted host key algorithms—e.g., `('ssh-ed25519')`.

**accepted-key-types** (type: maybe-string-list)  
The list of accepted user public key types.

**extra-content** (default: "") (type: raw-configuration-string)  
Extra content appended as-is to this **Host** block in `~/.ssh/config`.

The **parcimonie** service runs a daemon that slowly refreshes a GnuPG public key from a keyserver. It refreshes one key at a time; between every key update **parcimonie** sleeps a random amount of time, long enough for the previously used Tor circuit to expire. This process is meant to make it hard for an attacker to correlate the multiple key update.

As an example, here is how you would configure **parcimonie** to refresh the keys in your GnuPG keyring, as well as those keyrings created by Guix, such as when running **guix import**:

```
(service home-parcimonie-service-type
```

```
(home-parcimonie-configuration
 (refresh-guix-keyrings? #t)))
```

This assumes that the Tor anonymous routing daemon is already running on your system. On Guix System, this can be achieved by setting up `tor-service-type` (see Section 11.10.5 [Networking Services], page 322).

The service reference is given below.

`parcimonie-service-type` [Variable]

This is the service type for `parcimonie` (Parcimonie’s web site (<https://salsa.debian.org/intrigeri/parcimonie>)). Its value must be a `home-parcimonie-configuration`, as shown below.

`home-parcimonie-configuration` [Data Table]

Available `home-parcimonie-configuration` fields are:

`parcimonie` (default: `parcimonie`) (type: file-like)

The `parcimonie` package to use.

`verbose?` (default: `#f`) (type: boolean)

Whether to have more verbose logging from the service.

`gnupg-already-torified?` (default: `#f`) (type: boolean)

Whether GnuPG is already configured to pass all traffic through Tor (<https://torproject.org>).

`refresh-guix-keyrings?` (default: `#f`) (type: boolean)

Guix creates a few keyrings in the `$XDG_CONFIG_HOME`, such as when running `guix import` (see Section 9.5 [Invoking `guix import`], page 197). Setting this to `#t` will also refresh any keyrings which Guix has created.

`extra-content` (default: `#f`) (type: raw-configuration-string)

Raw content to add to the `parcimonie` command.

The OpenSSH package (<https://www.openssh.com>) includes a daemon, the `ssh-agent` command, that manages keys to connect to remote machines using the SSH (secure shell) protocol. With the `(gnu home services ssh)` service, you can configure the OpenSSH `ssh-agent` to run upon login. See Section 13.3.7 [GNU Privacy Guard], page 754, for an alternative to OpenSSH’s `ssh-agent`.

Here is an example of a service and its configuration that you could add to the `services` field of your `home-environment`:

```
(service home-ssh-agent-service-type
 (home-ssh-agent-configuration
 (extra-options '("-t" "1h30m"))))
```

`home-ssh-agent-service-type` [Variable]

This is the type of the `ssh-agent` home service, whose value is a `home-ssh-agent-configuration` object.

`home-ssh-agent-configuration` [Data Type]

Available `home-ssh-agent-configuration` fields are:

`openssh` (default: `openssh`) (type: file-like)

The OpenSSH package to use.

**socket-directory** (default: `XDG_RUNTIME_DIR/ssh-agent`) (type: gexp)

The directory to write the ssh-agent's `socket` file.

**extra-options** (default: `'()`)

Extra options will be passed to `ssh-agent`, please run `man ssh-agent` for more information.

### 13.3.7 GNU Privacy Guard

The (`gnu home services gnupg`) module provides services that help you set up the GNU Privacy Guard, also known as GnuPG or GPG, in your home environment.

The `gpg-agent` service configures and sets up GPG's agent, the program that is responsible for managing OpenPGP private keys and, optionally, OpenSSH (secure shell) private keys (see Section “Invoking GPG-AGENT” in *Using the GNU Privacy Guard*).

As an example, here is how you would configure `gpg-agent` with SSH support such that it uses the Emacs-based Pinentry interface when prompting for a passphrase:

```
(service home-gpg-agent-service-type
  (home-gpg-agent-configuration
    (pinentry-program
      (file-append pinentry-emacs "/bin/pinentry-emacs")))
    (ssh-support? #t)))
```

The service reference is given below.

**home-gpg-agent-service-type** [Variable]

This is the service type for `gpg-agent` (see Section “Invoking GPG-AGENT” in *Using the GNU Privacy Guard*). Its value must be a `home-gpg-agent-configuration`, as shown below.

**home-gpg-agent-configuration** [Data Type]

Available `home-gpg-agent-configuration` fields are:

**gnupg** (default: `gnupg`) (type: file-like)

The GnuPG package to use.

**pinentry-program** (type: file-like)

Pinentry program to use. Pinentry is a small user interface that `gpg-agent` delegates to anytime it needs user input for a passphrase or PIN (personal identification number) (see *Using the PIN-Entry*).

**ssh-support?** (default: `#f`) (type: boolean)

Whether to enable SSH (secure shell) support. When true, `gpg-agent` acts as a drop-in replacement for OpenSSH's `ssh-agent` program, taking care of OpenSSH secret keys and directing passphrase requests to the chosen Pinentry program.

**default-cache-ttl** (default: `600`) (type: integer)

Time a cache entry is valid, in seconds.

**max-cache-ttl** (default: `7200`) (type: integer)

Maximum time a cache entry is valid, in seconds. After this time a cache entry will be expired even if it has been accessed recently.

**default-cache-ttl-ssh** (default: 1800) (type: integer)  
 Time a cache entry for SSH keys is valid, in seconds.

**max-cache-ttl-ssh** (default: 7200) (type: integer)  
 Maximum time a cache entry for SSH keys is valid, in seconds.

**extra-content** (default: "") (type: raw-configuration-string)  
 Raw content to add to the end of `~/.gnupg/gpg-agent.conf`.

### 13.3.8 Desktop Home Services

The `(gnu home services desktop)` module provides services that you may find useful on “desktop” systems running a graphical user environment such as Xorg.

**home-x11-service-type** [Variable]

This is the service type representing the X Window graphical display server (also referred to as “X11”).

X Window is necessarily started by a system service; on Guix System, starting it is the responsibility of **gdm-service-type** and similar services (see Section 11.10.7 [X Window], page 355). At the level of Guix Home, as an unprivileged user, we cannot start X Window; all we can do is check whether it is running. This is what this service does.

As a user, you probably don’t need to worry or explicitly instantiate **home-x11-service-type**. Services that require an X Window graphical display, such as **home-redshift-service-type** below, instantiate it and depend on its corresponding **x11-display** Shepherd service (see Section 13.3.5 [Shepherd Home Service], page 747).

When X Window is running, the **x11-display** Shepherd service starts and sets the `DISPLAY` environment variable of the **shepherd** process, using its original value if it was already set; otherwise, it fails to start.

The service can also be forced to use a given value for `DISPLAY`, like so:

```
herd start x11-display :3
```

In the example above, **x11-display** is instructed to set `DISPLAY` to `:3`.

**home-redshift-service-type** [Variable]

This is the service type for Redshift (<https://github.com/jonls/redshift>), a program that adjusts the display color temperature according to the time of day. Its associated value must be a **home-redshift-configuration** record, as shown below. A typical configuration, where we manually specify the latitude and longitude, might look like this:

```
(service home-redshift-service-type
  (home-redshift-configuration
    (location-provider 'manual)
    (latitude 35.81)      ;northern hemisphere
    (longitude -0.80))) ;west of Greenwich
```

**home-redshift-configuration** [Data Type]

Available **home-redshift-configuration** fields are:

**redshift** (default: **redshift**) (type: file-like)  
 Redshift package to use.

**location-provider** (default: `geoclue2`) (type: symbol)  
 Geolocation provider—`'manual'` or `'geoclue2'`. In the former case, you must also specify the `latitude` and `longitude` fields so Redshift can determine daytime at your place. In the latter case, the Geoclue system service must be running; it will be queried for location information.

**adjustment-method** (default: `randr`) (type: symbol)  
 Color adjustment method.

**daytime-temperature** (default: `6500`) (type: integer)  
 Daytime color temperature (kelvins).

**nighttime-temperature** (default: `4500`) (type: integer)  
 Nighttime color temperature (kelvins).

**daytime-brightness** (type: maybe-inexact-number)  
 Daytime screen brightness, between 0.1 and 1.0, or left unspecified.

**nighttime-brightness** (type: maybe-inexact-number)  
 Nighttime screen brightness, between 0.1 and 1.0, or left unspecified.

**latitude** (type: maybe-inexact-number)  
 Latitude, when `location-provider` is `'manual'`.

**longitude** (type: maybe-inexact-number)  
 Longitude, when `location-provider` is `'manual'`.

**dawn-time** (type: maybe-string)  
 Custom time for the transition from night to day in the morning—`"HH:MM"` format. When specified, solar elevation is not used to determine the daytime/nighttime period.

**dusk-time** (type: maybe-string)  
 Likewise, custom time for the transition from day to night in the evening.

**extra-content** (default: `"`) (type: raw-configuration-string)  
 Extra content appended as-is to the Redshift configuration file. Run `man redshift` for more information about the configuration file format.

**home-dbus-service-type** [Variable]  
 This is the service type for running a session-specific D-Bus, for unprivileged applications that require D-Bus to be running.

**home-dbus-configuration** [Data Type]  
 The configuration record for `home-dbus-service-type`.

**dbus** (default: `dbus`)  
 The package providing the `/bin/dbus-daemon` command.

**home-unclutter-service-type** [Variable]  
 This is the service type for Unclutter, a program that runs on the background of an X11 session and detects when the X pointer hasn't moved for a specified idle timeout, after which it hides the cursor so that you can focus on the text underneath. Its associated value must be a `home-unclutter-configuration` record, as shown below.



A typical configuration, where we manually specify the idle timeout (in seconds), might look like this:

```
(service home-unclutter-service-type
  (home-unclutter-configuration
    (idle-timeout 2)))
```

**home-unclutter-configuration** [Data Type]

The configuration record for **home-unclutter-service-type**.

**unclutter** (default: **unclutter**) (type: file-like)  
Unclutter package to use.

**idle-timeout** (default: 5) (type: integer)  
A timeout in seconds after which to hide cursor.

**home-xmodmap-service-type** [Variable]

This is the service type for the xmodmap (<https://gitlab.freedesktop.org/xorg/app/xmodmap>) utility to modify keymaps and pointer button mappings under the Xorg display server. Its associated value must be a **home-xmodmap-configuration** record, as shown below.

The **key-map** field takes a list of objects, each of which is either a *statement* (a string) or an *assignment* (a pair of strings). As an example, the snippet below swaps around the *Caps\_Lock* and the *Control\_L* keys, by first removing the keysyms (on the right-hand side) from the corresponding modifier maps (on the left-hand side), re-assigning them by swapping each other out, and finally adding back the keysyms to the modifier maps.

```
(service home-xmodmap-service-type
  (home-xmodmap-configuration
    (key-map '(("remove Lock" . "Caps_Lock")
              ("remove Control" . "Control_L")
              ("keysym Control_L" . "Caps_Lock")
              ("keysym Caps_Lock" . "Control_L")
              ("add Lock" . "Caps_Lock")
              ("add Control" . "Control_L")))))
```

**home-xmodmap-configuration** [Data Type]

The configuration record for **home-xmodmap-service-type**. Its available fields are:

**xmodmap** (default: **xmodmap**) (type: file-like)  
The **xmodmap** package to use.

**key-map** (default: '()) (type: list)  
The list of expressions to be read by **xmodmap** on service startup.

**home-startx-command-service-type** [Variable]

Add **startx** to the home profile putting it onto **PATH**.

The value for this service is a **<xorg-configuration>** object which is passed to the **xorg-start-command-xinit** procedure producing the **startx** used. Default value is (**xorg-configuration**).

**home-darkman-service-type** [Variable]

This is the service type for the Darkman (<https://gitlab.com/WhyNotHugo/darkman>), a program for dark-mode and light-mode transitions on Unix-like desktops. It invokes scripts in `$XDG_DATA_DIRS/dark-mode.d` at sundown, and invokes scripts in `$XDG_DATA_DIRS/light-mode.d` at sunrise. Darkman tries to determine the user's current location from:

- The config file.
- The cache file from last time it ran.
- Using the system geoclue.

**home-darkman-configuration** [Data Type]

The configuration record for **home-darkman-service-type**. Its available fields are:

**darkman** (default: **darkman**) (type: file-like)

The **darkman** package to use.

**latitude** (type: maybe-number)

Latitude, used at startup or when **use-geoclue** is **#f**. More than one decimal point is generally not needed.

**longitude** (type: maybe-number)

Longitude, used at startup or when **use-geoclue** is **#f**. More than one decimal point is generally not needed.

**use-geoclue** (default: **#f**) (type: boolean)

Whether to use a local geoclue instance to determine the current location. Setting this to false without explicitly setting latitude and longitude disables automatic transitions entirely.

**dbus-server** (default: **#t**) (type: boolean)

Whether to expose the current mode via darkman's own D-Bus API.

**portal** (default: **#t**) (type: boolean)

Whether to expose the current mode via the XDG settings portal D-Bus API.

**home-gcr-ssh-agent-service-type** [Variable]

GNOME Keyring used to be able to act as a wrapper over **ssh-agent**. It would display a graphical password entry dialog, including a checkbox to remember the password, which, if selected, would allow fully passwordless use of that key.

The SSH functionality is disabled by default in **gnome-keyring-daemon** builds since version 1.46 (<https://gitlab.gnome.org/GNOME/gnome-keyring/-/commit/25c5a1982467802fa12c6852b03c57924553ba73>). It has been moved ([https://gitlab.gnome.org/GNOME/gcr/-/merge\\_requests/67](https://gitlab.gnome.org/GNOME/gcr/-/merge_requests/67)) into **gcr-ssh-agent**, which is part of the **gcr** package.

To enable the SSH agent functionality it is sufficient to add the following to your Home configuration.

```
(use-modules (gnu home services gnome) ;for 'home-gcr-ssh-agent-service-type'
...)
```

```
(home-environment
  (services
    (list
      ...
      (service home-gcr-ssh-agent-service-type))))
```

**gcr-ssh-agent-configuration** [Data Type]

The configuration record for `home-gcr-ssh-agent-service-type`. Its available fields are:

**package** (default: `gcr`) (type: `package`)  
The `gcr` package to use.

**log-file** (type: `maybe-string`)  
Where the service will write its logs. If unset, it defaults to `$HOME/.local/state/shepherd/gcr-ssh-agent.log`.

### 13.3.9 Guix Home Services

The `(gnu home services guix)` module provides services for user-specific Guix configuration.

**home-channels-service-type** [Variable]

This is the service type for managing `$XDG_CONFIG_HOME/guix/channels.scm`, the file that controls the channels received on `guix pull` (see Chapter 6 [Channels], page 68). Its associated value is a list of `channel` records, defined in the `(guix channels)` module.

Generally, it is better to extend this service than to directly configure it, as its default value is the default guix channel(s) defined by `%default-channels`. If you configure this service directly, be sure to include a guix channel. See Section 6.1 [Specifying Additional Channels], page 68, and Section 6.2 [Using a Custom Guix Channel], page 68, for more details.

A typical extension for adding a channel might look like this:

```
(simple-service 'variant-packages-service
  home-channels-service-type
  (list
    (channel
      (name 'variant-packages)
      (url "https://example.org/variant-packages.git")))))■
```

### 13.3.10 Backup Services

The `(gnu home services backup)` module offers services for backing up file system trees. For now, it provides the `home-restic-backup-service-type`.

With `home-restic-backup-service-type`, you can periodically back up directories and files with Restic (<https://restic.net/>), which supports end-to-end encryption and deduplication. Consider the following configuration:

```
(use-modules (gnu home services backup) ;for 'restic-backup-job', 'home-restic-backup-
```

```

        (gnu packages sync)          ;for 'rclone'
        ...)

(home-environment

  (packages (list rclone          ;for use by restic
                 ...))

  (services
    (list
      ...
      (simple-service 'backup-jobs
        home-restic-backup-service-type
        (list (restic-backup-job
          (name "remote-ftp")
          (repository "rclone:remote-ftp:backup/restic")
          (password-file "/home/alice/.restic")
          ;; Every day at 23.
          (schedule "0 23 * * *")
          (files '("/home/alice/.restic"
                  "/home/alice/.config/rclone"
                  "/home/alice/Pictures"))))))))

```

In general it is preferable to extend the `home-restic-backup-service-type`, as shown in the example above. This is because it takes care of wrapping everything with `for-home`, which enables the `home-restic-backup-service-type` and `restic-backup-service-type` to share the same codebase.

For a custom configuration, wrap your `restic-backup-configuration` in `for-home`, as in this example:

```

(use-modules (gnu services)          ;for 'for-home'
             (gnu services backup)   ;for 'restic-backup-job' and 'restic-backup-co
             (gnu home services backup) ;for 'home-restic-backup-service-type'
             (gnu packages sync)      ;for 'rclone'
             ...)

(home-environment

  (packages (list rclone          ;for use by restic
                 ...))

  (services
    (list
      ...
      (service home-restic-backup-service-type
        (for-home
          (restic-backup-configuration
            (jobs (list ...))))))

```

You can refer to see Section 11.10.39 [Miscellaneous Services], page 649, for details about `restic-backup-configuration` and `restic-backup-job`. The only difference is

that the `home-restic-backup-service-type` will ignore the `user` and `group` field of `restic-backup-job`.

### 13.3.11 Fonts Home Services

The `(gnu home services fontutils)` module provides services for user-specific Fontconfig setup. The Fontconfig (<https://www.freedesktop.org/wiki/Software/fontconfig>) library is used by many applications to access fonts on the system.

**home-fontconfig-service-type** [Variable]

This is the service type for generating configurations for Fontconfig. Its associated value is a list of either strings (or gexps) pointing to fonts locations, or SXML (see Section “SXML” in *GNU Guile Reference Manual*) fragments to be converted into XML and put inside the main `fontconfig` node.

Generally, it is better to extend this service than to directly configure it, as its default value is the default Guix Home’s profile font installation path (`~/.guix-home/profile/share/fonts`). If you configure this service directly, be sure to include the above directory.

Here’s how you’d extend it to include fonts installed with the Nix package manager, and to prefer your favourite monospace font:

```
(simple-service 'additional-fonts-service
               home-fontconfig-service-type
               (list "~/.nix-profile/share/fonts"
                    '(alias
                      (family "monospace")
                      (prefer
                        (family "Liberation Mono"))))))
```

### 13.3.12 Sound Home Services

The `(gnu home services sound)` module provides services related to sound support.

#### PulseAudio RTP Streaming Services

The following services dynamically reconfigure the PulseAudio sound server (<https://pulseaudio.org>): they let you toggle broadcast of audio output over the network using the RTP (real-time transport protocol) and, correspondingly, playback of sound received over RTP. Once `home-pulseaudio-rtp-sink-service-type` is among your home services, you can start broadcasting audio output by running this command:

```
herd start pulseaudio-rtp-sink
```

You can then run a PulseAudio-capable mixer, such as `pavucontrol` or `pulsemixer` (both from the same-named package) to control which audio stream(s) should be sent to the RTP “sink”.

By default, audio is broadcasted to a multicast address: any device on the LAN (local area network) receives it and may play it. Using multicast in this way puts a lot of pressure on the network and degrades its performance, so you may instead prefer sending to specifically one device. The first way to do that is by specifying the IP address of the target device when starting the service:

```
herd start pulseaudio-rtp-sink 192.168.1.42
```

The other option is to specify this IP address as the one to use by default in your home environment configuration:

```
(service home-pulseaudio-rtp-sink-service-type
  "192.168.1.42")
```

On the device where you intend to receive and play the RTP stream, you can use `home-pulseaudio-rtp-source-service-type`, like so:

```
(service home-pulseaudio-rtp-source-service-type)
```

This will then let you start the receiving module for PulseAudio:

```
herd start pulseaudio-rtp-source
```

Again, by default it will listen on the multicast address. If, instead, you'd like it to listen for direct incoming connections, you can do that by running:

```
(service home-pulseaudio-rtp-source-service-type
  "0.0.0.0")
```

The reference of these services is given below.

`home-pulseaudio-rtp-sink-service-type` [Variable]

`home-pulseaudio-rtp-source-service-type` [Variable]

This is the type of the service to send, respectively receive, audio streams over RTP (real-time transport protocol).

The value associated with this service is the IP address (a string) where to send, respectively receive, the audio stream. By default, audio is sent/received on multicast address `%pulseaudio-rtp-multicast-address`.

This service defines one Shepherd service: `pulseaudio-rtp-sink`, respectively `pulseaudio-rtp-source`. The service is not started by default, so you have to explicitly start it when you want to turn it on, as in this example:

```
herd start pulseaudio-rtp-sink
```

Stopping the Shepherd service turns off broadcasting.

`%pulseaudio-rtp-multicast-address` [Variable]

This is the multicast address used by default by the two services above.

## PipeWire Home Service

PipeWire (<https://pipewire.org>) provides a low-latency, graph-based audio and video processing service. In addition to its native protocol, it can also be used as a replacement for both JACK and PulseAudio.

While PipeWire provides the media processing and API, it does not, directly, know about devices such as sound cards, nor how you might want to connect applications, hardware, and media processing filters. Instead, PipeWire relies on a *session manager* to specify all these relationships. While you may use any session manager you wish, for most people the WirePlumber (<https://pipewire.pages.freedesktop.org/wireplumber/>) session manager, a reference implementation provided by the PipeWire project itself, suffices, and that is the one `home-pipewire-service-type` uses.

PipeWire can be used as a replacement for PulseAudio by setting `enable-pulseaudio?` to `#t` in `home-pipewire-configuration`, so that existing PulseAudio clients may use it without any further configuration.

In addition, JACK clients may connect to PipeWire by using the `pw-jack` program, which comes with PipeWire. Simply prefix the command with `pw-jack` when you run it, and audio data should go through PipeWire:

```
pw-jack mpv -ao=jack sound-file.wav
```

For more information on PulseAudio emulation, see <https://gitlab.freedesktop.org/pipewire/pipewire/-/wikis/Config-PulseAudio>, for JACK, see <https://gitlab.freedesktop.org/pipewire/pipewire/-/wikis/Config-JACK>.

As PipeWire does not use `dbus` to start its services on demand (as PulseAudio does), `home-pipewire-service-type` uses Shepherd to start services when logged in, provisioning the `pipewire`, `wireplumber`, and, if configured, `pipewire-pulseaudio` services. See Section 13.3.5 [Shepherd Home Service], page 747.

**home-pipewire-service-type** [Variable]

This provides the service definition for `pipewire`, which will run on login. Its value is a `home-pipewire-configuration` object.

To start the service, add it to the `service` field of your `home-environment`, such as:

```
(service home-pipewire-service-type)
```

**home-pipewire-configuration** [Data Type]

Available `home-pipewire-configuration` fields are:

**pipewire** (default: `pipewire`) (type: file-like)

The PipeWire package to use.

**wireplumber** (default: `wireplumber`) (type: file-like)

The WirePlumber package to use.

**enable-pulseaudio?** (default: `#t`) (type: boolean)

When true, enable PipeWire's PulseAudio emulation support, allowing PulseAudio clients to use PipeWire transparently.

**extra-content** (default: `"`) (type: string)

Extra content to add to the end of `~/.config/alsa/asoundrc`.

### 13.3.13 Mail Home Services

The (`gnu home services mail`) module provides services that help you set up the tools to work with emails in your home environment.

MSMTP (<https://marlam.de/msmtp>) is a SMTP (Simple Mail Transfer Protocol) client. It sends mail to a predefined SMTP server that takes care of proper delivery.

The service reference is given below.

**home-msmtp-service-type** [Variable]

This is the service type for `msmtp`. Its value must be a `home-msmtp-configuration`, as shown below. It provides the `~/.config/msmtp/config` file.

As an example, here is how you would configure `msmtp` for a single account:

```
(service home-msmtp-service-type
  (home-msmtp-configuration
    (accounts
```

```
(list
  (msmtp-account
    (name "alice")
    (configuration
      (msmtp-configuration
        (host "mail.example.org")
        (port 587)
        (user "alice")
        (password-eval "pass Mail/alice"))))))))
```

**home-msmtp-configuration** [Data Type]

Available **home-msmtp-configuration** fields are:

**defaults** (type: **msmtp-configuration**)

The configuration that will be set as default for all accounts.

**accounts** (default: '()') (type: **list-of-msmtp-accounts**)

A list of **msmtp-account** records which contain information about all your accounts.

**default-account** (type: **maybe-string**)

Set the default account.

**extra-content** (default: "") (type: **string**)

Extra content appended as-is to the configuration file. Run **man msmtp** for more information about the configuration file format.

**msmtp-account** [Data Type]

Available **msmtp-account** fields are:

**name** (type: **string**)

The unique name of the account.

**configuration** (type: **msmtp-configuration**)

The configuration for this given account.

**msmtp-configuration** [Data Type]

Available **msmtp-configuration** fields are:

**auth?** (type: **maybe-boolean**)

Enable or disable authentication.

**tls?** (type: **maybe-boolean**)

Enable or disable TLS (also known as SSL) for secured connections.

**tls-starttls?** (type: **maybe-boolean**)

Choose the TLS variant: start TLS from within the session ('on', default), or tunnel the session through TLS ('off').

**tls-trust-file** (type: **maybe-string**)

Activate server certificate verification using a list of trusted Certification Authorities (CAs).



**log-file** (type: maybe-string)  
 Enable logging to the specified file. An empty argument disables logging.  
 The file name '-' directs the log information to standard output.

**host** (type: maybe-string)  
 The SMTP server to send the mail to.

**port** (type: maybe-integer)  
 The port that the SMTP server listens on. The default is 25 ("smtp"),  
 unless TLS without STARTTLS is used, in which case it is 465 ("smtps").

**user** (type: maybe-string)  
 Set the user name for authentication.

**from** (type: maybe-string)  
 Set the envelope-from address.

**password-eval** (type: maybe-string)  
 Set the password for authentication to the output (stdout) of the com-  
 mand cmd.

**extra-content** (default: "") (type: string)  
 Extra content appended as-is to the configuration block. Run `man msmtplib`  
 for more information about the configuration file format.

### 13.3.14 Messaging Home Services

The ZNC bouncer (<https://znc.in>) can be run as a daemon to manage your IRC presence. With the (`gnu home services messaging`) service, you can configure ZNC to run upon login.

You will have to provide a `~/.znc/configs/znc.conf` separately.

Here is an example of a service and its configuration that you could add to the `services` field of your `home-environment`:

```
(service home-znc-service-type
```

[Variable]

This is the type of the ZNC home service, whose value is a `home-znc-configuration` object.

```
home-znc-configuration
```

[Data Type]

Available `home-znc-configuration` fields are:

**znc** (default: `znc`) (type: file-like)  
 The ZNC package to use.

**extra-options** (default: '()')  
 Extra options will be passed to `znc`, please run `man znc` for more information.

The Snuik IRC bot (<https://gitlab.com/janneke/snuik>) can be run as a daemon to aid talking to users that are currently offline. With the (`gnu home services messaging`) service, you can configure Snuik to run upon login.

**home-snuik-service-type** [Variable]

This is the type of the Snuik home service, whose value is a **home-snuik-configuration** object.

You may specify a custom configuration by providing a **snuik-configuration** record, exactly like for **snuik-service-type** and wrapping it in **for-home**. Here is an example of a service and its configuration that you could add to the **services** field of your **home-environment**:

```
(service home-snuik-service-type
  (for-home
    (snuik-configuration
      (server "irc.oftc.net")
      (nick "sneek")
      (password "snuik is sneaky")
      (channels '("#reproducible-builds")))))
```

### 13.3.15 Media Home Services

The Kodi media center (<https://kodi.tv>) can be run as a daemon on a media server. With the **(gnu home services kodi)** service, you can configure Kodi to run upon login.

Here is an example of a service and its configuration that you could add to the **services** field of your **home-environment**:

```
(service home-kodi-service-type
  (home-kodi-configuration
    (extra-options '("--settings=<settings-file>"))))
```

**home-kodi-service-type** [Variable]

This is the type of the Kodi home service, whose value is a **home-kodi-configuration** object.

**home-kodi-configuration** [Data Type]

Available **home-kodi-configuration** fields are:

**kodi** (default: **kodi**) (type: file-like)  
The Kodi package to use.

**extra-options** (default: **'()**)  
Extra options will be passed to **kodi**, please run **man kodi** for more information.

### 13.3.16 mpv Media Player

Configuring the mpv media player (<https://mpv.io/>) can be somewhat daunting task, due to the sheer amount of options available, especially if one wants to be able to inherit the configuration in their code. The **home-mpv-service-type** is provided to help with that problem. When the service is added to your home environment, file based on the passed configuration is generated and placed into the correct location.

Due to the bug [#74748](#), it does not use Guix Records to represent the configuration, but uses keyword arguments to achieve similar result. Example follows:

```
(service home-mpv-service-type
```

```
(make-home-mpv-configuration
 #:global (make-mpv-profile-configuration
           #:fullscreen? #t
           #:alang '("jpn" "eng"))))
```

**home-mpv-service-type** [Variable]

This is the type of the mpv home service, whose value is a **home-mpv-configuration** object.

**make-home-mpv-configuration** [Procedure]

Return a new instance of **home-mpv-configuration**. Available keyword arguments are:

**inherit** (default: **#t**)

Inherit fields from another instance.

**global** (default: **(make-mpv-profile-configuration)**)

The global configuration preceding all profiles.

**profiles** (default: **()**)

An alist containing configuration for any additional profiles to include in the configuration file.

```
(make-home-mpv-configuration
 #:profiles `((fullscreen . ,(make-mpv-profile-configuration
                               #:fullscreen? #t))))
```

**extra-config** (default: **#f**)

Additional content to include in the configuration file. It is placed at the end of the file.

**make-mpv-profile-configuration** [Procedure]

Return a new instance of **mpv-profile-configuration**. As above, it also supports the **#:inherit** argument. Additionally it supports keyword arguments named after the options of mpv. Hence **--fullscreen** (or **fullscreen** in the configuration file) becomes **#:fullscreen?**.

Few options are using their aliases instead. **audio** instead of **aid**, **video** instead of **vid**, **sub** instead of **sid**, **screenshot-directory** instead of **screenshot-dir** and **watch-later-directory** instead of **watch-later-dir**.

Valid values for the fields depend on their type.

**Flags** The value should be **#f** or **#t**.

**Numerical fields**

Integer and integer64 fields are validated by **integer?**, float and double fields by **real?**. Ranges are checked when applicable.

**Aspect** Same as integer.

**ByteSize** Same as integer64.

**Choice** The value should be a symbol representing member of the enumeration. If the choice has ‘**or ...**’ part, it can also be value of that alternative type.

*type list*    The value should be a list of the *type*.

Other types accept strings, with validation of the values where possible (e.g. type ‘Color’ is validated, but type ‘Audio channels or channel map’ is not). If an option is marked as taking a file, values satisfying *file-like?* can also be used.

The full list of currently supported keyword arguments is below. For the types, allowed values and full description please refer to the *mpv --list-options* and *mpv manual* (<https://mpv.io/manual/stable/#options>).

Only set fields are output to the configuration file. Accessors are provided for every field, returning either their value or a sentinel object *%unset-value* (from (gnu services configuration)).

ab-loop-a, ab-loop-b, ab-loop-count, access-references?, ad, ad-lavc-ac3drc, ad-lavc-downmix?, ad-lavc-o, ad-lavc-threads, ad-queue-enable?, ad-queue-max-bytes, ad-queue-max-samples, ad-queue-max-secs, af, audio, alang, allow-delayed-peak-detect?, alsa-buffer-time, alsa-ignore-chmap?, alsa-mixer-device, alsa-mixer-index, alsa-mixer-name, alsa-non-interleaved?, alsa-periods, alsa-resample?, ao, ao-null-broken-delay?, ao-null-broken-eof?, ao-null-buffer, ao-null-channel-layouts, ao-null-format, ao-null-latency, ao-null-outburst, ao-null-speed, ao-null-untimed?, ao-pcm-append?, ao-pcm-file, ao-pcm-waveheader?, archive-exts, audio-backward-batch, audio-backward-overlap, audio-buffer, audio-channels, audio-client-name, audio-delay, audio-demuxer, audio-device, audio-display, audio-exclusive?, audio-exts, audio-fallback-to-null?, audio-file-auto, audio-file-paths, audio-files, audio-format, audio-normalize-downmix?, audio-pitch-correction?, audio-resample-cutoff, audio-resample-filter-size, audio-resample-linear?, audio-resample-max-output-size, audio-resample-phase-shift, audio-reversal-buffer, audio-samplerate, audio-spdif, audio-stream-silence?, audio-swresample-o, audio-wait-open, auto-window-resize?, autocreate-playlist, autofit, autofit-larger, autofit-smaller, autoload-files?, autosync, background, background-color, blend-subtitles, bluray-device, border?, border-background, brightness, cache, cache-on-disk?, cache-pause?, cache-pause-initial?, cache-pause-wait, cache-secs, cdda-cdtext?, cdda-device, cdda-overlap, cdda-paranoia, cdda-sector-size, cdda-skip?, cdda-span-a, cdda-span-b, cdda-speed, cdda-toc-offset, chapter-merge-threshold, chapter-seek-threshold, chapters-file, clipboard-backends, clipboard-monitor?, config?, container-fps-override, contrast, cookies?, cookies-file, corner-rounding, correct-downscaling?, correct-pts?, cover-art-auto, cover-art-files, cover-art-whitelist, cscale, cscale-antiring, cscale-blur, cscale-clamp, cscale-param1, cscale-param2, cscale-radius, cscale-taper, cscale-window, cscale-wparam, cscale-wtaper, cursor-autohide, cursor-autohide-fs-only?, deband?, deband-grain, deband-iterations, deband-range, deband-threshold, deinterlace, deinterlace-field-parity, demuxer, demuxer-backward-playback-step, demuxer-cache-dir, demuxer-cache-unlink-files, demuxer-cache-wait?, demuxer-donate-buffer?, demuxer-hysteresis-secs, demuxer-lavf-allow-mimetype?, demuxer-lavf-analyzeduration, demuxer-lavf-buffersize, demuxer-lavf-format,

demuxer-lavf-hacks?, demuxer-lavf-linearize-timestamps, demuxer-lavf-o,  
 demuxer-lavf-probe-info, demuxer-lavf-probescore, demuxer-lavf-  
 probesize, demuxer-lavf-propagate-opts?, demuxer-max-back-bytes,  
 demuxer-max-bytes, demuxer-mkv-crop-compat?, demuxer-mkv-probe-start-  
 time?, demuxer-mkv-probe-video-duration, demuxer-mkv-subtitle-preroll,  
 demuxer-mkv-subtitle-preroll-secs, demuxer-mkv-subtitle-preroll-  
 secs-index, demuxer-rawaudio-channels, demuxer-rawaudio-format,  
 demuxer-rawaudio-rate, demuxer-rawvideo-codec, demuxer-rawvideo-format,  
 demuxer-rawvideo-fps, demuxer-rawvideo-h, demuxer-rawvideo-mp-format,  
 demuxer-rawvideo-size, demuxer-rawvideo-w, demuxer-readahead-secs,  
 demuxer-seekable-cache, demuxer-termination-timeout, demuxer-thread?,  
 directory-filter-types, directory-mode, display-fps-override, display-  
 tags, dither, dither-depth, dither-size-fruit, drag-and-drop, dscale,  
 dscale-antiring, dscale-blur, dscale-clamp, dscale-param1, dscale-param2,  
 dscale-radius, dscale-taper, dscale-window, dscale-wparam, dscale-wtaper,  
 dump-stats, dvbin-card, dvbin-channel-switch-offset, dvbin-file,  
 dvbin-full-transponder?, dvbin-prog, dvbin-timeout, dvd-angle, dvd-device,  
 dvd-speed, edition, egl-config-id, egl-output-format, embeddedfonts?, end,  
 error-diffusion, external-files, fbo-format, focus-on, force-media-title,  
 force-render?, force-rgba-osd-rendering?, force-seekable?, force-window,  
 force-window-position?, framedrop, frames, fs-screen, fs-screen-name,  
 fullscreen?, gamma, gamma-auto?, gamma-factor, gamut-mapping-mode, gapless-  
 audio, geometry, glsl-shader-opts, glsl-shaders, gpu-api, gpu-context,  
 gpu-debug?, gpu-dumb-mode, gpu-hwdec-interop, gpu-shader-cache?, gpu-  
 shader-cache-dir, gpu-sw?, gpu-tex-pad-x, gpu-tex-pad-y, hdr-compute-peak,  
 hdr-contrast-recovery, hdr-contrast-smoothness, hdr-peak-decay-rate,  
 hdr-peak-percentile, hdr-scene-threshold-high, hdr-scene-threshold-low,  
 hidpi-window-scale?, hls-bitrate, hr-seeking, hr-seeking-demuxer-offset,  
 hr-seeking-framedrop?, http-header-fields, http-proxy, hue, hwdec, hwdec-  
 codecs, hwdec-extra-frames, hwdec-image-format, hwdec-software-fallback,  
 icc-3dlut-size, icc-cache?, icc-cache-dir, icc-force-contrast, icc-intent,  
 icc-profile, icc-profile-auto?, icc-use-luma?, idle, ignore-path-in-  
 watch-later-config?, image-display-duration, image-exts, image-lut,  
 image-lut-type, image-sub-video-resolution?, include, index, initial-  
 audio-sync?, input-ar-delay, input-ar-rate, input-builtin-bindings?,  
 input-builtin-dragging?, input-commands, input-conf, input-cursor?, input-  
 cursor-passthrough?, input-default-bindings?, input-doubleclick-time,  
 input-dragging-deadzone, input-ime?, input-ipc-client, input-ipc-server,  
 input-key-fifo-size, input-media-keys?, input-preprocess-wheel?,  
 input-right-alt-gr?, input-terminal?, input-test?, input-touch-emulate-  
 mouse?, input-vo-keyboard?, interpolation?, interpolation-preserve?,  
 interpolation-threshold, inverse-tone-mapping?, jack-autostart?,  
 jack-connect?, jack-name, jack-port, jack-std-channel-layout, keep-open,  
 keep-open-pause?, keepaspect?, keepaspect-window?, lavfi-complex, length,  
 libplacebo-opts, linear-downscaling?, linear-upscaling?, load-auto-  
 profiles, load-commands?, load-console?, load-positioning?, load-scripts?,  
 load-select?, load-stats-overlay?, load-unsafe-playlists?, log-file,

loop-file, loop-playlist, lut, lut-type, mc, media-controls?, merge-files?,  
 metadata-codepage, mf-fps, mf-type, monitoraspect, monitorpixelaspect,  
 msg-color?, msg-level, msg-module?, msg-time?, mute?, native-fs?,  
 native-keyrepeat?, native-touch?, network-timeout, oac, oacopts,  
 ocopy-metadata?, of, ofopts, on-all-workspaces?, ontop?, ontop-level,  
 opengl-check-pattern-a, opengl-check-pattern-b, opengl-early-flush,  
 opengl-es, opengl-glfinish?, opengl-pbo?, opengl-rectangle-textures?,  
 opengl-swapinterval, opengl-waitvsync?, orawts?, ordered-chapters?,  
 ordered-chapters-files, oremove-metadata, osc?, osd-align-x, osd-align-y,  
 osd-back-color, osd-bar?, osd-bar-align-x, osd-bar-align-y, osd-bar-h,  
 osd-bar-marker-min-size, osd-bar-marker-scale, osd-bar-marker-style,  
 osd-bar-outline-size, osd-bar-w, osd-blur, osd-bold?, osd-border-style,  
 osd-color, osd-duration, osd-font, osd-font-provider, osd-font-size,  
 osd-fonts-dir, osd-fractions?, osd-italic?, osd-justify, osd-level,  
 osd-margin-x, osd-margin-y, osd-msg1, osd-msg2, osd-msg3, osd-on-seek,  
 osd-outline-color, osd-outline-size, osd-playing-msg, osd-playing-  
 msg-duration, osd-playlist-entry, osd-scale, osd-scale-by-window?,  
 osd-selected-color, osd-selected-outline-color, osd-shadow-offset,  
 osd-spacing, osd-status-msg, oset-metadata, ovc, ovcopts, panscan, pause?,  
 pipewire-buffer, pipewire-remote, pipewire-volume-mode, pitch, play-  
 direction, player-operation-mode, playlist-exts, playlist-start, prefetch-  
 playlist?, profile, pulse-allow-suspended?, pulse-buffer, pulse-host,  
 pulse-latency-hacks?, quiet?, really-quiet?, rebase-start-time?, referrer,  
 replaygain, replaygain-clip?, replaygain-fallback, replaygain-preamp,  
 reset-on-next-file, resume-playback?, resume-playback-check-mtime?,  
 rtsp-transport, saturation, save-position-on-quit?, save-watch-history?,  
 scale, scale-antiring, scale-blur, scale-clamp, scale-param1, scale-param2,  
 scale-radius, scale-taper, scale-window, scale-wparam, scale-wtaper,  
 scaler-resizes-only?, screen, screen-name, screenshot-avif-encoder,  
 screenshot-avif-opts, screenshot-avif-pixfmt, screenshot-directory,  
 screenshot-format, screenshot-high-bit-depth?, screenshot-jpeg-  
 quality, screenshot-jpeg-source-chroma?, screenshot-jxl-distance,  
 screenshot-jxl-effort, screenshot-png-compression, screenshot-png-  
 filter, screenshot-sw?, screenshot-tag-colorspace?, screenshot-template,  
 screenshot-webp-compression, screenshot-webp-lossless?, screenshot-webp-  
 quality, script-opts, scripts, secondary-sid, secondary-sub-ass-override,  
 secondary-sub-delay, secondary-sub-pos, secondary-sub-visibility?,  
 sharpen, show-in-taskbar?, shuffle?, sub, sigmoid-center, sigmoid-slope,  
 sigmoid-upscaling?, slang, snap-window?, speed, spirv-compiler, sstep,  
 start, stop-playback-on-init-failure?, stop-screensaver, stream-buffer-  
 size, stream-dump, stream-lavf-o, stream-record, stretch-dvd-subst?,  
 stretch-image-subst-to-screen?, sub-align-x, sub-align-y, sub-ass?,  
 sub-ass-force-margins?, sub-ass-justify?, sub-ass-override, sub-ass-  
 prune-delay, sub-ass-scale-with-window?, sub-ass-style-overrides,  
 sub-ass-styles, sub-ass-use-video-data, sub-ass-video-aspect-override,  
 sub-ass-vsfilter-color-compat, sub-auto, sub-auto-exts, sub-back-color,  
 sub-blur, sub-bold?, sub-border-style, sub-clear-on-seek?, sub-codepage,

sub-color, sub-create-cc-track?, sub-delay, sub-demuxer, sub-file-paths, sub-files, sub-filter-jsre, sub-filter-regex, sub-filter-regex-enable?, sub-filter-regex-plain?, sub-filter-regex-warn?, sub-filter-sdh?, sub-filter-sdh-enclosures, sub-filter-sdh-harder?, sub-fix-timing?, sub-font, sub-font-provider, sub-font-size, sub-fonts-dir, sub-forced-events-only?, sub-fps, sub-gauss, sub-gray?, sub-hinting, sub-italic?, sub-justify, sub-lavc-o, sub-line-spacing, sub-margin-x, sub-margin-y, sub-outline-color, sub-outline-size, sub-past-video-end?, sub-pos, sub-scale, sub-scale-by-window?, sub-scale-signs?, sub-scale-with-window?, sub-shadow-offset, sub-shaper, sub-spacing, sub-speed, sub-stretch-durations?, sub-use-margins?, sub-visibility?, sub-vsfilter-bidi-compat?, subs-fallback, subs-fallback-forced, subs-match-os-language?, subs-with-matching-audio, swapchain-depth, sws-allow-zimg?, sws-bitexact?, sws-cgb, sws-chs, sws-cs, sws-cvs, sws-fast?, sws-lgb, sws-ls, sws-scaler, target-colorspace-hint, target-contrast, target-gamut, target-lut, target-peak, target-prim, target-trc, taskbar-progress?, teletext-page, temporal-dither?, temporal-dither-period, term-osd, term-osd-bar?, term-osd-bar-chars, term-playing-msg, term-status-msg, term-title, terminal?, title, title-bar?, tls-ca-file, tls-cert-file, tls-key-file, tls-verify?, tone-mapping, tone-mapping-max-boost, tone-mapping-param, tone-mapping-visualize?, track-auto-selection?, tscale, tscale-antiring, tscale-blur, tscale-clamp, tscale-param1, tscale-param2, tscale-radius, tscale-taper, tscale-window, tscale-wparam, tscale-wtaper, untimed?, use-embedded-icc-profile?, use-filedir-conf?, user-agent, vaapi-device, vd, vd-apply-cropping?, vd-lavc-assume-old-x264?, vd-lavc-bitexact?, vd-lavc-check-hw-profile?, vd-lavc-dr, vd-lavc-fast?, vd-lavc-film-grain, vd-lavc-framedrop, vd-lavc-o, vd-lavc-show-all?, vd-lavc-skipframe, vd-lavc-skipidct, vd-lavc-skiploopfilter, vd-lavc-threads, vd-queue-enable?, vd-queue-max-bytes, vd-queue-max-samples, vd-queue-max-secs, vf, video, video-align-x, video-align-y, video-aspect-method, video-aspect-override, video-backward-batch, video-backward-overlap, video-crop, video-exts, video-latency-hacks?, video-margin-ratio-bottom, video-margin-ratio-left, video-margin-ratio-right, video-margin-ratio-top, video-osd?, video-output-levels, video-pan-x, video-pan-y, video-recenter?, video-reversal-buffer, video-rotate, video-scale-x, video-scale-y, video-sync, video-sync-max-audio-change, video-sync-max-factor, video-sync-max-video-change, video-timing-offset, video-unscaled, video-zoom, vlang, vo, vo-image-avif-encoder, vo-image-avif-opts, vo-image-avif-pixfmt, vo-image-format, vo-image-high-bit-depth?, vo-image-jpeg-quality, vo-image-jpeg-source-chroma?, vo-image-jxl-distance, vo-image-jxl-effort, vo-image-outdir, vo-image-png-compression, vo-image-png-filter, vo-image-tag-colorspace?, vo-image-webp-compression, vo-image-webp-lossless?, vo-image-webp-quality, vo-kitty-alt-screen?, vo-kitty-cols, vo-kitty-config-clear?, vo-kitty-height, vo-kitty-left, vo-kitty-rows, vo-kitty-top, vo-kitty-use-shm?, vo-kitty-width, vo-null-fps, vo-sixel-alt-screen?, vo-sixel-buffered?, vo-sixel-cols, vo-sixel-config-clear?, vo-sixel-dither, vo-sixel-fixedpalette?, vo-sixel-height, vo-sixel-left,

```
vo-sixel-pad-x, vo-sixel-pad-y, vo-sixel-reqcolors, vo-sixel-rows, vo-
sixel-threshold, vo-sixel-top, vo-sixel-width, vo-tct-256?, vo-tct-algo,
vo-tct-buffering, vo-tct-height, vo-tct-width, vo-vaapi-scaled-osd?,
vo-vaapi-scaling, vo-udpau-chroma-deint?, vo-udpau-colorkey, vo-udpau-
composite-detect?, vo-udpau-denoise, vo-udpau-force-yuv?, vo-udpau-fps,
vo-udpau-hqscaling, vo-udpau-output-surfaces, vo-udpau-pullup?, vo-
udpau-queuetime-fs, vo-udpau-queuetime-windowed, vo-udpau-sharpen,
volume, volume-gain, volume-gain-max, volume-gain-min, volume-max,
vulkan-async-compute?, vulkan-async-transfer?, vulkan-device, vulkan-
display-display, vulkan-display-mode, vulkan-display-plane, vulkan-
queue-count, vulkan-swap-mode, watch-history-path, watch-later-directory,
watch-later-options, wayland-app-id, wayland-configure-bounds, wayland-
content-type, wayland-disable-vsycn?, wayland-edge-pixels-pointer,
wayland-edge-pixels-touch, wayland-internal-vsycn, wayland-present?, wid,
window-dragging?, window-maximized?, window-minimized?, window-scale,
write-filename-in-watch-later-config?, x11-bypass-compositor, x11-name,
x11-netwm, x11-present, x11-wid-title?, xv-adaptor, xv-buffers, xv-ck,
xv-ck-method, xv-colorkey, xv-port, ytdl?, ytdl-format, ytdl-raw-options,
zing-dither, zimg-fast?, zimg-scaler, zimg-scaler-chroma, zimg-scaler-
chroma-param-a, zimg-scaler-chroma-param-b, zimg-scaler-param-a,
zing-scaler-param-b, and zimg-threads.
```

### 13.3.17 Niri window manager

The (gnu home services niri) module provides the `home-niri-service-type` service, which installs and configures Niri (<https://github.com/YaLTeR/niri>), a scrollable tiling Wayland compositor. Niri combines traditional tiling window management with smooth scrolling between workspaces.

This service starts Niri as a user-level desktop session with proper environment variables configured for Wayland compatibility, including `XDG_CURRENT_DESKTOP`, `XDG_SESSION_DESKTOP`, and `DESKTOP_SESSION`.

Here is a minimal example service declaration for your `home-environment` configuration:

```
(service home-niri-service-type)
```

For more advanced configurations, you can customize the service with a `niri-configuration` record to specify additional environment variables.

### 13.3.18 Sway window manager

The (gnu home services sway) module provides `home-sway-service-type`, a home service to configure the Sway window manager for Wayland (<https://github.com/swaywm/sway>) in a declarative way.

Here is an example of a service and its configuration that you could add to the `services` field of your `home-environment`:

```
(service home-sway-service-type
  (sway-configuration
    (gestures
      '((swipe:3:down . "move to scratchpad"))
```



```

        (swipe:3:up . "scratchpad show"))))
(outputs
  (list (sway-output
        (identifier '*))
        (background (file-append sway
                                "\
/share/backgrounds/sway/Sway_Wallpaper_Blue_1920x1080.png"))))))))

```

The above example describes a Sway configuration in which

- all monitors use a particular wallpaper whose `.png` is provided by the `sway` package;
- swiping down (resp. up) with three fingers moves the active window to the scratchpad (resp. shows/hides the scratchpad).

**Note:** This home service only sets up the configuration file and profile packages for Sway. It does *not* start Sway in any way. If you want to do so, you might be interested in using `greeterd-wlgreet-sway-session` instead.

The procedure `sway-configuration->file` defined below can be used to provide the value for the *optional* `sway-configuration` field of `greeterd-wlgreet-sway-session`.

`sway-configuration->file` *config* [Procedure]

This procedure takes one argument `config`, which must be a `sway-configuration` record (defined below), and returns a file-like object representing the serialized configuration.

`home-sway-service-type` [Variable]

This is a home service type to set up Sway. It takes care of:

- providing a `~/.config/sway/config` file,
- adding Sway-related packages to your profile.

It is possible to extend this service (see Section 11.19.1 [Service Composition], page 709). The resulting configuration file contains the concatenation of keybindings, modes, *etc.* found in all extensions.

**Note:** At the moment, only one bar configuration is allowed. If more are found, an error will be raised.

`%empty-sway-configuration` [Variable]

The variable `%empty-sway-configuration` contains an empty `sway-configuration` record. This value can help users write simple service extensions.

`sway-configuration` [Data Type]

This configuration record describes the Sway configuration (see `sway(5)`). Available fields are:

`variables` (default: `%sway-default-variables`)

The value of this field is an association list in which keys are symbols and values are either strings, G-expressions or file-like objects (see Section 8.12 [G-Expressions], page 165).

Example:

```
(variables `((mod . "Mod4") ; string
            (term ; file-append
              . ,(file-append foot "/bin/foot"))
            (Term ; G-expression
              . ,#~(string-append #foot "/bin/foot"))))
```

**Note:** Default keybindings assume the existence of variables named `$mod`, `$left`, `$right`, `$up` and `$down`. If you choose not to define these variables, make sure to remove keybindings referring to them.

**keybindings** (default: `%sway-default-keybindings`)

This field describes keybindings for the *default* mode. The value is an association list in which keys are symbols. Values can either be:

- strings or G-expressions,
- a cons-cell of a string or G-expression and a list of options. Options must be a string starting with “input-device=” or strings among “no-warn”, “whole-window”, “border”, “exclude-titlebar”, “release”, “locked”, “to-code”, “inhibited” and “no-repeat”.

The following snippet launches the terminal when pressing `$mod+t` and `$mod+Shift+t` (assuming that a variable `$term` is defined):

```
`(($mod+t . ,#~(string-append "exec " #foot "/bin/foot")))■
  ($mod+Shift+t . "exec $term")
  ($mod+q "exec $term" . ("to-code"))) ;; passes the --to-code option.■
```

**gestures** (default: `%sway-default-gestures`)

Similar to the previous field, but for finger-gestures. Options must start with “input-device=” or be among “exact” and “no-warn”.

The following snippet allows to navigate through workspaces by swiping right and left with three fingers:

```
'((swipe:3:right . "workspace next_on_output")
   (swipe:3:left . "workspace prev_on_output" . ("exact"))))■
```

**packages** (default: `%sway-default-packages`)

This field describes a list of packages to add to the user profile. At the moment, the default value only adds `sway` to the profile.

**inputs** (default: `'()`)

List of `sway-input` configuration records (described below).

**outputs** (default: `'()`)

List of `sway-output` configuration records (described below).

**bar** (optional `sway-bar` record)

Optional `sway-bar` record (described below) to configure a Sway bar.

**modes** (default: `%sway-default-modes`)

List of `sway-mode` records (described below) to add modes to the Sway configuration. The default value `%sway-default-modes` adds the “resize” mode of the default Sway configuration (as described below).

**startup+reload-programs** (default: '()')

Programs to execute at startup time *and* after every configuration reload. The value of this field is a list of strings, G-expressions or file-like objects (see Section 8.12 [G-Expressions], page 165).

**startup-programs** (default: %sway-default-startup-programs)

Programs to execute at startup time. As above, values of this field are a list of strings, G-expressions or file-like objects.

The default value, %sway-default-startup-programs, executes **swayidle** in order to lock the screen after 5 minutes of inactivity (displaying a background distributed with Sway) and turn the screen off after 10 minutes of inactivity.

**extra-content** (default: '()')

Lines to add to the configuration file. The value of this field is a list of strings or G-expressions.

**sway-input** [Data Type]

**sway-input** records describe input blocks (see *sway-input(5)*). For example, the following snippet makes all keyboards use a French layout, in which **capslock** has been remapped to **ctrl**:

```
(sway-input (identifier "type:keyboard")
  (layout
    (keyboard-layout "fr" #:options '("ctrl:nocaps"))))
```

Available fields for **sway-input** configuration records are:

**identifier** (default: '\*')

Identifier of the input. The field accepts symbols and strings. If the **identifier** is a symbol, it is inserted as is; if it is a string, it will be quoted in the configuration file.

**layout** (optional <keyboard-layout> record)

Keyboard specific option. Field specifying the layout to use for the input. The value must be a <keyboard-layout> record (see Section 11.8 [Keyboard Layout], page 271).

**Note:** (gnu home services sway) does not re-export the **keyboard-layout** procedure.

**disable-while-typing** (optional boolean)

If **#t** (resp. **#f**) enables (resp. disables) the “disable while typing” option for this input.

**disable-while-trackpointing** (optional boolean)

If **#t** (resp. **#f**), enables (resp. disables) the “disable while track-pointing” option for this input.

**tap** (optional boolean)

Enables or disables the “tap” option, which allows clicking by tapping on a touchpad.

**extra-content** (default: '()')

Lines to add to the input block. The value of this field must be a list whose elements are either strings or G-expressions.

**sway-output** [Data Type]

**sway-output** records describe Sway outputs (see *sway-output(5)*). Available fields are:

**identifier** (default: '\*')

Identifier of the monitor. The field accepts symbols and strings. If the **identifier** is a symbol, it is inserted as is; if it is a string, it will be quoted in the configuration file.

**resolution** (optional string)

This string defines the resolution of the monitor.

**position** (optional)

The (optional) value of this field must be a **point**. Example:

```
(position
 (point (x 1920)
        (y 0)))
```

**background** (optional)

The value of this field describes what wallpaper to use on this output. The field accepts the following types of values:

- a string,
- a G-expression,
- a file-like object,
- a pair. The first argument of this pair must be a string, a G-expression or a file-like object. The second element describes how the wallpaper will be displayed. It must be a symbol among **stretch**, **fill**, **fit**, **center** and **tile**.

If the second element is not specified (*i.e.* when the value is not a pair), the **fill** mode will be used.

**Note:** In order to use an SVG file, you must have **librsvg** in your profile (*e.g.* by adding it in the **packages** field of **sway-configuration**).

**extra-content** (default: '()')

List defining additional lines to add to the output configuration block. Elements of the list must be either strings or G-expressions.

**sway-border-color** [Data Type]

**border** Color of the border.

**background** Color of the background.

**text** Color of the text.

**sway-color** [Data Type]

- background** (optional string)  
Background color of the bar.
- statusline** (optional string)  
Text color of the status line.
- focused-background** (optional string)  
Background color of the bar on the currently focused monitor.
- focused-statusline** (optional string)  
Text color of the statusline on the currently focused monitor.
- focused-workspace** (optional **sway-border-color**)  
Color scheme for focused workspaces.
- active-workspace** (optional **sway-border-color**)  
Color scheme for active workspaces.
- inactive-workspace** (optional **sway-border-color**)  
Color scheme for inactive workspaces.
- urgent-workspace** (optional **sway-border-color**)  
Color scheme for workspaces containing “urgent windows”.
- binding-mode** (optional **sway-border-color**)  
Color scheme for the binding mode indicator.

**sway-bar** [Data Type]

- Describes the Sway bar (see *sway-bar(5)*).
- identifier** (default: 'bar0')  
Identifier of the bar. The value must be a symbol.
- position** (optional)  
Specify the position of the bar. Accepted values are 'top or 'bottom.
- hidden-state** (optional)  
Specify the appearance of the bar when it is hidden. Accepted values are 'hide or 'show.
- binding-mode-indicator** (optional)  
Boolean enabling or disabling the binding mode indicator.
- colors** (optional)  
An optional **sway-color** configuration record.
- status-command** (optional)  
This field accept strings, G-expressions and executable file-like values. The default value is a command (string) that prints the date and time every second.  
Each line printed on **stdout** by this command (or script) will be displayed on the status area of the bar.  
Below are a few examples using:
  - a string: "while date +%Y-%m-%d %X'; do sleep 1; done",

- a G-exp:

```
#~(string-append "while "
                  #coreutils "/bin/date"
                  " +'%Y-%m-%d %X'; do sleep 1; done")
```

- an executable file:

```
(program-file
 "sway-bar-status"
 #~(begin
   (use-modules (ice-9 format)
                 (srfi srfi-19))
   (let loop ()
     (let* ((date (date->string
                     (current-date)
                     "~d/~m/~Y (~a) • ~H:~M:~S"))
            (format #t "~a~%~!" date)
            (sleep 1)
            (loop))))))
```

`mouse-bindings` (default: '())

This field accepts an associative list. Keys are integers describing mouse events. Values are similar to that of key-bindings (except that “to-code” is not a valid option for mouse-bindings).

The module `(gnu home services sway)` exports constants `%ev-code-mouse-left`, `%ev-code-mouse-right` and `%ev-code-mouse-scroll-click` whose values are integers corresponding to left, right and scroll click respectively. For example, with `(mouse-bindings `((,%ev-code-mouse-left . "exec $term")))`, left clicks in the status bar open the terminal (assuming that the variable `$term` is bound to a terminal).

## sway-mode

[Data Type]

Describes a Sway mode (see `sway(5)`). For example, the following snippet defines the resize mode of the default Sway configuration:

```
(sway-mode
 (mode-name "resize")
 (keybindings
 ' (($left . "resize shrink width 10px")
   ($right . "resize grow width 10px")
   ($down . "resize grow height 10px")
   ($up . "resize shrink height 10px")
   (Left . "resize shrink width 10px")
   (Right . "resize grow width 10px")
   (Down . "resize grow height 10px")
   (Up . "resize shrink height 10px")
   (Return . "mode \"default\"")
   (Escape . "mode \"default\"")))))
```

**mode-name** (default: "default")  
 Name of the mode. This field accepts strings.

**keybindings** (default: '()')  
 This field describes keybindings. The value is an association list. As above, keys are symbols and values are either strings, G-expressions or cons-cells.

**mouse-bindings** (default: '()')  
 Ditto, but keys are mouse events (integers). Constants `%ev-code-mouse-*` described above can be used as helpers to define mouse bindings.

### 13.3.19 Networking Home Services

This section lists services somewhat networking-related that you may use with Guix Home.

#### Syncthing Service

The `(gnu home services syncthing)` module provides a service to set up the <https://syncthing.net> (Syncthing) continuous file backup service.

**home-syncthing-service-type** [Variable]

This is the service type for the `syncthing` daemon; it is the Home counterpart of the `syncthing-service-type` system service (see Section 11.10.5 [Networking Services], page 322). The value for this service type is a `syncthing-configuration`.

Here is how you would set it up with the default configuration:

```
(service home-syncthing-service-type
```

For a custom configuration, wrap your `syncthing-configuration` in `for-home`, as in this example:

```
(service home-syncthing-service-type
  (for-home
    (syncthing-configuration (logflags 5))))
```

For details about `syncthing-configuration`, check out the documentation of the system service (see Section 11.10.5 [Networking Services], page 322).

### 13.3.20 Secrets Home Services

The `(gnu home services secrets)` module provides services pertaining to storing secrets, notably including password managers.

#### Himitsu Services

Himitsu (<https://himitsustore.org/>) is a daemon for storing arbitrary secrets encrypted by a single password. The daemon operates alongside a prompter, which asks the user for permission when an application requests a secret, and for the store password on first use in a session. Neither Himitsu nor its underlying cryptography library have been independently audited.

**home-himitsu-service-type** [Variable]

This service provides the main Himitsu store daemon. It will not launch unless you have a store set up - run `himitsu-store -i` after the service is added to your profile

to create it, and then `herd start himitsu` to restart the service. Configuration is as follows:

**home-himitsu-configuration** [Data Type]

**package** (default: `himitsu`)

Package to find `himitsud` in.

**notify-reuse** (optional)

Interpreted shell command to run on application access to an already-approved secret.

**prompter** (default: `(wrap-himitsu-prompter (file-append hiprompt-gtk "/bin/hiprompt-gtk"))`)

Path to a prompter binary. Defaults to `hiprompt-gtk`, the reference implementation. If your prompter is graphical, you should pass it through the provided `wrap-himitsu-prompter` procedure, which will autodetect a running display to prompt on, due to the Himitsu daemon starting before the display server is online.

**extra-options** (optional)

List of extra strings or G-expressions to insert verbatim into the Himitsu configuration file. Each list element gets its own line.

Himitsu can be used through its command-line tool `hiq` or through any number of supporting services, providing compatability to other applications. Git credential support may be provided simply by adding the `himitsu-git` package to your profile, but other such supporting services are listed below:

**home-himitsu-ssh-service-type** [Variable]

This service provides an `ssh-agent` implementation that stores keys through Himitsu. You may interact with a running daemon through `hissh-import` and `hissh-export`.

**home-himitsu-ssh-configuration** [Data Type]

**package** (default: `himitsu-ssh`)

Package to find `hissh-agent` in.

**persist** (default: `'(session 300 refuse)'`)

List of options given when prompting to allow this service access to your list of keys. The option chosen decides how long `himitsu-ssh` has this access. Options can be either `'session` (access is retained until the daemon closes), `'refuse` (access is declined and the prompter will not ask again), `'skip` (ask again next use), or a numerical access timeout in seconds.

**disclose** (default: `'(skip session 300)'`)

List of options given when prompting to allow this service use of your secret keys. Has the same format as `persist`.

**home-himitsu-secret-service-type** [Variable]

This service provides a freedesktop Secret Service (<https://specifications.freedesktop.org/secret-service/latest/>) implementation, allowing Himitsu to



replace see Section 11.10.9 [Desktop Services], page 377, and see Section 11.10.9 [Desktop Services], page 377.

**home-himitsu-secret-service-configuration** [Data Type]  
 package (default: himitsu-secret-service)  
 Package to find hisecrets-agent in.

### 13.3.21 Miscellaneous Home Services

This section lists Home services that lack a better place.

#### Beets Service

The (gnu home services music) module provides the following service:

**home-beets-service-type** [Variable]  
 Beets (<https://beets.io>) is a music file and metadata manager that can be used via its command-line interface, **beet**. Beets requires a YAML configuration file and this Guix Home service is to create such file.

The service can be used as follows:

```
(service home-beets-service-type
  (home-beets-configuration (directory "/home/alice/music")))
```

Additional options can be specified via the service wild-card field **extra-options**:

```
(service home-beets-service-type
  (home-beets-configuration
    (directory "/home/alice/music")
    (extra-options '("
import:
  move: yes")))))
```

#### Dictionary Service

The (gnu home services dict) module provides the following service:

**home-dicod-service-type** [Variable]  
 This is the type of the service that runs the **dicod** daemon, an implementation of DICT server (see Section “Dicod” in *GNU Dico Manual*).  
 You can add **open localhost** to your **~/.dico** file to make **localhost** the default server for dico client (see Section “Initialization File” in *GNU Dico Manual*).

This service is a direct mapping of the **dicod-service-type** system service (see Section 11.10.39 [Miscellaneous Services], page 649). You can use it like this:

```
(service home-dicod-service-type)
```

You may specify a custom configuration by providing a **dicod-configuration** record, exactly like for **dicod-service-type**, but wrapping it in **for-home**:

```
(service home-dicod-service-type
  (for-home
    (dicod-configuration ...)))
```

## ReadyMedia Service

The `(gnu home services upnp)` module provides a service to set up the ReadyMedia (<https://minidlna.sourceforge.net/>) DLNA (Digital Living Network Alliance) UPnP (Universal Plug and Play) media server.

**home-readymedia-service-type** [Variable]

This is the service type for the `minidlnad` daemon; it is the Home counterpart of the `readymedia-service-type` system service (see Section 11.10.39 [Miscellaneous Services], page 649). The value for this service type is a `readymedia-configuration`.

The service can be used with a `readymedia-configuration` wrapped in `for-home` like this:

```
(service home-readymedia-service-type
  (for-home
    (readymedia-configuration
      (media-directories
        (list (readymedia-directory
              (path "/path/to/media/audio")
              (types '(A)))
            (readymedia-directory
              (path "/path/to/media/video")
              (types '(V)))
            (readymedia-directory
              (path "/path/to/media/misc"))))))))
```

For details about `readymedia-configuration`, check out the documentation of the system service (see Section 11.10.39 [Miscellaneous Services], page 649).

## OCI backed services

The `(gnu home services containers)` module provides the following service:

**home-oci-service-type** [Variable]

This is the type of the service that allows to manage your OCI containers with the same consistent interface you use for your other Home Shepherd services.

This service is a direct mapping of the `oci-service-type` system service (see Section 11.10.39 [Miscellaneous Services], page 649). You can use it like this:

```
(use-modules (gnu services containers)
  (gnu home services containers))

(simple-service 'home-oci-provisioning
  home-oci-service-type
  (oci-extension
    (volumes
      (list
        (oci-volume-configuration (name "prometheus"))
        (oci-volume-configuration (name "grafana"))))
    (networks
```

```

(list
  (oci-network-configuration (name "monitoring"))))
(containers
  (list
    (oci-container-configuration
      (network "monitoring")
      (image
        (oci-image
          (repository "guile")
          (tag "3")
          (value (specifications->manifest '("guile"))))
          (pack-options '(:symlinks ((/bin/guile -> "bin/guile"))
                        #:max-layers 2))))
      (entrypoint "/bin/guile")
      (command
        '("-c" "(display \"hello!\n\")"))))
    (oci-container-configuration
      (image "prom/prometheus")
      (network "monitoring")
      (ports
        '(("9000" . "9000")
          ("9090" . "9090")))
      (volumes
        (list
          '(("prometheus" . "/var/lib/prometheus")))))
    (oci-container-configuration
      (image "grafana/grafana:10.0.1")
      (network "monitoring")
      (volumes
        '(("grafana:/var/lib/grafana"))))))))

```

You may specify a custom configuration by providing a `oci-configuration` record, exactly like for `oci-service-type`, but wrapping it in `for-home`:

```

(use-modules (gnu services)
             (gnu services containers)
             (gnu home services containers))

(service home-oci-service-type
  (for-home
    (oci-configuration
      (runtime 'podman)
      (verbose? #t))))

(simple-service 'home-oci-provisioning
  home-oci-service-type
  (oci-extension

```

```

(volumes
  (list
    (oci-volume-configuration (name "prometheus"))
    (oci-volume-configuration (name "grafana"))))
(networks
  (list
    (oci-network-configuration (name "monitoring"))))
(containers
  (list
    (oci-container-configuration
      (network "monitoring")
      (image
        (oci-image
          (repository "guile")
          (tag "3")
          (value (specifications->manifest '("guile"))))
          (pack-options '(:symlinks ("/bin/guile" -> "bin/guile"))
                        #:max-layers 2))))
      (entrypoint "/bin/guile")
      (command
        '("-c" "(display \"hello!\n\")"))))
    (oci-container-configuration
      (image "prom/prometheus")
      (network "monitoring")
      (ports
        '(("9000" . "9000")
          ("9090" . "9090")))
      (volumes
        (list
          '(("prometheus" . "/var/lib/prometheus"))))
      (oci-container-configuration
        (image "grafana/grafana:10.0.1")
        (network "monitoring")
        (volumes
          '(("grafana:/var/lib/grafana"))))))))

```

## 13.4 Invoking guix home

Once you have written a home environment declaration (see Section 13.1 [Declaring the Home Environment], page 732, it can be *instantiated* using the `guix home` command. The synopsis is:

```
guix home options... action file
```

*file* must be the name of a file containing a `home-environment` declaration. *action* specifies how the home environment is instantiated, but there are few auxiliary actions which don't instantiate it. Currently the following values are supported:

**search**     Display available home service type definitions that match the given regular expressions, sorted by relevance:

```
$ guix home search shell
name: home-shell-profile
location: gnu/home/services/shells.scm:100:2
extends: home-files
description: Create '~/.profile', which is used for environment initializati
+ This service type can be extended with a list of file-like objects.█
relevance: 6

name: home-fish
location: gnu/home/services/shells.scm:640:2
extends: home-files home-profile
description: Install and configure Fish, the friendly interactive shell.█
relevance: 3

name: home-zsh
location: gnu/home/services/shells.scm:290:2
extends: home-files home-profile
description: Install and configure Zsh.
relevance: 1

name: home-bash
location: gnu/home/services/shells.scm:508:2
extends: home-files home-profile
description: Install and configure GNU Bash.
relevance: 1

...
```

As for `guix search`, the result is written in `recutils` format, which makes it easy to filter the output (see *GNU recutils manual*).

**container**

Spawn a shell in an isolated environment—a *container*—containing your home as specified by *file*.

For example, this is how you would start an interactive shell in a container with your home:

```
guix home container config.scm
```

This is a throw-away container where you can lightheartedly fiddle with files; any changes made within the container, any process started—all this disappears as soon as you exit that shell.

As with `guix shell`, several options control that container:

`--network`

`-N`            Enable networking within the container (it is disabled by default).

```
--expose=source[=target]
--share=source[=target]
```

As with `guix shell`, make directory *source* of the host system available as *target* inside the container—read-only if you pass `--expose`, and writable if you pass `--share` (see Section 7.1 [Invoking `guix shell`], page 78).

Additionally, you can run a command in that container, instead of spawning an interactive shell. For instance, here is how you would check which Shepherd services are started in a throw-away home container:

```
guix home container config.scm -- herd status
```

The command to run in the container must come after `--` (double hyphen).

**edit** Edit or view the definition of the given Home service types.

For example, the command below opens your editor, as specified by the `EDITOR` environment variable, on the definition of the `home-mcron` service type:

```
guix home edit home-mcron
```

**reconfigure**

Build the home environment described in *file*, and switch to it. Switching means that the activation script will be evaluated and (in basic scenario) symlinks to configuration files generated from `home-environment` declaration will be created in `~`. If the file with the same path already exists in home folder it will be moved to `~/timestamp-guix-home-legacy-configs-backup`, where *timestamp* is a current UNIX epoch time.

**Note:** It is highly recommended to run `guix pull` once before you run `guix home reconfigure` for the first time (see Section 5.7 [Invoking `guix pull`], page 56).

This effects all the configuration specified in *file*. The command starts Shepherd services specified in *file* that are not currently running; if a service is currently running, this command will arrange for it to be upgraded the next time it is stopped (e.g. by `herd stop service` or `herd restart service`).

This command creates a new generation whose number is one greater than the current generation (as reported by `guix home list-generations`). If that generation already exists, it will be overwritten. This behavior mirrors that of `guix package` (see Section 5.2 [Invoking `guix package`], page 37).

Upon completion, the new home is deployed under `~/.guix-home`. This directory contains *provenance meta-data*: the list of channels in use (see Chapter 6 [Channels], page 68) and *file* itself, when available. You can view the provenance information by running:

```
guix home describe
```

This information is useful should you later want to inspect how this particular generation was built. In fact, assuming *file* is self-contained, you can later rebuild generation *n* of your home environment with:

```
guix time-machine \
  -C /var/guix/profiles/per-user/USER/guix-home-n-link/channels.scm -- \
```

```
home reconfigure \
/var/guix/profiles/per-user/USER/guix-home-n-link/configuration.scm
```

You can think of it as some sort of built-in version control! Your home is not just a binary artifact: *it carries its own source*.

**Note:** If you're using Guix System, [guix-home-service-type], page 638, on how to embed your home configuration in your system configuration such that `guix system reconfigure` deploys both your system and your home.

#### switch-generation

Switch to an existing home generation. This action atomically switches the home profile to the specified home generation.

The target generation can be specified explicitly by its generation number. For example, the following invocation would switch to home generation 7:

```
guix home switch-generation 7
```

The target generation can also be specified relative to the current generation with the form `+N` or `-N`, where `+3` means “3 generations ahead of the current generation,” and `-1` means “1 generation prior to the current generation.” When specifying a negative value such as `-1`, you must precede it with `--` to prevent it from being parsed as an option. For example:

```
guix home switch-generation -- -1
```

This action will fail if the specified generation does not exist.

#### roll-back

Switch to the preceding home generation. This is the inverse of `reconfigure`, and it is exactly the same as invoking `switch-generation` with an argument of `-1`.

#### delete-generations

Delete home generations, making them candidates for garbage collection (see Section 5.6 [Invoking `guix gc`], page 53, for information on how to run the “garbage collector”).

This works in the same way as ‘`guix package --delete-generations`’ (see Section 5.2 [Invoking `guix package`], page 37). With no arguments, all home generations but the current one are deleted:

```
guix home delete-generations
```

You can also select the generations you want to delete. The example below deletes all the home generations that are more than two months old:

```
guix home delete-generations 2m
```

**build** Build the derivation of the home environment, which includes all the configuration files and programs needed. This action does not actually install anything.

**describe** Describe the current home generation: its file name, as well as provenance information when available.

To show installed packages in the current home generation's profile, the `--list-installed` flag is provided, with the same syntax that is used in `guix package`

`--list-installed` (see Section 5.2 [Invoking guix package], page 37). For instance, the following command shows a table of all the packages with “emacs” in their name that are installed in the current home generation’s profile:

```
guix home describe --list-installed=emacs
```

#### list-generations

List a summary of each generation of the home environment available on disk, in a human-readable way. This is similar to the `--list-generations` option of `guix package` (see Section 5.2 [Invoking guix package], page 37).

Optionally, one can specify a pattern, with the same syntax that is used in `guix package --list-generations`, to restrict the list of generations displayed. For instance, the following command displays generations that are up to 10 days old:

```
guix home list-generations 10d
```

The `--list-installed` flag may also be specified, with the same syntax that is used in `guix home describe`. This may be helpful if trying to determine when a package was added to the home profile.

#### import

Generate a *home environment* from the packages in the default profile and configuration files found in the user’s home directory. The configuration files will be copied to the specified directory, and a `home-configuration.scm` will be populated with the home environment. Note that not every home service that exists is supported (see Section 13.3 [Home Services], page 735).

```
$ guix home import ~/guix-config
```

```
guix home: '/home/alice/guix-config' populated with all the Home configurati
```

And there’s more! `guix home` also provides the following sub-commands to visualize how the services of your home environment relate to one another:

#### extension-graph

Emit to standard output the *service extension graph* of the home environment defined in *file* (see Section 11.19.1 [Service Composition], page 709, for more information on service extensions). By default the output is in Dot/Graphviz format, but you can choose a different format with `--graph-backend`, as with `guix graph` (see Section 9.10 [Invoking guix graph], page 222):

The command:

```
guix home extension-graph file | xdot -
```

shows the extension relations among services.

#### shepherd-graph

Emit to standard output the *dependency graph* of shepherd services of the home environment defined in *file*. See Section 11.19.4 [Shepherd Services], page 717, for more information and for an example graph.

Again, the default output format is Dot/Graphviz, but you can pass `--graph-backend` to select a different one.

*options* can contain any of the common build options (see Section 9.1.1 [Common Build Options], page 180). In addition, *options* can contain one of the following:



`--expression=expr`

`-e expr` Consider the home-environment `expr` evaluates to. This is an alternative to specifying a file which evaluates to a home environment.

`--allow-downgrades`

Instruct `guix home reconfigure` to allow system downgrades.

Just like `guix system`, `guix home reconfigure`, by default, prevents you from downgrading your home to older or unrelated revisions compared to the channel revisions that were used to deploy it—those shown by `guix home describe`. Using `--allow-downgrades` allows you to bypass that check, at the risk of downgrading your home—be careful!

## 14 Documentation

In most cases packages installed with Guix come with documentation. There are two main documentation formats: “Info”, a browsable hypertext format used for GNU software, and “manual pages” (or “man pages”), the linear documentation format traditionally found on Unix. Info manuals are accessed with the `info` command or with Emacs, and man pages are accessed using `man`.

You can look for documentation of software installed on your system by keyword. For example, the following command searches for information about “TLS” in Info manuals:

```
$ info -k TLS
"(emacs)Network Security" -- STARTTLS
"(emacs)Network Security" -- TLS
"(gnutls)Core TLS API" -- gnutls_certificate_set_verify_flags
"(gnutls)Core TLS API" -- gnutls_certificate_set_verify_function
...
```

The command below searches for the same keyword in man pages<sup>1</sup>:

```
$ man -k TLS
SSL (7)           - OpenSSL SSL/TLS library
certtool (1)      - GnuTLS certificate tool
...
```

These searches are purely local to your computer so you have the guarantee that documentation you find corresponds to what you have actually installed, you can access it off-line, and your privacy is respected.

Once you have these results, you can view the relevant documentation by running, say:

```
$ info "(gnutls)Core TLS API"
```

or:

```
$ man certtool
```

Info manuals contain sections and indices as well as hyperlinks like those found in Web pages. The `info` reader (see *Stand-alone GNU Info*) and its Emacs counterpart (see Section “Misc Help” in *The GNU Emacs Manual*) provide intuitive key bindings to navigate manuals. See Section “Getting Started” in *Info: An Introduction*, for an introduction to Info navigation.

---

<sup>1</sup> The database searched by `man -k` is only created in profiles that contain the `man-db` package.

## 15 Platforms

The packages and systems built by Guix are intended, like most computer programs, to run on a CPU with a specific instruction set, and under a specific operating system. Those programs are often also targeting a specific kernel and system library. Those constraints are captured by Guix in `platform` records.

### 15.1 platform Reference

The `platform` data type describes a *platform*: an ISA (instruction set architecture), combined with an operating system and possibly additional system-wide settings such as the ABI (application binary interface).

`platform` [Data Type]

This is the data type representing a platform.

**target** This field specifies the platform's GNU triplet as a string (see Section "Specifying Target Triplets" in *Autoconf*).

**system** This string is the system type as it is known to Guix and passed, for instance, to the `--system` option of most commands.

It usually has the form "*cpu-kernel*", where *cpu* is the target CPU and *kernel* the target operating system kernel.

It can be for instance "*aarch64-linux*" or "*armhf-linux*". You will encounter system types when you perform native builds (see Section 10.2 [Native Builds], page 240).

**linux-architecture** (default: `#false`)

This optional string field is only relevant if the kernel is Linux. In that case, it corresponds to the ARCH variable used when building Linux, "*mips*" for instance.

**rust-target** (default: `#false`)

This optional string field is used to determine which rust target is best supported by this platform. For example, the base level system targeted by *armhf-linux* system is closest to *armv7-unknown-linux-gnueabi*.

**glibc-dynamic-linker**

This field is the name of the GNU C Library dynamic linker for the corresponding system, as a string. It can be `"/lib/ld-linux-armhf.so.3"`.

### 15.2 Supported Platforms

The `(guix platforms ...)` modules export the following variables, each of which is bound to a `platform` record.

**armv7-linux** [Variable]

Platform targeting ARM v7 CPU running GNU/Linux.

**aarch64-linux** [Variable]

Platform targeting ARM v8 CPU running GNU/Linux.

<b>loongarch64-linux</b>	[Variable]
Platform targeting loongarch 64-bit CPU running GNU/Linux.	
<b>mips64-linux</b>	[Variable]
Platform targeting MIPS little-endian 64-bit CPU running GNU/Linux.	
<b>powerpc-linux</b>	[Variable]
Platform targeting PowerPC big-endian 32-bit CPU running GNU/Linux.	
<b>powerpc64le-linux</b>	[Variable]
Platform targeting PowerPC little-endian 64-bit CPU running GNU/Linux.	
<b>riscv64-linux</b>	[Variable]
Platform targeting RISC-V 64-bit CPU running GNU/Linux.	
<b>i686-linux</b>	[Variable]
Platform targeting x86 CPU running GNU/Linux.	
<b>x86_64-linux</b>	[Variable]
Platform targeting x86 64-bit CPU running GNU/Linux.	
<b>x86_64-linux-x32</b>	[Variable]
Platform targeting x86 64-bit CPU running GNU/Linux with the run-time using the X32 ABI.	
<b>i686-mingw</b>	[Variable]
Platform targeting x86 CPU running Windows, with run-time support from MinGW.	
<b>x86_64-mingw</b>	[Variable]
Platform targeting x86 64-bit CPU running Windows, with run-time support from MinGW.	
<b>i586-gnu</b>	[Variable]
Platform targeting x86 CPU running GNU/Hurd (also referred to as “GNU”).	
<b>avr</b>	[Variable]
Platform targeting AVR CPUs without an operating system, with run-time support from AVR Libc.	
<b>or1k-elf</b>	[Variable]
Platform targeting OpenRISC 1000 CPU without an operating system and without a C standard library.	
<b>xtensa-ath9k-elf</b>	[Variable]
Platform targeting Xtensa CPU used in the Qualcomm Atheros AR7010 and AR9271 USB 802.11n NICs (Network Interface Controllers).	

## 16 Creating System Images

When it comes to installing Guix System for the first time on a new machine, you can basically proceed in three different ways. The first one is to use an existing operating system on the machine to run the `guix system init` command (see Section 11.16 [Invoking guix system], page 692). The second one, is to produce an installation image (see Section 3.9 [Building the Installation Image], page 32). This is a bootable system which role is to eventually run `guix system init`. Finally, the third option would be to produce an image that is a direct instantiation of the system you wish to run. That image can then be copied on a bootable device such as an USB drive or a memory card. The target machine would then directly boot from it, without any kind of installation procedure.

The `guix system image` command is able to turn an operating system definition into a bootable image. This command supports different image types, such as `mbr-hybrid-raw`, `iso9660` and `docker`. Any modern `x86_64` machine will probably be able to boot from an `iso9660` image. However, there are a few machines out there that require specific image types. Those machines, in general using ARM processors, may expect specific partitions at specific offsets.

This chapter explains how to define customized system images and how to turn them into actual bootable images.

### 16.1 image Reference

The `image` record, described right after, allows you to define a customized bootable system image.

**image** [Data Type]

This is the data type representing a system image.

**name** (default: `#false`)

The image name as a symbol, `'my-iso9660` for instance. The name is optional and it defaults to `#false`.

**format** The image format as a symbol. The following formats are supported:

- `disk-image`, a raw disk image composed of one or multiple partitions.
- `compressed-qcow2`, a compressed qcow2 image composed of one or multiple partitions.
- `docker`, a Docker image.
- `iso9660`, an ISO-9660 image.
- `tarball`, a tar.gz image archive.
- `ws12`, a WSL2 image.

**platform** (default: `#false`)

The `platform` record the image is targeting (see Chapter 15 [Platforms], page 791), `aarch64-linux` for instance. By default, this field is set to `#false` and the image will target the host platform.

- size** (default: `'guess'`)  
The image size in bytes or `'guess'`. The `'guess'` symbol, which is the default, means that the image size will be inferred based on the image content.
- operating-system**  
The image's `operating-system` record that is instantiated.
- partition-table-type** (default: `'mbr'`)  
The image partition table type as a symbol. Possible values are `'mbr'` and `'gpt'`. It default to `'mbr'`.
- partitions** (default: `'()'`)  
The image partitions as a list of `partition` records (see Section 16.1.1 [partition Reference], page 794).
- compression?** (default: `#true`)  
Whether the image content should be compressed, as a boolean. It defaults to `#true` and only applies to `'iso9660` image formats.
- volatile-root?** (default: `#true`)  
Whether the image root partition should be made volatile, as a boolean. This is achieved by using a RAM backed file system (overlayfs) that is mounted on top of the root partition by the `initrd`. It defaults to `#true`. When set to `#false`, the image root partition is mounted as read-write partition by the `initrd`.
- shared-store?** (default: `#false`)  
Whether the image's store should be shared with the host system, as a boolean. This can be useful when creating images dedicated to virtual machines. When set to `#false`, which is the default, the image's `operating-system` closure is copied to the image. Otherwise, when set to `#true`, it is assumed that the host store will be made available at boot, using a `9p` mount for instance.
- shared-network?** (default: `#false`)  
Whether to use the host network interfaces within the image, as a boolean. This is only used for the `'docker` image format. It defaults to `#false`.
- substitutable?** (default: `#true`)  
Whether the image derivation should be substitutable, as a boolean. It defaults to `true`.

### 16.1.1 partition Reference

In `image` record may contain some partitions.

**partition** [Data Type]

This is the data type representing an image partition.

- size** (default: `'guess'`)  
The partition size in bytes or `'guess'`. The `'guess'` symbol, which is the default, means that the partition size will be inferred based on the partition content.

**offset** (default: 0)

The partition's start offset in bytes, relative to the image start or the previous partition end. It defaults to 0 which means that there is no offset applied.

**file-system** (default: "ext4")

The partition file system as a string, defaulting to "ext4".

The supported values are "btrfs", "ext4", "fat16", "f2fs", "fat32", "vfat", and "swap".

"vfat", "fat16", and "fat32" partitions without the 'esp flag are by default LBA compatible.

**file-system-options** (default: '()')

The partition file system creation options that should be passed to the partition creation tool, as a list of strings. This is only supported when creating "vfat", "fat16", "fat32" or "ext4" partitions.

See the "extended-options" man page section of the "mke2fs" tool for a more complete reference.

**label** The partition label as a mandatory string, "my-root" for instance.

**uuid** (default: #false)

The partition UUID as an uuid record (see Section 11.4 [File Systems], page 257). By default it is #false, which means that the partition creation tool will attribute a random UUID to the partition.

**flags** (default: '()')

The partition flags as a list of symbols. Possible values are 'boot and 'esp. The 'boot flags should be set if you want to boot from this partition. Exactly one partition should have this flag set, usually the root one. The 'esp flag identifies a UEFI System Partition.

**initializer** (default: #false)

The partition initializer procedure as a gexp. This procedure is called to populate a partition. If no initializer is passed, the `initialize-root-partition` procedure from the (`gnu build image`) module is used.

## 16.2 Instantiate an Image

Let's say you would like to create an MBR image with three distinct partitions:

- The ESP (EFI System Partition), a partition of 40 MiB at offset 1024 KiB with a vfat file system.
- an ext4 partition of 50 MiB data file, and labeled "data".
- an ext4 bootable partition containing the %simple-os operating-system.

You would then write the following image definition in a `my-image.scm` file for instance.

```
(use-modules (gnu)
  (gnu image)
  (gnu tests)
  (gnu system image))
```

```

(guix gexp))

(define MiB (expt 2 20))

(image
  (format 'disk-image)
  (operating-system %simple-os)
  (partitions
    (list
      (partition
        (size (* 40 MiB))
        (offset (* 1024 1024))
        (label "GNU-ESP")
        (file-system "vfat")
        (flags '(esp))
        (initializer (gexp initialize-efi-partition)))
      (partition
        (size (* 50 MiB))
        (label "DATA")
        (file-system "ext4")
        (initializer #~(lambda* (root . rest)
                        (mkdir root)
                        (call-with-output-file
                          (string-append root "/data")
                          (lambda (port)
                            (format port "my-data"))))))))
      (partition
        (size 'guess)
        (label root-label)
        (file-system "ext4")
        (flags '(boot))
        (initializer (gexp initialize-root-partition))))))

```

Note that the first and third partitions use generic initializers procedures, `initialize-efi-partition` and `initialize-root-partition` respectively. The `initialize-efi-partition` installs a GRUB EFI loader that is loading the GRUB bootloader located in the root partition. The `initialize-root-partition` instantiates a complete system as defined by the `%simple-os` operating-system.

You can now run:

```
guix system image my-image.scm
```

to instantiate the `image` definition. That produces a disk image which has the expected structure:

```

$ parted $(guix system image my-image.scm) print
...
Model: (file)
Disk /gnu/store/yhylv1bp5b2ypb97pd3bbhz6jk5nbhwx-disk-image: 1714MB
Sector size (logical/physical): 512B/512B

```



Partition Table: msdos  
Disk Flags:

Number	Start	End	Size	Type	File system	Flags
1	1049kB	43.0MB	41.9MB	primary	fat16	esp
2	43.0MB	95.4MB	52.4MB	primary	ext4	
3	95.4MB	1714MB	1619MB	primary	ext4	boot

The size of the boot partition has been inferred to 1619MB so that it is large enough to host the %simple-os operating-system.

You can also use existing `image` record definitions and inherit from them to simplify the image definition. The `(gnu system image)` module provides the following `image` definition variables.

**mbr-disk-image** [Variable]

An MBR disk-image composed of a single ROOT partition. The ROOT partition starts at a 1 MiB offset so that the bootloader can install itself in the post-MBR gap.

**mbr-hybrid-disk-image** [Variable]

An MBR disk-image composed of two partitions: a 64 bits ESP partition and a ROOT boot partition. The ESP partition starts at a 1 MiB offset so that a BIOS compatible bootloader can install itself in the post-MBR gap. The image can be used by `x86_64` and `i686` machines supporting only legacy BIOS booting. The ESP partition ensures that it can also be used by newer machines relying on UEFI booting, hence the *hybrid* denomination.

**efi-disk-image** [Variable]

A GPT disk-image composed of two partitions: a 64 bits ESP partition and a ROOT boot partition. This image can be used on most `x86_64` and `i686` machines, supporting BIOS or UEFI booting.

**efi32-disk-image** [Variable]

Same as `efi-disk-image` but with a 32 bits EFI partition.

**iso9660-image** [Variable]

An ISO-9660 image composed of a single bootable partition. This image can also be used on most `x86_64` and `i686` machines.

**docker-image** [Variable]

A Docker image that can be used to spawn a Docker container.

Using the `efi-disk-image` we can simplify our previous `image` declaration this way:

```
(use-modules (gnu)
              (gnu image)
              (gnu tests)
              (gnu system image)
              (guix gexp)
              (ice-9 match))
```

```

(define MiB (expt 2 20))

(define data
  (partition
    (size (* 50 MiB))
    (label "DATA")
    (file-system "ext4")
    (initializer #~(lambda* (root . rest)
                     (mkdir root)
                     (call-with-output-file
                      (string-append root "/data")
                      (lambda (port)
                        (format port "my-data"))))))))

(image
  (inherit efi-disk-image)
  (operating-system %simple-os)
  (partitions
    (match (image-partitions efi-disk-image)
      ((esp root)
       (list esp data root)))))

```

This will give the exact same `image` instantiation but the `image` declaration is simpler.

### 16.3 image-type Reference

The `guix system image` command can, as we saw above, take a file containing an `image` declaration as argument and produce an actual disk image from it. The same command can also handle a file containing an `operating-system` declaration as argument. In that case, how is the `operating-system` turned into an image?

That's where the `image-type` record intervenes. This record defines how to transform an `operating-system` record into an `image` record.

**image-type** [Data Type]

This is the data type representing an image-type.

**name**           The image-type name as a mandatory symbol, 'efi32-raw for instance.

**constructor**       The image-type constructor, as a mandatory procedure that takes an `operating-system` record as argument and returns an `image` record.

There are several `image-type` records provided by the `(gnu system image)` and the `(gnu system images ...)` modules.

**mbr-raw-image-type** [Variable]

Build an image based on the `mbr-disk-image` image.

**mbr-hybrid-raw-image-type** [Variable]

Build an image based on the `mbr-hybrid-disk-image` image.

- efi-raw-image-type** [Variable]  
Build an image based on the `efi-disk-image` image.
- efi32-raw-image-type** [Variable]  
Build an image based on the `efi32-disk-image` image.
- qcow2-image-type** [Variable]  
Build an image based on the `mbr-disk-image` image but with the `compressed-qcow2` image format. The resulting image will have an MBR embedded bootloader as well as an EFI bootloader. This image is not suitable for architectures that do not support ‘grub-pc’, such as AArch64. See `qcow2-gpt-image-type` for an alternative.
- qcow2-gpt-image-type** [Variable]  
Build an image based on the `efi-disk-image` image but with the `compressed-qcow2` image format. The resulting image will have only EFI bootloader, unlike `qcow2-image-type`. This image is suitable for architectures that do not support ‘grub-pc’, such as AArch64.
- iso-image-type** [Variable]  
Build a compressed image based on the `iso9660-image` image.
- uncompressed-iso-image-type** [Variable]  
Build an image based on the `iso9660-image` image but with the `compression?` field set to `#false`.
- docker-image-type** [Variable]  
Build an image based on the `docker-image` image.
- raw-with-offset-image-type** [Variable]  
Build an MBR image with a single partition starting at a 1024KiB offset. This is useful to leave some room to install a bootloader in the post-MBR gap.
- pinebook-pro-image-type** [Variable]  
Build an image that is targeting the Pinebook Pro machine. The MBR image contains a single partition starting at a 9MiB offset. The `u-boot-pinebook-pro-rk3399-bootloader` bootloader will be installed in this gap.
- rock64-image-type** [Variable]  
Build an image that is targeting the Rock64 machine. The MBR image contains a single partition starting at a 16MiB offset. The `u-boot-rock64-rk3328-bootloader` bootloader will be installed in this gap.
- novena-image-type** [Variable]  
Build an image that is targeting the Novena machine. It has the same characteristics as `raw-with-offset-image-type`.
- pine64-image-type** [Variable]  
Build an image that is targeting the Pine64 machine. It has the same characteristics as `raw-with-offset-image-type`.

**hurdl-image-type** [Variable]

Build an image that is targeting a i386 machine running the Hurd kernel. The MBR image contains a single ext2 partitions with specific `file-system-options` flags.

**hurdl-qcow2-image-type** [Variable]

Build an image similar to the one built by the `hurdl-image-type` but with the `format` set to `'compressed-qcow2'`.

**wsl2-image-type** [Variable]

Build an image for the WSL2 (Windows Subsystem for Linux 2). It can be imported by running:

```
wsl --import Guix ./guix ./wsl2-image.tar.gz
wsl -d Guix
```

So, if we get back to the `guix system image` command taking an `operating-system` declaration as argument. By default, the `mbr-raw-image-type` is used to turn the provided `operating-system` into an actual bootable image.

To use a different `image-type`, the `--image-type` option can be used. The `--list-image-types` option will list all the supported image types. It turns out to be a textual listing of all the `image-types` variables described just above (see Section 11.16 [Invoking `guix system`], page 692).

## 16.4 Image Modules

Let's take the example of the Pine64, an ARM based machine. To be able to produce an image targeting this board, we need the following elements:

- An `operating-system` record containing at least an appropriate kernel (`linux-libre-arm64-generic`) and bootloader `u-boot-pine64-lts-bootloader`) for the Pine64.
- Possibly, an `image-type` record providing a way to turn an `operating-system` record to an `image` record suitable for the Pine64.
- An actual `image` that can be instantiated with the `guix system image` command.

The `(gnu system images pine64)` module provides all those elements: `pine64-barebones-os`, `pine64-image-type` and `pine64-barebones-raw-image` respectively.

The module returns the `pine64-barebones-raw-image` in order for users to be able to run:

```
guix system image gnu/system/images/pine64.scm
```

Now, thanks to the `pine64-image-type` record declaring the `'pine64-raw image-type`, one could also prepare a `my-pine.scm` file with the following content:

```
(use-modules (gnu system images pine64))
(operating-system
  (inherit pine64-barebones-os)
  (timezone "Europe/Athens"))
```

to customize the `pine64-barebones-os`, and run:

```
$ guix system image --image-type=pine64-raw my-pine.scm
```

Note that there are other modules in the `gnu/system/images` directory targeting Novena, Pine64, PinebookPro and Rock64 machines.

## 17 Installing Debugging Files

Program binaries, as produced by the GCC compilers for instance, are typically written in the ELF format, with a section containing *debugging information*. Debugging information is what allows the debugger, GDB, to map binary code to source code; it is required to debug a compiled program in good conditions.

This chapter explains how to use separate debug info when packages provide it, and how to rebuild packages with debug info when it's missing.

### 17.1 Separate Debug Info

The problem with debugging information is that it takes up a fair amount of disk space. For example, debugging information for the GNU C Library weighs in at more than 60 MiB. Thus, as a user, keeping all the debugging info of all the installed programs is usually not an option. Yet, space savings should not come at the cost of an impediment to debugging—especially in the GNU system, which should make it easier for users to exert their computing freedom (see Section 1.2 [GNU Distribution], page 2).

Thankfully, the GNU Binary Utilities (Binutils) and GDB provide a mechanism that allows users to get the best of both worlds: debugging information can be stripped from the binaries and stored in separate files. GDB is then able to load debugging information from those files, when they are available (see Section “Separate Debug Files” in *Debugging with GDB*).

The GNU distribution takes advantage of this by storing debugging information in the `lib/debug` sub-directory of a separate package output unimaginatively called `debug` (see Section 5.4 [Packages with Multiple Outputs], page 51). Users can choose to install the `debug` output of a package when they need it. For instance, the following command installs the debugging information for the GNU C Library and for GNU Guile:

```
guix install glibc:debug guile:debug
```

GDB must then be told to look for debug files in the user's profile, by setting the `debug-file-directory` variable (consider setting it from the `~/.gdbinit` file, see Section “Startup” in *Debugging with GDB*):

```
(gdb) set debug-file-directory ~/.guix-profile/lib/debug
```

From there on, GDB will pick up debugging information from the `.debug` files under `~/.guix-profile/lib/debug`.

Below is an alternative GDB script which is useful when working with other profiles. It takes advantage of the optional Guile integration in GDB. This snippet is included by default on Guix System in the `~/.gdbinit` file.

```
guile
(use-modules (gdb))
(execute (string-append "set debug-file-directory "
                        (or (getenv "GDB_DEBUG_FILE_DIRECTORY")
                            "~/.guix-profile/lib/debug"))))
end
```

In addition, you will most likely want GDB to be able to show the source code being debugged. To do that, you will have to unpack the source code of the package of interest

(obtained with `guix build --source`, see Section 9.1 [Invoking guix build], page 180), and to point GDB to that source directory using the `directory` command (see Section “Source Path” in *Debugging with GDB*).

The `debug` output mechanism in Guix is implemented by the `gnu-build-system` (see Section 8.5 [Build Systems], page 121). Currently, it is opt-in—debugging information is available only for the packages with definitions explicitly declaring a `debug` output. To check whether a package has a `debug` output, use `guix package --list-available` (see Section 5.2 [Invoking guix package], page 37).

Read on for how to deal with packages lacking a `debug` output.

## 17.2 Rebuilding Debug Info

As we saw above, some packages, but not all, provide debugging info in a `debug` output. What can you do when debugging info is missing? The `--with-debug-info` option provides a solution to that: it allows you to rebuild the package(s) for which debugging info is missing—and only those—and to graft those onto the application you’re debugging. Thus, while it’s not as fast as installing a `debug` output, it is relatively inexpensive.

Let’s illustrate that. Suppose you’re experiencing a bug in Inkscape and would like to see what’s going on in GLib, a library that’s deep down in its dependency graph. As it turns out, GLib does not have a `debug` output and the backtrace GDB shows is all sadness:

```
(gdb) bt
#0  0x00007ffff5f92190 in g_getenv ()
    from /gnu/store/...-glib-2.62.6/lib/libglib-2.0.so.0
#1  0x00007ffff608a7d6 in gobject_init_ctor ()
    from /gnu/store/...-glib-2.62.6/lib/libgobject-2.0.so.0
#2  0x00007ffff7fe275a in call_init (l=<optimized out>, argc=argc@entry=1, argv=argv@entry=argv@entry, env=env@entry=0x7ffffffffffcfe8) at dl-init.c:72
#3  0x00007ffff7fe2866 in call_init (env=0x7ffffffffffcfe8, argv=0x7ffffffffffcfd8, argc=1, at dl-init.c:118
```

To address that, you install Inkscape linked against a variant GLib that contains debug info:

```
guix install inkscape --with-debug-info=glib
```

This time, debugging will be a whole lot nicer:

```
$ gdb --args sh -c 'exec inkscape'
...
(gdb) b g_getenv
Function "g_getenv" not defined.
Make breakpoint pending on future shared library load? (y or [n]) y
Breakpoint 1 (g_getenv) pending.
(gdb) r
Starting program: /gnu/store/...-profile/bin/sh -c exec\ inkscape
...
(gdb) bt
#0  g_getenv (variable=variable@entry=0x7ffff60c7a2e "GOBJECT_DEBUG") at ../glib-2.62.
#1  0x00007ffff608a7d6 in gobject_init () at ../glib-2.62.6/gobject/gtype.c:4380
```

```
#2 gobject_init_ctor () at ../glib-2.62.6/gobject/gtype.c:4493
#3 0x00007ffff7fe275a in call_init (l=<optimized out>, argc=argc@entry=3, argv=argv@entry=argv@entry,
    env=env@entry=0x7fffffd0a8) at dl-init.c:72
...
```

Much better!

Note that there can be packages for which `--with-debug-info` will not have the desired effect. See Section 9.1.2 [Package Transformation Options], page 182, for more information.

## 18 Using T<sub>E</sub>X and L<sup>A</sup>T<sub>E</sub>X

Guix provides packages for the T<sub>E</sub>X, L<sup>A</sup>T<sub>E</sub>X, ConTeXt, LuaTeX, and related typesetting systems, taken from the T<sub>E</sub>X Live distribution (<https://www.tug.org/texlive/>). Because T<sub>E</sub>X Live is so huge and because finding one's way in this maze is tricky, this section provides some guidance on how to deploy the relevant packages to compile T<sub>E</sub>X and L<sup>A</sup>T<sub>E</sub>X documents. Guix contains the full T<sub>E</sub>X Live distribution in a modular setup, in which you only install the packages you need, whose name is always prefixed with ‘texlive-’.

Building a coherent system that provides all the essential tools and, at the same time, satisfies all of its internal dependencies can be a difficult task. It is therefore recommended to start with sets of packages, called *collections*, and *schemes*, the name for collections of collections. The following command lists available schemes and collections (see [Invoking guix package], page 43):

```
guix search texlive-\(scheme\|collection\) | recsel -p name,description
```

If needed, you may then complete your system with individual packages, particularly when they belong to a large collection you're not otherwise interested in.

For instance, the following manifest is a reasonable, yet frugal starting point for a French L<sup>A</sup>T<sub>E</sub>X user:

```
(specifications->manifest
  ("rubber"

   "texlive-scheme-basic"
   "texlive-collection-latexrecommended"
   "texlive-collection-fontsrecommended"
   "texlive-babel-french"

   ;; From "latexextra" collection.
   "texlive-tabulararray"
   ;; From "binextra" collection.
   "texlive-texdoc"))
```

If you come across a document that does not compile in such a basic setting, the main difficulty is finding the missing packages. In this case, `pdflatex` and similar commands tend to fail with obscure error messages along the lines of:

```
doc.tex: File `tikz.sty' not found.
doc.tex:7: Emergency stop.
```

or, for a missing font:

```
kpathsea: Running mktexmf phvr7t
! I can't find file `phvr7t'.
```

How do you determine what the missing package is? In the first case, you will find the answer by running:

```
$ guix search texlive tikz
name: texlive-pgf
version: 59745
...
```



In the second case, `guix search` turns up nothing. Instead, you can search the T<sub>E</sub>X Live package database using the `tlmgr` command:

```
$ tlmgr info phvr7t
tlmgr: cannot find package phvr7t, searching for other matches:
```

```
Packages containing `phvr7t' in their title/description:
```

```
Packages containing files matching `phvr7t':
```

```
helvetic:
```

```
texmf-dist/fonts/tfm/adobe/helvetic/phvr7t.tfm
texmf-dist/fonts/tfm/adobe/helvetic/phvr7tn.tfm
texmf-dist/fonts/vf/adobe/helvetic/phvr7t.vf
texmf-dist/fonts/vf/adobe/helvetic/phvr7tn.vf
```

```
tex4ht:
```

```
texmf-dist/tex4ht/ht-fonts/alias/adobe/helvetic/phvr7t.htf
```

The file is available in the T<sub>E</sub>X Live `helvetic` package, which is known in Guix as `texlive-helvetic`. Quite a ride, but you found it!

## 19 Security Updates

Occasionally, important security vulnerabilities are discovered in software packages and must be patched. Guix developers try hard to keep track of known vulnerabilities and to apply fixes as soon as possible in the `master` branch of Guix (we do not yet provide a “stable” branch containing only security updates). The `guix lint` tool helps developers find out about vulnerable versions of software packages in the distribution:

```
$ guix lint -c cve
gnu/packages/base.scm:652:2: glibc@2.21: probably vulnerable to CVE-2015-1781, CVE-2015-7547
gnu/packages/gcc.scm:334:2: gcc@4.9.3: probably vulnerable to CVE-2015-5276
gnu/packages/image.scm:312:2: openjpeg@2.1.0: probably vulnerable to CVE-2016-1923, CVE-2016-1924
...
```

See Section 9.8 [Invoking `guix lint`], page 217, for more information.

Guix follows a functional package management discipline (see Chapter 1 [Introduction], page 1), which implies that, when a package is changed, *every package that depends on it* must be rebuilt. This can significantly slow down the deployment of fixes in core packages such as `libc` or `Bash`, since basically the whole distribution would need to be rebuilt. Using pre-built binaries helps (see Section 5.3 [Substitutes], page 46), but deployment may still take more time than desired.

To address this, Guix implements *grafts*, a mechanism that allows for fast deployment of critical updates without the costs associated with a whole-distribution rebuild. The idea is to rebuild only the package that needs to be patched, and then to “graft” it onto packages explicitly installed by the user and that were previously referring to the original package. The cost of grafting is typically very low, and order of magnitudes lower than a full rebuild of the dependency chain.

For instance, suppose a security update needs to be applied to `Bash`. Guix developers will provide a package definition for the “fixed” `Bash`, say `bash-fixed`, in the usual way (see Section 8.2 [Defining Packages], page 101). Then, the original package definition is augmented with a `replacement` field pointing to the package containing the bug fix:

```
(define bash
  (package
    (name "bash")
    ;; ...
    (replacement bash-fixed)))
```

From there on, any package depending directly or indirectly on `Bash`—as reported by `guix gc --requisites` (see Section 5.6 [Invoking `guix gc`], page 53)—that is installed is automatically “rewritten” to refer to `bash-fixed` instead of `bash`. This grafting process takes time proportional to the size of the package, usually less than a minute for an “average” package on a recent machine. Grafting is recursive: when an indirect dependency requires grafting, then grafting “propagates” up to the package that the user is installing.

Currently, the length of the name and version of the graft and that of the package it replaces (`bash-fixed` and `bash` in the example above) must be equal. This restriction mostly comes from the fact that grafting works by patching files, including binary files, directly. Other restrictions may apply: for instance, when adding a graft to a package providing a shared library, the original shared library and its replacement must have the same `SONAME` and be binary-compatible.

The `--no-grafts` command-line option allows you to forcefully avoid grafting (see Section 9.1.1 [Common Build Options], page 180). Thus, the command:

```
guix build bash --no-grafts
```

returns the store file name of the original Bash, whereas:

```
guix build bash
```

returns the store file name of the “fixed”, replacement Bash. This allows you to distinguish between the two variants of Bash.

To verify which Bash your whole profile refers to, you can run (see Section 5.6 [Invoking `guix gc`], page 53):

```
guix gc -R $(readlink -f ~/.guix-profile) | grep bash
```

... and compare the store file names that you get with those above. Likewise for a complete Guix system generation:

```
guix gc -R $(guix system build my-config.scm) | grep bash
```

Lastly, to check which Bash running processes are using, you can use the `lsuf` command:

```
lsuf | grep /gnu/store/.*bash
```

## 20 Bootstrapping

Bootstrapping in our context refers to how the distribution gets built “from nothing”. Remember that the build environment of a derivation contains nothing but its declared inputs (see Chapter 1 [Introduction], page 1). So there’s an obvious chicken-and-egg problem: how does the first package get built? How does the first compiler get compiled?

It is tempting to think of this question as one that only die-hard hackers may care about. However, while the answer to that question is technical in nature, its implications are wide-ranging. How the distribution is bootstrapped defines the extent to which we, as individuals and as a collective of users and hackers, can trust the software we run. It is a central concern from the standpoint of *security* and from a *user freedom* viewpoint.

The GNU system is primarily made of C code, with `libc` at its core. The GNU build system itself assumes the availability of a Bourne shell and command-line tools provided by GNU Coreutils, Awk, Findutils, ‘sed’, and ‘grep’. Furthermore, build programs—programs that run `./configure`, `make`, etc.—are written in Guile Scheme (see Section 8.10 [Derivations], page 158). Consequently, to be able to build anything at all, from scratch, Guix relies on pre-built binaries of Guile, GCC, Binutils, `libc`, and the other packages mentioned above—the *bootstrap binaries*.

These bootstrap binaries are “taken for granted”, though we can also re-create them if needed (see Section 20.2 [Preparing to Use the Bootstrap Binaries], page 811).

### 20.1 The Full-Source Bootstrap

Guix—like other GNU/Linux distributions—is traditionally bootstrapped from a set of bootstrap binaries: Bourne shell, command-line tools provided by GNU Coreutils, Awk, Findutils, ‘sed’, and ‘grep’ and Guile, GCC, Binutils, and the GNU C Library (see Chapter 20 [Bootstrapping], page 808). Usually, these bootstrap binaries are “taken for granted.”

Taking the bootstrap binaries for granted means that we consider them to be a correct and trustworthy “seed” for building the complete system. Therein lies a problem: the combined size of these bootstrap binaries is about 250MB (see Section “Bootstrappable Builds” in *GNU Mes*). Auditing or even inspecting these is next to impossible.

For `i686-linux` and `x86_64-linux`, Guix now features a *full-source bootstrap*. This bootstrap is rooted in `hex0-seed` from the Stage0 (<https://savannah.gnu.org/projects/stage0>) package. The `hex0` program is minimalist assembler: it reads space-separated hexadecimal digits (nibbles) from a file, possibly including comments, and emits on standard output the bytes corresponding to those hexadecimal numbers. The source code of this initial `hex0` program is a file called `hex0_x86.hex0` ([https://github.com/oriansj/bootstrap-seeds/blob/master/POSIX/x86/hex0\\_x86.hex0](https://github.com/oriansj/bootstrap-seeds/blob/master/POSIX/x86/hex0_x86.hex0)) and is written in the `hex0` language.

Hex0 is self-hosting, which means that it can build itself:

```
./hex0-seed hex0_x86.hex0 hex0
```

Hex0 is the ASCII-equivalent of the binary program and can be produced by doing something much like:

```
sed 's/[;#].*$/g' hex0_x86.hex0 | xxd -r -p > hex0
chmod +x hex0
```

It is because of this ASCII-binary equivalence that we can bless this initial 357-byte binary as source, and hence ‘full-source bootstrap’.

The bootstrap then continues: `hex0` builds `hex1` and then on to `M0`, `hex2`, `M1`, `mescc-tools` and finally `M2-Planet`. Then, using `mescc-tools`, `M2-Planet` we build Mes (see *GNU Mes*, a Scheme interpreter and C compiler in Scheme). From here on starts the more traditional C-based bootstrap of the GNU System.


Another step that Guix has taken is to replace the shell and all its utilities with implementations in Guile Scheme, the *Scheme-only bootstrap*. Gash (see Section “Gash” in *The Gash manual*) is a POSIX-compatible shell that replaces Bash, and it comes with Gash Utils which has minimalist replacements for Awk, the GNU Core Utilities, Grep, Gzip, Sed, and Tar.

Building the GNU System from source is currently only possible by adding some historical GNU packages as intermediate steps<sup>1</sup>. As Gash and Gash Utils mature, and GNU packages become more bootstrappable again (e.g., new releases of GNU Sed will also ship as gzipped tarballs again, as alternative to the hard to bootstrap `xz`-compression), this set of added packages can hopefully be reduced again.

The graph below shows the resulting dependency graph for `gcc-core-mesboot0`, the bootstrap compiler used for the traditional bootstrap of the rest of the Guix System.

---

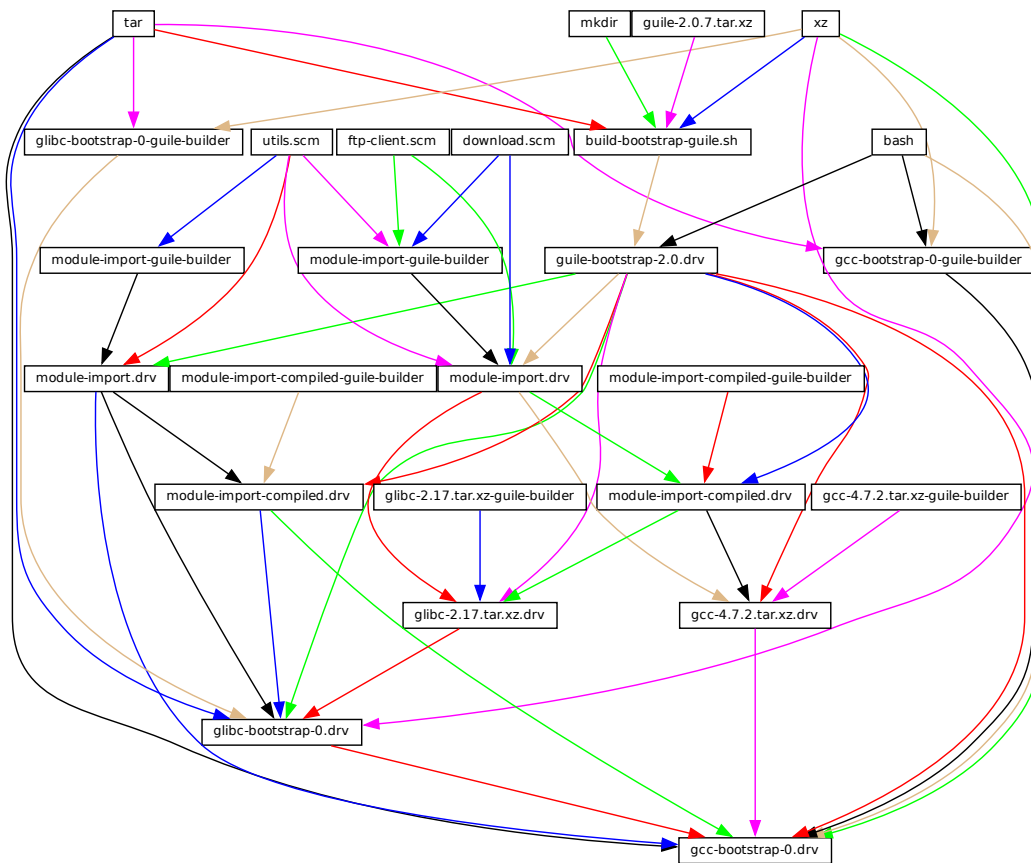
<sup>1</sup> Packages such as `gcc-2.95.3`, `binutils-2.14`, `glibc-2.2.5`, `gzip-1.2.4`, `tar-1.22`, and some others. For details, see `gnu/packages/commencement.scm`.



Work is ongoing to bring these bootstraps to the `arm-linux` and `aarch64-linux` architectures and to the Hurd.

If you are interested, join us on ‘#bootstrappable’ on the Libera.Chat IRC network or discuss on `bug-mes@gnu.org` or `gash-devel@nongnu.org`.

## 20.2 Preparing to Use the Bootstrap Binaries



The figure above shows the very beginning of the dependency graph of the distribution, corresponding to the package definitions of the `(gnu packages bootstrap)` module. A similar figure can be generated with `guix graph` (see Section 9.10 [Invoking `guix graph`], page 222), along the lines of:

```
guix graph -t derivation \
  -e '(@@ (gnu packages bootstrap) %bootstrap-gcc)' \
  | dot -Tps > gcc.ps
```

or, for the further Reduced Binary Seed bootstrap

```
guix graph -t derivation \
  -e '(@@ (gnu packages bootstrap) %bootstrap-mes)' \
  | dot -Tps > mes.ps
```

At this level of detail, things are slightly complex. First, Guile itself consists of an ELF executable, along with many source and compiled Scheme files that are dynamically loaded when it runs. This gets stored in the `guile-2.0.7.tar.xz` tarball shown in this graph.

This tarball is part of Guix’s “source” distribution, and gets inserted into the store with `add-to-store` (see Section 8.9 [The Store], page 156).

But how do we write a derivation that unpacks this tarball and adds it to the store? To solve this problem, the `guile-bootstrap-2.0.drv` derivation—the first one that gets built—uses `bash` as its builder, which runs `build-bootstrap-guile.sh`, which in turn calls `tar` to unpack the tarball. Thus, `bash`, `tar`, `xz`, and `mkdir` are statically-linked binaries, also part of the Guix source distribution, whose sole purpose is to allow the Guile tarball to be unpacked.

Once `guile-bootstrap-2.0.drv` is built, we have a functioning Guile that can be used to run subsequent build programs. Its first task is to download tarballs containing the other pre-built binaries—this is what the `.tar.xz.drv` derivations do. Guix modules such as `ftp-client.scm` are used for this purpose. The `module-import.drv` derivations import those modules in a directory in the store, using the original layout. The `module-import-compiled.drv` derivations compile those modules, and write them in an output directory with the right layout. This corresponds to the `#:modules` argument of `build-expression->derivation` (see Section 8.10 [Derivations], page 158).

Finally, the various tarballs are unpacked by the derivations `gcc-bootstrap-0.drv`, `glibc-bootstrap-0.drv`, or `bootstrap-mes-0.drv` and `bootstrap-mescc-tools-0.drv`, at which point we have a working C tool chain.

## Building the Build Tools

Bootstrapping is complete when we have a full tool chain that does not depend on the pre-built bootstrap tools discussed above. This no-dependency requirement is verified by checking whether the files of the final tool chain contain references to the `/gnu/store` directories of the bootstrap inputs. The process that leads to this “final” tool chain is described by the package definitions found in the `(gnu packages commencement)` module.

The `guix graph` command allows us to “zoom out” compared to the graph above, by looking at the level of package objects instead of individual derivations—remember that a package may translate to several derivations, typically one derivation to download its source, one to build the Guile modules it needs, and one to actually build the package from source. The command:

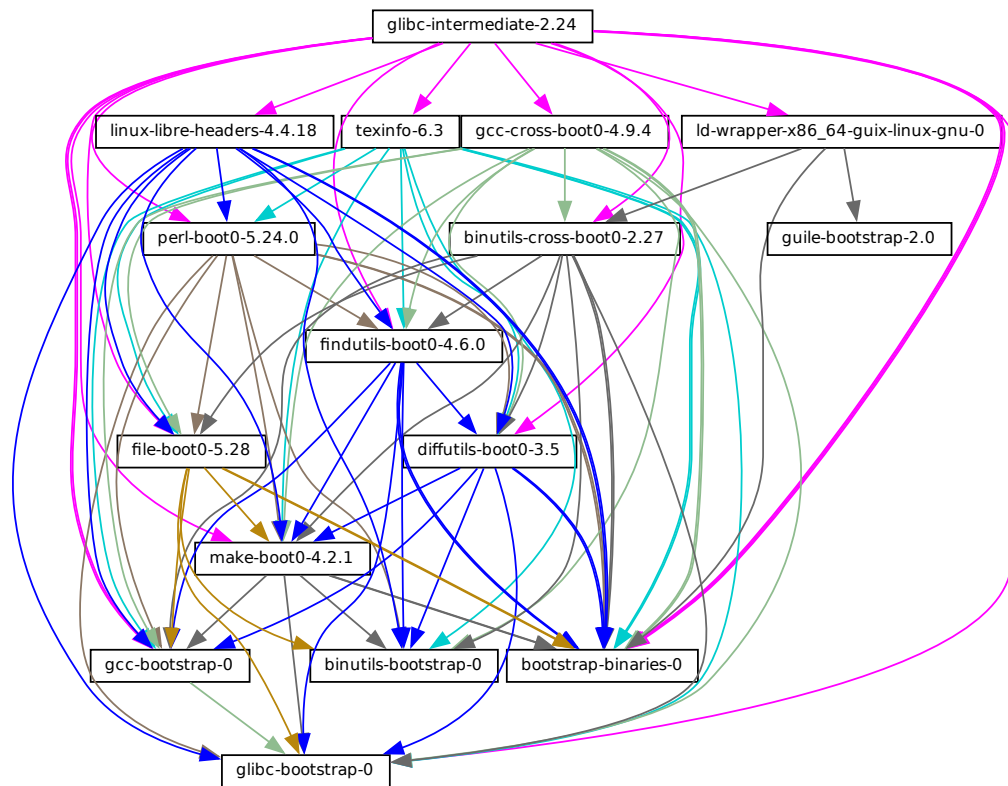
```
guix graph -t bag \
  -e '(@@ (gnu packages commencement)
        glibc-final-with-bootstrap-bash)' | xdot -
```

displays the dependency graph leading to the “final” C library<sup>2</sup>, depicted below.

---

<sup>2</sup> You may notice the `glibc-intermediate` label, suggesting that it is not *quite* final, but as a good approximation, we will consider it final.





The first tool that gets built with the bootstrap binaries is GNU Make—noted **make-boot0** above—which is a prerequisite for all the following packages. From there Findutils and Diffutils get built.

Then come the first-stage Binutils and GCC, built as pseudo cross tools—i.e., with `--target` equal to `--host`. They are used to build libc. Thanks to this cross-build trick, this libc is guaranteed not to hold any reference to the initial tool chain.

From there the final Binutils and GCC (not shown above) are built. GCC uses `ld` from the final Binutils, and links programs against the just-built libc. This tool chain is used to build the other packages used by Guix and by the GNU Build System: Guile, Bash, Coreutils, etc.

And voilà! At this point we have the complete set of build tools that the GNU Build System expects. These are in the `%final-inputs` variable of the `(gnu packages commencement)` module, and are implicitly used by any package that uses `gnu-build-system` (see Section 8.5 [Build Systems], page 121).

## Building the Bootstrap Binaries

Because the final tool chain does not depend on the bootstrap binaries, those rarely need to be updated. Nevertheless, it is useful to have an automated way to produce them, should an update occur, and this is what the `(gnu packages make-bootstrap)` module provides.

The following command builds the tarballs containing the bootstrap binaries (Binutils, GCC, glibc, for the traditional bootstrap and linux-libre-headers, bootstrap-mescc-tools, bootstrap-mes for the Reduced Binary Seed bootstrap, and Guile, and a tarball containing a mixture of Coreutils and other basic command-line tools):

```
guix build bootstrap-tarballs
```

The generated tarballs are those that should be referred to in the `(gnu packages bootstrap)` module mentioned at the beginning of this section.

Still here? Then perhaps by now you've started to wonder: when do we reach a fixed point? That is an interesting question! The answer is unknown, but if you would like to investigate further (and have significant computational and storage resources to do so), then let us know.

## Reducing the Set of Bootstrap Binaries

Our traditional bootstrap includes GCC, GNU Libc, Guile, etc. That's a lot of binary code! Why is that a problem? It's a problem because these big chunks of binary code are practically non-auditable, which makes it hard to establish what source code produced them. Every unauditable binary also leaves us vulnerable to compiler backdoors as described by Ken Thompson in the 1984 paper *Reflections on Trusting Trust*.

This is mitigated by the fact that our bootstrap binaries were generated from an earlier Guix revision. Nevertheless it lacks the level of transparency that we get in the rest of the package dependency graph, where Guix always gives us a source-to-binary mapping. Thus, our goal is to reduce the set of bootstrap binaries to the bare minimum.

The Bootstrappable.org web site (<https://bootstrappable.org>) lists on-going projects to do that. One of these is about replacing the bootstrap GCC with a sequence of assemblers, interpreters, and compilers of increasing complexity, which could be built from source starting from a simple and auditable assembler.

Our first major achievement is the replacement of GCC, the GNU C Library and Binutils by MesCC-Tools (a simple hex linker and macro assembler) and Mes (see *GNU Mes*, a Scheme interpreter and C compiler in Scheme). Neither MesCC-Tools nor Mes can be fully bootstrapped yet and thus we inject them as binary seeds. We call this the Reduced Binary Seed bootstrap, as it has halved the size of our bootstrap binaries! Also, it has eliminated the C compiler binary; i686-linux and x86\_64-linux Guix packages are now bootstrapped without any binary C compiler.

Work is ongoing to make MesCC-Tools and Mes fully bootstrappable and we are also looking at any other bootstrap binaries. Your help is welcome!

## 21 Porting to a New Platform

As discussed above, the GNU distribution is self-contained, and self-containment is achieved by relying on pre-built “bootstrap binaries” (see Chapter 20 [Bootstrapping], page 808). These binaries are specific to an operating system kernel, CPU architecture, and application binary interface (ABI). Thus, to port the distribution to a platform that is not yet supported, one must build those bootstrap binaries, and update the `(gnu packages bootstrap)` module to use them on that platform.

Fortunately, Guix can *cross compile* those bootstrap binaries. When everything goes well, and assuming the GNU tool chain supports the target platform, this can be as simple as running a command like this one:

```
guix build --target=armv5tel-linux-gnueabi bootstrap-tarballs
```

For this to work, it is first required to register a new platform as defined in the `(guix platform)` module. A platform is making the connection between a GNU triplet (see Section “Specifying Target Triplets” in *Autoconf*), the equivalent *system* in Nix notation, the name of the *glibc-dynamic-linker*, and the corresponding Linux architecture name if applicable (see Chapter 15 [Platforms], page 791).

Once the bootstrap tarball are built, the `(gnu packages bootstrap)` module needs to be updated to refer to these binaries on the target platform. That is, the hashes and URLs of the bootstrap tarballs for the new platform must be added alongside those of the currently supported platforms. The bootstrap Guile tarball is treated specially: it is expected to be available locally, and `gnu/local.mk` has rules to download it for the supported architectures; a rule for the new platform must be added as well.

In practice, there may be some complications. First, it may be that the extended GNU triplet that specifies an ABI (like the `eabi` suffix above) is not recognized by all the GNU tools. Typically, `glibc` recognizes some of these, whereas `GCC` uses an extra `--with-abi` configure flag (see `gcc.scm` for examples of how to handle this). Second, some of the required packages could fail to build for that platform. Lastly, the generated binaries could be broken for some reason.

## 22 Contributing

This project is a cooperative effort, and we need your help to make it grow! Please get in touch with us on [guix-devel@gnu.org](mailto:guix-devel@gnu.org) and [#guix](#) on the Libera Chat IRC network. We welcome ideas, bug reports, patches, and anything that may be helpful to the project. We particularly welcome help on packaging (see Section 22.8 [Packaging Guidelines], page 830).

We want to provide a warm, friendly, and harassment-free environment, so that anyone can contribute to the best of their abilities. To this end our project uses a “Contributor Covenant”, which was adapted from <https://contributor-covenant.org/>. You can find a local version in the `CODE-OF-CONDUCT` file in the source tree.

Contributors are not required to use their legal name in patches and on-line communication; they can use any name or pseudonym of their choice.

### 22.1 Requirements

You can easily hack on Guix itself using Guix and Git, which we use for version control (see Section 22.2 [Building from Git], page 817).

But when packaging Guix for foreign distros or when bootstrapping on systems without Guix, and if you decide to not just trust and install our readily made binary (see Section 2.1 [Binary Installation], page 4), you can download a release version of our reproducible source tarball and read on.

This section lists requirements when building Guix from source. The build procedure for Guix is the same as for other GNU software, and is not covered here. Please see the files `README` and `INSTALL` in the Guix source tree for additional details.

GNU Guix is available for download from its website at <https://guix.gnu.org>.

GNU Guix depends on the following packages:

- GNU Guile (<https://gnu.org/software/guile/>), version 3.0.x, version 3.0.9 or later;
- Guile-Gcrypt (<https://notabug.org/cwebber/guile-gcrypt>), version 0.1.0 or later;
- Guile-GnuTLS (<https://gitlab.com/gnutls/guile/>) (see Section “Guile Preparations” in *GnuTLS-Guile*)<sup>1</sup>;
- Guile-SQLite3 (<https://notabug.org/guile-sqlite3/guile-sqlite3>), version 0.1.0 or later;
- Guile-zlib (<https://notabug.org/guile-zlib/guile-zlib>), version 0.1.0 or later;
- Guile-lzlib (<https://notabug.org/guile-lzlib/guile-lzlib>);
- Guile-Avahi (<https://www.nongnu.org/guile-avahi/>);
- Guile-Git (<https://gitlab.com/guile-git/guile-git>), version 0.10.0 or later;
- Git (<https://git-scm.com>) (yes, both!);
- Guile-JSON (<https://github.com/aconchillo/guile-json>) 4.3.0 or later;
- GNU Make (<https://www.gnu.org/software/make/>).

---

<sup>1</sup> The Guile bindings to GnuTLS (<https://gnutls.org/>) were distributed as part of GnuTLS until version 3.7.8 included.

The following dependencies are optional:

- Support for build offloading (see Section 2.2.2 [Daemon Offload Setup], page 9) and `guix copy` (see Section 9.13 [Invoking `guix copy`], page 233) depends on Guile-SSH (<https://github.com/artiom-poptsov/guile-ssh>), version 0.13.0 or later.
- Guile-zstd (<https://notabug.org/guile-zstd/guile-zstd>), for zstd compression and decompression in `guix publish` and for substitutes (see Section 9.11 [Invoking `guix publish`], page 227).
- Guile-Semver (<https://ngyro.com/software/guile-semver.html>) for the `crate` importer (see Section 9.5 [Invoking `guix import`], page 197).
- Guile-Lib (<https://www.nongnu.org/guile-lib/doc/ref/htmlprag/>) for the `go` importer (see Section 9.5 [Invoking `guix import`], page 197) and for some of the “updaters” (see Section 9.6 [Invoking `guix refresh`], page 206).
- When `libbz2` (<http://www.bzip.org>) is available, `guix-daemon` can use it to compress build logs.

Unless `--disable-daemon` was passed to `configure`, the following packages are also needed:

- GNU `libcrypt` (<https://gnupg.org/>);
- SQLite 3 (<https://sqlite.org>);
- GCC’s `g++` (<https://gcc.gnu.org>), with support for the C++11 standard.

## 22.2 Building from Git

If you want to hack Guix itself, it is recommended to use the latest version from the Git repository<sup>2</sup>:

```
git clone https://git.guix.gnu.org/guix.git
cd guix
```

How do you ensure that you obtained a genuine copy of the repository? To do that, run `guix git authenticate`, passing it the commit and OpenPGP fingerprint of the *channel introduction* (see Section 7.5 [Invoking `guix git authenticate`], page 98):

```
git fetch origin keyring:keyring
guix git authenticate 9edb3f66fd807b096b48283debdcdccfea34bad \
  "BBB0 2DDF 2CEA F6A8 0D1D E643 A2A0 6DF2 A33A 54FA"
```

This command completes with exit code zero on success; it prints an error message and exits with a non-zero code otherwise.

As you can see, there is a chicken-and-egg problem: you first need to have Guix installed. Typically you would install Guix System (see Chapter 3 [System Installation], page 23) or Guix on top of another distro (see Section 2.1 [Binary Installation], page 4); in either case, you would verify the OpenPGP signature on the installation medium. This “bootstraps” the trust chain.

---

<sup>2</sup> Guix uses `master` as the default branch name. A proposal in 2025 to change the default branch name to `main` did not reach consensus. See Guix Consensus Document 003 (<https://codeberg.org/guix/guix-consensus-documents/src/branch/main/003-rename-default-branch.md>).

The easiest way to set up a development environment for Guix is, of course, by using Guix! The following command starts a new shell where all the dependencies and appropriate environment variables are set up to hack on Guix:

```
guix shell -D guix -CPW
```

or even, from within a Git worktree for Guix:

```
echo $(pwd) >> ~/.config/guix/shell-authorized-directories
guix shell -CPW
```

If `-C` (short for `--container`) is not supported on your system, try `--pure` instead of `-CPW`. See Section 7.1 [Invoking guix shell], page 78, for more information on that command.

Most of the test suite runs without an internet connection. The `-N` (short for `--network`) can be added to the `guix shell` command to run the tests that require internet access.

If you are unable to use Guix when building Guix from a checkout, the following are the required packages in addition to those mentioned in the installation instructions (see Section 22.1 [Requirements], page 816).

- GNU Autoconf (<https://gnu.org/software/autoconf/>);
- GNU Automake (<https://gnu.org/software/automake/>);
- GNU Gettext (<https://gnu.org/software/gettext/>);
- GNU Texinfo (<https://gnu.org/software/texinfo/>);
- Graphviz (<https://www.graphviz.org/>);
- GNU Help2man (optional) (<https://www.gnu.org/software/help2man/>).

On Guix, other development tools, like `strace`, can be included by giving them as extra arguments to `guix shell`:

```
guix shell -D guix -CPW strace
```

From there you can generate the build system infrastructure using Autoconf and Automake:

```
./bootstrap
```

If you get an error like this one:

```
configure.ac:46: error: possibly undefined macro: PKG_CHECK_MODULES
```

it probably means that Autoconf couldn't find `pkg.m4`, which is provided by `pkg-config`. Make sure that `pkg.m4` is available. The same holds for the `guile.m4` set of macros provided by Guile. For instance, if you installed Automake in `/usr/local`, it wouldn't look for `.m4` files in `/usr/share`. In that case, you have to invoke the following command:

```
export ACLOCAL_PATH=/usr/share/aclocal
```

See Section “Macro Search Path” in *The GNU Automake Manual*, for more information.

Then, run:

```
./configure
```

Optionally, `--localstatedir` and `--sysconfdir` can also be provided as arguments. By default, `localstatedir` is `/var` (see Section 8.9 [The Store], page 156, for information about this) and `sysconfdir` is `/etc`. Note that you will probably not run `make install` at the end (you don't have to) but it's still important to pass the right `localstatedir` and `sysconfdir` values, which get recorded in the `(guix config)` Guile module.

Finally, you can build Guix and, if you feel so inclined, run the tests (see Section 22.3 [Running the Test Suite], page 819):

```
make
make check
```

If anything fails, take a look at installation instructions (see Chapter 2 [Installation], page 4) or send a message to the mailing list.

From there on, you can authenticate all the commits included in your checkout by running:

```
guix git authenticate \
  9edb3f66fd807b096b48283debdcdccfea34bad \
  "BBB0 2DDF 2CEA F6A8 0D1D E643 A2A0 6DF2 A33A 54FA"
```

The first run takes a couple of minutes, but subsequent runs are faster. On subsequent runs, you can run the command without any arguments since the *introduction* (the commit ID and OpenPGP fingerprints above) will have been recorded<sup>3</sup>:

```
guix git authenticate
```

When your configuration for your local Git repository doesn't match the default one, you can provide the reference for the `keyring` branch *via* the `-k` option. The following example assumes that you have a Git remote called `'myremote'` pointing to the official repository:

```
guix git authenticate \
  -k myremote/keyring \
  9edb3f66fd807b096b48283debdcdccfea34bad \
  "BBB0 2DDF 2CEA F6A8 0D1D E643 A2A0 6DF2 A33A 54FA"
```

See Section 7.5 [Invoking `guix git authenticate`], page 98, for more information on this command.

**Note:** By default, hooks are installed such that `guix git authenticate` is invoked anytime you run `git pull` or `git push`.

After updating the repository, `make` might fail with an error similar to the following example:

```
error: failed to load 'gnu/packages/linux.scm':
ice-9/eval.scm:293:34: In procedure abi-check: #<record-type <origin>>: record ABI mis
```

This means that one of the record types that Guix defines (in this example, the `origin` record) has changed, and all of guix needs to be recompiled to take that change into account. To do so, run `make clean-go` followed by `make`.

Should `make` fail with an Automake error message after updating, you need to repeat the steps outlined in this section, commencing with `./bootstrap`.

## 22.3 Running the Test Suite

After a successful `configure` and `make` run, it is a good idea to run the test suite. It can help catch issues with the setup or environment, or bugs in Guix itself—and really, reporting test failures is a good way to help improve the software. To run the test suite, type:

```
make check
```

---

<sup>3</sup> This requires a recent version of Guix, from May 2024 or more recent.

Test cases can run in parallel: you can use the `-j` option of GNU make to speed things up. The first run may take a few minutes on a recent machine; subsequent runs will be faster because the store that is created for test purposes will already have various things in cache.

It is also possible to run a subset of the tests by defining the `TESTS` makefile variable as in this example:

```
make check TESTS="tests/store.scm tests/cpio.scm"
```

By default, tests results are displayed at a file level. In order to see the details of every individual test cases, it is possible to define the `SCM_LOG_DRIVER_FLAGS` makefile variable as in this example:

```
make check TESTS="tests/base64.scm" SCM_LOG_DRIVER_FLAGS="--brief=no"
```

The underlying SRFI 64 custom Automake test driver used for the 'check' test suite (located at `build-aux/test-driver.scm`) also allows selecting which test cases to run at a finer level, via its `--select` and `--exclude` options. Here's an example, to run all the test cases from the `tests/packages.scm` test file whose names start with "transaction-upgrade-entry":

```
export SCM_LOG_DRIVER_FLAGS="--select=~transaction-upgrade-entry"
make check TESTS="tests/packages.scm"
```

Those wishing to inspect the results of failed tests directly from the command line can add the `--errors-only=yes` option to the `SCM_LOG_DRIVER_FLAGS` makefile variable and set the `VERBOSE` Automake makefile variable, as in:

```
make check SCM_LOG_DRIVER_FLAGS="--brief=no --errors-only=yes" VERBOSE=1
```

The `--show-duration=yes` option can be used to print the duration of the individual test cases, when used in combination with `--brief=no`:

```
make check SCM_LOG_DRIVER_FLAGS="--brief=no --show-duration=yes"
```

See Section "Parallel Test Harness" in *GNU Automake* for more information about the Automake Parallel Test Harness.

Upon failure, please email [bug-guix@gnu.org](mailto:bug-guix@gnu.org) and attach the `test-suite.log` file. Please specify the Guix version being used as well as version numbers of the dependencies (see Section 22.1 [Requirements], page 816) in your message.

Guix also comes with a whole-system test suite that tests complete Guix System instances. It can only run on systems where Guix is already installed, using:

```
make check-system
```

or, again, by defining `TESTS` to select a subset of tests to run:

```
make check-system TESTS="basic mcron"
```

These system tests are defined in the `(gnu tests ...)` modules. They work by running the operating systems under test with lightweight instrumentation in a virtual machine (VM). They can be computationally intensive or rather cheap, depending on whether substitutes are available for their dependencies (see Section 5.3 [Substitutes], page 46). Some of them require a lot of storage space to hold VM images.

If you encounter an error like:

```
Compiling Scheme modules...
ice-9/eval.scm:142:16: In procedure compile-top-call:
```



```
error: all-system-tests: unbound variable
hint: Did you forget `(use-modules (gnu tests))'?
```

there may be inconsistencies in the work tree from previous builds. To resolve this, try running `make clean-go` followed by `make`.

Again in case of test failures, please send [bug-guix@gnu.org](mailto:bug-guix@gnu.org) all the details.

## 22.4 Running Guix Before It Is Installed

In order to keep a sane working environment, you will find it useful to test the changes made in your local source tree checkout without actually installing them. So that you can distinguish between your “end-user” hat and your “motley” costume.

To that end, all the command-line tools can be used even if you have not run `make install`. To do that, you first need to have an environment with all the dependencies available (see Section 22.2 [Building from Git], page 817), and then simply prefix each command with `./pre-inst-env` (the `pre-inst-env` script lives in the top build tree of Guix; see Section 22.2 [Building from Git], page 817, to generate it). As an example, here is how you would build the `hello` package as defined in your working tree (this assumes `guix-daemon` is already running on your system; it’s OK if it’s a different version):

```
$ ./pre-inst-env guix build hello
```

Similarly, an example for a Guile session using the Guix modules:

```
$ ./pre-inst-env guile -c '(use-modules (guix utils)) (pk (%current-system))'■
```

```
;;; ("x86_64-linux")
```

... and for a REPL (see Section 8.14 [Using Guix Interactively], page 176):

```
$ ./pre-inst-env guile
scheme@(guile-user)> ,use(guix)
scheme@(guile-user)> ,use(gnu)
scheme@(guile-user)> (define snakes
  (fold-packages
    (lambda (package lst)
      (if (string-prefix? "python"
                          (package-name package))
          (cons package lst)
          lst))
    '()))
scheme@(guile-user)> (length snakes)
$1 = 361
```

If you are hacking on the daemon and its supporting code or if `guix-daemon` is not already running on your system, you can launch it straight from the build tree<sup>4</sup>:

```
$ sudo -E ./pre-inst-env guix-daemon --build-users-group=guixbuild
```

The `pre-inst-env` script sets up all the environment variables necessary to support this, including `PATH` and `GUILE_LOAD_PATH`.

<sup>4</sup> The `-E` flag to `sudo` guarantees that `GUILE_LOAD_PATH` is correctly set such that `guix-daemon` and the tools it uses can find the Guile modules they need.

Note that `./pre-inst-env guix pull` does *not* upgrade the local source tree; it simply updates the `~/.config/guix/current` symlink (see Section 5.7 [Invoking `guix pull`], page 56). Run `git pull` instead if you want to upgrade your local source tree.

Sometimes, especially if you have recently updated your repository, running `./pre-inst-env` will print a message similar to the following example:

```
;;; note: source file /home/user/projects/guix/guix/progress.scm
;;;      newer than compiled /home/user/projects/guix/guix/progress.go
```

This is only a note and you can safely ignore it. You can get rid of the message by running `make -j4`. Until you do, Guile will run slightly slower because it will interpret the code instead of using prepared Guile object (`.go`) files.

You can run `make` automatically as you work using `watchexec` from the `watchexec` package. For example, to build again each time you update a package file, run `'watchexec -w gnu/packages -- make -j4'`.

## 22.5 The Perfect Setup

The Perfect Setup to hack on Guix is basically the perfect setup used for Guile hacking (see Section “Using Guile in Emacs” in *Guile Reference Manual*). First, you need more than an editor, you need Emacs (<https://www.gnu.org/software/emacs/>), empowered by the wonderful Geiser (<https://nongnu.org/geiser/>). To set that up, run:

```
guix install emacs guile emacs-geiser emacs-geiser-guile
```

Geiser allows for interactive and incremental development from within Emacs: code compilation and evaluation from within buffers, access to on-line documentation (docstrings), context-sensitive completion, `M-.`  to jump to an object definition, a REPL to try out your code, and more (see Section “Introduction” in *Geiser User Manual*). If you allow Emacs to load the `.dir-locals.el` file at the root of the project checkout, it will cause Geiser to automatically add the local Guix sources to the Guile load path.

To actually edit the code, Emacs already has a neat Scheme mode. But in addition to that, you must not miss Paredit (<https://www.emacswiki.org/emacs/ParEdit>). It provides facilities to directly operate on the syntax tree, such as raising an s-expression or wrapping it, swallowing or rejecting the following s-expression, etc.

We also provide templates for common git commit messages and package definitions in the `etc/snippets` directory. These templates can be used to expand short trigger strings to interactive text snippets. If you use YASnippet (<https://joaotavora.github.io/yasnipet/>), you may want to add the `etc/snippets/yas` snippets directory to the `yas-snippet-dirs` variable. If you use Tempel (<https://github.com/minad/tempel/>), you may want to add the `etc/snippets/tempel/*` path to the `tempel-path` variable in Emacs.

```
;; Assuming the Guix checkout is in ~/src/guix.
;; Yasnippet configuration
(with-eval-after-load 'yasnipet
  (add-to-list 'yas-snippet-dirs "~/src/guix/etc/snippets/yas"))
;; Tempel configuration
(with-eval-after-load 'tempel
  ;; Ensure tempel-path is a list -- it may also be a string.
  (unless (listp 'tempel-path)
```

```
(setq tempel-path (list tempel-path))
(add-to-list 'tempel-path "~/src/guix/etc/snippets/tempel/*"))
```

The commit message snippets depend on Magit (<https://magit.vc/>) to display staged files. When editing a commit message type `add` followed by *TAB* to insert a commit message template for adding a package; type `update` followed by *TAB* to insert a template for updating a package; type `https` followed by *TAB* to insert a template for changing the home page URI of a package to HTTPS.

The main snippet for `scheme-mode` is triggered by typing `package...` followed by *TAB*. This snippet also inserts the trigger string `origin...`, which can be expanded further. The `origin` snippet in turn may insert other trigger strings ending on `...`, which also can be expanded further.

We additionally provide insertion and automatic update of a copyright in `etc/copyright.el`. You may want to set your full name, mail, and load a file.

```
(setq user-full-name "Alice Doe")
(setq user-mail-address "alice@mail.org")
;; Assuming the Guix checkout is in ~/src/guix.
(load-file "~/src/guix/etc/copyright.el")
```

To insert a copyright at the current line invoke `M-x guix-copyright`.

To update a copyright you need to specify a `copyright-names-regexp`.

```
(setq copyright-names-regexp
  (format "%s <%s>" user-full-name user-mail-address))
```

You can check if your copyright is up to date by evaluating `M-x copyright-update`. If you want to do it automatically after each buffer save then add `(add-hook 'after-save-hook 'copyright-update)` in Emacs.

### 22.5.1 Viewing Bugs within Emacs

Emacs has a nice minor mode called `bug-reference`, which, when combined with `'emacs-debbugs'` (the Emacs package), can be used to open links such as `'<https://bugs.gnu.org/58697>'` or `'<https://issues.guix.gnu.org/58697>'` as bug report buffers. From there you can easily consult the email thread via the Gnus interface, reply or modify the bug status, all without leaving the comfort of Emacs! Below is a sample configuration to add to your `~/.emacs` configuration file:

```
;;; Bug references.
(require 'bug-reference)
(add-hook 'prog-mode-hook #'bug-reference-prog-mode)
(add-hook 'gnus-mode-hook #'bug-reference-mode)
(add-hook 'erc-mode-hook #'bug-reference-mode)
(add-hook 'gnus-summary-mode-hook #'bug-reference-mode)
(add-hook 'gnus-article-mode-hook #'bug-reference-mode)

;;; This extends the default expression (the top-most, first expression
;;; provided to 'or') to also match URLs such as
;;; <https://issues.guix.gnu.org/58697> or <https://bugs.gnu.org/58697>.
;;; It is also extended to detect "Fixes: #NNNNN" git trailers.
(setq bug-reference-bug-regexp
```

```

(rx (group (or (seq word-boundary
  (or (seq (char "Bb") "ug"
    (zero-or-one " ")
    (zero-or-one "#"))
    (seq (char "Pp") "atch"
      (zero-or-one " ")
      "#")
    (seq (char "Ff") "ixes"
      (zero-or-one ":")
      (zero-or-one " ") "#")
    (seq "RFE"
      (zero-or-one " ") "#")
    (seq "PR "
      (one-or-more (char "a-z+-")) "/"))
    (group (one-or-more (char "0-9"))
      (zero-or-one
        (seq "#" (one-or-more
          (char "0-9"))))))
    (seq (? "<") "https://bugs.gnu.org/"
      (group-n 2 (one-or-more (char "0-9")))
      (? ">"))
    (seq (? "<") "https://issues.gnu.org/"
      (? "issue/")
      (group-n 2 (one-or-more (char "0-9")))
      (? ">")))))
  (setq bug-reference-url-format "https://issues.gnu.org/%s")

(require 'debbugs)
(require 'debbugs-browse)
(add-hook 'bug-reference-mode-hook #'debbugs-browse-mode)
(add-hook 'bug-reference-prog-mode-hook #'debbugs-browse-mode)

;; The following allows Emacs Debbugs user to open the issue directly within
;; Emacs.
(setq debbugs-browse-url-regexp
  (rx line-start
    "http" (zero-or-one "s") "://"
    (or "debbugs" "issues.gnu" "bugs")
    ".gnu.org" (one-or-more "/")
    (group (zero-or-one "cgi/bugreport.cgi?bug="))
    (group-n 3 (one-or-more digit))
    line-end))

;; Change the default when run as 'M-x debbugs-gnu'.
(setq debbugs-gnu-default-packages '("guix" "guix-patches"))

;; Show feature requests.

```

```
(setq debugs-gnu-default-severities
  '("serious" "important" "normal" "minor" "wishlist"))
```

For more information, refer to Section “Bug Reference” in *The GNU Emacs Manual* and Section “Minor Mode” in *The Debugs User Guide*.

## 22.6 Alternative Setups

Alternative setups than Emacs may let you work on Guix with a similar development experience and they might work better with the tools you currently use or help you make the transition to Emacs.

The options listed below only provide the alternatives to the Emacs based setup, which is the most widely used in the Guix community. If you want to really understand how the perfect setup for Guix development is supposed to work, we encourage you to read the section before this regardless of the editor you choose to use.

### 22.6.1 Guile Studio

Guile Studio is a pre-configured Emacs with mostly everything you need to start hacking in Guile. If you are not familiar with Emacs it makes the transition easier for you.

```
guix install guile-studio
```

Guile Studio comes with Geiser preinstalled and prepared for action.

### 22.6.2 Vim and NeoVim

Vim (and NeoVim) are also packaged in Guix, just in case you decided to go for the evil path.

```
guix install vim
```

If you want to enjoy a similar development experience to that in the perfect setup, you should install several plugins to configure the editor. Vim (and NeoVim) have the equivalent to Paredit, `paredit.vim` ([https://www.vim.org/scripts/script.php?script\\_id=3998](https://www.vim.org/scripts/script.php?script_id=3998)), that will help you with the structural editing of Scheme files (the support for very large files is not great, though).

```
guix install vim-paredit
```

We also recommend that you run `:set autoindent` so that your code is automatically indented as you type.

For the interaction with Git, `fugitive.vim` ([https://www.vim.org/scripts/script.php?script\\_id=2975](https://www.vim.org/scripts/script.php?script_id=2975)) is the most commonly used plugin:

```
guix install vim-fugitive
```

And of course if you want to interact with Guix directly from inside of vim, using the built-in terminal emulator, we have our very own `guix.vim` package!

```
guix install vim-guix-vim
```

In NeoVim you can even make a similar setup to Geiser using Conjure (<https://conjure.fun/>) that lets you connect to a running Guile process and inject your code there live (sadly it’s not packaged in Guix yet).

## 22.7 Source Tree Structure

If you're willing to contribute to Guix beyond packages, or if you'd like to learn how it all fits together, this section provides a guided tour in the code base that you may find useful.

Overall, the Guix source tree contains almost exclusively Guile *modules*, each of which can be seen as an independent library (see Section “Modules” in *GNU Guile Reference Manual*).

The following table gives an overview of the main directories and what they contain. Remember that in Guile, each module name is derived from its file name—e.g., the module in file `guix/packages.scm` is called `(guix packages)`.

**guix** This is the location of core Guix mechanisms. To illustrate what is meant by “core”, here are a few examples, starting from low-level tools and going towards higher-level tools:

`(guix store)`

Connecting to and interacting with the build daemon (see Section 8.9 [The Store], page 156).

`(guix derivations)`

Creating derivations (see Section 8.10 [Derivations], page 158).

`(guix gexps)`

Writing G-expressions (see Section 8.12 [G-Expressions], page 165).

`(guix packages)`

Defining packages and origins (see Section 8.2.1 [package Reference], page 104).

`(guix download)`

`(guix git-download)`

The `url-fetch` and `git-fetch` origin download methods (see Section 8.2.2 [origin Reference], page 109).

`(guix swb)`

Fetching source code from the Software Heritage archive (<https://archive.softwareheritage.org>).

`(guix search-paths)`

Implementing search paths (see Section 8.8 [Search Paths], page 153).

`(guix build-system)`

The build system interface (see Section 8.5 [Build Systems], page 121).

`(guix profiles)`

Implementing profiles.

**guix/build-system**

This directory contains specific build system implementations (see Section 8.5 [Build Systems], page 121), such as:

`(guix build-system gnu)`

the GNU build system;

```
(guix build-system cmake)
    the CMake build system;

(guix build-system pyproject)
    The Python “pyproject” build system.
```

**guix/build**

This contains code generally used on the “build side” (see Section 8.12 [G-Expressions], page 165). This includes code used to build packages or other operating system components, as well as utilities:

```
(guix build utils)
    Utilities for package definitions and more (see Section 8.7 [Build
    Utilities], page 146).

(guix build gnu-build-system)
(guix build cmake-build-system)
(guix build pyproject-build-system)
    Implementation of build systems, and in particular definition of
    their build phases (see Section 8.6 [Build Phases], page 143).

(guix build syscalls)
    Interface to the C library and to Linux system calls.
```

**guix/scripts**

This contains modules corresponding to **guix** sub-commands. For example, the `(guix scripts shell)` module exports the `guix-shell` procedure, which directly corresponds to the `guix shell` command (see Section 7.1 [Invoking guix shell], page 78).

**guix/import**

This contains supporting code for the importers and updaters (see Section 9.5 [Invoking guix import], page 197, and see Section 9.6 [Invoking guix refresh], page 206). For example, `(guix import pypi)` defines the interface to PyPI, which is used by the `guix import pypi` command.

The directories we have seen so far all live under **guix/**. The other important place is the **gnu/** directory, which contains primarily package definitions as well as libraries and tools for Guix System (see Chapter 11 [System Configuration], page 242) and Guix Home (see Chapter 13 [Home Configuration], page 732), all of which build upon functionality provided by `(guix ...)` modules<sup>5</sup>.

**gnu/packages**

This is by far the most crowded directory of the source tree: it contains *package modules* that export package definitions (see Section 8.1 [Package Modules], page 100). A few examples:

```
(gnu packages base)
    Module providing “base” packages: glibc, coreutils, grep, etc.
```

---

<sup>5</sup> For this reason, `(guix ...)` modules must generally not depend on `(gnu ...)` modules, with notable exceptions: `(guix build-system ...)` modules may look up packages at run time—e.g., `(guix build-system cmake)` needs to access the `cmake` variable at run time—, `(guix scripts ...)` often rely on `(gnu ...)` modules, and the same goes for some of the `(guix import ...)` modules.

(gnu packages guile)

Guile and core Guile packages.

(gnu packages linux)

The Linux-libre kernel and related packages.

(gnu packages python)

Python and core Python packages.

(gnu packages python-xyz)

Miscellaneous Python packages (we were not very creative).

In any case, you can jump to a package definition using `guix edit` (see Section 9.2 [Invoking guix edit], page 195) and view its location with `guix show` (see Section 5.2 [Invoking guix package], page 37).

#### gnu/packages/patches

This directory contains patches applied against packages and obtained using the `search-patches` procedure.

#### gnu/services

This contains service definitions, primarily for Guix System (see Section 11.10 [Services], page 275) but some of them are adapted and reused for Guix Home as we will see below. Examples:

(gnu services)

The service framework itself, which defines the service and service type data types (see Section 11.19.1 [Service Composition], page 709).

(gnu services base)

“Base” services (see Section 11.10.1 [Base Services], page 276).

(gnu services desktop)

“Desktop” services (see Section 11.10.9 [Desktop Services], page 377).

(gnu services shepherd)

Support for Shepherd services (see Section 11.19.4 [Shepherd Services], page 717).

You can jump to a service definition using `guix system edit` and view its location with `guix system search` (see Section 11.16 [Invoking guix system], page 692).

#### gnu/system

These are core Guix System modules, such as:

(gnu system)

Defines `operating-system` (see Section 11.3 [operating-system Reference], page 253).

(gnu system file-systems)

Defines `file-system` (see Section 11.4 [File Systems], page 257).



(gnu system mapped-devices)

Defines `mapped-device` (see Section 11.5 [Mapped Devices], page 263).

#### gnu/build

These are modules that are either used on the “build side” when building operating systems or packages, or at run time by operating systems.

(gnu build accounts)

Creating `/etc/passwd`, `/etc/shadow`, etc. (see Section 11.7 [User Accounts], page 269).

(gnu build activation)

Activating an operating system at boot time or reconfiguration time.

(gnu build file-systems)

Searching, checking, and mounting file systems.

(gnu build linux-boot)

(gnu build hurd-boot)

Booting GNU/Linux and GNU/Hurd operating systems.

(gnu build linux-initrd)

Creating a Linux initial RAM disk (see Section 11.14 [Initial RAM Disk], page 682).

**gnu/home** This contains all things Guix Home (see Chapter 13 [Home Configuration], page 732); examples:

(gnu home services)

Core services such as `home-files-service-type`.

(gnu home services ssh)

SSH-related services (see Section 13.3.6 [Secure Shell], page 749).

#### gnu/installer

This contains the text-mode graphical system installer (see Section 3.5 [Guided Graphical Installation], page 25).

#### gnu/machine

These are the *machine abstractions* used by `guix deploy` (see Section 11.17 [Invoking `guix deploy`], page 701).

#### gnu/tests

This contains system tests—tests that spawn virtual machines to check that system services work as expected (see Section 22.3 [Running the Test Suite], page 819).

Last, there’s also a few directories that contain files that are *not* Guile modules:

**nix** This is the C++ implementation of `guix-daemon`, inherited from Nix (see Section 2.3 [Invoking `guix-daemon`], page 14).

<b>tests</b>	These are unit tests, each file corresponding more or less to one module, in particular ( <code>guix ...</code> ) modules (see Section 22.3 [Running the Test Suite], page 819).
<b>doc</b>	This is the documentation in the form of Texinfo files: this manual and the Cookbook. See Section “Writing a Texinfo File” in <i>GNU Texinfo</i> , for information on Texinfo markup language.
<b>po</b>	This is the location of translations of Guix itself, of package synopses and descriptions, of the manual, and of the cookbook. Note that <code>.po</code> files that live here are pulled directly from Weblate (see Section 22.19 [Translating Guix], page 859).
<b>etc</b>	Miscellaneous files: shell completions, support for <code>systemd</code> and other init systems, Git hooks, etc.

With all this, a fair chunk of your operating system is at your fingertips! Beyond `grep` and `git grep`, see Section 22.5 [The Perfect Setup], page 822, on how to navigate code from your editor, and see Section 8.14 [Using Guix Interactively], page 176, for information on how to use Scheme modules interactively. Enjoy!

## 22.8 Packaging Guidelines

The GNU distribution is nascent and may well lack some of your favorite packages. This section describes how you can help make the distribution grow.

Free software packages are usually distributed in the form of *source code tarballs*—typically `tar.gz` files that contain all the source files. Adding a package to the distribution means essentially two things: adding a *recipe* that describes how to build the package, including a list of other packages required to build it, and adding *package metadata* along with that recipe, such as a description and licensing information.

In Guix all this information is embodied in *package definitions*. Package definitions provide a high-level view of the package. They are written using the syntax of the Scheme programming language; in fact, for each package we define a variable bound to the package definition, and export that variable from a module (see Section 8.1 [Package Modules], page 100). However, in-depth Scheme knowledge is *not* a prerequisite for creating packages. For more information on package definitions, see Section 8.2 [Defining Packages], page 101.

Once a package definition is in place, stored in a file in the Guix source tree, it can be tested using the `guix build` command (see Section 9.1 [Invoking `guix build`], page 180). For example, assuming the new package is called `gnew`, you may run this command from the Guix build tree (see Section 22.4 [Running Guix Before It Is Installed], page 821):

```
./pre-inst-env guix build gnew --keep-failed
```

Using `--keep-failed` makes it easier to debug build failures since it provides access to the failed build tree. Another useful command-line option when debugging is `--log-file`, to access the build log.

If the package is unknown to the `guix` command, it may be that the source file contains a syntax error, or lacks a `define-public` clause to export the package variable. To figure it out, you may load the module from Guile to get more information about the actual error:

```
./pre-inst-env guile -c '(use-modules (gnu packages gnew))'
```

Once your package builds correctly, please send us a patch (see Section 22.10 [Submitting Patches], page 842). Well, if you need help, we will be happy to help you too. Once the patch is committed in the Guix repository, the new package automatically gets built on the supported platforms by our continuous integration system (<https://bordeaux.guix.gnu.org>).

Users can obtain the new package definition simply by running `guix pull` (see Section 5.7 [Invoking guix pull], page 56). When [bordeaux.guix.gnu.org](https://bordeaux.guix.gnu.org) is done building the package, installing the package automatically downloads binaries from there (see Section 5.3 [Substitutes], page 46). The only place where human intervention is needed is to review and apply the patch.

### 22.8.1 Software Freedom

The GNU operating system has been developed so that users can have freedom in their computing. GNU is *free software*, meaning that users have the four essential freedoms (<https://www.gnu.org/philosophy/free-sw.html>): to run the program, to study and change the program in source code form, to redistribute exact copies, and to distribute modified versions. Packages found in the GNU distribution provide only software that conveys these four freedoms.

In addition, the GNU distribution follow the free software distribution guidelines (<https://www.gnu.org/distros/free-system-distribution-guidelines.html>). Among other things, these guidelines reject non-free firmware, recommendations of non-free software, and discuss ways to deal with trademarks and patents.

Some otherwise free upstream package sources contain a small and optional subset that violates the above guidelines, for instance because this subset is itself non-free code. When that happens, the offending items are removed with appropriate patches or code snippets in the `origin` form of the package (see Section 8.2 [Defining Packages], page 101). This way, `guix build --source` returns the “freed” source rather than the unmodified upstream source.

### 22.8.2 Package Naming

A package actually has two names associated with it. First, there is the name of the *Scheme variable*, the one following `define-public`. By this name, the package can be made known in the Scheme code, for instance as input to another package. Second, there is the string in the `name` field of a package definition. This name is used by package management commands such as `guix package` and `guix build`.

Both are usually the same and correspond to the lowercase conversion of the project name chosen upstream, with underscores replaced with hyphens. For instance, GNUnet is available as `gnunet`, and SDL\_net as `sdl-net`.

A noteworthy exception to this rule is when the project name is only a single character, or if an older maintained project with the same name already exists—regardless of whether it has already been packaged for Guix. Use common sense to make such names unambiguous and meaningful. For example, Guix’s package for the shell called “s” upstream is `s-shell` and *not* `s`. Feel free to ask your fellow hackers for inspiration.

We do not add `lib` prefixes for library packages, unless these are already part of the official project name. But see Section 22.8.9 [Python Modules], page 836, and Section 22.8.10

[Perl Modules], page 837, for special rules concerning modules for the Python and Perl languages.

Font package names are handled differently, see Section 22.8.14 [Fonts], page 840.

### 22.8.3 Version Numbers

We usually package only the latest version of a given free software project. But sometimes, for instance for incompatible library versions, two (or more) versions of the same package are needed. These require different Scheme variable names. We use the name as defined in Section 22.8.2 [Package Naming], page 831, for the most recent version; previous versions use the same name, suffixed by `-` and the smallest prefix of the version number that may distinguish the two versions.

The name inside the package definition is the same for all versions of a package and does not contain any version number.

For instance, the versions 2.24.20 and 3.9.12 of GTK+ may be packaged as follows:

```
(define-public gtk+
  (package
    (name "gtk+")
    (version "3.9.12")
    ...))
(define-public gtk+-2
  (package
    (name "gtk+")
    (version "2.24.20")
    ...))
```

If we also wanted GTK+ 3.8.2, this would be packaged as

```
(define-public gtk+-3.8
  (package
    (name "gtk+")
    (version "3.8.2")
    ...))
```

Occasionally, we package snapshots of upstream’s version control system (VCS) instead of formal releases. This should remain exceptional, because it is up to upstream developers to clarify what the stable release is. Yet, it is sometimes necessary. So, what should we put in the `version` field?

Clearly, we need to make the commit identifier of the VCS snapshot visible in the version string, but we also need to make sure that the version string is monotonically increasing so that `guix package --upgrade` can determine which version is newer. Since commit identifiers, notably with Git, are not monotonically increasing, we add a revision number that we increase each time we upgrade to a newer snapshot. The resulting version string looks like this:

```
2.0.11-3.cabba9e
^      ^      ^
|      |      `-- upstream commit ID
|      |
|      `--- Guix package revision
```

```
|
latest upstream version
```

It is a good idea to strip commit identifiers in the `version` field to, say, 7 digits. It avoids an aesthetic annoyance (assuming aesthetics have a role to play here) as well as problems related to OS limits such as the maximum shebang length (127 bytes for the Linux kernel). There are helper functions for doing this for packages using `git-fetch` or `hg-fetch` (see below). It is best to use the full commit identifiers in `origins`, though, to avoid ambiguities. A typical package definition may look like this:

```
(define my-package
  (let ((commit "c3f29bc928d5900971f65965feaae59e1272a3f7")
        (revision "1"))
    ;Guix package revision
    (package
      (version (git-version "0.9" revision commit))
      (source (origin
                (method git-fetch)
                (uri (git-reference
                     (url "git://example.org/my-package.git")
                     (commit commit)))
                (sha256 (base32 "1mbikn..."))
                (file-name (git-file-name name version))))
      ;; ...
    )))
```

**git-version** *VERSION REVISION COMMIT* [Procedure]

Return the version string for packages using `git-fetch`.

```
(git-version "0.2.3" "0" "93818c936ee7e2f1ba1b315578bde363a7d43d05")
⇒ "0.2.3-0.93818c9"
```

**hg-version** *VERSION REVISION CHANGESET* [Procedure]

Return the version string for packages using `hg-fetch`. It works in the same way as `git-version`.

## 22.8.4 Synopses and Descriptions

As we have seen before, each package in GNU Guix includes a synopsis and a description (see Section 8.2 [Defining Packages], page 101). Synopses and descriptions are important: They are what `guix package --search` searches, and a crucial piece of information to help users determine whether a given package suits their needs. Consequently, packagers should pay attention to what goes into them.

Synopses must start with a capital letter and must not end with a period. They must not start with “a” or “the”, which usually does not bring anything; for instance, prefer “File-frobbing tool” over “A tool that frobs files”. The synopsis should say what the package is—e.g., “Core GNU utilities (file, text, shell)” —or what it is used for—e.g., the synopsis for GNU `grep` is “Print lines matching a pattern”.

Keep in mind that the synopsis must be meaningful for a very wide audience. For example, “Manipulate alignments in the SAM format” might make sense for a seasoned bioinformatics researcher, but might be fairly unhelpful or even misleading to a non-specialized

audience. It is a good idea to come up with a synopsis that gives an idea of the application domain of the package. In this example, this might give something like “Manipulate nucleotide sequence alignments”, which hopefully gives the user a better idea of whether this is what they are looking for.

Descriptions should take between five and ten lines. Use full sentences, and avoid using acronyms without first introducing them. Please avoid marketing phrases such as “world-leading”, “industrial-strength”, and “next-generation”, and avoid superlatives like “the most advanced”—they are not helpful to users looking for a package and may even sound suspicious. Instead, try to be factual, mentioning use cases and features.

Descriptions can include Texinfo markup, which is useful to introduce ornaments such as `@code` or `@dfn`, bullet lists, or hyperlinks (see Section “Overview” in *GNU Texinfo*). However you should be careful when using some characters for example ‘@’ and curly braces which are the basic special characters in Texinfo (see Section “Special Characters” in *GNU Texinfo*). User interfaces such as `guix show` take care of rendering it appropriately.

Synopses and descriptions are translated by volunteers at Weblate (<https://translate.codeberg.org/projects/guix/packages>) so that as many users as possible can read them in their native language. User interfaces search them and display them in the language specified by the current locale.

To allow `xgettext` to extract them as translatable strings, synopses and descriptions *must be literal strings*. This means that you cannot use `string-append` or `format` to construct these strings:

```
(package
  ;; ...
  (synopsis "This is translatable")
  (description (string-append "This is " "*not*" " translatable.")))
```

Translation is a lot of work so, as a packager, please pay even more attention to your synopses and descriptions as every change may entail additional work for translators. In order to help them, it is possible to make recommendations or instructions visible to them by inserting special comments like this (see Section “xgettext Invocation” in *GNU Gettext*):

```
;; TRANSLATORS: "X11 resize-and-rotate" should not be translated.
(description "ARandR is designed to provide a simple visual front end
for the X11 resize-and-rotate (RandR) extension. ...")
```

### 22.8.5 Snippets versus Phases

The boundary between using an origin snippet versus a build phase to modify the sources of a package can be elusive. Origin snippets are typically used to remove unwanted files such as bundled libraries, nonfree sources, or to apply simple substitutions. The source derived from an origin should produce a source that can be used to build the package on any system that the upstream package supports (i.e., act as the corresponding source). In particular, origin snippets must not embed store items in the sources; such patching should rather be done using build phases. Refer to the `origin` record documentation for more information (see Section 8.2.2 [origin Reference], page 109).

### 22.8.6 Cyclic Module Dependencies

While there cannot be circular dependencies between packages, Guile’s lax module loading mechanism allows circular dependencies between Guile modules, which doesn’t cause problems as long as the following conditions are followed for two modules part of a dependency cycle:

1. Macros are not shared between the co-dependent modules
2. Top-level variables are only referenced in delayed (*thunked*) package fields: `arguments`, `native-inputs`, `inputs`, `propagated-inputs` or `replacement`
3. Procedures referencing top-level variables from another module are not called at the top level of a module themselves.

Straying away from the above rules may work while there are no dependency cycles between modules, but given such cycles are confusing and difficult to troubleshoot, it is best to follow the rules to avoid introducing problems down the line.

Here is a common trap to avoid:

```
(define-public avr-binutils
  (package
    (inherit (cross-binutils "avr"))
    (name "avr-binutils")))
```

In the above example, the `avr-binutils` package was defined in the module `(gnu packages avr)`, and the `cross-binutils` procedure in `(gnu packages cross-base)`. Because the `inherit` field is not delayed (thunked), it is evaluated at the top level at load time, which is problematic in the presence of module dependency cycles. This could be resolved by turning the package into a procedure instead, like:

```
(define (make-avr-binutils)
  (package
    (inherit (cross-binutils "avr"))
    (name "avr-binutils")))
```

Care would need to be taken to ensure the above procedure is only ever used in a package delayed fields or within another procedure also not called at the top level.

### 22.8.7 Bulk Updates

Sometimes you may need to update a monotonic collection of grouped packages, in a particular module—e.g., `(gnu packages astronomy)`. Doing so requires refreshing package versions, creating commits following guidelines and testing them. Let’s say there are 200+ packages with fresh versions, you may deal with this task by applying the following technique.

Enter a development shell in your local Guix checkout, assuming you already ran `bootstrap` and `configure`:

```
~$ guix shell --development guix
[env: /gnu/store/...-profile]
> ./pre-inst-env guix refresh --select=module:astronomy --update
> ./pre-inst-env etc/committer.scm
```

The first command updates all packages in `(gnu packages astronomy)` to the latest versions and writes changes back to the file. The second invokes `etc/committer.scm` script, a

helper script for larger scale updates; it identifies each package diff and creates corresponding commits with the correct style.

If your machine is powerful enough, you may check that the whole collection has no regressions and can still be built successfully:

```
[env: /gnu/store/...-profile]
> ./pre-inst-env guix build $(./pre-inst-env guix package -A | awk '/astronomy.scm/{pr
```

### 22.8.8 Emacs Packages

Emacs packages should preferably use the Emacs build system (see [emacs-build-system], page 140), for uniformity and the benefits provided by its build phases, such as the auto-generation of the autoloads file and the byte compilation of the sources. Because there is no standardized way to run a test suite for Emacs packages, tests are disabled by default. When a test suite is available, it should be enabled by setting the `#:tests?` argument to `#true`. By default, the command to run the test is `make check`, but any command can be specified via the `#:test-command` argument. The `#:test-command` argument expects a list containing a command and its arguments, to be invoked during the `check` phase.

The `Elisp` dependencies of Emacs packages are typically provided as `propagated-inputs` when required at run time. As for other packages, build or test dependencies should be specified as `native-inputs`.

Emacs packages sometimes depend on resources directories that should be installed along the `Elisp` files. The `#:include` argument can be used for that purpose, by specifying a list of regexps to match. The best practice when using the `#:include` argument is to extend rather than override its default value (accessible via the `%default-include` variable). As an example, a `yasnipet` extension package typically include a `snippets` directory, which could be copied to the installation directory using:

```
#:include (cons "~snippets/" %default-include)
```

When encountering problems, it is wise to check for the presence of the `Package-Requires` extension header in the package main source file, and whether any dependencies and their versions listed therein are satisfied.

### 22.8.9 Python Modules

We currently package Python 2 and Python 3, under the Scheme variable names `python-2` and `python` as explained in Section 22.8.3 [Version Numbers], page 832. To avoid confusion and naming clashes with other programming languages, it seems desirable that the name of a package for a Python module contains the word `python`.

Some modules are compatible with only one version of Python, others with both. If the package `Foo` is compiled with Python 3, we name it `python-foo`. If it is compiled with Python 2, we name it `python2-foo`. Python 2 packages are being removed from the distribution; please do not submit any new Python 2 packages.

If a project already contains the word `python`, we drop this; for instance, the module `python-dateutil` is packaged under the names `python-dateutil` and `python2-dateutil`. If the project name starts with `py` (e.g. `pytz`), we keep it and prefix it as described above.

**Note:** Currently there are two different build systems for Python packages in Guix: *python-build-system* and *pyproject-build-system*. For the longest time, Python packages were built from an informally specified `setup.py` file. That



worked amazingly well, considering Python’s success, but was difficult to build tooling around. As a result, a host of alternative build systems emerged and the community eventually settled on a formal standard (<https://peps.python.org/pep-0517/>) for specifying build requirements. *pyproject-build-system* is Guix’s implementation of this standard. It is considered “experimental” in that it does not yet support all the various PEP-517 *build backends*, but you are encouraged to try it for new Python packages and report any problems. It will eventually be deprecated and merged into *python-build-system*.

### 22.8.9.1 Specifying Dependencies

Dependency information for Python packages is usually available in the package source tree, with varying degrees of accuracy: in the `pyproject.toml` file, the `setup.py` file, in `requirements.txt`, or in `tox.ini` (the latter mostly for test dependencies).

Your mission, when writing a recipe for a Python package, is to map these dependencies to the appropriate type of “input” (see Section 8.2.1 [package Reference], page 104). Although the `pypi` importer normally does a good job (see Section 9.5 [Invoking guix import], page 197), you may want to check the following check list to determine which dependency goes where.

- We currently package Python with `setuptools` and `pip` installed per default. This is about to change, and users are encouraged to use `python-toolchain` if they want a build environment for Python.  
`guix lint` will warn if `setuptools` or `pip` are added as native-inputs because they are generally not necessary.
- Python dependencies required at run time go into `propagated-inputs`. They are typically defined with the `install_requires` keyword in `setup.py`, or in the `requirements.txt` file.
- Python packages required only at build time—e.g., those listed under `build-system.requires` in `pyproject.toml` or with the `setup_requires` keyword in `setup.py`—or dependencies only for testing—e.g., those in `tests_require` or `tox.ini`—go into `native-inputs`. The rationale is that (1) they do not need to be propagated because they are not needed at run time, and (2) in a cross-compilation context, it’s the “native” input that we’d want.

Examples are the `pytest`, `mock`, and `nose` test frameworks. Of course if any of these packages is also required at run-time, it needs to go to `propagated-inputs`.

- Anything that does not fall in the previous categories goes to `inputs`, for example programs or C libraries required for building Python packages containing C extensions.
- If a Python package has optional dependencies (`extras_require`), it is up to you to decide whether to add them or not, based on their usefulness/overhead ratio (see Section 22.10 [Submitting Patches], page 842).

### 22.8.10 Perl Modules

Perl programs standing for themselves are named as any other package, using the lowercase upstream name. For Perl packages containing a single class, we use the lowercase class name, replace all occurrences of `::` by dashes and prepend the prefix `perl-`. So the class `XML::Parser` becomes `perl-xml-parser`. Modules containing several classes keep their

lowercase upstream name and are also prepended by `perl-`. Such modules tend to have the word `perl` somewhere in their name, which gets dropped in favor of the prefix. For instance, `libwww-perl` becomes `perl-libwww`.

### 22.8.11 Java Packages

Java programs standing for themselves are named as any other package, using the lowercase upstream name.

To avoid confusion and naming clashes with other programming languages, it is desirable that the name of a package for a Java package is prefixed with `java-`. If a project already contains the word `java`, we drop this; for instance, the package `ngsjava` is packaged under the name `java-ngs`.

For Java packages containing a single class or a small class hierarchy, we use the lowercase class name, replace all occurrences of `.` by dashes and prepend the prefix `java-`. So the class `apache.commons.cli` becomes package `java-apache-commons-cli`.

### 22.8.12 Rust Crates

Rust applications (binary crates) and libraries (library crates) are packaged separately. We put our main efforts into applications and only package libraries as sources, utilizing automation with a manual focus on unbundling vendored dependencies.

Rust applications are treated like any other package and named using the lowercase upstream name. When using the Cargo build system (see Section 8.5 [Build Systems], page 121), Rust applications should have the `#:install-source?` parameter set to `#f`, as this parameter only makes sense for libraries. When the package source is a Cargo workspace, `#:cargo-install-paths` must be set to enable relevant support.

Rust libraries are hidden from the user interface and managed in two modules:

(gnu packages rust-crates)

Source definitions imported from Rust packages' `Cargo.lock` via the `crate` importer (see Section 9.5 [Invoking guix import], page 197).

Imported definitions must be checked and have vendored dependencies unbundled before being contributed to Guix.

(gnu packages rust-sources)

More complex definitions that need to be full packages. This includes Rust libraries requiring external inputs to unbundle and Cargo workspaces.

These libraries should have the `#:skip-build?` parameter set to `#t`. For Cargo workspaces, `#:cargo-package-crates` must be set.

Since they are added manually, they follow the usual naming convention for Guix packages (see Section 22.8.2 [Package Naming], page 831), with a `rust-` prefix.

In the Rust community it is common for multiple incompatible versions of a package to be used at any given time, so all libraries should have a versioned suffix. The versioned suffix is the left-most non-zero digit (and any leading zeros, of course). This follows the “caret” version scheme intended by Cargo. Examples `rust-clap-2`, `rust-rand-0.6`.

In practice we often package development snapshots of Rust libraries and can't simply identify them by version. In this case, the complete version information can be turned into a version string via `git-version`, for example `rust-pipewire-0.8.0.fd3d8f7` and `rust-pubgrub-0.3.0.b70cf70`.

To avoid merge conflicts with changes from multiple branches, these two modules are maintained by the Rust team (see Section 22.12 [Teams], page 849).

Rust libraries are not referenced directly. (`guix build-system cargo`) provides a `cargo-inputs` procedure to lookup input lists generated by the lockfile importer.

`cargo-inputs name` [`#:module '(gnu packages rust-crates)`] [Procedure]

Lookup Cargo inputs for *name* in *module*, return an empty list if unavailable.

*name* must be consistent with the one used in lockfile importer invocation, usually matching the variable name of the package:

```
$ guix import -i gnu/packages/rust-crates.scm crate -f /path/to/Cargo.lock name
```

*module* must export a public interface `lookup-cargo-inputs`, a template is available in the `etc/teams/rust` directory of Guix source tree.

See Section “Packaging Rust Crates” in *GNU Guix Cookbook*, for packaging workflow.

### 22.8.13 Elm Packages

Elm applications can be named like other software: their names need not mention Elm.

Packages in the Elm sense (see `elm-build-system` under Section 8.5 [Build Systems], page 121) are required use names of the format *author/project*, where both the *author* and the *project* may contain hyphens internally, and the *author* sometimes contains uppercase letters.

To form the Guix package name from the upstream name, we follow a convention similar to Python packages (see Section 22.8.9 [Python Modules], page 836), adding an `elm-` prefix unless the name would already begin with `elm-`.

In many cases we can reconstruct an Elm package's upstream name heuristically, but, since conversion to a Guix-style name involves a loss of information, this is not always possible. Care should be taken to add the '`upstream-name`' property when necessary so that '`guix import elm`' will work correctly (see Section 9.5 [Invoking guix import], page 197). The most notable scenarios when explicitly specifying the upstream name is necessary are:

1. When the *author* is `elm` and the *project* contains one or more hyphens, as with `elm/virtual-dom`; and
2. When the *author* contains hyphens or uppercase letters, as with `Elm-Canvas/raster-shapes`—unless the *author* is `elm-explorations`, which is handled as a special case, so packages like `elm-explorations/markdown` do *not* need to use the '`upstream-name`' property.

The module (`guix build-system elm`) provides the following utilities for working with names and related conventions:

`elm-package-origin elm-name version hash` [Procedure]

Returns a Git origin using the repository naming and tagging regime required for a published Elm package with the upstream name *elm-name* at version *version* with sha256 checksum *hash*.

For example:

```
(package
  (name "elm-html")
  (version "1.0.0")
  (source
    (elm-package-origin
      "elm/html"
      version
        (base32 "15k1679ja57vvlpinpv06znmrxy09lbhzfkzdc89i01qa8c4gb4a"))))■
...)
```

**elm->package-name** *elm-name* [Procedure]  
Returns the Guix-style package name for an Elm package with upstream name *elm-name*.

Note that there is more than one possible *elm-name* for which **elm->package-name** will produce a given result.

**guix-package->elm-name** *package* [Procedure]  
Given an Elm *package*, returns the possibly-inferred upstream name, or **#f** if the upstream name is not specified via the '**upstream-name**' property and can not be inferred by **infer-elm-package-name**.

**infer-elm-package-name** *guix-name* [Procedure]  
Given the *guix-name* of an Elm package, returns the inferred upstream name, or **#f** if the upstream name can't be inferred. If the result is not **#f**, supplying it to **elm->package-name** would produce *guix-name*.

### 22.8.14 Fonts

For fonts that are in general not installed by a user for typesetting purposes, or that are distributed as part of a larger software package, we rely on the general packaging rules for software; for instance, this applies to the fonts delivered as part of the X.Org system or fonts that are part of TeX Live.

To make it easier for a user to search for fonts, names for other packages containing only fonts are constructed as follows, independently of the upstream package name.

The name of a package containing only one font family starts with **font-**; it is followed by the foundry name and a dash - if the foundry is known, and the font family name, in which spaces are replaced by dashes (and as usual, all upper case letters are transformed to lower case). For example, the Gentium font family by SIL is packaged under the name **font-sil-gentium**.

For a package containing several font families, the name of the collection is used in the place of the font family name. For instance, the Liberation fonts consist of three families, Liberation Sans, Liberation Serif and Liberation Mono. These could be packaged separately under the names **font-liberation-sans** and so on; but as they are distributed together under a common name, we prefer to package them together as **font-liberation**.

In the case where several formats of the same font family or font collection are packaged separately, a short form of the format, prepended by a dash, is added to the package name.

We use `-ttf` for TrueType fonts, `-otf` for OpenType fonts and `-type1` for PostScript Type 1 fonts.

## 22.9 Coding Style

In general our code follows the GNU Coding Standards (see *GNU Coding Standards*). However, they do not say much about Scheme, so here are some additional rules.

### 22.9.1 Programming Paradigm

Scheme code in Guix is written in a purely functional style. One exception is code that involves input/output, and procedures that implement low-level concepts, such as the `memoize` procedure.

### 22.9.2 Modules

Guile modules that are meant to be used on the builder side must live in the `(guix build ...)` name space. They must not refer to other Guix or GNU modules. However, it is OK for a “host-side” module to use a build-side module. As an example, the `(guix search-paths)` module should not be imported and used by a package since it isn’t meant to be used as a “build-side” module. It would also couple the module with the package’s dependency graph, which is undesirable.

Modules that deal with the broader GNU system should be in the `(gnu ...)` name space rather than `(guix ...)`.

### 22.9.3 Data Types and Pattern Matching

The tendency in classical Lisp is to use lists to represent everything, and then to browse them “by hand” using `car`, `cdr`, `cadr`, and `co`. There are several problems with that style, notably the fact that it is hard to read, error-prone, and a hindrance to proper type error reports.

Guix code should define appropriate data types (for instance, using `define-record-type*`) rather than abuse lists. In addition, it should use pattern matching, via Guile’s `(ice-9 match)` module, especially when matching lists (see Section “Pattern Matching” in *GNU Guile Reference Manual*); pattern matching for records is better done using `match-record` from `(guix records)`, which, unlike `match`, verifies field names at macro-expansion time.

When defining a new record type, keep the *record type descriptor* (RTD) private (see Section “Records” in *GNU Guile Reference Manual*, for more on records and RTDs). As an example, the `(guix packages)` module defines `<package>` as the RTD for package records but it does not export it; instead, it exports a type predicate, a constructor, and field accessors. Exporting RTDs would make it harder to change the application binary interface (because code in other modules might be matching fields by position) and would make it trivial for users to forge records of that type, bypassing any checks we may have in the official constructor (such as “field sanitizers”).

### 22.9.4 Formatting Code

When writing Scheme code, we follow common wisdom among Scheme programmers. In general, we follow the Riastradh’s Lisp Style Rules (<https://mumble.net/~campbell/>

`scheme/style.txt`). This document happens to describe the conventions mostly used in Guile’s code too. It is very thoughtful and well written, so please do read it.

Some special forms introduced in Guix, such as the `substitute*` macro, have special indentation rules. These are defined in the `.dir-locals.el` file, which Emacs automatically uses. Also note that Emacs-Guix provides `guix-devel-mode` mode that indents and highlights Guix code properly (see Section “Development” in *The Emacs-Guix Reference Manual*).

If you do not use Emacs, please make sure your editor knows these rules. To automatically indent a package definition, you can also run:

```
./pre-inst-env guix style package
```

See Section 9.7 [Invoking guix style], page 213, for more information.

We require all top-level procedures to carry a docstring. This requirement can be relaxed for simple private procedures in the `(guix build ...)` name space, though.

Procedures should not have more than four positional parameters. Use keyword parameters for procedures that take more than four parameters.

## 22.10 Submitting Patches

Development is done using the Git distributed version control system. Thus, access to the repository is not strictly necessary. We welcome contributions as *pull requests* (PR) at <https://codeberg.org/guix/guix/pulls/><sup>6</sup>.

There are two ways to create a pull request:

- By creating a personal “fork” of the repository (<https://codeberg.org/guix/guix>), creating a branch with your changes, and submitting a pull request for that branch (<https://docs.codeberg.org/collaborating/pull-requests-and-git-flow/>).
- By following the *AGit workflow* (<https://forgejo.org/docs/next/user/agit-support/>), which is somewhat more convenient and consumes less disk space at Codeberg since it removes the need to create a “fork”.

Assuming your local checkout is on the branch containing the changes you would like to submit for inclusion in the `master` branch, run:

```
git push origin HEAD:refs/for/master \
-o topic=topic \
-o title=title -o description=description
```

*topic* must be an identifier without whitespace, similar to a branch name; *title* and *description* are arbitrary text that can be omitted, in which case the message of the last commit on the branch fills that role.

To send an update to an AGit pull request of yours, run:

```
git push origin HEAD:refs/for/master \
-o topic=topic -o force-push=yes
```

Codeberg automatically figures out which pull request *topic* corresponds to and updates the associated branch.

<sup>6</sup> Before January 1st, 2026, changes would be submitted by email to [guix-patches@gnu.org](mailto:guix-patches@gnu.org) and could be browsed online (<https://issues.guix.gnu.org>). In accordance with Guix Consensus Document 002 (<https://codeberg.org/guix/guix-consensus-documents/src/branch/main/002-codeberg.md>), this method is no longer supported.

Contributors are encouraged to take a moment to set some Git repository options (see Section 22.10.1 [Configuring Git], page 845) first, which can improve the readability of patches. Seasoned Guix developers may also want to look at the section on commit access (see Section 22.14 [Commit Access], page 851).

Please write commit logs in the ChangeLog format (see Section “Change Logs” in *GNU Coding Standards*); you can check the commit history for examples.

You can help make the review process more efficient, and increase the chance that your patch will be reviewed quickly, by describing the context of your patch and the impact you expect it to have. For example, if your patch is fixing something that is broken, describe the problem and how your patch fixes it. Tell us how you have tested your patch. Will users of the code changed by your patch have to adjust their workflow at all? If so, tell us how. In general, try to imagine what questions a reviewer will ask, and answer those questions in advance.

Before submitting a patch that adds or modifies a package definition, please run through this check list:

1. If the authors of the packaged software provide a cryptographic signature for the release tarball, make an effort to verify the authenticity of the archive. For a detached GPG signature file this would be done with the `gpg --verify` command.
2. Take some time to provide an adequate synopsis and description for the package. See Section 22.8.4 [Synopses and Descriptions], page 833, for some guidelines.
3. Run `guix lint package`, where *package* is the name of the new or modified package, and fix any errors it reports (see Section 9.8 [Invoking guix lint], page 217).
4. Run `guix style package` to format the new package definition according to the project’s conventions (see Section 9.7 [Invoking guix style], page 213).
5. Make sure the package builds on your platform, using `guix build package`. Also build at least its direct dependents with `guix build --dependents=1 package` (see [build-dependents], page 191).
6. We recommend you also try building the package on other supported platforms. As you may not have access to actual hardware platforms, we recommend using the `qemu-binfmt-service-type` to emulate them. In order to enable it, add the `virtualization` service module and the following service to the list of services in your `operating-system` configuration:

```
(service qemu-binfmt-service-type
  (qemu-binfmt-configuration
    (platforms (lookup-qemu-platforms "arm" "aarch64"))))
```

Then reconfigure your system.

You can then build packages for different platforms by specifying the `--system` option. For example, to build the “hello” package for the armhf or aarch64 architectures, you would run the following commands, respectively:

```
guix build --system=armhf-linux --rounds=2 hello
guix build --system=aarch64-linux --rounds=2 hello
```

7. Make sure the package does not use bundled copies of software already available as separate packages.

Sometimes, packages include copies of the source code of their dependencies as a convenience for users. However, as a distribution, we want to make sure that such packages end up using the copy we already have in the distribution, if there is one. This improves resource usage (the dependency is built and stored only once), and allows the distribution to make transverse changes such as applying security updates for a given software package in a single place and have them affect the whole system—something that bundled copies prevent.

8. Take a look at the profile reported by `guix size` (see Section 9.9 [Invoking `guix size`], page 220). This will allow you to notice references to other packages unwillingly retained. It may also help determine whether to split the package (see Section 5.4 [Packages with Multiple Outputs], page 51), and which optional dependencies should be used. In particular, avoid adding `texlive` as a dependency: because of its extreme size, use `texlive-local-tree` procedure instead.
9. Check that dependent packages (if applicable) are not affected by the change; `guix refresh --list-dependent package` will help you do that (see Section 9.6 [Invoking `guix refresh`], page 206).
10. Check whether the package’s build process is deterministic. This typically means checking whether an independent build of the package yields the exact same result that you obtained, bit for bit.

A simple way to do that is by building the same package several times in a row on your machine (see Section 9.1 [Invoking `guix build`], page 180):

```
guix build --rounds=2 my-package
```

This is enough to catch a class of common non-determinism issues, such as timestamps or randomly-generated output in the build result.

Another option is to use `guix challenge` (see Section 9.12 [Invoking `guix challenge`], page 231). You may run it once the package has been committed and built by `bordeaux.guix.gnu.org` to check whether it obtains the same result as you did. Better yet: Find another machine that can build it and run `guix publish`. Since the remote build machine is likely different from yours, this can catch non-determinism issues related to the hardware—e.g., use of different instruction set extensions—or to the operating system kernel—e.g., reliance on `uname` or `/proc` files.

11. When writing documentation, please use gender-neutral wording when referring to people, such as singular “they”, “their”, “them” ([https://en.wikipedia.org/wiki/Singular\\_they](https://en.wikipedia.org/wiki/Singular_they)), and so forth.
12. Verify that your patch contains only one set of related changes. Bundling unrelated changes together makes reviewing harder and slower.  
Examples of unrelated changes include the addition of several packages, or a package update along with fixes to that package.
13. Please follow our code formatting rules, possibly running `guix style` script to do that automatically for you (see Section 22.9.4 [Formatting Code], page 841).
14. When possible, use mirrors in the source URL (see Section 9.3 [Invoking `guix download`], page 195). Use reliable URLs, not generated ones. For instance, GitHub archives are not necessarily identical from one generation to the next, so in this case it’s often better to clone the repository. Don’t use the `name` field in the URL: it is not very useful and if the name changes, the URL will probably be wrong.



15. Check if Guix builds (see Section 22.2 [Building from Git], page 817) and address the warnings, especially those about use of undefined symbols.
16. Make sure your changes do not break Guix and simulate a `guix pull` with:

```
guix pull --url=/path/to/your/checkout --profile=/tmp/guix.master --disable-authen
```

### 22.10.1 Configuring Git

If you have not done so already, you may wish to set a name and email that will be associated with your commits (see Section “Telling Git your name” in *Git User Manual*). If you wish to use a different name or email just for commits in this repository, you can use `git config --local`, or edit `.git/config` in the repository instead of `~/.gitconfig`.

Other important Git configuration will automatically be configured when building the project (see Section 22.2 [Building from Git], page 817). A `.git/hooks/commit-msg` hook will be installed that embeds ‘Change-Id’ Git *trailers* in your commit messages for traceability purposes. It is important to preserve these when editing your commit messages, particularly if a first version of your proposed changes was already submitted for review. If you have a `commit-msg` hook of your own you would like to use with Guix, you can place it under the `.git/hooks/commit-msg.d/` directory.

## 22.11 Tracking Bugs and Changes

This section describes how the Guix project tracks its bug reports, patch submissions and topic branches.

### 22.11.1 The Issue Tracker

Bug reports (*issues*) and patch submissions (*pull requests*) are currently tracked at the project’s home at Codeberg (<https://codeberg.org/guix/guix>).

Before January 1st, 2026, bug reports and patches were tracked using the Debbugs instance at <https://bugs.gnu.org>, which can also be browsed at <https://issues.guix.gnu.org>. This has been discontinued and Codeberg is now the sole option to report new bugs and propose new changes<sup>7</sup>.

### 22.11.2 Managing Patches and Branches

Submitted patches are picked up and tested by the quality assurance service; the result of that testing eventually shows up on the dashboard at ‘[https://qa.guix.gnu.org/issue/ISSUE\\_NUMBER](https://qa.guix.gnu.org/issue/ISSUE_NUMBER)’, where *ISSUE\_NUMBER* is the number assigned by the issue tracker, currently Debbugs (work is in progress to support Codeberg pull requests). Leave time for a review, without committing anything.

As an exception, some changes considered “trivial” or “obvious” may be pushed directly to the `master` branch. This includes changes to fix typos and reverting commits that caused immediate problems. This is subject to being adjusted, allowing individuals to commit directly on non-controversial changes on parts they’re familiar with.

Changes which affect more than 300 dependent packages (see Section 9.6 [Invoking `guix refresh`], page 206) should first be pushed to a topic branch other than

<sup>7</sup> Check out Guix Consensus Document 002 (<https://codeberg.org/guix/guix-consensus-documents/src/branch/main/002-codeberg.md>) for details on this decision.

**master**; the set of changes should be consistent—e.g., “GNOME update”, “NumPy update”, etc. This allows for testing: the branch will automatically show up at ‘<https://qa.guix.gnu.org/branch/branch>’, with an indication of its build status on various platforms.

To help coordinate the merging of branches, you must create a new guix-patches issue each time you create a branch (see Section 22.11.1 [The Issue Tracker], page 845). The title of the issue requesting to merge a branch should have the following format:

Request for merging "*name*" branch

The QA infrastructure (<https://qa.guix.gnu.org/>) recognizes such issues and lists the merge requests on its main page. The following points apply to managing these branches:

1. The commits on the branch should be a combination of the patches relevant to the branch. Patches not related to the topic of the branch should go elsewhere.
2. Any changes that can be made on the master branch, should be made on the master branch. If a commit can be split to apply part of the changes on master, this is good to do.
3. It should be possible to re-create the branch by starting from master and applying the relevant patches.
4. Avoid merging master in to the branch. Prefer rebasing or re-creating the branch on top of an updated master revision.
5. Minimise the changes on master that are missing on the branch prior to merging the branch in to master. This means that the state of the branch better reflects the state of master should the branch be merged.
6. If you don’t have commit access, create the “Request for merging” issue and request that someone creates the branch. Include a list of issues/patches to include on the branch.

Normally branches will be merged in a “first come, first merged” manner, tracked through the guix-patches issues. If you agree on a different order with those involved, you can track this by updating which issues block<sup>8</sup> which other issues. Therefore, to know which branch is at the front of the queue, look for the oldest issue, or the issue that isn’t *blocked* by any other branch merges. An ordered list of branches with the open issues is available at <https://qa.guix.gnu.org>.

Once a branch is at the front of the queue, wait until sufficient time has passed for the build farms to have processed the changes, and for the necessary testing to have happened. For example, you can check ‘<https://qa.guix.gnu.org/branch/branch>’ to see information on some builds and substitute availability.

Once the branch has been merged, the issue should be closed and the branch deleted.

Sometimes, a branch may be a work in progress, for example for larger efforts such as updating the GNOME desktop. In these cases, the branch name should reflect this by having the ‘wip-’ prefix. The QA infrastructure will avoid building work-in-progress branches, so that the available resources can be better focused on building the branches

---

<sup>8</sup> You can mark an issue as blocked by another by emailing [control@debbugs.gnu.org](mailto:control@debbugs.gnu.org) with the following line in the body of the email: `block XXXXX by YYYYY`. Where XXXXX is the number for the blocked issue, and YYYYY is the number for the issue blocking it.

that are ready to be merged. When the branch is no longer a work in progress, it should be renamed, with the ‘wip-’ prefix removed, and only then should the merge requests be created, as documented earlier.

### 22.11.3 Debbugs User Interfaces

**Warning:** This section is about bug reports and patches sent by email to Debbugs (<https://issues.guix.gnu.org>), which is discontinued. Since January 1st, 2026, bug reports and patches are submitted as issues and pull requests on Codeberg (<https://codeberg.org/guix/guix>). See Section 22.10 [Submitting Patches], page 842, for more info.

#### 22.11.3.1 Web interface

A web interface (actually *two* web interfaces!) are available to browse issues:

- <https://issues.guix.gnu.org> provides a pleasant interface powered by mumi<sup>9</sup> to browse bug reports and patches, and to participate in discussions; mumi also has a command-line interface as we will see below;
- <https://bugs.gnu.org/guix> lists bug reports;
- <https://bugs.gnu.org/guix-patches> lists patch submissions.

To view discussions related to issue number *n*, go to ‘<https://issues.guix.gnu.org/n>’ or ‘<https://bugs.gnu.org/n>’.

#### 22.11.3.2 Command-Line Interface

Mumi also comes with a command-line interface that can be used to search existing issues, open new issues, compose replies, apply and send patches. You do not need to use Emacs to use the mumi command-line client. You interact with it only on the command-line.

To use the mumi command-line interface, navigate to a local clone of the Guix git repository, and drop into a shell with mumi, git and git:send-email installed.

```
$ cd guix
~/guix$ guix shell mumi git git:send-email
```

To search for issues, say all open issues about "zig", run

```
~/guix [env]$ mumi search zig is:open
```

```
#60889 Add zig-build-system
opened on 17 Jan 17:37 Z by Ekaitz Zarraga
#61036 [PATCH 0/3] Update zig to 0.10.1
opened on 24 Jan 09:42 Z by Efraim Flashner
#39136 [PATCH] gnu: services: Add endlessh.
opened on 14 Jan 2020 21:21 by Nicol? Balzarotti
#60424 [PATCH] gnu: Add python-online-judge-tools
opened on 30 Dec 2022 07:03 by gemmaro
#45601 [PATCH 0/6] vlang 0.2 update
opened on 1 Jan 2021 19:23 by Ryan Prior
```

<sup>9</sup> Mumi is a nice piece of software written in Guile, and you can help! See <https://codeberg.org/guix/mumi>.

Pick an issue and make it the "current" issue.

```
~/guix [env]$ mumi current 61036
```

```
#61036 [PATCH 0/3] Update zig to 0.10.1
opened on 24 Jan 09:42 Z by Efraim Flashner
```

Once an issue is the current issue, you can open the issue in a web browser, compose replies, apply patches, send patches, etc. with short succinct commands.

Open the issue in your web browser using

```
~/guix [env]$ mumi www
```

Compose a reply using

```
~/guix [env]$ mumi compose
```

Compose a reply and close the issue using

```
~/guix [env]$ mumi compose --close
```

`mumi compose` opens your mail client by passing ‘mailto:’ URIs to `xdg-open`. So, you need to have `xdg-open` set up to open your mail client correctly.

Apply the latest patchset from the issue using

```
~/guix [env]$ mumi am
```

You may also apply a patchset of a specific version (say, v3) using

```
~/guix [env]$ mumi am v3
```

Or, you may apply a patch from a specific e-mail message. For example, to apply the patch from the 4th message (message index starts from 0), run

```
~/guix [env]$ mumi am @4
```

`mumi am` is a wrapper around `git am`. You can pass `git am` arguments to it after a ‘--’. For example, to add a Signed-off-by trailer, run

```
~/guix [env]$ mumi am -- -s
```

Create and send patches to the issue using

```
~/guix [env]$ git format-patch origin/master
~/guix [env]$ mumi send-email foo.patch bar.patch
```

Note that you do not have to pass in ‘--to’ or ‘--cc’ arguments to `git format-patch`. `mumi send-email` will put them in correctly when sending the patches.

To open a new issue, run

```
~/guix [env]$ mumi new
```

and send an email (using `mumi compose`) or patches (using `mumi send-email`).

`mumi send-email` is really a wrapper around `git send-email` that automates away all the nitty-gritty of sending patches. It uses the current issue state to automatically figure out the correct ‘To’ address to send to, other participants to ‘Cc’, headers to add, etc.

Also note that, unlike `git send-email`, `mumi send-email` works perfectly well with single and multiple patches alike. It automates away the debbugs dance of sending the first patch, waiting for a response from debbugs and sending the remaining patches. It does so by sending the first patch, polling the server for a response, and then sending the remaining patches. This polling can unfortunately take a few minutes. So, please be patient.

### 22.11.3.3 Emacs Interface

If you use Emacs, you may find it more convenient to interact with issues using `debbugs.el`, which you can install with:

```
guix install emacs-debbugs
```

For example, to list all open issues on `guix-patches`, hit:

```
C-u M-x debbugs-gnu RET RET guix-patches RET n y
```

For a more convenient (shorter) way to access both the bugs and patches submissions, you may want to configure the `debbugs-gnu-default-packages` and `debbugs-gnu-default-severities` Emacs variables (see Section 22.5.1 [Viewing Bugs within Emacs], page 823).

To search for bugs, ‘`M-x debbugs-gnu-guix-search`’ can be used.

See *Debbugs User Guide*, for more information on this nifty tool!

### 22.11.4 Cuirass Build Notifications

Cuirass includes RSS (Really Simple Syndication) feeds as one of its features (see Section “Notifications” in *cuirass*). Since Berlin (<https://ci.guix.gnu.org/>) runs an instance of Cuirass, this feature can be used to keep track of recently broken or fixed packages caused by changes pushed to the Guix git repository. Any RSS client can be used. A good one, included with Emacs, is See Section “Gnus” in *gnus*. To register the feed, copy its URL, then from the main Gnus buffer, ‘`*Group*`’, do the following:

```
G R https://ci.guix.gnu.org/events/rss/?specification=master RET
Guix CI - master RET Build events for specification master. RET
```

Then, back at the ‘`*Group*`’ buffer, press `s` to save the newly added RSS group. As for any other Gnus group, you can update its content by pressing the `g` key. You should now receive notifications that read like:

```
. [ ?: Cuirass ] Build tree-sitter-meson.aarch64-linux on master is fixed.■
. [ ?: Cuirass ] Build rust-pbkdf2.aarch64-linux on master is fixed.
. [ ?: Cuirass ] Build rust-pbkdf2.x86_64-linux on master is fixed.
```

where each RSS entry contains a link to the Cuirass build details page of the associated build.

## 22.12 Teams

To organize work on Guix, including but not just development efforts, the project has a set of *teams*. Each team has its own focus and interests and is the primary contact point for questions and contributions in those areas. A team’s primary mission is to coordinate and review the work of individuals in its scope (see Section 22.15 [Reviewing the Work of Others], page 854); it can make decisions within its scope, in agreement with other teams whenever there is overlap or a close connection, and in accordance with other project rules such as seeking consensus (see Section 22.13 [Making Decisions], page 851).

As an example, the Python team is responsible for core Python packaging matters; it can decide to upgrade core Python packages in a dedicated `python-team` branch, in collaboration with any team whose scope is directly dependent on Python—e.g., the Science team—and following branching rules (see Section 22.11.2 [Managing Patches and Branches], page 845). The Documentation team helps review changes to the documentation and can

initiate overarching documentation changes. The Translations team organizes translation of Guix and its manual and coordinates efforts in that area. The Core team is responsible for the development of core functionality and interfaces of Guix; because of its central nature, some of its work may require soliciting input from the community at large and seeking consensus before enacting decisions that would affect the entire community.

Teams are defined in the `etc/teams.scm` file in the Guix repository. The scope of each team is defined, when applicable, as a set of files or as a regular expression matching file names.

Anyone with interest in a team’s domain and willing to contribute to its work can apply to become a member by contacting current members by email; commit access is not a precondition. Membership is formalized by adding the person’s name and email address to `etc/teams.scm`. Members who have not been participating in the team’s work for one year or more may be removed; they are free to reapply for membership later.

One or more people may propose the creation of a new team by reaching out to the community by email at `guix-devel@gnu.org`, clarifying the intended scope and purpose. When consensus is reached on the creation of this team, someone with commit access formalizes its creation by adding it and its initial members to `etc/teams.scm`.

To list existing teams, run the following command from a Guix checkout:

```
$ ./etc/teams.scm list-teams
id: mentors
name: Mentors
description: A group of mentors who chaperone contributions by newcomers.
members:
+ Charlie Smith <charlie@example.org>
...
```

You can run the following command to have the Mentors team put in CC of a patch series:

```
$ git send-email --to=ISSUE_NUMBER@debbugs.gnu.org \
  --header-cmd='etc/teams.scm cc-mentors-header-cmd' *.patch
```

The appropriate team or teams can also be inferred from the modified files. For instance, if you want to send the two latest commits of the current Git repository to review, you can run:

```
$ guix shell -D guix
[env]$ git send-email --to=ISSUE_NUMBER@debbugs.gnu.org -2
```

To generate a `CODEOWNERS` file, which Forgejo—the software behind Codeberg—uses to determine which team or person should review changes to a given set of files, run:

```
./etc/teams.scm codeowners > CODEOWNERS
```

Teams defined in `etc/teams.scm` are reflected on Codeberg (<https://codeberg.org/guix/guix>); the “Owners” team on Codeberg, whose members have administration privileges over the entire “organization” (<https://codeberg.org/guix>), is the only one missing from this file. To synchronize team changes with Codeberg, a member of the “Owners” team can run:

```
./etc/teams.scm sync-codeberg-teams token
```

... where *token* is a token created via the Codeberg applications settings page (<https://codeberg.org/user/settings/applications>), granting read/write access to the ‘organization’ permission.

## 22.13 Making Decisions

It is expected from all contributors, and even more so from committers, to help build consensus and make decisions based on consensus. By using consensus, we are committed to finding solutions that everyone can live with. It implies that no decision is made against significant concerns and these concerns are actively resolved with proposals that work for everyone.

A contributor (who may or may not have commit access) wishing to block a proposal bears a special responsibility for finding alternatives, proposing ideas/code or explain the rationale for the status quo to resolve the deadlock. To learn what consensus decision making means and understand its finer details, you are encouraged to read <https://www.seedsforchange.org.uk/consensus>.

## 22.14 Commit Access

Everyone can contribute to Guix without having commit access (see Section 22.10 [Submitting Patches], page 842). However, for frequent contributors, having write access to the repository can be convenient. As a rule of thumb, a contributor should have accumulated fifty (50) reviewed commits to be considered as a committer and have sustained their activity in the project for at least 6 months. This ensures enough interactions with the contributor, which is essential for mentoring and assessing whether they are ready to become a committer. Commit access should not be thought of as a “badge of honor” but rather as a responsibility a contributor is willing to take to help the project.

Committers are in a position where they enact technical decisions. Such decisions must be made by *actively building consensus* among interested parties and stakeholders. See Section 22.13 [Making Decisions], page 851, for more on that.

The following sections explain how to get commit access, how to be ready to push commits, and the policies and community expectations for commits pushed upstream.

### 22.14.1 Applying for Commit Access

When you deem it necessary, consider applying for commit access by following these steps:

1. Find three committers who would vouch for you. Each of them should email a statement to [guix-maintainers@gnu.org](mailto:guix-maintainers@gnu.org) (a private alias for the collective of maintainers), signed with their OpenPGP key.

Committers are expected to have had some interactions with you as a contributor and to be able to judge whether you are sufficiently familiar with the project’s practices. It is *not* a judgment on the value of your work, so a refusal should rather be interpreted as “let’s try again later”.

2. Send [guix-maintainers@gnu.org](mailto:guix-maintainers@gnu.org) a message stating your intent, listing the three committers who support your application, signed with the OpenPGP key you will use to sign commits, and giving its fingerprint (see below). See <https://emailselfdefense.fsf.org/en/>, for an introduction to public-key cryptography with GnuPG.

Set up GnuPG such that it never uses the SHA1 hash algorithm for digital signatures, which is known to be unsafe since 2019, for instance by adding the following line to `~/.gnupg/gpg.conf` (see Section “GPG Esoteric Options” in *The GNU Privacy Guard Manual*):

```
digest-algo sha512
```

3. Maintainers ultimately decide whether to grant you commit access, usually following your referrals’ recommendation.
4. If and once you’ve been given access, please send a message to `guix-devel@gnu.org` to say so, again signed with the OpenPGP key you will use to sign commits (do that before pushing your first commit). That way, everyone can notice and ensure you control that OpenPGP key.

**Important:** Before you can push for the first time, maintainers must:

1. add your OpenPGP key to the `keyring` branch;
  2. add your OpenPGP fingerprint to the `.guix-authorizations` file of the branch(es) you will commit to.
5. Make sure to read the rest of this section and... profit!

**Note:** Maintainers are happy to give commit access to people who have been contributing for some time and have a track record—don’t be shy and don’t underestimate your work!

However, note that the project is working towards a more automated patch review and merging system, which, as a consequence, may lead us to have fewer people with commit access to the main repository. Stay tuned!

All commits that are pushed to the central repository on Codeberg must be signed with an OpenPGP key, and the public key should be uploaded to your user account on Codeberg and to public key servers, such as `keys.openpgp.org`. To configure Git to automatically sign commits, run:

```
git config commit.gpgsign true
```

```
# Substitute the fingerprint of your public PGP key.
git config user.signingkey CABBA6EA1DC0FF33
```

To check that commits are signed with correct key, use:

```
guix git authenticate
```

See Section 22.2 [Building from Git], page 817, for running the first authentication of a Guix checkout.

To avoid accidentally pushing unsigned or signed with the wrong key commits to Codeberg, make sure to configure Git according to See Section 22.10.1 [Configuring Git], page 845.

### 22.14.2 Commit Policy

If you get commit access, please make sure to follow the policy below (discussions of the policy can take place on `guix-devel@gnu.org`).

Ensure you’re aware of how the changes should be handled (see Section 22.11.2 [Managing Patches and Branches], page 845) prior to being pushed to the repository, especially for the `master` branch.



If you're committing and pushing your own changes, try and wait at least one week (two weeks for more significant changes, up to one month for changes such as removing a package—see [package-removal-policy], page 857) after you send them for review. After this, if no one else is available to review them and if you're confident about the changes, it's OK to commit.

When pushing a commit on behalf of somebody else, please add a **Signed-off-by** line at the end of the commit log message—e.g., with `git am --signoff`. This improves tracking of who did what.

When adding channel news entries (see Chapter 6 [Channels], page 68), make sure they are well-formed by running the following command right before pushing:

```
make check-channel-news
```

### 22.14.3 Addressing Issues

Peer review (see Section 22.10 [Submitting Patches], page 842) and tools such as `guix lint` (see Section 9.8 [Invoking guix lint], page 217) and the test suite (see Section 22.3 [Running the Test Suite], page 819) should catch issues before they are pushed. Yet, commits that “break” functionality might occasionally go through. When that happens, there are two priorities: mitigating the impact, and understanding what happened to reduce the chance of similar incidents in the future. The responsibility for both these things primarily lies with those involved, but like everything this is a group effort.

Some issues can directly affect all users—for instance because they make `guix pull` fail or break core functionality, because they break major packages (at build time or run time), or because they introduce known security vulnerabilities.

The people involved in authoring, reviewing, and pushing such commit(s) should be at the forefront to mitigate their impact in a timely fashion: by pushing a followup commit to fix it (if possible), or by reverting it to leave time to come up with a proper fix, and by communicating with other developers about the problem.

If these persons are unavailable to address the issue in time, other committers are entitled to revert the commit(s), explaining in the commit log and on the mailing list what the problem was, with the goal of leaving time to the original committer, reviewer(s), and author(s) to propose a way forward.

Once the problem has been dealt with, it is the responsibility of those involved to make sure the situation is understood. If you are working to understand what happened, focus on gathering information and avoid assigning any blame. Do ask those involved to describe what happened, do not ask them to explain the situation—this would implicitly blame them, which is unhelpful. Accountability comes from a consensus about the problem, learning from it and improving processes so that it's less likely to reoccur.

#### 22.14.3.1 Reverting commits

Like normal commits, the commit message should state why the changes are being made, which in this case would be why the commits are being reverted.

If the changes are being reverted because they led to excessive number of packages being affected, then a decision should be made whether to allow the build farms to build the changes, or whether to avoid this. For the bordeaux build farm, commits can be ignored by

adding them to the `ignore-commits` list in the `build-from-guix-data-service` record, found in the bayfront machine configuration.

#### 22.14.4 Commit Revocation

In order to reduce the possibility of mistakes, committers will have their Codeberg account removed from the Guix Codeberg project and their key removed from `.guix-authorizations` after 12 months of inactivity; they can ask to regain commit access by emailing the maintainers, without going through the vouching process.

Maintainers<sup>10</sup> may also revoke an individual's commit rights, as a last resort, if cooperation with the rest of the community has caused too much friction—even within the bounds of the project's code of conduct (see Chapter 22 [Contributing], page 816). They would only do so after public or private discussion with the individual and a clear notice. Examples of behavior that hinders cooperation and could lead to such a decision include:

- repeated violation of the commit policy stated above;
- repeated failure to take peer criticism into account;
- breaching trust through a series of grave incidents.

When maintainers resort to such a decision, they notify developers on `guix-devel@gnu.org`; inquiries may be sent to `guix-maintainers@gnu.org`. Depending on the situation, the individual may still be welcome to contribute.

#### 22.14.5 Helping Out

One last thing: the project keeps moving forward because committers not only push their own awesome changes, but also offer some of their time *reviewing* and pushing other people's changes. As a committer, you're welcome to use your expertise and commit rights to help other contributors, too!

### 22.15 Reviewing the Work of Others

Perhaps the biggest action you can do to help GNU Guix grow as a project is to review the work contributed by others. You do not need to be a committer to do so; applying, reading the source, building, linting and running other people's series and sharing your comments about your experience will give some confidence to committers. You must ensure the check list found in the Section 22.10 [Submitting Patches], page 842, section has been correctly followed. A reviewed patch series should give the best chances for the proposed change to be merged faster, so if a change you would like to see merged hasn't yet been reviewed, this is the most appropriate thing to do! If you would like to review changes in a specific area and to receive notifications for incoming patches relevant to that domain, consider joining the relevant team(s) (see Section 22.12 [Teams], page 849).

Review comments should be unambiguous; be as clear and explicit as you can about what you think should be changed, ensuring the author can take action on it. Please try to keep the following guidelines in mind during review:

---

<sup>10</sup> See <https://guix.gnu.org/en/about> for the current list of maintainers. You can email them privately at `guix-maintainers@gnu.org`.

1. *Be clear and explicit about changes you are suggesting*, ensuring the author can take action on it. In particular, it is a good idea to explicitly ask for new revisions when you want it.
2. *Remain focused: do not change the scope of the work being reviewed*. For example, if the contribution touches code that follows a pattern deemed unwieldy, it would be unfair to ask the submitter to fix all occurrences of that pattern in the code; to put it simply, if a problem unrelated to the patch at hand was already there, do not ask the submitter to fix it.
3. *Ensure progress*. As they respond to review, submitters may submit new revisions of their changes; avoid requesting changes that you did not request in the previous round of comments. Overall, the submitter should get a clear sense of progress; the number of items open for discussion should clearly decrease over time.
4. *Aim for finalization*. Reviewing code is time-consuming. Your goal as a reviewer is to put the process on a clear path towards integration, possibly with agreed-upon changes, or rejection, with a clear and mutually-understood reasoning. Avoid leaving the review process in a lingering state with no clear way out.
5. *Review is a discussion*. The submitter's and reviewer's views on how to achieve a particular change may not always be aligned. To lead the discussion, remain focused, ensure progress and aim for finalization, spending time proportional to the stakes<sup>11</sup>. As a reviewer, try hard to explain the rationale for suggestions you make, and to understand and take into account the submitter's motivation for doing things in a certain way. In other words, build consensus with everyone involved (see Section 22.13 [Making Decisions], page 851).

When you deem the proposed change adequate and ready for inclusion in Guix, you can explicitly say so for pull requests on Codeberg, for example by selecting “Approve changes” in the “Finish review” dialog of the web interface.

The following well understood/codified ‘Reviewed-by: Your Name <your-email@example.com>’<sup>12</sup> line should be used to sign off as a reviewer, meaning you have reviewed the change and that it looks good to you:

- If the *whole* series (containing multiple commits) looks good to you, reply with ‘Reviewed-by: Your Name <your-email@example.com>’ to the cover page if it has one, or to the last patch of the series otherwise, adding another ‘(for the whole series)’ comment on the line below to explicit this fact.
- If you instead want to mark a *single commit* as reviewed (but not the whole series), simply reply with ‘Reviewed-by: Your Name <your-email@example.com>’ to that commit message.

If you are not a committer, you can help others find a *series* you have reviewed more easily by adding a **reviewed-looks-good** usertag for the **guix** user.

<sup>11</sup> The tendency to discuss minute details at length is often referred to as “bikeshedding”, where much time is spent discussing each one’s preference for the color of the shed at the expense of progress made on the project to keep bikes dry.

<sup>12</sup> The ‘Reviewed-by’ Git trailer is used by other projects such as Linux, and is understood by third-party tools such as the ‘b4 am’ sub-command, which is able to retrieve the complete submission email thread from a public-inbox instance and add the Git trailers found in replies to the commit patches.

## 22.16 Updating the Guix Package

It is sometimes desirable to update the `guix` package itself (the package defined in `(gnu packages package-management)`), for example to make new daemon features available for use by the `guix-service-type` service type. In order to simplify this task, the following command can be used:

```
make update-guix-package
```

The `update-guix-package` make target will use the last known *commit* corresponding to `HEAD` in your Guix checkout, compute the hash of the Guix sources corresponding to that commit and update the `commit`, `revision` and hash of the `guix` package definition.

To validate that the updated `guix` package hashes are correct and that it can be built successfully, the following command can be run from the directory of your Guix checkout:

```
./pre-inst-env guix build guix
```

To guard against accidentally updating the `guix` package to a commit that others can't refer to, a check is made that the commit used has already been pushed to the Codeberg-hosted Guix Git repository.

This check can be disabled, *at your own peril*, by setting the `GUIX_ALLOW_ME_TO_USE_PRIVATE_COMMIT` environment variable. When this variable is set, the updated package source is also added to the store. This is used as part of the release process of Guix.

## 22.17 Deprecation Policy

As any lively project with a broad scope, Guix changes all the time and at all levels. Because it's user-extensible and programmable, incompatible changes can directly impact users and make their life harder. It is thus important to reduce user-visible incompatible changes to a minimum and, when such changes are deemed necessary, to clearly communicate them through a *deprecation period* so everyone can adapt with minimum hassle. This section defines the project's commitments for smooth deprecation and describes procedures and mechanisms to honor them.

There are several ways to use Guix; how to handle deprecation will depend on each use case. Those can be roughly categorized like this:

- package management exclusively through the command line;
- advanced package management using the manifest and package interfaces;
- Home and System management, using the `operating-system` and/or `home-environment` interfaces together with the service interfaces;
- development or use of external tools that use programming interfaces such as the `(guix ...)` modules.

These use cases form a spectrum with varying degrees of coupling—from “distant” to tightly coupled. Based on this insight, we define the following *deprecation policies* that we consider suitable for each of these levels.

### Command-line tools

Guix sub-commands should be thought of as remaining available “forever”. Once a Guix sub-command is to be removed, it should be deprecated first, and then remain available for **at least one year** after the first release that deprecated it.

Deprecation should first be announced in the manual and as an entry in `etc/news.scm`; additional communication such as a blog post explaining the rationale is welcome. Months before the scheduled removal date, the command should print a warning explaining how to migrate. An example of this is the replacement of `guix environment` by `guix shell`, started in October 2021<sup>13</sup>.

Because of the broad impact of such a change, we recommend conducting a user survey before enacting a plan.

#### Package name changes

When a package name changes, it must remain available under its old name for **at least one year**. For example, `go-ipfs` was renamed to `kubo` following a decision made upstream; to communicate the name change to users, the package module provided this definition:

```
(define-deprecated-package go-ipfs
  kubo)
```

That way, someone running `guix install go-ipfs` or similar sees a deprecation warning mentioning the new package name; likewise, someone referring to the `go-ipfs` *variable* from their manifest or package definition will see a warning.

Under the hood, `define-deprecated-package` combines `define-deprecated/public-alias` to mark the variable as deprecated (see below), and the `deprecated-package` procedure to mark the package name as deprecated.

#### Package removal

Packages whose upstream developers have declared as having reached “end of life” or being unmaintained may be removed; likewise, packages that have been **failing to build for two months or more** may be removed.

There is no formal deprecation mechanism for this case, unless a replacement exists, in which case the `define-deprecated-package` macro mentioned above can be used.

If the package being removed is a “leaf” (no other packages depend on it), it may be removed after a **one-month review period** of the patch removing it (this applies even when the removal has additional motivations such as security problems affecting the package).

**Note:** Learn about pending package removals by checking issues and pull requests with the ‘`deprecation`’ label (<https://codeberg.org/guix/guix/issues?labels=445131>).

If the package has many dependent packages—as is the case for example with Python version 2—the relevant team must propose a deprecation removal agenda and seek consensus with other packagers for **at least one month**. It may also invite feedback from the broader user community, for example through a survey. Removal of all impacted packages may be gradual, spanning multiple months, to accommodate all use cases.

<sup>13</sup> For more details on the `guix shell` transition, see <https://guix.gnu.org/en/blog/2021/from-guix-environment-to-guix-shell/>.

When the package being removed is considered popular, whether or not it is a leaf, its deprecation must be announced as an entry in `etc/news.scm`.

#### Package upgrade

In the case of packages with many dependents and/or many users, an upgrade may be treated like the *removal* of the previous version.

Examples include major version upgrades of programming language implementations, as we’ve seen above with Python, and major upgrades of “big” libraries such as Qt or GTK.

#### Services

Changes to services for Guix Home and Guix System have a direct impact on user configuration. For a user, adjusting to interface changes is rarely rewarding, which is why any such change must be clearly communicated in advance through deprecation warnings and documentation.

Renaming of variables related to service, home, or system configuration must be communicated for at least six months before removal using the `(guix deprecation)` mechanisms. For example, renaming of `murmur-configuration` to `mumble-server-configuration` was communicated through a series of definitions like this one:

```
(define-deprecated/public-alias
  murmur-configuration
  mumble-server-configuration)
```

Procedures slated for removal may be defined like this:

```
(define-deprecated (elogind-service #:key (config (elogind-configuration)))
  elogind-service-type
  (service elogind-service-type config))
```

Record fields, notably fields of service configuration records, must follow a similar deprecation period. This is usually achieved through *ad hoc* means though. For example, the `hosts-file` field of `operating-system` was deprecated by adding a `sanitized` property that would emit a warning:

```
(define-record-type* <operating-system>
  ;; ...
  (hosts-file %operating-system-hosts-file          ;deprecated
    (default #f)
    (sanitize warn-hosts-file-field-deprecation)))

(define-deprecated (operating-system-hosts-file os)
  hosts-service-type
  (%operating-system-hosts-file os))
```

When deprecating interfaces in `operating-system`, `home-environment`, `(gnu services)`, or any popular service, the deprecation must come with an entry in `etc/news.scm`.

#### Core interfaces

Core programming interfaces, in particular the `(guix ...)` modules, may be relied on by a variety of external tools and channels. Any incompatible change must be formally deprecated with `define-deprecated`, as shown above, for **at**

**least one year** before removal. The manual must clearly document the new interface and, except in obvious cases, explain how to migrate from the old one. As an example, the `build-expression->derivation` procedure was superseded by `gexp->derivation` and remained available as a deprecated symbol:

```
(define-deprecated (build-expression->derivation store name exp
  #:key ...)
  gexp->derivation
  ...)
```

Sometimes bindings are moved from one module to another. In those cases, bindings must be reexported from the original module for at least one year.

This section does not cover all possible situations but hopefully allows users to know what to expect and developers to stick to its spirit. Please email [guix-devel@gnu.org](mailto:guix-devel@gnu.org) for any questions.

## 22.18 Writing Documentation

Guix is documented using the Texinfo system. If you are not yet familiar with it, we accept contributions for documentation in most formats. That includes plain text, Markdown, Org, etc.

Documentation contributions can be sent to [guix-patches@gnu.org](mailto:guix-patches@gnu.org). Prepend `[DOCUMENTATION]` to the subject.

When you need to make more than a simple addition to the documentation, we prefer that you send a proper patch as opposed to sending an email as described above. See Section 22.10 [Submitting Patches], page 842, for more information on how to send your patches.

To modify the documentation, you need to edit `doc/guix.texi` and `doc/contributing.texi` (which contains this documentation section), or `doc/guix-cookbook.texi` for the cookbook. If you compiled the Guix repository before, you will have many more `.texi` files that are translations of these documents. Do not modify them, the translation is managed through Weblate (<https://translate.codeberg.org/projects/guix>). See Section 22.19 [Translating Guix], page 859, for more information.

To render documentation, you must first make sure that you ran `./configure` in your source tree (see Section 22.4 [Running Guix Before It Is Installed], page 821). After that you can run one of the following commands:

- `'make doc/guix.info'` to compile the Info manual. You can check it with `info doc/guix.info`.
- `'make doc/guix.html'` to compile the HTML version. You can point your browser to the relevant file in the `doc/guix.html` directory.
- `'make doc/guix-cookbook.info'` for the cookbook Info manual.
- `'make doc/guix-cookbook.html'` for the cookbook HTML version.

## 22.19 Translating Guix

Writing code and packages is not the only way to provide a meaningful contribution to Guix. Translating to a language you speak is another example of a valuable contribution

you can make. This section is designed to describe the translation process. It gives you advice on how you can get involved, what can be translated, what mistakes you should avoid and what we can do to help you!

Guix is a big project that has multiple components that can be translated. We coordinate the translation effort on a Weblate instance (<https://translate.codeberg.org/projects/guix/>) hosted by Codeberg. You will need an account to submit translations.

Some of the software packaged in Guix also contain translations. We do not host a translation platform for them. If you want to translate a package provided by Guix, you should contact their developers or find the information on their website. As an example, you can find the homepage of the `hello` package by typing `guix show hello`. On the “homepage” line, you will see <https://www.gnu.org/software/hello/> as the homepage.

Many GNU and non-GNU packages can be translated on the Translation Project (<https://translationproject.org>). Some projects with multiple components have their own platform. For instance, GNOME has its own platform, Damned Lies (<https://110n.gnome.org/>).

Guix has five components hosted on Weblate.

- `guix` contains all the strings from the Guix software (the guided system installer, the package manager, etc), excluding packages.
- `packages` contains the synopsis (single-sentence description of a package) and description (longer description) of packages in Guix.
- `website` contains the official Guix website, except for blog posts and multimedia content.
- `documentation-manual` corresponds to this manual.
- `documentation-cookbook` is the component for the cookbook.

## General Directions

Once you get an account, you should be able to select a component from the guix project (<https://translate.codeberg.org/projects/guix/>), and select a language. If your language does not appear in the list, go to the bottom and click on the “Start new translation” button. Select the language you want to translate to from the list, to start your new translation.

Like lots of other free software packages, Guix uses GNU Gettext (<https://www.gnu.org/software/gettext>) for its translations, with which translatable strings are extracted from the source code to so-called PO files.

Even though PO files are text files, changes should not be made with a text editor but with PO editing software. Weblate integrates PO editing functionality. Alternatively, translators can use any of various free-software tools for filling in translations, of which Poedit (<https://poedit.net/>) is one example, and (after logging in) upload (<https://docs.weblate.org/en/latest/user/files.html>) the changed file. There is also a special PO editing mode (<https://www.emacswiki.org/emacs/PoMode>) for users of GNU Emacs. Over time translators find out what software they are happy with and what features they need.

On Weblate, you will find various links to the editor, that will show various subsets (or all) of the strings. Have a look around and at the documentation (<https://docs.weblate.org/en/latest/>) to familiarize yourself with the platform.



## Translation Components

In this section, we provide more detailed guidance on the translation process, as well as details on what you should or should not do. When in doubt, please contact us, we will be happy to help!

**guix** Guix is written in the Guile programming language, and some strings contain special formatting that is interpreted by Guile. These special formatting should be highlighted by Weblate. They start with `~` followed by one or more characters.

When printing the string, Guile replaces the special formatting symbols with actual values. For instance, the string `'ambiguous package specification ~a'` would be substituted to contain said package specification instead of `~a`. To properly translate this string, you must keep the formatting code in your translation, although you can place it where it makes sense in your language. For instance, the French translation says `'spécification du paquet « ~a » ambiguë'` because the adjective needs to be placed in the end of the sentence.

If there are multiple formatting symbols, make sure to respect the order. Guile does not know in which order you intended the string to be read, so it will substitute the symbols in the same order as the English sentence.

As an example, you cannot translate `'package '~a' has been superseded by '~a'` by `'~a' supersedes package '~a'`, because the meaning would be reversed. If *foo* is superseded by *bar*, the translation would read `'foo' supersedes package 'bar'`. To work around this problem, it is possible to use more advanced formatting to select a given piece of data, instead of following the default English order. See Section “Formatted Output” in *GNU Guile Reference Manual*, for more information on formatting in Guile.

**packages**

Package descriptions occasionally contain Texinfo markup (see Section 22.8.4 [Synopses and Descriptions], page 833). Texinfo markup looks like `@code{rm-rf}`, `@emph{important}`, etc. When translating, please leave markup as is.

The characters after “@” form the name of the markup, and the text between “{” and “}” is its content. In general, you should not translate the content of markup like `@code`, as it contains literal code that do not change with language. You can translate the content of formatting markup such as `@emph`, `@i`, `@itemize`, `@item`. However, do not translate the name of the markup, or it will not be recognized. Do not translate the word after `@end`, it is the name of the markup that is closed at this position (e.g. `@itemize ... @end itemize`).

**documentation-manual and documentation-cookbook**

The first step to ensure a successful translation of the manual is to find and translate the following strings *first*:

- **version.texi**: Translate this string as `version-xx.texi`, where *xx* is your language code (the one shown in the URL on weblate).
- **contributing.texi**: Translate this string as `contributing.xx.texi`, where *xx* is the same language code.

- **Top:** Do not translate this string, it is important for Texinfo. If you translate it, the document will be empty (missing a Top node). Please look for it, and register **Top** as its translation.

Translating these strings first ensure we can include your translation in the guix repository without breaking the make process or the `guix pull` machinery.

The manual and the cookbook both use Texinfo. As for `packages`, please keep Texinfo markup as is. There are more possible markup types in the manual than in the package descriptions. In general, do not translate the content of `@code`, `@file`, `@var`, `@value`, etc. You should translate the content of formatting markup such as `@emph`, `@i`, etc.

The manual contains sections that can be referred to by name by `@ref`, `@xref` and `@pxref`. We have a mechanism in place so you do not have to translate their content. If you keep the English title, we will automatically replace it with your translation of that title. This ensures that Texinfo will always be able to find the node. If you decide to change the translation of the title, the references will automatically be updated and you will not have to update them all yourself.

When translating references from the cookbook to the manual, you need to replace the name of the manual and the name of the section. For instance, to translate `@pxref{Defining Packages,,, guix, GNU Guix Reference Manual}`, you would replace `Defining Packages` with the title of that section in the translated manual *only* if that title is translated. If the title is not translated in your language yet, do not translate it here, or the link will be broken. Replace `guix` with `guix.xx` where `xx` is your language code. `GNU Guix Reference Manual` is the text of the link. You can translate it however you wish.

## website

The website pages are written using SXML, an s-expression version of HTML, the basic language of the web. We have a process to extract translatable strings from the source, and replace complex s-expressions with a more familiar XML markup, where each markup is numbered. Translators can arbitrarily change the ordering, as in the following example.

```
#. TRANSLATORS: Defining Packages is a section name
#. in the English (en) manual.
#: apps/base/templates/about.scm:64
msgid "Packages are <1>defined<1.1>en</1.1><1.2>Defining-Packages.html</1.2>"
msgstr "Pakete werden als reine <2>Guile</2>-Module <1>definiert<1.1>de</1.1>"
```

Note that you need to include the same markups. You cannot skip any.

In case you make a mistake, the component might fail to build properly with your language, or even make `guix pull` fail. To prevent that, we have a process in place to check the content of the files before pushing to our repository. We will not be able to update the translation for your language in Guix, so we will notify you (through weblate and/or by email) so you get a chance to fix the issue.

## Outside of Weblate

Currently, some parts of Guix cannot be translated on Weblate, help wanted!

- `guix pull` news can be translated in `news.scm`, but is not available from Weblate. If you want to provide a translation, you can prepare a patch as described above, or simply send us your translation with the name of the news entry you translated and your language. See Section 6.12 [Writing Channel News], page 76, for more information about channel news.
- Guix blog posts cannot currently be translated.
- The installer script (for foreign distributions) is entirely in English.
- Some of the libraries Guix uses cannot be translated or are translated outside of the Guix project. Guile itself is not internationalized.
- Other manuals linked from this manual or the cookbook might not be translated.

## Conditions for Inclusion

There are no conditions for adding new translations of the `guix` and `guix-packages` components, other than they need at least one translated string. New languages will be added to Guix as soon as possible. The files may be removed if they fall out of sync and have no more translated strings.

Given that the web site is dedicated to new users, we want its translation to be as complete as possible before we include it in the language menu. For a new language to be included, it needs to reach at least 80% completion. When a language is included, it may be removed in the future if it stays out of sync and falls below 60% completion.

The manual and cookbook are automatically added in the default compilation target. Every time we synchronize translations, developers need to recompile all the translated manuals and cookbooks. This is useless for what is essentially the English manual or cookbook. Therefore, we will only include a new language when it reaches 10% completion in the component. When a language is included, it may be removed in the future if it stays out of sync and falls below 5% completion.

## Translation Infrastructure

Weblate is backed by a git repository from which it discovers new strings to translate and pushes new and updated translations. Normally, it would be enough to give it commit access to our repositories. However, we decided to use a separate repository for two reasons. First, we would have to give Weblate commit access and authorize its signing key, but we do not trust it in the same way we trust guix developers, especially since we do not manage the instance ourselves. Second, if translators mess something up, it can break the generation of the website and/or guix pull for all our users, independently of their language.

For these reasons, we use a dedicated repository to host translations, and we synchronize it with our guix and artworks repositories after checking no issue was introduced in the translation.

Developers can download the latest PO files from weblate in the Guix repository by running the `make download-po` command. It will automatically download the latest files from weblate, reformat them to a canonical form, and check they do not contain issues. The manual needs to be built again to check for additional issues that might crash Texinfo.

Before pushing new translation files, developers should add them to the make machinery so the translations are actually available. The process differs for the various components.

- New po files for the `guix` and `packages` components must be registered by adding the new language to `po/guix/LINGUAS` or `po/packages/LINGUAS`.
- New po files for the `documentation-manual` component must be registered by adding the file name to `DOC_PO_FILES` in `po/doc/local.mk`, the generated `%D%/guix.xx.texi` manual to `info_TEXINFOS` in `doc/local.mk` and the generated `%D%/guix.xx.texi` and `%D%/contributing.xx.texi` to `TRANSLATED_INFO` also in `doc/local.mk`.
- New po files for the `documentation-cookbook` component must be registered by adding the file name to `DOC_COOKBOOK_PO_FILES` in `po/doc/local.mk`, the generated `%D%/guix-cookbook.xx.texi` manual to `info_TEXINFOS` in `doc/local.mk` and the generated `%D%/guix-cookbook.xx.texi` to `TRANSLATED_INFO` also in `doc/local.mk`.
- New po files for the `website` component must be added to the `guix-artwork` repository, in `website/po/`. `website/po/LINGUAS` and `website/po/ietf-tags.scm` must be updated accordingly (see `website/i18n-howto.txt` for more information on the process).

## 22.20 Contributing to Guix’s Infrastructure

Since its inception, the Guix project has always valued its autonomy, and that reflects in its infrastructure: our servers run Guix System and exclusively free software, and are administered by volunteers.

Of course this comes at a cost and this is why we need contributions. Our hope is to make infrastructure-related activity more legible so that maybe you can picture yourself helping in one of these areas.

### 22.20.1 Coding

The project runs many Guix-specific services; this is all lovely Scheme code but it tends to receive less attention than Guix itself:

- Build Farm Front End (<https://git.cbaines.net/guix/bffe>)
- Cuirass (<https://codeberg.org/guix/cuirass/>)
- Goggles (<https://codeberg.org/guix/maintenance/src/branch/master/hydra/goggles.scm>) (IRC logger)
- Guix Build Coordinator (<https://codeberg.org/guix/build-coordinator>)
- Guix Data Service (<https://codeberg.org/guix/data-service>)
- Guix Packages Website (<https://codeberg.org/luis-felipe/guix-packages-website>)
- mumi (<https://codeberg.org/guix/mumi>)
- Nar Herder (<https://codeberg.org/guix/nar-herder>)
- QA Frontpage (<https://codeberg.org/guix/qa-frontpage>)

There is no time constraint on this coding activity: any improvement is welcome, whenever it comes. Most of these code bases are relatively small, which should make it easier to get started.

Prerequisites: Familiarity with Guile, HTTP, and databases.

If you wish to get started, check out the README of the project of your choice and get in touch with guix-devel and the primary developer(s) of the tool as per `git shortlog -s | sort -k1 -n`.

## 22.20.2 System Administration

Guix System configuration for all our systems is held in this repository:

<https://codeberg.org/guix/maintenance.git/src/branch/master/hydra/>

The two front-ends are `berlin.scm` (the machine behind `ci.guix.gnu.org`) and `bayfront.scm` (the machine behind `bordeaux.guix.gnu.org`, `guix.gnu.org`, `hpc.guix.info`, `qa.guix.gnu.org`, and more). Both connect to a number of build machines and helpers.

Without even having SSH access to the machine, you can help by posting patches to improve the configuration (you can test it with `guix system vm`). Here are ways you can help:

- Improve infra monitoring: set up a dashboard to monitor all the infrastructure, and an out-of-band channel to communicate about downtime.
- Implement web site redundancy: `guix.gnu.org` should be backed by several machines on different sites. Get in touch with us and/or send a patch!
- Implement substitute redundancy: likewise, `bordeaux.guix.gnu.org` and `ci.guix.gnu.org` should be backed by several head nodes.
- Improve backup: there's currently ad-hoc backup of selected pieces over rsync between the two head nodes; we can improve on that, for example with a dedicated backup site and proper testing of recoverability.
- Support mirroring: We'd like to make it easy for others to mirror substitutes from `ci.guix` and `bordeaux.guix`, perhaps by offering public rsync access.
- Optimize our web services: Monitor the performance of our services and tweak nginx config or whatever it takes to improve it.

There is no time constraint on this activity: any improvement is welcome, whenever you can work on it.

Prerequisite: Familiarity with Guix System administration and ideally with the infrastructure handbook:

Infra-Handbook (<https://codeberg.org/guix/maintenance.git/src/branch/master/doc/infra-handbook.org>)

## 22.20.3 Day-to-Day System Administration

We're also looking for people who'd be willing to have SSH access to some of the infrastructure to help with day-to-day maintenance: restarting a build, restarting the occasional service that has gone wild (that can happen), reconfiguring/upgrading a machine, rebooting, etc.

This day-to-day activity requires you to be available some of the time (during office hours or not, during the week-end or not), whenever is convenient for you, so you can react to issues reported on IRC, on the mailing list, or elsewhere, and synchronize with other sysadmins.

Prerequisite: Being a “known” member of the community, familiarity with Guix System administration, with some of the services/web sites being run, and with the infrastructure handbook:

Infra-Handbook (<https://codeberg.org/guix/maintenance.git/src/branch/master/doc/infra-handbook.org>)

#### 22.20.4 On-Site Intervention

The first front-end is currently generously hosted by the Max Delbrück Center (MDC), a research institute in Berlin, Germany. Only authorized personnel can physically access it.

The second one, [bordeaux.guix.gnu.org](https://bordeaux.guix.gnu.org), is hosted in Bordeaux, France, in a professional data center shared with non-profit ISP Aquilenet. If you live in the region of Bordeaux and would like to help out when we need to go on-site, please make yourself known by emailing [guix-sysadmin@gnu.org](mailto:guix-sysadmin@gnu.org).

On-site interventions are rare, but they’re usually in response to an emergency.

#### 22.20.5 Hosting

We’re looking for people who can host machines and help out whenever physical access is needed. More specifically:

- We need hosting of “small” machines such as single-board computers (AArch64, RISC-V) for use as build machines.
- We need hosting for front-ends and x86\_64 build machines in a data center where they can be racked and where, ideally, several local Guix sysadmins can physically access them.

The machines should be accessible over Wireguard VPN most of the time, so longer power or network interruptions should be the exception.

Prerequisites: Familiarity with installing and remotely administering Guix System.

#### 22.20.6 Administrative Tasks

The infra remains up and running thanks to crucial administrative tasks, which includes:

- Selecting and purchasing hardware, for example build machines.
- Renewing domain names.
- Securing funding, in particular via the Guix Foundation: Guix Foundation (<https://foundation.guix.info>)

Prerequisites: Familiarity with hardware, and/or DNS registrars, and/or sponsorship, and/or crowdfunding.

## 23 Acknowledgments

Guix is based on the Nix package manager (<https://nixos.org/nix/>), which was designed and implemented by Eelco Dolstra, with contributions from other people (see the `nix/AUTHORS` file in Guix). Nix pioneered functional package management, and promoted unprecedented features, such as transactional package upgrades and rollbacks, per-user profiles, and referentially transparent build processes. Without this work, Guix would not exist.

The Nix-based software distributions, Nixpkgs and NixOS, have also been an inspiration for Guix.

GNU Guix itself is a collective work with contributions from a number of people. See the `AUTHORS` file in Guix for more information on these fine people. The `THANKS` file lists people who have helped by reporting bugs, taking care of the infrastructure, providing artwork and themes, making suggestions, and more—thank you!

# Appendix A GNU Free Documentation License

Version 1.3, 3 November 2008

Copyright © 2000, 2001, 2002, 2007, 2008 Free Software Foundation, Inc.

<https://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

## 0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document *free* in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or non-commercially. Secondly, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of “copyleft”, which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

## 1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The “Document”, below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as “you”. You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A “Modified Version” of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A “Secondary Section” is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document’s overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The “Invariant Sections” are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released



under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The “Cover Texts” are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A “Transparent” copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not “Transparent” is called “Opaque”.

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The “Title Page” means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, “Title Page” means the text near the most prominent appearance of the work’s title, preceding the beginning of the body of the text.

The “publisher” means any person or entity that distributes copies of the Document to the public.

A section “Entitled XYZ” means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as “Acknowledgements”, “Dedications”, “Endorsements”, or “History”.) To “Preserve the Title” of such a section when you modify the Document means that it remains a section “Entitled XYZ” according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

## 2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

### 3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

### 4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

- A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any,

- be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.
- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
  - C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
  - D. Preserve all the copyright notices of the Document.
  - E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
  - F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
  - G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
  - H. Include an unaltered copy of this License.
  - I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
  - J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
  - K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
  - L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
  - M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
  - N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
  - O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their

titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

## 5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements."

## 6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

## 7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an “aggregate” if the copyright resulting from the compilation is not used to limit the legal rights of the compilation’s users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document’s Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

## 8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled “Acknowledgements”, “Dedications”, or “History”, the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

## 9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense, or distribute it is void, and will automatically terminate your rights under this License.

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, receipt of a copy of some or all of the same material does not give you any rights to use it.

## 10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <https://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License “or any later version” applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation. If the Document specifies that a proxy can decide which future versions of this License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Document.

## 11. RELICENSING

“Massive Multiauthor Collaboration Site” (or “MMC Site”) means any World Wide Web server that publishes copyrightable works and also provides prominent facilities for anybody to edit those works. A public wiki that anybody can edit is an example of such a server. A “Massive Multiauthor Collaboration” (or “MMC”) contained in the site means any set of copyrightable works thus published on the MMC site.

“CC-BY-SA” means the Creative Commons Attribution-Share Alike 3.0 license published by Creative Commons Corporation, a not-for-profit corporation with a principal place of business in San Francisco, California, as well as future copyleft versions of that license published by that same organization.

“Incorporate” means to publish or republish a Document, in whole or in part, as part of another Document.

An MMC is “eligible for relicensing” if it is licensed under this License, and if all works that were first published under this License somewhere other than this MMC, and subsequently incorporated in whole or in part into the MMC, (1) had no cover texts or invariant sections, and (2) were thus incorporated prior to November 1, 2008.

The operator of an MMC Site may republish an MMC contained in the site under CC-BY-SA on the same site at any time before August 1, 2009, provided the MMC is eligible for relicensing.

## ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

```
Copyright (C)  year  your name.
Permission is granted to copy, distribute and/or modify this document
under the terms of the GNU Free Documentation License, Version 1.3
or any later version published by the Free Software Foundation;
with no Invariant Sections, no Front-Cover Texts, and no Back-Cover
Texts.  A copy of the license is included in the section entitled ``GNU
Free Documentation License''.
```

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the “with...Texts.” line with this:

```
with the Invariant Sections being list their titles, with
the Front-Cover Texts being list, and with the Back-Cover Texts
being list.
```

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

# Concept Index

.

.local, host name lookup ..... 680  
 .NET ..... 199

/

/bin/sh ..... 276  
 /etc/hosts default entries ..... 277  
 /usr/bin/env ..... 276

## A

AArch64, bootloaders ..... 685  
 access control list (ACL), for substitutes ..... 47  
 accounts ..... 269  
 ACL (access control list), for substitutes ..... 47  
 actions, of Shepherd services ..... 719  
 agate ..... 515  
 AGit workflow, for pull requests ..... 842  
 aliases, for email addresses ..... 433  
 aliases, for **guix package** ..... 37  
 ALSA ..... 396  
 AMD GPUs ..... 184  
 Android distribution ..... 124  
 Android NDK build system ..... 124  
 AppArmor ..... 13  
 AppArmor, profile installation ..... 13  
 AppImage, create an AppImage file with **guix pack** ..... 92, 93  
 application bundle ..... 91  
 archival of source code, Software Heritage ..... 218  
 archive ..... 65  
 ARM, bootloaders ..... 685  
 Audit ..... 665  
 authentication, of a Guix checkout ..... 817  
 authentication, of channel code ..... 59, 71  
 authentication, of Git checkouts ..... 98  
 authenticity, of code obtained with **guix pull** .. 56  
 authorized keys, SSH ..... 343  
 authorizing, archives ..... 67  
 autofs ..... 548  
 automount ..... 548  
 AutoSSH ..... 345  
 availability of substitutes ..... 235

## B

Backup ..... 674  
 backup service, Syncthing ..... 334, 779  
 bag (low-level package representation) ..... 121  
 bash ..... 740, 785  
 Beets service, for Home ..... 781  
 bill of materials (manifests) ..... 117  
 bind mounts, service ..... 641  
**binfmt\_misc** ..... 588  
 Bioconductor ..... 200  
 BIOS boot, on Intel machines ..... 246  
 BIOS, bootloader ..... 685  
 black list, of kernel modules ..... 684  
 boot loader ..... 685  
 boot menu ..... 689  
 bootloader ..... 685  
 bootstrap binaries ..... 808, 814  
 bootstrapping ..... 808  
 Bounced Network Connection, BNC ..... 451  
 bouncer, IRC ..... 451  
 branching strategy ..... 845  
 build code quoting ..... 165  
 build daemon ..... 1, 5  
 build daemon, unprivileged ..... 7  
 build environment ..... 6, 14  
 build environment isolation ..... 8  
 build event notifications, RSS feed ..... 849  
 build failures, debugging ..... 193  
 build farm ..... 47  
 build hook ..... 9  
 build logs, access ..... 193  
 build logs, publication ..... 228  
 build logs, verbosity ..... 182  
 build phases ..... 122, 141, 151  
 build phases, customizing ..... 145  
 build phases, for packages ..... 143  
 build phases, modifying ..... 151  
 build system ..... 121  
 build system, directory structure ..... 826  
 build users ..... 6  
 build VMs ..... 591  
 build-side modules ..... 841  
 build-time dependencies ..... 158  
 building the dependents of a package ..... 191  
 bundle ..... 91  
 bundling ..... 843



## C

cache invalidation, `nsd` ..... 283  
 cache, sharing ..... 631  
`cached` ..... 646  
 caching, in `guix shell` ..... 84  
 caching, of profiles ..... 84  
 CalDAV ..... 433, 544  
 capabilities, POSIX ..... 678  
 CardDAV ..... 433  
 Cargo (Rust build system) ..... 125  
 cargo inputs ..... 839  
`cat-avatar-generator` ..... 512  
 CD image format ..... 699  
 certificates ..... 86  
 Cgit service ..... 610  
 challenge ..... 231  
 channel authorizations ..... 74  
 channel introduction ..... 75  
 channel news ..... 58  
 channels ..... 68  
 channels, for personal packages ..... 72  
 channels, for the default Guix ..... 70  
`channels.scm`, configuration file ..... 58, 68  
`childhurd` ..... 593  
`childhurd`, offloading ..... 595  
`chroot` ..... 8, 14  
`chroot`, guix system ..... 730  
`chrooting`, guix system ..... 730  
 Clojure (programming language) ..... 127  
 closure ..... 55, 220  
 clusters, daemon setup ..... 18, 157  
 code of conduct, of contributors ..... 816  
 code snippets ..... 822  
 code staging ..... 146, 165  
 code staging, in package definitions ..... 102  
 coding style ..... 841  
 colliding packages in profiles ..... 43  
 collisions, in a profile ..... 43  
 command modules ..... 827  
 command-line tools, as Guile modules ..... 827  
 commit access, for developers ..... 851  
`commit-msg` hook ..... 845  
 complex configurations ..... 723  
`Composer` ..... 204  
 composition of Guix revisions ..... 62  
 Compressed RAM-based block devices ..... 648  
 compressed swap ..... 648  
 configuration file for channels ..... 58, 68  
 configuration file, of a shepherd service ..... 275  
 configuration file, of Shepherd services ..... 721  
 configuration file, of the system ..... 242  
 configuration of `guix pull` ..... 68  
 configuration, action for shepherd services ..... 275  
 configure flags, changing them ..... 187  
`Connman` ..... 315  
 consensus seeking ..... 851  
 container ..... 82, 86, 89, 234  
 container management, `podman` ..... 655

container nesting, for `guix shell` ..... 84  
 container, build environment ..... 14  
 container, for `guix home` ..... 733, 785  
`containerd`, container runtime ..... 653  
 containers, subordinate IDs ..... 271  
`ContentDB` ..... 199  
 continuous integration ..... 186, 554  
 continuous integration, Forgejo ..... 559  
 continuous integration, statistics ..... 235  
 contributor covenant ..... 816  
 copy, of store items, over SSH ..... 233  
 corruption, recovering from ..... 56, 193  
 CPAN ..... 200  
 CPU frequency scaling with `thermald` ..... 568  
 CRAN ..... 200  
 crate ..... 203  
 Creating system images in various formats ..... 696  
 creating virtual machine images ..... 697  
`cron` ..... 299  
`cron`, per-user ..... 745  
 cross compilation ..... 166  
 cross compilation, package dependencies ..... 105  
 cross-compilation ..... 97, 103, 192  
 CTAN ..... 201  
 customization, of packages ..... 1, 100  
 customization, of services ..... 247  
 customizing packages ..... 113  
 CVE, Common Vulnerabilities and Exposures ..... 218

## D

daemon ..... 5  
 daemon, cluster setup ..... 18, 157  
 daemon, remote access ..... 18, 157  
 daemons ..... 709  
 DAG ..... 222  
`darkstat` ..... 474  
 database ..... 401  
`Debbugs`, issue tracking system ..... 845  
`Debbugs`, Mumi Web interface ..... 507  
 Debian, build a `.deb` package with `guix pack` ..... 92  
 debugging files ..... 801  
 debugging info, rebuilding ..... 185, 802  
 decision making ..... 851  
 deduplication ..... 17, 56  
 deleting home generations ..... 787  
 deleting system generations ..... 695  
 dependencies, build-time ..... 158  
 dependencies, channels ..... 74  
 dependencies, run-time ..... 158  
 dependency graph rewriting ..... 116  
 dependency graph, of Shepherd services ..... 252  
 dependents of a package, building them ..... 191  
 deprecation of programming interfaces ..... 858  
 deprecation policy ..... 856  
 derivation ..... 55, 100  
 derivation path ..... 158  
 derivations ..... 158

determinism, checking ..... 193  
 determinism, of build processes ..... 844  
 development environments ..... 78  
 development inputs, of a package ..... 108  
 device mapping ..... 263  
 DHCP ..... 28  
 DHCP, networking service ..... 311  
 DHCPd, networking service ..... 311  
 dhtproxy, for use with jami ..... 329  
 dictionary ..... 652  
 dictionary service, for Home ..... 781  
 digital signatures ..... 49  
 directories related to foreign distro ..... 4  
 disk encryption ..... 265  
 disk space ..... 53  
 display manager, lightdm ..... 360  
 DLNA/UPnP ..... 676  
 DNS (domain name system) ..... 522  
 Docker ..... 653  
 Docker, build an image with guix pack ..... 91  
 docker-image, creating docker images ..... 697  
 documentation ..... 51, 859  
 documentation, searching for ..... 790  
 domain name system (DNS) ..... 522  
 dot files in Guix Home ..... 737  
 downgrade attacks, protection against ..... 59  
 downloading package sources ..... 195  
 dual boot ..... 689  
 DVD image format ..... 699  
 Dynamic IP, with Wireguard ..... 542  
 dyndns, usage with Wireguard ..... 542

## E

early out of memory daemon ..... 643  
 earlyoom ..... 643  
 editing, service type definition ..... 693, 786  
 EFI boot ..... 246  
 EFI, bootloader ..... 685  
 EFI, installation ..... 28  
 egg ..... 206  
 elisp, packaging ..... 836  
 elm ..... 204  
 Elm ..... 839  
 elpa ..... 203  
 emacs ..... 21  
 emacs, packaging ..... 836  
 email ..... 407, 437  
 email aliases ..... 433  
 email workflow, deprecated ..... 845  
 emulation ..... 588, 589  
 encrypted disk ..... 29, 248  
 EndlessSSH ..... 345  
 entry point arguments, for docker images ..... 95  
 entry point, for Docker and Singularity images ..... 95  
 env, in /usr/bin ..... 276  
 environment variables ..... 735  
 environment, package build environment ..... 78

error strategy ..... 699  
 ESP, EFI system partition ..... 28  
 essential services ..... 256  
 exporting files from the store ..... 65  
 exporting store items ..... 65  
 extending the package collection (channels) ..... 68  
 extensibility of the distribution ..... 1  
 extension graph, of services ..... 252  
 extensions, for gexps ..... 167

## F

Fail2Ban ..... 668  
 fastcgi ..... 509  
 fc-cache ..... 20  
 fcgiwrap ..... 509  
 feature branches, coordination ..... 846  
 FHS (file system hierarchy standard) ..... 83  
 file search ..... 52, 400  
 file system hierarchy standard (FHS) ..... 83  
 file, searching ..... 149  
 file, searching in packages ..... 52  
 file-like objects ..... 170  
 fingerprint ..... 649  
 firewall, nftables ..... 325  
 firmware ..... 254  
 fixed-output derivations ..... 158  
 fixed-output derivations, for download ..... 109  
 font cache ..... 20  
 fonts ..... 20, 840  
 foreign architectures ..... 239  
 foreign distro ..... 4, 18  
 Forgejo, continuous integration ..... 559  
 format conventions ..... 213  
 format, code style ..... 213  
 formatting code ..... 841  
 formatting, of code ..... 842  
 free software ..... 831  
 fs-cache, file system caching (Linux) ..... 646  
 fstrim service ..... 644  
 functional package management ..... 1

## G

G-expression ..... 165  
 ganeti ..... 596  
 garbage collector ..... 53  
 garbage collector root, for environments ..... 85, 87  
 garbage collector root, for packs ..... 97  
 garbage collector roots ..... 17, 53  
 garbage collector roots, adding ..... 193  
 GC roots ..... 17, 53  
 GC roots, adding ..... 193  
 GCC ..... 98  
 GDM ..... 355  
 gem ..... 198  
 generations ..... 41, 45, 59, 694  
 Git checkout authentication ..... 98

git configuration ..... 845  
 Git, forge ..... 624  
 Git, hosting ..... 622  
 Git, using the latest commit ..... 186  
 Git, web interface ..... 610  
 Gitile service ..... 624  
 Gitolite service ..... 622  
 global security system ..... 547  
 gmnisrv ..... 514  
 GNOME, login manager ..... 355  
 GNU Build System ..... 102  
 GNU Mailutils IMAP4 Daemon ..... 433  
 GNU Privacy Guard, Home service ..... 754  
 Gnus, configuration to read CI RSS feeds ..... 849  
 go ..... 205  
 go-webdav ..... 504  
 gpg-agent, Home service ..... 754  
 GPG, Home service ..... 754  
 gpm ..... 292  
 GPUs, AMD compilation targets ..... 184  
 grafts ..... 806  
 groups ..... 269, 270  
 GSS ..... 547  
 GSSD ..... 547  
 guix archive ..... 65  
 guix build ..... 180  
 guix challenge ..... 231  
 guix container ..... 234  
 guix copy ..... 233  
 guix deploy ..... 701  
 guix describe ..... 64  
 guix download ..... 195  
 guix edit ..... 195  
 guix environment ..... 78, 85  
 guix gc ..... 53  
 guix git authenticate ..... 98  
 guix graph ..... 222  
 guix hash ..... 196  
 guix home ..... 784  
 guix lint ..... 217  
 guix pack ..... 90  
 guix package ..... 37  
 guix processes ..... 237  
 guix publish ..... 227  
 guix pull ..... 56  
 guix pull for the root user, on a foreign distro ..... 21  
 guix pull, configuration file ..... 68  
 guix refresh ..... 206  
 guix repl ..... 175  
 guix shell ..... 78  
 guix size ..... 220  
 guix style ..... 213  
 Guix System ..... 1, 2, 4  
 guix system ..... 692  
 Guix System Distribution, now Guix System ..... 1  
 guix system troubleshooting ..... 729  
 Guix System, installation ..... 23  
 guix time-machine ..... 60

guix weather ..... 235  
 guix-daemon ..... 14  
 guix-emacs-autoload-packages, refreshing Emacs  
   packages ..... 21  
 GuixSD, now Guix System ..... 1

## H

hackage ..... 202  
 hardware support on Guix System ..... 23  
 HDPI ..... 691, 692  
 hermetic build environment ..... 8  
 hexpm ..... 206  
 hibernation ..... 266  
 HiDPI ..... 691, 692  
 home configuration ..... 732  
 home generations ..... 787  
 home services ..... 735  
 host-side modules ..... 841  
 hostapd service, for Wi-Fi access points ..... 323  
 hpcguix-web ..... 513  
 HTTP ..... 494  
 HTTP proxy, for guix-daemon ..... 288  
 HTTP, HTTPS ..... 518  
 HTTPS, certificates ..... 679  
 hurd ..... 253, 593  
 Hurd, offloading ..... 595

## I

i18n ..... 859  
 idmapd ..... 547  
 image, creating disk images ..... 696  
 IMAP ..... 429  
 implicit inputs, of a package ..... 108  
 imported modules, for gexps ..... 166  
 importer modules ..... 827  
 importing files to the store ..... 65  
 importing packages ..... 197  
 incompatibility, of locale data ..... 274  
 indentation, of code ..... 842  
 inetd ..... 327  
 inferior packages ..... 62, 63  
 inferiors ..... 62, 175  
 Info, documentation format ..... 790  
 infrastructure ..... 864  
 inherit, for package definitions ..... 113  
 init system ..... 717  
 initial RAM disk ..... 254, 682, 684  
 initrd ..... 254, 682, 684  
 input rewriting ..... 116  
 inputattach ..... 651  
 inputrc ..... 743  
 inputs, for Python packages ..... 837  
 inputs, of packages ..... 105  
 insert or update copyright ..... 823  
 inspecting system services ..... 252  
 installation image ..... 32

installer script ..... 4  
 installing Guix ..... 4  
 installing Guix from binaries ..... 4  
 installing Guix System ..... 23  
 installing over SSH ..... 28  
 installing packages ..... 37  
 integrity checking ..... 56  
 integrity, of the store ..... 56  
 interactive shell ..... 740  
 interactive use ..... 176  
 introduction, for Git authentication ..... 98  
 invalid store items ..... 157  
 Invoking **guix import** ..... 197  
 invoking programs, from Scheme ..... 150  
 IPFS ..... 351  
 iptables ..... 324  
 irc-bot ..... 454, 765  
 IRC (Internet Relay Chat) ..... 444, 445, 451, 454  
 IRC bouncer configuration for Libera.Chat ..... 451  
 IRC bouncer configuration for OFTC ..... 451  
 IRC daemon service ..... 445  
 IRC gateway ..... 444  
 IRC server service ..... 445  
 IRCd, Internet Relay Chat daemon ..... 445  
 ISO-9660 format ..... 699  
 isolated build environment ..... 8  
 isolation ..... 1  
 issue tracker, for the email workflow ..... 845  
 iwd, iNet Wireless Daemon ..... 319

## J

jabber ..... 437  
 jackd ..... 294  
 java ..... 838  
 joycond ..... 626  
 JSON ..... 65  
 JSON, import ..... 201

## K

keepalived ..... 351  
 Kerberos ..... 483  
 kernel module loader ..... 645  
 kexec, for fast reboots ..... 693  
 keyboard layout ..... 27, 271  
 keyboard layout, configuration ..... 272  
 keyboard layout, definition ..... 271  
 keyboard layout, for the bootloader ..... 688  
 keyboard layout, for Xorg ..... 363  
 keymap ..... 271  
 keymap, for Xorg ..... 363  
 keyword arguments, for build systems ..... 102  
 kodi ..... 766

## L

l10n ..... 859  
 latest commit, building ..... 186  
 L<sup>A</sup>T<sub>E</sub>X packages ..... 804  
 ld-wrapper ..... 98  
 LDAP ..... 484  
 LDAP, server ..... 491  
 legacy boot, on Intel machines ..... 246  
 Let's Encrypt ..... 518  
 LGTM, Looks Good To Me ..... 855  
 Libera.Chat, IRC bouncer configuration ..... 451  
 license, GNU Free Documentation License ..... 868  
 license, of packages ..... 106  
 lightdm, graphical login manager ..... 360  
 linker wrapper ..... 98  
 lint, code style ..... 213  
 LIRC ..... 650  
 locale ..... 273  
 locale definition ..... 273  
 locale name ..... 274  
 locales, when not on Guix System ..... 18  
 localstatedir ..... 818  
 lock files ..... 69  
 log rotation ..... 304  
 log rotation, for user services ..... 748  
 logging ..... 304  
 logging, anonymization ..... 305  
 logging, syslog ..... 285  
 login manager ..... 355, 358  
 login shell ..... 740  
 loopback device ..... 310  
 lowering, of high-level objects in gexps ..... 166, 174  
 luanti ..... 627  
 LUKS ..... 265  
 LVM, logical volume manager ..... 265

## M

M-x **copyright-update** ..... 823  
 M-x **guix-copyright** ..... 823  
 mail ..... 407  
 mail transfer agent (MTA) ..... 428  
 man pages ..... 790  
 mandatory access control, SELinux ..... 12  
 manifest ..... 117  
 manifest, exporting ..... 45, 81  
 manual pages ..... 790  
 mapped devices ..... 263  
 maximum layers argument, for docker images ..... 96  
 mcron ..... 301  
 mcron, per-user ..... 745  
 mentoring ..... 850  
 merge requests, template ..... 846  
 Message broker ..... 517  
 message of the day ..... 278  
 Message streaming ..... 517  
 messaging ..... 437  
 meta-data, channels ..... 74

minetest .....	199
miniflux .....	516
ModemManager .....	321
Modeswitching .....	322
modprobe .....	645
module closure .....	167
module, black-listing .....	684
monad .....	161
monadic functions .....	161
monadic values .....	161
mouse .....	292
mpd .....	575
MPD, web interface .....	579
mpv, configuration .....	766
mpv, Home Service .....	766
msmtp .....	763
MTA (mail transfer agent) .....	428
multiple-output packages .....	51
Mumble .....	460
mumi am .....	847
mumi command-line interface .....	847
mumi compose .....	847
mumi send-email .....	847
mumi www .....	847
Mumi, Debbugs Web interface .....	507
mumi, web interface for issues .....	847
Murmur .....	460
myMPD service .....	579

## N

name mapper .....	547
name service switch .....	680
name service switch, glibc .....	19
nar bundle, archive format .....	66
nar, archive format .....	66
native language support .....	859
nested containers, for <b>guix shell</b> .....	84
Network information service (NIS) .....	19
network interface controller (NIC) .....	307
networking, with QEMU .....	311
NetworkManager .....	313
news, for channels .....	76
NFS .....	545
NFS, server .....	545
nftables .....	325
NIC, networking interface controller .....	307
nintendo controllers .....	626
niri, configuration .....	772
niri, Home Service .....	772
NIS (Network information service) .....	19
Nix .....	667
Node.js .....	204
nofile .....	294
non-determinism, in package builds .....	232
normalized archive (nar) .....	66
normalized codeset in locale names .....	274
notifications, build events .....	849

npm .....	204
nsd, cache invalidation .....	283
nsd (name service cache daemon) .....	19, 283
nsld, LDAP service .....	484
nsncd, replacement for nsd .....	20
<b>nss-certs</b> .....	20, 679
nss-mdns .....	680
<b>nsswitch.conf</b> .....	19
NSS .....	680
NSS (name service switch), glibc .....	19
NTP (Network Time Protocol), service .....	325
ntpd, service for the Network Time Protocol daemon .....	325
nuget .....	199

## O

OCaml .....	204
OCI-backed, for Home .....	782
OCI-backed, Shepherd services .....	657
official website .....	816
offload status .....	12
offload test .....	12
offloading .....	9, 15
OFTC, IRC bouncer configuration .....	451
on-error .....	699
on-error strategy .....	699
one-shot services, for the Shepherd .....	719
onion service, tor .....	331
onion services, for Tor .....	330
oom .....	643
OPAM .....	204
open file descriptors .....	294
opendht, distributed hash table network service .....	329
OpenNTPD .....	326
OpenPGP, signed commits .....	852
OpenSnitch .....	672
optimization, of package code .....	183
out of memory killer .....	643
outputs .....	51
overlay mounts, service .....	641

## P

pack .....	91
package building .....	180
package conversion .....	197
package definition, editing .....	195
package dependencies .....	55, 222
package deprecation .....	857
package description .....	833
package import .....	197
package installation .....	37
package module search path .....	100
package modules .....	827
package multi-versioning .....	183
package name .....	831
package outputs .....	51

package removal ..... 37  
 package removal policy ..... 857  
 package size ..... 220  
 package synopsis ..... 833  
 package tranformation ..... 182  
 package transformations ..... 115  
 package transformations, upgrades ..... 40  
 package variants ..... 113  
 package variants, transformation ..... 182  
 package version ..... 832  
 package, checking for errors ..... 217  
 packages ..... 36  
 packages, creating ..... 830  
 packaging, maintenance ..... 835  
 pam volume mounting ..... 630  
 pam-krb5 ..... 484  
 pam-mount ..... 628  
 PAM ..... 256  
 Parcimonie, Home service ..... 752  
 partial target version, guix refresh ..... 207  
 partial version specification, guix refresh ..... 207  
 password, for user accounts ..... 270  
 patch submission ..... 842  
 patches ..... 102  
 Patchwork ..... 505  
 pattern matching ..... 841  
 pcsd ..... 650  
 performance, tuning code ..... 183  
 perl ..... 837  
 persistent environment ..... 85, 87  
 personal packages (channels) ..... 72  
 php-fpm ..... 509  
 PHP ..... 204  
 PID 1 ..... 717  
 pinning, channel revisions of a profile ..... 46  
 pinning, channels ..... 60, 69  
 pipefs ..... 547  
 PipeWire, home service ..... 762  
 Platform Reliability, Availability and Serviceability  
   daemon ..... 647  
 pluggable authentication modules ..... 256  
 pluggable transports, tor ..... 331  
 podman, rootless ..... 655  
 POP ..... 429  
 postgis ..... 402  
 postgresql extension-packages ..... 402  
 power consumption tuning with PowerTOP ... 568  
 power management ..... 746  
 power management with TLP ..... 561  
 power-profiles-daemon ..... 560  
 PowerTOP ..... 568  
 pre-built binaries ..... 46  
 primary URL, channels ..... 76  
 printer support with CUPS ..... 365  
 priority ..... 294  
 privileged programs ..... 678  
 profile ..... 33, 37, 38  
 profile collisions ..... 43

profile declaration ..... 41  
 profile manifest ..... 41  
 profile, choosing ..... 42  
 program invocation, from Scheme ..... 150  
 program wrappers ..... 152  
 prometheus-node-exporter ..... 475  
 propagated inputs ..... 38  
 provenance tracking, of software artifacts ..... 37  
 provenance tracking, of the home environment . 786  
 provenance tracking, of the operating  
   system ..... 693, 698, 716  
 provenance, of the system ..... 243  
 proxy, during system installation ..... 28  
 proxy, for **guix-daemon** HTTP access ..... 288  
 pull ..... 56  
 pull requests ..... 842  
 PulseAudio, home service ..... 761  
 PulseAudio, sound support ..... 396  
 purpose ..... 1  
 pypi ..... 198  
 python ..... 836

## Q

QEMU ..... 707  
 QEMU, networking ..... 311

## R

rasdaemon ..... 647  
 read-eval-print loop ..... 821  
 readline ..... 743  
 ReadyMedia, DLNA/UPnP Service ..... 782  
 real time clock ..... 325  
 realm, kerberos ..... 483  
 realtime ..... 294  
 rebooting *via* Linux kexec ..... 693  
 rebuild scheduling strategy ..... 845  
 reconfiguring the system ..... 243, 253  
 reducing boilerplate ..... 822  
 references ..... 158  
 relative file name, in **local-file** ..... 171  
 relocatable binaries ..... 94  
 relocatable binaries, with **guix pack** ..... 91  
 remote access to the daemon ..... 18, 157  
 remote build ..... 556  
 removing packages ..... 37  
 repairing GRUB, via chroot ..... 730  
 repairing store items ..... 193  
 repairing the store ..... 56  
 REPL ..... 821  
 REPL service, for shepherd ..... 719  
 REPL, read-eval-print loop ..... 176  
 REPL, read-eval-print loop, script ..... 175  
 replacements of packages, for grafts ..... 806  
 replicating Guix ..... 60, 64, 69  
 replication, of software environments ..... 37  
 reproducibility ..... 36, 64

reproducibility, checking ..... 193  
 reproducibility, of Guix ..... 60, 69  
 reproducible build environments ..... 78  
 reproducible builds ..... 14, 36, 50, 231  
 reproducible builds, checking ..... 844  
 resize-file-system ..... 673  
 resolution ..... 691, 692  
 reverting commits ..... 853  
 review tags ..... 855  
 Reviewed-by, git trailer ..... 855  
 reviewing, guidelines ..... 854  
 ROCm/HIP GPU targets ..... 184  
 roll back, for the system ..... 244  
 rolling back ..... 41, 59, 694, 787  
 rootless build daemon ..... 7  
 rootless daemon, migration ..... 7  
 rootless **guix-daemon** ..... 285  
 Rootless Podman ..... 655  
 rootless podman, container management tool.. 655  
 rpc\_pipefs ..... 547  
 rpcbind ..... 546  
 RPM, build an RPM archive with guix pack... 93  
 rshiny ..... 666  
 RSS feeds, Gnus configuration ..... 849  
 RTP, for PulseAudio ..... 761  
 rules to cope with circular module dependencies. 835  
 run-time dependencies ..... 158  
 RUNPATH, validation ..... 123, 143  
 rust ..... 838  
 Rust programming language ..... 125  
 RYF, Respects Your Freedom ..... 23

## S

Samba ..... 552  
 saving space ..... 695, 787  
 scanner access ..... 386  
 scheduling jobs ..... 299  
 scheduling jobs per-user ..... 745  
 Scheme programming language, getting started. 103  
 scrollable tiling window manager ..... 772  
 search path ..... 153  
 search paths ..... 38, 42  
 searching for a file ..... 400  
 searching for documentation ..... 790  
 searching for packages ..... 43, 52  
 searching for packages, by file name ..... 52  
 secrets, Knot service ..... 529  
 secure shell client, configuration ..... 749  
 security ..... 47  
 security updates ..... 806  
 security vulnerabilities ..... 218, 806  
 security, **guix pull** ..... 56  
 security, **guix-daemon** ..... 12  
 SELinux, daemon policy ..... 12  
 SELinux, limitations ..... 13  
 SELinux, policy installation ..... 12  
 service deprecation ..... 858

service extension graph, of a home environment . 788  
 service extensions ..... 709  
 service type ..... 714  
 service type definition, editing ..... 693, 786  
 service types ..... 710  
 services ..... 247, 709  
 session limits ..... 294  
 session types ..... 355  
 setcap ..... 678  
 setgid programs ..... 678  
 setuid programs ..... 678  
**sh**, in **/bin** ..... 276  
 shared cache ..... 631  
 sharing store items across machines ..... 233  
 shebang, for **guix shell** ..... 78  
 shell ..... 740, 785  
 shell-profile ..... 785  
 Shepherd dependency graph, for a home  
   environment ..... 788  
 shepherd services ..... 717  
 shepherd services, for users ..... 747  
 Shepherd timers ..... 300  
 shortest path, between packages ..... 225  
 signing, archives ..... 66  
 SIMD support ..... 183  
 simple Clojure build system ..... 127  
 Singularity, build an image with guix pack ..... 91  
 Singularity, container service ..... 655  
 size ..... 220  
 SMB ..... 552  
 SMTP ..... 428  
 snippets, when to use ..... 834  
 snuik ..... 454, 765  
 socket activation, for **guix publish** ..... 227  
 socket activation, for **guix-daemon** ..... 14  
 software bundle ..... 91  
 software development ..... 78  
 Software Heritage, source code archive ..... 218  
 solid state drives, periodic trim ..... 644  
 solid state drives, trim ..... 644  
 sound support ..... 396  
 source, verification ..... 191  
 spam ..... 437  
 SPICE ..... 651  
 SQL ..... 401  
 SquashFS, build an image with guix pack ..... 91  
 ssh-agent ..... 753  
 SSH ..... 341, 708  
 SSH access to build daemons ..... 157  
 SSH agent, with gpg-agent ..... 754  
 SSH authorized keys ..... 343  
 SSH client, configuration ..... 749  
 SSH server ..... 341, 708  
 SSH, copy of store items ..... 233  
 stackage ..... 202  
 staging, for build system arguments ..... 102  
 staging, of code ..... 146, 165  
 state directory ..... 818

state monad ..... 163  
 statistics, for substitutes ..... 235  
 store ..... 1, 156  
 store items ..... 156  
 store paths ..... 156  
 Stow-like dot file management ..... 737  
 strata of code ..... 165  
 structure, of the source tree ..... 826  
 styling rules ..... 213  
 subdirectory, channels ..... 73  
 subid, subordinate user and group IDs ..... 663  
 subordinate user and group IDs ..... 663  
 substitute availability ..... 235  
 substitute servers, adding more ..... 48  
 substituter ..... 831  
 substitutes ..... 14, 37, 46  
 substitutes, authorization thereof ..... 5, 47, 286  
 substitutes, how to disable ..... 48  
 sudo vs. **guix pull** ..... 244  
 sudoers file ..... 256  
 suspend to disk ..... 266  
 swap devices ..... 255  
 swap encryption ..... 265  
 swap space ..... 266  
 sway, configuration ..... 772  
 sway, Home Service ..... 772  
 symbolic links, guix shell ..... 83  
 syncthing ..... 334  
 Syncthing, file synchronization service ..... 334, 779  
 sysconfd ..... 818  
 sysctl ..... 649  
 syslog, with external syslogd ..... 285  
 syslogd, Shepherd system log service ..... 723  
 system configuration ..... 242  
 system configuration directory ..... 818  
 system configuration file ..... 242  
 system images ..... 793  
 System images, creation in various formats ..... 696  
 system instantiation ..... 243, 253  
 system log service, from Shepherd ..... 723  
 system service ..... 710  
 system services ..... 275  
 system services, inspecting ..... 252  
 system services, upgrading ..... 243  
 system-wide Guix, customization ..... 70

## T

tablet input, for Xorg ..... 651  
 target version, guix refresh ..... 207  
 team creation ..... 850  
 team membership ..... 850  
 teams ..... 849  
 teams, on Codeberg ..... 850  
 telephony, services ..... 456  
 templates ..... 822  
 test suite ..... 819  
 test suite, skipping ..... 188

TeX Live ..... 201  
 TeX packages ..... 804  
 Texinfo markup, in package descriptions ..... 834  
 the Hurd ..... 593  
 thermald ..... 568  
 time traps ..... 591  
 timer service, Shepherd ..... 722  
 timers, per-user ..... 748  
 timers, Shepherd ..... 300  
 tlp ..... 561  
 TLS ..... 679  
 TLS certificates ..... 518  
 tool chain, changing the build tool chain of a  
     package ..... 185  
 tool chain, choosing a package's tool chain .... 108  
 toolchain, for C development ..... 98  
 toolchain, for Fortran development ..... 98  
 Tor ..... 330  
 touchscreen input, for Xorg ..... 651  
 transactions ..... 36, 37  
 transactions, undoing ..... 41, 59  
 transferring store items across machines ..... 233  
 transient service, Shepherd ..... 722  
 translation ..... 859  
 troubleshooting, for system services ..... 252  
 troubleshooting, Guix System ..... 729  
 trust, of pre-built binaries ..... 50  
 tunable packages ..... 183  
 tuning, of package code ..... 183

## U

UEFI boot ..... 246  
 UEFI, bootloader ..... 685  
 UEFI, installation ..... 28  
 ulimit ..... 294  
 unattended upgrades ..... 352  
 undoing transactions ..... 41, 59  
 uninstallation, of Guix ..... 5  
 uninstalling Guix ..... 5  
 unprivileged build daemon ..... 7  
 unprivileged daemon, migration ..... 7  
 unprivileged **guix-daemon** ..... 285  
 update-guix-package, updating the guix  
     package ..... 856  
 updater-extra-inputs, package property ..... 209  
 updater-ignored-inputs, package property ... 209  
 updating Guix ..... 56  
 upgrade, of the system ..... 244  
 upgrades, unattended ..... 352  
 upgrading Guix ..... 56  
 upgrading Guix for the root user, on a foreign  
     distro ..... 21  
 Upgrading Guix, on a foreign distro ..... 21  
 upgrading packages ..... 40  
 upgrading system services ..... 243  
 upgrading the Guix daemon, on a foreign distro . 21  
 upstream, latest version ..... 187



USB\_ModeSwitch..... 322  
 user accounts..... 269  
 user interfaces..... 1  
**user-processes** Shepherd service..... 718  
 users..... 269

## V

vacuum the store database..... 56  
 variant packages (channels)..... 68  
 variants, of packages..... 113  
 Varnish..... 502  
 verbosity, of the command-line tools..... 182  
 verifiable builds..... 231  
 version number, for VCS snapshots..... 832  
 VFS mapping..... 641  
 virtual build machines..... 591  
 Virtual File System mapping..... 641  
 virtual machine..... 695, 707  
 virtual machine, Guix System installation..... 31  
 virtual private network (VPN)..... 538  
 virtual private server (VPS)..... 31  
 VM..... 695  
 VMs, for offloading..... 591  
 VNC (virtual network computing)..... 535  
 vnstat..... 475  
 VoIP server..... 460  
 voxel-based games..... 627  
 VPN (virtual private network)..... 538  
 VPS (virtual private server)..... 31

## W

Wayland compositor..... 772  
 weather, substitute availability..... 235  
 web..... 494  
 Web..... 518  
 WebSSH..... 347  
 wesnothd..... 628  
 Whoogle Search..... 504

Wi-Fi access points, hostapd service..... 323  
 WiFi..... 27  
 WiFi, hardware support..... 23  
 window manager..... 355  
 wip branches..... 846  
 wireless..... 27  
 work-in-progress branches, wip..... 846  
 WPA Supplicant..... 319  
 wrapping programs..... 152  
 wsdd, Web service discovery daemon..... 553  
 www..... 494

## X

X Window System..... 355  
 X Window, for Guix Home services..... 755  
 X.509 certificates..... 679  
 X11..... 355  
 X11 login..... 358  
 X11, in Guix Home..... 755  
 XDMCP (x display manager control protocol)..... 535  
 Xen..... 590  
 Xen guest agent..... 590  
 XKB, keyboard layouts..... 272  
 xlsfonts..... 20  
 XMPP..... 437  
 Xorg, configuration..... 362  
 xterm..... 20

## Z

zabbix zabbix-agent..... 481  
 zabbix zabbix-front-end..... 482  
 zabbix zabbix-server..... 479  
 znc..... 765  
 zram..... 648  
 zsh..... 740, 785

# Programming Index

## #

#~ ..... 102  
#~exp..... 167

## %

%store-directory ..... 147

## (

(gexp..... 167

## >

>>=..... 162

## A

add-text-to-store..... 158  
ar-file?..... 147  
assume-source-relative-file-name ..... 171  
assume-valid-file-name..... 170

## B

base-initrd..... 685  
binary-file..... 164  
build..... 178  
build-derivations..... 158  
build-expression->derivation..... 160  
build-options ..... 178  
bzip-fetch..... 112

## C

cargo-inputs ..... 839  
cat-avatar-generator-service..... 513  
close-connection..... 157  
computed-file ..... 171  
concatenate-manifests ..... 120  
configuration->documentation..... 727  
configure-flags..... 178  
copy-recursively..... 148  
current-state ..... 163  
cvs-fetch..... 112

## D

debootstrap-os ..... 601  
debootstrap-variant ..... 601  
define-cargo-inputs ..... 839  
define-configuration..... 724  
define-deprecated..... 858  
define-deprecated-package..... 857  
define-deprecated/public-alias..... 858  
define-maybe ..... 726  
define-record-type\* ..... 841  
delayed ..... 174  
delete-file-recursively..... 148  
deprecated-package..... 857  
derivation..... 159  
dicod-freedit-database..... 653  
directory-exists?..... 147  
directory-union..... 173

## E

elf-file?..... 147  
elm->package-name..... 840  
elm-package-origin..... 839  
enter-store-monad..... 179  
evaluate-search-paths ..... 156  
executable-file?..... 147  
expression->initrd..... 685  
extra-special-file..... 276

## F

fail2ban-jail-service..... 669  
file->udev-hardware ..... 291  
file->udev-rule ..... 291  
file-append..... 173  
file-name-predicate ..... 149  
file-system-label..... 257, 260  
file-union..... 172  
find-files..... 149  
fold-services ..... 715

## G

generate-documentation..... 727  
geoclue-application ..... 387  
gexp ..... 102  
gexp->approximate-sexp..... 175  
gexp->derivation..... 169  
gexp->file..... 172  
gexp->script ..... 171  
gexp-input..... 174  
gexp?..... 169  
git-fetch..... 111  
git-fetch/lfs ..... 111

git-http-nginx-location-configuration... 609  
 git-version..... 833  
 graft?..... 178  
 grub-theme..... 691  
 guix-for-channels..... 71  
 guix-os..... 601  
 guix-package->elm-name..... 840  
 guix-variant..... 601  
 gzip-file?..... 147

## H

hg-fetch..... 111  
 hg-version..... 833  
 home-environment..... 733  
 host..... 278

## I

infer-elm-package-name..... 840  
 inferior-for-channels..... 63  
 inferior-package-description..... 63  
 inferior-package-home-page..... 63  
 inferior-package-inputs..... 63  
 inferior-package-location..... 63  
 inferior-package-name..... 63  
 inferior-package-native-inputs..... 63  
 inferior-package-native-search-paths..... 63  
 inferior-package-propagated-inputs..... 63  
 inferior-package-search-paths..... 63  
 inferior-package-synopsis..... 63  
 inferior-package-transitive-native-search-paths..... 63  
 inferior-package-transitive-propagated-inputs..... 63  
 inferior-package-version..... 63  
 inferior-package?..... 63  
 inferior-packages..... 63  
 install-file..... 148  
 interned-file..... 164  
 invoke..... 150  
 invoke-error-arguments..... 150  
 invoke-error-exit-status..... 150  
 invoke-error-program..... 150  
 invoke-error-stop-signal..... 150  
 invoke-error-term-signal..... 150  
 invoke-error?..... 150  
 invoke/quiet..... 150

## K

keyboard-layout..... 272

## L

let-system..... 173  
 literal-string..... 735  
 local-file..... 170  
 lookup-cargo-inputs..... 839  
 lookup-inferior-packages..... 63  
 lookup-package-direct-input..... 107  
 lookup-package-input..... 107  
 lookup-package-native-input..... 107  
 lookup-package-propagated-input..... 107  
 lookup-qemu-platforms..... 589  
 lower..... 178  
 lower-object..... 175

## M

make-file-writable..... 148  
 make-flags..... 178  
 make-home-mpv-configuration..... 767  
 make-mpv-profile-configuration..... 767  
 match-record..... 841  
 maybe-value-set?..... 726  
 mbegin..... 163  
 mdns-host-lookup-nss..... 681  
 mixed-text-file..... 172  
 mkdir-p..... 148  
 mlet..... 162  
 mlet\*..... 162  
 modify-inputs..... 114  
 modify-phases..... 151  
 modify-services..... 247, 713  
 munless..... 163  
 mwhen..... 163

## N

nginx-php-location..... 512

## O

open-connection..... 157  
 open-inferior..... 63  
 operating-system..... 244  
 operating-system-derivation..... 253  
 options->transformation..... 115

**P**

package->cross-derivation..... 165  
 package->derivation..... 165  
 package->development-manifest..... 120  
 package->manifest-entry..... 120  
 package-cross-derivation..... 104  
 package-derivation..... 103  
 package-development-inputs..... 108  
 package-file..... 165  
 package-input-rewriting..... 116  
 package-input-rewriting/spec..... 116  
 package-mapping..... 117  
 package-name->name+version..... 147  
 package-with-c-toolchain..... 108  
 packages->manifest..... 41, 120  
 phases..... 178  
 plain-file..... 171  
 program-file..... 171

**Q**

qemu-platform-name..... 589  
 qemu-platform?..... 589

**R**

raw-initrd..... 684  
 report-invoke-error..... 150  
 reset-gzip-timestamp..... 147  
 return..... 162  
 run-in-store..... 179  
 run-with-state..... 164  
 run-with-store..... 164

**S**

scheme-file..... 172  
 search-input-directory..... 149  
 search-input-file..... 149  
 search-patches..... 828  
 serialize-configuration..... 727  
 service..... 712  
 service-extension..... 714  
 service-extension?..... 714  
 service-kind..... 713  
 service-value..... 713  
 service?..... 713  
 set-current-state..... 164  
 set-xorg-configuration..... 272, 363  
 shepherd-configuration-action..... 721  
 shepherd-timer..... 300  
 simple-service..... 715

source-module-closure..... 167  
 specification->package..... 247  
 specifications->manifest..... 121  
 specifications->packages..... 247  
 state-pop..... 164  
 state-push..... 164  
 store-file-name?..... 147  
 strip-store-file-name..... 147  
 substitute\*..... 148  
 svn-fetch..... 112  
 sway-configuration->file..... 773  
 symbolic-link?..... 147  
 symlink-to..... 173

**T**

text-file..... 164  
 text-file\*..... 172  
 this-operating-system..... 256  
 this-package..... 107  
 transmission-password-hash..... 464  
 transmission-random-salt..... 465

**U**

udev-hardware..... 290  
 udev-hardware-service..... 291  
 udev-rule..... 290  
 udev-rules-service..... 290  
 url-fetch..... 110  
 uuid..... 257, 261

**V**

valid-path?..... 157  
 verbosity..... 178

**W**

which..... 149  
 with-directory-excursion..... 147  
 with-extensions..... 167, 169  
 with-imported-modules..... 166, 168  
 with-monad..... 162  
 with-parameters..... 174  
 wrap-program..... 152  
 wrap-script..... 153

**X**

xorg-start-command..... 363  
 xorg-start-command-xinit..... 363